

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication
Kind Code
Publication Date
Inventor(s)

20250267260
A1
August 21, 2025
STRÖM; Jacob et al.

ADAPTIVE LOOP FILTERING

Abstract

A method for decoding an image. The method includes obtaining a signed 12-bit input value, wherein the 12-bit input value is a function of at least a first sample value associated with the image. The method also includes producing a signed 19-bit output value that is equal to $i.\text{sub}.v \times 2.\text{sup}.n \times (a+1)$, where $i.\text{sub}.v$ is the 12-bit input value, n is an integer greater than or equal to 0, and a is either 2 or 4. The method further includes using the signed 19-bit output value to produce a filtered sample value.

Inventors: STRÖM; Jacob (Stockholm, SE), ZHANG; Zhi (Solna, SE), ANDERSSON; Kenneth (Gävle, SE)
Applicant: Telefonaktiebolaget LM Ericsson (publ) (Stockholm, SE)
Family ID: 1000008576845
Assignee: Telefonaktiebolaget LM Ericsson (publ) (Stockholm, SE)
Appl. No.: 19/198563
Filed: May 05, 2025

Related U.S. Application Data

parent US division 17783132 20220607 PENDING WO division PCT/SE2020/051096 20201116 child US 19198563
us-provisional-application US 62945489 20191209

Publication Classification

Int. Cl.: H04N19/117 (20140101); H04N19/132 (20140101); H04N19/82 (20140101)
U.S. Cl.:
CPC H04N19/117 (20141101); H04N19/132 (20141101); H04N19/82 (20141101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a divisional of U.S. patent application Ser. No. 17/783,132, having a 371(c) date of 2022 Jun. 7 (status pending), which is the 35 U.S.C. § 371 National Stage of International Patent Application No. PCT/SE2020/051096, filed 2020 Nov. 16, which claims priority to U.S. provisional patent application No. 62/945,489, filed on 2019 Dec. 9. The above identified applications are incorporated by reference.

TECHNICAL FIELD

[0002] This disclosure relates to video encoding and/or decoding.

BACKGROUND

[0003] This disclosure relates to the encoding and/or decoding of an image or a video sequence. A video sequence consists of several images. When viewed on a screen, the image consists of pixels, each pixel having a red, green and blue (RGB) value. However, when encoding and decoding a video sequence, the image is often not represented using RGB values but typically using another color space, including but not limited to YCbCr, IC.sub.TC.sub.P, non-constant-luminance YCbCr, and constant luminance YCbCr. If we take the example of YCbCr, it is made up of three components: luma (Y) which roughly represents luminance, and chroma (Cb, and Cr), both of which represents chrominance. It is often the case that Y is of full resolution, whereas the two other components, Cb and Cr, are of a smaller resolution. A typical example is a high definition (HD) video sequence containing 1920×1080 RGB pixels, which is often represented with a 1920×1080-resolution Y component, a 960×540 Cb component and a 960×540 Cr component. The elements in the components are called samples. In the example given above, there are therefore 1920×1080 samples in the Y component, and hence a direct relationship between samples and pixels. Therefore, in this document, we sometimes use the term pixels and samples interchangeably. For the Cb and Cr components, there is no direct relationship between samples and pixels; a single Cb sample typically influences several pixels.

[0004] In the draft for the Versatile Video Coding (VVC) standard, which is developed by the Joint Video Experts Team (JVET), the decoding of an image can be thought of as carried out in two stages: (1) prediction decoding and (2) loop filtering. In the prediction decoding stage, the samples of the components (Y, Cb and Cr) are partitioned into rectangular blocks. As an example, one block may be of size 4×8 samples, whereas another block may be of size 64×64 samples. The decoder obtains instructions for how to do a prediction for each block, for instance to copy samples from a previously decoded image (an example of temporal prediction), or copy samples from already decoded parts of the current image (an example of

intra prediction), or a combination thereof. To improve this prediction, the decoder may obtain a residual, often encoded using transform coding such as discrete sine transform (DST). This residual is added to the prediction, and the decoder can proceed to decode the subsequent block.

[0005] The output from the prediction decoding stage is the three components Y, Cb and Cr. However, it is possible to further improve the fidelity of these components, and this is done in the loop filtering stage. The loop filtering stage in the current draft of VVC consists of three sub-stages: (1) a deblocking filter stage, (2) a sample adaptive offset filter (SAO) sub-stage, and (3) an adaptive loop filter (ALF) sub-stage.

[0006] In the deblocking filter sub-stage, the decoder changes Y, Cb and Cr by smoothing edges near block boundaries when certain conditions are met. This increases perceptual quality (subjective quality) since the human visual system is very good at detecting regular edges such as block artifacts along block boundaries. In the SAO sub-stage, the decoder adds or subtracts a signaled value to samples that meet certain conditions, such as being in a certain value range (band offset SAO) or having a specific neighborhood (edge offset SAO). This can reduce ringing noise since such noise often aggregate in certain value ranges or in specific neighborhoods (e.g., in local maxima). In this document we will denote the reconstructed image component that are the result of this stage Y_SAO, Cb_SAO, Cr_SAO.

[0007] Embodiments of this disclosure relate to the third sub-stage (i.e., the ALF stage). The basic idea behind adaptive loop filtering is that the fidelity of the image components Y_SAO Cb_SAO and Cr_SAO can often be improved by filtering the image using a linear filter that is signaled from the encoder to the decoder. As an example, by solving a least-squares problem, the encoder can determine what coefficient values a linear filter should have in order to most efficiently lower the error between the reconstructed image components so far, Y_SAO, Cb_SAO, Cr_SAO, and the original image components Y_org, Cb_org and Cr_org. These coefficient values (or simply “coefficients” for short) can then be signaled from the encoder to the decoder. The decoder reconstructs the image as described above to get Y_SAO, Cb_SAO, and Cr_SAO, obtains the filter coefficients from the bit stream and then applies the filter to get the final output, which we will denote Y_ALF, Cb_ALF, Cr_ALF.

[0008] In VVC, the ALF is more advanced than this. To start with, it is observed that it is often advantageous to filter some samples with one set of coefficients, but avoid filtering other samples, or perhaps filter those other samples with another set of coefficients. To that end, VVC classifies every Y sample (i.e., every luma sample) into one of 25 classes. The class to which a sample belongs is decided based on the local neighborhood of that sample, specifically on the gradients of surrounding samples and the activity of surrounding samples. It is possible for the encoder to signal one set of coefficients for each of the 25 classes. The decoder will then first decide which class a sample belongs to, and then select the appropriate set of coefficients to filter the sample. However, signaling 25 sets of coefficients can be costly. Hence the VVC standard also allows that only a few of the 25 classes are filtered using unique sets of coefficients. The remaining classes may reuse a set of coefficients used in another class, or it may be determined that they should not be filtered at all.

[0009] Another way to reduce cost is to use what is called the fixed coefficient set. This is a set of 64 hard-coded filters (i.e., 64 groups of coefficient values) that are known to the decoder. It is possible for the encoder to signal the use of one of these fixed (i.e., hard-coded) filters to the decoder very inexpensively, since they are already known to the decoder. For example, the decoder stores a set of 16 different groups of N index values (e.g., N=25) and the encoder transmits an initial index value that points to one of the 16 groups of N index values, where each one of the index values included in the group of N index values is associated with a class and each one of the index values points to one of the 64 hard-coded filters. For example, the first of the N values in the group of index values points to the fixed filter that should be used for the first class, the second value points to the fixed filter that should be used for the second class, etc. Accordingly, the decoder obtains an index value for a particular filter based on the initial index value and the class. Although these filters are cheap, they may not match the desired filter perfectly and thus result in slightly worse quality. The 64 allowed fixed filter coefficient sets are listed in Table 4. For samples belonging to Cb or Cr, i.e., for chroma samples, no classification is used and the same set of coefficients is used for all samples.

[0010] Transmitting the filter coefficients is costly, and therefore the same coefficient value is used for two filter positions. For luma (samples in the Y-component), the coefficients are re-used in the way shown in FIG. 1.

[0011] Referring to FIG. 1, assume R(x,y) is the sample to be filtered, situated in the middle of FIG. 1. Then samples R(x,y-1) (the sample exactly above) and the sample R(x,y+1) (the sample exactly below) will be treated with the same coefficient C6. The filtered version of the sample in position (x, y), which we will denote R.sub.F(x, y), is calculated with the help of the variable sum which is in turn calculated as shown in Table 1 below:

TABLE-US-00001 TABLE 1
$$\text{sum} = C0 * [\text{clip}(s0, R(x, y - 3) - R(x, y)) + \text{clip}(s0, R(x, y + 3) - R(x, y))] + C1 * [\text{clip}(s1, R(x - 1, y - 2) - R(x, y)) + \text{clip}(s1, R(x + 1, y + 2) - R(x, y))] + C2 * [\text{clip}(s2, R(x, y - 2) - R(x, y)) + \text{clip}(s2, R(x, y + 2) - R(x, y))] + C3 * [\text{clip}(s3, R(x - 1, y - 2) - R(x, y))] + C4 * [\text{clip}(s4, R(x - 2, y - 1) - R(x, y)) + \text{clip}(s4, R(x + 2, y + 1) - R(x, y))] + C5 * [\text{clip}(s5, R(x - 1, y - 1) - R(x, y)) + \text{clip}(s5, R(x + 1, y + 1) - R(x, y))] + (\text{Eqn 1}) C6 * [\text{clip}(s6, R(x, y - 1) - R(x, y)) + \text{clip}(s6, R(x, y + 1) - R(x, y))] + C7 * [\text{clip}(s7, R(x + 1, y - 1) - R(x, y)) + \text{clip}(s7, R(x - 1, y + 1) - R(x, y))] + C8 * [\text{clip}(s8, R(x + 2, y - 1) - R(x, y)) + \text{clip}(s8, R(x - 2, y + 1) - R(x, y))] + C9 * [\text{clip}(s9, R(x - 3, y) - R(x, y)) + \text{clip}(s9, R(x + 3, y) - R(x, y))] + C10 * [\text{clip}(s10, R(x - 2, y) - R(x, y)) + \text{clip}(s10, R(x + 2, y) - R(x, y))] + C11 * [\text{clip}(s11, R(x - 1, y) - R(x, y)) + \text{clip}(s11, R(x + 1, y) - R(x, y))].$$

[0012] Here the clip (m, x) operation simply makes sure that the magnitude of the value x never exceeds m:

$$\text{clip}(m, x) = \begin{cases} \min(x, m) & \text{if } x \geq 0 \\ \max(x, -m) & \text{if } x < 0 \end{cases} \quad (\text{Eqn2})$$

[0013] The filtered value R.sub.F(x, y) is finally calculated as:

$$R_F(x, y) = R(x, y) + ((\text{sum} + 64) \gg 7) \quad (\text{Eqn3})$$

[0014] The magnitudes s0 through s11 are also be signaled from the encoder to the decoder. Note that coefficient C12 is not used in Equation 1 since the value clip(s12, R(x,y)-R(x, y)) is always zero.

SUMMARY

[0015] Certain challenges currently exist. For example, as can be seen in Equation (Eqn) 1, there are 12 multiplications per sample necessary to calculate the sum value. In hardware, multiplications can be expensive, especially if they must be dimensioned for large values, and if they have to produce a result every clock cycle. The allowed range of values for the coefficients C0 through C11 is [-127, 127], and the largest value for the clip parameters s0 through s11 is 1023, in a 10-bit implementation. The value that C0 is multiplied by is a sum of two clip outputs, and can therefore be at most 2046 and at smallest -2046. This means that a signed 12-bit number can hold this factor. This in turn means that an 8-bit×12-bit multiplier must be implemented. In a typical scenario, it is required to be able to filter one sample per clock. This would mean that twelve multipliers of this size must be implemented. This is quite big in terms of silicon surface area, and hence quite costly to implement.

[0016] FIG. 2 shows one way of implementing an 8-bit×12-bit multiplier in hardware. This particular multiplier calculates the multiplication of two positive numbers, but a multiplier capable of multiplying signed 8-bit×12-bit numbers is very similar. There are more efficient ways to implement multiplication than the shifted additions shown in FIG. 2. However, the implementation still gives a rough approximation of the hardware complexity that would be needed in order to implement this multiplication in one clock-cycle; about seven adders of a width between 13 and 20 bits. This is a large number given that the hardware needs to be replicated twelve times, once for each coefficient.

[0017] Accordingly, this disclosure proposes ways to lower the size of the silicon surface area needed to implement this in hardware. In one embodiment, the coefficients are constrained such that each coefficient is a sum of two power-of-two numbers. This means that every coefficient multiplication in the filter can be implemented using only one 13-bit wide addition, as well as some other logic that is roughly the size of one more addition.

[0018] In one aspect a method for decoding an image is provided. In one embodiment the method includes obtaining a set of sample values

associated with the image, the set of sample values comprising a first sample value. The method also includes employing an adaptive loop filter (ALF) to filter the first sample value. The ALF is operable to filter the first sample value using any set of N coefficient values in which each one of the N coefficient values is included in a set of M unique coefficient values, wherein N is greater than 1 and M is greater than 1. The set of M unique coefficient values consists of the following unique values or consists of a subset of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128, and the set of M unique coefficient values includes at least one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126$, or 127. Employing the ALF to filter the first sample value comprises the steps of: a) obtaining a first set of N coefficient values for use in filtering the first sample value and b) using the ALF to filter the first sample value using the obtained first set of N coefficient values and the set of sample values, thereby producing a first filtered sample value, and each coefficient value included in the obtained first set of N coefficient values is constrained such that the coefficient value must be equal to one of the values included in the set of M unique values.

[0019] In one embodiment the method includes obtaining a set of sample values associated with the image, the set of sample values comprising a first sample value, and also obtaining an index value that points to a particular coefficient value group included within a set of M predefined coefficient value groups (e.g., $M=64$). Each coefficient value group included in the set of predefined coefficient value groups consists of N coefficient values, N being greater than 1, and: i) for each coefficient value group included in the set of predefined coefficient value groups, each coefficient value included in the coefficient group is constrained such that the coefficient value must be equal to one of the following values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128 and ii) for at least one coefficient value group included in the set of predefined coefficient value groups, at least one of the coefficient values included in said at least one coefficient value group is equal to one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126$, or 127. The method also includes using the index value to select the particular coefficient value group from the set of predefined coefficient value groups and employing an adaptive loop filter (ALF) to filter the first sample value using the particular coefficient value group selected from the set of predefined coefficient value groups.

[0020] In another aspect a decoding apparatus is provided. The decoding apparatus is adapted to perform any one of the decoding methods disclosed herein. In some embodiments, the decoding apparatus includes processing circuitry and a memory, said memory containing instructions executable by said processing circuitry.

[0021] In another aspect there is provided a method performed by an encoder. The method includes the encoder selecting a set of coefficient values for use by a decoder in filtering a sample value, the selected set of coefficient values consisting of N coefficient values. Each one of the N coefficient values is included in a set of M unique coefficient values, wherein N is greater than 1 and M is greater than 1 and further wherein i) the set of M unique coefficient values consists of the following unique values or consists of a subset of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128 and ii) the set of M unique coefficient values includes at least one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126$, or 127. And each coefficient value included in the set of N coefficient values is constrained such that the coefficient value must be equal to one of the values included in the set of M unique values. The method also includes the encoder providing to a decoder the N coefficient values or an initial index value for use by the decoder to determine the set of N coefficient values.

[0022] In another aspect an encoding apparatus is provided. The encoding apparatus is adapted to perform any one of the encoding methods disclosed herein. In some embodiments, the encoding apparatus includes processing circuitry and a memory, said memory containing instructions executable by said processing circuitry.

[0023] In another aspect there is provided a computer program comprising instructions which when executed by processing circuitry causes the processing circuitry to perform any of the method disclosed herein. In another aspect a carrier containing the computer program is provided, wherein the carrier is one of an electronic signal, an optical signal, a radio signal, and a computer readable storage medium.

Advantages

[0024] As a rough estimate of the surface area needed to implement various functions, assume that we calculate the number of additions and multiply by the width of those additions. As can be seen in FIG. 2, four 13-bit adders, two 16-bit adders and one 20-bit adder are used to multiply the two numbers. This gives a surface cost of $13 \times 4 + 16 \times 2 + 20 = 104$. In the proposed embodiment, only two 13-bit adders are used. This gives a surface cost of $213 = 23$. Hence the cost in terms of surface has decreased by $(104 - 23) / 104 = 78\%$, a substantial reduction. Given that the multipliers in an ALF make up a major part of the surface area needed, this translates to a significant reduction in the surface area needed for implementing the entire ALF method. Not only can surface area be saved, but also there will be a saving in power consumption, which can increase the battery life of a hand-held device, because the operations that are carried out in the proposed embodiments are less complex than a multiplication, and hence consume less power.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0025] The accompanying drawings, which are incorporated herein and form part of the specification, illustrate various embodiments.

[0026] FIG. 1 provides an illustration of coefficient reuse.

[0027] FIG. 2 shows one way of implementing an 8-bit \times 11-bit multiplier in hardware.

[0028] FIG. 3 illustrates a system comprising an encoder and a decoder.

[0029] FIG. 4 illustrates an example encoder.

[0030] FIG. 5 illustrates an example decoder.

[0031] FIG. 6 illustrates the set of allowed coefficient values according to an embodiment.

[0032] FIG. 7 illustrates the set of allowed coefficient values according to an embodiment.

[0033] FIGS. 8-14 illustrate circuits according to various embodiments.

[0034] FIG. 15A shows a frequency plot of different coefficient code values.

[0035] FIG. 15B shows a frequency distribution of coefficient 11.

[0036] FIG. 16 illustrates a circuit according to an embodiment.

[0037] FIG. 17 shows the values of coefficient C0, C1, C6 and C11 respectively

[0038] FIG. 18A is a flow chart illustrating a process according to an embodiment.

[0039] FIG. 18B is a flow chart illustrating a process according to an embodiment.

[0040] FIG. 19 is a block diagram of an apparatus according to one embodiment.

[0041] FIG. 20 is a flow chart illustrating a process according to an embodiment.

DETAILED DESCRIPTION

[0042] FIG. 3 illustrates a system 300 according to an example embodiment. System 300 includes an encoder 302 and a decoder 304. In the example shown, decoder 304 can receive via a network 110 (e.g., the Internet or other network) encoded images produced by encoder 302.

[0043] FIG. 4 is a schematic block diagram of encoder 302. As illustrated in FIG. 4, The encoder 302 takes in an original image and subtracts a prediction 41 that is selected 51 from either previously decoded samples ("Intra Prediction" 49) or samples from previously decoded frames stored in

the frame buffer 48 through a method called motion compensation 50. The task of finding the best motion compensation samples is typically called motion estimation 50 and involves comparing against the original samples. After subtracting the prediction 41 the resulting difference is transformed 42 and subsequently quantized 43. The quantized results are entropy encoded 44 resulting in bits that can be stored, transmitted or further processed. The output from the quantization 43 is also inversely quantized 45 followed by an inverse transform 46. Then the prediction from 51 is added 47 and the result is forwarded to both the intra prediction unit 49 and to the Loopfilter Unit 100. The loopfilter unit 100 may do deblocking, SAO and/or ALF filtering. The result is stored in the frame buffer 48, which is used for future prediction. Not shown in the figure is that coding parameters for other blocks such as 42, 43, 49, 50, 51 and 100 also may also be entropy coded.

[0044] FIG. 5 is a corresponding schematic block diagram of decoder 304 according to some embodiments. The decoder 304 takes in entropy coded transform coefficients which are then decoded by decoder 61. The output of decoder 61 then undergoes inverse quantization 62 followed by inverse transform 63 to form a decoded residual. To this decoded residual, a prediction is added 64. The prediction is selected 68 from either a motion compensation unit 67 or from an intra prediction unit 66. After having added the prediction to the decoded residual 64, the samples can be forwarded for intra prediction of subsequent blocks. The samples are also forwarded to the loopfilter unit 100, which may do deblocking, SAO processing, and/or adaptive ALF processing. The output of the loopfilter unit 100 is forwarded to the frame buffer 65, which can be used for motion compensation prediction of subsequently decoded images 67. The output of the loopfilter unit 100 can also be output the decoded images for viewing or subsequent processing outside the decoder. Not shown in the figure is that parameters for other blocks such as 63, 67, 66 and 100 may also be entropy decoded. As an example, the coefficients for the ALF filter in block 100 may be entropy decoded.

[0045] In one embodiment, to achieve the advantages discussed above, the ALF part of the loopfilter unit 100 is configured such that the coefficients are restricted to certain values for which there is an inexpensive way to implement a multiplication. In one embodiment, the coefficients are restricted to pure powers-of-two, rather than allowing all values between -128 and 128. That is, the coefficients are constrained such that each coefficient must be equal to one of the following values: $\pm\{0, 1, 2, 4, 8, 16, 32, 64, 128\}$. Multiplication of $a*b$, where a is one of the allowed coefficient values, would then for positive values be implemented using $b \ll k$, where k is 0 through 7. For negative values the result would have to be sign-corrected also. This would substantially reduce the complexity when implementing the multiplications, but it would come at a cost in precision. As it turns out, this means that the quality in terms of the average bit rate difference (BD-rate) can go down substantially, by as much as 0.2%. This is not a good trade-off between complexity and image quality.

[0046] Accordingly, in another embodiment it is proposed to use a less severe restriction on the allowable coefficient values. Instead of allowing all values between -128 and 128, only values that can be written as a pure power-of-two number or as the sum of two power-of-two numbers of arbitrary sign are allowed. As an example, 6 would be allowed, since it can be written as $4+2$, and 7 would be allowed, since it can be written as $8-1$, but 22 would not be allowed, since it cannot be written as either $\pm 2^n$ or $(\pm 2.\text{sup}.n \pm 2.\text{sup}.m)$. Also zero would be allowed since it can be written as, for instance $2\{\text{circumflex over } (\)\}1-2\{\text{circumflex over } (\)\}1$.

[0047] The allowed coefficient values between -128 and 128 (excluding -128 and 128) are listed as set Z, $Z=\{-127, -126, -124, -120, -112, -96, -80, -72, -68, -66, -65, -64, -63, -62, 60, -56, -48, -40, -36, -34, -33, -32, -31, -30, -28, -24, -20, -18, -17, -16, -15, -14, -12, -10, -9, -8, -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127\}$.

[0048] In FIG. 6 we can see which numbers would be allowed in this case. As can be seen, most coefficients around zero would be allowed. For coefficients with larger magnitudes, more coefficients would be disallowed. This is reasonable, given the fact that the coefficients typically have a distribution where smaller coefficients are more common, i.e., similar to a Laplace distribution, and it is hence more important to be able to faithfully reproduce those. However, in the FIG. 6 it can also be seen that more values are allowed in some areas even though they have a high magnitude. As an example, around 64, most values are allowed, whereas around 48, very few are allowed. This does not make sense if the coefficients are really Laplace distributed; in that case a value around 48 is more probable than a value around 64. Therefore, in another embodiment the coefficients that can be used as ALF filter coefficients is further restricted.

[0049] That is, in one embodiment, a subset of Z is used, namely Z.sub.sub. In one example $Z.\text{sub}.\text{sub}=\{-40, -33, -28, -24, -20, -17, -15, -14, -12, -10, -9, -8, -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 17, 20, 24, 28, 33, 40\}$. In another embodiment, the coefficients are constrained so that the coefficients can must be written as either 0, $\pm 2.\text{sup}.n$ or $\pm(2.\text{sup}.n+2.\text{sup}.n-1)$ or $\pm(2.\text{sup}.n+2.\text{sup}.n-2)$. In another embodiment, the coefficients are constrained so that they can be written as either 0, $\pm 2.\text{sup}.n$ or $\pm(2.\text{sup}.n+2.\text{sup}.n-1)$. As can be seen in FIG. 7, much fewer coefficients are allowed in this last case. However, the distribution is better adapted to a Laplacian distribution, where the number of allowed coefficients decreases with magnitude.

[0050] The allowed values in this embodiment then belong to the following set S, $S=\{-128, -96, -64, -48, -32, -24, -16, -12, -8, -6, -4, -3, -2, -1, 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64, 96, 128\}$. The set of coefficients currently allowed in VVC is denoted herein as T, where $T=\{127, -126, -125, \dots, -2, -1, 0, 1, 2, \dots, 125, 126, 127\}$.

[0051] To calculate the sum value from Equation 1 (see table 1), we need to perform several multiplications of the form $a*b$, where a is an allowed coefficient, i.e., it belongs to the set S, and b is a sum of two clipped difference value, i.e., it can take any value in the range $[-2046, 2046]$, needing a signed 12-bit variable to hold it.

[0052] We can write $2.\text{sup}.n+2.\text{sup}.n-1$ as $2.\text{sup}.n-1(2+1)=2.\text{sup}.n-1*3$. Hence the value a is either 0, a pure power-of-two or a pure power-of-two multiplied by three. We can write this as:

$$[00003] a = \pm(k_1 * 2 + k_0 * 1) * 2^s, \quad (\text{Eqn. 4})$$

[0053] where $k.\text{sub}.0$ and $k.\text{sub}.1$ can take the values of 0 or 1.

[0054] In the case when we have a pure power-of-two, such as 128, we set $k.\text{sub}.1=1$, $k.\text{sub}.0=0$ and s to a suitable shift value, 6 in the case of 128. (Since $k.\text{sub}.1=1$ we multiply by two, hence we should use 6 to represent 128.) In the case when we have a power-of-two number multiplied by three, such as 96, we set both $k.\text{sub}.1$ and $k.\text{sub}.0$ to 1, and use a suitable shift value, such as 5 in the case of 96. Table 2 shows possible values for $k.\text{sub}.1$, $k.\text{sub}.0$ and s for the values in S. It also shows the value n , which indicates if the value should be negated.

TABLE-US-00002 TABLE 2 How to write the allowed coefficients on the form $(-1).\text{sup}.n(2k.\text{sub}.1 + k.\text{sub}.0)2.\text{sup}.s$ coefficient $k_1 k_0 s n$

-128	1	0	6	1
-96	1	1	5	1
-64	1	0	5	1
-48	1	1	4	1
-32	1	0	4	1
-24	1	1	3	1
-16	1	0	3	1
-12	1	1	2	1
-8	1	0	2	1
-6	1	1	1	1
-4	1	0	1	1
-3	1	1	0	1
-2	1	0	0	1
-1	0	1	0	0
0	0	0	0	0
1	0	1	0	0
2	1	0	0	0
3	1	1	0	0
4	1	0	0	1
6	1	1	0	1
8	1	0	1	1
12	1	1	1	1
16	1	0	1	1
24	1	1	1	1
32	1	0	1	1
48	1	1	1	1
64	1	0	5	0
96	1	1	5	0
128	1	0	6	0

[0055] The decoder can use Table 2 to determine the values of $k.\text{sub}.1$, $k.\text{sub}.0$, s and n from the coefficient. An alternative is to use the following pseudo code for a coefficient coeff:

TABLE-US-00003 TABLE 3 $k_1 = (\text{abs}(\text{coeff}) < 2 ? 0 : 1)$; $k_0 = \text{coeff} \& 1$; $s = \max(0, 6-\text{clz}(\text{abs}(\text{coeff})))$; $n = \text{sign}(\text{coeff})$;

[0056] Here $\text{abs}(x)$ denotes absolute value of x , $\&$ denotes bitwise AND, $\max(a,b)$ returns the largest value of a and b , $\text{clz}(x)$ counts the leading number of zeros in the binary representation of x , so the binary 8-bit number 0001111 (15 in decimal representation) will return 3, and $\text{sign}(x)$ returns the sign of x . $\text{clz}()$ is a common assembly instruction on most CPUs so it is inexpensive.

[0057] Note that this conversion only needs to happen when the coefficients are read from the Adaptive Parameter Set (APS). The APS consists of a set of parameters that the encoder transmits to the decoder. In particular, it contains the coefficient values used in ALF, and these are sent/received at most once per frame. Hence it is not critical that this conversion from coefficient to values is extremely fast or efficient. If, on the other hand, this conversion would have to happen every sample, it would be very important that it could be done quickly.

[0058] Once they have been converted, a hardware implementation can store them for later use during the filtering. Since of $k.\text{sub}.1$, $k.\text{sub}.0$, and n

are 1-bit values, and s is a 3-bit value, the total number of bits that need to be stored is 6 bits. This is less than the current implementation of ALE, which needs to store an 8-bit value between -127 and 127 for each coefficient.

[0059] The multiplication $a*b$ can be re-written as:

$$[00004] \ a * b = (-1)^n (2k_1 + k_0)2^s * b = (-1)^n (2 * b * k_1 + b * k_0)2^s = (-1)^n (((b * k_1) \ll 1) + b * k_0)2^s. \quad (\text{Eqn5a})$$

[0060] To evaluate the bottom-most expression, we can start by multiplying b by k.sub.1. Since k.sub.1 is either 0 or 1 this is the same as doing AND between every bit in b and k.sub.1. After this, we will shift it one step left. Likewise, we will do AND between b and k.sub.0. We add these two results together, negate it if necessary and shift it 0 to 6 steps. Because the multiplications can be replaced by ANDs, Equation 5 can be written as:

$$[00005] \ a * b = (-1)^n (((b \&_b k_1) \ll 1) + b \&_b k_0)2^s, \quad (\text{Eqn5b})$$

where $x \&_b y$ is used to denote that every bit in x is ANDed with the one-bit value y. Equation 5b can be efficiently implemented by the circuit shown in FIG. 8.

[0061] As can be seen in FIG. 8, the top left unit marked “bit-wise &” takes in the signed 12-bit number

$b = b_{\text{sub.11}}b_{\text{sub.10}}b_{\text{sub.9}}b_{\text{sub.8}}b_{\text{sub.7}}b_{\text{sub.6}}b_{\text{sub.5}}b_{\text{sub.4}}b_{\text{sub.3}}b_{\text{sub.2}}b_{\text{sub.1}}b_{\text{sub.0}}$, where $b_{\text{sub.11}}$ is the most significant bit, and does a bit-wise AND with $k_{\text{sub.1}}$ according to $\text{out}_{\text{sub.k}} = b_{\text{sub.k}} \& k_{\text{sub.1}}$. Such a unit can be constructed using 12 AND-gates for 12-bit input data. FIG. 9 shows how such a unit capable of 11-bit input can be implemented in hardware using only 11 AND-gates, which are one of the least expensive computational units available in hardware. In FIG. 9a is fed with the input k_i from FIG. 8. The output of the top left unit is shifted one bit, and a zero is inserted in the least significant bit position. This means that the resulting value is 13 bits.

[0062] In a similar manner, the value b is bit-wise AND:ed with $k_{\text{sub.0}}$ in the bottom-left unit marked “bit-wise &”. The output is not shifted, instead the sign bit is extended so that the result is also 13 bits. This is indicated by the wiring diagram between the lower “bit-wise &” unit and the adder. As can be seen in FIG. 10, this does not contain any logic. The top portion of FIG. 10 shows the wiring diagram from FIG. 8 and the bottom portion of FIG. 10 shows in further detail the same wiring, where the input bits are copied to the output, and the most significant bit in 11 is copied to the two most significant bits in the output out12 and out11.

[0063] These 13-bit values are then added together using a 13-bit adder. The output is 14 bits, since one bit may carry. This result is then input to the unit marked “conditional negate”, which implements the multiplication of $(-1)^{\text{sup.n}}$. FIG. 11 shows how such a unit can be implemented in hardware. That is, FIG. 11 shows an example of how to implement a conditional negater. If the value n is 1, the input value is negated. If the value is 0, the input value is left untouched.

[0064] As is well-known for a person skilled in the art, it is possible to negate a value by inverting all the bits and adding 1. This should only be done in the case when $n=1$. By using an XOR gate, each input bit is inverted in the case when $n=1$, and left untouched when $n=0$. The result is then fed to an adder, where the other input is zero, and where the carry-in is set to n. This means that it will leave the value untouched if $n=0$, but if $n=1$ it will add 1. The result is a 14-bit value which is negated in relation to the input if $n=1$ and left untouched otherwise.

[0065] Finally, the right-most box in FIG. 8 marked “variable bit shift” implements the multiplication by $2^{\text{sup.s}}$. This can be efficiently implemented using a barrel shifter, as shown in FIG. 12, which illustrates a barrel-shifter, that shifts the input up to 6 steps to the left (up in this figure), controlled by the parameter $s = s_{\text{sub.2}}s_{\text{sub.1}}s_{\text{sub.0}}$, where $s_{\text{sub.2}}$ is the most-significant bit in s.

[0066] The barrel-shifter in FIG. 12 consists of 51 1-bit multiplexors, and although it may seem complex, it can be implemented very efficiently in hardware. The input to the barrel shifter is shifted s steps to the left, (up in FIG. 12) where the binary representation of $s = s_{\text{sub.2}}s_{\text{sub.1}}s_{\text{sub.0}}$ and where $s_{\text{sub.2}}$ is the most significant bit. Note that although the input is a signed 14-bit value, and the maximum shift is 6 steps, the output is a 19-bit value and not a 20-bit value. This is due to the fact that the values with largest magnitude to come out are $-2046 * 128 = -261888$ and $2046 * 128 = 261888$, both of which can be represented by a 19-bit number that can hold numbers in the range $[-2^{\text{sup.18}}, 2^{\text{sup.18}} - 1] = [-262144, 262143]$.

[0067] The FIGS. 9-12 are just examples of how to implement this type of multiplication. Often there are less expensive methods to implement the various building blocks in FIG. 8. As an example, FIG. 11 uses an adder where one of the inputs are set to zero. For a person skilled in the art, it is clear that this does not need to be implemented using a full general adder, but it is possible to reduce the size of this adder since we already know that one input will be zero.

[0068] In fact, it is possible to fully remove the conditional negater from FIG. 8. The result will then have the wrong sign, but this can be taken care of later when the value is added to another product to form the value ‘sum’ in Equation 1. In detail, consider the first two terms in the sum in Equation 1, rewritten here for convenience:

$$[00006] \ \text{sum} = C0 * [\text{clip}(s0, R(x, y - 3) - R(x, y)) + \text{clip}(s0, R(x, y + 3) - R(x, y))] + C1 * [\text{clip}(s1, R(x - 1, y - 2) - R(x, y)) + \text{clip}(s1, R(x + 1, y + 2) - R(x, y))] + \dots$$

[0069] These two terms can be written as the addition of two products

$$[00007] \ \text{partsum} = a0 * b0 + a1 * b1, \quad (\text{Eqn6})$$

where $a0 = C0$ which belongs to set S and where $b0$ is the value in the first square bracket. Likewise, $a1 = C1$ and $b1$ is the value in the second square bracket. Assume we have calculated the correct value for $a0 * b0$, but that we have the incorrect sign for the term $a1 * b1$. We can then invert the bits in $a1 * b1$ and use the carry-in in the adder to add one to the expression without paying the penalty of another adder. In detail, we use

$$[00008] \ \text{partsum} = a0 * b0 + \text{bit_invert}(a1 * b1) + 1, \quad (\text{Eqn7})$$

where the 1 is added by setting carry-in to 1. If instead $a0 * b0$ has the incorrect sign, but $a1 * b1$ has the right sign, it is instead possible to use

$$[00009] \ \text{partsum} = \text{bit_invert}(a0 * b0) + a1 * b1 + 1, \quad (\text{Eqn8})$$

where again the extra 1 is added using the carry in. If both values have the correct sign, we simply use

$$[00010] \ \text{partsum} = a0 * b0 + a1 * b1. \quad (\text{Eqn9})$$

[0070] However, if both terms have the incorrect sign, it is not possible to get the correct output. However, in that case it is still possible to get the negative of the correct output by using Equation 9. And since this output is in turn going to be added to further partial sums in Equation 1, it is possible to again avoid the extra adder. In the end one only needs a conditional negater for the value sum itself, which is much better than having a conditional negater for every multiplication in Equation 1, of which there are 12.

[0071] By removing the conditional negater from FIG. 8 in this way, the two most expensive operations that remain are the 13-bit adder and the barrel shift. Assuming that the barrel-shifter is approximately as complex as the adder, the cost of implementation is down to approximately two 13-bit adders.

0. Encoder-Only Embodiments

[0072] It is possible to have the encoder voluntarily restrict the coefficients to those in set S, i.e., values that are 0, ± 1 , ± 2 , ± 3 , ± 4 , ± 6 , ± 8 , ± 12 , ± 16 , ± 24 , ± 32 , ± 48 , $\pm 64 \pm 96$ or ± 128 . However, such a variant would not be compatible with the current VVC decoder, since values of ± 128 are not allowed. Therefore, in one embodiment, it is possible for the encoder to voluntarily restrict the coefficients to an alternative set: $S_{\text{sub.96}} = \{0, \pm 1, \pm 2, \pm 3, \pm 4, \pm 6, \pm 8, \pm 12, \pm 16, \pm 24, \pm 32, \pm 48, \pm 64, \pm 96\}$, i.e., $S - \{-128, 128\}$. Therefore, in one embodiment, the encoder restricts the coefficients to those in set $S_{\text{sub.96}}$, for instance by quantizing every coefficient to the nearest allowed coefficient, such as changing -34 to the allowed value -32.

[0073] However, for the decoder to be able to take advantage of the fact that the computation can be implemented less expensively, it has to be able to make sure that only values in $S_{\text{sub.96}}$ are used for the coefficients. This can be done by checking during decoding; if all coefficients in all filters belong to $S_{\text{sub.96}}$, then the less expensive (faster) implementation can be used. If there are one or more coefficients that do not belong to $S_{\text{sub.96}}$,

then the more expensive implementation is used. This solution has the advantage of being changed, since the expensive implementation (which is currently used) can always be used. For decoders that want to take advantage of the possibility of processing the data faster, or using less power, it can do so by checking the coefficients against S.sub.96.

[0074] It should be noted that many of the 64 fixed filters that are currently defined have coefficients outside S.sub.96, for instance the filter with index 0, see TABLE 4. Hence a decoder would also need to check that these are not used if the fast data processing should be used.

[0075] In one embodiment the encoder signals whether it uses coefficients in S.sub.96 or if it allows all types of coefficients. In particular, it may also mean that the encoder guarantees that none of the 64 fixed filters that use coefficients outside S.sub.96 are used. The decoder can then know which method to use without having to test every coefficient.

1. Embodiments that Change the Decoder Normatively

[0076] Using the fixed filters can be very effective, especially for low bit rates. Therefore it is a great handicap not to be able to use the 64 fixed filters. An alternative is therefore to change the fixed filters. This can be done by quantizing every filter coefficient to the nearest allowed coefficient.

In Table 4 the fixed filters that are used in the current version of VVC are shown and Table 5 shows a quantized version that only uses values in S.

TABLE-US-00004 TABLE 4 index fixed coefficients in VVC

0 0 0 2 -3 1 -4 1 7 -1 1 -5	1 0 0 0 0 -1 0 1 0 0 -1 2	2 0 0 0 0 0 0 1 0 0 0 0
3 0 0 0 0 0 0 0 0 0 -1 1	4 2 2 -7 -3 0 -5 13 22 12 -3 -3 17	5 -1 0 6 -8 1 -5 1 23 0 2 -5 10
6 0 0 -1 -1 0 -1 2 1 0 0 -1 4	7 0 0 3 -11 1 0 -1	8 0 0 8 -8 -2 -7 4 4 2 1 -1 25
9 0 0 1 -1 0 -3 1 3 -1 1 -1 3	10 0 0 3 -3 0 -6 5 -1 2 1 -4 21 11	-7 1 5 4 -3 5 11 13 12 -8 11 12 12
-5 -3 6 -2 -3 8 14 15 2 -7 11 16 13 2 -1	-6 -5 -2 -2 20 14 -4 0 -3 25 14 3 1 -8	-4 0 -8 22 5 -3 2 -10 29 15 2 1 -7 -1 2 -11 23 -5 0 2 -10 29
16 -6 -3 8 9 -4 8 9 7 14 -2 8 9 17 2 1 -4	-7 0 -8 17 22 1 -1 -4 23 18 3 0 -5 -7 0 -7 15 18 -5 0 -5 27 19 2 0 0 -7 1 -10 13 13 -4 2 -7 24 20 3 3	-13 4 -2 -5 9 21 25 -2 -3 12 21 -5 -2 7 -3 -7 9 8 9 16 -2 15 12 22 0 -1 0 -7 -5 4 11 11 8 -6 12 21 23 3 -2 -3 -8 -4 -1 16 15 -2 -3 3 26 24 2 1
-5 -4 -1 -8 16 4 -2 1 -7 33 25 2 1 -4 -2 1 -10 17 -2 0 2 -11 33 26 1 -2 7 -15 -16 10 8 8 20 11 14 11 27 2 2 3 -13 -13 4 8 12 2 -3 16 24 28 1 4	0 -7 -8 -4 9 9 -2 -2 8 29 29 1 1 2 -4 -1 -6 6 3 -1 -1 -3 30 30 -7 3 2 10 -2 3 7 11 19 -7 8 10 31 0 -2 -5 -3 -2 4 20 15 -1 -3 -1 22 32 3 -1 -8	-4 -1 -4 22 8 -4 2 -8 28 33 0 3 -14 3 0 1 19 17 8 -3 -7 20 34 0 2 -1 -8 3 -6 5 21 1 1 -9 13 35 -4 -2 8 20 -2 2 3 5 21 4 6 1 36 2 -2 -3 -9 -4 2
14 16 3 -6 8 24 37 2 1 5 -16 -7 2 3 11 15 -3 11 22 38 1 2 3 -11 -2 -5 4 8 9 -3 -2 26 39 0 -1 10 -9 -1 -8 2 3 4 0 0 29 40 1 2 0 -5 1 -9 9 3 0 1 -7	20 41 -2 8 -6 -4 3 -9 -8 45 14 2 -13 7 42 1 -1 16 -19 -8 -4 -3 2 19 0 4 30 43 1 1 -3 0 2 -11 15 -5 1 2 -9 24 44 0 1 -2 0 1 -4 4 0 0 1 -4 7 45 0 1	2 -5 1 -6 4 10 -2 1 -4 10 46 3 0 -3 -6 -2 -6 14 8 -1 -1 -3 31 47 0 1 0 -2 1 -6 5 1 0 1 -5 13 48 3 1 9 -19 -21 9 7 6 13 5 15 21 49 2 4 3 -12 -13 1
7 8 3 0 12 26 50 3 1 -8 -2 0 -6 18 2 -2 3 -10 23 51 1 1 -4 -1 1 -5 8 1 -1 2 -5 10 52 0 1 -1 0 0 -2 2 0 0 1 -2 3 53 1 1 -2 -7 1 -7 14 18 0 0 -7 21	54 0 1 0 -2 0 -7 8 1 -2 0 -3 24 55 0 1 1 -2 2 -10 10 0 -2 1 -7 23 56 0 2 2 -11 2 -4 -3 39 7 1 -10 9 57 1 0 13 -16 -5 -6 -1 8 6 0 6 29 58 1 3 1 -6	-4 -7 9 6 -3 -2 3 33 59 4 0 -17 -1 -1 5 26 8 -2 3 -15 30 60 0 1 -2 0 2 -8 12 -6 1 1 -6 16 61 0 0 0 -1 1 -4 4 0 0 0 -3 11 62 0 1 2 -8 2 -6 5 15 0 2
-7 9 6 3 1 -1 12 -15 -7 -2 3 6 6 -1 7 30		

TABLE-US-00005 TABLE 5 index proposed fixed coefficients

0 0 0 2 -3 1 -4 1 6 -1 1 -1 4	1 0 0 0 0 -1 0 1 0 0 -1 2	2 0 0 0 0 0 0 1 0 0 0 0
3 0 0 0 0 0 0 0 0 0 -1 1	4 2 2 -6 -3 0 -4 12 24 12 -3 -3 16	5 -1 0 6 -8 1 -4 1 24 0 2 -4 8
6 0 0 -1 -1 0 -1 2 1 0 0 -1 4	7 0 0 3 -12 1 0 -1	8 0 0 8 -8 -2 -6 4 4 2 1 -1 24
9 0 0 1 -1 0 -3 1 3 -1 1 -1 3	10 0 0 3 -3 0 -6 4 -1 2 1 -4 24 11	-6 1 4 4 -3 4 12 12 12 -8 12 12 12 -4
-3 6 -2 -3 8 12 16 2 -6 12 16 13 2 -1	-6 -4 -2 -2 16 12 -4 0 -3 24 14 3 1 -8	-4 0 -8 24 4 -3 2 -8 32 15 2 1 -6 -1 2 -12 24 -4 0 2 -8 32 16 -6
-3 8 8 -4 8 8 6 12 -2 8 8 17 2 1 -4 -6 0 -8 16 24 1 -1	-4 24 18 3 0 -4 -6 0 -6 16 16 -4 0 -4 24 19 2 0 0 -6 1 -8 12 12 -4 2 -6 24 20 3 3 -12 4	-2 -4 8 24 24 -2 -3 12 21 -4 -2 6 -3 -6 8 8 16 -2 16 12 22 0 -1 0 -6 -4 4 12 12 8 -6 12 24 23 3 -2 -3 -8 -4 -1 16 16 -2 -3 3 24 24 2 1 -4 -4
-1 -8 16 4 -2 1 -6 32 25 2 1 -4 -2 1 -8 16 -2 0 2 -12 32 26 1 -2 6 -16 -16 8 8 8 16 12 12 12 27 2 2 3 -12 -12 4 8 12 2 -3 16 24 28 1 4 0 -6 -8	-4 8 8 -2 -2 8 32 29 1 1 2 -4 -1 -6 6 3 -1 -1 -3 32 30 -6 3 2 8 -2 3 6 12 16 -6 8 8 31 0 -2 -4 -3 -2 4 16 16 -1 -3 -1 24 32 3 -1 -8 -4 -1 -4 24	8 -4 2 -8 24 33 0 3 -12 3 0 1 16 16 8 -3 -6 16 34 0 2 -1 -8 3 -6 4 24 1 1 -8 12 35 -4 -2 8 16 -2 2 3 4 24 4 6 1 36 2 -2 -3 -8 -4 2 12 16 3 -6 8
24 37 2 1 4 -16 -6 2 3 12 16 -3 12 24 38 1 2 3 -12 -2 -4 4 8 8 -3 -2 24 39 0 -1 8 -8 -1 -8 2 3 4 0 0 32 40 1 2 0 -4 1 -8 8 3 0 1 -6 16 41 -2 8 -6	-4 3 -8 -8 48 12 2 -12 6 42 1 -1 16 -16 -8 -4 -3 2 16 0 4 32 43 1 1 -3 0 2 -12 16 -4 1 2 -8 24 44 0 1 -2 0 1 -4 4 0 0 1 -4 6 45 0 1 2 -4 1 -6 4 8	-2 1 -4 8 46 3 0 -3 -6 -2 -6 12 8 -1 -1 -3 32 47 0 1 0 -2 1 -6 4 1 0 1 -4 12 48 3 1 8 -16 -24 8 6 6 12 4 16 24 49 2 4 3 -12 -12 1 6 8 3 0 12 24
50 3 1 -8 -2 0 -6 16 2 -2 3 -8 24 51 1 1 -4 -1 1 -4 8 1 -1 2 -4 8 52 0 1 -1 0 0 -2 2 0 0 1 -2 3 53 1 1 -2 -6 1 -6 12 16 0 0 -6 24 54 0 1 0 -2 0 -6	8 1 -2 0 -3 24 55 0 1 1 -2 2 -8 8 0 -2 1 -6 24 56 0 2 2 -12 2 -4 -3 32 6 1 -8 8 57 1 0 12 -16 -4 -6 -1 8 6 0 6 32 58 1 3 1 -6 -4 -6 8 6 -3 -2 3	32 59 4 0 -16 -1 -1 4 24 8 -2 3 -16 32 60 0 1 -2 0 2 -8 12 -6 1 1 -6 16 61 0 0 0 -1 1 -4 4 0 0 0 -3 12 62 0 1 2 -8 2 -6 4 16 0 2 -6 8 63 1 -1 12
-16 -6 -2 3 6 6 -1 6 32		

[0077] Note that if one uses a representation where the fixed coefficients have already been converted to k.sub.1, k.sub.0, s and n, one can store the entire Table 5 using 6 bits per coefficients. Since the largest magnitude is 45, a 7-bit number (capable of holding values in the range [-64, 63]) would otherwise be needed. Hence one bit per stored value can be saved this way.

[0078] An alternative to using S.sub.96 and S is to use S.sub.127={0, ±1, ±2, ±3, ±4, ±6, ±8, ±12, ±16, ±24, ±32, ±48, ±64, ±96, ±127}. S.sub.127 is similar to S but uses ±127 instead of ±128. This, however, would make a hardware implementation more difficult because it would have to be able to handle multiplication a*b where a=127, which can be done relatively cheaply using (b<<7)-b. This could be added as a step after FIG. 8. This embodiment has been tested under the common test conditions (CTC) for VTM 6.0, the reference software for VVC version 6. The decoder was changed only by changing the fixed coefficients according to Table 5. The encoder was changed so that it quantized every coefficient to the nearest one in S.sub.127. The result was an increase in average bit rate difference (BD-rate) of about 0.1%, meaning that for the same quality, the bit rate increased 0.1%. The exact numbers for the different configurations were +0.09% (all intra) +0.10% (random access) +0.07% (low-delay B) +0.15% (low-delay P).

[0079] Although 0.1% may not seem as a big increase in bit rate, it would be better to have a smaller BD-rate penalty for the simplification. There are much fewer coefficients in S, S.sub.96 and S.sub.127 than in the currently allowed coefficient set T that includes every number between -127 and 127. Since the number of allowed coefficients differs so much, it makes sense to code them differently in the case of S and T. However, that means that the decoder must be changed in a normative way. This may be advantageous from another perspective as well: in the encoder-only implementations, the decoder always has to be able to fall back to a solution that can handle all types of coefficients if the encoder has not constrained the coefficients. This means that we need two implementations: one for non-restricted coefficients (set T) and one for the restricted coefficients (e.g., S, S.sub.96 or S.sub.127 dependent on implementation). Hence in hardware one would have to implement more hardware than if only restricted coefficients were used. In summary, an encoder-only solution might not provide many benefits.

1.1 More Efficient Coefficient Encoding

[0080] In one embodiment, the encoder is forced to always restrict the coefficients, for instance to S. This way, a hardware implementation can lower the complexity by implementing only the solution described in FIGS. 8-12. Likewise, a software implementation can gain speed if the software architecture is one where this is possible.

[0081] Since the decoder has to be changed anyway, it is possible to use a different encoding of the coefficients than is used in the current VVC draft. Currently, the magnitude abs(coeff) is first coded using 3-Exponential-Golomb coding as shown in Table 6.

TABLE-US-00006 TABLE 6 magnitude magnitude bits sign bit

0 1000 1 1001 0/1 2 1010 0/1 3 1011 0/1 4 1100 0/1 5 1101 0/1 6 1110 0/1 7 1111
0/1 8 010000 0/1 9 010001 0/1 10 010010 0/1 11 010011 0/1 12 010100 0/1 13 010101 0/1 14 010110 0/1 15 010111 0/1 16 011000 0/1 17 011001
0/1 18 011010 0/1 19 011011 0/1 20 011100 0/1 21 011101 0/1 22 011110 0/1 23 011111 0/1 24 00100000 0/1

[0082] Apart from the magnitude, the sign will also be encoded/decoded for all values except 0. This means that 0 is represented by four bits, ±1, ±2, ±3, ±4, ±5, ±6 and ±7 are represented by five bits, values from ±8 through ±24 are represented by seven bits, etc.

[0083] Since we do not need to represent most of these values when we restrict the coefficients to the set S, in one embodiment we instead use the

truncated binary coding according to the index of the coefficients magnitude according to Table 7:

TABLE-US-00007 TABLE 7 magnitude index index bits sign bit 0 0 000 1 1 0010 0/1 2 2 0011 0/1 3 3 0100 0/1 4 4 0101 0/1 6 5 0110 0/1 8 6 0111 0/1 12 7 1000 0/1 16 8 1001 0/1 24 9 1010 0/1 32 10 1011 0/1 48 11 1100 0/1 64 12 1101 0/1 96 13 1110 0/1 128 14 1111 0/1

[0084] By comparing Table 6 with Table 7, we see that the encoding in Table 7 always uses the same number of bits or fewer bits. Hence, we will always save bits if encoding and decoding according to Table 7. This also turns out to be the case in practice; when testing on the CTC for VTM6.0 we get the following BD-rate numbers: +0.03% (all intra) +0.04% (random access) +0.00% (low-delay B) +0.02% (low-delay P). Thus most of the penalty is gone.

1.2 Further Reducing the Allowed Set of Coefficients

[0085] When analyzing the bit streams obtained in the previous test, it is clear that the two largest magnitudes, 96 and 128, are very rarely used. Therefore, in an alternative embodiment, a further restriction is used, allowing only coefficients in the following set: S.sub.64={-64, -48, -32, -24, -16, -12, -8, -6, -4, -3, -2, -1, 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64}. Since there are now fewer magnitudes, it is possible to reduce the number of bits for the smallest magnitudes, as is shown in Table 8.

TABLE-US-00008 TABLE 8 magnitude index index bits sign bit 0 0 000 1 1 001 0/1 2 2 010 0/1 3 3 0110 0/1 4 4 0111 0/1 6 5 1000 0/1 8 6 1001 0/1 12 7 1010 0/1 16 8 1011 0/1 24 9 1100 0/1 32 10 1101 0/1 48 11 1110 0/1 64 12 1111 0/1

[0086] Table 8 shows that in one embodiment magnitudes 96 and 128 are not allowed. This makes it possible to use shorter codes for magnitude 0, 1, and 2.

[0087] Trying this version on the CTC for VTM7.0 gives the following BD-rate figures: +0.02% (all intra) +0.03% (random access). Thus the penalty for quantizing coefficients has been further reduced in terms of BD-rate.

[0088] Obtaining the values for the variables k.sub.0, k.sub.1, n and s can be done using the C-like pseudo-code in Table 9.

TABLE-US-00009 TABLE 9 char s_from_index[7] = {0, 0, 1, 2, 3, 4, 5}; xReadTruncBinCode(index, 13); //read index
if(index==0) READ_FLAG(n, "variable n"); s = s_from_index[index>>1]; k1 = (index < 2 ? 0 : 1); k0 = index % 1;

[0089] As an example, the xReadTruncBinCode() function reads the bits 1010 which, according to Table 8 gives an index of 7. (This, according to Table 8, is indicative of a magnitude of 12.) Hence the value 7 is put into the index variable. Since index is not 0, the code proceeds to read one bit using READ_FLAG(n, "variable n") and puts the result in n. Assume it gets n=1. That indicates that the sign is negative. The coefficient to use is thus -12. The shift value s used is found in the array s_from_index, specifically by using the 7>>1=3 as index to the array. This means that s will become 2. Next, since index>2, k1 will be 1. Finally, k0 will be set to 7%1 which equals 1. We thus have finished decoding the necessary values n=1, s=2, k0=1 and k1=1. We can now double-check with Equation 5a that this indeed gives the correct coefficient value -12: coeff=(-1).sup.n(2k.sub.1+k.sub.0) 2.sup.s=(-1).sup.1(2*1+1)*2.sup.2=(-1)*3*4=-12.

1.3a Using Signed Truncated Coding for the Coefficients

[0090] In another embodiment it is possible to use signed truncated coding for the coefficients. Table 10A shows how the coefficients may be coded in such an embodiment.

TABLE-US-00010 TABLE 10A bit signed representation index value coefficient 0000 0 0 0 0001 1 -1 -1 0010 2 1 1 0011 3 -2 -2 0100 4 2 2 0101 5 -3 -3 0110 6 3 3 0111 7 -4 -4 0111 8 4 4 10000 9 -5 -6 10001 10 5 6 10010 11 -6 -8 10011 12 6 8 10100 13 -7 -12 10101 14 7 12 10110 15 -8 -16 10111 16 8 16 11000 17 -9 -24 11001 18 9 24 11010 19 -10 -32 11011 20 10 32 11100 21 -11 -48 11101 22 11 48 11110 23 -12 -64 11111 24 12 64

[0091] The coefficients could be recovered using the following pseudo-code:

TABLE-US-00011 TABLE 11A char magtab[13] = {0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64};
xReadTruncBinCode(index, 25); // read index sign = (-1) * (index & 1); magnitude = (index+1) >> 1; coefficient = sign*magtab [magnitude];

1.3b Using Fixed Length Coding for the Coefficients

[0092] In another embodiment it is possible to use fixed length coding for the coefficients. Table 10B shows how the coefficients may be coded in such an embodiment.

TABLE-US-00012 TABLE 10B magnitude index index bits sign bit 0 0 0000 1 1 0001 0/1 2 2 0010 0/1 3 3 0011 0/1 4 4 0100 0/1 6 5 0101 0/1 8 6 0110 0/1 12 7 0111 0/1 16 8 1000 0/1 24 9 1001 0/1 32 10 1010 0/1 48 11 1011 0/1 64 12 1100 0/1 96 13 1101 0/1

[0093] The coefficients could be recovered using the following pseudo-code:

TABLE-US-00013 TABLE 11B char s_from_index[7] = {0, 0, 1, 2, 3, 4, 5}; xReadFixedLength(index, 4); // read 4 bits
if(index==0) READ_FLAG(n, "variable n"); s = s_from_index[index>>1]; k1 = (index < 2 ? 0 : 1); k0 = index % 1;

[0094] Since there are two more possible codewords (1110 and 1111) it would be possible to also accommodate a magnitude of 128 and even 192. It is also possible to restrict the coding so that 64 becomes the largest magnitude.

[0095] Alternatively, the variables k.sub.0, k.sub.1, n and s can be directly recovered from the index using the following pseudo code:

TABLE-US-00014 TABLE 12 char s_from_index[7] = {0, 0, 1, 2, 3, 4, 5}; xReadTruncBinCode(index, 25); // read index
n = (index & 1); s = s_from_index[index>>2]; k1 = (index < 3 ? 0 : 1); k0 = ((index+1)>>1) % 1;

1.4 Allowing Power-of-Two Multiples of 0, 1, 3 and 5.

[0096] In some circumstances it may be limiting to constrain the coefficients to only be of the form $\pm\{0, 1, 3\} \times 2^{\text{sup.n}}$. Most coefficients are close to zero, which means that it is most important to be able to represent coefficients close to zero, such as $\{0, \pm 1, \pm 2, \pm 3, \pm 4, \pm 5, \pm 6, \pm 7, \pm 8, \pm 9, \pm 10\}$. Out of these only $\{0, \pm 1, \pm 2, \pm 3, \pm 4, \pm 6, \pm 8\}$ are possible to represent on the form $\pm\{0, 1, 3\} \times 2^{\text{sup.n}}$. However, if we also allow $5 \times 2^{\text{sup.n}}$, we can also represent ± 5 and ± 10 . As it turns out, it is not much more expensive to create hardware that allows for $\pm\{0, 1, 3, 5\} \times 2^{\text{sup.n}}$ than it is to create hardware that allows for $\pm\{0, 1, 3\} \times 2^{\text{sup.n}}$. The reason for this is that, just as for the factor 3, multiplying a number by 5 can also be implemented using a single addition and shifts, since $5x=4x+x=(x<<2)+x$.

[0097] In general, we can modify Equation 5b so that we will be able to incorporate also a multiplication $a*b$ when $a=5$:

$$[00011] a * b = (-1)^n (((b \&_b k_1) \ll s_0) + b \&_b k_0) 2^{s_1} \quad (\text{Eqn9b})$$

[0098] The difference compared to Equation 5 is that, instead of always shifting 1 step, we now shift 1 or 2 steps, controlled by the variable s.sub.0. Another change compared to Equation 5 is that the variable s has changed name to s.sub.1. FIG. 13 shows how such a hardware implementation can be constructed.

[0099] When comparing the diagram in FIG. 13 to the one in FIG. 8 we see that the major difference is the box marked 803. This now shifts the value b &sub.b k.sub.1 left either one or two steps, whereas in FIG. 8 it always shifts left one step. This means that both $3*b=(b<<1)+b$ and $5*b=(b<<2)+b$ can be constructed as output of the adder 805. This shift is controlled by the shift value s.sub.0, which shifts 1 step if s.sub.0=0 and two steps if s.sub.0=1. The box 803 can be implemented inexpensively using 13 1-bit muxes as shown in FIG. 14.

[0100] The other difference against FIG. 8 is that the output of 802 is now sign-extended two bits to go from a signed 12-bit number to a signed 14-bit number using the wiring marked with 804. It is now possible to use a coefficient set such as S.sub.135= $\pm\{0, 1, 2, 3, 4, 5, 6, 8, 10, 12, 16, 20, 24, 32, 40, 48, 64\}$.

[0101] Table 13 shows what values to set for k.sub.1, k.sub.2, s.sub.0 and s.sub.1 to obtain the positive coefficients in S.sub.135. (The value n is 0 for

the positive coefficients.) The values for k.sub.1, k.sub.2, s.sub.0 and s.sub.1 for the negative coefficients are the same as for the positive coefficients, but n=1.

TABLE-US-00015 TABLE 13 coefficient k₁ k₀ s₀ s₁ 0 0 0 0 0 1 0 1 0 0 2 1 0 0 0 3 1 1 0 0 4 1 0 0 1 5 1 1 1 0 6 1 1 0 1 8 1 0 0 2 10 1 1 1 1 12 1 1 0 2 16 1 0 0 3 20 1 1 1 2 24 1 1 0 3 32 1 0 0 4 40 1 1 1 3 48 1 1 0 4 64 1 0 0 5

[0102] This gives the following results when evaluated using the CTC for VTM7.0: 0.00% (all intra) 0.00% (random access). Hence there is no longer any penalty compared to the original ALF. On the other hand, the gain of about 0.03% may not be enough to motivate the extra hardware needed.

1.5 Embodiments that Take Coefficient Statistics into Account

[0103] As described earlier, in the original VVC version of ALF, coefficients of smaller magnitudes are typically more common than coefficients of larger magnitudes. However, now that we have used a representation with increasing gaps between the allowed coefficient magnitudes, it may very well be the case that a value of 48 is as common as a value of 4. This is due to the fact that all values that are between 40 and 56 before quantization assume the value 48, whereas only values 4 and 5 may be quantized to 4. In FIG. 15A we give the results for the frequency distribution for coefficient C0 from embodiment 1.1. A shorthand is used for the different coefficient values described in Table 14.

TABLE-US-00016 TABLE 14 sign + shorthand magnitude shorthand bits sign bit -128 -14 1111 1 -96 -13 1110 1 -64 -12 1101 1 -48 -11 1100 1 -32 -10 1011 1 -24 -9 1010 1 -16 -8 1001 1 -12 -7 1000 1 -8 -6 0111 1 -6 -5 0110 1 -4 -4 0101 1 -3 -3 0100 1 -2 -2 0011 1 -1 -1 0010 1 0 0 000 1 1 0010 0 2 2 0011 0 3 3 0100 0 4 4 0101 0 6 5 0110 0 8 6 0111 0 12 7 1000 0 16 8 1001 0 24 9 1010 0 32 10 1011 0 48 11 1100 0 64 12 1101 0 96 13 1110 0 128 14 1111 0

[0104] Plotting the values used for coefficient C0 in embodiment 1.1 gives the result shown in FIG. 15A. That is, FIG. 15A shows a frequency plot of different coefficient code values. Here -14 represents -128, -13 represents -96, . . . , 14 represents +128. Note that values -1 and 0 are the most common coefficients.

[0105] As can be seen in FIG. 15A, small magnitudes are still much more common than large ones for C0. However, this picture changes when looking at C11, which has the distribution shown in FIG. 15B. That is, FIG. 15B shows a frequency distribution of coefficient 11. Note that the most common shorthand value is +9, which corresponds to a coefficient value of +24. Other common values are +16 and +32.

[0106] Here it is seen that the most common coefficient values are 12, 16, 24 and 32. Even so the coefficient value that requires the fewest number of bits to code is still 0, which is rather uncommon for coefficient 11. It would be better if 24 (shorthand +9) would instead have the shortest codeword. That would be especially good if combined with the embodiment described in 1.2, since it has more codewords that are short. Hence in yet another embodiment, we subtract 9 from the shorthand before encoding it according to Table 15:

TABLE-US-00017 TABLE 15 Shorthand Shorthand Shorthand sign magnitude before shift after shift bits bit -64 -12 4 0111 0 -48 -11 5 1000 0 -32 -10 6 1001 0 -24 -9 7 1010 0 -16 -8 8 1011 0 -12 -7 9 1100 0 -8 -6 10 1101 0 -6 -5 11 1110 0 -4 -4 12 1111 0 -3 -3 -12 1111 1 -2 -2 -11 1110 1 -1 -1 -10 1101 1 0 0 -9 1100 1 1 1 -8 1011 1 2 2 -7 1010 1 3 3 -6 1001 1 4 4 -5 1000 1 6 5 -4 0111 1 8 6 -3 0110 1 12 7 -2 010 1 16 8 -1 001 1 24 9 0 000 32 10 1 001 0 48 11 2 010 0 64 12 3 0110 0

[0107] Table 15 shows that shifting the shorthand makes it possible to assign shorter codewords to more likely coefficients, such as the value +24 for coefficient 11, which gets encoded using 3 bits.

[0108] As can be seen in the table 15, the value +32, which is very common for coefficient 11, first gets assigned shorthand value 10. However, this value is shifted by subtracting 9 (the most common value for C11), using modulus calculation:

[00012] shiftedshorthand = (((shorthand + 12) - 9) mod 25) - 12. (Eqn10)

[0109] For the shorthand value 10, this becomes ((10+12-9) mod 25)-12=(13 mod 25)-12=1. This shorthand value is encoded as 001 0, which is only four bits. The modulus calculation is used to avoid values outside the range [-12, 12].

[0110] Since the statistics differ between different coefficients, it is best to subtract a different value depending upon which coefficient we are encoding. Hence one may use:

[00013] shiftedshorthand = (((shorthand + 12) - offset_k) mod 25) - 12, (Eqn11)

where k depends on which coefficient we are encoding (k=0 for C0 etc) and offset.sub.k={0, 0, 6, -1, 0, -5, 7, 8, 7, 0, -6, 8}. This gave the following results when evaluated using the CTC for VTM6.0: 0.01% (all intra) 0.02% (random access) -0.03% (low delay B) 0.01% (low-delay P)

[0111] For the chroma components, the current version of ALF only uses 6 coefficients. Hence the best shift value to use is different between luma and chroma. As an example one can use the following shift values for the chroma coefficients: {-5, 8, 9, 8, -6, 8}. One then gets the following (luma) BD-rate results using the CTC for VTM6.0: 0.00% (all intra) 0.01% (random access) -0.03% (low-delay B) 0.00% (low-delay P). This has completely eliminated the penalty for all intra, low-delay B and low-delay P and has only a small penalty for random access.

1.6 Embodiments where Coefficients are not Encoded Using Magnitude Plus Sign

[0112] The ALF coefficient coding from embodiment 1.1 to embodiment 1.3 codes the coefficient magnitude (or magnitude index) and the coefficient sign separately. Here, the coefficient which has a magnitude of 0 is coded with shorter code (fewer bits) compared to a coefficient which has a magnitude that is larger than 0. Considering the coefficient statistic in embodiment 1.3, there is another way to code the ALF coefficient more efficiently by coding the index of the signed magnitude.

[0113] The index (shorthand) of the signed magnitude before shift ranges from 0, 1, . . . to 24, which represents the signed magnitude {0, 1, 2, . . . , 48, 64, -64, -48, . . . , -2, -1}. The index ranges from 0, 1, . . . to 24 are coded by truncated binary code with a maximum symbol of 25.

TABLE-US-00018 TABLE 16 shorthand value Shorthand bits 0 0 0000 1 1 0001 2 2 0010 3 3 0011 4 4 0100 5 5 0101 6 6 0110 7 7 0111 8 8 1000 9 9 1001 10 10 1010 11 11 1011 12 12 1100 13 13 1101 14 14 1110 15 15 1111 16 16 12 17 1200 13 13 1201 14 14 1210 15 15 1211 16 16 13 17 1300 14 14 1301 15 15 1310 16 16 1311 17 17 14 18 1400 15 15 1401 16 16 1410 17 17 1411 18 18 15 19 1500 16 16 1501 17 17 1510 18 18 1511 19 19 16 20 1600 17 17 1601 18 18 1610 19 19 1611 20 20 17 21 1700 18 18 1701 19 19 1710 20 20 1711 21 21 18 22 1800 19 19 1801 20 20 1810 21 21 1811 22 22 19 23 1900 20 20 1901 21 21 1910 22 22 1911 23 23 20 24 2000 21 21 2001 22 22 2010 23 23 2011 24 24 21 25 2100 22 22 2101 23 23 2110 24 24 2111 25 25 22 26 2200 23 23 2201 24 24 2210 25 25 2211 26 26 23 27 2300 24 24 2301 25 25 2310 26 26 2311 27 27 24 28 2400 25 25 2401 26 26 2410 27 27 2411 28 28 25 29 2500 26 26 2501 27 27 2510 28 28 2511 29 29 26 30 2600 27 27 2601 28 28 2610 29 29 2611 30 30 27 31 2700 28 28 2701 29 29 2710 30 30 2711 31 31 28 32 2800 29 29 2801 30 30 2810 31 31 2811 32 32 29 33 2900 30 30 2901 31 31 2910 32 32 2911 33 33 30 34 3000 31 31 3001 32 32 3010 33 33 3011 34 34 31 35 3100 32 32 3101 33 33 3110 34 34 3111 35 35 32 36 3200 33 33 3201 34 34 3210 35 35 3211 36 36 33 37 3300 34 34 3301 35 35 3310 36 36 3311 37 37 34 38 3400 35 35 3401 36 36 3410 37 37 3411 38 38 35 39 3500 36 36 3501 37 37 3510 38 38 3511 39 39 36 40 3600 37 37 3601 38 38 3610 39 39 3611 40 40 37 41 3700 38 38 3701 39 39 3710 40 40 3711 41 41 38 42 3800 39 39 3801 40 40 3810 41 41 3811 42 42 39 43 3900 40 40 3901 41 41 3910 42 42 3911 43 43 40 44 4000 41 41 4001 42 42 4010 43 43 4011 44 44 41 45 4100 42 42 4101 43 43 4110 44 44 4111 45 45 42 46 4200 43 43 4201 44 44 4210 45 45 4211 46 46 43 47 4300 44 44 4301 45 45 4310 46 46 4311 47 47 44 48 4400 45 45 4401 46 46 4410 47 47 4411 48 48 45 49 4500 46 46 4501 47 47 4510 48 48 4511 49 49 46 50 4600 47 47 4601 48 48 4610 49 49 4611 50 50 47 51 4700 48 48 4701 49 49 4710 50 50 4711 51 51 48 52 4800 49 49 4801 50 50 4810 51 51 4811 52 52 49 53 4900 50 50 4901 51 51 4910 52 52 4911 53 53 50 54 5000 51 51 5001 52 52 5010 53 53 5011 54 54 51 55 5100 52 52 5101 53 53 5110 54 54 5111 55 55 52 56 5200 53 53 5201 54 54 5210 55 55 5211 56 56 53 57 5300 54 54 5301 55 55 5310 56 56 5311 57 57 54 58 5400 55 55 5401 56 56 5410 57 57 5411 58 58 55 59 5500 56 56 5501 57 57 5510 58 58 5511 59 59 56 60 5600 57 57 5601 58 58 5610 59 59 5611 60 60 57 61 5700 58 58 5701 59 59 5710 60 60 5711 61 61 58 62 5800 59 59 5801 60 60 5810 61 61 5811 62 62 59 63 5900 60 60 5901 61 61 5910 62 62 5911 63 63 60 64 6000 61 61 6001 62 62 6010 63 63 6011 64 64 61 65 6100 62 62 6101 63 63 6110 64 64 6111 65 65 62 66 6200 63 63 6201 64 64 6210 65 65 6211 66 66 63 67 6300 64 64 6301 65 65 6310 66 66 6311 67 67 64 68 6400 65 65 6401 66 66 6410 67 67 6411 68 68 65 69 6500 66 66 6501 67 67 6510 68 68 6511 69 69 66 70 6600 67 67 6601 68 68 6610 69 69 6611 70 70 67 71 6700 68 68 6701 69 69 6710 70 70 6711 71 71 68 72 6800 69 69 6801 70 70 6810 71 71 6811 72 72 69 73 6900 70 70 6901 71 71 6910 72 72 6911 73 73 70 74 7000 71 71 7001 72 72 7010 73 73 7011 74 74 71 75 7100 72 72 7101 73 73 7110 74 74 7111 75 75 72 76 7200 73 73 7201 74 74 7210 75 75 7211 76 76 73 77 7300 74 74 7301 75 75 7310 76 76 7311 77 77 74 78 7400 75 75 7401 76 76 7410 77 77 7411 78 78 75 79 7500 76 76 7501 77 77 7510 78 78 7511 79 79 76 80 7600 77 77 7601 78 78 7610 79 79 7611 80 80 77 81 7700 78 78 7701 79 79 7710 80 80 7711 81 81 78 82 7800 79 79 7801 80 80 7810 81 81 7811 82 82 79 83 7900 80 80 7901 81 81 7910 82 82 7911 83 83 80 84 8000 81 81 8001 82 82 8010 83 83 8011 84 84 81 85 8100 82 82 8101 83 83 8110 84 84 8111 85 85 82 86 8200 83 83 8201 84 84 8210 85 85 8211 86 86 83 87 8300 84 84 8301 85 85 8310 86 86 8311 87 87 84 88 8400 85 85 8401 86 86 8410 87 87 8411 88 88 85 89 8500 86 86 8501 87 87 8510 88 88 8511 89 89 86 90 8600 87 87 8601 88 88 8610 89 89 8611 90 90 87 91 8700 88 88 8701 89 89 8710 90 90 8711 91 91 88 92 8800 89 89 8801 90 90 8810 91 91 8811 92 92 89 93 8900 90 90 8901 91 91 8910 92 92 8911 93 93 90 94 9000 91 91 9001 92 92 9010 93 93 9011 94 94 91 95 9100 92 92 9101 93 93 9110 94 94 9111 95 95 92 96 9200 93 93 9201 94 94 9210 95 95 9211 96 96 93 97 9300 94 94 9301 95 95 9310 96 96 9311 97 97 94 98 9400 95 95 9401 96 96 9410 97 97 9411 98 98 95 99 9500 96 96 9501 97 97 9510 98 98 9511 99 99 96 100 9600 97 97 9601 98 98 9610 99 99 9611 100 100 97 101 9700 98 98 9701 99 99 9710 100 100 9711 101 101 98 102 9800 99 99 9801 100 100 9810 101 101 9811 102 102 99 103 9900 100 100 9901 101 101 9910 102 102 9911 103 103 100 104 10000 101 101 10001 102 102 10010 103 103 10011 104 104 101 105 10100 102 102 10101 103 103 10110 104 104 10111 105 105 102 106 10200 103 103 10201 104 104 10210 105 105 10211 106 106 103 107 10300 104 104 10301 105 105 10310 106 106 10311 107 107 104 108 10400 105 105 10401 106 106 10410 107 107 10411 108 108 105 109 10500 106 106 10501 107 107 10510 108 108 10511 109 109 106 110 10600 107 107 10601 108 108 10610 109 109 10611 110 110 107 111 10700 108 108 10701 109 109 10710 110 110 10711 111 111 108 112 10800 109 109 10801 110 110 10810 111 111 10811 112 112 109 113 10900 110 110 10901 111 111 10910 112 112 10911 113 113 110 114 11000 111 111 11001 112 112 11010 113 113 11011 114 114 111 115 11100 112 112 11101 113 113 11110 114 114 11111 115 115 112 116 11200 113 113 11201 114 114 11210 115 115 11211 116 116 113 117 11300 114 114 11301 115 115 11310 116 116 11311 117 117 114 118 11400 115 115 11401 116 116 11410 117 117 11411 118 118 115 119 11500 116 116 11501 117 117 11510 118 118 11511 119 119 116 120 11600 117 117 11601 118 118 11610 119 119 11611 120 120 117 121 11700 118 118 11701 119 119 11710 120 120 11711 121 121 118 122 11800 119 119 11801 120 120 11810 121 121 11811 122 122 119 123 11900 120 120 11901 121 121 11910 122 122 11911 123 123 120 124 12000 121 121 12001 122 122 12010 123 123 12011 124 124 121 125 12100 122 122 12101 123 123 12110 124 124 12111 125 125 122 126 12200 123 123 12201 124 124 12210 125 125 12211 126 126 123 127 12300 124 124 12301 125 125 12310 126 126 12311 127 127 124 128 12400 125 125 12401 126 126 12410 127 127 12411 128 128 125 129 12500 126 126 12501 127 127 12510 128 128 12511 129 129 126 130 12600 127 127 12601 128 128 12610 129 129 12611 130 130 127 131 12700 128 128 12701 129 129 12710 130 130 12711 131 131 128 132 12800 129 129 12801 130 130 12810 131 131 12811 132 132 129 133 12900 130 130 12901 131 131 12910 132 132 12911 133 133 130 134 13000 131 131 13001 132 132 13010 133 133 13011 134 134 131 135 13100 132 132 13101 133 133 13110 134 134 13111 135 135 132 136 13200 133 133 13201 134 134 13210 135 135 13211 136 136 133 137 13300 134 134 13301 135 135 13310 136 136 13311 137 137 134 138 13400 135 135 13401 136 136 13410 137 137 13411 138 138 135 139 13500 136 136 13501 137 137 13510 138 138 13511 139 139 136 140 13600 137 137 13601 138 138 13610 139 139 13611 140 140 137 141 13700 138 138 13701 139 139 13710 140 140 13711 141 141 138 142 13800 139 139 13801 140 140 13810 141 141 13811 142 142 139 143 13900 140 140 13901 141 141 13910 142 142 13911 143 143 140 144 14000 141 141 14001 142 142 14010 143 143 14011 144 144 141 145 14100 142 142 14101 143 143 14110 144 144 14111 145 145 142 146 14200 143 143 14201 144 144 14210 145 145 14211 146 146 143 147 14300 144 144 14301 145 145 14310 146 146 14311 147 147 144 148 14400 145 145 14401 146 146 14410 147 147 14411 148 148 145 149 14500 146 146 14501 147 147 14510 148 148 14511 149 149 146 150 14600 147 147 14601 148 148 14610 149 149 14611 150 150 147 151 14700 148 148 14701 149 149 14710 150 150 14711 151 151 148 152 14800 149 149 14801 150 150 14810 151 151 14811 152 152 149 153 14900 150 150 14901 151 151 14910 152 152 14911 153 153 150 154 15000 151 151 15001 152 152 15010 153 153 15011 154 154 151 155 15100 152 152 15101 153 153 15110 154 154 15111 155 155 152 156 15200 153 153 15201 154 154 15210 155 155 15211 156 156 153 157 15300 154 154 15301 155 155 15310 156 156 15311 157 157 154 158 15400 155 155 15401 156 156 15410 157 157 15411 158 158 155 159 15500 156 156 15501 157 157 15510 158 158 15511 159 159 156 160 15600 157 157 15601 158 158 15610 159 159 15611 160 160 157 161 15700 158 158 15701 159 159 15710 160 160 15711 161 161 158 162 15800 159 159 15801 160 160 15810 161 161 15811 162 162 159 163 15900 160 160 15901 161 161 15910 162 162 15911 163 163 160 164 16000 161 161 16001 162 162 16010 163 163 16011 164 164 161 165 16100 162 162 16101 163 163 16110 164 164 16111 165 165 162 166 16200 163 163 16201 164 164 16210 165 165 16211 166 166 163 167 16300 164 164 16301 165 165 16310 166 166 16311 167 167 164 168 16400 165 165 16401 166 166 16410 167 167 16411 168 168 165 169 16500 166 166 16501 167 167 16510 168 168 16511 169 169 166 170 16600 167 167 16601 168 168 16610 169 169 16611 170 170 167 171 16700 168 168 16701 169 169 16710 170 170 16711 171 171 168 172 16800 169 169 16801 170 170 16810 171 171 16811 172 172 169 173 16900 170 170 16901 171 171 16910 172 172 16911 173 173 170 174 17000 171 171 17001 172 172 17010 173 173 17011 174 174 171 175 17100 172 172 17101 173 173 17110 174 174 17111 175 175 172 176 17200 173 173 17201 174 174 17210 175 175 17211 176 1

-62, -56, -50, -44, -38, -32, -26, -20, -14, -8, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127}. The APS ALF coefficients coding is same as VTM7.0, as: a) 3-order Exponential-Golomb coding for coefficients magnitude; b) 1 bit coefficient sign coding if the coefficient is not equal to 0.

[0120] A multiplication $a*b$ where a belongs to set Z can be written as

$$[00014] \ a * b = (-1)^n (((b \&_b k_1) \ll s_0) + (-1)^c (b \&_b k_0)) 2^{s_1} \quad (\text{Eqn12})$$

[0121] As an example, $62*b$ can be written as: $(-1).sup.0 (((b \&.sub.b 1) \ll 5) + (-1).sup.1 (b \&.sub.b 1)) 2.sup.1$ since that evaluates to $((b \ll 5) - b) 2.sup.1 = (32b - b) * 2 = 31b * 2 = 62b$. Equation 12 can be inexpensively be implemented using the hardware depicted in FIG. 16.

[0122] Compared to FIG. 13, there are two major differences. The unit **903** (compare to **803** in FIG. 13) can now shift between 1 and 7 steps to the left, instead of just 1 or 2 steps. Furthermore, there is a new box in **908**. This box inverts all bits if $c=1$, otherwise it just passes them through. This can be inexpensively implemented by XORing with c , just as is done in the left part of FIG. 11. Furthermore the bit c is used as carry in to the adder **905**. Together with the conditional inverter **908**, this has the effect of negating the expression $(b \&.sub.b k.sub.0)$ if $c=1$, and hence implements $(-1).sup.c (b \&.sub.b k.sub.0)$ from Equation 12.

[0123] This gives the following results when evaluated using the CTC for VTM7.0: 0.01% (all intra) 0.01% (random access).

1.8 Embodiments where Coefficients Belong to a Subset of Z are Used as ALF Filter Coefficients

[0124] In all previous embodiments, the ALF filter coefficients belong to a subset of Z .

[0125] One example in this embodiment, the ALF filter coefficients belong to set $Z.sub.sub=\{-40, -33, -28, -24, -20, -17, -15, -14, -12, -10, -9, -8, -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 17, 20, 24, 28, 33, 40\}$. The APS ALF coefficients coding uses the truncated binary coding to code the index of the coefficient magnitude and 1-bit coefficient sign coding if the coefficient is not equal to 0:

TABLE-US-00021 TABLE 19 magnitude index index bits sign bit 0 0000 1 0001 0/1 2 0010 0/1 3 0011 0/1 4 0100 0/1 5 0101 0/1 6 0110 0/1 7 0111 0/1 8 1000 0/1 9 1001 0/1 10 1010 0/1 11 1011 0/1 12 1100 0/1 13 1101 0/1 14 11010 0/1 15 11011 0/1 16 11100 0/1 17 11101 0/1 18 11110 0/1 19 11111 0/1

[0126] Since $Z.sub.sub$ is a subset of Z , it is possible to use the hardware implementation in FIG. 16 to implement the multiplications. The variables $k.sub.1$, $k.sub.0$, $s.sub.0$, $s.sub.1$, c and n can be obtained using look-up from Table 20. Note that this can take place at the time of decoding the transform coefficients and hence only needs to happen once per CTU, not once per sample.

TABLE-US-00022 TABLE 20 (Values that can be used for k_1 , k_0 , s_0 , s_1 , c and n when the coefficient belongs to $Z.sub.sub$.) coefficient k_1 k_0 s_0 s_1 c n -40 1 1 1 3 0 1 -33 1 1 4 0 0 1 -28 1 1 2 2 1 1 -24 1 1 0 3 0 1 -20 1 1 1 2 0 1 -17 1 1 3 0 0 1 -15 1 1 3 0 1 1 -14 1 1 2 1 1 1 -12 1 1 0 2 0 1 -10 1 1 1 1 0 1 -9 1 1 1 1 1 1 -8 1 0 1 1 0 1 -7 1 1 2 0 1 1 -6 1 1 0 1 0 1 -5 1 1 1 0 0 1 -4 1 0 1 0 0 1 -3 1 1 0 0 0 1 -2 1 0 0 0 0 1 -1 0 1 0 0 0 1 0 0 0 0 0 0 1 0 1 0 0 0 0 2 1 0 0 0 0 0 3 1 1 0 0 0 0 4 1 0 1 0 0 0 5 1 1 1 0 0 0 6 1 1 0 1 0 0 7 1 1 2 0 1 0 8 1 0 1 1 0 0 9 1 1 1 1 0 10 1 1 1 1 0 11 1 1 0 0 12 1 1 0 2 0 0 14 1 1 2 1 1 0 15 1 1 3 0 1 0 17 1 1 3 0 0 0 20 1 1 1 2 0 0 24 1 1 0 3 0 0 28 1 1 2 2 1 0 33 1 1 4 0 0 0 40 1 1 1 3 0 0

[0127] This gives the following results when evaluated using the CTC for VTM7.0: -0.01% (all intra) -0.01% (random access).

1.9 Embodiments that Treat Coefficients Differently

[0128] In FIG. 15A the frequencies of different coefficient sizes for an embodiment that already has restricted the coefficients to belong to the set S is plotted. It is also possible to plot the frequencies for the non-altered code in VTM-6.0, where the coefficients are allowed all values between $[-127, 127]$, i.e., they belong to the set T . FIG. 17 shows the values of coefficient C_0 , C_1 , C_6 and C_{11} respectively. FIG. 17: Two of the coefficients that are far away from the center are C_0 and C_1 (top), and they have a rather different statistics than the two coefficients closest to the center, C_6 and C_{11} (bottom).

[0129] The coefficients C_0 and C_1 , whose frequency plots are depicted in the top half of FIG. 17, are far away from the center sample, as can be seen in FIG. 1.

[0130] Note that their statistics is quite different from the statistics of C_6 and C_{11} (bottom part of FIG. 17), which are the closest ones to the center sample. Two things stand out: First, whereas C_0 and C_1 are clustered around the value 0, C_6 has its peak at $C_6=+7$ and C_{11} has its peak at $C_{11}=+20$. Second—the distributions of C_6 and C_{11} are much flatter than for C_0 and C_1 —the peak goes up to around 2000 compared to over 7000 for C_0 and C_1 .

[0131] The first fact, that the average value of C_6 and C_{11} are larger than 0, means that the most common values will be heavily quantized. As an example, if we can only represent coefficient C_{11} with a value from $S.sub.64=\{0, \pm 1, \pm 2, \pm 3, \pm 4, \pm 6, \pm 8, \pm 12, \pm 16, \pm 24, \pm 32, \pm 48, \pm 64\}$, we cannot reach the most common value 20. This means that we have to choose between too weak a filtering using $C_{11}=16$ or too strong a filtering using $C_{11}=24$.

[0132] The second fact, i.e., that the distributions of C_6 and C_{11} are flatter than for the other coefficients, means that higher values are more common in general. Unfortunately, if we can only represent these coefficients with values from the set $S.sub.64$, this means that the error will on average be larger for C_6 and C_{11} than for C_0 and C_1 . As an example, if C_0 would never go beyond $[-4, 4]$, there would be no error compared to the VTM-6.0 version, since all values between $[-4, 4]$ are available in $S.sub.64$. For C_6 and C_{11} the opposite is true—these coefficients are almost never near the zero-error region of $[-4, 4]$. Hence forcing C_6 and C_{11} to be a value in $S.sub.64$ will contribute much more to the error than forcing C_0 and C_1 to belong to $S.sub.64$.

[0133] Therefore, in one embodiment of the present invention, C_6 and C_{11} are allowed to assume any value in T , i.e., any value in $[-127, 127]$. All the other coefficients, i.e., C_0 - C_5 and C_7 - C_{10} will have to take a value a restricted subset of T , such as $S.sub.64$. This means that for a hardware implementation, the hardware circuit handling the multiplication by C_6 and C_{11} may be different from the hardware circuit handling the multiplication by C_0 - C_5 and C_7 - C_{10} . Hence, instead of replacing all 12 multiplications in Equation 1 with inexpensive addition-based hardware such as that depicted in FIG. 8, only 10 of these multiplications can be replaced. The remaining two multiplications, i.e., for C_6 and C_{11} , will have to be full multiplications capable of multiplying by any number in T . Implementing this in VTM-7.0 will give the following BDR figures: +0.01% (AI) and +0.02% (RA).

[0134] In another embodiment, only C_{11} will use a full multiplication, whereas C_0 - C_{10} (i.e., including C_6) will use restricted multiplication capable of only a subset such as $S.sub.64$.

[0135] In one embodiment, C_6 and C_{11} can take any number in T whereas C_0 - C_5 and C_7 - C_{10} will be restricted to $S.sub.POT=\{0, \pm 1, \pm 2, \pm 4, \pm 8, \pm 16, \pm 32, \pm 64, \pm 128\}$, i.e., only numbers that are either zero or can be written as a power of two. Implementing this in VTM-7.0 will give the following BDR figures: +0.07% (AI) and +0.10% (RA).

[0136] In yet another embodiment, C_6 and C_{11} can take a number in a restricted set such as $S.sub.64$ whereas C_0 - C_5 and C_7 - C_{10} can take a number in an even more restricted set such as $S.sub.POT$.

1.10 Embodiments that Use an Average Value

[0137] In another embodiment, instead of representing values close to 0 with a higher accuracy, values close to the average value for C_6 and C_{11} are represented with a higher accuracy.

[0138] As an example, take again Equation 1 and assume all coefficients except for C_{11} are zero. Then:

$$[00015] \ \text{sum} = C_{11} * [\text{clip}(s_{11}, R(x-1, y) - R(x, y)) + \text{clip}(s_{11}, R(x+1, y) - R(x, y))], \quad (\text{Eqn17})$$

and by letting b be the expression in square brackets, one gets:

$$[00016] \ \text{sum} = C_{11} * b \quad (\text{Eqn18})$$

[0139] As described above, the value C11 is constrained to be in a certain set, such as S.sub.64, while allowing b to take any value. However, assume that one uses $C11 = 16 + \Delta \cdot \text{sub.11}$, and that it is $\Delta \cdot \text{sub.11}$ that is signaled instead of C11. This means that one can write Equation 18 as

$$[00017] \text{ sum} = (16 + \text{sub.11}) * b = 16 * b + \text{sub.11} * b = (b \ll 4) + \text{sub.11} * b. \quad (\text{Eqn19})$$

[0140] Now, if $\Delta \cdot \text{sub.11}$ is restricted to S.sub.64, one can use the inexpensive hardware in FIG. 8 to calculate $\Delta \cdot \text{sub.11} * b$. One must then add ($b \ll 4$) to get the correct term, and that requires another addition, which is expensive. However, even when adding the cost of this extra addition to the hardware in FIG. 8, the total cost is still much less than a general multiplication.

[0141] Forcing C11 to be $16 + \Delta \cdot \text{sub.11}$ is equivalent to forcing C11 to be in the subset $S \cdot \text{sub.64} + 16 = \{-48, -32, -16, -8, 0, 4, 8, 10, 12, 13, 14, 15, 16, 17, 18, 19, 20, 22, 24, 28, 32, 40, 48, 64, 80\}$. Since this subset contains many more values close to 20 than does S.sub.64, the average error induced by forcing C11 to $S \cdot \text{sub.64} + 16$ will be much smaller than the average error induced by forcing C11 to S.sub.64.

[0142] Similarly, C6 may be set to $8 + \Delta \cdot 8$, which can be implemented similarly inexpensively. By implementing this approach for C6 and C11 in VTM-7.0 it is possible to reach the following BDR figures: +0.01% (AI) and +0.06% (RA). In one solution, every coefficient Cx is set to $i + \Delta \cdot \text{sub.x}$ where we have a bias value i that is either a power-of-two $\pm 2 \cdot \text{sup.nx}$ (positive or negative) or zero. In other implementations, it may be sufficient to have some of these bias values being non-zero.

[0143] FIG. 18A is a flow chart illustrating a process 1800, according to one embodiment, for decoding an image. Process 1800 may begin in steps s1802.

[0144] Step s1802 comprises obtaining a set of sample values associated with the image, the set of sample values comprising a first sample value.

[0145] Step s1804 comprises employing an adaptive loop filter (ALF) to filter the first sample value, wherein the ALF is operable to filter the first sample value using any set of N coefficient values in which each one of the N coefficient values is included in a set of M unique coefficient values, wherein N is greater than 1 and M is greater than or equal to N and further wherein i) the set of M unique coefficient values consists of the following unique values or consists of a subset of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128 (i.e., $Z+128$) and ii) the set of M unique coefficient values includes at least one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$.

[0146] Employing the ALF to filter the first sample value comprises the steps of: a) obtaining a first set of N coefficient values for use in filtering the first sample value and b) using the ALF to filter the first sample value using the obtained first set of N coefficient values and the set of sample values, thereby producing a first filtered sample value, and each coefficient value included in the obtained first set of N coefficient values is constrained such that the coefficient value must be equal to one of the values included in the set of M unique values.

[0147] In one embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48$, or 64 (i.e., S.sub.64).

[0148] In another embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64$, or 96 (i.e., S.sub.96).

[0149] In another embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64, 96$, or 127 (i.e., S.sub.127).

[0150] In another embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 8, 10, 12, 16, 20, 24, 32, 40, 48$, or 64 (i.e., S.sub.135).

[0151] In another embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64, 96$, or 128 (i.e., S).

[0152] In another embodiment, the set of M unique coefficient values consists of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 17, 20, 24, 28, 33$, or 40 (i.e., Ssub).

[0153] FIG. 18B is a flow chart illustrating a process 1850, according to one embodiment, for decoding an image. Process 1850 may begin in steps s1852.

[0154] Step s1852 comprises obtaining a set of sample values associated with the image, the set of sample values comprising a first sample value.

[0155] Step s1854 comprises obtaining an index value that points to a particular coefficient value group included within a set of M predefined coefficient value groups (e.g., $M=64$). Each coefficient value group included in the set of predefined coefficient value groups consists of N coefficient values, N being greater than 1. For each coefficient value group included in the set of predefined coefficient value groups, each coefficient value included in the coefficient group is constrained such that the coefficient value must be equal to one of the following values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128 (i.e., $Z+128$). Also, for at least one coefficient value group included in the set of predefined coefficient value groups, at least one of the coefficient values included in said at least one coefficient value group is equal to one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126$, or 127.

[0156] Step s1856 comprises using the index value to select the particular coefficient value group from the set of predefined coefficient value groups.

[0157] Step s1858 comprises employing an adaptive loop filter (ALF) to filter the first sample value using the particular coefficient value group selected from the set of predefined coefficient value groups.

[0158] FIG. 20 is a flow chart illustrating a process 2000, according to one embodiment, that is performed by encoder 302. Process 2000 may begin in steps s2002.

[0159] Step s2002 comprises the encoder selecting a set of coefficient values for use by a decoder in filtering a sample value, the selected set of coefficient values consisting of N coefficient values. Each one of the N coefficient values is included in a set of M unique coefficient values, wherein N is greater than 1 and M is greater than 1 and further wherein i) the set of M unique coefficient values consists of the following unique values or consists of a subset of the following unique values: $\pm 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 14, 15, 16, 17, 18, 20, 24, 28, 30, 31, 32, 33, 34, 36, 40, 48, 56, 60, 62, 63, 64, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126, 127$, or 128 (i.e., $Z+128$) and ii) the set of M unique coefficient values includes at least one of the following values: $\pm 3, 5, 6, 7, 9, 10, 12, 14, 15, 17, 18, 20, 24, 28, 30, 31, 33, 34, 36, 40, 48, 56, 60, 62, 63, 65, 66, 68, 72, 80, 96, 112, 120, 124, 126$, or 127, and each coefficient value included in the set of N coefficient values is constrained such that the coefficient value must be equal to one of the values included in the set of M unique values.

[0160] Step s2004 comprises the encoder providing to a decoder (304) the N coefficient values or an initial index value for use by the encoder to determine the set of N coefficient values.

[0161] In some embodiment process 2000 also includes the step of determining a class to which the first sample value belongs, and the step of obtaining the index value comprises obtaining the index value using an initial index value signaled by an encoder and information identifying the determined class. For example, the initial index value may point to a particular set of N index values, where each one of the N index values is associated with a different class, and the decoder obtains the index value by obtaining the index value from the set of N index value that is associated with the determined class.

[0162] FIG. 19 is a block diagram of an apparatus 1901 for implementing encoder 302 or decoder 304, according to some embodiments. That is, apparatus 1901 can be configured to perform the methods disclosed herein. In embodiments where apparatus 1901 implements encoder 302, apparatus 1901 may be referred to as “encoding apparatus 1901,” and in embodiments where apparatus 1901 implements decoder 304, apparatus 1901 may be referred to as a “decoding apparatus 1901.” As shown in FIG. 19, apparatus 1901 may comprise: processing circuitry (PC) 1902, which

may include one or more processors (P) **1955** (e.g., one or more general purpose microprocessors and/or one or more other processors, such as an application specific integrated circuit (ASIC), field-programmable gate arrays (FPGAs), and the like), which processors may be co-located in a single housing or in a single data center or may be geographically distributed; one or more network interfaces **1948** (which may be co-located or geographically distributed) where each network interface includes a transmitter (Tx) **1945** and a receiver (Rx) **1947** for enabling apparatus **1901** to transmit data to and receive data from other nodes connected to network **110** (e.g., an Internet Protocol (IP) network) to which network interface **1948** is connected; and one or more storage units (a.k.a., “data storage systems”) **1908** which may be co-located or geographically distributed and which may include one or more non volatile storage devices and/or one or more volatile storage devices. In embodiments where PC **1902** includes a programmable processor, a computer program product (CPP) **1941** may be provided. CPP **1941** includes a computer readable medium (CRM) **1942** storing a computer program (CP) **1943** comprising computer readable instructions (CRI) **1944**. CRM **1942** may be a non-transitory computer readable medium, such as, magnetic media (e.g., a hard disk), optical media, memory devices (e.g., random access memory, flash memory), and the like. In some embodiments, the CRI **1944** of computer program **1943** is configured such that when executed by PC **1902**, the CRI causes apparatus **1901** to perform steps described herein (e.g., steps described herein with reference to the flow charts). In other embodiments, apparatus **1901** may be configured to perform steps described herein without the need for code. That is, for example, PC **1902** may consist merely of one or more ASICs. Hence, the features of the embodiments described herein may be implemented in hardware and/or software.

[0163] While various embodiments are described herein, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of this disclosure should not be limited by any of the above-described exemplary embodiments. Moreover, any combination of the above-described elements in all possible variations thereof is encompassed by the disclosure unless otherwise indicated herein or otherwise clearly contradicted by context.

[0164] Additionally, while the processes described above and illustrated in the drawings are shown as a sequence of steps, this was done solely for the sake of illustration. Accordingly, it is contemplated that some steps may be added, some steps may be omitted, the order of the steps may be re-arranged, and some steps may be performed in parallel.

Claims

1. A method for decoding an image, the method comprising: obtaining a signed 12-bit input value, wherein the signed 12-bit input value is a function of at least a first sample value associated with the image; producing a signed 19-bit output value that is equal to $i.sub.v \times 2.sup.n \times (a+1)$, where $i.sub.v$ is the signed 12-bit input value, n is an integer greater than or equal to 0, and a is either 2 or 4; and using the signed 19-bit output value to produce a filtered sample value.

2. The method of claim 1, wherein producing the signed 19-bit output value comprises: inputting a first signed 15-bit value into a conditional negate unit that outputs a second signed 15-bit value, wherein the second 15-bit value is either equal to the first signed 15-bit value of the negative of the first signed 15-bit value; and producing the signed 19-bit output value by shifting the second signed 15-bit value n times.

3. The method of claim 2, wherein producing the signed 19-bit output value further comprises: producing a first signed 14-bit value using the 12-bit input value ($i.sub.v$), wherein the first signed 14-bit value is equal to: $i.v \times a$; producing a second signed 14-bit value using the 12-bit input value ($i.sub.v$); and producing the first signed 15-bit value using an adder that adds the first and second signed 14-bit values.

4. An apparatus for decoding an image, the apparatus comprising: memory; and processing circuitry, wherein the apparatus is configured to preform a method comprising: obtaining a signed 12-bit input value, wherein the signed 12-bit input value is a function of at least a first sample value associated with the image; producing a signed 19-bit output value that is equal to $i.sub.v \times 2.sup.n \times (a+1)$, where $i.sub.v$ is the signed 12-bit input value, n is an integer greater than or equal to 0, and a is either 2 or 4; and using the signed 19-bit output value to produce a filtered sample value.

5. The apparatus of claim 4, wherein the processing circuitry comprises a conditional negate unit, and producing the signed 19-bit output value comprises: inputting a first signed 15-bit value into the conditional negate unit which then outputs a second signed 15-bit value, wherein the second 15-bit value is either equal to the first signed 15-bit value of the negative of the first signed 15-bit value; and producing the signed 19-bit output value by shifting the second signed 15-bit value n times.

6. The apparatus of claim 5, wherein producing the signed 19-bit output value further comprises: producing a first signed 14-bit value using the 12-bit input value ($i.sub.v$), wherein the first signed 14-bit value is equal to: $i.v \times a$; producing a second signed 14-bit value using the 12-bit input value ($i.sub.v$); and producing the first signed 15-bit value using an adder that adds the first and second signed 14-bit values.
