

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent Application Publication

20250267300

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

YANG; Shengqi

---

### **VIDEO DECODER WITH CONFIGURABLE REFERENCE CACHE AND ON-CHIP MEMORY BUFFER**

---

#### **Abstract**

A device includes an on-chip memory and one or more processors coupled to the on-chip memory. The one or more processors are configured to configure a reference cache of the on-chip memory to have a reference cache size based on video coding information. The on-chip memory includes an on-chip memory region that is external to the reference cache. The one or more processors are also configured to assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

---

**Inventors:** YANG; Shengqi (Carlsbad, CA)

**Applicant:** QUALCOMM Incorporated (San Diego, CA)

**Family ID:** 1000007686333

**Appl. No.:** 18/443946

**Filed:** February 16, 2024

---

#### **Publication Classification**

**Int. Cl.:** H04N19/423 (20140101); H04N19/167 (20140101); H04N19/80 (20140101)

**U.S. Cl.:**

**CPC** H04N19/423 (20141101); H04N19/167 (20141101); H04N19/80 (20141101);

---

#### **Background/Summary**

I. FIELD

[0001] The present disclosure is generally related to decoding video data.

## II. DESCRIPTION OF RELATED ART

[0002] Digital video capabilities can be incorporated into a wide range of devices, including digital televisions, digital direct broadcast systems, wireless broadcast systems, personal digital assistants (PDAs), laptop or desktop computers, tablet computers, e-book readers, digital cameras, digital recording devices, digital media players, video gaming devices, video game consoles, cellular or satellite radio telephones, so-called “smart phones,” video teleconferencing devices, video streaming devices, and the like. Digital video devices implement video coding techniques, such as those described in the standards defined by MPEG-2, MPEG-4, ITU-T H.263, ITU-T H.264/MPEG-4, Part 10, Advanced Video Coding (AVC), ITU-T H.265/High Efficiency Video Coding (HEVC), ITU-T H.266/Versatile Video Coding (VVC) and extensions of such standards. Such video devices may transmit, receive, encode, decode, and/or store digital video information more efficiently by implementing such video coding techniques.

[0003] Conventionally, a video device can implement a video playback dataflow that includes a video decoder receiving an input data frame via a bitstream and generating a reconstructed frame (a “decoded frame”) of the video. The video decoder reconstructs the frame portion-by-portion. For example, the video decoder generates reconstructed image portions from left-to-right and top-to-bottom of the reconstructed frame.

[0004] In some examples, an image portion of the reconstructed frame has dependency on surrounding image portions of the reconstructed frame. To illustrate, the video decoder generates a left image portion that has a dependency on a right image portion that is adjacent to and on the right of the left image portion. Right boundary pixels of the left image portion are to be adjusted based on left boundary pixels of the right image portion to preserve continuity between the image portions. The adjusted left image portion and the adjusted right image portion are provided to a display memory and provided to a display device via a display refresh. In some examples, the adjusted image portions are also stored in a reference cache to be used, e.g., during inter-frame prediction, to reconstruct additional image portions. Reconstructing surrounding image portions and performing the adjustment causes delay in displaying an image portion. On the other hand, using the reconstructed image portions without adjustment as reference image portions can reduce quality of subsequently reconstructed image portions.

## III. Summary

[0005] According to a particular implementation of the techniques disclosed herein, a device includes an on-chip memory and one or more processors coupled to the on-chip memory. The one or more processors are configured to configure a reference cache of the on-chip memory to have a reference cache size based on video coding information. The on-chip memory includes an on-chip memory region that is external to the reference cache. The one or more processors are also configured to assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0006] According to a particular implementation of the techniques disclosed herein, a method includes configuring, at a device, a reference cache of an on-chip memory to have a reference cache size based on video coding information. The on-chip memory includes an on-chip memory region that is external to the reference cache. The method also includes assigning, at the device, a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0007] According to a particular implementation of the techniques disclosed herein, a non-transitory computer-readable medium includes instructions that, when executed by one or more processors, cause the one or more processors to configure a reference cache of an on-chip memory to have a reference cache size based on video coding information. The on-chip memory includes an on-chip memory region that is external to the reference cache. The instructions, when executed by one or more processors, also cause the one or more processors to assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0008] Other implementations, advantages, and features of the present disclosure will become

## Description

### IV. BRIEF DESCRIPTION OF THE DRAWINGS

[0009] FIG. 1 is a block diagram illustrating an example of an implementation of a system operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0010] FIG. 2 is a block diagram illustrating an example of components of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0011] FIG. 3 is a block diagram illustrating an example of a configured reference cache that can be implemented in the system of FIG. 1, in accordance with some examples of the present disclosure.

[0012] FIG. 4 is a diagram illustrating an example of components of a video decoder that can be implemented in the system of FIG. 1, in accordance with some examples of the present disclosure.

[0013] FIG. 5 is a diagram illustrating an example of a video decoding operation that can be implemented in the system of FIG. 1, in accordance with some examples of the present disclosure.

[0014] FIG. 6 is a block diagram illustrating an example of a video decoding operation that can be implemented in the system of FIG. 1, in accordance with some examples of the present disclosure.

[0015] FIG. 7 is a block diagram illustrating an implementation of an integrated circuit operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0016] FIG. 8 is a diagram of an implementation of a portable electronic device operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0017] FIG. 9 is a diagram of a camera operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0018] FIG. 10 is a diagram of a wearable electronic device operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0019] FIG. 11 is a diagram of an extended reality device, such as augmented reality glasses, operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0020] FIG. 12 is a diagram of a headset, such as a virtual reality, mixed reality, or augmented reality headset, operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0021] FIG. 13 is a diagram of a voice-controlled speaker system operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0022] FIG. 14 is a diagram of a first example of a vehicle operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0023] FIG. 15 is a diagram of a second example of a vehicle operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0024] FIG. 16 is a diagram of a particular implementation of a method of processing video data

using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

[0025] FIG. 17 is a block diagram of a particular illustrative example of a device that is operable to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer, in accordance with some examples of the present disclosure.

## V. DETAILED DESCRIPTION

[0026] Systems and methods to perform video decoding using a video decoder with a configurable reference cache and on-chip memory buffer are disclosed. In conventional video decoding techniques, a video decoder reconstructs an image frame portion-by-portion. For example, the video decoder generates reconstructed image portions from left-to-right and top-to-bottom of the reconstructed frame. In some examples, when an image portion of the reconstructed frame has dependency on surrounding image portions of the reconstructed frame, the surrounding image portions are reconstructed prior to adjusting a particular image portion and the adjusted image portion is provided to the display memory, which causes delay that can adversely impact user experience. When the adjusted image portion corresponds to a reference image portion, the adjusted image portion is stored in a reference cache to be used to generate additional image portions.

[0027] The disclosed systems and methods include techniques to configure a reference cache and an on-chip memory buffer so that the on-chip memory buffer can be used to store a reconstructed image portion concurrently with providing the reconstructed image portion to a display memory. The reconstructed image portion is available in the on-chip memory buffer for any adjustment to be performed after reconstruction of surrounding image portions. The adjusted image portion can be stored in the reference cache to be used for reconstruction of additional image portions. In some examples, a size of the reference cache is configured based on video coding information (e.g., codec and coding unit information) of the video decoder. The on-chip memory includes the reference cache and an on-chip memory region that is external to the reference cache. One or more portions of the on-chip memory region are assigned as the on-chip memory buffer that can be used to store reconstructed image portions. Providing the reconstructed image portions (e.g., prior to adjustment) to the display memory reduces latency, while storing the adjusted image portions in the reference cache maintains quality of subsequent reconstruction of dependent image portions.

[0028] Particular aspects of the present disclosure are described below with reference to the drawings. In the description, common features are designated by common reference numbers. As used herein, various terminology is used for the purpose of describing particular implementations only and is not intended to be limiting of implementations. For example, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. Further, some features described herein are singular in some implementations and plural in other implementations. To illustrate, FIG. 1 depicts a device 102 including one or more processors (“processor(s)” 116 of FIG. 1), which indicates that in some implementations the device 102 includes a single processor 116 and in other implementations the device 102 includes multiple processors 116. For ease of reference herein, such features are generally introduced as “one or more” features and are subsequently referred to in the singular or optional plural (as indicated by “(s)” in the name of the feature) unless aspects related to multiple of the features are being described.

[0029] It may be further understood that the terms “comprise,” “comprises,” and “comprising” may be used interchangeably with “include,” “includes,” or “including.” Additionally, it will be understood that the term “wherein” may be used interchangeably with “where.” As used herein, “exemplary” may indicate an example, an implementation, and/or an aspect, and should not be construed as limiting or as indicating a preference or a preferred implementation. As used herein, an ordinal term (e.g., “first,” “second,” “third,” etc.) used to modify an element, such as a structure, a component, an operation, etc., does not by itself indicate any priority or order of the element with

respect to another element, but rather merely distinguishes the element from another element having a same name (but for use of the ordinal term). As used herein, the term “set” refers to one or more of a particular element, and the term “plurality” refers to multiple (e.g., two or more) of a particular element.

[0030] As used herein, “coupled” may include “communicatively coupled,” “electrically coupled,” or “physically coupled,” and may also (or alternatively) include any combinations thereof. Two devices (or components) may be coupled (e.g., communicatively coupled, electrically coupled, or physically coupled) directly or indirectly via one or more other devices, components, wires, buses, networks (e.g., a wired network, a wireless network, or a combination thereof), etc. Two devices (or components) that are electrically coupled may be included in the same device or in different devices and may be connected via electronics, one or more connectors, or inductive coupling, as illustrative, non-limiting examples. In some implementations, two devices (or components) that are communicatively coupled, such as in electrical communication, may send and receive signals (e.g., digital signals or analog signals) directly or indirectly, via one or more wires, buses, networks, etc. As used herein, “directly coupled” may include two devices that are coupled (e.g., communicatively coupled, electrically coupled, or physically coupled) without intervening components.

[0031] In the present disclosure, terms such as “obtaining,” “determining,” “calculating,” “estimating,” “shifting,” “adjusting,” etc. may be used to describe how one or more operations are performed. It should be noted that such terms are not to be construed as limiting and other techniques may be utilized to perform similar operations. Additionally, as referred to herein, “obtaining,” “generating,” “calculating,” “estimating,” “using,” “selecting,” “accessing,” and “determining” may be used interchangeably. For example, “obtaining,” “generating,” “calculating,” “estimating,” or “determining” a parameter (or a signal) may refer to actively generating, estimating, calculating, or determining the parameter (or the signal) or may refer to using, selecting, retrieving, receiving, or accessing the parameter (or signal) that is already generated, such as by another component or device.

[0032] Referring to FIG. 1, a particular illustrative aspect of a system **100** is depicted that includes a device **102** that is coupled to a display device **104** and that is configured to perform video decoding using a video decoder **124** with a configurable reference cache **152** and one or more memory buffers **156**. For example, the device **102** is configured to use a memory configurer **190** to configure an on-chip memory **150** to adjust a size of the reference cache **152**, to designate one or more portions of the on-chip memory **150** as the memory buffer(s) **156**, or a combination thereof.

[0033] The device **102** includes a memory **110** coupled to one or more processors **116** and configured to store instructions **112** and video data. For example, the memory **110** may include encoded video data **122**, adjusted pixels **130** corresponding to image portions of the encoded video data **122**, or any combination thereof, as described in further detail below. In a particular implementation, the memory **110** corresponds to a dynamic random access memory (DRAM) of a double data rate (DDR) memory subsystem. It should be understood that an “image portion” can refer to a portion of a larger image frame or to an entire image frame.

[0034] The on-chip memory **150** is integrated directly into a processor chip or within the same package as the one or more processors **116**. The on-chip memory **150** is generally characterized by its proximity to the one or more processors **116** (e.g., processor core(s)) enabling extremely fast access compared to off-chip memory (e.g., the memory **110**). The on-chip memory **150** can include various types of caches, such as the reference cache **152**. The reference cache **152** is typically used to store data that can be repeatedly accessed quickly during execution. For example, the reference cache **152** is configured to store data corresponding to decoded reference pixels that are to be used to reconstruct additional pixels.

[0035] The one or more processors **116** are configured to execute the instructions **112** to perform operations associated with configuring the on-chip memory **150** at the memory configurer **190**,

decoding encoded video data **122** at the video decoder **124**, or both. In various implementations, some or all of the functionality associated with the memory configurer **190**, the video decoder **124**, or both, is performed via execution of the instructions **112** by the one or more processors **116**, performed by processing circuitry of the one or more processors **116** in a hardware implementation, or a combination thereof.

[0036] The one or more processors **116** include the video decoder **124** coupled to an encoded data source **120**. The video decoder **124** is also coupled to the on-chip memory **150** (e.g., an on-chip cache). The video decoder **124** is configured to obtain the encoded video data **122** from the encoded data source **120**. For example, the encoded data source **120** may correspond to a portion of one or more of media files (e.g., a media file including the encoded video data **122** that is retrieved from the memory **110**), a game engine, one or more other sources of video information, such as a remote media server, or a combination thereof.

[0037] In a particular implementation, the on-chip memory **150** and the video decoder **124** are integrated into a single substrate **196** (e.g., a single chip). Although the on-chip memory **150** is illustrated as distinct from and coupled to the video decoder **124**, in other examples the on-chip memory **150** is integrated in the video decoder **124**. According to an aspect, the on-chip memory **150** includes a static random access memory (SRAM).

[0038] The on-chip memory **150** includes a reference cache **152** configured to store decoded pixels of reference image portions that can be used as reference pixels to reconstruct additional pixels. The one or more processors **116** include the memory configurer **190** that is configured to adjust a size of the reference cache **152**, and to add one or more memory buffer(s) **156** in the on-chip memory **150**, as further described with reference to FIG. 2. For example, the memory configurer **190** determines, based on video coding information **192** (e.g., codec information, coding unit information, or both) of the video decoder **124**, that the reference cache **152** is to have a reference cache size **172**. In some implementations, the memory configurer **190** updates the reference cache size **172** in a register **170** to adjust the size of the reference cache **152**, as further described with reference to FIG. 3. The memory configurer **190** also determines that the on-chip memory **150** includes a memory region **154** that is external to the reference cache **152**, and designates one or more portions of the memory region **154** as one or more memory buffers **156** having corresponding one or more buffer sizes **182**. The memory buffer(s) **156** are configured to buffer data (such as pixel data, line buffer data, streaming meta data, or any other type of data) associated with operation of the video decoder **124**. The memory-buffer(s) **156** can reduce traffic to the memory **110** (e.g., DDR) by buffering data associated with pixel adjustment, data associated with other operations of the video decoder **124**, or a combination thereof. In some implementations, the memory configurer **190** is configured to store, in a register **180**, one or more address ranges **184** of the one or more portions of the memory region **154** to assign the one or more portions to the memory buffer(s) **156**.

[0039] The video decoder **124** is configured to decode an image portion of the encoded video data **122** to generate first pixels, illustrated as reconstructed pixels **132** and to store the reconstructed pixels **132** in the memory buffer(s) **156** for display output. For example, the video decoder **124** may be configured to receive a first image portion **162** in the encoded video data **122** and to process the first image portion **162** to generate the reconstructed pixels **132**.

[0040] The reconstructed pixels **132** may be output by the video decoder **124** to the on-chip memory **150** for storage in the memory buffer(s) **156** and later retrieval by a display unit **140** (e.g., a data processing unit (DPU)) for storage in a display memory **144** prior to output to a display device **104** as video data output **142**. The display unit **140** is configured to receive the reconstructed pixels **132**, to generate the video data output **142** based on the reconstructed pixels **132**, and to provide the video data output **142**, such as a display refresh, to the display device **104**. To illustrate, the display unit **140** is configured to receive the reconstructed pixels **132** from the on-chip memory **150** and to store the reconstructed pixels **132** in the display memory **144**. In some implementations,

the display unit **140** may also retrieve additional data from the memory **110** for use in conjunction with processing the reconstructed pixels **132** (e.g., layer composition) to generate the video data output **142**.

[0041] The display device **104** is configured to display the video data output **142**, which is based on the reconstructed pixels **132**. For example, the video data output **142** can include at least a portion of a reconstructed frame that is based on the reconstructed pixels **132** for viewing by a user of the device **102**.

[0042] In some aspects, an image portion has dependence on one or more other image portions (e.g., surrounding image portions). For example, the video decoder **124** is configured to adjust the reconstructed pixels **132** of the first image portion **162** based on reconstructed pixels **134** of a second image portion **164** to generate adjusted pixels **130** of the first image portion **162**. In a particular aspect, the video decoder **124** is configured to generate the adjusted pixels **130** concurrently with or subsequent to storing the reconstructed pixels **134** in the memory buffer(s) **156**. The video decoder **124** is configured to store the adjusted pixels **130** in the reference cache **152**, the memory **110**, or both, so that the adjusted pixels **130** are available to be used as reference pixels to reconstruct additional pixels of other image portions.

[0043] Optionally, the video decoder **124** may be configured to selectively generate the adjusted pixels **130** and store the adjusted pixels **130** into the reference cache **152**, the memory **110**, or both, based on whether the adjusted pixels **130** correspond to reference pixels. To illustrate, the encoded video data **122** may include information that indicates which image portions are reference image portions, and the video decoder **124** may use the information to determine whether to generate and store the adjusted pixels **130** or refrain from generating the adjusted pixels **130**. For example, the adjusted pixels **130** may be generated based on a determination that the first image portion **162** is a reference image portion, and generation and storage of non-reference image portions to the reference cache **152** and the memory **110** may be skipped. Skipping generation and storage of non-reference image portions enables memory access bandwidth, power consumption, and storage capacity of the memory **110** and the reference cache **152** that is used during decoding, to be reduced.

[0044] The device **102** optionally includes a modem **118** that is coupled to the one or more processors **116** and configured to enable communication with one or more other devices, such as via one or more wireless networks. According to some aspects, the modem **118** is configured to receive the encoded video data **122** from a second device, such as video data that is streamed via a wireless transmission **194** from a remote device **198** (e.g., a remote server) for playback at the device **102**.

[0045] During operation, the memory configurer **190** configures the reference cache **152** to have a reference cache size **172** based on video coding information **192** (e.g., codec information, coding unit information, or both) of the video decoder **124**, as further described with reference to FIGS. 2-3. For example, the memory configurer **190** determines the reference cache size **172** based on a recommended reference cache size for the video coding information **192**. In some examples, the recommended reference cache size is associated with a target cache miss rate. In a particular aspect, the memory configurer **190** updates the register **170** to indicate the reference cache size **172** of the reference cache **152**.

[0046] The memory configurer **190** assigns one or more portions of the memory region **154** to be used as memory buffer(s) **156** to buffer data associated with operation of the video decoder **124**, as further described with reference to FIG. 2. For example, the memory configurer **190** identifies a memory region **154** of the on-chip memory **150** that is external to the reference cache **152**. To illustrate, the memory configurer **190** determines a size of the memory region **154** (e.g., 512 kilobytes (KB)-40 KB=472 KB) based on a difference between a first size (e.g., 512 KB) of the on-chip memory **150** and the reference cache size **172** (e.g., 40 KB) of the reference cache **152**. The memory configurer **190** assigns memory buffer(s) **156** from the memory region **154** of buffer

size(s) **182** that fit within the size of the memory region **154** (e.g., 472 KB). In a particular aspect, the memory configurer **190** updates the register **180** to indicate the address range(s) **184**, the buffer size(s) **182**, or a combination thereof, of the memory buffer(s) **156**. In some aspects, the memory configurer **190** configures the reference cache **152** and assigns the memory buffer(s) **156** during execution of the video decoder **124**.

[0047] In a particular aspect, a default size (e.g., 512 KB) of the reference cache **152** can include all of the space available in the on-chip memory **150** and the default location of the memory buffer(s) **156** can be off-chip memory (e.g., the memory **110**). For some codecs and coding unit sizes, the reference cache **152** is typically not fully utilized and the size of the reference cache **152** can be reduced with little or no adverse impact on performance. Configuring the reference cache **152** to have the reference cache size **172** (e.g., 40 KB) that is less than the default size (e.g., 512 KB) can enable space (e.g., the memory region **154**) to become available in the on-chip memory **150** for the memory buffer(s) **156**. A technical advantage of having the memory buffer(s) **156** on the on-chip memory **150** instead of on the off-chip memory includes reduction in traffic to the off-chip memory, thereby resulting in bandwidth savings and power savings.

[0048] The encoded video data **122** may be received at the video decoder **124** as a bitstream that includes a sequence of image portions including a first image portion **162**, a second image portion **164**, and one or more additional image portions including an Nth image portion **166** (N is a positive integer). The encoded video data **122** is processed by the video decoder **124** to generate reconstructed pixels for each of the image portions **162-166**. For example, the video decoder **124** processes the first image portion **162** (e.g., the encoded version of the first image portion **162**) to generate reconstructed pixels **132** of the first image portion **162**. The video decoder **124** sends the reconstructed pixels **132** to the display memory **144**. For example, the video decoder **124** stores the reconstructed pixels **132** into the memory buffer(s) **156**, and the reconstructed pixels **132** are retrieved from the on-chip memory **150** and provided to the display unit **140** for output at the display device **104**. In a particular aspect, the display unit **140** stores the reconstructed pixels **132** in the display memory **144**, uses the reconstructed pixels **132** to generate video data output **142**, and provides the video data output **142** to the display device **104**.

[0049] The reconstructed pixels **132** may be generated at the video decoder **124** and transferred to the display unit **140** via the on-chip memory **150** without being stored at the memory **110**. However, in some cases, such as based on the size of the memory buffer(s) **156** and a management policy used by the device **102**, the reconstructed pixels **132** may be evicted from the on-chip memory **150** and stored into the memory **110**. In such cases, in response to a request for the reconstructed pixels **132** resulting in a cache miss at the memory buffer(s) **156**, the reconstructed pixels **132** may be retrieved from the memory **110** and provided to the display memory **144** of the display unit **140**. The display unit **140** generates the video data output **142** based on the reconstructed pixels **132** and provides the video data output **142** to the display device **104** for playout, such as to a user of the device **102**. In a particular aspect, the video decoder **124** sends the reconstructed pixels **132** to the display memory **144** independently of dependency on later-reconstructed image portions. The reconstructed pixels **132** are thus sent to the display memory **144** prior to any adjustment of the reconstructed pixels **132** based on later-reconstructed image portions.

[0050] Storage of the reconstructed pixels **132** at the on-chip memory **150** enables the reconstructed pixels **132** to be conveyed from the video decoder **124** to the display unit **140** without being stored into, and later retrieved from, the memory **110**. As a result, usage of the storage capacity of the memory **110**, memory bandwidth associated with data transfer into and out of the memory **110**, and power consumption associated with the memory **110** are reduced as compared to storing the reconstructed pixels **132** into the memory **110** and later reading the reconstructed pixels **132** from the memory **110**.

[0051] In a particular implementation, the video decoder **124**, in response to determining that the reconstructed pixels **132** do not have a dependency on later-reconstructed pixels and determining



that the reconstructed pixels **132** correspond to reference pixels that are to be used to generate other reconstructed pixels, stores the reconstructed pixels **132** of the first image portion **162** in the reference cache **152**. Alternatively, the video decoder **124**, in response to determining that the reconstructed pixels **132** have a dependency on later-reconstructed pixels (e.g., of surrounding image portions), refrains from storing the reconstructed pixels **132** in the reference cache **152**. [0052] The video decoder **124** processes additional image portions of the encoded video data **122** to generate additional reconstructed pixels. For example, the video decoder **124** processes the second image portion **164** (e.g., the encoded version of the second image portion **164**) to generate reconstructed pixels **134** of the second image portion **164**.

[0053] The video decoder **124**, in response to determining that the reconstructed pixels **132** of the first image portion **162** have a dependency on later-reconstructed image portions and that the later-reconstructed image portions are available, adjusts the reconstructed pixels **132** based on the later-reconstructed image portions to generate adjusted pixels **130** of the first image portion **162** for use as reference pixels, as further described with reference to FIGS. **4** and **6**. For example, the video decoder **124**, in response to determining that the reconstructed pixels **132** of the first image portion **162** have a dependency on at least the reconstructed pixels **134** of the second image portion **164**, adjusts the reconstructed pixels **132** based at least in part on the reconstructed pixels **134** to generate the adjusted pixels **130** of the first image portion **162** for use as reference pixels. To illustrate, the video decoder **124**, in response to determining that first boundary pixels of the reconstructed pixels **132** are adjacent to second boundary pixels of the reconstructed pixels **134**, generates adjusted boundary pixels (e.g., the adjusted pixels **130**) that are based on the first boundary pixels and the second boundary pixels (e.g., using a deblocking filter operation). The reconstructed pixels **132** thus are adjusted subsequent to sending the reconstructed pixels **132** to the display memory **144** and subsequent to generating the reconstructed pixels **134**.

[0054] The video decoder **124**, in response to determining that the first image portion **162** corresponds to a reference image portion that is to be used to reconstruct additional image portions, stores the adjusted pixels **130** of the first image portion **162** in the reference cache **152**, the memory **110**, or both, as further described with reference to FIG. **4**. In a particular aspect, determining that the first image portion **162** corresponds to a reference image portion is equivalent to determining that the reconstructed pixels **132**, the adjusted pixels **130**, or both, correspond to reference pixels.

[0055] The video decoder **124** processes one or more additional image portions of the encoded video data **122** based on the adjusted pixels **130** to generate additional reconstructed pixels. For example, the video decoder **124** retrieves the adjusted pixels **130** from the reference cache **152** or the memory **110**, and uses intra-frame prediction or inter-frame prediction based on the adjusted pixels **130** of the first image portion **162** to generate reconstructed pixels **136** of the Nth image portion **166**. Generating the reconstructed pixels **136** based on the adjusted pixels **130**, instead of the reconstructed pixels **132**, may result in the reconstructed pixels **136** corresponding to a more accurate representation of the Nth image portion **166**.

[0056] The adjusted pixels **130** may be generated at the video decoder **124**, stored at the reference cache **152** of the on-chip memory **150**, and used to generate the reconstructed pixels **136** without being stored at the memory **110**. However, in some cases, such as based on the size of the on-chip memory **150** and a management policy used by the device **102**, the adjusted pixels **130** may be evicted from the on-chip memory **150** and stored into the memory **110**. In such cases, in response to a request for the adjusted pixels **130** resulting in a cache miss at the reference cache **152**, the adjusted pixels **130** may be retrieved from the memory **110** and provided to the video decoder **124**. The video decoder **124** generates the reconstructed pixels **136** based on the adjusted pixels **130**.

[0057] A technical advantage of sending the reconstructed pixels **132** from the memory buffer(s) **156** to the display memory **144** prior to adjustment of the reconstructed pixels **132** based on dependency on later-constructed image portions can include reduced latency associated with displaying a representation of the first image portion **162**. For example, the reconstructed pixels

**132** can be provided to the display unit **140** prior to or concurrently with generating the reconstructed pixels **134**. A technical advantage of using the adjusted pixels **130**, instead of the reconstructed pixels **132**, to generate the reconstructed pixels **136** can include improved reconstruction accuracy.

[0058] According to some aspects, the one or more processors **116** are integrated in at least one of a mobile phone or a tablet computer device, such as illustrated in FIG. **8**, or a wearable electronic device, such as illustrated in FIG. **10**. According to some aspects, the one or more processors **116** are integrated in a camera device, such as illustrated in FIG. **9**, or a voice-controlled speaker system, such as illustrated in FIG. **13**. According to some aspects, the one or more processors **116** are integrated in an extended reality headset device that is configured to display an output based on the reconstructed pixels **132**, such as illustrated in FIG. **11** and FIG. **12**. According to some aspects, the one or more processors **116** are integrated in a vehicle that also includes a display device configured to display an output based on the reconstructed pixels **132**, such as illustrated in FIG. **14** and FIG. **15**.

[0059] Although the display device **104** is illustrated as included in (e.g., integrated with) the device **102**, in other implementations the display device **104** may be coupled to, but not included in, the device **102**. Although the modem **118** is illustrated as included in the device **102**, in other examples the modem **118** may be omitted.

[0060] FIG. **2** depicts an illustrative example **200** including components of the system **100** of FIG. **1**, in accordance with some examples of the present disclosure. To illustrate, the example **200** illustrates the memory configurer **190** coupled to the register **170** and the register **180**.

[0061] The memory configurer **190** is configured to adjust a size of the reference cache **152** and to designate the memory buffer(s) **156** in the on-chip memory **150**. For example, the memory configurer **190** determines the reference cache size **172** based on the video coding information **192** (e.g., codec information, coding unit information, or both). The video decoder **124** is configured to support different reference cache sizes for different combinations of codecs and coding unit sizes.

[0062] In a particular implementation, the memory configurer **190** has access to cache size mapping data that maps codecs (e.g., H264, H265, AV1, VVC, etc.), coding unit sizes (e.g., 16×16, 32×32, 64×64, 128×128 pixels, etc.), or a combination thereof, to recommended reference cache sizes (e.g., 40 KB, 60 KB, 80 KB, 120 KB, 140 KB, and so on) that are supported by the video decoder **124**.

[0063] Table 1 provides illustrative, non-limiting examples of recommended reference cache sizes corresponding to codecs and coding unit sizes.

TABLE-US-00001  
TABLE 1 Recommended Coding Reference Cache Memory Region  
Codec Unit Size Size (KB) Size (KB) H264 16 × 16 40 472 H265 32 × 32 60 452 H265/VP9/ 64 × 64 80 432 AV1/VVC AV1 128 × 128 120 392 VVC 128 × 128 140 372

[0064] As shown in Table 1, for a H264 codec with a 16×16 coding unit size, the recommended reference cache size is 40 KB with 472 KB remaining in a 512 KB on-chip memory for allocation as the memory buffer(s) **156**. Table 1 provides illustrative, non-limiting examples of recommended reference cache sizes for some codecs and coding unit sizes. In other examples, the memory configurer **190** can support different reference cache sizes for different combinations of codecs and coding unit sizes. For example, the cache size mapping data can include mappings for one or more different codecs in addition to, or as an alternative to, the codecs listed in Table 1. The total size of “recommended reference cache size” plus “memory region size” of 512 KB is also provided as an illustrative non-limiting example. In other examples, the total size can be more than 512 KB or less than 512 KB.

[0065] In a particular aspect, the recommended reference cache sizes are associated with a target cache miss rate. For example, the memory configurer **190** has access to first cache size mapping data that maps video coding information (e.g., codecs, coding unit sizes, or a combination thereof) to first recommended reference cache sizes that are expected to satisfy a first target cache miss rate.

As another example, the memory configurer **190** has access to second cache size mapping data that maps video coding information (e.g., codecs, coding unit sizes, or a combination thereof) to second recommended reference cache sizes that are expected to satisfy a second target cache miss rate. In this aspect, the memory configurer **190** selects cache size mapping data corresponding to a target cache miss rate **294** and uses the selected cache size mapping data to determine the reference cache size **172** that is expected to satisfy the target cache miss rate **294**. In a particular aspect, the target cache miss rate **294** is based on a configuration setting, default data, user input, or a combination thereof.

[0066] The memory configurer **190**, in response to determining that cache size mapping data indicates that the video coding information **192** of the video decoder **124** maps to a recommended reference cache size (e.g., 40 KB), determines the reference cache size **172** based on the recommended reference cache size.

[0067] The memory configurer **190** determines, based on a difference between a size (e.g., 512 KB) of the on-chip memory **150** and the reference cache size **172** (e.g., 40 KB) that the on-chip memory **150** includes a memory region **154** (e.g., 512 KB–40 KB=472 KB) that is external to the reference cache **152**. In a particular aspect, the memory configurer **190** assigns one or more portions of the memory region **154** as the memory buffer(s) **156** to buffer data associated with operation of the video decoder **124**. In some aspects, the memory configurer **190** has access to traffic mapping data that indicates bandwidth savings, power savings, or a combination thereof, associated with sizes of memory buffers in the on-chip memory **150** for different types of traffic to the memory **110** associated with operation of the video decoder **124**.

[0068] Table 2 provides illustrative, non-limiting examples of recommended on-chip memory buffer sizes corresponding to traffic types, and associated bandwidth and power savings.

TABLE-US-00002

TABLE 2	On-Chip	Memory	Bandwidth	Power	Buffer	Size	Savings	Savings
Traffic Type	(Kilobytes)	(Megabytes/second)	(milliWatts)	8K30	AV1D	TLB	256	256
H265d	256	200	20	8K30	VVCd	TLB	256	156
ROI	512	265	66	Reference				

[0069] As shown in Table 2, a top line buffer (TLB) in the on-chip memory **150** to store data for 8K30 AV1D with a size of 256 KB is expected to result in 256 megabytes/second (MB/s) of bandwidth savings and 64 milliWatts (mW) of power savings. Table 2 provides illustrative, non-limiting examples of bandwidth savings and power savings for some traffic types. In other examples, buffers in the on-chip memory **150** for various types of traffic can result in different combinations of bandwidth savings and power savings. To illustrate, non-limiting illustrative examples of buffers for various types of traffic can include a line buffer (e.g., top line buffer, left line buffer, or prediction unit buffer), a bitstream buffer, a video slice and tile buffer, a bin header and coefficient stream buffer, a collocated motion vector buffer, a coder control buffer, a film grain data buffer, or a combination thereof.

[0070] The memory configurer **190**, based on a size (e.g., 472 KB) of the memory region **154** and traffic type at the device **102**, allocates buffers for traffic types that would satisfy a first buffer allocation criterion. The first buffer allocation criterion can include allocating a memory buffer that has a recommended buffer size that corresponds to highest expected bandwidth savings, highest expected power savings, highest read and write latency reduction, or a highest combination of expected savings among recommended buffer sizes that fit within an available portion of the memory region **154**. In an example, the memory configurer **190**, in response to determining that a TLB to store data for 8K30 AV1D with a size of 256 KB is expected to result in bandwidth savings (e.g., 256 MB/s) and power savings (e.g., 64 mW) that are highest for a recommended buffer size that would fit in the memory region **154**, allocates a first memory buffer **156** to store data for 8K30 AV1D having a first buffer size **182** that is based on a recommended buffer size (e.g., 256 KB).

[0071] The memory configurer **190** determines that after allocating the first buffer size **182** to the first memory buffer **156**, there is a remaining portion of the memory region **154** based on a

difference between the size of the memory region **154** and the allocated buffer size(s) **182** (e.g., 472 KB-256 KB=216 KB). In some implementations, the remaining portion of the memory region **154** remains unallocated. In some implementations, at least some of the remaining portion of the memory region **154** is added to the reference cache **152** so the reference cache size **172** is updated (e.g., to 40 KB+216 KB=256 KB).

[0072] In some implementations, at least some of the remaining portion of the memory region **154** is assigned to a second memory buffer **156** to store data of a second traffic type (e.g., 4K60 H265d region-of-interest (ROI) of Reference) although the remaining portion (e.g., 216 KB) of the memory region **154** is less than the recommended on-chip memory buffer size (e.g., 512 KB) for the second traffic type. To illustrate, a second buffer size **182** (e.g., less than or equal to 216 KB) of the second memory buffer **156** is less than the recommended on-chip memory buffer size (e.g., 512 KB) for the second traffic type. In some aspects, the memory configurer **190** allocates the remaining portion of the memory region **154** as a second memory buffer **156** to store data of a second traffic type that satisfies a second buffer allocation criterion. In a first example, the second buffer allocation criterion can include allocating a memory buffer that has a recommended buffer size that corresponds to highest expected bandwidth savings, highest expected power savings, highest read and write latency saving, or a highest combination of expected savings among unallocated buffers for traffic types at the device **102**. In the first example, the second buffer allocation criterion is independent of the recommended buffer size. In a second example, the second buffer allocation criterion can include allocating a memory buffer that has a recommended buffer size that corresponds to highest expected bandwidth savings, highest expected power savings, highest read and write latency saving, or a highest combination of expected savings among unallocated buffers for traffic types at the device **102** that have a recommended buffer size (e.g., 256 KB) that is closest to the size (e.g., 216 KB) of the remaining portion of the memory region **154**.

[0073] In a particular embodiment, the memory configurer **190** stores an indication of the reference cache size **172** in the register **170**. In some implementations, the memory configurer **190** stores data indicating one or more address ranges of the on-chip memory **150** in the register **170**. One or more portions of the on-chip memory corresponding to the address range(s) are allocated to the reference cache **152**.

[0074] In a particular embodiment, the memory configurer **190** stores data indicating buffer size(s) **182** and address range(s) **184** of the memory buffer(s) **156** in the register **180**. For example, the memory configurer **190** stores, in the register **180**, data indicating the first buffer size (e.g., 256 KB) of the first memory buffer **156** and one or more first address range(s) of the memory region **154** allocated to the first memory buffer **156**. In a particular aspect, the memory configurer **190** also stores, in the register **180**, data indicating that the first memory buffer **156** is allocated to store data of the first traffic type. As another example, the memory configurer **190** stores, in the register **180**, data indicating the second buffer size (e.g., 216 KB) of the second memory buffer **156** and one or more second address range(s) of the memory region **154** allocated to the second memory buffer **156**. In a particular aspect, the memory configurer **190** also stores, in the register **180**, data indicating that the second memory buffer **156** is allocated to store data of the second traffic type.

[0075] It should be understood that the memory configurer **190** updating two registers (e.g., the register **170** and the register **180**) to update the reference cache size **172** and to allocate the memory buffer(s) **156** is provided as an illustrative example. In other implementations, the memory configurer **190** may use other techniques to update the reference cache size **172**, to allocate the memory buffer(s) **156**, or both. For example, the memory configurer **190** can update a single combined register to update the reference cache size **172**, to allocate the memory buffer(s) **156**, both. In another example, the memory configurer **190** can update three or more registers to update the reference cache size **172**, to allocate the memory buffer(s) **156**, both.

[0076] FIG. 3 depicts an illustrative example **300** of a configured reference cache **152** that can be

implemented in the system **100** of FIG. **1**, in accordance with some examples of the present disclosure.

[0077] The memory configurer **190** of FIG. **1** configures, during execution of the video decoder **124**, the reference cache **152** to have the reference cache size **172**. In the example **300**, the memory configurer **190** updates the register **170** to indicate a count of cache ways **350**, a cache height **352**, a cache width **354**, or a combination thereof, of the reference cache **152** for the reference cache **152** to have the reference cache size **172**.

[0078] In an example, the on-chip memory **150** includes a plurality of memory units (e.g., bytes). A count of memory units allocated to the reference cache **152** corresponds to the count of cache ways **350**, the cache height **352**, the cache width **354**, or a combination thereof (e.g., count of memory units=count of cache ways **350**×the cache height **352**×the cache width **354**).

[0079] The memory configurer **190** designates sets of memory units (e.g., bytes) as a plurality of cache lines **304**. For example, each cache line **304** corresponds to a particular count of memory units corresponding to the cache width **354**.

[0080] The memory configurer **190** designates sets of cache lines **304** as cache ways **302** corresponding to the count of cache ways **350**. In an illustrative example, the count of cache ways **350** indicates **2** cache ways and the reference cache **152** includes a cache way **302A** and a cache way **302B**. In other examples, the reference cache **152** can include fewer than **2** or more than **2** cache ways corresponding to the count of cache ways **350**. Each cache way **302** corresponds to a particular count of cache lines **304** corresponding to the cache height **352**. For example, the cache way **302A** includes cache lines **304A**, such as a cache line **304AA**, a cache line **304AB**, and so on. A count of the cache lines **304A** is based on the cache height **352**. As another example, the cache way **302B** includes cache lines **304B**, such as a cache line **304BA**, a cache line **304BB**, and so on. A count of the cache lines **304B** is based on the cache height **352**.

[0081] In a particular aspect, the memory configurer **190** assigns a plurality of tags **314** to the plurality of cache lines **304**. For example, each cache line **304** is associated with a respective tag **314**. To illustrate, the cache line **304AA**, the cache line **304AB**, the cache line **304BA**, and the cache line **304BB** are associated with a tag **314AA**, a tag **314AB**, a tag **314BA**, and a tag **314BB**, respectively.

[0082] Although illustrated as a 2-way cache, in other implementations the reference cache **152** may have one or more other configurations. To illustrate, the reference cache **152** may be configured as a single-way cache having a cache width **354** that is based on (e.g., matches) a coding unit size, as an illustrative, non-limiting example. In some optional examples, the tags **314** may be omitted.

[0083] FIG. **4** depicts an illustrative example **400** including components that may be implemented in the video decoder **124**, in accordance with some examples of the present disclosure. In the example **400**, the video decoder **124** includes a bitstream parsing unit **410**, a pixel prediction processing unit **412**, an inverse transform processing unit **414**, a pixel reconstruction and inloop filtering unit **416**, and a pixel adjustment unit **418**. In a particular implementation, the bitstream parsing unit **410**, the pixel prediction processing unit **412**, the inverse transform processing unit **414**, the pixel reconstruction and inloop filtering unit **416**, the pixel adjustment unit **418**, or any combination thereof, may be implemented in one or more processors or in processing circuitry.

[0084] The various units shown in FIG. **4** are illustrated to assist with understanding the operations performed by the video decoder **124** in accordance with some implementations. The units may be implemented as fixed-function circuits, programmable circuits, or a combination thereof. Fixed-function circuits refer to circuits that provide particular functionality, and are preset on the operations that can be performed. Programmable circuits refer to circuits that can be programmed to perform various tasks, and provide flexible functionality in the operations that can be performed. For instance, programmable circuits may execute software or firmware that cause the programmable circuits to operate in the manner defined by instructions of the software or firmware.

Fixed-function circuits may execute software instructions (e.g., to receive parameters or output parameters), but the types of operations that the fixed-function circuits perform are generally immutable. In some examples, one or more of the units may be distinct circuit blocks (fixed-function or programmable), and in some examples, the one or more units may be integrated circuits.

[0085] In general, the video decoder **124** reconstructs a picture on a portion-by-portion basis. The video decoder **124** may perform a reconstruction operation on each image portion individually (where the image portion currently being reconstructed, i.e., decoded, may be referred to as a “current image portion”).

[0086] The bitstream parsing unit **410** receives encoded video data **122** and may entropy decode the encoded video data **122** to reproduce syntax elements. The pixel prediction processing unit **412**, inverse transform processing unit **414**, and the pixel reconstruction and inloop filtering unit **416** may generate decoded video data based on the syntax elements extracted from the bitstream **422**. In some implementations, the bitstream parsing unit **410** may decode information indicating which image portions in the bitstream **422** are reference image portions, which the video decoder **124** may use to determine which decoded image portions to store into the reference cache **152**.

[0087] The bitstream parsing unit **410** may entropy decode syntax elements defining quantized transform coefficients of a quantized transform coefficient image portion, as well as transform information, such as a quantization parameter (QP) and/or transform mode indication(s). The QP associated with the quantized transform coefficient image portion may be used to determine a degree of quantization and a degree of inverse quantization to apply. In an example, a bitwise left-shift operation may be performed to inverse quantize the quantized transform coefficients, and a transform coefficient image portion including transform coefficients may be formed.

[0088] The pixel prediction processing unit **412** may include one or more units to perform prediction in accordance with one or more prediction modes. As examples, the pixel prediction processing unit **412** may include a motion compensation unit, an inter-prediction unit, an intra-prediction unit, a palette unit, an affine unit, a linear model (LM) unit, one or more other units configured to prediction, or a combination thereof.

[0089] In addition, the pixel prediction processing unit **412** generates a prediction image portion according to prediction information syntax elements that were entropy decoded by the bitstream parsing unit **410**. For example, if the prediction information syntax elements indicate that the current image portion is inter-predicted, a motion compensation unit (not shown) may generate the prediction image portion. In this case, the prediction information syntax elements may indicate a reference image portion (e.g., reconstructed pixels or adjusted pixels of the reference image portion) to be retrieved from the reference cache **152**, as well as a motion vector identifying a location of the reference image portion in a reference image relative to the location of the current image portion in a current image. In some aspects, the reference image portion (e.g., the reconstructed pixels or the adjusted pixels of the reference image portion) may have been evicted from the reference cache **152** and may be retrieved from the memory **110**.

[0090] As another example, if the prediction information syntax elements indicate that the current image portion is intra-predicted, an intra-prediction unit (not shown) may generate the prediction image portion according to an intra-prediction mode indicated by the prediction information syntax elements. The pixel prediction processing unit **412** may retrieve data of neighboring samples (e.g., reconstructed pixels or adjusted pixels of neighboring image portions) to the current image block from the reference cache **152**. In some aspects, the neighboring image portions (e.g., the reconstructed pixels or the adjusted pixels of the neighboring image portions) may have been evicted from the reference cache **152** and may be retrieved from the memory **110**.

[0091] The pixel prediction processing unit **412** may also determine to decode image portions of video data using an intra block copy (IBC) mode. In general, in IBC mode, the video decoder **124** may determine predictive image portions for a current image portion, where the predictive image

portions are in the same frame as the current image portion. The predictive image portions may be identified by an image portion vector (e.g., a motion vector) and limited to the locations of image portions that have already been decoded.

[0092] The inverse transform processing unit **414** may apply one or more inverse transforms to the transform coefficient image portion from the bitstream parsing unit **410** to generate a residual image portion associated with the current image portion. For example, the inverse transform processing unit **414** may apply an inverse discrete cosine transform (DCT), an inverse integer transform, an inverse Karhunen-Loeve transform (KLT), an inverse rotational transform, an inverse directional transform, or another inverse transform to the coefficient image portion.

[0093] The pixel reconstruction and inloop filtering unit **416** may reconstruct the current image portion using the prediction image portion and the residual image portion. For example, the pixel reconstruction and inloop filtering unit **416** may add samples of the residual image portion to corresponding samples of the prediction image portion to reconstruct the current image portion.

[0094] In some implementations, the pixel reconstruction and inloop filtering unit **416** may perform one or more filter operations on reconstructed image portions. For example, the pixel reconstruction and inloop filtering unit **416** may access reconstructed pixels of image portions and perform deblocking operations to generate adjusted pixels with reduced blockiness artifacts along edges. Operations of the pixel reconstruction and inloop filtering unit **416** are not necessarily performed in all examples.

[0095] The video decoder **124** may store the reconstructed pixels of the image portions in the memory buffer(s) **156** on the on-chip memory **150**, the off-chip memory **110**, or both, for larger storage capacity. The memory buffer(s) **156** generally store reconstructed pixels of image portions, illustrated as reconstructed pixels **132** and reconstructed pixels **134**, which the video decoder **124** may output to the display memory **144**. The video decoder **124** may store reconstructed pixels of reference image portions that do not have dependence on later-reconstructed image portions in the reference cache **152**, which the video decoder **124** may use as reference video data when decoding subsequent data or image portions of the encoded video bitstream.

[0096] The pixel adjustment unit **418** is configured to process reconstructed pixels of image portions that have dependence on later-reconstructed image portions to generate adjusted pixels. For example, the pixel adjustment unit **418**, based on determining that the first image portion **162** has a dependence on one or more image portions (e.g., the second image portion **164**), retrieves the reconstructed pixels **132** of the first image portion **162** and the reconstructed pixels (e.g., the reconstructed pixels **134**) of the one or more image portions, and generates adjusted pixels **130** based on the retrieved reconstructed pixels **132** and **134**. In a particular implementation, the pixel adjustment unit **418** applies an image filter to adjust the reconstructed pixels **132** based on the reconstructed pixels **134** to generate the adjusted pixels **130**. The image filter includes at least one of a deblocking filter, a sample adaptive offset filter, an adaptive loop filter, a cross-component adaptive loop filter, or a luma mapping filter. Sending the reconstructed pixels **132** to the display memory **144** prior to applying the image filter reduces a latency associated with displaying a representation of the first image portion **162**. In some implementations, the pixel adjustment unit **418** can be included in the pixel reconstruction and inloop filtering unit **416**. In some implementations, some of the functionality described with reference to the pixel reconstruction and inloop filtering unit **416** can be performed in the pixel adjustment unit **418**.

[0097] The video decoder **124** may store the adjusted pixels of a reference image portion, illustrated as the adjusted pixels **130**, in the reference cache **152**, which the video decoder **124** may use as reference video data when decoding subsequent data or image portions of the encoded video bitstream. In a particular implementation, the pixel adjustment unit **418** generates adjusted pixels of reference image portions that have dependence on later-reconstructed image portions, and refrains from generating adjusted pixels of image portions that are non-reference image portions independently of whether the image portions have dependence on later-reconstructed image

portions. A technical advantage of selectively adjusting the reconstructed pixels for reference image portions that have dependence includes reduced resource usage (e.g., memory, computation cycles, or both).

[0098] As discussed above, the reference cache **152** may provide reference information, such as samples of a current image portion for intra-prediction and previously decoded image portions for subsequent motion compensation, to the pixel prediction processing unit **412**. Such reference information may be provided from the reference cache **152** of the on-chip memory **150** to the pixel prediction processing unit **412**. In some aspects, when the reference information is stored at the memory **110** (e.g., after eviction from the reference cache **152**), the reference information may be provided from the memory **110** to the pixel prediction processing unit **412** via a path indicated by the dotted line.

[0099] The video decoder **124** may provide reconstructed pixels from the memory buffer(s) **156** to the display memory **144** for subsequent presentation on a display device. For example, the reconstructed pixels **132** are output from the memory buffer(s) **156** to the display memory **144**, and are retrieved from or provided to the display memory **144** by the display unit **140** to generate the video data output **142** provided to the display device **104** of FIG. 1. In some implementations, the reconstructed pixels **132** may be saved to the memory **110** through the on-chip memory **150**, depending on a particular usage and/or configuration of the on-chip memory **150**, and may be retrieved from the memory **110** in response to a cache miss in the memory buffer(s) **156** by the display unit **140**.

[0100] In some implementations, the video decoder **124** stores a first subset of the reconstructed pixels **132** in the memory buffer(s) **156**, stores a second subset of the reconstructed pixels **134** in the memory buffer(s) **156**, and generates a third subset of the adjusted pixels **130** based on the first subset and the second subset. For example, the first subset corresponds to a first boundary (e.g., a right boundary) of the first image portion **162**, the second subset corresponds to a second boundary (e.g., a left boundary) of the second image portion **164**, and the third subset corresponds to the first boundary (e.g., a right boundary) of an adjusted and reconstructed version of the first image portion **162**. To illustrate, the video decoder **124** updates the first subset of the reconstructed pixels **132** based on the third subset to generate the adjusted pixels **130**.

[0101] FIG. 5 is a diagram illustrating an example of a video decoding operation that can be implemented in the system **100** of FIG. 1, in accordance with some examples of the present disclosure.

[0102] An example **500** depicts that an image frame **502** can be logically divided into a plurality of image portions **504**. An image portion **504** includes at least one of an image tile, an image region or an image subpicture of the image frame **502**. The encoded video data **122** can include data representing the image portions **504** (e.g., high resolution versions), data representing image portions **514** that correspond to low resolution versions of the image portions **504**, or both. In an example, an image portion **514A** corresponds to a low resolution version of the image portion **504A**. As another example, an image portion **514B** corresponds to a low resolution version of the image portion **504B**. In some aspects, if decoded pixels of the image portion **504A** are unavailable in time, the display unit **140** uses decoded pixels of the image portion **514A** to generate the video data output **142** of FIG. 1.

[0103] In some examples, the image frame **502** includes a ROI **510** associated with a viewing direction of a user. The ROI **510** includes one or more image portions **504**, such as the image portion **504A**, one or more additional image portions, or a combination thereof. One or more image portions **514** are also considered included in the ROI **510** that correspond to the image portion(s) **504** that are included in the ROI **510**.

[0104] Delay associated with displaying a representation of image portions included in the ROI **510** has a greater adverse impact on user experience. In some implementations, the display unit **140** of FIG. 1 uses the decoded pixels of image portions **504** (e.g., higher resolution image portions)



that are included in the ROI **510**, and uses decoded pixels of image portions **514** (e.g., lower resolution image portions) that are not included in the ROI **510** to generate the video data output **142**.

[0105] An example **550** depicts an image frame **502** including a plurality of image portions **524**. An image portion **524** corresponds to an image portion **504** or an image portion **514** (e.g., a lower resolution version of the image portion **504**). The example **550** depicts a left-to-right and top-to-bottom reconstruction of image portions **524** of the image frame **502**. For example, an image portion **524A**, an image portion **524B**, an image portion **524C**, an image portion **524D**, and an image portion **524E** of the encoded video data **122** are processed in sequential order to generate first reconstructed pixels, second reconstructed pixels, third reconstructed pixels, fourth reconstructed pixels, and fifth reconstructed pixels, respectively.

[0106] If the image portion **524A** depends on the image portion **524E**, the video decoder **124** typically has to wait for reconstructed pixels of the image portion **524E** to adjust the image portion **524A**. For example, the video decoder **124** generates the first reconstructed pixels and then has to wait for the second reconstructed pixels, the third reconstructed pixels, the fourth reconstructed pixels, and the fifth reconstructed pixels to be generated and the first reconstructed pixels to be adjusted before the adjusted pixels are provided as a representation of the image portion **524A** to the display memory **144**. The delay between generating the first reconstructed pixels and generating the adjusted pixels corresponds to a latency associated with displaying a representation of the image portion **524A** that can adversely impact user experience.

[0107] An example **590** depicts the video decoder **124** providing the reconstructed pixels to the display memory **144** prior to any adjustment of the reconstructed pixels. For example, the video decoder **124** provides the first reconstructed pixels, illustrated as reconstructed pixels **132**, to the display memory **144** prior to generating the second reconstructed pixels. The display unit **140** can use the reconstructed pixels **132** from the display memory **144** to generate the video data output **142** of FIG. **1** for output to the display device **104**. Sending the reconstructed pixels **132** to the display memory **144** prior to adjustment of the reconstructed pixels **132** based on later-reconstructed image portions reduces a latency associated with displaying a representation of the image portion **524A**.

[0108] FIG. **6** is a block diagram illustrating an example **600** of a video decoding operation that can be implemented in the system of FIG. **1**, in accordance with some examples of the present disclosure.

[0109] The memory buffer(s) **156** include a top line buffer **620** to store top line buffer data, a left line buffer **622** to store left line buffer data, or both. The video decoder **124** processes (e.g., decodes) an image portion **524A** of the encoded video data **122** to generate reconstructed pixels **132**. The video decoder **124** sends the reconstructed pixels **132** to the display memory **144**, as described with reference to FIG. **1**.

[0110] In a particular aspect, the video decoder **124** stores at least a portion of the reconstructed pixels **132** corresponding to a bottom boundary (e.g., bottom boundary pixels) of the image portion **524A** in the top line buffer **620**. In a particular aspect, the video decoder **124** stores at least a portion (e.g., right line data **604**) of the reconstructed pixels **132** corresponding to right boundary pixels of a right boundary **614** of the image portion **524A** in the left line buffer **622**. In an example, the right line data **604** corresponds to uncompressed boundary pixel data of the right boundary **614**.

[0111] The video decoder **124** processes (e.g., decodes) an image portion **524B** of the encoded video data **122** to generate reconstructed pixels **134**. In the example **600**, the image portion **524B** is adjacent and to the right of the image portion **524A** in an image frame **502**. The video decoder **124**, in response to determining that the image portion **524A** has a dependence on the image portion **524B**, adjusts the reconstructed pixels **132** based on the reconstructed pixels **134** to generate adjusted pixels **130**. The adjusted pixels **130** are thus generated subsequent to sending the reconstructed pixels **132** to the display memory **144**.

[0112] In an example, the video decoder **124** adjusts the right line data **604** in the left line buffer **622** based on left line data **606** (e.g., uncompressed boundary pixel data) corresponding to left boundary pixels of a left boundary **616** of the image portion **524** to generate adjusted right line data. In a particular aspect, the video decoder **124** stores the right line data **604** (e.g., uncompressed boundary pixel data) corresponding to the right boundary pixels of the Image portion **524A** in a first portion of a line buffer (e.g., the left line buffer **622**). The video decoder **124**, after generating the left boundary pixels of the image portion **524B**, stores the left line data **606** (e.g., uncompressed boundary pixel data) corresponding to the left boundary pixels in a second portion of the same line buffer or another line buffer of the memory buffers **156** to enable adjustment (e.g., filtering) of the right boundary pixels and the left boundary pixels.

[0113] In a particular aspect, the video decoder **124** applies an image filter to the right line data **604** and the left line data **606** to generate adjusted right line data. In a particular aspect, the right line data **604** is adjusted in the line buffer (e.g., the left line buffer **622**) and converted to the adjusted right line data. Adjusting the right line data **604** in a line buffer (e.g., the left line buffer **622**) at the on-chip memory **150**, subsequent to sending the reconstructed pixels **132** to the display memory **144**, reduces a latency associated with availability of the adjusted pixels **130** as reference pixels. For example, a latency associated with retrieving the right line data **604** from the memory **110** to perform the adjustment is removed. In some implementations, reconstructed pixels are compressed for storage in the memory **110**, adjusting the right line data **604** that is stored as uncompressed data in the line buffer (e.g., the left line buffer **622**) at the on-chip memory **150** also removes a latency associated with uncompressing the right line data **604** retrieved from the memory **110** to perform the adjustment.

[0114] The video decoder **124** stores the adjusted pixels **130** in the on-chip memory **150**, the memory **110**, or both. In a particular aspect, the adjusted pixels **130** are stored into the reference cache **152** and are available for use by the pixel prediction processing unit **412** of FIG. 4 as reference data to decode additional image portions. However, in some cases, such as based on the size of the reference cache **152** and a management policy used by the device **102**, the adjusted pixels **130** may be evicted from the reference cache **152** and stored into the memory **110**, and the pixel prediction processing unit **412** may (e.g., after a cache miss in the reference cache **152**) retrieve the adjusted pixels **130** from the memory **110** to use as reference data.

[0115] The top line buffer **620** and the left line buffer **622** are provided as illustrative non-limiting examples of the memory buffer(s) **156**, in other examples the memory buffer(s) **156** can include various other types of buffers in addition to, or as an alternative to, the top line buffer **620**, the left line buffer **622**, or both, that fit in the memory region **154**. As an example, the memory buffer(s) **156** can include a ROI reference buffer to store ROI reference data, a prediction unit buffer to store prediction unit data, one or more buffers to store other types of data, or a combination thereof. Having the various types of buffers in the on-chip memory **150** can result in power savings, bandwidth saving, read and write latency reduction, or a combination thereof.

[0116] FIG. 7 is a block diagram illustrating an implementation **700** of the device **102** as an integrated circuit **702** for performing video decoding using a video decoder with a configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**). The integrated circuit **702** includes the one or more processors **116**, which include the video decoder **124**, the memory configurer **190**, and the on-chip memory **150** with the reference cache **152** and the memory buffer(s) **156**. Optionally, the integrated circuit **702** also includes the memory **110**, the encoded data source **120**, the display unit **140**, or any combination thereof. The integrated circuit **702** also includes a signal input **704**, such as a bus interface, to enable the encoded video data **122** to be received. The integrated circuit **702** includes a signal output **706**, such as a bus interface, to enable outputting a video data output **726**, such as the video data output **142** or a sequence of image portions including the reconstructed pixels **132**. The integrated circuit **702** enables implementation of video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the

memory buffer(s) **156**) as a component in a system that performs video decoding playback, such as depicted in FIG. **1**.

[0117] FIG. **8** depicts an implementation **800** in which the device **102** includes a mobile device **802**, such as a phone or tablet, as illustrative, non-limiting examples. The mobile device **802** includes a display screen **804**. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the mobile device **802**, such as in the integrated circuit **702** that is illustrated using dashed lines to indicate internal components that are not generally visible to a user of the mobile device **802**. In a particular example, the video decoder **124** operates to perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**). For example, the mobile device **802** may receive encoded video data from a remote device (e.g., a phone or computer device of another participant on a video conference), decode the encoded video data using the video decoder **124**, output the reconstructed pixels to a display memory, and display the resulting decoded video at the display screen **804**.

[0118] FIG. **9** depicts an implementation **900** in which the device **102** includes a portable electronic device that corresponds to a camera device **902**. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the camera device **902**, such as in the integrated circuit **702**. During operation, the video decoder **124** performs video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) during playback of video data via a display of the camera device **902**, such as video data that is captured by the camera device **902** and stored as encoded video data at a memory of the camera device **902**.

[0119] FIG. **10** depicts an implementation **1000** of a wearable electronic device **1002**, illustrated as a “smart watch.” In a particular aspect, the wearable electronic device **1002** includes the device **102**. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the wearable electronic device **1002**, such as in the integrated circuit **702**. In a particular aspect, the wearable electronic device **1002** is coupled to or includes a display screen **1004** to display video data decoded by the video decoder **124**, and the video decoder **124** operates to perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**). In a particular example, the wearable electronic device **1002** includes a haptic device that provides a haptic notification (e.g., vibrates) associated with playback of decoded video data via the display screen **1004**. For example, the haptic notification can cause a user to look at the wearable electronic device **1002** to watch video playback, such as a video announcement of an incoming video phone call or a video message received at the wearable electronic device **1002**.

[0120] FIG. **11** depicts an implementation **1100** in which the device **102** includes a portable electronic device that corresponds to an extended reality device, such as augmented reality or mixed reality glasses **1102**. The glasses **1102** include a holographic projection unit **1104** configured to project visual data onto a surface of a lens **1106** or to reflect the visual data off of a surface of the lens **1106** and onto the wearer's retina. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the glasses **1102**, such as in the integrated circuit **702**. In a particular aspect, the video decoder **124** operates to perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) during playback of video data via a projection onto the surface of the lens **1106** (e.g., the display device **104**) to enable display of video associated with augmented reality, mixed reality, or virtual reality scenes to the user while the glasses **1102** are worn.

[0121] FIG. **12** depicts an implementation **1200** of a portable electronic device that corresponds to a virtual reality, augmented reality, or mixed reality headset **1202**. In a particular aspect, the headset **1202** includes the device **102** of FIG. **1**. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the headset

**1202**, such as in the integrated circuit **702**. In a particular aspect, the video decoder **124** operates to perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) during playback of video data via a visual interface device **1204** (e.g., the display device **104**). The visual interface device **1204** is positioned in front of the user's eyes to enable display of video associated with augmented reality, mixed reality, or virtual reality scenes to the user while the headset **1202** is worn.

[0122] FIG. **13** is an implementation **1300** of a wireless speaker and voice activated device **1302**. In a particular aspect, the wireless speaker and voice activated device **1302** includes the device **102** of FIG. **1**. The wireless speaker and voice activated device **1302** can have wireless network connectivity and is configured to execute an assistant operation. The one or more processors **116** are included in the wireless speaker and voice activated device **1302** and include the video decoder **124** and the memory configurer **190**. In a particular aspect, the wireless speaker and voice activated device **1302** includes one or more microphones **1310** and one or more speakers **1304**, and also includes or is coupled to a display device **1320** for playback of video that is output by the video decoder **124**. During operation, the video decoder **124** performs video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) during playback of video data via the display device **1320**. In response to receiving a verbal command via one or more microphones **1310**, the wireless speaker and voice activated device **1302** can execute assistant operations, such as via execution of a voice activation system (e.g., an integrated assistant application). The assistant operations can include adjusting a temperature, playing media content such as stored or streaming audio and video content, turning on lights, etc. For example, the assistant operations are performed responsive to receiving a command after a keyword or key phrase (e.g., "hello assistant").

[0123] FIG. **14** depicts an implementation **1400** in which the device **102** corresponds to or is integrated within a vehicle **1402**, illustrated as a manned or unmanned aerial device (e.g., a package delivery drone). The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the vehicle **1402**, such as in the integrated circuit **702**. The vehicle **1402** also includes a display device **1404** configured to display an output based on reconstructed pixels of image portions generated by the video decoder **124**, such as the video data output **142**.

[0124] The video decoder **124** operates to perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) during playback of video data that is decoded by the video decoder **124** and played back via a display device **1404**. In some implementations, the vehicle **1402** is manned (e.g., carries a pilot, one or more passengers, or both), the display device **1404** is internal to a cabin of the vehicle **1402**, and the video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is performed during playback to a pilot or a passenger of the vehicle **1402**. In another implementation, the vehicle **1402** is unmanned, the display device **1404** is mounted to an external surface of the vehicle **1402**, and the video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is performed during video playback to one or more viewers external to the vehicle **1402**. For example, the vehicle **1402** may move (e.g., circle an outdoor audience during a concert) while playing out video such as advertisements or steaming video of the concert stage, and the one or more processors **116** (e.g., including the video decoder **124**) may perform video decoding using the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) to generate the video from an encoded video stream.

[0125] FIG. **15** depicts an implementation **1500** in which the device **102** corresponds to, or is integrated within, a vehicle **1502**, illustrated as a car. The video decoder **124** with the configurable reference cache **152** and on-chip memory buffer (e.g., the memory buffer(s) **156**) is integrated in the vehicle **1502**, such as in the integrated circuit **702**. The vehicle **1502** also includes a display device **1520** and one or more speakers **1510**. In some implementations, the display device **1520** is

configured to display video data output based on reconstructed pixels generated by the video decoder **124**, such as the video data output **142**. For example, the video data may correspond to streaming video data from a remote source (e.g., a remote media server), video stored at the vehicle **1502**, such as entertainment content or instructional videos regarding operation of the vehicle **1502**, or video captured via one or more camera sensors of the vehicle **1502**, such as a backup camera. [0126] FIG. **16** illustrates an example of a method **1600** of decoding video data. One or more operations of the method **1600** may be performed by the system **100** of FIG. **1** (e.g., the device **102**, the one or more processors **116**, the memory configurer **190**, or the video decoder **124**), as an illustrative, non-limiting example.

[0127] The method **1600** includes, at block **1602**, configuring, at a device, a reference cache of an on-chip memory to have a reference cache size based on video coding information. For example, the memory configurer **190** of FIG. **1** configures the reference cache **152** of the on-chip memory **150** to have a reference cache size **172** based on video coding information **192** of the video decoder **124**, as described with reference to FIGS. **1-3**. The reference cache **152** is configured to buffer pixel data (e.g., the reconstructed pixels **132**) of one or more reference image portions (e.g., the first image portion **162**). The on-chip memory **150** includes an on-chip memory region (e.g., the memory region **154**) that is external to the reference cache **152**.

[0128] The method **1600** includes, at block **1604**, assigning, at the device, a portion of the on-chip memory region to buffer data associated with operation of a video decoder. For example, the video decoder **124** assigns a portion of the memory region **154** to the memory buffer(s) **156** to buffer data associated with operation of the video decoder **124**, as described with reference to FIGS. **1-2**.

[0129] The method **1600** of FIG. **16** may be implemented by a field-programmable gate array (FPGA) device, an application-specific integrated circuit (ASIC), a processing unit such as a central processing unit (CPU), a digital signal processor (DSP), a controller, another hardware device, firmware device, or any combination thereof. As an example, the method **1600** of FIG. **16** may be performed by a processor that executes instructions, such as described with reference to FIG. **17**.

[0130] Referring to FIG. **17**, a block diagram of a particular illustrative implementation of a device is depicted and generally designated **1700**. In various implementations, the device **1700** may have more or fewer components than illustrated in FIG. **17**. In an illustrative implementation, the device **1700** may correspond to the device **102** of FIG. **1**. In an illustrative implementation, the device **1700** may perform one or more operations described with reference to FIGS. **1-16**.

[0131] In a particular implementation, the device **1700** includes a processor **1706** (e.g., a CPU). The device **1700** may include one or more additional processors **1710** (e.g., one or more DSPs). In a particular implementation, the one or more processors **116** of FIG. **1** correspond to the processor **1706**, the processors **1710**, or a combination thereof. For example, the processors **1710** may include the video decoder **124**, the memory configurer **190**, the on-chip memory **150**, the display unit **140**, and a speech and music coder-decoder (CODEC) **1708**. The speech and music CODEC **1708** may include a voice coder (“vocoder”) encoder **1736**, a vocoder decoder **1738**, or a combination thereof.

[0132] The device **1700** may include a memory **1786** and a CODEC **1734**. The memory **1786** may include instructions **1756** that are executable by the one or more additional processors **1710** (or the processor **1706**) to implement the functionality described with reference to the video decoder **124**, the memory configurer **190**, or both. In a particular example, the memory **1786** corresponds to the memory **110** and the instructions **1756** correspond to the instructions **112** of FIG. **1**. The device **1700** may include the modem **118** coupled, via a transceiver **1750**, to an antenna **1752**.

[0133] The device **1700** may include a display **1728**, such as the display device **104**, coupled to a display controller **1726**. One or more speakers **1792**, one or more microphones **1790**, or a combination thereof, may be coupled to the CODEC **1734**. The CODEC **1734** may include a digital-to-analog converter (DAC) **1702** and an analog-to-digital converter (ADC) **1704**. In a

particular implementation, the CODEC **1734** may receive analog signals from the microphones **1790**, convert the analog signals to digital signals using the analog-to-digital converter **1704**, and send the digital signals to the speech and music codec **1708**. In a particular implementation, the speech and music codec **1708** may provide digital signals to the CODEC **1734**. The CODEC **1734** may convert the digital signals to analog signals using the digital-to-analog converter **1702** and may provide the analog signals to the speakers **1792**.

[0134] In a particular implementation, the device **1700** may be included in a system-in-package or system-on-chip device **1722**. In a particular implementation, the memory **1786**, the processor **1706**, the processors **1710**, the display controller **1726**, the CODEC **1734**, and the modem **118** are included in a system-in-package or system-on-chip device **1722**. In a particular implementation, an input device **1730** (e.g., a keyboard, a touchscreen, or a pointing device) and a power supply **1744** are coupled to the system-in-package or system-on-chip device **1722**. Moreover, in a particular implementation, as illustrated in FIG. **17**, the display **1728**, the input device **1730**, the speakers **1792**, the microphones **1790**, the antenna **1752**, and the power supply **1744** are external to the system-in-package or system-on-chip device **1722**. In a particular implementation, each of the display **1728**, the input device **1730**, the speakers **1792**, the microphones **1790**, the antenna **1752**, and the power supply **1744** may be coupled to a component of the system-in-package or system-on-chip device **1722**, such as an interface or a controller.

[0135] The device **1700** may include a smart speaker, a speaker bar, a mobile communication device, a smart phone, a cellular phone, a laptop computer, a computer, a tablet, a personal digital assistant, a display device, a television, a gaming console, a music player, a radio, a digital video player, a digital video disc (DVD) player, a tuner, a camera, a navigation device, a vehicle, a headset, an augmented reality headset, a mixed reality headset, a virtual reality headset, an extended reality headset, an aerial vehicle, a home automation system, a voice-activated device, a wireless speaker and voice activated device, a portable electronic device, a car, a vehicle, a computing device, a communication device, an internet-of-things (IoT) device, a virtual reality (VR) device, a base station, a mobile device, or any combination thereof.

[0136] In conjunction with the described techniques, an apparatus includes means for configuring a reference cache of an on-chip memory to have a reference cache size based on video coding information, where the on-chip memory includes an on-chip memory region that is external to the reference cache. In an example, the means for configuring can include the memory configurer **190**, the one or more processors **116**, the device **102**, the system **100**, one or more other circuits or devices to configure the reference cache, or a combination thereof.

[0137] The apparatus also includes means for assigning a portion of the on-chip memory region to buffer data associated with operation of a video decoder. In an example, the means for assigning can include the memory configurer **190**, the one or more processors **116**, the device **102**, the system **100**, one or more other circuits or devices to assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder, or a combination thereof.

[0138] In some implementations, a non-transitory computer-readable medium (e.g., a computer-readable storage device, such as the memory **110**) includes instructions (e.g., the instructions **112**) that, when executed by one or more processors (e.g., the one or more processors **116**), cause the one or more processors to perform operations corresponding to at least a portion of any of the techniques described with reference to FIGS. **1-15**, the method of FIG. **16**, or any combination thereof. For example, the instructions, when executed by the one or more processors, cause the one or more processors to configure a reference cache (e.g., the reference cache **152**) of an on-chip memory (e.g., the on-chip memory **150**) to have a reference cache size (e.g., the reference cache size **172**) based on video coding information (e.g., the video coding information **192**). The on-chip memory includes an on-chip memory region (e.g., the memory region **154**) that is external to the reference cache. The instructions, when executed by the one or more processors, also cause the one or more processors to assign a portion of the on-chip memory region to buffer data associated with

operation of a video decoder (e.g., the video decoder **124**).

[0139] Particular aspects of the disclosure are described below in the following sets of interrelated Examples:

[0140] According to Example 1, a device includes an on-chip memory; and one or more processors coupled to the on-chip memory, the one or more processors configured to: configure a reference cache of the on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0141] Example 2 includes the device of Example 1, wherein the video coding information includes codec or coding unit information.

[0142] Example 3 includes the device of Example 1 or Example 2, wherein the one or more processors are configured to update a register to indicate a cache width, a cache height, or a count of cache ways of the reference cache to configure the reference cache to have the reference cache size.

[0143] Example 4 includes the device of any of Examples 1 to 3, wherein the one or more processors are configured to update a register to assign the portion of the on-chip memory region to buffer the data associated with the operation of the video decoder.

[0144] Example 5 includes the device of any of Examples 1 to 4, wherein the portion of the on-chip memory region is assigned based on a size of the on-chip memory region.

[0145] Example 6 includes the device of any of Examples 1 to 5, wherein the data associated with the operation of the video decoder comprises top line buffer data, left line buffer data, or prediction unit data.

[0146] Example 7 includes the device of any of Examples 1 to 6, wherein assignment of the portion of the on-chip memory region to buffer the data produced during the operation of the video decoder reduces traffic to off-chip memory.

[0147] Example 8 includes the device of any of Examples 1 to 7, wherein the reference cache size is selected to satisfy a target cache miss rate.

[0148] Example 9 includes the device of any of Examples 1 to 8, wherein the one or more processors are configured to, based on the video coding information, configure the reference cache of the on-chip memory to have the reference cache size of 40 kilobytes, 60 kilobytes, 80 kilobytes, 120 kilobytes, or 140 kilobytes.

[0149] Example 10 includes the device of any of Examples 1 to 9, wherein the one or more processors are configured to generate reconstructed pixels of an image; and send the reconstructed pixels to a display memory prior to adjustment of the reconstructed pixels based on dependency on a later-reconstructed image.

[0150] Example 11 includes the device of Example 10, wherein sending the reconstructed pixels to the display memory prior to the adjustment of the reconstructed pixels reduces a display latency.

[0151] Example 12 includes the device of Example 10 or Example 11, wherein the image includes at least one of an image tile, an image region, or an image subpicture.

[0152] Example 13 includes the device of any of Example 10 to 12, wherein the image is within a region of interest associated with a viewing direction of a user.

[0153] Example 14 includes the device of any of Examples 10 to 13, wherein the one or more processors are configured to adjust the reconstructed pixels subsequent to sending the reconstructed pixels to the display memory.

[0154] Example 15 includes the device of any of Examples 10 to 14, wherein the one or more processors are configured to store the reconstructed pixels in a line buffer included in the portion of the on-chip memory region; subsequent to sending the reconstructed pixels to the display memory, adjust the reconstructed pixels to generate adjusted pixels; and store the adjusted pixels in the reference cache or off-chip memory.

[0155] Example 16 includes the device of Example 15, wherein adjusting the reconstructed pixels in the line buffer, subsequent to sending the reconstructed pixels to the display memory, reduces a latency associated with availability of the adjusted pixels as reference pixels.

[0156] Example 17 includes the device of any of Examples 1 to 16, wherein the one or more processors are configured to send reconstructed pixels of an image to a display memory independently of dependency on a later-reconstructed image.

[0157] Example 18 includes the device of Example 17, wherein the one or more processors are configured to, subsequent to sending the reconstructed pixels to the display memory, adjust the reconstructed pixels based on the dependency on the later-reconstructed image to generate adjusted pixels for use as reference pixels.

[0158] Example 19 includes the device of Example 17 or Example 18, wherein the one or more processors are configured to apply an image filter to adjust the reconstructed pixels.

[0159] Example 20 includes the device of Example 19, wherein sending the reconstructed pixels to the display memory prior to applying the image filter reduces a latency associated with displaying the image.

[0160] Example 21 includes the device of Example 19 or Example 20, wherein the image filter includes at least one of a deblocking filter, a sample adaptive offset filter, an adaptive loop filter, a cross-component adaptive loop filter, or a luma mapping filter.

[0161] Example 22 includes the device of any of Examples 1 to 21, wherein the one or more processors are configured to generate first boundary pixels of an image; store first uncompressed pixel data corresponding to the first boundary pixels to a first portion of a line buffer, the line buffer included in the portion of the on-chip memory region; and after generating second boundary pixels of a later-reconstructed image, store second uncompressed boundary pixel data corresponding to the second boundary pixels to a second portion of the line buffer to enable filtering of the first boundary pixels and the second boundary pixels.

[0162] Example 23 includes the device of any of Examples 1 to 22, wherein the on-chip memory and the one or more processors are integrated into an extended reality (XR) device.

[0163] Example 24 includes the device of any of Examples 1 to 23, wherein the on-chip memory and the one or more processors are integrated into at least one of a mobile device, an internet-of-things (IoT) device, a computer, or a vehicle.

[0164] According to Example 25, a method includes configuring, at a device, a reference cache of an on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assigning, at the device, a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0165] Example 26 includes the method of Example 25, wherein the video coding information includes codec or coding unit information.

[0166] Example 27 includes the method of Example 25 or Example 26, further comprising updating, at the device, a register to indicate a cache width, a cache height, or a count of cache ways of the reference cache to configure the reference cache to have the reference cache size.

[0167] Example 28 includes the method of any of Examples 25 to 27, further comprising updating, at the device, a register to assign the portion of the on-chip memory region to buffer the data associated with the operation of the video decoder.

[0168] Example 29 includes the method of any of Examples 25 to 28, wherein the portion of the on-chip memory region is assigned based on a size of the on-chip memory region.

[0169] Example 30 includes the method of any of Examples 25 to 29, wherein data associated with the operation of the video decoder comprises top line buffer data, left line buffer data, or prediction unit data.

[0170] Example 31 includes the method of any of Examples 25 to 30, wherein assignment of the portion of the on-chip memory region to buffer the data produced during the operation of the video



decoder reduces traffic to off-chip memory.

[0171] Example 32 includes the method of any of Examples 25 to 31, wherein the reference cache size is selected to satisfy a target cache miss rate.

[0172] Example 33 includes the method of any of Examples 25 to 32, wherein the reference cache of the on-chip memory is configured, based on the video coding information, to have the reference cache size of 40 kilobytes, 60 kilobytes, 80 kilobytes, 120 kilobytes, or 140 kilobytes.

[0173] Example 34 includes the method of any of Examples 25 to 33, and further includes generating, at the device, reconstructed pixels of an image; and sending the reconstructed pixels to a display memory prior to adjustment of the reconstructed pixels based on dependency on a later-reconstructed image.

[0174] Example 35 includes the method of Example 34, wherein sending the reconstructed pixels to the display memory prior to the adjustment of the reconstructed pixels reduces a display latency.

[0175] Example 36 includes the method of Example 34 or Example 35, wherein the image includes at least one of an image tile, an image region, or an image subpicture.

[0176] Example 37 includes the method of any of Examples 34 to 36, wherein the image is within a region of interest associated with a viewing direction of a user.

[0177] Example 38 includes the method of any of Examples 34 to 37, and further includes adjusting, at the device, the reconstructed pixels subsequent to sending the reconstructed pixels to the display memory.

[0178] Example 39 includes the method of any of Examples 34 to 38, and further includes storing the reconstructed pixels in a line buffer included in the portion of the on-chip memory region; subsequent to sending the reconstructed pixels to the display memory, adjust the reconstructed pixels to generate adjusted pixels; and storing the adjusted pixels in the reference cache or off-chip memory.

[0179] Example 40 includes the method of Example 39, wherein adjusting the reconstructed pixels in the line buffer, subsequent to sending the reconstructed pixels to the display memory, reduces a latency associated with availability of the adjusted pixels as reference pixels.

[0180] Example 41 includes the method of any of Examples 25 to 40, and further includes sending reconstructed pixels of an image to a display memory independently of dependency on a later-reconstructed image.

[0181] Example 42 includes the method of Example 41, and further includes, subsequent to sending the reconstructed pixels to the display memory, adjusting the reconstructed pixels based on the dependency on the later-reconstructed image to generate adjusted pixels for use as reference pixels.

[0182] Example 43 includes the method of Example 41 or Example 42, and further includes applying an image filter to adjust the reconstructed pixels.

[0183] Example 44 includes the method of Example 43, wherein sending the reconstructed pixels to the display memory prior to applying the image filter reduces a latency associated with displaying the image.

[0184] Example 45 includes the method of Example 43 or Example 44, wherein the image filter includes at least one of a deblocking filter, a sample adaptive offset filter, an adaptive loop filter, a cross-component adaptive loop filter, or a luma mapping filter.

[0185] Example 46 includes the method of any of Examples 25 to 45, and further includes generating first boundary pixels of an image; storing first uncompressed pixel data corresponding to the first boundary pixels to a first portion of a line buffer, the line buffer included in the portion of the on-chip memory region; and after generating second boundary pixels of a later-reconstructed image, storing second uncompressed boundary pixel data corresponding to the second boundary pixels to a second portion of the line buffer to enable filtering of the first boundary pixels and the second boundary pixels.

[0186] Example 47 includes the method of any of Examples 25 to 46, wherein the on-chip memory

is integrated into an extended reality (XR) device.

[0187] Example 48 includes the method of any of Examples 25 to 47, wherein the on-chip memory is integrated into at least one of a mobile device, an internet-of-things (IoT) device, a computer, or a vehicle.

[0188] According to Example 49, a non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause the one or more processors to configure a reference cache of an on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0189] Example 50 includes the non-transitory computer-readable medium of Example 49, wherein the instructions, when executed by the one or more processors, cause the one or more processors to update a register to indicate a cache width, a cache height, or a count of cache ways of the reference cache to configure the reference cache to have the reference cache size.

[0190] According to Example 51, an apparatus includes: means for configuring a reference cache of an on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and means for assigning a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

[0191] Example 52 includes the apparatus of Example 51, wherein the means for configuring and the means for assigning are integrated into at least one of a mobile communication device, a computer, a display device, a gaming console, a camera, a navigation device, a vehicle, a headset, an extended reality headset, an aerial vehicle, a vehicle, or an internet-of-things (IoT) device.

[0192] Those of skill would further appreciate that the various illustrative logical blocks, configurations, circuits, and algorithm steps described in connection with the implementations disclosed herein may be implemented as electronic hardware, computer software executed by a processing device such as a hardware processor, or combinations of both. Various illustrative components, blocks, configurations, circuits, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware or executable software depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the present disclosure.

[0193] The steps of a method or algorithm described in connection with the implementations disclosed herein may be embodied directly in hardware, in a software module executed by a processor, or in a combination of the two. A software module may reside in a memory device, such as random access memory (RAM), magnetoresistive random access memory (MRAM), spin-torque transfer MRAM (STT-MRAM), flash memory, read-only memory (ROM), programmable read-only memory (PROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), registers, hard disk, a removable disk, a compact disc read-only memory (CD-ROM), or any other form of non-transient storage medium known in the art. An exemplary memory device is coupled to the processor such that the processor can read information from, and write information to, the memory device. In the alternative, the memory device may be integral to the processor. The processor and the storage medium may reside in an application-specific integrated circuit (ASIC). The ASIC may reside in a computing device or a user terminal. In the alternative, the processor and the storage medium may reside as discrete components in a computing device or a user terminal.

[0194] The previous description of the disclosed implementations is provided to enable a person skilled in the art to make or use the disclosed implementations. Various modifications to these implementations will be readily apparent to those skilled in the art, and the principles defined

herein may be applied to other implementations without departing from the scope of the disclosure. Thus, the present disclosure is not intended to be limited to the implementations shown herein but is to be accorded the widest scope possible consistent with the principles and novel features as defined by the following claims.

## Claims

1. A device comprising: an on-chip memory; and one or more processors coupled to the on-chip memory, the one or more processors configured to: configure a reference cache of the on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.
2. The device of claim 1, wherein the video coding information includes codec or coding unit information.
3. The device of claim 1, wherein the one or more processors are configured to update a register to indicate a cache width, a cache height, or a count of cache ways of the reference cache to configure the reference cache to have the reference cache size.
4. The device of claim 1, wherein the one or more processors are configured to update a register to assign the portion of the on-chip memory region to buffer the data associated with the operation of the video decoder.
5. The device of claim 1, wherein the portion of the on-chip memory region is assigned based on a size of the on-chip memory region.
6. The device of claim 1, wherein the data associated with the operation of the video decoder comprises top line buffer data, left line buffer data, or prediction unit data.
7. The device of claim 1, wherein the reference cache size is selected to satisfy a target cache miss rate.
8. The device of claim 1, wherein the one or more processors are configured to, based on the video coding information, configure the reference cache of the on-chip memory to have the reference cache size of 40 kilobytes, 60 kilobytes, 80 kilobytes, 120 kilobytes, or 140 kilobytes.
9. The device of claim 1, wherein the one or more processors are configured to: generate reconstructed pixels of an image; and send the reconstructed pixels to a display memory prior to adjustment of the reconstructed pixels based on dependency on a later-reconstructed image.
10. The device of claim 9, wherein the image is within a region of interest associated with a viewing direction of a user.
11. The device of claim 9, wherein the one or more processors are configured to: store the reconstructed pixels in a line buffer included in the portion of the on-chip memory region; subsequent to sending the reconstructed pixels to the display memory, adjust the reconstructed pixels to generate adjusted pixels; and store the adjusted pixels in the reference cache or off-chip memory.
12. The device of claim 1, wherein the one or more processors are configured to send reconstructed pixels of an image to a display memory independently of dependency on a later-reconstructed image.
13. The device of claim 12, wherein the one or more processors are configured to, subsequent to sending the reconstructed pixels to the display memory, adjust the reconstructed pixels based on the dependency on the later-reconstructed image to generate adjusted pixels for use as reference pixels.
14. The device of claim 12, wherein the one or more processors are configured to apply an image filter to adjust the reconstructed pixels.
15. The device of claim 1, wherein the one or more processors are configured to: generate first boundary pixels of an image; store first uncompressed pixel data corresponding to the first boundary pixels to a first portion of a line buffer, the line buffer included in the portion of the on-

chip memory region; and after generating second boundary pixels of a later-reconstructed image, storing second uncompressed boundary pixel data corresponding to the second boundary pixels to a second portion of the line buffer to enable filtering of the first boundary pixels and the second boundary pixels.

**16.** The device of claim 1, wherein the on-chip memory and the one or more processors are integrated into at least one of an extended reality (XR) device, a mobile device, an internet-of-things (IoT) device, a computer, or a vehicle.

**17.** A method comprising: configuring, at a device, a reference cache of an on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assigning, at the device, a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

**18.** The method of claim 17, wherein the video coding information includes codec or coding unit information.

**19.** A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause the one or more processors to: configure a reference cache of an on-chip memory to have a reference cache size based on video coding information, wherein the on-chip memory includes an on-chip memory region that is external to the reference cache; and assign a portion of the on-chip memory region to buffer data associated with operation of a video decoder.

**20.** The non-transitory computer-readable medium of claim 19, wherein the instructions, when executed by the one or more processors, cause the one or more processors to update a register to indicate a cache width, a cache height, or a count of cache ways of the reference cache to configure the reference cache to have the reference cache size.

---