



US 20250272848A1

(19) **United States**

(12) **Patent Application Publication**
Barnes et al.

(10) **Pub. No.: US 2025/0272848 A1**

(43) **Pub. Date: Aug. 28, 2025**

(54) **PIXEL GENERATION TECHNIQUE**

Publication Classification

(71) Applicant: **NVIDIA Corporation**, Santa Clara, CA (US)

(51) **Int. Cl.**
G06T 7/13 (2017.01)

(72) Inventors: **Levi Daniel Barnes**, Lake Oswego, OR (US); **Kumara Narasimha Sastry Kunigal**, Portland, OR (US); **Justin Paul Luitjens**, Sandy, UT (US); **John Arthur Swanson**, Banks, OR (US); **Benjamin Frank Pinzone**, Santa Clara, CA (US)


(52) **U.S. Cl.**
CPC **G06T 7/13** (2017.01)

(57) **ABSTRACT**

Apparatuses, systems, and techniques are to represent polygon data as pixels as part of a rasterization process. In at least one embodiment, a processor causes identification of pixels within a polygon based, at least in part, on one or more prefix sums of amounts of edges of pixels. covered by that polygon.

(21) Appl. No.: **18/589,214**

(22) Filed: **Feb. 27, 2024**






FIG. 1




FIG. 2




FIG. 3

400 →




FIG. 4




FIG. 5

600 →




FIG. 6

700 →




FIG. 7




FIG. 8




FIG. 9




FIG. 10




FIG. 11




FIG. 12




FIG. 13




FIG. 14

1500 →




FIG. 15

1600 →




FIG. 16




FIG. 17




FIG. 18

1900 →




FIG. 19




FIG. 20




FIG. 21




FIG. 22




FIG. 23




FIG. 24




FIG. 25




FIG. 26




FIG. 27




FIG. 28




FIG. 29




FIG. 30




FIG. 31




FIG. 32




FIG. 33A




FIG. 33B




FIG. 34A



FIG. 34B




FIG. 35A




FIG. 35B




FIG. 35C




FIG. 36




FIG. 37





FIG. 38






FIG. 40




FIG. 41




FIG. 42




FIG. 43




FIG. 44




FIG. 45




FIG. 46




FIG. 47




FIG. 48




FIG. 49




FIG. 50




FIG. 51A





FIG. 51B






FIG. 52




FIG. 53




FIG. 54

Programmable Processing
Unit 5320(1)




FIG. 55




FIG. 56

PIXEL GENERATION TECHNIQUE

FIELD

[0001] At least one embodiment pertains to processing resources used during rasterization to convert data into pixels. For example, one or more processors comprising one or more circuits are to convert polygon data into pixels based, at least in part, on performing a prefix sum on an amount of pixel edges covered by a polygon edge.

BACKGROUND

[0002] Rasterization uses significant computing resources. For example, a processor may use thousands of processing cores to convert polygon data into pixels. Therefore, improvements can be made to rasterization performance.

BRIEF DESCRIPTION OF THE DRAWINGS

[0003] FIG. 1 illustrates a system to identify one or more pixels within one or more polygons, in accordance with at least one embodiment;

[0004] FIG. 2 illustrates a system to identify one or more pixels within one or more polygons by using a prefix sum partial coverage module, in accordance with at least one embodiment;

[0005] FIG. 3 illustrates a system to identify one or more pixels within one or more polygons by using a prefix sum partial coverage module, in accordance with at least one embodiment;

[0006] FIG. 4 illustrates a system to identify one or more pixels within one or more polygons by using amounts of pixels covered by a polygon, in accordance with at least one embodiment;

[0007] FIG. 5 illustrates a system to identify one or more pixels within one or more polygons by using amounts of pixels covered by a polygon, in accordance with at least one embodiment;

[0008] FIG. 6 illustrates a system to identify one or more pixels within one or more polygons by using amounts of pixels covered by a polygon, in accordance with at least one embodiment;

[0009] FIG. 7 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment, in accordance with at least one embodiment;

[0010] FIG. 8 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0011] FIG. 9 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0012] FIG. 10 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0013] FIG. 11 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0014] FIG. 12 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0015] FIG. 13 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0016] FIG. 14 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages, in accordance with at least one embodiment;

[0017] FIG. 15 illustrates a system to identify one or more pixels within one or more polygons by using bottom pixel edge coverages, in accordance with at least one embodiment;

[0018] FIG. 16 illustrates a system to identify one or more pixels within one or more polygons by using bottom pixel edge coverages, in accordance with at least one embodiment;

[0019] FIG. 17 illustrates a system to identify one or more pixels within one or more polygons by using one or more prefix sums of bottom pixel edge coverages, in accordance with at least one embodiment;

[0020] FIG. 18 illustrates a system to identify one or more pixels within one or more polygons by using partial coverages and prefix summed bottom pixel edge coverages, in accordance with at least one embodiment;

[0021] FIG. 19 illustrates a system to identify one or more pixels by using sums of partial coverages and prefix summed bottom pixel edge coverages, in accordance with at least one embodiment;

[0022] FIG. 20 illustrates a system to identify one or more pixels by using sums of partial coverages and prefix summed bottom pixel edge coverages, in accordance with at least one embodiment;

[0023] FIG. 21 illustrates a block diagram of a process to output coverage fractions of pixels in accordance with at least one embodiment;

[0024] FIG. 22 illustrates a block diagram of a process to cause an application programming interface (API) to identify coverage fractions of pixels, in accordance with at least one embodiment;

[0025] FIG. 23 illustrates a block diagram of a driver and/or runtime used to identify pixels within a polygon, according to at least one embodiment;

[0026] FIG. 24 illustrates an exemplary data center, in accordance with at least one embodiment;

[0027] FIG. 25 illustrates a processing system, in accordance with at least one embodiment;

[0028] FIG. 26 illustrates a computer system, in accordance with at least one embodiment;

[0029] FIG. 27 illustrates a system, in accordance with at least one embodiment;

[0030] FIG. 28 illustrates an exemplary integrated circuit, in accordance with at least one embodiment;

[0031] FIG. 29 illustrates a computing system, according to at least one embodiment;

[0032] FIG. 30 illustrates an APU, in accordance with at least one embodiment;

[0033] FIG. 31 illustrates a CPU, in accordance with at least one embodiment;

[0034] FIG. 32 illustrates an exemplary accelerator integration slice, in accordance with at least one embodiment;

[0035] FIGS. 33A-33B illustrate exemplary graphics processors, in accordance with at least one embodiment;

[0036] FIG. 34A illustrates a graphics core, in accordance with at least one embodiment;

[0037] FIG. 34B illustrates a GPGPU, in accordance with at least one embodiment;

[0038] FIG. 35A illustrates a parallel processor, in accordance with at least one embodiment;

[0039] FIG. 35B illustrates a processing cluster, in accordance with at least one embodiment;

[0040] FIG. 35C illustrates a graphics multiprocessor, in accordance with at least one embodiment;

[0041] FIG. 36 illustrates a graphics processor, in accordance with at least one embodiment;

[0042] FIG. 37 illustrates a processor, in accordance with at least one embodiment;

[0043] FIG. 38 illustrates a processor, in accordance with at least one embodiment;

[0044] FIG. 39 illustrates a graphics processor core, in accordance with at least one embodiment;

[0045] FIG. 40 illustrates a PPU, in accordance with at least one embodiment;

[0046] FIG. 41 illustrates a GPC, in accordance with at least one embodiment;

[0047] FIG. 42 illustrates a streaming multiprocessor, in accordance with at least one embodiment;

[0048] FIG. 43 illustrates a software stack of a programming platform, in accordance with at least one embodiment;

[0049] FIG. 44 illustrates a CUDA implementation of a software stack of FIG. 43, in accordance with at least one embodiment;

[0050] FIG. 45 illustrates a ROCm implementation of a software stack of FIG. 43, in accordance with at least one embodiment;

[0051] FIG. 46 illustrates an OpenCL implementation of a software stack of FIG. 43, in accordance with at least one embodiment;

[0052] FIG. 47 illustrates software that is supported by a programming platform, in accordance with at least one embodiment;

[0053] FIG. 48 illustrates compiling code to execute on programming platforms of FIGS. 43-46, in accordance with at least one embodiment;

[0054] FIG. 49 illustrates in greater detail compiling code to execute on programming platforms of FIGS. 43-46, in accordance with at least one embodiment;

[0055] FIG. 50 illustrates translating source code prior to compiling source code, in accordance with at least one embodiment;

[0056] FIG. 51A illustrates a system configured to compile and execute CUDA source code using different types of processing units, in accordance with at least one embodiment;

[0057] FIG. 51B illustrates a system configured to compile and execute CUDA source code of FIG. 51A using a CPU and a CUDA-enabled GPU, in accordance with at least one embodiment;

[0058] FIG. 51C illustrates a system configured to compile and execute CUDA source code of FIG. 51A using a CPU and a non-CUDA-enabled GPU, in accordance with at least one embodiment;

[0059] FIG. 52 illustrates an exemplary kernel translated by CUDA-to-HIP translation tool of FIG. 51C, in accordance with at least one embodiment;

[0060] FIG. 53 illustrates non-CUDA-enabled GPU of FIG. 51C in greater detail, in accordance with at least one embodiment;

[0061] FIG. 54 illustrates how threads of an exemplary CUDA grid are mapped to different compute units of FIG. 53, in accordance with at least one embodiment;

[0062] FIG. 55 illustrates how to migrate existing CUDA code to Data Parallel C++ code, in accordance with at least one embodiment; and

[0063] FIG. 56 illustrates components of a system to access a large language model, according to at least one embodiment.

DETAILED DESCRIPTION

[0064] In the following description, numerous specific details are set forth to provide a more thorough understanding of at least one embodiment. However, it will be apparent to one skilled in the art that inventive concepts described herein may be practiced without one or more of these specific details, and that two or more aspects of any one or more embodiments described herein may be combined.

[0065] In at least one embodiment, one or more processors comprising one or more circuits are to identify one or more pixels within one or more polygons based, at least in part, on whether those one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries. In at least one embodiment, one or more processors comprising one or more circuits perform operations to identify one or more portions (fractions) of pixels that are within a polygon, but do not cross with an edge of a polygon, by using information about edges of other pixels that do cross with one or more edges of a polygon. In at least one embodiment, one or more processors comprising one or more circuits identify portions of pixel edges that are within a polygon to use them as width values to calculate areas of other pixels within a polygon. In at least one embodiment, information about pixel edges used as width values correspond to common pixel edges shared between adjacent pixels. In at least one embodiment, one or more processors comprising one or more circuits perform operations that use information about pixel edges as width values with respect to pixels within a column or row of a pixel grid. In at least one embodiment, one or more processors comprising one or more circuits use prefix sums of width values, at least in part, to output a total amount of each pixel that is within a polygon.

[0066] In at least one embodiment, for example, Pixel A crosses an edge of a polygon, while Pixel B does not. In at least one embodiment, a common pixel edge shared between Pixel A and Pixel B is completely enclosed by a polygon, and that polygon edge completely covers that common pixel edge as it crosses Pixel A. In at least one embodiment, one or more processors comprising one or more circuits perform operations to identify that Pixel B must be entirely within that polygon because that polygon edge completely covers a common pixel edge as it crosses Pixel A. In at least one embodiment, a processor comprising one or more circuits perform operations that use information about how a polygon edge covers a common edge to also calculate portions of other pixels within that polygon that share a column or row with Pixels A and B. In at least one embodiment, one or more processors comprising one or more circuits, by performing techniques described herein, are able to calculate how much of Pixel B is within a polygon without comparing Pixel B with every edge of a polygon.

[0067] In at least one embodiment, a pixel is a two-dimensional (2D) element that represents corresponding display element in a display device, such as a screen of a mobile device, a computer monitor, or a television. In at least one embodiment, a pixel represents a corresponding printing element in a printing device, such as a mask writer

used in computational lithography. In at least one embodiment, one or more processors comprising one or more circuits are to identify one or more pixels within one or more polygons to generate, display, or otherwise activate, on a display and/or printing device. In at least one embodiment, a pixel is within a polygon if an amount of that pixel that is inside that polygon meets or exceeds a threshold value (e.g., 0.50 or 50%). In at least one embodiment, a pixel that is adjacent to another pixel shares a common pixel edge (pixel border) with that other pixel. In at least one embodiment, a polygon boundary (polygon edge) crosses a pixel if that polygon traverses, intersects, and/or enters a pixel if that polygon were overlaid or placed within grid of pixels (e.g., pixel grid).

[0068] In at least one embodiment, one or more processors comprising one or more circuits are to perform mathematical operations, including prefix sum operations, of an algorithm to identify how much of a pixel is within a polygon. In at least one embodiment, one or more processors comprising one or more circuits are to identify an amount of one or more pixels within a polygon based, at least in part, on an amount of one or more edges of those one or more pixels covered by one or more edges of that polygon. In at least one embodiment, an amount of one or more pixels within a polygon is a value representing an area of a pixel within that polygon. In at least one embodiment, an amount of a pixel within a polygon is referred to as an amount inside and/or covered by that polygon. In at least one embodiment, an amount of an edge of a pixel (e.g., pixel edge) covered by one or more edges of a polygon is a portion of that pixel edge output by, or similar to, one or more mathematical operations used to perform scalar and/or vector projections onto a horizontal pixel edge. In at least one embodiment, an amount of an edge of a pixel (pixel edge or pixel boundary) covered by one or more edges of a polygon is referred to as a pixel edge coverage value, pixel edge coverage, or edge coverage. In at least one embodiment, an amount of a pixel edge covered by a polygon edge is a value expressed as a fraction, percentage, length, or some combination thereof.

[0069] In at least one embodiment, a pixel edge coverage is an amount of a pixel's edge that is shared with an adjacent pixel above and that is covered by one or more portions of one or more polygon edges that intersect with an adjacent pixel above. In at least one embodiment, a pixel adjacent to another pixel is a pixel next to that other pixel. In at least one embodiment, a pixel edge coverage value of a pixel is used, at least in part, to identify amounts within a polygon of other pixels beneath that pixel, adjacent or otherwise, and within that pixel's column.

[0070] In at least one embodiment, advantages of techniques described herein include reading data values of pixel edge coverages to, at least in part, identify an amount of a pixel within that polygon, even though that pixel does not intersect with an edge of that polygon. In at least one embodiment, advantages of techniques described herein include reading data values of pixel edge coverages to, at least in part, identify an amount of a pixel within that polygon without having to refer to information about a polygon edge (e.g., direction, beginning coordinates, ending coordinates) to make that calculation. In at least one embodiment, advantages of techniques described herein include parallelization of computing operations based on each segment of a polygon edge within a pixel. In at least one embodiment, advantages of techniques described herein

include a precision in calculations dependent on a precision of data values used by mathematical operations (e.g., add, multiply), and not dependent on a resolution of a pixel grid.

[0071] FIG. 1 illustrates a block diagram of a system 100 that includes one or more processors comprising one or more circuits to identify one or more pixels within one or more polygons based, at least in part, on whether one or more pixels are adjacent to one or more pixels that cross (traverse) one or more polygon edges, according to at least one embodiment. In at least one embodiment, one or more aspects of one or more embodiments described herein in conjunction with FIG. 1 are combined with one or more aspects of one or more embodiments described herein, including those described at least in conjunction with FIGS. 2-23. In at least one embodiment, one or more processors perform one or more operations of system 100. In at least one embodiment, any one or more processors described herein comprise one or more circuits. In at least one embodiment, one or more processors that perform one or more operations of system 100 are any one processor, or combination of processors, described herein, including processor(s) 104 of FIG. 1, CPU 3100 of FIG. 31, graphics processor 3340 of FIG. 33B, general purpose graphics processing unit (GPGPU) 3430 of FIG. 34B, and parallel processing unit (PPU) 4000 of FIG. 40. In at least one embodiment, processor(s) 104 perform an operation used by system 100, such as loading/storing pixel edge coverages in arithmetic logic unit(s) (ALUs), such as ALUs 3716 and 3718 of FIG. 37. In at least one embodiment, processor(s) 104 perform one or more operations described in conjunction with process 2100 of FIG. 21, such as calculating bottom edge pixel coverages. In at least one embodiment, processor(s) 104 perform one or more operations described in conjunction with process 2200 of FIG. 22, such as performing API function(s) to calculate pixel coverages using prefix sums of bottom pixel edge coverages. In at least one embodiment, processor(s) 104 perform one or more operations described in conjunction with FIG. 23, such as performing driver/runtime 2304.

[0072] In at least one embodiment, system 100 includes and/or otherwise obtains polygon data as input data, which is depicted as polygon input data 102. In at least one embodiment, polygon input data 102 includes vertex data and/or vector data. In at least one embodiment, processor(s) 104 are to perform operations to convert any representations of an environment, model, object, or some combination thereof, into pixel representations. In at least one embodiment, when processor(s) 104 perform operations to, at least in part, convert any data representation of an environment, model, object, or some combination thereof, into a pixel representation, performing such operations is referred to as rasterization. In at least one embodiment, one or more aspects of a rasterization process is referred to as a shading. In at least one embodiment, processor(s) 104 perform operations to, at least in part, identify one or more pixels within a polygon to be used with rasterization. In at least one embodiment, one or more operations described herein are performed by one or more shader cores, such as shader cores 3355A-3355N of graphics processor 3340 of FIG. 33B.

[0073] In at least one embodiment, polygon input data 102 includes information about polygons, such as triangles, used to represent a 3D model and/or object. In at least one embodiment, polygon input data 102 includes data representing edges and vertices of polygons. In at least one embodiment, polygon input data 102 defines all polygon

edges to be represented by pixels of a pixel grid. In at least one embodiment, polygon input data **102** includes data that, at least in part, define polygon edges based, at least in part, on properties (characteristics) of those polygon edges such as position, start point, end point, direction, attributes, or some combination thereof, and as otherwise described herein. In at least one embodiment, polygon input data **102** includes information about vertices such as their position, color, reflectance, texture, or some combination thereof. In at least one embodiment, polygon input data **102** uses vertices to represent edges that make up polygons to be represented as pixels. In at least one embodiment, polygon input data **102** represents polygon edges as vectors with directions. In at least one embodiment, polygon input data **102** includes information about edge directions, which are used, according to a convention, to indicate whether space and/or pixels lying in one direction or one side of that edge is inside or outside a polygon containing that edge. In at least one embodiment, polygon input data **102** is stored as one or more tensors.

[0074] In at least one embodiment, polygon input data **102** includes data representing one or more polygons that are each represented by one or more triangles. In at least one embodiment, polygon input data **102** includes two-dimensional (2D) coordinates of vertices of polygon, such as:

[0075] Polygon 1: (0,10), (0,20), (10,20);

[0076] Polygon 2: (-200, -200), (-200, -100), (-100, -100), (-100, -200); and

[0077] Polygon 3: (-150, -150), (-150, -100), (-120, -150).

[0078] In at least one embodiment, data input into processor(s) **104** includes one or more pixel grids. In at least one embodiment, a pixel grid uses a convention where an origin of that pixel grid is located in a lower-left position of that grid. In at least one embodiment, input data that represents a pixel grid includes information about spacing between pixel points, and dimensions such as width and height.

[0079] In at least one embodiment, processor(s) **104** are any one processor, or combination of processors, described herein, including processor(s) **104** of FIG. 1, CPU **3100** of FIG. 31, graphics processor **3340** of FIG. 33B, general purpose graphics processing unit (GPGPU) **3430** of FIG. 34B, and parallel processing unit (PPU) **4000** of FIG. 40. In at least one embodiment, any modules depicted as being implemented on processor(s) **104** are implemented on any one or combination of processors. In at least one embodiment, processor(s) **104** perform any operation used, at least in part, in rasterization. In at least one embodiment, any one or more modules of processor(s) **104** are implemented as a part of any one or more other modules of processor(s) **104**. In at least one embodiment, any one or more processor(s) **104** are implemented as a part of any one or more modules depicted in FIG. 1.

[0080] In at least one embodiment, as used in any embodiment described herein, unless otherwise clear from context or stated explicitly to contrary, terms such as “system,” “device,” “component,” or “module,” and nominalized verbs (e.g., compiler, shader, and/or other terms) each refers to any combination of software logic, firmware logic, hardware logic, and/or circuitry configured to provide functionality described herein. In at least one embodiment, any system, device, component, and module, described herein are combined and/or communicatively connected with at least one other component, system, device, component, and

module regardless of how such components are described to be combined and/or communicatively connected in other embodiments. In at least one embodiment, software may be embodied as a software package, code, and/or instruction set or instructions. In at least one embodiment, hardware includes, singly or in any combination, hardwired circuitry, programmable circuitry, state machine circuitry, fixed function circuitry, execution unit circuitry, and/or firmware that stores instructions executed by programmable circuitry. In at least one embodiment, modules may, collectively or individually, be embodied as circuitry that forms part of a larger system, for example, an integrated circuit (IC), system on-chip (SoC), and so forth. In at least one embodiment, any one or more architectures of any circuits of one or more modules are represented as a register-transfer level (RTL) representation and/or another fabless representation that may be licensed and/or used in tape-out, a final phase in IC design before being used in manufacturing an IC.

[0081] In at least one embodiment, processor(s) **104** performs operations of polygon transformation module **106** to transform (or modify) polygon input data **102** into arrays (or tensors) of data that can later be converted into pixels. In at least one embodiment, processor(s) **104** perform operations of polygon transformation module **106** to convert three-dimensional (3D) position coordinates of vertices of polygon input data **102** into two-dimensional (2D) position coordinates that correspond to a pixel grid, a display device, a printing device, or some combination thereof. In at least one embodiment, processor(s) **104** performs operations of polygon transformation module **106** to transform polygon input data **102** into formats suitable to be processed with other tasks of a rasterization process, tasks such as lighting, projection, and clipping. In at least one embodiment, processor(s) **104** performs operations of polygon transformation module **106** and outputs transformed polygon information to be received and/or otherwise obtained by prefix sum pixel coverage module **108**.

[0082] In at least one embodiment, processor(s) **104** perform operations of prefix sum pixel coverage module **108** to identify one or more pixels within a polygon based, at least in part, on whether those one or more pixels are adjacent to one or more pixels that cross (traverse) one or more polygon boundaries. In at least one embodiment, processor(s) **104** perform operations of prefix sum pixel coverage module **108** to identify portions of one or more edges of a polygon that traverse one or more pixels. In at least one embodiment, prefix sum pixel coverage module **108** is, at least in part, any one or combination of prefix sum pixel coverage modules described herein in conjunction with FIGS. 2-20. In at least one embodiment, a reference to a module, such as prefix sum pixel coverage module **108**, performing one or more operations refers to processor(s) **104** performing one or more operations of that module. In at least one embodiment, a reference to a module, such as prefix sum pixel coverage module **108**, performing one or more operations refers to processor(s) implemented on that module to perform those one or more operations.

[0083] In at least one embodiment, prefix sum pixel coverage module **108** performs a prefix sum pixel coverage algorithm that comprises a set of operations used to identify one or more pixels within a polygon. In at least one embodiment, to identify one or more values and/or objects refers to processor(s) **104** that are to perform operations to calculate, compute, and/or otherwise determine those one or

more values and/or objects. In at least one embodiment, identifying a value or object includes calculating and/or computing one or more mathematical and/or logical operations.

[0084] In at least one embodiment, a prefix sum pixel coverage algorithm performed by prefix sum pixel coverage module **108** comprises steps such as:

[0085] Step 1: Compute partial coverages;

[0086] Step 2: Compute pixel edge coverage;

[0087] Step 3: Prefix sum down columns of pixels;

[0088] Step 4: Sum results of Steps 1 and 3;

[0089] In at least one embodiment, a prefix sum pixel coverage module performs an algorithm to output values that represent each represent an amount of a pixel that lies within a pixel, where such values are used to generate a rasterized an image polygon input data **102**.

[0090] In at least one embodiment, steps of an algorithm are performed in any order. In at least one embodiment, any one or more steps of an algorithm are omitted. In at least one embodiment, any one or more steps of an algorithm include one or more operations that are performed in any order. In at least one embodiment, prefix sum pixel coverage module **108** parallelizes computations described herein, such as computing bottom pixel edge coverages, over two or more portions of one or more edges of that polygon, and as described further herein. In at least one embodiment, prefix sum pixel coverage module **108** parallelizes computations of an amount of one or more pixels covered by a polygon over two or more edges of that polygon, and as described further herein.

[0091] In at least one embodiment, prefix sum pixel coverage module **108** performs step 1 of a prefix sum pixel coverage algorithm, which includes operations to compute partial coverages. In at least one embodiment, prefix sum pixel coverage module **108** performs operations to identify an area (an amount) of one or more pixels within a polygon based, at least in part, on one or more segments of one or more edges of that polygon that cross those one or more pixels, and as described further herein. In at least one embodiment, an area of one or more pixels within a polygon based, at least in part, on one or more segments of one or more edges of that polygon that cross those one or more pixels is referred to as a partial coverage. In at least one embodiment, computing a partial coverages includes parallelizing computations of one or more amounts of one or more pixels within a polygon over two or more portions of one or more edges of that polygon. In at least one embodiment, identifying partial coverages means computing partial coverages. In at least one embodiment, partial coverages are computed for pixels that have one or more polygon edges (polygon boundaries) that cross (traverse) those pixels. In at least one embodiment, a partial coverage is an amount of a pixel, such as an area or portion of that pixel, that is covered by a polygon. In at least one embodiment, an amount of a pixel covered by a polygon refers to an amount of a pixel that is within or inside that polygon. In at least one embodiment, prefix sum pixel coverage module **108** performs operations to compute partial coverages of pixels, where each of those pixels cross one or more polygon edges, and as further described herein. In at least one embodiment, pixels that cross one or more polygon edges are pixels that have one or more polygon edges that cross those pixels. In at least one embodiment, prefix sum pixel coverage module **108** computes pixel coverages by using one or more polygon

edges and one or more pixel edges (pixel boundaries) as boundaries used to calculate areas covered by that polygon.

[0092] In at least one embodiment, polygon edges that point from right-to-left are used to generate partial coverages, in accordance with a convention that defines characteristics of polygons. In at least one embodiment, polygon edges that point from left-to-right are used to identify negative partial coverages, which are areas outside of a polygon. In at least one embodiment, partial coverages are summed together with other pixel areas covered by a polygon to output one or more total areas of one or more pixels covered by a polygon. In at least one embodiment, prefix sum pixel coverage module **108** calculates other pixel areas when respective pixels do not have a polygon edge that cross those pixels.

[0093] In at least one embodiment, prefix sum pixel coverage module **108** identifies a portion of a pixel crossed by a polygon edge by identifying one or more locations or points where that polygon edge intersects with any pixel edge. In at least one embodiment, a boundary used to calculate a partial coverage of a pixel spans a length between a location, or point, of intersection between a polygon edge and a pixel edge and a pixel's bottom edge, as depicted at least in FIGS. **8** and **13**. In at least one embodiment, operations performed by prefix sum pixel coverage module **108** to calculate partial coverages are described further herein at least in conjunction with FIGS. **7-13**.

[0094] In at least one embodiment, prefix sum pixel coverage module **108** performs step 2 of a prefix sum pixel coverage algorithm, which includes operations to compute bottom pixel edge coverages. In at least one embodiment, prefix sum pixel coverage module **108** identifies portions, such as bottom edge coverages, of one or more edges of a polygon based, at least in part, on intersections between one or more edges of that polygon and one or more adjacent pixels, and as described further herein at least in conjunction with FIG. **13**. In at least one embodiment, prefix sum pixel coverage module **108** identifies portions, such as bottom edge coverages, of one or more edges of a polygon based, at least in part, on one or more locations, such as vertices, where two or more edges of that polygon meet, and as described further herein at least in conjunction with FIG. **13**. In at least one embodiment, prefix sum pixel coverage module **108** performs operations to identify an area of one or more pixels within a polygon based, at least in part, on information, such as bottom pixel edge coverages, of one or more other pixels in a column, and as described further herein. In at least one embodiment, prefix sum pixel coverage module **108** perform operations to identify an amount of one or more pixels within a polygon based, at least in part, on an amount of one or more edges of those one or more pixels covered by one or more edges of that polygon, and as described further herein. In at least one embodiment, a bottom pixel edge coverage is an amount of a pixel edge covered by a polygon. In at least one embodiment, a bottom pixel edge is a bottom pixel edge of pixel that crosses one or more polygon edges. In at least one embodiment, an amount of a bottom pixel edge is a length or fraction of that bottom pixel edge that is covered by a portion of a polygon edge that crosses that pixel.

[0095] In at least one embodiment, a bottom pixel edge coverage is based on a polygon edge that points from right-to-left. In at least one embodiment, a bottom pixel edge coverage based on a polygon edge that points from right-

to-left is used to calculate a pixel area covered by a polygon, in accordance with a convention that defines characteristics of polygons based on directions in which polygon edges point, and as described further herein. In at least one embodiment, a bottom pixel edge coverage based on a polygon edge that points from left-to-right is used to calculate a pixel area outside polygon. In at least one embodiment, a bottom pixel edge coverage based on a polygon edge that points from left-to-right is referred to as a negative bottom pixel edge coverage. In at least one embodiment, an amount of a bottom pixel edge covered by a polygon edge is calculated by using a scalar projection of a portion (segment) of that polygon edge onto that bottom pixel edge. In at least one embodiment, such a projection is a projection of one or more portions of one or more edges of a polygon onto an edge of one or more adjacent pixels, and as described further herein. In at least one embodiment, a bottom pixel edge coverage is equivalent to a scalar x-component of a vector that represents a portion of a polygon edge that crosses a pixel. In at least one embodiment, a bottom pixel edge coverage is based on a portion of a polygon edge that is within a pixel. In at least one embodiment, a bottom pixel edge coverage is based on a portion of a polygon edge that is defined by one or more locations at which that polygon edge intersects with one or more pixel edges, which is depicted at least in FIGS. 15 and 16. In at least one embodiment, a bottom pixel edge coverage is based on a portion of a polygon edge that is defined by one or more locations at which that polygon edge meets another polygon edge, such as at a vertex, which is depicted at least in FIGS. 15 and 16.

[0096] In at least one embodiment, a bottom pixel edge coverage that is used, in part, to identify an amount of a pixel covered by that polygon is a common pixel edge that is shared by an adjacent pixel. In at least one embodiment, an algorithm defines an adjacent pixel as being above another pixel. In at least one embodiment, an algorithm defines an adjacent pixel by being below, left, or right of another pixel. In at least one embodiment, prefix sum pixel coverage module 108 identifies whether an adjacent pixel is crossed by a polygon edge. In at least one embodiment, prefix sum pixel coverage module 108 performs operations to identify whether one or more pixels are adjacent to one or more other pixels that cross one or more polygon edges (polygon boundaries), and as described further herein. In at least one embodiment, identifying whether an adjacent pixel is crossed by a polygon edge is used to determine whether a bottom pixel edge coverage of a pixel edge common to or shared between that adjacent pixel and another pixel will be used to calculate an amount of that other pixel covered by that polygon. In at least one embodiment, a bottom pixel edge coverage is used to calculate an amount of a pixel covered by a polygon, where that pixel is below but not adjacent to a pixel corresponding to that bottom pixel edge coverage. In at least one embodiment, operations used to calculate a bottom pixel edge coverage is described further herein at least in conjunction with FIGS. 15-16 and 18.

[0097] In at least one embodiment, prefix sum pixel coverage module 108 performs operations that use a bottom pixel edge coverage to calculate a width value, which is then used to calculate an area of a pixel that is covered by a polygon. In at least one embodiment, prefix sum pixel coverage module 108 performs operations to identify an amount of one or more pixels within a polygon by using one

or more width values by using one or more portions of one or more edges of that polygon, and as otherwise described herein. In at least one embodiment, prefix sum pixel coverage module 108 performs operations to identify an amount of one or more pixels within a polygon based, at least in part, on one or more portions of one or more edges of that polygon to calculate width value as described further herein. In at least one embodiment, prefix sum pixel coverage module 108 performs operations that use a negative bottom pixel edge coverage to calculate a width value, which is then used to calculate an area of a pixel outside of a polygon. In at least one embodiment, a bottom pixel edge coverage value that is expressed as a fraction of a bottom pixel edge covered by a polygon also serves as a width value when a pixel width is equal to 1. In at least one embodiment, a bottom pixel edge coverage is propagated to one or more calculations of pixel areas covered by a polygon, where those pixels are located below that bottom pixel edge. In at least one embodiment, a bottom pixel edge coverage is used as a width value to calculate areas that correspond to one or more pixels below a pixel edge that corresponds to that bottom pixel edge coverage. In at least one embodiment, operations used to calculate areas using a bottom pixel edge coverage are described further herein at least in conjunction with FIGS. 15-17.

[0098] In at least one embodiment, a pixel area covered by a polygon and calculated by using a partial edge coverage is referred to as an edge-based pixel coverage. In at least one embodiment, a pixel areas outside of a polygon calculated by using a negative bottom pixel edge coverage is referred to as a negative edge-based pixel coverage. In at least one embodiment, pixel areas calculated by using a bottom pixel edge coverage are equal to that bottom pixel edge coverage when a pixel width is 1 and a pixel height is 1. In at least one embodiment, when a pixel width and height are both 1, pixel areas calculated by using bottom pixel edge coverages are referred to as bottom pixel edge coverages because both values are equal to each other.

[0099] In at least one embodiment, prefix sum pixel coverage module 108 performs operations of step 3 of a prefix sum pixel coverage algorithm, which includes operations used to perform a prefix sum of bottom pixel edge coverages down a column of pixels. In at least one embodiment, prefix sum pixel coverage module 108 performs operations to use one or more prefix sums along one or more dimensions to identify an amount of one or more pixels within a polygon, and as described further herein. In at least one embodiment, a dimension refers to a column or a row of a pixel grid. In at least one embodiment, prefix sum pixel coverage module 108 performs prefix sum operations on edge-based pixel coverages (or bottom pixel edge coverages) corresponding to pixels in a column of pixels, from top to bottom. In at least one embodiment, operations used to calculate edge-based pixel coverages and perform prefix sums are described further herein at least in conjunction with FIGS. 16-19.

[0100] In at least one embodiment, prefix sum pixel coverage module 108 performs step 4 of a prefix sum pixel coverage algorithm, which includes operations that sum together partial coverages, negative areas, edge-based pixel coverages (or bottom pixel edge coverages), negative edge-based pixel coverages (or negative bottom pixel edge coverages) of each pixel. In at least one embodiment, a sum of partial coverages, negative areas, edge-based pixel coverages (or bottom pixel edge coverages), negative edge-based

pixel coverages (or negative bottom pixel edge coverages) of each pixel equals a total amount of that pixel that is covered by polygon. In at least one embodiment, a total amount of a pixel covered by a polygon is referred to as a coverage fraction of a pixel or a total pixel coverage.

[0101] In at least one embodiment, prefix sum pixel coverage module **108** performs operations that compare a total pixel coverage with a threshold value (e.g., 0.5, 50%) to identify whether a pixel is covered (within) a polygon. In at least one embodiment, processor(s) **104** perform operations of prefix sum pixel coverage module **108** to use one or more pixels identified as being within a polygon to, at least in part, perform rasterization and/or computational lithography tasks. In at least one embodiment, pixels identified as being within a polygon are generated, displayed, printed, and/or otherwise activated on displays and/or printers to output a rasterized image of polygon input data **102**. In at least one embodiment, a rasterized image of polygon input data **102** is used, at least in part, to generate lithography masks of a computational lithography process.

[0102] In at least one embodiment, processor(s) **104** perform operations of prefix sum pixel coverage module **108** to receive and/or otherwise obtain data about polygon edges that make up one or more polygons. In at least one embodiment, processor(s) **104** perform operations of prefix sum pixel coverage module **108** to receive and/or otherwise obtain data about edges that make up one or more polygons, where such data indicates whether an area to a left or right of an edge is inside or outside a polygon. In at least one embodiment, processor(s) perform operations of prefix sum pixel coverage module **108** to identify each portion (segment) of a polygon edge that traverses a pixel of a pixel grid. In at least one embodiment, processor(s) perform operations of prefix sum pixel coverage module **108** to identify portions of pixel edges covered by polygon edges that traverse respective pixels. In at least one embodiment, a pixel grid is a grid of pixel points. In at least one embodiment, a pixel point is a point of a pixel that represents a portion of a pixel. In at least one embodiment, a pixel point is a corner of that pixel. In at least one embodiment, any number of pixel points are used to identify portions of a pixel. In at least one embodiment, for example, a pixel is represented by 144 equally spaced points, with 16 points positioned in 9 equally-sized sections of that pixel. In at least one embodiment, a pixel grid resolution refers to a number of pixel points used to represent portions of each pixel of that pixel grid. In at least one embodiment, representing portions of a pixel by points is referred to as sub-pixel mapping, and a pixel grid that depicts pixel points is referred to as a pixel point grid or a sub-pixel map. In at least one embodiment, a pixel point is referred to as a grid point.

[0103] In at least one embodiment, processor(s) **104** call, invoke, or otherwise perform APIs of prefix sum pixel coverage API(s) module **110**. In at least one embodiment, APIs are API functions as described further herein at least in conjunction with FIG. 23. In at least one embodiment, APIs are one or more APIs **2310** of FIG. 23. In at least one embodiment, processors perform operations of prefix sum pixel coverage module **108** to perform prefix sum pixel coverage API(s) module **110**. In at least one embodiment, prefix sum pixel coverage API(s) module **110** is implemented as part of prefix sum pixel coverage module **108**. In at least one embodiment, prefix sum pixel coverage API(s) module **110** includes a library of API functions used to

perform any one or more operations described herein. In at least one embodiment, a user or software, such as software program **2302** of FIG. 23, causes processor(s) **104** to call, invoke, and/or otherwise perform APIs **2310** of FIG. 23. In at least one embodiment, processor(s) **104** call, invoke, and/or otherwise perform APIs by performing operations and/or functions of those APIs. In at least one embodiment, processor(s) **104** perform operations and/or functions of APIs to receive and/or otherwise obtain as input, polygon input data **102**. In at least one embodiment, processor(s) **104** perform operations and/or functions of APIs to output at least a pixel grid, partial coverage, negative area, edge coverage, negative edge coverage, coverage fraction of a pixel, indication of a pixel within a polygon, or some combination thereof.

[0104] In at least one embodiment, processor(s) **104** output coverage fractions of pixels **112**, which is data generated by prefix sum pixel coverage module **108** and/or prefix sum pixel coverage API(s) module **110** to be used to generate a rasterized image of polygon input data **102**. In at least one embodiment, polygon input data **102** is referred to as a polygon dataset. In at least one embodiment, coverage fractions of pixels **112** is data that includes, for each pixel, one or more indications of what portion and/or what amount of a pixel is covered by one or more polygons.

[0105] FIG. 2 illustrates a block diagram of system **200** that includes a prefix sum pixel coverage module used, at least in part, to identify one or more pixels within one or more polygons based, at least in part, on whether those one or more pixels are adjacent to one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, one or more figures depicted herein are not too scale. In at least one embodiment, prefix sum pixel coverage module **202** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, prefix sum pixel coverage module **202** receives and/or otherwise obtains, as input data, pixel point grid **204**. In at least one embodiment, prefix sum pixel coverage module **202** generates pixel point grid **204**.

[0106] In at least one embodiment, prefix sum pixel coverage module **202** uses data of a pixel point grid with a polygon to identify locations, orientations, positions, or some combination thereof, of portions of polygon edges that traverse pixels. In at least one embodiment, a polygon comprises directional polygon edges, such as polygon edge **206**.

[0107] In at least one embodiment, pixel point grid **204** comprises pixel points, such as pixel point **208**, that represent corners of a pixel. In at least one embodiment, prefix sum pixel coverage module **202** identifies a pixel edge of each pixel of pixel grid **204**. In at least one embodiment, pixel edges (or pixel boundaries) of a pixel are depicted as dark lines of pixel **210**.

[0108] In at least one embodiment, FIG. 2 depicts a polygon represented by polygon input data **102** of FIG. 1 that follows a convention that defines regions to a left of a polygon edge as being inside that polygon. In at least one embodiment, a convention defines regions completely enclosed by polygon edges that are arranged in a counter-clockwise manner to be inside that polygon. In at least one embodiment, a convention defines regions completely enclosed by polygon edges that are arranged in a clockwise manner to be outside that polygon, and therefore, those

regions represent a hole. In at least one embodiment, a convention defines a polygon arranged in a counterclockwise manner as a positive polygon. In at least one embodiment, a convention defines a polygon arranged in a clockwise manner as a negative polygon. In at least one embodiment, a left side of a polygon edge is determined from a point of view from an origin of a polygon edge and looking at an end, or arrow end, of that polygon edge. In at least one embodiment, a input polygon data defines a left side of an upward-pointing polygon edge to be inside a polygon and a left side of a downward-point edge to be outside a polygon. In at least one embodiment, any one or more aspects of one or more conventions used to characterize polygon data can be used with techniques described herein. In at least one embodiment, FIG. 2 visualizes a polygon as being overlaid pixel point grid 204 for explanatory purposes. In at least one embodiment, prefix sum pixel coverage module 202 performs operations of system 200 to generate objects such as pixel point grid 204 and one or more polygons, using data representing such objects, without visualizing or displaying those objects.

[0109] FIG. 3 illustrates another version of block diagram of system 300 that includes a prefix sum pixel coverage module used, at least in part, to identify one or more pixels within one or more polygons (FIG. 3 only depicts one polygon, but more than polygon can be used to identify said one or more pixels) based, at least in part, on whether those one or more pixels are adjacent to one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module 302 is, at least in part, prefix sum pixel coverage module 108 performed by processor(s) 104 of FIG. 1. In at least one embodiment, FIG. 3 depicts another visualization of a pixel point grid, pixel point grid 312, that uses dotted lines as pixel edges, such as pixel edge 306, and pixel corners as pixel points, such as pixel point 308. In at least one embodiment, a pixel point grid, such as one depicted in FIG. 3, is used to explain techniques as further described herein at least in conjunction with FIGS. 13, 14, and 20. In at least one embodiment, pixel point grid 304 comprises pixels with N columns and N rows and is not limited to said columns and rows of pixels as illustrated in FIG. 3. In at least one embodiment, prefix sum pixel coverage module 302 performs calculations that identify pixel coverages in parallel over each portion of a polygon edge that traverses a pixel.

[0110] FIG. 4 illustrates a block diagram of system 400 that includes a prefix sum pixel coverage module used, at least in part, to identify an amount of one or more pixels within a polygon by using, at least in part, prefix sums of pixel edge coverages, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module 402 is, at least in part, prefix sum pixel coverage module 108 performed by processor(s) 104 of FIG. 1. In at least one embodiment, FIG. 4 depicts an overview of how areas are identified as being covered by a polygon based, at least in part, on portions of polygon edges that traverse a pixel, which is described further herein. In at least one embodiment, prefix sum pixel coverage module 402 receives or generates one or more pixel grids. In at least one embodiment, prefix sum pixel coverage module 402 receives or generates two or more blank pixel grids, with one of those pixel grids being used to calculate partial coverages as described further herein. In at least one embodiment, a

pixel grid used to calculate partial coverages is referred to as a cover pixel grid. In at least one embodiment, another blank pixel grid is used to calculate pixel edge coverages, as described further herein at least in conjunction with FIGS. 15-17. In at least one embodiment, a blank pixel grid used to calculate pixel edge coverages is referred to by a pixel edge used, such as a left pixel grid or bottom pixel grid, where pixel edges used is described further herein at least in conjunction with FIGS. 15-17.

[0111] In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm in combination with, or as an alternative to that described in conjunction by FIG. 1, by first beginning with two blank (zeroed) pixel grids referred to as cover and left pixel grids. In at least one embodiment, prefix sum pixel coverage module 402 determines, for each polygon edge, which pixels those polygon edges traverse. In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm to determine, for each of those pixels determined to traverse a polygon edge, what fraction of that pixel is left of that polygon edge, and then add that fraction to a cover pixel grid. In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm to determine, for each of those pixels determined to traverse a polygon edge, what fraction of a left edge of that pixel is left of that polygon edge, and then add that fraction to a left pixel grid. In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm by computing a prefix sum of each fraction of each horizontal row of a left grid. In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm to sum (add) together fractions and/or other values corresponding to each respective pixel in a cover pixel grid and left pixel grid. In at least one embodiment, prefix sum pixel coverage module 402 performs an algorithm that allows a rasterization method to parallelize a step, such as an initial step, over polygon edges. In at least one embodiment, prefix sum pixel coverage module 402 performs operations, such as adds, in a cover and left grid are atomic so that each operation based on a polygon edge is independent of each other. In at least one embodiment, prefix sum pixel coverage module 402 performs operations along a dimensions, such as rows or columns, that are independent of any operations of other rows or columns.

[0112] In at least one embodiment, as described further herein, an algorithm defines an order in which a prefix sum pixel coverage module is to make determinations, calculations, computations, identifications, or some combination thereof, by using any combination of dimensions, directions, and/or orientations, of pixels, pixel edges, polygon edges, or some combination thereof. In at least one embodiment, dimensions, directions, and/or orientations include left, right, top, bottom, column, row, or some combination thereof. In at least one embodiment, a prefix sum pixel coverage module makes determinations not based on fractions of pixels left of a polygon edge or rows of pixels, but rather are based on fractions of pixels below a polygon edge and columns of pixels.

[0113] In at least one embodiment, prefix sum pixel coverage module 402 identifies different pixel coverages as depicted in pixel point grids 404a-c. In at least one embodiment, a partial coverage refers to an area of a pixel bounded by: a portion of a polygon edge that traverses that pixel; any pixel edge directly underneath that polygon edge; and a vertical boundary extended from that polygon edge to that

pixel's bottom pixel edge, as depicted in FIG. 4. In at least one embodiment, areas (depicted as darkened regions) identified in pixel point grids **404a** and **404b** are positive contributions because they are based on polygon edges that point right to left. In at least one embodiment, an area identified in pixel point grid **404c** is a negative contribution because it is based on a polygon edge that points from left to right. In at least one embodiment, a positive contribution area is referred to as a positive coverage area, positive pixel coverage, positive coverage, or positive area. In at least one embodiment, a negative contribution area is referred to as a negative coverage area, negative pixel coverage, negative coverage, or a negative area.

[0114] FIG. 5 illustrates a block diagram of system **500** that includes a prefix sum pixel coverage module used, at least in part, to identify one or more pixels within a polygon by using, at least in part, prefix sums of pixel edge coverages, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **502** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 5 depicts an overview of how areas are identified as being covered by a polygon based, at least in part, on portions of polygon edges that traverse a pixel, which is described further herein. In at least one embodiment, prefix sum pixel coverage module **502** identifies different pixel coverages as depicted in pixel point grids **504a-d**. In at least one embodiment, a negative contribution area is depicted in a pixel of pixel point grid **504a** as empty space. In at least one embodiment, as shown in pixel point grids **504b-d**, a positive contribution area fills in, or adds to, a negative contribution area of pixel point grid **504a** based on a left-to-right polygon edge that shares a vertex with a right-to-left polygon edge.

[0115] FIG. 6 illustrates a block diagram of system **600** that includes a prefix sum pixel coverage module used, at least in part, to identify an amount of one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **602** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 6 depicts an overview of how areas identified as being covered by a polygon are used to calculate a total area of a pixel covered by a polygon, which is described further herein. In at least one embodiment, prefix sum pixel coverage module **602** sums together different positive and negative contribution areas, such as those depicted in FIGS. 4 and 5, to identify a total amount of area covered by a polygon, which is depicted as a darkened area in pixel point grid **604**.

[0116] FIG. 7 illustrates a block diagram of system **700** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **702** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, an area covered by a polygon edge that traverses a pixel is referred to as a partial coverage. In at least one embodiment, one or more partial coverages are summed

with other pixel coverage values described further herein to identify a total area of one or more pixels covered by a polygon.

[0117] In at least one embodiment, prefix sum pixel coverage module **702** performs one or more operations of step 1 of an algorithm that comprises, at least, one or more steps as follows:

[0118] Step 1: Compute partial coverages;

[0119] Step 2: Compute pixel edge coverage;

[0120] Step 3: Prefix sum down columns of pixels;

[0121] Step 4: Sum results of Steps 1 and 3;

[0122] In at least one embodiment, steps of an algorithm include one or more operations that are performed in any order within that step or with any step. In at least one embodiment, one or more steps, and/or one or more operations thereof, of an algorithm described herein is performed in any order. In at least one embodiment, one or more steps, and/or one or more operations thereof, of an algorithm described herein are omitted from being performed.

[0123] In at least one embodiment, prefix sum pixel coverage module **702** computes partial coverages corresponding to a polygon edge depicted as a darker, thicker arrow in pixel point grid **704a**. In at least one embodiment, portions of that polygon edge are shown as separate edges in pixel point grid **704b**. In at least one embodiment, those separate edges are based on locations where a polygon edge intersects with (or crosses with) pixel edges. In at least one embodiment, prefix sum pixel coverage module **702** calculates partial coverages depicted in pixel point grid **704c** based on a formula used to calculate areas of trapezoids and those separate edges depicted in pixel point grid **704b**. In at least one embodiment, a partial coverage formula is as follows:

$$\text{Partial Coverage} = \frac{\sigma_{\text{traps}}(h1 + h2) * w/2}{\text{pixel width} * \text{pixel height}}$$

[0124] In at least one embodiment, σ_{traps} refers to a sum of trapezoidal areas within a pixel. In at least one embodiment, $h1+h2$ refers to height values of a trapezoidal area being calculated. In at least one embodiment, w refers to a width value of a trapezoidal area being calculated. In at least one embodiment, a height and width of a pixel is equal to 1 to simplify calculations. In at least one embodiment, partial coverages are calculated only for pixels that cross one or more polygon edges.

[0125] FIG. 8 illustrates a block diagram of system **800** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **802** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 8 illustrates a prefix sum pixel coverage module **802** that calculates positive contribution areas and negative contribution areas of a pixel as shown in pixel point grids **804a-c**. In at least one embodiment, a portion of a left-to-right polygon edge creates, at least in part, a negative coverage areas of -0.16 in one pixel and -0.05 in another pixel, as depicted in pixel point grid **804a**. In at least one embodiment, a positive coverage area of 0.7 is created, at least in part, by a portion of a right-to-left polygon edge. In

at least one embodiment, a total area of a pixel covered by a polygon is calculated by adding positive coverage area **806** to a positive coverage area of 0.7 and a negative coverage area -0.16 as depicted in pixel point grid **804c**. In at least one embodiment, positive coverage area **806** is calculated by using a pixel edge coverage value described further herein.

[0126] FIG. 9 illustrates a block diagram of system **900** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **902** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 9 illustrates exemplary partial coverage values. In at least one embodiment, a partial coverage value of 0.29 of pixel **906** is based on areas bounded boundaries represented by pixel edges and arrows. In at least one embodiment, bounds represented by arrows originate at locations where a polygon edge intersects with a pixel edge or where two polygon edges meet at a vertex, as depicted in pixel point grid **904**. In at least one embodiment, bounds represented by arrows end at a bottom pixel edge as depicted in pixel point grid **904**. In at least one embodiment, partial coverages are stored in an array.

[0127] FIG. 10 illustrates a block diagram of system **1000** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1002** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 10 illustrates exemplary partial coverage values. In at least one embodiment, a partial coverage value does not indicate a total amount of a pixel that is covered by a polygon, as indicated, for example, by question marks depicted in pixel point grid **1004**. In at least one embodiment, other portions of a pixel covered by a polygon are calculated by using, at least in part, techniques described at least in conjunction with FIGS. 15-17.

[0128] FIG. 11 illustrates a block diagram of system **1100** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1102** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 11 illustrates exemplary partial coverage values and exemplary partial coverages, such as partial coverages **1106a**, **1106b**, and **1106c**. In at least one embodiment, partial coverages **1106a**, **1106b**, and **1106c** are identified based, at least in part, on bounds that are represented as arrows in FIG. 10. In at least one embodiment, FIG. 11 depicts a negative coverage area **1108** that overlaps with a positive coverage **1106c**.

[0129] FIG. 12 illustrates a block diagram of system **1200** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part,

on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1202** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 12 illustrates exemplary partial coverage values and an exemplary negative partial coverage, such as negative partial coverage **1206** created by a left-to-right polygon edge. In at least one embodiment, negative partial coverage **1206** has a value of -0.28 .

[0130] FIG. 13 illustrates a block diagram of system **1300** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1302** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 13 illustrates partial coverages created by a right-to-left polygon edges. In at least one embodiment, partial coverages depicted in FIG. 13, such as partial coverages **1306a** and **1306b**, are positive contributions. In at least one embodiment, pixel point grid **1304** depicts each partial coverage based on right-to-left polygon edges of a polygon as patterned lines, which is also described in a key of FIG. 13. In at least one embodiment, downward-directed dashed arrows represent bounds used to calculate partial coverages and are equivalent to those downward arrows of FIG. 9.

[0131] FIG. 14 illustrates a block diagram of system **1400** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon based, at least in part, on one or more pixels that cross one or more polygon edges, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1402** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 14 illustrates negative areas created by a left-to-right polygon edges. In at least one embodiment, negative areas identified in FIG. 14, such as negative areas **1406a** and **1406b**, are based on left-to-right polygon edges that traverse pixels.

[0132] FIG. 15 illustrates a block diagram of system **1500** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon by using an amount of one or more edges of one or more pixels covered by one or more edges of that polygon, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1502** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 15 illustrates bottom pixel edge coverages and a negative bottom pixel edge coverage. In at least one embodiment, a pixel crosses two portions of two different polygon edges that point from right-to-left, and therefore, includes two partial coverages 0.36 and 0.11 as depicted in pixel point grid **1504a**. In at least one embodiment, 0.36 and 0.11 represent fractions of a total width of a pixel. In at least one embodiment, each pixel of a pixel grid has a height and width of 1.00, as depicted in pixel point grid **1504a**. In at least one embodiment, prefix sum pixel coverage module **1502** performs

operations to use bottom pixel edge coverages as width values to, in turn, calculate edge-based pixel coverages. In at least one embodiment, prefix sum pixel coverage module **1502** performs operations to generate a projection of a scalar x-component of polygon edge portion **1506** onto a bottom pixel edge, which is used as a bottom pixel edge coverage. In at least one embodiment, prefix sum pixel coverage module **1502** performs operations to multiply together a bottom pixel edge coverage (0.36) and a pixel height (1.00) to generate an edge-based pixel coverage of 0.36. In at least one embodiment, prefix sum pixel coverage module **1502** performs operations to calculate a negative edge-based pixel coverage by using operations identical to those used to calculate an edge-based pixel coverage area, except for multiplying a -1 to a value such as a negative bottom pixel edge coverage or a negative edge-based pixel coverage. In at least one embodiment, pixel point grid **1504b** depicts bottom pixel edge coverages as solid black lines and a negative bottom pixel edge coverage as a dashed line.

[0133] FIG. 16 illustrates a block diagram of system **1600** that includes a prefix sum pixel coverage module used, at least in part, to perform steps of an algorithm used to identify one or more pixels within a polygon by using an amount of one or more edges of one or more pixels covered by one or more edges of that polygon, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1602** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, FIG. 16 illustrates a negative bottom pixel edge coverage of -0.2 and a respective negative edge-based pixel coverage. In at least one embodiment, pixel point grid **1604b** depicts partial coverages and edge-based pixel coverages calculated as a result of two polygon edge portions **1606a** and **1606b**.

[0134] FIG. 17 illustrates a block diagram of system **1700** that includes a prefix sum pixel coverage module used, at least in part, to perform operations of one or more prefix sums along one or more dimensions to identify an amount of one or more pixels within a polygon, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1702** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, non-italicized numbers shown in pixel point grid **1704** represent bottom pixel edge coverages of each pixel. In at least one embodiment, italicized numbers shown in pixel point grid **1704** represent prefix sums of bottom pixel edge coverages. In at least one embodiment, a prefix sum pixel coverage module **1702** performs a prefix sum of bottom pixel edge coverages by progressing down each column of pixels, which is shown by dashed lines that point down. In at least one embodiment, a prefix sum pixel coverage module **1702** performs prefix sums on bottom edge coverages by performing prefix sums of amounts of edges of pixels covered by edges of a polygon.

[0135] In at least one embodiment, a prefix sum is an algorithm that includes one or more summing operations. In at least one embodiment, a prefix sum is referred to as a cumulative sum, inclusive scan, or scan. In at least one embodiment, a prefix sum (or prefix sums) is a sequence of numbers (values) where, at least in part, each number is a sum of numbers before it. In at least one embodiment, prefix sum pixel coverage module **1702** performs any type of prefix sums, such as inclusive and/or exclusive prefix sums.

In at least one embodiment, prefix sum pixel coverage module **1702** performs any one or more techniques of any type of prefix sum.

[0136] In at least one embodiment, prefix sum pixel coverage module **1704** stores bottom pixel edge coverages in an indexed array such that each bottom pixel edge coverage is correlated with its respective pixel. In at least one embodiment, prefix sum pixel coverage module **1704** stores prefix sum values in an indexed array such that each prefix sum value is correlated with its respective pixel. In at least one embodiment, a bottom pixel edge coverage is referred to as a raw input while a prefix sum value is referred to as a summed output. In at least one embodiment, prefix sum pixel coverage module **1702** assigns a default, initial prefix sum value of 0.0 to each pixel located on top of each column of pixels. In at least one embodiment, with each pixel located on top of each column, prefix sum pixel coverage module **1702** adds together a bottom pixel edge coverage with a default, initial prefix sum value of 0.0. In at least one embodiment, for a top-left pixel of pixel point grid **1704**, prefix sum pixel coverage module **1702** adds together a bottom pixel edge coverage (0.2) with default, initial prefix sum value (0.0) to generate another, next prefix sum value of 0.2. In at least one embodiment, for a pixel adjacent to and below a top-left pixel of pixel point grid **1704**, prefix sum pixel coverage module **1702** adds together a negative bottom pixel edge coverage (-0.2) with prefix sum value 0.2 to generate another, next prefix sum value of 0.0. In at least one embodiment, prefix sum pixel coverage module **1702** calculates a prefix sum value based on bottom pixel edge coverages with each pixel, wherein those prefix sum values indicate, at least in part, an amount of a pixel that is covered by a polygon. In at least one embodiment, prefix sum values (summed outputs) calculated by prefix sum pixel coverage module **1702** each represent a fraction of a top edge of each pixel that is covered by one or more polygon edge portions above each pixel.

[0137] FIG. 18 illustrates a block diagram of system **1800** that includes a prefix sum pixel coverage module used, at least in part, to perform operations that sum together partial coverages and summed outputs of prefix sum operations to identify an amount of one or more pixels within a polygon, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **1802** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. 1. In at least one embodiment, prefix sum pixel coverage module **1802** receives and/or otherwise obtains values of partial coverages and prefix summed edge coverages stored as one or more arrays in one or more storage locations. In at least one embodiment, summed outputs are depicted as italicized numbers and are results of prefix sums performed, at least in part, on bottom pixel edge coverages as described herein at least in conjunction with FIG. 17. In at least one embodiment, non-italicized numbers shown in pixel point grid **1804** represent partial coverages of each pixel. In at least one embodiment, a prefix sum pixel coverage module **1802** performs sums of partial coverages and previously calculated summed outputs (prefix sum values) of bottom pixel edge coverages as described in conjunction with FIG. 17. In at least one embodiment, summing together a partial coverage and a summed output with each pixel generates a total pixel coverage, as described further herein. In at least one

embodiment, prefix sum pixel coverage module **1802** performs summing operations in parallel per pixel.

[0138] FIG. **19** illustrates a block diagram of system **1900** that includes a prefix sum pixel coverage module used, at least in part, to perform operations to generate total pixel coverages used to identify an amount of one or more pixels within a polygon, according to at least one embodiment. In at least one embodiment, a total pixel coverage is referred to as a coverage fraction of a pixel. In at least one embodiment, values depicted in FIG. **19** are coverage fractions that have been rounded up and/or down according to a rounding convention. In at least one embodiment, values depicted in FIG. **19** are coverage fractions that have not been rounded up or down. In at least one embodiment, prefix sum pixel coverage module **1902** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. **1**. In at least one embodiment, prefix sum pixel coverage module **1902** performs operations to sum (add) together partial coverages and prefix summed edge coverages. In at least one embodiment, summing together partial coverages and prefix summed edge coverages includes a summation of negative areas and prefix summed negative pixel edge coverages. In at least one embodiment, numbers depicted in pixel point grid **1904** are total pixel coverages. In at least one embodiment, non-italicized numbers shown in pixel point grid **1804** represent partial coverages of each pixel. In at least one embodiment, a prefix sum pixel coverage module **1802** performs sums of partial coverages and previously calculated summed outputs (prefix sum values) of bottom pixel edge coverages as described in conjunction with FIG. **17**. In at least one embodiment, summing together a partial coverage and a summed output with each pixel generates a total pixel coverage, as described further herein. In at least one embodiment, prefix sum pixel coverage module **1802** performs summing operations in parallel per pixel.

[0139] FIG. **20** illustrates a block diagram of system **2000** that includes a prefix sum pixel coverage module used, at least in part, to perform operations to calculate total pixel coverages by combining partial coverages, negative areas, edge-based pixel coverages, and negative edge-based pixel coverages, according to at least one embodiment. In at least one embodiment, prefix sum pixel coverage module **2002** is, at least in part, prefix sum pixel coverage module **108** performed by processor(s) **104** of FIG. **1**. In at least one embodiment, pixel point grids **2004a-c** depict a column of pixels in which prefix sum pixel coverage module **2002** has calculated one or more pixel coverages, negative areas, edge-based pixel coverages, and negative edge-based pixel coverages, or some combination thereof. In at least one embodiment, pixel point grids **2004a-c** depict at least a portion of a pixel point grid and at least a portion of a polygon as depicted in FIGS. **3**, **13**, and **14**. In at least one embodiment, pixel point grids **2004a-c** depict at least a portion of a pixel point grid and at least a portion of a polygon using one or more visual representations of one or more aspects of one or more embodiments used in FIGS. **3**, **13**, and **14**.

[0140] In at least one embodiment, pixel point grid **2004a** depicts a portion of a pixel that has a positive contribution with a positive (plus) symbol. In at least one embodiment, pixel point grid **2004a** depicts partial coverages calculated by prefix sum pixel coverage module **2002**, such as partial coverage **2006**. In at least one embodiment, partial coverages are depicted with a horizontal line pattern in FIG. **20**.

In at least one embodiment, pixel point grid **2004a** depicts edge-based pixel coverages calculated by prefix sum pixel coverage module **2002**, such as edge-based pixel coverage **2008**. In at least one embodiment, edge-based pixel coverages are depicted with a slanted-line pattern in FIG. **20**. In at least one embodiment, pixel point grid **2004a** depicts a portion of a pixel that has at least two positive contributions with at least two positive symbols. In at least one embodiment, portions of a pixel that have at least two positive contributions are depicted with a crisscross line pattern in FIG. **20**. In at least one embodiment, prefix sum pixel coverage module **2002** calculates a positive coverage comprising both an edge-based pixel coverage and a partial coverage based on right-to-left polygon edge **2012b**, such as positive coverage **2010**. In at least one embodiment, prefix sum pixel coverage module **2002** calculates a positive coverage comprising two separate edge-based pixel coverages, such as positive coverage **2011**, with one such coverage based on right-to-left polygon edge **2012a** and another such coverage based on right-to-left polygon edge **2012b**.

[0141] In at least one embodiment, pixel point grid **2004b** depicts a portion of a pixel that has a negative contribution with a negative (minus) symbol. In at least one embodiment, negative contributions are depicted with a vertical line pattern in FIG. **20**. In at least one embodiment, pixel point grid **2004b** depicts negative areas calculated by prefix sum pixel coverage module **2010**, such as negative area **2014**. In at least one embodiment, pixel point grid **2004b** depicts negative edge-based pixel coverages. In at least one embodiment, pixel point grid **2004b** depicts negative edge-based pixel coverages, such as negative-based pixel coverage **2015**. In at least one embodiment, pixel point grid **2004b** depicts areas where two or more negative edge-based pixel coverages that overlap each other, such as negative edge-based pixel coverage **2016**, which is indicated with two negative (minus) symbols.

[0142] In at least one embodiment, pixel point grid **2004c** depicts total pixel coverages for each pixel in a column of a pixel grid. In at least one embodiment, prefix sum pixel coverage module **2002** calculates total pixel coverages, such as total pixel coverages **2018a** and **2018b**, by summing any partial coverages, negative areas, edge-based pixel coverages, and negative edge-based pixel coverages, that have been calculated in conjunction with a pixel. In at least one embodiment, a result of summing together positive and negative coverages of pixel point grids **2004a** and **2004b** is a single positive contribution value, or zero, that corresponds with each pixel of a pixel point grid. In at least one embodiment, each single positive contribution of pixel point grid **2004c** represents a total pixel coverage of each pixel to be used to identify whether each pixel is within a polygon. In at least one embodiment, a sum or combination of two or more positive and/or negative coverages based on edge-based pixel coverages is referred to a summed edge coverage, such as summed edge coverage **2020**.

[0143] FIG. **21** illustrates a block diagram of a process **2100** to identify one or more pixels within a polygon by determining whether those one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries, according to at least one embodiment. In at least one embodiment, any operation of process **2100** comprises one or more operations. In at least one embodiment, operations of process **2100** are performed in a different order. In at least one embodiment, at least one operation of process

2100 is omitted. In at least one embodiment, one or more aspects of one or more embodiments described herein in conjunction with FIG. 21 are combined with one or more aspects of one or more embodiments described herein, including those described at least in conjunction with FIGS. 1-20 and 22-23. In at least one embodiment, one or more processors perform one or more operations of process **2100**. In at least one embodiment, any one or more processors described herein comprise one or more circuits. In at least one embodiment, one or more processors that perform one or more operations of process **2100** are any one processor, or combination of processors, described herein, including processor(s) **104** of FIG. 1, CPU **3100** of FIG. 31, graphics processor **3340** of FIG. 33B, general purpose graphics processing unit (GPGPU) **3430** of FIG. 34B, and parallel processing unit (PPU) **4000** of FIG. 40. In at least one embodiment, processor(s) **104** perform an operation used by system **100**, such as loading/storing summed outputs of prefix summed bottom edge pixel coverages calculated with operation **2106** in arithmetic logic unit(s) (ALUs), such as ALUs **3716** and **3718** of FIG. 37. In at least one embodiment, processor(s) **104** perform at least one operation of process **2100**, such those used to calculate bottom pixel edge coverages as described with operation **2106**. In at least one embodiment, one or more processors perform one or more operations of process **2100** by using one or more APIs of prefix sum pixel coverage API(s) module **110** of FIG. 1. In at least one embodiment, one or more processors perform one or more operations of process **2100** by using one or more API functions to calculate coverages using prefix sums of coverages as described in conjunction with at least FIGS. 22 and 23.

[0144] In at least one embodiment, one or more processors begin process **2100** by performing operation **2102** to receive polygon input data. In at least one embodiment, polygon data includes data representing one or more polygons to be rasterized into a pixel representation. In at least one embodiment, polygon data is to be rasterized to be displayed on a display device and/or be printed by a printing device. In at least one embodiment, polygon data is used in designing and/or manufacturing of semiconductors that use computational lithography. In at least one embodiment, one or more operations of operation **2102** are described further herein at least in conjunction with FIG. 1.

[0145] In at least one embodiment, one or more processors continue process **2100** by performing operation **2104** to calculate partial coverages of pixels. In at least one embodiment, calculating partial coverages of pixels includes calculating negative areas. In at least one embodiment, calculating partial coverages of pixels is described further herein at least in conjunction with FIGS. 8-14.

[0146] In at least one embodiment, one or more processors continue process **2100** by performing operation **2106** to calculate bottom pixel edge coverages. In at least one embodiment, one or more processors use bottom pixel edge coverages to calculate pixel coverages of pixels that do not have a polygon edge that cross those pixels. In at least one embodiment, calculating bottom pixel edge coverages is described further herein at least in conjunction with FIGS. 15-17.

[0147] In at least one embodiment, one or more processors continue process **2100** by performing operation **2108** to perform prefix sums of bottom pixel edge coverages by column. In at least one embodiment, one or more processors

perform operation **2018** by calculating prefix sums by a dimension other than column, such as a row dimension. In at least one embodiment, one or more processors perform operation **2018** to perform prefix sums of pixel edge coverages of pixel edges other than bottom pixel edges, such as right pixel edges, left pixel edges, or top pixel edges. In at least one embodiment, performing prefix sums of bottom pixel edge coverages is described further herein at least in conjunction with FIG. 17.

[0148] In at least one embodiment, one or more processors continue process **2100** by performing operation **2010** to sum together partial coverages and prefix summed bottom pixel edge coverages. In at least one embodiment, operation **2010** includes sums of negative areas with partial coverages. In at least one embodiment, operation **2010** includes prefix sums that include negative bottom pixel edge coverages and bottom pixel edge coverages. In at least one embodiment, summing together partial coverages and prefix summed bottom pixel edge coverages is described further herein at least in conjunction with FIGS. 19 and 20.

[0149] In at least one embodiment, one or more processors continue process **2100** by performing operation **2012** to output coverage fractions of pixels. In at least one embodiment, coverage fractions of pixels are referred to as total pixel coverages. In at least one embodiment, one or more processors output a coverage fraction of each pixel of a pixel grid. In at least one embodiment, each coverage fraction of a pixel is used to identify whether a pixel is within a polygon, and therefore, should be generated, displayed, printed, and/or otherwise activated as part of a process of rasterizing polygon data. In at least one embodiment, outputting coverage fractions of pixels is described further herein at least in conjunction with FIGS. 1 and 19.

[0150] FIG. 22 illustrates a block diagram of a process **2200** using a processor to perform one or more API functions that cause one or more processors to perform operations to identify one or more pixels within one or more polygons by determining whether one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries. In at least one embodiment, one or more aspects of one or more embodiments described herein in conjunction with FIG. 22 are combined with one or more aspects of one or more embodiments described herein, including those described at least in conjunction with FIGS. 1-21 and 23. In at least one embodiment, one or more processors perform one or more operations of process **2200**. In at least one embodiment, any one or more processors described herein comprise one or more circuits. In at least one embodiment, one or more processors that perform one or more operations of process **2200** are any one processor, or combination of processors, described herein, including processor(s) **104** of FIG. 1, CPU **3100** of FIG. 31, graphics processor **3340** of FIG. 33B, general purpose graphics processing unit (GPGPU) **3430** of FIG. 34B, and parallel processing unit (PPU) **4000** of FIG. 40. In at least one embodiment, processor(s) **104** of FIG. 1 performs one or more operations of process **2200**, such as performing API functions to calculate prefix sums of bottom edge prefix coverages as described further herein. In at least one embodiment, one or more processors used to perform process **2200** perform one or more operations of system **100** to identify one or more pixels within one or more polygons as described further herein. In at least one embodiment, one or more processors used to perform process **2200** perform one

or more operations of prefix sum pixel coverage API(s) module **110** to perform API functions to calculate pixel coverages using prefix sums of bottom pixel edge coverages of operation **2204**.

[0151] In at least one embodiment, one or more processors being process **2200** by performing operation **2202** to input polygon data into an API function invoked by software or by a user. In at least one embodiment, polygon data includes data representing one or more polygons to be rasterized into a pixel representation. In at least one embodiment, polygon data is to be rasterized to be displayed on a display device and/or be printed by a printing device. In at least one embodiment, polygon data is used in designing and/or manufacturing of semiconductors that use computational lithography. In at least one embodiment, one or more operations of operation **2102** are described further herein at least in conjunction with FIG. 1.

[0152] In at least one embodiment, one or more processors continue process **2200** with operation **2204**, by causing performance of one or more API function(s) calculate pixel coverages using prefix sums of bottom pixel edge coverages. In at least one embodiment, one or more API function(s) to calculate pixel coverages using prefix sums of bottom pixel edge coverages, are one or more functions performed by one or more processors to identify partial coverages, negative areas, bottom pixel edge coverages, negative bottom pixel edge coverages, total pixel coverages, prefix sums, or some combination thereof. In at least one embodiment, one or more API function(s) to calculate pixel coverages using prefix sums of bottom pixel edge coverages causes one or more processors to allocate memory locations in one or more data storage locations to store partial coverages, negative bottom pixel edge coverages, total pixel coverages, prefix sums, or some combination thereof. In at least one embodiment, one or more API functions performed with operation **2204** perform any one or more operations used, at least in part, to identify one or more pixels within a polygon as described further herein.

[0153] In at least one embodiment, one or more processors continue process **2200**, by performing one or more operations to output coverage fractions of pixels. In at least one embodiment, one or more processors perform one or more API functions to output coverage fractions of pixels to be used by those one or more processors to identify whether a pixel is within a polygon. In at least one embodiment, one or more processors that identify one or more pixels that are within a polygon cause one or more pixels to be generated, displayed, printed, and/or otherwise, as otherwise described herein.

[0154] FIG. 23 illustrates a block diagram of a driver and/or runtime comprising one or more libraries to provide one or more application programming interfaces (APIs), according to at least one embodiment. In at least one embodiment, any one processor, or combination of processors, perform API(s) **2310**, including processor(s) **104** of FIG. 1, CPU **3100** of FIG. 31, graphics processor **3340** of FIG. 33B, general purpose graphics processing unit (GPGPU) **3430** of FIG. 34B, and parallel processing unit (PPU) **4000** of FIG. 40. In at least one embodiment, API(s) **2310** are described further herein. In at least one embodiment, an invocation of API(s) **2310** cause any one or more operations of any one or more modules described herein, such as prefix sum pixel coverage module **108**, to be performed by one or more processors. In at least one

embodiment, an invocation of API(s) **2310** causes one or more processors to perform any one or more operations described in conjunction with FIGS. 1-22. In at least one embodiment, API(s) **2310** receives as input, polygon input data, or an indication thereof, and causes prefix sum pixel coverage module **108** of FIG. 1, and therefore any prefix sum pixel coverage module described herein, to perform operations used, at least in part, to identify one or more pixels within a polygon by, at least in part, determining whether adjacent pixels cross one or more polygon edges. In at least one embodiment, one or more API(s) **2310** and/or function(s) to calculate pixel coverages using prefix sums **2312** are performed by one or more processor(s) to perform one or more operations of process **2200**, including operation **2204**.

[0155] In at least one embodiment, a software program **2302** is a software module. In at least one embodiment, a software program **2302** comprises one or more software modules. In at least one embodiment, one or more APIs **2310** are sets of software instructions that, if executed, cause one or more processors to perform one or more computational operations. In at least one embodiment, one or more APIs **2310** are distributed or otherwise provided as a part of one or more libraries **2306**, runtimes **2304**, drivers **2304**, and/or any other grouping of software and/or executable code further described herein. In at least one embodiment, one or more APIs **2310** perform one or more computational operations in response to invocation by software programs **2302**. In at least one embodiment, a software program **2302** is a collection of software code, commands, instructions, or other sequences of text to instruct a computing device to perform one or more computational operations and/or invoke one or more other sets of instructions, such as APIs **2310** or function(s) **2312**, to be executed. In at least one embodiment, functionality provided by one or more APIs **2310** include software functions, such as those usable to accelerate one or more portions of software programs **2302** using one or more parallel processing units (PPUs), such as graphics processing units (GPUs). In at least one embodiment, a software program is a compiler.

[0156] In at least one embodiment, APIs **2310** are hardware interfaces to one or more circuits to perform one or more computational operations. In at least one embodiment, one or more software APIs **2310** described herein are implemented as one or more circuits to perform one or more techniques described herein. In at least one embodiment, one or more software programs **2302** comprise instructions that, if executed, cause one or more hardware devices and/or circuits to perform one or more techniques further described herein.

[0157] In at least one embodiment, software programs **2302**, such as user-implemented software programs, utilize one or more application programming interfaces (APIs) **230** to perform various computing operations, such as memory reservation, matrix multiplication, arithmetic operations, or any computing operation performed by parallel processing units (PPUs), such as graphics processing units (GPUs), as further described herein. In at least one embodiment, one or more APIs **2310** provide a set of callable function(s) **2312**, referred to herein as APIs, API functions, and/or functions, that individually perform one or more computing operations, such as computing operations related to parallel computing. For example, in an embodiment, one or more APIs **2310** provide function(s) **2312** to cause a scheduler to schedule

instructions to be performed by processors based on latency of interconnects coupled to these processors. In at least one embodiment, API(s) **2310** provide one or more function(s) **2312** that are one or more neural networks, such as a neural network trained to improve efficiency of processor use during rasterization processes. In at least one embodiment, APIs **2310** are stored in memory **2314**. In at least one embodiment, APIs **2310** include an API function that causes one or more processor(s) to receive and/or otherwise obtain input polygon data **102** of FIG. 1. In at least one embodiment, an API function of APIs **2310** causes one or more processor(s) to receive and/or otherwise obtain vertex data from a vertex buffer, such as API function `vkCmdBindVertexBuffers()` of a Khronos® Vulkan® API library. In at least one embodiment, an API function of APIs **2310** causes one or more processor(s) to output a rasterized image of input polygon data and/or geometric primitives by, at least in part, writing data to a buffer and/or an array that indicates which geometric primitives should be represented by pixels and/or how such primitives should be represented by pixels. In at least one embodiment, an API function of APIs **2310** causes one or more processor(s) to output a rasterized image of input polygon data by using, at least in part, an API of a fragment shader module of a Khronos® Vulkan® API library.

[0158] In at least one embodiment, one or more software programs **2302** interact or otherwise communicate with one or more APIs **2310** to perform one or more computing operations using one or more PPUs, such as GPUs. In at least one embodiment, one or more computing operations using one or more PPUs comprise at least one or more groups of computing operations to be accelerated by execution at least in part by said one or more PPUs. In at least one embodiment, one or more software programs **2302** interact with one or more APIs **2310** to facilitate parallel computing using a remote or local interface.

[0159] In at least one embodiment, an interface is software instructions that, if executed, provide access to one or more function(s) **2312** provided by one or more APIs **2310**. In at least one embodiment, a software program **2302** uses a local interface when a software developer compiles one or more software programs **2302** in conjunction with one or more libraries **2306** comprising or otherwise providing access to one or more APIs **2310**. In at least one embodiment, one or more libraries **2306** are API libraries used in graphics processing (e.g., Microsoft® DirectX®, Khronos® Vulkan®, and Khronos® OpenGL®). In at least one embodiment, one or more software programs **2302** are compiled statically in conjunction with pre-compiled libraries **2306** or uncompiled source code comprising instructions to perform one or more APIs **2310**. In at least one embodiment, one or more software programs **2302** are compiled dynamically and said one or more software programs utilize a linker to link to one or more pre-compiled libraries **2306** comprising one or more APIs **2310**.

[0160] In at least one embodiment, a software program **2302** uses a remote interface when a software developer executes a software program that utilizes or otherwise communicates with a library **2306** comprising one or more APIs **2310** over a network or other remote communication medium. In at least one embodiment, one or more libraries **2306** comprising one or more APIs **2310** are to be performed by a remote computing service, such as a computing resource services provider. In another embodiment, one or

more libraries **2306** comprising one or more APIs **2310** are to be performed by any other computing host providing said one or more APIs **2310** to one or more software programs **2302**.

[0161] In at least one embodiment, a processor performing or using one or more software programs **2302** calls, uses, performs, or otherwise implements one or more APIs **2310** to allocate and otherwise manage memory to be used by said software programs **2302**. In at least one embodiment, one or more software programs **2302** utilize one or more APIs **2310** to allocate and otherwise manage memory to be used by one or more portions of said software programs **2302** to be accelerated using one or more PPUs, such as GPUs or any other accelerator or processor further described herein. Those software programs **2302** may be performed by one or more processors based, at least in part, on latency of interconnects coupled to one or more processors using function(s) **2312** provided, in an embodiment, by one or more APIs **2310**.

[0162] In at least one embodiment, an API **2310** is an API to facilitate parallel computing. In at least one embodiment, an API **2310** is any other API further described herein. In at least one embodiment, an API **2310** is provided by a driver and/or runtime **2304**. In at least one embodiment, an API **2310** is provided by a CUDA user-mode driver. In at least one embodiment, an API **2310** is provided by a CUDA runtime. In at least one embodiment, a driver **2304** is data values and software instructions that, if executed, perform or otherwise facilitate operation of one or more function(s) **2312** of an API **2310** during load and execution of one or more portions of a software program **2302**. In at least one embodiment, a runtime **2304** is data values and software instructions that, if executed, perform or otherwise facilitate operation of one or more function(s) **2312** of an API **2310** during execution of a software program **2302**. In at least one embodiment, one or more software programs **2302** utilize one or more APIs **2310** implemented or otherwise provided by a driver and/or runtime **2304** to perform combined arithmetic operations by said one or more software programs **2302** during execution by one or more PPUs, such as GPUs.

[0163] In at least one embodiment, one or more software programs **2302** utilize one or more APIs **2310** provided by a driver and/or runtime **2304** to perform combined arithmetic operations of one or more PPUs, such as GPUs. In at least one embodiment, one or more APIs **2310** provide combined arithmetic operations through a driver and/or runtime **2304**, as described above. In at least one embodiment, one or more software programs **2302** utilize one or more APIs **2310** provided by a driver and/or runtime **2304** to allocate or otherwise reserve one or more blocks of memory **2314** of one or more PPUs, such as GPUs. In at least one embodiment, one or more software programs **2302** utilize one or more APIs **2310** provided by a driver and/or runtime **2304** to allocate or otherwise reserve blocks of memory. In at least one embodiment, one or more APIs **2310** are to perform combined mathematical functions as described herein.

[0164] In at least one embodiment, to improve software programs **2302** usability and/or optimization of one or more portions of said software programs **2302** to be accelerated by one or more PPUs, such as GPUs, one or more APIs **2310** provide one or more API function(s) **2312** to perform a scheduling system usable or used by one or more computing devices as described herein. In at least one embodiment, a processor performs one or more software programs to com-

bine two or more application programming interfaces (APIs) into a single API. In at least one embodiment, a processor uses an API to cause a scheduler to select a thread selection mechanism and/or otherwise perform operations described herein. In at least one embodiment, an API invokes a scheduler to cause a resource allocation. In at least one embodiment, a processor uses an exemplary API to schedule one or more instructions to be performed by one or more processors based, at least in part, on latency of one or more interconnects coupled to these one or more processors.

[0165] In at least one embodiment, memory **2314** is system memory **2904** of computing stem **2900**. In at least one embodiment, memory **2314** is any form of hardware that stores data and is referred to as storage or data storage. In at least one embodiment, memory **2314** is a component of system **100** of FIG. **1**. In at least one embodiment, memory **2314** is a component installed on processor(s) **104** of FIG. **1**. In at least one embodiment, memory **2314** is a component installed on prefix sum pixel coverage API(s) module **110** of FIG. **1**. In at least one embodiment, memory **2314** stores data used in various operations described herein, including polygon input data **102** of FIG. **1**. In at least one embodiment, memory **2314** stores data used in various operations described herein, including partial coverages, negative areas, bottom pixel edge coverages, negative bottom pixel edge coverages, total pixel coverages, and as otherwise described herein at least in conjunction with FIGS. **1**, and **7-22**.

[0166] In at least one embodiment, memory **2314** is a computer readable storage medium and/or code stored on said computer readable storage medium in a form of a computer program including a plurality of computer readable instructions executable by one or more processors. In at least one embodiment, a computer readable storage medium is a non-transitory computer readable medium. In at least one embodiment, at least some computer readable instructions usable to perform operations described in relation to FIG. **1** are not stored solely using transitory signals (e.g., a propagating transient electric or electromagnetic transmission). In at least one embodiment, a non-transitory computer readable medium does not necessarily include non-transitory data storage circuitry (e.g., buffers, caches, and queues) within transceivers of transitory signals. In at least one embodiment, memory **2314** is implemented as a non-transitory computer readable storage medium storing executable instructions that, if executed by one or more processors of a computer system, cause said computer system to infer computer system architecture designs used to perform one or more operations described at least in conjunction with FIGS. **1-22**.

Data Center

[0167] FIG. **24** illustrates an exemplary data center **2400**, in accordance with at least one embodiment. In at least one embodiment, data center **2400** includes, without limitation, a data center infrastructure layer **2410**, a framework layer **2420**, a software layer **2430** and an application layer **2440**.

[0168] In at least one embodiment, as shown in FIG. **24**, data center infrastructure layer **2410** may include a resource orchestrator **2412**, grouped computing resources **2414**, and node computing resources ("node C.R.s") **2416(1)-2416(N)**, where "N" represents any whole, positive integer. In at least one embodiment, node C.R.s **2416(1)-2416(N)** may include, but are not limited to, any number of central processing units

("CPUs") or other processors (including accelerators, field programmable gate arrays ("FPGAs"), data processing units ("DPUs") in network devices, graphics processors, etc.), memory devices (e.g., dynamic read-only memory), storage devices (e.g., solid state or disk drives), network input/output ("NW I/O") devices, network switches, virtual machines ("VMs"), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s **2416(1)-2416(N)** may be a server having one or more of above-mentioned computing resources.

[0169] In at least one embodiment, at least one component shown or described with respect to FIG. **24** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, node computing resources ("node C.R.s") **716(1)-716(N)** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0170] In at least one embodiment, grouped computing resources **2414** may include separate groupings of node C.R.s housed within one or more racks (not shown), or many racks housed in data centers at various geographical locations (also not shown). Separate groupings of node C.R.s within grouped computing resources **2414** may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s including CPUs or processors may grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks may also include any number of power modules, cooling modules, and network switches, in any combination.

[0171] In at least one embodiment, resource orchestrator **2412** may configure or otherwise control one or more node C.R.s **2416(1)-2416(N)** and/or grouped computing resources **2414**. In at least one embodiment, resource orchestrator **2412** may include a software design infrastructure ("SDI") management entity for data center **2400**. In at least one embodiment, resource orchestrator **2412** may include hardware, software or some combination thereof.

[0172] In at least one embodiment, as shown in FIG. **24**, framework layer **2420** includes, without limitation, a job scheduler **2432**, a configuration manager **2434**, a resource manager **2436** and a distributed file system **2438**. In at least one embodiment, framework layer **2420** may include a framework to support software **2452** of software layer **2430** and/or one or more application(s) **2442** of application layer **2440**. In at least one embodiment, software **2452** or application(s) **2442** may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. In at least one embodiment, framework layer **2420** may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter "Spark") that may utilize distributed file system **2438** for large-scale data processing (e.g., "big data"). In at least one embodiment, job scheduler **2432** may include a Spark driver to facilitate scheduling of workloads supported by various layers of data center **2400**. In at least one embodiment, configuration manager **2434** may be capable of configuring different layers such as software layer **2430** and framework

layer **2420**, including Spark and distributed file system **2438** for supporting large-scale data processing. In at least one embodiment, resource manager **2436** may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **2438** and job scheduler **2432**. In at least one embodiment, clustered or grouped computing resources may include grouped computing resource **2414** at data center infrastructure layer **2410**. In at least one embodiment, resource manager **2436** may coordinate with resource orchestrator **2412** to manage these mapped or allocated computing resources.

[0173] In at least one embodiment, software **2452** included in software layer **2430** may include software used by at least portions of node C.R.s **2416(1)-2416(N)**, grouped computing resources **2414**, and/or distributed file system **2438** of framework layer **2420**. One or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

[0174] In at least one embodiment, application(s) **2442** included in application layer **2440** may include one or more types of applications used by at least portions of node C.R.s **2416(1)-2416(N)**, grouped computing resources **2414**, and/or distributed file system **2438** of framework layer **2420**. In at least one or more types of applications may include, without limitation, CUDA applications.

[0175] In at least one embodiment, any of configuration manager **2434**, resource manager **2436**, and resource orchestrator **2412** may implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions may relieve a data center operator of data center **2400** from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor performing portions of a data center.

Computer-Based Systems

[0176] The following figures set forth, without limitation, exemplary computer-based systems that can be used to implement at least one embodiment.

[0177] FIG. 25 illustrates a processing system **2500**, in accordance with at least one embodiment. In at least one embodiment, processing system **2500** includes one or more processors **2502** and one or more graphics processors **2508**, and may be a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **2502** or processor cores **2507**. In at least one embodiment, processing system **2500** is a processing platform incorporated within a system-on-a-chip (“SoC”) integrated circuit for use in mobile, handheld, or embedded devices. In at least one embodiment, a processors core **2507** is referred to as a computing unit or compute unit.

[0178] In at least one embodiment, at least one component shown or described with respect to FIG. 25 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, processing system **2500** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0179] In at least one embodiment, processing system **2500** can include, or be incorporated within a server-based

gaming platform, a game console, a media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, processing system **2500** is a mobile phone, smart phone, tablet computing device or mobile Internet device. In at least one embodiment, processing system **2500** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, smart eyewear device, augmented reality device, or virtual reality device. In at least one embodiment, processing system **2500** is a television or set top box device having one or more processors **2502** and a graphical interface generated by one or more graphics processors **2508**.

[0180] In at least one embodiment, one or more processors **2502** each include one or more processor cores **2507** to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, each of one or more processor cores **2507** is configured to process a specific instruction set **2509**. In at least one embodiment, instruction set **2509** may facilitate Complex Instruction Set Computing (“CISC”), Reduced Instruction Set Computing (“RISC”), or computing via a Very Long Instruction Word (“VLIW”). In at least one embodiment, processor cores **2507** may each process a different instruction set **2509**, which may include instructions to facilitate emulation of other instruction sets. In at least one embodiment, processor core **2507** may also include other processing devices, such as a digital signal processor (“DSP”).

[0181] In at least one embodiment, processor **2502** includes cache memory (“cache”) **2504**. In at least one embodiment, processor **2502** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory is shared among various components of processor **2502**. In at least one embodiment, processor **2502** also uses an external cache (e.g., a Level 3 (“L3”) cache or Last Level Cache (“LLC”)) (not shown), which may be shared among processor cores **2507** using known cache coherency techniques. In at least one embodiment, register file **2506** is additionally included in processor **2502** which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **2506** may include general-purpose registers or other registers.

[0182] In at least one embodiment, one or more processor(s) **2502** are coupled with one or more interface bus(es) **2510** to transmit communication signals such as address, data, or control signals between processor **2502** and other components in processing system **2500**. In at least one embodiment interface bus **2510**, in one embodiment, can be a processor bus, such as a version of a Direct Media Interface (“DMI”) bus. In at least one embodiment, interface bus **2510** is not limited to a DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., “PCI,” PCI Express (“PCIe”)), memory buses, or other types of interface buses. In at least one embodiment processor(s) **2502** include an integrated memory controller **2516** and a platform controller hub **2530**. In at least one embodiment, memory controller **2516** facilitates communication between a memory device and other components of processing system **2500**, while platform controller hub (“PCH”) **2530** provides connections to Input/Output (“I/O”) devices via a local I/O bus. In at least one embodiment, one or more

Peripheral Component Interconnect buses include PCIe Gen 5, which provides an interface for processors.

[0183] In at least one embodiment, memory device 2520 can be a dynamic random access memory (“DRAM”) device, a static random access memory (“SRAM”) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as processor memory. In at least one embodiment memory device 2520 can operate as system memory for processing system 2500, to store data 2522 and instructions 2521 for use when one or more processors 2502 executes an application or process. In at least one embodiment, memory controller 2516 also couples with an optional external graphics processor 2512, which may communicate with one or more graphics processors 2508 in processors 2502 to perform graphics and media operations. In at least one embodiment, a display device 2511 can connect to processor(s) 2502. In at least one embodiment display device 2511 can include one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device 2511 can include a head mounted display (“HMD”) such as a stereoscopic display device for use in virtual reality (“VR”) applications or augmented reality (“AR”) applications.

[0184] In at least one embodiment, platform controller hub 2530 enables peripherals to connect to memory device 2520 and processor 2502 via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, an audio controller 2546, a network controller 2534, a firmware interface 2528, a wireless transceiver 2526, touch sensors 2525, a data storage device 2524 (e.g., hard disk drive, flash memory, etc.). In at least one embodiment, data storage device 2524 can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as PCI, or PCIe. In at least one embodiment, touch sensors 2525 can include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver 2526 can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (“LTE”) transceiver. In at least one embodiment, firmware interface 2528 enables communication with system firmware, and can be, for example, a unified extensible firmware interface (“UEFI”). In at least one embodiment, network controller 2534 can enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not shown) couples with interface bus 2510. In at least one embodiment, audio controller 2546 is a multi-channel high definition audio controller. In at least one embodiment, processing system 2500 includes an optional legacy I/O controller 2540 for coupling legacy (e.g., Personal System 2 (“PS/2”)) devices to processing system 2500. In at least one embodiment, platform controller hub 2530 can also connect to one or more Universal Serial Bus (“USB”) controllers 2542 connect input devices, such as keyboard and mouse 2543 combinations, a camera 2544, or other USB input devices.

[0185] In at least one embodiment, an instance of memory controller 2516 and platform controller hub 2530 may be integrated into a discreet external graphics processor, such as external graphics processor 2512. In at least one embodiment, platform controller hub 2530 and/or memory controller 2516 may be external to one or more processor(s) 2502.

For example, in at least one embodiment, processing system 2500 can include an external memory controller 2516 and platform controller hub 2530, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with processor(s) 2502.

[0186] FIG. 26 illustrates a computer system 2600, in accordance with at least one embodiment. In at least one embodiment, computer system 2600 may be a system with interconnected devices and components, an SOC, or some combination. In at least one embodiment, computer system 2600 is formed with a processor 2602 that may include execution units to execute an instruction. In at least one embodiment, computer system 2600 may include, without limitation, a component, such as processor 2602 to employ execution units including logic to perform algorithms for processing data. In at least one embodiment, computer system 2600 may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system 2600 may execute a version of WINDOWS® operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux for example), embedded software, and/or graphical user interfaces, may also be used.

[0187] In at least one embodiment, at least one component shown or described with respect to FIG. 26 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, computer system 2600 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0188] In at least one embodiment, computer system 2600 may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants (“PDAs”), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (DSP), an SoC, network computers (“NetPCs”), set-top boxes, network hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions.

[0189] In at least one embodiment, computer system 2600 may include, without limitation, processor 2602 that may include, without limitation, one or more execution units 2608 that may be configured to execute a Compute Unified Device Architecture (“CUDA”) (CUDA® is developed by NVIDIA Corporation of Santa Clara, CA) program. In at least one embodiment, a CUDA program is at least a portion of a software application written in a CUDA programming language. In at least one embodiment, computer system 2600 is a single processor desktop or server system. In at least one embodiment, computer system 2600 may be a multiprocessor system. In at least one embodiment, processor 2602 may include, without limitation, a CISC microprocessor, a RISC microprocessor, a VLIW microprocessor, a processor implementing a combination of instruction sets,

or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **2602** may be coupled to a processor bus **2610** that may transmit data signals between processor **2602** and other components in computer system **2600**.

[0190] In at least one embodiment, processor **2602** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **2604**. In at least one embodiment, processor **2602** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **2602**. In at least one embodiment, processor **2602** may also include a combination of both internal and external caches. In at least one embodiment, a register file **2606** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and instruction pointer register.

[0191] In at least one embodiment, execution unit **2608**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **2602**. Processor **2602** may also include a microcode (“ucode”) read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **2608** may include logic to handle a packed instruction set **2609**. In at least one embodiment, by including packed instruction set **2609** in an instruction set of a general-purpose processor **2602**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **2602**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using full width of a processor’s data bus for performing operations on packed data, which may eliminate a need to transfer smaller units of data across a processor’s data bus to perform one or more operations one data element at a time.

[0192] In at least one embodiment, execution unit **2608** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **2600** may include, without limitation, a memory **2620**. In at least one embodiment, memory **2620** may be implemented as a DRAM device, an SRAM device, flash memory device, or other memory device. Memory **2620** may store instruction (s) **2619** and/or data **2621** represented by data signals that may be executed by processor **2602**.

[0193] In at least one embodiment, a system logic chip may be coupled to processor bus **2610** and memory **2620**. In at least one embodiment, the system logic chip may include, without limitation, a memory controller hub (“MCH”) **2616**, and processor **2602** may communicate with MCH **2616** via processor bus **2610**. In at least one embodiment, MCH **2616** may provide a high bandwidth memory path **2618** to memory **2620** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **2616** may direct data signals between processor **2602**, memory **2620**, and other components in computer system **2600** and to bridge data signals between processor bus **2610**, memory **2620**, and a system I/O **2622**. In at least one embodiment, system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **2616** may be coupled to memory **2620** through high bandwidth memory path **2618**

and graphics/video card **2612** may be coupled to MCH **2616** through an Accelerated Graphics Port (“AGP”) interconnect **2614**.

[0194] In at least one embodiment, computer system **2600** may use system I/O **2622** that is a proprietary hub interface bus to couple MCH **2616** to I/O controller hub (“ICH”) **2630**. In at least one embodiment, ICH **2630** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **2620**, a chipset, and processor **2602**. Examples may include, without limitation, an audio controller **2629**, a firmware hub (“flash BIOS”) **2628**, a wireless transceiver **2626**, a data storage **2624**, a legacy I/O controller **2623** containing a user input interface **2625** and a keyboard interface, a serial expansion port **2627**, such as a USB, and a network controller **2634**. Data storage **2624** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

[0195] In at least one embodiment, FIG. 26 illustrates a system, which includes interconnected hardware devices or “chips.” In at least one embodiment, FIG. 26 may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. 26 may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe), or some combination thereof. In at least one embodiment, one or more components of system **2600** are interconnected using compute express link (“CXL”) interconnects.

[0196] FIG. 27 illustrates a system **2700**, in accordance with at least one embodiment. In at least one embodiment, system **2700** is an electronic device that utilizes a processor **2710**. In at least one embodiment, system **2700** may be, for example and without limitation, a notebook, a tower server, a rack server, a blade server, an edge device communicatively coupled to one or more on-premise or cloud service providers, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, or any other suitable electronic device.

[0197] In at least one embodiment, at least one component shown or described with respect to FIG. 27 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, system **2700** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0198] In at least one embodiment, system **2700** may include, without limitation, processor **2710** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **2710** is coupled using a bus or interface, such as an I2C bus, a System Management Bus (“SMBus”), a Low Pin Count (“LPC”) bus, a Serial Peripheral Interface (“SPI”), a High Definition Audio (“HDA”) bus, a Serial Advance Technology Attachment (“SATA”) bus, a USB (versions 1, 2, 3), or a Universal Asynchronous Receiver/Transmitter (“UART”) bus. In at least one embodiment, FIG. 27 illustrates a system which includes interconnected hardware devices or “chips.” In at least one embodiment, FIG. 27 may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. 27 may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at

least one embodiment, one or more components of FIG. 27 are interconnected using CXL interconnects.

[0199] In at least one embodiment, FIG. 27 may include a display 2724, a touch screen 2725, a touch pad 2730, a Near Field Communications unit (“NFC”) 2745, a sensor hub 2740, a thermal sensor 2746, an Express Chipset (“EC”) 2735, a Trusted Platform Module (“TPM”) 2738, BIOS/firmware/flash memory (“BIOS, FW Flash”) 2722, a DSP 2760, a Solid State Disk (“SSD”) or Hard Disk Drive (“HDD”) 2720, a wireless local area network unit (“WLAN”) 2750, a Bluetooth unit 2752, a Wireless Wide Area Network unit (“WWAN”) 2756, a Global Positioning System (“GPS”) 2755, a camera (“USB 3.0 camera”) 2754 such as a USB 3.0 camera, or a Low Power Double Data Rate (“LPDDR”) memory unit (“LPDDR3”) 2715 implemented in, for example, LPDDR3 standard. These components may each be implemented in any suitable manner.

[0200] In at least one embodiment, other components may be communicatively coupled to processor 2710 through components discussed above. In at least one embodiment, an accelerometer 2741, an Ambient Light Sensor (“ALS”) 2742, a compass 2743, and a gyroscope 2744 may be communicatively coupled to sensor hub 2740. In at least one embodiment, a thermal sensor 2739, a fan 2737, a keyboard 2736, and a touch pad 2730 may be communicatively coupled to EC 2735. In at least one embodiment, a speaker 2763, a headphones 2764, and a microphone (“mic”) 2765 may be communicatively coupled to an audio unit (“audio codec and class d amp”) 2762, which may in turn be communicatively coupled to DSP 2760. In at least one embodiment, audio unit 2762 may include, for example and without limitation, an audio coder/decoder (“codec”) and a class D amplifier. In at least one embodiment, a SIM card (“SIM”) 2757 may be communicatively coupled to WWAN unit 2756. In at least one embodiment, components such as WLAN unit 2750 and Bluetooth unit 2752, as well as WWAN unit 2756 may be implemented in a Next Generation Form Factor (“NGFF”).

[0201] FIG. 28 illustrates an exemplary integrated circuit 2800, in accordance with at least one embodiment. In at least one embodiment, exemplary integrated circuit 2800 is an SoC that may be fabricated using one or more IP cores. In at least one embodiment, integrated circuit 2800 includes one or more application processor(s) 2805 (e.g., CPUs, DPUs), at least one graphics processor 2810, and may additionally include an image processor 2815 and/or a video processor 2820, any of which may be a modular IP core. In at least one embodiment, integrated circuit 2800 includes peripheral or bus logic including a USB controller 2825, a UART controller 2830, an SPI/SDIO controller 2835, and an I²S/I²C controller 2840. In at least one embodiment, integrated circuit 2800 can include a display device 2845 coupled to one or more of a high-definition multimedia interface (“HDMI”) controller 2850 and a mobile industry processor interface (“MIPI”) display interface 2855. In at least one embodiment, storage may be provided by a flash memory subsystem 2860 including flash memory and a flash memory controller. In at least one embodiment, a memory interface may be provided via a memory controller 2865 for access to SDRAM or SRAM memory devices. In at least one embodiment, some integrated circuits additionally include an embedded security engine 2870.

[0202] In at least one embodiment, at least one component shown or described with respect to FIG. 28 is used to

implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, IC 2800 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0203] FIG. 29 illustrates a computing system 2900, according to at least one embodiment; In at least one embodiment, computing system 2900 includes a processing subsystem 2901 having one or more processor(s) 2902 and a system memory 2904 communicating via an interconnection path that may include a memory hub 2905. In at least one embodiment, memory hub 2905 may be a separate component within a chipset component or may be integrated within one or more processor(s) 2902. In at least one embodiment, memory hub 2905 couples with an I/O subsystem 2911 via a communication link 2906. In at least one embodiment, I/O subsystem 2911 includes an I/O hub 2907 that can enable computing system 2900 to receive input from one or more input device(s) 2908. In at least one embodiment, I/O hub 2907 can enable a display controller, which may be included in one or more processor(s) 2902, to provide outputs to one or more display device(s) 2910A. In at least one embodiment, one or more display device(s) 2910A coupled with I/O hub 2907 can include a local, internal, or embedded display device.

[0204] In at least one embodiment, at least one component shown or described with respect to FIG. 29 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, computing system 2900 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0205] In at least one embodiment, processing subsystem 2901 includes one or more parallel processor(s) 2912 coupled to memory hub 2905 via a bus or other communication link 2913. In at least one embodiment, communication link 2913 may be one of any number of standards based communication link technologies or protocols, such as, but not limited to PCIe, or may be a vendor specific communications interface or communications fabric. In at least one embodiment, one or more parallel processor(s) 2912 form a computationally focused parallel or vector processing system that can include a large number of processing cores and/or processing clusters, such as a many integrated core processor or compute units. In at least one embodiment, one or more parallel processor(s) 2912 form a graphics processing subsystem that can output pixels to one of one or more display device(s) 2910A coupled via I/O Hub 2907. In at least one embodiment, one or more parallel processor(s) 2912 can also include a display controller and display interface (not shown) to enable a direct connection to one or more display device(s) 2910B.

[0206] In at least one embodiment, a system storage unit 2914 can connect to I/O hub 2907 to provide a storage mechanism for computing system 2900. In at least one embodiment, an I/O switch 2916 can be used to provide an interface mechanism to enable connections between I/O hub 2907 and other components, such as a network adapter 2918 and/or wireless network adapter 2919 that may be integrated into a platform, and various other devices that can be added

via one or more add-in device(s) **2920**. In at least one embodiment, network adapter **2918** can be an Ethernet adapter or another wired network adapter. In at least one embodiment, wireless network adapter **2919** can include one or more of a Wi-Fi, Bluetooth, NFC, or other network device that includes one or more wireless radios.

[0207] In at least one embodiment, computing system **2900** can include other components not explicitly shown, including USB or other port connections, optical storage drives, video capture devices, and the like, that may also be connected to I/O hub **2907**. In at least one embodiment, communication paths interconnecting various components in FIG. **29** may be implemented using any suitable protocols, such as PCI based protocols (e.g., PCIe), or other bus or point-to-point communication interfaces and/or protocol(s), such as NVLink high-speed interconnect, or interconnect protocols.

[0208] In at least one embodiment, one or more parallel processor(s) **2912** incorporate circuitry optimized for graphics and video processing, including, for example, video output circuitry, and constitutes a graphics processing unit (“GPU”). In at least one embodiment, one or more parallel processor(s) **2912** incorporate circuitry optimized for general purpose processing. In at least one embodiment, components of computing system **2900** may be integrated with one or more other system elements on a single integrated circuit. For example, in at least one embodiment, one or more parallel processor(s) **2912**, memory hub **2905**, processor(s) **2902**, and I/O hub **2907** can be integrated into an SoC integrated circuit. In at least one embodiment, components of computing system **2900** can be integrated into a single package to form a system in package (“SIP”) configuration. In at least one embodiment, at least a portion of the components of computing system **2900** can be integrated into a multi-chip module (“MCM”), which can be interconnected with other multi-chip modules into a modular computing system. In at least one embodiment, I/O subsystem **2911** and display devices **2910B** are omitted from computing system **2900**. In at least one embodiment, one or more parallel processor(s) **2912** include one or more tensor memory accelerators (TMA) units that can transfer blocks of data between global memory and shared memory. In at least one embodiment, one or more processors uses or access one or more TMAs to perform bi-directional copy operations, e.g., from global to shared memory and vice versa.

Processing Systems

[0209] The following figures set forth, without limitation, exemplary processing systems that can be used to implement at least one embodiment.

[0210] FIG. **30** illustrates an accelerated processing unit (“APU”) **3000**, in accordance with at least one embodiment. In at least one embodiment, APU **3000** is developed by AMD

[0211] Corporation of Santa Clara, CA. In at least one embodiment, APU **3000** can be configured to execute an application program, such as a CUDA program. In at least one embodiment, APU **3000** includes, without limitation, a core complex **3010**, a graphics complex **3040**, fabric **3060**, I/O interfaces **3070**, memory controllers **3080**, a display controller **3092**, and a multimedia engine **3094**. In at least one embodiment, APU **3000** may include, without limitation, any number of core complexes **3010**, any number of graphics complexes **3050**, any number of display controllers

3092, and any number of multimedia engines **3094** in any combination. For explanatory purposes, multiple instances of like objects are denoted herein with reference numbers identifying the object and parenthetical numbers identifying the instance where needed.

[0212] In at least one embodiment, at least one component shown or described with respect to FIG. **30** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, APU **3000** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0213] In at least one embodiment, core complex **3010** is a CPU, graphics complex **3040** is a GPU, and APU **3000** is a processing unit that integrates, without limitation, **3010** and **3040** onto a single chip. In at least one embodiment, some tasks may be assigned to core complex **3010** and other tasks may be assigned to graphics complex **3040**. In at least one embodiment, core complex **3010** is configured to execute main control software associated with APU **3000**, such as an operating system. In at least one embodiment, core complex **3010** is the master processor of APU **3000**, controlling and coordinating operations of other processors. In at least one embodiment, core complex **3010** issues commands that control the operation of graphics complex **3040**. In at least one embodiment, core complex **3010** can be configured to execute host executable code derived from CUDA source code, and graphics complex **3040** can be configured to execute device executable code derived from CUDA source code.

[0214] In at least one embodiment, core complex **3010** includes, without limitation, cores **3020(1)-3020(4)** and an L3 cache **3030**. In at least one embodiment, core complex **3010** may include, without limitation, any number of cores **3020** and any number and type of caches in any combination. In at least one embodiment, cores **3020** are configured to execute instructions of a particular instruction set architecture (“ISA”). In at least one embodiment, each core **3020** is a CPU core. In at least one embodiment, core **3020** is referred to as a computing unit or compute unit.

[0215] In at least one embodiment, each core **3020** includes, without limitation, a fetch/decode unit **3022**, an integer execution engine **3024**, a floating point execution engine **3026**, and an L2 cache **3028**. In at least one embodiment, fetch/decode unit **3022** fetches instructions, decodes such instructions, generates micro-operations, and dispatches separate micro-instructions to integer execution engine **3024** and floating point execution engine **3026**. In at least one embodiment, fetch/decode unit **3022** can concurrently dispatch one micro-instruction to integer execution engine **3024** and another micro-instruction to floating point execution engine **3026**. In at least one embodiment, integer execution engine **3024** executes, without limitation, integer and memory operations. In at least one embodiment, floating point engine **3026** executes, without limitation, floating point and vector operations. In at least one embodiment, fetch-decode unit **3022** dispatches micro-instructions to a single execution engine that replaces both integer execution engine **3024** and floating point execution engine **3026**.

[0216] In at least one embodiment, each core **3020(i)**, where *i* is an integer representing a particular instance of core **3020**, may access L2 cache **3028(i)** included in core

3020(i). In at least one embodiment, each core **3020** included in core complex **3010(j)**, where *j* is an integer representing a particular instance of core complex **3010**, is connected to other cores **3020** included in core complex **3010(j)** via L3 cache **3030(j)** included in core complex **3010(j)**. In at least one embodiment, cores **3020** included in core complex **3010(j)**, where *j* is an integer representing a particular instance of core complex **3010**, can access all of L3 cache **3030(j)** included in core complex **3010(j)**. In at least one embodiment, L3 cache **3030** may include, without limitation, any number of slices.

[0217] In at least one embodiment, graphics complex **3040** can be configured to perform compute operations in a highly-parallel fashion. In at least one embodiment, graphics complex **3040** is configured to execute graphics pipeline operations such as draw commands, pixel operations, geometric computations, and other operations associated with rendering an image to a display. In at least one embodiment, graphics complex **3040** is configured to execute operations unrelated to graphics. In at least one embodiment, graphics complex **3040** is configured to execute both operations related to graphics and operations unrelated to graphics.

[0218] In at least one embodiment, graphics complex **3040** includes, without limitation, any number of compute units **3050** and an L2 cache **3042**. In at least one embodiment, compute units **3050** share L2 cache **3042**. In at least one embodiment, L2 cache **3042** is partitioned. In at least one embodiment, graphics complex **3040** includes, without limitation, any number of compute units **3050** and any number (including zero) and type of caches. In at least one embodiment, graphics complex **3040** includes, without limitation, any amount of dedicated graphics hardware.

[0219] In at least one embodiment, each compute unit **3050** includes, without limitation, any number of SIMD units **3052** and a shared memory **3054**. In at least one embodiment, each SIMD unit **3052** implements a SIMD architecture and is configured to perform operations in parallel. In at least one embodiment, each compute unit **3050** may execute any number of thread blocks, but each thread block executes on a single compute unit **3050**. In at least one embodiment, a thread block includes, without limitation, any number of threads of execution. In at least one embodiment, a workgroup is a thread block. In at least one embodiment, each SIMD unit **3052** executes a different warp. In at least one embodiment, a warp is a group of threads (e.g., 16 threads), where each thread in the warp belongs to a single thread block and is configured to process a different set of data based on a single set of instructions. In at least one embodiment, predication can be used to disable one or more threads in a warp. In at least one embodiment, a lane is a thread. In at least one embodiment, a work item is a thread. In at least one embodiment, a wavefront is a warp. In at least one embodiment, different wavefronts in a thread block may synchronize together and communicate via shared memory **3054**. In at least one embodiment, each compute unit **3050** includes one or more thread block clusters, where a thread block cluster can enable programmatic control of locality at a granularity larger than a single thread block of a single streaming multiprocessor (SM). In at least one embodiment, thread block clusters (also referred to as “clusters”) enables multiple thread blocks running concurrently across streaming multiprocessors to synchronize and collaboratively fetch, exchange, or otherwise use data.

[0220] In at least one embodiment, fabric **3060** is a system interconnect that facilitates data and control transmissions across core complex **3010**, graphics complex **3040**, I/O interfaces **3070**, memory controllers **3080**, display controller **3092**, and multimedia engine **3094**. In at least one embodiment, APU **3000** may include, without limitation, any amount and type of system interconnect in addition to or instead of fabric **3060** that facilitates data and control transmissions across any number and type of directly or indirectly linked components that may be internal or external to APU **3000**. In at least one embodiment, I/O interfaces **3070** are representative of any number and type of I/O interfaces (e.g., PCI, PCI-Extended (“PCI-X”), PCIe, gigabit Ethernet (“GBE”), USB, etc.). In at least one embodiment, various types of peripheral devices are coupled to I/O interfaces **3070**. In at least one embodiment, peripheral devices that are coupled to I/O interfaces **3070** may include, without limitation, keyboards, mice, printers, scanners, joysticks or other types of game controllers, media recording devices, external storage devices, network interface cards, and so forth.

[0221] In at least one embodiment, display controller **AMD92** displays images on one or more display device(s), such as a liquid crystal display (“LCD”) device. In at least one embodiment, multimedia engine **3094** includes, without limitation, any amount and type of circuitry that is related to multimedia, such as a video decoder, a video encoder, an image signal processor, etc. In at least one embodiment, memory controllers **3080** facilitate data transfers between APU **3000** and a unified system memory **3090**. In at least one embodiment, core complex **3010** and graphics complex **3040** share unified system memory **3090**.

[0222] In at least one embodiment, APU **3000** implements a memory subsystem that includes, without limitation, any amount and type of memory controllers **3080** and memory devices (e.g., shared memory **3054**) that may be dedicated to one component or shared among multiple components. In at least one embodiment, APU **3000** implements a cache subsystem that includes, without limitation, one or more cache memories (e.g., L2 caches **3128**, L3 cache **3030**, and L2 cache **3042**) that may each be private to or shared between any number of components (e.g., cores **3020**, core complex **3010**, SIMD units **3052**, compute units **3050**, and graphics complex **3040**).

[0223] FIG. 31 illustrates a CPU **3100**, in accordance with at least one embodiment. In at least one embodiment, CPU **3100** is developed by AMD Corporation of Santa Clara, CA. In at least one embodiment, CPU **3100** can be configured to execute an application program. In at least one embodiment, CPU **3100** is configured to execute main control software, such as an operating system. In at least one embodiment, CPU **3100** issues commands that control the operation of an external GPU (not shown). In at least one embodiment, CPU **3100** can be configured to execute host executable code derived from CUDA source code, and an external GPU can be configured to execute device executable code derived from such CUDA source code. In at least one embodiment, CPU **3100** includes, without limitation, any number of core complexes **3110**, fabric **3160**, I/O interfaces **3170**, and memory controllers **3180**.

[0224] In at least one embodiment, at least one component shown or described with respect to FIG. 31 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, CPU **3100**

performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0225] In at least one embodiment, core complex 3110 includes, without limitation, cores 3120(1)-3120(4) and an L3 cache 3130. In at least one embodiment, core complex 3110 may include, without limitation, any number of cores 3120 and any number and type of caches in any combination. In at least one embodiment, cores 3120 are configured to execute instructions of a particular ISA. In at least one embodiment, each core 3120 is a CPU core.

[0226] In at least one embodiment, each core 3120 includes, without limitation, a fetch/decode unit 3122, an integer execution engine 3124, a floating point execution engine 3126, and an L2 cache 3128. In at least one embodiment, fetch/decode unit 3122 fetches instructions, decodes such instructions, generates micro-operations, and dispatches separate micro-instructions to integer execution engine 3124 and floating point execution engine 3126. In at least one embodiment, fetch/decode unit 3122 can concurrently dispatch one micro-instruction to integer execution engine 3124 and another micro-instruction to floating point execution engine 3126. In at least one embodiment, integer execution engine 3124 executes, without limitation, integer and memory operations. In at least one embodiment, floating point engine 3126 executes, without limitation, floating point and vector operations. In at least one embodiment, fetch-decode unit 3122 dispatches micro-instructions to a single execution engine that replaces both integer execution engine 3124 and floating point execution engine 3126.

[0227] In at least one embodiment, each core 3120(*i*), where *i* is an integer representing a particular instance of core 3120, may access L2 cache 3128(*i*) included in core 3120(*i*). In at least one embodiment, each core 3120 included in core complex 3110(*j*), where *j* is an integer representing a particular instance of core complex 3110, is connected to other cores 3120 in core complex 3110(*j*) via L3 cache 3130(*j*) included in core complex 3110(*j*). In at least one embodiment, cores 3120 included in core complex 3110(*j*), where *j* is an integer representing a particular instance of core complex 3110, can access all of L3 cache 3130(*j*) included in core complex 3110(*j*). In at least one embodiment, L3 cache 3130 may include, without limitation, any number of slices.

[0228] In at least one embodiment, fabric 3160 is a system interconnect that facilitates data and control transmissions across core complexes 3110(1)-3110(*N*) (where *N* is an integer greater than zero), I/O interfaces 3170, and memory controllers 3180. In at least one embodiment, CPU 3100 may include, without limitation, any amount and type of system interconnect in addition to or instead of fabric 3160 that facilitates data and control transmissions across any number and type of directly or indirectly linked components that may be internal or external to CPU 3100. In at least one embodiment, I/O interfaces 3170 are representative of any number and type of I/O interfaces (e.g., PCI, PCI-X, PCIe, GBE, USB, etc.). In at least one embodiment, various types of peripheral devices are coupled to I/O interfaces 3170. In at least one embodiment, peripheral devices that are coupled to I/O interfaces 3170 may include, without limitation, displays, keyboards, mice, printers, scanners, joysticks or

other types of game controllers, media recording devices, external storage devices, network interface cards, and so forth.

[0229] In at least one embodiment, memory controllers 3180 facilitate data transfers between CPU 3100 and a system memory 3190. In at least one embodiment, core complex 3110 and graphics complex 3140 share system memory 3190. In at least one embodiment, CPU 3100 implements a memory subsystem that includes, without limitation, any amount and type of memory controllers 3180 and memory devices that may be dedicated to one component or shared among multiple components. In at least one embodiment, CPU 3100 implements a cache subsystem that includes, without limitation, one or more cache memories (e.g., L2 caches 3128 and L3 caches 3130) that may each be private to or shared between any number of components (e.g., cores 3120 and core complexes 3110).

[0230] FIG. 32 illustrates an exemplary accelerator integration slice 3290, in accordance with at least one embodiment. As used herein, a “slice” comprises a specified portion of processing resources of an accelerator integration circuit. In at least one embodiment, the accelerator integration circuit provides cache management, memory access, context management, and interrupt management services on behalf of multiple graphics processing engines included in a graphics acceleration module. The graphics processing engines may each comprise a separate GPU. Alternatively, the graphics processing engines may comprise different types of graphics processing engines within a GPU such as graphics execution units, media processing engines (e.g., video encoders/decoders), samplers, and blit engines. In at least one embodiment, the graphics acceleration module may be a GPU with multiple graphics processing engines. In at least one embodiment, the graphics processing engines may be individual GPUs integrated on a common package, line card, or chip.

[0231] In at least one embodiment, at least one component shown or described with respect to FIG. 32 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, accelerated integration slice 3290 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0232] An application effective address space 3282 within system memory 3214 stores process elements 3283. In one embodiment, process elements 3283 are stored in response to GPU invocations 3281 from applications 3280 executed on processor 3207. A process element 3283 contains process state for corresponding application 3280. A work descriptor (“WD”) 3284 contained in process element 3283 can be a single job requested by an application or may contain a pointer to a queue of jobs. In at least one embodiment, WD 3284 is a pointer to a job request queue in application effective address space 3282.

[0233] Graphics acceleration module 3246 and/or individual graphics processing engines can be shared by all or a subset of processes in a system. In at least one embodiment, an infrastructure for setting up process state and sending WD 3284 to graphics acceleration module 3246 to start a job in a virtualized environment may be included.

[0234] In at least one embodiment, a dedicated-process programming model is implementation-specific. In this model, a single process owns graphics acceleration module 3246 or an individual graphics processing engine. Because graphics acceleration module 3246 is owned by a single process, a hypervisor initializes an accelerator integration circuit for an owning partition and an operating system initializes accelerator integration circuit for an owning process when graphics acceleration module 3246 is assigned.

[0235] In operation, a WD fetch unit 3291 in accelerator integration slice 3290 fetches next WD 3284 which includes an indication of work to be done by one or more graphics processing engines of graphics acceleration module 3246. Data from WD 3284 may be stored in registers 3245 and used by a memory management unit (“MMU”) 3239, interrupt management circuit 3247 and/or context management circuit 3248 as illustrated. For example, one embodiment of MMU 3239 includes segment/page walk circuitry for accessing segment/page tables 3286 within OS virtual address space 3285. Interrupt management circuit 3247 may process interrupt events (“INT”) 3292 received from graphics acceleration module 3246. When performing graphics operations, an effective address 3293 generated by a graphics processing engine is translated to a real address by MMU 3239.

[0236] In one embodiment, a same set of registers 3245 are duplicated for each graphics processing engine and/or graphics acceleration module 3246 and may be initialized by a hypervisor or operating system. Each of these duplicated registers may be included in accelerator integration slice 3290. Exemplary registers that may be initialized by a hypervisor are shown in Table 1.

TABLE 1

Hypervisor Initialized Registers	
1	Slice Control Register
2	Real Address (RA) Scheduled Processes Area Pointer
3	Authority Mask Override Register
4	Interrupt Vector Table Entry Offset
5	Interrupt Vector Table Entry Limit
6	State Register
7	Logical Partition ID
8	Real address (RA) Hypervisor Accelerator Utilization Record Pointer
9	Storage Description Register

[0237] Exemplary registers that may be initialized by an operating system are shown in Table 2.

TABLE 2

Operating System Initialized Registers	
1	Process and Thread Identification
2	Effective Address (EA) Context Save/Restore Pointer
3	Virtual Address (VA) Accelerator Utilization Record Pointer
4	Virtual Address (VA) Storage Segment Table Pointer
5	Authority Mask
6	Work descriptor

[0238] In one embodiment, each WD 3284 is specific to a particular graphics acceleration module 3246 and/or a particular graphics processing engine. It contains all information required by a graphics processing engine to do work or it can be a pointer to a memory location where an application has set up a command queue of work to be completed.

[0239] FIGS. 33A-33B illustrate exemplary graphics processors, in accordance with at least one embodiment. In at least one embodiment, any of the exemplary graphics processors may be fabricated using one or more IP cores. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores. In at least one embodiment, the exemplary graphics processors are for use within an SoC.

[0240] FIG. 33A illustrates an exemplary graphics processor 3310 of an SoC integrated circuit that may be fabricated using one or more IP cores, in accordance with at least one embodiment. FIG. 33B illustrates an additional exemplary graphics processor 3340 of an SoC integrated circuit that may be fabricated using one or more IP cores, in accordance with at least one embodiment. In at least one embodiment, graphics processor 3310 of FIG. 33A is a low power graphics processor core. In at least one embodiment, graphics processor 3340 of FIG. 33B is a higher performance graphics processor core. In at least one embodiment, each of graphics processors 3310, 3340 can be variants of graphics processor 2810 of FIG. 28.

[0241] In at least one embodiment, at least one component shown or described with respect to FIGS. 33A-33B is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, graphics processor 3340 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0242] In at least one embodiment, graphics processor 3310 includes a vertex processor 3305 and one or more fragment processor(s) 3315A-3315N (e.g., 3315A, 3315B, 3315C, 3315D, through 3315N-1, and 3315N). In at least one embodiment, graphics processor 3310 can execute different shader programs via separate logic, such that vertex processor 3305 is optimized to execute operations for vertex shader programs, while one or more fragment processor(s) 3315A-3315N execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. In at least one embodiment, vertex processor 3305 performs a vertex processing stage of a 3D graphics pipeline and generates primitives and vertex data. In at least one embodiment, fragment processor(s) 3315A-3315N use primitive and vertex data generated by vertex processor 3305 to produce a framebuffer that is displayed on a display device. In at least one embodiment, fragment processor(s) 3315A-3315N are optimized to execute fragment shader programs as provided for in an OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in a Direct 3D API.

[0243] In at least one embodiment, graphics processor 3310 additionally includes one or more MMU(s) 3320A-3320B, cache(s) 3325A-3325B, and circuit interconnect(s) 3330A-3330B. In at least one embodiment, one or more MMU(s) 3320A-3320B provide for virtual to physical address mapping for graphics processor 3310, including for vertex processor 3305 and/or fragment processor(s) 3315A-3315N, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in one or more cache(s) 3325A-3325B. In at least one

embodiment, one or more MMU(s) **3320A-3320B** may be synchronized with other MMUs within a system, including one or more MMUs associated with one or more application processor(s) **2805**, image processors **2815**, and/or video processors **2820** of FIG. **28**, such that each processor **2805-2820** can participate in a shared or unified virtual memory system. In at least one embodiment, one or more circuit interconnect(s) **3330A-3330B** enable graphics processor **3310** to interface with other IP cores within an SoC, either via an internal bus of the SoC or via a direct connection.

[0244] In at least one embodiment, graphics processor **3340** includes one or more MMU(s) **3320A-3320B**, caches **3325A-3325B**, and circuit interconnects **3330A-3330B** of graphics processor **3310** of FIG. **33A**. In at least one embodiment, graphics processor **3340** includes one or more shader core(s) **3355A-3355N** (e.g., **3355A**, **3355B**, **3355C**, **3355D**, **3355E**, **3355F**, through **3355N-1**, and **3355N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. In at least one embodiment, a number of shader cores can vary. In at least one embodiment, graphics processor **3340** includes an inter-core task manager **3345**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **3355A-3355N** and a tiling unit **3358** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

[0245] FIG. **34A** illustrates a graphics core **3400**, in accordance with at least one embodiment. In at least one embodiment, graphics core **3400** may be included within graphics processor **2810** of FIG. **28**. In at least one embodiment, graphics core **3400** may be a unified shader core **3355A-3355N** as in FIG. **33B**. In at least one embodiment, graphics core **3400** includes a shared instruction cache **3402**, a texture unit **3418**, and a cache/shared memory **3420** that are common to execution resources within graphics core **3400**. In at least one embodiment, graphics core **3400** can include multiple slices **3401A-3401N** or partition for each core, and a graphics processor can include multiple instances of graphics core **3400**. Slices **3401A-3401N** can include support logic including a local instruction cache **3404A-3404N**, a thread scheduler **3406A-3406N**, a thread dispatcher **3408A-3408N**, and a set of registers **3410A-3410N**. In at least one embodiment, slices **3401A-3401N** can include a set of additional function units (“AFUs”) **3412A-3412N**, floating-point units (“FPUs”) **3414A-3414N**, integer arithmetic logic units (“ALUs”) **3416-3416N**, address computational units (“ACUs”) **3413A-3413N**, double-precision floating-point units (“DPFPUs”) **3415A-3415N**, and matrix processing units (“MPUs”) **3417A-3417N**. In at least one embodiment, a graphics core **3400** is referred to as a compute unit or computing unit.

[0246] In at least one embodiment, FPUs **3414A-3414N** can perform single-precision (32-bit) and half-precision (16-bit) floating point operations, while DPFPUs **3415A-3415N** perform double precision (64-bit) floating point operations. In at least one embodiment, ALUs **3416A-3416N** can perform variable precision integer operations at 8-bit, 16-bit, and 32-bit precision, and can be configured for mixed precision operations. In at least one embodiment, MPUs **3417A-3417N** can also be configured for mixed

precision matrix operations, including half-precision floating point and 8-bit integer operations. In at least one embodiment, MPUs **3417-3417N** can perform a variety of matrix operations to accelerate CUDA programs, including enabling support for accelerated general matrix to matrix multiplication (“GEMM”). In at least one embodiment, AFUs **3412A-3412N** can perform additional logic operations not supported by floating-point or integer units, including trigonometric operations (e.g., Sine, Cosine, etc.).

[0247] FIG. **34B** illustrates a general-purpose graphics processing unit (“GPGPU”) **3430**, in accordance with at least one embodiment. In at least one embodiment, GPGPU **3430** is highly-parallel and suitable for deployment on a multi-chip module. In at least one embodiment, GPGPU **3430** can be configured to enable highly-parallel compute operations to be performed by an array of GPUs. In at least one embodiment, GPGPU **3430** can be linked directly to other instances of GPGPU **3430** to create a multi-GPU cluster to improve execution time for CUDA programs. In at least one embodiment, GPGPU **3430** includes a host interface **3432** to enable a connection with a host processor. In at least one embodiment, host interface **3432** is a PCIe interface. In at least one embodiment, host interface **3432** can be a vendor specific communications interface or communications fabric. In at least one embodiment, GPGPU **3430** receives commands from a host processor and uses a global scheduler **3434** to distribute execution threads associated with those commands to a set of compute clusters **3436A-3436H**. In at least one embodiment, compute clusters **3436A-3436H** share a cache memory **3438**. In at least one embodiment, cache memory **3438** can serve as a higher-level cache for cache memories within compute clusters **3436A-3436H**.

[0248] In at least one embodiment, at least one component shown or described with respect to FIGS. **34A-34B** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, GPGPU **3430** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0249] In at least one embodiment, GPGPU **3430** includes memory **3444A-3444B** coupled with compute clusters **3436A-3436H** via a set of memory controllers **3442A-3442B**. In at least one embodiment, memory **3444A-3444B** can include various types of memory devices including DRAM or graphics random access memory, such as synchronous graphics random access memory (“SGRAM”), including graphics double data rate (“GDDR”) memory.

[0250] In at least one embodiment, compute clusters **3436A-3436H** each include a set of graphics cores, such as graphics core **3400** of FIG. **34A**, which can include multiple types of integer and floating point logic units that can perform computational operations at a range of precisions including suited for computations associated with CUDA programs. For example, in at least one embodiment, at least a subset of floating point units in each of compute clusters **3436A-3436H** can be configured to perform 16-bit or 32-bit floating point operations, while a different subset of floating point units can be configured to perform 64-bit floating point operations.

[0251] In at least one embodiment, multiple instances of GPGPU **3430** can be configured to operate as a compute

cluster. Compute clusters **3436A-3436H** may implement any technically feasible communication techniques for synchronization and data exchange. In at least one embodiment, multiple instances of GPGPU **3430** communicate over host interface **3432**. In at least one embodiment, GPGPU **3430** includes an I/O hub **3439** that couples GPGPU **3430** with a GPU link **3440** that enables a direct connection to other instances of GPGPU **3430**. In at least one embodiment, GPU link **3440** is coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of GPGPU **3430**. In at least one embodiment GPU link **3440** couples with a high speed interconnect to transmit and receive data to other GPGPUs **3430** or parallel processors. In at least one embodiment, multiple instances of GPGPU **3430** are located in separate data processing systems and communicate via a network device that is accessible via host interface **3432**. In at least one embodiment GPU link **3440** can be configured to enable a connection to a host processor in addition to or as an alternative to host interface **3432**. In at least one embodiment, GPGPU **3430** can be configured to execute a CUDA program.

[0252] FIG. 35A illustrates a parallel processor **3500**, in accordance with at least one embodiment. In at least one embodiment, various components of parallel processor **3500** may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (“ASICs”), or FPGAs.

[0253] In at least one embodiment, parallel processor **3500** includes a parallel processing unit **3502**. In at least one embodiment, parallel processing unit **3502** includes an I/O unit **3504** that enables communication with other devices, including other instances of parallel processing unit **3502**. In at least one embodiment, I/O unit **3504** may be directly connected to other devices. In at least one embodiment, I/O unit **3504** connects with other devices via use of a hub or switch interface, such as memory hub **3505**. In at least one embodiment, connections between memory hub **3505** and I/O unit **3504** form a communication link. In at least one embodiment, I/O unit **3504** connects with a host interface **3506** and a memory crossbar **3516**, where host interface **3506** receives commands directed to performing processing operations and memory crossbar **3516** receives commands directed to performing memory operations.

[0254] In at least one embodiment, when host interface **3506** receives a command buffer via I/O unit **3504**, host interface **3506** can direct work operations to perform those commands to a front end **3508**. In at least one embodiment, front end **3508** couples with a scheduler **3510**, which is configured to distribute commands or other work items to a processing array **3512**. In at least one embodiment, scheduler **3510** ensures that processing array **3512** is properly configured and in a valid state before tasks are distributed to processing array **3512**. In at least one embodiment, scheduler **3510** is implemented via firmware logic executing on a microcontroller. In at least one embodiment, microcontroller implemented scheduler **3510** is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, enabling rapid preemption and context switching of threads executing on processing array **3512**. In at least one embodiment, host software can prove workloads for scheduling on processing array **3512** via one of multiple graphics processing doorbells. In at least one embodiment, workloads can then be automatically distrib-

uted across processing array **3512** by scheduler **3510** logic within a microcontroller including scheduler **3510**.

[0255] In at least one embodiment, processing array **3512** can include up to “N” clusters (e.g., cluster **3514A**, cluster **3514B**, through cluster **3514N**). In at least one embodiment, each cluster **3514A-3514N** of processing array **3512** can execute a large number of concurrent threads. In at least one embodiment, scheduler **3510** can allocate work to clusters **3514A-3514N** of processing array **3512** using various scheduling and/or work distribution algorithms, which may vary depending on the workload arising for each type of program or computation. In at least one embodiment, scheduling can be handled dynamically by scheduler **3510**, or can be assisted in part by compiler logic during compilation of program logic configured for execution by processing array **3512**. In at least one embodiment, different clusters **3514A-3514N** of processing array **3512** can be allocated for processing different types of programs or for performing different types of computations.

[0256] In at least one embodiment, processing array **3512** can be configured to perform various types of parallel processing operations. In at least one embodiment, processing array **3512** is configured to perform general-purpose parallel compute operations. For example, in at least one embodiment, processing array **3512** can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

[0257] In at least one embodiment, processing array **3512** is configured to perform parallel graphics processing operations. In at least one embodiment, processing array **3512** can include additional logic to support execution of such graphics processing operations, including, but not limited to texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. In at least one embodiment, processing array **3512** can be configured to execute graphics processing related shader programs such as, but not limited to vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. In at least one embodiment, parallel processing unit **3502** can transfer data from system memory via I/O unit **3504** for processing. In at least one embodiment, during processing, transferred data can be stored to on-chip memory (e.g., a parallel processor memory **3522**) during processing, then written back to system memory.

[0258] In at least one embodiment, when parallel processing unit **3502** is used to perform graphics processing, scheduler **3510** can be configured to divide a processing workload into approximately equal sized tasks, to better enable distribution of graphics processing operations to multiple clusters **3514A-3514N** of processing array **3512**. In at least one embodiment, portions of processing array **3512** can be configured to perform different types of processing. For example, in at least one embodiment, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. In at least one embodiment, intermediate data produced by one or more of clusters **3514A-3514N** may be stored in buffers to allow intermediate data to be transmitted between clusters **3514A-3514N** for further processing.

[0259] In at least one embodiment, processing array 3512 can receive processing tasks to be executed via scheduler 3510, which receives commands defining processing tasks from front end 3508. In at least one embodiment, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and commands defining how data is to be processed (e.g., what program is to be executed). In at least one embodiment, scheduler 3510 may be configured to fetch indices corresponding to tasks or may receive indices from front end 3508. In at least one embodiment, front end 3508 can be configured to ensure processing array 3512 is configured to a valid state before a workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

[0260] In at least one embodiment, each of one or more instances of parallel processing unit 3502 can couple with parallel processor memory 3522. In at least one embodiment, parallel processor memory 3522 can be accessed via memory crossbar 3516, which can receive memory requests from processing array 3512 as well as I/O unit 3504. In at least one embodiment, memory crossbar 3516 can access parallel processor memory 3522 via a memory interface 3518. In at least one embodiment, memory interface 3518 can include multiple partition units (e.g., a partition unit 3520A, partition unit 3520B, through partition unit 3520N) that can each couple to a portion (e.g., memory unit) of parallel processor memory 3522. In at least one embodiment, a number of partition units 3520A-3520N is configured to be equal to a number of memory units, such that a first partition unit 3520A has a corresponding first memory unit 3524A, a second partition unit 3520B has a corresponding memory unit 3524B, and an Nth partition unit 3520N has a corresponding Nth memory unit 3524N. In at least one embodiment, a number of partition units 3520A-3520N may not be equal to a number of memory devices.

[0261] In at least one embodiment, memory units 3524A-3524N can include various types of memory devices, including DRAM or graphics random access memory, such as SGRAM, including GDDR memory. In at least one embodiment, memory units 3524A-3524N may also include 3D stacked memory, including but not limited to high bandwidth memory (“HBM”). In at least one embodiment, render targets, such as frame buffers or texture maps may be stored across memory units 3524A-3524N, allowing partition units 3520A-3520N to write portions of each render target in parallel to efficiently use available bandwidth of parallel processor memory 3522. In at least one embodiment, a local instance of parallel processor memory 3522 may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

[0262] In at least one embodiment, any one of clusters 3514A-3514N of processing array 3512 can process data that will be written to any of memory units 3524A-3524N within parallel processor memory 3522. In at least one embodiment, memory crossbar 3516 can be configured to transfer an output of each cluster 3514A-3514N to any partition unit 3520A-3520N or to another cluster 3514A-3514N, which can perform additional processing operations on an output. In at least one embodiment, each cluster 3514A-3514N can communicate with memory interface 3518 through memory crossbar 3516 to read from or write to various external memory devices. In at least one embodiment, memory crossbar 3516 has a connection to memory

interface 3518 to communicate with I/O unit 3504, as well as a connection to a local instance of parallel processor memory 3522, enabling processing units within different clusters 3514A-3514N to communicate with system memory or other memory that is not local to parallel processing unit 3502. In at least one embodiment, memory crossbar 3516 can use virtual channels to separate traffic streams between clusters 3514A-3514N and partition units 3520A-3520N.

[0263] In at least one embodiment, multiple instances of parallel processing unit 3502 can be provided on a single add-in card, or multiple add-in cards can be interconnected. In at least one embodiment, different instances of parallel processing unit 3502 can be configured to inter-operate even if different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. For example, in at least one embodiment, some instances of parallel processing unit 3502 can include higher precision floating point units relative to other instances. In at least one embodiment, systems incorporating one or more instances of parallel processing unit 3502 or parallel processor 3500 can be implemented in a variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems.

[0264] FIG. 35B illustrates a processing cluster 3594, in accordance with at least one embodiment. In at least one embodiment, processing cluster 3594 is included within a parallel processing unit. In at least one embodiment, processing cluster 3594 is one of processing clusters 3514A-3514N of FIG. 35. In at least one embodiment, processing cluster 3594 can be configured to execute many threads in parallel, where the term “thread” refers to an instance of a particular program executing on a particular set of input data. In at least one embodiment, single instruction, multiple data (“SIMD”) instruction issue techniques are used to support parallel execution of a large number of threads without providing multiple independent instruction units. In at least one embodiment, single instruction, multiple thread (“SIMT”) techniques are used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of processing engines within each processing cluster 3594.

[0265] In at least one embodiment, at least one component shown or described with respect to FIG. 35A-35B is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, processing cluster 3594 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0266] In at least one embodiment, operation of processing cluster 3594 can be controlled via a pipeline manager 3532 that distributes processing tasks to SIMT parallel processors. In at least one embodiment, pipeline manager 3532 receives instructions from scheduler 3510 of FIG. 35 and manages execution of those instructions via a graphics multiprocessor 3534 and/or a texture unit 3536. In at least one embodiment, graphics multiprocessor 3534 is an exemplary instance of a SIMT parallel processor. However, in at least one embodiment, various types of SIMT parallel pro-

processors of differing architectures may be included within processing cluster 3594. In at least one embodiment, one or more instances of graphics multiprocessor 3534 can be included within processing cluster 3594. In at least one embodiment, graphics multiprocessor 3534 can process data and a data crossbar 3540 can be used to distribute processed data to one of multiple possible destinations, including other shader units. In at least one embodiment, pipeline manager 3532 can facilitate distribution of processed data by specifying destinations for processed data to be distributed via data crossbar 3540.

[0267] In at least one embodiment, each graphics multiprocessor 3534 within processing cluster 3594 can include an identical set of functional execution logic (e.g., arithmetic logic units, load/store units (“LSUs”), etc.). In at least one embodiment, functional execution logic can be configured in a pipelined manner in which new instructions can be issued before previous instructions are complete. In at least one embodiment, functional execution logic supports a variety of operations including integer and floating point arithmetic, comparison operations, Boolean operations, bit-shifting, and computation of various algebraic functions. In at least one embodiment, same functional-unit hardware can be leveraged to perform different operations and any combination of functional units may be present.

[0268] In at least one embodiment, instructions transmitted to processing cluster 3594 constitute a thread. In at least one embodiment, a set of threads executing across a set of parallel processing engines is a thread group. In at least one embodiment, a thread group executes a program on different input data. In at least one embodiment, each thread within a thread group can be assigned to a different processing engine within graphics multiprocessor 3534. In at least one embodiment, a thread group may include fewer threads than a number of processing engines within graphics multiprocessor 3534. In at least one embodiment, when a thread group includes fewer threads than a number of processing engines, one or more of the processing engines may be idle during cycles in which that thread group is being processed. In at least one embodiment, a thread group may also include more threads than a number of processing engines within graphics multiprocessor 3534. In at least one embodiment, when a thread group includes more threads than the number of processing engines within graphics multiprocessor 3534, processing can be performed over consecutive clock cycles. In at least one embodiment, multiple thread groups can be executed concurrently on graphics multiprocessor 3534.

[0269] In at least one embodiment, graphics multiprocessor 3534 includes an internal cache memory to perform load and store operations. In at least one embodiment, graphics multiprocessor 3534 can forego an internal cache and use a cache memory (e.g., L1 cache 3548) within processing cluster 3594. In at least one embodiment, each graphics multiprocessor 3534 also has access to Level 2 (“L2”) caches within partition units (e.g., partition units 3520A-3520N of FIG. 35A) that are shared among all processing clusters 3594 and may be used to transfer data between threads. In at least one embodiment, graphics multiprocessor 3534 may also access off-chip global memory, which can include one or more of local parallel processor memory and/or system memory. In at least one embodiment, any memory external to parallel processing unit 3502 may be used as global memory. In at least one embodiment, processing cluster 3594 includes multiple instances of graphics

multiprocessor 3534 that can share common instructions and data, which may be stored in L1 cache 3548.

[0270] In at least one embodiment, each processing cluster 3594 may include an MMU 3545 that is configured to map virtual addresses into physical addresses. In at least one embodiment, one or more instances of MMU 3545 may reside within memory interface 3518 of FIG. 35. In at least one embodiment, MMU 3545 includes a set of page table entries (“PTEs”) used to map a virtual address to a physical address of a tile and optionally a cache line index. In at least one embodiment, MMU 3545 may include address translation lookaside buffers (“TLBs”) or caches that may reside within graphics multiprocessor 3534 or L1 cache 3548 or processing cluster 3594. In at least one embodiment, a physical address is processed to distribute surface data access locality to allow efficient request interleaving among partition units. In at least one embodiment, a cache line index may be used to determine whether a request for a cache line is a hit or miss.

[0271] In at least one embodiment, processing cluster 3594 may be configured such that each graphics multiprocessor 3534 is coupled to a texture unit 3536 for performing texture mapping operations, e.g., determining texture sample positions, reading texture data, and filtering texture data. In at least one embodiment, texture data is read from an internal texture L1 cache (not shown) or from an L1 cache within graphics multiprocessor 3534 and is fetched from an L2 cache, local parallel processor memory, or system memory, as needed. In at least one embodiment, each graphics multiprocessor 3534 outputs a processed task to data crossbar 3540 to provide the processed task to another processing cluster 3594 for further processing or to store the processed task in an L2 cache, a local parallel processor memory, or a system memory via memory crossbar 3516. In at least one embodiment, a pre-raster operations unit (“preROP”) 3542 is configured to receive data from graphics multiprocessor 3534, direct data to ROP units, which may be located with partition units as described herein (e.g., partition units 3520A-3520N of FIG. 35). In at least one embodiment, PreROP 3542 can perform optimizations for color blending, organize pixel color data, and perform address translations.

[0272] FIG. 35C illustrates a graphics multiprocessor 3596, in accordance with at least one embodiment. In at least one embodiment, graphics multiprocessor 3596 is graphics multiprocessor 3534 of FIG. 35B. In at least one embodiment, graphics multiprocessor 3596 couples with pipeline manager 3532 of processing cluster 3594. In at least one embodiment, graphics multiprocessor 3596 has an execution pipeline including but not limited to an instruction cache 3552, an instruction unit 3554, an address mapping unit 3556, a register file 3558, one or more GPGPU cores 3562, and one or more LSUs 3566. GPGPU cores 3562 and LSUs 3566 are coupled with cache memory 3572 and shared memory 3570 via a memory and cache interconnect 3568.

[0273] In at least one embodiment, instruction cache 3552 receives a stream of instructions to execute from pipeline manager 3532. In at least one embodiment, instructions are cached in instruction cache 3552 and dispatched for execution by instruction unit 3554. In at least one embodiment, instruction unit 3554 can dispatch instructions as thread groups (e.g., warps), with each thread of a thread group assigned to a different execution unit within GPGPU core 3562. In at least one embodiment, an instruction can access

any of a local, shared, or global address space by specifying an address within a unified address space. In at least one embodiment, address mapping unit **3556** can be used to translate addresses in a unified address space into a distinct memory address that can be accessed by LSUs **3566**.

[0274] In at least one embodiment, register file **3558** provides a set of registers for functional units of graphics multiprocessor **3596**. In at least one embodiment, register file **3558** provides temporary storage for operands connected to data paths of functional units (e.g., GPGPU cores **3562**, LSUs **3566**) of graphics multiprocessor **3596**. In at least one embodiment, register file **3558** is divided between each of functional units such that each functional unit is allocated a dedicated portion of register file **3558**. In at least one embodiment, register file **3558** is divided between different thread groups being executed by graphics multiprocessor **3596**.

[0275] In at least one embodiment, GPGPU cores **3562** can each include FPUs and/or integer ALUs that are used to execute instructions of graphics multiprocessor **3596**. GPGPU cores **3562** can be similar in architecture or can differ in architecture. In at least one embodiment, a first portion of GPGPU cores **3562** include a single precision FPU and an integer ALU while a second portion of GPGPU cores **3562** include a double precision FPU. In at least one embodiment, FPUs can implement IEEE 754-2008 standard for floating point arithmetic or enable variable precision floating point arithmetic. In at least one embodiment, graphics multiprocessor **3596** can additionally include one or more fixed function or special function units to perform specific functions such as copy rectangle or pixel blending operations. In at least one embodiment one or more of GPGPU cores **3562** can also include fixed or special function logic.

[0276] In at least one embodiment, GPGPU cores **3562** include SIMD logic capable of performing a single instruction on multiple sets of data. In at least one embodiment GPGPU cores **3562** can physically execute SIMD4, SIMD8, and SIMD16 instructions and logically execute SIMD1, SIMD2, and SIMD32 instructions. In at least one embodiment, SIMD instructions for GPGPU cores **3562** can be generated at compile time by a shader compiler or automatically generated when executing programs written and compiled for single program multiple data (“SPMD”) or SIMT architectures. In at least one embodiment, multiple threads of a program configured for an SIMT execution model can be executed via a single SIMD instruction. For example, in at least one embodiment, eight SIMT threads that perform the same or similar operations can be executed in parallel via a single SIMD8 logic unit.

[0277] In at least one embodiment, memory and cache interconnect **3568** is an interconnect network that connects each functional unit of graphics multiprocessor **3596** to register file **3558** and to shared memory **3570**. In at least one embodiment, memory and cache interconnect **3568** is a crossbar interconnect that allows LSU **3566** to implement load and store operations between shared memory **3570** and register file **3558**. In at least one embodiment, register file **3558** can operate at a same frequency as GPGPU cores **3562**, thus data transfer between GPGPU cores **3562** and register file **3558** is very low latency. In at least one embodiment, shared memory **3570** can be used to enable communication between threads that execute on functional units within graphics multiprocessor **3596**. In at least one

embodiment, cache memory **3572** can be used as a data cache for example, to cache texture data communicated between functional units and texture unit **3536**. In at least one embodiment, shared memory **3570** can also be used as a program managed cache. In at least one embodiment, threads executing on GPGPU cores **3562** can programmatically store data within shared memory in addition to automatically cached data that is stored within cache memory **3572**.

[0278] In at least one embodiment, a parallel processor or GPGPU as described herein is communicatively coupled to host/processor cores to accelerate graphics operations, machine-learning operations, pattern analysis operations, and various general purpose GPU (GPGPU) functions. In at least one embodiment, a GPU may be communicatively coupled to host processor/cores over a bus or other interconnect (e.g., a high speed interconnect such as PCIe or NVLink). In at least one embodiment, a GPU may be integrated on the same package or chip as cores and communicatively coupled to cores over a processor bus/interconnect that is internal to a package or a chip. In at least one embodiment, regardless of the manner in which a GPU is connected, processor cores may allocate work to the GPU in the form of sequences of commands/instructions contained in a WD. In at least one embodiment, the GPU then uses dedicated circuitry/logic for efficiently processing these commands/instructions.

[0279] FIG. 36 illustrates a graphics processor **3600**, in accordance with at least one embodiment. In at least one embodiment, graphics processor **3600** includes a ring interconnect **3602**, a pipeline front-end **3604**, a media engine **3637**, and graphics cores **3680A-3680N**. In at least one embodiment, ring interconnect **3602** couples graphics processor **3600** to other processing units, including other graphics processors or one or more general-purpose processor cores. In at least one embodiment, graphics processor **3600** is one of many processors integrated within a multi-core processing system.

[0280] In at least one embodiment, at least one component shown or described with respect to FIG. 36 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, graphics processor **3600** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0281] In at least one embodiment, graphics processor **3600** receives batches of commands via ring interconnect **3602**. In at least one embodiment, incoming commands are interpreted by a command streamer **3603** in pipeline front-end **3604**. In at least one embodiment, graphics processor **3600** includes scalable execution logic to perform 3D geometry processing and media processing via graphics core(s) **3680A-3680N**. In at least one embodiment, for 3D geometry processing commands, command streamer **3603** supplies commands to geometry pipeline **3636**. In at least one embodiment, for at least some media processing commands, command streamer **3603** supplies commands to a video front end **3634**, which couples with a media engine **3637**. In at least one embodiment, media engine **3637** includes a Video Quality Engine (“VQE”) **3630** for video and image post-processing and a multi-format encode/decode (“MFX”) engine **3633** to provide hardware-accelerated media data

encode and decode. In at least one embodiment, geometry pipeline **3636** and media engine **3637** each generate execution threads for thread execution resources provided by at least one graphics core **3680A**.

[0282] In at least one embodiment, graphics processor **3600** includes scalable thread execution resources featuring modular graphics cores **3680A-3680N** (sometimes referred to as core slices), each having multiple sub-cores **3650A-3650N**, **3660A-3660N** (sometimes referred to as core sub-slices). In at least one embodiment, graphics processor **3600** can have any number of graphics cores **3680A** through **3680N**. In at least one embodiment, graphics processor **3600** includes a graphics core **3680A** having at least a first sub-core **3650A** and a second sub-core **3660A**. In at least one embodiment, graphics processor **3600** is a low power processor with a single sub-core (e.g., sub-core **3650A**). In at least one embodiment, graphics processor **3600** includes multiple graphics cores **3680A-3680N**, each including a set of first sub-cores **3650A-3650N** and a set of second sub-cores **3660A-3660N**. In at least one embodiment, each sub-core in first sub-cores **3650A-3650N** includes at least a first set of execution units (“EUs”) **3652A-3652N** and media/texture samplers **3654A-3654N**. In at least one embodiment, each sub-core in second sub-cores **3660A-3660N** includes at least a second set of execution units **3662A-3662N** and samplers **3664A-3664N**. In at least one embodiment, each sub-core **3650A-3650N**, **3660A-3660N** shares a set of shared resources **3670A-3670N**. In at least one embodiment, shared resources **3670** include shared cache memory and pixel operation logic.

[0283] FIG. 37 illustrates a processor **3700**, in accordance with at least one embodiment. In at least one embodiment, processor **3700** may include, without limitation, logic circuits to perform instructions. In at least one embodiment, processor **3700** may perform instructions, including x86 instructions, ARM instructions, specialized instructions for ASICs, etc. In at least one embodiment, processor **3710** may include registers to store packed data, such as 64-bit wide MMX™ registers in microprocessors enabled with MMX technology from Intel Corporation of Santa Clara, Calif. In at least one embodiment, MMX registers, available in both integer and floating point forms, may operate with packed data elements that accompany SIMD and streaming SIMD extensions (“SSE”) instructions. In at least one embodiment, 128-bit wide XMM registers relating to SSE2, SSE3, SSE4, AVX, or beyond (referred to generically as “SSEx”) technology may hold such packed data operands. In at least one embodiment, processors **3710** may perform instructions to accelerate CUDA programs.

[0284] In at least one embodiment, at least one component shown or described with respect to FIG. 37 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, processor **3700** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0285] In at least one embodiment, processor **3700** includes an in-order front end (“front end”) **3701** to fetch instructions to be executed and prepare instructions to be used later in processor pipeline. In at least one embodiment, front end **3701** may include several units. In at least one embodiment, an instruction prefetcher **3726** fetches instruc-

tions from memory and feeds instructions to an instruction decoder **3728** which in turn decodes or interprets instructions. For example, in at least one embodiment, instruction decoder **3728** decodes a received instruction into one or more operations called “micro-instructions” or “micro-operations” (also called “micro ops” or “uops”) for execution. In at least one embodiment, instruction decoder **3728** parses instruction into an opcode and corresponding data and control fields that may be used by micro-architecture to perform operations. In at least one embodiment, a trace cache **3730** may assemble decoded uops into program ordered sequences or traces in a uop queue **3734** for execution. In at least one embodiment, when trace cache **3730** encounters a complex instruction, a microcode ROM **3732** provides uops needed to complete an operation.

[0286] In at least one embodiment, some instructions may be converted into a single micro-op, whereas others need several micro-ops to complete full operation. In at least one embodiment, if more than four micro-ops are needed to complete an instruction, instruction decoder **3728** may access microcode ROM **3732** to perform instruction. In at least one embodiment, an instruction may be decoded into a small number of micro-ops for processing at instruction decoder **3728**. In at least one embodiment, an instruction may be stored within microcode ROM **3732** should a number of micro-ops be needed to accomplish operation. In at least one embodiment, trace cache **3730** refers to an entry point programmable logic array (“PLA”) to determine a correct micro-instruction pointer for reading microcode sequences to complete one or more instructions from microcode ROM **3732**. In at least one embodiment, after microcode ROM **3732** finishes sequencing micro-ops for an instruction, front end **3701** of machine may resume fetching micro-ops from trace cache **3730**.

[0287] In at least one embodiment, out-of-order execution engine (“out of order engine”) **3703** may prepare instructions for execution. In at least one embodiment, out-of-order execution logic has a number of buffers to smooth out and re-order the flow of instructions to optimize performance as they go down a pipeline and get scheduled for execution. Out-of-order execution engine **3703** includes, without limitation, an allocator/register renamer **3740**, a memory uop queue **3742**, an integer/floating point uop queue **3744**, a memory scheduler **3746**, a fast scheduler **3702**, a slow/general floating point scheduler (“slow/general FP scheduler”) **3704**, and a simple floating point scheduler (“simple FP scheduler”) **3706**. In at least one embodiment, fast scheduler **3702**, slow/general floating point scheduler **3704**, and simple floating point scheduler **3706** are also collectively referred to herein as “uop schedulers **3702**, **3704**, **3706**.” Allocator/register renamer **3740** allocates machine buffers and resources that each uop needs in order to execute. In at least one embodiment, allocator/register renamer **3740** renames logic registers onto entries in a register file. In at least one embodiment, allocator/register renamer **3740** also allocates an entry for each uop in one of two uop queues, memory uop queue **3742** for memory operations and integer/floating point uop queue **3744** for non-memory operations, in front of memory scheduler **3746** and uop schedulers **3702**, **3704**, **3706**. In at least one embodiment, uop schedulers **3702**, **3704**, **3706**, determine when a uop is ready to execute based on readiness of their dependent input register operand sources and availability of execution resources uops need to complete their operation. In at least

one embodiment, fast scheduler **3702** of at least one embodiment may schedule on each half of main clock cycle while slow/general floating point scheduler **3704** and simple floating point scheduler **3706** may schedule once per main processor clock cycle. In at least one embodiment, uop schedulers **3702**, **3704**, **3706** arbitrate for dispatch ports to schedule uops for execution.

[0288] In at least one embodiment, execution block **3711** includes, without limitation, an integer register file/bypass network **3708**, a floating point register file/bypass network (“FP register file/bypass network”) **3710**, address generation units (“AGUs”) **3712** and **3714**, fast ALUs **3716** and **3718**, a slow ALU **3720**, a floating point ALU (“FP”) **3722**, and a floating point move unit (“FP move”) **3724**. In at least one embodiment, integer register file/bypass network **3708** and floating point register file/bypass network **3710** are also referred to herein as “register files **3708**, **3710**.” In at least one embodiment, AGUs **3712** and **3714**, fast ALUs **3716** and **3718**, slow ALU **3720**, floating point ALU **3722**, and floating point move unit **3724** are also referred to herein as “execution units **3712**, **3714**, **3716**, **3718**, **3720**, **3722**, and **3724**.” In at least one embodiment, an execution block may include, without limitation, any number (including zero) and type of register files, bypass networks, address generation units, and execution units, in any combination.

[0289] In at least one embodiment, register files **3708**, **3710** may be arranged between uop schedulers **3702**, **3704**, **3706**, and execution units **3712**, **3714**, **3716**, **3718**, **3720**, **3722**, and **3724**. In at least one embodiment, integer register file/bypass network **3708** performs integer operations. In at least one embodiment, floating point register file/bypass network **3710** performs floating point operations. In at least one embodiment, each of register files **3708**, **3710** may include, without limitation, a bypass network that may bypass or forward just completed results that have not yet been written into register file to new dependent uops. In at least one embodiment, register files **3708**, **3710** may communicate data with each other. In at least one embodiment, integer register file/bypass network **3708** may include, without limitation, two separate register files, one register file for low-order thirty-two bits of data and a second register file for high order thirty-two bits of data. In at least one embodiment, floating point register file/bypass network **3710** may include, without limitation, 128-bit wide entries because floating point instructions typically have operands from 64 to 128 bits in width.

[0290] In at least one embodiment, execution units **3712**, **3714**, **3716**, **3718**, **3720**, **3722**, **3724** may execute instructions. In at least one embodiment, register files **3708**, **3710** store integer and floating point data operand values that micro-instructions need to execute. In at least one embodiment, processor **3700** may include, without limitation, any number and combination of execution units **3712**, **3714**, **3716**, **3718**, **3720**, **3722**, **3724**. In at least one embodiment, floating point ALU **3722** and floating point move unit **3724** may execute floating point, MMX, SIMD, AVX and SSE, or other operations. In at least one embodiment, floating point ALU **3722** may include, without limitation, a 64-bit by 64-bit floating point divider to execute divide, square root, and remainder micro ops. In at least one embodiment, instructions involving a floating point value may be handled with floating point hardware. In at least one embodiment, ALU operations may be passed to fast ALUs **3716**, **3718**. In at least one embodiment, fast ALUs **3716**, **3718** may

execute fast operations with an effective latency of half a clock cycle. In at least one embodiment, most complex integer operations go to slow ALU **3720** as slow ALU **3720** may include, without limitation, integer execution hardware for long-latency type of operations, such as a multiplier, shifts, flag logic, and branch processing. In at least one embodiment, memory load/store operations may be executed by AGUs **3712**, **3714**. In at least one embodiment, fast ALU **3716**, fast ALU **3718**, and slow ALU **3720** may perform integer operations on 64-bit data operands. In at least one embodiment, fast ALU **3716**, fast ALU **3718**, and slow ALU **3720** may be implemented to support a variety of data bit sizes including sixteen, thirty-two, 128, 256, etc. In at least one embodiment, floating point ALU **3722** and floating point move unit **3724** may be implemented to support a range of operands having bits of various widths. In at least one embodiment, floating point ALU **3722** and floating point move unit **3724** may operate on 128-bit wide packed data operands in conjunction with SIMD and multimedia instructions.

[0291] In at least one embodiment, uop schedulers **3702**, **3704**, **3706** dispatch dependent operations before parent load has finished executing. In at least one embodiment, as uops may be speculatively scheduled and executed in processor **3700**, processor **3700** may also include logic to handle memory misses. In at least one embodiment, if a data load misses in a data cache, there may be dependent operations in flight in pipeline that have left a scheduler with temporarily incorrect data. In at least one embodiment, a replay mechanism tracks and re-executes instructions that use incorrect data. In at least one embodiment, dependent operations might need to be replayed and independent ones may be allowed to complete. In at least one embodiment, schedulers and replay mechanisms of at least one embodiment of a processor may also be designed to catch instruction sequences for text string comparison operations.

[0292] In at least one embodiment, the term “registers” may refer to on-board processor storage locations that may be used as part of instructions to identify operands. In at least one embodiment, registers may be those that may be usable from outside of a processor (from a programmer’s perspective). In at least one embodiment, registers might not be limited to a particular type of circuit. Rather, in at least one embodiment, a register may store data, provide data, and perform functions described herein. In at least one embodiment, registers described herein may be implemented by circuitry within a processor using any number of different techniques, such as dedicated physical registers, dynamically allocated physical registers using register renaming, combinations of dedicated and dynamically allocated physical registers, etc. In at least one embodiment, integer registers store 32-bit integer data. A register file of at least one embodiment also contains eight multimedia SIMD registers for packed data.

[0293] FIG. 38 illustrates a processor **3800**, in accordance with at least one embodiment. In at least one embodiment, processor **3800** includes, without limitation, one or more processor cores (“cores”) **3802A-3802N**, an integrated memory controller **3814**, and an integrated graphics processor **3808**. In at least one embodiment, processor **3800** can include additional cores up to and including additional processor core **3802N** represented by dashed lined boxes. In at least one embodiment, each of processor cores **3802A-3802N** includes one or more internal cache units **3804A-**

3804N. In at least one embodiment, each processor core also has access to one or more shared cache units **3806**. In at least one embodiment, one or more processor cores **3802A-3802N** are referred to as one or more compute units or computing units.

[0294] In at least one embodiment, at least one component shown or described with respect to FIG. 38 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, processor **3800** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0295] In at least one embodiment, internal cache units **3804A-3804N** and shared cache units **3806** represent a cache memory hierarchy within processor **3800**. In at least one embodiment, cache memory units **3804A-3804N** may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as an L2, L3, Level 4 (“L4”), or other levels of cache, where a highest level of cache before external memory is classified as an LLC. In at least one embodiment, cache coherency logic maintains coherency between various cache units **3806** and **3804A-3804N**.

[0296] In at least one embodiment, processor **3800** may also include a set of one or more bus controller units **3816** and a system agent core **3810**. In at least one embodiment, one or more bus controller units **3816** manage a set of peripheral buses, such as one or more PCI or PCI express buses. In at least one embodiment, system agent core **3810** provides management functionality for various processor components. In at least one embodiment, system agent core **3810** includes one or more integrated memory controllers **3814** to manage access to various external memory devices (not shown).

[0297] In at least one embodiment, one or more of processor cores **3802A-3802N** include support for simultaneous multi-threading. In at least one embodiment, system agent core **3810** includes components for coordinating and operating processor cores **3802A-3802N** during multi-threaded processing. In at least one embodiment, system agent core **3810** may additionally include a power control unit (“PCU”), which includes logic and components to regulate one or more power states of processor cores **3802A-3802N** and graphics processor **3808**.

[0298] In at least one embodiment, processor **3800** additionally includes graphics processor **3808** to execute graphics processing operations. In at least one embodiment, graphics processor **3808** couples with shared cache units **3806**, and system agent core **3810**, including one or more integrated memory controllers **3814**. In at least one embodiment, system agent core **3810** also includes a display controller **3811** to drive graphics processor output to one or more coupled displays. In at least one embodiment, display controller **3811** may also be a separate module coupled with graphics processor **3808** via at least one interconnect, or may be integrated within graphics processor **3808**.

[0299] In at least one embodiment, a ring based interconnect unit **3812** is used to couple internal components of processor **3800**. In at least one embodiment, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, or other techniques.

In at least one embodiment, graphics processor **3808** couples with ring interconnect **3812** via an I/O link **3813**.

[0300] In at least one embodiment, I/O link **3813** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **3818**, such as an eDRAM module. In at least one embodiment, each of processor cores **3802A-3802N** and graphics processor **3808** use embedded memory modules **3818** as a shared LLC.

[0301] In at least one embodiment, processor cores **3802A-3802N** are homogeneous cores executing a common instruction set architecture. In at least one embodiment, processor cores **3802A-3802N** are heterogeneous in terms of ISA, where one or more of processor cores **3802A-3802N** execute a common instruction set, while one or more other cores of processor cores **3802A-3802N** executes a subset of a common instruction set or a different instruction set. In at least one embodiment, processor cores **3802A-3802N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more cores having a lower power consumption. In at least one embodiment, processor **3800** can be implemented on one or more chips or as an SoC integrated circuit.

[0302] FIG. 39 illustrates a graphics processor core **3900**, in accordance with at least one embodiment described. In at least one embodiment, graphics processor core **3900** is included within a graphics core array. In at least one embodiment, graphics processor core **3900**, sometimes referred to as a core slice, can be one or multiple graphics cores within a modular graphics processor. In at least one embodiment, graphics processor core **3900** is exemplary of one graphics core slice, and a graphics processor as described herein may include multiple graphics core slices based on target power and performance envelopes. In at least one embodiment, each graphics core **3900** can include a fixed function block **3930** coupled with multiple sub-cores **3901A-3901F**, also referred to as sub-slices, that include modular blocks of general-purpose and fixed function logic.

[0303] In at least one embodiment, at least one component shown or described with respect to FIG. 39 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, graphics core **3900** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0304] In at least one embodiment, fixed function block **3930** includes a geometry/fixed function pipeline **3936** that can be shared by all sub-cores in graphics processor **3900**, for example, in lower performance and/or lower power graphics processor implementations. In at least one embodiment, geometry/fixed function pipeline **3936** includes a 3D fixed function pipeline, a video front-end unit, a thread spawner and thread dispatcher, and a unified return buffer manager, which manages unified return buffers.

[0305] In at least one embodiment, fixed function block **3930** also includes a graphics SoC interface **3937**, a graphics microcontroller **3938**, and a media pipeline **3939**. Graphics SoC interface **3937** provides an interface between graphics core **3900** and other processor cores within an SoC integrated circuit. In at least one embodiment, graphics micro-

controller **3938** is a programmable sub-processor that is configurable to manage various functions of graphics processor **3900**, including thread dispatch, scheduling, and pre-emption. In at least one embodiment, media pipeline **3939** includes logic to facilitate decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. In at least one embodiment, media pipeline **3939** implements media operations via requests to compute or sampling logic within sub-cores **3901-3901F**.

[0306] In at least one embodiment, SoC interface **3937** enables graphics core **3900** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC, including memory hierarchy elements such as a shared LLC memory, system RAM, and/or embedded on-chip or on-package DRAM. In at least one embodiment, SoC interface **3937** can also enable communication with fixed function devices within an SoC, such as camera imaging pipelines, and enables use of and/or implements global memory atomics that may be shared between graphics core **3900** and CPUs within an SoC. In at least one embodiment, SoC interface **3937** can also implement power management controls for graphics core **3900** and enable an interface between a clock domain of graphic core **3900** and other clock domains within an SoC. In at least one embodiment, SoC interface **3937** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. In at least one embodiment, commands and instructions can be dispatched to media pipeline **3939**, when media operations are to be performed, or a geometry and fixed function pipeline (e.g., geometry and fixed function pipeline **3936**, geometry and fixed function pipeline **3914**) when graphics processing operations are to be performed.

[0307] In at least one embodiment, graphics microcontroller **3938** can be configured to perform various scheduling and management tasks for graphics core **3900**. In at least one embodiment, graphics microcontroller **3938** can perform graphics and/or compute workload scheduling on various graphics parallel engines within execution unit (EU) arrays **3902A-3902F**, **3904A-3904F** within sub-cores **3901A-3901F**. In at least one embodiment, host software executing on a CPU core of an SoC including graphics core **3900** can submit workloads one of multiple graphic processor doorbells, which invokes a scheduling operation on an appropriate graphics engine. In at least one embodiment, scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In at least one embodiment, graphics microcontroller **3938** can also facilitate low-power or idle states for graphics core **3900**, providing graphics core **3900** with an ability to save and restore registers within graphics core **3900** across low-power state transitions independently from an operating system and/or graphics driver software on a system.

[0308] In at least one embodiment, graphics core **3900** may have greater than or fewer than illustrated sub-cores **3901A-3901F**, up to N modular sub-cores. For each set of N sub-cores, in at least one embodiment, graphics core **3900** can also include shared function logic **3910**, shared and/or

cache memory **3912**, a geometry/fixed function pipeline **3914**, as well as additional fixed function logic **3916** to accelerate various graphics and compute processing operations. In at least one embodiment, shared function logic **3910** can include logic units (e.g., sampler, math, and/or inter-thread communication logic) that can be shared by each N sub-cores within graphics core **3900**. Shared and/or cache memory **3912** can be an LLC for N sub-cores **3901A-3901F** within graphics core **3900** and can also serve as shared memory that is accessible by multiple sub-cores. In at least one embodiment, geometry/fixed function pipeline **3914** can be included instead of geometry/fixed function pipeline **3936** within fixed function block **3930** and can include same or similar logic units.

[0309] In at least one embodiment, graphics core **3900** includes additional fixed function logic **3916** that can include various fixed function acceleration logic for use by graphics core **3900**. In at least one embodiment, additional fixed function logic **3916** includes an additional geometry pipeline for use in position only shading. In position-only shading, at least two geometry pipelines exist, whereas in a full geometry pipeline within geometry/fixed function pipeline **3916**, **3936**, and a cull pipeline, which is an additional geometry pipeline which may be included within additional fixed function logic **3916**. In at least one embodiment, cull pipeline is a trimmed down version of a full geometry pipeline. In at least one embodiment, a full pipeline and a cull pipeline can execute different instances of an application, each instance having a separate context. In at least one embodiment, position only shading can hide long cull runs of discarded triangles, enabling shading to be completed earlier in some instances. For example, in at least one embodiment, cull pipeline logic within additional fixed function logic **3916** can execute position shaders in parallel with a main application and generally generates critical results faster than a full pipeline, as a cull pipeline fetches and shades position attribute of vertices, without performing rasterization and rendering of pixels to a frame buffer. In at least one embodiment, a cull pipeline can use generated critical results to compute visibility information for all triangles without regard to whether those triangles are culled. In at least one embodiment, a full pipeline (which in this instance may be referred to as a replay pipeline) can consume visibility information to skip culled triangles to shade only visible triangles that are finally passed to a rasterization phase.

[0310] In at least one embodiment, additional fixed function logic **3916** can also include general purpose processing acceleration logic, such as fixed function matrix multiplication logic, for accelerating CUDA programs.

[0311] In at least one embodiment, each graphics sub-core **3901A-3901F** includes a set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. In at least one embodiment, graphics sub-cores **3901A-3901F** include multiple EU arrays **3902A-3902F**, **3904A-3904F**, thread dispatch and inter-thread communication (“TD/IC”) logic **3903A-3903F**, a 3D (e.g., texture) sampler **3905A-3905F**, a media sampler **3906A-3906F**, a shader processor **3907A-3907F**, and shared local memory (“SLM”) **3908A-3908F**. EU arrays **3902A-3902F**, **3904A-3904F** each include multiple execution units, which are GPGPUs capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media,

or compute operation, including graphics, media, or compute shader programs. In at least one embodiment, TD/IC logic **3903A-3903F** performs local thread dispatch and thread control operations for execution units within a sub-core and facilitate communication between threads executing on execution units of a sub-core. In at least one embodiment, 3D sampler **3905A-3905F** can read texture or other 3D graphics related data into memory. In at least one embodiment, 3D sampler can read texture data differently based on a configured sample state and texture format associated with a given texture. In at least one embodiment, media sampler **3906A-3906F** can perform similar read operations based on a type and format associated with media data. In at least one embodiment, each graphics sub-core **3901A-3901F** can alternately include a unified 3D and media sampler. In at least one embodiment, threads executing on execution units within each of sub-cores **3901A-3901F** can make use of shared local memory **3908A-3908F** within each sub-core, to enable threads executing within a thread group to execute using a common pool of on-chip memory.

[0312] FIG. 40 illustrates a parallel processing unit (“PPU”) **4000**, in accordance with at least one embodiment. In at least one embodiment, PPU **4000** is configured with machine-readable code that, if executed by PPU **4000**, causes PPU **4000** to perform some or all of processes and techniques described herein. In at least one embodiment, PPU **4000** is a multi-threaded processor that is implemented on one or more integrated circuit devices and that utilizes multithreading as a latency-hiding technique designed to process computer-readable instructions (also referred to as machine-readable instructions or simply instructions) on multiple threads in parallel. In at least one embodiment, a thread refers to a thread of execution and is an instantiation of a set of instructions configured to be executed by PPU **4000**. In at least one embodiment, PPU **4000** is a GPU configured to implement a graphics rendering pipeline for processing three-dimensional (“3D”) graphics data in order to generate two-dimensional (“2D”) image data for display on a display device such as an LCD device. In at least one embodiment, PPU **4000** is utilized to perform computations such as linear algebra operations and machine-learning operations. FIG. 40 illustrates an example parallel processor for illustrative purposes only and should be construed as a non-limiting example of a processor architecture that may be implemented in at least one embodiment.

[0313] In at least one embodiment, at least one component shown or described with respect to FIG. 40 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, PPU **4000** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0314] In at least one embodiment, one or more PPUs **4000** are configured to accelerate High Performance Computing (“HPC”), data center, and machine learning applications. In at least one embodiment, one or more PPUs **4000** are configured to accelerate CUDA programs. In at least one embodiment, PPU **4000** includes, without limitation, an I/O unit **4006**, a front-end unit **4010**, a scheduler unit **4012**, a work distribution unit **4014**, a hub **4016**, a crossbar (“Xbar”) **4020**, one or more general processing clusters (“GPCs”)

4018, and one or more partition units (“memory partition units”) **4022**. In at least one embodiment, PPU **4000** is connected to a host processor or other PPUs **4000** via one or more high-speed GPU interconnects (“GPU interconnects”) **4008**. In at least one embodiment, PPU **4000** is connected to a host processor or other peripheral devices via a system bus or interconnect **4002**. In at least one embodiment, PPU **4000** is connected to a local memory comprising one or more memory devices (“memory”) **4004**. In at least one embodiment, memory devices **4004** include, without limitation, one or more dynamic random access memory (DRAM) devices. In at least one embodiment, one or more DRAM devices are configured and/or configurable as high-bandwidth memory (“HBM”) subsystems, with multiple DRAM dies stacked within each device.

[0315] In at least one embodiment, high-speed GPU interconnect **4008** may refer to a wire-based multi-lane communications link that is used by systems to scale and include one or more PPUs **4000** combined with one or more CPUs, supports cache coherence between PPUs **4000** and CPUs, and CPU mastering. In at least one embodiment, data and/or commands are transmitted by high-speed GPU interconnect **4008** through hub **4016** to/from other units of PPU **4000** such as one or more copy engines, video encoders, video decoders, power management units, and other components which may not be explicitly illustrated in FIG. 40.

[0316] In at least one embodiment, I/O unit **4006** is configured to transmit and receive communications (e.g., commands, data) from a host processor (not illustrated in FIG. 40) over system bus **4002**. In at least one embodiment, I/O unit **4006** communicates with host processor directly via system bus **4002** or through one or more intermediate devices such as a memory bridge. In at least one embodiment, I/O unit **4006** may communicate with one or more other processors, such as one or more of PPUs **4000** via system bus **4002**. In at least one embodiment, I/O unit **4006** implements a PCIe interface for communications over a PCIe bus. In at least one embodiment, I/O unit **4006** implements interfaces for communicating with external devices.

[0317] In at least one embodiment, I/O unit **4006** decodes packets received via system bus **4002**. In at least one embodiment, at least some packets represent commands configured to cause PPU **4000** to perform various operations. In at least one embodiment, I/O unit **4006** transmits decoded commands to various other units of PPU **4000** as specified by commands. In at least one embodiment, commands are transmitted to front-end unit **4010** and/or transmitted to hub **4016** or other units of PPU **4000** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly illustrated in FIG. 40). In at least one embodiment, I/O unit **4006** is configured to route communications between and among various logical units of PPU **4000**.

[0318] In at least one embodiment, a program executed by host processor encodes a command stream in a buffer that provides workloads to PPU **4000** for processing. In at least one embodiment, a workload comprises instructions and data to be processed by those instructions. In at least one embodiment, buffer is a region in a memory that is accessible (e.g., read/write) by both a host processor and PPU **4000**—a host interface unit may be configured to access buffer in a system memory connected to system bus **4002** via memory requests transmitted over system bus **4002** by I/O unit **4006**. In at least one embodiment, a host processor

writes a command stream to a buffer and then transmits a pointer to the start of the command stream to PPU 4000 such that front-end unit 4010 receives pointers to one or more command streams and manages one or more command streams, reading commands from command streams and forwarding commands to various units of PPU 4000.

[0319] In at least one embodiment, front-end unit 4010 is coupled to scheduler unit 4012 that configures various GPCs 4018 to process tasks defined by one or more command streams. In at least one embodiment, scheduler unit 4012 is configured to track state information related to various tasks managed by scheduler unit 4012 where state information may indicate which of GPCs 4018 a task is assigned to, whether task is active or inactive, a priority level associated with task, and so forth. In at least one embodiment, scheduler unit 4012 manages execution of a plurality of tasks on one or more of GPCs 4018.

[0320] In at least one embodiment, scheduler unit 4012 is coupled to work distribution unit 4014 that is configured to dispatch tasks for execution on GPCs 4018. In at least one embodiment, work distribution unit 4014 tracks a number of scheduled tasks received from scheduler unit 4012 and work distribution unit 4014 manages a pending task pool and an active task pool for each of GPCs 4018. In at least one embodiment, pending task pool comprises a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular GPC 4018; active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by GPCs 4018 such that as one of GPCs 4018 completes execution of a task, that task is evicted from active task pool for GPC 4018 and one of other tasks from pending task pool is selected and scheduled for execution on GPC 4018. In at least one embodiment, if an active task is idle on GPC 4018, such as while waiting for a data dependency to be resolved, then the active task is evicted from GPC 4018 and returned to a pending task pool while another task in the pending task pool is selected and scheduled for execution on GPC 4018.

[0321] In at least one embodiment, work distribution unit 4014 communicates with one or more GPCs 4018 via XBar 4020. In at least one embodiment, XBar 4020 is an interconnect network that couples many units of PPU 4000 to other units of PPU 4000 and can be configured to couple work distribution unit 4014 to a particular GPC 4018. In at least one embodiment, one or more other units of PPU 4000 may also be connected to XBar 4020 via hub 4016.

[0322] In at least one embodiment, tasks are managed by scheduler unit 4012 and dispatched to one of GPCs 4018 by work distribution unit 4014. GPC 4018 is configured to process task and generate results. In at least one embodiment, results may be consumed by other tasks within GPC 4018, routed to a different GPC 4018 via XBar 4020, or stored in memory 4004. In at least one embodiment, results can be written to memory 4004 via partition units 4022, which implement a memory interface for reading and writing data to/from memory 4004. In at least one embodiment, results can be transmitted to another PPU 4004 or CPU via high-speed GPU interconnect 4008. In at least one embodiment, PPU 4000 includes, without limitation, a number U of partition units 4022 that is equal to number of separate and distinct memory devices 4004 coupled to PPU 4000.

[0323] In at least one embodiment, a host processor executes a driver kernel that implements an application programming interface (“API”) that enables one or more

applications executing on host processor to schedule operations for execution on PPU 4000. In at least one embodiment, multiple compute applications are simultaneously executed by PPU 4000 and PPU 4000 provides isolation, quality of service (“QoS”), and independent address spaces for multiple compute applications. In at least one embodiment, an application generates instructions (e.g., in the form of API calls) that cause a driver kernel to generate one or more tasks for execution by PPU 4000 and the driver kernel outputs tasks to one or more streams being processed by PPU 4000. In at least one embodiment, each task comprises one or more groups of related threads, which may be referred to as a warp. In at least one embodiment, a warp comprises a plurality of related threads (e.g., 32 threads) that can be executed in parallel. In at least one embodiment, cooperating threads can refer to a plurality of threads including instructions to perform a task and that exchange data through shared memory.

[0324] FIG. 41 illustrates a GPC 4100, in accordance with at least one embodiment. In at least one embodiment, GPC 4100 is GPC 4018 of FIG. 40. In at least one embodiment, each GPC 4100 includes, without limitation, a number of hardware units for processing tasks and each GPC 4100 includes, without limitation, a pipeline manager 4102, a pre-raster operations unit (“PROP”) 4104, a raster engine 4108, a work distribution crossbar (“WDX”) 4116, an MMU 4118, one or more Data Processing Clusters (“DPCs”) 4106, and any suitable combination of parts.

[0325] In at least one embodiment, at least one component shown or described with respect to FIG. 41 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, GPC 4100 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0326] In at least one embodiment, at least one component shown or described with respect to FIG. 41 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, GPC 4100 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0327] In at least one embodiment, operation of GPC 4100 is controlled by pipeline manager 4102. In at least one embodiment, pipeline manager 4102 manages configuration of one or more DPCs 4106 for processing tasks allocated to GPC 4100. In at least one embodiment, pipeline manager 4102 configures at least one of one or more DPCs 4106 to implement at least a portion of a graphics rendering pipeline. In at least one embodiment, DPC 4106 is configured to execute a vertex shader program on a programmable streaming multiprocessor (“SM”) 4114. In at least one embodiment, pipeline manager 4102 is configured to route packets received from a work distribution unit to appropriate logical units within GPC 4100 and, in at least one embodiment, some packets may be routed to fixed function hardware units in PROP 4104 and/or raster engine 4108 while other packets may be routed to DPCs 4106 for processing by a primitive engine 4112 or SM 4114. In at least one embodiment, pipeline manager 4102 configures at least one of DPCs 4106

to implement a computing pipeline. In at least one embodiment, pipeline manager **4102** configures at least one of DPCs **4106** to execute at least a portion of a CUDA program.

[0328] In at least one embodiment, PROP unit **4104** is configured to route data generated by raster engine **4108** and DPCs **4106** to a Raster Operations (“ROP”) unit in a partition unit, such as memory partition unit **4022** described in more detail above in conjunction with FIG. 40. In at least one embodiment, PROP unit **4104** is configured to perform optimizations for color blending, organize pixel data, perform address translations, and more. In at least one embodiment, raster engine **4108** includes, without limitation, a number of fixed function hardware units configured to perform various raster operations and, in at least one embodiment, raster engine **4108** includes, without limitation, a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, a tile coalescing engine, and any suitable combination thereof. In at least one embodiment, a setup engine receives transformed vertices and generates plane equations associated with geometric primitive defined by vertices; plane equations are transmitted to a coarse raster engine to generate coverage information (e.g., an x, y coverage mask for a tile) for a primitive; the output of the coarse raster engine is transmitted to a culling engine where fragments associated with a primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. In at least one embodiment, fragments that survive clipping and culling are passed to a fine raster engine to generate attributes for pixel fragments based on plane equations generated by a setup engine. In at least one embodiment, the output of raster engine **4108** comprises fragments to be processed by any suitable entity such as by a fragment shader implemented within DPC **4106**.

[0329] In at least one embodiment, each DPC **4106** included in GPC **4100** comprise, without limitation, an M-Pipe Controller (“MPC”) **4110**; primitive engine **4112**; one or more SMs **4114**; and any suitable combination thereof. In at least one embodiment, MPC **4110** controls operation of DPC **4106**, routing packets received from pipeline manager **4102** to appropriate units in DPC **4106**. In at least one embodiment, packets associated with a vertex are routed to primitive engine **4112**, which is configured to fetch vertex attributes associated with vertex from memory; in contrast, packets associated with a shader program may be transmitted to SM **4114**.

[0330] In at least one embodiment, SM **4114** comprises, without limitation, a programmable streaming processor that is configured to process tasks represented by a number of threads. In at least one embodiment, SM **4114** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently and implements a SIMD architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on same set of instructions. In at least one embodiment, all threads in group of threads execute same instructions. In at least one embodiment, SM **4114** implements a SMT architecture wherein each thread in a group of threads is configured to process a different set of data based on same set of instructions, but where individual threads in group of threads are allowed to diverge during execution. In at least one embodiment, a program counter, a call stack, and an execution state is maintained for each warp, enabling concurrency between

warps and serial execution within warps when threads within a warp diverge. In another embodiment, a program counter, a call stack, and an execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. In at least one embodiment, an execution state is maintained for each individual thread and threads executing the same instructions may be converged and executed in parallel for better efficiency. At least one embodiment of SM **4114** is described in more detail in conjunction with FIG. 42.

[0331] In at least one embodiment, MMU **4118** provides an interface between GPC **4100** and a memory partition unit (e.g., partition unit **4022** of FIG. 40) and MMU **4118** provides translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In at least one embodiment, MMU **4118** provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in memory.

[0332] FIG. 42 illustrates a streaming multiprocessor (“SM”) **4200**, in accordance with at least one embodiment. In at least one embodiment, SM **4200** is SM **4114** of FIG. 41. In at least one embodiment, SM **4200** includes, without limitation, an instruction cache **4202**; one or more scheduler units **4204**; a register file **4208**; one or more processing cores (“cores”) **4210**; one or more special function units (“SFUs”) **4212**; one or more LSUs **4214**; an interconnect network **4216**; a shared memory/L1 cache **4218**; and any suitable combination thereof. In at least one embodiment, a work distribution unit dispatches tasks for execution on GPCs of parallel processing units (PPUs) and each task is allocated to a particular Data Processing Cluster (DPC) within a GPC and, if a task is associated with a shader program, then the task is allocated to one of SMs **4200**. In at least one embodiment, scheduler unit **4204** receives tasks from a work distribution unit and manages instruction scheduling for one or more thread blocks assigned to SM **4200**. In at least one embodiment, scheduler unit **4204** schedules thread blocks for execution as warps of parallel threads, wherein each thread block is allocated at least one warp. In at least one embodiment, each warp executes threads. In at least one embodiment, scheduler unit **4204** manages a plurality of different thread blocks, allocating warps to different thread blocks and then dispatching instructions from a plurality of different cooperative groups to various functional units (e.g., processing cores **4210**, SFUs **4212**, and LSUs **4214**) during each clock cycle. In at least one embodiment, SM **4200** includes one or more thread block clusters, where a thread block cluster can enable programmatic control of locality at a granularity larger than a single thread block of a single streaming multiprocessor (SM). In at least one embodiment, thread block clusters (also referred to as “clusters”) enables multiple thread blocks running concurrently across streaming multiprocessors to synchronize and collaboratively fetch, exchange, or otherwise use data.

[0333] In at least one embodiment, at least one component shown or described with respect to FIG. 42 is used to implement techniques and/or functions described in conjunction with FIGS. 1-23. In at least one embodiment, SM **4200** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0334] In at least one embodiment, “cooperative groups” may refer to a programming model for organizing groups of communicating threads that allows developers to express granularity at which threads are communicating, enabling expression of richer, more efficient parallel decompositions. In at least one embodiment, cooperative launch APIs support synchronization amongst thread blocks for execution of parallel algorithms. In at least one embodiment, APIs of conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., `syncthreads( )` function). However, in at least one embodiment, programmers may define groups of threads at smaller than thread block granularities and synchronize within defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces. In at least one embodiment, cooperative groups enable programmers to define groups of threads explicitly at sub-block and multi-block granularities, and to perform collective operations such as synchronization on threads in a cooperative group. In at least one embodiment, a sub-block granularity is as small as a single thread. In at least one embodiment, a programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. In at least one embodiment, cooperative group primitives enable new patterns of cooperative parallelism, including, without limitation, producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

[0335] In at least one embodiment, a dispatch unit **4206** is configured to transmit instructions to one or more of functional units and scheduler unit **4204** includes, without limitation, two dispatch units **4206** that enable two different instructions from same warp to be dispatched during each clock cycle. In at least one embodiment, each scheduler unit **4204** includes a single dispatch unit **4206** or additional dispatch units **4206**.

[0336] In at least one embodiment, each SM **4200**, in at least one embodiment, includes, without limitation, register file **4208** that provides a set of registers for functional units of SM **4200**. In at least one embodiment, register file **4208** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of register file **4208**. In at least one embodiment, register file **4208** is divided between different warps being executed by SM **4200** and register file **4208** provides temporary storage for operands connected to data paths of functional units. In at least one embodiment, each SM **4200** comprises, without limitation, a plurality of L processing cores **4210**. In at least one embodiment, SM **4200** includes, without limitation, a large number (e.g., 128 or more) of distinct processing cores **4210**. In at least one embodiment, each processing core **4210** includes, without limitation, a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes, without limitation, a floating point arithmetic logic unit and an integer arithmetic logic unit. In at least one embodiment, floating point arithmetic logic units implement IEEE 754-2008 standard for floating point arithmetic. In at least one embodiment, processing cores **4210** include, without limitation, 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

[0337] In at least one embodiment, tensor cores are configured to perform matrix operations. In at least one embodiment, one or more tensor cores are included in processing cores **4210**. In at least one embodiment, tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In at least one embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D=A \times B + C$, where A, B, C, and D are 4×4 matrices.

[0338] In at least one embodiment, matrix multiply inputs A and B are 16-bit floating point matrices and accumulation matrices C and D are 16-bit floating point or 32-bit floating point matrices. In at least one embodiment, tensor cores operate on 16-bit floating point input data with 32-bit floating point accumulation. In at least one embodiment, 16-bit floating point multiply uses 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with other intermediate products for a 4×4×4 matrix multiply. Tensor cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements, in at least one embodiment. In at least one embodiment, an API, such as a CUDA-C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use tensor cores from a CUDA-C++ program. In at least one embodiment, at the CUDA level, a warp-level interface assumes 16×16 size matrices spanning all 32 threads of a warp.

[0339] In at least one embodiment, each SM **4200** comprises, without limitation, M SFUs **4212** that perform special functions (e.g., attribute evaluation, reciprocal square root, and like). In at least one embodiment, SFUs **4212** include, without limitation, a tree traversal unit configured to traverse a hierarchical tree data structure. In at least one embodiment, SFUs **4212** include, without limitation, a texture unit configured to perform texture map filtering operations. In at least one embodiment, texture units are configured to load texture maps (e.g., a 2D array of texels) from memory and sample texture maps to produce sampled texture values for use in shader programs executed by SM **4200**. In at least one embodiment, texture maps are stored in shared memory/L1 cache **4218**. In at least one embodiment, texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In at least one embodiment, each SM **4200** includes, without limitation, two texture units.

[0340] In at least one embodiment, each SM **4200** comprises, without limitation, N LSUs **4214** that implement load and store operations between shared memory/L1 cache **4218** and register file **4208**. In at least one embodiment, each SM **4200** includes, without limitation, interconnect network **4216** that connects each of the functional units to register file **4208** and LSU **4214** to register file **4208** and shared memory/L1 cache **4218**. In at least one embodiment, interconnect network **4216** is a crossbar that can be configured to connect any of the functional units to any of the registers in register file **4208** and connect LSUs **4214** to register file **4208** and memory locations in shared memory/L1 cache **4218**.

[0341] In at least one embodiment, shared memory/L1 cache **4218** is an array of on-chip memory that allows for data storage and communication between SM **4200** and a primitive engine and between threads in SM **4200**. In at least

one embodiment, shared memory/L1 cache **4218** comprises, without limitation, 128 KB of storage capacity and is in a path from SM **4200** to a partition unit. In at least one embodiment, shared memory/L1 cache **4218** is used to cache reads and writes. In at least one embodiment, one or more of shared memory/L1 cache **4218**, L2 cache, and memory are backing stores.

[0342] In at least one embodiment, combining data cache and shared memory functionality into a single memory block provides improved performance for both types of memory accesses. In at least one embodiment, capacity is used or is usable as a cache by programs that do not use shared memory, such as if shared memory is configured to use half of capacity, texture and load/store operations can use remaining capacity. In at least one embodiment, integration within shared memory/L1 cache **4218** enables shared memory/L1 cache **4218** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data. In at least one embodiment, when configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. In at least one embodiment, fixed function GPUs are bypassed, creating a much simpler programming model. In at least one embodiment and in a general purpose parallel computation configuration, a work distribution unit assigns and distributes blocks of threads directly to DPCs. In at least one embodiment, threads in a block execute the same program, using a unique thread ID in a calculation to ensure each thread generates unique results, using SM **4200** to execute a program and perform calculations, shared memory/L1 cache **4218** to communicate between threads, and LSU **4214** to read and write global memory through shared memory/L1 cache **4218** and a memory partition unit. In at least one embodiment, when configured for general purpose parallel computation, SM **4200** writes commands that scheduler unit **4204** can use to launch new work on DPCs. In at least one embodiment, SM **4200** includes one or more distributed shared memories (or distributed shared memory) that enable direct SM-to-SM operations such as loading, storing, and performing atomics across multiple SM shared memory blocks.

[0343] In at least one embodiment, SM **4200** includes one or more asynchronous execution functions that include a tensor memory accelerator (TMA) unit that can transfer blocks of data between global memory and shared memory. In at least one embodiment, one or more processors uses or access one or more TMAs to perform bi-directional copy operations, e.g., from global to shared memory and vice versa. In at least one embodiment, SM **4200** includes one or more TMAs to asynchronously copy between thread blocks in a cluster. In at least one embodiment, SM **4200** includes one or more asynchronous transaction barriers to perform atomic data movement and synchronization. In at least one embodiment, SM **4200** includes a tensor core transformer engine, which includes software and one or more cores to accelerate transformer model training and inferencing. In at least one embodiment, a transformer one or more processor cores performing one or more tensor core transformer engines manage and dynamically choose between FP8 and 16-bit calculations by re-casting and scaling between FP8 and 16-bit in each layer of one or more neural networks.

[0344] In at least one embodiment, PPU is included in or coupled to a desktop computer, a laptop computer, a tablet

computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), a PDA, a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and more. In at least one embodiment, PPU is embodied on a single semiconductor substrate. In at least one embodiment, PPU is included in an SoC along with one or more other devices such as additional PPUs, memory, a RISC CPU, an MMU, a digital-to-analog converter (“DAC”), and like.

[0345] In at least one embodiment, PPU may be included on a graphics card that includes one or more memory devices. In at least one embodiment, a graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In at least one embodiment, PPU may be an integrated GPU (“iGPU”) included in chipset of motherboard.

Software Constructions for General-Purpose Computing

[0346] The following figures set forth, without limitation, exemplary software constructs for implementing at least one embodiment.

[0347] FIG. 43 illustrates a software stack of a programming platform, in accordance with at least one embodiment. In at least one embodiment, a programming platform is a platform for leveraging hardware on a computing system to accelerate computational tasks. A programming platform may be accessible to software developers through libraries, compiler directives, and/or extensions to programming languages, in at least one embodiment. In at least one embodiment, a programming platform may be, but is not limited to, CUDA, Radeon Open Compute Platform (“ROCm”), OpenCL (OpenCL™ is developed by Khronos group), SYCL, or Intel One API.

[0348] In at least one embodiment, at least one component shown or described with respect to FIG. 43 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, software stack **4300** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0349] In at least one embodiment, a software stack **4300** of a programming platform provides an execution environment for an application **4301**. In at least one embodiment, application **4301** may include any computer software capable of being launched on software stack **4300**. In at least one embodiment, application **4301** may include, but is not limited to, an artificial intelligence (“AI”)/machine learning (“ML”) application, a high performance computing (“HPC”) application, a virtual desktop infrastructure (“VDI”), or a data center workload.

[0350] In at least one embodiment, application **4301** and software stack **4300** run on hardware **4307**. Hardware **4307** may include one or more GPUs, CPUs, FPGAs, AI engines, and/or other types of compute devices that support a programming platform, in at least one embodiment. In at least one embodiment, such as with CUDA, software stack **4300** may be vendor specific and compatible with only devices from particular vendor(s). In at least one embodiment, such as in with OpenCL, software stack **4300** may be used with devices from different vendors. In at least one embodiment, hardware **4307** includes a host connected to one more devices that can be accessed to perform computational tasks

via application programming interface (“API”) calls. A device within hardware **4307** may include, but is not limited to, a GPU, FPGA, AI engine, or other compute device (but may also include a CPU) and its memory, as opposed to a host within hardware **4307** that may include, but is not limited to, a CPU (but may also include a compute device) and its memory, in at least one embodiment.

[0351] In at least one embodiment, software stack **4300** of a programming platform includes, without limitation, a number of libraries **4303**, a runtime **4305**, and a device kernel driver **4306**. Each of libraries **4303** may include data and programming code that can be used by computer programs and leveraged during software development, in at least one embodiment. In at least one embodiment, libraries **4303** may include, but are not limited to, pre-written code and subroutines, classes, values, type specifications, configuration data, documentation, help data, and/or message templates. In at least one embodiment, libraries **4303** include functions that are optimized for execution on one or more types of devices. In at least one embodiment, libraries **4303** may include, but are not limited to, functions for performing mathematical, deep learning, and/or other types of operations on devices. In at least one embodiment, libraries **4303** are associated with corresponding APIs **4302**, which may include one or more APIs, that expose functions implemented in libraries **4303**. In at least one embodiment, a processor (e.g., CPU, GPU) performs, calls, or otherwise uses one or more APIs to prioritize kernels. For example, a first kernel (e.g., parent) can launch a second kernel (e.g., child kernel), and said second kernel can be used by a processor to launch additional kernels (e.g., grandchildren kernels) independent of said first kernel. In at least one embodiment, a processor performs an API or calls an API from memory to be performed to support dynamic stream priority (e.g., updating priority while a stream is being used to perform operations). For example, when a processor performs said API, it allows a programmer to copy stream priority from one stream to one or more other streams.

[0352] In at least one embodiment, software stack **4300** includes an API to support dynamic stream priority (e.g., updating priority while a stream is being used to perform operations), which allows a programmer to set priority of a stream at any time after creation. In at least one embodiment, software stack **4300** includes an API to support dynamic stream priority (e.g., updating priority while the stream is being used to perform operations), which allows a programmer to obtain current priority of a stream, where the priority is one of a plurality of attributes of a stream. In at least one embodiment, software stack **4300** includes an API to support dynamic stream priority (e.g., updating priority while the stream is being used to perform operations), which allows a programmer to obtain current priority of a stream as a single attribute. In at least one embodiment, software stack **4300** includes an API to support dynamic stream priority (e.g., updating priority while the stream is being used to perform operations), which allows a programmer to launch a kernel to perform operations on a stream at a set priority, which may be different from the stream priority. In at least one embodiment, software stack **4300** includes an API to indicate whether an object (e.g., a thread synchronization object such as a barrier) tracks whether all data movement operations for a set of threads operating on a GPU are complete has a specified state after a specified period of time, where a specified state can be a state indicating that data has been

moved and is ready for use, and is specified using an expected parity value as an input to the API.

[0353] In at least one embodiment, software stack **4300** includes one or more APIs to update kernels. In at least one embodiment, a processor performs an API or calls an API from memory to be performed to update to an existing API is to support context-free kernels, which allows a programmer to add a kernel node to a graph without a graphics context, so that a graphics context can be dynamically associated with a kernel at runtime. In at least one embodiment, software stack **4300** includes one or more APIs to allow a programmer to obtain a kernel identifier and a graphics context as separate parameters from a kernel node, so that parameters to be obtained from kernels and from context-free kernels. In at least one embodiment, software stack **4300** includes one or more APIs to use parallel processor(s), such as one or more graphics processing units, to launch task graphs (e.g., task graphs) and to execute one or more task graphs (e.g., including one or more programs).

[0354] In at least one embodiment, software stack **4300** includes one or more APIs to associate one or more instructions with one or more memory ordering operations, such as a fence or membar operation. In at least one embodiment, instructions are associated with one or more domains such that a memory ordering operation is executed in association to one or more particular domains without interfering with instructions of other domains. an API to indicate a thread has arrived (e.g., at a thread synchronization barrier), or finished a stage of work in relation to asynchronous data movement operations on a GPU. In at least one embodiment, software stack **4300** includes one or more to allow programmers to manually indicate an expected transaction count when a thread has finished a stage of work, which is used to update an object that tracks whether all data movement operations for a set of threads are complete.

[0355] In at least one embodiment, application **4301** is written as source code that is compiled into executable code, as discussed in greater detail below in conjunction with FIGS. **48-50**. Executable code of application **4301** may run, at least in part, on an execution environment provided by software stack **4300**, in at least one embodiment. In at least one embodiment, during execution of application **4301**, code may be reached that needs to run on a device, as opposed to a host. In such a case, runtime **4305** may be called to load and launch requisite code on the device, in at least one embodiment. In at least one embodiment, runtime **4305** may include any technically feasible runtime system that is able to support execution of application **S01**.

[0356] In at least one embodiment, runtime **4305** is implemented as one or more runtime libraries associated with corresponding APIs, which are shown as API(s) **4304**. One or more of such runtime libraries may include, without limitation, functions for memory management, execution control, device management, error handling, and/or synchronization, among other things, in at least one embodiment. In at least one embodiment, memory management functions may include, but are not limited to, functions to allocate, deallocate, and copy device memory, as well as transfer data between host memory and device memory. In at least one embodiment, execution control functions may include, but are not limited to, functions to launch a function (sometimes referred to as a “kernel” when a function is a global function

callable from a host) on a device and set attribute values in a buffer maintained by a runtime library for a given function to be executed on a device.

[0357] Runtime libraries and corresponding API(s) **4304** may be implemented in any technically feasible manner, in at least one embodiment. In at least one embodiment, one (or any number of) API may expose a low-level set of functions for fine-grained control of a device, while another (or any number of) API may expose a higher-level set of such functions. In at least one embodiment, a high-level runtime API may be built on top of a low-level API. In at least one embodiment, one or more of runtime APIs may be language-specific APIs that are layered on top of a language-independent runtime API.

[0358] In at least one embodiment, one or more processors disclosed in “processing systems” can perform, access, or otherwise use software stack **4300**. For example, APU **3000**, CPU **3100**, **33A-33B** exemplary graphics processors, general-purpose graphics processing unit (“GPGPU”) **3430**, parallel processor **3500**, processing cluster **3594**, graphics multiprocessor **3534**, graphics multiprocessor **3596**, graphics processor **3600**, processor **3700**, processor **3800**, parallel processing unit (“PPU”) **4000**, GPC **4100**, and/or streaming multiprocessor (“SM”) **4200** can perform, use, call, or otherwise implement (e.g., through accessing a memory) one or more APIs included in software stack **4300**.

[0359] In at least one embodiment, device kernel driver **4306** is configured to facilitate communication with an underlying device. In at least one embodiment, device kernel driver **4306** may provide low-level functionalities upon which APIs, such as API(s) **4304**, and/or other software relies. In at least one embodiment, device kernel driver **4306** may be configured to compile intermediate representation (“IR”) code into binary code at runtime. For CUDA, device kernel driver **4306** may compile Parallel Thread Execution (“PTX”) IR code that is not hardware specific into binary code for a specific target device at runtime (with caching of compiled binary code), which is also sometimes referred to as “finalizing” code, in at least one embodiment. Doing so may permit finalized code to run on a target device, which may not have existed when source code was originally compiled into PTX code, in at least one embodiment. Alternatively, in at least one embodiment, device source code may be compiled into binary code offline, without requiring device kernel driver **4306** to compile IR code at runtime.

[0360] FIG. **44** illustrates a CUDA implementation of software stack **4300** of FIG. **43**, in accordance with at least one embodiment. In at least one embodiment, a CUDA software stack **4400**, on which an application **4401** may be launched, includes CUDA libraries **4403**, a CUDA runtime **4405**, a CUDA driver **4407**, and a device kernel driver **4408**. In at least one embodiment, CUDA software stack **4400** executes on hardware **4409**, which may include a GPU that supports CUDA and is developed by NVIDIA Corporation of Santa Clara, CA.

[0361] In at least one embodiment, at least one component shown or described with respect to FIG. **44** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, application **4401** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross

one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0362] In at least one embodiment, application **4401**, CUDA runtime **4405**, and device kernel driver **4408** may perform similar functionalities as application **4301**, runtime **4305**, and device kernel driver **4306**, respectively, which are described above in conjunction with FIG. **43**. In at least one embodiment, CUDA driver **4407** includes a library (libcudart.so) that implements a CUDA driver API **4406**. Similar to a CUDA runtime API **4404** implemented by a CUDA runtime library (cudart), CUDA driver API **4406** may, without limitation, expose functions for memory management, execution control, device management, error handling, synchronization, and/or graphics interoperability, among other things, in at least one embodiment. In at least one embodiment, CUDA driver API **4406** differs from CUDA runtime API **4404** in that CUDA runtime API **4404** simplifies device code management by providing implicit initialization, context (analogous to a process) management, and module (analogous to dynamically loaded libraries) management. In contrast to high-level CUDA runtime API **4404**, CUDA driver API **4406** is a low-level API providing more fine-grained control of the device, particularly with respect to contexts and module loading, in at least one embodiment. In at least one embodiment, CUDA driver API **4406** may expose functions for context management that are not exposed by CUDA runtime API **4404**. In at least one embodiment, CUDA driver API **4406** is also language-independent and supports, e.g., OpenCL in addition to CUDA runtime API **4404**. Further, in at least one embodiment, development libraries, including CUDA runtime **4405**, may be considered as separate from driver components, including user-mode CUDA driver **4407** and kernel-mode device driver **4408** (also sometimes referred to as a “display” driver).

[0363] In at least one embodiment, CUDA libraries **4403** may include, but are not limited to, mathematical libraries, deep learning libraries, parallel algorithm libraries, and/or signal/image/video processing libraries, which parallel computing applications such as application **4401** may utilize. In at least one embodiment, CUDA libraries **4403** may include mathematical libraries such as a cuBLAS library that is an implementation of Basic Linear Algebra Subprograms (“BLAS”) for performing linear algebra operations, a cuFFT library for computing fast Fourier transforms (“FFTs”), and a cuRAND library for generating random numbers, among others. In at least one embodiment, CUDA libraries **4403** may include deep learning libraries such as a cuDNN library of primitives for deep neural networks and a TensorRT platform for high-performance deep learning inference, among others.

[0364] FIG. **45** illustrates a ROCm implementation of software stack **4300** of FIG. **43**, in accordance with at least one embodiment. In at least one embodiment, a ROCm software stack **4500**, on which an application **4501** may be launched, includes a language runtime **4503**, a system runtime **4505**, a thunk **4507**, and a ROCm kernel driver **4508**. In at least one embodiment, ROCm software stack **4500** executes on hardware **4509**, which may include a GPU that supports ROCm and is developed by AMD Corporation of Santa Clara, CA.

[0365] In at least one embodiment, at least one component shown or described with respect to FIG. **45** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, ROCm

software stack **4500** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0366] In at least one embodiment, application **4501** may perform similar functionalities as application **4301** discussed above in conjunction with FIG. 43. In addition, language runtime **4503** and system runtime **4505** may perform similar functionalities as runtime **4305** discussed above in conjunction with FIG. 43, in at least one embodiment. In at least one embodiment, language runtime **4503** and system runtime **4505** differ in that system runtime **4505** is a language-independent runtime that implements a ROCr system runtime API **4504** and makes use of a Heterogeneous System Architecture (“HSA”) Runtime API. HSA runtime API is a thin, user-mode API that exposes interfaces to access and interact with an AMD GPU, including functions for memory management, execution control via architected dispatch of kernels, error handling, system and agent information, and runtime initialization and shutdown, among other things, in at least one embodiment. In contrast to system runtime **4505**, language runtime **4503** is an implementation of a language-specific runtime API **4502** layered on top of ROCr system runtime API **4504**, in at least one embodiment. In at least one embodiment, language runtime API may include, but is not limited to, a Heterogeneous compute Interface for Portability (“HIP”) language runtime API, a Heterogeneous Compute Compiler (“HCC”) language runtime API, or an OpenCL API, among others. HIP language in particular is an extension of C++ programming language with functionally similar versions of CUDA mechanisms, and, in at least one embodiment, a HIP language runtime API includes functions that are similar to those of CUDA runtime API **4404** discussed above in conjunction with FIG. 44, such as functions for memory management, execution control, device management, error handling, and synchronization, among other things.

[0367] In at least one embodiment, thunk (ROct) **4507** is an interface **4506** that can be used to interact with underlying ROCm driver **4508**. In at least one embodiment, ROCm driver **4508** is a ROCK driver, which is a combination of an AMDGPU driver and a HSA kernel driver (amdkfd). In at least one embodiment, AMDGPU driver is a device kernel driver for GPUs developed by AMD that performs similar functionalities as device kernel driver **4306** discussed above in conjunction with FIG. 43. In at least one embodiment, HSA kernel driver is a driver permitting different types of processors to share system resources more effectively via hardware features.

[0368] In at least one embodiment, various libraries (not shown) may be included in ROCm software stack **4500** above language runtime **4503** and provide functionality similarity to CUDA libraries **4403**, discussed above in conjunction with FIG. 44. In at least one embodiment, various libraries may include, but are not limited to, mathematical, deep learning, and/or other libraries such as a hipBLAS library that implements functions similar to those of CUDA cuBLAS, a rocFFT library for computing FFTs that is similar to CUDA cuFFT, among others.

[0369] FIG. 46 illustrates an OpenCL implementation of software stack **4300** of FIG. 43, in accordance with at least one embodiment. In at least one embodiment, an OpenCL

software stack **4600**, on which an application **4601** may be launched, includes an OpenCL framework **4610**, an OpenCL runtime **4606**, and a driver **4607**. In at least one embodiment, OpenCL software stack **4600** executes on hardware **4409** that is not vendor-specific. As OpenCL is supported by devices developed by different vendors, specific OpenCL drivers may be required to interoperate with hardware from such vendors, in at least one embodiment.

[0370] In at least one embodiment, at least one component shown or described with respect to FIG. 46 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, application **4601** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0371] In at least one embodiment, application **4601**, OpenCL runtime **4606**, device kernel driver **4607**, and hardware **4608** may perform similar functionalities as application **4301**, runtime **4305**, device kernel driver **4306**, and hardware **4307**, respectively, that are discussed above in conjunction with FIG. 43. In at least one embodiment, application **4601** further includes an OpenCL kernel **4602** with code that is to be executed on a device.

[0372] In at least one embodiment, OpenCL defines a “platform” that allows a host to control devices connected to the host. In at least one embodiment, an OpenCL framework provides a platform layer API and a runtime API, shown as platform API **4603** and runtime API **4605**. In at least one embodiment, runtime API **4605** uses contexts to manage execution of kernels on devices. In at least one embodiment, each identified device may be associated with a respective context, which runtime API **4605** may use to manage command queues, program objects, and kernel objects, share memory objects, among other things, for that device. In at least one embodiment, platform API **4603** exposes functions that permit device contexts to be used to select and initialize devices, submit work to devices via command queues, and enable data transfer to and from devices, among other things. In addition, OpenCL framework provides various built-in functions (not shown), including math functions, relational functions, and image processing functions, among others, in at least one embodiment.

[0373] In at least one embodiment, a compiler **4604** is also included in OpenCL framework **4610**. Source code may be compiled offline prior to executing an application or online during execution of an application, in at least one embodiment. In contrast to CUDA and ROCm, OpenCL applications in at least one embodiment may be compiled online by compiler **4604**, which is included to be representative of any number of compilers that may be used to compile source code and/or IR code, such as Standard Portable Intermediate Representation (“SPIR-V”) code, into binary code. Alternatively, in at least one embodiment, OpenCL applications may be compiled offline, prior to execution of such applications.

[0374] FIG. 47 illustrates software that is supported by a programming platform, in accordance with at least one embodiment. In at least one embodiment, a programming platform **4704** is configured to support various programming models **4703**, middlewares and/or libraries **4702**, and frameworks **4701** that an application **4700** may rely upon. In at least one embodiment, application **4700** may be an AI/ML

application implemented using, for example, a deep learning framework such as MXNet, PyTorch, or TensorFlow, which may rely on libraries such as cuDNN, NVIDIA Collective Communications Library (“NCCL”), and/or NVIDIA Developer Data Loading Library (“DALI”) CUDA libraries to provide accelerated computing on underlying hardware.

[0375] In at least one embodiment, at least one component shown or described with respect to FIG. 47 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, application 4702 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0376] In at least one embodiment, programming platform 4704 may be one of a CUDA, ROCm, or OpenCL platform described above in conjunction with FIG. 44, FIG. 45, and FIG. 46, respectively. In at least one embodiment, programming platform 4704 supports multiple programming models 4703, which are abstractions of an underlying computing system permitting expressions of algorithms and data structures. Programming models 4703 may expose features of underlying hardware in order to improve performance, in at least one embodiment. In at least one embodiment, programming models 4703 may include, but are not limited to, CUDA, HIP, OpenCL, C++ Accelerated Massive Parallelism (“C++ AMP”), Open Multi-Processing (“OpenMP”), Open Accelerators (“OpenACC”), and/or Vulkan Compute.

[0377] In at least one embodiment, libraries and/or middlewares 4702 provide implementations of abstractions of programming models 4704. In at least one embodiment, such libraries include data and programming code that may be used by computer programs and leveraged during software development. In at least one embodiment, such middlewares include software that provides services to applications beyond those available from programming platform 4704. In at least one embodiment, libraries and/or middlewares 4702 may include, but are not limited to, cuBLAS, cuFFT, cuRAND, and other CUDA libraries, or rocBLAS, rocFFT, rocRAND, and other ROCm libraries. In addition, in at least one embodiment, libraries and/or middlewares 4702 may include NCCL and ROCm Communication Collectives Library (“RCCL”) libraries providing communication routines for GPUs, a MIOpen library for deep learning acceleration, and/or an Eigen library for linear algebra, matrix and vector operations, geometrical transformations, numerical solvers, and related algorithms.

[0378] In at least one embodiment, application frameworks 4701 depend on libraries and/or middlewares 4702. In at least one embodiment, each of application frameworks 4701 is a software framework used to implement a standard structure of application software. Returning to the AI/ML example discussed above, an AI/ML application may be implemented using a framework such as Caffe, Caffe2, TensorFlow, Keras, PyTorch, or MxNet deep learning frameworks, in at least one embodiment.

[0379] FIG. 48 illustrates compiling code to execute on one of programming platforms of FIGS. 43-46, in accordance with at least one embodiment. In at least one embodiment, a compiler 4801 receives source code 4800 that includes both host code as well as device code. In at least one embodiment, compiler 4801 is configured to convert source code 4800 into host executable code 4802 for execu-

tion on a host and device executable code 4803 for execution on a device. In at least one embodiment, source code 4800 may either be compiled offline prior to execution of an application, or online during execution of an application. In at least one embodiment, compiler 4801 includes or has access to one or more libraries to recognize a sequence of API calls to perform a single fused API, where a single fused API is a combined API for two or more APIs.

[0380] In at least one embodiment, at least one component shown or described with respect to FIG. 48 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, compiler 4801 performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0381] In at least one embodiment, source code 4800 may include code in any programming language supported by compiler 4801, such as C++, C, Fortran, etc. In at least one embodiment, source code 4800 may be included in a single-source file having a mixture of host code and device code, with locations of device code being indicated therein. In at least one embodiment, a single-source file may be a .cu file that includes CUDA code or a .hip.cpp file that includes HIP code. Alternatively, in at least one embodiment, source code 4800 may include multiple source code files, rather than a single-source file, into which host code and device code are separated.

[0382] In at least one embodiment, compiler 4801 is configured to compile source code 4800 into host executable code 4802 for execution on a host and device executable code 4803 for execution on a device. In at least one embodiment, compiler 4801 performs operations including parsing source code 4800 into an abstract system tree (AST), performing optimizations, and generating executable code. In at least one embodiment in which source code 4800 includes a single-source file, compiler 4801 may separate device code from host code in such a single-source file, compile device code and host code into device executable code 4803 and host executable code 4802, respectively, and link device executable code 4803 and host executable code 4802 together in a single file, as discussed in greater detail below with respect to FIG. 49.

[0383] In at least one embodiment, host executable code 4802 and device executable code 4803 may be in any suitable format, such as binary code and/or IR code. In the case of CUDA, host executable code 4802 may include native object code and device executable code 4803 may include code in PTX intermediate representation, in at least one embodiment. In the case of ROCm, both host executable code 4802 and device executable code 4803 may include target binary code, in at least one embodiment.

[0384] FIG. 49 is a more detailed illustration of compiling code to execute on one of programming platforms of FIGS. 43-46, in accordance with at least one embodiment. In at least one embodiment, a compiler 4901 is configured to receive source code 4900, compile source code 4900, and output an executable file 4910. In at least one embodiment, source code 4900 is a single-source file, such as a .cu file, a .hip.cpp file, or a file in another format, that includes both host and device code. In at least one embodiment, compiler 4901 may be, but is not limited to, an NVIDIA CUDA

compiler (“NVCC”) for compiling CUDA code in .cu files, or a HCC compiler for compiling HIP code in .hip.cpp files.

[0385] In at least one embodiment, at least one component shown or described with respect to FIG. 49 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, compiler 4901 compiles code that causes one or more processors to perform one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0386] In at least one embodiment, compiler 4901 includes a compiler front end 4902, a host compiler 4905, a device compiler 4906, and a linker 4909. In at least one embodiment, compiler front end 4902 is configured to separate device code 4904 from host code 4903 in source code 4900. Device code 4904 is compiled by device compiler 4906 into device executable code 4908, which as described may include binary code or IR code, in at least one embodiment. Separately, host code 4903 is compiled by host compiler 4905 into host executable code 4907, in at least one embodiment. For NVCC, host compiler 4905 may be, but is not limited to, a general purpose C/C++ compiler that outputs native object code, while device compiler 4906 may be, but is not limited to, a Low Level Virtual Machine (“LLVM”)–based compiler that forks a LLVM compiler infrastructure and outputs PTX code or binary code, in at least one embodiment. For HCC, both host compiler 4905 and device compiler 4906 may be, but are not limited to, LLVM-based compilers that output target binary code, in at least one embodiment.

[0387] Subsequent to compiling source code 4900 into host executable code 4907 and device executable code 4908, linker 4909 links host and device executable code 4907 and 4908 together in executable file 4910, in at least one embodiment. In at least one embodiment, native object code for a host and PTX or binary code for a device may be linked together in an Executable and Linkable Format (“ELF”) file, which is a container format used to store object code.

[0388] FIG. 50 illustrates translating source code prior to compiling source code, in accordance with at least one embodiment. In at least one embodiment, source code 5000 is passed through a translation tool 5001, which translates source code 5000 into translated source code 5002. In at least one embodiment, a compiler 5003 is used to compile translated source code 5002 into host executable code 5004 and device executable code 5005 in a process that is similar to compilation of source code 4800 by compiler 4801 into host executable code 4802 and device executable 4803, as discussed above in conjunction with FIG. 48.

[0389] In at least one embodiment, at least one component shown or described with respect to FIG. 50 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, translation tool 5001 translates code that causes one or more processors to perform one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0390] In at least one embodiment, a translation performed by translation tool 5001 is used to port source 5000 for execution in a different environment than that in which it

was originally intended to run. In at least one embodiment, translation tool 5001 may include, but is not limited to, a HIP translator that is used to “hipify” CUDA code intended for a CUDA platform into HIP code that can be compiled and executed on a ROCm platform. In at least one embodiment, translation of source code 5000 may include parsing source code 5000 and converting calls to API(s) provided by one programming model (e.g., CUDA) into corresponding calls to API(s) provided by another programming model (e.g., HIP), as discussed in greater detail below in conjunction with FIGS. 51A-52. Returning to the example of hipifying CUDA code, calls to CUDA runtime API, CUDA driver API, and/or CUDA libraries may be converted to corresponding HIP API calls, in at least one embodiment. In at least one embodiment, automated translations performed by translation tool 5001 may sometimes be incomplete, requiring additional, manual effort to fully port source code 5000.

Configuring GPUS for General-Purpose Computing

[0391] The following figures set forth, without limitation, exemplary architectures for compiling and executing compute source code, in accordance with at least one embodiment.

[0392] FIG. 51A illustrates a system 5100 configured to compile and execute CUDA source code 5110 using different types of processing units, in accordance with at least one embodiment. In at least one embodiment, system 5100 includes, without limitation, CUDA source code 5110, a CUDA compiler 5150, host executable code 5170(1), host executable code 5170(2), CUDA device executable code 5184, a CPU 5190, a CUDA-enabled GPU 5194, a GPU 5192, a CUDA to HIP translation tool 5120, HIP source code 5130, a HIP compiler driver 5140, an HCC 5160, and HCC device executable code 5182.

[0393] In at least one embodiment, CUDA source code 5110 is a collection of human-readable code in a CUDA programming language. In at least one embodiment, CUDA code is human-readable code in a CUDA programming language. In at least one embodiment, a CUDA programming language is an extension of the C++ programming language that includes, without limitation, mechanisms to define device code and distinguish between device code and host code. In at least one embodiment, device code is source code that, after compilation, is executable in parallel on a device. In at least one embodiment, a device may be a processor that is optimized for parallel instruction processing, such as CUDA-enabled GPU 5190, GPU 51192, or another GPGPU, etc. In at least one embodiment, host code is source code that, after compilation, is executable on a host. In at least one embodiment, a host is a processor that is optimized for sequential instruction processing, such as CPU 5190.

[0394] In at least one embodiment, CUDA source code 5110 includes, without limitation, any number (including zero) of global functions 5112, any number (including zero) of device functions 5114, any number (including zero) of host functions 5116, and any number (including zero) of host/device functions 5118. In at least one embodiment, global functions 5112, device functions 5114, host functions 5116, and host/device functions 5118 may be mixed in CUDA source code 5110. In at least one embodiment, each of global functions 5112 is executable on a device and callable from a host. In at least one embodiment, one or

more of global functions **5112** may therefore act as entry points to a device. In at least one embodiment, each of global functions **5112** is a kernel. In at least one embodiment and in a technique known as dynamic parallelism, one or more of global functions **5112** defines a kernel that is executable on a device and callable from such a device. In at least one embodiment, a kernel is executed N (where N is any positive integer) times in parallel by N different threads on a device during execution.

[0395] In at least one embodiment, each of device functions **5114** is executed on a device and callable from such a device only. In at least one embodiment, each of host functions **5116** is executed on a host and callable from such a host only. In at least one embodiment, each of host/device functions **5116** defines both a host version of a function that is executable on a host and callable from such a host only and a device version of the function that is executable on a device and callable from such a device only.

[0396] In at least one embodiment, CUDA source code **5110** may also include, without limitation, any number of calls to any number of functions that are defined via a CUDA runtime API **5102**. In at least one embodiment, CUDA runtime API **5102** may include, without limitation, any number of functions that execute on a host to allocate and deallocate device memory, transfer data between host memory and device memory, manage systems with multiple devices, etc. In at least one embodiment, CUDA source code **5110** may also include any number of calls to any number of functions that are specified in any number of other CUDA APIs. In at least one embodiment, a CUDA API may be any API that is designed for use by CUDA code. In at least one embodiment, CUDA APIs include, without limitation, CUDA runtime API **5102**, a CUDA driver API, APIs for any number of CUDA libraries, etc. In at least one embodiment and relative to CUDA runtime API **5102**, a CUDA driver API is a lower-level API but provides finer-grained control of a device. In at least one embodiment, examples of CUDA libraries include, without limitation, cuBLAS, cuFFT, cuRAND, cuDNN, etc.

[0397] In at least one embodiment, CUDA compiler **5150** compiles input CUDA code (e.g., CUDA source code **5110**) to generate host executable code **5170(1)** and CUDA device executable code **5184**. In at least one embodiment, CUDA compiler **5150** is NVCC. In at least one embodiment, host executable code **5170(1)** is a compiled version of host code included in input source code that is executable on CPU **5190**. In at least one embodiment, CPU **5190** may be any processor that is optimized for sequential instruction processing.

[0398] In at least one embodiment, CUDA device executable code **5184** is a compiled version of device code included in input source code that is executable on CUDA-enabled GPU **5194**. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, binary code. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, IR code, such as PTX code, that is further compiled at runtime into binary code for a specific target device (e.g., CUDA-enabled GPU **5194**) by a device driver. In at least one embodiment, CUDA-enabled GPU **5194** may be any processor that is optimized for parallel instruction processing and that supports CUDA. In at least one embodiment, CUDA-enabled GPU **5194** is developed by NVIDIA Corporation of Santa Clara, CA.

[0399] In at least one embodiment, CUDA to HIP translation tool **5120** is configured to translate CUDA source code **5110** to functionally similar HIP source code **5130**. In a least one embodiment, HIP source code **5130** is a collection of human-readable code in a HIP programming language. In at least one embodiment, HIP code is human-readable code in a HIP programming language. In at least one embodiment, a HIP programming language is an extension of the C++ programming language that includes, without limitation, functionally similar versions of CUDA mechanisms to define device code and distinguish between device code and host code. In at least one embodiment, a HIP programming language may include a subset of functionality of a CUDA programming language. In at least one embodiment, for example, a HIP programming language includes, without limitation, mechanism(s) to define global functions **5112**, but such a HIP programming language may lack support for dynamic parallelism and therefore global functions **5112** defined in HIP code may be callable from a host only.

[0400] In at least one embodiment, HIP source code **5130** includes, without limitation, any number (including zero) of global functions **5112**, any number (including zero) of device functions **5114**, any number (including zero) of host functions **5116**, and any number (including zero) of host/device functions **5118**. In at least one embodiment, HIP source code **5130** may also include any number of calls to any number of functions that are specified in a HIP runtime API **5132**. In at least one embodiment, HIP runtime API **5132** includes, without limitation, functionally similar versions of a subset of functions included in CUDA runtime API **5102**. In at least one embodiment, HIP source code **5130** may also include any number of calls to any number of functions that are specified in any number of other HIP APIs. In at least one embodiment, a HIP API may be any API that is designed for use by HIP code and/or ROCm. In at least one embodiment, HIP APIs include, without limitation, HIP runtime API **5132**, a HIP driver API, APIs for any number of HIP libraries, APIs for any number of ROCm libraries, etc.

[0401] In at least one embodiment, CUDA to HIP translation tool **5120** converts each kernel call in CUDA code from a CUDA syntax to a HIP syntax and converts any number of other CUDA calls in CUDA code to any number of other functionally similar HIP calls. In at least one embodiment, a CUDA call is a call to a function specified in a CUDA API, and a HIP call is a call to a function specified in a HIP API. In at least one embodiment, CUDA to HIP translation tool **5120** converts any number of calls to functions specified in CUDA runtime API **5102** to any number of calls to functions specified in HIP runtime API **5132**.

[0402] In at least one embodiment, CUDA to HIP translation tool **5120** is a tool known as hipify-perl that executes a text-based translation process. In at least one embodiment, CUDA to HIP translation tool **5120** is a tool known as hipify-clang that, relative to hipify-perl, executes a more complex and more robust translation process that involves parsing CUDA code using clang (a compiler front-end) and then translating resulting symbols. In at least one embodiment, properly converting CUDA code to HIP code may require modifications (e.g., manual edits) in addition to those performed by CUDA to HIP translation tool **5120**.

[0403] In at least one embodiment, HIP compiler driver **5140** is a front end that determines a target device **5146** and

then configures a compiler that is compatible with target device **5146** to compile HIP source code **5130**. In at least one embodiment, target device **5146** is a processor that is optimized for parallel instruction processing. In at least one embodiment, HIP compiler driver **5140** may determine target device **5146** in any technically feasible fashion.

[0404] In at least one embodiment, if target device **5146** is compatible with CUDA (e.g., CUDA-enabled GPU **5194**), then HIP compiler driver **5140** generates a HIP/NVCC compilation command **5142**. In at least one embodiment and as described in greater detail in conjunction with FIG. **51B**, HIP/NVCC compilation command **5142** configures CUDA compiler **5150** to compile HIP source code **5130** using, without limitation, a HIP to CUDA translation header and a CUDA runtime library. In at least one embodiment and in response to HIP/NVCC compilation command **5142**, CUDA compiler **5150** generates host executable code **5170(1)** and CUDA device executable code **5184**.

[0405] In at least one embodiment, if target device **5146** is not compatible with CUDA, then HIP compiler driver **5140** generates a HIP/HCC compilation command **5144**. In at least one embodiment and as described in greater detail in conjunction with FIG. **51C**, HIP/HCC compilation command **5144** configures HCC **5160** to compile HIP source code **5130** using, without limitation, an HCC header and a HIP/HCC runtime library. In at least one embodiment and in response to HIP/HCC compilation command **5144**, HCC **5160** generates host executable code **5170(2)** and HCC device executable code **5182**. In at least one embodiment, HCC device executable code **5182** is a compiled version of device code included in HIP source code **5130** that is executable on GPU **5192**. In at least one embodiment, GPU **5192** may be any processor that is optimized for parallel instruction processing, is not compatible with CUDA, and is compatible with HCC. In at least one embodiment, GPU **5192** is developed by AMD Corporation of Santa Clara, CA. In at least one embodiment GPU, **5192** is a non-CUDA-enabled GPU **5192**.

[0406] For explanatory purposes only, three different flows that may be implemented in at least one embodiment to compile CUDA source code **5110** for execution on CPU **5190** and different devices are depicted in FIG. **51A**. In at least one embodiment, a direct CUDA flow compiles CUDA source code **5110** for execution on CPU **5190** and CUDA-enabled GPU **5194** without translating CUDA source code **5110** to HIP source code **5130**. In at least one embodiment, an indirect CUDA flow translates CUDA source code **5110** to HIP source code **5130** and then compiles HIP source code **5130** for execution on CPU **5190** and CUDA-enabled GPU **5194**. In at least one embodiment, a CUDA/HCC flow translates CUDA source code **5110** to HIP source code **5130** and then compiles HIP source code **5130** for execution on CPU **5190** and GPU **5192**.

[0407] A direct CUDA flow that may be implemented in at least one embodiment is depicted via dashed lines and a series of bubbles annotated A1-A3. In at least one embodiment and as depicted with bubble annotated A1, CUDA compiler **5150** receives CUDA source code **5110** and a CUDA compile command **5148** that configures CUDA compiler **5150** to compile CUDA source code **5110**. In at least one embodiment, CUDA source code **5110** used in a direct CUDA flow is written in a CUDA programming language that is based on a programming language other than C++ (e.g., C, Fortran, Python, Java, etc.). In at least one

embodiment and in response to CUDA compile command **5148**, CUDA compiler **5150** generates host executable code **5170(1)** and CUDA device executable code **5184** (depicted with bubble annotated A2). In at least one embodiment and as depicted with bubble annotated A3, host executable code **5170(1)** and CUDA device executable code **5184** may be executed on, respectively, CPU **5190** and CUDA-enabled GPU **5194**. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, binary code. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, PTX code and is further compiled into binary code for a specific target device at runtime.

[0408] An indirect CUDA flow that may be implemented in at least one embodiment is depicted via dotted lines and a series of bubbles annotated B1-B6. In at least one embodiment and as depicted with bubble annotated B1, CUDA to HIP translation tool **5120** receives CUDA source code **5110**. In at least one embodiment and as depicted with bubble annotated B2, CUDA to HIP translation tool **5120** translates CUDA source code **5110** to HIP source code **5130**. In at least one embodiment and as depicted with bubble annotated B3, HIP compiler driver **5140** receives HIP source code **5130** and determines that target device **5146** is CUDA-enabled.

[0409] In at least one embodiment and as depicted with bubble annotated B4, HIP compiler driver **5140** generates HIP/NVCC compilation command **5142** and transmits both HIP/NVCC compilation command **5142** and HIP source code **5130** to CUDA compiler **5150**. In at least one embodiment and as described in greater detail in conjunction with FIG. **51B**, HIP/NVCC compilation command **5142** configures CUDA compiler **5150** to compile HIP source code **5130** using, without limitation, a HIP to CUDA translation header and a CUDA runtime library. In at least one embodiment and in response to HIP/NVCC compilation command **5142**, CUDA compiler **5150** generates host executable code **5170(1)** and CUDA device executable code **5184** (depicted with bubble annotated B5). In at least one embodiment and as depicted with bubble annotated B6, host executable code **5170(1)** and CUDA device executable code **5184** may be executed on, respectively, CPU **5190** and CUDA-enabled GPU **5194**. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, binary code. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, PTX code and is further compiled into binary code for a specific target device at runtime.

[0410] A CUDA/HCC flow that may be implemented in at least one embodiment is depicted via solid lines and a series of bubbles annotated C1-C6. In at least one embodiment and as depicted with bubble annotated C1, CUDA to HIP translation tool **5120** receives CUDA source code **5110**. In at least one embodiment and as depicted with bubble annotated C2, CUDA to HIP translation tool **5120** translates CUDA source code **5110** to HIP source code **5130**. In at least one embodiment and as depicted with bubble annotated C3, HIP compiler driver **5140** receives HIP source code **5130** and determines that target device **5146** is not CUDA-enabled.

[0411] In at least one embodiment, HIP compiler driver **5140** generates HIP/HCC compilation command **5144** and transmits both HIP/HCC compilation command **5144** and HIP source code **5130** to HCC **5160** (depicted with bubble annotated C4). In at least one embodiment and as described in greater detail in conjunction with FIG. **51C**, HIP/HCC

compilation command **5144** configures HCC **5160** to compile HIP source code **5130** using, without limitation, an HCC header and a HIP/HCC runtime library. In at least one embodiment and in response to HIP/HCC compilation command **5144**, HCC **5160** generates host executable code **5170(2)** and HCC device executable code **5182** (depicted with bubble annotated C5). In at least one embodiment and as depicted with bubble annotated C6, host executable code **5170(2)** and HCC device executable code **5182** may be executed on, respectively, CPU **5190** and GPU **5192**.

[0412] In at least one embodiment, after CUDA source code **5110** is translated to HIP source code **5130**, HIP compiler driver **5140** may subsequently be used to generate executable code for either CUDA-enabled GPU **5194** or GPU **5192** without re-executing CUDA to HIP translation tool **5120**. In at least one embodiment, CUDA to HIP translation tool **5120** translates CUDA source code **5110** to HIP source code **5130** that is then stored in memory. In at least one embodiment, HIP compiler driver **5140** then configures HCC **5160** to generate host executable code **5170(2)** and HCC device executable code **5182** based on HIP source code **5130**. In at least one embodiment, HIP compiler driver **5140** subsequently configures CUDA compiler **5150** to generate host executable code **5170(1)** and CUDA device executable code **5184** based on stored HIP source code **5130**.

[0413] FIG. 51B illustrates a system **5104** configured to compile and execute CUDA source code **5110** of FIG. 51A using CPU **5190** and CUDA-enabled GPU **5194**, in accordance with at least one embodiment. In at least one embodiment, system **5104** includes, without limitation, CUDA source code **5110**, CUDA to HIP translation tool **5120**, HIP source code **5130**, HIP compiler driver **5140**, CUDA compiler **5150**, host executable code **5170(1)**, CUDA device executable code **5184**, CPU **5190**, and CUDA-enabled GPU **5194**.

[0414] In at least one embodiment and as described previously herein in conjunction with FIG. 51A, CUDA source code **5110** includes, without limitation, any number (including zero) of global functions **5112**, any number (including zero) of device functions **5114**, any number (including zero) of host functions **5116**, and any number (including zero) of host/device functions **5118**. In at least one embodiment, CUDA source code **5110** also includes, without limitation, any number of calls to any number of functions that are specified in any number of CUDA APIs.

[0415] In at least one embodiment, CUDA to HIP translation tool **5120** translates CUDA source code **5110** to HIP source code **5130**. In at least one embodiment, CUDA to HIP translation tool **5120** converts each kernel call in CUDA source code **5110** from a CUDA syntax to a HIP syntax and converts any number of other CUDA calls in CUDA source code **5110** to any number of other functionally similar HIP calls.

[0416] In at least one embodiment, HIP compiler driver **5140** determines that target device **5146** is CUDA-enabled and generates HIP/NVCC compilation command **5142**. In at least one embodiment, HIP compiler driver **5140** then configures CUDA compiler **5150** via HIP/NVCC compilation command **5142** to compile HIP source code **5130**. In at least one embodiment, HIP compiler driver **5140** provides access to a HIP to CUDA translation header **5152** as part of configuring CUDA compiler **5150**. In at least one embodiment, HIP to CUDA translation header **5152** translates any number of mechanisms (e.g., functions) specified in any

number of HIP APIs to any number of mechanisms specified in any number of CUDA APIs. In at least one embodiment, CUDA compiler **5150** uses HIP to CUDA translation header **5152** in conjunction with a CUDA runtime library **5154** corresponding to CUDA runtime API **5102** to generate host executable code **5170(1)** and CUDA device executable code **5184**. In at least one embodiment, host executable code **5170(1)** and CUDA device executable code **5184** may then be executed on, respectively, CPU **5190** and CUDA-enabled GPU **5194**. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, binary code. In at least one embodiment, CUDA device executable code **5184** includes, without limitation, PTX code and is further compiled into binary code for a specific target device at runtime.

[0417] FIG. 51C illustrates a system **5106** configured to compile and execute CUDA source code **5110** of FIG. 51A using CPU **5190** and non-CUDA-enabled GPU **5192**, in accordance with at least one embodiment. In at least one embodiment, system **5106** includes, without limitation, CUDA source code **5110**, CUDA to HIP translation tool **5120**, HIP source code **5130**, HIP compiler driver **5140**, HCC **5160**, host executable code **5170(2)**, HCC device executable code **5182**, CPU **5190**, and GPU **5192**.

[0418] In at least one embodiment, at least one component shown or described with respect to FIGS. 51A-51C is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, system **5106** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0419] In at least one embodiment and as described previously herein in conjunction with FIG. 51A, CUDA source code **5110** includes, without limitation, any number (including zero) of global functions **5112**, any number (including zero) of device functions **5114**, any number (including zero) of host functions **5116**, and any number (including zero) of host/device functions **5118**. In at least one embodiment, CUDA source code **5110** also includes, without limitation, any number of calls to any number of functions that are specified in any number of CUDA APIs.

[0420] In at least one embodiment, CUDA to HIP translation tool **5120** translates CUDA source code **5110** to HIP source code **5130**. In at least one embodiment, CUDA to HIP translation tool **5120** converts each kernel call in CUDA source code **5110** from a CUDA syntax to a HIP syntax and converts any number of other CUDA calls in source code **5110** to any number of other functionally similar HIP calls.

[0421] In at least one embodiment, HIP compiler driver **5140** subsequently determines that target device **5146** is not CUDA-enabled and generates HIP/HCC compilation command **5144**. In at least one embodiment, HIP compiler driver **5140** then configures HCC **5160** to execute HIP/HCC compilation command **5144** to compile HIP source code **5130**. In at least one embodiment, HIP/HCC compilation command **5144** configures HCC **5160** to use, without limitation, a HIP/HCC runtime library **5158** and an HCC header **5156** to generate host executable code **5170(2)** and HCC device executable code **5182**. In at least one embodiment, HIP/HCC runtime library **5158** corresponds to HIP runtime API **5132**. In at least one embodiment, HCC header **5156** includes, without limitation, any number and type of

interoperability mechanisms for HIP and HCC. In at least one embodiment, host executable code **5170(2)** and HCC device executable code **5182** may be executed on, respectively, CPU **5190** and GPU **5192**.

[0422] FIG. **52** illustrates an exemplary kernel translated by CUDA-to-HIP translation tool **5120** of FIG. **51C**, in accordance with at least one embodiment. In at least one embodiment, CUDA source code **5110** partitions an overall problem that a given kernel is designed to solve into relatively coarse sub-problems that can independently be solved using thread blocks. In at least one embodiment, each thread block includes, without limitation, any number of threads. In at least one embodiment, each sub-problem is partitioned into relatively fine pieces that can be solved cooperatively in parallel by threads within a thread block. In at least one embodiment, threads within a thread block can cooperate by sharing data through shared memory and by synchronizing execution to coordinate memory accesses.

[0423] In at least one embodiment, at least one component shown or described with respect to FIG. **52** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, a kernel of FIG. **52** comprises instructions that, when launched, performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0424] In at least one embodiment, CUDA source code **5110** organizes thread blocks associated with a given kernel into a one-dimensional, a two-dimensional, or a three-dimensional grid of thread blocks. In at least one embodiment, each thread block includes, without limitation, any number of threads, and a grid includes, without limitation, any number of thread blocks.

[0425] In at least one embodiment, a kernel is a function in device code that is defined using a “_global_” declaration specifier. In at least one embodiment, the dimension of a grid that executes a kernel for a given kernel call and associated streams are specified using a CUDA kernel launch syntax **5210**. In at least one embodiment, CUDA kernel launch syntax **5210** is specified as “KernelName <<<GridSize, BlockSize, SharedMemorySize, Stream>>> (KernelArguments);”. In at least one embodiment, an execution configuration syntax is a “<<<. . . >>>” construct that is inserted between a kernel name (“KernelName”) and a parenthesized list of kernel arguments (“KernelArguments”). In at least one embodiment, CUDA kernel launch syntax **5210** includes, without limitation, a CUDA launch function syntax instead of an execution configuration syntax.

[0426] In at least one embodiment, “GridSize” is of a type dim3 and specifies the dimension and size of a grid. In at least one embodiment, type dim3 is a CUDA-defined structure that includes, without limitation, unsigned integers x, y, and z. In at least one embodiment, if z is not specified, then z defaults to one. In at least one embodiment, if y is not specified, then y defaults to one. In at least one embodiment, the number of thread blocks in a grid is equal to the product of GridSize.x, GridSize.y, and GridSize.z. In at least one embodiment, “BlockSize” is of type dim3 and specifies the dimension and size of each thread block. In at least one embodiment, the number of threads per thread block is equal to the product of BlockSize.x, BlockSize.y, and BlockSize.z. In at least one embodiment, each thread that executes a

kernel is given a unique thread ID that is accessible within the kernel through a built-in variable (e.g., “threadIdx”).

[0427] In at least one embodiment and with respect to CUDA kernel launch syntax **5210**, “SharedMemorySize” is an optional argument that specifies a number of bytes in a shared memory that is dynamically allocated per thread block for a given kernel call in addition to statically allocated memory. In at least one embodiment and with respect to CUDA kernel launch syntax **5210**, SharedMemorySize defaults to zero. In at least one embodiment and with respect to CUDA kernel launch syntax **5210**, “Stream” is an optional argument that specifies an associated stream and defaults to zero to specify a default stream. In at least one embodiment, a stream is a sequence of commands (possibly issued by different host threads) that execute in order. In at least one embodiment, different streams may execute commands out of order with respect to one another or concurrently.

[0428] In at least one embodiment, CUDA source code **5110** includes, without limitation, a kernel definition for an exemplary kernel “MatAdd” and a main function. In at least one embodiment, main function is host code that executes on a host and includes, without limitation, a kernel call that causes kernel MatAdd to execute on a device. In at least one embodiment and as shown, kernel MatAdd adds two matrices A and B of size N×N, where N is a positive integer, and stores the result in a matrix C. In at least one embodiment, main function defines a threadsPerBlock variable as 16 by 16 and a numBlocks variable as N/16 by N/16. In at least one embodiment, main function then specifies kernel call “MatAdd<<<numBlocks, threadsPerBlock>>> (A, B, C);”. In at least one embodiment and as per CUDA kernel launch syntax **5210**, kernel MatAdd is executed using a grid of thread blocks having a dimension N/16 by N/16, where each thread block has a dimension of 16 by 16. In at least one embodiment, each thread block includes 256 threads, a grid is created with enough blocks to have one thread per matrix element, and each thread in such a grid executes kernel MatAdd to perform one pair-wise addition.

[0429] In at least one embodiment, while translating CUDA source code **5110** to HIP source code **5130**, CUDA to HIP translation tool **5120** translates each kernel call in CUDA source code **5110** from CUDA kernel launch syntax **5210** to a HIP kernel launch syntax **5220** and converts any number of other CUDA calls in source code **5110** to any number of other functionally similar HIP calls. In at least one embodiment, HIP kernel launch syntax **5220** is specified as “hipLaunchKernelGGL (KernelName, GridSize, BlockSize, SharedMemorySize, Stream, KernelArguments);”. In at least one embodiment, each of KernelName, GridSize, BlockSize, SharedMemorySize, Stream, and KernelArguments has the same meaning in HIP kernel launch syntax **5220** as in CUDA kernel launch syntax **5210** (described previously herein). In at least one embodiment, arguments SharedMemorySize and Stream are required in HIP kernel launch syntax **5220** and are optional in CUDA kernel launch syntax **5210**.

[0430] In at least one embodiment, a portion of HIP source code **5130** depicted in FIG. **52** is identical to a portion of CUDA source code **5110** depicted in FIG. **52** except for a kernel call that causes kernel MatAdd to execute on a device. In at least one embodiment, kernel MatAdd is defined in HIP source code **5130** with the same “global” declaration specifier with which kernel MatAdd is defined in CUDA source code **5110**. In at least one embodiment, a

kernel call in HIP source code **5130** is “hipLaunchKernel-GGL (MatAdd, numBlocks, threadsPerBlock, 0, 0, A, B, C);”, while a corresponding kernel call in CUDA source code **5110** is “MatAdd<<<numBlocks, threadsPerBlock>>> (A, B, C);”.

[0431] FIG. **53** illustrates non-CUDA-enabled GPU **5192** of FIG. **51C** in greater detail, in accordance with at least one embodiment. In at least one embodiment, GPU **5192** is developed by AMD corporation of Santa Clara. In at least one embodiment, GPU **5192** can be configured to perform compute operations in a highly-parallel fashion. In at least one embodiment, GPU **5192** is configured to execute graphics pipeline operations such as draw commands, pixel operations, geometric computations, and other operations associated with rendering an image to a display. In at least one embodiment, GPU **5192** is configured to execute operations unrelated to graphics. In at least one embodiment, GPU **5192** is configured to execute both operations related to graphics and operations unrelated to graphics. In at least one embodiment, GPU **5192** can be configured to execute device code included in HIP source code **5130**.

[0432] In at least one embodiment, at least one component shown or described with respect to FIG. **53** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, non-CUDA-enabled GPU **5192** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0433] In at least one embodiment, GPU **5192** includes, without limitation, any number of programmable processing units **5320**, a command processor **5310**, an L2 cache **5322**, memory controllers **5370**, DMA engines **5380(1)**, system memory controllers **5382**, DMA engines **5380(2)**, and GPU controllers **5384**. In at least one embodiment, each programmable processing unit **5320** includes, without limitation, a workload manager **5330** and any number of compute units **5340**. In at least one embodiment, command processor **5310** reads commands from one or more command queues (not shown) and distributes commands to workload managers **5330**. In at least one embodiment, for each programmable processing unit **5320**, associated workload manager **5330** distributes work to compute units **5340** included in programmable processing unit **5320**. In at least one embodiment, each compute unit **5340** may execute any number of thread blocks, but each thread block executes on a single compute unit **5340**. In at least one embodiment, a workgroup is a thread block.

[0434] In at least one embodiment, each compute unit **5340** includes, without limitation, any number of SIMD units **5350** and a shared memory **5360**. In at least one embodiment, each SIMD unit **5350** implements a SIMD architecture and is configured to perform operations in parallel. In at least one embodiment, each SIMD unit **5350** includes, without limitation, a vector ALU **5352** and a vector register file **5354**. In at least one embodiment, each SIMD unit **5350** executes a different warp. In at least one embodiment, a warp is a group of threads (e.g., 16 threads), where each thread in the warp belongs to a single thread block and is configured to process a different set of data based on a single set of instructions. In at least one embodiment, predication can be used to disable one or more threads in a

warp. In at least one embodiment, a lane is a thread. In at least one embodiment, a work item is a thread. In at least one embodiment, a wavefront is a warp. In at least one embodiment, different wavefronts in a thread block may synchronize together and communicate via shared memory **5360**. In at least one embodiment, compute unit **5340** includes one or more distributed shared memories (or distributed shared memory) that enable direct streaming multiprocessor (SM) to streaming multiple processor (SM) for operations related to loading, storing, and performing atomics across multiple SM shared memory blocks. compute unit **5340** includes one or more cluster distributed shared memories (DSMEM), which are blocks of memory within a cluster that enabled to access each other's shared memory directly.

[0435] In at least one embodiment, programmable processing units **5320** are referred to as “shader engines.” In at least one embodiment, each programmable processing unit **5320** includes, without limitation, any amount of dedicated graphics hardware in addition to compute units **5340**. In at least one embodiment, each programmable processing unit **5320** includes, without limitation, any number (including zero) of geometry processors, any number (including zero) of rasterizers, any number (including zero) of render back ends, workload manager **5330**, and any number of compute units **5340**.

[0436] In at least one embodiment, compute units **5340** share L2 cache **5322**. In at least one embodiment, L2 cache **5322** is partitioned. In at least one embodiment, a GPU memory **5390** is accessible by all compute units **5340** in GPU **5192**. In at least one embodiment, memory controllers **5370** and system memory controllers **5382** facilitate data transfers between GPU **5192** and a host, and DMA engines **5380(1)** enable asynchronous memory transfers between GPU **5192** and such a host. In at least one embodiment, memory controllers **5370** and GPU controllers **5384** facilitate data transfers between GPU **5192** and other GPUs **5192**, and DMA engines **5380(2)** enable asynchronous memory transfers between GPU **5192** and other GPUs **5192**.

[0437] In at least one embodiment, GPU **5192** includes, without limitation, any amount and type of system interconnect that facilitates data and control transmissions across any number and type of directly or indirectly linked components that may be internal or external to GPU **5192**. In at least one embodiment, GPU **5192** includes, without limitation, any number and type of I/O interfaces (e.g., PCIe) that are coupled to any number and type of peripheral devices. In at least one embodiment, GPU **5192** may include, without limitation, any number (including zero) of display engines and any number (including zero) of multimedia engines. In at least one embodiment, GPU **5192** implements a memory subsystem that includes, without limitation, any amount and type of memory controllers (e.g., memory controllers **5370** and system memory controllers **5382**) and memory devices (e.g., shared memories **5360**) that may be dedicated to one component or shared among multiple components. In at least one embodiment, GPU **5192** implements a cache subsystem that includes, without limitation, one or more cache memories (e.g., L2 cache **5322**) that may each be private to or shared between any number of components (e.g., SIMD units **5350**, compute units **5340**, and programmable processing units **5320**).

[0438] FIG. **54** illustrates how threads of an exemplary CUDA grid **5420** are mapped to different compute units **5340** of FIG. **53**, in accordance with at least one embodi-

ment. In at least one embodiment and for explanatory purposes only, grid **5420** has a GridSize of BX by BY by 1 and a BlockSize of TX by TY by 1. In at least one embodiment, grid **5420** therefore includes, without limitation, (BX*BY) thread blocks **5430** and each thread block **5430** includes, without limitation, (TX*TY) threads **5440**. Threads **5440** are depicted in FIG. **54** as squiggly arrows.

[0439] In at least one embodiment, at least one component shown or described with respect to FIG. **54** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, CUDA grid **5420** performs one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0440] In at least one embodiment, grid **5420** is mapped to programmable processing unit **5320(1)** that includes, without limitation, compute units **5340(1)-5340(C)**. In at least one embodiment and as shown, (BJ*BY) thread blocks **5430** are mapped to compute unit **5340(1)**, and the remaining thread blocks **5430** are mapped to compute unit **5340(2)**. In at least one embodiment, each thread block **5430** may include, without limitation, any number of warps, and each warp is mapped to a different SIMD unit **5350** of FIG. **53**.

[0441] In at least one embodiment, warps in a given thread block **5430** may synchronize together and communicate through shared memory **5360** included in associated compute unit **5340**. For example and in at least one embodiment, warps in thread block **5430(BJ,1)** can synchronize together and communicate through shared memory **5360(1)**. For example and in at least one embodiment, warps in thread block **5430(BJ+1,1)** can synchronize together and communicate through shared memory **5360(2)**.

[0442] FIG. **55** illustrates how to migrate existing CUDA code to Data Parallel C++ code, in accordance with at least one embodiment. Data Parallel C++ (DPC++) may refer to an open, standards-based alternative to single-architecture proprietary languages that allows developers to reuse code across hardware targets (CPUs and accelerators such as GPUs and FPGAs) and also perform custom tuning for a specific accelerator. DPC++ use similar and/or identical C and C++ constructs in accordance with ISO C++ which developers may be familiar with. DPC++ incorporates standard SYCL from The Khronos Group to support data parallelism and heterogeneous programming. SYCL refers to a cross-platform abstraction layer that builds on underlying concepts, portability and efficiency of OpenCL that enables code for heterogeneous processors to be written in a “single-source” style using standard C++. SYCL may enable single source development where C++ template functions can contain both host and device code to construct complex algorithms that use OpenCL acceleration, and then re-use them throughout their source code on different types of data.

[0443] In at least one embodiment, at least one component shown or described with respect to FIG. **55** is used to implement techniques and/or functions described in connection with FIGS. **1-23**. In at least one embodiment, CUDA code migrated to Data Parallel C++ code, when performed, causes one or more processors to perform one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one

or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. **15**, and as otherwise described herein.

[0444] In at least one embodiment, a DPC++ compiler is used to compile DPC++ source code which can be deployed across diverse hardware targets. In at least one embodiment, a DPC++ compiler is used to generate DPC++ applications that can be deployed across diverse hardware targets and a DPC++ compatibility tool can be used to migrate CUDA applications to a multiplatform program in DPC++. In at least one embodiment, a DPC++ base tool kit includes a DPC++ compiler to deploy applications across diverse hardware targets; a DPC++ library to increase productivity and performance across CPUs, GPUs, and FPGAs; a DPC++ compatibility tool to migrate CUDA applications to multiplatform applications; and any suitable combination thereof.

[0445] In at least one embodiment, a DPC++ programming model is utilized to simply one or more aspects relating to programming CPUs and accelerators by using modern C++ features to express parallelism with a programming language called Data Parallel C++. DPC++ programming language may be utilized to code reuse for hosts (e.g., a CPU) and accelerators (e.g., a GPU or FPGA) using a single source language, with execution and memory dependencies being clearly communicated. Mappings within DPC++ code can be used to transition an application to run on a hardware or set of hardware devices that best accelerates a workload. A host may be available to simplify development and debugging of device code, even on platforms that do not have an accelerator available.

[0446] In at least one embodiment, CUDA source code **5500** is provided as an input to a DPC++ compatibility tool **5502** to generate human readable DPC++ **5504**. In at least one embodiment, human readable DPC++ **5504** includes inline comments generated by DPC++ compatibility tool **5502** that guides a developer on how and/or where to modify DPC++ code to complete coding and tuning to desired performance **5506**, thereby generating DPC++ source code **5508**.

[0447] In at least one embodiment, CUDA source code **5500** is or includes a collection of human-readable source code in a CUDA programming language. In at least one embodiment, CUDA source code **5500** is human-readable source code in a CUDA programming language. In at least one embodiment, a CUDA programming language is an extension of the C++ programming language that includes, without limitation, mechanisms to define device code and distinguish between device code and host code. In at least one embodiment, device code is source code that, after compilation, is executable on a device (e.g., GPU or FPGA) and may include or more parallelizable workflows that can be executed on one or more processor cores of a device. In at least one embodiment, a device may be a processor that is optimized for parallel instruction processing, such as CUDA-enabled GPU, GPU, or another GPGPU, etc. In at least one embodiment, host code is source code that, after compilation, is executable on a host. In at least one embodiment, some or all of host code and device code can be executed in parallel across a CPU and GPU/FPGA. In at least one embodiment, a host is a processor that is optimized for sequential instruction processing, such as CPU. CUDA source code **5500** described in connection with FIG. **55** may be in accordance with those discussed elsewhere in this document.

[0448] In at least one embodiment, DPC++ compatibility tool **5502** refers to an executable tool, program, application, or any other suitable type of tool that is used to facilitate migration of CUDA source code **5500** to DPC++ source code **5508**. In at least one embodiment, DPC++ compatibility tool **5502** is a command-line-based code migration tool available as part of a DPC++ tool kit that is used to port existing CUDA sources to DPC++. In at least one embodiment, DPC++ compatibility tool **5502** converts some or all source code of a CUDA application from CUDA to DPC++ and generates a resulting file that is written at least partially in DPC++, referred to as human readable DPC++ **5504**. In at least one embodiment, human readable DPC++ **5504** includes comments that are generated by DPC++ compatibility tool **5502** to indicate where user intervention may be necessary. In at least one embodiment, user intervention is necessary when CUDA source code **5500** calls a CUDA API that has no analogous DPC++ API; other examples where user intervention is required are discussed later in greater detail.

[0449] In at least one embodiment, a workflow for migrating CUDA source code **5500** (e.g., application or portion thereof) includes creating one or more compilation database files; migrating CUDA to DPC++ using a DPC++ compatibility tool **5502**; completing migration and verifying correctness, thereby generating DPC++ source code **5508**; and compiling DPC++ source code **5508** with a DPC++ compiler to generate a DPC++ application. In at least one embodiment, a compatibility tool provides a utility that intercepts commands used when Makefile executes and stores them in a compilation database file. In at least one embodiment, a file is stored in JSON format. In at least one embodiment, an intercept-built command converts Makefile command to a DPC compatibility command.

[0450] In at least one embodiment, intercept-build is a utility script that intercepts a build process to capture compilation options, macro defs, and include paths, and writes this data to a compilation database file. In at least one embodiment, a compilation database file is a JSON file. In at least one embodiment, DPC++ compatibility tool **5502** parses a compilation database and applies options when migrating input sources. In at least one embodiment, use of intercept-build is optional, but highly recommended for Make or CMake based environments. In at least one embodiment, a migration database includes commands, directories, and files; command may include necessary compilation flags; directory may include paths to header files; file may include paths to CUDA files.

[0451] In at least one embodiment, DPC++ compatibility tool **5502** migrates CUDA code (e.g., applications) written in CUDA to DPC++ by generating DPC++ wherever possible. In at least one embodiment, DPC++ compatibility tool **5502** is available as part of a tool kit. In at least one embodiment, a DPC++ tool kit includes an intercept-build tool. In at least one embodiment, an intercept-built tool creates a compilation database that captures compilation commands to migrate CUDA files. In at least one embodiment, a compilation database generated by an intercept-built tool is used by DPC++ compatibility tool **5502** to migrate CUDA code to DPC++. In at least one embodiment, non-CUDA C++ code and files are migrated as is. In at least one embodiment, DPC++ compatibility tool **5502** generates human readable DPC++ **5504** which may be DPC++ code that, as generated by DPC++ compatibility tool **5502**, cannot

be compiled by DPC++ compiler and requires additional plumbing for verifying portions of code that were not migrated correctly, and may involve manual intervention, such as by a developer. In at least one embodiment, DPC++ compatibility tool **5502** provides hints or tools embedded in code to help developers manually migrate additional code that could not be migrated automatically. In at least one embodiment, migration is a one-time activity for a source file, project, or application.

[0452] In at least one embodiment, DPC++ compatibility tool **5502** is able to successfully migrate all portions of CUDA code to DPC++ and there may simply be an optional step for manually verifying and tuning performance of DPC++ source code that was generated. In at least one embodiment, DPC++ compatibility tool **5502** directly generates DPC++ source code **5508** which is compiled by a DPC++ compiler without requiring or utilizing human intervention to modify DPC++ code generated by DPC++ compatibility tool **5502**. In at least one embodiment, DPC++ compatibility tool generates compile-able DPC++ code which can be optionally tuned by a developer for performance, readability, maintainability, other various considerations; or any combination thereof.

[0453] In at least one embodiment, one or more CUDA source files are migrated to DPC++ source files at least partially using DPC++ compatibility tool **5502**. In at least one embodiment, CUDA source code includes one or more header files which may include CUDA header files. In at least one embodiment, a CUDA source file includes a <cuda.h> header file and a <stdio.h> header file which can be used to print text. In at least one embodiment, a portion of a vector addition kernel CUDA source file may be written as or related to:

```
#include <cuda.h>
#include <stdio.h>
#define VECTOR_SIZE 256
[ ] global__ void VectorAddKernel(float* A, float* B, float* C)
{
    A[threadIdx.x] = threadIdx.x + 1.0f;
    B[threadIdx.x] = threadIdx.x + 1.0f;
    C[threadIdx.x] = A[threadIdx.x] + B[threadIdx.x];
}
int main()
{
    float *d_A, *d_B, *d_C;
    cudaMalloc(&d_A, VECTOR_SIZE*sizeof(float));
    cudaMalloc(&d_B, VECTOR_SIZE*sizeof(float));
    cudaMalloc(&d_C, VECTOR_SIZE*sizeof(float));
    VectorAddKernel<<<1, VECTOR_SIZE>>>(&d_A, &d_B, &d_C);
    float Result[VECTOR_SIZE] = { };
    cudaMemcpy(Result, d_C, VECTOR_SIZE*sizeof(float),
    cudaMemcpyDeviceToHost);
    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);
    for (int i=0; i<VECTOR_SIZE; i++) {
        if (i % 16 == 0) {
            printf("\n");
        }
        printf("%f", Result[i]);
    }
    return 0;
}
```

[0454] In at least one embodiment and in connection with CUDA source file presented above, DPC++ compatibility tool **5502** parses a CUDA source code and replaces header files with appropriate DPC++ and SYCL header files. In at

least one embodiment, DPC++ header files includes helper declarations. In CUDA, there is a concept of a thread ID and correspondingly, in DPC++ or SYCL, for each element there is a local identifier.

[0455] In at least one embodiment and in connection with CUDA source file presented above, there are two vectors A and B which are initialized and a vector addition result is put into vector C as part of VectorAddKernel(). In at least one embodiment, DPC++ compatibility tool **5502** converts CUDA thread IDs used to index work elements to SYCL standard addressing for work elements via a local ID as part of migrating CUDA code to DPC++ code. In at least one embodiment, DPC++ code generated by DPC++ compatibility tool **5502** can be optimized—for example, by reducing dimensionality of an nd_item, thereby increasing memory and/or processor utilization.

[0456] In at least one embodiment and in connection with CUDA source file presented above, memory allocation is migrated. In at least one embodiment, cudaMalloc() is migrated to a unified shared memory SYCL call malloc_device() to which a device and context is passed, relying on SYCL concepts such as platform, device, context, and queue. In at least one embodiment, a SYCL platform can have multiple devices (e.g., host and GPU devices); a device may have multiple queues to which jobs can be submitted; each device may have a context; and a context may have multiple devices and manage shared memory objects.

[0457] In at least one embodiment and in connection with CUDA source file presented above, a main() function invokes or calls VectorAddKernel() to add two vectors A and B together and store result in vector C. In at least one embodiment, CUDA code to invoke VectorAddKernel() is replaced by DPC++ code to submit a kernel to a command queue for execution. In at least one embodiment, a command group handler cgh passes data, synchronization, and computation that is submitted to the queue, parallel_for is called for a number of global elements and a number of work items in that work group where VectorAddKernel() is called.

[0458] In at least one embodiment and in connection with CUDA source file presented above, CUDA calls to copy device memory and then free memory for vectors A, B, and C are migrated to corresponding DPC++ calls. In at least one embodiment, C++ code (e.g., standard ISO C++ code for printing a vector of floating point variables) is migrated as is, without being modified by DPC++ compatibility tool **5502**. In at least one embodiment, DPC++ compatibility tool **5502** modify CUDA APIs for memory setup and/or host calls to execute kernel on the acceleration device. In at least one embodiment and in connection with CUDA source file presented above, a corresponding human readable DPC++ **5504** (e.g., which can be compiled) is written as or related to:

```
#include <CL/sycl.hpp>
#include <dpct/dpct.hpp>
#define VECTOR_SIZE 256
void VectorAddKernel(float* A, float* B, float* C,
    sycl::nd_item<3> item_ct1)
{
    A[item_ct1.get_local_id(2)] = item_ct1.get_local_id(2) + 1.0f;
    B[item_ct1.get_local_id(2)] = item_ct1.get_local_id(2) + 1.0f;
    C[item_ct1.get_local_id(2)] =
        A[item_ct1.get_local_id(2)] + B[item_ct1.get_local_id(2)];
}
int main( )
```

-continued

```
{
    float *d_A, *d_B, *d_C;
    d_A = (float *)sycl::malloc_device(VECTOR_SIZE * sizeof(float),
        dpct::get_current_device( ),
        dpct::get_default_context( ));
    d_B = (float *)sycl::malloc_device(VECTOR_SIZE * sizeof(float),
        dpct::get_current_device( ),
        dpct::get_default_context( ));
    d_C = (float *)sycl::malloc_device(VECTOR_SIZE * sizeof(float),
        dpct::get_current_device( ),
        dpct::get_default_context( ));
    dpct::get_default_queue__wait( ).submit([&](sycl::handler &cgh) {
        cgh.parallel_for(
            sycl::nd_range<3>(sycl::range<3>(1, 1, 1) *
                sycl::range<3>(1, 1, VECTOR_SIZE) *
                sycl::range<3>(1, 1, VECTOR_SIZE)),
            [=](sycl::nd_item<3> item_ct1) {
                VectorAddKernel(d_A, d_B, d_C, item_ct1);
            });
    });
    float Result[VECTOR_SIZE] = { };
    dpct::get_default_queue__wait( )
        .memcpy(Result, d_C, VECTOR_SIZE * sizeof(float))
        .wait( );
    sycl::free(d_A, dpct::get_default_context( ));
    sycl::free(d_B, dpct::get_default_context( ));
    sycl::free(d_C, dpct::get_default_context( ));
    for (int i=0; i<VECTOR_SIZE; i++) {
        if (i % 16 == 0) {
            printf("\n");
        }
        printf("%f", Result[i]);
    }
    return 0;
}
```

[0459] In at least one embodiment, human readable DPC++ **5504** refers to output generated by DPC++ compatibility tool **5502** and may be optimized in one manner or another. In at least one embodiment, human readable DPC++ **5504** generated by DPC++ compatibility tool **5502** can be manually edited by a developer after migration to make it more maintainable, performance, or other considerations. In at least one embodiment, DPC++ code generated by DPC++ compatibility tool **5502** such as DPC++ disclosed can be optimized by removing repeat calls to get_current_device() and/or get_default_context() for each malloc_device() call. In at least one embodiment, DPC++ code generated above uses a 3 dimensional nd_range which can be refactored to use only a single dimension, thereby reducing memory usage. In at least one embodiment, a developer can manually edit DPC++ code generated by DPC++ compatibility tool **5502** replace uses of unified shared memory with accessors. In at least one embodiment, DPC++ compatibility tool **5502** has an option to change how it migrates CUDA code to DPC++ code. In at least one embodiment, DPC++ compatibility tool **5502** is verbose because it is using a general template to migrate CUDA code to DPC++ code that works for a large number of cases.

[0460] In at least one embodiment, a CUDA to DPC++ migration workflow includes steps to: prepare for migration using intercept-build script; perform migration of CUDA projects to DPC++ using DPC++ compatibility tool **5502**; review and edit migrated source files manually for completion and correctness; and compile final DPC++ code to generate a DPC++ application. In at least one embodiment, manual review of DPC++ source code may be required in one or more scenarios including but not limited to: migrated API does not return error code (CUDA code can return an

error code which can then be consumed by the application but SYCL uses exceptions to report errors, and therefore does not use error codes to surface errors); CUDA compute capability dependent logic is not supported by DPC++; statement could not be removed. In at least one embodiment, scenarios in which DPC++ code requires manual intervention may include, without limitation: error code logic replaced with `(*,0)` code or commented out; equivalent DPC++ API not available; CUDA compute capability-dependent logic; hardware-dependent API (`clock()`); missing features unsupported API; execution time measurement logic; handling built-in vector type conflicts; migration of cuBLAS API; and more.

[0461] In at least one embodiment, one or more techniques described herein utilize a oneAPI programming model. In at least one embodiment, a oneAPI programming model refers to a programming model for interacting with various compute accelerator architectures. In at least one embodiment, oneAPI refers to an application programming interface (API) designed to interact with various compute accelerator architectures. In at least one embodiment, a oneAPI programming model utilizes a DPC++ programming language. In at least one embodiment, a DPC++ programming language refers to a high-level language for data parallel programming productivity. In at least one embodiment, a DPC++ programming language is based at least in part on C and/or C++ programming languages. In at least one embodiment, a oneAPI programming model is a programming model such as those developed by Intel Corporation of Santa Clara, CA.

[0462] In at least one embodiment, oneAPI and/or oneAPI programming model is utilized to interact with various accelerator, GPU, processor, and/or variations thereof, architectures. In at least one embodiment, oneAPI includes a set of libraries that implement various functionalities. In at least one embodiment, oneAPI includes at least a oneAPI DPC++ library, a oneAPI math kernel library, a oneAPI data analytics library, a oneAPI deep neural network library, a oneAPI collective communications library, a oneAPI threading building blocks library, a oneAPI video processing library, and/or variations thereof.

[0463] In at least one embodiment, a oneAPI DPC++ library, also referred to as oneDPL, is a library that implements algorithms and functions to accelerate DPC++ kernel programming. In at least one embodiment, oneDPL implements one or more standard template library (STL) functions. In at least one embodiment, oneDPL implements one or more parallel STL functions. In at least one embodiment, oneDPL provides a set of library classes and functions such as parallel algorithms, iterators, function object classes, range-based API, and/or variations thereof. In at least one embodiment, oneDPL implements one or more classes and/or functions of a C++ standard library. In at least one embodiment, oneDPL implements one or more random number generator functions.

[0464] In at least one embodiment, a oneAPI math kernel library, also referred to as oneMKL, is a library that implements various optimized and parallelized routines for various mathematical functions and/or operations. In at least one embodiment, oneMKL implements one or more basic linear algebra subprograms (BLAS) and/or linear algebra package (LAPACK) dense linear algebra routines. In at least one embodiment, oneMKL implements one or more sparse BLAS linear algebra routines. In at least one embodiment,

oneMKL implements one or more random number generators (RNGs). In at least one embodiment, oneMKL implements one or more vector mathematics (VM) routines for mathematical operations on vectors. In at least one embodiment, oneMKL implements one or more Fast Fourier Transform (FFT) functions.

[0465] In at least one embodiment, a oneAPI data analytics library, also referred to as oneDAL, is a library that implements various data analysis applications and distributed computations. In at least one embodiment, oneDAL implements various algorithms for preprocessing, transformation, analysis, modeling, validation, and decision making for data analytics, in batch, online, and distributed processing modes of computation. In at least one embodiment, oneDAL implements various C++ and/or Java APIs and various connectors to one or more data sources. In at least one embodiment, oneDAL implements DPC++ API extensions to a traditional C++ interface and enables GPU usage for various algorithms.

[0466] In at least one embodiment, a oneAPI deep neural network library, also referred to as oneDNN, is a library that implements various deep learning functions. In at least one embodiment, oneDNN implements various neural network, machine learning, and deep learning functions, algorithms, and/or variations thereof.

[0467] In at least one embodiment, a oneAPI collective communications library, also referred to as oneCCL, is a library that implements various applications for deep learning and machine learning workloads. In at least one embodiment, oneCCL is built upon lower-level communication middleware, such as message passing interface (MPI) and libfabrics. In at least one embodiment, oneCCL enables a set of deep learning specific optimizations, such as prioritization, persistent operations, out of order executions, and/or variations thereof. In at least one embodiment, oneCCL implements various CPU and GPU functions.

[0468] In at least one embodiment, a oneAPI threading building blocks library, also referred to as oneTBB, is a library that implements various parallelized processes for various applications. In at least one embodiment, oneTBB is utilized for task-based, shared parallel programming on a host. In at least one embodiment, oneTBB implements generic parallel algorithms. In at least one embodiment, oneTBB implements concurrent containers. In at least one embodiment, oneTBB implements a scalable memory allocator. In at least one embodiment, oneTBB implements a work-stealing task scheduler. In at least one embodiment, oneTBB implements low-level synchronization primitives. In at least one embodiment, oneTBB is compiler-independent and usable on various processors, such as GPUs, PUs, CPUs, and/or variations thereof.

[0469] In at least one embodiment, a oneAPI video processing library, also referred to as oneVPL, is a library that is utilized for accelerating video processing in one or more applications. In at least one embodiment, oneVPL implements various video decoding, encoding, and processing functions. In at least one embodiment, one VPL implements various functions for media pipelines on CPUs, GPUs, and other accelerators. In at least one embodiment, one VPL implements device discovery and selection in media centric and video analytics workloads. In at least one embodiment, one VPL implements API primitives for zero-copy buffer sharing.

[0470] In at least one embodiment, a oneAPI programming model utilizes a DPC++ programming language. In at least one embodiment, a DPC++ programming language is a programming language that includes, without limitation, functionally similar versions of CUDA mechanisms to define device code and distinguish between device code and host code. In at least one embodiment, a DPC++ programming language may include a subset of functionality of a CUDA programming language. In at least one embodiment, one or more CUDA programming model operations are performed using a oneAPI programming model using a DPC++ programming language.

[0471] FIG. 56 is a system diagram illustrating system 5600 for interfacing with an application 5602 to process data, according to at least one embodiment. In at least one embodiment, application 5602 uses large language model (LLM) 5612 to generate output data 5620 based, at least in part, on input data 5610. In at least one embodiment, input data 5610 is a text prompt. In at least one embodiment, input data 5610 includes unstructured text. In at least one embodiment, input data 5610 includes a sequence of tokens. In at least one embodiment, a token is a portion of input data. In at least one embodiment, a token is a word. In at least one embodiment, a token is a character. In at least one embodiment, a token is a subword. In at least one embodiment, input data 5610 is formatted in Chat Markup Language (ChatML). In at least one embodiment, input data 5610 is an image. In at least one embodiment, input data 5610 is one or more video frames. In at least one embodiment, input data 5610 is any other expressive medium.

[0472] In at least one embodiment, at least one component shown or described with respect to FIG. 56 is used to implement techniques and/or functions described in connection with FIGS. 1-23. In at least one embodiment, LLM 5612 generates algorithms customized for a user to perform one or more operations to identify one or more pixels within a polygon based on one or more pixels that are adjacent to those one or more pixels and that cross one or more polygon boundaries as described in conjunction with FIG. 15, and as otherwise described herein.

[0473] In at least one embodiment, large language model 5612 comprises a deep neural network. In at least one embodiment, a deep neural network is a neural network with two or more layers. In at least one embodiment, large language model 5612 comprises a transformer model. In at least one embodiment, large language model 5612 comprises a neural network configured to perform natural language processing. In at least one embodiment, large language model 5612 is configured to process one or more sequences of data. In at least one embodiment, large language model 5612 is configured to process text. In at least one embodiment, weights and biases of a large language model 5612 are configured to process text. In at least one embodiment, large language model 5612 is configured to determine patterns in data to perform one or more natural language processing tasks. In at least one embodiment, a natural language processing task comprises text generation. In at least one embodiment, a natural language processing task comprises question answering. In at least one embodiment, performing a natural language processing task results in output data 5620.

[0474] In at least one embodiment, a processor uses input data 5610 to query retrieval database 5614. In at least one embodiment, retrieval database 5614 is a key-value store. In

at least one embodiment, retrieval database 5614 is a corpus used to train large language model 5612. In at least one embodiment, a processor uses retrieval database 5614 to provide large language model 5612 with updated information. In at least one embodiment, retrieval database 5614 comprises data from an internet source. In at least one embodiment, large language model 5612 does not use retrieval database 5614 to perform inferencing.

[0475] In at least one embodiment, an encoder encodes input data 5610 into one or more feature vectors. In at least one embodiment, an encoder encodes input data 5610 into a sentence embedding vector. In at least one embodiment, a processor uses said sentence embedding vector to perform a nearest neighbor search to generate one or more neighbors 5616. In at least one embodiment, one or more neighbors 5616 is value in retrieval database 5614 corresponding to a key comprising input data 5610. In at least one embodiment, one or more neighbors 5616 comprise text data. In at least one embodiment, encoder 5618 encodes one or more neighbors 5616. In at least one embodiment, encoder 5618 encodes one or more neighbors 5616 into a text embedding vector. In at least one embodiment, encoder 5618 encodes one or more neighbors 5616 into a sentence embedding vector. In at least one embodiment, large language model 5616 uses input data 5610 and data generated by encoder 5618 to generate output data 5620. In at least one embodiment, processor 5606 interfaces with application 5602 using large language model (LLM) application programming interface(s) (API(s)) 5604. In at least one embodiment, processor 5606 accesses large language model 5616 using large language model (LLM) application programming interface(s) (API(s)) 5604.

[0476] In at least one embodiment, output data 5620 comprise computer instructions. In at least one embodiment, output data 5620 comprise instructions written in CUDA programming language. In at least one embodiment, output data 5620 comprise instructions to be performed by processor 5606. In at least one embodiment, output data 5620 comprise instructions to control execution of one or more algorithm modules 5608. In at least one embodiment, one or more algorithm modules 5608 comprise, for example, one or more neural networks to perform pattern recognition. In at least one embodiment, one or more algorithm modules 5608 comprise, for example, one or more neural networks to perform frame generation. In at least one embodiment, one or more algorithm modules 5608 comprise, for example, one or more neural networks to generate a drive path. In at least one embodiment, one or more algorithm modules 5608 comprise, for example, one or more neural networks to generate a 5G signal. In at least one embodiment, processor 5606 interfaces with application 5602 using large language model (LLM) application programming interface(s) (API(s)) 5604. In at least one embodiment, processor 5606 may use one or more parallel computing platforms and/or programming models (e.g., NVIDIA's CUDA model).

[0477] In at least one embodiment, aspects of systems and techniques described herein in relation to FIG. 56 are incorporated into aspects of preceding figure(s). For example, in at least one embodiment, an apparatus depicted in preceding figure(s) includes processor 5606.

[0478] For example, in at least one embodiment, system 5600 uses ChatGPT to write CUDA code. For example, in at least one embodiment, system 5600 uses ChatGPT to train an object classification neural network. For example, in at

least one embodiment, system **5600** uses ChatGPT and a neural network to identify a driving path. For example, in at least one embodiment, system **5600** uses ChatGPT and a neural network to generate a 5G signal.

[0479] It should be noted that, while example embodiments described herein may relate to a CUDA programming model, techniques described herein can be utilized with any suitable programming model, such as HIP, oneAPI (e.g., using oneAPI-based programming to perform or implement a method disclosed herein), and/or variations thereof.

[0480] In at least one embodiment, one or more components of systems and/or processors disclosed above can communicate with one or more CPUs, ASICs, GPUs, FPGAs, or other hardware, circuitry, or integrated circuit components that include, e.g., an upscaler or upsampler to upscale an image, an image blender or image blender component to blend, mix, or add images together, a sampler to sample an image (e.g., as part of a DSP), a neural network circuit that is configured to perform an upscaler to upscale an image (e.g., from a low resolution image to a high resolution image), or other hardware to modify or generate an image, frame, or video to adjust its resolution, size, or pixels; one or more components of systems and/or processors disclosed above can use components described in this disclosure to perform methods, operations, or instructions that generate or modify an image.

[0481] At least one embodiment of the disclosure can be described in view of the following clauses:

[0482] 1. A processor comprising: one or more circuits to identify one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.

[0483] 2. The processor of clause 1, wherein the one or more circuits are to identify an amount of one or more pixels within the polygon based, at least in part, on an amount of one or more edges of the one or more pixels covered by one or more edges of the polygon.

[0484] 3. The processor of any of clauses 1-2, wherein the one or more circuits are to use one or more prefix sums along one or more dimensions to identify an amount of one or more pixels within the polygon.

[0485] 4. The processor of any of clauses 1-3, wherein the one or more circuits are to parallelize computations of an amount of one or more pixels covered by the polygon over two or more edges of the polygon.

[0486] 5. The processor of any of clauses 1-4, wherein the one or more circuits are to use the identified one or more pixels to, at least in part, perform computational lithography tasks.

[0487] 6. The processor of any of clauses 1-5, wherein the one or more circuits are to identify one or more portions of one or more edges of the polygon based, at least in part, on one or more intersections between the one or more edges and one or more pixels adjacent to the one or more pixels.

[0488] 7. The processor of any of clauses 1-6, wherein the one or more circuits are to identify an amount of one or more pixels within a polygon based, at least in part, on one or more portions of one or more edges of the polygon.

[0489] 8. A system, comprising:

[0490] one or more processors to identify one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.

[0491] 9. The system of clause 8, wherein the one or more processors are to identify an amount of one or more pixels within the polygon based, at least in part, on a fraction of one or more edges of the one or more pixels covered by one or more edges of the polygon.

[0492] 10. The system of any of clauses 8-9, wherein the one or more processors are to use one or more prefix sums along one or more rows of one or more pixels to identify an amount of one or more pixels within the polygon.

[0493] 11. The system of any of clauses 8-10, wherein the one or more processors are to parallelize computations of an amount of one or more edges of one or more pixels covered by the polygon over two or more edges of the polygon.

[0494] 12. The system of any of clauses 8-11, wherein the one or more processors are to identify the one or more pixels within a polygon to be used with computational lithography.

[0495] 13. The system of any of clauses 8-12, wherein the one or more processors are to identify one or more portions of one or more edges of the polygon based, at least in part, on one or more locations where two or more edges of the polygon meet.

[0496] 14. The system of any of clauses 8-13, wherein the one or more processors are to identify an amount of one or more pixels within a polygon by using one or more width values based, at least in part, on one or more portions of one or more edges of the polygon.

[0497] 15. A method, comprising:

[0498] identifying one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.

[0499] 16. The method of clause 15, further comprising identifying an amount of one or more pixels within the polygon based, at least in part, on a projection of one or more portions of one or more edges of the polygon onto an edge of one or more adjacent pixels.

[0500] 17. The method of any of clauses 15-16, further comprising using one or more prefix sums of amounts of one or more edges of the one or more pixels covered by one or more edges of the polygon.

[0501] 18. The method of any of clauses 15-17, further comprising parallelizing computations of one or more amounts of one or more pixels within the polygon over two or more portions of one or more edges of the polygon.

[0502] 19. The method of any of clauses 15-18, further comprising identifying the one or more pixels within a polygon to be used with rasterization.

[0503] 20. The method of any of clauses 15-19, further comprising identifying one or more portions of one or more edges of the polygon that traverse one or more pixels.

[0504] Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and

have been described above in detail. It should be understood, however, that there is no intention to limit disclosure to specific form or forms disclosed, but on contrary, intention is to cover all modifications, alternative constructions, and equivalents falling within spirit and scope of disclosure, as defined in appended claims.

[0505] Use of terms “a” and “an” and “the” and similar referents in context of describing disclosed embodiments (especially in context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. term “connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within range, unless otherwise indicated herein and each separate value is incorporated into specification as if it were individually recited herein. Use of term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, term “subset” of a corresponding set does not necessarily denote a proper subset of corresponding set, but subset and corresponding set may be equal.

[0506] Conjunctive language, such as phrases of form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of set of A and B and C. For instance, in illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). A number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, phrase “based on” means “based at least in part on” and not “based solely on.”

[0507] Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in form of a computer program comprising a plurality of instructions

executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (e.g., as a result of being executed) by one or more processors of a computer system, cause computer system to perform operations described herein. A set of non-transitory computer-readable storage media, in at least one embodiment, comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-transitory computer-readable storage media lack all of code while multiple non-transitory computer-readable storage media collectively store all of code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors—for example, a non-transitory computer-readable storage medium store instructions and a main central processing unit (“CPU”) executes some of instructions while a graphics processing unit (“GPU”) executes other instructions. In at least one embodiment, different components of a computer system have separate processors and different processors execute different subsets of instructions.

[0508] Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

[0509] Use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of disclosure and does not pose a limitation on scope of disclosure unless otherwise claimed. No language in specification should be construed as indicating any non-claimed element as essential to practice of disclosure.

[0510] All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

[0511] In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may be not intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

[0512] Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system’s registers and/or memories into other data similarly represented as physical quantities within computing system’s memories, registers or other such information storage, transmission or display devices.

[0513] In a similar manner, term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting examples, “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. Terms “system” and “method” are used herein interchangeably insofar as system may embody one or more methods and methods may be considered a system.

[0514] In at least one embodiment, an arithmetic logic unit is a set of combinational logic circuitry that takes one or more inputs to produce a result. In at least one embodiment, an arithmetic logic unit is used by a processor to implement mathematical operation such as addition, subtraction, or multiplication. In at least one embodiment, an arithmetic logic unit is used to implement logical operations such as logical AND/OR or XOR. In at least one embodiment, an arithmetic logic unit is stateless, and made from physical switching components such as semiconductor transistors arranged to form logical gates. In at least one embodiment, an arithmetic logic unit may operate internally as a stateful logic circuit with an associated clock. In at least one embodiment, an arithmetic logic unit may be constructed as an asynchronous logic circuit with an internal state not maintained in an associated register set. In at least one embodiment, an arithmetic logic unit is used by a processor to combine operands stored in one or more registers of the processor and produce an output that can be stored by the processor in another register or a memory location.

[0515] In at least one embodiment, as a result of processing an instruction retrieved by the processor, the processor presents one or more inputs or operands to an arithmetic logic unit, causing the arithmetic logic unit to produce a result based at least in part on an instruction code provided to inputs of the arithmetic logic unit. In at least one embodiment, the instruction codes provided by the processor to the ALU are based at least in part on the instruction executed by the processor. In at least one embodiment combinational logic in the ALU processes the inputs and produces an output which is placed on a bus within the processor. In at least one embodiment, the processor selects a destination register, memory location, output device, or output storage location on the output bus so that clocking the processor causes the results produced by the ALU to be sent to the desired location.

[0516] In present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital

data into a subsystem, computer system, or computer-implemented machine. Process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In some implementations, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In another implementation, process of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. References may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, process of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or interprocess communication mechanism.

[0517] Although discussion above sets forth example implementations of described techniques, other architectures may be used to implement described functionality, and are intended to be within scope of this disclosure. Furthermore, although specific distributions of responsibilities are defined above for purposes of discussion, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

[0518] Furthermore, although subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

What is claimed is:

1. A processor comprising: one or more circuits to identify one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.
2. The processor of claim 1, wherein the one or more circuits are to identify an amount of one or more pixels within the polygon based, at least in part, on an amount of one or more edges of the one or more pixels covered by one or more edges of the polygon.
3. The processor of claim 1, wherein the one or more circuits are to use one or more prefix sums along one or more dimensions to identify an amount of one or more pixels within the polygon.
4. The processor of claim 1, wherein the one or more circuits are to parallelize computations of an amount of one or more pixels covered by the polygon over two or more edges of the polygon.
5. The processor of claim 1, wherein the one or more circuits are to use the identified one or more pixels to, at least in part, perform computational lithography tasks.
6. The processor of claim 1, wherein the one or more circuits are to identify one or more portions of one or more edges of the polygon based, at least in part, on one or more intersections between the one or more edges and one or more pixels adjacent to the one or more pixels.
7. The processor of claim 1, wherein the one or more circuits are to identify an amount of one or more pixels within a polygon based, at least in part, on one or more portions of one or more edges of the polygon.

- 8.** A system, comprising:
one or more processors to identify one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.
- 9.** The system of claim **8**, wherein the one or more processors are to identify an amount of one or more pixels within the polygon based, at least in part, on a fraction of one or more edges of the one or more pixels covered by one or more edges of the polygon.
- 10.** The system of claim **8**, wherein the one or more processors are to use one or more prefix sums along one or more rows of one or more pixels to identify an amount of one or more pixels within the polygon.
- 11.** The system of claim **8**, wherein the one or more processors are to parallelize computations of an amount of one or more edges of one or more pixels covered by the polygon over two or more edges of the polygon.
- 12.** The system of claim **8**, wherein the one or more processors are to identify the one or more pixels within a polygon to be used with computational lithography.
- 13.** The system of claim **8**, wherein the one or more processors are to identify one or more portions of one or more edges of the polygon based, at least in part, on one or more locations where two or more edges of the polygon meet.
- 14.** The system of claim **8**, wherein the one or more processors are to identify an amount of one or more pixels

within a polygon by using one or more width values based, at least in part, on one or more portions of one or more edges of the polygon.

- 15.** A method, comprising:
identifying one or more pixels within a polygon based, at least in part, on whether the one or more pixels are adjacent to one or more pixels that cross one or more polygon boundaries.
- 16.** The method of claim **15**, further comprising identifying an amount of one or more pixels within the polygon based, at least in part, on a projection of one or more portions of one or more edges of the polygon onto an edge of one or more adjacent pixels.
- 17.** The method of claim **15**, further comprising using one or more prefix sums of amounts of one or more edges of the one or more pixels covered by one or more edges of the polygon.
- 18.** The method of claim **15**, further comprising parallelizing computations of one or more amounts of one or more pixels within the polygon over two or more portions of one or more edges of the polygon.
- 19.** The method of claim **15**, further comprising identifying the one or more pixels within a polygon to be used with rasterization.
- 20.** The method of claim **15**, further comprising identifying one or more portions of one or more edges of the polygon that traverse one or more pixels.

* * * * *