

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication
Kind Code
Publication Date
Inventor(s)

20250278187
A1
September 04, 2025
Papa; Paolo et al.

CLUSTERED PARITY FOR NAND DATA PLACEMENT SCHEMA

Abstract

Disclosed in some examples are improvements to data placement architectures in NAND that provide additional data protection through an improved NAND data placement schema that allows for recovery from certain failure scenarios. The present disclosure stripes data diagonally across page lines and planes to enhance the data protection. Parity bits are stored in SLC blocks for extra protection until the block is finished writing and then the parity bits may be deleted.

Inventors: Papa; Paolo (Naples, IT), Manganelli; Carminatonio (Benevento, IT), Iaculo; Massimo (San Marco Evangelista, IT), D'Eliseo; Giuseppe (Caserta, IT), Sassari; Alberto (Napoli, IT)

Applicant: Micron Technology, Inc. (Boise, ID)

Family ID: 67908475

Appl. No.: 19/212220

Filed: May 19, 2025

Related U.S. Application Data

parent US continuation 18137820 20230421 parent-grant-document US 12321594 child US 19212220
parent US continuation 16488696 20190826 parent-grant-document US 11635894 US continuation PCT/US2019/022576 20190315 child US 18137820
us-provisional-application US 62675451 20180523
us-provisional-application US 62644282 20180316

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/0608 (20130101); G06F3/0655 (20130101); G06F3/0679 (20130101);

Background/Summary

PRIORITY APPLICATIONS [0001] The current application is a continuation of U.S. application Ser. No. 18/137,820, filed Apr. 21, 2023, which is a continuation of U.S. application Ser. No. 16/488,696, filed Aug. 26, 2019, now issued as U.S. Pat. No. 11,635,894, which is a U.S. National Stage Application under 35 U.S.C. 371 from International Application No. PCT/US2019/022576, filed Mar. 15, 2019, which claims the benefit of priority to U.S. Provisional Application Ser. No. 62/675,451, filed May 23, 2018 and to U.S. Provisional Application Ser. No. 62/644,282, filed Mar. 16, 2018, all of which are incorporated herein by reference in their entireties.

BACKGROUND

[0002] Memory devices are typically provided as internal, semiconductor, integrated circuits in computers or other electronic devices. There are many different types of memory, including volatile and non-volatile memory.

[0003] Volatile memory requires power to maintain its data, and includes random-access memory (RAM), dynamic random-access memory (DRAM), or synchronous dynamic random-access memory (SDRAM), among others.

[0004] Non-volatile memory can retain stored data when not powered, and includes flash memory, read-only memory (ROM), electrically erasable programmable ROM (EEPROM), static RAM (SRAM), erasable programmable ROM (EPROM), resistance variable memory, such as phase-change random-access memory (PCRAM), resistive random-access memory (RRAM), magnetoresistive random-access memory (MRAM), or 3D XPoint™ memory, among others.

[0005] Flash memory is utilized as non-volatile memory for a wide range of electronic applications. Flash memory devices typically include one or more groups of one-transistor, floating gate or charge trap memory cells that allow for high memory densities, high reliability, and low power consumption.

[0006] Two common types of flash memory array architectures include NAND and NOR architectures, named after the logic form in which the basic memory cell configuration of each is arranged. The memory cells of the memory array are typically arranged in a matrix. In an example, the gates of each floating gate memory cell in a row of the array are coupled to an access line (e.g., a word line). In a NOR architecture, the drains of each memory cell in a column of the array are coupled to a data line (e.g., a bit line). In a NAND architecture, the drains of each memory cell in a string of

the array are coupled together in series, source to drain, between a source line and a bit line.

[0007] Both NOR and NAND architecture semiconductor memory arrays are accessed through decoders that activate specific memory cells by selecting the word line coupled to their gates. In a NOR architecture semiconductor memory array, once activated, the selected memory cells place their data values on bit lines, causing different currents to flow depending on the state at which a particular cell is programmed. In a NAND architecture semiconductor memory array, a high bias voltage is applied to a drain-side select gate (SGD) line. Word lines coupled to the gates of the unselected memory cells of each group are driven at a specified pass voltage (e.g., Vpass) to operate the unselected memory cells of each group as pass transistors (e.g., to pass current in a manner unrestricted by their stored data values). Current then flows from the source line to the bit line through each series coupled group, restricted only by the selected memory cells of each group, placing current encoded data values of selected memory cells on the bit lines.

[0008] Each flash memory cell in a NOR or NAND architecture semiconductor memory array can be programmed individually or collectively to one or a number of programmed states. For example, a single-level cell (SLC) can represent one of two programmed states (e.g., 1 or 0), representing one bit of data.

[0009] However, flash memory cells can also represent one of more than two programmed states, allowing the manufacture of higher density memories without increasing the number of memory cells, as each cell can represent more than one binary digit (e.g., more than one bit). Such cells can be referred to as multi-state memory cells, multi-digit cells, or multi-level cells (MLCs). In certain examples, MLC can refer to a memory cell that can store two bits of data per cell (e.g., one of four programmed states), a triple-level cell (TLC) can refer to a memory cell that can store three bits of data per cell (e.g., one of eight programmed states), and a quad-level cell (QLC) can store four bits of data per cell. MLC is used herein in its broader context, to can refer to any memory cell that can store more than one bit of data per cell (i.e., that can represent more than two programmed states).

[0010] Traditional memory arrays are two-dimensional (2D) structures arranged on a surface of a semiconductor substrate. To increase memory capacity for a given area, and to decrease cost, the size of the individual memory cells has decreased. However, there is a technological limit to the reduction in size of the individual memory cells, and thus, to the memory density of 2D memory arrays. In response, three-dimensional (3D) memory structures, such as 3D NAND architecture semiconductor memory devices, are being developed to further increase memory density and lower memory cost.

[0011] Such 3D NAND devices often include strings of storage cells, coupled in series (e.g., drain to source), between one or more source-side select gates (SGSs) proximate a source, and one or more drain-side select gates (SGDs) proximate a bit line. In an example, the SGSs or the SGDs can include one or more field-effect transistors (FETs) or metal-oxide semiconductor (MOS) structure devices, etc. In some examples, the strings will extend vertically, through multiple vertically spaced tiers containing respective word lines. A semiconductor structure (e.g., a polysilicon structure) may extend adjacent a string of storage cells to form a channel for the storages cells of the string. In the example of a vertical string, the polysilicon structure may be in the form of a vertically extending pillar. In some examples the string may be “folded,” and thus arranged relative to a U-shaped pillar. In other examples, multiple vertical structures may be stacked upon one another to form stacked arrays of storage cell strings.

[0012] Memory arrays or devices can be combined together to form a storage volume of a memory system, such as a solid-state drive (SSD), a Universal Flash Storage (UFS™) device, a MultiMediaCard (MMC) solid-state storage device, an embedded MMC device (eMMC™), etc. An SSD can be used as, among other things, the main storage device of a computer, having advantages over traditional hard drives with moving parts with respect to, for example, performance, size, weight, ruggedness, operating temperature range, and power consumption. For example, SSDs can have reduced seek time, latency, or other delay associated with magnetic disk drives (e.g., electromechanical, etc.). SSDs use non-volatile memory cells, such as flash memory cells to obviate internal battery supply requirements, thus allowing the drive to be more versatile and compact.

[0013] An SSD can include a number of memory devices, including a number of dies or logical units (e.g., logical unit numbers or LUNs), and can include one or more processors or other controllers performing logic functions required to operate the memory devices or interface with external systems. Such SSDs may include one or more flash memory die, including a number of memory arrays and peripheral circuitry thereon. The flash memory arrays can include a number of blocks of memory cells organized into a number of physical pages. In many examples, the SSDs will also include DRAM or SRAM (or other forms of memory die or other memory structures). The SSD can receive commands from a host in association with memory operations, such as read or write operations to transfer data (e.g., user data and associated integrity data, such as error data and address data, etc.) between the memory devices and the host, or erase operations to erase data from the memory devices.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0014] In the drawings, which are not necessarily drawn to scale, like numerals may describe similar components in different views. Like numerals having different letter suffixes may represent different instances of similar components. The drawings illustrate generally, by way of example, but not by way of limitation, various embodiments discussed in the present document.

[0015] FIG. 1 illustrates an example of an environment including a memory device.

[0016] FIGS. 2-3 illustrate schematic diagrams of an example of a 3D NAND architecture semiconductor memory array.

[0017] FIG. 4 illustrates an example block diagram of a memory module.

[0018] FIG. 5 illustrates possible effects of a programming failure of a NAND.

[0019] FIG. 6 illustrates possible effects an asynchronous power loss during programming of a NAND.

[0020] FIGS. 7 and 8 illustrate an improved NAND data placement schema of data on a TLC NAND array according to some examples of the present disclosure.

[0021] FIG. 9 illustrates a logical placement of the parity pages in a volatile memory of the controller or other component of the NAND according to some examples of the present disclosure.

[0022] FIG. 10 illustrates a schematic of a parity value storage used to store the parity pages in nonvolatile NAND.

[0023] FIG. 11 illustrates a flowchart of a method 1100 of applying an improved NAND data placement schema to storing data on a NAND die of a NAND device according to some examples of the present disclosure.

[0024] FIG. 12 illustrates an example of an uncorrectable error of a NAND according to some examples of the present disclosure.

[0025] FIG. 13 illustrates storage of a plurality of parity values that are arranged in a plurality of clusters according to some examples of the present disclosure.

[0026] FIG. 14 illustrates storage of a plurality of compressed parity values according to some examples of the present disclosure.

[0027] FIG. 15 illustrates a flow chart of a method of clustering parity values in an improved NAND data placement schema of a NAND device according to some examples of the present disclosure.

[0028] FIG. 16 illustrates a schematic of a memory controller according to some examples of the present disclosure.

[0029] FIG. 17 is a block diagram illustrating an example of a machine upon which one or more embodiments may be implemented.

DETAILED DESCRIPTION

[0030] Disclosed in some examples are methods of organizing data written to a memory device (such as a NAND memory device) in order to protect against certain types of faults. For example, a first portion is programmed on a first page line and a first plane, a second portion is programmed on a second page line and a second plane, a third portion is programmed on a third page line and a third plane, and a fourth portion is programmed on a

fourth page line and a fourth plane. The first page line, first plane, first plane, second plane, third plane, and third plane are selected such that the first portion, second portion and third portion are stored in memory cells that are on different page lines and different planes with respect to each other.

[0031] Electronic devices, such as mobile electronic devices (e.g., smart phones, tablets, etc.), electronic devices for use in automotive applications (e.g., automotive sensors, control units, driver-assistance systems, passenger safety or comfort systems, etc.), and internet-connected appliances or devices (e.g., internet-of-things (IoT) devices, etc.), have varying storage needs depending on, among other things, the type of electronic device, use environment, performance expectations, etc.

[0032] Electronic devices can be broken down into several main components: a processor (e.g., a central processing unit (CPU) or other main processor); memory (e.g., one or more volatile or non-volatile random-access memory (RAM) memory device, such as dynamic RAM (DRAM), mobile or low-power double-data-rate synchronous DRAM (DDR SDRAM), etc.); and a storage device (e.g., non-volatile memory (NVM) device, such as flash memory, read-only memory (ROM), an SSD, an MMC, or other memory card structure or assembly, etc.). In certain examples, electronic devices can include a user interface (e.g., a display, touch-screen, keyboard, one or more buttons, etc.), a graphics processing unit (GPU), a power management circuit, a baseband processor or one or more transceiver circuits, etc.

[0033] FIG. 1 illustrates an example of an environment **100** including a host device **105** and a memory device **110** configured to communicate over a communication interface. The host device **105** or the memory device **110** may be included in a variety of products **150**, such as Internet of Things (IoT) devices (e.g., a refrigerator or other appliance, sensor, motor or actuator, mobile communication device, automobile, drone, etc.) to support processing, communications, or control of the product **150**.

[0034] The memory device **110** includes a memory controller **115** and a memory array **120** including, for example, a number of individual memory die (e.g., a stack of three-dimensional (3D) NAND die). In 3D architecture semiconductor memory technology, vertical structures are stacked, increasing the number of tiers, physical pages, and accordingly, the density of a memory device (e.g., a storage device). In an example, the memory device **110** can be a discrete memory or storage device component of the host device **105**. In other examples, the memory device **110** can be a portion of an integrated circuit (e.g., system on a chip (SOC), etc.), stacked or otherwise included with one or more other components of the host device **105**.

[0035] One or more communication interfaces can be used to transfer data between the memory device **110** and one or more other components of the host device **105**, such as a Serial Advanced Technology Attachment (SATA) interface, a Peripheral Component Interconnect Express (PCIe) interface, a Universal Serial Bus (USB) interface, a Universal Flash Storage (UFS) interface, an eMMC™ interface, or one or more other connectors or interfaces. The host device **105** can include a host system, an electronic device, a processor, a memory card reader, or one or more other electronic devices external to the memory device **110**. In some examples, the host **105** may be a machine having some portion, or all, of the components discussed in reference to the machine **1700** of FIG. 17.

[0036] The memory controller **115** can receive instructions from the host **105**, and can communicate with the memory array, such as to transfer data to (e.g., write or erase) or from (e.g., read) one or more of the memory cells, planes, sub-blocks, blocks, or pages of the memory array. The memory controller **115** can include, among other things, circuitry or firmware, including one or more components or integrated circuits. For example, the memory controller **115** can include one or more memory control units, circuits, or components configured to control access across the memory array **120** and to provide a translation layer between the host **105** and the memory device **110**. The memory controller **115** can include one or more input/output (I/O) circuits, lines, or interfaces to transfer data to or from the memory array **120**. The memory controller **115** can include a memory manager **125** and an array controller **135**.

[0037] The memory manager **125** can include, among other things, circuitry or firmware, such as a number of components or integrated circuits associated with various memory management functions. For purposes of the present description example memory operation and management functions will be described in the context of NAND memory. Persons skilled in the art will recognize that other forms of non-volatile memory may have analogous memory operations or management functions. Such NAND management functions include wear leveling (e.g., garbage collection or reclamation), error detection or correction, block retirement, or one or more other memory management functions. The memory manager **125** can parse or format host commands (e.g., commands received from a host) into device commands (e.g., commands associated with operation of a memory array, etc.), or generate device commands (e.g., to accomplish various memory management functions) for the array controller **135** or one or more other components of the memory device **110**.

[0038] The memory manager **125** can include a set of management tables **130** configured to maintain various information associated with one or more component of the memory device **110** (e.g., various information associated with a memory array or one or more memory cells coupled to the memory controller **115**). For example, the management tables **130** can include information regarding block age, block erase count, error history, or one or more error counts (e.g., a write operation error count, a read bit error count, a read operation error count, an erase error count, etc.) for one or more blocks of memory cells coupled to the memory controller **115**. In certain examples, if the number of detected errors for one or more of the error counts is above a threshold, the bit error can be referred to as an uncorrectable bit error. The management tables **130** can maintain a count of correctable or uncorrectable bit errors, among other things.

[0039] The array controller **135** can include, among other things, circuitry or components configured to control memory operations associated with writing data to, reading data from, or erasing one or more memory cells of the memory device **110** coupled to the memory controller **115**. The memory operations can be based on, for example, host commands received from the host **105**, or internally generated by the memory manager **125** (e.g., in association with wear leveling, error detection or correction, etc.).

[0040] The array controller **135** can include an error correction code (ECC) component **140**, which can include, among other things, an ECC engine or other circuitry configured to detect or correct errors associated with writing data to or reading data from one or more memory cells of the memory device **110** coupled to the memory controller **115**. The memory controller **115** can be configured to actively detect and recover from error occurrences (e.g., bit errors, operation errors, etc.) associated with various operations or storage of data, while maintaining integrity of the data transferred between the host **105** and the memory device **110**, or maintaining integrity of stored data (e.g., using redundant RAID storage, etc.), and can remove (e.g., retire) failing memory resources (e.g., memory cells, memory arrays, pages, blocks, etc.) to prevent future errors.

[0041] In some examples, the memory array may comprise a number of NAND dies and one or more functions of the memory controller **115** for a particular NAND die may be implemented on an on-die controller on that particular die. Other organizations and delineations of control functionality may also be utilized, such as a controller for each die, plane, superblock, block, page, and the like.

[0042] The memory array **120** can include several memory cells arranged in, for example, a number of devices, semi-conductor dies, planes, sub-blocks, blocks, or pages. As one example, a 48 GB TLC NAND memory device can include 18,592 bytes (B) of data per page (16,384+2208 bytes), 1536 pages per block, 548 blocks per plane, and 4 or more planes per device. As another example, a 32 GB MLC memory device (storing two bits of data per cell (i.e., 4 programmable states)) can include 18,592 bytes (B) of data per page (16,384+2208 bytes), 1024 pages per block, 548 blocks per plane, and 4 planes per device, but with half the required write time and twice the program/erase (P/E) cycles as a corresponding TLC memory device. Other examples can include other numbers or arrangements. In some examples, a memory device, or a portion thereof, may be selectively operated in SLC mode, or in a desired MLC mode (such as TLC, QLC, etc.).

[0043] In operation, data is typically written to or read from the NAND memory device **110** in pages, and erased in blocks. However, one or more memory operations (e.g., read, write, erase, etc.) can be performed on larger or smaller groups of memory cells, as desired. The data transfer size of a NAND memory device **110** is typically referred to as a page, whereas the data transfer size of a host is typically referred to as a sector.

[0044] Although a page of data can include a number of bytes of user data (e.g., a data payload including a number of sectors of data) and its corresponding metadata, the size of the page often refers only to the number of bytes used to store the user data. As an example, a page of data

having a page size of 4 KB may include 4 KB of user data (e.g., 8 sectors assuming a sector size of 512 B) as well as a number of bytes (e.g., 32 B, 54 B, 224 B, etc.) of metadata corresponding to the user data, such as integrity data (e.g., error detecting or correcting code data), address data (e.g., logical address data, etc.), or other metadata associated with the user data.

[0045] Different types of memory cells or memory arrays **120** can provide for different page sizes, or may require different amounts of metadata associated therewith. For example, different memory device types may have different bit error rates, which can lead to different amounts of metadata necessary to ensure integrity of the page of data (e.g., a memory device with a higher bit error rate may require more bytes of error correction code data than a memory device with a lower bit error rate). As an example, a multi-level cell (MLC) NAND flash device may have a higher bit error rate than a corresponding single-level cell (SLC) NAND flash device. As such, the MLC device may require more metadata bytes for error data than the corresponding SLC device.

[0046] FIG. 2 illustrates an example schematic diagram of a 3D NAND architecture semiconductor memory array **200** including a number of strings of memory cells (e.g., first-third A.sub.0 memory strings **205A.sub.0-207A.sub.0**, first-third A.sub.n memory strings **205A.sub.n-207A.sub.n**, first-third B.sub.0 memory strings **205B.sub.0-207B.sub.0**, first-third B.sub.n memory strings **205B.sub.n-207B.sub.n**, etc.), organized in blocks (e.g., block A **201A**, block B **201B**, etc.) and sub-blocks (e.g., sub-block A.sub.0 **201A.sub.0**, sub-block A.sub.n **201A.sub.n**, sub-block B.sub.0 **201B.sub.0**, sub-block B.sub.n **201B.sub.n**, etc.). The memory array **200** represents a portion of a greater number of similar structures that would typically be found in a block, device, or other unit of a memory device.

[0047] Each string of memory cells includes a number of tiers of charge storage transistors (e.g., floating gate transistors, charge-trapping structures, etc.) stacked in the Z direction, source to drain, between a source line (SRC) **235** or a source-side select gate (SGS) (e.g., first-third A.sub.0 SGS **231A.sub.0-233A.sub.0**, first-third A.sub.n SGS **231A.sub.n-233A.sub.n**, first-third B.sub.0 SGS **231B.sub.0-233B.sub.0**, first-third B.sub.n SGS **231B.sub.n-233B.sub.n**, etc.) and a drain-side select gate (SGD) (e.g., first-third A.sub.0 SGD **226A.sub.0-228A.sub.0**, first-third A.sub.n SGD **226A.sub.n-228A.sub.n**, first-third B.sub.0 SGD **226B.sub.0-228B.sub.0**, first-third B.sub.n SGD **226B.sub.n-228B.sub.n**, etc.). Each string of memory cells in the 3D memory array can be arranged along the X direction as data lines (e.g., bit lines (BL) **BL0-BL2 220-222**), and along the Y direction as physical pages.

[0048] Within a physical page, each tier represents a row of memory cells, and each string of memory cells represents a column. A sub-block can include one or more physical pages. A block can include a number of sub-blocks (or physical pages) (e.g., **128, 256, 384**, etc.). Although illustrated herein as having two blocks, each block having two sub-blocks, each sub-block having a single physical page, each physical page having three strings of memory cells, and each string having 8 tiers of memory cells, in other examples, the memory array **200** can include more or fewer blocks, sub-blocks, physical pages, strings of memory cells, memory cells, or tiers. For example, each string of memory cells can include more or fewer tiers (e.g., 16, 32, 64, 128, etc.), as well as one or more additional tiers of semiconductor material above or below the charge storage transistors (e.g., select gates, data lines, etc.), as desired. As an example, a 48 GB TLC NAND memory device can include 18,592 bytes (B) of data per page (16,384+2208 bytes), 1536 pages per block, 548 blocks per plane, and 4 or more planes per device.

[0049] Each memory cell in the memory array **200** includes a control gate (CG) coupled to (e.g., electrically or otherwise operatively connected to) an access line (e.g., word lines (WL) **WL0.sub.0-WL7.sub.0, 210A-217A, WL0.sub.1-WL7.sub.1 210B-217B**, etc.), which collectively couples the control gates (CGs) across a specific tier, or a portion of a tier, as desired. Specific tiers in the 3D memory array, and accordingly, specific memory cells in a string, can be accessed or controlled using respective access lines. Groups of select gates can be accessed using various select lines. For example, first-third A.sub.0 SGD **226A.sub.0-228A.sub.0** can be accessed using an A.sub.0 SGD line **SGDA.sub.0 225A.sub.0**, first-third A.sub.n SGD **226A.sub.n-228A.sub.n** can be accessed using an A.sub.n SGD line **SGDA.sub.n 225A.sub.n**, first-third B.sub.0 SGD **226B.sub.0-228B.sub.0** can be accessed using an B.sub.0 SGD line **SGDB.sub.0 225B.sub.0**, and first-third B.sub.n SGD **226B.sub.n-228B.sub.n** can be accessed using an B.sub.n SGD line **SGDB.sub.n 225B.sub.n**. First-third A.sub.0 SGS **231A.sub.0-233A.sub.0**, and first-third A.sub.n SGS **231A.sub.n-233A.sub.n** can be accessed using a gate select line **SGS.sub.0 230A.sub.0**, and first-third B.sub.0 SGS **231B.sub.0-233B.sub.0** and first-third B.sub.n SGS **231B.sub.n-233B.sub.n** can be accessed using a gate select line **SGS.sub.1 230B**.

[0050] In an example, the memory array **200** can include a number of levels of semiconductor material (e.g., polysilicon, etc.) configured to couple the control gates (CGs) of each memory cell or select gate (or a portion of the CGs or select gates) of a respective tier of the array. Specific strings of memory cells in the array can be accessed, selected, or controlled using a combination of bit lines (BLs) and select gates, etc., and specific memory cells at one or more tiers in the specific strings can be accessed, selected, or controlled using one or more access lines (e.g., word lines).

[0051] FIG. 3 illustrates an example schematic diagram of a portion of a NAND architecture semiconductor memory array **300** including a plurality of memory cells **302** arranged in a two-dimensional array of strings (e.g., first-third strings **305-307**) and tiers (e.g., illustrated as respective word lines (WL) **WL0-WL7 310-317**, a drain-side select gate (SGD) line **325**, a source-side select gate (SGS) line **330**, etc.), and sense amplifiers or devices **360**. For example, the memory array **300** can illustrate an example schematic diagram of a portion of one physical page of memory cells of a 3D NAND architecture semiconductor memory device, such as illustrated in FIG. 2.

[0052] Each string of memory cells is coupled to a source line (SRC) using a respective source-side select gate (SGS) (e.g., first-third SGS **331-333**), and to a respective data line (e.g., first-third bit lines (BL) **BL0-BL2 320-322**) using a respective drain-side select gate (SGD) (e.g., first-third SGD **326-328**). Although illustrated with 8 tiers (e.g., using word lines (WL) **WL0-WL7 310-317**) and three data lines (**BL0-BL2 326-328**) in the example of FIG. 3, other examples can include strings of memory cells having more or fewer tiers or data lines, as desired.

[0053] In a NAND architecture semiconductor memory array, such as the example memory array **300**, the state of a selected memory cell **302** can be accessed by sensing a current or voltage variation associated with a particular data line containing the selected memory cell. The memory array **300** can be accessed (e.g., by a control circuit, one or more processors, digital logic, etc.) using one or more drivers. In an example, one or more drivers can activate a specific memory cell, or set of memory cells, by driving a particular potential to one or more data lines (e.g., bit lines **BL0-BL2**), access lines (e.g., word lines **WL0-WL7**), or select gates, depending on the type of operation desired to be performed on the specific memory cell or set of memory cells.

[0054] To program or write data to a memory cell, a programming voltage (V_{pgm}) (e.g., one or more programming pulses, etc.) can be applied to selected word lines (e.g., **WL4**), and thus, to a control gate of each memory cell coupled to the selected word lines (e.g., first-third control gates (CGs) **341-343** of the memory cells coupled to **WL4**). Programming pulses can begin, for example, at or near 15V, and, in certain examples, can increase in magnitude during each programming pulse application. While the program voltage is applied to the selected word lines, a potential, such as a ground potential (e.g., V_{ss}), can be applied to the data lines (e.g., bit lines) and substrates (and thus the channels, between the sources and drains) of the memory cells targeted for programming, resulting in a charge transfer (e.g., direct injection or Fowler-Nordheim (FN) tunneling, etc.) from the channels to the floating gates of the targeted memory cells.

[0055] In contrast, a pass voltage (V_{pass}) can be applied to one or more word lines having memory cells that are not targeted for programming, or an inhibit voltage (e.g., V_{cc}) can be applied to data lines (e.g., bit lines) having memory cells that are not targeted for programming, for example, to inhibit charge from being transferred from the channels to the floating gates of such non-targeted memory cells. The pass voltage can be variable, depending, for example, on the proximity of the applied pass voltages to a word line targeted for programming. The inhibit voltage can include a supply voltage (V_{cc}), such as a voltage from an external source or supply (e.g., a battery, an AC-to-DC converter, etc.), relative to a ground potential (e.g., V_{ss}).

[0056] As an example, if a programming voltage (e.g., 15V or more) is applied to a specific word line, such as **WL4**, a pass voltage of 10V can be applied to one or more other word lines, such as **WL3, WL5**, etc., to inhibit programming of non-targeted memory cells, or to retain the values stored on such memory cells not targeted for programming. As the distance between an applied program voltage and the non-targeted memory cells

increases, the pass voltage required to refrain from programming cells can decrease. For example, where a programming voltage of 15V is applied to WL4, a pass voltage of 10V can be applied to WL3 and WL5, a pass voltage of 8V can be applied to WL2 and WL6, a pass voltage of 7V can be applied to WL1 and WL7, etc. In other examples, the pass voltages, or number of word lines, etc., can be higher or lower, or more or less.

[0057] The sense amplifiers **360**, coupled to one or more of the data lines (e.g., first, second, or third bit lines (BL0-BL2) **320-322**), can detect the state of each memory cell in respective data lines by sensing a voltage or current on a particular data line.

[0058] Between applications of one or more programming pulses (e.g., Vpgm), a verify operation can be performed to determine if a selected memory cell has reached its intended programmed state. If the selected memory cell has reached its intended programmed state, it can be inhibited from further programming. If the selected memory cell has not reached its intended programmed state, additional programming pulses can be applied. If the selected memory cell has not reached its intended programmed state after a particular number of programming pulses (e.g., a maximum number), the selected memory cell, or a string, block, or page associated with such selected memory cell, can be marked as defective.

[0059] To erase a memory cell or a group of memory cells (e.g., erasure is typically performed in blocks or sub-blocks), an erasure voltage (Vers) (e.g., typically Vpgm) can be applied to the substrates (and thus the channels, between the sources and drains) of the memory cells targeted for erasure (e.g., using one or more bit lines, select gates, etc.), while the word lines of the targeted memory cells are kept at a potential, such as a ground potential (e.g., Vss), resulting in a charge transfer (e.g., direct injection or Fowler-Nordheim (FN) tunneling, etc.) from the floating gates of the targeted memory cells to the channels.

[0060] FIG. 4 illustrates an example block diagram of a memory device **400** including a memory array **402** having a plurality of memory cells **404**, and one or more circuits or components to provide communication with, or perform one or more memory operations on, the memory array **402**. The memory device **400** can include a row decoder **412**, a column decoder **414**, sense amplifiers **420**, a page buffer **422**, a selector **424**, an input/output (I/O) circuit **426**, and a memory control unit **430**.

[0061] The memory cells **404** of the memory array **402** can be arranged in blocks, such as first and second blocks **402A**, **402B**. Each block can include sub-blocks. For example, the first block **402A** can include first and second sub-blocks **402A.sub.0**, **402A.sub.n**, and the second block **402B** can include first and second sub-blocks **402B.sub.0**, **402B.sub.n**. Each sub-block can include a number of physical pages, each page including a number of memory cells **404**. Although illustrated herein as having two blocks, each block having two sub-blocks, and each sub-block having a number of memory cells **404**, in other examples, the memory array **402** can include more or fewer blocks, sub-blocks, memory cells, etc. In other examples, the memory cells **404** can be arranged in a number of rows, columns, pages, sub-blocks, blocks, etc., and accessed using, for example, access lines **406**, first data lines **410**, or one or more select gates, source lines, etc.

[0062] The memory control unit **430** can control memory operations of the memory device **400** according to one or more signals or instructions received on control lines **432**, including, for example, one or more clock signals or control signals that indicate a desired operation (e.g., write, read, erase, etc.), or address signals (A0-AX) received on one or more address lines **416**. One or more devices external to the memory device **400** can control the values of the control signals on the control lines **432**, or the address signals on the address line **416**. Examples of devices external to the memory device **400** can include, but are not limited to, a host, a memory controller, a processor, or one or more circuits or components not illustrated in FIG. 4.

[0063] The memory device **400** can use access lines **406** and first data lines **410** to transfer data to (e.g., write or erase) or from (e.g., read) one or more of the memory cells **404**. The row decoder **412** and the column decoder **414** can receive and decode the address signals (A0-AX) from the address line **416**, can determine which of the memory cells **404** are to be accessed, and can provide signals to one or more of the access lines **406** (e.g., one or more of a plurality of word lines (WL0-WLn)) or the first data lines **410** (e.g., one or more of a plurality of bit lines (BL0-BLn)), such as described above.

[0064] The memory device **400** can include sense circuitry, such as the sense amplifiers **420**, configured to determine the values of data on (e.g., read), or to determine the values of data to be written to, the memory cells **404** using the first data lines **410**. For example, in a selected string of memory cells **404**, one or more of the sense amplifiers **420** can read a logic level in the selected memory cell **404** in response to a read current flowing in the memory array **402** through the selected string to the data lines **410**.

[0065] One or more devices external to the memory device **400** can communicate with the memory device **400** using the I/O lines (DQ0-DQN) **408**, address lines **416** (A0-AX), or control lines **432**. The input/output (I/O) circuit **426** can transfer values of data in or out of the memory device **400**, such as in or out of the page buffer **422** or the memory array **402**, using the I/O lines **408**, according to, for example, the control lines **432** and address lines **416**. The page buffer **422** can store data received from the one or more devices external to the memory device **400** before the data is programmed into relevant portions of the memory array **402**, or can store data read from the memory array **402** before the data is transmitted to the one or more devices external to the memory device **400**.

[0066] The column decoder **414** can receive and decode address signals (A0-AX) into one or more column select signals (CSEL1-CSELn). The selector **424** (e.g., a select circuit) can receive the column select signals (CSEL1-CSELn) and select data in the page buffer **422** representing values of data to be read from or to be programmed into memory cells **404**. Selected data can be transferred between the page buffer **422** and the I/O circuit **426** using second data lines **418**.

[0067] The memory control unit **430** can receive positive and negative supply signals, such as a supply voltage (Vcc) **434** and a negative supply (Vss) **436** (e.g., a ground potential), from an external source or supply (e.g., an internal or external battery, an AC-to-DC converter, etc.). In certain examples, the memory control unit **430** can include a regulator **428** to internally provide positive or negative supply signals.

[0068] ECC and other techniques have increased the reliability of NAND devices substantially. Nonetheless, there are certain circumstances in which additional protection against data loss is desired. For example, as shown in FIG. 5, a write operation that has a programming failure while programming page line X, **510**, may corrupt many pages inside that plane. As shown in FIG. 5, all pages in the plane (e.g., plane **1**) have been corrupted by the programming failure of page line X, **510**. Similarly, and as shown in FIG. 6, an asynchronous power loss (loss of power to the NAND device without warning) during programming of a first page line X, **607** may also corrupt a different page line Z, **605**.

[0069] As used herein, a page line is a logical construct that identifies a group of pages comprising pages at a same position in each plane in a group of planes. Thus, for example, the first page in planes **0-3** is identified by page line **0**. A page is made up of memory cells belonging to the same word line. A block is a group of pages—i.e., all NAND strings that share the same group of word lines (a NAND string is a group of NAND cells connected in series). In some NAND configurations, a block is a smallest erasable unit. A page is a smallest addressable unit for reading and writing. A plane is a group of physical blocks on a single NAND die, configured for operation such that physical blocks from each of multiple planes can be erased in parallel (i.e., during a given time interval the physical blocks can be erased essentially simultaneously, or in overlap with one another), but only a single physical block in any individual plane can be erased at any one time. There may be multiple planes per NAND die. As shown in FIGS. **5-8**, **10**, and **12-14** a plane is represented by a single physical block (chosen from the list of physical blocks of that plane)—so for example, in FIG. 5, the plane depicts the list of pages in that selected physical block, but there are additional physical blocks not shown for purposes of clarity.

[0070] Disclosed in some examples are improvements to NAND devices that provide additional data protection through an improved NAND data placement schema that allows for recovery from the failure scenarios described in FIG. 5 and FIG. 6. The present disclosure stripes data portions across page lines and planes to ensure that a power failure or programming error affecting an entire page line or plane does not corrupt the entire data item and at most corrupts a single portion of the data item. In some examples, parity information may be calculated and stored until the programming is finished. This parity information may be utilized to recover from corruption of a portion of the data item.

[0071] For example, the NAND may receive a data item from a host device. This data item may be split into a number of portions. For purposes of

the parity values of these data items: **902-924** can be calculated and stored in volatile storage such as RAM, as shown in FIG. 9. [0082] At time **T1**, the portions of data items P.sub.13-P.sub.24 are written to page lines 12-23 and the corresponding parities are calculated and stored in RAM, as shown in FIG. 9. In some examples, the parity values **902-924** are overwritten with the parity values **926-948**. In some examples, the parity values **902-924** may be written to NAND before they are overwritten, for example, to a reliable SLC block. Similarly, at time **T2** parity values for pages 25-36 can be calculated and stored as those pages are written.

[0083] FIG. 10 shows a schematic of a parity page storage **1000** used to store the parity pages in nonvolatile NAND. In the example of FIG. 10, the parity page storage **1000** may be configured as SLC blocks for increased reliability. The parity page storage **1000** as shown in FIG. 10 may be in a separate location from the location that stores the user data (e.g., TLC NAND array **700** from FIG. 7) from which the parity pages are generated. In other examples, the parity pages may be stored on a same die as the user data from which the parity pages are generated. Parity pages 1-12 are stored at time **T0**. Flushed page line count (FPC) indicates the number of page lines of programmed user data (e.g., on die **700**). Parity pages 13-24 are stored at time **T1**, and so on until parity pages 205-216 are stored at time **T17**.

[0084] The parity data stored in RAM or in the SLC **1000** may be used to recover a page of user data. Turning back to FIG. 7, if plane **1** is corrupted (as shown in FIG. 5), only a single portion of the P.sub.n data items may be corrupted by virtue of the failure on plane **1**. The system may utilize the parity data to recover these portions. An individual portion of the user data may be recoverable by applying an XOR operation on the remaining pages and the parity page. For example, if Plane **1** is corrupted, the second portion of the P.sub.1 data item **715** may be corrupted. This portion may be recovered by XORing the first portion of P.sub.1 **710**, third portion P.sub.1 **720**, fourth portion P.sub.1, and the parity **P1 902**. Similarly, if a page line is corrupted as shown in FIG. 6, only a single portion of the user data page may be lost. For example, if any one of page lines 0-3 of FIG. 7 are corrupted, only a single portion of a particular user page may be lost. For example, if the pages on page line 1 are corrupted, then the second portion of P.sub.1 **715** is corrupted, but the first, second, and fourth portions are not and as a result, the third portion may be reconstructed using the parity value. For example, by:

[00002] Second Portion $P_1 = \text{First Portion } P_1 \oplus \text{Third Portion } P_1 \oplus \text{Fourth Portion } P_1 \oplus \text{Parity } P_1$

[0085] Thus, both horizontal corruption across planes on a same page line and vertical corruption affecting all page lines of a plane can be recoverable as a result of the positional rotation of the portions of the user data.

[0086] FIG. 11 shows a flowchart of a method **1100** of applying an improved NAND data placement schema to storing data on a NAND die of a NAND device according to some examples of the present disclosure. A host device may send a data item to write to the NAND. This data item may be programmed to the NAND in a number of portions—a first portion, second portion, third portion, and fourth portion. Example portions may correspond to an upper page, lower page, and extra page of a TLC NAND. In other examples, the data item may be portioned into only two portions (corresponding to a lower page and upper page of an MLC NAND), or four portions (corresponding to a lower page, upper page, middle page, and extra page for a QLC NAND). In other examples, the portions may not correspond to the programming phases of the NAND, but may be divided in other ways (e.g., the most significant bits, least significant bits, and the like). In some examples, the programming page of the NAND that corresponds to the portion may shift based upon a grouping of data items. Thus, in a first grouping of data items, the lower page is written on the first page line of the group and on a later grouping of data items, the extra page may be written on the first page line of the group and so on.

[0087] At operation **1102** the controller may program a first portion of the received data into the NAND array at a first page line and a first plane. At operation **1104** the controller may program a second portion of the received data into the NAND array at a second page line and a second plane. At operation **1106** the controller may program a third portion of the received data into the NAND array at a third page line and a third plane. At operation **1107** the controller may program a fourth portion of the received data into the NAND array at a fourth page line and a fourth plane. In some examples, the fourth portion is a copy of one of the first, second, third, or fourth portions. For example, the fourth portion may be a copy of the portion written first (e.g., a first portion) for the group. The first page line, second page line, third page line, fourth page line, first plane, second plane, third plane, and fourth plane may be selected such that the first portion, second portion, third portion, and fourth portions of a particular data item are programmed into memory cells that are on different page lines and are on different planes with respect to each other. In some examples, all the locations are on the same die. At operation **1108**, the NAND memory device may calculate a parity value for the data item using the first, second, and third portions. This parity value may be stored in volatile, or non-volatile memory.

[0088] While the above data placement schema minimizes data loss by strategic placement of data and the use of parity values, the parity values may be discarded upon completion of the NAND block programming for the data to which the parity corresponds. While this saves the overhead of storing the parity data, this parity data may be leveraged along with the same data placement schema to recover from an Uncorrectable Error Correction Code (UECC) situation in which the NAND cannot recover the data with traditional ECC. For example, as shown in FIG. 12, plane **1**, page lines **Z** to **Z+5** may suffer from an unrecoverable ECC error.

[0089] By storing the data in the improved NAND data placement schema above, losing a single portion of the data is recoverable with the parity value. This may be costly in terms of additional overhead for storing the parity values. Disclosed in some examples, are methods, systems, memory devices, and machine-readable mediums for clustered parity storage for NAND devices utilizing the above disclosed improved data placement schemas. Rather than storing each parity value, parity values from multiple data items may be combined using an XOR operation and stored in a compressed form to reduce the overprovisioning utilized to save the parity values. A compressed parity value may be a consolidation of multiple parity values, for example, an XOR combination of two or more parity values.

[0090] For example, FIG. 13 shows an SLC NAND block **1300** storing a plurality of parity values that are arranged in a plurality of clusters. In FIG. 13, they are arranged in 16 clusters of 54 page lines of four planes for each cluster. Each cluster thus contains 216 parity values. Each parity value corresponds to a data item (e.g., from FIG. 10). A first compressed parity value may be created by XORing parity values in a same position from each different cluster. For example, a compressed parity value may be created from an XOR of the parity value 1, parity value 217, and so on until parity value 3241. This is represented by the dark line in FIG. 13. A second compressed parity value may be created by XORing the second parity value in each cluster. For example, parity 2, parity 218, and so on until parity value 3242. Clusters **2-14** are not shown for compactness and clarity but are included in the XOR calculations.

[0091] More generally, a set of compressed parity values may be created. Each parity value in each cluster may be referred to by using a relative position in the cluster using a notation: (cluster, page, plane). For example, parity value **1** in FIG. 13 may be (0, 0, 0) which indicates this parity value is in cluster **0**, page **0**, and plane **0**. The page values may be relative to the cluster i.e., parity 3241 may be addressed by (15, 0, 0) as it is in the first page and the first plane of cluster **15** even though it is in page line **810** overall. The compressed parity values may be calculated by XORing the parity values of same relative positions of all clusters. For example:

[00003] Compressed Parity $(X, Y) = (0, X, Y) \oplus (1, X, Y) \oplus (2, X, Y) \oplus (3, X, Y) \oplus (4, X, Y) \oplus (5, X, Y) \oplus (6, X, Y) \oplus (7, X, Y) \oplus (8, X, Y) \oplus (9, X, Y) \oplus (10, X, Y) \oplus \dots$

As shown in FIG. 13, values of **X** is of a range 0-53 and **Y** is a range of 0-3.

[0092] As shown in FIG. 14, these compressed parity values may be stored in the NAND and may be denoted **PARITYX.sub.a**. In some examples, the compressed parity and/or the uncompressed parity values may be stored in the original NAND block before the original block is closed. The uncompressed values may then be deleted and the space freed up for other values (e.g., the compressed parity values). The group of data items P.sub.n that were used to calculate the parity values that were used to calculate a compressed parity value may be called a compressed parity data item group G.sub.n. For example, if FIG. 13 shows parity values that correspond to data items P.sub.n from FIG. 7, all portions of data items that were used to produce the constituent parity values (e.g., parity 1, parity 217, and so on until parity 3241) for a first compressed parity value may be part of compressed parity data item group G.sub.1. For parity value 1, the portions of the data items may include first P.sub.1, second P.sub.1, third

P.sub.1, and fourth P.sub.1 from FIG. 7.

[0093] Should an uncorrectable ECC error happen, such as shown in FIG. 12, the NAND device may utilize the compressed parity value to recreate the portion of the data item lost. By XORing the compressed parity value and the values of the data item portions in the group G.sub.n (excluding of course the corrupted portion), the corrupted portion of the data item may be recovered. For example, in the case of a first portion of data item P.sub.1 from (FIG. 7) being unrecoverable, the system may XOR the compressed parity value PARITYX1, and the data portions in the group G.sub.1 (including the second, third, and fourth portions of data item P.sub.1) with the exception of the first portion of data item P.sub.1 to recover the data item P.sub.1.

[0094] FIG. 15 illustrates a flow chart of a method of clustering parity values according to some examples of the present disclosure. At operation 1502, incoming data may be stored according to the disclosed improved NAND data placement schema. For example, storing a received data item in memory cells of the NAND array such that a first portion, second portion, third portion, and fourth portion of the data item are stored in memory cells in the array that are on different page lines and different planes with respect to each other as illustrated in FIG. 7. At operation 1504, a parity value may be calculated for a received data item and stored in volatile or non-volatile storage. As noted, the parity value may be the XOR of the portions of the data item. At operation 1506, the NAND device may assign the parity value calculated at operation 1504 into a position of a parity cluster. For example, a first position of the parity cluster. At operation 1508, using a same position of each of the clusters, a compressed parity value may be calculated. For example, by applying an XOR operator to the parity value in a same position in each cluster of parity values.

[0095] This compressed parity value may be utilized to recover a data portion as previously described. For example, by XORing the data portions in the compressed parity data item group (with the exception of the data portion that was corrupted) and the compressed parity, the data portion that was corrupted may be recovered.

[0096] For example if we have the following data items P.sub.n belonging to the first parity value (e.g., first page line and the first plane) in each parity value cluster:

TABLE-US-00002 First Second Third Fourth P.sub.n portion portion portion Portion Parity.sub.n 1 010 111 000 010 111 54 101 001 110 101 111 810 000 110 010 000 100

[0097] Then the compressed parity for the group of P.sub.n data items in the table (G.sub.1) is the XOR of 111, 111, and 100 which is 100. To recover a lost portion of data, say the second portion of P.sub.54 the system calculates the XOR of 010, 111, 000, 010, 101, 110, 101, 000, 110, 010, 000 and 100 (which is the compressed parity). This results in: 001 which is the correct value. Note that the example table and example above is simplified by having only three P values (1, 54, and 810) and leaving out values for P.sub.108 that would be in group G.sub.1 as shown in FIG. 13.

[0098] FIG. 16 illustrates a schematic of a memory controller 1615 according to some examples of the present disclosure. Memory controller 1615 is an example of memory controller 115, memory manager 1625 is an example of memory manager 125, management tables 1630 may be an example of management table 130. Controller 1635 and ECC 1640 may be an example of controller 135 and ECC 140 of FIG. 1. Controller 1635 includes a schema controller 1642 that may determine where to store portions of data items. In some examples, the positioning is in accordance with the NAND data placement schema disclosed herein. For example, placing a first portion of a received data item into the array at a first page line of a plurality of page lines of a NAND and at a first plane of a plurality of planes; programming a second portion of the received data item into the array at a second page line of the plurality of page lines and at a second plane of the plurality of planes; programming a third portion of the received data item into the array at a third page line of the plurality of page lines and at a third plane of the plurality of planes; calculating a parity value for the data item using the first portion, second portion, and third portion; and wherein the first page line, second page line, third page line, first plane, second plane, and third plane are selected such that the first portion, second portion and third portion are stored in memory cells that are on different page lines and different planes with respect to each other. For example, schema controller 1642 may position data in the NAND as shown in FIGS. 7 and 8.

[0099] In some examples, the schema controller 1642 may also calculate one or more parity values. In some examples, the parity values may be calculated by a hardware XOR processor.

[0100] The schema controller 1642 may allocate and assign parity values to volatile memory locations (as shown in FIG. 9), non-volatile memory locations (as shown in FIG. 10) and the like. Schema controller 1642 may assign the parity values to clusters and utilize a parity value from each cluster to create a compressed parity value and that compressed parity value may be stored. The parity value utilized from each cluster may be selected based upon a formula or other schema. For example, each parity value may be assigned a relative position in each cluster (e.g., a first parity value in a cluster, a second parity value, and so on) and a same position in each cluster may be utilized to create the compressed parity value. The schema controller 1642 may store the compressed parity values in volatile or non-volatile memory (e.g., as shown in FIG. 14). Schema controller 1642 may implement the methods of FIG. 11 and FIG. 15.

[0101] FIG. 17 illustrates a block diagram of an example machine 1700 upon which any one or more of the techniques (e.g., methodologies) discussed herein may perform. In alternative embodiments, the machine 1700 may operate as a standalone device or may be connected (e.g., networked) to other machines. In a networked deployment, the machine 1700 may operate in the capacity of a server machine, a client machine, or both in server-client network environments. In an example, the machine 1700 may act as a peer machine in peer-to-peer (P2P) (or other distributed) network environment. The machine 1700 may be a personal computer (PC), a tablet PC, a set-top box (STB), a personal digital assistant (PDA), a mobile telephone, a web appliance, an IoT device, automotive system, or any machine capable of executing instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single machine is illustrated, the term "machine" shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein, such as cloud computing, software as a service (SaaS), other computer cluster configurations.

[0102] Examples, as described herein, may include, or may operate by, logic, components, devices, packages, or mechanisms. Circuitry is a collection (e.g., set) of circuits implemented in tangible entities that include hardware (e.g., simple circuits, gates, logic, etc.). Circuitry membership may be flexible over time and underlying hardware variability. Circuitries include members that may, alone or in combination, perform specific tasks when operating. In an example, hardware of the circuitry may be immutably designed to carry out a specific operation (e.g., hardwired). In an example, the hardware of the circuitry may include variably connected physical components (e.g., execution units, transistors, simple circuits, etc.) including a computer readable medium physically modified (e.g., magnetically, electrically, moveable placement of invariant massed particles, etc.) to encode instructions of the specific operation. In connecting the physical components, the underlying electrical properties of a hardware constituent are changed, for example, from an insulator to a conductor or vice versa. The instructions enable participating hardware (e.g., the execution units or a loading mechanism) to create members of the circuitry in hardware via the variable connections to carry out portions of the specific tasks when in operation. Accordingly, the computer readable medium is communicatively coupled to the other components of the circuitry when the device is operating. In an example, any of the physical components may be used in more than one member of more than one circuitry. For example, under operation, execution units may be used in a first circuit of a first circuitry at one point in time and reused by a second circuit in the first circuitry, or by a third circuit in a second circuitry at a different time.

[0103] The machine (e.g., computer system) 1700 (e.g., the host device 105, the memory device 110, etc.) may include a hardware processor 1702 (e.g., a central processing unit (CPU), a graphics processing unit (GPU), a hardware processor core, or any combination thereof, such as the memory controller 115, etc.), a main memory 1704 and a static memory 1706, some or all of which may communicate with each other via an interlink (e.g., bus) 1708. The machine 1700 may further include a display unit 1710, an alphanumeric input device 1712 (e.g., a keyboard), and a user interface (UI) navigation device 1714 (e.g., a mouse). In an example, the display unit 1710, input device 1712 and UI navigation device 1714 may be a touch screen display. The machine 1700 may additionally include a storage device (e.g., drive unit) 1716, a signal generation device 1718 (e.g., a speaker), a network interface device 1720, and one or more sensors 1716, such as a global positioning system (GPS) sensor, compass, accelerometer, or other

sensor. The machine **1700** may include a communication controller **1728**, such as a serial bus (USB), parallel, or other wired or wireless (e.g., infrared (IR), near field communication (NFC), etc.) connection to communicate or control one or more peripheral devices (e.g., a printer, card reader, etc.).

[0104] The storage device **1716** may include a machine readable medium **1722** on which is stored one or more sets of data structures or instructions **1724** (e.g., software) embodying or utilized by any one or more of the techniques or functions described herein. The instructions **1724** may also reside, completely or at least partially, within the main memory **1704**, within static memory **1706**, or within the hardware processor **1702** during execution thereof by the machine **1700**. In an example, one or any combination of the hardware processor **1702**, the main memory **1704**, the static memory **1706**, or the storage device **1716** may constitute the machine readable medium **1722**.

[0105] While the machine readable medium **1722** is illustrated as a single medium, the term “machine readable medium” may include a single medium or multiple media (e.g., a centralized or distributed database, or associated caches and servers) configured to store the one or more instructions **1724**.

[0106] The term “machine readable medium” may include any medium capable of storing, encoding, or carrying instructions for execution by the machine **1700** and that cause the machine **1700** to perform any one or more of the techniques of the present disclosure, or capable of storing, encoding or carrying data structures used by or associated with such instructions. Non-limiting machine readable medium examples may include solid-state memories, and optical and magnetic media. In an example, a massed machine readable medium comprises a machine-readable medium with a plurality of particles having invariant (e.g., rest) mass. Accordingly, massed machine-readable media are not transitory propagating signals. Specific examples of massed machine readable media may include: non-volatile memory, such as semiconductor memory devices (e.g., Electrically Programmable Read-Only Memory (EPROM), Electrically Erasable Programmable Read-Only Memory (EEPROM)) and flash memory devices; magnetic disks, such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks.

[0107] The instructions **1724** (e.g., software, programs, an operating system (OS), etc.) or other data are stored on the storage device **1721**, can be accessed by the memory **1704** for use by the processor **1702**. The memory **1704** (e.g., DRAM) is typically fast, but volatile, and thus a different type of storage than the storage device **1721** (e.g., an SSD), which is suitable for long-term storage, including while in an “off” condition. The instructions **1724** or data in use by a user or the machine **1700** are typically loaded in the memory **1704** for use by the processor **1702**. When the memory **1704** is full, virtual space from the storage device **1721** can be allocated to supplement the memory **1704**; however, because the storage **1721** device is typically slower than the memory **1704**, and write speeds are typically at least twice as slow as read speeds, use of virtual memory can greatly reduce user experience due to storage device latency (in contrast to the memory **1704**, e.g., DRAM). Further, use of the storage device **1721** for virtual memory can greatly reduce the usable lifespan of the storage device **1721**.

[0108] In contrast to virtual memory, virtual memory compression (e.g., the Linux® kernel feature “ZRAM”) uses part of the memory as compressed block storage to avoid paging to the storage device **1721**. Paging takes place in the compressed block until it is necessary to write such data to the storage device **1721**. Virtual memory compression increases the usable size of memory **1704**, while reducing wear on the storage device **1721**.

[0109] Storage devices optimized for mobile electronic devices, or mobile storage, traditionally include MMC solid-state storage devices (e.g., micro Secure Digital (microSD™) cards, etc.). MMC devices include a number of parallel interfaces (e.g., an 8-bit parallel interface) with a host device, and are often removable and separate components from the host device. In contrast, eMMC™ devices are attached to a circuit board and considered a component of the host device, with read speeds that rival serial ATA™ (Serial AT (Advanced Technology) Attachment, or SATA) based SSD devices. However, demand for mobile device performance continues to increase, such as to fully enable virtual or augmented-reality devices, utilize increasing networks speeds, etc. In response to this demand, storage devices have shifted from parallel to serial communication interfaces. Universal Flash Storage (UFS) devices, including controllers and firmware, communicate with a host device using a low-voltage differential signaling (LVDS) serial interface with dedicated read/write paths, further advancing greater read/write speeds.

[0110] The instructions **1724** may further be transmitted or received over a communications network **1726** using a transmission medium via the network interface device **1720** utilizing any one of a number of transfer protocols (e.g., frame relay, internet protocol (IP), transmission control protocol (TCP), user datagram protocol (UDP), hypertext transfer protocol (HTTP), etc.). Example communication networks may include a local area network (LAN), a wide area network (WAN), a packet data network (e.g., the Internet), mobile telephone networks (e.g., cellular networks), Plain Old Telephone (POTS) networks, and wireless data networks (e.g., Institute of Electrical and Electronics Engineers (IEEE) 802.11 family of standards known as Wi-Fi®, IEEE 802.16 family of standards known as WiMax®, IEEE 802.15.4 family of standards, peer-to-peer (P2P) networks, among others. In an example, the network interface device **1720** may include one or more physical jacks (e.g., Ethernet, coaxial, or phone jacks) or one or more antennas to connect to the communications network **1726**. In an example, the network interface device **1720** may include a plurality of antennas to wirelessly communicate using at least one of single-input multiple-output (SIMO), multiple-input multiple-output (MIMO), or multiple-input single-output (MISO) techniques. The term “transmission medium” shall be taken to include any intangible medium capable of storing, encoding or carrying instructions for execution by the machine **1700**, and includes digital or analog communications signals or other intangible medium to facilitate communication of such software.

[0111] The above detailed description includes references to the accompanying drawings, which form a part of the detailed description. The drawings show, by way of illustration, specific embodiments in which the invention can be practiced. These embodiments are also referred to herein as “examples”. Such examples can include elements in addition to those shown or described. However, the present inventors also contemplate examples in which only those elements shown or described are provided. Moreover, the present inventors also contemplate examples using any combination or permutation of those elements shown or described (or one or more aspects thereof), either with respect to a particular example (or one or more aspects thereof), or with respect to other examples (or one or more aspects thereof) shown or described herein.

[0112] In this document, the terms “a” or “an” are used, as is common in patent documents, to include one or more than one, independent of any other instances or usages of “at least one” or “one or more.” In this document, the term “or” is used to refer to a nonexclusive or, such that “A or B” may include “A but not B,” “B but not A,” and “A and B,” unless otherwise indicated. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein”. Also, in the following claims, the terms “including” and “comprising” are open-ended, i.e., a system, device, article, or process that includes elements in addition to those listed after such a term in a claim are still deemed to fall within the scope of that claim. Moreover, in the following claims, the terms “first,” “second,” and “third,” etc. are used merely as labels, and are not intended to impose numerical requirements on their objects.

[0113] In various examples, the components, controllers, processors, units, engines, or tables described herein can include, among other things, physical circuitry or firmware stored on a physical device. As used herein, “processor” means any type of computational circuit such as, but not limited to, a microprocessor, a microcontroller, a graphics processor, a digital signal processor (DSP), or any other type of processor or processing circuit, including a group of processors or multi-core devices.

[0114] The term “horizontal” as used in this document is defined as a plane parallel to the conventional plane or surface of a substrate, such as that underlying a wafer or die, regardless of the actual orientation of the substrate at any point in time. The term “vertical” refers to a direction perpendicular to the horizontal as defined above. Prepositions, such as “on,” “over,” and “under” are defined with respect to the conventional plane or surface being on the top or exposed surface of the substrate, regardless of the orientation of the substrate; and while “on” is intended to suggest a direct contact of one structure relative to another structure which it lies “on” (in the absence of an express indication to the contrary); the terms “over” and “under” are expressly intended to identify a relative placement of structures (or layers, features, etc.), which expressly includes—but is not limited to—direct contact between the identified structures unless specifically identified as such. Similarly, the terms “over” and “under” are not limited to horizontal orientations, as a structure may be “over” a referenced structure if it is, at some point in time, an outermost portion of the

under discussion, even if such structure extends vertically relative to the referenced structure, rather than in a horizontal orientation. [0115] The terms “wafer” and “substrate” are used herein to refer generally to any structure on which integrated circuits are formed, and also to such structures during various stages of integrated circuit fabrication. The following detailed description is, therefore, not to be taken in a limiting sense, and the scope of the various embodiments is defined only by the appended claims, along with the full scope of equivalents to which such claims are entitled.

[0116] Various embodiments according to the present disclosure and described herein include memory utilizing a vertical structure of memory cells (e.g., NAND strings of memory cells). As used herein, directional adjectives will be taken relative a surface of a substrate upon which the memory cells are formed (i.e., a vertical structure will be taken as extending away from the substrate surface, a bottom end of the vertical structure will be taken as the end nearest the substrate surface and a top end of the vertical structure will be taken as the end farthest from the substrate surface).

[0117] As used herein, directional adjectives, such as horizontal, vertical, normal, parallel, perpendicular, etc., can refer to relative orientations, and are not intended to require strict adherence to specific geometric properties, unless otherwise noted. For example, as used herein, a vertical structure need not be strictly perpendicular to a surface of a substrate, but may instead be generally perpendicular to the surface of the substrate, and may form an acute angle with the surface of the substrate (e.g., between 60 and 120 degrees, etc.).

[0118] In some embodiments described herein, different doping configurations may be applied to a source-side select gate (SGS), a control gate (CG), and a drain-side select gate (SGD), each of which, in this example, may be formed of or at least include polysilicon, with the result such that these tiers (e.g., polysilicon, etc.) may have different etch rates when exposed to an etching solution. For example, in a process of forming a monolithic pillar in a 3D semiconductor device, the SGS and the CG may form recesses, while the SGD may remain less recessed or even not recessed. These doping configurations may thus enable selective etching into the distinct tiers (e.g., SGS, CG, and SGD) in the 3D semiconductor device by using an etching solution (e.g., tetramethylammonium hydroxide (TMCH)).

[0119] Operating a memory cell, as used herein, includes reading from, writing to, or erasing the memory cell. The operation of placing a memory cell in an intended state is referred to herein as “programming,” and can include both writing to or erasing from the memory cell (e.g., the memory cell may be programmed to an erased state).

[0120] According to one or more embodiments of the present disclosure, a memory controller (e.g., a processor, controller, firmware, etc.) located internal or external to a memory device, is capable of determining (e.g., selecting, setting, adjusting, computing, changing, clearing, communicating, adapting, deriving, defining, utilizing, modifying, applying, etc.) a quantity of wear cycles, or a wear state (e.g., recording wear cycles, counting operations of the memory device as they occur, tracking the operations of the memory device it initiates, evaluating the memory device characteristics corresponding to a wear state, etc.)

[0121] According to one or more embodiments of the present disclosure, a memory access device may be configured to provide wear cycle information to the memory device with each memory operation. The memory device control circuitry (e.g., control logic) may be programmed to compensate for memory device performance changes corresponding to the wear cycle information. The memory device may receive the wear cycle information and determine one or more operating parameters (e.g., a value, characteristic) in response to the wear cycle information.

[0122] It will be understood that when an element is referred to as being “on,” “connected to” or “coupled with” another element, it can be directly on, connected, or coupled with the other element or intervening elements may be present. In contrast, when an element is referred to as being “directly on,” “directly connected to” or “directly coupled with” another element, there are no intervening elements or layers present. If two elements are shown in the drawings with a line connecting them, the two elements can be either be coupled, or directly coupled, unless otherwise indicated.

[0123] Method examples described herein can be machine or computer-implemented at least in part. Some examples can include a computer-readable medium or machine-readable medium encoded with instructions operable to configure an electronic device to perform methods as described in the above examples. An implementation of such methods can include code, such as microcode, assembly language code, a higher-level language code, or the like. Such code can include computer readable instructions for performing various methods. The code may form portions of computer program products. Further, the code can be tangibly stored on one or more volatile or non-volatile tangible computer-readable media, such as during execution or at other times. Examples of these tangible computer-readable media can include, but are not limited to, hard disks, removable magnetic disks, removable optical disks (e.g., compact discs and digital video disks), magnetic cassettes, memory cards or sticks, random access memories (RAMs), read only memories (ROMs), solid state drives (SSDs), Universal Flash Storage (UFS) device, embedded MMC (eMMC) device, and the like.

[0124] The above description is intended to be illustrative, and not restrictive. For example, the above-described examples (or one or more aspects thereof) may be used in combination with each other. Other embodiments can be used, such as by one of ordinary skill in the art upon reviewing the above description. It is submitted with the understanding that it will not be used to interpret or limit the scope or meaning of the claims. Also, in the above Detailed Description, various features may be grouped together to streamline the disclosure. This should not be interpreted as intending that an unclaimed disclosed feature is essential to any claim. Rather, inventive subject matter may lie in less than all features of a particular disclosed embodiment. Thus, the following claims are hereby incorporated into the Detailed Description, with each claim standing on its own as a separate embodiment, and it is contemplated that such embodiments can be combined with each other in various combinations or permutations. The scope of the invention should be determined with reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

OTHER NOTES AND EXAMPLES

[0125] Example 1 is a NAND memory device comprising: an array of NAND memory cells organized into multiple planes and addressable by multiple page lines; and a controller configured to perform operations comprising: programming a first portion of a received data item into the array at a first page line of the multiple page lines and at a first plane of the multiple planes; programming a second portion of the received data item into the array at a second page line of the multiple page lines and at a second plane of the multiple planes; programming a third portion of the received data item into the array at a third page line of the multiple page lines and at a third plane of the multiple planes.

[0126] In Example 2, the subject matter of Example 1, wherein the operations comprise storing a parity value by applying an XOR operator to the first portion, second portion, and third portion.

[0127] In Example 3, the subject matter of Example 2, wherein the operations further comprise storing the parity value in the array of NAND memory cells.

[0128] In Example 4, the subject matter of Example 3, wherein the parity value is stored in a location of the array of NAND memory cells that are configured as Single Level Cell (SLC) memory cells.

[0129] In Example 5, the subject matter of any of Examples 3-4, wherein the first portion, second portion, and third portion are programmed into a location of the array of NAND memory cells that are configured as Triple Level Cell (TLC) memory cells.

[0130] In Example 6, the subject matter of any of Examples 2-5, wherein the operations comprise: determining that one of the first portion, second portion, or third portion was corrupted; and recovering the one of the first portion, second portion, or third portion by utilizing the parity value.

[0131] In Example 7, the subject matter of any of Examples 2-6 includes, a volatile memory in communication with the controller; and wherein the operations further comprise: initially storing the parity value in the volatile memory; and moving the parity value to the array of NAND memory cells.

[0132] In Example 8, the subject matter of any of Examples 1-7, wherein the operations further comprise: programming a fourth portion of the received data item into the array at a fourth page line of the multiple page lines and at a fourth plane of the multiple planes.

[0133] In Example 9, the subject matter of any of Examples 1-8, wherein the first portion, the second portion and the third portion comprise individual bits within the received data.

[0134] In Example 10, the subject matter of any of Examples 1-9, wherein the third page line and third plane are greater than the second page line and second plane, and the second page line and second plane are greater than the first page line and first plane.

[0135] Example 11 is a method of storing data on a NAND device that includes, an array of NAND memory cells organized into multiple planes and addressable by multiple page lines, the method comprising: programming a first portion of a received data item into the array at a first page line of the multiple page lines and at a first plane of the multiple planes; programming a second portion of the received data item into the array at a second page line of the multiple page lines and at a second plane of the multiple planes; programming a third portion of the received data item into the array at a third page line of the multiple page lines and at a third plane of the multiple planes.

[0136] In Example 12, the subject matter of Example 11 includes, storing a parity value by applying an XOR operator to the first portion, second portion, and third portion.

[0137] In Example 13, the subject matter of Example 12 includes, storing the parity value in the array of NAND memory cells.

[0138] In Example 14, the subject matter of Example 13, wherein the parity value is stored in a location of the array of NAND memory cells that are configured as Single Level Cell (SLC) memory cells.

[0139] In Example 15, the subject matter of any of Examples 13-14, wherein the first portion, second portion, and third portion are programmed into a location of the array of NAND memory cells that are configured as Triple Level Cell (TLC) memory cells.

[0140] In Example 16, the subject matter of any of Examples 12-15 includes, determining that one of the first portion, second portion, or third portion was corrupted; and recovering the one of the first portion, second portion, or third portion by utilizing the parity value.

[0141] In Example 17, the subject matter of any of Examples 12-16 includes, initially storing the parity value in a volatile memory; and moving the parity value to the array of NAND memory cells.

[0142] In Example 18, the subject matter of any of Examples 11-17 includes, programming a fourth portion of the received data item into the array at a fourth page line of the multiple page lines and at a fourth plane of the multiple planes.

[0143] In Example 19, the subject matter of any of Examples 11-18, wherein the first portion, the second portion and the third portion comprise individual bits within the received data.

[0144] In Example 20, the subject matter of any of Examples 11-19, wherein the third page line and third plane are greater than the second page line and second plane, and the second page line and second plane are greater than the first page line and first plane.

[0145] Example 21 is a machine-readable medium, comprising instructions, which when executed by a machine, cause the machine to perform operations comprising: programming a first portion of a received data item into a NAND array at a first page line of multiple page lines and at a first plane of multiple planes of the NAND array; programming a second portion of the received data item into the array at a second page line of the multiple page lines and at a second plane of the multiple planes; programming a third portion of the received data item into the array at a third page line of the multiple page lines and at a third plane of the multiple planes.

[0146] In Example 22, the subject matter of Example 21, wherein the operations comprise storing a parity value by applying an XOR operator to the first portion, second portion, and third portion.

[0147] In Example 23, the subject matter of Example 22, wherein the operations further comprise storing the parity value in the array of NAND memory cells.

[0148] In Example 24, the subject matter of Example 23, wherein the parity value is stored in a location of the array of NAND memory cells that are configured as Single Level Cell (SLC) memory cells.

[0149] In Example 25, the subject matter of any of Examples 23-24, wherein the first portion, second portion, and third portion are programmed into a location of the array of NAND memory cells that are configured as Triple Level Cell (TLC) memory cells.

[0150] In Example 26, the subject matter of any of Examples 22-25, wherein the operations comprise: determining that one of the first portion, second portion, or third portion was corrupted; and recovering the one of the first portion, second portion, or third portion by utilizing the parity value.

[0151] In Example 27, the subject matter of any of Examples 22-26 includes, a volatile memory in communication with the controller; and wherein the operations further comprise: initially storing the parity value in the volatile memory; and moving the parity value to the array of NAND memory cells.

[0152] In Example 28, the subject matter of any of Examples 21-27, wherein the operations further comprise: programming a fourth portion of the received data item into the array at a fourth page line of the multiple page lines and at a fourth plane of the multiple planes.

[0153] In Example 29, the subject matter of any of Examples 21-28, wherein the first portion, the second portion and the third portion comprise individual bits within the received data.

[0154] In Example 30, the subject matter of any of Examples 21-29, wherein the third page line and third plane are greater than the second page line and second plane, and the second page line and second plane are greater than the first page line and first plane.

[0155] Example 31 is a NAND memory device comprising: an array of NAND memory cells organized into multiple planes and addressable by multiple page lines; and a controller configured to perform operations comprising: means for programming a first portion of a received data item into the array at a first page line of the multiple page lines and at a first plane of the multiple planes; means for programming a second portion of the received data item into the array at a second page line of the multiple page lines and at a second plane of the multiple planes; means for programming a third portion of the received data item into the array at a third page line of the multiple page lines and at a third plane of the multiple planes.

[0156] In Example 32, the subject matter of Example 31 includes, means for storing a parity value by applying an XOR operator to the first portion, second portion, and third portion.

[0157] In Example 33, the subject matter of Example 32 includes, means for storing the parity value in the array of NAND memory cells.

[0158] In Example 34, the subject matter of Example 33, wherein the parity value is stored in a location of the array of NAND memory cells that are configured as Single Level Cell (SLC) memory cells.

[0159] In Example 35, the subject matter of any of Examples 33-34, wherein the first portion, second portion, and third portion are programmed into a location of the array of NAND memory cells that are configured as Triple Level Cell (TLC) memory cells.

[0160] In Example 36, the subject matter of any of Examples 32-35 includes, means for determining that one of the first portion, second portion, or third portion was corrupted; and means for recovering the one of the first portion, second portion, or third portion by utilizing the parity value.

[0161] In Example 37, the subject matter of any of Examples 32-36 includes, means for initially storing the parity value in a volatile memory; and means for moving the parity value to the array of NAND memory cells.

[0162] In Example 38, the subject matter of any of Examples 31-37 includes, means for programming a fourth portion of the received data item into the array at a fourth page line of the multiple page lines and at a fourth plane of the multiple planes.

[0163] In Example 39, the subject matter of any of Examples 31-38, wherein the first portion, the second portion and the third portion comprise individual bits within the received data.

[0164] In Example 40, the subject matter of any of Examples 31-39, wherein the third page line and third plane are greater than the second page line and second plane, and the second page line and second plane are greater than the first page line and first plane.

[0165] Example 41 is a NAND memory device comprising: an array of NAND memory cells organized into multiple planes and addressable by multiple page lines; and a controller configured to perform operations comprising: storing a received data item in memory cells of the NAND array such that a first portion, second portion and third portion of the data item are stored in memory cells in the array such that are on different page lines and different planes with respect to each other; calculating a parity value for the received data item using the first portion, second portion, and third

portion; assigning the parity value for the received data item to a first position of a parity cluster of multiple parity clusters; and calculating a compressed parity value based upon the parity value and a second parity value of a second parity cluster of the multiple parity clusters.

[0166] In Example 42, the subject matter of Example 41, wherein the operations of calculating the compressed parity value based upon the parity value and the second parity value of the second parity cluster comprises selecting the second parity value from the second parity cluster based upon a relative position of the second parity value in the second parity cluster matching the relative position of the parity value in the first parity cluster.

[0167] In Example 43, the subject matter of any of Examples 41-42, wherein the parity value is stored in a block of the NAND memory cells.

[0168] In Example 44, the subject matter of any of Examples 42-43, wherein the operations comprise storing the compressed parity value in a block of the NAND memory cells.

[0169] In Example 45, the subject matter of any of Examples 43-44, wherein the operations comprise overwriting the parity value with the compressed parity values.

[0170] In Example 46, the subject matter of any of Examples 41-45, wherein the operations comprise: calculating the second parity value based upon first, second, and third portions of a second data item; receiving an indication that a first portion of the data item read from the array of NAND memory cells failed an Error Correction Code check; recovering the first portion using the compressed parity value, the second and third portions of the data item, and the first, second, and third portions of the second data item.

[0171] In Example 47, the subject matter of any of Examples 41-46, wherein the operations of calculating the compressed parity value comprises applying an XOR operation to the parity value and the second parity value.

[0172] In Example 48, the subject matter of any of Examples 41-47, wherein the operations of calculating the parity value for the received data item using the first portion, second portion, and third portion comprises applying an XOR operator to the first portion, second portion, and third portion.

[0173] Example 49 is a machine-readable medium, comprising instructions, which when executed by a machine, cause the machine to perform operations comprising: storing a received data item in memory cells of a NAND array such that a first portion, second portion and third portion of the data item are stored in memory cells in the array that they are on different page lines and different planes with respect to each other; calculating a parity value for the received data item using the first portion, second portion, and third portion; assigning the parity value for the received data item into a first position of a parity cluster of multiple parity clusters; and calculating a compressed parity value based upon the parity value and a second parity value of a second parity cluster of the multiple parity clusters.

[0174] In Example 50, the subject matter of Example 49, wherein the operations of calculating the compressed parity value based upon the parity value and the second parity value of the second parity cluster comprises selecting the second parity value from the second parity cluster based upon a relative position of the second parity value in the second parity cluster matching the relative position of the parity value in the first parity cluster.

[0175] In Example 51, the subject matter of any of Examples 49-50, wherein the parity value is stored in a block of the NAND memory cells.

[0176] In Example 52, the subject matter of any of Examples 50-51, wherein the operations further comprise storing the compressed parity value in a block of the NAND memory cells.

[0177] In Example 53, the subject matter of any of Examples 51-52, wherein the operations further comprise overwriting the parity value with the compressed parity values.

[0178] In Example 54, the subject matter of any of Examples 49-53, wherein the operations further comprise: calculating the second parity value based upon first, second, and third portions of a second data item; receiving an indication that a first portion of the data item read from the array of NAND memory cells failed an Error Correction Code check; recovering the first portion using the compressed parity value, the second and third portions of the data item, and the first, second, and third portions of the second data item.

[0179] In Example 55, the subject matter of any of Examples 49-54, wherein the operations of calculating the compressed parity value comprises applying an XOR operation to the parity value and the second parity value.

[0180] In Example 56, the subject matter of any of Examples 49-55, wherein the operations of calculating the parity value for the received data item using the first portion, second portion, and third portion comprises applying an XOR operator to the first portion, second portion, and third portion.

[0181] Example 57 is a method of storing data on a NAND device that includes, an array of NAND memory cells organized into multiple planes and addressable by multiple page lines, the method comprising: storing a received data item in memory cells of a NAND array such that a first portion, second portion and third portion of the data item are stored in memory cells in the array that they are on different page lines and different planes with respect to each other; calculating a parity value for the received data item using the first portion, second portion, and third portion; assigning the parity value for the received data item into a first position of a parity cluster of multiple parity clusters; and calculating a compressed parity value based upon the parity value and a second parity value of a second parity cluster of the multiple parity clusters.

[0182] In Example 58, the subject matter of Example 57, wherein calculating the compressed parity value based upon the parity value and the second parity value of the second parity cluster comprises selecting the second parity value from the second parity cluster based upon a relative position of the second parity value in the second parity cluster matching the relative position of the parity value in the first parity cluster.

[0183] In Example 59, the subject matter of any of Examples 57-58, wherein the parity value is stored in a block of the NAND memory cells.

[0184] In Example 60, the subject matter of any of Examples 58-59 includes, storing the compressed parity value in a block of the NAND memory cells.

[0185] In Example 61, the subject matter of any of Examples 59-60 includes, overwriting the parity value with the compressed parity values.

[0186] In Example 62, the subject matter of any of Examples 57-61 includes, calculating the second parity value based upon first, second, and third portions of a second data item; receiving an indication that a first portion of the data item read from the array of NAND memory cells failed an Error Correction Code check; recovering the first portion using the compressed parity value, the second and third portions of the data item, and the first, second, and third portions of the second data item.

[0187] In Example 63, the subject matter of any of Examples 57-62, wherein calculating the compressed parity value comprises applying an XOR operation to the parity value and the second parity value.

[0188] In Example 64, the subject matter of any of Examples 57-63, wherein calculating the parity value for the received data item using the first portion, second portion, and third portion comprises applying an XOR operator to the first portion, second portion, and third portion.

[0189] Example 65 is a NAND memory device comprising: an array of NAND memory cells organized into multiple planes and addressable by multiple page lines; and a controller configured to perform operations comprising: means for storing a received data item in memory cells of a NAND array such that a first portion, second portion and third portion of the data item are stored in memory cells in the array that they are on different page lines and different planes with respect to each other; means for calculating a parity value for the received data item using the first portion, second portion, and third portion; means for assigning the parity value for the received data item into a first position of a parity cluster of multiple parity clusters; and means for calculating a compressed parity value based upon the parity value and a second parity value of a second parity cluster of the multiple parity clusters.

[0190] In Example 66, the subject matter of Example 65, wherein the means for calculating the compressed parity value based upon the parity value and the second parity value of the second parity cluster comprises means for selecting the second parity value from the second parity cluster based upon a relative position of the second parity value in the second parity cluster matching the relative position of the parity value in the first parity cluster.

[0191] In Example 67, the subject matter of any of Examples 65-66, wherein the parity value is stored in a block of the NAND memory cells.

[0192] In Example 68, the subject matter of any of Examples 66-67 includes, means for storing the compressed parity value in a block of the NAND memory cells.

[0193] In Example 69, the subject matter of any of Examples 67-68 includes, means for overwriting the parity value with the compressed parity values.

[0194] In Example 70, the subject matter of any of Examples 65-69 includes, means for calculating the second parity value based upon first, second, and third portions of a second data item; means for receiving an indication that a first portion of the data item read from the array of NAND memory cells failed an Error Correction Code check; and means for recovering the first portion using the compressed parity value, the second and third portions of the data item, and the first, second, and third portions of the second data item.

[0195] In Example 71, the subject matter of any of Examples 65-70, wherein the means for calculating the compressed parity value comprises means for applying an XOR operation to the parity value and the second parity value.

[0196] In Example 72, the subject matter of any of Examples 65-71, wherein the means for calculating the parity value for the received data item using the first portion, second portion, and third portion comprises means for applying an XOR operator to the first portion, second portion, and third portion.

[0197] Example 73 is at least one machine-readable medium including instructions that, when executed by processing circuitry, cause the processing circuitry to perform operations to implement any of Examples 1-72.

[0198] Example 74 is an apparatus comprising means to implement any of Examples 1-72.

[0199] Example 75 is a system to implement any of Examples 1-72.

[0200] Example 76 is a method to implement any of Examples 1-72.

Claims

1. A NAND memory device comprising: an array of NAND memory cells organized into multiple planes and addressable by multiple page lines; and a controller configured to perform operations comprising: receiving a data item from a host over a host interface; diagonally striping the data item across pages and dies of NAND memory cells by: programming a first portion of the received data item into an array of NAND memory cells at a first page of a first die; programming a second portion of the received data item into the array of NAND memory cells at a second page of a second die; programming a third portion of the received data item into the array of NAND memory cells at a third page of a third die; and wherein each data portion of the received data item is programmed on a different page line than every other data portion; calculating a first parity value for the received data item using the first portion, the second portion, and the third portion; assigning the first parity value for the received data item into a first position of a first parity cluster; calculating a compressed parity value using the first parity value and a second parity value of a second parity cluster, the second parity value occupying a same first position in the second parity cluster as the first parity value, the second parity value a parity value of a first, a second, and a third portion of a second data item; and storing the compressed parity value in a portion of a NAND block that was used to store either the first or second parity value.
2. The NAND memory device of claim 1, wherein calculating the first parity value comprises applying an XOR operation to the first portion, second portion, and third portion of the received data item.
3. The NAND memory device of claim 1, wherein calculating the compressed parity value comprises applying an XOR operation to the first parity value and the second parity value.
4. The NAND memory device of claim 1, wherein the operations further comprise: storing the first parity value initially in volatile memory; and moving the first parity value to a block of NAND memory cells configured as Single Level Cell (SLC) memory cells.
5. The NAND memory device of claim 1, wherein the operations further comprise: detecting an uncorrectable error indicating that one of the first portion, second portion, or third portion stored in the array of NAND memory cells was corrupted; identifying a compressed parity data item group comprising the received data item and the second data item whose parity values were used to calculate the compressed parity value; retrieving uncorrupted portions of both the received data item and the second data item from the compressed parity data item group; retrieving the compressed parity value; recovering the corrupted portion by performing a parity operation on the compressed parity value, all uncorrupted portions of the received data item; and all portions of all other data items in the compressed parity data item group including the second parity value; and storing the recovered portion back to the array of NAND memory cells to replace the corrupted portion.
6. The NAND memory device of claim 1, wherein the operations further comprise organizing multiple parity values into clusters, wherein each cluster contains parity values for a predetermined number of data items.
7. The NAND memory device of claim 1, wherein the operations further comprise: calculating multiple compressed parity values by applying XOR operations to parity values occupying matching relative positions across different parity clusters; and storing the multiple compressed parity values in the NAND block.
8. The NAND memory device of claim 1, wherein the operations further comprise: detecting a programming failure affecting a first plane of NAND memory cells; identifying affected portions of data items stored in the first plane; and recovering each affected portion using corresponding compressed parity values and unaffected portions of data items in respective compressed parity data item groups.
9. A method for managing data in a NAND memory device, the method comprising: using one or more computer processors: receiving a data item from a host over a host interface; diagonally striping the data item across pages and dies of NAND memory cells by: programming a first portion of the received data item into an array of NAND memory cells at a first page of a first die; programming a second portion of the received data item into the array of NAND memory cells at a second page of a second die; programming a third portion of the received data item into the array of NAND memory cells at a third page of a third die; and wherein each data portion of the received data item is programmed on a different page line than every other data portion; calculating a first parity value for the received data item using the first portion, the second portion, and the third portion; assigning the first parity value for the received data item into a first position of a first parity cluster; calculating a compressed parity value using the first parity value and a second parity value of a second parity cluster, the second parity value occupying a same first position in the second parity cluster as the first parity value, the second parity value a parity value of a first, a second, and a third portion of a second data item; and storing the compressed parity value in a portion of a NAND block that was used to store either the first or second parity value.
10. The method of claim 9, wherein the method further comprises: applying an XOR operation to the first portion, second portion, and third portion of the received data item to calculate the first parity value.
11. The method of claim 9, wherein the method further comprises: applying an XOR operation to the first parity value and the second parity value to calculate the compressed parity value.
12. The method of claim 9, wherein the method further comprises: storing the first parity value initially in volatile memory; and moving the first parity value to a block of NAND memory cells configured as Single Level Cell (SLC) memory cells.
13. The method of claim 9, wherein the method further comprises: detecting an uncorrectable error indicating that one of the first portion, second portion, or third portion stored in the array of NAND memory cells was corrupted; identifying a compressed parity data item group comprising the received data item and the second data item whose parity values were used to calculate the compressed parity value; retrieving uncorrupted portions of both the received data item and the second data item from the compressed parity data item group; retrieving the compressed parity value; recovering the corrupted portion by performing a parity operation on the compressed parity value, all uncorrupted portions of the received data item, and all portions of all other data items in the compressed parity data item group including the second parity value; and storing the recovered portion back to the array of NAND memory cells to replace the corrupted portion.
14. The method of claim 9, wherein the method further comprises: organizing multiple parity values into clusters, wherein each cluster contains

parity values for a predetermined number of data items.

15. The method of claim 9, wherein the method further comprises: calculating multiple compressed parity values by applying XOR operations to parity values occupying matching relative positions across different parity clusters; and storing the multiple compressed parity values in the NAND block.

16. The method of claim 9, wherein the method further comprises: detecting a programming failure affecting a first plane of NAND memory cells; identifying affected portions of data items stored in the first plane; and recovering each affected portion using corresponding compressed parity values and unaffected portions of data items in respective compressed parity data item groups.

17. A non-transitory machine-readable medium, storing instructions for managing data in a NAND memory device, the instructions, which when executed, cause a machine to perform operations comprising: receiving a data item from a host over a host interface; diagonally striping the data item across pages and dies of NAND memory cells by: programming a first portion of the received data item into an array of NAND memory cells at a first page of a first die; programming a second portion of the received data item into the array of NAND memory cells at a second page of a second die; programming a third portion of the received data item into the array of NAND memory cells at a third page of a third die; and wherein the operation of programming each data portion of the received data item further comprises programming each portion on a different page line than every other data portion; calculating a first parity value for the received data item using the first portion, the second portion, and the third portion; assigning the first parity value for the received data item into a first position of a first parity cluster; calculating a compressed parity value using the first parity value and a second parity value of a second parity cluster, the second parity value occupying a same first position in the second parity cluster as the first parity value, the second parity value a parity value of a first, a second, and a third portion of a second data item; and storing the compressed parity value in a portion of a NAND block that was used to store either the first or second parity value.

18. The non-transitory machine-readable medium of claim 17, wherein the operation of calculating the first parity value further comprises applying an XOR operation to the first portion, second portion, and third portion of the received data item.

19. The non-transitory machine-readable medium of claim 17, wherein the operation of calculating the compressed parity value further comprises applying an XOR operation to the first parity value and the second parity value.

20. The non-transitory machine-readable medium of claim 17, wherein the operations further comprise: storing the first parity value initially in volatile memory; and moving the first parity value to a block of NAND memory cells configured as Single Level Cell (SLC) memory cells.
