

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250278278

Kind Code

A1

Publication Date

September 04, 2025

Inventor(s)

Prasky; Brian Robert et al.

BRANCH PREDICTION USING SPECULATIVE INDEXING AND INTRALINE COUNT

Abstract

A method of branch prediction in a processor includes: generating a new line index and an intraline index; generating a first output of a pattern-based predictor structure using the new line index; generating a second output of the pattern-based predictor structure using the intraline index; selecting one of the first output and the second output based on a result of a prediction pipeline; and predicting a direction of a branch using the selected one of the first output and the second output.

Inventors: Prasky; Brian Robert (Campbell Hall, NY), Bonanno; James (Liberty Hill, TX), Collura; Adam Benjamin (Hopewell Junction, NY), Malley; Edward Thomas (Eastchester, NY), Bhattacharjee; Deepankar (Bangalore, IN)

Applicant: INTERNATIONAL BUSINESS MACHINES CORPORATION (Armonk, NY)

Family ID: 88095872

Appl. No.: 19/210296

Filed: May 16, 2025

Related U.S. Application Data

parent US division 17703063 20220324 parent-grant-document US 12327122 child US 19210296

Publication Classification

Int. Cl.: G06F9/38 (20180101)

U.S. Cl.:

CPC G06F9/3844 (20130101);

Background/Summary

BACKGROUND

[0001] Aspects of the present invention relate generally to computer processing and, more particularly, to branch prediction in a processor.

[0002] Branch prediction is a type of metadata prediction that is a performance-critical component of a pipelined high frequency microprocessor. It is used to predict the direction (taken vs. not taken) and the target address of branch instructions. This is beneficial because it allows processing to continue along a branch's predicted path rather than having to wait for the outcome of the branch to be determined. A penalty is incurred only if a branch is mis-predicted.

[0003] A branch target buffer (BTB) is a structure that stores branch and target information. Other structures such as a branch history table (BHT), pattern history table (PHT), and multiple target table (MTT), can be included to store additional information used for branch direction and target prediction.

[0004] The BTB can be searched in parallel to and independently from instruction fetching to find upcoming branches, in which case it is called lookahead branch prediction. Alternatively, it can be accessed simultaneously with or after fetching instructions and determining instruction boundaries in order to provide a prediction for each encountered branch instruction, in which case it is called synchronous branch prediction. In either case the performance benefit of the BTB is a function of the accuracy of the prediction provided from the BTB and the latency required to access the BTB.

[0005] The output of a predictor, like the BTB, can be read and cached for reuse to make predictions repeatedly while staying in the same section of code that can be called a line of code. This is a prediction cache or line output buffer (LOB). Such a design consists of an index pipeline to index the BTB and read its content and put it into the LOB, and a prediction pipeline to make predictions from the LOB content.

[0006] Pattern-based metadata predictors utilize history of prior instructions to make predictions about current ones. Traditionally this history consists of the branch direction history or taken branch instruction addresses. The history vector can be combined with the branch instruction address to form an index into a pattern history table (PHT) or multiple target table (MTT).

SUMMARY

[0007] In a first aspect of the invention, there is a method of branch prediction in a processor, the method including: obtaining, by the processor, a branch instruction for which a direction of a branch is to be predicted; generating, by the processor, an index based on an instruction address, a global path vector (GPV), and a counter; selecting, by the processor, an entry from a data structure using the index; and predicting, by the processor, the direction of the branch using information included in the selected entry.

[0008] In another aspect of the invention, there is a method of branch prediction in a processor, the method including: generating a new line index and an intraline index; generating a first output of a pattern-based predictor structure using the new line index; generating a second output of the pattern-based predictor structure using the intraline index; selecting one of the first output and the second output based on a result of a prediction pipeline; and predicting a direction of a branch using the selected one of the first output and the second output.

[0009] In another aspect of the invention, there is computer system for branch prediction in a processing pipeline, the system including a memory and a processor operatively coupled to the memory. The computer system is configured to: generate a first output of a pattern-based predictor structure and a second output of the pattern-based predictor structure by generating a new line index of the pattern-based predictor structure and an intraline index of the pattern-based predictor structure using a line entry instruction address, a global path vector (GPV) leading up to the line

entry instruction address, and an intraline counter of predicted taken branches, wherein the pattern-based predictor structure has a first read port associated with the new line index and a second read port associated with the intraline index.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0010] Aspects of the present invention are described in the detailed description which follows, in reference to the noted plurality of drawings by way of non-limiting examples of exemplary embodiments of the present invention.

[0011] FIG. 1 depicts one example of a computing environment to incorporate and use one or more aspects of a branch prediction capability in accordance with aspects of the invention.

[0012] FIG. 2 depicts another example of a computing environment to incorporate and use one or more aspects of a branch prediction capability in accordance with aspects of the invention.

[0013] FIG. 3A depicts yet another example of a computing environment to incorporate and use one or more aspects of a branch prediction capability in accordance with aspects of the invention.

[0014] FIG. 3B depicts further details of the memory of FIG. 3A.

[0015] FIG. 4 depicts one example of a processing system to use branch prediction in accordance with aspects of the invention.

[0016] FIG. 5 shows an example of logic for generating an index used with a pattern history table in accordance with aspects of the invention.

[0017] FIG. 6 shows an example of an anchor point, global path vector, and counter in accordance with aspects of the invention.

[0018] FIG. 7 shows an example of modifying a global path vector in accordance with aspects of the invention.

[0019] FIG. 8 shows an example of logic for generating a speculative intraline index of pattern-based predictors in accordance with aspects of the invention.

[0020] FIG. 9 shows a block diagram for generating an index used with a pattern history table in accordance with aspects of the invention.

[0021] FIG. 10 shows a block diagram for generating a new line index and an intraline index used with a pattern history table in accordance with aspects of the invention.

[0022] FIG. 11 shows a table of cycles of a prediction pipeline in accordance with aspects of the invention.

[0023] FIG. 12 shows a flowchart of an exemplary method in accordance with aspects of the invention.

[0024] FIG. 13 shows a flowchart of an exemplary method in accordance with aspects of the invention.

DETAILED DESCRIPTION

[0025] Aspects of the present invention relate generally to computer processing and, more particularly, to branch prediction in a processor. Pattern history tables (PHT) are often used in branch prediction. In general, a PHT selects a static random short length of branch history for table indexing and records the outcome of a given branch based on the pattern to get to the given branch. Another system used in branch prediction is TAGE (TAGged GEometric length predictor), which stores a tag along with the history prediction for each branch entry. TAGE supports multiple tables each with a different pattern history. In TAGE, longer history lengths may hash reduce into fewer bits of indexing, which introduces aliasing within the given history table.

[0026] In such systems, history lengths are picked at coordinated randomness and are static for the life of the hardware. The static nature is less than optimal, however, and it would be advantageous to use a dynamic history length and not be limited to a fixed history length.

[0027] Using more tables results in a larger storage footprint and slower learning. It would be advantageous to provide a system that uses a smaller footprint and learns faster.

[0028] Some systems employ plural tables to make plural different predictions for a single branch and then use logic to select which one of the predictions to use. In such systems, a first table may be based on a global history path and a second table may be based on a counter value past an anchoring point. The concept of counting past a given anchor point has been shown to provide benefit. However, the use of plural tables to make plural predictions for a single branch is still inefficient. Accordingly, it would be advantageous to provide a system that defines anchor points in a manner that leverages the counting scheme but that does not utilize plural tables to make plural predictions for a single branch.

[0029] Moreover, no conventional system provides a way to incorporate pattern-based predictors into a system with a primary predictor cache. Because pattern-based predictors typically require history updated with all previous branch predictions, they cannot be indexed until after such predictions are made. A straightforward implementation would require waiting for such predictions to be made before indexing the pattern tables.

[0030] An advantage of a metadata prediction cache is to avoid latency of having to re-index the primary predictor, such as a BTB. But it does not provide a way to incorporate pattern-based tables without incurring extra latency for re-indexing them after a prediction is made.

[0031] Implementations of the invention address these issues by indexing pattern-based predictors as a function of history state and address information upon entering into a line of code along with an intraline count updated while making predictions that stay within the line. This provides a way to speculatively index the pattern-based structures with multiple read ports before it is known whether current predictions will stay intraline or require a new line. This provides for using prediction pipeline indications to select among the outputs from the multiple read ports and either use that information immediately in the prediction pipeline or hold it for later when needed. In this manner, implementations of the invention advantageously allow the incorporation of pattern-based predictors in systems including a line-based primary predictor cache without requiring extra latency for using the pattern-based predictors.

[0032] Implementations of the invention also address the issues described above by defining an anchor point as the entry point into a region of code (such as a line of code), and by generating a PHT index based on three factors including: an instruction address (IA) of the entry point; a count value past the anchoring point; and a global history of taken branches leading into the entry point. In embodiments, global history knowledge from within the line is not accounted for in forming the index for non-zero count values. In embodiments, the system modifies the tag contained with the PHT from being only branch IA bits to including knowledge of the branch path taken to the anchor point. This may include the most recent bits of a global path vector (GPV) or may be a hash of taken IA (e.g., bits **50-62**) of all taken branches since that of the anchor point.

[0033] In this manner, implementations of the invention provide a branch predictor that generates a PHT index based on: a function of line entry; leveraging a GPV that is up to the branches that got to the anchoring line of interest; and introducing a count value that is XOR'd into the above values. In embodiments, as the count value increases, the amount of GPV may be reduced. In embodiments, the system modifies the PHT tag to include path vector knowledge up to the branch of interest that is being predicted, which may be a GPV of the last "X" number of branches. In some embodiments, the anchor point is based on static regions of code, such as a 128B line of code. The region size of code may be larger than the line size. In some embodiments, the region size is not static in nature but is defined by dynamic events such as a multi-target branch and/or a flush denoting a new boundary.

[0034] Computing environments of different architectures may incorporate and use one or more aspects of the branch prediction capability provided herein. For instance, environments based on the z/Architecture, offered by International Business Machines Corporation (IBM®) may include

one or more aspects, as well as computing environments of other architectures, including but not limited to, the PowerPC architecture, also referred to as Power ISA, also offered by International Business Machines Corporation. IBM, Z/ARCHITECTURE, Z/OS, Z/VM, z/VSE, POWER, and POWERPC (referenced herein) are registered trademarks of International Business Machines Corporation, Armonk, N.Y. Other names used herein may be registered trademarks, trademarks or product names of International Business Machines Corporation or other companies.

[0035] One example of a computing environment to incorporate and use one or more aspects of the branch prediction capability is described with reference to FIG. 1. In one example, a computing environment **100** includes a processor (e.g., central processing unit, CPU, etc.) **102** that includes at least one memory management unit **104** to manage the memory and provide, for instance, address translation; one or more caches **106**; a branch predictor **107** to facilitate predicting the direction of branches; and one or more branch prediction tables **109** used by the branch predictor to predict the direction of a branch (referred to herein as branch prediction). In one embodiment, the branch predictor includes the branch prediction logic, while the branch prediction tables are separate. However, in other embodiments, the branch predictor includes the branch prediction logic and the branch prediction tables.

[0036] Processor **102** is communicatively coupled to a memory portion **108** having one or more caches **110**, and to an input/output (I/O) subsystem **112**. I/O subsystem **112** is communicatively coupled to external I/O devices and data **114** that may include, for example, data input devices, sensors and/or output devices, such as displays.

[0037] Another embodiment of a computing environment to incorporate and use one or more aspects of the branch prediction capability is described with reference to FIG. 2. Referring to FIG. 2, in one example, a computing environment **200** is based on the z/Architecture, offered by International Business Machines Corporation, Armonk, N.Y. Computing environment **200** includes, for example, a central processor complex (CPC) **202** coupled to one or more input/output (I/O) devices **206** via one or more control units **208**. Central processor complex **202** includes, for instance, a processor memory **204** (a.k.a., main memory, main storage, central storage) coupled to one or more central processors (a.k.a., central processing units (CPUs)) **210**, and an input/output subsystem **211**, each of which is described below.

[0038] Processor memory **204** includes, for example, one or more partitions **212** (e.g., logical partitions), and processor firmware **213**, which includes, for instance, a logical partition hypervisor **214** and other processor firmware **215**. One example of logical partition hypervisor **214** is the Processor Resource/System Manager (PR/SM), offered by International Business Machines Corporation, Armonk, N.Y.

[0039] A logical partition functions as a separate system and has one or more applications **220**, and optionally, a resident operating system **222** therein, which may differ for each logical partition. In one embodiment, the operating system is the z/OS operating system, the z/VM operating system, the z/VSE operating system, or the TPF operating system, offered by International Business Machines Corporation, Armonk, N.Y.

[0040] Logical partitions **212** are managed by logical partition hypervisor **214**, which is implemented by firmware running on processors **210**. As used herein, firmware includes, e.g., the microcode and/or millicode of the processor. It includes, for instance, the hardware-level instructions and/or data structures used in implementation of higher-level machine code. In one embodiment, it includes, for instance, proprietary code that is typically delivered as microcode that includes trusted software or microcode specific to the underlying hardware and controls operating system access to the system hardware.

[0041] Central processors **210** are physical processor resources allocated to the logical partitions. In particular, each logical partition **212** has one or more logical processors, each of which represents all or a share of a physical processor **210** allocated to the partition. The logical processors of a particular partition **212** may be either dedicated to the partition, so that the underlying processor

resource **210** is reserved for that partition; or shared with another partition, so that the underlying processor resource is potentially available to another partition.

[0042] Central processors **210** include, in one embodiment, a branch predictor **224** used along with branch prediction tables **226** to predict the direction of branches included in instruction streams. Again, the branch prediction tables may be considered a part of the branch predictor or separate therefrom.

[0043] Input/output subsystem **211** directs the flow of information between input/output devices **206** and main storage **204**. It is coupled to the central processing complex, in that it can be a part of the central processing complex or separate therefrom. The I/O subsystem relieves the central processors of the task of communicating directly with the input/output devices and permits data processing to proceed concurrently with input/output processing. To provide communications, the I/O subsystem employs I/O communications adapters. There are various types of communications adapters including, for instance, channels, I/O adapters, PCI cards, Ethernet cards, Small Computer Storage Interface (SCSI) cards, etc. In the particular example described herein, the I/O communications adapters are channels, and therefore, the I/O subsystem is referred to herein as a channel subsystem. However, this is only one example. Other types of I/O subsystems can be used.

[0044] The I/O subsystem uses one or more input/output paths as communication links in managing the flow of information to or from input/output devices **206**. In this particular example, these paths are called channel paths, since the communication adapters are channels.

[0045] Another embodiment of a computing environment to incorporate and use one or more aspects of the branch prediction capability is described with reference to FIG. 3A. In this example, a computing environment **300** includes, for instance, a native central processing unit (CPU) **302**, a memory **304**, and one or more input/output devices and/or interfaces **306** coupled to one another via, for example, one or more buses **308** and/or other connections. As examples, computing environment **300** may include a PowerPC processor, a pSeries server or an xSeries server offered by International Business Machines Corporation, Armonk, N.Y.

[0046] Native central processing unit **302** includes one or more native registers **310**, such as one or more general purpose registers and/or one or more special purpose registers used during processing within the environment. These registers include information that represent the state of the environment at any particular point in time.

[0047] Moreover, native central processing unit **302** executes instructions and code that are stored in memory **304**. In one particular example, the central processing unit executes emulator code **312** stored in memory **304**. This code enables the computing environment configured in one architecture to emulate another architecture. For instance, emulator code **312** allows machines based on architectures other than the z/Architecture, such as PowerPC processors, pSeries servers, xSeries servers or others, to emulate the z/Architecture and to execute software and instructions developed based on the z/Architecture.

[0048] Further details relating to emulator code **312** are described with reference to FIG. 3B. Guest instructions **350** stored in memory **304** comprise software instructions (e.g., correlating to machine instructions) that were developed to be executed in an architecture other than that of native CPU **302**. For example, guest instructions **350** may have been designed to execute on a z/Architecture processor **102**, but instead, are being emulated on native CPU **302**, which may be, for example, an Intel Itanium II processor. In one example, emulator code **312** includes an instruction fetching routine **352** to obtain one or more guest instructions **350** from memory **304**, and to optionally provide local buffering for the instructions obtained. It also includes an instruction translation routine **354** to determine the type of guest instruction that has been obtained and to translate the guest instruction into one or more corresponding native instructions **356**. This translation includes, for instance, identifying the function to be performed by the guest instruction and choosing the native instruction(s) to perform that function.

[0049] Further, emulator code **312** includes an emulation control routine **360** to cause the native

instructions to be executed. Emulation control routine **360** may cause native CPU **302** to execute a routine of native instructions that emulate one or more previously obtained guest instructions and, at the conclusion of such execution, return control to the instruction fetch routine to emulate the obtaining of the next guest instruction or a group of guest instructions. Execution of the native instructions **356** may include loading data into a register from memory **304**; storing data back to memory from a register; or performing some type of arithmetic or logic operation, as determined by the translation routine.

[0050] Each routine is, for instance, implemented in software, which is stored in memory and executed by native central processing unit **302**. In other examples, one or more of the routines or operations are implemented in firmware, hardware, software or some combination thereof. The registers of the emulated processor may be emulated using registers **310** of the native CPU or by using locations in memory **304**. In embodiments, guest instructions **350**, native instructions **356** and emulator code **312** may reside in the same memory or may be disbursed among different memory devices.

[0051] The computing environments described above are only examples of computing environments that can be used. Other environments, including but not limited to, other non-partitioned environments, other partitioned environments, and/or other emulated environments, may be used; embodiments are not limited to any one environment.

[0052] Each of the computing environments described above may include a processing system that may be included within a computer processor or otherwise distributed within a computing environment. In one example, referring to FIG. **4**, a processing system **400** includes a memory **402**, an instruction cache **404**, an instruction fetch unit **408**, and a processing pipeline **406**. Instructions and data may be stored in memory **402**, and instruction cache **404** may access instructions in memory **402** and store the instructions to be fetched. Memory **402** may include any type of volatile or non-volatile memory, such as cache memory. Memory **402** and instruction cache **404** may include multiple cache levels. A data cache (not depicted) may also be included in processing system **400**.

[0053] Instruction fetch unit **408** includes a branch predictor **418** that includes the branch prediction logic and, in one embodiment, the branch prediction tables of one or more aspects described herein. Alternatively, branch predictor **418** may be located apart from instruction fetch unit **408**, and/or the tables may be separate therefrom. In one example, the branch predictor **418** comprises a lookahead asynchronous branch predictor.

[0054] Processing pipeline **406** includes a decode unit **410**, an issue unit **412**, an execution stage **414**, and write-back logic **416**. The entire instruction fetch unit **408** or branch predictor **418** may also be part of processing pipeline **406**. Processing pipeline **406** can include other features, such as error checking and handling logic, one or more parallel paths through processing pipeline **406**, and other features known in the art. Further, processing system **400** can include multiple processing pipelines **406** and instruction fetch units **408**. While a forward path through processing system **400** is depicted in FIG. **4**, other feedback and signaling paths may be included between elements of processing system **400**. Decode unit **410** decodes instructions and passes the decoded instructions, portions of the instructions, or other decoded data to issue unit **412**. Issue unit **412** analyzes the instructions or other data and transmits the decoded instructions, portions of the instructions, or other data to one or more execution units in execution stage **414** based on the analysis. Execution stage **414** executes the instructions, and may include a plurality of execution units, such as fixed-point execution units, floating point execution units, load/store execution units, and vector execution units, as examples. Write-back logic **416** writes results of instruction execution back to a destination resource **420**. Destination resource **420** may be any type of resource, including, for instance, registers, cache memory, other memory, I/O circuitry to communicate with other devices, other processing circuits, or any other type of destination for executed instructions or data.

[0055] In accordance with one or more aspects, a branch predictor (e.g., branch predictor **107**,

branch predictor **224**, or branch predictor **418**, as examples) uses a branch prediction structure to predict the direction of a branch. In embodiments, the branch prediction structure comprises a pattern history table (PHT). In embodiments, the branch predictor may include plural PHTs; however, according to aspects of the invention, only a single PHT is used to make any given branch prediction. Each PHT is implemented, in one example, as an n-way (e.g., 4-way) set associative structure used to hold a saturation counter for each branch to be predicted. A saturation counter records a 4-state value to enable branch direction to be predicted. The four states include, for instance: 0—strongly not taken; 1—weakly not taken; 2—weakly taken; and 3—strongly taken. As each branch is taken, the saturation counter is incremented; an increment to a value of, e.g., 3 saturates and remains at 3. When a branch that is encountered is not taken, the saturation counter is decremented; a decrement to a value of, e.g., 0 saturates and remains at 0.

[0056] In embodiments, the PHT is used to predict whether a branch is to be taken or not taken. It is indexed using a function of the address of the branch to be predicted and the branch history (i.e., the outcome of the immediately preceding branches). In embodiments, the index is generated using a function of: (i) an address of a region of code that includes the branch being predicted, (ii) a global history of branches leading into the region of code that includes the branch being predicted, and (iii) a counter that is based on a number of predicted taken branches within the region of code that includes the branch being predicted.

[0057] FIG. 5 shows an embodiment of logic **500** for generating an index used with a pattern history table (PHT) when making a branch prediction in accordance with aspects of the invention. The logic may be included in a branch predictor (e.g., branch predictor **107**, branch predictor **224**, or branch predictor **418**, as examples) and used to predict the direction (e.g., taken versus not taken) of a branch instruction that is encountered in the processing pipeline. In one example, the branch predictor comprises a lookahead asynchronous branch predictor.

[0058] In accordance with aspects of the invention, the logic **500** includes a PHT **505** and an index generator **508**. In embodiments, the PHT **505** includes plural entries (e.g., 0 through 1023) where each entry includes a saturation counter (SAT CTR) and a tag (TAG). The saturation counter may comprise a 2-bit value that indicates one of four states including strongly not taken, weakly not taken, weakly taken, and strongly taken, as described herein. In embodiments, the tag of a given entry is used to compare against the address of the branch, which is being predicted, hashed with equivalent content at present to which was used to generate the tag initially. Based on a tag match, the table is deemed to have an entry for the branch that is to be predicted. Prediction logic **509** may use the output of the PHT **505** to make the prediction for a particular branch, e.g., taken versus not taken, using any suitable technique.

[0059] In accordance with aspects of the invention, the index generator **508** generates an index **510** at branch prediction time for reading the PHT **505** or at update time for writing to the PHT **505**. In embodiments, the index generator **508** generates the index **510** using function of: (i) an address of a region of code that includes the branch being predicted, represented by address **511**; (ii) a global history of branches leading into the region of code that includes the branch being predicted, represented by global path vector (GPV) **512**; and (iii) a counter that is based on a number of predicted taken branches within the region of code, represented by counter **513**. In embodiments, the function used by the index generator **508** is a 3-way exclusive OR (XOR) function that receives the address **511**, GPV **512**, and counter **513** as inputs and generates an index **510** as the output, although implementations are not limited to an XOR and other functions (e.g., such as other hash functions) may be used.

[0060] In embodiments, the address **511** comprises a portion of an instruction address of a region of code that includes the branch being predicted. In one example, the region is defined as a 128-byte (128B) line of the code that includes the branch being predicted, although implementations are not limited to this particular region and other regions may be used as described herein. In one example, the address **511** comprises a portion of the instruction address of the 128B line itself, for example,

when the line is arrived at sequentially. In another example, the address **511** comprises a portion of the instruction address of an entry point into the 128B line, for example, when the line is arrived at due to a branching operation. In embodiments, the portion of an instruction address is a predefined number of bits of the instruction address. In one example, the instruction address is a 64-bit address of the 128B line, and the predefined number of bits includes bits **53:62** of the 64-bit address. As such, using this particular example, the address **511** comprises bits **53:62** of the 64-bit address of the entry point into the 128B line of code that contains the branch being predicted. In embodiments, a branch target buffer (BTB) is a structure that stores branch and target information including the definition of each 128B line of the code.

[0061] In embodiments, the GPV **512** comprises a vector that contains information representing a history of taken branches leading into the region of code. In one example, when a taken branch is encountered in the pipeline during execution of an application, an identifier of the taken branch, which may comprise, for example, the instruction address or the target instruction address of the taken branch, is input to a history generator function. In this example, the history generator function generates an output that is based on the instruction address or the target instruction address of the taken branch and this output is then shifted into the GPV **512**. In one example, the output of the history generator function comprises a number of bits of the instruction address or the target instruction address of the taken branch. In a particular example, the target instruction address of the taken branch is a 64-bit address, and the output of the history generator function comprises bit **62** of the target instruction address of the taken branch. In this particular example, the GPV **512** includes “X” number of bits where each bit is bit **62** of a target instruction address of a taken branch leading into the region of code that includes the branch being predicted. In embodiments, older bits in the GPV **512** are shifted out as new bits are shifted in. In this manner, the GPV **512** comprises a vector of “X” number of bits that represent a sequential history of the last (e.g., most recent) “X” number of taken branches to arrive at the region of code that includes the branch being predicted. According to aspects of the invention, the GPV **512** includes information only for taken branches leading into the region of code that includes the branch being predicted, and does not include information regarding taken branches within the region of code that includes the branch being predicted.

[0062] In embodiments, the counter **513** comprises a count of a number of predicted taken branches within the region of code that includes the branch being predicted. In embodiments, the counter starts at a value of zero and is incremented each time a branch within this same region of code is predicted taken. In this manner, the counter **513** represents a number of branches past a point of interest. In embodiments, the point of interest is referred to as an anchor point and may comprise, for example, a beginning of the 128B line containing the branch being predicted or an entry point into the 128B line containing the branch being predicted. In some embodiments, as described herein, the anchor point is dynamically based on an event.

[0063] FIG. **6** shows an example of an anchor point, GPV, and counter in accordance with aspects of the invention. In embodiments, data structure **605** represents a history of taken branches in sequential order from oldest (on the left) to most recent (on the right). Each cell in the data structure **605** may contain a bit that represents a taken branch, such as bit **62** of a target instruction address of a taken branch as described above. In the example shown in FIG. **6**, anchor point **610** represents an entry point into the region of code that includes the branch being predicted. In this example, the portion of the data structure **605** to the left of the anchor point **610** represents the history of taken branches leading into the region of code that includes the branch being predicted, and the portion of the data structure to the right of the anchor point **610** represents predicted taken branches within the region of code that includes the branch being predicted. In this example, the number “X” of the GPV **512** equals 10.

[0064] In the example shown in FIG. **6**, at time **t1**, no taken branches have been predicted past the anchor point **610**. As such, at time **t1** the GPV **512** includes the 10 bits immediately preceding the

anchor point **610** and the counter **513** equals 0.

[0065] In this example, at time **t2**, the system has predicted 1 taken branch past the anchor point **610**. As such, at time **t2** the GPV **512** includes the 10 bits immediately preceding the anchor point **610** and the counter **513** equals 1.

[0066] In this example, at time **t3**, the system has predicted 2 taken branches past the anchor point **610**. As such, at time **t3** the GPV **512** includes the 10 bits immediately preceding the anchor point **610** and the counter **513** equals 2.

[0067] As illustrated by the example shown in FIG. 6, in aspects the GPV **512** stays static while the counter **513** increments for each predicted taken branch with the region of code past the anchor point. This is in contrast to systems where the GPV is changed with each branch prediction. In embodiments, when a new anchor point is established (e.g., for a different 128B line of code), the counter **513** resets to zero and the GPV **512** is modified to include the “X” number of bits that represent a sequential history of the last “X” number of taken branches to arrive at the new anchor point.

[0068] In accordance with further aspects of the invention, a second GPV referred to as the TagGPV **515** represents the last “X” number of taken branches inclusive of those after the anchor point **610**. As illustrated in the example shown in FIG. 6, the TagGPV **515** changes with each successive predicted taken branch with the region of code (i.e., after the anchor point **610**).

[0069] Referring back to FIG. 5, in embodiments the logic **500** includes a tag generator **520** that operates to generate a tag (TAG) that is stored in the PHT **505**. In embodiments, the tag generator **520** generates the tag using a function of the branch address **516** and the TagGPV **515**. In embodiments, the function is an XOR function although different hash functions may be used. In embodiments, the branch address **516** is based on the instruction address of the branch being predicted. In one example, using a 64-bit address, the branch address **516** comprises bits **50:62** of the instruction address of the branch being predicted. In this manner, the tag includes information on both the branch being predicted and the path taken to arrive at the branch being predicted, as opposed to containing only information about the branch being predicted. In some embodiments, the counter **513** is not limited such that it may achieve values that exceed the number of entries in the PHT **505**. In these embodiments, the use of a tag generated by the tag generator **520** using the TagGPV **515** alleviates aliasing that might otherwise occur if the tag was based only on the branch being predicted.

[0070] In accordance with some aspects, the GPV **512** is modified as the counter **513** increments. In one example, the number of entries included in the GPV **512** is reduced as the counter **513** increments. FIG. 7 illustrates an example in which the GPV **512'** is reduced by one unit each time the counter increments. For example, the GPV **512'** includes 10 bits when the counter **513** equals 0 at time **t1**. In this example, the GPV **512'** is reduced in size to 9 bits when the counter **513** increments to 1 at time **t2**. Continuing this example, the GPV **512'** is reduced in size to 8 bits when the counter increments to 8 at time **t3**, and so on. In this example, modifying the GPV **512'** includes removing (or zeroing) the bit that corresponds to the oldest taken branch contained in the GPV **512'** (e.g., as shown at the left side of the GPV **512'** in FIG. 7) for each increment of the counter **513**. Other modifications of the GPV **512** based on the incrementing of the counter **513** may be used.

[0071] Aspects of the invention have been described thus far with each region of code being defined in a static nature, e.g., a 128B line of code as defined in the BTB. However, implementations may also utilize dynamically determined regions of code. In one example, the anchor point for a region of code may be defined by a resolved branch that was predicted incorrectly which resulted in a branch flush restart of the prediction pipeline. In this example, the address **511** would be based on the instruction address of the resolved branch, the GPV **512** would include “X” number of bits that represent a sequential history of the last “X” number of taken branches to arrive at the new anchor point (i.e., the resolved branch), and the counter **513** would reset to zero beginning at the new anchor point and increment for each predicted taken branch after

the new anchor point. In another example, the anchor point for a dynamically determined region of code may be defined by a point of restart after a branch flush restart of the prediction pipeline. In this example, the point of restart may be a next instruction address after a resolved branch that was predicted incorrectly and that resulted in the branch flush restart. In this example, the next instruction address after the resolved branch may be a next sequential instruction address or a target of the resolved branch. Another example of a dynamically determined region of code uses an anchor point that is based on saturation of the counter **513**. In this example, when the counter **513** reaches a predefined upper limit, the system creates a new anchor point at that branch, adjusts the GPV **512** based on the new anchor point, and resets the counter **513** to zero.

[0072] Aspects of the invention may utilize plural different PHTs each having a different way of defining its anchoring point. For example, a first PHT may be associated with an anchoring point that is based on the most recent branch wrong prediction, a second PHT may be associated with an anchoring point that is based on the second most recent branch wrong prediction, a third PHT may be associated with an anchoring point that is based on the third most recent branch wrong prediction, etc. The example described above includes different tables all using a same type of dynamically determined anchoring point, e.g., an anchoring point based on a branch wrong prediction. However, different tables may use different types of dynamically determined anchoring points. For example, a first PHT may be associated with an anchoring point that is based on the most recent branch wrong prediction, and a second PHT may be associated with an anchoring point that is based on function calls that exhibit multiple targets. Such implementations provide for the use of dynamic history lengths instead of static history lengths, such as those used in TAGE (TAGged GEometric length predictor) systems. Using plural tables with different anchoring points, and thus using dynamic history lengths, provides faster learning and greater accuracy while minimizing the total number of required tables.

[0073] FIG. **8** shows an example of logic for generating a speculative intraline index of pattern-based predictors in accordance with aspects of the invention. The embodiment shown in FIG. **8** is in the context of a branch predictor with a primary branch target buffer (BTB) **805** predictor with a line output buffer (LOB) **810** prediction cache. A pattern history table (PHT) **815** is used as the example of a pattern-based predictor. Embodiments of the invention are not limited to only these types of predictors. Other types of metadata predictors (beyond branch predictors) can be employed in implementations of the invention. Other types of pattern-based predictors are possible. Multiple PHTs are possible in a TAGE configuration. Multiple target table (MTT) predictors may also be included as pattern-based predictors utilizing aspects described here.

[0074] In a lookahead asynchronous branch predictor, the branch prediction logic gets restarted from the processor pipeline, then runs ahead on its own providing prediction information to the pipeline and redirecting itself upon finding taken branches.

[0075] A primary predictor such as a branch predictor's BTB **805** is read on a line basis and its output is written into a prediction cache, the LOB **810**, to be reused while making predictions staying within that line.

[0076] In embodiments, line addresses are put into a line index buffer (LIB) **820**. They come from restart addresses (from the processor pipeline or the prediction pipeline), a line index accelerator **825**, or sequential addresses. The index pipeline (e.g., comprising LIB **820**, BTB **805**, and LOB **810**) uses addresses from the LIB **820** to index and read content from the BTB **805** and put the results of those reads into the LOB **810**.

[0077] In embodiments, the prediction pipeline comprises PHT index logic **830**, PHT **815**, PHT output buffer **835**, and prediction logic **840**. In embodiments, the prediction pipeline uses content of the LOB **810** to make predictions between a starting search address and the first predicted taken branch (if there is one). If there is a predicted taken branch that stays within the same line, the next pass in the prediction pipeline reuses the current LOB entry. Otherwise, the next LOB entry is used for the next pass of the prediction pipeline, and it is either the content for the target line following a

predicted taken branch to a new line, or the content for a sequential line following either no predictions or only not-taken predictions.

[0078] In embodiments, and as described herein, the PHT index logic **830** generates two indexes for a branch prediction: a new line index **841** and an intraline index **842**. In embodiments, the PHT **815** includes two read ports, one for each index. In embodiments, the output data of the PHT **815** includes two PHT outputs: PHT output **851** based on the new line index **841** and PHT output **852** based on the intraline index **842**. In embodiments, the system generates the two indexes and the two corresponding PHT outputs because often, at the time of generating the indexes and the outputs, the prediction logic **840** does not yet know whether it will utilize a prediction based on staying within the line or a prediction based on a new line. By generating both indexes and both PHT outputs ahead of time in this manner, both PHT outputs are ready for use by the prediction logic **840** immediately when the prediction logic **840** determines whether to stay in the line or go to a new line, and the system selects which of the PHT outputs to use at this stage. In this manner, implementations of the invention increase the speed of the prediction pipeline.

[0079] FIG. **9** shows a block diagram for generating an index used with a pattern history table (e.g., PHT **815** of FIG. **8**) in accordance with aspects of the invention. In implementations, the PHT index logic **905** (which corresponds to PHT index logic **830** of FIG. **8**) generates a PHT index as a function of the path taken to get to the branch. In this manner, the PHT is indexed as a function of the branch's instruction address, as represented by element **910**, and a global path vector of hashed instruction addresses of prior branches. However instead of including all prior branches in the history, the state of the history upon entering into a line of code is used, as represented by element **915**. Then when in that line, an additional factor, an intraline count, as represented by element **920**, is incorporated. This count can be the count of taken predicted branches within the line. Upon entering the line, it is 0. It is incremented by 1 upon each subsequent predicted taken branch. One advantage of using the line-entry history status in combination with an intraline count is to make it possible to update the index by incrementing the count while staying within a line. In this manner, the index is not dependent on specific information about the branches predicted while in the line, but only dependent on the number of taken intraline branches predicted. In one example, elements **905**, **910**, **915**, and **920** of FIG. **9** are implemented using index generator **508**, address **511**, GPV **512**, and counter **513**, respectively, of FIG. **5**. In embodiments, the PHT index **925** generated by the index generator **905** is the new line index **841** or the intraline index **842**, depending on which data is used for elements **910**, **915**, and **920** as described next with respect to FIG. **10**.

[0080] FIG. **10** shows a block diagram for generating the new line index **841** and the intraline index **842** used with the PHT **815** (of FIG. **8**) in accordance with aspects of the invention. In embodiments, when generating the new line index **841**, the index generator **905** uses information from the post-prediction state **1001**, prediction pipeline activity **1002**, restarts **1003**, and LIB content **1004** to determine the content of element **910** (line entry instruction address), element **915** (GPV upon line entry), and element **920** (intraline count of taken branch predictions). In embodiments, when generating the intraline line index **842**, the index generator **905** uses information from the post-prediction state **1001** and prediction pipeline activity **1002** to determine the content of element **910** (line entry instruction address), element **915** (GPV upon line entry), and element **920** (intraline count of taken branch predictions).

[0081] In embodiments, the post-prediction state **1001** is information about the state of the prediction pipeline that is updated upon making predictions at the end of the prediction pipeline. This information may include a new line entry address, a new GPV, and a new count, for example.

[0082] In embodiments, the prediction pipeline activity **1002** is information about the activity of predictions that are currently being made in the prediction pipeline. It may include, for example, information about b2, b3, and b4 cycles, as described herein.

[0083] In embodiments, restarts **1003** are events that cause speculation to be rolled back and processing to resume at a particular instruction address with history state based on instructions

older than the restart point. As a result, restarts can cause a change to the line entry address, GPV, and count that are indicated by the post-prediction state **1001**.

[0084] In embodiments, the LIB content **1004** is information about the line instruction addresses and global path history to get there. In embodiments, the LIB content **1004** includes the line instruction addresses. It also contains indications of how it was installed: sequential lines or target lines from the line index accelerator or prediction pipeline. When from target lines, the line index accelerator also includes information about the taken branch that it predicts jumps into that line, e.g., the GPV value that taken branch shifts into the history vector.

[0085] Referring back to FIG. **8**, because the pattern-based predictors are indexed in advance of being used, the status of prior prediction(s) still in the prediction pipeline may be unknown at the time a new index is needed. Embodiments solve this problem by providing pattern predictors (e.g., PHT **815**) that include multiple read ports. In embodiments, at least one read port is configured to assume intraline indexing by assuming predictions will be made with intraline targets from the current/oldest LOB entry, and at least one other read port is configured to assume new line indexing assuming advancement to the next/younger LOB entry. Depending on the implementation, different numbers of read ports may be used. The example shown in FIG. **8** and described below uses two read ports: one read port for a new line index **841** and one read port for an intraline index **842**.

[0086] In embodiments, the pattern-based predictors are indexed in two circumstances. In a first circumstance, indexing is triggered by index pipeline restarts (either external restarts from the processor pipeline or restarts from the prediction pipeline that redirects the branch prediction search differently from what line index accelerator or sequential LIB population assumed). In a second circumstance, indexing is triggered by a valid cycle in the prediction pipeline (for example, every b2 cycle as described herein).

[0087] With continued reference to FIG. **8**, in embodiments the pattern predictor outputs (e.g., PHT outputs **851**, **852**) may not be needed by the prediction pipeline immediately. For example, the pipeline may need to stall if downstream consumers of branch prediction information are not ready to receive any. Therefore, in embodiments, a capability to capture the pattern predictor output and save it for later is provided. In embodiments, the PHT output buffer **835** serves this purpose.

[0088] In embodiments, when feeding pattern predictor data into the prediction logic **840** of the prediction pipeline, it can come from the array (e.g., PHT **815**) or the saved state (e.g., PHT output buffer **835**). In embodiments, the array outputs from the various ports are multiplexed at **860** to select the appropriate port's output based on the status of older branch predictions in the pipeline, e.g., whether they had an intraline prediction or not.

[0089] FIG. **11** shows a table **1100** that describes an exemplary use case of cycles of a prediction pipeline in accordance with aspects of the invention. Cycles 0-6 are depicted in columns. A current prediction pipeline pass is depicted in row **1101**. Row **1102** depicts actions triggered by activities in row **1101**. Row **1103** depicts a first alternative case in which the next prediction pipeline pass is immediately ready for the PHT output. Row **1104** depicts a second alternative case in which the next prediction pipeline pass is not immediately ready for the PHT output.

[0090] In this example, each prediction pipeline pass starts in a b2 cycle, and the prediction for this pass status will be known in the b4 cycle. In the b4 cycle, the concurrent b2 cycle (the next possible b2 cycle for this thread) is either from the same line (oldest LOB entry) or the next line (second oldest LOB entry). Which one to use in the b2 cycle of the next pass is based on the newline vs intraline status of the b4 cycle of the current pass of the prediction pipeline.

[0091] In this example, a valid b2 cycle in the prediction pipeline is the trigger for indexing the pattern predictor (e.g., PHT **815** of FIG. **8**).

[0092] In this example, cycle 0 is the b2 cycle of the current prediction pipeline pass (row **1101**), and this triggers the next PHT index, called a "b2-2" cycle (row **1102**).

[0093] In this example, cycle 1 is the b3 cycle of the current prediction pipeline pass (row **1101**),

and it is the b2-1 cycle of the PHT array access (row **1102**).

[0094] In this example, cycle 2 is the b4 cycle of the current prediction pipeline pass (row **1101**). In this cycle it is known whether or not there is an intraline taken prediction. That indication allows selection of the PHT array output data from either the new line or intraline read port. In embodiments, the selected data is written into the PHT output buffer **835** of FIG. **8**.

[0095] In this example, in case 1 (row **1103**), if the prediction pipeline is ready for a valid b2 cycle in cycle 2, then the PHT output data is chosen from the array output after the multiplexing (e.g., at **860** in FIG. **8**). In this case, the PHT output data (e.g., one of PHT output **851** and PHT output **852** of FIG. **8**) is selected by the multiplexer **860** and then provided directly to the prediction logic **840** of FIG. **8**.

[0096] In this example, if the prediction logic **840** of the prediction pipeline is not ready for the PHT output data at cycle 2 (e.g., is delayed), then there is no valid b2 cycle in cycle 2 and the next valid b2 cycle happens at some time in the future. This is depicted in case 2 (row **1104**), which shows the next valid b2 cycle occurring in cycle 4 in which case the PHT output data is taken from the PHT output buffer **835**.

[0097] According to aspects of the invention, the pattern-based predictors can be written (e.g., to the PHT **815**) non-speculatively at completion time with an index generated based on information of new line and intraline-based state. In embodiments, to facilitate this, during the completion stage the GPV and address are snapshot at line entry and the current state is maintained.

[0098] In embodiments, during completion stages, there can be multiple groups completing at a same cycle. Each group can have limitations on the number of instructions, branches, and types of branches present. Different combinations of new line vs intraline are possible depending on the “N” number of branches. In embodiments, to enable the intraline count in index generation, the GPV and entry instruction address at line entry are taken as a snapshot. In embodiments, current values are also updated upon all completing branches and maintained. In embodiments, the intraline count is incremented as long as the subsequent taken branch remains intraline. In embodiments, the intraline count, line entry GPV, and line entry address are used to generate the index. Furthermore, upon restarts the state can be updated as if entering into a new line to be consistent with the prediction logic restarting in the same way.

[0099] According to aspects of the invention, the GPV for the line entry is snapshot in following cases: there is a restart and a last GPV calculated for a taken branch is used to snapshot; there is a restart and group crosses a line sequentially; and there is branch to a new line.

[0100] In embodiments, knowledge of first instruction after restarts can be tracked through the processor pipeline in whatever manner instructions are being tracked. One example is to utilize a global completion table.

[0101] In embodiments, for all groups completing in a single cycle, there can be different combinations of taken branches in a group, the group crossing a line, and a branch position within the line. The following are examples in a system with two groups completing and at most one taken branch in a group. In one example, the first group crosses a line sequentially and the branch within the group goes to a new line. In another example, the first group has a branch and causes line crossing, and the subsequent group completing in same cycle crosses the line sequentially. In another example, the first group crosses a line sequentially and the branch within the current group crosses the line again and branch in subsequent group remain within the line. In another example, the first group crosses a line and the branch within the current group goes back to previous line. In another example, both groups completing from same thread with different combination of branch line crossing and groups crossing sequentially. Correct detection of the line entry and snapshot of the corresponding GPV, in a manner described herein, enables the efficiency of PHT index generation during completion cycle to feed the PHT.

[0102] In embodiments, depending on specifics of the number of different pattern-based predictors, and the index functions of them, certain GPV bits can be combined with certain intraline count bits

before transmitting to the index generation logic and array structures themselves. This can minimize the number of writes required. For example, bit **0** of the GPV can be XOR'ed with intraline count bit **0**, and subsequently transmitted to write queues to be used for subsequent index generation for multiple structures, for example PHT and CTB.

[0103] As described herein, a prediction pipeline does not know whether a taken branch has an intraline target or new line target early enough to affect the next pattern-predictor index.

Embodiments of the invention solve this problem by providing a pattern-predictor (e.g., PHT **815**) that has two read ports in which one port is configured based on an assumption there will be a new line target, and one port is configured based on an assumption there will be an intraline target. In embodiments, generating a new intraline index requires only incrementing the counter. This is because the line entry address and line entry GPV are stable and don't depend on the result of the prediction pipeline which is not available yet at the time the index is needed. In this manner, the intraline index is not dependent on information from intraline predictions, and instead is only based on the number of intraline predictions. Using this technique, embodiments are able to speculatively index by generating two indexes (i.e., the new line index and the intraline index) and corresponding PHT outputs before it is known which output will be used, and then selecting the appropriate output to use when the new line versus intraline determination is made in a later cycle. In this manner, implementations of the invention provide a method of using an index function incorporating intraline branch count and line-entry pattern history, triggered by the prediction pipeline, and multiple read ports to index speculatively assuming entering new line and assuming staying in same line before the actual predictions are made. In embodiments, the method includes using prediction pipeline results of new line vs intraline status to select the appropriate pattern predictor output port. In embodiments, the method includes retaining pattern-predictor output and using it later if the prediction pipeline is delayed.

[0104] FIG. **12** shows a flowchart of an exemplary method of branch prediction in a processor in accordance with aspects of the present invention. Steps of the method may be carried out using the logic of FIG. **5** and are described with reference to elements depicted in FIG. **5**. In embodiments, the method comprises: obtaining a branch instruction for which a direction of a branch is to be predicted (step **1201**); generating an index based on an instruction address, a global path vector (GPV), and a counter (step **1202**); selecting an entry from a data structure using the index (step **1203**); and predicting the direction of the branch using information included in the selected entry (step **1204**). The instruction address may be address **511** of FIG. **5** and may comprise an address of a region of code that includes the branch being predicted. The GPV may be GPV **512** and may comprise information defining a history of branches taken to an anchor point in the region of code. The counter may be counter **513** and may comprise a count of predicted taken branches within the region of code after the anchor point. The method may include modifying the GPV based on a value of the counter, and the modifying the GPV comprises reducing a size of the GPV in response to a value of the counter increasing (e.g., as described at FIG. **7**). The data structure may comprise a pattern history table such as PHT **505** of FIG. **5**. The information included in the selected entry of the PHT is, for instance, a saturation counter. The method may include modifying a tag (e.g., TAG of FIG. **5**) in the selected entry in the PHT based on a tag GPV (e.g., TagGPV **515** of FIG. **5**). The predicting the direction at step **1204** may be performed using prediction logic **509** of FIG. **5**. The method may include using the predicted direction of the branch for executing one or more instructions by the processor.

[0105] FIG. **13** shows a flowchart of an exemplary method of branch prediction in a processor in accordance with aspects of the present invention. Steps of the method may be carried out using the logic of FIG. **8** and are described with reference to elements depicted in FIG. **8**. In embodiments, the method comprises: generating a new line index and an intraline index (step **1301**); generating a first output of a pattern-based predictor structure using the new line index (step **1302**); generating a second output of the pattern-based predictor structure using the intraline index (step **1303**);

selecting one of the first output and the second output based on a result of a prediction pipeline (step **1304**); and predicting a direction of a branch using the selected one of the first output and the second output (step **1305**). In embodiments, the generating the new line index and the intraline index is triggered at a first cycle of a pass of the prediction pipeline, and the result of the prediction pipeline is a result of a second cycle of the pass of the prediction pipeline after the first cycle of a pass of the prediction pipeline. In embodiments, the predicting the direction of the branch using the selected one of the first output and the second output is performed in another pass of the prediction pipeline. In embodiments, the pattern-based predictor structure has a first read port associated with the new line index and a second read port associated with the intraline index. The method may include storing the selected one of the first output and the second output in a buffer based on the prediction pipeline being delayed. In embodiments, the new line index and the intraline index are each generated using a line entry instruction address, a global path vector (GPV) leading up to the line entry instruction address, and an intraline counter of predicted taken branches. The method may include using the predicted direction of the branch for executing one or more instructions by the processor.

[0106] The present invention may be a system, a method, and/or a computer program product at any possible technical detail level of integration. The computer program product may include a computer readable storage medium (or media) having computer readable program instructions thereon for causing a processor to carry out aspects of the present invention.

[0107] The computer readable storage medium can be a tangible device that can retain and store instructions for use by an instruction execution device. The computer readable storage medium may be, for example, but is not limited to, an electronic storage device, a magnetic storage device, an optical storage device, an electromagnetic storage device, a semiconductor storage device, or any suitable combination of the foregoing. A non-exhaustive list of more specific examples of the computer readable storage medium includes the following: a portable computer diskette, a hard disk, a random access memory (RAM), a read-only memory (ROM), an erasable programmable read-only memory (EPROM or Flash memory), a static random access memory (SRAM), a portable compact disc read-only memory (CD-ROM), a digital versatile disk (DVD), a memory stick, a floppy disk, a mechanically encoded device such as punch-cards or raised structures in a groove having instructions recorded thereon, and any suitable combination of the foregoing. A computer readable storage medium or media, as used herein, is not to be construed as being transitory signals per se, such as radio waves or other freely propagating electromagnetic waves, electromagnetic waves propagating through a waveguide or other transmission media (e.g., light pulses passing through a fiber-optic cable), or electrical signals transmitted through a wire.

[0108] Computer readable program instructions described herein can be downloaded to respective computing/processing devices from a computer readable storage medium or to an external computer or external storage device via a network, for example, the Internet, a local area network, a wide area network and/or a wireless network. The network may comprise copper transmission cables, optical transmission fibers, wireless transmission, routers, firewalls, switches, gateway computers and/or edge servers. A network adapter card or network interface in each computing/processing device receives computer readable program instructions from the network and forwards the computer readable program instructions for storage in a computer readable storage medium within the respective computing/processing device.

[0109] Computer readable program instructions for carrying out operations of the present invention may be assembler instructions, instruction-set-architecture (ISA) instructions, machine instructions, machine dependent instructions, microcode, firmware instructions, state-setting data, configuration data for integrated circuitry, or either source code or object code written in any combination of one or more programming languages, including an object oriented programming language such as Smalltalk, C++, or the like, and procedural programming languages, such as the “C” programming language or similar programming languages. The computer readable program instructions may

execute entirely on the user's computer, partly on the user's computer, as a stand-alone software package, partly on the user's computer and partly on a remote computer or entirely on the remote computer or server. In the latter scenario, the remote computer may be connected to the user's computer through any type of network, including a local area network (LAN) or a wide area network (WAN), or the connection may be made to an external computer (for example, through the Internet using an Internet Service Provider). In some embodiments, electronic circuitry including, for example, programmable logic circuitry, field-programmable gate arrays (FPGA), or programmable logic arrays (PLA) may execute the computer readable program instructions by utilizing state information of the computer readable program instructions to personalize the electronic circuitry, in order to perform aspects of the present invention.

[0110] Aspects of the present invention are described herein with reference to flowchart illustrations and/or block diagrams of methods, apparatus (systems), and computer program products according to embodiments of the invention. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, can be implemented by computer readable program instructions.

[0111] These computer readable program instructions may be provided to a processor of a computer, or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor of the computer or other programmable data processing apparatus, create means for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks. These computer readable program instructions may also be stored in a computer readable storage medium that can direct a computer, a programmable data processing apparatus, and/or other devices to function in a particular manner, such that the computer readable storage medium having instructions stored therein comprises an article of manufacture including instructions which implement aspects of the function/act specified in the flowchart and/or block diagram block or blocks.

[0112] The computer readable program instructions may also be loaded onto a computer, other programmable data processing apparatus, or other device to cause a series of operational steps to be performed on the computer, other programmable apparatus or other device to produce a computer implemented process, such that the instructions which execute on the computer, other programmable apparatus, or other device implement the functions/acts specified in the flowchart and/or block diagram block or blocks.

[0113] The flowchart and block diagrams in the Figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the present invention. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of instructions, which comprises one or more executable instructions for implementing the specified logical function(s). In some alternative implementations, the functions noted in the blocks may occur out of the order noted in the Figures. For example, two blocks shown in succession may, in fact, be accomplished as one step, executed concurrently, substantially concurrently, in a partially or wholly temporally overlapping manner, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts or carry out combinations of special purpose hardware and computer instructions.

[0114] The descriptions of the various embodiments of the present invention have been presented for purposes of illustration, but are not intended to be exhaustive or limited to the embodiments disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the described embodiments. The terminology used

herein was chosen to best explain the principles of the embodiments, the practical application or technical improvement over technologies found in the marketplace, or to enable others of ordinary skill in the art to understand the embodiments disclosed herein.

Claims

1. A method of branch prediction in a processor, the method comprising: generating a new line index and an intraline index; generating a first output of a pattern-based predictor structure using the new line index; generating a second output of the pattern-based predictor structure using the intraline index; selecting one of the first output and the second output based on a result of a prediction pipeline; and predicting a direction of a branch using the selected one of the first output and the second output.
2. The method of claim 1, wherein: the generating the new line index and the intraline index is triggered at a first cycle of a pass of the prediction pipeline; and the result of the prediction pipeline is a result of a second cycle of the pass of the prediction pipeline after the first cycle of a pass of the prediction pipeline.
3. The method of claim 2, wherein the predicting the direction of the branch using the selected one of the first output and the second output is performed in another pass of the prediction pipeline.
4. The method of claim 1, wherein the pattern-based predictor structure has a first read port associated with the new line index and a second read port associated with the intraline index.
5. The method of claim 1, further comprising storing the selected one of the first output and the second output in a buffer based on the prediction pipeline being delayed.
6. The method of claim 1, wherein the new line index and the intraline index are each generated using a line entry instruction address, a global path vector (GPV) leading up to the line entry instruction address, and an intraline counter of predicted taken branches.
7. The method of claim 6, wherein the GPV is a vector that defines a history of branches taken leading into an anchor point in a region of code that includes the branch.
8. The method of claim 7, wherein the intraline counter is a count of predicted taken branches within the region of code after the anchor point.
9. The method of claim 7, further comprising dynamically determining the region of code by selecting the anchor point based on an event.
10. The method of claim 9, wherein the event is one of a branch flush restart and saturation of the intraline counter.
11. The method of claim 7, wherein respective ones of plural bits in the GPV have a value equal to a predefined bit of a respective target instruction address of a respective taken branch leading into the region of code.
12. The method of claim 6, further comprising reducing a size of the GPV in response to a value of the intraline counter increasing.
13. A computer system for branch prediction in a processing pipeline, the system comprising: a memory; and a processor operatively coupled to the memory, the computer system configured to: generate a first output of a pattern-based predictor structure and a second output of the pattern-based predictor structure by generating a new line index of the pattern-based predictor structure and an intraline index of the pattern-based predictor structure using a line entry instruction address, a global path vector (GPV) leading up to the line entry instruction address, and an intraline counter of predicted taken branches, wherein the pattern-based predictor structure has a first read port associated with the new line index and a second read port associated with the intraline index.
14. The computer system of claim 13, wherein the computer system is further configured to: speculatively generate the first output and the second output before it is known whether the first output or the second output will be used in a prediction pipeline; select one of the first output and the second output based on a result of the prediction pipeline; and predict a direction of a branch

using the selected one of the first output and the second output.

15. The computer system of claim 14, wherein the GPV is a vector that defines a history of branches taken leading into an anchor point in a region of code that includes the branch.

16. The computer system of claim 15, wherein the intraline counter is a count of predicted taken branches within the region of code after the anchor point.

17. The computer system of claim 15, wherein the computer system is further configured to dynamically determine the region of code by selecting the anchor point based on an event.

18. The computer system of claim 17, wherein the event is one of a branch flush restart and saturation of the intraline counter.

19. The computer system of claim 15, wherein respective ones of plural bits in the GPV have a value equal to a predefined bit of a respective target instruction address of a respective taken branch leading into the region of code.

20. The computer system of claim 13, wherein the computer system is further configured to reduce a size of the GPV in response to a value of the intraline counter increasing.
