



US 20250278394A1

(19) **United States**

(12) **Patent Application Publication**

Servia Rodriguez et al.

(10) **Pub. No.: US 2025/0278394 A1**

(43) **Pub. Date: Sep. 4, 2025**

(54) **PRIVATE DATA SET INTERSECTION WITH
MUTUAL DEVICE ANONYMITY**

(52) **U.S. Cl.**
CPC **G06F 16/23** (2019.01); **G06F 21/6245**
(2013.01)

(71) Applicant: **CrowdStrike, Inc.**, Sunnyvale, CA
(US)

(72) Inventors: **Sandra Servia Rodriguez**, Cambridge
(GB); **Calin-Bogdan Miron**,
Nottingham (GB); **Marian Radu**,
Bucharest (RO)

(57) **ABSTRACT**

A method for detecting a private set intersection includes receiving, at a third computing device, a first plurality of transformed data elements from a first computing device; receiving, at the third computing device, a second plurality of transformed data elements from a second computing device, wherein an identity of the first computing device is unknown to the second computing device and an identity of the second computing device is unknown to the first computing device; and transmitting, by a processing device executing on the third computing device to the first computing device and the second computing device, an indication of a subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements.

(21) Appl. No.: **19/213,640**


(22) Filed: **May 20, 2025**

Related U.S. Application Data

(63) Continuation of application No. 18/358,728, filed on
Jul. 25, 2023, now Pat. No. 12,326,849.

Publication Classification

(51) **Int. Cl.**
G06F 16/23 (2019.01)
G06F 21/62 (2013.01)






FIG. 1




FIG. 2A




FIG. 2B




FIG. 2C




FIG. 3




FIG. 4




FIG. 5




FIG. 6





FIG. 7



PRIVATE DATA SET INTERSECTION WITH MUTUAL DEVICE ANONYMITY

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] This application is a continuation of U.S. patent application Ser. No. 18/358,728, filed on Jul. 25, 2023, the entire contents of which are incorporated herein by reference.

TECHNICAL FIELD

[0002] Aspects of the present disclosure relate to determining set intersection, and more particularly, to determining set intersection in a secure environment while maintaining mutual device anonymity.

BACKGROUND

[0003] Private set intersection (PSI) allows two parties, each having a set of data, to find elements that the sets have in common (e.g., an intersection) without revealing the data elements of their sets to each other. PSI may be used in a number of technological environments. For example, PSI may be utilized between two servers to determine if they share a common user and/or a common password. In such a scenario, it may be important to be able to determine the common intersection (e.g., a same user is using the same password on two different systems) without divulging the contents that are being compared (e.g., the usernames and/or passwords).

[0004] PSI can be useful in scenarios in which the data is security- and/or privacy-related. For example, PSI could be utilized to determine if a particular patient's DNA includes genes in common with certain diseases without revealing the patient's DNA. Similarly, technological diagnostics programs may wish to identify whether certain characteristics of a user's computer system are indicative of a larger problem without revealing the full contents of the user's particular configuration. As the amount of data that is maintained on systems, and the interaction of those systems, increases, the need for an effective PSI protocol also increases.

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] The described embodiments and the advantages thereof may best be understood by reference to the following description taken in conjunction with the accompanying drawings. These drawings in no way limit any changes in form and detail that may be made to the described embodiments by one skilled in the art without departing from the scope of the described embodiments.

[0006] FIG. 1 is a block diagram that illustrates an example system, according to some embodiments of the present disclosure.

[0007] FIG. 2A is a flow diagram that illustrates operations of the transformation engine and/or the evaluation engine, in accordance with some embodiments of the present disclosure.

[0008] FIG. 2B is a modified flow diagram that illustrates the use of blinding for the transformation engine, in accordance with some embodiments of the present disclosure.

[0009] FIG. 2C is a modified flow diagram that illustrates the use of the OPRF protocol for the transformation engine.

[0010] FIG. 3 is a block diagram of a system in which the transformation engine and the evaluation engine may be provided on separate computing devices, in accordance with some embodiments of the present disclosure.

[0011] FIG. 4 is a block diagram of an example of an evaluation engine utilizing a hash structure for each client computing device, in accordance with some embodiments of the present disclosure.

[0012] FIG. 5 is a block diagram of an example of an evaluation engine utilizing a common hash structure, in accordance with some embodiments of the present disclosure.

[0013] FIG. 6 is a flow diagram of a method for performing a delegated PSI, in accordance with some embodiments of the present disclosure.

[0014] FIG. 7 is a component diagram of an example of a device architecture for private set intersection (PSI), in accordance with embodiments of the disclosure.

[0015] FIG. 8 is a block diagram of an example computing device that may perform one or more of the operations described herein, in accordance with embodiments of the disclosure.

DETAILED DESCRIPTION

[0016] As described above, PSI refers to the problem in which two parties, each holding a set, wish to find the intersection of their sets without revealing any additional information about the sets themselves. Several PSI protocols have been proposed to date. They allow one party to learn the result of the intersection whereas the other party learns nothing. Solutions for delegated PSI, which involves the use of a third party to aid in the computation of the intersection, have also been proposed, and in some cases, they allow parties to learn the intersection amongst all parties' sets, a problem known as multiparty PSI.

[0017] However, little attention has been paid to the problem in which each party needs to learn whether items in their set are also present in any other party's set without learning the identity of the parties with whom they have a match. Note that this problem is more complex than the traditional two-party PSI problem whose goal is for one party to learn the intersection of its set with the other party's set, as well as from the multiparty PSI problem whose goal is for either a designated party or all parties to learn which elements appear simultaneously in all parties' sets. Executing a two-party PSI protocol between each two parties may not be a feasible solution in some scenarios for a number of reasons. For example, a two-party PSI protocol violates the requirement that the parties should not learn with which parties they have matches because, if there is a match, it must be the other party. In addition, a two-party PSI protocol may be computationally inefficient when a large number of parties are involved, as each party would execute a PSI protocol with each one of the other parties. Similarly, some multiparty PSI protocols may not be suitable because they may compute the intersection with all other parties and they also may not hide the matching party to the involved parties.

[0018] The present disclosure addresses the above-noted and other deficiencies by providing techniques that use two non-colluding entities. For example, one entity (e.g. a transformation engine) may provide transformation services (e.g., parties can query this service to obtain a pseudorandom function (PRF) on every element of their set without having access to a key used by the transformation engine, so

that the output can be securely shared with untrusted third-parties), while another entity (e.g., an evaluation engine) evaluates the transformed values. In some embodiments, the transformation engine may act as the sender in an oblivious pseudorandom function (OPRF) protocol that enables each party to compute a PRF on every element of its set without having access to a key used by the transformation engine, and more importantly, without disclosing their private set to the transformation engine.

[0019] In some embodiments, a multi-point OPRF protocol may be utilized that enforces the use of the same key in consecutive executions of the protocol so that the same transformation is applied to each party's set, allowing for direct computation of the intersection in the PRF output domain. In order to reduce the computation costs in the parties, as well as for impeding and/or preventing them from learning with which other party or parties they have a match, embodiments of the present disclosure may outsource the detection of the intersection to the evaluation engine. For example, after the PRF transformation, each party may send its PRF output set to the evaluation engine which may check for matches between elements in the provided sets and notify respective parties about the elements matching in their sets. In some embodiments, the evaluation engine may store each set into a hash table (e.g., a Cuckoo hash) for fast lookups. In some embodiments, once the sets of every party have been checked, the transformation engine may delete its key and the evaluation engine may delete its hash tables in order to enhance security.

[0020] The embodiments described herein provide improvements both to the operations of the respective computing devices implementing the protocol, and to the technology associated with computer security. By offloading the calculation of the set intersection to a centralized server function, the determination of the intersection of the sets of large numbers of computing devices (e.g., thousands and/or millions) may be more efficiently processed. For example, the calculations associated with a pairwise comparison of sets between a large number of computing devices may vastly exceed the computation required for one centralized server to check the intersections between all of the computing devices. Moreover, the use of the centralized server may allow for the source of the set intersection to be masked. For example, a particular computing device may learn that its set has a member that is duplicated elsewhere, but not the identity of where it is duplicated. This quality may protect from bad actors within the computing devices attempting to learn which systems have a similar member in their sets.

[0021] Embodiments of the present disclosure may improve the technology associated with computer security by allowing for the private and secure comparison of set members, while hiding the identity of the owners of the sets and the set contents themselves. As a non-limiting example, embodiments of present disclosure may provide the ability to determine if a particular password of a computing device is shared. By determining if a password in its set of passwords exists on another device, a computing device may be able to alert the user associated with the password to change it. As noted above, this may be accomplished without learning which system has the same password, reducing an opportunity to exploit this knowledge.

[0022] FIG. 1 is a block diagram that illustrates an example system 100, according to some embodiments of the present disclosure. FIG. 1 and the other figures may use like

reference numerals to identify like elements. A letter after a reference numeral, such as "120A," indicates that the text refers specifically to the element having that particular reference numeral. A reference numeral in the text without a following letter, such as "120," refers to all of the elements in the figures bearing that reference numeral.

[0023] As illustrated in FIG. 1, the system 100 includes a first computing device 110 (also referred to herein as an intersection computing device 110) and a plurality of second computing devices 120 (also referred to herein as client computing devices 120). In FIG. 1, a first client computing device 120A, a second client computing device 120B, up to an Nth client computing device 120N are illustrated merely as examples. It will be understood that more or fewer client computing devices 120 may be present in some embodiments of the present disclosure. The intersection computing device 110 and the client computing devices 120 may each include hardware such as processing device 122 (e.g., processors, central processing units (CPUs)), memory 124 (e.g., random access memory (RAM), storage devices 126 (e.g., hard-disk drive (HDD)), and solid-state drives (SSD), etc.), and other hardware devices (e.g., sound card, video card, etc.).

[0024] In some embodiments, memory 124 may be volatile memory that loses contents when the power to the computing device is removed or non-volatile memory that retains its contents when power is removed. In some embodiments, memory 124 may be non-uniform access (NUMA), such that memory access time depends on the memory location relative to processing device 122.

[0025] Processing device 122 may include a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. Processing device 122 may also include one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. It should be noted that although, for simplicity, a single processing device 122 is depicted in the intersection computing device 110 and the client computing devices 120 depicted in FIG. 1, other embodiments of the intersection computing device 110 and/or the client computing devices 120 may include multiple processing devices 122, storage devices 126, or other devices.

[0026] The storage device 126 may comprise a persistent storage that is capable of storing data. A persistent storage may be a local storage unit or a remote storage unit. Persistent storage may be a magnetic storage unit, optical storage unit, solid state storage unit, electronic storage units (main memory), or similar storage unit. Persistent storage may also be a monolithic/single device or a distributed set of devices.

[0027] The intersection computing device 110 and/or the client computing devices 120 may comprise any suitable type of computing device or machine that has a programmable processor including, for example, server computers, desktop computers, laptop computers, tablet computers, smartphones, set-top boxes, etc. In some examples, the intersection computing device 110 and/or the client computing devices 120 may comprise a single machine or may include multiple interconnected machines (e.g., multiple

servers configured in a cluster). The intersection computing device **110** and/or the client computing devices **120** may be implemented by a common entity/organization or may be implemented by different entities/organizations.

[0028] The intersection computing device **110** and the client computing devices **120** may be coupled to each other (e.g., may be operatively coupled, communicatively coupled, may communicate data/messages with each other) via network **102**. Network **102** may be a public network (e.g., the internet), a private network (e.g., a local area network (LAN) or wide area network (WAN)), or a combination thereof. In one embodiment, network **102** may include a wired or a wireless infrastructure, which may be provided by one or more wireless communications systems, such as a WiFi™ hotspot connected with the network **102** and/or a wireless carrier system that can be implemented using various data processing equipment, communication towers (e.g., cell towers), etc. The network **102** may carry communications (e.g., data, message, packets, frames, etc.) between the intersection computing device **110** and/or the client computing devices **120**.

[0029] The intersection computing device **110** may be configured to provide a transformation engine **130** and an evaluation engine **140**. The configuration of the transformation engine **130** and/or the evaluation engine **140** illustrated in FIG. 1 is merely for the purposes of illustration, and is not intended to limit the embodiments of the present disclosure.

[0030] The transformation engine **130** may be configured to perform a transformation operation on sets **150** that are transmitted (e.g., over the network **102**) to the intersection computing device **110** to generate a transformed set **152**. In some embodiments, the first client computing device **120A** may provide values corresponding to a first set **150A** to the transformation engine **130**, and the transformation engine **130** may return a first transformed set **152A**. A second client computing device **120B** may provide values corresponding to a second set **150B** to the transformation engine **130**, and the transformation engine **130** may return a second transformed set **152B**. An N^{th} client computing device **120N** may provide values corresponding to an n^{th} set **150N** to the transformation engine **130**, and the transformation engine **130** may return an n^{th} transformed set **152N**, and so on.

[0031] In some embodiments, a first client computing device **120A** may execute a protocol (e.g., an OPRF protocol) with the transformation engine **130** utilizing values corresponding to a first set **150A** to generate a first transformed set **152A**. In some embodiments, the first transformed set **152A** may be generated as a result of an OPRF protocol executed between the first client computing device **120A** and the transformation engine **130**. In some embodiments, the second transformed set **152B** may be generated as a result of an OPRF protocol executed between the second client computing device **120B** and the transformation engine **130**. In some embodiments, the N^{th} transformed set **152N** may be generated as a result of an OPRF protocol executed between the N^{th} client computing device **120N** and the transformation engine **130**.

[0032] The set **150** may include one or more data elements. For example, the set **150** may include usernames, passwords, or other data values for which an intersection is to be determined with respect to the other client computing devices **120**. The contents of the data elements of the set **150** are not limited by the examples of FIG. 1, but may instead

include any data value that may be manipulated mathematically utilizing a pseudorandom function, as described further herein.

[0033] In some embodiments, the transformation engine **130** may utilize a pseudorandom function (PRF) to generate the transformed set **152** of data values from the initial set **150** of data values. A PRF is a deterministic function of a key and an input that is indistinguishable from a truly random function of the input. For example, if s is a security parameter, k is a key of length s bits, and $F(k, x)$ (or $F_k(x)$) is a function of key k and input x , then F is a pseudorandom function if F can be computed in polynomial time in s and if k is random, then F cannot be distinguished from a random function in polynomial time.

[0034] In some embodiments, the transformation engine **130** and the client **120A** may execute an oblivious pseudorandom function (OPRF) protocol to generate the transformed set **152** of data values from the initial set of data values **150**. An OPRF according to some embodiments of the present disclosure is a cryptographic primitive that allows a sender to choose or learn a PRF key k and a receiver to learn the PRF output $\text{OPRF}_k(x_1), \dots, \text{OPRF}_k(x_n)$ on its inputs $x_1, \dots, x_n \in X$. Nothing about the receiver's inputs is revealed to the sender and nothing about the key k is revealed to the receiver. In many OPRF applications, such as PSI, one may obviously evaluate a PRF more than once. If the OPRF allows for consecutive evaluations with respect to the same key, it may be referred to as a multi-point OPRF.

[0035] FIG. 2A is a flow diagram that illustrates operations of the transformation engine **130** and/or the evaluation engine **140**, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 2A that have been previously described will be omitted for brevity. The operations illustrated in FIG. 2A may be utilized, for example, by the system **100** of FIG. 1.

[0036] Referring to FIGS. 1 and 2A, at operation **210**, the transformation engine **130** generates a key **132**. In some embodiments, the key **132** may be a digital value, and may be randomly, or pseudorandomly, generated. The key **132** may be known to the transformation engine **130**, but unknown to and/or concealed from the evaluation engine **140** and the client computing devices **120**. In some embodiments, the transformation engine **130** and the evaluation engine **140** may be isolated from one another on the intersection computing device **110** to ensure that the transformation engine **130** and the evaluation engine **140** are unable to collude with one another. For example, containers and/or virtual machines may be utilized to isolate the transformation engine **130** from the evaluation engine **140**. In some embodiments, as will be described further herein, the transformation engine **130** and the evaluation engine **140** may be provided on separate computing devices.

[0037] At some point after the key **132** is generated, the first client computing device **120A** may send its set **150A** of data values to the transformation engine **130**. The transformation engine **130** may compute the pseudorandom function at operation **212** utilizing the key **132** and the set of values **150A**.

[0038] An example of a pseudorandom function (PRF) may be such that for a given input set of $\{a_1, a_2, \dots, a_n\}$ and a key α , it generates a set of transformed values **152** of $\{a_1', a_2', \dots, a_n'\}$ given by:

$$\forall i, 1 \leq i \leq n$$

$$a'_i = (a_i)^\alpha$$

The key α is chosen in such a way that it is computationally difficult to obtain a'_i without knowing α .

[0039] The transformed values **152A** that correspond to the set **150A** from the first client computing device **120A** may be transmitted back (e.g., over network **102**) to the first client computing device **120A**, and may be stored on the first client computing device **120A** (e.g., in storage device **126**).

[0040] The process of generating transformed values **152** may be repeated with the second client computing device **120B**. Namely, a set **150B** of data values may be transmitted to the transformation engine **130**, which may perform an PRF operation **214** utilizing the key **132** to generate transformed values **152B**. The transformed values **152B** may be sent back to the second client computing device **120B**.

[0041] In some embodiments, the use of the key **132** to generate the transformed values **152** from the set **150** of data values using the PRF operation (e.g., PRF operation **214**) may be part of an OPRF protocol **260** between the client computing device **120** and the transformation engine **130**. The OPRF protocol **260** is jointly executed between the client computing device **120** and the transformation engine **130** in order to transform the set **150** of data values in a way that the transformation engine **130** does not get to learn either the data values of the original set **150** or the data values of the transformed set **152** and the client computing device **120** does not get to learn the key **132**.

[0042] The first client computing device **120A** and the second client computing device **120B** may both transmit their transformed values **152A**, **152B** to the evaluation engine **140**. Because the pseudorandom functions performed by the transformation engine **130** were performed on the set values **150A**, **150B** using a same key **132**, the evaluation engine **140** is able to compare the transformed values **152A**, **152B** to determine if there is an intersection. In some embodiments, the evaluation engine **140** calculates **220** an intersection between the transformed values **152** received from the client computing devices **120**. Calculating an intersection between the two sets of transformed values may include, for each data element of the first transformed set **152A**, determining if that data element is also present in the second transformed set **152B**. For example, if the first set of transformed values **152A** from the first client computing device **120A** contains a same value as the second set of transformed values **152B** from the second client computing device **120B**, the evaluation engine **140** may transmit first intersection results **154A** to the first client computing device **120A** that indicates the subset of one or more shared values that are present in a set from another client computing device **120**. For example, the first intersection results **154A** may identify a subset of the first set of transformed values **152A** that are also present in at least one other set of transformed values **152**. In some embodiments, the first intersection results **154A** may not identify the second client computing device **120B** (or any other client computing device **120**) as the source of the subset of intersected values. Similarly, the evaluation engine **140** may transmit second intersection results **154B** to the second client computing device **120B** that indicates a subset of one or more shared values that are present in a set from another client computing device **120**.

In some embodiments, the second intersection results **154B** may not identify the first client computing device **120A** (or any other client computing device **120**) as the source of the subset of intersected values.

[0043] In some embodiments, in response to determining that there is an intersection between the first set **150A** of the first client computing device **120A** and another set **150** from another client computing device **120**, the first client computing device **120A** may perform remediation. For example, if the set **150** of data elements correspond to passwords, the first client computing device may generate an alert to the user associated with the password to warn them that their password is repeated elsewhere in another client computing device **120**. Other types of remediation, such as locking or otherwise restricting the user account, may be possible without deviating from the embodiments of the present disclosure.

[0044] In FIG. 2A, the client computing device **120** provides its set directly to the transformation engine **130**, but embodiments of the present disclosure are not limited to such a configuration. FIG. 2B is a modified flow diagram that illustrates the use of blinding for the transformation engine **130**, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 2B that have been previously described will be omitted for brevity. The operations illustrated in FIG. 2B may be utilized, for example, by the system **100** of FIG. 1.

[0045] The operations of FIG. 2B differ from those of FIG. 2A in that, rather than sending the data values of the set **150** to the transformation engine **130**, the transformation engine **130** and the client computing device **120A** jointly execute a 2-party computation protocol (e.g., OPRF protocol **260**) by which the client computing device **120A** obtains as a result the pseudorandom function computed over its set (transformed set), and the transformation engine **130** does not learn either the original client set **150** or its transformed version **152**. In some embodiments, the OPRF protocol **260** may utilize the Diffie-Hellman OPRF protocol to blind and/or transform the client set **150A**.

[0046] In the Diffie-Hellman Oblivious PRF, the first client computing device **120A** has an input set **150A** of $\{x_1, x_2, \dots, x_n\}$ and generates a set of random numbers β such that a hash function H_1 is defined as:

$$\forall i, 1 \leq i \leq n$$

$$\beta_i \xleftarrow{R} \mathbb{Z}_q$$

$$a_i = H_1(x_i)^{\beta_i}$$

[0047] Using the hash function H_1 , a set of intermediate values **250** of $\{a_1, a_2, \dots, a_n\}$ may be generated from the original set **150** of $\{x_1, x_2, \dots, x_n\}$. The intermediate values **250** (e.g., first intermediate values **250A** from the first client computing device **120A**) may be sent to the transformation engine **130** rather than the set **150** itself. The transformation engine **130** may perform an OPRF protocol **260** with the first client computing device **120A** that incorporates a transformation **222** on the first intermediate values **250A** from the first computing device **120A** based on the key **132** to return first transformed intermediate values **252A** of $\{a'_1, a'_2, \dots, a'_n\}$. The first transformed intermediate values **252A** may be based on a PRF operation that is a combination of the

intermediate values **250A** and a second transformation performed by the transformation engine **130**.

[0048] At operation **240A**, the first transformed values **152A** may be generated from the first transformed intermediate values **252A**. That is to say that the blinding done in operation **230A** may be reversed using a second hash function H_2 . In some embodiments, the second hash function H_2 may also be set up as a random oracle. Thus, for a given set of first transformed intermediate values **252A** from the transformation engine **130** of $\{a_1', a_2', \dots, a_n'\}$, the first client computing device **120A** may perform the following operation:

$$\forall i, 1 \leq i \leq n$$

$$a_i'' = H_2((a_i')^{1/\beta_i})$$

to generate the first transformed values **152A**. The first transformed values **152A** may then be sent to the evaluation engine **140** in a similar manner as described above with respect to FIG. 2A.

[0049] The operations of generating the second transformed values **152B** may be repeated by the second client computing device **120B**. Namely, at operation **230B**, the second client computing device **120B** may generate second intermediate values **250B** from its set **150B**, e.g., using a hash function such as the first hash function H_1 described above. The second intermediate values **250B** may then be sent to the transformation engine **130**, which may perform a transformation **224** on the second intermediate values **250B** from the second computing device **120B** based on the key **132** to generate second transformed intermediate values **252B**. The second transformed intermediate values may then be converted to the second transformed values **152B** at operation **240B**, e.g., using a hash function such as the second hash function H_2 described above.

[0050] The use of the blinding described in FIG. 2B may allow for the operations of the transformation engine **130** to be utilized without transferring the actual values of the set **150** to the transformation engine **130**. The use of the two hash functions H_1 and H_2 allows for the resulting transformed values **152** to be generated in a way that still allows the results to be compared by the evaluation engine **140**. Thus, the operations of the evaluation engine **140** illustrated in FIG. 2A may still be performed on the transformed values **152A**, **152B** to determine if an intersection exists between the sets **150A**, **150B** of the first and second client computing devices **120A**, **120B**.

[0051] FIG. 2C is a modified flow diagram that illustrates the use of the OPRF protocol for the transformation engine **130**, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 2C that have been previously described will be omitted for brevity. The operations illustrated in FIG. 2C may be utilized, for example, by the system **100** of FIG. 1.

[0052] The operations of FIG. 2C illustrates that the embodiments of the present disclosure are not limited to an OPRF protocol **260** similar to the Diffie-Hellman OPRF protocol illustrated with respect to FIG. 2B. In some embodiments, any OPRF protocol **260** may be utilized between the client computing devices **120** and the transformation engine **130**. For example, as shown in FIG. 2C in conjunction with FIG. 1, the client computing devices **120**

and the transformation engine **130** may utilize the OPRF protocol **260** to transform a set **150** into a transformed set **152** of values.

[0053] For example, referring to FIG. 2C, the first client computing device **120A** may perform a first OPRF protocol **260A** in conjunction with the transformation engine **130** on its set **150A** and the key **132** of the transformation engine **130**. The first OPRF protocol **260A** may result in the set of transformed values **152A**, as described herein.

[0054] The first client computing device **120A** may provide the transformed values **152A** to the evaluation engine **140** at operation **270A**. The evaluation engine **140** may compare the transformed values from the first client computing device **120A** with other transformed values **152** from other client computing devices **120** to determine if values within the transformed values **152A** received from the first OPRF protocol **260A** are present in the transformed values **152** of other client computing devices **120**.

[0055] Similarly, the second client computing device **120B** may perform a second OPRF protocol **260B** in conjunction with the transformation engine **130** on its set **150B** and the key **132** of the transformation engine **130**. The second OPRF protocol **260A** may result in the second set of transformed values **152B**, as described herein. The second client computing device **120** may provide the received second set of transformed values **152B** to the evaluation engine **140** at operation **270B**, for comparison with the transformed values **152** of other client computing devices **120** (such as the first client computing device **120A**).

[0056] The embodiments of FIGS. 2A, 2B and 2C provide a number of benefits. First, the client computing devices **120** are able to learn about intersections (e.g., repeats) within their sets **150** of data values without having to identify themselves to the other client computing devices **120**. By performing a PSI that utilizes the OPRF protocol **260**, the client computing devices **120** may learn both that there is an intersection with their set **150** and another set **150** of another client computing device **120** as well as which of the elements of the set **150** are within the intersection (e.g., which values are repeated). Using the example of passwords, the first client computing device **120A** may learn that one of its passwords is in use on another client computing device **120**, but may not learn which client computing device **120** it is.

[0057] In addition, since the evaluation engine **140** is only provided transformed values **152**, it may be able to calculate the intersection of the various sets **150** without seeing the original values of the set **150**. In the example of passwords, this may mean that the evaluation engine **140** may be able to calculate the intersection without seeing any of the passwords. As a result, in cases where a blinded set is used, as in FIG. 2B, a client computing device **120** may be able to receive information related to the intersection of its set **150** with other client computing devices **120** without providing the actual values of the set **150** to another entity (e.g., the transformation engine **130**), which may further enhance security.

[0058] Referring back to FIG. 1, the system **100** may provide the ability for the various client computing devices **120** to determine intersections in their various sets **150** utilizing a PRF, which could be obtained and/or computed via an OPRF protocol, as described herein. Though an OPRF protocol utilizing blinded exponentiation is described with respect to FIGS. 2A and 2B, the embodiments of the present disclosure are not limited to such an embodiment. In some

embodiments, other types of OPRFs may be used, such as constructions based on oblivious transfer and homomorphic encryption, to name just two examples.

[0059] As transformed sets **152** (e.g., transformed sets **152A**, **152B**, . . . , **152N**) are generated by the transformation engine **130** and provided to the evaluation engine **140**, they may be stored within a transformation store **142**. Thus, as a new client computing device **120** is added to the network **102**, it may communicate with the transformation engine **130** to generate a transformed set **152** based on its set **150** of data values, and may provide this transformed set **152** to the evaluation engine **140**, which may compare the values of the transformed set **152** to the values within its transformation store **142**.

[0060] In some embodiments, the key **132** and/or the transformation store **142** may be deleted and/or regenerated. Regenerating the key **132** may render prior computations moot, such that any new transformed sets **152** using the new key **132** may not be compared against prior transformed set **152** utilized the prior key **132**. However, generating a new key **132** and/or transformation store **142** may prevent and/or reduce security issues that may arise from having the transformed sets **152** in storage. Also, regenerating the key **132** periodically may reduce a time window over which the key **132** can be deduced/cracked and utilized to attempt to decrypt the original values of the sets **150**. In some embodiments, the key **132** and/or the transformation store **142** may be deleted and/or regenerated periodically (e.g., every day, every week, etc.). In some embodiments, the key **132** and/or the transformation store **142** may be deleted and/or regenerated in response to events, such as with the detection of a new client computing device **120** within the network **102**.

[0061] In the example of FIG. 1, the transformation engine **130** and the evaluation engine **140** are provided on a same intersection computing device **110**. However, the embodiments of the present disclosure are not limited to such a configuration. FIG. 3 is a block diagram of a system **300** in which the transformation engine **130** and the evaluation engine **140** may be provided on separate computing devices, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 3 that have been previously described will be omitted for brevity.

[0062] Referring to FIG. 3, the transformation engine **130** may be provided as part of a transformation computing device **330** and the evaluation engine **140** may be provided as part of an evaluation computing device **340**. In some embodiments, the transformation computing device **330** and the evaluation computing device **340** may be different physical machines, though the embodiments of the present disclosure are not limited to such a configuration. In some embodiments, the transformation computing device **330** and the evaluation computing device **340** may be different virtual machines on a same physical device. In some embodiments, the transformation computing device **330** and the evaluation computing device **340** may be implemented by a common entity/organization or may be implemented by different entities/organizations.

[0063] Each of the transformation computing device **330** and the evaluation computing device **340** may include a memory **124**, a processing device **122**, and/or storage **126**, as described herein with respect to FIG. 1. In some embodiments, the configuration of the memory **124**, the processing device **122**, and/or the storage **126** of the transformation

computing device **330** may be different from that of the evaluation computing device **340**.

[0064] The transfer of the sets **150** (or intermediate sets **250**) to the transformation engine **130** may be performed to the transformation computing device **330** in a similar manner as described herein with respect to FIGS. 2A to 2C. Namely, the sets **150** (or intermediate sets **250**) may be transferred to the transformation engine **130** of the transformation computing device **330** over the network **102**, and an OPRF protocol incorporating a PRF operation may be performed on the transferred values based on the key **132**.

[0065] Similarly, the transfer of the transformed sets **152** to the evaluation engine **140** may be performed to the evaluation computing device **340** in a similar manner as described herein with respect to FIGS. 2A to 2C. Namely, transformed sets **152** may be transferred to the evaluation engine **140** of the evaluation computing device **340** over the network **102**, and an intersection **154** between the transformed sets **152** and the transformation store **142** may be identified and returned to the client computing device **120**.

[0066] The separated configuration of the system **300** may provide additional security benefits. By separating the transformation engine **130** and the evaluation engine **140** onto different computing devices **330**, **340**, the key **132** may be separated (e.g., physically) from the transformation store **142**. In some embodiments, the transformed sets **152** stored in the transformation store **142** may be possibly decrypted if the key **132** is known. Thus, if the transformation engine **130** and the evaluation engine **140** are on a same computing device, it may be possible for the values of the transformed sets **152** to be determined if that computing device is accessed by an unauthorized entity, or is a bad actor. Separating the key **132** from the transformation store **142** may reduce and/or prevent a possibility of the transformed sets **152** of the transformation store **142** from being decrypted using the key **132**. As an example, even if one of the transformation engine **130** or the evaluation engine **140** is compromised, the intruder still would not have enough information to determine information specific to the sets **150** of the client computing devices **120**.

[0067] In some embodiments, the sets **150** (and therefore the transformed sets **152**) may be very large. As a result, it may be useful to provide additional performance enhancements to increase the speed of operations. FIG. 4 is a block diagram of an example of an evaluation engine **140** utilizing a hash structure for each client computing device **120**, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 4 that have been previously described will be omitted for brevity.

[0068] FIG. 4 illustrates a portion of the communication discussed herein with respect to FIGS. 1 to 3. Namely, FIG. 4 illustrates scenarios in which the evaluation engine **140** receives a transformed set **152** of values (e.g., a first transformed set **152A**) from a client computing device **120**. In the embodiment of FIG. 4, the evaluation engine **140** is illustrated as executing within the separate evaluation computing device **340** described with respect to FIG. 3. However, it will be understood that the hashing techniques discussed with respect to FIG. 4 may also be implemented in solutions in which the evaluation engine **140** and the transformation engine **130** share a common computing device (e.g., the intersection computing device **110** of FIG. 1) but are logically isolated from one another.

[0069] Referring to FIG. 4, in some embodiments, the transformed set 152 may be transmitted to the evaluation engine 140 for each of the client computing devices 120, and the evaluation engine 140 may insert the transformed sets 152 into separate hash structures 452 within its transformation store 142. In some embodiments, the client computing devices 120 may store the transformed sets 152 in separate hash structures 452 (e.g., within the storage device 126 of the client computing device 120), and the separate hash structures 452 themselves may be sent directly from the client computing devices 120 to the evaluation engine 140.

[0070] The respective hash structures 452 may allow the transformed sets 152 to be stored more efficiently for searching. In some embodiments, the hash structure 452 may instantiate a hash filter. A hash filter creates a hash sum from data and compares the sum against other previously defined sums. Depending on the purpose of the filter, the data can then be included or excluded in a function based on whether it matches an existing sum. A hash filter may allow for large amounts of data to be stored efficiently and searched quickly. For example, a hash filter may allow for a given value to be quickly tested to determine if it is in the hash filter, without requiring that every member of the data of the hash filter be individually compared.

[0071] For example, referring to FIGS. 1 and 3, a first client computing device 120A may take the first transformed set 152A of data elements based on the pseudorandom function of the transformation engine 130, and then transmit (e.g., over the network 102) the first transformed set 152A to the evaluation engine 140. The evaluation engine 140 may store the first transformed set 152 into a first hash structure 452A that is associated with the first client computing device 120A. The first hash structure 452A may allow for all of the data elements of the first transformed set 152A to be searched quickly and efficiently using a hash filter.

[0072] The evaluation engine 140 may receive a transformed set 152 from each of the client computing devices 120. For example, the evaluation engine 140 may receive the first transformed set 152A of data values from the first client computing device 120A, which may be stored in the first hash structure 452A, up to an Nth transformed set 152N of data values from an Nth client computing device 120N, which may be stored in an Nth hash structure 452N. In some embodiments, rather than receiving the transformed set 152 from each of the client computing devices 120, the evaluation engine 140 may receive the hash structure 452 itself. For example, the first hash structure 452A may be received from the first client computing device 120, up to the Nth hash structure 452N from the Nth client computing device 120N.

[0073] When computing the intersection of a received transformed set 152 of data elements with the transformed sets 152 of the transformation store 142, the evaluation engine 140 may compare each data value of the received transformed set 152 with each of the separate hash structures 452. For example, each data value of the received transformed set 152 may be tested against the first hash structure 452A of the first client computing device 120. The comparison may be repeated for each of the hash structures 452 of the transformation store 142 up to the Nth hash structure 452N. Thus, for N hash structures 452, each value of the incoming transformed set 152 may be compared N times (e.g., once per hash structure 452). Once compared, the received data value of the incoming transformed set 152

may be added to a hash structure 452 in the transformation store 142 that is associated with the client computing device 120 from which the incoming transformed set 152 is received, and utilized in subsequent operations.

[0074] In some embodiments, the hash structure 452 may be a Cuckoo hash filter. Cuckoo filters are data structures for compact set representation that allow for fast membership testing with controllable false possibility probability (FPP). They have fast lookups ($O(1)$), deletions ($O(1)$) and insertions ($O(1)$) amortized expected, with reasonable high probability if load is well managed). Cuckoo filters consist of a table of buckets with fixed bucket size b . Inside the buckets, so-called tags are stored. Tags are small bit strings obtained by hashing items. For example, to represent an item x in a Cuckoo filter, its tag $t_x = H_v(x)$, where H_v is a hash function with output bit length v . This tag may be stored in one out of two possible buckets. The position of the first possible bucket is calculated as $p_1 = H(x)$, where H is another hash function that maps the input to a position in the table of buckets. In case this bucket is already full, the tag is stored in the second possible bucket at position $p_2 = p_1 \oplus H(t_x)$. Note that it is always possible to determine the other candidate bucket p_2 just from knowing its tag t_x and the current position p_1 : $p_2 = p_1 \oplus H(t_x)$. If both buckets are full, one tag in one of the buckets is chosen at random, removed from that bucket, and moved to its other possible bucket. This procedure is repeated recursively until no more relocations are necessary.

[0075] To check whether an item is contained in the Cuckoo filter, one computes its tag and both possible bucket locations and compares the tags stored there for equality. For deleting the item, the matching tag is removed from the filter. Due to hash collisions, two items may produce equal tags. As a consequence, lookups can lead to false positives. The false positive probability e_{max} is mainly dependent on the tag size v and also slightly on the bucket size b since larger buckets result in more possible collisions within each bucket.

[0076] As an example of an embodiment in which the hash structure 452 is a Cuckoo filter, pseudocode for an operation to test an intersection of n data elements of a transformed set 152 against m other hash structures 452 (denoted as cf) from m client computing devices 120 may be of the form:

```

 $\forall j, 1 \leq j \leq m$ 
 $\forall i, 1 \leq i \leq n$ 
  If  $cf_j.Check(\alpha''_i)$  then
    Return  $j, i$ 

```

[0077] When using a Cuckoo filter, operations to add data elements of a transformed set 152 of values into a respective hash associated with an $(m+1)^{st}$ client computing device 120 may be implemented by the following example pseudocode:

```

 $\forall i, 1 \leq i \leq n$ 
   $cf_{m+1}.Insert(\alpha''_i)$  then
  CF.add( $cf_{m+1}$ )

```

[0078] Though some embodiments of the hash structure 452 may be a Cuckoo hash filter, the embodiments of the present disclosure are not limited to this configuration. In some embodiments the hash structure 452 may be another type of hash filter, such as a bloom filter. Other types of hash

structures 452 may be possible without deviating from the scope of the present disclosure.

[0079] In some embodiments, it may be useful to implement a common hash filter, rather than individual hash filters for each of the computing devices 120. FIG. 5 is a block diagram of an example of an evaluation engine 140 utilizing a common hash structure, in accordance with some embodiments of the present disclosure. A description of elements of FIG. 5 that have been previously described will be omitted for brevity.

[0080] FIG. 5 illustrates a portion of the communication discussed herein with respect to FIGS. 1 to 3. Namely, FIG. 5 illustrates scenarios in which the evaluation engine 149 receives a transformed set 152 of values from a client computing device 120. In the embodiment of FIG. 5, the evaluation engine 140 is illustrated as executing within the separate evaluation computing device 340 described with respect to FIG. 3. However, it will be understood that the hashing techniques discussed with respect to FIG. 5 may also be implemented in solutions in which the evaluation engine 140 and the transformation engine 130 share a common computing device (e.g., the intersection computing device 110 of FIG. 1) but are logically isolated from one another.

[0081] Referring to FIG. 5, in some embodiments, the transformed set 152 may be transmitted to the evaluation engine 140 for each of the client computing devices 120, and the evaluation engine 140 may insert the transformed set 152 into a single common hash structure 552 within its transformation store 142. For example, each of the transformed sets 152 from the plurality of client computing devices 120 may each be sent to the evaluation engine 140 and inserted into a common hash structure 552 that may be searched for an intersection. As with the hash structure 452 of FIG. 4, the common hash structure 552 may allow all of the transformed set 152 to be stored more efficiently for searching.

[0082] The evaluation engine 140 may receive a transformed set 152 from each of the client computing devices 120. For example, the evaluation engine 140 may receive the first transformed set 152A of data values from the first client computing device 120A up an Nth transformed set 152N of data values from an Nth client computing device 120A. Each of these transformed sets 152 may be inserted into the common hash structure 552. In some embodiments, the values of the transformed sets 152 may be inserted into the common hash structure along with an identifier of the client computing device 120 from which the values of the transformed sets were received. When computing the intersection of a received transformed set 152 of data elements with the transformed sets 152 of the transformation store 142, the evaluation engine 140 may compare each data value of the received transformed set 152 with the common hash structure 552. Thus, each value of the incoming transformed set 152 must be compared only a single time against the common hash structure 552. If a matching value is found within the common hash structure 552, the identifier(s) associated with the client computing devices 120 of the matching value may be utilized to determine which client computing devices 120 have an intersection with the incoming transformed set 152. Once compared, the received transformed set 152 of data elements may be added to the common hash structure 552 of the transformation store 142 and utilized in subsequent operations.

[0083] FIG. 6 is a flow diagram of a method 600 for performing a delegated PSI, in accordance with some embodiments of the present disclosure. Method 600 may be performed by processing logic that may comprise hardware (e.g., circuitry, dedicated logic, programmable logic, a processor, a processing device, a central processing unit (CPU), a system-on-chip (SoC), etc.), software (e.g., instructions running/executing on a processing device), firmware (e.g., microcode), or a combination thereof. In some embodiments, the method 600 may be performed by a computing device (e.g., intersection computing device 110, transformation computing device 330, and/or evaluation computing device 340).

[0084] With reference to FIG. 6, method 600 illustrates example functions used by various embodiments. Although specific function blocks (“blocks”) are disclosed in method 600, such blocks are examples. That is, embodiments are well suited to performing various other blocks or variations of the blocks recited in method 600. It is appreciated that the blocks in method 600 may be performed in an order different than presented, and that not all of the blocks in method 600 may be performed.

[0085] Referring simultaneously to the prior figures as well, the method 600 begins at block 610, which includes receiving a first plurality of transformed data elements from a first computing device, the first plurality of transformed data elements representing a transform by a pseudorandom function of a first plurality of data elements. In some embodiments, the first plurality of transformed data elements from a first computing device may be similar to the transformed set of values 152 (e.g., the first transformed set of values 152A) from a client computing device 120 (e.g., the first client computing device 120A) described herein with respect to FIGS. 1 to 5. In some embodiments, the first plurality of data elements may be similar to the set of values 150 (e.g., the first set of values 150A) described herein with respect to FIGS. 1 to 5. In some embodiments, the first plurality of transformed data elements and the second plurality of transformed data elements are obtained via an oblivious pseudorandom function (OPRF) protocol. In some embodiments, the first plurality of transformed data elements and the second plurality of transformed data elements correspond to a first and second set of passwords of the first and second computing devices, respectively.

[0086] At block 620, operations of the method 600 may include receiving a second plurality of transformed data elements from a second computing device, the second plurality of transformed data elements representing a transform by the pseudorandom function of a second plurality of data elements. In some embodiments, the second plurality of transformed data elements from a second computing device may be similar to the transformed set of values 152 (e.g., the second transformed set of values 152B) from a client computing device 120 (e.g., the second client computing device 120B) described herein with respect to FIGS. 1 to 5. In some embodiments, the second plurality of data elements may be similar to the set of values 150 (e.g., the second set of values 150B) described herein with respect to FIGS. 1 to 5.

[0087] In some embodiments, the OPRF protocol comprises a Diffie-Hellman OPRF protocol. In some embodiments, the method further includes generating a key for the pseudorandom function of the OPRF protocol, wherein the key is concealed from the first and second computing

devices; receiving a first plurality of intermediate transformed data elements from the first computing device, the first plurality of intermediate transformed data elements representing a transform of the first plurality of data elements by the first computing device; providing the key and the first plurality of intermediate transformed data elements to the pseudorandom function to generate the first plurality of transformed data elements; and transmitting the first plurality of transformed data elements to the first computing device. In some embodiments, the key may be similar to the key **132** described herein with respect to FIGS. **1** to **5**. In some embodiments, the first plurality of intermediate transformed data elements may be similar to the transformed intermediate values **252** (e.g., the first transformed intermediate values **252A**) described herein with respect to FIG. **2B**. In some embodiments, a third computing device that generates the key for the pseudorandom function is different from a fourth computing device that receives the first and second plurality of transformed data elements. In some embodiments, the third computing device may be similar to the transformation computing device **330** described herein with respect to FIG. **3**. In some embodiments, the fourth computing device may be similar to the evaluation computing device **340** described herein with respect to FIG. **3**.

[0088] At block **630**, operations of the method **600** may include transmitting, to the first computing device and the second computing device, an indication of a subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements. In some embodiments, the indication of a subset of transformed data elements may be similar to the intersection **154** (e.g., the first intersection **154A** and/or the second intersection **154B**) described herein with respect to FIGS. **1** to **5**.

[0089] In some embodiments, the method further includes inserting the first plurality of transformed data elements into a first hash structure and the second plurality of transformed data elements into a second hash structure; and determining the subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements by analyzing the first and second hash structures. In some embodiments, the first hash structure and the second hash structure may be similar to the hash structures **452** (e.g., **452A**, . . . , **452N**) described herein with respect to FIG. **4**.

[0090] In some embodiments, the method further includes inserting the first and second plurality of transformed data elements into a common hash structure, and determining the subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements by analyzing the common hash structure. In some embodiments, the common hash structure may be similar to the common hash structure **552** described herein with respect to FIG. **5**.

[0091] FIG. **7** is a component diagram of an example of a device architecture **700** for private set intersection (PSI), in accordance with embodiments of the disclosure. The device architecture **700** includes computing device **710** having processing device **122** and memory **124**, as described herein with respect to FIGS. **1** to **6**.

[0092] Referring to FIG. **7**, the computing device **710** may receive a first plurality of transformed data elements **752A** from a first computing device **720A**, the first plurality of transformed data elements **752A** representing a transform by

a pseudorandom function of a first plurality of data elements. In some embodiments, the first plurality of transformed data elements **752A** from a first computing device **720A** may be similar to the transformed set of values **152** (e.g., the first transformed set of values **152A**) from a client computing device **120** (e.g., the first client computing device **120A**) described herein with respect to FIGS. **1** to **6**. In some embodiments, the first plurality of data elements may be similar to the set of values **150** (e.g., the first set of values **150A**) described herein with respect to FIGS. **1** to **6**. In some embodiments, the first plurality of transformed data elements **752A** are obtained via an oblivious pseudorandom function (OPRF) protocol.

[0093] The computing device **710** may receive a second plurality of transformed data elements **752B** from a second computing device **720B**, the second plurality of transformed data elements **752B** representing a transform of a second plurality of data elements by the pseudorandom function. In some embodiments, the second plurality of transformed data elements **752B** from a second computing device **720B** may be similar to the transformed set of values **152** (e.g., the second transformed set of values **152B**) from a client computing device **120** (e.g., the second client computing device **120B**) described herein with respect to FIGS. **1** to **6**. In some embodiments, the second plurality of transformed data elements **752B** are obtained via the oblivious pseudorandom function (OPRF) protocol.

[0094] The computing device **710** may transmit, to the first computing device **720A** and the second computing device **720B**, an indication of a subset of transformed data elements **775** that are present in both the first plurality of transformed data elements **752A** and the second plurality of transformed data elements **752B**. In some embodiments, the indication of a subset of transformed data elements **775** may be similar to the intersection **154** (e.g., the first intersection **154A** and/or the second intersection **154B**) described herein with respect to FIGS. **1** to **6**.

[0095] The device architecture **700** of FIG. **7** provides an improved capability for detection of an intersection of a plurality of private sets. The device architecture **700** may provide a mechanism by which multiple computing devices **720A**, **720B** are able to determine if any element in their set of private data is shared with another computing device, without sharing the contents of the private data. Moreover, the intersection may be determined without identifying the party with the shared value. In some embodiments, the generation of the first and second transformed data elements **752A**, **752B** may be separated from the evaluation of the intersection of the subset of transformed data elements **775**, which may reduce a possibility that the underlying data elements may be discovered. In some embodiments, the underlying data elements may be passwords, or other secure data, and the embodiments of the present disclosure may allow for them to be compared without risking a breach of the underlying information. Thus, embodiments of the device architecture **700** improve the technology associated with computer security, and improve the operation of the computing device **710** by allowing for the subset of transformed data elements **775** to be discovered more efficiently than typical techniques.

[0096] FIG. **8** is a block diagram of an example computing device **800** that may perform one or more of the operations described herein, in accordance with some embodiments of the disclosure. Computing device **800** may be connected to

other computing devices in a LAN, an intranet, an extranet, and/or the Internet. The computing device may operate in the capacity of a server machine in client-server network environment or in the capacity of a client in a peer-to-peer network environment. The computing device may be provided by a personal computer (PC), a set-top box (STB), a server, a network router, switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single computing device is illustrated, the term “computing device” shall also be taken to include any collection of computing devices that individually or jointly execute a set (or multiple sets) of instructions to perform the methods discussed herein.

[0097] The example computing device 800 may include a processing device (e.g., a general purpose processor, a PLD, etc.) 802, a main memory 804 (e.g., synchronous dynamic random access memory (DRAM), read-only memory (ROM)), a static memory 806 (e.g., flash memory) and a data storage device 818, which may communicate with each other via a bus 830.

[0098] Processing device 802 may be provided by one or more general-purpose processing devices such as a microprocessor, central processing unit, or the like. In an illustrative example, processing device 802 may include a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. Processing device 802 may also include one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device 802 may execute the operations described herein, in accordance with one or more aspects of the present disclosure, for performing the operations and steps discussed herein.

[0099] Computing device 800 may further include a network interface device 808 which may communicate with a network 820. The computing device 800 also may include a video display unit 810 (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)), an alphanumeric input device 812 (e.g., a keyboard), a cursor control device 814 (e.g., a mouse) and an acoustic signal generation device 816 (e.g., a speaker). In one embodiment, video display unit 810, alphanumeric input device 812, and cursor control device 814 may be combined into a single component or device (e.g., an LCD touch screen).

[0100] Data storage device 818 may include a computer-readable storage medium 828 on which may be stored one or more sets of instructions 825 that may include instructions for a transformation engine 130 and/or an evaluation engine 140 for carrying out the operations described herein, in accordance with one or more aspects of the present disclosure. Instructions 825 may also reside, completely or at least partially, within main memory 804 and/or within processing device 802 during execution thereof by computing device 800, main memory 804 and processing device 802 also constituting computer-readable media. The instructions 825 may further be transmitted or received over a network 820 via network interface device 808.

[0101] While computer-readable storage medium 828 is shown in an illustrative example to be a single medium, the

term “computer-readable storage medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database and/or associated caches and servers) that store the one or more sets of instructions. The term “computer-readable storage medium” shall also be taken to include any medium that is capable of storing, encoding or carrying a set of instructions for execution by the machine and that cause the machine to perform the methods described herein. The term “computer-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media and magnetic media.

[0102] Unless specifically stated otherwise, terms such as “receiving,” “transmitting,” “generating,” “providing,” “inserting,” “determining,” or the like, refer to actions and processes performed or implemented by computing devices that manipulates and transforms data represented as physical (electronic) quantities within the computing device’s registers and memories into other data similarly represented as physical quantities within the computing device memories or registers or other such information storage, transmission or display devices. Also, the terms “first,” “second,” “third,” “fourth,” etc., as used herein are meant as labels to distinguish among different elements and may not necessarily have an ordinal meaning according to their numerical designation.

[0103] Examples described herein also relate to an apparatus for performing the operations described herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general purpose computing device selectively programmed by a computer program stored in the computing device. Such a computer program may be stored in a computer-readable non-transitory storage medium.

[0104] The methods and illustrative examples described herein are not inherently related to any particular computer or other apparatus. Various general purpose systems may be used in accordance with the teachings described herein, or it may prove convenient to construct more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will appear as set forth in the description above.

[0105] The above description is intended to be illustrative, and not restrictive. Although the present disclosure has been described with references to specific illustrative examples, it will be recognized that the present disclosure is not limited to the examples described. The scope of the disclosure should be determined with reference to the following claims, along with the full scope of equivalents to which the claims are entitled.

[0106] As used herein, the term “engine” is not intended to be limiting of any particular implementation for accomplishing and/or performing the actions, operations, processes, etc., attributable to and/or performed by the engine. An engine may be, but is not limited to, software, hardware and/or firmware or any combination thereof that performs the specified functions including, but not limited to, any use of a general and/or specialized processing device in combination with appropriate software loaded or stored in a machine readable memory and executed by the processing device. Further, any name associated with a particular engine is, unless otherwise specified, for purposes of convenience of reference and not intended to be limiting to a specific implementation. Additionally, any functionality

attributed to an engine may be equally performed by multiple engines, incorporated into and/or combined with the functionality of another engine of the same or different type, or distributed across one or more engines of various configurations.

[0107] As used herein, the singular forms “a”, “an” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “comprises”, “comprising”, “includes”, and/or “including”, when used herein, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof. Therefore, the terminology used herein is for the purpose of describing particular embodiments only and is not intended to be limiting. As used herein, the term “and/or” includes any and all combinations of one or more of the associated listed items.

[0108] It should also be noted that in some alternative implementations, the functions/acts noted may occur out of the order noted in the figures. For example, two figures shown in succession may in fact be executed substantially concurrently or may sometimes be executed in the reverse order, depending upon the functionality/acts involved.

[0109] Although the method operations were described in a specific order, it should be understood that other operations may be performed in between described operations, described operations may be adjusted so that they occur at slightly different times or the described operations may be distributed in a system which allows the occurrence of the processing operations at various intervals associated with the processing.

[0110] Various units, circuits, or other components may be described or claimed as “configured to” or “configurable to” perform a task or tasks. In such contexts, the phrase “configured to” or “configurable to” is used to connote structure by indicating that the units/circuits/components include structure (e.g., circuitry) that performs the task or tasks during operation. As such, the unit/circuit/component can be said to be configured to perform the task, or configurable to perform the task, even when the specified unit/circuit/component is not currently operational (e.g., is not on). The units/circuits/components used with the “configured to” or “configurable to” language include hardware—for example, circuits, memory storing program instructions executable to implement the operation, etc. Reciting that a unit/circuit/component is “configured to” perform one or more tasks, or is “configurable to” perform one or more tasks, is expressly intended not to invoke 35 U.S.C. § 112(f) for that unit/circuit/component. Additionally, “configured to” or “configurable to” can include generic structure (e.g., generic circuitry) that is manipulated by software and/or firmware (e.g., an FPGA or a general-purpose processor executing software) to operate in manner that is capable of performing the task(s) at issue. “Configured to” may also include adapting a manufacturing process (e.g., a semiconductor fabrication facility) to fabricate devices (e.g., integrated circuits) that are adapted to implement or perform one or more tasks. “Configurable to” is expressly intended not to apply to blank media, an unprogrammed processor or unprogrammed generic computer, or an unprogrammed programmable logic device, programmable gate array, or other unprogrammed device, unless accompanied by programmed

media that confers the ability to the unprogrammed device to be configured to perform the disclosed function(s).

[0111] The foregoing description, for the purpose of explanation, has been described with reference to specific embodiments. However, the illustrative discussions above are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The embodiments were chosen and described in order to best explain the principles of the embodiments and its practical applications, to thereby enable others skilled in the art to best utilize the embodiments and various modifications as may be suited to the particular use contemplated. Accordingly, the present embodiments are to be considered as illustrative and not restrictive, and the invention is not to be limited to the details given herein, but may be modified within the scope and equivalents of the appended claims.

What is claimed is:

1. A method comprising:

receiving, at a third computing device, a first plurality of transformed data elements from a first computing device;

receiving, at the third computing device, a second plurality of transformed data elements from a second computing device, wherein an identity of the first computing device is unknown to the second computing device and an identity of the second computing device is unknown to the first computing device; and

transmitting, by a processing device executing on the third computing device to the first computing device and the second computing device, an indication of a subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements.

2. The method of claim 1, further comprising:

generating, by a fourth computing device, a pseudorandom key;

receiving, at the fourth computing device, a set of intermediate values from the first computing device, wherein the set of intermediate values represents a transformation of a set of data elements by the first computing device using a first hash function;

computing, by the fourth computing device, a first pseudorandom function using the pseudorandom key and the set of intermediate values;

producing, by the fourth computing device, the first plurality of transformed data elements based on the first pseudorandom function; and

providing, by the fourth computing device, the first plurality of transformed data elements to the first computing device.

3. The method of claim 1, further comprising:

storing, by the third computing device, the first plurality of transformed data elements in a first hash structure, wherein the first hash structure is devoid of information pertaining to the first computing device; and

storing, by the third computing device, the second plurality of transformed data elements in a second hash structure, wherein the second hash structure is devoid of information pertaining to the second computing device and is separate from the first hash structure;

determining, by the third computing device, an intersection of the first hash structure and the second hash structure; and

identifying the subset of transformed data elements based in the intersection of the first hash structure and the second hash structure.

4. The method of claim 1, further comprising:
generating the first plurality of transformed data elements based on a first oblivious pseudorandom function (OPRF) protocol executed between the first computing device and a fourth computing device using a key of the fourth computing device; and
generating the second plurality of transformed data elements based on a second OPRF protocol executed between the second computing device and the fourth computing device using the key of the fourth computing device.

5. The method of claim 4, wherein the third computing device and the fourth computing device are logically isolated on separate virtual machines.

6. The method of claim 1, further comprising:
inserting the first plurality of transformed data elements and the second plurality of transformed data elements into a common hash structure; and
determining the subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements by analyzing the common hash structure.

7. The method of claim 1, wherein the first plurality of transformed data elements and the second plurality of transformed data elements correspond to a first and second set of passwords of the first computing device and the second computing device, respectively.

8. A system comprising:
a memory; and
a processing device, operatively coupled to the memory, to:
receive a first plurality of transformed data elements from a first computing device;
receive a second plurality of transformed data elements from a second computing device, wherein an identity of the first computing device is unknown to the second computing device and an identity of the second computing device is unknown to the first computing device; and
transmit, to the first computing device and the second computing device, an indication of a subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements.

9. The system of claim 8, wherein the processing device is further to:
generate a pseudorandom key;
receive a set of intermediate values from the first computing device, wherein the set of intermediate values represents a transformation of a set of data elements by the first computing device using a first hash function;
compute a first pseudorandom function using the pseudorandom key and the set of intermediate values;
produce the first plurality of transformed data elements based on the first pseudorandom function; and
provide the first plurality of transformed data elements to the first computing device.

10. The system of claim 9, wherein a third computing device that generates the pseudorandom key is logically isolated from a fourth computing device that receives the

first plurality of transformed data elements and the second plurality of transformed data elements.

11. The system of claim 8, wherein the processing device is further to:

store the first plurality of transformed data elements in a first hash structure, wherein the first hash structure is devoid of information pertaining to the first computing device; and

store the second plurality of transformed data elements in a second hash structure, wherein the second hash structure is devoid of information pertaining to the second computing device and is separate from the first hash structure;

determine an intersection of the first hash structure and the second hash structure; and

identify the subset of transformed data elements based in the intersection of the first hash structure and the second hash structure.

12. The system of claim 8, wherein the processing device is further to:

generate the first plurality of transformed data elements based on a first oblivious pseudorandom function (OPRF) protocol executed between the first computing device and a fourth computing device using a key of the fourth computing device; and

generate the second plurality of transformed data elements based on a second OPRF protocol executed between the second computing device and the fourth computing device using the key of the fourth computing device.

13. The system of claim 8, wherein the processing device is further to:

insert the first plurality of transformed data elements and the second plurality of transformed data elements into a common hash structure; and

determine the subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements by analyzing the common hash structure.

14. The system of claim 8, wherein the first plurality of transformed data elements and the second plurality of transformed data elements correspond to a first and second set of passwords of the first computing device and the second computing device, respectively.

15. A non-transitory computer-readable storage medium including instructions that, when executed by a processing device, cause the processing device to:

receive a first plurality of transformed data elements from a first computing device;

receive a second plurality of transformed data elements from a second computing device, wherein an identity of the first computing device is unknown to the second computing device and an identity of the second computing device is unknown to the first computing device; and

transmit, by the processing device to the first computing device and the second computing device, an indication of a subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements.

16. The non-transitory computer-readable storage medium of claim 15, wherein the processing device is further to:

generate a pseudorandom key;
receive a set of intermediate values from the first computing device, wherein the set of intermediate values represents a transformation of a set of data elements by the first computing device using a first hash function;
compute a first pseudorandom function using the pseudorandom key and the set of intermediate values;
produce the first plurality of transformed data elements based on the first pseudorandom function; and
provide the first plurality of transformed data elements to the first computing device.

17. The non-transitory computer-readable storage medium of claim **15**, wherein the processing device is further to:

store the first plurality of transformed data elements in a first hash structure, wherein the first hash structure is devoid of information pertaining to the first computing device; and
store the second plurality of transformed data elements in a second hash structure, wherein the second hash structure is devoid of information pertaining to the second computing device and is separate from the first hash structure;
determine an intersection of the first hash structure and the second hash structure; and
identify the subset of transformed data elements based in the intersection of the first hash structure and the second hash structure.

18. The non-transitory computer-readable storage medium of claim **15**, wherein the processing device is further to:

generate the first plurality of transformed data elements based on a first oblivious pseudorandom function (OPRF) protocol executed between the first computing device and a fourth computing device using a key of the fourth computing device; and
generate the second plurality of transformed data elements based on a second OPRF protocol executed between the second computing device and the fourth computing device using the key of the fourth computing device.

19. The non-transitory computer-readable storage medium of claim **15**, wherein the processing device is further to:

insert the first plurality of transformed data elements and the second plurality of transformed data elements into a common hash structure; and
determine the subset of transformed data elements that are present in both the first plurality of transformed data elements and the second plurality of transformed data elements by analyzing the common hash structure.

20. The non-transitory computer-readable storage medium of claim **15**, wherein the first plurality of transformed data elements and the second plurality of transformed data elements correspond to a first and second set of passwords of the first computing device and the second computing device, respectively.

* * * * *