

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250278410

Kind Code

A1

Publication Date

September 04, 2025

Inventor(s)

Pandis; Ippokratis et al.

HYBRID TRANSACTIONAL AND ANALYTICAL PROCESSING ARCHITECTURE FOR OPTIMIZATION OF REAL-TIME ANALYTICAL QUERYING

Abstract

Methods for replicating transactional tables of a transactional database to an analytical database and maintaining updates to those transactional table representations are disclosed. Snapshots of the transactional tables are provided to the analytical database via a transport mechanism, such as a data storage service or a data streaming service, and stored at the analytical database. Then, checkpoints comprising portions of a change-data-capture log that has recorded transactional changes to the transactional tables of the transactional database are provided to the analytical database via the same or different transport mechanism and used to commit those transactional changes to the snapshot representations. The snapshot representations may be used to respond to incoming analytical queries in order to provide real-time querying results.

Inventors: Pandis; Ippokratis (Menlo Park, CA), Soundararajan; Gokul (San Jose, CA), Paliwal; Gopal (Milpitas, CA), Rajgaria; Punit (Saratoga, CA), Basu; Sanuj (San Mateo, CA), Green; Todd Jeffrey (Davis, CA), Saxena; Gaurav (Cupertino, CA), Skipin; Vadim (Berlin, DE), Wust; Johannes (Kleinmachnow, DE), Satyanarayana; Hemanth (Santa Clara, CA), Abrams; Matthew Perry (Las Cruces, NM), Brahmadesam; Murali (Tiruchirappalli, IN)

Applicant: Amazon Technologies, Inc. (Seattle, WA)

Family ID: 89433263

Assignee: Amazon Technologies, Inc. (Seattle, WA)

Appl. No.: 19/079910

Filed: March 14, 2025

Related U.S. Application Data

Publication Classification

Int. Cl.: **G06F16/27** (20190101); **G06F16/23** (20190101); **G06F16/25** (20190101); **G06F16/28** (20190101)

U.S. Cl.:

CPC **G06F16/27** (20190101); **G06F16/23** (20190101); **G06F16/2379** (20190101); **G06F16/25** (20190101); **G06F16/273** (20190101); **G06F16/275** (20190101); **G06F16/278** (20190101); **G06F16/283** (20190101);

Background/Summary

PRIORITY CLAIM [0001] This application is a continuation of U.S. patent application Ser. No. 17/810,314, filed Jun. 30, 2022, which is hereby incorporated by reference herein in its entirety.

BACKGROUND

[0002] As the technological capacity for organizations to create, track, and retain information continues to grow, a variety of different technologies for managing and storing the rising tide of information have been developed. Database systems, for example, provide clients with many different specialized or customized configurations of hardware and software to manage stored information. However, the increasing amounts of data that organizations must store and manage often correspondingly increases both the size and complexity of data storage and management technologies, like database systems, which in turn escalate the cost of maintaining the information. [0003] New technologies more and more seek to reduce both the complexity and storage requirements of maintaining data while simultaneously improving the efficiency of data processing and querying. Challenges in obtaining the right configuration of data storage, processing, and querying, such that these database systems may be efficiently configured to perform various functions for different workloads occurs frequently.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 illustrates a service provider network that includes at least a transactional database service and an analytical database service such that clients of the service provider network may both maintain transactional data and run analytical queries against the transactional data, according to some embodiments.

[0005] FIG. 2 illustrates various components of a transactional database system that includes a database engine and a separate distributed database storage service, according to some embodiments.

[0006] FIG. 3 illustrates a distributed storage system of a transactional database system, according to some embodiments.

[0007] FIG. 4A illustrates various components of an analytical database system configured to use warm and cold storage tiers to store data blocks for clients of an analytical database service, wherein the warm storage tier comprises one or more node clusters associated with said clients,

according to some embodiments.

[0008] FIG. 4B illustrates an example of a node cluster of an analytical database system performing queries against transactional database data, according to some embodiments.

[0009] FIG. 5 is a flow diagram illustrating a process of maintaining, within an analytical database system, a representation of portions of a transactional data table from a transactional database system, according to some embodiments.

[0010] FIG. 6 illustrates the process of a handshake protocol, used to negotiate and define the configurations and parameters for maintaining, at an analytical database, a representation of a table stored in a transactional database, according to some embodiments.

[0011] FIG. 7 is a flow diagram illustrating a process of initiating and performing a handshake protocol, used to negotiate and define the configurations and parameters for maintaining, at an analytical database, a representation of a table stored in a transactional database, according to some embodiments.

[0012] FIG. 8 illustrates the use of a transport mechanism to provide snapshots and checkpoints from a transactional database to an analytical database, according to some embodiments.

[0013] FIG. 9A illustrates a process of treating, at a transactional database, transactional changes to a transactional table and providing said transactional changes in a checkpoint to a transport mechanism for replication at a representation of the transactional table at an analytical database, according to some embodiments.

[0014] FIG. 9B illustrates a process of receiving, at an analytical database, a checkpoint of transactional changes that are then implemented to a representation of a transactional table maintained at the analytical database, according to some embodiments.

[0015] FIG. 10 is a flow diagram illustrating a process of implementing transactional changes to a representation of a transactional table in which delete events are first buffered and committed to a shadow table before being applied and committed as a batch to the representation, according to some embodiments.

[0016] FIGS. 11A-D illustrate the state of a table representation maintained at an analytical database and a corresponding shadow table, checkpoint, and table as viewed by an external client of the analytical database just prior to the implementation of a given transaction number **1180**, according to some embodiments.

[0017] FIGS. 12A-D illustrate the state of the table representation, shadow table, checkpoint, and table view of FIGS. 11A-D just after the implementation of transaction number **1180**, according to some embodiments.

[0018] FIGS. 13A-D illustrate the state of the table representation, shadow table, checkpoint, and table view of FIGS. 11A-D at a point after the implementation of transaction number **1180**, according to some embodiments.

[0019] FIG. 14 is a block diagram illustrating an example computing device that may be used in at least some embodiments.

[0020] While embodiments are described herein by way of example for several embodiments and illustrative drawings, those skilled in the art will recognize that embodiments are not limited to the embodiments or drawings described. It should be understood, that the drawings and detailed description thereto are not intended to limit embodiments to the particular form disclosed, but on the contrary, the intention is to cover all modifications, equivalents and alternatives falling within the spirit and scope as defined by the appended claims. The headings used herein are for organizational purposes only and are not meant to be used to limit the scope of the description or the claims. As used throughout this application, the word “may” is used in a permissive sense (i.e., meaning having the potential to), rather than the mandatory sense (i.e., meaning must). Similarly, the words “include,” “including,” and “includes” mean including, but not limited to.

[0021] It will also be understood that, although the terms first, second, etc. may be used herein to describe various elements, these elements should not be limited by these terms. These terms are

only used to distinguish one element from another. For example, a first contact could be termed a second contact, and, similarly, a second contact could be termed a first contact, without departing from the scope of the present invention. The first contact and the second contact are both contacts, but they are not the same contact.

DETAILED DESCRIPTION OF EMBODIMENTS

[0022] Various techniques pertaining to a hybrid transactional and analytical processing (HTAP) service are described. In some embodiments, a hybrid transactional and analytical processing system, which may implement at least a transactional database and an analytical database, may be used to maintain tables of transactional data at the transactional database, and maintain replicas of said tables at the analytical database. Such a hybrid transactional and analytical processing service may be optimized for both online transaction processing (OLTP) and online analytical processing (OLAP) related services, according to some embodiments. In order to maintain the replicas, or representations, of the transactional tables at the analytical database, a change-data-capture log of transactional changes made to the tables at the transactional database may be provided to the analytical database and incrementally applied and committed to the representations.

[0023] Running analytical queries against a transactional data store of the transactional database may impact the performance of the transactional queries, impact the performance of the computing resources of the transactional database, and, in some cases which may require leveraging materialized views and/or special indices, lead to a complex and/or challenging organization of database resources. In addition, scaling the structure of the transactional data stores of the transactional database such that they may be configured to treat analytical queries may be costly. On the other hand, “offloading” transactional data to an analytical database that is more optimized for analytical queries and analytical query management may be difficult to manage manually and/or lead to a lag (e.g., stale data). Techniques proposed herein, however, overcome these challenges by making use of the analytical database for running analytical queries against transactional data while minimizing the lag between the transactional data stored on the transactional database and the “offloaded” transactional data replications maintained on the analytical database, resulting in real-time analytics on data.

[0024] This specification continues with a general description of a service provider network that implements a hybrid transactional and analytical processing service, including a transactional database service and an analytical database service, that is configured to maintain transactional data, allow for querying against the transactional data, and support multiversion concurrency control (MVCC). Then, various examples of the hybrid transactional and analytical processing service, including different components/modules, or arrangements of components/module that may be employed as part of implementing the services are discussed. A number of different methods and techniques to maintain a representation in the analytical database service of a transactional table of the transactional database service are then discussed, some of which are illustrated in accompanying flowcharts. For example, methods and techniques for performing a handshake protocol that may define parameters and functionalities of the hybrid transactional and analytical processing service are described. In addition, methods and techniques for buffering delete events of transactional changes received by the analytical database such that they may be applied and committed to the representations as a background operation of the analytical database are described. Finally, a description of an example computing system upon which the various components, modules, systems, devices, and/or nodes may be implemented is provided. Various examples are provided throughout the specification. A person having ordinary skill in the art should also understand that the previous and following description of a hybrid transactional and analytical processing service is a logical description and thus is not to be construed as limiting as to the implementation of the hybrid transactional and analytical processing service, or portions thereof.

[0025] FIG. 1 illustrates a service provider network that includes at least a transactional database service and an analytical database service such that clients of the service provider network may

both maintain transactional data and run analytical queries against the transactional data, according to some embodiments.

[0026] In some embodiments, a hybrid transactional and analytical processing service may be implemented within a service provider network, such as service provider network **100**. In some embodiments, service provider network **100** may implement various computing resources or services, such as database service(s), (e.g., relational database services, non-relational database services, a map reduce service, a data warehouse service, data storage services, such as data storage service **120** (e.g., object storage services or block-based storage services that may implement a centralized data store for various types of data), and/or any other type of network based services (which may include a virtual compute service and various other types of storage, processing, analysis, communication, event handling, visualization, and security services not illustrated).

[0027] In some embodiments, a transactional database service, such as transactional database service **110**, may be configured to store and maintain tables of transactional data items for client(s) of the transactional database service. For some clients of the transactional database service, further optimization of both transactional data processing and query processing against said transactional data may be made if tables of transactional data items are replicated and maintained in an analytical database service, such as analytical database service **150**. In such a manner, processing and/or computing resources of the transactional database service may remain focused on processing transactional data without interference from potentially compute-intensive analytical query processing. By “outsourcing” such analytical query requests to an analytical database service, clients of the transactional database service may obtain near real-time analytical query results from the replicated tables in the analytical database service without limiting or taking away the computing resources of the transactional database service from transactional data processing.

[0028] In order to provide both initial replicas (e.g., snapshots) of the tables to the analytical database service and subsequent updates (e.g., checkpoints, or segments/portions of a change-data-capture log) that should be applied to the snapshots in order to maintain them at the analytical database service, one or more additional services of service provider network **100** may be used as transport mechanisms for the hybrid transactional and analytical processing service. For example, a data storage service, such as data storage service **120**, may be used to provide access to such snapshots and/or checkpoints for the analytical database system. In addition to (or instead of) the data storage service, a data streaming service, such as data streaming service **130** may be used to stream the snapshots and/or checkpoints to the analytical database system. A person having ordinary skill in the art should understand that additional embodiments using other transport mechanisms may similarly result in the transport of snapshots and checkpoints from the transactional database to the analytical database, and may include the use of other service(s) **140** of service provider network **100**.

[0029] As shown in the figure, multiple access points (e.g., client endpoints) may be used such that clients may access the different services of service provider network **100** more directly. For example, a client of clients **170** may have accounts with at least transactional database service **110** and analytical database service **150**, and may be able to access these services of service provider network **110** through network **160**. In some embodiments, a same or different network connection may be used at these different access points. Network **160** may represent the same network connection or multiple different network connections, according to some embodiments. For example, network **160** may generally encompass the various telecommunications networks and service providers that collectively implement the Internet. Network **160** may also include private networks such as local area networks (LANs) or wide area networks (WANs) as well as public or private wireless networks. For example, both a given client of clients **170** and/or **180** and the various network-based services of service provider network **100** may be respectively provisioned within enterprises having their own internal networks. In such an embodiment, network **160** may include the hardware (e.g., modems, routers, switches, load balancers, proxy servers, etc.) and

software (e.g., protocol stacks, accounting software, firewall/security software, etc.) necessary to establish a networking link between the given client and the Internet as well as between the Internet and the various network-based services of service provider network **100**. It is noted that in some embodiments, clients **170** and/or **180** may communicate with services of service provider network **100** using a private network rather than the public Internet. For example, clients **170**, **180**, and/or **190** may be provisioned within the same enterprise as various services of service provider network **100**. In such a case, clients **170**, **180**, and/or **190** may communicate with the various services of service provider network **100** entirely through a private network **160** (e.g., a LAN or WAN that may use Internet-based communication protocols but which is not publicly accessible).

[0030] The systems described herein may, in some embodiments, implement a network-based services that enables clients (e.g., subscribers) to operate a data storage system in a cloud computing environment. In some embodiments, the data storage system may be an enterprise-class database system that is highly scalable and extensible. In some embodiments, queries may be directed to database storage that is distributed across multiple physical resources, and the database system may be scaled up or down on an as needed basis. The database system may work effectively with database schemas of various types and/or organizations, in different embodiments. In some embodiments, clients/subscribers may submit queries in a number of ways, e.g., interactively via an SQL interface to the database system. In other embodiments, external applications and programs may submit queries using Open Database Connectivity (ODBC) and/or Java Database Connectivity (JDBC) driver interfaces to the database system.

[0031] FIG. 2 illustrates various components of a transactional database system that includes a database engine and a separate distributed database storage service, according to some embodiments.

[0032] In some embodiments, transactional database system **200** described herein may, in some embodiments, implement a service-oriented architecture in which various functional components of the database system are intrinsically distributed. For example, transactional database system **200** may be organized such that the basic operations of a database (e.g., query processing, transaction management, caching and storage) are managed in tiers that may be individually and independently scalable. For example, in some embodiments, transactional database system **200** may include a database tier (which may include a single database engine head node and a client-side storage system driver), and a separate, distributed storage system (which may include multiple storage nodes that collectively perform some of the operations traditionally performed in the database tier of existing systems).

[0033] As described in more detail herein, in some embodiments, some of the lowest level operations of a database (e.g., backup, restore, snapshot, recovery, log record manipulation, and/or various space management operations) may be offloaded from the database engine to the storage layer (or tier), such as a distributed storage system, and distributed across multiple nodes and storage devices. For example, in some embodiments, rather than the database engine applying changes to a database (or data pages thereof) and then sending the modified data pages to the storage layer, the application of changes to the stored database (and data pages thereof) may be the responsibility of the storage layer itself. In such embodiments, redo log records, rather than modified data pages, may be sent to the storage layer, after which redo processing (e.g., the application of the redo log records) may be performed in a distributed manner (e.g., by a background process). Log sequence numbers may be assigned to the redo log records from a log sequence number space. In some embodiments, crash recovery (e.g., the rebuilding of data pages from stored redo log records) may also be performed by the storage layer and may also be performed by a distributed background process.

[0034] In some embodiments, because only redo logs (and not modified data pages) are sent to the storage layer, there may be much less network traffic between the database tier and the storage layer than in existing database systems. In some embodiments, each redo log may be on the order

of one-tenth the size of the corresponding data page for which it specifies a change. Note that requests sent from the database tier and the distributed storage system may be asynchronous and that multiple such requests may be in flight at a time.

[0035] In general, after being given a piece of data, a primary requirement of a database is that it can eventually give that piece of data back. To do this, the database may include several different components (or tiers), each of which performs a different function. For example, a traditional database may be thought of as having three tiers: a first tier for performing query parsing, optimization and execution; a second tier for providing transactionality, recovery, and durability; and a third tier that provides storage, either on locally attached disks or on network-attached storage. As noted above, previous attempts to scale a traditional database have typically involved replicating all three tiers of the database and distributing those replicated database instances across multiple machines.

[0036] In some embodiments, the transactional database system described herein may partition functionality of a database system differently than in a traditional database, and may distribute only a subset of the functional components (rather than a complete database instance) across multiple machines in order to implement scaling. For example, in some embodiments, a client-facing tier may be configured to receive a request specifying what data is to be stored or retrieved, but not how to store or retrieve the data. This tier may perform request parsing and/or optimization (e.g., SQL parsing and optimization), while another tier may be responsible for query execution. In some embodiments, a third tier may be responsible for providing transactionality and consistency of results. For example, this tier may be configured to enforce some of the so-called ACID properties, in particular, the Atomicity of transactions that target the database, maintaining Consistency within the database, and ensuring Isolation between the transactions that target the database. In some embodiments, a fourth tier may then be responsible for providing Durability of the stored data in the presence of various sorts of faults. For example, this tier may be responsible for change logging, recovery from a database crash, managing access to the underlying storage volumes and/or space management in the underlying storage volumes.

[0037] In various embodiments, a database instance may include multiple functional components (or layers), each of which provides a portion of the functionality of the database instance. In one such example, a database instance may include a query parsing and query optimization layer, a query execution layer, a transactionality and consistency management layer, and a durability and space management layer. As noted above, in some existing database systems, scaling a database instance may involve duplicating the entire database instance one or more times (including all of the example layers), and then adding glue logic to stitch them together. In some embodiments, the systems described herein may instead offload the functionality of durability and space management layer from the database tier to a separate storage layer, and may distribute that functionality across multiple storage nodes in the storage layer.

[0038] In some embodiments, the transactional database systems described herein may retain much of the structure of the upper half of the database instance, such as query parsing and query optimization layer, a query execution layer, and a transactionality and consistency management layer, but may redistribute responsibility for at least portions of the backup, restore, snapshot, recovery, and/or various space management operations to the storage tier. Redistributing functionality in this manner and tightly coupling log processing between the database tier and the storage tier may improve performance, increase availability and reduce costs, when compared to previous approaches to providing a scalable database. For example, network and input/output bandwidth requirements may be reduced, since only redo log records (which are much smaller in size than the actual data pages) may be shipped across nodes or persisted within the latency path of write operations. In addition, the generation of data pages can be done independently in the background on each storage node (as foreground processing allows), without blocking incoming write operations. In some embodiments, the use of log-structured, non-overwrite storage may allow

backup, restore, snapshots, point-in-time recovery, and volume growth operations to be performed more efficiently, e.g., by using metadata manipulation rather than movement or copying of a data page. In some embodiments, the storage layer may also assume the responsibility for the replication of data stored on behalf of clients (and/or metadata associated with that data, such as redo log records) across multiple storage nodes. For example, data (and/or metadata) may be replicated locally (e.g., within a single “availability zone” in which a collection of storage nodes executes on its own physically distinct, independent infrastructure) and/or across availability zones in a single region or in different regions.

[0039] In various embodiments, the database systems described herein may support a standard or custom application programming interface (API) for a variety of database operations. For example, the API may support operations for creating a database, creating a table, altering a table, creating a user, dropping a user, inserting one or more rows in a table, copying values, selecting data from within a table (e.g., querying a table), canceling or aborting a query, creating a snapshot, and/or other operations.

[0040] In some embodiments, the database tier of a database instance may include a database engine head node server that receives read and/or write requests from various client programs (e.g., applications) and/or subscribers (users), then parses them and develops an execution plan to carry out the associated database operation(s). For example, the database engine head node may develop the series of steps necessary to obtain results for complex queries and joins. In some embodiments, the database engine head node may manage communications between the database tier of the database system and clients/subscribers, as well as communications between the database tier and a separate distributed storage system.

[0041] In some embodiments, the database engine head node may be responsible for receiving SQL requests from end clients through a JDBC or ODBC interface and for performing SQL processing and transaction management (which may include locking) locally. However, rather than generating data pages locally, the database engine head node (or various components thereof) may generate redo log records and may ship them to the appropriate nodes of a separate distributed storage system. In some embodiments, a client-side driver for the distributed storage system may be hosted on the database engine head node and may be responsible for routing redo log records to the storage system node (or nodes) that store the segments (or data pages thereof) to which those redo log records are directed. For example, in some embodiments, each segment may be mirrored (or otherwise made durable) on multiple storage system nodes that form a protection group. In such embodiments, the client-side driver may keep track of the nodes on which each segment is stored and may route redo logs to all of the nodes on which a segment is stored (e.g., asynchronously and in parallel, at substantially the same time), when a client request is received. As soon as the client-side driver receives an acknowledgement back from a write quorum of the storage nodes in the protection group (which may indicate that the redo log record has been written to the storage node), it may send an acknowledgement of the requested change to the database tier (e.g., to the database engine head node). For example, in embodiments in which data is made durable through the use of protection groups, the database engine head node may not be able to commit a transaction until and unless the client-side driver receives a reply from enough storage node instances to constitute a write quorum, as may be defined in a protection group policy for the data.

[0042] In some embodiments, the database tier (or more specifically, the database engine head node) may include a cache in which recently accessed data pages are held temporarily. In such embodiments, if a write request is received that targets a data page held in such a cache, in addition to shipping a corresponding redo log record to the storage layer, the database engine may apply the change to the copy of the data page held in its cache. However, unlike in other database systems, a data page held in this cache may not ever be flushed to the storage layer, and it may be discarded at any time (e.g., at any time after the redo log record for a write request that was most recently applied to the cached copy has been sent to the storage layer and acknowledged). The cache may

implement any of various locking mechanisms to control access to the cache by at most one writer (or multiple readers) at a time, in different embodiments. Note, however, that in embodiments that include such a cache, the cache may not be distributed across multiple nodes, but may exist only on the database engine head node for a given database instance. Therefore, there may be no cache coherency or consistency issues to manage.

[0043] In some embodiments, the client-side driver running on the database engine head node may expose a private interface to the storage tier. In some embodiments, it may also expose a traditional iSCSI interface to one or more other components (e.g., other database engines or virtual computing services components). In some embodiments, storage for a database instance in the storage tier may be modeled as a single volume that can grow in size without limits, and that can have an unlimited number of IOPS associated with it. When a volume is created, it may be created with a specific size, with a specific availability/durability characteristic (e.g., specifying how it is replicated), with an IOPS rate associated with it (e.g., both peak and sustained), and/or user specified encryption. For example, in some embodiments, a variety of different durability models may be supported, and users/subscribers may be able to specify, for their database, a number of replication copies, zones, or regions and/or whether replication is synchronous or asynchronous based upon their durability, performance and cost objectives.

[0044] In some embodiments, the client side driver may maintain metadata about the volume and may directly send asynchronous requests to each of the storage nodes necessary to fulfill read requests and write requests without requiring additional hops between storage nodes. The volume metadata may indicate which protection groups, and their respective storage nodes, maintain which partitions of the volume. For example, in some embodiments, in response to a request to make a change to a database, the client-side driver may be configured to determine the protection group, and its one or more nodes that are implementing the storage for the targeted data page, and to route the redo log record(s) specifying that change to those storage nodes. These read and write requests or separate synchronization requests may also include encrypted data which includes keys to decrypt data store of the protection group to which the storage node belongs. The storage nodes may then be responsible for applying the change specified in the redo log record to the targeted data page at some point in the future, including decrypting the received keys and using the received keys to decrypt stored data to apply the change specified in the redo log record. As writes are acknowledged back to the client-side driver, the client-side driver may advance the point at which the volume is durable and may acknowledge commits back to the database tier. As previously noted, in some embodiments, the client-side driver may not ever send data pages to the storage node servers. This may not only reduce network traffic, but may also remove the need for the checkpoint or background writer threads that constrain foreground-processing throughput in previous database systems.

[0045] In some embodiments, because accesses to the log-structured data storage for the redo log records may consist of a series of sequential input/output operations (rather than random input/output operations), the changes being made may be tightly packed together. It should also be noted that, in contrast to existing systems in which each change to a data page results in two input/output operations to persistent data storage (one for the redo log and one for the modified data page itself), in some embodiments, the systems described herein may avoid this “write amplification” by coalescing data pages at the storage nodes of the distributed storage system based on receipt of the redo log records. Coalesce operations performed upon data pages may utilize received keys to access encrypted data pages and redo log records at the storage nodes.

[0046] As previously noted, in some embodiments, the storage tier of the transactional database system may be responsible for taking database snapshots. However, because the storage tier implements log-structured storage, taking a snapshot of a data page (e.g., a data block) may include recording a timestamp associated with the redo log record that was most recently applied to the data page/block (or a timestamp associated with the most recent operation to coalesce multiple redo

log records to create a new version of the data page/block), and preventing garbage collection of the previous version of the page/block and any subsequent log entries up to the recorded point in time. In such embodiments, taking a database snapshot may not require reading, copying, or writing the data block, as would be required when employing an off-volume backup strategy. In some embodiments, the space requirements for snapshots may be minimal, since only modified data would require additional space, although user/subscribers may be able to choose how much additional space they want to keep for on-volume snapshots in addition to the active data set. In different embodiments, snapshots may be discrete (e.g., each snapshot may provide access to all of the data in a data page as of a specific point in time) or continuous (e.g., each snapshot may provide access to all versions of the data that existing in a data page between two points in time). In some embodiments, reverting to a prior snapshot may include recording a log record to indicate that all redo log records and data pages since that snapshot are invalid and garbage collectable, and discarding all database cache entries after the snapshot point. In such embodiments, no roll-forward may be required since the storage system will, on a block-by-block basis, apply redo log records to data blocks as requested and in the background across all nodes, just as it does in normal forward read/write processing, which again may utilize a received encryption keys. Crash recovery may thereby be made parallel and distributed across nodes.

[0047] One embodiment of a service system architecture that may be configured to implement a network-based services-based transactional database service is illustrated in FIG. 2. In the illustrated embodiment, a number of clients (shown as clients **250a-250n**) may be configured to interact with transactional database system **200** via a network **260**. Distributed storage service **210** may be implemented as log-structured storage using a single log sequence number space. It is noted that where one or more instances of a given component may exist, reference to that component herein may be made in either the singular or the plural. However, usage of either form is not intended to preclude the other.

[0048] In various embodiments, the components illustrated in FIG. 2 may be implemented directly within computer hardware, as instructions directly or indirectly executable by computer hardware (e.g., a microprocessor or computer system), or using a combination of these techniques. For example, the components of FIG. 2 may be implemented by a system that includes a number of computing nodes (or simply, nodes), each of which may be similar to the computer system embodiment illustrated in FIG. 14 and described below. In various embodiments, the functionality of a given service system component (e.g., a component of the database service or a component of the storage service) may be implemented by a particular node or may be distributed across several nodes. In some embodiments, a given node may implement the functionality of more than one service system component (e.g., more than one database service system component).

[0049] Generally speaking, clients **250** may encompass any type of client configurable to submit network-based services requests to a network-based services platform via network **260**, including requests for database services. For example, a given client **250** may include a suitable version of a web browser, or may include a plug-in module or other type of code module configured to execute as an extension to or within an execution environment provided by a web browser. Alternatively, a client **250** (e.g., a database service clients **250a** through **250n**) may encompass an application such as a database application (or user interface thereof), a media application, an office application or any other application that may make use of persistent storage resources to store and/or access one or more databases. In some embodiments, such an application may include sufficient protocol support (e.g., for a suitable version of Hypertext Transfer Protocol (HTTP)) for generating and processing network-based services requests without necessarily implementing full browser support for all types of network-based data. That is, client **250** may be an application configured to interact directly with the network-based services of service provider network **100**. In some embodiments, client **250** may be configured to generate network-based services requests according to a Representational State Transfer (REST)-style network-based services architecture, a document- or

message-based network-based services architecture, or another suitable network-based services architecture.

[0050] In some embodiments, a client **250** (e.g., a database service client) may be configured to provide access to network-based services-based storage of databases to other applications in a manner that is transparent to those applications. For example, client **250** may be configured to integrate with an operating system or file system to provide storage in accordance with a suitable variant of the storage models described herein. However, the operating system or file system may present a different storage interface to applications, such as a conventional file system hierarchy of files, directories and/or folders. In such an embodiment, applications may not need to be modified to make use of the storage system service model. Instead, the details of interfacing to network-based services of service provider network **100** may be coordinated by client **250** and the operating system or file system on behalf of applications executing within the operating system environment.

[0051] Clients **250** may convey network-based services requests (e.g., read/restore data requests, etc.) to and receive responses from service provider network **100** via network **260**. In various embodiments, network **260** may encompass any suitable combination of networking hardware and protocols necessary to establish network-based-based communications between clients **250** and transactional database system **200** of transactional database service **110**.

[0052] Generally speaking, transactional database service **110** may be configured to implement one or more service endpoints for clients of service provider network **100** which may be configured to receive and process network-based services requests, such as requests to access data pages (or records thereof). For example, transactional database service **110** may include hardware and/or software configured to implement a particular endpoint, such that an HTTP-based network-based services request directed to that endpoint is properly received and processed. In some embodiments, transactional database service **110** may be implemented as a server system configured to receive network-based services requests from clients **250** and to forward them to components of transactional database system **200**. In other embodiments, transactional database service **110** may be configured as a number of distinct systems (e.g., in a distributed topology, such as shown in FIGS. 2 and 3 for transactional database system **200**) implementing load balancing and other request management features configured to dynamically manage large-scale network-based services request processing loads. In various embodiments, transactional database service **110** may be configured to support REST-style or document-based (e.g., SOAP-based) types of network-based services requests.

[0053] In addition to functioning as an addressable endpoint for clients' network-based services requests, in some embodiments, transactional database service **110** may implement various client management features. For example, transactional database service **110** may coordinate the metering and accounting of client usage of network-based services, including storage resources, such as by tracking the identities of requesting clients **250**, the number and/or frequency of client requests, the size of data tables (or records thereof) stored or retrieved on behalf of clients **250**, overall storage bandwidth used by clients **250**, class of storage requested by clients **250**, or any other measurable client usage parameter. Transactional database service **110** may also implement financial accounting and billing systems, or may maintain a database of usage data that may be queried and processed by external systems for reporting and billing of client usage activity. In certain embodiments, transactional database service **110** may be configured to collect, monitor and/or aggregate a variety of storage service system operational metrics, such as metrics reflecting the rates and types of requests received from clients **250**, bandwidth utilized by such requests, system processing latency for such requests, system component utilization (e.g., network bandwidth and/or storage utilization within the storage service system), rates and types of errors resulting from requests, characteristics of stored and requested data pages or records thereof (e.g., size, data type, etc.), or any other suitable metrics. In some embodiments such metrics may be used by system administrators to tune and maintain system components, while in other embodiments such metrics

(or relevant portions of such metrics) may be exposed to clients **250** to enable such clients to monitor their usage of transactional database service **110**.

[0054] In some embodiments, transactional database service **110** may also implement user authentication and access control procedures. For example, for a given network-based services request to access a particular database, transactional database service **110** may be configured to ascertain whether the client **250** associated with the request is authorized to access the particular database. Transactional database service **110** may determine such authorization by, for example, evaluating an identity, password or other credential against credentials associated with the particular database, or evaluating the requested access to the particular database against an access control list for the particular database. For example, if a client **250** does not have sufficient credentials to access the particular database, transactional database service **110** may reject the corresponding network-based services request, for example by returning a response to the requesting client **250** indicating an error condition. Various access control policies may be stored as records or lists of access control information by transactional database service **110**.

[0055] It is noted that while service provider network **100** may represent the primary interface through which clients **250** may access the features of transactional database system **200** that implements transactional database service **110**, it need not represent the sole interface to such features. For example, an alternate API that may be distinct from a network-based services interface may be used to allow clients internal to the enterprise providing transactional database system **200** to bypass network-based services.

[0056] Although not illustrated, in various embodiments distributed storage service **210** may be configured to interface with backup data store, system, service, or device. Various data, such as data pages, log records, and/or any other data maintained by distributed storage service internal clients, such as transactional database service **110** or other virtual computing services, and/or external clients such as clients **250a** through **250n**, may be sent to a backup data store.

[0057] Note that in various embodiments, different storage policies may be implemented by transactional database service **110** and/or distributed storage service **210**. Examples of such storage policies may include a durability policy (e.g., a policy indicating the number of instances of a database (or data page thereof) that will be stored and the number of different nodes on which they will be stored) and/or a load balancing policy (which may distribute databases, or data pages thereof, across different nodes, volumes and/or disks in an attempt to equalize request traffic). In addition, different storage policies may be applied to different types of stored items by various one of the services. For example, in some embodiments, distributed storage service **210** may implement a higher durability for redo log records than for data pages.

[0058] In some embodiments, such as shown in FIG. 2, transactional database system **200** includes a respective database engine head node **220** for each of several databases and a distributed storage service **210** (which may or may not be visible to the clients of the database system, shown as database clients **250a-250n**). As illustrated in this example, one or more of database clients **250a-250n** may access a database engine head node **220** (e.g., head node **220a**, head node **220b**, or head node **220c**, each of which is a component of a respective database instance) via network **260** (e.g., these components may be network-addressable and accessible to the database clients **250a-250n**). However, distributed storage service **210**, which may be employed by transactional database system **200** to store a database volume (such as data pages of one or more databases, as well as redo log records and/or other metadata associated therewith) on behalf of database clients **250a-250n**, and to perform other functions of transactional database system **200** as described herein, may or may not be network-addressable and accessible to the storage clients **250a-250n**, in different embodiments. For example, in some embodiments, distributed storage service **210** may perform various storage, access, change logging, recovery, log record manipulation, and/or space management operations in a manner that is invisible to storage clients **250a-250n**.

[0059] As previously noted, each database instance may include a single database engine head node

220 that receives requests (e.g., a snapshot request, etc.) from various client programs (e.g., applications) and/or subscribers (users), then parses them, optimizes them, and develops an execution plan to carry out the associated database operation(s). In the example illustrated in FIG. 2, a query parsing, optimization, and execution component **205** of database engine head node **220a** may perform these functions for queries that are received from database client **250a** and that target the database instance of which database engine head node **220a** is a component. In some embodiments, query parsing, optimization, and execution component **205** may return query responses to database client **250a**, which may include write acknowledgements, requested data pages (or portions thereof), error messages, and or other responses, as appropriate. As illustrated in this example, database engine head node **220a** may also include a client-side storage service driver **225**, which may route read requests and/or redo log records to various storage nodes within distributed storage service **210**, receive write acknowledgements from distributed storage service **210**, receive requested data pages from distributed storage service **210**, and/or return data pages, error messages, or other responses to query parsing, optimization, and execution component **205** (which may, in turn, return them to database client **250a**). Client-side storage service driver **225** may maintain mapping information about the database volume stored in distributed storage service **210**, such that a particular protection group maintaining a partition of the database volume may be determined. Read requests and redo log records may then be routed to storage nodes that are members of the protection group according to the partition of user data to which the read request is directed or to which the redo log record pertains. Client-side storage service driver **225** may obtain encrypted keys for the different protection groups of a data volume from distributed storage service **210** and provide them to individual storage nodes of the respective protection groups so that the client-side storage service driver sends the key for each protection group to the storage nodes of the protection group when communicating with storage nodes of the protection group.

[0060] In this example, database engine head node **220a** includes a data page cache **235**, in which data pages that were recently accessed may be temporarily held. As illustrated in FIG. 2, database engine head node **220a** may also include a transaction and consistency management component **230**, which may be responsible for providing transactionality and consistency in the database instance of which database engine head node **220a** is a component. For example, this component may be responsible for ensuring the Atomicity, Consistency, and Isolation properties of the database instance and the transactions that are directed to the database instance. As illustrated in FIG. 2, database engine head node **220a** may also include a transaction log **240** and an undo log **245**, which may be employed by transaction and consistency management component **230** to track the status of various transactions and roll back any locally cached results of transactions that do not commit.

[0061] Note that each of the other database engine head nodes **220** illustrated in FIG. 2 (e.g., **220b** and **220c**) may include similar components and may perform similar functions for queries received by one or more of database clients **250a-250n** and directed to the respective database instances of which it is a component.

[0062] In some embodiments, the distributed storage system of distributed storage service **210** described herein (e.g., distributed storage system **310** shown in FIG. 3) may organize data in various logical data volumes, extents (which may include partitions of the user data space in the volume and a segmentation of the log for the volume) made durable among a protection group of storage nodes, segments (which may be data stored on an individual storage node of a protection group) and pages for storage on one or more storage nodes. For example, in some embodiments, each database is represented by a logical volume, and each logical volume is partitioned over a collection of storage nodes into extents. A protection group may be composed of different storage nodes in the distributed storage service that together make an extent durable. Multiple segments, each of which lives on a particular one of the storage nodes in a protection group, are used to make the extent durable.

[0063] In some embodiments, each data page is stored in a segment, such that each segment stores a collection of one or more data pages and a change log (also referred to as a redo log) (e.g., a log of redo log records) for each data page that it stores. Thus, change logs may be log records segmented to the protection group of which the segment is a member. As described in detail herein, the storage nodes may be configured to receive redo log records (which may also be referred to herein as ULRs) and to coalesce them to create new versions of the corresponding data pages and/or additional or replacement log records (e.g., lazily and/or in response to a request for a data page or a database crash). If the data volume is encrypted, such a coalesce operation may need to obtain the encryption key to perform the coalesce operation. In some embodiments, data pages and/or change logs may be mirrored across multiple storage nodes, according to a variable configuration, such as in a protection group (which may be specified by the client on whose behalf the databases are being maintained in the database system). For example, in different embodiments, one, two, or three copies of the data or change logs may be stored in each of one, two, or three different availability zones or regions, according to a default configuration, an application-specific durability preference, or a client-specified durability preference.

[0064] As used herein, the following terms may be used to describe the organization of data by a distributed storage system such as distributed storage system **310**, according to various embodiments.

[0065] Volume: A volume may be a logical concept representing a highly durable unit of storage that a user/client/application of the storage system understands. More specifically, a volume is a distributed store that appears to the user/client/application as a single consistent ordered log of write operations to various user pages of a database. Each write operation may be encoded in a User Log Record (ULR), which represents a logical, ordered mutation to the contents of a single user page within the volume. As noted above, a ULR may also be referred to herein as a redo log record. Each ULR may include a unique identifier (e.g., a Logical Sequence Number (LSN)) assigned from a log sequence number space. Each ULR may be persisted to one or more synchronous segments in the log-structured distributed store that form a Protection Group (PG) maintaining the partition of user data space (i.e. extent) to which the update indicated by the log record pertains in order to provide high durability and availability for the ULR. A volume may provide an LSN-type read/write interface for a variable-size contiguous range of bytes.

[0066] In some embodiments, a volume may consist of multiple extents, each made durable through a protection group. In such embodiments, a volume may represent a unit of storage composed of a mutable contiguous sequence of Volume Extents. Reads and writes that are directed to a volume may be mapped into corresponding reads and writes to the constituent volume extents. In some embodiments, the size of a volume may be changed by adding or removing volume extents from the end of the volume.

[0067] Segment: A segment may be a limited-durability unit of storage assigned to a single storage node. Multiple segments may be implemented in a protection group to persist an extent. More specifically, a segment provides limited best-effort durability (e.g., a persistent, but non-redundant single point of failure that is a storage node) for a specific fixed-size byte range of data. This data may in some cases be a mirror of user-addressable data, or it may be other data, such as volume metadata or erasure coded bits, in various embodiments. A given segment may live on exactly one storage node. Within a storage node, multiple segments may live on each SSD, and each segment may be restricted to one SSD (e.g., a segment may not span across multiple SSDs). In some embodiments, a segment may not be required to occupy a contiguous region on an SSD; rather there may be an allocation map in each SSD describing the areas that are owned by each of the segments. As noted above, a protection group may consist of multiple segments spread across multiple storage nodes. In some embodiments, a segment may provide an LSN-type read/write interface for a fixed-size contiguous range of bytes (where the size is defined at creation). In some embodiments, each segment may be identified by a Segment UUID (e.g., a universally unique

identifier of the segment).

[0068] Storage page: A storage page may be a block of memory, generally of fixed size. In some embodiments, each page is a block of memory (e.g., of virtual memory, disk, or other physical memory) of a size defined by the operating system, and may also be referred to herein by the term “data block”. More specifically, a storage page may be a set of contiguous sectors. It may serve as the unit of allocation in SSDs, as well as the unit in log pages for which there is a header and metadata. In some embodiments, and in the context of the database systems described herein, the term “page” or “storage page” may refer to a similar block of a size defined by the database configuration, which may typically be a multiple of 2, such as 4096, 8192, 16384, or 32768 bytes.

[0069] Log page: A log page may be a type of storage page that is used to store log records (e.g., redo log records or undo log records). In some embodiments, log pages may be identical in size to storage pages. Each log page may include a header containing metadata about that log page, e.g., metadata identifying the segment to which it belongs. Note that a log page is a unit of organization and may not necessarily be the unit of data included in write operations. For example, in some embodiments, during normal forward processing, write operations may write to the tail of the log one sector at a time.

[0070] Log Records: Log records (e.g., the individual elements of a log page) may be of several different classes. For example, User Log Records (ULRs), which are created and understood by users/clients/applications of the storage system, may be used to indicate changes to user data in a volume. Log records may include metadata, such as pointers or back links, that indicate a previous LSN for log record maintained at a particular segment and/or the previous LSN in the log sequence number space. Control Log Records (CLRs), which are generated by the storage system, may also contain control information used to keep track of metadata such as the current unconditional volume durable LSN (VDL). Null Log Records (NLRs) may in some embodiments be used as padding to fill in unused space in a log sector or log page. In some embodiments, there may be various types of log records within each of these classes, and the type of a log record may correspond to a function that needs to be invoked to interpret the log record. For example, one type may represent all the data of a user page in compressed format using a specific compression format; a second type may represent new values for a byte range within a user page; a third type may represent an increment operation to a sequence of bytes interpreted as an integer; and a fourth type may represent copying one byte range to another location within the page. In some embodiments, log record types may be identified by GUIDs (rather than by integers or enums), which may simplify versioning and development, especially for ULRs.

[0071] Payload: The payload of a log record may be the data or parameter values that are specific to the log record or to log records of a particular type. For example, in some embodiments, there may be a set of parameters or attributes that most (or all) log records include, and that the storage system itself understands. These attributes may be part of a common log record header/structure, which may be relatively small compared to the sector size. In addition, most log records may include additional parameters or data specific to that log record type, and this additional information may be considered the payload of that log record. In some embodiments, if the payload for a particular ULR is larger than the user page size, it may be replaced by an absolute ULR (an AULR) whose payload includes all the data for the user page. This may enable the storage system to enforce an upper limit on the size of the payload for ULRs that is equal to the size of user pages.

[0072] Note that when storing log records in the segment log, the payload may be stored along with the log header, in some embodiments. In other embodiments, the payload may be stored in a separate location, and pointers to the location at which that payload is stored may be stored with the log header. In still other embodiments, a portion of the payload may be stored in the header, and the remainder of the payload may be stored in a separate location. If the entire payload is stored with the log header, this may be referred to as in-band storage; otherwise the storage may be referred to as being out-of-band. In some embodiments, the payloads of most large AULRs may be stored out-

of-band in the cold zone of log (which is described below).

[0073] User pages: User pages may be the byte ranges (of a fixed size) and alignments thereof for a particular volume that are visible to users/clients of the storage system. User pages are a logical concept, and the bytes in particular user pages may or not be stored in any storage page as-is. The size of the user pages for a particular volume may be independent of the storage page size for that volume. In some embodiments, the user page size may be configurable per volume, and different segments on a storage node may have different user page sizes. In some embodiments, user page sizes may be constrained to be a multiple of the sector size (e.g., 4 KB), and may have an upper limit (e.g., 64 KB). The storage page size, on the other hand, may be fixed for an entire storage node and may not change unless there is a change to the underlying hardware.

[0074] Data page: A data page may be a type of storage page that is used to store user page data in compressed form. In some embodiments every piece of data stored in a data page is associated with a log record, and each log record may include a pointer to a sector within a data page (also referred to as a data sector). In some embodiments, data pages may not include any embedded metadata other than that provided by each sector. There may be no relationship between the sectors in a data page. Instead, the organization into pages may exist only as an expression of the granularity of the allocation of data to a segment.

[0075] Storage node: A storage node may be a single virtual machine that on which storage node server code is deployed. Each storage node may contain multiple locally attached SSDs, and may provide a network API for access to one or more segments. In some embodiments, various nodes may be on an active list or on a degraded list (e.g., if they are slow to respond or are otherwise impaired, but are not completely unusable). In some embodiments, the client-side driver may assist in (or be responsible for) classifying nodes as active or degraded, for determining if and when they should be replaced, and/or for determining when and how to redistribute data among various nodes, based on observed performance. Multiple storage nodes may together implement a protection group, in some embodiments.

[0076] SSD: As referred to herein, the term “SSD” may refer to a local block storage volume as seen by the storage node, regardless of the type of storage employed by that storage volume, e.g., disk, a solid-state drive, a battery-backed RAM, a non-volatile RAM device (e.g., one or more NV-DIMMs) or another type of persistent storage device. An SSD is not necessarily mapped directly to hardware. For example, a single solid-state storage device might be broken up into multiple local volumes where each volume is split into and striped across multiple segments, and/or a single drive may be broken up into multiple volumes simply for ease of management, in different embodiments. In some embodiments, each SSD may store an allocation map at a single fixed location. This map may indicate which storage pages that are owned by particular segments, and which of these pages are log pages (as opposed to data pages). In some embodiments, storage pages may be pre-allocated to each segment so that forward processing may not need to wait for allocation. Any changes to the allocation map may need to be made durable before newly allocated storage pages are used by the segments.

[0077] FIG. 3 illustrates a distributed storage system of a transactional database system, according to some embodiments.

[0078] In some embodiments, a distributed storage system of a transactional database system may resemble distributed storage system 310 of transactional database system 300, as shown in FIG. 3. In at least some embodiments, storage nodes 330-350 may store data for different clients as part of a multi-tenant storage service. In some embodiments, transactional database system 300 may be a client of distributed storage system 310, which communicates with a database engine head node 320 over interconnect 360. As in the example illustrated in FIG. 3, database engine head node 320 may include a client-side storage service driver 325. In this example, distributed storage system 310 includes multiple storage system server nodes (including those shown as 330, 340, and 350), each of which includes storage for data pages and redo logs for the segment(s) it stores, and

hardware and/or software configured to perform various segment management functions. For example, each storage system server node may include hardware and/or software configured to perform at least a portion of any or all of the following operations: replication (locally, e.g., within the storage node), coalescing of redo logs to generate data pages, snapshots (e.g., creating, restoration, deletion, etc.), log management (e.g., manipulating log records), crash recovery (e.g., determining candidate log records for volume recovery), and/or space management (e.g., for a segment). Each storage system server node may also have multiple attached storage devices (e.g., SSDs) on which data blocks may be stored on behalf of clients (e.g., users, client applications, and/or database service subscribers). In some embodiments, storage system server nodes may provide multi-tenant storage, storing portions of different data volumes for different clients in the attached storage devices. Correspondingly, storage system server nodes may maintain numerous permanent keys for the different data volumes.

[0079] In some embodiments, as shown in FIG. 3, storage system server node **330** includes data page(s) **333**, segment redo log(s) **335**, segment management functions **337**, monitor **349** and attached SSDs **371-378**. Monitor **339** may report health statistics, performance measurements and any other information to volume manager **312** for distributed storage system management. A person having ordinary skill in the art should understand that the label “SSD” may or may not refer to a solid-state drive, but may more generally refer to a local block storage volume, regardless of its underlying hardware. Similarly, storage system server node **340** includes data page(s) **343**, segment redo log(s) **345**, segment management functions **347**, and attached SSDs **381-388**; and storage system server node **350** includes data page(s) **353**, segment redo log(s) **355**, segment management functions **357**, and attached SSDs **391-398**.

[0080] In some embodiments, each of the storage system server nodes in distributed storage system **310** may implement a set of processes running on the node server's operating system that manage communication with database engine head node **320**, e.g., to receive redo logs, send back data pages, etc. In some embodiments, all data blocks written to the distributed storage system may be backed up to long-term and/or archival storage (e.g., in a remote key-value durable backup storage system).

[0081] Distributed storage system **310** may also implement a storage control plane. Storage control plane may be one or more compute nodes configured to perform a variety of different storage system management functions. For example, storage control plane may implement a volume manager **312**, which may be configured to maintain mapping information or other metadata for a volume, such as current volume state, current writer, truncation tables or other truncation information, or any other information for a volume as it is persisted in varying different, extents, segments, and protection groups. Volume manager **312** may be configured to communicate with a client of storage system **310**, such as client-side driver **325** in order to “mount” or “open” the volume for the client, providing client-side driver **325** with mapping information, protection group policies, and various other information necessary to send write and read requests to storage nodes **330-350**. Volume manager **312** may be configured to provide the maintained information, including keys, to storage clients, such as database engine head node **320** or client-side driver **325** or to other system components such as recovery service agents implemented for volume recovery/repair service **314**. For example, volume manager **312** may provide a current volume state (e.g., clean, dirty or recovery), current epoch indicator and/or any other information about the data volume.

[0082] In some embodiments, distributed storage system may implement volume recovery/repair service **314**. Volume recovery/repair service **314** may be configured to monitor or evaluate the current data volume state and based, at least in part, on the evaluation, direct recovery service agent(s) to perform a recovery operation with respect to a data volume. For example, volume recovery/repair **314** service may determine that the state of a particular data volume has remained in recovery state in excess of a time threshold and, in response direct one or more recovery service agent(s) to perform a recovery operation for the data volume.

[0083] In some embodiments, volume recovery/repair **314** may detect recovery events for a data volume. For example, volume recovery/repair **314** may receive an indication from client-side storage service driver **325** or database engine head node **320** that the data volume may not be re-opened or used again by a writing client. Additionally, recovery event detection module may determine other recovery events that may include detecting that a lease has expired for data volume client-side storage service driver **325** or database engine head node **320**, or may receive a release request or command to evict the database engine head node **320** from another component of the database service **300**. More generally, volume recovery/repair **314** may detect any event which may trigger or require the performance of a recovery operation.

[0084] Volume recovery/repair **314** may also implement recovery service agent(s) (not illustrated) which may be individual or collections of computing systems, nodes or devices that are configured to perform recovery operations. Recovery service agent(s) may be provisioned (from other computing resources of distributed storage system **310** or another system of service provider network **100**) or directed to perform a recovery operation by volume recovery/repair **314**. Recovery service agent(s) may abort the performance of a recovery operation, if it is determined that a recovery operation for the same data volume is already complete. In order to repair data volumes, may need to enable access to encrypted data for the data stored at the different storage nodes. Thus, similar to client-side storage service driver **325**, recovery service agent(s) may receive an encrypted version of permanent key(s) to a data volume and send them to those storage nodes storing the data volume.

[0085] FIG. 4A illustrates various components of an analytical database system configured to use warm and cold storage tiers to store data blocks for clients of an analytical database service, wherein the warm storage tier comprises one or more node clusters associated with said clients, according to some embodiments.

[0086] In various embodiments, the components illustrated in at least FIGS. 4A and 4B may be implemented directly within computer hardware, as instructions directly or indirectly executable by computer hardware (e.g., a microprocessor or computer system), or using a combination of these techniques. For example, the components of shown in FIGS. 4A and 4B may be implemented by a system that includes a number of computing nodes (or simply, nodes), each of which may be similar to the computer system embodiment illustrated in FIG. 14 and described below. In various embodiments, the functionality of a given system or service component (e.g., a component of analytical database system **400**) may be implemented by a particular node or may be distributed across several nodes. In some embodiments, a given node may implement the functionality of more than one service system component (e.g., more than one data store component).

[0087] Analytical database service **150** may be various types of data processing services that perform general or specialized data processing functions (e.g., querying transactional data tables, anomaly detection, machine learning, data mining, big data querying, or any other type of data processing operation). For example, analytical database service **150** may include various types of database services (both relational and non-relational) for storing, querying, updating, and maintaining data such as transactional data tables. Such services may be enterprise-class database systems that are highly scalable and extensible. Queries may be directed to a database in analytical database service **150** that is distributed across multiple physical resources, and the analytical database system may be scaled up or down on an as needed basis.

[0088] Analytical database service **150** may work effectively with database schemas of various types and/or organizations, in different embodiments. In some embodiments, clients/subscribers may submit queries in a number of ways, e.g., interactively via an SQL interface to the database system. In other embodiments, external applications and programs may submit queries using Open Database Connectivity (ODBC) and/or Java Database Connectivity (JDBC) driver interfaces to the database system. For instance, analytical database service **150** may implement, in some embodiments, a data warehouse service, that utilizes one or more of the additional services of

service provider network **100**, to execute portions of queries or other access requests with respect to data that is stored in a remote data store, such as cold storage tier **406** (or another data store within data storage services **120**, etc.) to implement query processing for distributed data sets.

[0089] In at least some embodiments, analytical database service **150** may be a data warehouse service. Thus, in the description that follows, analytical database service **150** may be discussed according to the various features or components that may be implemented as part of a data warehouse service, including a control plane, such as control plane **402**, and processing node clusters **420**, **430**, and **440**. Note that such features or components may also be implemented in a similar fashion for other types of data processing services and thus the following examples may be applicable to other types of data processing services, such as database services. Analytical database service **150** may implement one (or more) processing clusters that are attached to a database (e.g., a data warehouse). In some embodiments, these processing clusters may be designated as a primary and secondary (or concurrent, additional, or burst processing clusters) that perform queries to an attached database warehouse.

[0090] In embodiments where analytical database service **150** is a data warehouse service, the data warehouse service may offer clients a variety of different data management services, according to their various needs. In some cases, clients may wish to store and maintain large amounts of data, such as transactional records, website analytics and metrics, sales records marketing, management reporting, business process management, budget forecasting, financial reporting, or many other types or kinds of data. A client's use for the data may also affect the configuration of the data management system used to store the data. For instance, for certain types of data analysis and other operations, such as those that aggregate large sets of data from small numbers of columns within each row, a columnar database table may provide more efficient performance. In other words, column information from database tables may be stored into data blocks on disk, rather than storing entire rows of columns in each data block (as in traditional database schemes). The following discussion describes various embodiments of a relational columnar database system implemented as a data warehouse. However, various versions of the components discussed below as may be equally adapted to implement embodiments for various other types of relational database systems, such as row-oriented database systems. Therefore, the following examples are not intended to be limiting as to various other types or formats of database systems.

[0091] In some embodiments, storing table data in such a columnar fashion may reduce the overall disk I/O requirements for various queries and may improve analytic query performance. For example, storing database table information in a columnar fashion may reduce the number of disk I/O requests performed when retrieving data into memory to perform database operations as part of processing a query (e.g., when retrieving all of the column field values for all of the rows in a table) and may reduce the amount of data that needs to be loaded from disk when processing a query. Conversely, for a given number of disk requests, more column field values for rows may be retrieved than is necessary when processing a query if each data block stored entire table rows. In some embodiments, the disk requirements may be further reduced using compression methods that are matched to the columnar storage data type. For example, since each block contains uniform data (i.e., column field values that are all of the same data type), disk storage and retrieval requirements may be further reduced by applying a compression method that is best suited to the particular column data type. In some embodiments, the savings in space for storing data blocks containing only field values of a single column on disk may translate into savings in space when retrieving and then storing that data in system memory (e.g., when analyzing or otherwise processing the retrieved data).

[0092] Analytical database system **400** may be implemented by a large collection of computing devices, such as customized or off-the-shelf computing systems, servers, or any other combination of computing systems or devices, such as the various types of systems **1400** described below with regard to FIG. **14**. Different subsets of these computing devices may be controlled by a control

plane of the analytical database system **400**. Control plane **402**, for example, may provide a cluster control interface to clients or users who wish to interact with the processing clusters, such as node cluster(s) **420**, **430**, and **440** managed by control plane **402**. For example, control plane **402** may generate one or more graphical user interfaces (GUIs) for clients, which may then be utilized to select various control functions offered by the control interface for the processing clusters **420**, **430**, and **440** hosted in the data processing service **100**. Control plane **402** may provide or implement access to various metrics collected for the performance of different features of data processing service **100**, including processing cluster performance, in some embodiments.

[0093] As discussed above, various clients (or customers, organizations, entities, or users) may wish to store and manage data using a data processing service **100**. Processing clusters **420**, **430**, and **440** may respond to various requests, including write/update/store requests (e.g., to write data into storage) or queries for data (e.g., such as a Server Query Language request (SQL) for particular data). For example, multiple users or clients may access a processing cluster to obtain data warehouse services.

[0094] Processing clusters, such as node clusters **420**, **430**, and **440**, hosted by analytical database service **150** may provide an enterprise-class database query and management system that allows users to send data processing requests to be executed by the clusters, such as by sending a query. Processing clusters **420**, **430**, and **440** may perform data processing operations with respect to data stored locally in a processing cluster, as well as remotely stored data. For example, cold storage tier **406** may comprise backups or other data of a database stored in a cluster. In some embodiments, database data may not be stored locally in a processing cluster **420**, **430**, or **440** but instead may be stored in cold storage tier **406** (e.g., with data being partially or temporarily stored in processing cluster **420**, **430**, or **440** to perform queries). Queries sent to a processing cluster **420**, **430**, or **440** (or routed/redirect/assigned/allocated to processing cluster(s)) may be directed to local data stored in the processing cluster and/or remote data. Therefore, processing clusters may implement local data processing, such as local data processing, to plan and execute the performance of queries with respect to local data in the processing cluster, as well as a remote data processing client.

[0095] Analytical database system **400** of analytical database service **150** may implement different types or configurations of processing clusters. For example, different configurations **420**, **430**, or **440**, may utilize various different configurations of computing resources, including, but not limited to, different numbers of computational nodes, different processing capabilities (e.g., processor size, power, custom or task-specific hardware, such as hardware accelerators to perform different operations, such as regular expression searching or other data processing operations), different amounts of memory, different networking capabilities, and so on. Thus, for some queries, different configurations of processing cluster **420**, **430**, **440**, etc. may offer different execution times. As shown in FIG. **4A**, node cluster **420** comprises nodes **422**, **424**, and **426**, node cluster **430** comprises nodes **432**, **434**, **436**, and **438**, and node cluster **440** comprises node **442** and **444**. Different configurations of processing clusters may be maintained in different pools of available processing clusters to be attached to a database. Attached processing clusters may then be made exclusively assigned or allocated for the use of performing queries to the attached database, in some embodiments. The number of processing clusters attached to a database may change over time according to the selection techniques discussed below.

[0096] In some embodiments, analytical database service **150** may have at least one processing cluster attached to a database, which may be the “primary cluster.” Primary clusters may be reserved, allocated, permanent, or otherwise dedicated processing resources that store and/or provide access to a database for a client, in some embodiments. Primary clusters, however, may be changed. For example, a different processing cluster may be attached to a database and then designated as the primary database (e.g., allowing an old primary cluster to still be used as a “secondary” processing cluster or released to a pool of processing clusters made available to be attached to a different database). Techniques to resize or change to a different configuration of a

primary cluster may be performed, in some embodiments. The available processing clusters that may also be attached, as determined, to a database may be maintained (as noted earlier) in different configuration type pools, which may be a set of warmed, pre-configured, initialized, or otherwise prepared clusters which may be on standby to provide additional query performance capacity in addition to that provided by a primary cluster. Control plane **402** may manage cluster pools by managing the size of cluster pools (e.g., by adding or removing processing clusters based on demand to use the different processing clusters).

[0097] As databases are created, updated, and/or otherwise modified, snapshots, copies, or other replicas of the database at different states may be stored in cold storage tier **406**, according to some embodiments. For example, a leader node, or other processing cluster component, may implement a backup agent or system that creates and store database backups for a database to be stored as database data in cold storage tier **406** and/or data storage service **120**. Database data may include user data (e.g., tables, rows, column values, etc.) and database metadata (e.g., information describing the tables which may be used to perform queries to a database, such as schema information, data distribution, range values or other content descriptors for filtering out portions of a table from a query, a superblock, etc.). A timestamp or other sequence value indicating the version of database data may be maintained in some embodiments, so that the latest database data may, for instance, be obtained by a processing cluster in order to perform queries. In at least some embodiments, database data (e.g., cold storage tier **406** data) may be treated as the authoritative version of data, and data stored in processing clusters **420**, **430**, and **440** for local processing (e.g., warm storage tier **404**) as a cached version of data.

[0098] Cold storage tier **406** may implement different types of data stores for storing, accessing, and managing data on behalf of clients **410**, **412**, **414**, etc. as a network-based service that enables clients **410**, **412**, **414**, etc. to operate a data storage system in a cloud or network computing environment. Cold storage tier **406** may also include various kinds of object or file data stores for putting, updating, and getting data objects or files. For example, one cold storage tier **406** may be an object-based data store that allows for different data objects of different formats or types of data, such as structured data (e.g., database data stored in different database schemas), unstructured data (e.g., different types of documents or media content), or semi-structured data (e.g., different log files, human-readable data in different formats like JavaScript Object Notation (JSON) or Extensible Markup Language (XML)) to be stored and managed according to a key value or other unique identifier that identifies the object. In at least some embodiments, cold storage tier **406** may be treated as a data lake. For example, an organization may generate many different kinds of data, stored in one or multiple collections of data objects in a cold storage tier **406**. The data objects in the collection may include related or homogenous data objects, such as database partitions of sales data, as well as unrelated or heterogeneous data objects, such as audio files and web site log files. Cold storage tier **406** may be accessed via programmatic interfaces (e.g., APIs) or graphical user interfaces. For example, format independent analytical database service **400** may access data objects stored in data storage services via the programmatic interfaces.

[0099] As described above with regard to clients **170-190** and **250**, clients **410**, **412**, **414**, etc. may encompass any type of client that can submit network-based requests to service provider network **200** via network **408** (e.g., also network **160** and **260**), including requests for storage services (e.g., a request to query data analytical service **150**, or a request to create, read, write, obtain, or modify data in cold storage tier **406** and/or data storage service **120**, etc.).

[0100] FIG. **4B** illustrates an example of a node cluster of an analytical database system performing queries against transactional database data, according to some embodiments. As illustrated in this example, a processing node cluster **430** may include a leader node **432** and compute nodes **434**, **436**, **438**, etc., which may communicate with each other over an interconnect (not illustrated). Leader node **432** may implement query planning **452** to generate query plan(s), query execution **454** for executing queries on processing node cluster **430** that perform data

processing that can utilize remote query processing resources for remotely stored data (e.g., by utilizing one or more query execution slot(s)/queue(s) **458**). As described herein, each node in a primary processing cluster **430** may include attached storage, such as attached storage **468a**, **468b**, and **468n**, on which a database (or portions thereof) may be stored on behalf of clients (e.g., users, client applications, and/or storage service subscribers).

[0101] Note that in at least some embodiments, query processing capability may be separated from compute nodes, and thus in some embodiments, additional components may be implemented for processing queries. Additionally, it may be that in some embodiments, no one node in processing cluster **430** is a leader node as illustrated in FIG. **4B**, but rather different nodes of the nodes in processing cluster **430** may act as a leader node or otherwise direct processing of queries to data stored in processing cluster **430**. While nodes of processing cluster may be implemented on separate systems or devices, in at least some embodiments, some or all of processing cluster may be implemented as separate virtual nodes or instance on the same underlying hardware system (e.g., on a same server).

[0102] Leader node **432** may manage communications with clients, such as clients **410**, **412**, and **414** discussed above with regard to FIG. **4A**. Leader node **432** may receive query **450** and return query results **476** to clients **410**, **412**, **414**, etc. or to a proxy service (instead of communicating directly with a client application).

[0103] Leader node **432** may be a node that receives a query **450** from various client programs (e.g., applications) and/or subscribers (users) (either directly or routed to leader node **432** from a proxy service), then parses them and develops an execution plan (e.g., query plan(s)) to carry out the associated database operation(s)). More specifically, leader node **432** may develop the series of steps necessary to obtain results for the query. Query **450** may be directed to data that is stored both locally within a warm tier implementing using local storage of processing cluster **430** (e.g., at one or more of compute nodes **434**, **436**, or **438**) and data stored remotely, such as in cold storage tier **406** (which may be implemented as part of data storage service **120**, according to some embodiments). Leader node **432** may also manage the communications among compute nodes **434**, **436**, and **438** instructed to carry out database operations for data stored in the processing cluster **430**. For example, node-specific query instructions **460** may be generated or compiled code by query execution **454** that is distributed by leader node **432** to various ones of the compute nodes **434**, **436**, and **438** to carry out the steps needed to perform query **450**, including executing the code to generate intermediate results of query **450** at individual compute nodes may be sent back to the leader node **432**. Leader node **432** may receive data and query responses or results from compute nodes **434**, **436**, and **438** in order to determine a final result **476** for query **450**.

[0104] A database schema, data format and/or other metadata information for the data stored among the compute nodes, such as the data tables stored in the cluster, may be managed and stored by leader node **432**. Query planning **452** may account for remotely stored data by generating node-specific query instructions that include remote operations to be directed by individual compute node(s). Although not illustrated, in some embodiments, a leader node may implement burst manager to send a query plan generated by query planning **452** to be performed at another attached processing cluster and return results received from the burst processing cluster to a client as part of results **476**.

[0105] In at least some embodiments, a result cache **456** may be implemented as part of leader node **432**. For example, as query results are generated, the results may also be stored in result cache **456** (or pointers to storage locations that store the results either in primary processing cluster **430** or in external storage locations), in some embodiments. Result cache **456** may be used instead of other processing cluster capacity, in some embodiments, by recognizing queries which would otherwise be sent to another attached processing cluster to be performed that have results stored in result cache **456**. Various caching strategies (e.g., LRU, FIFO, etc.) for result cache **456** may be implemented, in some embodiments. Although not illustrated in FIG. **4B**, result cache **456** could be

stored in other storage systems (e.g., other storage services, such as a NoSQL database, and/or data storage service **120**) and/or could store sub-query results.

[0106] Processing node cluster **430** may also include compute nodes, such as compute nodes **434**, **436**, and **438**. Compute nodes, may for example, be implemented on servers or other computing devices, such as those described below with regard to computer system **1400** in FIG. **14**, and each may include individual query processing “slices” defined, for example, for each core of a server's multi-core processor, one or more query processing engine(s), such as query engine(s) **462a**, **462b**, and **464n**, to execute the instructions **460** or otherwise perform the portions of the query plan assigned to the compute node. Query engine(s) **462** may access a certain memory and disk space in order to process a portion of the workload for a query (or other database operation) that is sent to one or more of the compute nodes **434**, **436**, or **438**. Query engine **462** may access attached storage, such as **468a**, **468b**, and **468n**, to perform local operation(s), such as local operations **466a**, **466b**, and **466n**. For example, query engine **462** may scan data in attached storage **468**, access indexes, perform joins, semi joins, aggregations, or any other processing operation assigned to the compute node **434**, **436**, or **438**.

[0107] Query engine **462a** may also direct the execution of remote data processing operations, by providing remote operation(s), such as remote operations **464a**, **464b**, and **464n**, to remote data processing clients, such as remote data processing **470a**, **470b**, and **470n**. Remote data processing **470** may be implemented by a client library, plugin, driver or other component that sends request sub-queries to be performed by cold storage tier **406** or requests to for data, **472a**, **472b**, and **472n**. As noted above, in some embodiments, Remote data processing **470** may read, process, or otherwise obtain data **474a**, **474b**, and **474n**, in response from cold storage tier **406**, which may further process, combine, and or include them with results of location operations **466**.

[0108] Compute nodes **434**, **436**, and **438** may send intermediate results from queries back to leader node **432** for final result generation (e.g., combining, aggregating, modifying, joining, etc.). Remote data processing clients **470** may retry data requests **472** that do not return within a retry threshold.

[0109] Attached storage **468** may be implemented as one or more of any type of storage devices and/or storage system suitable for storing data accessible to the compute nodes, including, but not limited to: redundant array of inexpensive disks (RAID) devices, disk drives (e.g., hard disk drives or solid state drives) or arrays of disk drives such as Just a Bunch Of Disks (JBOD), (used to refer to disks that are not implemented according to RAID), optical storage devices, tape drives, RAM disks, Storage Area Network (SAN), Network Access Storage (NAS), or combinations thereof. In various embodiments, disks may be formatted to store database tables (e.g., in column-oriented data formats or other data formats).

[0110] Although FIGS. **4A** and **4B** have been described and illustrated in the context of a service provider network implementing an analytical database service, like a data warehousing service, the various components illustrated and described in FIGS. **4A** and **4B** may be easily applied to other database services that can utilize the methods and systems described herein. As such, FIGS. **4A** and **4B** are not intended to be limiting as to other embodiments maintaining and querying representations of transactional data tables for managed databases.

[0111] FIG. **5** is a flow diagram illustrating a process of maintaining, within an analytical database system, a representation of portions of a transactional data table from a transactional database system, according to some embodiments.

[0112] In some embodiments, maintaining representations of transactional tables at an analytical database, such as for the embodiments described herein, may include the following procedure steps. In the following embodiments shown in FIG. **5**, it may be assumed that services of a provider network, such as transactional database service **110** and analytical database service **150** of service provider network **100**, and the functionalities and techniques described for said services herein, may be used to implement a hybrid transactional and analytical processing service. However, a

person having ordinary skill in the art should understand that other implementations and/or embodiments that fulfill the following procedure steps may also be incorporated to the description herein.

[0113] In block **500**, portion(s) of a table that are being maintained at a transactional database, such as transactional database system **200** of transactional database service **110**, may be replicated to an analytical database, such as analytical database system **400** of analytical database service **150**, and subsequently maintained at the analytical database. In some embodiments, the means for maintaining a representation (e.g., a replica of portion(s) of a table from the transactional database) at the analytical database may use the procedure described in blocks **502-510**.

[0114] In block **502**, transactional changes that are made to a transactional table that is stored and maintained in the transactional database are written to a change-data-capture log, such as transaction log **240** (see also the description for at least change-data-capture logs **610** described herein with regard to FIG. **6**). In block **504**, portion(s) of the table that have been chosen to be replicated into the analytical database are partitioned into segments such that the portion(s) may be provided to the analytical database. In some embodiments, such segments may be referred to as snapshots, as they refer to the state of the table at a given moment (e.g., at a certain transaction number in embodiments in which the table contains transactional data). The snapshots may be provided to the analytical database via a transport mechanism. In some embodiments, the transport mechanism may resemble a data storage service, such as data storage service **120**, or a data streaming service, such as data streaming service **130**, of service provider network **100** (see also the description of transport mechanism **826** herein). A person having ordinary skill in the art should understand that “snapshots” may be plural or singular depending upon given embodiments. For example, if only one portion of one transactional table is being replicated to the analytical database and may be provided as a unit (e.g., without being further partitioned) via the transport mechanism, “snapshot” may refer to the sum of the segments, according to some embodiments. In a second example, if a given portion of a given transactional table is partitioned into more than one segment, “snapshots” may refer to the segments that sum to the portion of the table being provided via the transport mechanism. Additional example embodiments may be given and the above examples should not be misconstrued as restrictive.

[0115] In block **506**, checkpoints are also provided to the analytical database. In some embodiments, checkpoints may resemble portions of transactional changes listed in the change-data-capture log of the transactional database for the given table portion(s) being replicated. In some embodiments in which more than one snapshot has been stored to respective compute nodes of a node cluster in the analytical database, respective checkpoints may also be partitioned based on this same mapping. In some embodiments, checkpoints may be provided to the analytical database by the same or different transport mechanism as the snapshots. For example, the snapshots may be provided via a data storage service, and the subsequent checkpoints may be streamed to the analytical database via a data streaming service.

[0116] In block **508**, the snapshots and checkpoints are stored in the analytical database. In some embodiments, the snapshots and their related checkpoints may be stored across multiple compute nodes of a node cluster of the analytical database (e.g., compute nodes **434-438** of node cluster **430**). The stored snapshots at the analytical database may now be referred to as the representation of the transactional portion(s) of the table maintained at the transactional database (see the description of transactional table representations **820** described herein).

[0117] In block **510**, transactional changes that have been provided in the checkpoints are applied and committed to the representation (see also the description herein with regard to FIGS. **11A-13D**), such that the representation is updated and maintained as a replica of the table stored in the transactional database. The process of receiving, applying, and committing additional checkpoints may continue as long as the hybrid transactional and analytical processing service maintains the representation in the analytical database. In addition, at any point after the storage of the first set of

snapshots to the analytical database, a client of the hybrid transactional and analytical processing service may run a query against the transactional data in the representation, as the analytical database is configured to have simultaneous read/write properties (e.g., responding to the query and writing, applying, and/or committing new transactional changes to the representation).

[0118] FIG. 6 illustrates the process of a handshake protocol, used to negotiate and define the configurations and parameters for maintaining, at an analytical database, a representation of a table stored in a transactional database, according to some embodiments.

[0119] In some embodiments, a handshake protocol between the computing devices of the transactional database and the compute nodes of the analytical database may be used to determine the logistics of how a representation of a transactional table of the transactional database is going to be maintained at the analytical database. By determining such parameters and defining the procedures for providing and mapping the snapshots and checkpoints to compute nodes of a node cluster in the analytical database in advance of providing the initial snapshot(s), the transactional database and the analytical database may remain loosely coupled during the maintenance of the representation at the analytical database.

[0120] In some embodiments, transactional database **600** may resemble transactional database **200** and transactional database **300** of transactional database service **110**, and their functionalities described herein. Computing devices **602** may resemble respective database engine head nodes **220**. Interface **604** (e.g., SQL interface to the database system) may be used as a submission platform for database clients **250** providing incoming transactions to transactional database **600**. Storage **606** may resemble distributed storage system **310** in which transactional tables **608** and corresponding change-data-capture logs **610** for transactional tables **608** are stored.

[0121] Analytical database **612** may resemble analytical database system **400** of analytical database service **150**, according to some embodiments. Compute nodes **614** may represent compute nodes of a given node cluster, such as compute nodes **434-438** of node cluster **430**. Interface **616** (e.g., SQL interface to the database system) may be used as a client endpoint for client(s) **170** and **190**, wherein said clients may submit queries such as query **450**, according to some embodiments. Storage **618** may resemble attached storage **468** of compute nodes **434-438**, and/or remote storage such as cold storage tier **406**. Storage **618** may be configured such that it may store one or more of snapshots of transactional tables **608** in order to maintain respective representation(s) at the analytical database.

[0122] In some embodiments, computing devices **602** and/or compute nodes **614** may initiate a handshake procedure in preparation for maintaining one or more of transactional tables **608** at analytical database **612**. Maintaining the representations of transactional tables **608** may follow the methods described in at least blocks **500** and **1000**, according to some embodiments. In order to efficiently and effectively maintain the representations of transactional tables **608** at analytical database **612**, the handshake procedure may include negotiations between computing devices **602** and compute nodes **614** in order to determine data-type mappings, topology requirements, compatible/incompatible data definition language commands that may be written to change-data-capture logs **610** and/or interpreted by compute nodes **614**. Negotiations **622-654** may represent examples of the information that may be exchanged and/or determined via computing devices **602** and compute nodes **614**, according to some embodiments. A person having ordinary skill in the art should understand that handshake protocol **620** is meant to be a visual representation of negotiations between computing devices **602** at transactional database **600** and compute nodes **614** at analytical database **612**. Other negotiations of handshake protocol **620** besides negotiations **622-654** may additionally be included in performing handshake protocol **620**, and negotiations **622-654** are meant to be example embodiments of the methods and techniques described herein pertaining to performing a handshake protocol (see also the description of FIG. 7 herein). In addition, handshake protocol **620** may occur at computing devices **602**, compute nodes **614**, or at both computing devices **602** and compute nodes **614** through the interactions described in the following

paragraphs. Furthermore, handshake protocol **620** may involve a first stage in which computing devices **602** provides information from all or parts of negotiations **622-640** to compute nodes **614**, and then a second stage in which compute nodes **614** may respond with all or parts of negotiations **642-654**, or vice versa. In other embodiments, handshake protocol **620** may resemble a more iterative process. For example, computing devices **602** may provide list of utilized data definition language commands **624**, and compute nodes **614** may respond with list of known data definition language commands **642**, and another iteration pertaining to data definition language commands may occur in order to determine and/or confirm the results of the handshake protocol pertaining to data definition language commands. Then, a similar process may occur for negotiations **626** and **644**, etc., until the handshake protocol is complete.

[0123] In some embodiments, computing devices **602** may provide a list of portion(s) of table(s) to replicate **622**, wherein the portions are portions of transactional tables **608** to be stored and maintained by storage **618** and compute nodes **614** at analytical database **612**. By consequence of determining the portion(s) of transactional tables **608** to be maintained at analytical database **612**, list of portion(s) of table(s) to replicate **622** may also be used to determine a list of the corresponding change-data-capture logs of change-data-capture logs **610** that will be sent as checkpoints in order to maintain the transactional table representations at analytical database **612**, based on the information in negotiation **622**. Furthermore, computing devices **602** may provide information about the primary keys that correspond to the list of portion(s) of table(s) to replicate **622** in primary key(s) information **628**, according to some embodiments. In some embodiments, primary keys may correspond to unique row identifiers of respective transactional tables **608**, such that respective rows may be identified by compute nodes **614** when applying transactional changes to representations of transactional tables **608**. For example, a given transactional table of transactional tables **608** may contain an additional column of the table with respective row identifiers (e.g., row 1, row 2, row 3, etc. for each row in the given transactional table) that may be used as primary keys. In a second example, a concatenation of some subset of the columns for each row may be used as primary keys (e.g., a concatenation of the data items in column 1, column 2, and column 3 of the table). In a third example, primary keys of the given transactional table may be a concatenation of all columns in each row (e.g., a “hash” of all data items in each row). In some embodiments, computing devices **602** may use negotiation **628** to inform compute nodes **614** that there is no current primary keys scheme for the list of portion(s) of table(s) to replicate **622**. In such embodiments, computing devices **602** and compute nodes **614** may determine, during handshake protocol **620**, to use a concatenation of all columns in each row (e.g., the “hash” example described above) as the method of communicating information (e.g., transactional changes) about rows of the given transactional tables. In some embodiments, negotiation **628** may be referred to as determining a logic for generating respective primary keys, either via a provided primary key scheme or by determining to use a concatenation of all columns in each row, etc.

[0124] The results of the negotiation pertaining to primary key(s) information **628** may be used during the maintenance of the representations at analytical database **612**, as, when providing checkpoints to compute nodes **614**, the primary keys scheme may be trusted as an agreed upon form of communication when referencing respective rows to which compute nodes **614** should apply transactional changes of the checkpoints, according to some embodiments.

[0125] Computing devices **602** may also provide a list of data definition language commands **624** that are used when writing transactional changes to change-data-capture logs **610**, and compute nodes **614** may provide a list of known data definition language commands **642**. Negotiations **624** and **642** may be used to generate a list of compatible and/or incompatible data definition language commands, according to some embodiments. Such a list of compatible/incompatible data definition language commands may be used with regard to providing/receiving checkpoints of portions of change-data-capture logs **610** during the maintenance of representations of transactional tables **608** at analytical database **612**. As the compatible/incompatible data definition language commands

may be written to handshake results **660** at the end of handshake protocol **620**, computing devices **602** may preemptively trigger a new snapshot of a given table (e.g., provide snapshots **828** discussed with regard to FIG. **8** below) if computing devices **602** determine that an incompatible data definition language command is included in a given checkpoint (or determine that there is a data definition language command in a given checkpoint that is not part of the list of compatible data definition language commands) during maintenance of the representations of transactional tables **608**. Alternatively, compute nodes **614** may reactively request a new snapshot (e.g., request new snapshots **832** discussed with regard to FIG. **8** below) if compute nodes **614** determine that an incompatible data definition language command has been received as part of a given checkpoint. [0126] In some embodiments, computing devices **602** may also provide transactional database sharding policy **626**, in which computing devices **602** may propose a method of how to proportion the list of portion(s) of table(s) to replicate **622** for storage in storage **618** of analytical database **612**. A person having ordinary skill in the art should understand that the storage capacity and/or the way that the storage capacity is distributed across SSDs **371-398** of distributed storage system **310** at transactional database **600** may differ from the storage capacity and/or the way that the storage capacity is distributed across attached storage **468** and cold storage tier **406** at analytical database **612**, and therefore a negotiation pertaining to a mapping of the storage of transactional tables **608** at transactional database **600** to the storage of the representations of transactional tables **608** at analytical database **612** may be included in handshake protocol **620**. As part of said negotiation, compute nodes **614** may additionally, or alternatively, propose analytical database slicing policy **644**, pertaining to the storage capacity and/or the way that the storage capacity is distributed across attached storage **468** and cold storage tier **406** at analytical database **612**.

[0127] Computing devices **602** may additionally provide a proposed snapshots procedure **630**, which may also be based on other negotiations **622-654**, according to some embodiments. For example, a mapping procedure determined via transactional database sharding policy **626** and analytical database slicing policy **664** may further determine the way that snapshots of portion(s) of table(s) to replicate **622** are proportioned (e.g., in preparation for storing and maintaining the portion(s) of the table(s) at multiple compute nodes of compute nodes **614**). In a second example, proposed snapshots procedure **630** may be based on proposed transport mechanisms **638** and proposed transport mechanisms **652** (see continued description in the following paragraphs), in which partitioning and/or size constraints of the determined transport mechanisms may determine proposed snapshots procedure **630**, according to some embodiments. Furthermore, compute nodes **614** may also or alternatively propose snapshots procedure **646** as part of the negotiations of handshake protocol **620**. For example, depending upon how the representations of portion(s) of table(s) to replicate **622** are going to be distributed across compute nodes **614** of a given node cluster, compute nodes **614** may propose an optimized method of receiving snapshots from transactional database **600**. Compute nodes **614** may also provide information pertaining to the structure of analytical database **612** via mapping of compute node structure **648** (e.g., compute nodes **614** may propose one or more node clusters that could be used to store and maintain list of portion(s) of table(s) to replicate **622**). For example, mapping of compute node structure **648** may provide information about the number of compute nodes in a given node cluster, information about the storage capacity of said compute nodes, etc. Such information about the mapping of compute node structure **648** may be used to determine methods for providing snapshots and checkpoints to analytical database **612**, according to some embodiments.

[0128] In some embodiments, computing devices **602** may propose checkpoints procedure **632** based at least in part on negotiations **630** and **646**. Once a mapping of providing snapshots to compute nodes **614** of the given node cluster at analytical database **612** has been determined, computing devices **602** may propose a corresponding mapping for checkpoints. Such mapping for checkpoints may include methods for partitioning change-data-capture logs **610** into checkpoints that may be provided using the determined transport mechanisms (see description for proposed

transport mechanisms **638** and proposed transport mechanisms **652** herein).

[0129] Handshake protocol **620** may also include negotiations pertaining to the client whose transactional tables of transactional tables **608** are going to be maintained at analytical database **612** via the methods and techniques described herein. For example, if list of portion(s) of table(s) to replicate **622** pertain to client **170** of transactional database service **110**, client information **634** (e.g., a username, list of database engine head nodes **220** that pertain to said client, and/or account information with service provider network **100**) for client **170** may be written to handshake results **660**. Continuing with this example, client **170** may also be a client of analytical database service **150**, in which case compute nodes **614** may provide list of client's node cluster(s) **650**, according to some embodiments. Furthermore, in addition to (or in response to) determining that client **170** is a client of transactional database service **110** and of analytical database service **150**, computing devices **602** may request a certain node cluster **636** from the list of client's node cluster(s) **650**. A person having ordinary skill in the art should understand that negotiations **634**, **636**, and **650** may take place iteratively, in combination with one another, or separately during the overarching process of performing handshake protocol **620**.

[0130] In some embodiments, one or more transport mechanisms may be used to provide snapshots and checkpoints to analytical database **612** (see the description for blocks **504** and **506** herein). As part of handshake protocol **620**, computing devices **602** may propose transport mechanisms **638** and/or compute nodes **614** may propose transport mechanisms **652**. For example, computing devices **602** may propose to provide snapshots (e.g., provide snapshots **828**) via a given service of service provider network **100**, such as data storage service **120**. Compute nodes **614** may similarly propose to receive snapshots (e.g., receive snapshots **830**) via a given service of service provider network **100** (e.g., a same or different transport mechanism than those proposed by proposed transport mechanisms **638**) in proposed transport mechanisms **652**. During negotiations **638** and **652**, an agreed upon transport mechanism for providing snapshots to analytical database **612** may be determined and written to handshake results **660**. In some embodiments, negotiations **638** and **652** may also be used to determine an agreed upon transport mechanism for providing checkpoints to analytical database **612** in which said transport mechanism may be the same or different transport mechanism determined for providing snapshots. For example, computing devices **602** and compute nodes **614** may determine, via negotiations **638** and **652**, that snapshots may be provided to analytical database **612** via a first transport mechanism (e.g., data storage service **120**), and that checkpoints may be provided to analytical database **612** via a second transport mechanism (e.g., data streaming service **130**).

[0131] A person having ordinary skill in the art should understand that additional information (e.g., other information **640** and other information **654**) may additionally be used to determine results of handshake protocol **620**, and that negotiations **622-654** are meant to be example embodiments rather than an exhaustive list of negotiations that may take place during the performance of handshake protocol **620**.

[0132] In some embodiments, after performing handshake protocol **620**, results of the determined parameters for maintaining representations of transactional tables of transactional database **600** at analytical database **612** may be written, via write results of handshake **656**, to a data store, such as data store **658**, that is made accessible to transactional database **600** and analytical database **612**. As discussed above, computing devices **602** may write a portion of handshake results **660** and compute nodes **614** may write an additional portion of handshake results via write results of handshake **656**, resulting in handshake results **660**, or, alternatively, either computing devices **602** or compute nodes **614** may write the results of handshake protocol **620**, resulting in handshake results **660**. Data store **658** may be located at a storage that is accessible to transactional database **600** and analytical database **612**. For example, data store **658** may be located in a given storage of data storage service **120**, according to some embodiments. In a second example, data store **658** may be located in storage of transactional database service **110**, and data store **658** may be made accessible

to analytical database service **150** such that read access is given to both transactional database **600** and analytical database **612**, according to some embodiments.

[0133] As shown in FIG. **6**, transactional database **600** may have read access to data store **662** and analytical database **612** may have read access to data store **664**. Read access to data store **662** and **664** may allow computing devices **602** and compute nodes **614** to refer to handshake results **660** during the process of maintaining representations of transactional tables **608** at analytical database **612**, according to some embodiments. For example, computing devices **602** may use handshake results **660** to verify that transactional changes in a given checkpoint that it will provide to analytical database **612** (e.g., via provide checkpoints **834**) do not contain incompatible data definition language commands listed in handshake results **660**. In a second example, compute nodes **614** may use handshake results **660** to verify the determined transport mechanism by which analytical database **612** may expect to receive snapshots by (e.g., via receive snapshots **830**), according to some embodiments. Such example embodiments describe the “loose coupling” of transactional database **600** with analytical database **612** after the completion of handshake protocol **620**. By establishing standard processes and procedures for maintaining representations of transactional tables at analytical database **612** during handshake protocol **620**, subsequent processes (e.g., provide snapshots **828**, provide checkpoints **834**, request new snapshots **832**, request new checkpoints **838**, etc.) may be automated.

[0134] Furthermore, determined results of handshake results **660** may not be edited and/or written to by computing devices **602** or compute nodes **614** via read access to data store **662** or **664**, according to some embodiments. In some embodiments, updates or changes to the structure and/or configurations of transactional database **600**, analytical database **612**, and/or any other services of service provider network **100** that are utilized as transport mechanisms (e.g., data storage service **120**, data streaming service **130**, etc.) may cause handshake results **660** to become out-of-date. In some embodiments in which it is determined that one or more of the determined results of handshake results **660** should be updated or changed, either computing devices **602**, compute nodes **614**, or both computing devices **602** and compute nodes **614** may re-initiate a new handshake protocol **620**. One or more of the determined results may then be updated, modified, or changed based on performing the new handshake protocol **620**, and handshake results **660** may be overwritten by updated determined results of the new handshake protocol **620**, according to some embodiments.

[0135] FIG. **7** is a flow diagram illustrating a process of initiating and performing a handshake protocol, used to negotiate and define the configurations and parameters for maintaining, at an analytical database, a representation of a table stored in a transactional database, according to some embodiments.

[0136] In some embodiments, the methods and techniques for performing handshake protocol **620** may resemble the embodiments shown in FIG. **7** via blocks **700-712**. In block **700**, a handshake protocol may be initiated in order to determine parameters for maintaining representations of transactional tables of the transactional database at the analytical database, according to some embodiments. As described above with regard to FIG. **6**, the handshake protocol may be initiated by either the transactional database side of the hybrid transactional and analytical processing service, the analytical database side, or both. In block **702**, the handshake protocol is performed. Blocks **704-710** may represent embodiments of negotiations that may take place during performance of the handshake protocol, according to some embodiments. Blocks **704-710** are not meant to be an exhaustive list of negotiations, and additional negotiations not shown in FIG. **7** may take place during performance of the handshake protocol (e.g., block **702**).

[0137] In block **704**, a mapping for distributing checkpoints across compute nodes of a given node cluster of the analytical database may be determined. In some embodiments, block **704** may resemble at least negotiations **632** and **648** and their descriptions herein. In block **706**, one or more transport mechanisms for providing snapshots and checkpoints to the analytical database may be

determined. In some embodiments, block **706** may resemble at least negotiations **638** and **652** and their descriptions herein. In block **708**, a list of data definition language commands may be agreed upon by the computing devices of the transactional database and the compute nodes of the analytical database during performance of the handshake protocol. In some embodiments, block **708** may resemble at least negotiations **624** and **642**. In block **710**, additional information may be determined during the performance of the handshake protocol such that additional parameters and/or functionalities for maintaining representations of transactional tables at the analytical database may be defined. In some embodiments, block **710** may pertain to any additional negotiations of negotiations **622-654** that have not already been determined. Block **710** may additionally refer to any negotiations that will promote a “loose coupling” of the transactional database and the analytical database after the completion of the handshake protocol and autonomous/automatic functionalities pertaining to maintaining representations of transactional tables at the analytical database.

[0138] In block **712**, the determined results of at least blocks **704-710** may be stored in a data store that is made accessible to the transactional database and the analytical database. In some embodiments, the data store of block **712** may resemble data store **658**, which stores handshake results **660**.

[0139] FIG. **8** illustrates the use of a transport mechanism to provide snapshots and checkpoints from a transactional database to an analytical database, according to some embodiments.

[0140] In some embodiments, transactional database **800**, computing devices **802**, interface **804**, storage **806**, transactional tables **808**, and change-data-capture logs **810** may resemble transactional database **600**, computing devices **602**, and interface **604** and their functionalities described herein. Similarly, analytical database **812**, compute nodes **814**, and interface **816** may resemble analytical database **612**, compute nodes **614**, and interface **616** and their functionalities described herein. In some embodiments, FIG. **8** may describe embodiments in which one or more snapshots corresponding to transactional tables **808** have been provided to analytical database **812**, and maintenance of the respective representations at analytical database **812** has already begun (e.g., via the methods and techniques described by at least block **500**). As shown in FIG. **8**, portions of transactional tables **808** may be provided to analytical database **812** via provide snapshots **828**, transport mechanism **826**, and receive snapshots **830**. The portions of transactional tables **808** may then be stored in transactional table representations **820** within storage **818**, according to some embodiments.

[0141] In order to maintain transactional table representations **820**, checkpoints containing transactional changes of change-data-capture logs **810** may be provided to analytical database **812** via provide checkpoints **834**, transport mechanism **826**, and receive checkpoints **836**, according to some embodiments. The checkpoints may then be stored in checkpoints **822** within storage **818**, according to some embodiments. In addition, storage **818** may be used to maintain shadow tables **824**, in which delete events of the transactional changes of checkpoints **822** may be buffered and committed to before being later committed to transactional table representations **820** (see description of at least FIGS. **10-13D** herein).

[0142] In some embodiments, snapshots and checkpoints may be provided to analytical database **812** via one or more transport mechanisms **826** that have been determined during performance of handshake protocol **620**. As shown in FIG. **8**, transactional database **800** and analytical database **812** may have read access to data store **846** (e.g., via read access **842** and **844**, respectively) in order to read handshake results **848**. In a first example in which transport mechanism **826** is implemented using data storage service **120**, snapshots and/or checkpoints may be provided to transport mechanism **826** via provide snapshots **828** and provide checkpoints **834** and stored in data stores of data storage service **120**. In some embodiments, analytical database **812** may be provided access to said data stores of data storage service **120** such that the snapshots and checkpoints may be received to analytical database **812** via receive snapshots **830** and receive checkpoints **836**. In a

second example in which transport mechanism **826** is implemented using data streaming service **130**, snapshots and/or checkpoints may be streamed to analytical database **812**.

[0143] In some embodiments, transport mechanism **826** may resemble a “push” or “pull” mechanism. For example, transport mechanism **826** may be implemented by transactional database service **110**, and computing devices **802** may be configured to “push” snapshots and checkpoints to analytical database **812** via a direct connection. Alternatively, transport mechanism **826** may be implemented by analytical database service **150**, and compute nodes **814** may be configured to “pull” snapshots and checkpoints to analytical database **812** from transactional database **800** via a direction connection. In some embodiments, transport mechanism **826** may also be configured to have read access to data store **840**, such that client account information, mapping, topology, and/or other relevant information pertaining to the transport of snapshots and checkpoints may be viewed by transport mechanism **826**.

[0144] FIG. **9A** illustrates a process of treating, at a transactional database, transactional changes to a transactional table and providing said transactional changes in a checkpoint to a transport mechanism for replication at a representation of the transactional table at an analytical database, according to some embodiments.

[0145] In some embodiments, transactional database **900** may be configured to maintain transactional tables, such as Table A, for a client of transactional database service **110**.

Transactional database **900** may resemble the configurations and have the functionalities described in at least FIGS. **2** and **3** for transactional database **200** and transactional database **300**, according to some embodiments. In addition, computing device(s) **904** may resemble computing devices **602** and/or computing devices **802**, and their functionalities described for said computing devices herein. A person having ordinary skill in the art should understand that FIG. **9A** is meant for illustrative purposes, and that transactional database **900** may store additional transactional tables in addition to table A. Transactional database **900** may also include additional computing device(s) besides computing device(s) **904**, and/or computing device(s) **904** may be configured to maintain additional transactional tables in addition to table A.

[0146] In some embodiments, transactional database **900** may receive new transactions/transactional changes for table A **902**. New transactions may include one or more additional rows of data items that may be inserted to Table A, and transactional changes may include one or more changes to Table A (e.g., the deletion of a row of data items, an update to a row of data items, etc.). Computing device(s) **904** may then perform transactional changes **906** to Table A **908**, wherein portions of Table A may be distributed across multiple storage system server nodes of distributed storage system **310**, according to some embodiments. In addition, for each transactional change **906** that is applied to Table A **908**, the transactional changes are written to change-data-capture log **912** via write transactional changes to change-data-capture log **910**. At a given point during the maintenance of Table A, a checkpoint containing a portion of change-data-capture log **912** may be provided, via provide checkpoint pertaining to Table A **914**, to transport mechanism **916**. In some embodiments, the checkpoint is provided to transport mechanism **916** via the methods and techniques described herein with regard to at least FIG. **8**. In addition, the frequency of checkpoints that are provided, the transport mechanism that is used to provide the checkpoints, and additional configurations of the checkpoints may be determined via the methods and techniques described herein with regard to at least FIGS. **6** and **7**, and to the results of the handshake protocol.

[0147] FIG. **9B** illustrates a process of receiving, at an analytical database, a checkpoint of transactional changes that are then implemented to a representation of a transactional table maintained at the analytical database, according to some embodiments.

[0148] In some embodiments, the checkpoint pertaining to Table A **914** that was provided to transport mechanism **916** from transactional database **900** may be received (e.g., receive checkpoint pertaining to Table A **920**) by leader node **932** of node cluster **930** of an analytical

database. Node cluster **930** may resemble node cluster **430** and the functionalities of node cluster **430** described herein (e.g., query processing and execution). In addition, node cluster **930** may be a node cluster of analytical database system **400** within analytical database service **150**, according to some embodiments.

[0149] Leader node **932** may partition checkpoint pertaining to Table A **920** according to the slices of the representation of Table A that are stored at respective compute nodes (e.g., compute nodes **934-938**) of node cluster **930** using change-data-capture log manager **934**. Leader node **932** may then provide node-specific checkpoint instructions **960** to respective compute nodes (e.g., compute nodes **934-938**) such that the transactional changes of the checkpoint may be applied to the representation of Table A being maintained at node cluster **930**. In other embodiments, incoming the transactional changes of checkpoint **920** may already be partitioned by compute node, and leader node **932** provides the node-specific checkpoint instructions **960**.

[0150] Once node-specific checkpoint instructions **960** have been received at the respective compute nodes of node cluster **930**, computing device(s) **962** may perform transactional changes of the checkpoint **964** to transactional Table A representation **970**, according to some embodiments. Perform transactional changes of checkpoint **964** may include the processes described in at least block **1006**, such as applying and committing “insert” events to the representation of Table A. In some embodiments, node-specific checkpoint instructions **960** may include “delete” events, which may be buffered and committed to shadow table **968** via buffer delete events **966**. Buffered delete events in shadow table **968** may later be applied and committed to transactional Table A representation **970** as a background operation of node cluster **930** (see description for FIGS. **11A-13D** herein). Such processes of receiving checkpoints and performing the transactional changes to a representation of a given transactional table may be repeated via the methods and techniques described herein (e.g., processes described in block **500**).

[0151] Leader node **932** may also be configured to maintain node-specific transaction number tracking **936**, according to some embodiments. In order to understand a possible lag between the maximum transaction number applied to a given transactional table at the transactional database and the maximum transaction number applied to the corresponding representation of the transactional table at the analytical database, node-specific transaction number tracking **936** may include information about the maximum transaction number that has been committed at respective compute nodes of node cluster **930**. Maintaining such information at leader node **932** may allow leader node **932** to organize node-specific checkpoint instructions **960** that correspond to the transactional changes at respective compute nodes that still need to be applied to transactional Table A representation **970**. In some embodiments, based on information stored in node-specific transaction number tracking **936**, leader node **932** may request an additional checkpoint (e.g., poll the transactional database via request new checkpoints **838**) upon confirming that one or more compute nodes of the node cluster has committed all transactional changes in the current given checkpoint. In addition, in some embodiments in which node cluster **930** temporarily goes offline and/or an error in performing transactional changes to Table A representation occurs, leader node **932** may use information stored in node-specific transaction number tracking **936** to determine the transactional number that one or more of the compute nodes of node cluster **930** should be reset to (e.g., the respective maximum committed transaction numbers of the corresponding compute nodes).

[0152] FIG. **10** is a flow diagram illustrating a process of implementing transactional changes to a representation of a transactional table in which delete events are first buffered and committed to a shadow table before being applied and committed as a batch to the representation, according to some embodiments.

[0153] In some embodiments, the implementation of the transactional changes that are received by the analytical database and used to both update the representation and keep it maintained as a replication of a transactional table that is stored in the transactional database may include the steps

shown in FIG. 10. In addition, block 510, shown in FIG. 5, may resemble the process shown in blocks 1006, 1008, 1010, and 1012, according to some embodiments.

[0154] In block 1000, a representation of at least some portions of a transactional table, stored at a transactional database, is maintained at an analytical database. In some embodiments, the process for maintaining the representation may include the processes described in blocks 1002, 1004, 1006, 1008, 1010, and 1012, and may also include methods and techniques described in at least blocks 500-510 of FIG. 5. In some embodiments, the transactional database described in block 1000 may resemble transactional database system 200 of transactional database service 110, and the analytical database may resemble analytical database system 400 of analytical database service 150, including the functionalities of said components described herein. Furthermore, FIGS. 11A-13D provide example embodiments of the maintenance of a representation of a given transactional Table A at an analytical database.

[0155] In block 1002, segments of the portion(s) of the transactional table are received as snapshots to the analytical database. As described above, such snapshots may be stored in one or more compute nodes of a node cluster at the analytical database, such as compute nodes 934-938 of node cluster 930, according to some embodiments. Continuing with the example embodiments shown in FIGS. 11A-13D, block 1002 may refer to a snapshot of Table A, which is being maintained as a representation at the analytical database. In block 1004, checkpoints that correspond to the snapshots described in block 1002 are also received at the analytical database. In some embodiments in which the snapshot of block 1002 refers to Table A, the checkpoints of block 1004 may resemble transactional changes that are to be applied to the representation of Table A in order to maintain the replication of Table A at the analytical database (e.g., checkpoint 1140).

[0156] In block 1006, the transactional changes of the checkpoints of block 1004 are implemented in transactional order to the representation of the portion(s) of the table. Continuing with the example embodiments shown in FIGS. 11A-13D, block 1006 (and blocks 1008, 1010, and 1012) may refer to the implementation of at least transactional changes contained in checkpoint 1140 (e.g., transactions 1150-1180) to Table A representation 1100, according to some embodiments. Block 1008 may include the process of applying and committing “insert” events (see the description of transactional changes within checkpoint 1140 below), while blocks 1010 and 1012 may include the process of buffering and committing “delete” events to a shadow table, and then later applying and committing them to the representation of Table A (see at least the description of shadow table 1230 and FIGS. 13A-13D below). As shown by blocks 1008, 1010, and 1012, insert operations of the transactional changes within the given checkpoint described by block 1006 may be applied and committed to the transactional table representation without being buffered or committed to a shadow table, while delete operations may be buffered to a shadow table for implementation to the representation of Table A at a later point. The buffered delete operations may be applied to the representation of Table A in transactional order as a “lazy” operation (e.g., at moments in which there is available processing power, and/or once a threshold limit of the shadow table has been reached), and then committed as a batch.

[0157] FIGS. 11A-D illustrate the state of a table representation maintained at an analytical database and a corresponding shadow table, checkpoint, and table as viewed by an external client of the analytical database just prior to the implementation of a given transaction number 1180, according to some embodiments.

[0158] FIG. 11A provides an illustrative example of a given transactional table, Table A, that is stored at the transactional database and whose representation (e.g., a replication of Table A) is being maintained at the analytical database via the methods and techniques described herein. Table A representation 1100 is meant to be a conceptual example for illustrative purposes. A person having ordinary skill in the art should understand that table A representation 1100 in FIG. 11A may represent a visual reconstruction of the portions of the Table A representation that may be stored across multiple compute nodes of a given node cluster at the analytical database. For example,

portions (e.g., slices, shards, etc.) of Table A representation **1100** may be maintained across compute nodes **934-938**, and stored in transactional table A representation **970** and/or cold storage tier **406**, according to some embodiments. Furthermore, Table A contains three columns in FIG. **11A** (column **1102**, column **1104**, and column **1106**), but a person having ordinary skill in the art should also understand that Table A and/or other table representations maintained at the analytical database may contain more or less columns than in the example given in FIG. **11A** and still utilize the methods and techniques described herein. In addition, one or more columns may be added or removed from Table A representation **1100** during the process of maintaining Table A representation **1100** at the analytical database. Table A representation **1100** may also have columns that are for internal use by the analytical database, namely inserts tracking **1108** and deletes tracking **1110**, according to some embodiments. Inserts tracking **1108** and deletes tracking **1110** columns will be further detailed in the following paragraphs.

[0159] In some embodiments, Table A representation **1100** may contain data items that may be organized in a columnar, tabular format. At the moment in time illustrated by FIGS. **11A-D** (e.g., after the application of transaction number **1170** but before beginning the application of transaction number **1180**), Table A representation **1100** may contain four rows, each comprising three data items. For example, row **1112** contains a string 'Alpha' in column **1102**, an integer '1' in column **1104**, and another string 'Value_1' in column **1106**. Similarly, row **1114** contains string 'Alpha' in column **1102**, integer '2' in column **1104**, and string 'Value_1' in column **1106**; row **1116** contains string 'Beta' in column **1102**, integer '1' in column **1104**, and string 'Value_1' in column **1106**; and row **1118** contains string 'Beta' in column **1102**, integer '2' in column **1104**, and string 'Value_1' in column **1106**. As a shorthand for the description herein, the data items in a given row, such as row **1112**, may be referred to as 'Alpha|1|Value_1.' Again, Table A representation **1100** is meant for illustrative purposes, and additional columns containing other data types (e.g., floats, characters, Boolean variables, etc.) may also be incorporated into other embodiments of the examples described herein.

[0160] In some embodiments, Table A representation **1100** may also contain columns for internal use, such as inserts tracking **1108** and deletes tracking **1110**, that may track the transaction number that a given row was added to and/or deleted from the table. For example, as shown in FIG. **11A**, row **1112** has been added during the implementation of transaction number **1150**, and then deleted during transaction number **1160**. In another example, row **1114** has been added during transaction number **1155**, and, at the moment in time depicted by FIG. **11A** (e.g., after the application of transaction **1170** but before beginning the application of transaction **1180**), row **1114** has not been deleted. "Hidden" columns, such as inserts tracking **1108** and deletes tracking **1110**, may serve at least several purposes: Firstly, such transaction number tracking allows the analytical database to return query results based on transaction number. For example, an incoming query may request data items pertaining to the state of Table A representation at transaction number **1175**. At transaction number **1175**, the representation of Table A may resemble transaction **1175** Table A view **1142**, if viewed by a client of the analytical database service, as shown in FIG. **11D**, in which rows **1114** and **1118** are visible. The incoming query may return results containing data items from rows **1114** and **1118**, but not rows **1112** or **1116**, as, from the perspective of Table A at transaction number **1175**, rows **1112** and **1116** have already been deleted from Table A. Secondly, the hidden columns may also serve as a record in the event that Table A representation needs to revert back to the state it was at during a previously committed transaction. For example, if transactional changes for Table A representation **1100** have been both applied and committed up through transaction number **1160**, but transaction numbers **1165** and **1170** have been applied but not yet committed, in the event that the analytical database goes temporarily offline at this moment (e.g., due to a temporary network malfunction, power outage affecting servers of the analytical database, etc.), Table A representation **1100** will be reverted back to the state it was at just after the last transaction commit (e.g., transaction number **1160** in the given example).

[0161] In some embodiments, examples of transactional changes of a given checkpoint (e.g., checkpoint **1140**) that may be implemented to Table A representation **1100** are shown in FIG. **11C**. Transactional changes may include “insert,” “delete,” and “update” events (e.g., operations), which are further explained in the following paragraphs.

[0162] In a first example of “insert” events, transaction **1150** contains the instructions to insert (e.g., append, add, etc.) a new row to Table A representation **1100**, wherein the new row contains the following data items for column **1102**, column **1104**, and column **1106**, respectively: ‘Alpha,’ ‘1,’ and ‘Value_1.’ As shown in row **1112** of Table A representation **1100**, data items ‘Alpha,’ ‘1,’ and ‘Value_1’ are entered into Table A representation **1100**, and transaction number **1150** is recorded in the inserts tracking **1108** column for row **1112** as the transaction number during which row **1112** was inserted. After committing (or at least applying) the transactional changes contained in transaction **1150**, the compute nodes of the node cluster at the analytical database will advance to the instructions of the next transaction number contained in checkpoint **1140** (e.g., transaction **1155** in this example). Similarly, transaction **1155** contains the instructions to insert row **1114** into Table A representation **1100** using data items ‘Alpha,’ ‘2,’ and ‘Value_1,’ and transaction number **1155** is recorded in the inserts tracking **1108** column for row **1114** as the transaction number during which row **1114** was inserted.

[0163] In a second example containing an “update” event, transaction **1160** contains the instructions to update an already-existing row of Table A representation **1100**. In some embodiments, an “update” event may contain at least one “insert” operation and one “delete” operation. In addition, transaction **1160** uses a primary key to inform the compute nodes of the node cluster which row the update event pertains to. As shown in FIG. **11C**, the instructions for transaction **1160** pertain to the row with a primary key of ‘{Alpha, 1},’ which the compute nodes of the node cluster then understand to mean row **1112**, according to some embodiments. The instructions for the given update event then instruct the compute nodes of the node cluster to replace ‘Alpha’ in column **1102** of row **1112** with ‘Beta.’ In some embodiments, this update event may then comprise two steps: (1) deletion of row **1112** containing data items ‘Alpha|1|Value_1’ from Table A representation **1100**, and (2) insertion of row **1116** containing data items ‘Beta|1|Value_1’ into Table A representation **1100**. As shown in Table A representation **1100**, a record of the deletion of row **1112** is recorded into the deletes tracking **1110** column of row **1112**, and a record of the insertion of row **1116** is recorded into the inserts tracking **1108** column of row **1116**, according to some embodiments. A person having ordinary skill in the art should understand that, according to different embodiments of the implementation of the given update event, the compute nodes of the node cluster may first delete row **1112** and then add row **1116**, or first add row **1116** and then delete row **1112**, as both embodiments still result in the application of the update event. Furthermore, the instructions for the deletion of row **1112** may instead be buffered and committed to shadow table **1130** (e.g., with primary key ‘{Alpha,1}’ being stored in the primary key **1132** column and transaction number **1160** being stored in the buffered deletes tracking **1134** column), which is discussed below with regard to FIGS. **12A-12D**.

[0164] In some embodiments, checkpoint **1140** may then contain instructions in transaction **1165** to insert a new row to Table A representation **1100** containing data items ‘Beta|2|Value_1,’ which is shown in Table A representation **1100** as row **1118**. Similarly, a record of the insertion of row **1118** into Table A representation **1100** is recorded in the inserts tracking **1108** column of row **1118** as being inserted during the implementation of transaction number **1165**. Following the committing (or at least application) of the transactional changes contained in transaction **1165**, checkpoint **1140** may then contain instructions in transaction **1170** to delete the row of Table A representation **1100** with the primary key ‘{Beta,1}.’ As shown in Table A representation **1100**, row **1116**, which has data items ‘Beta’ in column **1102** and ‘1’ in column **1104**, is deleted from Table A representation **1100**, as recorded in the deletes tracking **1100** column of row **1116**. In some embodiments, after the implementation of transactional changes with lower transaction numbers than transaction **1180**, the

compute nodes of the node cluster may move on to the implementation of the transactional changes in transaction **1180**. Transaction **1180** is discussed below in further detail with regard to FIGS. **12A-D**.

[0165] Checkpoint **1140** is meant to be an illustrative example of a checkpoint that contains a portion of the transactional changes of the change-data-capture log received from the transactional database via the methods and techniques described herein, and a person having ordinary skill in the art should understand that checkpoints such as checkpoint **1140** may be written in other forms (e.g., structured query language (SQL)) as long as the data definition language (DDL) commands of the given checkpoint may be interpreted by the analytical database. Furthermore, checkpoint **1140** may be received to compute nodes **934-938** via node-specific checkpoint instructions **960**, according to some embodiments. In addition, for illustrative purposes, transactions **1150**, **1155**, **1160**, **1165**, and **1170** have been crossed out in FIG. **11C** to represent the concept that transactions **1150**, **1155**, **1160**, **1165**, and **1170** have already been applied (and, in some embodiments, committed) to Table A representation **1100**. The arrow pointing to transaction **1180** in FIG. **11C** similarly represents the next transaction number that the compute nodes of the node cluster are going to apply to Table A representation **1100** at the moment in time depicted by FIGS. **11A-11D**.

[0166] FIGS. **12A-D** illustrate the state of the table representation, shadow table, checkpoint, and table view of FIGS. **11A-D** just after the implementation of transaction number **1180**, according to some embodiments.

[0167] Continuing with the example embodiments shown in FIGS. **11A-D**, FIGS. **12A-D** demonstrate the effects of the implementation of transaction number **1180** to Table A representation **1100** and shadow table **1130**, in which Table A representation **1100** and shadow table **1130** may represent the states of Table A representation and the corresponding shadow table just prior to the implementation of transaction number **1180**, and Table A representation **1200** and shadow table **1230** may represent the states of Table A representation and the corresponding shadow table just after the implementation of transaction number **1180**, according to some embodiments. In some embodiments, transaction number **1180** may include an “update” event that applies to all non-deleted rows of Table A representation **1200** (e.g., at the moment of time that transaction **1180** is being applied to the Table A representation). As discussed above with regard to FIGS. **11A-D**, the update event of transaction number **1180** may include respective insert and delete operations for each of the respective rows that the update event pertains to. For example, the non-deleted rows of Table A representation **1200** at the moment in time just prior to the implementation of the instructions for transaction number **1180** may be rows **1114** and **1118**, according to some embodiments. Therefore, individual insert and delete operations pertaining to rows **1114** and **1118**, respectively, may be applied to Table A representation as part of implementing the transactional changes of transaction number **1180**, according to some embodiments. As shown in Table A representation **1200**, rows **1220** and **1222** are inserted into Table A representation **1200** in order to apply the respective insert operations of transaction number **1180** in which the data items in column **1106** of rows **1114** and **1118** have both been respectively changed from ‘Value_1’ to ‘Value_2’ in column **1106** of rows **1220** and **1222**. In some embodiments, an additional “hidden” column of Table A representation **1200** may be used to track operations of a given transaction. For example, in embodiments in which row **1220** is inserted into Table A representation **1200** before row **1222**, both as part of the implementation of transaction number **1180**, row **1220** may be listed as operation 1 of transaction **1180** (e.g., **1180.1**) and row **1222** may be listed as operation 2 of transaction **1180** (e.g., **1180.2**) in “hidden” column operations tracking. A tracking of operations within a given transaction number may also be referred to as assigning Statement IDs to the operations of a given transaction number, according to some embodiments.

[0168] In some embodiments, as part of implementing the transactional changes of transaction number **1180**, the delete operations (e.g., the deletions of rows **1114** and **1118**) may be buffered and committed to shadow table **1230**, as shown in FIG. **12B**, via the methods and techniques described

herein with regard to at least FIG. 9B and FIG. 10. As shown in shadow table 1230, rows 1236 and 1238 respectively represent row 1114 and 1118 of Table A representation 1200, wherein primary keys ‘{Alpha,2}’ and ‘{Beta,2}’ are respectively stored in primary key 1132 column of rows 1236 and 1238, and transaction number 1180 is recorded in buffered deletes tracking 1134 as the transaction number at which rows 1114 and 1118 were deleted from Table A representation 1200. As discussed with regard to block 1012 of FIG. 10, rows 1236 and 1238 may remain in shadow table 1230 for a given amount of time before being applied and committed to Table A representation 1200. The “batch” commit of rows 1236 and 1238 of shadow table 1230 to Table A representation 1200 is discussed below with regard to FIGS. 13A-D. In addition, rows 1112 and 1116 of Table A representation 1200 remain unchanged between the embodiments shown in FIG. 11A and the embodiments shown in FIG. 12A, as the instructions contained in transaction number 1180 do not apply to row 1112 or row 1116.

[0169] As discussed above with regard to Table A representation 1100, shadow table 1230 is meant to be a conceptual example for illustrative purposes. A person having ordinary skill in the art should understand that shadow table 1230 may represent a visual reconstruction of portions of the shadow table that may be stored across multiple compute nodes of a given node cluster at the analytical database. For example, portions of shadow table 1230 may be maintained across compute nodes 934-938, and stored in shadow table 968, according to some embodiments.

[0170] As shown in FIG. 12C, following the insertion of rows 1220 and 1222 to Table A representation 1200 and the buffering and committing of rows 1236 and 1238 in shadow table 1230, the instructions contained in transaction number 1180 may be considered to be complete, and the compute nodes of the node cluster may move on to a transaction number that follows transaction number 1180 (as indicated by the arrow on checkpoint 1140). In some embodiments, this may be transaction number 1181. In other embodiments, transaction number 1180 may be the final transaction number in node-specific checkpoint instructions 960, at which point leader node 932 may provide additional node-specific checkpoint instructions to compute nodes 934-938, and/or poll the analytical database for additional transactional changes to apply to Table A representation.

[0171] In some embodiments, FIG. 12D may represent Table A representation 1200 if viewed by a client of the analytical database service just after the implementation of transaction number 1180. For example, transaction 1181 Table A view 1242 may include two rows which correspond to rows 1220 and 1222 of Table A representation. In some embodiments, rows that have been indicated as deleted from Table A representation 1200 on or before transaction number 1180 via deletes tracking 1110 column (e.g., row 1112 and row 1114) may not be included in transaction 1181 Table A view 1242. In addition, rows that have been marked for future deletion from Table A representation 1200 may be buffered to shadow table 1230 (e.g., row 1236 and row 1238), and therefore may also not be included in transaction 1181 Table A view 1242, according to some embodiments. To further emphasize this point, this may mean that even though rows 1114 and 1118 may not have been recorded as being deleted at transaction number 1180 in the deletes tracking 1110 column of Table A representation 1200, if they have been buffered and committed to shadow table 1230 (e.g., as shown by rows 1236 and 1238), then they may be treated as deleted rows of Table A representation 1200 from the perspective of a client of the analytical database service. Furthermore, if an incoming query were to request results of Table A representation 1200 based on the state of Table A representation just after the completion of instructions contained in transaction number 1180, results of rows 1112, 1114, 1116, and 1118 may not be included in the results of the query, based on the state of Table A representation 1200 and shadow table 1230.

[0172] FIGS. 13A-D illustrate the state of the table representation, shadow table, checkpoint, and table view of FIGS. 11A-D at a point after the implementation of transaction number 1180, according to some embodiments.

[0173] In some embodiments, the “batch” commit of rows 1236 and 1238 of shadow table 1230 to

Table A representation **1300** may resemble embodiments shown in FIGS. **13A-D**. Following the embodiments shown in FIGS. **12A-D** in which the respective insert operations of transaction number **1180** have been inserted into Table A representation **1200** and the respective delete operations of transaction number **1180** have been buffered and committed to shadow table **1230**, FIGS. **13A-D** illustrate a later point in time in which rows **1236** and **1238** of shadow table **1230** are applied and committed to Table A representation **1300**. As discussed with regard to FIG. **10**, delete operations that have been buffered and committed to the shadow table throughout the process of implementing transactional changes of given checkpoint(s) (e.g., checkpoint **1140** as shown in FIG. **13C**) to the representation of Table A may then be applied in transactional order to the representation of Table A and committed as a batch. In some embodiments, buffered delete events of shadow table **1330** may be applied to representation of Table A **1300** as a “lazy,” background operation, in which the application of delete events to Table A representation **1300** are delayed and subsequently applied as a background operation (e.g., when computing/processing power may be available for such delete operation applications). In other embodiments, delete events will continue to be buffered and committed to the shadow table until the shadow table reaches a given threshold capacity, at which point a “batch” application of delete events to Table A representation is triggered. A person having ordinary skill in the art should understand that different combinations of such reasons for first buffering delete events to a shadow table and subsequently committing them to Table A representation may also be included in such embodiments. For example, if the shadow table has not yet reached the given threshold capacity, but processing power is available to apply the delete events to Table A representation **1300**, the delete events may still be applied using the available processing power. After the buffered delete events have been applied to Table A representation **1300**, they are committed as a batch and removed from shadow table **1330**, as described in the following paragraphs.

[0174] As shown in Table A representation **1300**, rows **1112**, **1116**, **1220**, and **1222** remain unchanged between the embodiments shown in FIG. **12A** and the embodiments shown in FIG. **13A**, as applying and committing rows **1236** and **1238** of shadow table **1230** to Table A representation **1300** does not concern rows **1112**, **1116**, **1220**, and **1222**. However, the application of rows **1236** and **1238** of shadow table **1230** to Table A representation **1300** may involve recording transaction number **1180** into the deletes tracking **1110** column of rows **1314** and **1318**.

[0175] In some embodiments, the application of a given buffered delete event of the shadow table to Table A representation may include three steps: (1) using the primary key in the primary key **1132** column, scan (or parse, search, etc.) the representation of Table A for the row at which to apply the delete event, (2) apply (or write, etc.) the delete event to the given row of Table A representation, and (3) commit the delete event. This process may be repeated, in transactional order, for a number of buffered delete events of the shadow table. For example, row **1236** of shadow table **1230** contains information that a delete event at transaction number **1180** has been buffered and committed to shadow table **1230**, and pertains to a row of Table A representation **1300** that has a primary key of ‘{Alpha, 2}.’ Upon scanning Table A representation **1300** for a row with a primary key of ‘{Alpha, 2}’ and confirming that primary key ‘{Alpha, 2}’ corresponds to row **1314**, transaction number **1180** is recorded into the deletes tracking **1110** column of row **1314**. In some embodiments, the process may repeat for row **1238** of shadow table **1230** in which primary key ‘{Beta, 2}’ is used to record transaction number **1180** into the deletes tracking **1110** column of row **1318**. These transactional changes to rows **1314** and **1318** of Table A representation **1300** may then be committed as a batch. Upon committing said transactional changes, leader node **932** may update node-specific transaction number tracking **936** to having committed transaction number **1180** for the compute nodes pertaining to the embodiments shown in Table A representation **1300**, according to some embodiments.

[0176] In some embodiments, FIG. **13D** may represent Table A representation as viewed from a client of the analytical database service just after the commit of transaction number **1180**. It may be

noted that transaction **1181** Table A view **1342**, as shown in FIG. **13D**, remains unchanged from transaction **1181** Table A view **1232**, as shown in FIG. **12D**, since, from the perspective of a client of the analytical database service, rows **1314** and **1318** were either marked as a buffered delete event of the shadow table (e.g., FIG. **12B**) or committed delete events of the representation of Table A (e.g., FIG. **13B**).

[0177] Embodiments of the hybrid transactional and analytical processing methods and systems described herein may be executed on one or more computer systems, which may interact with various other devices. One such computer system is illustrated by FIG. **14**. FIG. **14** is a block diagram illustrating a computer system that may implement at least a portion of the systems described herein, according to various embodiments. For example, computer system **1400** may implement a database engine head node of a database tier, or one of a plurality of storage nodes of a separate distributed storage system that stores databases and associated metadata on behalf of clients of the database tier, in different embodiments. Computer system **1400** may be any of various types of devices, including, but not limited to, a personal computer system, desktop computer, laptop or notebook computer, mainframe computer system, handheld computer, workstation, network computer, a consumer device, application server, storage device, telephone, mobile telephone, or in general any type of computing device.

[0178] Computer system **1400** includes one or more processors **1410** (any of which may include multiple cores, which may be single or multi-threaded) coupled to a system memory **1420** via an input/output (I/O) interface **1430**. Computer system **1400** further includes a network interface **1440** coupled to I/O interface **1430**. In various embodiments, computer system **1400** may be a uniprocessor system including one processor **1410**, or a multiprocessor system including several processors **1410** (e.g., two, four, eight, or another suitable number). Processors **1410** may be any suitable processors capable of executing instructions. For example, in various embodiments, processors **1410** may be general-purpose or embedded processors implementing any of a variety of instruction set architectures (ISAs), such as the x86, PowerPC, SPARC, or MIPS ISAs, or any other suitable ISA. In multiprocessor systems, each of processors **1410** may commonly, but not necessarily, implement the same ISA. The computer system **1400** also includes one or more network communication devices (e.g., network interface **1440**) for communicating with other systems and/or components over a communications network (e.g. Internet, LAN, etc.). For example, a client application executing on system **1400** may use network interface **1440** to communicate with a server application executing on a single server or on a cluster of servers that implement one or more of the components of the database systems described herein. In another example, an instance of a server application executing on computer system **1400** may use network interface **1440** to communicate with other instances of the server application (or another server application) that may be implemented on other computer systems (e.g., computer systems **1490**).

[0179] In the illustrated embodiment, computer system **1400** also includes one or more persistent storage devices **1460** and/or one or more I/O devices **1480**. In various embodiments, persistent storage devices **1460** may correspond to disk drives, tape drives, solid state memory, other mass storage devices, or any other persistent storage device. Computer system **1400** (or a distributed application or operating system operating thereon) may store instructions and/or data in persistent storage devices **660**, as desired, and may retrieve the stored instruction and/or data as needed. For example, in some embodiments, computer system **1400** may host a storage node, and persistent storage **1460** may include the SSDs attached to that server node.

[0180] Computer system **1400** includes one or more system memories **1420** that may store instructions and data accessible by processor(s) **1410**. In various embodiments, system memories **1420** may be implemented using any suitable memory technology, (e.g., one or more of cache, static random-access memory (SRAM), DRAM, RDRAM, EDO RAM, DDR 10 RAM, synchronous dynamic RAM (SDRAM), Rambus RAM, EEPROM, non-volatile/Flash-type memory, or any other type of memory). System memory **1420** may contain program instructions

1425 that are executable by processor(s) **1410** to implement the methods and techniques described herein. In various embodiments, program instructions **1425** may be encoded in platform native binary, any interpreted language such as Java™ byte-code, or in any other language such as C/C++, Java™, etc., or in any combination thereof. For example, in the illustrated embodiment, program instructions **1425** include program instructions executable to implement the functionality of a database engine head node of a database tier, or one of a plurality of storage nodes of a separate distributed storage system that stores databases and associated metadata on behalf of clients of the database tier, in different embodiments. In some embodiments, program instructions **1425** may implement multiple separate clients, server nodes, and/or other components.

[0181] In some embodiments, program instructions **1425** may include instructions executable to implement an operating system (not shown), which may be any of various operating systems, such as UNIX, LINUX, Solaris™, MacOS™, Windows™, etc. Any or all of program instructions **1425** may be provided as a computer program product, or software, that may include a non-transitory computer-readable storage medium having stored thereon instructions, which may be used to program a computer system (or other electronic devices) to perform a process according to various embodiments. A non-transitory computer-readable storage medium may include any mechanism for storing information in a form (e.g., software, processing application) readable by a machine (e.g., a computer). Generally speaking, a non-transitory computer-accessible medium may include computer-readable storage media or memory media such as magnetic or optical media, e.g., disk or DVD/CD-ROM coupled to computer system **1400** via I/O interface **1430**. A non-transitory computer-readable storage medium may also include any volatile or non-volatile media such as RAM (e.g. SDRAM, DDR SDRAM, RDRAM, SRAM, etc.), ROM, etc., that may be included in some embodiments of computer system **1400** as system memory **1420** or another type of memory. In other embodiments, program instructions may be communicated using optical, acoustical or other form of propagated signal (e.g., carrier waves, infrared signals, digital signals, etc.) conveyed via a communication medium such as a network and/or a wireless link, such as may be implemented via network interface **1440**.

[0182] In some embodiments, system memory **1420** may include data store **1445**, which may be implemented as described herein. For example, the information described herein as being stored by the database tier (e.g., on a database engine head node), such as a transaction log, an undo log, cached page data, or other information used in performing the functions of the database tiers described herein may be stored in data store **1445** or in another portion of system memory **1420** on one or more nodes, in persistent storage **1460**, and/or on one or more remote storage devices **1470**, at different times and in various embodiments. Similarly, the information described herein as being stored by the storage tier (e.g., redo log records, coalesced data pages, and/or other information used in performing the functions of the distributed storage systems described herein) may be stored in data store **1445** or in another portion of system memory **1420** on one or more nodes, in persistent storage **1460**, and/or on one or more remote storage devices **1470**, at different times and in various embodiments. In general, system memory **1420** (e.g., data store **1445** within system memory **1420**), persistent storage **1460**, and/or remote storage **1470** may store data blocks, replicas of data blocks, metadata associated with data blocks and/or their state, database configuration information, and/or any other information usable in implementing the methods and techniques described herein.

[0183] In one embodiment, I/O interface **1430** may coordinate I/O traffic between processor **1410**, system memory **1420** and any peripheral devices in the system, including through network interface **1440** or other peripheral interfaces. In some embodiments, I/O interface **1430** may perform any necessary protocol, timing or other data transformations to convert data signals from one component (e.g., system memory **1420**) into a format suitable for use by another component (e.g., processor **1410**). In some embodiments, I/O interface **1430** may include support for devices attached through various types of peripheral buses, such as a variant of the Peripheral Component Interconnect (PCI) bus standard or the Universal Serial Bus (USB) standard, for example. In some

embodiments, the function of I/O interface **1430** may be split into two or more separate components, such as a north bridge and a south bridge, for example. Also, in some embodiments, some or all of the functionality of I/O interface **1430**, such as an interface to system memory **1420**, may be incorporated directly into processor **1410**.

[0184] Network interface **1440** may allow data to be exchanged between computer system **1400** and other devices attached to a network, such as other computer systems **1490** (which may implement one or more storage system server nodes, database engine head nodes, and/or clients of the database systems described herein), for example. In addition, network interface **1440** may allow communication between computer system **1400** and various I/O devices **1450** and/or remote storage **1470**. Input/output devices **1450** may, in some embodiments, include one or more display terminals, keyboards, keypads, touchpads, scanning devices, voice or optical recognition devices, or any other devices suitable for entering or retrieving data by one or more computer systems **1400**. Multiple input/output devices **1450** may be present in computer system **1400** or may be distributed on various nodes of a distributed system that includes computer system **1400**. In some embodiments, similar input/output devices may be separate from computer system **1400** and may interact with one or more nodes of a distributed system that includes computer system **1400** through a wired or wireless connection, such as over network interface **1440**. Network interface **1440** may commonly support one or more wireless networking protocols (e.g., Wi-Fi/IEEE 802.11, or another wireless networking standard). However, in various embodiments, network interface **1440** may support communication via any suitable wired or wireless general data networks, such as other types of Ethernet networks, for example. Additionally, network interface **1440** may support communication via telecommunications/telephony networks such as analog voice networks or digital fiber communications networks, via storage area networks such as Fibre Channel SANs, or via any other suitable type of network and/or protocol. In various embodiments, computer system **1400** may include more, fewer, or different components than those illustrated in FIG. 20 (e.g., displays, video cards, audio cards, peripheral devices, other network interfaces such as an ATM interface, an Ethernet interface, a Frame Relay interface, etc.)

[0185] It is noted that any of the distributed system embodiments described herein, or any of their components, may be implemented as one or more web services. For example, a database engine head node within the database tier of a database system may present database services and/or other types of data storage services that employ the distributed storage systems described herein to clients as web services. In some embodiments, a web service may be implemented by a software and/or hardware system designed to support interoperable machine-to-machine interaction over a network. A web service may have an interface described in a machine-processable format, such as the Web Services Description Language (WSDL). Other systems may interact with the web service in a manner prescribed by the description of the web service's interface. For example, the web service may define various operations that other systems may invoke, and may define a particular application programming interface (API) to which other systems may be expected to conform when requesting the various operations.

[0186] In various embodiments, a web service may be requested or invoked through the use of a message that includes parameters and/or data associated with the web services request. Such a message may be formatted according to a particular markup language such as Extensible Markup Language (XML), and/or may be encapsulated using a protocol such as Simple Object Access Protocol (SOAP). To perform a web services request, a web services client may assemble a message including the request and convey the message to an addressable endpoint (e.g., a Uniform Resource Locator (URL)) corresponding to the web service, using an Internet-based application layer transfer protocol such as Hypertext Transfer Protocol (HTTP).

[0187] In some embodiments, web services may be implemented using Representational State Transfer (“RESTful”) techniques rather than message-based techniques. For example, a web service implemented according to a RESTful technique may be invoked through parameters

included within an HTTP method such as PUT, GET, or DELETE, rather than encapsulated within a SOAP message.

[0188] The various methods as illustrated in the FIGs. and described herein represent example embodiments of methods. The methods may be implemented in software, hardware, or a combination thereof. The order of method may be changed, and various elements may be added, reordered, combined, omitted, modified, etc.

[0189] Various modifications and changes may be made as would be obvious to a person skilled in the art having the benefit of this disclosure. It is intended that the invention embrace all such modifications and changes and, accordingly, the above description to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A system, comprising: one or more computing devices, wherein the one or more computing devices are configured to: implement a transactional database; maintain, at the transactional database, a table comprising data items; and write, at the transactional database, transactional changes made to the table to a change-data-capture log; one or more compute nodes organized into a node cluster, wherein the one or more compute nodes are configured to: implement an analytical database; and maintain, at the analytical database, a representation of at least one portion of the table, wherein to maintain comprises committing respective ones of the transactional changes of the change-data-capture log, received at the analytical database, to the representation; at least one transport mechanism configured to: receive respective snapshots of segments of the at least one portion of the table, wherein each snapshot comprises a segment of the at least one portion of the table; make the respective snapshots accessible to the analytical database; receive checkpoints relative to the respective snapshots, wherein the checkpoints comprise respective portions of the transactional changes of the change-data-capture log; and make the checkpoints accessible to the analytical database; and wherein the one or more compute nodes are further configured to, responsive to receiving an incoming analytical query pertaining to at least one data item in the representation, provide the results of the incoming analytical query.
2. The system of claim 1, wherein: the transactional changes of a given checkpoint comprise at least two delete events; and the maintain the representation of the at least one portion of the table further comprises: committing, at the analytical database, the at least two delete events into a shadow table; and responsive to the shadow table having a given threshold of committed delete events, committing the committed delete events of the shadow table to the representation.
3. The system of claim 1, wherein the snapshots are respectively distributed across respective ones of the one or more compute nodes of the node cluster.
4. The system of claim 1, wherein: the transactional changes are labeled with respective primary keys corresponding to respective rows that the transactional changes occur at in the transactional database; and the committing the respective ones of the transactional changes to the representation is based, at least in part, on use of the respective primary keys.
5. The system of claim 1, wherein: the transactional database and the analytical database are part of a service provider network; a data storage service of the service provider network is used to implement a respective one of the at least one transport mechanism; and the respective snapshots or the checkpoints are stored in at least one data store of the data storage service, wherein the at least one data store is accessible to the transactional database and the analytical database.
6. The system of claim 1, wherein: the transactional database and the analytical database are part of a service provider network; and a data streaming service of the provider network is used to implement a respective one of the at least one transport mechanism; and the respective snapshots or the checkpoints are communicated from the transactional database to the analytical database via the data streaming service.

7. A method, comprising: maintaining, at an analytical database implemented by one or more compute nodes organized into a node cluster, a representation of at least one portion of a table comprising data items, wherein: the table is stored at a transactional database, implemented by one or more computing devices; and the maintaining comprises: writing, at the transactional database, transactional changes made to the table to a change-data-capture log; making, via at least one transport mechanism, snapshots of segments of the at least one portion of the table accessible to the analytical database; making, via a respective one of the at least one transport mechanism, checkpoints, relative to the respective snapshots, accessible to the analytical database, wherein the checkpoints comprise respective portions of the transactional changes of the change-data-capture log; storing, at the analytical database, the respective snapshots and the checkpoints across the one or more compute nodes of a node cluster; and committing the transactional changes of the respective checkpoints to the representation; and responsive to receiving an incoming query pertaining to at least one data item in the representation, providing the results of the incoming query.

8. The method of claim 7, wherein the applying the transactional changes comprises: committing at least two delete events of the transactional changes of a given checkpoint of the respective checkpoints into a shadow table; and responsive to the shadow table having a given threshold of committed delete events, committing the committed delete events in the shadow table to the representation.

9. The method of claim 7, wherein the method further comprises: storing, in relation to committing the transactional changes of the respective checkpoints to the representation, an indication of a transaction applied at the analytical database having a latest transaction identifier; and polling the transactional database to determine whether an additional checkpoint that comprises one or more transactional changes that have been applied at the transactional database have a later transaction identifier than the latest transaction identifier committed at the analytical database.

10. The method of claim 9, wherein, responsive to the polling the transactional database, making, via the respective one of the at least one transport mechanism, the additional checkpoint accessible to the analytical database; and committing the one or more transactional changes that have the later transaction identifier than the latest transaction identifier committed at the analytical database to the representation.

11. The method of claim 7, wherein the method further comprises: determining, via a negotiation between the one or more compute nodes of the node cluster and the one or more computing devices, one or more data definition language commands that are incompatible with the checkpoints.

12. The method of claim 11, wherein the method further comprises: responsive to determining, at the analytical database, that a given transactional change of a given checkpoint comprises a respective one of the one or more incompatible data definition language commands, providing an error message.

13. The method of claim 7, wherein: the transactional database and the analytical database are part of a service provider network; and the making the checkpoints accessible to the analytical database comprises storing the checkpoints in at least one data store of a data storage service of the service provider network.

14. The method of claim 7, wherein: the transactional database and the analytical database are part of a service provider network; and the making the checkpoints accessible to the analytical database comprises communicating the checkpoints from the transactional database to the analytical database via a data streaming service of the service provider network.

15. The method of claim 7, wherein: the checkpoints are stored in a storage used by the transactional database; the making the checkpoints accessible to the analytical database occurs via a direct connection between the transactional database and the analytical database; and the analytical database requests the checkpoints.

- 16.** The method of claim 7, wherein: the checkpoints are stored in a storage used by the transactional database; the making the checkpoints accessible to the analytical database occurs via a direct connection between the transactional database and the analytical database; and the transactional database provides the checkpoints.
- 17.** The method of claim 7, the method further comprising: performing a handshake protocol, wherein the performing the handshake protocol comprises: determining a mapping for partitioning the respective snapshots and the relative checkpoints across the one or more compute nodes of the node cluster; determining the at least one transport mechanism to be used for the making the snapshots accessible to the analytical database and the making the checkpoints accessible to the analytical database; and making determined results of the handshake protocol accessible to the transactional database and the analytical database.
- 18.** One or more non-transitory, computer-readable storage media storing program instructions that when executed on or across one or more processors, cause the one or more processors to: maintain, at an analytical database, a representation of at least one portion of a table comprising data items, wherein: the table is stored at a transactional database; and to maintain comprises: write, at the transactional database, transactional changes made to the table to a change-data-capture log; make, via at least one transport mechanism, snapshots of segments of the at least one portion of the table accessible to an analytical database; make, via a respective one of the at least one transport mechanism, checkpoints, relative to the respective snapshots, accessible to the analytical database, wherein the checkpoints comprise respective portions of the transactional changes of the change-data-capture log; store, at the analytical database, the snapshots and the checkpoints; and commit the transactional changes of the respective checkpoints to the representation; and responsive to receiving an incoming query pertaining to at least one data item in the representation, provide the results of the incoming query.
- 19.** The one or more non-transitory, computer-readable storage media of claim 18, wherein the program instructions, when executed on or across the one or more processors, further cause the one or more processors to: cause a leader node of the one or more compute nodes to store, in relation to the commit the transactional changes of the respective checkpoints to the representation at the analytical database, an indication of a transaction committed at the analytical database having a latest transaction identifier; and cause the leader node to poll the transactional database to determine whether an additional checkpoint that comprises one or more transactional changes that have been applied on the transactional database have a later transaction identifier than the latest transaction identifier committed at the analytical database.
- 20.** The one or more non-transitory, computer-readable storage media of claim 19, wherein, responsive to the cause the leader node to poll the transactional database, the program instructions, when executed on or across the one or more processors, further cause the one or more processors to: access the additional checkpoint; and commit the one or more transactional changes that have the later transaction identifier than the latest transaction identifier committed at the analytical database.
-