



(19) **United States**

(12) **Patent Application Publication**

(10) **Pub. No.: US 2025/0278533 A1**

(43) **Pub. Date: Sep. 4, 2025**

(12) **Morini et al.**

(54) **DISTINCT PATHS COUNTEREXAMPLES AND STIMULI GENERATION**

(52) **U.S. Cl.**
CPC **G06F 30/20** (2020.01); **G06T 15/005** (2013.01)

(71) Applicant: **Intel Corporation**, Santa Clara, CA (US)

(57) **ABSTRACT**

(72) Inventors: **Emiliano Morini**, El Dorado Hills, CA (US); **Theo Drane**, El Dorado Hills, CA (US); **William Zorn**, Folsom, CA (US); **Amol Bhinge**, Austin, TX (US)

(73) Assignee: **Intel Corporation**, Santa Clara, CA (US)


An apparatus to facilitate distinct paths counterexamples and stimuli generation is disclosed. The apparatus includes processing circuitry to: generate a graph representation structure from a data flow of a hardware design; for each leaf node of the graph representation structure, generate a cover property that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; execute, as a single run of a verification tool, a set of cover properties comprising the cover property for each of the leaf nodes; and report results of execution of the set of cover properties as a set of counterexamples for the hardware design.


(21) Appl. No.: **18/592,868**

(22) Filed: **Mar. 1, 2024**

Publication Classification

(51) **Int. Cl.**
G06F 30/20 (2020.01)
G06T 15/00 (2011.01)





100

FIG. 1




FIG. 2A

219




FIG. 2B




FIG. 2C




FIG. 2D




FIG. 3A




FIG. 3B




FIG. 3C




FIG. 4




FIG. 5A




FIG. 5B




FIG. 5C




FIG. 6




FIG. 7




FIG. 8

FIG. 9A

GRAPHICS PROCESSOR COMMAND FORMAT

900





FIG. 9B

GRAPHICS PROCESSOR COMMAND SEQUENCE

910



DATA PROCESSING SYSTEM
1000




FIG. 10

IP CORE DEVELOPMENT - 1100




FIG. 11A




FIG. 11B




FIG. 11C

1194




FIG. 11D




FIG. 12




FIG. 13A




FIG. 13B




FIG. 14




FIG. 15

1600




FIG. 16A




FIG. 16B

1610




FIG. 16C




FIG. 16D

1700

Extract a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges

1710

Generate, for each leaf node of the leaf nodes of the graph representation structure, a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the implementation design having the set of space constraints applied, where the set of space constraints define a path through the graph representation structure from a root node to the leaf node

1720


Execute, as a single run of a verification tool, a set of cover properties comprising the cover property for each leaf node of the graph representation structure

1730

Report results of execution of the set of cover properties as a set of counterexamples for the implementation design, where the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the implementation design

1740

FIG. 17

1800**FIG. 18**

DISTINCT PATHS COUNTEREXAMPLES AND STIMULI GENERATION

FIELD

[0001] This disclosure relates generally to data processing and more particularly to data processing via a general-purpose graphics processing unit, including distinct paths counterexamples and stimuli generation.

BACKGROUND OF THE DISCLOSURE

[0002] Current parallel graphics data processing includes systems and methods developed to perform specific operations on graphics data such as, for example, linear interpolation, tessellation, rasterization, texture mapping, depth testing, etc. Traditionally, graphics processors used fixed function computational units to process graphics data. However, more recently, portions of graphics processors have been made programmable, enabling such processors to support a wider variety of operations for processing vertex and fragment data.

[0003] To further increase performance, graphics processors typically implement processing techniques such as pipelining that attempt to process, in parallel, as much graphics data as possible throughout the different parts of the graphics pipeline. Parallel graphics processors with single instruction, multiple thread (SIMT) architectures are designed to maximize the amount of parallel processing in the graphics pipeline. In a SIMT architecture, groups of parallel threads attempt to execute program instructions synchronously together as often as possible to increase processing efficiency. A general overview of software and hardware for SIMT architectures can be found in Shane Cook, CUDA Programming Chapter 3, pages 37-51 (2013).

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] Embodiments described herein are illustrated by way of example and not limitation in the figures of the accompanying drawings in which like references indicate similar elements, and in which:

[0005] FIG. 1 is a block diagram of a processing system, according to an embodiment.

[0006] FIG. 2A is a block diagram of an embodiment of a processor having one or more processor cores, an integrated memory controller, and an integrated graphics processor.

[0007] FIG. 2B is a block diagram of hardware logic of a graphics processor core block, according to some embodiments described herein.

[0008] FIG. 2C illustrates a graphics processing unit (GPU) that includes dedicated sets of graphics processing resources arranged into multi-core groups.

[0009] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein.

[0010] FIG. 3A is a block diagram of a graphics processor, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces.

[0011] FIG. 3B illustrates a graphics processor having a tiled architecture, according to embodiments described herein.

[0012] FIG. 3C illustrates a compute accelerator, according to embodiments described herein.

[0013] FIG. 4 is a block diagram of a graphics processing engine of a graphics processor in accordance with some embodiments.

[0014] FIG. 5A illustrates graphics core cluster, according to an embodiment.

[0015] FIG. 5B illustrates a vector engine of a graphics core, according to an embodiment.

[0016] FIG. 5C illustrates a matrix engine of a graphics core, according to an embodiment.

[0017] FIG. 6 illustrates a tile of a multi-tile processor, according to an embodiment.

[0018] FIG. 7 is a block diagram illustrating graphics processor instruction formats according to some embodiments.

[0019] FIG. 8 is a block diagram of another embodiment of a graphics processor.

[0020] FIG. 9A is a block diagram illustrating a graphics processor command format that may be used to program graphics processing pipelines according to some embodiments.

[0021] FIG. 9B is a block diagram illustrating a graphics processor command sequence according to an embodiment.

[0022] FIG. 10 illustrates an example graphics software architecture for a data processing system according to some embodiments.

[0023] FIG. 11A is a block diagram illustrating an IP core development system that may be used to manufacture an integrated circuit to perform operations according to an embodiment.

[0024] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly 1170, according to some embodiments described herein.

[0025] FIG. 11C illustrates a package assembly that includes multiple units of hardware logic chiplets connected to a substrate.

[0026] FIG. 11D illustrates a package assembly including interchangeable chiplets, according to an embodiment.

[0027] FIG. 12 is a block diagram illustrating an example system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0028] FIG. 13A illustrates an example graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0029] FIG. 13B illustrates an additional example graphics processor 1340 of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0030] FIG. 14 is a block diagram illustrating an example system for distinct paths counterexample and stimuli generation, according to embodiments.

[0031] FIG. 15 is a block diagram illustrating a detailed view of an example formal verification tool facilitating multiple distinct paths counterexamples and stimuli generation, in accordance with implementations herein.

[0032] FIGS. 16A-16D illustrates example graph representation structures of a hardware design provided by sample register transfer level (RTL) code for an adder module, in accordance with implementations herein.

[0033] FIG. 17 is a flow diagram illustrating an embodiment of a method for generating multiple distinct paths counterexamples.

[0034] FIG. 18 is a flow diagram illustrating an embodiment of a method for stimuli generation based on multiple distinct paths counterexamples.

DETAILED DESCRIPTION

[0035] A graphics processing unit (GPU) is communicatively coupled to host/processor cores to accelerate, for example, graphics operations, machine-learning operations, pattern analysis operations, and/or various general-purpose GPU (GPGPU) functions. The GPU may be communicatively coupled to the host processor/cores over a bus or another interconnect (e.g., a high-speed interconnect such as PCIe or NVLink). Alternatively, the GPU may be integrated on the same package or chip as the cores and communicatively coupled to the cores over an internal processor bus/interconnect (i.e., internal to the package or chip). Regardless of the manner in which the GPU is connected, the processor cores may allocate work to the GPU in the form of sequences of commands/instructions contained in a work descriptor. The GPU then uses dedicated circuitry/logic for efficiently processing these commands/instructions.

[0036] In the following description, numerous specific details are set forth to provide a more thorough understanding. However, it will be apparent to one of skill in the art that the embodiments described herein may be practiced without one or more of these specific details. In other instances, well-known features have not been described to avoid obscuring the details of the present embodiments.

System Overview

[0037] FIG. 1 is a block diagram of a processing system 100, according to an embodiment. Processing system 100 may be used in a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors 102 or processor cores 107. In one embodiment, the processing system 100 is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices such as within Internet-of-things (IoT) devices with wired or wireless connectivity to a local or wide area network.

[0038] In one embodiment, processing system 100 can include, couple with, or be integrated within: a server-based gaming platform; a game console, including a game and media console; a mobile gaming console, a handheld game console, or an online game console. In some embodiments the processing system 100 is part of a mobile phone, smart phone, tablet computing device or mobile Internet-connected device such as a laptop with low internal storage capacity. Processing system 100 can also include, couple with, or be integrated within: a wearable device, such as a smart watch wearable device; smart eyewear or clothing enhanced with augmented reality (AR) or virtual reality (VR) features to provide visual, audio or tactile outputs to supplement real world visual, audio or tactile experiences or otherwise provide text, audio, graphics, video, holographic images or video, or tactile feedback; other augmented reality (AR) device; or other virtual reality (VR) device. In some embodiments, the processing system 100 includes or is part of a television or set top box device. In one embodiment, processing system 100 can include, couple with, or be integrated within a self-driving vehicle such as a bus, tractor trailer, car, motor or electric power cycle, plane, or glider (or

any combination thereof). The self-driving vehicle may use processing system 100 to process the environment sensed around the vehicle.

[0039] In some embodiments, the one or more processors 102 each include one or more processor cores 107 to process instructions which, when executed, perform operations for system or user software. In some embodiments, at least one of the one or more processor cores 107 is configured to process a specific instruction set 109. In some embodiments, instruction set 109 may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). One or more processor cores 107 may process a different instruction set 109, which may include instructions to facilitate the emulation of other instruction sets. Processor core 107 may also include other processing devices, such as a Digital Signal Processor (DSP).

[0040] In some embodiments, the processor 102 includes cache memory 104. Depending on the architecture, the processor 102 can have a single internal cache or multiple levels of internal cache. In some embodiments, the cache memory is shared among various components of the processor 102. In some embodiments, the processor 102 also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores 107 using known cache coherency techniques. A register file 106 can be additionally included in processor 102 and may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). Some registers may be general-purpose registers, while other registers may be specific to the design of the processor 102.

[0041] In some embodiments, one or more processor(s) 102 are coupled with one or more interface bus(es) 110 to transmit communication signals such as address, data, or control signals between processor 102 and other components in the processing system 100. The interface bus 110, in one embodiment, can be a processor bus, such as a version of the Direct Media Interface (DMI) bus. However, processor busses are not limited to the DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI express), memory busses, or other types of interface busses. In one embodiment the processor(s) 102 include a memory controller 116 and a platform controller hub 130. The memory controller 116 facilitates communication between a memory device and other components of the processing system 100, while the platform controller hub (PCH) 130 provides connections to I/O devices via a local I/O bus.

[0042] The memory device 120 can be a dynamic random-access memory (DRAM) device, a static random-access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In one embodiment the memory device 120 can operate as system memory for the processing system 100, to store data 122 and instructions 121 for use when the one or more processors 102 executes an application or process. The memory controller 116 also couples with an optional external graphics processor 118, which may communicate with the one or more graphics processors 108 in processors 102 to perform graphics and media operations. In some embodiments, graphics, media, and or compute operations may be

assisted by an accelerator **112** which is a coprocessor that can be configured to perform a specialized set of graphics, media, or compute operations. For example, in one embodiment the accelerator **112** is a matrix multiplication accelerator used to optimize machine learning or compute operations. In one embodiment the accelerator **112** is a ray-tracing accelerator that can be used to perform ray-tracing operations in concert with the graphics processor **108**. In one embodiment, an external accelerator **119** may be used in place of or in concert with the accelerator **112**.

[0043] In some embodiments a display device **111** can connect to the processor(s) **102**. The display device **111** can be one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In one embodiment the display device **111** can be a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

[0044] In some embodiments the platform controller hub **130** enables peripherals to connect to memory device **120** and processor **102** via a high-speed I/O bus. The I/O peripherals include, but are not limited to, an audio controller **146**, a network controller **134**, a firmware interface **128**, a wireless transceiver **126**, touch sensors **125**, a data storage device **124** (e.g., non-volatile memory, volatile memory, hard disk drive, flash memory, NAND, 3D NAND, 3D XPoint, etc.). The data storage device **124** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI express). The touch sensors **125** can include touch screen sensors, pressure sensors, or fingerprint sensors. The wireless transceiver **126** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, 5G, or Long-Term Evolution (LTE) transceiver. The firmware interface **128** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). The network controller **134** can enable a network connection to a wired network. In some embodiments, a high-performance network controller (not shown) couples with the interface bus **110**. The audio controller **146**, in one embodiment, is a multi-channel high-definition audio controller. In one embodiment the processing system **100** includes an optional legacy I/O controller **140** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the system. The platform controller hub **130** can also connect to one or more Universal Serial Bus (USB) controllers **142** connect input devices, such as keyboard and mouse **143** combinations, a camera **144**, or other USB input devices.

[0045] It will be appreciated that the processing system **100** shown is example and not limiting, as other types of data processing systems that are differently configured may also be used. For example, an instance of the memory controller **116** and platform controller hub **130** may be integrated into a discreet external graphics processor, such as the external graphics processor **118**. In one embodiment the platform controller hub **130** and/or memory controller **116** may be external to the one or more processor(s) **102** and reside in a system chipset that is in communication with the processor (s) **102**.

[0046] For example, circuit boards (“sleds”) can be used on which components such as CPUs, memory, and other components are placed are designed for increased thermal

performance. In some examples, processing components such as the processors are located on a top side of a sled while near memory, such as DIMMs, are located on a bottom side of the sled. As a result of the enhanced airflow provided by this design, the components may operate at higher frequencies and power levels than in typical systems, thereby increasing performance. Furthermore, the sleds are configured to blindly mate with power and data communication cables in a rack, thereby enhancing their ability to be quickly removed, upgraded, reinstalled, and/or replaced. Similarly, individual components located on the sleds, such as processors, accelerators, memory, and data storage drives, are configured to be easily upgraded due to their increased spacing from each other. In the illustrative embodiment, the components additionally include hardware attestation features to prove their authenticity.

[0047] A data center can utilize a single network architecture (“fabric”) that supports multiple other network architectures including Ethernet and Omni-Path. The sleds can be coupled to switches via optical fibers, which provide higher bandwidth and lower latency than typical twisted pair cabling (e.g., Category 5, Category 5e, Category 6, etc.). Due to the high bandwidth, low latency interconnections and network architecture, the data center may, in use, pool resources, such as memory, accelerators (e.g., GPUs, graphics accelerators, FPGAs, ASICs, neural network and/or artificial intelligence accelerators, etc.), and data storage drives that are physically disaggregated, and provide them to compute resources (e.g., processors), enabling the compute resources to access the pooled resources as if they were local.

[0048] A power supply or source can provide voltage and/or current to processing system **100** or any component or system described herein. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based power supply, solar power supply, or fuel cell source.

[0049] FIGS. 2A-2D illustrate computing systems and graphics processors provided by embodiments described herein. The elements of FIGS. 2A-2D having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0050] FIG. 2A is a block diagram of an embodiment of a processor **200** having one or more processor cores **202A-202N**, an integrated memory controller **214**, and an integrated graphics processor **208**. Processor **200** can include additional cores up to and including additional core **202N** represented by the dashed lined boxes. Each of processor cores **202A-202N** includes one or more internal cache units **204A-204N**. In some embodiments each processor core also has access to one or more shared cache units **206**. The internal cache units **204A-204N** and shared cache units **206** represent a cache memory hierarchy within the processor **200**. The cache memory hierarchy may include at least one level of instruction and data cache within each processor

core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where the topmost level of cache before external memory is classified as the LLC. In some embodiments, cache coherency logic maintains coherency between the various cache units **206** and **204A-204N**.

[0051] In some embodiments, processor **200** may also include a set of one or more bus controller units **216** and a system agent core **210**. The one or more bus controller units **216** manage a set of peripheral buses, such as one or more PCI or PCI express busses. System agent core **210** provides management functionality for the various processor components. In some embodiments, system agent core **210** includes one or more integrated memory controllers **214** to manage access to various external memory devices (not shown).

[0052] In some embodiments, one or more of the processor cores **202A-202N** include support for simultaneous multi-threading. In such embodiment, the system agent core **210** includes components for coordinating and operating cores **202A-202N** during multi-threaded processing. System agent core **210** may additionally include a power control unit (PCU), which includes logic and components to regulate the power state of processor cores **202A-202N** and graphics processor **208**.

[0053] In some embodiments, processor **200** additionally includes graphics processor **208** to execute graphics processing operations. In some embodiments, the graphics processor **208** couples with the set of shared cache units **206**, and the system agent core **210**, including the one or more integrated memory controllers **214**. In some embodiments, the system agent core **210** also includes a display controller **211** to drive graphics processor output to one or more coupled displays. In some embodiments, display controller **211** may also be a separate module coupled with the graphics processor via at least one interconnect, or may be integrated within the graphics processor **208**.

[0054] In some embodiments, a ring-based interconnect **212** is used to couple the internal components of the processor **200**. However, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, a mesh interconnect, or other techniques, including techniques well known in the art. In some embodiments, graphics processor **208** couples with the ring-based interconnect **212** via an I/O link **213**.

[0055] The example I/O link **213** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **218**, such as an eDRAM module or a high-bandwidth memory (HBM) module. In some embodiments, each of the processor cores **202A-202N** and graphics processor **208** can use the embedded memory module **218** as a shared Last Level Cache.

[0056] In some embodiments, processor cores **202A-202N** are homogenous cores executing the same instruction set architecture. In another embodiment, processor cores **202A-202N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **202A-202N** execute a first instruction set, while at least one of the other cores executes a subset of the first instruction set or a different instruction set. In one embodiment, processor cores **202A-202N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher

power consumption couple with one or more power cores having a lower power consumption.

[0057] In one embodiment, processor cores **202A-202N** are heterogeneous in terms of computational capability. Additionally, processor **200** can be implemented on one or more chips or as an SoC integrated circuit having the illustrated components, in addition to other components.

[0058] FIG. 2B is a block diagram of hardware logic of a graphics processor core block **219**, according to some embodiments described herein. In some embodiments, elements of FIG. 2B having the same reference numbers (or names) as the elements of any other figure herein may operate or function in a manner similar to that described elsewhere herein. The graphics processor core block **219** is example of one partition of a graphics processor. The graphics processor core block **219** can be included within the integrated graphics processor **208** of FIG. 2A or a discrete graphics processor, parallel processor, and/or compute accelerator. A graphics processor as described herein may include multiple graphics core blocks based on target power and performance envelopes. Each graphics processor core block **219** can include a function block **230** coupled with multiple graphics cores **221A-221F** that include modular blocks of fixed function logic and general-purpose programmable logic. The graphics processor core block **219** also includes shared/cache memory **236** that is accessible by all graphics cores **221A-221F**, rasterizer logic **237**, and additional fixed function logic **238**.

[0059] In some embodiments, the function block **230** includes a geometry/fixed function pipeline **231** that can be shared by all graphics cores in the graphics processor core block **219**. In various embodiments, the geometry/fixed function pipeline **231** includes a 3D geometry pipeline a video front-end unit, a thread spawner and global thread dispatcher, and a unified return buffer manager, which manages unified return buffers. In one embodiment the function block **230** also includes a graphics SoC interface **232**, a graphics microcontroller **233**, and a media pipeline **234**. The graphics SoC interface **232** provides an interface between the graphics processor core block **219** and other core blocks within a graphics processor or compute accelerator SoC. The graphics microcontroller **233** is a programmable sub-processor that is configurable to manage various functions of the graphics processor core block **219**, including thread dispatch, scheduling, and pre-emption. The media pipeline **234** includes logic to facilitate the decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. The media pipeline **234** implement media operations via requests to compute or sampling logic within the graphics cores **221-221F**. One or more pixel backends **235** can also be included within the function block **230**. The pixel backends **235** include a cache memory to store pixel color values and can perform blend operations and lossless color compression of rendered pixel data.

[0060] In one embodiment the graphics SoC interface **232** enables the graphics processor core block **219** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC or a system host CPU that is coupled with the SoC via a peripheral interface. The graphics SoC interface **232** also enables communication with off-chip memory hierarchy elements such as a shared last level cache memory, system RAM, and/or embedded on-chip or on-package DRAM. The

SoC interface **232** can also enable communication with fixed function devices within the SoC, such as camera imaging pipelines, and enables the use of and/or implements global memory atomics that may be shared between the graphics processor core block **219** and CPUs within the SoC. The graphics SoC interface **232** can also implement power management controls for the graphics processor core block **219** and enable an interface between a clock domain of the graphics processor core block **219** and other clock domains within the SoC. In one embodiment the graphics SoC interface **232** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. The commands and instructions can be dispatched to the media pipeline **234** when media operations are to be performed, the geometry and fixed function pipeline **231** when graphics processing operations are to be performed. When compute operations are to be performed, compute dispatch logic can dispatch the commands to the graphics cores **221A-221F**, bypassing the geometry and media pipelines.

[0061] The graphics microcontroller **233** can be configured to perform various scheduling and management tasks for the graphics processor core block **219**. In one embodiment the graphics microcontroller **233** can perform graphics and/or compute workload scheduling on the various vector engines **222A-222F**, **224A-224F** and matrix engines **223A-223F**, **225A-225F** within the graphics cores **221A-221F**. In this scheduling model, host software executing on a CPU core of an SoC including the graphics processor core block **219** can submit workloads to one of multiple graphics processor doorbells, which invokes a scheduling operation on the appropriate graphics engine. Scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In one embodiment the graphics microcontroller **233** can also facilitate low-power or idle states for the graphics processor core block **219**, providing the graphics processor core block **219** with the ability to save and restore registers within the graphics processor core block **219** across low-power state transitions independently from the operating system and/or graphics driver software on the system.

[0062] The graphics processor core block **219** may have greater than or fewer than the illustrated graphics cores **221A-221F**, up to N modular graphics cores. For each set of N graphics cores, the graphics processor core block **219** can also include shared/cache memory **236**, which can be configured as shared memory or cache memory, rasterizer logic **237**, and additional fixed function logic **238** to accelerate various graphics and compute processing operations.

[0063] Within each graphics cores **221A-221F** is set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. The graphics cores **221A-221F** include multiple vector engines **222A-222F**, **224A-224F**, matrix acceleration units **223A-223F**, **225A-225D**, cache/shared local memory (SLM), a sampler **226A-226F**, and a ray tracing unit **227A-227F**.

[0064] The vector engines **222A-222F**, **224A-224F** are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation,

including graphics, media, or compute/GPGPU programs. The vector engines **222A-222F**, **224A-224F** can operate at variable vector widths using single instruction multiple data (SIMD), single instruction multiple thread (SIMT), or SIMT+SIMD execution modes. The matrix acceleration units **223A-223F**, **225A-225D** include matrix-matrix and matrix-vector acceleration logic that improves performance on matrix operations, particularly low and mixed precision (e.g., INT8, FP16, BF16) matrix operations used for machine learning. In one embodiment, each of the matrix acceleration units **223A-223F**, **225A-225D** includes one or more systolic arrays of processing elements that can perform concurrent matrix multiply or dot product operations on matrix elements.

[0065] The sampler **226A-226F** can read media or texture data into memory and can sample data differently based on a configured sampler state and the texture/media format that is being read. Threads executing on the vector engines **222A-222F**, **224A-224F** or matrix acceleration units **223A-223F**, **225A-225D** can make use of the cache/SLM **228A-228F** within each execution core. The cache/SLM **228A-228F** can be configured as cache memory or as a pool of shared memory that is local to each of the respective graphics cores **221A-221F**. The ray tracing units **227A-227F** within the graphics cores **221A-221F** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. In one embodiment the ray tracing units **227A-227F** include circuitry for performing depth testing and culling (e.g., using a depth buffer or similar arrangement). In one implementation, the ray tracing units **227A-227F** perform traversal and intersection operations in concert with image denoising, at least a portion of which may be performed using an associated matrix acceleration unit **223A-223F**, **225A-225D**.

[0066] FIG. 2C illustrates a graphics processing unit (GPU) **239** that includes dedicated sets of graphics processing resources arranged into multi-core groups **240A-240N**. The details of multi-core group **240A** are illustrated. Multi-core groups **240B-240N** may be equipped with the same or similar sets of graphics processing resources.

[0067] As illustrated, a multi-core group **240A** may include a set of graphics cores **243**, a set of tensor cores **244**, and a set of ray tracing cores **245**. A scheduler/dispatcher **241** schedules and dispatches the graphics threads for execution on the various cores **243**, **244**, **245**. In one embodiment the tensor cores **244** are sparse tensor cores with hardware to enable multiplication operations having a zero-value input to be bypassed. The graphics cores **243** of the GPU **239** of FIG. 2C differ in hierarchical abstraction level relative to the graphics cores **221A-221F** of FIG. 2B, which are analogous to the multi-core groups **240A-240N** of FIG. 2C. The graphics cores **243**, tensor cores **244**, and ray tracing cores **245** of FIG. 2C are analogous to, respectively, the vector engines **222A-222F**, **224A-224F**, matrix engines **223A-223F**, **225A-225F**, and ray tracing units **227A-227F** of FIG. 2B.

[0068] A set of register files **242** can store operand values used by the cores **243**, **244**, **245** when executing the graphics threads. These may include, for example, integer registers for storing integer values, floating point registers for storing floating point values, vector registers for storing packed data elements (integer and/or floating-point data elements) and

tile registers for storing tensor/matrix values. In one embodiment, the tile registers are implemented as combined sets of vector registers.

[0069] One or more combined level 1 (L1) caches and shared memory units **247** store graphics data such as texture data, vertex data, pixel data, ray data, bounding volume data, etc., locally within each multi-core group **240A**. One or more texture units **247** can also be used to perform texturing operations, such as texture mapping and sampling. A Level 2 (L2) cache **253** shared by all or a subset of the multi-core groups **240A-240N** stores graphics data and/or instructions for multiple concurrent graphics threads. As illustrated, the L2 cache **253** may be shared across a plurality of multi-core groups **240A-240N**. One or more memory controllers **248** couple the GPU **239** to a memory **249** which may be a system memory (e.g., DRAM) and/or a dedicated graphics memory (e.g., GDDR6 memory).

[0070] Input/output (I/O) circuitry **250** couples the GPU **239** to one or more I/O devices **252** such as digital signal processors (DSPs), network controllers, or user input devices. An on-chip interconnect may be used to couple the I/O devices **252** to the GPU **239** and memory **249**. One or more I/O memory management units (IOMMUs) **251** of the I/O circuitry **250** couple the I/O devices **252** directly to the memory **249**. In one embodiment, the IOMMU **251** manages multiple sets of page tables to map virtual addresses to physical addresses in memory **249**. In this embodiment, the I/O devices **252**, CPU(s) **246**, and GPU **239** may share the same virtual address space.

[0071] In one implementation, the IOMMU **251** supports virtualization. In this case, it may manage a first set of page tables to map guest/graphics virtual addresses to guest/graphics physical addresses and a second set of page tables to map the guest/graphics physical addresses to system/host physical addresses (e.g., within memory **249**). The base addresses of each of the first and second sets of page tables may be stored in control registers and swapped out on a context switch (e.g., so that the new context is provided with access to the relevant set of page tables). While not illustrated in FIG. 2C, each of the cores **243**, **244**, **245** and/or multi-core groups **240A-240N** may include translation lookaside buffers (TLBs) to cache guest virtual to guest physical translations, guest physical to host physical translations, and guest virtual to host physical translations.

[0072] In one embodiment, the CPUs **246**, GPU **239**, and I/O devices **252** are integrated on a single semiconductor chip and/or chip package. The memory **249** may be integrated on the same chip or may be coupled to the memory controllers **248** via an off-chip interface. In one implementation, the memory **249** comprises GDDR6 memory which shares the same virtual address space as other physical system-level memories, although the underlying principles of the embodiments described herein are not limited to this specific implementation.

[0073] In one embodiment, the tensor cores **244** include a plurality of functional units specifically designed to perform matrix operations, which are the basic compute operation used to perform deep learning operations. For example, simultaneous matrix multiplication operations may be used for neural network training and inferencing. The tensor cores **244** may perform matrix processing using a variety of operand precisions including single precision floating-point (e.g., 32 bits), half-precision floating point (e.g., 16 bits), integer words (16 bits), bytes (8 bits), and half-bytes (4 bits).

In one embodiment, a neural network implementation extracts features of each rendered scene, potentially combining details from multiple frames, to construct a high-quality final image.

[0074] In deep learning implementations, parallel matrix multiplication work may be scheduled for execution on the tensor cores **244**. The training of neural networks, in particular, utilizes a significant number of matrix dot product operations. In order to process an inner-product formulation of an $N \times N \times N$ matrix multiply, the tensor cores **244** may include at least N dot-product processing elements. Before the matrix multiply begins, one entire matrix is loaded into tile registers and at least one column of a second matrix is loaded each cycle for N cycles. Each cycle, there are N dot products that are processed.

[0075] Matrix elements may be stored at different precisions depending on the particular implementation, including 16-bit words, 8-bit bytes (e.g., INT8) and 4-bit half-bytes (e.g., INT4). Different precision modes may be specified for the tensor cores **244** to ensure that the most efficient precision is used for different workloads (e.g., such as inferencing workloads which can tolerate quantization to bytes and half-bytes).

[0076] In one embodiment, the ray tracing cores **245** accelerate ray tracing operations for both real-time ray tracing and non-real-time ray tracing implementations. In particular, the ray tracing cores **245** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. The ray tracing cores **245** may also include circuitry for performing depth testing and culling (e.g., using a Z buffer or similar arrangement). In one implementation, the ray tracing cores **245** perform traversal and intersection operations in concert with the image denoising techniques described herein, at least a portion of which may be executed on the tensor cores **244**. For example, in one embodiment, the tensor cores **244** implement a deep learning neural network to perform denoising of frames generated by the ray tracing cores **245**. However, the CPU(s) **246**, graphics cores **243**, and/or ray tracing cores **245** may also implement all or a portion of the denoising and/or deep learning algorithms.

[0077] In addition, as described above, a distributed approach to denoising may be employed in which the GPU **239** is in a computing device coupled to other computing devices over a network or high-speed interconnect. In this embodiment, the interconnected computing devices share neural network learning/training data to improve the speed with which the overall system learns to perform denoising for different types of image frames and/or different graphics applications.

[0078] In one embodiment, the ray tracing cores **245** process all BVH traversal and ray-primitive intersections, saving the graphics cores **243** from being overloaded with thousands of instructions per ray. In one embodiment, each ray tracing core **245** includes a first set of specialized circuitry for performing bounding box tests (e.g., for traversal operations) and a second set of specialized circuitry for performing the ray-triangle intersection tests (e.g., intersecting rays which have been traversed). Thus, in one embodiment, the multi-core group **240A** can simply launch a ray probe, and the ray tracing cores **245** independently perform ray traversal and intersection and return hit data (e.g., a hit, no hit, multiple hits, etc.) to the thread context.

The other cores **243**, **244** are freed to perform other graphics or compute work while the ray tracing cores **245** perform the traversal and intersection operations.

[0079] In one embodiment, each ray tracing core **245** includes a traversal unit to perform BVH testing operations and an intersection unit which performs ray-primitive intersection tests. The intersection unit generates a “hit”, “no hit”, or “multiple hit” response, which it provides to the appropriate thread. During the traversal and intersection operations, the execution resources of the other cores (e.g., graphics cores **243** and tensor cores **244**) are freed to perform other forms of graphics work.

[0080] In one particular embodiment described below, a hybrid rasterization/ray tracing approach is used in which work is distributed between the graphics cores **243** and ray tracing cores **245**.

[0081] In one embodiment, the ray tracing cores **245** (and/or other cores **243**, **244**) include hardware support for a ray tracing instruction set such as Microsoft’s DirectX Ray Tracing (DXR) which includes a DispatchRays command, as well as ray-generation, closest-hit, any-hit, and miss shaders, which enable the assignment of sets of shaders and textures for each object. Another ray tracing platform which may be supported by the ray tracing cores **245**, graphics cores **243** and tensor cores **244** is Vulkan 1.1.85. Note, however, that the underlying principles of the embodiments described herein are not limited to any particular ray tracing ISA.

[0082] In general, the various cores **245**, **244**, **243** may support a ray tracing instruction set that includes instructions/functions for ray generation, closest hit, any hit, ray-primitive intersection, per-primitive and hierarchical bounding box construction, miss, visit, and exceptions. More specifically, one embodiment includes ray tracing instructions to perform the following functions:

- [0083]** Ray Generation—Ray generation instructions may be executed for each pixel, sample, or other user-defined work assignment.
- [0084]** Closest Hit—A closest hit instruction may be executed to locate the closest intersection point of a ray with primitives within a scene.
- [0085]** Any Hit—An any hit instruction identifies multiple intersections between a ray and primitives within a scene, potentially to identify a new closest intersection point.
- [0086]** Intersection—An intersection instruction performs a ray-primitive intersection test and outputs a result.
- [0087]** Per-primitive Bounding box Construction—This instruction builds a bounding box around a given primitive or group of primitives (e.g., when building a new BVH or other acceleration data structure).
- [0088]** Miss—Indicates that a ray misses all geometry within a scene, or specified region of a scene.
- [0089]** Visit—Indicates the child volumes a ray will traverse.
- [0090]** Exceptions—Includes various types of exception handlers (e.g., invoked for various error conditions).

[0091] In one embodiment the ray tracing cores **245** may be adapted to accelerate general-purpose compute operations that can be accelerated using computational techniques that are analogous to ray intersection tests. A compute framework can be provided that enables shader programs to be compiled into low level instructions and/or primitives

that perform general-purpose compute operations via the ray tracing cores. Example computational problems that can benefit from compute operations performed on the ray tracing cores **245** include computations involving beam, wave, ray, or particle propagation within a coordinate space. Interactions associated with that propagation can be computed relative to a geometry or mesh within the coordinate space. For example, computations associated with electromagnetic signal propagation through an environment can be accelerated via the use of instructions or primitives that are executed via the ray tracing cores. Diffraction and reflection of the signals by objects in the environment can be computed as direct ray-tracing analogies.

[0092] Ray tracing cores **245** can also be used to perform computations that are not directly analogous to ray tracing. For example, mesh projection, mesh refinement, and volume sampling computations can be accelerated using the ray tracing cores **245**. Generic coordinate space calculations, such as nearest neighbor calculations can also be performed. For example, the set of points near a given point can be discovered by defining a bounding box in the coordinate space around the point. BVH and ray probe logic within the ray tracing cores **245** can then be used to determine the set of point intersections within the bounding box. The intersections constitute the origin point and the nearest neighbors to that origin point. Computations that are performed using the ray tracing cores **245** can be performed in parallel with computations performed on the graphics cores **243** and tensor cores **244**. A shader compiler can be configured to compile a compute shader or other general-purpose graphics processing program into low level primitives that can be parallelized across the graphics cores **243**, tensor cores **244**, and ray tracing cores **245**.

[0093] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) **270** that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein. The GPGPU **270** can interconnect with host processors (e.g., one or more CPU(s) **246**) and memory **271**, **272** via one or more system and/or memory busses. In one embodiment the memory **271** is system memory that may be shared with the one or more CPU(s) **246**, while memory **272** is device memory that is dedicated to the GPGPU **270**. In one embodiment, components within the GPGPU **270** and memory **272** may be mapped into memory addresses that are accessible to the one or more CPU(s) **246**. Access to memory **271** and **272** may be facilitated via a memory controller **268**. In one embodiment the memory controller **268** includes an internal direct memory access (DMA) controller **269** or can include logic to perform operations that would otherwise be performed by a DMA controller.

[0094] The GPGPU **270** includes multiple cache memories, including an L2 cache **253**, L1 cache **254**, an instruction cache **255**, and shared memory **256**, at least a portion of which may also be partitioned as a cache memory. The GPGPU **270** also includes multiple compute units **260A-260N**, which represent a hierarchical abstraction level analogous to the graphics cores **221A-221F** of FIG. 2B and the multi-core groups **240A-240N** of FIG. 2C. Each compute unit **260A-260N** includes a set of vector registers **261**, scalar registers **262**, vector logic units **263**, and scalar logic units **264**. The compute units **260A-260N** can also include local shared memory **265** and a program counter **266**. The compute units **260A-260N** can couple with a constant cache **267**,

which can be used to store constant data, which is data that will not change during the run of kernel or shader program that executes on the GPGPU 270. In one embodiment the constant cache 267 is a scalar data cache and cached data can be fetched directly into the scalar registers 262.

[0095] During operation, the one or more CPU(s) 246 can write commands into registers or memory in the GPGPU 270 that has been mapped into an accessible address space. The command processors 257 can read the commands from registers or memory and determine how those commands will be processed within the GPGPU 270. A thread dispatcher 258 can then be used to dispatch threads to the compute units 260A-260N to perform those commands. Each compute unit 260A-260N can execute threads independently of the other compute units. Additionally, each compute unit 260A-260N can be independently configured for conditional computation and can conditionally output the results of computation to memory. The command processors 257 can interrupt the one or more CPU(s) 246 when the submitted commands are complete.

[0096] FIGS. 3A-3C illustrate block diagrams of additional graphics processor and compute accelerator architectures provided by embodiments described herein. The elements of FIGS. 3A-3C having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0097] FIG. 3A is a block diagram of a graphics processor 300, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces. In some embodiments, the graphics processor communicates via a memory mapped I/O interface to registers on the graphics processor and with commands placed into the processor memory. In some embodiments, graphics processor 300 includes a memory interface 314 to access memory. Memory interface 314 can be an interface to local memory, one or more internal caches, one or more shared external caches, and/or to system memory.

[0098] In some embodiments, graphics processor 300 also includes a display controller 302 to drive display output data to a display device 318. Display controller 302 includes hardware for one or more overlay planes for the display and composition of multiple layers of video or user interface elements. The display device 318 can be an internal or external display device. In one embodiment the display device 318 is a head mounted display device, such as a virtual reality (VR) display device or an augmented reality (AR) display device. In some embodiments, graphics processor 300 includes a video codec engine 306 to encode, decode, or transcode media to, from, or between one or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, H.265/HEVC, Alliance for Open Media (AOMedia) VP8, VP9, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

[0099] In some embodiments, graphics processor 300 includes a block image transfer (BLIT) engine to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in one

embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE) 310. In some embodiments, GPE 310 is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

[0100] In some embodiments, GPE 310 includes a 3D pipeline 312 for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). The 3D pipeline 312 includes programmable and fixed function elements that perform various tasks within the element and/or spawn execution threads to a 3D/Media subsystem 315. While 3D pipeline 312 can be used to perform media operations, an embodiment of GPE 310 also includes a media pipeline 316 that is specifically used to perform media operations, such as video post-processing and image enhancement.

[0101] In some embodiments, media pipeline 316 includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration, video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine 306. In some embodiments, media pipeline 316 additionally includes a thread spawning unit to spawn threads for execution on 3D/Media subsystem 315. The spawned threads perform computations for the media operations on one or more graphics cores included in 3D/Media subsystem 315.

[0102] In some embodiments, 3D/Media subsystem 315 includes logic for executing threads spawned by 3D pipeline 312 and media pipeline 316. In one embodiment, the pipelines send thread execution requests to 3D/Media subsystem 315, which includes thread dispatch logic for arbitrating and dispatching the various requests to available thread execution resources. The execution resources include an array of graphics cores to process the 3D and media threads. In some embodiments, 3D/Media subsystem 315 includes one or more internal caches for thread instructions and data. In some embodiments, the subsystem also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

[0103] FIG. 3B illustrates a graphics processor 320 having a tiled architecture, according to embodiments described herein. In one embodiment the graphics processor 320 includes a graphics processing engine cluster 322 having multiple instances of the graphics processing engine 310 of FIG. 3A within a graphics engine tile 310A-310D. Each graphics engine tile 310A-310D can be interconnected via a set of tile interconnects 323A-323F. Each graphics engine tile 310A-310D can also be connected to a memory module or memory device 326A-326D via memory interconnects 325A-325D. The memory devices 326A-326D can use any graphics memory technology. For example, the memory devices 326A-326D may be graphics double data rate (GDDR) memory. The memory devices 326A-326D, in one embodiment, are HBM modules that can be on-die with their respective graphics engine tile 310A-310D. In one embodiment the memory devices 326A-326D are stacked memory devices that can be stacked on top of their respective graphics engine tile 310A-310D. In one embodiment, each graphics engine tile 310A-310D and associated memory 326A-326D reside on separate chiplets, which are bonded to a base die or base substrate, as described on further detail in FIGS. 11B-11D.

[0104] The graphics processor 320 may be configured with a non-uniform memory access (NUMA) system in which memory devices 326A-326D are coupled with associated graphics engine tiles 310A-310D. A given memory device may be accessed by graphics engine tiles other than the tile to which it is directly connected. However, access latency to the memory devices 326A-326D may be lowest when accessing a local tile. In one embodiment, a cache coherent NUMA (ccNUMA) system is enabled that uses the tile interconnects 323A-323F to enable communication between cache controllers within the graphics engine tiles 310A-310D to maintain a consistent memory image when more than one cache stores the same memory location.

[0105] The graphics processing engine cluster 322 can connect with an on-chip or on-package fabric interconnect 324. In one embodiment the fabric interconnect 324 includes a network processor, network on a chip (NoC), or another switching processor to enable the fabric interconnect 324 to act as a packet switched fabric interconnect that switches data packets between components of the graphics processor 320. The fabric interconnect 324 can enable communication between graphics engine tiles 310A-310D and components such as the video codec engine 306 and one or more copy engines 304. The copy engines 304 can be used to move data out of, into, and between the memory devices 326A-326D and memory that is external to the graphics processor 320 (e.g., system memory). The fabric interconnect 324 can also couple with one or more of the tile interconnects 323A-323F to facilitate or enhance the interconnection between the graphics engine tiles 310A-310D. The fabric interconnect 324 is also configurable to interconnect multiple instances of the graphics processor 320 (e.g., via the host interface 328), enabling tile-to-tile communication between graphics engine tiles 310A-310D of multiple GPUs. In one embodiment, the graphics engine tiles 310A-310D of multiple GPUs can be presented to a host system as a single logical device.

[0106] The graphics processor 320 may optionally include a display controller 302 to enable a connection with the display device 318. The graphics processor may also be configured as a graphics or compute accelerator. In the accelerator configuration, the display controller 302 and display device 318 may be omitted.

[0107] The graphics processor 320 can connect to a host system via a host interface 328. The host interface 328 can enable communication between the graphics processor 320, system memory, and/or other system components. The host interface 328 can be, for example a PCI express bus or another type of host system interface. For example, the host interface 328 may be an NVLink or NVSwitch interface. The host interface 328 and fabric interconnect 324 can cooperate to enable multiple instances of the graphics processor 320 to act as single logical device. Cooperation between the host interface 328 and fabric interconnect 324 can also enable the individual graphics engine tiles 310A-310D to be presented to the host system as distinct logical graphics devices.

[0108] FIG. 3C illustrates a compute accelerator 330, according to embodiments described herein. The compute accelerator 330 can include architectural similarities with the graphics processor 320 of FIG. 3B and is optimized for compute acceleration. A compute engine cluster 332 can include a set of compute engine tiles 340A-340D that include execution logic that is optimized for parallel or

vector-based general-purpose compute operations. In some embodiments, the compute engine tiles 340A-340D do not include fixed function graphics processing logic, although in one embodiment one or more of the compute engine tiles 340A-340D can include logic to perform media acceleration. The compute engine tiles 340A-340D can connect to memory 326A-326D via memory interconnects 325A-325D. The memory 326A-326D and memory interconnects 325A-325D may be similar technology as in graphics processor 320 or can be different. The compute engine tiles 340A-340D can also be interconnected via a set of tile interconnects 323A-323F and may be connected with and/or interconnected by a fabric interconnect 324. Cross-tile communications can be facilitated via the fabric interconnect 324. The fabric interconnect 324 (e.g., via the host interface 328) can also facilitate communication between compute engine tiles 340A-340D of multiple instances of the compute accelerator 330. In one embodiment the compute accelerator 330 includes a large L3 cache 336 that can be configured as a device-wide cache. The compute accelerator 330 can also connect to a host processor and memory via a host interface 328 in a similar manner as the graphics processor 320 of FIG. 3B.

[0109] The compute accelerator 330 can also include an integrated network interface 342. In one embodiment the network interface 342 includes a network processor and controller logic that enables the compute engine cluster 332 to communicate over a physical layer interconnect 344 without utilizing data to traverse memory of a host system. In one embodiment, one of the compute engine tiles 340A-340D is replaced by network processor logic and data to be transmitted or received via the physical layer interconnect 344 may be transmitted directly to or from memory 326A-326D. Multiple instances of the compute accelerator 330 may be joined via the physical layer interconnect 344 into a single logical device. Alternatively, the various compute engine tiles 340A-340D may be presented as distinct network accessible compute accelerator devices.

Graphics Processing Engine

[0110] FIG. 4 is a block diagram of a graphics processing engine 410 of a graphics processor in accordance with some embodiments. In one embodiment, the graphics processing engine (GPE) 410 is a version of the GPE 310 shown in FIG. 3A and may also represent a graphics engine tile 310A-310D of FIG. 3B. Elements of FIG. 4 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. For example, the 3D pipeline 312 and media pipeline 316 of FIG. 3A are illustrated. The media pipeline 316 is optional in some embodiments of the GPE 410 and may not be explicitly included within the GPE 410. For example and in at least one embodiment, a separate media and/or image processor is coupled to the GPE 410.

[0111] In some embodiments, GPE 410 couples with or includes a command streamer 403, which provides a command stream to the 3D pipeline 312 and/or media pipelines 316. Alternatively or additionally, the command streamer 403 may be directly coupled to a unified return buffer 418. The unified return buffer 418 may be communicatively coupled to a graphics core cluster 414. In some embodiments, command streamer 403 is coupled with memory, which can be system memory, or one or more of internal

cache memory and shared cache memory. In some embodiments, command streamer **403** receives commands from the memory and sends the commands to 3D pipeline **312** and/or media pipeline **316**. The commands are directives fetched from a ring buffer, which stores commands for the 3D pipeline **312** and media pipeline **316**. In one embodiment, the ring buffer can additionally include batch command buffers storing batches of multiple commands. The commands for the 3D pipeline **312** can also include references to data stored in memory, such as but not limited to vertex and geometry data for the 3D pipeline **312** and/or image data and memory objects for the media pipeline **316**. The 3D pipeline **312** and media pipeline **316** process the commands and data by performing operations via logic within the respective pipelines or by dispatching one or more execution threads to a graphics core cluster **414**. In one embodiment the graphics core cluster **414** include one or more blocks of graphics cores (e.g., graphics core block **415A**, graphics core block **415B**), each block including one or more graphics cores. Each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic, such as matrix or AI acceleration logic.

[0112] In various embodiments the 3D pipeline **312** can include fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel shaders, fragment shaders, compute shaders, or other shader and/or GPGPU programs, by processing the instructions and dispatching execution threads to the graphics core cluster **414**. The graphics core cluster **414** provides a unified block of execution resources for use in processing these shader programs. Multi-purpose execution logic within the graphics core blocks **415A-415B** of the graphics core cluster **414** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

[0113] In some embodiments, the graphics core cluster **414** includes execution logic to perform media functions, such as video and/or image processing. In one embodiment, the graphics cores include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations. The general-purpose logic can perform processing operations in parallel or in conjunction with general-purpose logic within the processor core(s) **107** of FIG. 1 or core **202A-202N** as in FIG. 2A.

[0114] Output data generated by threads executing on the graphics core cluster **414** can output data to memory in a unified return buffer (URB) **418**. The URB **418** can store data for multiple threads. In some embodiments the URB **418** may be used to send data between different threads executing on the graphics core cluster **414**. In some embodiments the URB **418** may additionally be used for synchronization between threads on the graphics core array and fixed function logic within the shared function logic **420**.

[0115] In some embodiments, graphics core cluster **414** is scalable, such that the cluster includes a variable number of graphics cores, each having a variable number of graphics cores based on the target power and performance level of GPE **410**. In one embodiment the execution resources are dynamically scalable, such that execution resources may be enabled or disabled.

[0116] The graphics core cluster **414** couples with shared function logic **420** that includes multiple resources that are shared between the graphics cores in the graphics core array. The shared functions within the shared function logic **420** are hardware logic units that provide specialized supplemental functionality to the graphics core cluster **414**. In various embodiments, shared function logic **420** may include, but is not limited to sampler **421**, math **422**, and inter-thread communication (ITC) **423** logic. Additionally, some embodiments implement one or more cache(s) **425** within the shared function logic **420**. The shared function logic **420** can implement the same or similar functionality as the additional fixed function logic **238** of FIG. 2B.

[0117] A shared function is implemented at least in a case where the demand for a given specialized function is insufficient for inclusion within the graphics core cluster **414**. Instead, a single instantiation of that specialized function is implemented as a stand-alone entity in the shared function logic **420** and shared among the execution resources within the graphics core cluster **414**. The precise set of functions that are shared between the graphics core cluster **414** and included within the graphics core cluster **414** varies across embodiments. In some embodiments, specific shared functions within the shared function logic **420** that are used extensively by the graphics core cluster **414** may be included within shared function logic **416** within the graphics core cluster **414**. In various embodiments, the shared function logic **416** within the graphics core cluster **414** can include some or all logic within the shared function logic **420**. In one embodiment, all logic elements within the shared function logic **420** may be duplicated within the shared function logic **416** of the graphics core cluster **414**. In one embodiment the shared function logic **420** is excluded in favor of the shared function logic **416** within the graphics core cluster **414**.

Graphics Processing Resources

[0118] FIG. 5A-5C illustrate execution logic including an array of processing elements employed in a graphics processor, according to embodiments described herein. FIG. 5A illustrates graphics core cluster, according to an embodiment. FIG. 5B illustrates a vector engine of a graphics core, according to an embodiment. FIG. 5C illustrates a matrix engine of a graphics core, according to an embodiment. Elements of FIG. 5A-5C having the same reference numbers as the elements of any other figure herein may operate or function in any manner similar to that described elsewhere herein, but are not limited as such. For example, the elements of FIG. 5A-5C can be considered in the context of the graphics processor core block **219** of FIG. 2B, and/or the graphics core blocks **415A-415B** of FIG. 4. In one embodiment, the elements of FIG. 5A-5C have similar functionality to equivalent components of the graphics processor **208** of FIG. 2A, the GPU **239** of FIG. 2C or the GPGPU **270** of FIG. 2D.

[0119] As shown in FIG. 5A, in one embodiment the graphics core cluster **414** includes a graphics core block **415**, which may be graphics core block **415A** or graphics core block **415B** of FIG. 4. The graphics core block **415** can include any number of graphics cores (e.g., graphics core **515A**, graphics core **515B**, through graphics core **515N**). Multiple instances of the graphics core block **415** may be included. In one embodiment the elements of the graphics cores **515A-515N** have similar or equivalent functionality as the elements of the graphics cores **221A-221F** of FIG. 2B.

In such embodiment, the graphics cores **515A-515N** each include circuitry including but not limited to vector engines **502A-502N**, matrix engines **503A-503N**, memory load/store units **504A-504N**, instruction caches **505A-505N**, data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, samplers **510A-2710N**. The circuitry of the graphics cores **515A-515N** can additionally include fixed function logic **512A-512N**. The number of vector engines **502A-502N** and matrix engines **503A-503N** within the graphics cores **515A-515N** of a design can vary based on the workload, performance, and power targets for the design.

[0120] With reference to graphics core **515A**, the vector engine **502A** and matrix engine **503A** are configurable to perform parallel compute operations on data in a variety of integer and floating-point data formats based on instructions associated with shader programs. Each vector engine **502A** and matrix engine **503A** can act as a programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. The vector engine **502A** and matrix engine **503A** support the processing of variable width vectors at various SIMD widths, including but not limited to SIMD8, SIMD16, and SIMD32. Input data elements can be stored as a packed data type in a register and the vector engine **502A** and matrix engine **503A** can process the various elements based on the data size of the elements. For example, when operating on a 256-bit wide vector, the 256 bits of the vector are stored in a register and the vector is processed as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, different vector widths and register sizes are possible. In one embodiment, the vector engine **502A** and matrix engine **503A** are also configurable for SIMT operation on warps or thread groups of various sizes (e.g., 8, 16, or 32 threads).

[0121] Continuing with graphics core **515A**, the memory load/store unit **504A** services memory access requests that are issued by the vector engine **502A**, matrix engine **503A**, and/or other components of the graphics core **515A** that have access to memory. The memory access request can be processed by the memory load/store unit **504A** to load or store the requested data to or from cache or memory into a register file associated with the vector engine **502A** and/or matrix engine **503A**. The memory load/store unit **504A** can also perform prefetching operations. In one embodiment, the memory load/store unit **504A** is configured to provide SIMT scatter/gather prefetching or block prefetching for data stored in memory **610**, from memory that is local to other tiles via the tile interconnect **608**, or from system memory. Prefetching can be performed to a specific L1 cache (e.g., data cache/shared local memory **506A**), the L2 cache **604** or the L3 cache **606**. In one embodiment, a prefetch to the L3 cache **606** automatically results in the data being stored in the L2 cache **604**.

[0122] The instruction cache **505A** stores instructions to be executed by the graphics core **515A**. In one embodiment, the graphics core **515A** also includes instruction fetch and prefetch circuitry that fetches or prefetches instructions into the instruction cache **505A**. The graphics core **515A** also includes instruction decode logic to decode instructions within the instruction cache **505A**. The data cache/shared

local memory **506A** can be configured as a data cache that is managed by a cache controller that implements a cache replacement policy and/or configured as explicitly managed shared memory. The ray tracing unit **508A** includes circuitry to accelerate ray tracing operations. The sampler **510A** provides texture sampling for 3D operations and media sampling for media operations. The fixed function logic **512A** includes fixed function circuitry that is shared between the various instances of the vector engine **502A** and matrix engine **503A**. Graphics cores **515B-515N** can operate in a similar manner as graphics core **515A**.

[0123] Functionality of the instruction caches **505A-505N**, data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, samplers **510A-2710N**, and fixed function logic **512A-512N** corresponds with equivalent functionality in the graphics processor architectures described herein. For example, the instruction caches **505A-505N** can operate in a similar manner as instruction cache **255** of FIG. 2D. The data caches/shared local memory **506A-506N**, ray tracing units **508A-508N**, and samplers **510A-2710N** can operate in a similar manner as the cache/SLM **228A-228F**, ray tracing units **227A-227F**, and samplers **226A-226F** of FIG. 2B. The fixed function logic **512A-512N** can include elements of the geometry/fixed function pipeline **231** and/or additional fixed function logic **238** of FIG. 2B. In one embodiment, the ray tracing units **508A-508N** include circuitry to perform ray tracing acceleration operations performed by the ray tracing cores **245** of FIG. 2C.

[0124] As shown in FIG. 5B, in one embodiment the vector engine **502** includes an instruction fetch unit **537**, a general register file array (GRF) **524**, an architectural register file array (ARF) **526**, a thread arbiter **522**, a send unit **530**, a branch unit **532**, a set of SIMD floating point units (FPUs) **534**, and in one embodiment a set of integer SIMD ALUs **535**. The GRF **524** and ARF **526** includes the set of general register files and architecture register files associated with each hardware thread that may be active in the vector engine **502**. In one embodiment, per thread architectural state is maintained in the ARF **526**, while data used during thread execution is stored in the GRF **524**. The execution state of each thread, including the instruction pointers for each thread, can be held in thread-specific registers in the ARF **526**.

[0125] In one embodiment the vector engine **502** has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). The architecture has a modular configuration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per graphics core, where graphics core resources are divided across logic used to execute multiple simultaneous threads. The number of logical threads that may be executed by the vector engine **502** is not limited to the number of hardware threads, and multiple logical threads can be assigned to each hardware thread.

[0126] In one embodiment, the vector engine **502** can co-issue multiple instructions, which may each be different instructions. The thread arbiter **522** can dispatch the instructions to one of the send unit **530**, branch unit **532**, or SIMD FPU(s) **534** for execution. Each execution thread can access **128** general-purpose registers within the GRF **524**, where each register can store 32 bytes, accessible as a variable width vector of 32-bit data elements. In one embodiment,

each thread has access to 4 Kbytes within the GRF 524, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In one embodiment the vector engine 502 is partitioned into seven hardware threads that can independently perform computational operations, although the number of threads per vector engine 502 can also vary according to embodiments. For example, in one embodiment up to 16 hardware threads are supported. In an embodiment in which seven threads may access 4 Kbytes, the GRF 524 can store a total of 28 Kbytes. Where 16 threads may access 4 Kbytes, the GRF 524 can store a total of 64 Kbytes. Flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

[0127] In one embodiment, memory operations, sampler operations, and other longer-latency system communications are dispatched via “send” instructions that are executed by the message passing send unit 530. In one embodiment, branch instructions are dispatched to a dedicated branch unit 532 to facilitate SIMD divergence and eventual convergence.

[0128] In one embodiment the vector engine 502 includes one or more SIMD floating point units (FPU(s)) 534 to perform floating-point operations. In one embodiment, the FPU(s) 534 also support integer computation. In one embodiment the FPU(s) 534 can execute up to M number of 32-bit floating-point (or integer) operations, or execute up to 2M 16-bit integer or 16-bit floating-point operations. In one embodiment, at least one of the FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In some embodiments, a set of 8-bit integer SIMD ALUs 535 are also present and may be specifically optimized to perform operations associated with machine learning computations. In one embodiment, the SIMD ALUs are replaced by an additional set of SIMD FPUs 534 that are configurable to perform integer and floating-point operations. In one embodiment, the SIMD FPUs 534 and SIMD ALUs 535 are configurable to execute SIMT programs. In one embodiment, combined SIMD+SIMT operation is supported.

[0129] In one embodiment, arrays of multiple instances of the vector engine 502 can be instantiated in a graphics core. For scalability, product architects can choose the exact number of vector engines per graphics core grouping. In one embodiment the vector engine 502 can execute instructions across a plurality of execution channels. In a further embodiment, each thread executed on the vector engine 502 is executed on a different channel.

[0130] As shown in FIG. 5C, in one embodiment the matrix engine 503 includes an array of processing elements that are configured to perform tensor operations including vector/matrix and matrix/matrix operations, such as but not limited to matrix multiply and/or dot product operations. The matrix engine 503 is configured with M rows and N columns of processing elements (PE 552AA-PE 552MN) that include multiplier and adder circuits organized in a pipelined fashion. In one embodiment, the processing elements 552AA-PE 552MN make up the physical pipeline stages of an N wide and M deep systolic array that can be used to perform vector/matrix or matrix/matrix operations in a data-parallel manner, including matrix multiply, fused multiply-add, dot product or other general matrix-matrix multiplication (GEMM) operations. In one embodiment the

matrix engine 503 supports 16-bit floating point operations, as well as 8-bit, 4-bit, 2-bit, and binary integer operations. The matrix engine 503 can also be configured to accelerate specific machine learning operations. In such embodiments, the matrix engine 503 can be configured with support for the bfloat (brain floating point) 16-bit floating point format or a tensor float 32-bit floating point format (TF32) that have different numbers of mantissa and exponent bits relative to Institute of Electrical and Electronics Engineers (IEEE) 754 formats.

[0131] In one embodiment, during each cycle, each stage can add the result of operations performed at that stage to the output of the previous stage. In other embodiments, the pattern of data movement between the processing elements 552AA-552MN after a set of computational cycles can vary based on the instruction or macro-operation being performed. For example, in one embodiment partial sum loopback is enabled and the processing elements may instead add the output of a current cycle with output generated in the previous cycle. In one embodiment, the final stage of the systolic array can be configured with a loopback to the initial stage of the systolic array. In such embodiment, the number of physical pipeline stages may be decoupled from the number of logical pipeline stages that are supported by the matrix engine 503. For example, where the processing elements 552AA-552MN are configured as a systolic array of M physical stages, a loopback from stage M to the initial pipeline stage can enable the processing elements 552AA-PE552MN to operate as a systolic array of, for example, 2M, 3M, 4M, etc., logical pipeline stages.

[0132] In one embodiment, the matrix engine 503 includes memory 541A-541N, 542A-542M to store input data in the form of row and column data for input matrices. Memory 542A-542M is configurable to store row elements (A0-Am) of a first input matrix and memory 541A-541N is configurable to store column elements (B0-Bn) of a second input matrix. The row and column elements are provided as input to the processing elements 552AA-552MN for processing. In one embodiment, row and column elements of the input matrices can be stored in a systolic register file 540 within the matrix engine 503 before those elements are provided to the memory 541A-541N, 542A-542M. In one embodiment, the systolic register file 540 is excluded and the memory 541A-541N, 542A-542M is loaded from registers in an associated vector engine (e.g., GRF 524 of vector engine 502 of FIG. 5B) or other memory of the graphics core that includes the matrix engine 503 (e.g., data cache/shared local memory 506A for matrix engine 503A of FIG. 5A). Results generated by the processing elements 552AA-552MN are then output to an output buffer and/or written to a register file (e.g., systolic register file 540, GRF 524, data cache/shared local memory 506A-506N) for further processing by other functional units of the graphics processor or for output to memory.

[0133] In some embodiments, the matrix engine 503 is configured with support for input sparsity, where multiplication operations for sparse regions of input data can be bypassed by skipping multiply operations that have a zero-value operand. In one embodiment, the processing elements 552AA-552MN are configured to skip the performance of certain operations that have zero value input. In one embodiment, sparsity within input matrices can be detected and operations having known zero output values can be bypassed before being submitted to the processing elements

552AA-552MN. The loading of zero value operands into the processing elements can be bypassed and the processing elements **552AA-552MN** can be configured to perform multiplications on the non-zero value input elements. The matrix engine **503** can also be configured with support for output sparsity, such that operations with results that are pre-determined to be zero are bypassed. For input sparsity and/or output sparsity, in one embodiment, metadata is provided to the processing elements **552AA-552MN** to indicate, for a processing cycle, which processing elements and/or data channels are to be active during that cycle.

[0134] In one embodiment, the matrix engine **503** includes hardware to enable operations on sparse data having a compressed representation of a sparse matrix that stores non-zero values and metadata that identifies the positions of the non-zero values within the matrix. Example compressed representations include but are not limited to compressed tensor representations such as compressed sparse row (CSR), compressed sparse column (CSC), compressed sparse fiber (CSF) representations. Support for compressed representations enable operations to be performed on input in a compressed tensor format without utilizing the compressed representation to be decompressed or decoded. In such embodiment, operations can be performed on non-zero input values and the resulting non-zero output values can be mapped into an output matrix. In some embodiments, hardware support is also provided for machine-specific lossless data compression formats that are used when transmitting data within hardware or across system busses. Such data may be retained in a compressed format for sparse input data and the matrix engine **503** can use the compression metadata for the compressed data to enable operations to be performed on non-zero values, or to enable blocks of zero data input to be bypassed for multiply operations.

[0135] In various embodiments, input data can be provided by a programmer in a compressed tensor representation, or a codec can compress input data into the compressed tensor representation or another sparse data encoding. In addition to support for compressed tensor representations, streaming compression of sparse input data can be performed before the data is provided to the processing elements **552AA-552MN**. In one embodiment, compression is performed on data written to a cache memory associated with the graphics core cluster **414**, with the compression being performed with an encoding that is supported by the matrix engine **503**. In one embodiment, the matrix engine **503** includes support for input having structured sparsity in which a pre-determined level or pattern of sparsity is imposed on input data. This data may be compressed to a known compression ratio, with the compressed data being processed by the processing elements **552AA-552MN** according to metadata associated with the compressed data.

[0136] FIG. 6 illustrates a tile **600** of a multi-tile processor, according to an embodiment. In one embodiment, the tile **600** is representative of one of the graphics engine tiles **310A-310D** of FIG. 3B or compute engine tiles **340A-340D** of FIG. 3C. The tile **600** of the multi-tile graphics processor includes an array of graphics core clusters (e.g., graphics core cluster **414A**, graphics core cluster **414B**, through graphics core cluster **414N**), with each graphics core cluster having an array of graphics cores **515A-515N**. The tile **600** also includes a global dispatcher **602** to dispatch threads to processing resources of the tile **600**.

[0137] The tile **600** can include or couple with an L3 cache **606** and memory **610**. In various embodiments, the L3 cache **606** may be excluded or the tile **600** can include additional levels of cache, such as an L4 cache. In one embodiment, each instance of the tile **600** in the multi-tile graphics processor has an associated memory **610**, such as in FIG. 3B and FIG. 3C. In one embodiment, a multi-tile processor can be configured as a multi-chip module in which the L3 cache **606** and/or memory **610** reside on separate chiplets than the graphics core clusters **414A-414N**. In this context, a chiplet is an at least partially packaged integrated circuit that includes distinct units of logic that can be assembled with other chiplets into a larger package. For example, the L3 cache **606** can be included in a dedicated cache chiplet or can reside on the same chiplet as the graphics core clusters **414A-414N**. In one embodiment, the L3 cache **606** can be included in an active base die or active interposer, as illustrated in FIG. 11C.

[0138] A memory fabric **603** enables communication among the graphics core clusters **414A-414N**, L3 cache **606**, and memory **610**. An L2 cache **604** couples with the memory fabric **603** and is configurable to cache transactions performed via the memory fabric **603**. A tile interconnect **608** enables communication with other tiles on the graphics processors and may be one of tile interconnects **323A-323F** of FIGS. 3B and 3C. In embodiments in which the L3 cache **606** is excluded from the tile **600**, the L2 cache **604** may be configured as a combined L2/L3 cache. The memory fabric **603** is configurable to route data to the L3 cache **606** or memory controllers associated with the memory **610** based on the presence or absence of the L3 cache **606** in a specific implementation. The L3 cache **606** can be configured as a per-tile cache that is dedicated to processing resources of the tile **600** or may be a partition of a GPU-wide L3 cache.

[0139] FIG. 7 is a block diagram illustrating graphics processor instruction formats **700** according to some embodiments. In one or more embodiment, the graphics processor cores support an instruction set having instructions in multiple formats. The solid lined boxes illustrate the components that are generally included in a graphics core instruction, while the dashed lines include components that are optional or that are included in a sub-set of the instructions. In some embodiments, the graphics processor instruction format **700** described and illustrated are macro-instructions, in that they are instructions supplied to the graphics core, as opposed to micro-operations resulting from instruction decode once the instruction is processed. Thus, a single instruction may cause hardware to perform multiple micro-operations.

[0140] In some embodiments, the graphics processor natively supports instructions in a 128-bit instruction format **710**. A 64-bit compacted instruction format **730** is available for some instructions based on the selected instruction, instruction options, and number of operands. The native 128-bit instruction format **710** provides access to all instruction options, while some options and operations are restricted in the 64-bit format **730**. The native instructions available in the 64-bit format **730** vary by embodiment. In some embodiments, the instruction is compacted in part using a set of index values in an index field **713**. The graphics core hardware references a set of compaction tables based on the index values and uses the compaction table

outputs to reconstruct a native instruction in the 128-bit instruction format **710**. Other sizes and formats of instruction can be used.

[0141] For each format, instruction opcode **712** defines the operation that the graphics core is to perform. The graphics cores execute each instruction in parallel across the multiple data elements of each operand. For example, in response to an add instruction the graphics core performs a simultaneous add operation across each color channel representing a texture element or picture element. By default, the graphics core performs each instruction across all data channels of the operands. In some embodiments, instruction control field **714** enables control over certain execution options, such as channels selection (e.g., predication) and data channel order (e.g., swizzle). For instructions in the 128-bit instruction format **710** an exec-size field **716** limits the number of data channels that will be executed in parallel. In some embodiments, exec-size field **716** is not available for use in the 64-bit compact instruction format **730**.

[0142] Some graphics core instructions have up to three operands including two source operands, src0 **720**, src1 **722**, and one destination **718**. In some embodiments, the graphics cores support dual destination instructions, where one of the destinations is implied. Data manipulation instructions can have a third source operand (e.g., SRC2 **724**), where the instruction opcode **712** determines the number of source operands. An instruction's last source operand can be an immediate (e.g., hard-coded) value passed with the instruction.

[0143] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726** specifying, for example, whether direct register addressing mode or indirect register addressing mode is used. When direct register addressing mode is used, the register address of one or more operands is directly provided by bits in the instruction.

[0144] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726**, which specifies an address mode and/or an access mode for the instruction. In one embodiment the access mode is used to define a data access alignment for the instruction. Some embodiments support access modes including a 16-byte aligned access mode and a 1-byte aligned access mode, where the byte alignment of the access mode determines the access alignment of the instruction operands. For example, when in a first mode, the instruction may use byte-aligned addressing for source and destination operands and when in a second mode, the instruction may use 16-byte-aligned addressing for all source and destination operands.

[0145] In one embodiment, the address mode portion of the access/address mode field **726** determines whether the instruction is to use direct or indirect addressing. When direct register addressing mode is used bits in the instruction directly provide the register address of one or more operands. When indirect register addressing mode is used, the register address of one or more operands may be computed based on an address register value and an address immediate field in the instruction.

[0146] In some embodiments instructions are grouped based on opcode **712** bit-fields to simplify Opcode decode **740**. For an 8-bit opcode, bits **4**, **5**, and **6** allow the graphics core to determine the type of opcode. The precise opcode grouping shown is merely an example. In some embodiments, a move and logic opcode group **742** includes data

movement and logic instructions (e.g., move (mov), compare (cmp)). In some embodiments, move and logic group **742** shares the five most significant bits (MSB), where move (mov) instructions are in the form of 0000xxxxb and logic instructions are in the form of 0001xxxxb. A flow control instruction group **744** (e.g., call, jump (jmp)) includes instructions in the form of 0010xxxxb (e.g., 0x20). A miscellaneous instruction group **746** includes a mix of instructions, including synchronization instructions (e.g., wait, send) in the form of 0011xxxxb (e.g., 0x30). A parallel math instruction group **748** includes component-wise arithmetic instructions (e.g., add, multiply (mul)) in the form of 0100xxxxb (e.g., 0x40). The parallel math instruction group **748** performs the arithmetic operations in parallel across data channels. The vector math group **750** includes arithmetic instructions (e.g., dp4) in the form of 0101xxxxb (e.g., 0x50). The vector math group performs arithmetic such as dot product calculations on vector operands. The illustrated opcode decode **740**, in one embodiment, can be used to determine which portion of a graphics core will be used to execute a decoded instruction. For example, some instructions may be designated as systolic instructions that will be performed by a systolic array. Other instructions, such as ray-tracing instructions (not shown) can be routed to a ray-tracing core or ray-tracing logic within a slice or partition of execution logic.

Graphics Pipeline

[0147] FIG. **8** is a block diagram of another embodiment of a graphics processor **800**. Elements of FIG. **8** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0148] In some embodiments, graphics processor **800** includes a geometry pipeline **820**, a media pipeline **830**, a display engine **840**, thread execution logic **850**, and a render output pipeline **870**. In some embodiments, graphics processor **800** is a graphics processor within a multi-core processing system that includes one or more general-purpose processing cores. The graphics processor is controlled by register writes to one or more control registers (not shown) or via commands issued to graphics processor **800** via a ring interconnect **802**. In some embodiments, ring interconnect **802** couples graphics processor **800** to other processing components, such as other graphics processors or general-purpose processors. Commands from ring interconnect **802** are interpreted by a command streamer **803**, which supplies instructions to individual components of the geometry pipeline **820** or the media pipeline **830**.

[0149] In some embodiments, command streamer **803** directs the operation of a vertex fetcher **805** that reads vertex data from memory and executes vertex-processing commands provided by command streamer **803**. In some embodiments, vertex fetcher **805** provides vertex data to a vertex shader **807**, which performs coordinate space transformation and lighting operations to each vertex. In some embodiments, vertex fetcher **805** and vertex shader **807** execute vertex-processing instructions by dispatching execution threads to graphics cores **852A-852B** via a thread dispatcher **831**.

[0150] In some embodiments, graphics cores **852A-852B** are an array of vector processors having an instruction set for performing graphics and media operations. In some embodi-

ments, graphics cores **852A-852B** have an attached L1 cache **851** that is specific for each array or shared between the arrays. The cache can be configured as a data cache, an instruction cache, or a single cache that is partitioned to contain data and instructions in different partitions.

[0151] In some embodiments, geometry pipeline **820** includes tessellation components to perform hardware-accelerated tessellation of 3D objects. In some embodiments, a programmable hull shader **811** configures the tessellation operations. A programmable domain shader **817** provides back-end evaluation of tessellation output. A tessellator **813** operates at the direction of hull shader **811** and contains special purpose logic to generate a set of detailed geometric objects based on a coarse geometric model that is provided as input to geometry pipeline **820**. In some embodiments, if tessellation is not used, tessellation components (e.g., hull shader **811**, tessellator **813**, and domain shader **817**) can be bypassed. The tessellation components can operate based on data received from the vertex shader **807**.

[0152] In some embodiments, complete geometric objects can be processed by a geometry shader **819** via one or more threads dispatched to graphics cores **852A-852B** or can proceed directly to the clipper **829**. In some embodiments, the geometry shader operates on entire geometric objects, rather than vertices or patches of vertices as in previous stages of the graphics pipeline. If the tessellation is disabled the geometry shader **819** receives input from the vertex shader **807**. In some embodiments, geometry shader **819** is programmable by a geometry shader program to perform geometry tessellation if the tessellation units are disabled.

[0153] Before rasterization, a clipper **829** processes vertex data. The clipper **829** may be a fixed function clipper or a programmable clipper having clipping and geometry shader functions. In some embodiments, a rasterizer and depth test component **873** in the render output pipeline **870** dispatches pixel shaders to convert the geometric objects into per pixel representations. In some embodiments, pixel shader logic is included in thread execution logic **850**. In some embodiments, an application can bypass the rasterizer and depth test component **873** and access un-rasterized vertex data via a stream out unit **823**.

[0154] The graphics processor **800** has an interconnect bus, interconnect fabric, or some other interconnect mechanism that allows data and message passing amongst the major components of the processor. In some embodiments, graphics cores **852A-852B** and associated logic units (e.g., L1 cache **851**, sampler **854**, texture cache **858**, etc.) interconnect via a data port **856** to perform memory access and communicate with render output pipeline components of the processor. In some embodiments, sampler **854**, caches **851**, **858** and graphics cores **852A-852B** each have separate memory access paths. In one embodiment the texture cache **858** can also be configured as a sampler cache.

[0155] In some embodiments, render output pipeline **870** contains a rasterizer and depth test component **873** that converts vertex-based objects into an associated pixel-based representation. In some embodiments, the rasterizer logic includes a windower/masker unit to perform fixed function triangle and line rasterization. An associated render cache **878** and depth cache **879** are also available in some embodiments. A pixel operations component **877** performs pixel-based operations on the data, though in some instances, pixel operations associated with 2D operations (e.g., bit block image transfers with blending) are performed by the 2D

engine **841**, or substituted at display time by the display controller **843** using overlay display planes. In some embodiments, a shared L3 cache **875** is available to all graphics components, allowing the sharing of data without the use of main system memory.

[0156] In some embodiments, media pipeline **830** includes a media engine **837** and a video front-end **834**. In some embodiments, video front-end **834** receives pipeline commands from the command streamer **803**. In some embodiments, media pipeline **830** includes a separate command streamer. In some embodiments, video front-end **834** processes media commands before sending the command to the media engine **837**. In some embodiments, media engine **837** includes thread spawning functionality to spawn threads for dispatch to thread execution logic **850** via thread dispatcher **831**.

[0157] In some embodiments, graphics processor **800** includes a display engine **840**. In some embodiments, display engine **840** is external to processor **800** and couples with the graphics processor via the ring interconnect **802**, or some other interconnect bus or fabric. In some embodiments, display engine **840** includes a 2D engine **841** and a display controller **843**. In some embodiments, display engine **840** contains special purpose logic capable of operating independently of the 3D pipeline. In some embodiments, display controller **843** couples with a display device (not shown), which may be a system integrated display device, as in a laptop computer, or an external display device attached via a display device connector.

[0158] In some embodiments, the geometry pipeline **820** and media pipeline **830** are configurable to perform operations based on multiple graphics and media programming interfaces and are not specific to any one application programming interface (API). In some embodiments, driver software for the graphics processor translates API calls that are specific to a particular graphics or media library into commands that can be processed by the graphics processor. In some embodiments, support is provided for the Open Graphics Library (OpenGL), Open Computing Language (OpenCL), and/or Vulkan graphics and compute API, all from the Khronos Group. In some embodiments, support may also be provided for the Direct3D library from the Microsoft Corporation. In some embodiments, a combination of these libraries may be supported. Support may also be provided for the Open Source Computer Vision Library (OpenCV). A future API with a compatible 3D pipeline would also be supported if a mapping can be made from the pipeline of the future API to the pipeline of the graphics processor.

Graphics Pipeline Programming

[0159] FIG. 9A is a block diagram illustrating a graphics processor command format **900** that may be used to program graphics processing pipelines according to some embodiments. FIG. 9B is a block diagram illustrating a graphics processor command sequence **910** according to an embodiment. The solid lined boxes in FIG. 9A illustrate the components that are generally included in a graphics command while the dashed lines include components that are optional or that are included in a sub-set of the graphics commands. The example graphics processor command format **900** of FIG. 9A includes data fields to identify a client **902**, a command operation code (opcode) **904**, and a data field **906**

for the command. A sub-opcode **905** and a command size **908** are also included in some commands.

[0160] In some embodiments, client **902** specifies the client unit of the graphics device that processes the command data. In some embodiments, a graphics processor command parser examines the client field of each command to condition the further processing of the command and route the command data to the appropriate client unit. In some embodiments, the graphics processor client units include a memory interface unit, a render unit, a 2D unit, a 3D unit, and a media unit. Each client unit has a corresponding processing pipeline that processes the commands. Once the command is received by the client unit, the client unit reads the opcode **904** and, if present, sub-opcode **905** to determine the operation to perform. The client unit performs the command using information in data field **906**. For some commands an explicit command size **908** is expected to specify the size of the command. In some embodiments, the command parser automatically determines the size of at least some of the commands based on the command opcode. In some embodiments commands are aligned via multiples of a double word. Other command formats can be used.

[0161] The flow diagram in FIG. 9B illustrates an example graphics processor command sequence **910**. In some embodiments, software or firmware of a data processing system that features an embodiment of a graphics processor uses a version of the command sequence shown to set up, execute, and terminate a set of graphics operations. A sample command sequence is shown and described for purposes of example as embodiments are not limited to these specific commands or to this command sequence. Moreover, the commands may be issued as batch of commands in a command sequence, such that the graphics processor will process the sequence of commands in at least partially concurrence.

[0162] In some embodiments, the graphics processor command sequence **910** may begin with a pipeline flush command **912** to cause any active graphics pipeline to complete the currently pending commands for the pipeline. In some embodiments, the 3D pipeline **922** and the media pipeline **924** do not operate concurrently. The pipeline flush is performed to cause the active graphics pipeline to complete any pending commands. In response to a pipeline flush, the command parser for the graphics processor will pause command processing until the active drawing engines complete pending operations and the relevant read caches are invalidated. Optionally, any data in the render cache that is marked 'dirty' can be flushed to memory. In some embodiments, pipeline flush command **912** can be used for pipeline synchronization or before placing the graphics processor into a low power state.

[0163] In some embodiments, a pipeline select command **913** is used when a command sequence utilizes the graphics processor to explicitly switch between pipelines. In some embodiments, a pipeline select command **913** is utilized once within an execution context before issuing pipeline commands unless the context is to issue commands for both pipelines. In some embodiments, a pipeline flush command **912** is utilized before a pipeline switch via the pipeline select command **913**.

[0164] In some embodiments, a pipeline control command **914** configures a graphics pipeline for operation and is used to program the 3D pipeline **922** and the media pipeline **924**. In some embodiments, pipeline control command **914** con-

figures the pipeline state for the active pipeline. In one embodiment, the pipeline control command **914** is used for pipeline synchronization and to clear data from one or more cache memories within the active pipeline before processing a batch of commands.

[0165] In some embodiments, commands related to the return buffer state **916** are used to configure a set of return buffers for the respective pipelines to write data. Some pipeline operations utilize the allocation, selection, or configuration of one or more return buffers into which the operations write intermediate data during processing. In some embodiments, the graphics processor also uses one or more return buffers to store output data and to perform cross thread communication. In some embodiments, the return buffer state **916** includes selecting the size and number of return buffers to use for a set of pipeline operations.

[0166] The remaining commands in the command sequence differ based on the active pipeline for operations. Based on a pipeline determination **920**, the command sequence is tailored to the 3D pipeline **922** beginning with the 3D pipeline state **930** or the media pipeline **924** beginning at the media pipeline state **940**.

[0167] The commands to configure the 3D pipeline state **930** include 3D state setting commands for vertex buffer state, vertex element state, constant color state, depth buffer state, and other state variables that are to be configured before 3D primitive commands are processed. The values of these commands are determined at least in part based on the particular 3D API in use. In some embodiments, 3D pipeline state **930** commands are also able to selectively disable or bypass certain pipeline elements if those elements will not be used.

[0168] In some embodiments, 3D primitive **932** command is used to submit 3D primitives to be processed by the 3D pipeline. Commands and associated parameters that are passed to the graphics processor via the 3D primitive **932** command are forwarded to the vertex fetch function in the graphics pipeline. The vertex fetch function uses the 3D primitive **932** command data to generate vertex data structures. The vertex data structures are stored in one or more return buffers. In some embodiments, 3D primitive **932** command is used to perform vertex operations on 3D primitives via vertex shaders. To process vertex shaders, 3D pipeline **922** dispatches shader programs to the graphics cores.

[0169] In some embodiments, 3D pipeline **922** is triggered via an execute **934** command or event. In some embodiments, a register write triggers command execution. In some embodiments execution is triggered via a 'go' or 'kick' command in the command sequence. In one embodiment, command execution is triggered using a pipeline synchronization command to flush the command sequence through the graphics pipeline. The 3D pipeline will perform geometry processing for the 3D primitives. Once operations are complete, the resulting geometric objects are rasterized and the pixel engine colors the resulting pixels. Additional commands to control pixel shading and pixel back-end operations may also be included for those operations.

[0170] In some embodiments, the graphics processor command sequence **910** follows the media pipeline **924** path when performing media operations. In general, the specific use and manner of programming for the media pipeline **924** depends on the media or compute operations to be performed. Specific media decode operations may be offloaded

to the media pipeline during media decode. In some embodiments, the media pipeline can also be bypassed and media decode can be performed in whole or in part using resources provided by one or more general-purpose processing cores. In one embodiment, the media pipeline also includes elements for general-purpose graphics processor unit (GPGPU) operations, where the graphics processor is used to perform SIMD vector operations using computational shader programs that are not explicitly related to the rendering of graphics primitives.

[0171] In some embodiments, media pipeline 924 is configured in a similar manner as the 3D pipeline 922. A set of commands to configure the media pipeline state 940 are dispatched or placed into a command queue before the media object commands 942. In some embodiments, commands for the media pipeline state 940 include data to configure the media pipeline elements that will be used to process the media objects. This includes data to configure the video decode and video encode logic within the media pipeline, such as encode or decode format. In some embodiments, commands for the media pipeline state 940 also support the use of one or more pointers to “indirect” state elements that contain a batch of state settings.

[0172] In some embodiments, media object commands 942 supply pointers to media objects for processing by the media pipeline. The media objects include memory buffers containing video data to be processed. In some embodiments, all media pipeline states should be valid before issuing a media object command 942. Once the pipeline state is configured and media object commands 942 are queued, the media pipeline 924 is triggered via an execute command 944 or an equivalent execute event (e.g., register write). Output from media pipeline 924 may then be post processed by operations provided by the 3D pipeline 922 or the media pipeline 924. In some embodiments, GPGPU operations are configured and executed in a similar manner as media operations.

Graphics Software Architecture

[0173] FIG. 10 illustrates an example graphics software architecture for a data processing system 1000 according to some embodiments. In some embodiments, software architecture includes a 3D graphics application 1010, an operating system 1020, and at least one processor 1030. In some embodiments, processor 1030 includes a graphics processor 1032 and one or more general-purpose processor core(s) 1034. The graphics application 1010 and operating system 1020 each execute in the system memory 1050 of the data processing system.

[0174] In some embodiments, 3D graphics application 1010 contains one or more shader programs including shader instructions 1012. The shader language instructions may be in a high-level shader language, such as the High-Level Shader Language (HLSL) of Direct3D, the OpenGL Shader Language (GLSL), and so forth. The application also includes executable instructions 1014 in a machine language suitable for execution by the general-purpose processor core 1034. The application also includes graphics objects 1016 defined by vertex data.

[0175] In some embodiments, operating system 1020 is a Microsoft® Windows® operating system from the Microsoft Corporation, a proprietary UNIX-like operating system, or an open source UNIX-like operating system using a variant of the Linux kernel. The operating system

1020 can support a graphics API 1022 such as the Direct3D API, the OpenGL API, or the Vulkan API. When the Direct3D API is in use, the operating system 1020 uses a front-end shader compiler 1024 to compile any shader instructions 1012 in HLSL into a lower-level shader language. The compilation may be a just-in-time (JIT) compilation or the application can perform shader pre-compilation. In some embodiments, high-level shaders are compiled into low-level shaders during the compilation of the 3D graphics application 1010. In some embodiments, the shader instructions 1012 are provided in an intermediate form, such as a version of the Standard Portable Intermediate Representation (SPIR) used by the Vulkan API.

[0176] In some embodiments, user mode graphics driver 1026 contains a back-end shader compiler 1027 to convert the shader instructions 1012 into a hardware specific representation. When the OpenGL API is in use, shader instructions 1012 in the GLSL high-level language are passed to a user mode graphics driver 1026 for compilation. In some embodiments, user mode graphics driver 1026 uses operating system kernel mode functions 1028 to communicate with a kernel mode graphics driver 1029. In some embodiments, kernel mode graphics driver 1029 communicates with graphics processor 1032 to dispatch commands and instructions.

IP Core Implementations

[0177] One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

[0178] FIG. 11A is a block diagram illustrating an IP core development system 1100 that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The IP core development system 1100 may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., an SOC integrated circuit). A design facility 1130 can generate a software simulation 1110 of an IP core design in a high-level programming language (e.g., C/C++). The software simulation 1110 can be used to design, test, and verify the behavior of the IP core using a simulation model 1112. The simulation model 1112 may include functional, behavioral, and/or timing simulations. A register transfer level (RTL) design 1115 can then be created or synthesized from the simulation model 1112. The RTL design 1115 is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL

design 1115, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

[0179] The RTL design 1115 or equivalent may be further synthesized by the design facility into a hardware model 1120, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a 3rd party fabrication facility 1165 using non-volatile memory 1140 (e.g., hard disk, flash memory, or any non-volatile storage medium). Alternatively, the IP core design may be transmitted (e.g., via the Internet) over a wired connection 1150 or wireless connection 1160. The fabrication facility 1165 may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit can be configured to perform operations in accordance with at least one embodiment described herein.

[0180] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly 1170, according to some embodiments described herein. The integrated circuit package assembly 1170 illustrates an implementation of one or more processor or accelerator devices as described herein. The package assembly 1170 includes multiple units of hardware logic 1172, 1174 connected to a substrate 1180. The logic 1172, 1174 may be implemented at least partly in configurable logic or fixed-functionality logic hardware, and can include one or more portions of any of the processor core(s), graphics processor(s), or other accelerator devices described herein. Each unit of logic 1172, 1174 can be implemented within a semiconductor die and coupled with the substrate 1180 via an interconnect structure 1173. The interconnect structure 1173 may be configured to route electrical signals between the logic 1172, 1174 and the substrate 1180, and can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure 1173 may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic 1172, 1174. In some embodiments, the substrate 1180 is an epoxy-based laminate substrate. The substrate 1180 may include other suitable types of substrates in other embodiments. The package assembly 1170 can be connected to other electrical devices via a package interconnect 1183. The package interconnect 1183 may be coupled to a surface of the substrate 1180 to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0181] In some embodiments, the units of logic 1172, 1174 are electrically coupled with a bridge 1182 that is configured to route electrical signals between the logic 1172, 1174. The bridge 1182 may be a dense interconnect structure that provides a route for electrical signals. The bridge 1182 may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic 1172, 1174.

[0182] Although two units of logic 1172, 1174 and a bridge 1182 are illustrated, embodiments described herein may include more or fewer logic units on one or more dies. The one or more dies may be connected by zero or more bridges, as the bridge 1182 may be excluded when the logic

is included on a single die. Alternatively, multiple dies or units of logic can be connected by one or more bridges. Additionally, multiple logic units, dies, and bridges can be connected together in other possible configurations, including three-dimensional configurations.

[0183] FIG. 11C illustrates a package assembly 1190 that includes multiple units of hardware logic chiplets connected to a substrate 1180. A graphics processing unit, parallel processor, and/or compute accelerator as described herein can be composed from diverse silicon chiplets that are separately manufactured. A diverse set of chiplets with different IP core logic can be assembled into a single device. Additionally, the chiplets can be integrated into a base die or base chiplet using active interposer technology. The concepts described herein enable the interconnection and communication between the different forms of IP within the GPU. IP cores can be manufactured using different process technologies and composed during manufacturing, which avoids the complexity of converging multiple IPs, especially on a large SoC with several flavors IPs, to the same manufacturing process. Enabling the use of multiple process technologies improves the time to market and provides a cost-effective way to create multiple product SKUs. Additionally, the disaggregated IPs are more amenable to being power gated independently, components that are not in use on a given workload can be powered off, reducing overall power consumption.

[0184] In various embodiments a package assembly 1190 can include components and chiplets that are interconnected by a fabric 1185 and/or one or more bridges 1187. The chiplets within the package assembly 1190 may have a 2.5D arrangement using Chip-on-Wafer-on-Substrate stacking in which multiple dies are stacked side-by-side on a silicon interposer 1189 that couples the chiplets with the substrate 1180. The substrate 1180 includes electrical connections to the package interconnect 1183. In one embodiment the silicon interposer 1189 is a passive interposer that includes through-silicon vias (TSVs) to electrically couple chiplets within the package assembly 1190 to the substrate 1180. In one embodiment, silicon interposer 1189 is an active interposer that includes embedded logic in addition to TSVs. In such embodiment, the chiplets within the package assembly 1190 are arranged using 3D face to face die stacking on top of the active interposer 1189. The active interposer 1189 can include hardware logic for I/O 1191, cache memory 1192, and other hardware logic 1193, in addition to interconnect fabric 1185 and a silicon bridge 1187. The fabric 1185 enables communication between the various logic chiplets 1172, 1174 and the logic 1191, 1193 within the active interposer 1189. The fabric 1185 may be an NoC interconnect or another form of packet switched fabric that switches data packets between components of the package assembly. For complex assemblies, the fabric 1185 may be a dedicated chiplet enables communication between the various hardware logic of the package assembly 1190.

[0185] Bridge structures 1187 within the active interposer 1189 may be used to facilitate a point-to-point interconnect between, for example, logic or I/O chiplets 1174 and memory chiplets 1175. In some implementations, bridge structures 1187 may also be embedded within the substrate 1180. The hardware logic chiplets can include special purpose hardware logic chiplets 1172, logic or I/O chiplets 1174, and/or memory chiplets 1175. The hardware logic chiplets 1172 and logic or I/O chiplets 1174 may be imple-

mented at least partly in configurable logic or fixed-functionality logic hardware and can include one or more portions of any of the processor core(s), graphics processor(s), parallel processors, or other accelerator devices described herein. The memory chiplets **1175** can be DRAM (e.g., GDDR, HBM) memory or cache (SRAM) memory. Cache memory **1192** within the active interposer **1189** (or substrate **1180**) can act as a global cache for the package assembly **1190**, part of a distributed global cache, or as a dedicated cache for the fabric **1185**.

[**0186**] Each chiplet can be fabricated as separate semiconductor die and coupled with a base die that is embedded within or coupled with the substrate **1180**. The coupling with the substrate **1180** can be performed via an interconnect structure **1173**. The interconnect structure **1173** may be configured to route electrical signals between the various chiplets and logic within the substrate **1180**. The interconnect structure **1173** can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **1173** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic, I/O, and memory chiplets. In one embodiment, an additional interconnect structure couples the active interposer **1189** with the substrate **1180**.

[**0187**] In some embodiments, the substrate **1180** is an epoxy-based laminate substrate. The substrate **1180** may include other suitable types of substrates in other embodiments. The package assembly **1190** can be connected to other electrical devices via a package interconnect **1183**. The package interconnect **1183** may be coupled to a surface of the substrate **1180** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[**0188**] In some embodiments, a logic or I/O chiplet **1174** and a memory chiplet **1175** can be electrically coupled via a bridge **1187** that is configured to route electrical signals between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may be a dense interconnect structure that provides a route for electrical signals. The bridge **1187** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may also be referred to as a silicon bridge or an interconnect bridge. For example, the bridge **1187**, in some embodiments, is an Embedded Multi-die Interconnect Bridge (EMIB). In some embodiments, the bridge **1187** may simply be a direct connection from one chiplet to another chiplet.

[**0189**] FIG. 11D illustrates a package assembly **1194** including interchangeable chiplets **1195**, according to an embodiment. The interchangeable chiplets **1195** can be assembled into standardized slots on one or more base chiplets **1196**, **1198**. The base chiplets **1196**, **1198** can be coupled via a bridge interconnect **1197**, which can be similar to the other bridge interconnects described herein and may be, for example, an EMIB. Memory chiplets can also be connected to logic or I/O chiplets via a bridge interconnect. I/O and logic chiplets can communicate via an interconnect fabric. The base chiplets can each support one or more slots in a standardized format for one of logic or I/O or memory/cache.

[**0190**] In one embodiment, SRAM and power delivery circuits can be fabricated into one or more of the base chiplets **1196**, **1198**, which can be fabricated using a different process technology relative to the interchangeable chiplets **1195** that are stacked on top of the base chiplets. For example, the base chiplets **1196**, **1198** can be fabricated using a larger process technology, while the interchangeable chiplets can be manufactured using a smaller process technology. One or more of the interchangeable chiplets **1195** may be memory (e.g., DRAM) chiplets. Different memory densities can be selected for the package assembly **1194** based on the power, and/or performance targeted for the product that uses the package assembly **1194**. Additionally, logic chiplets with a different number of type of functional units can be selected at time of assembly based on the power, and/or performance targeted for the product. Additionally, chiplets containing IP logic cores of differing types can be inserted into the interchangeable chiplet slots, enabling hybrid processor designs that can mix and match different technology IP blocks.

Example System on a Chip Integrated Circuit

[**0191**] FIGS. 12-13B illustrate example integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

[**0192**] FIG. 12 is a block diagram illustrating an example system on a chip integrated circuit **1200** that may be fabricated using one or more IP cores, according to an embodiment. Example integrated circuit **1200** includes one or more application processor(s) **1205** (e.g., CPUs), at least one graphics processor **1210**, and may additionally include an image processor **1215** and/or a video processor **1220**, any of which may be a modular IP core from the same or multiple different design facilities. Integrated circuit **1200** includes peripheral or bus logic including a USB controller **1225**, UART controller **1230**, an SPI/SDIO controller **1235**, and an I2S/I2C controller **1240**. Additionally, the integrated circuit can include a display device **1245** coupled to one or more of a high-definition multimedia interface (HDMI) controller **1250** and a mobile industry processor interface (MIPI) display interface **1255**. Storage may be provided by a flash memory subsystem **1260** including flash memory and a flash memory controller. Memory interface may be provided via a memory controller **1265** for access to SDRAM or SRAM memory devices. Some integrated circuits additionally include an embedded security engine **1270**.

[**0193**] FIGS. 13A-13B are block diagrams illustrating example graphics processors for use within an SoC, according to embodiments described herein. FIG. 13A illustrates an example graphics processor **1310** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. FIG. 13B illustrates an additional example graphics processor **1340** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. Graphics processor **1310** of FIG. 13A is an example of a low power graphics processor core. Graphics processor **1340** of FIG. 13B is an example of a higher performance graphics processor.

cessor core. Each of graphics processor **1310** and graphics processor **1340** can be variants of the graphics processor **1210** of FIG. 12.

[0194] As shown in FIG. 13A, graphics processor **1310** includes a vertex processor **1305** and one or more fragment processor(s) **1315A-1315N** (e.g., **1315A**, **1315B**, **1315C**, **1315D**, through **1315N-1**, and **1315N**). Graphics processor **1310** can execute different shader programs via separate logic, such that the vertex processor **1305** is optimized to execute operations for vertex shader programs, while the one or more fragment processor(s) **1315A-1315N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. The vertex processor **1305** performs the vertex processing stage of the 3D graphics pipeline and generates primitives and vertex data. The fragment processor(s) **1315A-1315N** use the primitive and vertex data generated by the vertex processor **1305** to produce a frame-buffer that is displayed on a display device. In one embodiment, the fragment processor(s) **1315A-1315N** are optimized to execute fragment shader programs as provided for in the OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in the Direct 3D API.

[0195] Graphics processor **1310** additionally includes one or more memory management units (MMUs) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B**. The one or more MMU(s) **1320A-1320B** provide for virtual to physical address mapping for the graphics processor **1310**, including for the vertex processor **1305** and/or fragment processor(s) **1315A-1315N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in the one or more cache(s) **1325A-1325B**. In one embodiment the one or more MMU(s) **1320A-1320B** may be synchronized with other MMUs within the system, including one or more MMUs associated with the one or more application processor(s) **1205**, image processor **1215**, and/or video processor **1220** of FIG. 12, such that each processor **1205-1220** can participate in a shared or unified virtual memory system. The one or more circuit interconnect(s) **1330A-1330B** enable graphics processor **1310** to interface with other IP cores within the SoC, either via an internal bus of the SoC or via a direct connection, according to embodiments.

[0196] As shown FIG. 13B, graphics processor **1340** includes the one or more MMU(s) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B** of the graphics processor **1310** of FIG. 13A. Graphics processor **1340** includes one or more shader core(s) **1355A-1355N** (e.g., **1355A**, **1355B**, **1355C**, **1355D**, **1355E**, **1355F**, through **1355N-1**, and **1355N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. The unified shader core architecture is also configurable to execute direct compiled high-level GPGPU programs (e.g., CUDA). The exact number of shader cores present can vary among embodiments and implementations. Additionally, graphics processor **1340** includes an inter-core task manager **1345**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **1355A-1355N** and a tiling unit **1358** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in

image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

Parallel Computing

[0197] Parallel computing is a type of computation in which many calculations or the execution of processes are carried out simultaneously. Parallel computing may come in a variety of forms, including, but not limited to, SIMD or SIMT. SIMD describes computers with multiple processing elements that perform the same operation on multiple data points simultaneously. In one example, the figures discussed above refer to SIMD and its implementation in a general processor in terms of EUs, FPUs, and ALUs. In a common SIMD machine, data is packaged into registers, each containing an array of channels. Instructions operate on the data found in channel *n* of a register with the data found in the same channel of another register. SIMD machines are advantageous in areas where a single sequence of instructions can be simultaneously applied to high amounts of data. For example, in one embodiment, a graphics processor (e.g., GPGPU, GPU, etc.) can be used to perform SIMD vector operations using computational shader programs.

[0198] Various embodiments can also apply to use execution by use of Single Instruction Multiple Thread (SIMT) as an alternate to use of SIMD or in addition to use of SIMD. Reference to a SIMD core or operation can apply also to SIMT or apply to SIMD in combination with SIMT. The following description is discussed in terms of SIMD machines. However, embodiments herein are not solely limited to application in the SIMD context and may apply in other parallel computing paradigms, such as SIMT, for example. For ease of discussion and explanation, the following description generally focuses on a SIMD implementation. However, embodiments can similarly apply to SIMT machines with no modifications to the described techniques and methodologies. With respect to SIMT machines, similar patterns as discussed below can be followed to provide instructions to the systolic array and execute the instructions on the SIMT machine. Other types of parallel computing machines may also utilize embodiments herein as well.

[0199] Parallel rendering graphics architectures are frequently utilized to perform matrix multiplication operations. For example, matrix multiplication operations are a large part of artificial intelligence (AI)/machine learning (ML) workloads. As such, enhancements to the parallel rendering graphics architecture that can reduce the power requirements (e.g., dynamic capacitance), increase read and write throughout, and/or reduce bandwidth when performing the matrix multiplication operations would provide improved performance of the parallel rendering graphics architecture.

Distinct Paths Counterexamples and Stimuli Generation

[0200] The design of hardware circuits, such as the graphics architectures described herein, is often performed through different levels of abstraction, from high-level hardware description languages (HDLs) to the low-level design of the transistors and other components. As used herein, a hardware design may refer to a description of an integrated circuit that can be used to generate a hardware manifestation of the integrated circuit (e.g., the hardware design may be synthesized into silicon or used to program an FPGA). One abstraction level is the register transfer level (RTL) abstraction level, which is used in HDLs such as Verilog or VHDL

(Very High Speed Integrated Circuits Hardware Description Language). A representation of a circuit on the RTL abstraction level is subsequently synthesized to a netlist, and then to a circuit design to be used for manufacturing the integrated circuit.

[0201] Hardware languages, such as the HDLs discussed above, provide a multitude of possibilities of implementing a given functionality, leading to a large design space. Such a large design space may render any improvement or optimization of the circuit design more difficult, as the number of possible implementations is large, and an estimation of the hardware implementation cost is not straight forward. Moreover, it is not uncommon to introduce mistakes in the HDL code while writing the code.

[0202] As such, a process referred to as verification may be implemented to make sure that models generated via HDLs are functionally correct. The verification process may utilize verification methodologies that can be grouped into two categories: formal verification or simulation-based verification. Formal verification is a systematic process that uses mathematical reasoning to verify a property and/or formal equivalence of a hardware design. Formal verification may be classified further into two categories: equivalence checking and property verification. Simulation-based verification (also referred to herein as simulation) is a process in which a hardware design is tested by applying stimuli to the hardware design and monitoring the output of the hardware design in response to the stimuli.

[0203] Formal verification and simulation-based verification are often considered as two complementary techniques for checking correctness of hardware designs. The most prominent distinction between formal verification and simulation-based verification is that the latter utilizes input vectors and the former does not. The process in simulation-based verification is first to generate input vectors and then to derive reference outputs. The process is reversed in the formal verification process. The user starts out by stating what output behavior is desirable and then lets the formal checker prove or disprove it. Thus, in simulation-based verification, users do not concern themselves with input stimuli at all. As such, the simulation-based methodology may be considered input driven and the formal methodology may be considered output driven.

[0204] With respect to formal verification, formal verification tools may be utilized to apply formal verification techniques to a hardware design to prove the correctness of such a design. In conventional approaches, formal verification tools report one type of failure for a specific check. However, there can often be multiple reasons causing a same check to fail. In many cases, designers utilize the single counterexample to debug, fix the code, and re-run the verification tool with the hope that there are no other causes (within the design) of failure for the specific check. Otherwise, if there are other causes for the check in the design, the designer has to repeat the process when another portion of the code generates a failure for the same specific check.

[0205] In the context of simulation-based verification, a simulation tool may be utilized to process input vectors to derive reference outputs. The reference outputs from the simulation may be compared to outputs from a golden specification (golden model) for purposes of verification. In some cases, the input stimuli for simulation-based verification may be generated prior to the simulation and can be read into the design under test (e.g., from a database during

simulation), or may be generated during a simulation run. During simulation-based verification, the simulation tool generates simulation test cases that stress different scenarios in the design under test, which should match the golden specification. However, a problem encountered with conventional simulation-based verification approaches is that, in terms of input stimuli generation, it can be difficult and time-consuming to identify all of the corner cases (of input stimuli) that can be used to stress the design under test. There are existing techniques to generate input stimuli, but none of the techniques cover all of the corner cases.

[0206] Implementations herein seek to improve the verification process in the design of hardware circuits by facilitating distinct paths counterexample generation and stimuli generation. In implementations, the generation of multiple distinct counterexamples and/or the generation of stimuli may occur in a single run of a verification tool. With respect to counterexample generation, implementations herein can, for a given graph representation (e.g., multiplexer (mux) tree structure) representing the design under test, generate a distinct counterexample for each possible cause of a test failure for a given property. Implementations herein utilize a set of cover properties generated for each possible path from root node to leaf node in the graph representation of the design to determine if there is a test failure for a given property along each particular path. For those cover properties that pass (i.e., the design failed), a counterexample is produced for debug purposes. This allows for producing multiple counterexamples for specific check in a single run of the verification tool.

[0207] This approach can be similarly used to generate stimuli for the design under test by designing cover properties for each possible output (or subset of outputs) of a property to be tested. The inputs used when a cover property passes can be used as stimuli inputs for a simulation tool in the verification process.

[0208] Embodiments provide a technical advantage of improving performance of the verification process implemented by a computing device. The approaches discussed to implement distinct paths counterexample generation and stimuli generation can reduce the time engineers spend in the process of identifying and fixing bugs, in addition to save resources associated with compute costs. By producing a counterexample for each possible cause of the failure, designers can identify sooner the issues in their code and fix them all at the same time, avoiding to re-run the formal tool after each individual fix, running it instead once after all the fixes. This reduces processing time and compute resource usage. Furthermore, when this solution is applied to generate stimuli, it reduces the time to identify relevant simulation input values, without the risk of missing the identification of corner cases.

[0209] FIG. 14 is a block diagram illustrating an example system 1400 for distinct paths counterexample and stimuli generation, according to embodiments. In one implementation, system 1400 may include a GPGPU or GPU, such as the example GPGPUs and/or GPUs described herein with respect to FIGS. 1-13B. In one implementation, system 1400 may include a processing core of a GPGPU or GPU, such as the example graphics cores described herein with respect to FIGS. 1-13B. The elements of FIG. 14 having the same or similar names as the elements of any other figure herein describe the same elements as in the other figures, can operate or function in a manner similar to that, can comprise

the same components, and can be linked to other entities, as those described elsewhere herein, but are not limited to such.

[0210] Example system 1400 may include variant of the graphics processor 108, 208, 1210, or of any graphics processor described herein and may be used in place of any graphics processor described. Example system 1400 may also include a variant of the graphics core 515 described herein and may be used in place of any graphics core described.

[0211] As shown in FIG. 14, the system 1400 includes a simulation engine 1410, a hardware design (to implement function) 1420, verification property 1430, and a formal verification tool. In one implementation, the hardware design 1420 may refer to a description of an integrated circuit that can be used to generate a hardware manifestation of an integrated circuit. In one example, the integrated circuit may be GPGPU or GPU, such as the example GPGPUs and/or GPUs described herein with respect to FIGS. 1-13B. The hardware design 1420 may be implemented in a high-level HDL, such as, but not limited to, an RTL language. Examples of RTL languages include, but are not limited to, VHDL and Verilog. In some implementations, other HDLs may also be utilized for hardware design 1420, such as proprietary HDLs.

[0212] In implementations herein, the hardware design 1420 and at least one verification property 1430 are loaded into the formal verification tool 1440. The formal verification tool 1440 is hardware, software, firmware, or some combination of the above that is capable of performing formal verification of the hardware design 1420 using the verification properties 1430.

[0213] The verification property 1430 may refer to a desired behavior of the hardware design 1420, such as input/output behavior or internal design behavior, for example. The verification property 1430 can specify what the hardware design 1420 should do or how it should behave.

[0214] The verification property 1430 may be processed by the formal verification tool 1440 using a variety of formats. In some implementations, the formats may include, but are not limited to, coverage property (covers) and assertion properties (assert or assertions). The covers and/or assertions are used to capture temporal and combinational behavior of the hardware design in a formal and unambiguous way. The hardware design can then be verified to determine that it conforms to the constraints as captured by the cover(s)/assertion(s).

[0215] The difference between a cover and an assertion is that a cover may fail for some inputs but pass at least once, whereas for an assertion to pass it cannot have any inputs for which the assertion fails. The cover or assertion can be a statement embedded in the formal verification tool 1440 to check the validity of the verification property at a specific point in the design's execution.

[0216] Covers and/or assertions are typically expressed in property language that captures the hardware design behavior spread across multiple hardware design cycles (e.g. clock cycles) in a concise, unambiguous manner. While traditional HDLs, such as an RTL language, have the ability to capture individual cycle behavior, they are too detailed to describe properties at a higher level. In particular, property languages provide means to express temporal relationships and complex hardware design behaviors in a concise manner.

[0217] In some implementations, the verification properties 1430 may be provided to the formal verification tool 1440 pre-configured in one of the formats. In some embodiments, the formal verification tool 1440 may process the verification property 1430 to express the verification property in one of these formats.

[0218] The hardware design 1420 (e.g. RTL) and verification properties 1430 are loaded into the formal verification tool 1440. The formal verification tool 1440 is a software tool that is capable of performing formal verification of a hardware design 1420. As previously discussed, formal verification is a systematic process that uses mathematical reasoning to verify a property in a hardware design.

[0219] In formal verification, the hardware design 1420 is transformed into a mathematical model (e.g. a state-transition system) and the verification properties 1430 are expressed using mathematical logic using a precise syntax or a language with a precise mathematical syntax and semantics.

[0220] In some implementations, such as some of the conventional approaches to the verification process, a verification property 1430 can be verified by searching the entire reachable state space of the hardware design 1420 without explicitly traversing the state machine. The search is done by, for example, encoding the states using efficient Boolean encodings using Binary decision diagrams (BDDs), or using advanced SAT (satisfiability-based bounded model checking)-based techniques. In some cases, tools can be used to implement techniques, such as, but not limited to, abstraction, symmetry, symbolic indexing, and invariants to improve performance and achieve scalability.

[0221] The formal verification tool 1440 may output an indication 1445 of whether the verification property 1430 holds for the hardware design 1420 (i.e., the hardware design 1420 behaves as expected). The output may be confirming that the verification properties 1430 is true for the hardware design 1420; at least one of the verification properties 1430 is not true for the hardware design 1420; or the formal verification was inconclusive. The formal verification may be inconclusive, for example, because the computing-based device running the formal verification tool 1440 has run out of memory or because the formal verification tool 1440 has determined that a certain amount of progress has not been made after a predefined period of time.

[0222] In some cases, where at least one verification property 1430 fails, the formal verification tool 1440 may also output information indicating when and/or where the assertion failed to aid designers in correcting or amending the hardware design 1420. For example, such information may be in the form of a counterexample. A counterexample may be a sequence of input assignments bringing the hardware design 1420 from the initial state into a state where the verification property 1430 fails (e.g., a trace that demonstrates the failure of the specification in the model of the hardware design 1420).

[0223] Formal verification can be contrasted to simulation-based verification in which a hardware design is verified by applying stimuli to the hardware design and monitoring the output of the hardware design in response to the stimuli. System 1400 includes the simulation engine 1410 which may be a tool (implemented via hardware, software, firmware, and/or some combination of the above) capable of performing simulation-based verification of the hardware design 1420. In particular, the simulation engine 1410

monitors the output of the hardware design **1420** in response to each input provided to the simulation engine **1410** to determine if a specific property is met.

[0224] The simulation engine **1410** may perform the simulation-based verification using any known method. For example, the simulation engine **1410** may receive the hardware design **1420** in, for example HDL, convert the HDL to another language, such as C, and perform the simulation on the C code; the simulation engine **1410** may receive the hardware design **1420** as, for example, HDL and perform the simulation directly on the HDL; or the simulation engine **1410** may implement the hardware design **1420** in hardware and perform the simulation on the hardware.

[0225] Once the simulation is complete, the simulation engine **1410** may output an indication **1415** of whether or not the hardware design passed the simulation (i.e., the outputs for all inputs satisfied the property).

[0226] Once the verification (formal and/or simulation-based) is complete the hardware design **1420** may be modified based on the outcome of the verification (e.g., the information output by the formal verification tool **1440** and/or the simulation engine **1410**). For example, if the verification identifies that the hardware design **1420** is not operating as expected, then the hardware design **1420** may be modified to correct its operation. The modified hardware design **1420** may then be re-verified, and once verified to be operating as expected the modified hardware design may be implemented in hardware to generate an integrated circuit, for example.

[0227] In implementations herein, the formal verification tool **1440** can facilitate the generation of multiple distinct counterexamples **1450** as part of evaluating a verification property **1430**, as well as to generate stimuli **1460** for use in simulation engine **1410**. The generation of multiple distinct counterexamples **1450** and/or the generation of stimuli **1460** may occur in a single run of the formal verification tool **1440**, rather than conventional approach where multiple runs of the formal verification tool **1440** would generate the multiple counterexamples.

[0228] With respect to multiple counterexample generation **1450**, implementations herein can, for a given graph representation (e.g., mux tree structure) representing the design under test, generate a distinct counterexample for each possible cause of a test failure for a given property. Implementations herein utilize a set of cover properties generated for each possible path from root node to leaf node in the graph representation of the design to determine if there is a test failure for a given property along each particular path. For those cover properties that pass (i.e., the design failed), a counterexample is produced for debug purposes. This allows for producing multiple counterexamples **1450** for a specific verification property **1430** check in a single run of the formal verification tool **1440**.

[0229] With respect to stimuli generation, the techniques for multiple counterexample generation **1450** can be similarly implemented to generate stimuli **1460** for the hardware design **1420**. In one implementation, the formal verification tool **1440** can generate and execute cover properties for each possible output (or subset of outputs) of a verification property **1430** to be tested. The inputs that were utilized when one of the cover properties passes can then be used as stimuli inputs for the simulation engine **1410** in the verification process. FIGS. **15-18** discussed below provide further

details on the approach to facilitate multiple distinct counterexample and stimuli generation of implementations herein.

[0230] FIG. **15** is a block diagram illustrating a detailed view of an example formal verification tool **1500** facilitating multiple distinct paths counterexamples and stimuli generation, in accordance with implementations herein. In one implementation, formal verification tool **1500** may be the same as formal verification tool **1440** described with respect to FIG. **14**. For example, formal verification tool **1500** may provide a more detailed view of the components of formal verification tool **1440**. As such, the elements of FIG. **15** having the same or similar names as the elements of any other figure herein describe the same elements as in the other figures, can operate or function in a manner similar to that, can comprise the same components, and can be linked to other entities, as those described elsewhere herein, but are not limited to such. In one implementation, formal verification tool **1500** may be part of a GPGPU and/or GPU as described in any of FIGS. **1-13B** herein.

[0231] In one example implementation, formal verification tool **1500** may include a structure extractor **1510**, a cover generator **1520**, a counterexample generator **1530**, and a stimuli generator **1540**. More or less components than those illustrated and described in formal verification tool **1500** may be part of formal verification tool **1500**.

[0232] Formal verification tool **1500** may perform a verification process for a given hardware design, where the verification process is capable of generating multiple distinct paths counterexamples, and/or generating stimuli based on multiple distinct paths. As an example, the following case illustrates a situation where multiple counterexamples may be produced when checking a single verification property. The below example is RTL of a hardware design for an adder module, where the RTL as shown that computes the wrong result in more than one case.

1	module adder (A,B,S);
2	input logic [7:0] A, B;
3	output [8:0] S;
4	
5	always_comb
6	begin
7	if (A[1:0]==0)
8	S = B + 1
9	else if (A[1]==1)
10	S = B
11	else
12	S = A + B;
13	end
14	
15	endmodule

[0233] In the example, the RTL code provides an adder module that produces a sum, S, from variables A and B, where S=A+B. However, as shown in the example code above, a wrong results can be computed when A[1:0]==0 (in this case, if A=0, then S should be =B, rather than B+1 as provided in line 8) or when A[1]==1 (in this case, if A=1, then S=B+1, rather than B as provided in line 10). Referring to this example, in a conventional formal verification tool, if the verification property is to check (i.e., verification property) that the output of this sample adder module is equal to the addition of the two input values (e.g., check against a golden specification model), the conventional formal veri-

fication tool would return a failure related to one of the two possible wrong cases, without identifying the other wrong case.

[0234] However, implementations herein provide an enhanced formal verification tool **1500** that can produce a counterexample for each (all) of the cases that are incorrectly encoded. For example, the enhanced formal verification tool **1500** of implementations herein generates a counterexample for when $A[1:0]=0$ (targeting line number 8 of the example code), and also generates a counterexample for when $A[1]=1$ (targeting line 10 of the example code).

[0235] Referring to formal verification tool **1500** of FIG. **15**, the structure extractor **1510** may receive a hardware design **1502**. Hardware design **1502** may be the same as hardware design **1420** described with respect to FIG. **14**. In some implementations, the hardware design **1502** may also be referred to as an implementation design. Structure extractor **1510** may include a graph representation generator **1515**. The structure extractor **1510** may determine a data flow between components of the hardware design **1502** and then the graph representation generator **1515** may generate a graph representation structure **1518**, such as a mux tree structure, from modeling the data flow. The structure extractor **1510** may utilize any source, including a higher-abstraction level representation of the hardware design **1502** (such as System Verilog) or from another RTL representation of the hardware design **1502** (e.g., as defined in the Verilog or VHDL hardware description language), to identify the data flow between the components of the hardware design **1502**. The graph representation may then model the identified data flow of the hardware design **1502** to generate the graph representation structure **1518**.

[0236] In one implementation, the graph representation structure **1518** may include a tree structure (e.g., a mux tree structure) having two types of vertices (i.e., nodes): a first set of vertices representing operators which may be referred to herein as operator nodes, and second vertices representing operands which may be referred to herein as leaf nodes. The second set of vertices representing operands (leaf nodes) are connected to the first set of vertices representing operators (operator nodes) via edges of the graph structure. Operator nodes may be connected to other operator nodes as well, with the result of an operation performed by an operator node being used as operand by the other operator node. Thus, the output of an operation performed by an operator node may be provided as an operand to another operator node, or may be provided as an output of the circuit as a leaf node.

[0237] In implementations herein, the graph representation structure **1518** may also be referred to as a mux tree or a mux tree structure having operator nodes and leaf nodes, connected by edges. In implementations herein, case statements can be treated as cascade muxes in the graph representation structure **1518**. In this case, it is assumed that the graph representation (e.g., mux tree) is binary. When generating the graph representation structure **1518**, the graph representation generator **1515** may label each of the internal nodes (i.e., root node and operator nodes), such as from $N_1, \dots, N_m$. Each leaf node of the graph structure may also be labeled (e.g., such as from $L_1, \dots, L_m$) is identified by a specific path starting from a root operator node of the graph representation. As such, for each leaf of the graph representation,

it is reach with the values (i.e., operation result) of the muxes, where such path can be written as $N_j=b_j, \dots, N_k=b_k$, where each $b_i \in \{0,1\}$.

[0238] FIG. **16A** illustrates an example graph representation structure **1600** of the hardware design provided by the sample RTL code for the adder module provided above, in accordance with implementations herein. In one implementation, graph representation structure **1600** may be one example of graph representation structure **1518** of FIG. **15**.

[0239] As shown in FIG. **16A**, graph representation structure **1600** includes operator nodes 1 and 2, and leaf nodes L1, L2, L3. Operator node 1 represents line 7 of the sample RTL code where the if statement determines the value of $A[1:0]$. Operation node 2 represents line 9 of the sample RTL code where the else if statement determines the value of $A[1]$. The edges extending from operator nodes 1 and 2 may include labels, such as “0” and “1”, indicating the operand of the respective operator node. Note that when the result of the operation of operator node 1 is 1, the edge node labeled 1 proceeds to operator node 2. At this point, operator node 2 performs its operation in response to the operand of operator node 1 equaling “1”.

[0240] Leaf node L1 represents the result when the operation of operator node 1 is 0, and represents line 8 of the sample RTL code, where “ $S=B+1$ ”. Leaf node L2 represents the result when the operation of operator node 2 is 0, and represents line 12 of the sample RTL code, where “ $S=A+B$ ”. Leaf node L3 represents the result when the operation of operator node 2 is 1, and represents line 10 of the sample RTL code, where “ $S=B$ ”.

[0241] Referring back to FIG. **15**, the graph representation structure **1518** extracted from hardware design **1502** may be passed to cover generator **1520**. Cover generator **1520** may generate cover properties to apply to the graph representation structure **1518** to test one or more verification properties **1504**. In implementations herein, the cover generator writes a cover property for each leaf node of the graph representation structure **1518** to test for the verification property along the path through the graph representation structure **1518** from a root node to the leaf node. As such, for each leaf node in the graph representation structure, the cover generator **1520** writes a cover property that is to provide a set of space constraints on the graph representation structure **1518** and to query whether the particular verification property is satisfied by the hardware design **1502** having the set of space constraints applied. In this case, the set of space constraints are those that define a path through the graph representation structure **1518** from the root node to the leaf node.

[0242] With reference to the sample RTL code of the adder module provided above, if it is assumed that the verification property **1504** is that an output of the RTL code equals an expected output (e.g., a golden truth, golden model, etc.). This verification property **1504** can be written by the cover generator **1520** as “output=expected_output”. Utilizing this verification property **1504**, the cover generator **1520** can a cover property for each leaf node of the graph representation tree as follows. For each leaf, if it is reached with the values of the muxes $N_j=b_j, \dots, N_k=b_k$, where each $b_i \in \{0,1\}$, write the cover property as:

$$\text{Cover}((\text{output}=\text{expected_Output}) \& (N_j=b_j) \& \dots \& (N_k=b_k))$$

[0243] In the example cover property above, the cover is written as a function of the verification property (output-

$\neq$ expected_output) and the defined path through the graph representation structure from the root node to the leaf node ($N_j=b_j \& \dots \& N_k=b_k$). Other verification properties can be tested for each leaf node by changing the verification property portion of the cover property.

[0244] The cover generator 1520 can then execute all of the generated covers for a verification property 1504 in a single run of the formal verification tool 1500. The cover results 1525 produced for each cover property can be passed to counterexample generator 1530 and/or stimuli generator 1540.

[0245] In one implementation, the counterexample generator 1530 can report the cover results 1525 produced for each passing cover property as counterexamples, such as multiple generated counterexamples 1535. These counterexamples 1535 provide distinct cases of how the output of the hardware design 1502 is not computed correctly.

[0246] Referring to the sample RTL code for the adder module discussed above, three cover properties may be generated by cover generator 1520 in this example. These three cover properties corresponding to each leaf node (leaf nodes L1, L2, L3 of FIG. 16A) of the graph representation structure 1518 of the sample RTL code. The three cover properties are as shown below:

$\text{cover}((S \neq A+B) \& (A[1:0] = 0))$	Cover1:
$\text{cover}((S \neq A+B) \& (A[1:0] = 1) \& (A[1] = 0))$	Cover2:
$\text{cover}((S \neq A+B) \& (A[1:0] = 1) \& (A[1] = 1))$	Cover3:

[0247] In these example cover properties, Cover1 correlates to leaf 1 (L1) and line 8 of the sample RTL code, Cover2 correlates to leaf 3 (L3) and line 12 of the sample RTL code, and Cover3 correlates to leaf 2 (L2) and line 10 of the sample RTL code. Running the formal verification tool on each of these example cover properties provides the following results. Cover1 passes so there is a counterexample to report (which correspond to the bug at line 8). Cover2 fails because the implementation is correct in this path and there is no counterexample to report. Cover3 passes, so there is a counterexample to report (which correspond to the bug at line 10). As a result, at the end of the execution of a single formal run of the formal verification tool 1500, the counterexample generator 1530 can provide two counterexamples 1535 for the two different bugs that are in the sample RTL code for the hardware design 1502.

[0248] With reference to the stimuli generator 1540, the cover results 1525 can also be utilized to generate stimuli, which is provided as generated stimuli 1545. In the case of generating stimuli, the structure extractor 1510 may utilize both the hardware design 1502 as well as a specification design 1506 to determine the data flow between components of both designs. The specification design 1506 may a model that is trusted to be a correct design and can be referred to as the golden reference (or golden model), which is used for equivalence checking purposes with respect to the hardware design 1502. As previously noted, the hardware 1502 design may also be referred to as an implementation design. The graph representation generator 1515 then generates the graph representation structure 1518 as discussed above (using the data flow identified from both the hardware design 1502 and specification design 1506).

[0249] The cover generator may utilize the graph representation structure 1518 to write, for each leaf node of the graph representation structure, cover properties for each

possible output bit value of each output port of the leaf node. In one example implementation, where the verification property 1504 is to determine whether a particular bit (e.g., a flag) is set or not (e.g., is a value 0 or a value 1), the cover generator 1520 can write the cover properties for each leaf node, if the leaf node is reached with the values of the muxes $N_j=b_j, \dots, N_k=b_k$ (where each $b_i \in \{0,1\}$), for each bit O_i of each output port of the leaf node as follows:

$$\text{Cover1}((O_i=0) \& (N_j=b_j) \& \dots \& (N_k=b_k))$$

$$\text{Cover2}((O_i=1) \& (N_j=b_j) \& \dots \& (N_k=b_k))$$

[0250] In one example, if the hardware design 1502 is adding floating point (FP) numbers, along with the addition results, some components will produce other flags, such as overflow, underflow, etc. Verification properties 1504 can be used to design test cases to find different cases where the value of a particular flag, such as the overflow flag, is either 0 or 1 (binary case). As such, the cover properties above can be written to consider each of the possible binary outputs (0 or 1) of the overflow flag. Each of the Cover1 and Cover2 above would check down each path of the graph representation structure, if an overflow flag=0 is achieved, as well as checking down each path if an overflow flag=1 is achieved. The overflow flag is one example of a binary output value that can be checked. Other binary outputs may be similarly checked using the cover properties described above.

[0251] If a verification property is to be checked for outputs that are more than a binary value at a leaf node, then it may be difficult to write covers for every possible bit of output of the leaf node. In this case, the cover properties can be written to consider a threshold value v , such that the output ports with up to v bits are considered to generate the covers (e.g., if the output is less than (or equal to) the threshold value v , then write the cover properties for each of the v bits of the output port; if the output is more than threshold value v , then do not write any covers). In another example, the threshold value v may be considered so that the top (or bottom) v bits of the output ports are considered to generate the covers (e.g., the most significant v bits, the last v significant bits, etc.).

[0252] After the cover properties are executed a cover generator 1520, the cover results 1525 are passed to stimuli generator 1540 to report examples produced for each of the passing covers as generated stimuli 1545. The examples may be the input bits that were utilized to generate the passing covers. As such, the covers can be used to identify potential cases (e.g., corner cases) for testing and provide example input bits that can then be used as stimuli to stress these potential cases.

[0253] In some implementations, possible optimization can be applied to the formal verification tool 1500 to potentially improve the efficiency of the formal verification tool 1500. In one example, frameworks could be rewritten to introduce additional muxes into the hardware design 1502 (e.g., into the RTL code) without changing the functionality of the hardware design 1502. This may be performed as preprocessing step.

[0254] In another example, for the generation of different counterexamples 1535, the cover generator 1520 may write cover properties at every operator node, including internal operator nodes, of the graph representation structure 1518, instead of at the leaf nodes. The cover generator could then perform a preorder traversal of the graph structure tree to assess the cover properties at each internal node starting

from the root node down. If the cover generator 1520 determines that a cover fails for an internal node, then the cover generator 1520 can discontinue analyzing the subtree rooted at that internal node. This is because, if the cover fails at an internal node, then all of the covers generated in the subtree of the internal node would fail as well, as they are adding extra conditions with an AND operator. As a result, implementations do not generate covers at the other nodes in the subtree as they would not provide any useful additional counterexamples.

[0255] For example, FIG. 16B illustrates a graph representation structure 1610 where a cover property written for an internal node of the graph representation structure fails, in accordance with implementations herein. In one implementation, the graph representation structure 1610 may be the same as graph representation structure 1518 of FIG. 15. Graph representation structure 1610 may include operator nodes 1, 2, 3, 4, 5, and 6. In this case operator node 1 is the root node of the graph representation structure 1610. Graph representation structure 1610 may also include leaf nodes L1, L2, L3, L4, L5, L6, and L7. The operator nodes are connected to the leaf nodes via edges extending indicating the operand of the originating operator node. In the example shown in FIG. 16C, cover properties may be written for each internal operator node 1-6 of the graph representation structure 1610. As depicted, the cover property written for operator node 3, shown as Cover ($S \neq A+B$) & ($A[1:0] = 1$) fails. As such, all of the operator nodes and leaf nodes of the subtree of the internal operator node 3, shown in the dashed oval 1615 in FIG. 16B, are also considered as having failing covers.

[0256] A similar approach can be used for the generation of stimuli. If the cover fails for an internal node, implementations do not continue to analyze the same cover on the subtree rooted at that node. For example, FIG. 16C illustrates a graph representation structure 1620 where a cover property written for a possible output bit at an internal node of the graph representation structure fails, in accordance with implementations herein. In one implementation, the graph representation structure 1620 may be the same as graph representation structure 1518 of FIG. 15. In one implementation, graph representation structure 1620 may be the same as graph representation structure 1610 of FIG. 16B. Graph representation structure 1620 may include operator nodes 1, 2, 3, 4, 5, and 6. In this case operator node 1 is the root node of the graph representation structure 1620. Graph representation structure 1620 may also include leaf nodes L1, L2, L3, L4, L5, L6, and L7. The operator nodes are connected to the leaf nodes via edges extending indicating the operand of the originating operator node.

[0257] In the example shown in FIG. 16C, cover properties may be written for possible output bits of each internal operator node 1-6 of the graph representation structure 1620. As depicted, the cover properties written for operator node 3 include a cover property for the output bit $O_t = 1$ and $O_t = 0$. The cover property for $O_t = 1$ is depicted as failing. As such, all of the operator nodes and leaf nodes of the subtree of the internal operator node 3, shown in the dashed oval 1625 in FIG. 16B, are also considered as having failing covers for the cover written for $O_t = 1$. Note however, that the cover property written for $O_t = 0$ passes at internal operator node 3. As such, the cover generator should con-

tinue to generator this cover property for the other operator nodes and leaf nodes of the subtree originating at internal operator node 3.

[0258] In another example case for optimizing the formal verification tool 1500, there may be cases when all of the cover properties of the leaf nodes in a subtree pass. When this outcome is detected, the formal verification tool 1500 can provide additional information, such as a summary listing all of the passing covers and identifying the internal operator node at the root of the subtree having all of the passing cover properties. In this case, it may be that the leaf nodes of the subtree may have different root causes to generate a counterexample, or on the other hand, may be all affected by a wrong value propagated down from an upper level (e.g., the root internal operator node) of the subtree.

[0259] For example, FIG. 16D illustrates a graph representation structure 1630 where all cover properties written for leaf nodes of a subtree of the graph representation structure pass, in accordance with implementations herein. In one implementation, the graph representation structure 1630 may be the same as graph representation structure 1518 of FIG. 15, and may be the same as graph representation structure 1610 of FIG. 16B. Graph representation structure 1630 may include operator nodes 1, 2, 3, 4, 5, and 6. Operator node 1 is the root node of the graph representation structure 1630. Graph representation structure 1630 may also include leaf nodes L1, L2, L3, L4, L5, L6, and L7. The operator nodes are connected to the leaf nodes via edges extending indicating the operand of the originating operator node.

[0260] In the example shown in FIG. 16D, it is determined that cover properties written for leaf nodes L2, L6, and L7 fail. Furthermore, it is determined that cover properties written for leaf nodes L1, L3, L4, and L5 pass. In this case, the passing covers can be analyzed to identify if a subset (or the entire set) of the passing leaf nodes are part of a same subtree of the graph representation structure 1630. As shown, leaf nodes L1, L3, and L4 are all part of the same subtree shown in dashed oval 1635 originating at internal operator node 2. As such, the formal verification tool may provide additional information, such as a summary listing L1, L3, and L4 as passing leaf nodes that are all part of a same subtree originating at internal operator node 2. A designer may be able to utilize this information for debug purposes as an initial starting point to identify if a root cause of a failure of the verification property tested by the cover property occurs at the internal operator node.

[0261] FIG. 17 is a flow diagram illustrating an embodiment of a method 1700 for generating multiple distinct paths counterexamples. Method 1700 may be performed by processing logic that may comprise hardware (e.g., circuitry, dedicated logic, programmable logic, etc.), software (such as instructions run on a processing device), or a combination thereof. The process of method 1700 is illustrated in linear sequences for brevity and clarity in presentation; however, it is contemplated that any number of them can be performed in parallel, asynchronously, or in different orders. Further, for brevity, clarity, and ease of understanding, many of the components and processes described with respect to FIGS. 1-16 may not be repeated or discussed hereafter. In one implementation, a processor executing a formal verification tool, such as formal verification tool 1500 of FIG. 15, may perform method 1700.

[0262] Method 1700 begins at processing block 1710 where a processor may extract a graph representation structure from a data flow of a hardware design of an integrated circuit. In one implementation, the graph representation structure includes operator nodes and leaf nodes connected by edges. Then, at block 1720, the processor may generate, for each leaf node of the leaf nodes of the graph representation structure, a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the implementation design having the set of space constraints applied. In one implementation, the set of space constraints define a path through the graph representation structure from a root node to the leaf node.

[0263] Subsequently, at block 1730, the processor may execute, as a single run of a verification tool, a set of cover properties comprising the cover property for each leaf node of the graph representation structure. Lastly, at block 1740, the processor may report results of execution of the set of cover properties as a set of counterexamples for the implementation design. In one implementation, the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the implementation design.

[0264] FIG. 18 is a flow diagram illustrating an embodiment of a method 1800 for stimuli generation based on multiple distinct paths counterexamples. Method 1800 may be performed by processing logic that may comprise hardware (e.g., circuitry, dedicated logic, programmable logic, etc.), software (such as instructions run on a processing device), or a combination thereof. The process of method 1800 is illustrated in linear sequences for brevity and clarity in presentation; however, it is contemplated that any number of them can be performed in parallel, asynchronously, or in different orders. Further, for brevity, clarity, and ease of understanding, many of the components and processes described with respect to FIGS. 1-17 may not be repeated or discussed hereafter. In one implementation, a processor executing a formal verification tool, such as formal verification tool 1500 of FIG. 15, may perform method 1800.

[0265] Method 1800 begins at processing block 1810 where a processor may extract a graph representation structure from a data flow of a hardware design of an integrated circuit. In one implementation, the graph representation structure includes operator nodes and leaf nodes connected by edges. Then, at block 1820, the processor may generate, for each leaf node of the leaf nodes of the graph representation structure, at least one cover property for the leaf that is to provide a set of space constraints of the graph representation structure and to query whether a particular output of the leaf node is satisfied by the implementation design when having the set of space constraints applied. In one implementation, the set of space constraints define a path through the graph representation structure from a root node to the leaf node.

[0266] Subsequently, at block 1830, the processor may execute, as a single run of a verification tool, a set of cover properties comprising the cover property for each leaf node of the graph representation structure. At block 1840, the processor may identify, for each of the set of cover properties, a set of inputs for the at least one cover property that generates a passing result responsive to execution of the set of cover properties. Lastly, at block 1850, the processor may report results of execution of the set of cover properties. In one implementation, the results may include the set of inputs

identified for each of the at least one cover properties that generates the passing results, where the set of inputs are utilized as stimuli for the implementation design.

[0267] The following examples pertain to further embodiments. Example 1 is an apparatus to facilitate distinct paths counterexamples and stimuli generation. The apparatus of Example 1 includes a processor comprising processing circuitry to: generate a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges; for each leaf node of the leaf nodes of the graph representation structure, generate a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; execute, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and report results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

[0268] In Example 2, the subject matter of Example 1 can optionally include wherein the processing circuitry to generate the graph representation structure further comprises the processing circuitry to assign identification to each of the operator nodes and the leaf nodes of the graph representation structure. In Example 3, the subject matter of any one of Examples 1-2 can optionally include wherein the processing circuitry to generate the cover property further comprises the processing circuitry to: generate at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and identify a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties; wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

[0269] In Example 4, the subject matter of any one of Examples 1-3 can optionally include wherein the at least one cover property comprises multiple cover properties generated for every possible outcome of an output of the leaf node. In Example 5, the subject matter of any one of Examples 1-4 can optionally include wherein the at least one cover property comprises multiple cover properties generated for possible outcomes of an output of the leaf node as bounded by a threshold.

[0270] In Example 6, the subject matter of any one of Examples 1-5 can optionally include wherein the processing circuitry is further to: execute the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and responsive to the cover property failing on one of the internal operator nodes on the path, discontinue the execution of the cover property on the operator nodes and the leaf node of the path. In Example 7, the subject matter of any one of Examples 1-6 can optionally include wherein the processing circuitry is further to: execute the cover property for

each internal operator node of the graph representation structure; and responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicate as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

[0271] In Example 8, the subject matter of any one of Examples 1-7 can optionally include wherein the graph representation structure comprises a multiplexer (mux) tree structure. In Example 9, the subject matter of any one of Examples 1-8 can optionally include wherein the processor comprises a graphics processing unit (GPU).

[0272] Example 10 is a method for facilitating distinct paths counterexamples and stimuli generation. The method of Example 10 can include generating, by processing circuitry, a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges; for each leaf node of the leaf nodes of the graph representation structure, generating, by the processing circuitry, a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; executing, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and reporting results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

[0273] In Example 11, the subject matter of Example 10 can optionally include wherein the processing circuitry to generate the graph representation structure further comprises the processing circuitry to assign identification to each of the operator nodes and the leaf nodes of the graph representation structure. In Example 12, the subject matter of Examples 10-11 can optionally include further comprising: generating at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and identifying a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties; wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

[0274] In Example 13, the subject matter of Examples 10-12 can optionally include wherein the at least one cover property comprises multiple cover properties generated for at least one of every possible outcome of an output of the leaf node or all possible outcomes of the output of the leaf node as bounded by a threshold.

[0275] In Example 14, the subject matter of Examples 10-13 can optionally include further comprising: executing the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and responsive to the cover property failing on one of the internal operator nodes on the path, discontinuing the execution of the cover property on the operator nodes and the leaf node of the path. In Example 15, the

subject matter of Examples 10-14 can optionally include further comprising: executing the cover property for each internal operator node of the graph representation structure; and responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicating as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

[0276] Example 16 is a non-transitory computer-readable storage medium for facilitating distinct paths counterexamples and stimuli generation. The non-transitory computer-readable storage medium of Example 16 having instructions stored thereon, which when executed by one or more processors, cause the processors to: generating, by the one or more processors, a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges; for each leaf node of the leaf nodes of the graph representation structure, generating a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; executing, as a single run of a verification tool provided by the one or more processors, a set of cover properties comprising the cover property for each of the leaf nodes; and reporting results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

[0277] In Example 17, the subject matter of Example 16 can optionally include wherein the operations further comprise: generating at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and identifying a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties; wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

[0278] In Example 18, the subject matter of Examples 16-17 can optionally include wherein the at least one cover property comprises multiple cover properties generated for at least one of every possible outcome of an output of the leaf node or all possible outcomes of the output of the leaf node as bounded by a threshold. In Example 19, the subject matter of Examples 16-18 can optionally include wherein the operations further comprise: executing the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and responsive to the cover property failing on one of the internal operator nodes on the path, discontinuing the execution of the cover property on the operator nodes and the leaf node of the path.

[0279] In Example 20, the subject matter of Examples 16-19 can optionally include wherein the operations further comprise: executing the cover property for each internal operator node of the graph representation structure; and responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes,

indicating as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

[0280] Example 21 is a system for facilitating distinct paths counterexamples and stimuli generation. The system of Example 21 can optionally include a memory; and a processor communicably coupled to the memory and comprising processing circuitry to: generate a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges; for each leaf node of the leaf nodes of the graph representation structure, generate a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; execute, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and report results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

[0281] In Example 22, the subject matter of Example 21 can optionally include wherein the processing circuitry to generate the graph representation structure further comprises the processing circuitry to assign identification to each of the operator nodes and the leaf nodes of the graph representation structure. In Example 23, the subject matter of any one of Examples 21-22 can optionally include wherein the processing circuitry to generate the cover property further comprises the processing circuitry to: generate at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and identify a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties; wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

[0282] In Example 24, the subject matter of any one of Examples 21-23 can optionally include wherein the at least one cover property comprises multiple cover properties generated for every possible outcome of an output of the leaf node. In Example 25, the subject matter of any one of Examples 21-24 can optionally include wherein the at least one cover property comprises multiple cover properties generated for possible outcomes of an output of the leaf node as bounded by a threshold.

[0283] In Example 26, the subject matter of any one of Examples 21-25 can optionally include wherein the processing circuitry is further to: execute the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and responsive to the cover property failing on one of the internal operator nodes on the path, discontinue the execution of the cover property on the operator nodes and the leaf node of the path. In Example 27, the subject matter of any one of Examples 21-26 can optionally include wherein the processing circuitry is further to: execute the cover property

for each internal operator node of the graph representation structure; and responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicate as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

[0284] In Example 28, the subject matter of any one of Examples 21-27 can optionally include wherein the graph representation structure comprises a multiplexer (mux) tree structure. In Example 29, the subject matter of any one of Examples 21-28 can optionally include wherein the processor comprises a graphics processing unit (GPU).

[0285] Example 30 is an apparatus for facilitating distinct paths counterexamples and stimuli generation, comprising means for generating a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges; for each leaf node of the leaf nodes of the graph representation structure, means for generating a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node; means for executing, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and means for reporting results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design. In Example 31, the subject matter of Example 30 can optionally include the apparatus further configured to perform the method of any one of the Examples 11 to 15.

[0286] Example 32 is at least one machine readable medium comprising a plurality of instructions that in response to being executed on a computing device, cause the computing device to carry out a method according to any one of Examples 10-15. Example 33 is an apparatus for facilitating distinct paths counterexamples and stimuli generation, configured to perform the method of any one of Examples 10 to 15. Example 34 is an apparatus for facilitating distinct paths counterexamples and stimuli generation, comprising means for performing the method of any one of Examples 10 to 15. Specifics in the Examples may be used anywhere in one or more embodiments.

[0287] The foregoing description and drawings are to be regarded in an illustrative rather than a restrictive sense. Persons skilled in the art will understand that various modifications and changes may be made to the embodiments described herein without departing from the broader spirit and scope of the features set forth in the appended claims.

What is claimed is:

1. A processor comprising:
processing circuitry to:

generate a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges;

for each leaf node of the leaf nodes of the graph representation structure, generate a cover property for the leaf node that is to provide a set of space

constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node;

execute, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and

report results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

2. The processor of claim 1, wherein the processing circuitry to generate the graph representation structure further comprises the processing circuitry to assign identification to each of the operator nodes and the leaf nodes of the graph representation structure.

3. The processor of claim 1, wherein the processing circuitry to generate the cover property further comprises the processing circuitry to:

- generate at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and
- identify a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties;

wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

4. The processor of claim 3, wherein the at least one cover property comprises multiple cover properties generated for every possible outcome of an output of the leaf node.

5. The processor of claim 3, wherein the at least one cover property comprises multiple cover properties generated for possible outcomes of an output of the leaf node as bounded by a threshold.

6. The processor of claim 1, wherein the processing circuitry is further to:

- execute the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and
- responsive to the cover property failing on one of the internal operator nodes on the path, discontinue the execution of the cover property on the operator nodes and the leaf node of the path.

7. The processor of claim 1, wherein the processing circuitry is further to:

- execute the cover property for each internal operator node of the graph representation structure; and
- responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicate as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

8. The processor of claim 1, wherein the graph representation structure comprises a multiplexer (mux) tree structure.

9. The processor of claim 1, wherein the processor comprises a graphics processing unit (GPU).

10. A method comprising:

- generating, by processing circuitry, a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges;

- for each leaf node of the leaf nodes of the graph representation structure, generating, by the processing circuitry, a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node;

- executing, as a single run of a verification tool provided by the processing circuitry, a set of cover properties comprising the cover property for each of the leaf nodes; and

- reporting results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

11. The method of claim 10, wherein the processing circuitry to generate the graph representation structure further comprises the processing circuitry to assign identification to each of the operator nodes and the leaf nodes of the graph representation structure.

12. The method of claim 10, further comprising:

- generating at least one cover property for the leaf node to identify whether a particular output of the leaf node is satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and

- identifying a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties;

- wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

13. The method of claim 12, wherein the at least one cover property comprises multiple cover properties generated for at least one of every possible outcome of an output of the leaf node or all possible outcomes of the output of the leaf node as bounded by a threshold.

14. The method of claim 10, further comprising:

- executing the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and

- responsive to the cover property failing on one of the internal operator nodes on the path, discontinuing the execution of the cover property on the operator nodes and the leaf node of the path.

15. The method of claim 10, further comprising:

- executing the cover property for each internal operator node of the graph representation structure; and

- responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicating as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

16. A non-transitory computer-readable medium having instructions stored thereon, which when executed by one or more processors, cause the processors perform operations comprising:

generating, by the one or more processors, a graph representation structure from a data flow of a hardware design of an integrated circuit, the graph representation structure comprising operator nodes and leaf nodes connected by edges;

for each leaf node of the leaf nodes of the graph representation structure, generating a cover property for the leaf node that is to provide a set of space constraints of the graph representation structure and to query whether a verification property is satisfied by the hardware design having the set of space constraints applied, wherein the set of space constraints define a path through the graph representation structure from a root node to the leaf node;

executing, as a single run of a verification tool provided by the one or more processors, a set of cover properties comprising the cover property for each of the leaf nodes; and

reporting results of execution of the set of cover properties as a set of counterexamples for the hardware design, wherein the set of counterexamples is to provide a plurality of distinct cases where the verification property failed in the hardware design.

17. The non-transitory computer-readable medium of claim **16**, wherein the operations further comprise:

generating at least one cover property for the leaf node to identify whether a particular output of the leaf node is

satisfied, where the at least one cover property to query whether the particular output occurs at the leaf node; and

identifying a set of inputs for the at least one cover property that generates a passing result responsive to the execution of the set of cover properties;

wherein reporting results of the execution of the set of cover properties comprises providing the set of inputs as stimuli for the hardware design.

18. The non-transitory computer-readable medium of claim **17**, wherein the at least one cover property comprises multiple cover properties generated for at least one of every possible outcome of an output of the leaf node or all possible outcomes of the output of the leaf node as bounded by a threshold.

19. The non-transitory computer-readable medium of claim **16**, wherein the operations further comprise:

executing the cover property for each internal operator node of the graph representation structure on the path from the root node to the leaf node; and

responsive to the cover property failing on one of the internal operator nodes on the path, discontinuing the execution of the cover property on the operator nodes and the leaf node of the path.

20. The non-transitory computer-readable medium of claim **16**, wherein the operations further comprise:

executing the cover property for each internal operator node of the graph representation structure; and

responsive to the cover property passing on all of the leaf nodes corresponding to one of the internal operator nodes, indicating as part of the results of the execution a summary of passing cover properties associated with the one of the internal operator nodes.

* * * * *