

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250279964

Kind Code

A1

Publication Date

September 04, 2025

Inventor(s)

Williams; Matthew Robert et al.

SYSTEM AND METHOD FOR REDUCING BANDWIDTH USAGE OF A NETWORK

Abstract

A method of reducing the bandwidth usage of a network comprises intercepting traffic between a TCP server and a TCP client using TCP protocols that use client acknowledgements; identifying client acknowledgements from the TCP protocols; identifying the sequence number of a last received client acknowledgements from the intercepted traffic; identifying the sequence number of a last sent client acknowledgement from the intercepted traffic; calculating an unacknowledged byte value based on the difference between the last received client acknowledgement sequence number and the last sent client acknowledgement sequence number; comparing the calculated unacknowledged byte value with a predetermined threshold value, to determine whether the calculated unacknowledged byte value is at least as great as the predetermined threshold value; and transmitting the identified client acknowledgements into the network when the compared unacknowledged byte value is at least as great as the predetermined threshold value.

Inventors: Williams; Matthew Robert (Kanata, CA), Nebogatov; Yuri (Kanata, CA), Klemetti; Miika Anttoni (Kanata, CA)

Applicant: Adeia Media Holdings LLC (San Jose, CA)

Family ID: 53173203

Appl. No.: 18/977708

Filed: December 11, 2024

Related U.S. Application Data

parent US continuation 18131046 20230405 parent-grant-document US 12206588 child US 18977708

parent US continuation 16829308 20200325 parent-grant-document US 11671377 child US 18131046

parent US continuation 16044695 20180725 parent-grant-document US 10645012 child US 16829308

parent US continuation 15413485 20170124 parent-grant-document US 10057178 child US 16044695
parent US continuation 14606834 20150127 parent-grant-document US 9590913 child US 15413485
parent US continuation 13022073 20110207 parent-grant-document US 8717900 child US 14219151
parent US continuation-in-part 14219151 20140319 parent-grant-document US 9647945 child US 14606834

Publication Classification

Int. Cl.: **H04L47/2466** (20220101); **H04L47/193** (20220101); **H04L47/27** (20220101); **H04L47/32** (20220101); **H04L47/40** (20220101); **H04L67/59** (20220101); **H04L69/163** (20220101); **H04L69/22** (20220101); **H04W28/02** (20090101); **H04W80/06** (20090101)

U.S. Cl.:

CPC **H04L47/2466** (20130101); **H04L47/193** (20130101); **H04L47/27** (20130101); **H04L47/323** (20130101); **H04L47/40** (20130101); **H04L67/59** (20220501); **H04L69/163** (20130101); **H04W28/0273** (20130101); **H04L69/22** (20130101); **H04W80/06** (20130101)

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION [0001] This application is a continuation of, and claims priority to, U.S. patent application Ser. No. 16/829,308, filed Mar. 25, 2020, now allowed, which is a continuation of, and claims priority to, U.S. patent application Ser. No. 16/044,695, filed Jul. 25, 2018, now U.S. Pat. No. 10,645,012, which is a continuation of, and claims priority to, U.S. patent application Ser. No. 15/413,485, filed Jan. 24, 2017, now U.S. Pat. No. 10,057,178, which is a continuation of, and claims priority to, U.S. patent application Ser. No. 14/606,834, filed Jan. 27, 2015, now U.S. Pat. No. 9,590,913, which is a continuation-in-part of U.S. patent application Ser. No. 14/219,151, filed Mar. 19, 2014, now U.S. Pat. No. 9,647,945, which is a continuation of U.S. patent application Ser. No. 13/022,073, filed Feb. 7, 2011, now U.S. Pat. No. 8,717,900, each of which is hereby incorporated by reference herein in its entirety.

FIELD OF THE INVENTION

[0002] The present invention relates to communication data networks. More specifically, the present invention relates to systems and methods for increasing the throughput of data transmissions through a network as seen from the edge of the network.

BACKGROUND OF THE INVENTION

[0003] The Transmission Control Protocol (TCP) is one of the core protocols of the Internet Protocol suite. The TCP provides reliable, ordered delivery of a stream of bytes from one computer application to another computer application and provides the transport layer for many fundamental Internet applications such a web traffic, mail, and file transfer. TCP has been optimized for wired transmission. As such, packet loss is assumed to be the result of network congestion, and TCP has many mechanisms geared to deal with congestion that is triggered by packet loss. These mechanisms may be erroneously applied in a wireless networking scenario. Packet loss is more likely to have occurred due to some sporadic event such as interference or shading. If a computer application leveraging TCP for transport experiences packet loss in a wireless network, congestion

control mechanisms actually lead to network underutilization.

[0004] In addition to the basic issues described above, some networks have additional challenges. Some networks may be bandwidth constrained, requiring minimization of bandwidth usage. Some networks, such as power-line, wireless or WIFI, are half-duplex such that each end uses the same frequency to transmit and receive. In this case, a handshake or other back-off/contention mechanism is required to ensure that each member of the half-duplex network can transmit without colliding with another end station. These collision avoidance mechanisms are generally optimized for larger packet sizes, which influence the setting of their retry and back-off timers. This is particularly problematic for TCP Acknowledgement (ACK) packets, which are sent often by a computer application as it is receiving data from another computer application. If such a constrained network is involved in the path of this data stream, the collision avoidance handshake is invoked often as data is acknowledged. Since the TCP ACKs are small, they violate the assumptions of the collision avoidance protocols leading to additional underutilization.

SUMMARY

[0005] In accordance with one embodiment, a method of reducing the bandwidth usage of a network comprises intercepting traffic between a TCP server and a TCP client using TCP protocols that use client acknowledgements; identifying client acknowledgements from the TCP protocols; identifying the sequence number of a last received client acknowledgements from the intercepted traffic; identifying the sequence number of a last sent client acknowledgement from the intercepted traffic; calculating an unacknowledged byte value based on the difference between the last received client acknowledgement sequence number and the last sent client acknowledgement sequence number; comparing the calculated unacknowledged byte value with a predetermined threshold value, to determine whether the calculated unacknowledged byte value is at least as great as the predetermined threshold value; and transmitting the identified client acknowledgements into the network when the compared unacknowledged byte value is at least as great as the predetermined threshold value. In one implementation, the intercepting step is performed at an intercepting device located between a TCP server and at least one TCP client. The last sent client acknowledgement is preferably sent from at least one of the TCP clients.

[0006] In accordance with another embodiment, a system for reducing the bandwidth usage of a network, the method comprises a TCP server and a TCP client using TCP protocols that use client acknowledgements; and an intercepting network device located between the TCP server and the TCP client and adapted to receive the client acknowledgements flowing between the TCP server and the TCP client, and to transmit all the client acknowledgements into the network. In one implementation, the sequence numbers of a last received client acknowledgement and a last sent client acknowledgement sent from a client are identified at the intercepting network device, and an unacknowledged byte value based on the difference between the sequence numbers is compared to a predetermined threshold value. The client acknowledgements are preferably transmitted into the network from the intercepting device when the unacknowledged byte value is equal to or greater than the predetermined threshold value.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0007] Embodiments of the present invention will now be described, by way of example only, with reference to the attached Figures, wherein:

[0008] FIG. 1 is a diagrammatic illustration of a system in which a client and a server communicate via a network using TCP.

[0009] FIG. 2 illustrates a TCP Segment and TCP ACK exchange between the client and server in the system of FIG. 1.

[0010] FIG. 3 is a diagrammatic illustration of a modification of the system of FIG. 1 that includes wireless network elements.

[0011] FIG. 4 is an illustration of wireless collision avoidance messages transmitted in the system of FIG. 3.

[0012] FIG. 5 is a diagrammatic illustration of a modification of the system of FIG. 3 that includes intercepting network devices.

[0013] FIG. 6 is a block diagram illustrating the components in a intercepting network device used in the system of FIG. 5.

[0014] FIG. 7 is a diagrammatic illustration of an intercepting network device aggregating ACKs from a TCP connection into a single ACK packet.

[0015] FIG. 8 is an illustration of a state machine that implements the aggregation of TCP ACKs into a single ACK packet.

[0016] FIG. 9 is a flow chart for the state entry logic of the Idle State shown in FIG. 8.

[0017] FIG. 10 is a flow chart of the event handling code for the Idle State shown in FIG. 8.

[0018] FIG. 11 is a flow chart for the state entry logic of the ACK Pending State shown in FIG. 8.

[0019] FIG. 12 is a flow chart of the event handling code for the ACK Pending State shown in FIG. 8.

[0020] FIG. 13 is a flow chart for the state entry logic of the Send ACK State shown in FIG. 8.

[0021] FIG. 14A illustrates the handling of Selective ACK packets by an intercepting network device encoder.

[0022] FIG. 14B is a flow chart showing unAckedBytes count calculations used in the state machine shown in FIG. 8.

[0023] FIG. 15 illustrates the encoding of multiple ACKs from a TCP connection into a single packet by an intercepting network device.

[0024] FIG. 16 illustrates the decoding of multiple ACKs from a TCP connection into a single packet by an intercepting network device.

[0025] FIG. 17 is an illustration of a state machine that encodes multiple TCP ACKs from a TCP connection into a single packet.

[0026] FIG. 18 is a flow chart of the state entry logic of the ACK Pending State shown in FIG. 17.

[0027] FIG. 19 is a flow chart of the event handling logic of the ACK Pending State shown in FIG. 17.

[0028] FIG. 20 is a flow chart of the state entry logic of the Send ACK State shown in FIG. 17.

[0029] FIG. 21 is a modification of the system of FIG. 5 that includes multiple clients able to use IND 70.

[0030] FIG. 22 is an illustration of the encoding algorithm for multiple TCP connections combined with the aggregation of TCP ACKs from a single connection.

[0031] FIG. 23 is a flow chart of the algorithm for an IND encoder for splitting up a single TCP ACK into multiple ACKs.

[0032] FIG. 24 is a flow chart of a mechanism for splitting a TCP ACK into multiple TCP ACKs.

[0033] FIG. 25 is a continuation of the flow chart of FIG. 24.

[0034] FIG. 26 illustrates an encoding algorithm for splitting up a single TCP ACK into multiple ACKs using the network of FIG. 5.

DETAILED DESCRIPTION OF THE ILLUSTRATED EMBODIMENTS

[0035] Although the invention will be described in connection with certain preferred embodiments, it will be understood that the invention is not limited to those particular embodiments. On the contrary, the invention is intended to cover all alternatives, modifications, and equivalent arrangements as may be included within the spirit and scope of the invention as defined by the appended claims. The invention will be described in the context of a wireless network, but it will be understood that it applies to any bandwidth constrained networks or to half-duplex networks.

[0036] The Transmission Control Protocol (TCP) is one of the core Internet protocols providing

reliable, ordered delivery of a stream of bytes from one computer application to another computer application. FIG. 1 illustrates the network elements involved in TCP operation. The first element is the TCP Server **10**, which is characterized as a device that listens to a partially opened TCP port. The server sits on this port until other network elements, such as TCP clients, interact with the server using the open TCP port. A TCP Client is defined as a device, which requires reliable, ordered delivery of a stream of bytes from the server **10**. The client **30** initiates a connection with the server **10** that is uniquely identified by the IP address and TCP port number of the server **10** in combination with the IP address and TCP port number of the client **30**. Once the connection is established, reliable data flow is provided in both directions. The direction actually used is determined by the computer applications on the server **10** and the client **30**. In the case of web traffic, the flow is in general from server to client. File Transfer Protocol (FTP) allows for the transfer in either direction. The client **30** can receive a file from the server **10** or put a file on the server **10**.

[0037] The attributes of the network **20** are intentionally ambiguous. The only requirement is to ensure that TCP can work between the server **10** and the client **30**. This means that the client **30** is able to determine the IP Address of the server **10**, and the client **30** knows which TCP port on the server **10** is open for the client's application. Then, the network **20** needs to provide IP connectivity between the server **10** and the client **30**.

[0038] In FIG. 2, the application of the server **10** is sending a flow of bytes to the application of the client **10**. This flow begins with the client advertising a receive window size of 3500 bytes to the server in message **200**. This message is part of the flow control mechanism provided by TCP to ensure that the client **30** is not overwhelmed with data from the server **10**. When the message **200** arrives at the server **10**, three TCP Segments are sent to the client **30** (messages **210** to **212**) to fill the open receive window of the client **30**. To provide reliable transfer of the data from the server **10**, the client **30** acknowledges the arrival of the data using TCP ACK messages. All of the bytes from the server are numbered, and these additional bytes are acknowledged using the sequence number in the ACK message. In this case, the client **30** creates three TCP ACK messages **220** to **222**. These messages are the focus of the mechanisms documented in this patent.

[0039] TCP ACK messages are problematic in half-duplex wireless networks that require collision avoidance mechanisms to ensure reliable data transfer. FIGS. 3 and 4 illustrate how these mechanisms affect TCP connections. The system of FIG. 3 adds a wireless LAN in the path between the server **10** and the client **30**. A wireless client **50** provides access to the wireless LAN on behalf of the network devices on its wired side including the client **30**. A wireless Access Point (AP) **40** connects many wireless clients (including the wireless client **50**) to the network **20**. Given the assumption that there are many wireless clients, a collision avoidance protocol is required to ensure that each wireless client is provided with equitable service on the wireless LAN. FIG. 4 illustrates an example of how this protocol can operate, leveraging an example implementation from Wi-Fi, IEEE 802.11. In this figure, the message **220** from FIG. 2 is traveling from the client **30** to the server **10**, which is traversing the wireless LAN between the wireless AP **40** and the wireless client **50**. When the wireless client **50** receives the TCP ACK in a message **220**, it may announce that it is "Ready to Send," using a message **400**. Assuming that this request is valid, the wireless AP **40** sends a "Clear to Send" indication to the wireless client **50** using a message **401** that allows the wireless client **50** to start sending the TCP ACK in a message **402**. When the TCP ACK is received by the wireless AP **40**, it sends a wireless Acknowledgement Frame in a message **403**. At this point, the TCP ACK proceeds onto the server **10** through the network **20**.

[0040] FIG. 4 illustrates a successful exchange allowing a wireless client to send the TCP ACK message **402** quickly. In many cases, the wireless client **50** will need to wait to send the message **402**. For example, another wireless client could already be sending when the message **220** arrives at the wireless client **50**. The wireless client **50** will need to wait until the Maximum Transmission Unit (MTU) length is sent by the other wireless client, which may be sending a smaller packet than

the MTU, or has already started sending when the wireless client **50** receives the TCP ACK message **220**. In either case, network utilization will suffer while access to the wireless LAN is negotiated by the wireless clients. Note that the wireless LAN between the wireless client **50** and the wireless AP **40** may not be on the only wireless link in the path between the client **30** and the server **10**. If multiple half-duplex wireless links are in this the path, network utilization will suffer further.

[0041] In this situation, an Intercepting Network Device (IND) can be added to the path book-ending the wireless LAN to improve network utilization. The system of FIG. 5 includes an IND on the path between the client **30** and the server **10**. An IND **70** is deployed between the client **30** and the wireless client **50** to provide functions that improve the network utilization for packets heading toward the wireless AP **40**. A second IND **60** between the network **20** and the wireless AP **40** provides functions to improve network utilization toward the wireless clients, such as wireless client **50**.

[0042] FIG. 6 illustrates an exemplary embodiment of the INDs shown in FIG. 5. FIG. 6 decomposes an IND **70** into modules necessary to implement the encoding protocol to improve wireless network performance. These modules can be resident in a single device, or can be distributed among several devices. As can be seen in FIG. 6, the IND **70** has a native interface **600**, an encoded interface **620**, a decoder **630** and an encoder **610**. The native interface **600** sends and receives standard packets to and from the client **30**. The encoded interface **620** sends and receives encoded packets to and from the wireless client **50**.

[0043] The decoder **630** receives encoded packets from the encoded interface **620** and produces standard packets for transmission to the client **30** by way of the native interface **600**. The encoder **610**, on the other hand, receives standard packets from the native interface **600** and produces encoded packets for transmission to the wireless client **50** by way of the encoded interface **620**. The specific algorithms used by the encoder and decoder are optional and programmable. Several embodiments are described below.

[0044] In one implementation, the IND **60** and the IND **70** buffer ACKs from a TCP connection and summarize the information in a single ACK message for the connection. Since the ACK information in TCP messages is cumulative, the algorithm can be very simple. FIG. 7 demonstrates how this operation is accomplished for the packet sequence illustrated in FIG. 2. In this example, the three ACKs sent from the client **30** (messages **220** to **222**) are buffered at the IND **70**. The information in the three ACKs is summarized in a single ACK which is then forwarded to the wireless client **50** in a message **701**. Then, the wireless client **50** negotiates with the wireless access point **40** to transmit this ACK over the wireless LAN. This negotiation, contained in the messages **710** to **713**, is the same as the flow shown in FIG. 4. Then the TCP ACK is forwarded into the network **20** using a message **714**. This flow demonstrates an important optimization provided by this algorithm. The number of collision avoidance negotiations is reduced by three to one, which reduces the overhead messaging, resulting in higher throughput.

[0045] FIG. 7 assumes that the client **30** is downloading a file from the server **10**, which is a valid assumption for many computer applications such as web traffic. However, the client may also be uploading data onto the server using TCP. For example, the client **10** may be uploading a photo onto a web server using FTP to backup or to share. In this case, the IND **60** shown in FIG. 5 can perform the same functions as the IND **70**, but in the reverse direction. The encoded interface for the IND **60** is connected to the wireless AP **40**, and its native interface is connected to the network **20**. The IND **60** then can summarize the TCP ACKs flowing from the server **10** toward the client **30** to improve the utilization of the wireless LAN between the wireless AP **40** and the wireless client **50**.

[0046] To summarize the TCP ACKs for a connection, the encoders of the INDs (**60** and **70**) implement a state machine for each TCP connection. An embodiment of this state machine is illustrated in FIG. 8, which contains three states. The Idle state **800** is used to handle the time when

the IND has no ACKs to handle for this TCP connection. In the ACK Pending State, the IND has received one (Event **810**) or more (Event **811**) ACKs, but the conditions have not been achieved to summarize the ACKs for transmission. Once enough time has expired, or a threshold number of ACKs has been achieved (now defined as MAX ACKs), the state machine transitions to the Send ACK state **802** and transmits the ACK. This is a transitory state and once the transmission of the ACK is complete, the state machine transitions to the Idle state **800**.

[0047] FIGS. **9** and **10** illustrate detailed handling of the Idle State **800**. FIG. **9** provides the state entry functions for the Idle State **800**. It simply sets the current state of the state machine to the Idle State (**901**). FIG. **10** is a flowchart of the algorithm for event handling in the Idle State **800**. The algorithm starts by determining what type of event was received at step **1001**. If the event is an ACK being received for this connection, the state machine transitions to ACK Pending state **801** at step **1002**. Otherwise, an unexpected was received and an error event is generated at step **1003**.

[0048] FIGS. **11** and **12** illustrate detailed handling of the ACK Pending State **801**. FIG. **11** provides the state entry functions for the ACK Pending State **801**. It sets the current state to the ACK Pending State **801** at step **1101**, sets the current count of received ACKs to 1 at step **1102** and starts the ACK Timer at step **1103**. FIG. **12** is a flowchart of the algorithm for event handling in the ACK Pending State **801**. The algorithm starts by determining what type of event was received at step **1201**. If the event is an ACK being received for this connection, the current count of received ACKs is increased by 1 at step **1202** and is compared to MAX ACKs at step **1203**. If MAX ACKs has been reached, the state machine transitions to the Send ACK State **802** at step **1204**. If the timer expired event has been received, the state machine transitions to the Send ACK State **802** at step **1204**. Otherwise, an unexpected was received and an error event is generated at step **1204**.

[0049] FIG. **13** illustrates the functions performed in the Send ACK State **802**. When the state machine reaches this state, all the ACKs received are summarized into a single ACK. Assuming no packet loss has been experienced, the last sequence number received is placed in the summarized ACK. Then the summarized ACK is queued for transmission at step **1303**, and the state machine transitions to the Idle State **800**.

[0050] When the ACK is summarized, the INDs (**60** or **70**) may have to deal with selective ACKs (or SACKs). This option is defined in RFC 2018 and allows the receiver to acknowledge discontinuous blocks of packets that were received correctly, in addition to the sequence number of the last contiguous byte received successively, as in the basic TCP acknowledgment. The acknowledgement can specify a number of SACK blocks, where each SACK block is conveyed by the starting and ending sequence numbers of a contiguous range that the receiver correctly received. FIG. **14A** demonstrates how the encoder can deal with a SACK. The first ACK **1400** arrives from the client **30** with a sequence number of 4000 and a SACK block of 8500 to 10000. This is stored as the basis for the ACK which is forwarded onwards. The second ACK **1410** arrives with a selective ACK block of 5500 to 7000. The third ACK **1420** arrives with a SACK block of sequence number of 7000 which fills in the hole between 4000 and 5500. At this point the encoder **610** is ready to send the summarized TCP ACK which has a sequence number of 7000 and a SACK block of 8500 to 10000.

[0051] Note that the resulting encoded packet created by the IND **60** and the IND **70** is a standard TCP/IP packet. Therefore, no decoding of the packet is required, and the decoder function in the INDs (**60** and **70**) is NULL. Summarized TCP ACKs that arrive on the encoded interface of the IND **60** or the IND **70** can pass through unchanged. As well, no explicit signaling is required to create a connection between the INDs; they can proceed independently to improve network performance on the wireless LAN.

[0052] An implementation of this algorithm was used to measure the network utilization improvements leveraging a deployment similar to FIG. **5**. The only difference was that the server **10** was wired directly into the IND **60**, i.e., the network **20** was implemented as a physical wire. Then TCP flows were generated using the server **30**, and the performance was measured over the

wireless LAN. Flows were tested first with the INDs disabled and then with the INDs enabled with MAX ACK equal to 2. The resulting network utilization improvement is shown in the table below.

TABLE-US-00001	INDs Disabled (Mbps)	INDs Enabled (Mbps)	Improvement %
17.8	19.3	8.4	
18.2	20.3	11.5	
18.7	20.2	8.0	

[0053] As another embodiment, the timer **812** (FIG. **8**) and the ACK Count are complemented with another condition to move the state machine into the send state. FIG. **14B** illustrates a method using the sequence number difference of the last sent ACK and previously received ACK, which may be used to move the state machine into the send state before the timer or the ACK Count conditions are met. This embodiment uses the unacknowledged bytes count, referred to as unAckedBytes, wherein the packet sequence number difference between a previously sent ACK packet and the last received ACK from step **1440**) is used to calculate the unAckedBytes value at step **1442**. Once calculated, a query at step **1443** is used to determine if the unAckedBytes exceeds a predefined threshold maxUnAckedBytes. If the unAckedbytes are equal to or greater than this threshold, the state machine moves to the send ACK state at step **1444** allowing the packet to be forwarded into the network. If it is below the threshold, then the state machine stays in the Pending state **1441**, only moving to the send ACK state at step **1444** when the timer times out. The following formula is used to calculate the unAckedBytes:

[00001] $\text{unAckedBytes} = \text{lastReceivedAck} - \text{prevSentAck}$,

wherein the lastReceivedAck refers to the sequence number of the last received ACK, and the prevSentAck refers to the sequence number of the last sent ACK.

[0054] Therefore, the query at step **1443** is based on the following algorithm: if $\text{unAckedBytes} \geq \text{maxUnAckedBytes}$, go to send ACK state, otherwise go to PEND state, wherein maxUnAckedBytes refers to the maximum unAcked byte count threshold and is a function of the measured RTT and BDP.

[0055] The motivation for unacknowledged byte count is to optimize a TCP stack's sending rate. If there are too few summarizing ACKs out of the buffered ACKs, the sending rate of the TCP stack (that receives the ACKs) may be affected.

[0056] As another embodiment, summarizing ACKs can be delayed long enough that the cumulative ACK can be piggybacked with a data packet flowing in the reverse direction (all data packets carry the cumulative ACK number).

[0057] An alternate implementation of the INDs uses a method of aggregating the TCP ACK packets instead of summarizing them as described above. FIG. **15** illustrates how these aggregate packets can be assembled from multiple TCP ACKs in a single TCP connection. Step **1500** shows an initial buffer which holds the header **1501-H** for the next aggregate TCP packet produced. This header contains addressing to route packets between IND **60** and IND **70**. It also contains information on the TCP connection that is associated with this pairing. This allows the decoder to understand what is inside and to break apart the aggregate packet into its original components. In step **1510**, an entire TCP ACK is copied into the packet just after the aggregate packet header. This is completed in two parts. First, a header **1511-H** is added to identify this as the start of a TCP ACK; this header can contain a marker identifying the contents and the length of the packet. Second, the TCP ACK packet (**1511-P**) is copied. This process is repeated in step **1520** by adding another header **1521-H** to identify the ACK and packet **1521-P**. Once the packet is ready to be sent, an optional trailer **1531-T** can be included to carry association information to keep the INDs in sync or pertinent information for TCP ACKs included in the aggregate packet itself. This scheme is one of many that can be used to package multiple TCP ACKs into a single packet. Another option is to place all the ACKs back-to-back with the aggregate packet header containing markers to allow the packets to be split apart. The key is the aggregation of TCP ACKs.

[0058] FIG. **16** demonstrates how the aggregate packet is broken apart by the decoder on the other end of the wireless LAN. In step **1600**, the original packet is received, and then the header and trailer are removed in step **1610**. Then using the individual ACK headers **1101-H** and **1121-H**, the

individual TCP ACKs are identified at step **1620**. Once this is complete, the packets are queued for egress at step **1630**.

[0059] Since functionality is required on the encoder of one end of the connection and the decoder of the other end of the connection, the IND **50** and the IND **60** must create an association with each other to provide the aggregation function. Signaling mechanisms as described in copending U.S. application Ser. No. 14/718,650, filed Mar. 5, 2010, and incorporated by reference herein, can be used to create this association. The only difference is the mechanism which is being signaled. Instead of signaling a session to improve network reliability, it is signaling a TCP ACK aggregation session.

[0060] A state machine as shown in FIG. **17** implements the encoder functions to aggregate packets. It leverages the same states that were defined in FIG. **8**, but the state handling and transitions are different. The Idle State **1700** processing is identical to FIG. **8** and is illustrated in FIG. **9** and FIG. **10**. The ACK Pending State **1701** handling is different, and the entry function is illustrated in FIG. **18**. When a TCP ACK is received, the encoder enters the ACK Pending state for the associated TCP connection at step **1801**. The ACK Count is set to one at step **1802**, and the encoder prepares the aggregated ACK buffer for the ACK by adding a header at step **1803**. Then the ACK is added to the buffer at step **1804**. Following what was prescribed in FIG. **15**, a header is added before the TCP ACK packet, and the TCP ACK is added to the frame. This specific step may change depending on the packing algorithm. At this point, the final packet size is calculated at step **1805**, and the ACK Timer is started at step **1806**.

[0061] How subsequent ACKs are handled by the state machine in the ACK Pending State **1702** is illustrated in FIG. **19**. The state handling function starts by determining which event is being processed at step **1901**. If an ACK has been received, the state machine increases the ACK count by 1 at step **1902** and adds the ACK to the buffer at step **1903**. This follows the same procedure as described in connection with FIG. **15**. Then, the final size of the current aggregated packet is determined at step **1904**, and a comparison to the allowed MTU (MAX) on this TCP association between INDs is made at step **1905**. This MTU can be provisioned as part of policy for TCP connections, with the MTU for the specific interface used for transmission of the aggregated packet, negotiated as part of signaling or discovered as part of a Path MTU discovery process. If the aggregated packet exceeds the allowed MTU, the state machine proceeds to the Send ACK State **1701**; otherwise, the state machine goes back to sleep, waiting for the next event. Another valid event is the timer expiry. This timer length indicates that the IND encoder has waited long enough for subsequent ACKs and must send the packet. The duration of this timer can be provisioned, estimated from the round trip times of all packets snooped from this connection, or signaled as part of establishment of the IND's TCP association. If this event is received, the state machine proceeds to the Send ACK State **1702** at step **1907**. If any other event is received, it is unexpected and an error message is printed at step **1908**.

[0062] Upon entering the Send ACK state **1702**, the state machine needs to determine how to finalize the aggregated packet. As illustrated in FIG. **20**, this is determined by calculating the final MTU size at step **2002**, and then determining whether that size exceeds MAX at step **2003**. If the packet exceeds the allowed MTU, an ACK is removed from the packet at step **2007**, and the packet is then finalized as per FIG. **15** at step **2008** and sent at step **2009**. Then, since an additional ACK has been received, a new ACK buffer is created by transitioning back to the ACK Pending State **1701**. If the packet does not exceed the allowed MTU for this TCP association, the state machine can proceed directly to finalizing the aggregate packet as per FIG. **15** at step **2004**, and then sends the packet at step **2005**. Then the state machine transitions to the Idle State **1700** at step **2006**.

[0063] This is one implementation of the aggregation of TCP ACKs. A variation of this algorithm is to aggregate TCP ACKs from multiple clients instead of a single client. This scenario is illustrated in FIG. **21** where multiple clients **30-1** to **30-C** are speaking with multiple servers **10-1** to **10-S**. An IND **70** aggregates ACKs from all the clients **30** which are sending traffic across the wireless LAN

via IND **60**. Since the method for aggregation illustrated in FIG. **15** does not make any changes to the TCP packets themselves, the algorithms described for creating aggregate packets for a single TCP connection also applies to creating aggregate packets for multiple TCP connections. Therefore, FIGS. **17-20** completely specify the encoder for this variation. The only difference is that TCP connections associated with this aggregation channel will come and go. The signaling which controls the association must also handle the addition and deletion of TCP connections. [0064] Another variation allows for the combination of summarization of TCP ACKs with the aggregation of TCP ACKs for multiple TCP connections. FIG. **22** illustrates how this encoder operates. In this case, the encoder buffers TCP ACKs for multiple connections related to the clients **30-1** to **30-C**. A single TCP ACK (**2201**) is derived from each connection, and these ACKs are labeled **2201-1** to **2201-C**. An ACK for each active connection is then aggregated into a single packet and sent to the peer IND. The encoding of these packets is a variation of the state machine described in connection with FIGS. **17-20**. The only difference is how the procedure for adding a packet to the buffer is handled in FIG. **19** at step **1803**. Instead of just appending the ACK to the packet, the state machine must look for another ACK from the same connection. If one already exists, they are summarized together and placed into the packet. If a TCP ACK from the specific connection is not in the packet, then the TCP ACK is appended. Otherwise, the algorithm is identical. The decoding of the packets is the same as illustrated in FIG. **16**.

[0065] The methods described above improve the wireless LAN network utilization by influencing the behavior of the client and server by manipulating the flow control algorithm. An alternate method influences the congestion control mechanisms provided by TCP. One option is to break-up an ACK transmitted by the client **30** into several ACKs. This stimulates the congestion window on the server through its additive increase/multiplicative decrease algorithm to increase the congestion window. TCP implementations increase the congestion window by a fraction of the TCP maximum segment size for every ACK that is received. By increasing the number of ACKs, the congestion window opens faster, putting more TCP segments in flight.

[0066] This encoding algorithm can increase bandwidth utilization from the client **30** to the server **10**. Several parameters are defined to control the bandwidth used by the encoding algorithm. The first parameter is the MAXIMUM ACKs; which defines the maximum number of ACKs that can be generated from a single ACK. The MINIMUM SEGMENT SIZE determines the minimum difference between the sequence number of sequential ACKs. For example, MAXIMUM ACKs can be defined as 3 and MINIMUM SEGMENT SIZE can be defined as 500. If a packet arrives where the sequence difference between the current ACK and the previous ACK seen by the encoder is 2100 bytes, the encoder will produce 3 TCP ACKs each with a sequential difference of 700 bytes each. This encoding is limited by the MAXIMUM ACK parameter. If a TCP ACK arrives with a sequential difference of 400 bytes, this ACK is transmitted unchanged since it is less than the MINIMUM SEGMENT SIZE. Another method to control bandwidth is to define a duty cycle for the algorithm. This means that the encoding algorithm is enabled and disabled for defined periods. One method is to define these periods in terms of time defining a total period and a duty factor in percent. The duty factor defines the percentage of the period in which the encoding algorithm is operating. Another option is to define the periods in terms of packets. The total period is defined as PACKET_PERIOD. Within the PACKET_PERIOD, ENCODE_LIMIT defines the number of packets which are encoded.

[0067] The basic algorithm required to implement this function is provided in FIGS. **23-25**. FIG. **23** illustrates the generic algorithm for splitting TCP ACKs. It starts by inspecting the LAST_VALID_ACK parameter at step **2301**. This is used as the basis for splitting up the ACK before it is transmitted and ensures that each ACK is carrying valid information. Therefore, if this parameter is not valid, it is initialized to the current sequence number at step **2306**, and the packet period counter, ACK COUNT, is initialized to zero at step **2309**, and the ACK is transmitted at step **2307**. Assuming the LAST_VALID_ACK parameter is valid, the current ACK sequence number is

compared to the LAST_VALID_ACK at step **2302**. If the current ACK sequence is equal to the LAST_VALID_ACK parameter, the associated TCP connection has experience lost, and this is signaled back to the server **10**. Since this is the case, the ACK needs to get back to the server to initiate the fast recovery algorithms. Therefore, this ACK is sent immediately at step **2307**. If the current sequence is not equal to the LAST_VALID_ACK, the ACK Count is incremented at step **2303** and compared to the ENCODE_LIMIT at step **2304**. If the ACK Count is less than or equal to the ENCODE_LIMIT, the ACK is split up at step **2305**, the LAST_VALID_ACK is set to the current sequence number at step **2306**, and the ACKs are sent at step **2307**. If the ACK Count is determined to be greater than ENCODE_LIMIT at step **2304**, then the ACK Count is compared to the PACKET_PERIOD at step **2308**. If the ACK Count is less than the PACKET_PERIOD, then LAST_VALID_ACK is set to the current sequence number at step **2306** and the ACK is sent at step **2307**. Otherwise, the packet period is reset by setting ACK Count to Oat step **2309**.

[0068] The method for sending ACK in step **2307** is represented as a procedure because how this is handled is optional. In a wired environment, packets can be just enqueued directly on an interface for transmission. The increase of bandwidth by the splitting of ACKs may be easily handled by the bandwidth provided by wired LANs. Any loss can be handled by the natural assumptions of TCP operation. No decoding is required in this mode of operation since the encoded packets are still standard TCP/IP packets. However, if this method is used on its own in a half-duplex wireless environment, the increased number of packets can decrease wireless LAN utilization due to the collision avoidance algorithm running on the wireless LAN. To avoid this issue, the splitting up of a TCP ACK can be combined with the aggregation of TCP ACKs into a single packet sent across the wireless LAN. That way, the additional packets do not incur additional wireless transmission overhead to reach their goal. Therefore, the procedure in step **2307** in FIG. **23** can be implemented using the state machine in FIG. **17**. As illustrated in FIG. **16**, this mode requires an association between 2 INDs so these packets can be decoded. This association can be signaled dynamically as the INDs discover the TCP associations between a client and a server.

[0069] A method of splitting up a TCP ACK is illustrated in FIGS. **24** and **25**. The procedure starts by calculating the difference between the current sequence number and LAST_VALID_ACK at step **2401**. If the difference is determined to be less than the MINIMUM_SEGMENT_SIZE at step **2402**, only one ACK is produced at step **2403** by setting ACKs to 1 and Sequence Interval to Difference. If the difference is determined to be greater than MINIMUM_SEGMENT_SIZE at step **2402**, the number of ACKs is determined by dividing the difference by the MINIMUM_SEGMENT_SIZE at step **2404**. If the number of ACKs is determined to exceed the MAXIMUM_ACKs at step **2405**, then ACKs is set to MAXIMUM ACKs and the sequence interval is set to the difference divided by MAXIMUM ACKs at step **2406**. Otherwise, the sequence interval is set to MINIMUM_SEGMENT_SIZE at step **2407**.

[0070] Now that the number of TCP ACKs and the sequence interval between the ACKs has been determined, the TCP ACK can be segments. This algorithm is illustrated in FIG. **25**. To separate the ACK into multiple ACKs, a loop counter, current ACK is initialized to one, and the Last sequence is set to the LAST_VALID_ACK at step **2501**. Then the loop condition is checked by comparing Current ACK to ACK, the total number of TCP ACKs to be generated at step **2502**. If Current ACK does not equal ACK, then the sequence number of the first ACK is calculated using the Last sequence number plus the sequence interval at step **2503**, and this ACK packet is generated at step **2504**. Then the ACK is added to the list of ACKs to be sent at step **2505**, and the ACK Counter is incremented at step **2506**. This procedure repeats until Current ACK equals ACK at step **2502**. At this point, the loop exits, and the final ACK is generated by setting the sequence number of the final ACK to the current sequence number at step **2507** and adding this ACK to the list of ACKS to be transmitted at step **2508**.

[0071] FIG. **26** shows an example of functionality using a message **220** from FIG. **2** traveling from the client **20** to the server **10** in the network shown in FIG. **5**. FIG. **26** illustrates the algorithms of

FIGS. 23-25, using the algorithm in FIG. 17 to aggregate the ACKs generated. In this example, ENCODE_LIMIT and PACKET_PERIOD have been set to one, having every TCP ACK split into multiple ACKs. MINIMUM_SEGMENT_SIZE is set to 500 bytes, and the MAXIMUM_ACKs is set to 4. The LAST_VALID_ACK is one. As a result, when the message 220 reaches IND 70, this message is split into 3 ACKs with sequence numbers 501, 1001 and 1501. These messages are aggregated into a single encoded message 2601 and sent to the wireless client 50. The wireless client 50 and the wireless AP 40 negotiate the transmission of the encoded packet at step 2602 using the collision avoidance algorithm, and this is passed onto IND 60 in a message 2603. The IND 60 unpacks the encoded TCP ACKs in the packet following the algorithm illustrated in FIG. 16. When the decoder completes its work, the 3 ACKs are forwarded to the server 30 which then increases its congestion window.

[0072] Depending on the Wireless LAN technology, transmission rates over the LAN may be limited. For example, an 802.1b network has only an 11-Mbps transmission rate versus an 802.1g network which has a 54-Mbps rate. Detection of the LAN type can be estimated and shared among the associated INDs, which can influence the level of encoding selected by the INDs for a particular association. If the link is discerned to be limited in either transmission rate or in packets per second, the encoding rate may be increased. If the INDs are summarizing or aggregating TCP ACKs, the encoding may be increased by increasing MAX_ACKs or increasing the timeout period of the ACK Pending Timer.

[0073] To ensure that the TCP ACKs are sent in a timely fashion, all of the encoders, irrespective of the algorithm being used, can set the TOS/DSCP bits of the outgoing packets to ensure that are delivered in a timely fashion. When the encoded packet arrives at the wireless LAN, 802.11e can be employed to ensure that the TCP ACKs are delivered in a timely fashion. It can also be leveraged to reduce the overhead of the TCP ACK over the wireless LAN.

[0074] Embodiments of the invention may be implemented in a network having endpoints, such as servers or other computing devices, and associated encoding components. The encoding components, and the described methods, can be implemented in hardware, software or a combination thereof. Those portions that are implemented in software can be represented as a software product stored in a machine-readable medium (also referred to as a computer-readable medium, a processor-readable medium, or a computer usable medium having a computer readable program code embodied therein). The machine-readable medium may be any suitable tangible medium, including magnetic, optical, or electrical storage medium including a diskette, compact disk read only memory (CD-ROM), memory device (volatile or non-volatile), or similar storage mechanism. The machine-readable medium may contain various sets of instructions, code sequences, configuration information, or other data, which, when executed, cause a processor to perform steps in a method according to an embodiment of the invention. Those of ordinary skill in the art will appreciate that other instructions and operations necessary to implement the described invention may also be stored on the machine-readable medium. Software running from the machine readable medium may interface with circuitry to perform the described tasks.

[0075] While particular embodiments and applications of the present invention have been illustrated and described, it is to be understood that the invention is not limited to the precise construction and compositions disclosed herein and that various modifications, changes, and variations may be apparent from the foregoing descriptions without departing from the spirit and scope of the invention as defined in the appended claims.

Claims

1.-6. (canceled)

7. A computer implemented method comprising: identifying one or more TCP acknowledgements among data packets transmitted between a source Transmission Control Protocol (TCP) device and

a destination TCP device; preventing transmission of the one or more TCP acknowledgements; and transmitting a summary TCP acknowledgement based at least in part on a determination that an acknowledgement time indicating time elapsed since identifying of a first TCP acknowledgement of the one or more TCP acknowledgements is expired or that a threshold number of the one or more TCP acknowledgements has been prevented from being transmitted.

8. The method of claim 7, wherein the preventing transmission of the one or more TCP acknowledgements comprises buffering the one or more TCP acknowledgements.

9. The method of claim 7, wherein the one or more TCP acknowledgements are transmitted from a client TCP device to a wireless client or a wireless access point.

10. The method of claim 7, wherein the one or more TCP acknowledgements are transmitted as part of a single TCP session.

11. The method of claim 7, wherein the acknowledgment time is based at least in part on a measured round trip time between the source TCP device and the destination TCP device.

12. The method of claim 7, wherein the acknowledgment time is based at least in part on a measured network delay.

13. The method of claim 7, wherein the acknowledgment time is based at least in part on a type of wireless network segment between the source TCP device and the destination TCP device.

14. The method of claim 7, wherein the preventing transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between the client TCP device and a wireless client.

15. The method of claim 7, wherein the preventing transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between the client TCP device and a wireless access point.

16. The method of claim 7, wherein the preventing the transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between a wired network and a wireless access point.

17. A system comprising: a memory configured to store one or more TCP acknowledgements; and a controller configured to: identify the one or more TCP acknowledgements among data packets transmitted between a source Transmission Control Protocol (TCP) device and a destination TCP device; store the one or more TCP acknowledgements in the memory and prevent transmission of the one or more TCP acknowledgements; and transmit a summary TCP acknowledgement based at least in part on a determination that an acknowledgement time indicating time elapsed since identifying of a first TCP acknowledgement of the one or more TCP acknowledgements is expired or that a threshold number of the one or more TCP acknowledgements has been prevented from being transmitted.

18. The system of claim 17, wherein the preventing transmission of the one or more TCP acknowledgements comprises buffering the one or more TCP acknowledgements.

19. The system of claim 17, wherein the one or more TCP acknowledgements are transmitted from a client TCP device to a wireless client or a wireless access point.

20. The system of claim 17, wherein the one or more TCP acknowledgements are transmitted as part of a single TCP session.

21. The system of claim 17, wherein the acknowledgment time is based at least in part on a measured round trip time between the source TCP device and the destination TCP device.

22. The system of claim 17, wherein the acknowledgment time is based at least in part on a measured network delay.

23. The system of claim 17, wherein the acknowledgment time is based at least in part on a type of wireless network segment between the source TCP device and the destination TCP device.

24. The system of claim 17, wherein the preventing transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between the client TCP device and a wireless client.

- 25.** The system of claim 17, wherein the preventing transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between the client TCP device and a wireless access point.
- 26.** The system of claim 17, wherein the preventing the transmission of the one or more TCP acknowledgements and the transmitting of the summary TCP acknowledgement are performed between a wired network and a wireless access point.
-