



US 20250279988A1

(19) **United States**

(12) **Patent Application Publication**
SCHMITZ et al.

(10) **Pub. No.: US 2025/0279988 A1**

(43) **Pub. Date: Sep. 4, 2025**

(54) **CONNECTOR DEPLOYMENT WITHIN A
CONNECTIVITY FRAMEWORK**

(52) **U.S. Cl.**

CPC **H04L 63/08** (2013.01); **H04L 63/029**
(2013.01)

(71) Applicant: **SAP SE**, Walldorf (DE)

(72) Inventors: **Christian SCHMITZ**, Porto Alegre
(BR); **Peter SCHOENAU**, Stuttgart
(DE); **Silvio NORMEY GOMEZ**,
Santana do Livramento (BR); **Mitchell
GUDMUNDSON**, Colorado Springs,
CO (US)

(57)

ABSTRACT

Systems and methods include execution of a first instance of a connector registry service to receive a first registration request from a first instance of a first connector to a first data system, the first registration request comprising a first connector identifier, a first connection type and a first endpoint type, store the first connector identifier, the first connection type and the first endpoint type in a data store, read a plurality of associated connector identifiers, connection types and endpoint types from the data store, the plurality of associated connector identifiers, connection types and endpoint types including the first connector identifier, the first connection type and the first endpoint type, and transmit the plurality of associated connector identifiers, connection types and endpoint types to a first instance of a gateway service.

(21) Appl. No.: **18/772,498**

(22) Filed: **Jul. 15, 2024**


Related U.S. Application Data

(60) Provisional application No. 63/560,189, filed on Mar. 1, 2024.

Publication Classification

(51) **Int. Cl.**
H04L 9/40

(2022.01)






FIG. 1




FIG. 2




FIG. 3




FIG. 4




FIG. 5

CONNECTOR DEPLOYMENT WITHIN A CONNECTIVITY FRAMEWORK

Attorney Docket No.: 240091US02

Sheet 6/8




FIG. 6




FIG. 7




FIG. 8

CONNECTOR DEPLOYMENT WITHIN A CONNECTIVITY FRAMEWORK

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] The present application claims priority to and the benefit of U.S. Provisional Patent Application No. 63/560,189, filed Mar. 1, 2024, under 35 U.S.C. § 119(e). The contents of U.S. Provisional Patent Application No. 63/560,189 are incorporated herein for all purposes.

BACKGROUND

[0002] Modern organizations generate, process and store vast amounts of data. This data may be stored in various data stores and accessed by various applications. Different data stores are differently suited for storing different types of data. Moreover, a given application may use different types of data associated with different requirements.

[0003] The particular data stores in which data of an application are stored may be selected based on many factors, including but not limited to a desired redundancy, a desired accessibility, a desired processing speed, and a desired security protocol. Due to these considerations, application data may be stored across multiple databases, file storages, file systems, and external data sources such as third-party vendors, cloud services, and cloud databases.

[0004] The diversity of data stores increases the complexity of data integration. For example, the more types of data stores which are used within an application landscape, the more difficult it is to seamlessly pull and push data and metadata from the data stores. Problems arise at least due to the difficulty in properly configuring the application to connect to and interact with the data stores. In addition, existing middleware for routing communication between an application and its data stores suffers from poor scalability and extensibility.

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] FIG. 1 illustrates a framework for connecting to data systems according to some embodiments.

[0006] FIG. 2 is a sequence diagram of registering a connector within a framework for connecting to data systems according to some embodiments.

[0007] FIG. 3 is an entity relationship diagram for a connector registry persistency model according to some embodiments.

[0008] FIG. 4 illustrates entities of a gateway control plane according to some embodiments.

[0009] FIG. 5 illustrates clusters providing implementing components of a framework for connecting to data systems according to some embodiments.

[0010] FIG. 6 illustrates distribution of connector information among instances of components of a framework according to some embodiments.

[0011] FIG. 7 is a sequence diagram of using a registered connector to connect to a data system according to some embodiments.

[0012] FIG. 8 illustrates a hardware implementation of a cluster-implemented framework according to some embodiments.

DETAILED DESCRIPTION

[0013] Embodiments may provide scalable, resilient, observable, extensible, and high-performance connectivity with disparate data systems.

[0014] FIG. 1 illustrates landscape **100** according to some embodiments. Landscape **100** may comprise any number of hardware and software components which provide functionality to one or more users (not shown). The components may be implemented using any suitable combinations of computing hardware and/or software that are or become known. Such combinations may include on-premise servers, cloud-based servers, and/or elastically-allocated virtual machines. In some embodiments, two or more components are implemented by a single computing device.

[0015] Landscape **100** comprises connectors **110** and **115**, which provide integration with data and metadata stored by data systems of different connection and endpoint types. Examples of connection types include but are not limited to ABAP, Google BigQuery, and S4/HANA Cloud, and examples of endpoint types include but are not limited to Web Socket RFC (WSRFC), REST, and Content Distribution Internetworking (CDI). Gateway **120** receives an incoming request for access to a data system of a particular connection type and endpoint type and routes the incoming request to a connector corresponding to the connection type, endpoint type and tenant. Connector registry **130** provides routing information to gateway **120** and connection details to deployed connectors **110** and **115**. Connection details may include information needed to connect to a data system, such as hostname, port, username, password, certificates, etc. Connector registry **130** may also manage connector registration, cache connection details and notify connectors **110**, **115** of connection detail updates and connection invalidations.

[0016] Embodiments may implement any number of connectors, each of which is associated with one or more respective connection types, endpoint types and tenants. A connector implements logic for reading and writing data and metadata from and to a data system instance of the connection type and endpoint type with which the connector is associated. Each connector **110**, **115** implements and exposes the same well-defined and parametrizable connector API which may be called by a consumer to interact with a data system of the connection type and endpoint type associated with the connector. For example, via the API, connector **110** may receive requests from gateway **120** for data read and write, metadata read and write, and data definition. Data system client **114** may convert the requests from a data model of connectors **110**, **115** to a data model of a connection type and endpoint type with which connector **110** is associated and transmit the request to a data system using drivers, service invocations, or other methods of the data system. Similarly, data system client **114** may receive responses from a data system and convert the responses from the data model of the data system to the data model of connectors **110**, **115**.

[0017] A connector may provide a consumer with information specifying its characteristics and capabilities. Connectors **110**, **115** may conform to a capability model which obliges a consumer to send commands to read or write data within the set of operations delimited by the connector's capabilities.

[0018] According to some embodiments, each connector **110**, **115** is a separate microservice that exposes and imple-

ments the connector API. Each microservice (i.e., connector) may be provided by a plurality of pods executing connector instances in one or more nodes of a container orchestration system as is known in the art. Landscape **100** may be a cloud-native and decoupled architecture in which gateway **120** and connector registry **130** each also comprise separate microservices implemented by a respective one or more nodes. Consequently and advantageously, each of connector **110**, **115**, gateway **120** and connector registry **130** may scale elastically (e.g., by adding/deleting pods and/or nodes) according to their respective workloads. All microservices described herein may communicate with one another and with other unshown microservices using lightweight network communication mechanisms such as a resource API via Hyper Text Transfer Protocol (HTTP) request-response messages, but embodiments are not limited thereto.

[0019] According to some embodiments, landscape **100** uses gRPC for communication among the components, Protocol Buffers as an interface description language for message, service, and API definitions, and a data wire format such as Apache Arrow. Connectors **110**, **115**, gateway **120** and registry **130** may execute on an application server that runs and operates applications on Kubernetes.

[0020] Connector registry **130** comprises connector registry API **132**, gateway control plane API **134**, and credentials client **136**. Connector registry **130** maintains, monitors, and exposes connectivity to the connection type and endpoint type registered by each individual connector. Connector registry **130** keeps track of the available connectors serving connectivity and the runtime instances of the connectors.

[0021] Connector registry API **132** may comprise a bi-directional gRPC API. During bootstrapping, connector registry client **112** of connector **110** invokes connector registry API **132** to register its information with connector registry **130**. The information may include, for example, connector identifier, connector version, configuration, connection type, endpoint type, and supported tenants. Connector registry API **132** calls data access object **138** to persist the connector-specific information in registry persistency **145** of persistent storage **140**, which may in some embodiments consist of cloud-based persistent storage. Registry persistency **145** may synchronize with the infrastructure layer of landscape **100** to also maintain a list of the available instances (e.g., running pods) for each registered connector.

[0022] Gateway control plane API **134** provides routing information to gateway **120**. Gateway **120** caches the routing information and uses it to route all incoming traffic to the proper connector. Gateway **120** routes incoming requests to connectors based on connection information (e.g., connection identifier, connection type, endpoint type, tenant) specified by the requests. Routing may be based on other types of information specified in a request. Gateway **120** may comprise an L7 load balancer that supports HTTP/2 and connections multiplexing, which allows the reuse of the same TCP socket for multiple HTTP connections.

[0023] Gateway **120** can survive outages of gateway control plane API **134** using cached routing information, at least until the routing information stored in registry persistency **145** changes. Routing information stored in registry persistency **145** may change due to asynchronous registration or de-registration of connectors. To keep the routing information of gateway **120** up-to-date, gateway **120** may maintain a live gRPC stream to gateway control plane API **134**. Accordingly, gateway control plane API **134** may regularly

pull routing information from registry persistency **145** of persistent storage **140** via data access object **138** and determine whether any of the routing information has changed. If so, all current routing information (or just the changed routing information) stored in registry persistency **145** is transmitted to gateway **120**. According to some embodiments, prior to transmission of the routing information stored gateway control plane API **134** translates routing information entities corresponding to the connector data model into entities (e.g., clusters, endpoints, listeners, and routes) of the data model of gateway **120**.

[0024] As mentioned above, gateway **120** may be implemented using multiple instances. When a new instance is deployed, the instance calls gateway control plane API **134** to create a gRPC bi-directional stream and thereafter calls the gateway control plane API **134** to retrieve a routing configuration including all current routing information from registry persistency **145**.

[0025] Connector registry API **132** provides RPCs for connectors to retrieve connection details such as hostname, port, username, and password associated with a connection ID. Connection details may also include certificates and connection-specific drivers. Connector registry API **132** retrieves the connection details from credentials client **136**, which connects to credentials component **150** to obtain the connection details. Credentials component **150** may comprise a secure storage system in which a user (e.g., a tenant administrator) stores connection details for data systems to which the user has access, and where each data system (and the connection details thereof) is associated with a connection type, endpoint type and tenant. In some embodiments, credentials client **136** includes an in-memory cache mechanism to buffer connection details to avoid unnecessary calls to credentials component **150** and provide resilience in case of communication issues with credentials component **150**.

[0026] For example, in response to receipt of a request for connection details from a connector, credentials client **136** may determine whether the requested details are in its cache. If not, the details are requested from credentials component **150**, cached by credentials client **136**, and returned to the requesting connector via API **132**. If the requested details are found in the cache of credentials client **136**, credentials client **136** forwards an audit record to audit log service **160** to fulfill security and auditing requirements and returns the details to the requesting connector.

[0027] FIG. 2 is a sequence diagram illustrating registration of a connector within a framework according to some embodiments. Process **200** and the other processes described herein may be performed using any suitable combination of hardware and software. Program code embodying these processes may be stored by any non-transitory tangible medium, including a fixed disk, a volatile or non-volatile random-access memory, a DVD, a Flash drive, or a magnetic tape, and executed by any number of processing units, including but not limited to processors, processor cores, and processor threads. Such processors, processor cores, and processor threads may be implemented by a virtual machine provisioned in a cloud-based architecture.

[0028] FIG. 2 will be described with respect to the components of landscape **100**, but embodiments are not limited thereto. During bootstrap of connector **110**, connector registry client **112** requests a connection with connector registry API **132** to open a long-lived gRPC channel therewith. This

channel acts as a corridor for actions to be executed during bootstrap and across the life of connector **110**.

[0029] Using the corridor, connector registry client **112** calls connector registry API **132** and provides connector information. The connector information may include a connector identifier the connection type and endpoint type to which connector **110** provides connectivity, the associated tenant (if not multi-tenant), connector version, the endpoint of the registering connector instance, and any remaining configuration information needed for operation of connector **110** and routing of incoming requests. API **132** then translates the connector information into native entities of a connector registry persistency data model.

[0030] FIG. 3 shows an entity relationship diagram for a connector registry persistency data model according to some embodiments. Each of entities Connector **310**, Connector-Version **320** and ConnectionAndEndpoint **330** includes a namespace field to express ownership. A connector may be owned by a provider of landscape **100**, a partner thereof, or a tenant who developed connector **110**, for example.

[0031] The translated connector information is persisted in persistent storage **140**. Next, connector registry API **132** reads all the persisted connector information from persistent storage **140** and converts the entities of the read connector information to a data model of gateway **120**. FIG. 4 illustrates entities **410** of a connector registry persistency data model which may be persisted in cloud storage, and corresponding entities **420** of a gateway data model converted from entities **410** according to some embodiments. FIG. 4 also illustrates the population of various fields of entities **420** with values **411-417** of corresponding fields of entities **410**.

[0032] Connector registry API **132** sends the converted connector information to gateway **120** to update its internal routing configuration. Accordingly, gateway **120** is provided with a routing configuration based on the connector information currently stored in persistency **145** of persistent storage **140**.

[0033] FIG. 5 illustrates microservice implementations of connectors **510**, **514** and **516** gateway **520** and connector registry **530** according to some embodiments. Each of services **510**, **514**, **516**, **520** and **530** may be implemented by a plurality of pods of a container orchestration platform such as but not limited to Kubernetes. As such, each service may include a node executing several pods, in which each pod executes a separate instance of the service. Each node **512**, **522** and **532** is a virtual or physical machine including kubelet for node and container management, kube-proxy for a network proxy, and a container runtime (e.g., Docker) to run container.

[0034] Although FIG. 5 illustrates one node **512**, **522** and **532** per service, any service may include any number of nodes, each executing its own set of one or more pods. A master node (not shown) of each service may adjust the number of pods, the number of nodes and/or the computing resources of each node as desired. The adjustment may be based on expected workload or any other factors. For example, if an expected workload is greater than a first threshold, one or more additional pods are created. If the expected workload is less than a second threshold, one or more of pods are terminated.

[0035] Accordingly, each of connector services **510**, **514**, **516** may be implemented by several connector instances and connector registry service **530** may be implemented by several connector registry instances, in which all instances

execute independently and in parallel. If a connector registry client of a connector instance registers its connector information with a connector registry API of a connector registry instance as described with respect to FIG. 2, all other connector registry instances will initially be unaware of the registration. Accordingly, to establish a consistent state of connection information throughout the landscape, each running connector registry instance may periodically read the connector information from the shared cloud storage and push the connector information to the gateway instance(s) to which it is connected.

[0036] FIG. 6 illustrates connections between a connector instance, various instances of a connector registry and various instances of a gateway according to some embodiments. Specifically, FIG. 6 shows one running connector instance **610**, three running connector registry instances **620**, **622** and **624**, and five running gateway instances **630-634**. Each of gateway instances **630-634** is connected to one of connector registry instances **620**, **622** and **624** and maintains a gRPC long-lived channel therewith. Connector instance **620** sends its connection information to connector registry instance **620** for storage in cloud storage **640** and connector registry instance **620** updates the connector information of gateway instance **630**. To avoid a globally inconsistent state, remaining connector registry instances **622** and **624** periodically pull connector information from cloud storage **640** and push the connector information to gateway instances **630-634** through the gRPC long-lived channels. Specifically, connector registry instance **622** pushes the connector information to gateway instances **631** and **632** using its gateway control plane API and connector registry instance **623** pushes the connector information to gateway instances **633** and **634** using its gateway control plane API.

[0037] FIG. 7 is a sequence diagram illustrating a request from a consumer to a connector. The request may comprise a request to read data/metadata, a request to write data, or a Data Description Language request.

[0038] Consumer **710** may comprise a computing device which uses a client component library to communicate with gateway **720**. The client component library may encapsulate the communication and actions between consumer **710** and gateway **720**. The client component library may be tasked with managing gRPC connections, checking the backend and connectors' health status, security, resilience, and other aspects.

[0039] Initially, consumer **710** sends a request to create a channel to gateway **720** via the client component library. A gRPC channel is established and consumer **710** consumer uses the client component library to send a request that will read or modify the state of the database in target data system **770**. The request includes connection information (e.g., connection type, endpoint type, connection identifier and tenant). Advantageously, the request need not specify any particular connector or any connection details for connecting to target data system **770**.

[0040] Based on the received connection information and its stored routing information, gateway **720** routes the request to corresponding connector **730**. After receiving the request, connector **730** determines whether connection details corresponding to the connection ID of the request are stored in its in-memory buffer. If not, connector **730** calls connector registry **740** (e.g., using a connector registry client as described above) to ask for the connection details based on the connection ID and the requester tenant. Connector

registry **740** then determines whether connection details corresponding to the connection ID and tenant is in its in-memory buffer. If so, connector registry **740** sends an audit log record to audit log service **760** and returns the connection details to connector **730**. If the connection details are not in its buffer, connector registry **740** calls credentials component **750** to request the connection details, receives the connection details therefrom, and sends the connection details back to connector **730**.

[0041] Connector **730** uses the connection details (and, for example, a data system client component) to establish a connection with target data system **770**. Next, connector **730** sends the request to data system **770**, receives a result, and returns the result to consumer **710** via gateway **720**.

[0042] FIG. **8** illustrates a cloud-based deployment according to some embodiments. The illustrated components may comprise cloud-based resources residing in one or more public or private clouds providing self-service and immediate provisioning, autoscaling, security, compliance and identity management features.

[0043] Each illustrated node may comprise a physical or virtual machine of a Kubernetes cluster. Nodes **810** may execute instances of a gateway while nodes **820** may execute instances of a connector registry. Nodes **830** may comprise a database service for storing connection information. Nodes **840** may execute instances of a first connector and nodes **850** may execute instances of a second connector. The instances of the first connector may communicate with data system **845** and the instances of the second connector may communicate with data system **855**.

[0044] The foregoing diagrams represent logical architectures for describing processes according to some embodiments, and actual implementations may include more, or different components arranged in other manners. Other topologies may be used in conjunction with other embodiments. Moreover, each component or device described herein may be implemented by any number of devices in communication via any number of other public and/or private networks. Two or more of such computing devices may be located remote from one another and may communicate with one another via any known manner of networks and/or a dedicated connection. Each component or device may comprise any number of hardware and/or software elements suitable to provide the functions described herein as well as any other functions. For example, any computing device used in an implementation of a system according to some embodiments may include a processor to execute program code such that the computing device operates as described herein.

[0045] All systems and processes discussed herein may be embodied in program code stored on one or more non-transitory computer-readable media. Such media may include, for example, a hard disk, a DVD-ROM, a Flash drive, magnetic tape, and solid-state Random Access Memory (RAM) or Read Only Memory (ROM) storage units. Embodiments are therefore not limited to any specific combination of hardware and software.

[0046] Embodiments described herein are solely for the purpose of illustration. Those in the art will recognize other embodiments may be practiced with modifications and alterations to that described above.

What is claimed is:

1. A system comprising:
 - a memory storing program code;
 - one or more processing units to execute the program code to:
 - execute a first instance of a connector registry service to:
 - receive a first registration request from a first instance of a first connector to a first data system, the first registration request comprising a first connector identifier, a first connection type and a first endpoint type;
 - store the first connector identifier, the first connection type and the first endpoint type in a data store;
 - read a plurality of associated connector identifiers, connection types and endpoint types from the data store, the plurality of associated connector identifiers, connection types and endpoint types including the first connector identifier, the first connection type and the first endpoint type; and
 - transmit the plurality of associated connector identifiers, connection types and endpoint types to a first instance of a gateway service.
2. A system according to claim 1, the one or more processing units to execute the program code to cause the system to:
 - convert the plurality of associated connector identifiers, connection types and endpoint types from a data model of the first connector to a data model of the gateway service,
 - wherein transmission of the plurality of associated connector identifiers and connection types comprises transmission of the converted plurality of associated connector identifiers and connection types.
3. A system according to claim 1, the one or more processing units to execute the program code to cause the system to:
 - execute a second instance of the connector registry service to:
 - read the plurality of associated connector identifiers, connection types and endpoint types from the data store; and
 - transmit the plurality of associated connector identifiers, connection types and endpoint types to a second instance of the gateway service and to a third instance of the gateway service.
4. A system according to claim 3, the one or more processing units to execute the program code to cause the system to:
 - receive a second request from a second instance of the first connector for connection details associated with the first connector identifier;
 - in response to the second request, determine a hostname, port and user credentials associated with the first connector identifier; and
 - return the hostname, port and user credentials to the second instance of the first connector.
5. A system according to claim 4, wherein determination of the hostname, port and user credentials associated with the first connector identifier comprises:
 - determination of whether the hostname, port and user credentials associated with the first connector identifier are not stored in a cache of the system; and
 - in response to the determination, requesting of the hostname, port and user credentials associated with the first connector identifier from a second system.

6. A system according to claim 1, the one or more processing units to execute the program code to cause the system to:

receive a second request from a second instance of the first connector for connection details associated with the first connector identifier;

in response to the second request, determine a hostname, port and user credentials associated with the first connector identifier; and

return the hostname, port and user credentials to the second instance of the first connector.

7. A system according to claim 6, wherein determination of the hostname, port and user credentials associated with the first connector identifier comprises:

determination that the hostname, port and user credentials associated with the first connector identifier are not stored in a cache of the system; and

in response to the determination that the hostname, port and user credentials associated with the first connector identifier are not stored in the cache of the system, requesting of the hostname, port and user credentials associated with the first connector identifier from a second system.

8. A system according to claim 1, the one or more processing units to execute the program code to:

execute a second instance of the connector registry service to:

read the plurality of associated connector identifiers, connection types and endpoint types from the data store; and

transmit the plurality of associated connector identifiers, connection types and endpoint types to a second instance of the gateway service; and

execute a third instance of the connector registry service to:

receive a second registration request from a first instance of a second connector to a second data system, the second registration request comprising a second connector identifier, a second connection type and a second endpoint type;

store the second connector identifier, the second connection type and the second endpoint type in the data store;

read a second plurality of associated connector identifiers, connection types and endpoint types from the data store, the second plurality of associated connector identifiers, connection types and endpoint types including the second connector identifier, the second connection type and the second endpoint type; and

transmit the second plurality of associated connector identifiers, connection types and endpoint types to a third instance of the gateway service.

9. A system according to claim 8, the one or more processing units to execute the program code to:

execute the second instance of the connector registry service to:

read the second plurality of associated connector identifiers, connection types and endpoint types from the data store; and

transmit the second plurality of associated connector identifiers, connection types and endpoint types to the second instance of the gateway service; and

execute the first instance of the connector registry service to:

read the second plurality of associated connector identifiers, connection types and endpoint types from the data store; and

transmit the second plurality of associated connector identifiers, connection types and endpoint types to the first instance of the gateway service.

10. A method comprising:

receiving, at a first instance of a connector registry service, a first registration request from a first instance of a first connector to a first data system, the first registration request comprising a first connector identifier and a first connection type;

storing, by the first instance of the connector registry service, the first connector identifier and the first connection type in a data store;

reading, by the first instance of the connector registry service, a plurality of associated connector identifier and connection types from the data store, the plurality of associated connector identifiers and connection types including the first connector identifier and the first connection type; and

transmitting, by the first instance of the connector registry service, the plurality of associated connector identifiers and connection types to a first instance of a gateway service and to a second instance of the gateway service.

11. A method according to claim 10, further comprising:

converting, by the first instance of the connector registry service, the plurality of associated connector identifiers and connection types from a data model of the first connector to a data model of the gateway service,

wherein transmitting the plurality of associated connector identifiers and connection types comprises transmitting, by the first instance of the connector registry service, the converted plurality of associated connector identifiers and connection types.

12. A method according to claim 10, further comprising: executing a second instance of the connector registry service to:

read the plurality of associated connector identifiers and connection types from the data store; and

transmit the plurality of associated connector identifiers and connection types to a third instance of the gateway service.

13. A method according to claim 12, further comprising:

receiving, at the second instance of the connector registry service, a second request from a second instance of the first connector for connection details associated with the first connector identifier;

in response to the second request, determining, by the second instance of the connector registry service, a hostname, port and user credentials associated with the first connector identifier; and

returning, by the second instance of the connector registry service, the hostname, port and user credentials to the second instance of the first connector.

14. A method according to claim 13, wherein determining the hostname, port and user credentials associated with the first connector identifier comprises:

determining, by the second instance of the connector registry service, whether the hostname, port and user credentials associated with the first connector identifier are not stored in a cache; and

in response to the determination that the hostname, port and user credentials associated with the first connector identifier are not stored in the cache, requesting, by the second instance of the connector registry service, the hostname, port and user credentials associated with the first connector identifier from a system.

15. A method according to claim **10**, further comprising: receiving, at the first instance of the connector registry service, a second request from a second instance of the first connector for connection details associated with the first connector identifier;

in response to the second request, determining, by the first instance of the connector registry service, a hostname, port and user credentials associated with the first connector identifier; and

returning, by the first instance of the connector registry service, the hostname, port and user credentials to the second instance of the first connector.

16. A method according to claim **15**, wherein determining the hostname, port and user credentials associated with the first connector identifier comprises:

determining that the hostname, port and user credentials associated with the first connector identifier are not stored in a cache; and

in response to the determination that the hostname, port and user credentials associated with the first connector identifier are not stored in the cache, requesting the hostname, port and user credentials associated with the first connector identifier from a system.

17. A method according to claim **10**, further comprising: executing a second instance of the connector registry service to:

read the plurality of associated connector identifiers, connection types and endpoint types from the data store; and

transmit the plurality of associated connector identifiers and connection types to a third instance of the gateway service; and

execute a third instance of the connector registry service to:

receive a second registration request from a first instance of a second connector to a second data system, the second registration request comprising a second connector identifier and a second connection type;

store the second connector identifier and the second connection type in the data store;

read a second plurality of associated connector identifiers and connection types from the data store, the second plurality of associated connector identifiers and connection types including the second connector identifier and the second connection type; and

transmit the second plurality of associated connector identifiers and connection types to a fourth instance of the gateway service.

18. A method according to claim **17**, further comprising: executing the second instance of the connector registry service to:

read the second plurality of associated connector identifiers and connection types from the data store; and

transmit the second plurality of associated connector identifiers and connection types to the third instance of the gateway service; and

execute the first instance of the connector registry service to:

read the second plurality of associated connector identifiers and connection types from the data store; and

transmit the second plurality of associated connector identifiers and connection types to the first instance of the gateway service and to the second instance of the gateway service.

19. One or more non-transitory computer-readable media storing program code that, when executed by a computing system, cause the computing system to perform operations comprising:

receiving, at a first instance of a connector registry service, a first registration request from a first instance of a first connector to a first data system, the first registration request comprising a first connector identifier and a first connection type;

storing, by the first instance of the connector registry service, the first connector identifier and the first connection type in a data store;

reading, by the first instance of the connector registry service, a plurality of associated connector identifier and connection types from the data store, the plurality of associated connector identifiers and connection types including the first connector identifier and the first connection type;

transmitting, by the first instance of the connector registry service, the plurality of associated connector identifiers and connection types to a first instance of a gateway service and to a second instance of the gateway service;

receiving, at a second instance of the connector registry service, a second registration request from a first instance of a second connector to a second data system, the second registration request comprising a second connector identifier and a second connection type;

storing, by the second instance of the connector registry service, the second connector identifier and the second connection type in the data store;

reading, by the second instance of the connector registry service, a second plurality of associated connector identifiers and connection types from the data store, the second plurality of associated connector identifiers and connection types including the second connector identifier and the second connection type; and

transmitting, by the second instance of the connector registry service, the second plurality of associated connector identifiers and connection types to a third instance of the gateway service,

wherein the first instance of the connector registry service reads the second plurality of associated connector identifiers and connection types from the data store and transmits the second plurality of associated connector identifiers and connection types to the first instance and the second instance of the gateway service.

20. One or more non-transitory computer-readable media according to claim **19**, where the program code, when executed by a computing system, causes the computing system to perform operations comprising:

receiving, at the first instance of the connector registry service, a second request from a second instance of the first connector for connection details associated with the first connector identifier;

in response to the second request, determining, by the first instance of the connector registry service, a hostname, port and user credentials associated with the first connector identifier; and
returning, by the first instance of the connector registry service, the hostname, port and user credentials to the second instance of the first connector.

* * * * *