

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250280039

Kind Code

A1

Publication Date

September 04, 2025

Inventor(s)

GANGWANI; Tvisha Rajesh et al.

AUTOMATED IDENTIFICATION OF BENCHMARK SECURITY CONTROLS

Abstract

System, methods, apparatuses, and computer program products are disclosed for automatically identifying benchmark security controls. Service information related to a service of a customer is provided to a policy predictor model that is trained to predict a control mapping based on the service information. A predicted control mapping associated with a predicted security policy and a predicted security control are received from the policy predictor model. Implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy.

Inventors: GANGWANI; Tvisha Rajesh (Philadelphia, PA), KARABEY; Bugra (Redmond, WA), CHENG; Cheng (Kirkland, WA), NATARAJAN; Sasikumar (Redmond, WA), ALFRED; Andre Oneal (Redmond, WA), MANSUETO; Jenna Sara (Chicago, IL)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 94768033

Appl. No.: 18/591572

Filed: February 29, 2024

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/20 (20130101);

Background/Summary

BACKGROUND

[0001] Security control benchmarks are standards or sets of guidelines that organizations may use to assess the effectiveness of their security controls. Many security control benchmarks are based on widely accepted industry standards and frameworks. These benchmarks provide a framework for evaluating and improving an organization's security by establishing a baseline of security best practices. Security control benchmarks may cover various aspects of information security, including, but not limited to, policies, processes, technologies, and/or personnel. Identifying and implementing the proper security control benchmarks can help organizations enhance their overall security, demonstrate compliance with regulations, and/or mitigate the risks associated with cyber threats.

SUMMARY

[0002] This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

[0003] System, methods, apparatuses, and computer program products are disclosed for automatically identifying benchmark security controls. Service information related to a service of a customer is provided to a policy predictor model that is trained to predict a control mapping based on the service information. A predicted control mapping associated with a predicted security policy and a predicted security control are received from the policy predictor model. Implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy.

[0004] Further features and advantages of the embodiments, as well as the structure and operation of various embodiments, are described in detail below with reference to the accompanying drawings. It is noted that the claimed subject matter is not limited to the specific embodiments described herein. Such embodiments are presented herein for illustrative purposes only. Additional embodiments will be apparent to persons skilled in the relevant art(s) based on the teachings contained herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS/FIGURES

[0005] The accompanying drawings, which are incorporated herein and form a part of the specification, illustrate embodiments of the present application and, together with the description, further serve to explain the principles of the embodiments and to enable a person skilled in the pertinent art to make and use the embodiments.

[0006] FIG. 1 shows a block diagram of an example system for automatically identifying a control mapping using a policy predictor model, in accordance with an embodiment.

[0007] FIG. 2 depicts a block diagram of an example system for training a policy predictor model for identifying a control mapping and using the policy predictor model to identify a control mapping, in accordance with an embodiment.

[0008] FIG. 3 depicts a flowchart of a process for automatically identifying a control mapping, in accordance with an embodiment.

[0009] FIG. 4 depicts a flowchart of a process for generating a textual description of a service based on an architectural diagram of the service, in accordance with an embodiment.

[0010] FIG. 5 depicts a flowchart of a process for training a policy predictor model for predicting a

control mapping for a service based on a textual description of the service, in accordance with an embodiment.

[0011] FIG. 6 depicts a block diagram of an example system for automatically identifying a control mapping for a service deployed in a multi-cloud environment and implanting the identified control mapping in the multi-cloud environment, in accordance with an embodiment.

[0012] FIG. 7 depicts a flowchart of a process for automatically determining service information associated with a multi-cloud service operating in a multi-cloud environment, in accordance with an embodiment.

[0013] FIG. 8 shows a block diagram of an example computer system in which embodiments may be implemented.

[0014] The subject matter of the present application will now be described with reference to the accompanying drawings. In the drawings, like reference numbers indicate identical or functionally similar elements. Additionally, the left-most digit(s) of a reference number identifies the drawing in which the reference number first appears.

DETAILED DESCRIPTION

I. Introduction

[0015] The following detailed description discloses numerous example embodiments. The scope of the present patent application is not limited to the disclosed embodiments, but also encompasses combinations of the disclosed embodiments, as well as modifications to the disclosed embodiments. It is noted that any section/subsection headings provided herein are not intended to be limiting. Embodiments are described throughout this document, and any type of embodiment may be included under any section/subsection. Furthermore, embodiments disclosed in any section/subsection may be combined with any other embodiments described in the same section/subsection and/or a different section/subsection in any manner.

II. Example Embodiments

[0016] Cloud service platforms release new cloud capabilities and features on a regular basis. As these new cloud capabilities and features are released, cloud customers develop and publish new customer services that incorporate the new cloud capabilities and/or features. In order to secure the customer services, associated resources need to be properly configured. Cloud security control benchmarks have been developed to provide guidelines and/or recommendations to help cloud customers evaluate the effectiveness of their security measures and/or ensure compliance with industry and government standards.

[0017] Various cloud security control benchmarks are available to cloud customers to help cloud customers enhance the overall security of their customer services, demonstrate compliance with regulations, and/or mitigate the risks associated with cyber threats. Such cloud security control benchmarks have been developed by various entities, such as, but not limited to, technology companies, cloud providers, standards setting organizations, government regulators, industry groups, industry regulators, and/or the like.

[0018] Identifying the relevant security controls to properly secure a customer service is an arduous process that requires specialized knowledge of the various cloud services and/or industry standards. This process is further complicated when customer services are deployed in multi-cloud and/or hybrid cloud environments. In such scenarios, securing the customer service may need to account for configurations of resources in many possible environments. Manual identification of control mappings rely on the knowledge of human beings, is time consuming, and is likely to result in imperfect and/or incorrect control mappings. While automatic identification of control mappings may be an improvement over manual identification, current solutions are unable to account for nuances in natural language descriptions associated with customer services, policy descriptions, policy rationales, and/or control mappings.

[0019] Embodiments disclosed herein are directed to automatic identification of control mappings using a machine-learning model. For example, a policy predictor model trained using machine-

learning techniques can account for nuances in natural language descriptions and find the likeliest control mappings. Furthermore, the policy predictor model may be trained with control mappings applicable to various cloud environments to enable the policy predictor model to identify proper control mappings for services deployed in multi-cloud environments. Additionally, results provided by a policy predictor model may include a confidence score to enable automated implementation of control mappings in high-confidence scenarios.

[0020] In embodiments, subject matter experts may develop control mappings between security policies and corresponding security controls. Such control mappings facilitate compliance with security policies by facilitating the identification of security controls that, when implemented, result in compliance with the security policy. In order to improve identification of the security policies relevant to a particular customer service, embodiments disclosed herein employ a policy predictor model that is trained to predict a control mapping based on a similarity between a textual input and textual descriptions associated with the security policies.

[0021] In embodiments, the policy predictor model is trained based a labeled training dataset that includes a large number of training samples that comprise a textual description (e.g., policy description, control description, rationale, etc.) of a control mapping that maps a control identifier to a policy and/or rationale, and a label that identifies the control mapping. In embodiments, a first subset of the labeled training dataset may be used to train the policy predictor model, while a second subset may be reserved for validation of the trained policy predictor model. Various methods may be employed to train the policy predictor model, such as, but not limited to, the use of decision trees, a random forest algorithm, support vector machines, neural networks, k-Nearest Neighbors, K-means clustering, Naïve Bayes, gradient boosting, transfer learning, and/or the like.

[0022] In embodiments, the policy predictor model may be trained iteratively starting with a model initialized with initial weights. During an iteration, features are extracted from the textual descriptions of the security policies through various methods, such as, but not limited to, word embeddings, sentence embeddings, and/or the like. The extracted features may then be provided to the model to determine predicted control mapping identifiers. The predicted control mapping identifiers are compared, using a loss function, to the expected control mapping identifiers provided as labels in the labeled training dataset. The weights may be adjusted in iterations to reduce an error determined by the loss function. In embodiments, the policy predictor model may be trained until the error determined by the loss function satisfies an accuracy requirement and/or converges to a particular error value.

[0023] In embodiments, the policy predictor model may be trained using an ensemble method that combines the predictions of a plurality of machine-learning models. For example, the policy predictor model may be trained using a gradient boosting technique by building a series of weak learners and combining their predictions to create a robust and accurate model. This involves training a new learner to correct the errors made by the ensemble up to that point. Gradient boosting minimizes a predefined loss function by adjusting the weights of misclassified instances, where a new tree is fitted to the negative gradient of the loss function.

[0024] In embodiments, the trained policy predictor model may be validated based on a validation dataset. For instance, a subset of the labeled training dataset may be reserved for validation purposes. During validation, the trained policy predictor model is evaluated to determine how well the policy predictor model, which has been trained on a first subset of the labeled training dataset, can make accurate predictions on unseen data, such as a second subset of the labeled training dataset. In embodiments, the training and validation process may be performed through cross-validation, where iterations in a series of iterations includes training the model using a training dataset and validating the model using a separate validation dataset.

[0025] Once trained, the policy predictor model may be employed to predict, based on a textual description of a customer service, a relevant control mapping associated with a security policy and a corresponding security control. In embodiments, automatic identification of security controls

begins by determining service information associated with the customer service. For instance, the service information may be determined by receiving, from a customer via an interface, a textual description of their service, such as, but not limited to, service requirements, a possible security threat, an attack vector associated with the security threat, a description of the security threat, and/or the like. In embodiments, the textual description of the service may be determined by analyzing an architectural diagram of the customer service using a threat analysis tool. In embodiments, the service information may be determined based on information from a plurality of sources. For example, the service information may be compiled from a plurality of cloud environments and/or on-premises resources to which the customer service is deployed.

[0026] In embodiments, the service information is provided as a textual input to the policy predictor model to determine a relevant control mapping. The policy predictor model may, in embodiments, extract features from the input text by, for example, but not limited to, generating a word embedding and/or sentence embedding representative of the semantic meaning of the input text. In embodiments, the policy predictor engine may output one or more predicted control mappings based on the extracted features. For instance, the policy predictor model may output the predicted control mapping with the highest confidence and/or similarity score, and/or output a ranked list of control mappings ranked based on their confidence and/or similarity scores. In embodiments, the policy predictor model may output one or more control mapping identifiers that identify a security policy and/or a corresponding security control. For instance, control mapping identifiers may, in embodiments, be determined based on a security policy identifier that identifies the security policy associated with the control mapping and a security control identifier that identifies the security control associated with the control mapping.

[0027] Various responsive actions may be performed based on a predicted control mapping that is associated with a predicted security policy and predicted security control. In embodiments, action handler may determine the responsive action to be performed based on one or more factors, such as, but not limited to, the predicted security policy, the predicted security control, customer preferences, the confidence and/or similarity score associated with the prediction, and/or the like. For instance, in embodiments, the predicted control mapping may be provided to the customer to enable the customer to implement the predicted security control on the customer service to comply with the predicted security policy. In embodiments, the predicted security control may be automatically implemented on behalf of the customer to comply with the predicted security policy. Automatic implementation of the predicted security control may, in embodiments, be performed by issuing one or more instructions and/or commands to one or more cloud and/or on-premises environments.

[0028] These and further embodiments are disclosed herein that enable the functionality described above and additional functionality. Such embodiments are described in further detail as follows.

[0029] For instance, FIG. 1 shows a block diagram of an example system **100** for automatically identifying a control mapping using a policy predictor model, in accordance with an embodiment. As shown in FIG. 1, system **100** includes one or more customer devices **102** communicatively coupled to a server infrastructure **104** via one or more networks **106**. Furthermore, customer device(s) **102** includes an application **118**, and server infrastructure **104** includes a service information processor **108**, a policy predictor model **110**, an action handler **112**, and one or more cloud environments **114**. Moreover, cloud environment(s) **114** may include one or more customer services **116** deployed thereon. System **100** is described in further detail as follows.

[0030] Customer device(s) **102** may comprise any type of stationary or mobile processing device, including, but not limited to, a desktop computer, a server, a mobile or handheld device (e.g., a tablet, a personal data assistant (PDA), a smart phone, a laptop, etc.), an Internet-of-Things (IoT) device, etc. As shown in FIG. 1, customer device(s) **102** includes an application **118** that can transmit to server infrastructure **104** one or more requests **120** for requesting one or more predicted control mappings **124**. Various example implementations of customer device(s) **102** are described

below in reference to FIG. 8 (e.g., computing device **802**, network-based server infrastructure **870**, on-premises servers **892**, and/or components thereof).

[0031] In embodiments, application **118** may comprise various applications, such as, but not limited to, mobile applications, desktop applications, a web browser, server applications, scripts, and/or the like, configured to generate request(s) **120** for requesting for predicted control mapping(s) **124**. In embodiments, application **118** may provide a user interface to enable a user to interact with policy predictor model **110** on server infrastructure **104** to request predicted control mapping(s) **124**. In embodiments, application **118** may automatically interact with policy predictor model **110** via an application programming interface (API) to request predicted control mapping(s) **124**. Various example implementations of application **118** are described below in reference to FIG. 8 (e.g., application programs **814**, application programs **876**, application programs **896**, and/or components thereof).

[0032] Network(s) **106** may comprise one or more networks such as local area networks (LANs), wide area networks (WANs), enterprise networks, the Internet, etc., and may include one or more wired and/or wireless portions. Various example implementations of network(s) **106** are described below in reference to FIG. 8 (e.g., network **804**, and/or components thereof).

[0033] Server infrastructure **104** may be a network-accessible server set (e.g., cloud-based environment or platform). In an embodiment, the underlying resources of server infrastructure **104** may be co-located (e.g., housed in one or more nearby buildings with associated components such as backup power supplies, redundant data communications, environmental controls, etc.) to form a datacenter, may be distributed across different regions, and/or may be arranged in other manners. In embodiments, the underlying resources of server infrastructure **104** may be owned and/or operated by a plurality of different entities (e.g., cloud providers). Various example implementations of server infrastructure **104** are described below in reference to FIG. 8 (e.g., network-based server infrastructure **870**, and/or components thereof).

[0034] Customer service(s) **116** may include any computing resource and/or application of a cloud customer that is hosted in a hybrid-cloud and/or multi-cloud environment. For instance, customer service(s) **116** may be deployed on one or more of cloud environment(s) **114** and may utilize computing resources, storage resources, applications, and/or cloud services provided by cloud environment(s) **114**. In embodiments, customer service(s) **116** may be duplicated across cloud environment(s) **114** in order to provide redundancy and improve service availability. In embodiments, customer service(s) **116** may combine computing resources, storage resources, applications, and/or cloud services provided by different cloud environment(s) **114** to provide customer service(s) **116**. In embodiments, customer service(s) **116** may also utilize computing resources, storage resources, applications, and/or other services provided by an on-premises resource (not depicted).

[0035] Cloud environment(s) **114** may include cloud environments owned and/or operated by different entities (e.g., cloud providers) that provide cloud resources and/or services, such as, but not limited to, Infrastructure as a Service (IaaS), Platform as a Service (PaaS), Software as a Service (SaaS), Anything (or everything) as a service (XaaS), and/or the like. In embodiments, cloud environment(s) **114** may provide customer interfaces to provision, configure, deploy, monitor, and/or manage customer service(s) **116**. For instance, cloud environment(s) **114** may, in embodiments, provide one or more APIs to enable customers to provision, configure, deploy, monitor, and/or manage customer service(s) **116**.

[0036] Service information processor **108** is configured to accept, as input, request(s) **120** comprising service information associated with a service of a customer and process the input to produce processed service information **122** that is in a format that enhances the quality and/or effectiveness of policy predictor model **110**. In embodiments, the input provided to service information processor **108** may include, but is not limited to, a description of the service, a service requirement associated with the service, resources associated with the service, entities associated

with the service, a threat assessment associated with the service, a threat vector, an architectural diagram associated with an architecture of the service, and/or the like. In embodiments, the processing performed by service information processor **108** may include, but is not limited to, cleaning of raw text data by removing irrelevant characters, symbols, and/or formatting inconsistencies, tokenization to break down sentences into individual words and/or subword units, removal of stop words and/or commonly occurring terms that may not contribute significantly to the semantic meaning, stemming and/or lemmatization to normalize word forms, feature extraction to encode the input into numerical vectors that capture semantic meaning of the input, dimensionality reduction to reduce the number of features needed to represent the input, and/or the like. In embodiments, feature extraction may be accomplished through various techniques, such as, but not limited to, bag-of-words to represent the input based on word frequencies, TF-IDF (Term Frequency-Inverse Document Frequency), to weigh the importance of terms in the input to a larger corpus, word and/or sentence embeddings, such as, but not limited to, Word2Vec or GloVe, to capture semantic relationships between words. In embodiments, service information processor **108** may convert non-textual input into a textual format. For instance, when provided with an input comprising an architectural diagram of the service, service information processor **108** may analyze the architectural diagram to generate a textual description of the service, including, but not limited to, a service requirement associated with the service, resources associated with the service, entities associated with the service, a threat assessment associated with the service, a threat vector, and/or the like. In embodiments, service information processor **108** may provide processed service information **122** to policy predictor model **110** for classification.

[0037] Policy predictor model **110** is configured to receive processed service information **122** from service information processor **108** and determine predicted control mapping(s) **124** that are relevant to processed service information **122**. In embodiments, predicted control mapping(s) **124** provided by policy predictor model **110** may include a confidence and/or similarity score that indicates a degree of relevancy or similarity between predicted control mapping(s) **124** and processed service information **122**. In embodiments, policy predictor model **110** is trained using a supervised machine learning algorithm that will be described in greater detail below in conjunction with FIG. 2.

[0038] Action handler **112** is configured to receive predicted control mapping(s) **124** from policy predictor model **110** and perform a responsive action based on the predicted control mapping(s) **124**. In embodiments, the responsive action performed may be determined based on one or more factors, such as, but not limited to, a predicted security policy associated with predicted control mapping(s) **124**, a predicted security control associated with predicted control mapping(s) **124**, customer preferences, the confidence and/or similarity score associated with predicted control mapping(s) **124**, and/or the like. For instance, action handler **112** may, in embodiments, provide the predicted control mapping(s) **124** to the application **118** to enable the customer to implement a predicted security control associated with predicted control mapping(s) **124** on customer service(s) **116** to comply with a predicted security policy associated with predicted control mapping(s) **124**. In embodiments, action handler **112** may automatically implement, on behalf of the customer, a predicted security control associated with predicted control mapping(s) **124** on customer service(s) **116** to comply with a predicted security policy associated with predicted control mapping(s) **124**. For instance, action handler **112** may, in embodiments, automatically implement a predicted security control associated with predicted control mapping(s) **124** by issuing one or more instructions and/or commands **126** to one or more on-premises resources (not depicted) and/or cloud environment(s) **114**. In embodiments, action handler **112** may issue instructions and/or commands **126** to cloud environment(s) **114** using one or more interfaces (e.g., APIs) provided by cloud environment(s) **114**. In embodiments, action handler **112** may issue instructions and/or commands **126** to cloud environment(s) **114** to modify a configuration and/or operation of customer service(s) **116**.

[0039] Embodiments described herein may operate in various ways to train a policy predictor model for identifying a control mapping and using the policy predictor model to identify a control mapping. For instance, FIG. 2 depicts a block diagram of an example system **200** for training a policy predictor model for identifying a control mapping and using the policy predictor model to identify a control mapping, in accordance with an embodiment. As shown in FIG. 2, system **200** includes customer device(s) **102**, server infrastructure **104**, network(s) **106**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, customer service(s) **116**, and application **118**. In system **200**, server infrastructure **104** further includes training data **202**, and a model trainer **204**. System **200** is described in further detail as follows.

[0040] In embodiments, training data **202** may include a labeled training dataset with a plurality of training samples that comprise a textual description (e.g., policy description, control description, rationale, etc.) of a control mapping that maps a control identifier to a policy and/or rationale, and a label that identifies the control mapping. In embodiments, the label may include a control mapping identifier that is determined based on a policy identifier and a control identifier associated with the policy and control, respectively, that constitute the control mapping. For instance, a control mapping identifier may, in embodiments, be determined by performing an operation, such as, but not limited to, a logical operation, a mathematical operation, a string operation, and/or the like, on the policy identifier and the control identifier.

[0041] In embodiments, training data **202** may be extracted and/or derived from security frameworks and/or benchmarks developed by a variety of entities, such as, but not limited to, technology companies, standards setting organizations (e.g., NIST, ISO, etc.), industry organizations, and/or regulatory agencies. For instance, textual descriptions of policies and/or controls may, in embodiments, be extracted from the security frameworks and/or benchmarks to form the textual description of the control mappings in the training samples. Similarly, policy identifier and/or control identifiers may, in embodiments, be extracted from the security frameworks and/or benchmarks to form the basis for the control mapping identifier in the training samples.

[0042] In embodiments, training data **202** is provided to service information processor **108** to produce processed training data **206** that is in a format that enhances the quality and/or effectiveness of policy predictor model **110**. In embodiments, training data **202** provided as input to service information processor **108** may include, but is not limited to, a description of the service, a service requirement associated with the service, resources associated with the service, entities associated with the service, a threat assessment associated with the service, a threat vector, and/or the like. In embodiments, the processing performed by service information processor **108** may include, but is not limited to, cleaning of raw text data by removing irrelevant characters, symbols, and/or formatting inconsistencies, tokenization to break down sentences into individual words and/or subword units, removal of stop words and/or commonly occurring terms that may not contribute significantly to the semantic meaning, stemming and/or lemmatization to normalize word forms, feature extraction to encode the input into numerical vectors that capture semantic meaning of the input, dimensionality reduction to reduce the number of features needed to represent the input, and/or the like. In embodiments, feature extraction may be accomplished through various techniques, such as, but not limited to, bag-of-words to represent the input based on word frequencies, TF-IDF (Term Frequency-Inverse Document Frequency), to weigh the importance of terms in the input to a larger corpus, word and/or sentence embeddings, such as, but not limited to, Word2Vec or GloVe, to capture semantic relationships between words. In embodiments, service information processor **108** may provide processed training data **206** to model trainer **204** for training and/or validating policy predictor model **110**.

[0043] Model trainer **204** is configured to train, validate, and/or deploy policy predictor model **110** based on processed training data **206**. In embodiments, model trainer **204** may train policy predictor model **110** by iteratively training and validating policy predictor model **110** using separate

subsets of processed training data **206**. For instance, during an iteration, a first subset of processed training data **206** may be used to train policy predictor model **110**, while a second subset of processed training data **206** may be reserved for validation of policy predictor model **110**. Various methods may be employed to train policy predictor model **110**, such as, but not limited to, the use of decision trees, a random forest algorithm, support vector machines, neural networks, k-Nearest Neighbors, K-means clustering, Naïve Bayes, gradient boosting, transfer learning, and/or the like. [0044] In embodiments, policy predictor model **110** may be trained iteratively starting with a model that is initialized with initial weights. During an iteration, a subset of processed training data **206** is provided to the model to determine predicted control mapping identifiers. The predicted control mapping identifiers are compared, using a loss function, to the expected control mapping identifiers provide as labels in the labeled training dataset. In embodiments, the weights may be adjusted in an iteration to reduce an error determined by the loss function.

[0045] In embodiments, model trainer **204** may train policy predictor model **110** using an ensemble method that combines the predictions of a plurality of machine-learning models. For example, model trainer **204** may train policy predictor model **110** using a gradient boosting technique by building a series of weak learners and combining their predictions to create a robust and accurate model. By training a new learner to correct the errors made by the ensemble up to that point, gradient boosting minimizes a predefined loss function by adjusting the weights of misclassified instances, where a new tree is fitted to the negative gradient of the loss function.

[0046] In embodiments, model trainer **204** may validate policy predictor model **110** based on a second subset of processed training data **206** that is reserved for validation purposes. During validation, model trainer **204** may evaluate the ability of policy predictor model **110** to predict a control mapping based on unseen data (e.g., the second subset of the labeled training dataset). In embodiments, the training and validation process may be performed through cross-validation, where iterations in a series of iterations include training the model using a training dataset and validating the model using a separate validation dataset.

[0047] In embodiments, model trainer **204** may train policy predictor model **110** until an error determined by the loss function satisfies an accuracy requirement and/or converges to a particular error value. Once policy predictor model **110** is trained, model trainer **204** may deploy, as trained model **208**, policy predictor model **110** to enable policy predictor model **110** to classify customer data in a production environment. For example, model trainer **204** may, in embodiments, deploy one or more instances of policy predictor model **110** to one or more servers of infrastructure **104**.

[0048] Embodiments described herein may operate in various ways to automatically identify a control mapping. For instance, FIG. 3 depicts a flowchart **300** of a process for automatically identifying a control mapping, in accordance with an embodiment. Server infrastructure **104**, service information processor **108**, policy predictor model **110**, and/or action handler **112** of FIGS. 1-2 may operate according to flowchart **300**, for example. Note that not all steps of flowchart **300** may need to be performed in all embodiments, and in some embodiments, the steps of flowchart **300** may be performed in different orders than shown. Flowchart **300** is described as follows with respect to FIGS. 1-2 for illustrative purposes.

[0049] Flowchart **300** starts at step **302**. In step **302**, service information related to a service of a customer is determined. For example, service information processor **108** may determine service information related to customer service(s) **116**. In embodiments, service information processor **108** may determine service information related to customer service(s) **116** based on information provided by customer device(s) **102** via request(s) **120**. In embodiments, service information processor **108** may determine service information related to customer service(s) **116** by receiving service information from cloud environment(s) **114** via one or more interfaces (e.g., APIs) provided by cloud environment(s) **114**. In embodiments, service information processor **108** may process the determined service information to generate processed service information **122**, and provide processed service information **122** to policy predictor model **110** for classification.

[0050] In step **304**, the service information is provided to a policy predictor model trained to predict a control mapping based on the service information. For example, service information processor **108** may provide processed service information **122** to policy predictor model **110** to predict a control mapping based on processed service information

[0051] In step **306**, a predicted control mapping associated with a predicted security policy and a predicted security control is received from the policy predictor model, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy. For example, action handler **112** may receive, from policy predictor model **110**, predicted control mapping(s) **124**. In embodiments, action handler **112** may receive, for predicted control mapping(s) **124**, confidence and/or similarity score(s) indicative of a degree of relevance between processed service information **122** and predicted control mapping(s) **124**.

[0052] In step **308**, an action is performed based on the predicted control mapping. For example, action handler **112** may perform an action based on predicted control mapping(s) **124**. As discussed above, action handler **112** may determine the responsive action to be performed based on one or more factors, such as, but not limited to, a predicted security policy associated with predicted control mapping(s) **124**, a predicted security control associated with predicted control mapping(s) **124**, customer preferences, the confidence and/or similarity score associated with predicted control mapping(s) **124**, and/or the like. For instance, action handler **112** may, in embodiments, provide the predicted control mapping(s) **124** to the application **118** to enable the customer to implement a predicted security control associated with predicted control mapping(s) **124** on customer service(s) **116** to comply with a predicted security policy associated with predicted control mapping(s) **124**. In embodiments, action handler **112** may automatically implement, on behalf of the customer, a predicted security control associated with predicted control mapping(s) **124** on customer service(s) **116** to comply with a predicted security policy associated with predicted control mapping(s) **124**. For instance, action handler **112** may, in embodiments, automatically implement a predicted security control associated with predicted control mapping(s) **124** by issuing one or more instructions and/or commands to one or more on-premises resources (not depicted) and/or cloud environment(s) **114**. In embodiments, action handler **112** may issue instructions and/or commands to cloud environment(s) **114** using one or more interfaces (e.g., APIs) provided by cloud environment(s) **114**. In embodiments, action handler **112** may issue instructions and/or commands to cloud environment(s) **114** to modify a configuration and/or operation of customer service(s) **116**.

[0053] Embodiments described herein may operate in various ways to generate a textual description of a service based on an architectural diagram of the service. For instance, FIG. 4 depicts a flowchart **400** of a process for generating a textual description of a service based on an architectural diagram of the service, in accordance with an embodiment. Server infrastructure **104**, service information processor **108**, policy predictor model **110**, and/or action handler **112** of FIGS. 1-2 may operate according to flowchart **400**, for example. Note that not all steps of flowchart **400** may need to be performed in all embodiments, and in some embodiments, the steps of flowchart **400** may be performed in different orders than shown. Flowchart **400** is described as follows with respect to FIGS. 1-2 for illustrative purposes.

[0054] Flowchart **400** starts at step **402**. In step **402**, an architectural diagram indicative of a service architecture associated with the service of the customer is received, the architecture diagram comprising requirements of the service of the customer. For example, service information processor **108** may receive an architectural diagram indicative of a service architecture of customer service(s) **116** via request(s) **120**.

[0055] In step **404**, the architecture diagram is analyzed to determine a security threat associated with the service architecture. For example, service information processor **108** may analyze the architecture diagram to determine a security threat associated with the service architecture. In embodiments, service information processor **108** may analyze the architectural diagram using a

threat assessment tool (not depicted).

[0056] In step **406**, a textual description of the architectural diagram is determined, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat. For example, service information processor **108** may determine a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat. In embodiments, service information processor **108** may determine the textual description of the architectural diagram using a threat assessment tool (not depicted).

[0057] Embodiments described herein may operate in various ways to train a policy predictor model for predicting a control mapping for a service based on a textual description of the service. For instance, FIG. 5 depicts a flowchart **500** of a process for training a policy predictor model for predicting a control mapping for a service based on a textual description of the service, in accordance with an embodiment. Server infrastructure **104**, service information processor **108**, policy predictor model **110**, action handler **112**, training data **202**, and/or model trainer **204** of FIGS. 1-2 may operate according to flowchart **500**, for example. Note that not all steps of flowchart **500** may need to be performed in all embodiments, and in some embodiments, the steps of flowchart **500** may be performed in different orders than shown. Flowchart **500** is described as follows with respect to FIGS. 1-2 for illustrative purposes.

[0058] Flowchart **500** starts at step **502**. In step **502**, a labeled training dataset is received, the labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control. For example, model trainer **204** may receive processed training data **206**. As discussed above, processed training data **206** may include a plurality of training samples that comprise a textual description (e.g., policy description, control description, rationale, etc.) of a control mapping that maps a control identifier to a policy and/or rationale, and a label that identifies the control mapping. In embodiments, the label may include a control mapping identifier that is determined based on a policy identifier and a control identifier associated with the policy and control, respectively, that constitute the control mapping.

[0059] In step **504**, the policy predictor model is iteratively trained based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy. For example, model trainer **204** may train policy predictor model **110** based on a first subset of processed training data **206**. As discussed above, various methods may be employed to train policy predictor model **110**, such as, but not limited to, the use of decision trees, a random forest algorithm, support vector machines, neural networks, k-Nearest Neighbors, K-means clustering, Naïve Bayes, gradient boosting, transfer learning, and/or the like.

[0060] In step **506**, a second subset of the labeled training dataset is used to validate the policy predictor model. For example, model trainer **204** may validate policy predictor model **110** using a second subset of processed training data **206** reserved for validation purposes.

[0061] In step **508**, the policy predictor model is deployed. For example, model trainer **204** may deploy policy predictor model **110**. In embodiments, model trainer **204** may train policy predictor model **110** until an error determined by the loss function satisfies an accuracy requirement and/or converges to a particular error value. Once policy predictor model **110** is trained, model trainer **204** may deploy policy predictor model **110** to enable policy predictor model **110** to classify customer data in a production environment. For example, model trainer **204** may, in embodiments, deploy policy predictor model **110** to one or more servers of infrastructure **104**.

[0062] Embodiments described herein may operate in various ways to automatically identify a control mapping for a multi-cloud service and implement the identified control mapping in the multi-cloud environment. For instance, FIG. 6 depicts a block diagram of an example system **600**

for automatically identifying a control mapping for a multi-cloud service and implanting the identified control mapping in the multi-cloud environment, in accordance with an embodiment. As shown in FIG. 6, system **600** includes server infrastructure **104**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, and customer service(s) **116**. System **600** is described in further detail as follows.

[0063] As shown in FIG. 6, service information processor **108** may receive service information **602** from cloud environment(s) **114**. For instance, service information processor **108** may, in embodiments, determine service information related to customer service(s) **116** that is deployed in a hybrid-cloud and/or multi-cloud environment by receiving service information from cloud environment(s) **114** via one or more interfaces (e.g., APIs) provided by cloud environment(s) **114**. As discussed above, customer service(s) **116** may be deployed on one or more of cloud environment(s) **114** and may utilize computing resources, storage resources, applications, and/or cloud services provided by cloud environment(s) **114**. In embodiments, customer service(s) **116** may be duplicated across cloud environment(s) **114** in order to provide redundancy and improve service availability. In embodiments, customer service(s) **116** may combine computing resources, storage resources, applications, and/or cloud services provided by different cloud environment(s) **114** to provide customer service(s) **116**. In embodiments, customer service(s) **116** may also utilize computing resources, storage resources, applications, and/or other services provided by an on-premises resource (not depicted).

[0064] Moreover, as shown in FIG. 6, action handler **112** may, in embodiments, action handler **112** may automatically implement, on behalf of the customer, a predicted security control associated with predicted control mapping(s) **124** on customer service(s) **116** to comply with a predicted security policy associated with predicted control mapping(s) **124**. For instance, action handler **112** may, in embodiments, issue one or more instructions and/or commands **126** to one or more cloud environment(s) **114**. In embodiments, action handler **112** may issue instructions and/or commands **126** to cloud environment(s) **114** using one or more interfaces (e.g., APIs) provided by cloud environment(s) **114** to modify a configuration and/or operation of customer service(s) **116**.

[0065] Embodiments described herein may operate in various ways to automatically determine service information associated with a multi-cloud service operating in a multi-cloud environment. For instance, FIG. 7 depicts a flowchart **700** of a process for automatically determining service information associated with a multi-cloud service operating in a multi-cloud environment, in accordance with an embodiment. Server infrastructure **104**, service information processor **108**, policy predictor model **110**, and/or action handler **112**, of FIGS. 1-2 and 6 may operate according to flowchart **700**, for example. Note that not all steps of flowchart **700** may need to be performed in all embodiments, and in some embodiments, the steps of flowchart **700** may be performed in different orders than shown. Flowchart **700** is described as follows with respect to FIGS. 1-2 and 6 for illustrative purposes.

[0066] Flowchart **700** starts at step **702**. In step **702**, a first subset of the service information associated with a first resource of the service operating in the first cloud environment is received from the first cloud environment. For example, service information processor **108** may receive, from a first cloud environment of cloud environment(s) **114**, a first subset of service information **602** associated with a first resource (not shown) operating in the first cloud environment of cloud environment(s) **114**.

[0067] In step **704**, a second subset of the service information associated with a second resource of the service operating in the second cloud environment is received from the second cloud environment. For example, service information processor **108** may receive, from a second cloud environment of cloud environment(s) **114**, a second subset of service information **602** associated with a second resource (not shown) operating in the second cloud environment of cloud environment(s) **114**.

III. Example Mobile Device and Computer System Implementation

[0068] The systems and methods described above in reference to FIGS. 1-7, including customer device(s) **102**, server infrastructure **104**, network(s) **106**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, customer service(s) **116**, application **118**, training data **202**, model trainer **204**, and/or each of the components described therein, and/or the steps of flowcharts **300**, **400** and/or **500** may be implemented in hardware, or hardware combined with one or both of software and/or firmware. For example, customer device(s) **102**, server infrastructure **104**, network(s) **106**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, customer service(s) **116**, application **118**, training data **202**, model trainer **204**, and/or each of the components described therein, and/or the steps of flowcharts **300**, **400** and/or **500** may be each implemented as computer program code/instructions configured to be executed in one or more processors and stored in a computer readable storage medium. Alternatively, customer device(s) **102**, server infrastructure **104**, network(s) **106**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, customer service(s) **116**, application **118**, training data **202**, model trainer **204**, and/or each of the components described therein, and/or the steps of flowcharts **300**, **400** and/or **500** may be each implemented in one or more SoCs (system on chip). An SoC may include an integrated circuit chip that includes one or more of a processor (e.g., a central processing unit (CPU), microcontroller, microprocessor, digital signal processor (DSP), etc.), memory, one or more communication interfaces, and/or further circuits, and may optionally execute received program code and/or include embedded firmware to perform functions.

[0069] Embodiments disclosed herein may be implemented in one or more computing devices that may be mobile (a mobile device) and/or stationary (a stationary device) and may include any combination of the features of such mobile and stationary computing devices. Examples of computing devices in which embodiments may be implemented are described as follows with respect to FIG. 8. FIG. 8 shows a block diagram of an exemplary computing environment **800** that includes a computing device **802**. Computing device **802** is an example of customer device(s) **102** shown in FIGS. 1 and 2, which may each include one or more of the components of computing device **802**. In some embodiments, computing device **802** is communicatively coupled with devices (not shown in FIG. 8) external to computing environment **800** via network **804**. Network **804** comprises one or more networks such as local area networks (LANs), wide area networks (WANs), enterprise networks, the Internet, etc., and may include one or more wired and/or wireless portions. Network **804** may additionally or alternatively include a cellular network for cellular communications. Computing device **802** is described in detail as follows.

[0070] Computing device **802** can be any of a variety of types of computing devices. For example, computing device **802** may be a mobile computing device such as a handheld computer (e.g., a personal digital assistant (PDA)), a laptop computer, a tablet computer, a hybrid device, a notebook computer, a netbook, a mobile phone (e.g., a cell phone, a smart phone, etc.), a wearable computing device (e.g., a head-mounted augmented reality and/or virtual reality device including smart glasses), or other type of mobile computing device. Computing device **802** may alternatively be a stationary computing device such as a desktop computer, a personal computer (PC), a stationary server device, a minicomputer, a mainframe, a supercomputer, etc.

[0071] As shown in FIG. 8, computing device **802** includes a variety of hardware and software components, including a processor **810**, a storage **820**, one or more input devices **850**, one or more output devices **850**, one or more wireless modems **860**, one or more wired interfaces **860**, a power supply **862**, a location information (LI) receiver **864**, and an accelerometer **866**. Storage **820** includes memory **856**, which includes non-removable memory **822** and removable memory **824**, and a storage device **890**. Storage **820** also stores an operating system **812**, application programs **814**, and application data **816**. Wireless modem(s) **860** include a Wi-Fi modem **862**, a Bluetooth modem **864**, and a cellular modem **866**. Output device(s) **850** includes a speaker **852** and a display **854**. Input device(s) **850** includes a touch screen **852**, a microphone **854**, a camera **856**, a physical

keyboard **858**, and a trackball **840**. Not all components of computing device **802** shown in FIG. **8** are present in all embodiments, additional components not shown may be present, and any combination of the components may be present in a particular embodiment. These components of computing device **802** are described as follows.

[0072] A single processor **810** (e.g., central processing unit (CPU), microcontroller, a microprocessor, signal processor, ASIC (application specific integrated circuit), and/or other physical hardware processor circuit) or multiple processors **810** may be present in computing device **802** for performing such tasks as program execution, signal coding, data processing, input/output processing, power control, and/or other functions. Processor **810** may be a single-core or multi-core processor, and each processor core may be single-threaded or multithreaded (to provide multiple threads of execution concurrently). Processor **810** is configured to execute program code stored in a computer readable medium, such as program code of operating system **812** and application programs **814** stored in storage **820**. The program code is structured to cause processor **810** to perform operations, including the processes/methods disclosed herein. Operating system **812** controls the allocation and usage of the components of computing device **802** and provides support for one or more application programs **814** (also referred to as “applications” or “apps”). Application programs **814** may include common computing applications (e.g., e-mail applications, calendars, contact managers, web browsers, messaging applications), further computing applications (e.g., word processing applications, mapping applications, media player applications, productivity suite applications), one or more machine learning (ML) models, as well as applications related to the embodiments disclosed elsewhere herein. Processor(s) **810** may include one or more general processors (e.g., CPUs) configured with or coupled to one or more hardware accelerators, such as one or more NPUs and/or one or more GPUs.

[0073] Any component in computing device **802** can communicate with any other component according to function, although not all connections are shown for ease of illustration. For instance, as shown in FIG. **8**, bus **806** is a multiple signal line communication medium (e.g., conductive traces in silicon, metal traces along a motherboard, wires, etc.) that may be present to communicatively couple processor **810** to various other components of computing device **802**, although in other embodiments, an alternative bus, further buses, and/or one or more individual signal lines may be present to communicatively couple components. Bus **806** represents one or more of any of several types of bus structures, including a memory bus or memory controller, a peripheral bus, an accelerated graphics port, and a processor or local bus using any of a variety of bus architectures.

[0074] Storage **820** is physical storage that includes one or both of memory **856** and storage device **890**, which store operating system **812**, application programs **814**, and application data **816** according to any distribution. Non-removable memory **822** includes one or more of RAM (random access memory), ROM (read only memory), flash memory, a solid-state drive (SSD), a hard disk drive (e.g., a disk drive for reading from and writing to a hard disk), and/or other physical memory device type. Non-removable memory **822** may include main memory and may be separate from or fabricated in a same integrated circuit as processor **810**. As shown in FIG. **8**, non-removable memory **822** stores firmware **818**, which may be present to provide low-level control of hardware. Examples of firmware **818** include BIOS (Basic Input/Output System, such as on personal computers) and boot firmware (e.g., on smart phones). Removable memory **824** may be inserted into a receptacle of or otherwise coupled to computing device **802** and can be removed by a user from computing device **802**. Removable memory **824** can include any suitable removable memory device type, including an SD (Secure Digital) card, a Subscriber Identity Module (SIM) card, which is well known in GSM (Global System for Mobile Communications) communication systems, and/or other removable physical memory device type. One or more of storage device **890** may be present that are internal and/or external to a housing of computing device **802** and may or may not be removable. Examples of storage device **890** include a hard disk drive, a SSD, a thumb

drive (e.g., a USB (Universal Serial Bus) flash drive), or other physical storage device.

[0075] One or more programs may be stored in storage **820**. Such programs include operating system **812**, one or more application programs **814**, and other program modules and program data. Examples of such application programs may include, for example, computer program logic (e.g., computer program code/instructions) for implementing customer device(s) **102**, server infrastructure **104**, network(s) **106**, service information processor **108**, policy predictor model **110**, action handler **112**, cloud environment(s) **114**, customer service(s) **116**, application **118**, training data **202**, model trainer **204**, and/or each of the components described therein, as well as any of flowcharts **300**, **400**, **500**, and/or any individual steps thereof.

[0076] Storage **820** also stores data used and/or generated by operating system **812** and application programs **814** as application data **816**. Examples of application data **816** include web pages, text, images, tables, sound files, video data, and other data, which may also be sent to and/or received from one or more network servers or other devices via one or more wired or wireless networks. Storage **820** can be used to store further data including a subscriber identifier, such as an International Mobile Subscriber Identity (IMSI), and an equipment identifier, such as an International Mobile Equipment Identifier (IMEI). Such identifiers can be transmitted to a network server to identify users and equipment.

[0077] A user may enter commands and information into computing device **802** through one or more input devices **850** and may receive information from computing device **802** through one or more output devices **850**. Input device(s) **850** may include one or more of touch screen **852**, microphone **854**, camera **856**, physical keyboard **858** and/or trackball **840** and output device(s) **850** may include one or more of speaker **852** and display **854**. Each of input device(s) **850** and output device(s) **850** may be integral to computing device **802** (e.g., built into a housing of computing device **802**) or external to computing device **802** (e.g., communicatively coupled wired or wirelessly to computing device **802** via wired interface(s) **860** and/or wireless modem(s) **860**). Further input devices **850** (not shown) can include a Natural User Interface (NUI), a pointing device (computer mouse), a joystick, a video game controller, a scanner, a touch pad, a stylus pen, a voice recognition system to receive voice input, a gesture recognition system to receive gesture input, or the like. Other possible output devices (not shown) can include piezoelectric or other haptic output devices. Some devices can serve more than one input/output function. For instance, display **854** may display information, as well as operating as touch screen **852** by receiving user commands and/or other information (e.g., by touch, finger gestures, virtual keyboard, etc.) as a user interface. Any number of each type of input device(s) **850** and output device(s) **850** may be present, including multiple microphones **854**, multiple cameras **856**, multiple speakers **852**, and/or multiple displays **854**.

[0078] One or more wireless modems **860** can be coupled to antenna(s) (not shown) of computing device **802** and can support two-way communications between processor **810** and devices external to computing device **802** through network **804**, as would be understood to persons skilled in the relevant art(s). Wireless modem **860** is shown generically and can include a cellular modem **866** for communicating with one or more cellular networks, such as a GSM network for data and voice communications within a single cellular network, between cellular networks, or between the mobile device and a public switched telephone network (PSTN). Wireless modem **860** may also or alternatively include other radio-based modem types, such as a Bluetooth modem **864** (also referred to as a “Bluetooth device”) and/or Wi-Fi modem **862** (also referred to as an “wireless adaptor”). Wi-Fi modem **862** is configured to communicate with an access point or other remote Wi-Fi-capable device according to one or more of the wireless network protocols based on the IEEE (Institute of Electrical and Electronics Engineers) 802.11 family of standards, commonly used for local area networking of devices and Internet access. Bluetooth modem **864** is configured to communicate with another Bluetooth-capable device according to the Bluetooth short-range wireless technology standard(s) such as IEEE 802.15.1 and/or managed by the Bluetooth Special

Interest Group (SIG).

[0079] Computing device **802** can further include power supply **862**, LI receiver **864**, accelerometer **866**, and/or one or more wired interfaces **860**. Example wired interfaces **860** include a USB port, IEEE 1394 (Fire Wire) port, a RS-232 port, an HDMI (High-Definition Multimedia Interface) port (e.g., for connection to an external display), a DisplayPort port (e.g., for connection to an external display), an audio port, and/or an Ethernet port, the purposes and functions of each of which are well known to persons skilled in the relevant art(s). Wired interface(s) **860** of computing device **802** provide for wired connections between computing device **802** and network **804**, or between computing device **802** and one or more devices/peripherals when such devices/peripherals are external to computing device **802** (e.g., a pointing device, display **854**, speaker **852**, camera **856**, physical keyboard **858**, etc.). Power supply **862** is configured to supply power to each of the components of computing device **802** and may receive power from a battery internal to computing device **802**, and/or from a power cord plugged into a power port of computing device **802** (e.g., a USB port, an A/C power port). LI receiver **864** may be used for location determination of computing device **802** and may include a satellite navigation receiver such as a Global Positioning System (GPS) receiver or may include other type of location determiner configured to determine location of computing device **802** based on received information (e.g., using cell tower triangulation, etc.). Accelerometer **866** may be present to determine an orientation of computing device **802**.

[0080] Note that the illustrated components of computing device **802** are not required or all-inclusive, and fewer or greater numbers of components may be present as would be recognized by one skilled in the art. For example, computing device **802** may also include one or more of a gyroscope, barometer, proximity sensor, ambient light sensor, digital compass, etc. Processor **810** and memory **856** may be co-located in a same semiconductor device package, such as being included together in an integrated circuit chip, FPGA, or system-on-chip (SOC), optionally along with further components of computing device **802**.

[0081] In embodiments, computing device **802** is configured to implement any of the above-described features of flowcharts herein. Computer program logic for performing any of the operations, steps, and/or functions described herein may be stored in storage **820** and executed by processor **810**.

[0082] In some embodiments, server infrastructure **870** may be present in computing environment **800** and may be communicatively coupled with computing device **802** via network **804**. Server infrastructure **870**, when present, may be a network-accessible server set (e.g., a cloud-based environment or platform). As shown in FIG. 8, server infrastructure **870** includes clusters **872**. Each of clusters **872** may comprise a group of one or more compute nodes and/or a group of one or more storage nodes. For example, as shown in FIG. 8, cluster **872** includes nodes **874**. Each of nodes **874** are accessible via network **804** (e.g., in a “cloud-based” embodiment) to build, deploy, and manage applications and services. Any of nodes **874** may be a storage node that comprises a plurality of physical storage disks, SSDs, and/or other physical storage devices that are accessible via network **804** and are configured to store data associated with the applications and services managed by nodes **874**. For example, as shown in FIG. 8, nodes **874** may store application data **878**.

[0083] Each of nodes **874** may, as a compute node, comprise one or more server computers, server systems, and/or computing devices. For instance, a node **874** may include one or more of the components of computing device **802** disclosed herein. Each of nodes **874** may be configured to execute one or more software applications (or “applications”) and/or services and/or manage hardware resources (e.g., processors, memory, etc.), which may be utilized by users (e.g., customers) of the network-accessible server set. For example, as shown in FIG. 8, nodes **874** may operate application programs **876**. In an implementation, a node of nodes **874** may operate or comprise one or more virtual machines, with each virtual machine emulating a system architecture

(e.g., an operating system), in an isolated manner, upon which applications such as application programs **876** may be executed.

[0084] In an embodiment, one or more of clusters **872** may be co-located (e.g., housed in one or more nearby buildings with associated components such as backup power supplies, redundant data communications, environmental controls, etc.) to form a datacenter, or may be arranged in other manners. Accordingly, in an embodiment, one or more of clusters **872** may be a datacenter in a distributed collection of datacenters. In embodiments, exemplary computing environment **800** comprises part of a cloud-based platform.

[0085] In an embodiment, computing device **802** may access application programs **876** for execution in any manner, such as by a client application and/or a browser at computing device **802**.

[0086] For purposes of network (e.g., cloud) backup and data security, computing device **802** may additionally and/or alternatively synchronize copies of application programs **814** and/or application data **816** to be stored at network-based server infrastructure **870** as application programs **876** and/or application data **878**. For instance, operating system **812** and/or application programs **814** may include a file hosting service client configured to synchronize applications and/or data stored in storage **820** at network-based server infrastructure **870**.

[0087] In some embodiments, on-premises servers **892** may be present in computing environment **800** and may be communicatively coupled with computing device **802** via network **804**. On-premises servers **892**, when present, are hosted within an organization's infrastructure and, in many cases, physically onsite of a facility of that organization. On-premises servers **892** are controlled, administered, and maintained by IT (Information Technology) personnel of the organization or an IT partner to the organization. Application data **898** may be shared by on-premises servers **892** between computing devices of the organization, including computing device **802** (when part of an organization) through a local network of the organization, and/or through further networks accessible to the organization (including the Internet). Furthermore, on-premises servers **892** may serve applications such as application programs **896** to the computing devices of the organization, including computing device **802**. Accordingly, on-premises servers **892** may include storage **894** (which includes one or more physical storage devices such as storage disks and/or SSDs) for storage of application programs **896** and application data **898** and may include one or more processors for execution of application programs **896**. Still further, computing device **802** may be configured to synchronize copies of application programs **814** and/or application data **816** for backup storage at on-premises servers **892** as application programs **896** and/or application data **898**.

[0088] Embodiments described herein may be implemented in one or more of computing device **802**, network-based server infrastructure **870**, and on-premises servers **892**. For example, in some embodiments, computing device **802** may be used to implement systems, clients, or devices, or components/subcomponents thereof, disclosed elsewhere herein. In other embodiments, a combination of computing device **802**, network-based server infrastructure **870**, and/or on-premises servers **892** may be used to implement the systems, clients, or devices, or components/subcomponents thereof, disclosed elsewhere herein.

[0089] As used herein, the terms “computer program medium,” “computer-readable medium,” “computer-readable storage medium,” and “computer-readable storage device,” etc., are used to refer to physical hardware media. Examples of such physical hardware media include any hard disk, optical disk, SSD, other physical hardware media such as RAMs, ROMs, flash memory, digital video disks, zip disks, MEMs (microelectronic machine) memory, nanotechnology-based storage devices, and further types of physical/tangible hardware storage media of storage **820**. Such computer-readable media and/or storage media are distinguished from and non-overlapping with communication media and propagating signals (do not include communication media and propagating signals). Communication media embodies computer-readable instructions, data structures, program modules or other data in a modulated data signal such as a carrier wave. The

term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, communication media includes wireless media such as acoustic, RF, infrared, and other wireless media, as well as wired media. Embodiments are also directed to such communication media that are separate and non-overlapping with embodiments directed to computer-readable storage media.

[0090] As noted above, computer programs and modules (including application programs **814**) may be stored in storage **820**. Such computer programs may also be received via wired interface(s) **860** and/or wireless modem(s) **860** over network **804**. Such computer programs, when executed or loaded by an application, enable computing device **802** to implement features of embodiments discussed herein. Accordingly, such computer programs represent controllers of the computing device **802**.

[0091] Embodiments are also directed to computer program products comprising computer code or instructions stored on any computer-readable medium or computer-readable storage medium. Such computer program products include the physical storage of storage **820** as well as further physical storage types.

IV. Additional Example Embodiments

[0092] In embodiments, a method comprises: determining service information related to a service of a customer; providing the service information to a policy predictor model trained to predict a control mapping based on the service information; receiving, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and performing an action based on the predicted control mapping.

[0093] In embodiments, determining service information comprises: receiving an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyzing the architecture diagram to determine a security threat associated with the service architecture; and determining, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.

[0094] In embodiments, the service of the customer is deployed in a multi-cloud environment, and determining service information comprises: receiving, from a first cloud environment of the multi-cloud environment, a first subset of the service information comprising service information associated with a first resource of the service, the first resource operating in the first cloud environment; and receiving, from a second cloud environment of the multi-cloud environment, a second subset of the service information comprising service information associated with a second resource of the service, the second resource operating in the second cloud environment.

[0095] In embodiments, performing an action comprises at least one of: providing the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically applying the predicted security control to the service of the customer.

[0096] In embodiments, the method further comprises: receiving a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively training the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validating the policy predictor model based on a second subset of the labeled training dataset; and deploying the policy predictor model.

[0097] In embodiments, the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack vector associated

with the security threat; a description of the security threat; or a description of mitigation of the security threat.

[0098] In embodiments, the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.

[0099] In embodiments, a system comprises: a processor; and a memory device comprising program code structured to cause the processor to: determine service information related to a service of a customer; provide the service information to a policy predictor model trained to predict a control mapping based on the service information; receive, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and perform an action based on the predicted control mapping.

[0100] In embodiments, to determine the service information, the program code is structured to further cause the processor to: receive an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyze the architecture diagram to determine a security threat associated with the service architecture; and determine, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.

[0101] In embodiments, the service of the customer is deployed in a multi-cloud environment, and, to determine the service information, the program code is structured to further cause the processor to: receive, from a first cloud environment of the multi-cloud environment, a first subset of the service information comprising service information associated with a first resource of the service, the first resource operating in the first cloud environment; and receive, from a second cloud environment of the multi-cloud environment, a second subset of the service information comprising service information associated with a second resource of the service, the second resource operating in the second cloud environment.

[0102] In embodiments, to perform the action, the program code is structured to further cause the processor to perform at least one of: provide the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically apply the predicted security control to the service of the customer.

[0103] In embodiments, the program code is structured to further cause the processor to: receive a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively train the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validate the policy predictor model based on a second subset of the labeled training dataset; and deploy the policy predictor model.

[0104] In embodiments, the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack vector associated with the security threat; a description of the security threat; or a description of mitigation of the security threat.

[0105] In embodiments, the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.

[0106] In embodiments, a computer-readable storage medium comprises computer-executable instructions that, when executed by a processor, cause the processor to: receive service information

related to a service of a customer, the service deployed in a multi-cloud environment, at least a first portion of the service information received from a first cloud environment of the multi-cloud environment and a second portion of the service information received from a second cloud environment of the multi-cloud environment; provide the service information to a policy predictor model trained to predict a control mapping based on the service information; receive, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and perform an action based on the predicted control mapping.

[0107] In embodiments, to determine the service information, the computer-executable instructions, when executed by the processor, further cause the processor to: receive an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyze the architecture diagram to determine a security threat associated with the service architecture; and determine, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.

[0108] In embodiments, to perform the action, the computer-executable instructions, when executed by the processor, further cause the processor to perform at least one of: provide the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically apply the predicted security control to the service of the customer.

[0109] In embodiments, the computer-executable instructions, when executed by the processor, further cause the processor to: receive a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively train the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validate the policy predictor model based on a second subset of the labeled training dataset; and deploy the policy predictor model.

[0110] In embodiments, the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack vector associated with the security threat; a description of the security threat; or a description of mitigation of the security threat.

[0111] In embodiments, the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.

V. Conclusion

[0112] References in the specification to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to effect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.

[0113] In the discussion, unless otherwise stated, adjectives such as “substantially” and “about” modifying a condition or relationship characteristic of a feature or features of an embodiment of the disclosure, are understood to mean that the condition or characteristic is defined to within tolerances that are acceptable for operation of the embodiment for an application for which it is

intended. Furthermore, where “based on” is used to indicate an effect being a result of an indicated cause, it is to be understood that the effect is not required to only result from the indicated cause, but that any number of possible additional causes may also contribute to the effect. Thus, as used herein, the term “based on” should be understood to be equivalent to the term “based at least on.” [0114] While various embodiments of the present disclosure have been described above, it should be understood that they have been presented by way of example only, and not limitation. It will be understood by those skilled in the relevant art(s) that various changes in form and details may be made therein without departing from the spirit and scope of the invention as defined in the appended claims. Accordingly, the breadth and scope of the present invention should not be limited by any of the above-described exemplary embodiments, but should be defined only in accordance with the following claims and their equivalents.

Claims

1. A method comprising: determining service information related to a service of a customer; providing the service information to a policy predictor model trained to predict a control mapping based on the service information; receiving, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and performing an action based on the predicted control mapping.
2. The method of claim 1, wherein said determining service information comprises: receiving an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyzing the architecture diagram to determine a security threat associated with the service architecture; and determining, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.
3. The method of claim 1, wherein the service of the customer is deployed in a multi-cloud environment, and wherein said determining service information comprises: receiving, from a first cloud environment of the multi-cloud environment, a first subset of the service information comprising service information associated with a first resource of the service, the first resource operating in the first cloud environment; and receiving, from a second cloud environment of the multi-cloud environment, a second subset of the service information comprising service information associated with a second resource of the service, the second resource operating in the second cloud environment.
4. The method of claim 1, wherein said performing an action comprises at least one of: providing the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically applying the predicted security control to the service of the customer.
5. The method of claim 1, further comprising: receiving a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively training the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validating the policy predictor model based on a second subset of the labeled training dataset; and deploying the policy predictor model.
6. The method of claim 1, wherein the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack

vector associated with the security threat; a description of the security threat; or a description of mitigation of the security threat.

7. The method of claim 1, wherein the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.

8. A system comprising: a processor; and a memory device comprising program code structured to cause the processor to: determine service information related to a service of a customer; provide the service information to a policy predictor model trained to predict a control mapping based on the service information; receive, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and perform an action based on the predicted control mapping.

9. The system of claim 8, wherein, to determine the service information, the program code is structured to further cause the processor to: receive an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyze the architecture diagram to determine a security threat associated with the service architecture; and determine, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.

10. The system of claim 8, wherein the service of the customer is deployed in a multi-cloud environment, and wherein, to determine the service information, the program code is structured to further cause the processor to: receive, from a first cloud environment of the multi-cloud environment, a first subset of the service information comprising service information associated with a first resource of the service, the first resource operating in the first cloud environment; and receive, from a second cloud environment of the multi-cloud environment, a second subset of the service information comprising service information associated with a second resource of the service, the second resource operating in the second cloud environment.

11. The system of claim 8, wherein, to perform the action, the program code is structured to further cause the processor to perform at least one of: provide the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically apply the predicted security control to the service of the customer.

12. The system of claim 8, wherein the program code is structured to further cause the processor to: receive a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively train the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validate the policy predictor model based on a second subset of the labeled training dataset; and deploy the policy predictor model.

13. The system of claim 8, wherein the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack vector associated with the security threat; a description of the security threat; or a description of mitigation of the security threat.

14. The system of claim 8, wherein the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.

15. A computer-readable storage medium comprising computer-executable instructions that, when executed by a processor, cause the processor to: receive service information related to a service of a

customer, the service deployed in a multi-cloud environment, at least a first portion of the service information received from a first cloud environment of the multi-cloud environment and a second portion of the service information received from a second cloud environment of the multi-cloud environment; provide the service information to a policy predictor model trained to predict a control mapping based on the service information; receive, from the policy predictor model, a predicted control mapping associated with a predicted security policy and a predicted security control, an implementation of the predicted security control enables the service of the customer to comply with a service requirement associated with the predicted security policy; and perform an action based on the predicted control mapping.

16. The computer-readable storage medium of claim 15, wherein, to determine the service information, the computer-executable instructions, when executed by the processor, further cause the processor to: receive an architectural diagram indicative of a service architecture associated with the service of the customer, the architecture diagram comprising requirements of the service of the customer; analyze the architecture diagram to determine a security threat associated with the service architecture; and determine, as the service information, a textual description of the architectural diagram, the textual description comprising at least one of: an attack vector associated with the security threat, a description of the security threat, or a description of mitigation of the security threat.

17. The computer-readable storage medium of claim 15, wherein, to perform the action, the computer-executable instructions, when executed by the processor, further cause the processor to perform at least one of: provide the predicted control mapping to the customer, the customer enabled to apply the predicted security control to the service; or automatically apply the predicted security control to the service of the customer.

18. The computer-readable storage medium of claim 15, wherein the computer-executable instructions, when executed by the processor, further cause the processor to: receive a labeled training dataset comprising a textual description of a first security policy that is labeled with a first control mapping identifier associated with the first security policy and a first security control; iteratively train the policy predictor model based on a first subset of the labeled training dataset using a supervised machine learning algorithm to predict the first control mapping identifier based on a similarity to the textual description of the first security policy; validate the policy predictor model based on a second subset of the labeled training dataset; and deploy the policy predictor model.

19. The computer-readable storage medium of claim 15, wherein the service information comprises at least one of: a service requirement associated with the service; a security threat associated with the service; an attack vector associated with the security threat; a description of the security threat; or a description of mitigation of the security threat.

20. The computer-readable storage medium of claim 15, wherein the predicted security control comprises a control associated with at least one of: a framework developed by a subject matter expert; a standard developed by a standards setting organization; or a regulation developed by a regulatory agency.
