

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250280043

Kind Code

A1

Publication Date

September 04, 2025

Inventor(s)

Baker; Brandon S. et al.

SECURE EXECUTION OF USER-DEFINED FUNCTIONS

Abstract

The subject technology receives, by one or more hardware processors, a request to execute a user-defined function (UDF) within a sandbox process. The subject technology establishes a secure egress path for the UDF using an overlay network, where the overlay network includes a dedicated DNS resolver at a proxy service. The subject technology receives, from the UDF, a DNS request to resolve a hostname. The subject technology validates, by the proxy service, that the hostname is included in an allowed host list associated with the UDF. The subject technology resolves, by the dedicated DNS resolver, the hostname to an IP address using a UDP listener configured to handle DNS protocol traffic on a designated port of the proxy service. The subject technology enables the UDF to communicate with a host at the resolved IP address via the secure egress path.

Inventors: Baker; Brandon S. (Redmond, WA), Denny-Brown; Derek (Seattle, WA), Halcrow; Michael A. (Kirkland, WA), Konigsmark; Sven Tenzing Choden (Seattle, WA), Sharma; Niranjana Kumar (Redmond, WA), Sharma; Nitya Kumar (Bellevue, WA), Yu; Haowei (Newark, CA), Zhan; Andong (San Mateo, CA)

Applicant: Snowflake Inc. (Bozeman, MT)

Family ID: 90184485

Appl. No.: 19/213325

Filed: May 20, 2025

Related U.S. Application Data

parent US continuation 18428890 20240131 parent-grant-document US 12323468 child US 19213325

parent US continuation 18309024 20230428 parent-grant-document US 11930045 child US 18428890

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/20 (20130101); H04L63/0245 (20130101); H04L63/101 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is Continuation of U.S. patent application Ser. No. 18/428,890, filed on Jan. 31, 2024, which is a Continuation of U.S. patent application Ser. No. 18/309,024, filed on Apr. 28, 2023, which claims the benefit of priority thereto, the contents of each are incorporated herein by reference in their entireties for all purposes.

TECHNICAL FIELD

[0002] The subject matter disclosed herein generally relates to methods, systems, machine storage mediums, and computer programs for providing secure access for sandboxed user-defined functions.

BACKGROUND

[0003] Network-based database systems may be provided through a cloud data platform, which allows organizations, customers, and users to store, manage, and retrieve data from the cloud. With respect to this type of data processing, a cloud data platform could implement online transactional processing, online analytical processing, and/or another type of data processing. Moreover, a cloud data platform could be or include a relational database management system and/or one or more other types of database management systems.

[0004] Cloud-based data warehouses and other database systems sometimes provide support for custom user functions, such as a User-Defined Function (UDF) or stored procedures that enable such systems to perform operations that are not available through the built-in, system-defined functions. Existing techniques for the execution of UDFs, however, may lack robust security mechanisms for mitigating the associated security risks and ensuring that the user code is executed securely and with sufficient visibility for auditing.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] The present disclosure will be apparent from the following more particular description of examples of embodiments of the technology, as illustrated in the accompanying drawings. The drawings are not necessarily to scale, emphasis instead being placed upon illustrating embodiments of the present disclosure. In the drawings, like numerals may describe similar components in different views. Like numerals having different letter suffixes may represent different instances of similar components. Various ones of the appended drawings merely illustrate example embodiments of the present disclosure and should not be considered as limiting its scope.

[0006] FIG. 1 illustrates an example computing environment that includes a cloud data platform in communication with a cloud storage provider system, according to some example embodiments.

[0007] FIG. 2 is a block diagram illustrating components of a compute service manager, according to some example embodiments.

[0008] FIG. 3 is a block diagram illustrating components of an execution platform, according to

some example embodiments.

[0009] FIG. 4 is a computing environment conceptually illustrating an example software architecture executing a user-defined function (UDF) by a process running on a given execution node of the execution platform, according to some example embodiments.

[0010] FIG. 5 is a block diagram illustrating a User-Defined Function (UDx) external access system architecture, according to some example embodiments.

[0011] FIG. 6 is a block diagram illustrating an external function access system architecture, according to some example embodiments.

[0012] FIG. 7 is a block diagram illustrating a system architecture depicting an overlay network, according to some example embodiments.

[0013] FIG. 8 is a block diagram illustrating GENEVE encapsulation, according to some example embodiments.

[0014] FIG. 9 is a block diagram illustrating a high-level overview depicting the flow of metrics and events, according to some example embodiments.

[0015] FIG. 10 is a block diagram illustrating a high-level overview depicting eBPF mapping, according to some example embodiments.

[0016] FIG. 11 is a block diagram illustrating eBPF-based packet routing, according to some example embodiments.

[0017] FIG. 12 is a block diagram illustrating a packet routing diagram, according to some example embodiments.

[0018] FIG. 13 is a block diagram illustrating a packet routing diagram depicting how a returned packet is routed, according to some example embodiments.

[0019] FIG. 14 is a block diagram illustrating an example service proxy process used by the cloud data platform to route and manage network traffic, according to some example embodiments.

[0020] FIG. 15 is a block diagram illustrating an example methodology used by the cloud data platform to route and manage network traffic, according to some example embodiments.

[0021] FIG. 16 is a block diagram illustrating an example methodology used by the cloud data platform to manage network traffic, according to some example embodiments.

[0022] FIG. 17 is a flow diagram illustrating operations of a cloud data platform performing an example method for permitting any sandboxed user-defined function code to access the Internet, according to some example embodiments.

[0023] FIG. 18 illustrates a diagrammatic representation of a machine in the form of a computer system within which a set of instructions can be executed for causing the machine to perform any one or more of the methodologies discussed herein, according to some example embodiments.

DETAILED DESCRIPTION

[0024] The description that follows includes systems, methods, techniques, instruction sequences, and computing machine program products that embody illustrative embodiments of the disclosure. In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide an understanding of various embodiments of the inventive subject matter. It will be evident, however, to those skilled in the art, that embodiments of the inventive subject matter can be practiced without these specific details. In general, well-known instruction instances, protocols, structures, and techniques are not necessarily shown in detail.

[0025] From the perspective of a user, external code can be made a seamless part of a cloud data platform functionality similar to any internally defined function or procedure. That is, the cloud data platform can access systems that are, by their nature, external to the cloud data platform (e.g., geocoding systems). Further, users are enabled to use any arbitrary external code regardless of the language used to author the code or the system on which the code executes. Moreover, the cloud data platform described herein allows users to invoke external functionality while avoiding security concerns that can arise from executing the code within the cloud data platform and do so in a manner that is orthogonal to the functionality of the cloud data platform. In addition, the technique

for invoking external functionality from the cloud data platform externalizes security-sensitive authentication information as this information is handled by the cloud data platform.

[0026] Example embodiments of the present disclosure include a cloud data platform (e.g., a cloud-based data warehousing platform) that enables users to have direct access to the Internet from a User-Defined Function (UDx), where the UDx includes code that runs in a sandbox process or sandboxed environment. More specifically, example embodiments allow users to directly connect to external resources from sandboxed code. External access from the UDx is a critical requirement for users (e.g., customers) of the cloud data platform that cannot be met with previous technical solutions, such as external functions.

[0027] For purposes of this description, example embodiments can apply to a User-Defined Function (UDF), User-Defined Table Function (UDTF), User-Defined Aggregation Function (UDAF), external functions, Streamlit®, or other stored procedures used in relational databases for performing complex data processing tasks, enforcing business rules, and the like can be applied or employed according to the present disclosure. However, for simplicity, the detailed embodiments will describe examples of providing secure external access to the UDx executing within a sandbox environment directly to the Internet using familiar programming languages (e.g., Java, Scala, Python, etc.), but it will be understood that the same principles may be used for other types of database logic and programmatic constructs from a sandboxed environment or a non-sandboxed environment. For example, although example embodiments describe external access of user-defined functions in a sandboxed environment, similar logic can be applied to non-sandboxed environments, such as external access of user-defined functions in containerized environments, or other constructs of the cloud data platform.

[0028] In computer security, a sandbox (e.g., sandbox environment) is a security mechanism for separating running programs, usually to prevent system failures or prevent exploitation of software vulnerabilities. A sandbox can be used to execute untested or untrusted packages, programs, functions, or code, possibly from unverified or untrusted third parties, suppliers, users, or websites, without risking harm to the host machine or operating system. A sandbox can provide a tightly controlled set of resources for guest programs to run in, such as storage and memory scratch space. Network access, the ability to inspect the host system or read from input devices can be disallowed or restricted. UDxs typically can run in a sandbox environment. Some example embodiments described herein can be run within a sandbox environment, which is described and depicted in more detail in connection with FIG. 4.

[0029] Traditional approaches for accessing the Internet required the use of external functions (described in more detail below); however, sandboxed UDFs have no way of accessing the Internet directly. Currently, user code running in a sandboxed environment or a sandbox process cannot make network calls. The sandbox does not have a network device, and the required system calls are blocked. However, simply enabling networking in the sandbox is not enough, as the cloud data platform must ensure that the traffic can only go to allowed destinations, does not impact other workloads, and is well-monitored. For example, currently for users who want their workloads in the cloud data platform to access external services must use the external function feature, which constrains the service to be something that must implement a Hypertext Transfer Protocol Secure (HTTPS) endpoint that is accessible from the cloud data platform and must accept JavaScript Object Notation (JSON) inputs and return JSON outputs representing batches of rows. While the external function feature is important, it requires that customers must take care in the way they use batches of rows with external functions. And, while customers can mitigate the JSON issue with a customer serializer/deserializer tied to JavaScript, such prior solutions are cumbersome. Thus, while prior solutions make it possible for customers to cobble together an end-to-end solution for many of their data access needs using the prior external function feature, it is unwieldy and less streamlined of an experience for the customer.

[0030] For example, external functions are a type of UDF, which call executable code that is

developed, maintained, stored, and executed outside of the cloud data platform, where the remotely executed code is known as a remote service. Information sent to a remote service is usually relayed through a proxy service. Inside the cloud data platform, the external function is stored as a database object that contains information that the cloud data platform uses to call the remote service. This stored information can include a Uniform Resource Locator (URL) of the proxy service that relays information to and from the remote service (e.g., information specified as part of a CREATE EXTERNAL FUNCTION command). The remotely executed code is known as a remote service, which can act like a function (e.g., it must return a value). Examples of a remote service can be implemented as an AWS Lambda® function, a Microsoft Azure® function, an HTTPS server (e.g., Node.js) running on an EC2® instance, and the like. The cloud data platform calls a proxy service instead of a remote service directly, which relays the data to the remote device. The proxy service can increase security by authenticating requests to the remote service. The proxy service can also relay the response from the remote service back to the cloud data platform. Examples of a proxy service can include Amazon® Application Programming Interface (API) Gateway, Microsoft Azure® API Management service, and the like. Users can write and call their own remote services, or call remote services written by third parties. From the perspective of a user running a Structured Query Language (SQL) statement, an external function behaves like any other UDF.

[0031] According to prior existing solutions, external functions can be burdensome to implement. In order to develop an external function, developers typically had to write code to convert the batches into a format the ultimate service they are using can understand, where that code needed to live externally (e.g., in a Lambda function, or in a request translator in the cloud data platform). In either case, that translation logic ends up being separated from other logic (e.g., error handling, conditional business logic, etc.) that should belong together, making the end-to-end hard to maintain.

[0032] Prior existing solutions can have additional complications that make technological advancements necessary. For example, APIs can be non-standard, such that the code needed to produce or handle batches of records in the cloud data platform's format can be substantial and can be unique to a cloud data platform implementation. According to some prior existing solutions, the cloud data platform currently exposes a batched scalar function interface, such that users wishing to implement a table function or stored procedure require additional interfaces to be supported. Further existing limitations can include limited interfaces only interacting via HTTP. In order to use protocols not mediated by HTTP (e.g., like those underlying most Java Database Connectivity (JDBC) connections), users must build and host an HTTP frontend outside of the cloud data platform.

[0033] Examples of the present disclosure overcome the current technological challenges by providing a generic and orthogonal solution that adds external access as an option to all sandboxed code in the cloud data platform. With external access as an orthogonal component, users are able to implement external variants of any contract the cloud data platform exposes, as it is exposed. Additional improvements enable users to write standard code. For example, users already know how to interact with remote services, and they do so by using standard libraries and policies (e.g. Python requests library, JDBC, etc.). Example embodiments enable users to incorporate or bring along these libraries without modification and enable users to use them with little to no changes in their code. Examples of the present disclosure overcome the current technological challenges by retaining security and providing control to the account administrators. For example, administrators on the machine must explicitly control the ability to reach out of the system to prevent unintentional data egress. Examples of the system strictly enforce these controls and provide an agent-based approach to reach data stored in private networks.

[0034] Example embodiments of the present disclosure overcome the technical challenges relating to Internet access by providing strict egress control and network isolation. The methods, systems, and computer programs described herein allow for direct UDX external access using the advantages

of a secure egress path from the UDx running in a sandbox process to the Internet. According to some examples, systems, methods, and computer programs provide a mechanism to permit any sandboxed UDx code to reach out securely to the hosts explicitly allowed (e.g., via an allowlist) by an account administrator. Examples provide characteristics and requirements to allow administrators and developers to write standard code and use familiar libraries to perform external access for a wide variety of scenarios, while leaving the account administrators in control of data egress.

[0035] The new solution presented herein overcomes the prior existing technological challenge of accessing the Internet by providing a secure egress path from UDx executing in sandboxes to the Internet. This new solution allows a UDx to be written in standard code and use other third-party libraries overcoming the usability challenges associated with external functions. For example, the code for the remote service can be written in languages that other UDFs cannot be written in, the remote services can use functions and libraries that cannot be accessed by internal UDFs (e.g., remote services can interface with commercially available third-party libraries), developers can write remote services that can be called from both the cloud data platform and from other software written to use the same interface, and other such technological advantages.

[0036] According to some example embodiments, allowing external access from a UDx will unblock both external customers of the cloud data platform and internal users of the cloud data platform (e.g., native applications, Connectors, Streamlit®, etc.). For example, customers and users work with some data that may not always exist within the cloud data platform, such as metadata, reference data, or even low latency cache. Enabling direct external access according to examples presented herein provides a path for users to connect these databases via tools (e.g., JDBC, database Software Development Kits (SDKs)) in order to pull in that data during a larger query or data pipeline process. It will also be possible to write a code to interact with Kafka topics from inside the cloud data platform once Internet access is enabled.

[0037] In some example embodiments, architectural improvements (e.g., functional requirements and non-functional requirements) can include support for standard code and libraries, strict egress control, strict network isolation, performance, and/or additional improvements are provided. For example, by supporting standard code and libraries, customers are able to bring their code and libraries from local environments or other hosted environments without modifications, such as customizations or specialized APIs required by the cloud data platform. Therein overcoming one of the largest usability challenges from prior solutions such as external functions.

[0038] For example, by providing strict egress control, any egress traffic from sandboxes must go through one or more egress proxies, where any egress traffic can be allowed if it is authorized by the database administrator. Strict egress control can be guaranteed by the cloud data platform as violations will be blocked and reported. Strict network isolation is provided as any egress traffic from a sandbox is not allowed to access any other sandboxes or any other execution platform instances. In addition, the cloud data platform can manage performance, observability, and debuggability by providing efficient external access while ensuring the egress traffic is well-monitored, traced, and providing feedback to the cloud data platform and users.

[0039] Example embodiments include technical solutions over prior external function features by providing direct access to services that inherently live outside of the cloud data platform. Additional examples include enabling users to store their API key as a secret in the cloud data platform, configure an integration object that permits use of this API key and access to the remote API, and create a standard or relatively-standard Java or Python function that uses the stored credential to make the call using a standard API (e.g., a request).

[0040] According to some examples, users may desire to ingest data from third-party platforms (e.g., Software-as-a-Service (SaaS) data ingestion) that effectively function as databases and expose a web API. While this can be accomplished by existing scalar external functions, such prior methods are complicated and require additional steps. For example, using existing scalar external

functions results in collections that are packaged and returned as variant objects that require being broken out in the calling SQL, thereby complicating the calling query, requiring multiple calls to be made, and/or requiring special handling for variable return types or error conditions. Example embodiments of the present disclosure overcome the technical challenges of prior methods by providing direct sandbox access including table function interfacing to be wrapped around standard methods for calling and retrieving batches from an external resource, thereby reducing the amount of one-off code needed to implement such functions.

[0041] Example embodiments further provide account administrators with additional functionalities to enhance security, privileges, and abilities. For example, administrators can specify where sandboxed functions can connect, such as basing these specifications on Domain Name System (DNS) names or wildcard-based DNS patterns. Administrators can grant and revoke use of these access usage rights to developers in a granular fashion and allow developers to apply granted access rights to specific functions that the developers author. Administrators can additionally selectively grant developers the ability to use secrets to connect to a remote system, and administrators can easily revoke access and prevent code from reaching out.

[0042] Example embodiments provide developer users with the ability to use authorization they have been granted by an admin to create a function that can open a Transmission Control Protocol (TCP) or HTTP connection to an authorized host. Developers can use standard APIs for making these TCP or HTTP connections without needing to rewrite a JDBC connector to make it work using the cloud data platform's TCP system. Further examples enable external access for any sandboxed code of any type, for example, UDF, UDTF, UDAF, User Datagram Protocol (UDP), etc. Additional example embodiments provide developer users with the ability to enable batched requests. For example, if a user wants to call a batched scalar function (e.g., the HERE geocoder), the user can write a batched UDx to do so.

[0043] Additional examples of the system disclosed herein prevent code from making any connection to a host not authorized by its specification and prevent a developer from authorizing any code to connect to a host they have not been granted rights to access. The system can prevent any user from seeing the contents of a secret unless this permission has been explicitly granted; for example, a caller of a function should not gain access to a secret simply because they can call the function. Examples of the system can further prevent code from making any DNS requests for hosts it is not authorized to connect to (e.g., exfiltration risk). The system can ensure external control for any network access, whereby such requests can route through the cloud data platform egress proxy and a compromised execution platform node should not be able to bypass such restrictions. Additional examples of the system disclosed herein can capture and report to the account administrator attempts to open a connection to or perform a DNS lookup that was blocked by permissions.

[0044] FIG. 1 illustrates an example computing environment **100** that includes a database system in the example form of a cloud data platform **102**, in accordance with some embodiments of the present disclosure. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from FIG. 1. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the computing environment **100** to facilitate additional functionality that is not specifically described herein. In other embodiments, the computing environment may comprise another type of network-based database system or a cloud data platform.

[0045] As shown, the computing environment **100** comprises the cloud data platform **102** in communication with a cloud storage platform **104** (e.g., AWS®, Microsoft Azure Blob Storage®, or Google® Cloud Storage). The cloud data platform **102** is a network-based system used for reporting and analysis of integrated data from one or more disparate sources including one or more storage locations within the cloud storage platform **104**. The cloud data platform **102** can be a

network-based data platform or network-based data system. The cloud storage platform **104** comprises a plurality of computing machines and provides on-demand computer system resources such as data storage and computing power to the cloud data platform **102**.

[0046] The cloud data platform **102** comprises a compute service manager **108**, an execution platform **110**, a proxy service **115**, and one or more metadata databases **112**. The cloud data platform **102** hosts and provides data reporting and analysis services to multiple client accounts.

[0047] The compute service manager **108** coordinates and manages operations of the cloud data platform **102**. The compute service manager **108** also performs query optimization and compilation as well as managing clusters of computing services that provide compute resources (also referred to as “virtual warehouses”). The compute service manager **108** can support any number of client accounts such as end users providing data storage and retrieval requests, system administrators managing the systems and methods described herein, and other components/devices that interact with compute service manager **108**.

[0048] The compute service manager **108** is also in communication with a client device **114**. The client device **114** corresponds to a user of one of the multiple client accounts supported by the cloud data platform **102**. A user may utilize the client device **114** to submit data storage, retrieval, and analysis requests to the compute service manager **108**.

[0049] The compute service manager **108** is also coupled to one or more metadata databases **112** that store metadata pertaining to various functions and aspects associated with the cloud data platform **102** and its users. For example, a metadata database **112** may include a summary of data stored in remote data storage systems as well as data available from a local cache. Additionally, a metadata database **112** may include information regarding how data is organized in remote data storage systems (e.g., the cloud storage platform **104**) and the local caches. Information stored by a metadata database **112** allows systems and services to determine whether a piece of data needs to be accessed without loading or accessing the actual data from a storage device.

[0050] The compute service manager **108** is further coupled to the execution platform **110**, which provides multiple computing resources that execute various data storage and data retrieval tasks. The execution platform **110** is coupled to cloud storage platform **104**. The cloud storage platform **104** comprises multiple data storage devices **120-1** to **120-N**. In some embodiments, the data storage devices **120-1** to **120-N** are cloud-based storage devices located in one or more geographic locations. For example, the data storage devices **120-1** to **120-N** can be part of a public cloud infrastructure or a private cloud infrastructure. The data storage devices **120-1** to **120-N** may be hard disk drives (HDDs), solid state drives (SSDs), storage clusters, Amazon S3™ storage systems, or any other data storage technology. Additionally, the cloud storage platform **104** may include distributed file systems (e.g., as Hadoop Distributed File Systems (HDFS)), object storage systems, and the like.

[0051] The execution platform **110** comprises a plurality of compute nodes. A set of processes on a compute node executes a query plan compiled by the compute service manager **108**. The set of processes can include: a first process to execute the query plan; a second process to monitor and delete cache files using a least recently used (LRU) policy and implement an out of memory (OOM) error mitigation process; a third process that extracts health information from process logs and status to send back to the compute service manager **108**; a fourth process to establish communication with the compute service manager **108** after a system boot; and a fifth process to handle all communication with a compute cluster for a given job provided by the compute service manager **108** and to communicate information back to the compute service manager **108** and other compute nodes of the execution platform **110**.

[0052] The compute service manager **108**, metadata database(s) **112**, proxy service **115**, and execution platform **110** are operatively connected to a platform agent **109**, which provides for an agent in the execution platform **110** as a long running service to handle extended Berkeley Packet Filter (eBPF) related operations. The platform agent **109** can include a Remote Procedure Call

(RPC) server via a Unix domain socket that can handle requests sent from execution platform worker processes. Sample requests can include load specific eBPF programs, read/write to BPF maps, configure network devices, and the like. The platform agent **109** can further handle external access BPF code and can be extended to capture more BPF uses cases, while receiving relevant cloud data platform information from any of the compute service manager **108**, metadata database(s) **112**, proxy service **115**, execution platform **110**, or alternative operatively connected modules from within the cloud data platform **102**, or externally connected data sources. The platform agent **109** is depicted and described in combination with FIG. 6.

[0053] In some embodiments, communication links between elements of the computing environment **100** are implemented via one or more data communication networks. These data communication networks may utilize any communication protocol and any type of communication medium. In some embodiments, the data communication networks are a combination of two or more data communication networks (or sub-Networks) coupled to one another. In alternate embodiments, these communication links are implemented using any type of communication medium and any communication protocol.

[0054] The compute service manager **108**, metadata database(s) **112**, execution platform **110**, platform agent **109**, proxy service **115**, and cloud storage platform **104** are shown in FIG. 1 as individual discrete components. However, each of the compute service manager **108**, metadata database(s) **112**, proxy service **115**, execution platform **110**, platform agent **109**, and cloud storage platform **104** can be implemented as a distributed system (e.g., distributed across multiple systems/platforms at multiple geographic locations). Additionally, each of the compute service manager **108**, metadata database(s) **112**, execution platform **110**, platform agent **109**, proxy service **115**, and cloud storage platform **104** can be scaled up or down (independently of one another) depending on changes to the requests received and the changing needs of the cloud data platform **102**. Thus, in the described embodiments, the cloud data platform **102** is dynamic and supports regular changes to meet the current data processing needs.

[0055] During typical operation, the cloud data platform **102** processes multiple jobs determined by the compute service manager **108**. These jobs are scheduled and managed by the compute service manager **108** to determine when and how to execute the job. For example, the compute service manager **108** may divide the job into multiple discrete tasks and may determine what data is needed to execute each of the multiple discrete tasks. The compute service manager **108** may assign each of the multiple discrete tasks to one or more nodes of the execution platform **110** to process the task. The compute service manager **108** may determine what data is needed to process a task and further determine which nodes within the execution platform **110** are best suited to process the task. Some nodes may have already cached the data needed to process the task and, therefore, be a suitable candidate for processing the task. Metadata stored in a metadata database **112** assists the compute service manager **108** in determining which nodes in the execution platform **110** have already cached at least a portion of the data needed to process the task. One or more nodes in the execution platform **110** process the task using data cached by the nodes and, if necessary, data retrieved from the cloud storage platform **104**. It is desirable to retrieve as much data as possible from caches within the execution platform **110** because the retrieval speed is typically much faster than retrieving data from the cloud storage platform **104**.

[0056] As shown in FIG. 1, the computing environment **100** separates the execution platform **110** from the cloud storage platform **104**. In this arrangement, the processing resources and cache resources in the execution platform **110** operate independently of the data storage devices **120-1** to **120-N** in the cloud storage platform **104**. Thus, the computing resources and cache resources are not restricted to specific data storage devices **120-1** to **120-N**. Instead, all computing resources and all cache resources may retrieve data from, and store data to, any of the data storage resources in the cloud storage platform **104**.

[0057] The platform agent **109** is illustrated as a component of execution platform **110**; however,

additional example embodiments of the platform agent **109** can be implemented by any of the virtual warehouses of the execution platform **110**, such as the execution node **302-1**, compute service manager **108**, the request processing service **208**, the security manager **422**, and/or external components of the cloud data platform **102** in accordance with some embodiments of the present disclosure.

[0058] FIG. **2** is a block diagram **200** illustrating components of the compute service manager **108**, in accordance with some embodiments of the present disclosure. As shown in FIG. **2**, the compute service manager **108** includes an access manager **202** and a credential management system **204** coupled to access data storage device **206**, which is an example of the metadata database(s) **112**. Access manager **202** handles authentication and authorization tasks for the systems described herein. The credential management system **204** facilitates use of remote stored credentials to access external resources such as data resources in a remote storage device. As used herein, the remote storage devices may also be referred to as “persistent storage devices” or “shared storage devices.” For example, the credential management system **204** may create and maintain remote credential store definitions and credential objects (e.g., in the data storage device **206**). A remote credential store definition identifies a remote credential store and includes access information to access security credentials from the remote credential store. A credential object identifies one or more security credentials using non-sensitive information (e.g., text strings) that are to be retrieved from a remote credential store for use in accessing an external resource. When a request invoking an external resource is received at run time, the credential management system **204** and access manager **202** use information stored in the data storage device **206** (e.g., a credential object and a credential store definition) to retrieve security credentials used to access the external resource from a remote credential store.

[0059] A request processing service **208** manages received data storage requests and data retrieval requests (e.g., jobs to be performed on database data). For example, the request processing service **208** may determine the data to process a received query (e.g., a data storage request or data retrieval request). The data can be stored in a cache within the execution platform **110** or in a data storage device in cloud storage platform **104**.

[0060] A management console service **210** supports access to various systems and processes by administrators and other system managers. Additionally, the management console service **210** may receive a request to execute a job and monitor the workload on the system.

[0061] The compute service manager **108** also includes a job compiler **212**, a job optimizer **214**, and a job executor **216**. The job compiler **212** parses a job into multiple discrete tasks and generates the execution code for each of the multiple discrete tasks. The job optimizer **214** determines the best method to execute the multiple discrete tasks based on the data that needs to be processed. The job optimizer **214** also handles various data pruning operations and other data optimization techniques to improve the speed and efficiency of executing the job. The job executor **216** executes the execution code for jobs received from a queue or determined by the compute service manager **108**.

[0062] A job scheduler and coordinator **218** sends received jobs to the appropriate services or systems for compilation, optimization, and dispatch to the execution platform **110**. For example, jobs can be prioritized and then processed in that prioritized order. In an embodiment, the job scheduler and coordinator **218** determines a priority for internal jobs that are scheduled by the compute service manager **108** with other “outside” jobs such as user queries that can be scheduled by other systems in the database but may utilize the same processing resources in the execution platform **110**. In some embodiments, the job scheduler and coordinator **218** identifies or assigns particular nodes in the execution platform **110** to process particular tasks. A virtual warehouse manager **220** manages the operation of multiple virtual warehouses implemented in the execution platform **110**. For example, the virtual warehouse manager **220** may generate query plans for executing received queries.

[0063] Additionally, the compute service manager **108** includes a configuration and metadata manager **222**, which manages the information related to the data stored in the remote data storage devices and in the local buffers (e.g., the buffers in execution platform **110**). The configuration and metadata manager **222** uses metadata to determine which data files need to be accessed to retrieve data for processing a particular task or job. A monitor and workload analyzer **224** oversees processes performed by the compute service manager **108** and manages the distribution of tasks (e.g., workload) across the virtual warehouses and execution nodes in the execution platform **110**. The monitor and workload analyzer **224** also redistributes tasks, as needed, based on changing workloads throughout the cloud data platform **102** and may further redistribute tasks based on a user (e.g., “external”) query workload that may also be processed by the execution platform **110**. The configuration and metadata manager **222** and the monitor and workload analyzer **224** are coupled to a data storage device **226**. Data storage device **226** in FIG. 2 represents any data storage device within the cloud data platform **102**. For example, data storage device **226** may represent buffers in execution platform **110**, storage devices in cloud storage platform **104**, or any other storage device.

[0064] As described in embodiments herein, the compute service manager **108** validates all communication from an execution platform (e.g., the execution platform **110**) to validate that the content and context of that communication are consistent with the task(s) known to be assigned to the execution platform. For example, an instance of the execution platform executing a query A should not be allowed to request access to data-source D (e.g., data storage device **226**) that is not relevant to query A. Similarly, a given execution node (e.g., execution node **302-1**) may need to communicate with another execution node (e.g., execution node **302-2**), and should be disallowed from communicating with a third execution node (e.g., execution node **312-1**) and any such illicit communication can be recorded (e.g., in a log or other location). Also, the information stored on a given execution node is restricted to data relevant to the current query and any other data is unusable, rendered so by destruction or encryption where the key is unavailable.

[0065] FIG. 3 is a block diagram **300** illustrating components of the execution platform **110**, in accordance with some embodiments of the present disclosure. As shown in FIG. 3, the execution platform **110** includes multiple virtual warehouses, including virtual warehouse **1**, virtual warehouse **2**, and virtual warehouse N. Each virtual warehouse includes multiple execution nodes that each include a data cache and a processor. The virtual warehouses can execute multiple tasks in parallel by using the multiple execution nodes. As discussed herein, the execution platform **110** can add new virtual warehouses and drop existing virtual warehouses in real-time based on the current processing needs of the systems and users. This flexibility allows the execution platform **110** to quickly deploy large amounts of computing resources when needed without being forced to continue paying for those computing resources when they are no longer needed. All virtual warehouses can access data from any data storage device (e.g., any storage device in cloud storage platform **104**).

[0066] Although each virtual warehouse shown in FIG. 3 includes three execution nodes, a particular virtual warehouse may include any number of execution nodes. Further, the number of execution nodes in a virtual warehouse is dynamic, such that new execution nodes are created when additional demand is present, and existing execution nodes are deleted when they are no longer useful.

[0067] Each virtual warehouse is capable of accessing any of the data storage devices **120-1** to **120-N** shown in FIG. 1. Thus, the virtual warehouses are not necessarily assigned to a specific data storage device **120-1** to **120-N** and, instead, can access data from any of the data storage devices **120-1** to **120-N** within the cloud storage platform **104**. Similarly, each of the execution nodes shown in FIG. 3 can access data from any of the data storage devices **120-1** to **120-N**. In some embodiments, a particular virtual warehouse or a particular execution node can be temporarily assigned to a specific data storage device, but the virtual warehouse or execution node may later

access data from any other data storage device.

[0068] In the example of FIG. 3, virtual warehouse **1** includes three execution nodes **302-1**, **302-2**, and **302-N**. Execution node **302-1** includes a cache **304-1** and a processor **306-1**. Execution node **302-2** includes a cache **304-2** and a processor **306-2**. Execution node **302-N** includes a cache **304-N** and a processor **306-N**. Each execution node **302-1**, **302-2**, and **302-N** is associated with processing one or more data storage and/or data retrieval tasks. For example, a virtual warehouse may handle data storage and data retrieval tasks associated with an internal service, such as a clustering service, a materialized view refresh service, a file compaction service, a storage procedure service, or a file upgrade service. In other implementations, a particular virtual warehouse may handle data storage and data retrieval tasks associated with a particular data storage system or a particular category of data.

[0069] Similar to virtual warehouse **1** discussed above, virtual warehouse **2** includes three execution nodes **312-1**, **312-2**, and **312-N**. Execution node **312-1** includes a cache **314-1** and a processor **316-1**. Execution node **312-2** includes a cache **314-2** and a processor **316-2**. Execution node **312-N** includes a cache **314-N** and a processor **316-N**. Additionally, virtual warehouse **N** includes three execution nodes **322-1**, **322-2**, and **322-N**. Execution node **322-1** includes a cache **324-1** and a processor **326-1**. Execution node **322-2** includes a cache **324-2** and a processor **326-2**. Execution node **322-N** includes a cache **324-N** and a processor **326-N**.

[0070] In some embodiments, the execution nodes shown in FIG. 3 are stateless with respect to the data being cached by the execution nodes. For example, these execution nodes do not store or otherwise maintain state information about the execution node, or the data being cached by a particular execution node. Thus, in the event of an execution node failure, the failed node can be transparently replaced by another node. Since there is no state information associated with the failed execution node, the new (replacement) execution node can easily replace the failed node without concern for recreating a particular state.

[0071] Although the execution nodes shown in FIG. 3 each includes one data cache and one processor, alternate embodiments may include execution nodes containing any number of processors and any number of caches. Additionally, the caches may vary in size among the different execution nodes. The caches shown in FIG. 3 store, in the local execution node, data that was retrieved from one or more data storage devices in cloud storage platform **104**. Thus, the caches reduce or eliminate the bottleneck problems occurring in platforms that consistently retrieve data from remote storage systems. Instead of repeatedly accessing data from the remote storage devices, the systems and methods described herein access data from the caches in the execution nodes, which is significantly faster and avoids the bottleneck problem discussed above. In some embodiments, the caches are implemented using high-speed memory devices that provide fast access to the cached data. Each cache can store data from any of the storage devices in the cloud storage platform **104**.

[0072] Further, the cache resources and computing resources may vary between different execution nodes. For example, one execution node may contain significant computing resources and minimal cache resources, making the execution node useful for tasks that require significant computing resources. Another execution node may contain significant cache resources and minimal computing resources, making this execution node useful for tasks that require caching of large amounts of data. Yet another execution node may contain cache resources providing faster input-output operations, useful for tasks that require fast scanning of large amounts of data. In some embodiments, the cache resources and computing resources associated with a particular execution node are determined when the execution node is created, based on the expected tasks to be performed by the execution node.

[0073] Additionally, the cache resources and computing resources associated with a particular execution node may change over time based on changing tasks performed by the execution node. For example, an execution node may be assigned more processing resources if the tasks performed

by the execution node become more processor intensive. Similarly, an execution node may be assigned more cache resources if the tasks performed by the execution node require a larger cache capacity.

[0074] Although virtual warehouses **1**, **2**, and **N** are associated with the same execution platform **110**, the virtual warehouses can be implemented using multiple computing systems at multiple geographic locations. For example, virtual warehouse **1** can be implemented by a computing system at a first geographic location, while virtual warehouses **2** and **N** are implemented by another computing system at a second geographic location. In some embodiments, these different computing systems are cloud-based computing systems maintained by one or more different entities.

[0075] Additionally, each virtual warehouse is shown in FIG. **3** as having multiple execution nodes. The multiple execution nodes associated with each virtual warehouse can be implemented using multiple computing systems at multiple geographic locations. For example, an instance of virtual warehouse **1** implements execution nodes **302-1** and **302-2** on one computing platform at a geographic location and implements execution node **302-N** at a different computing platform at another geographic location. Selecting particular computing systems to implement an execution node may depend on various factors, such as the level of resources needed for a particular execution node (e.g., processing resource requirements and cache requirements), the resources available at particular computing systems, communication capabilities of networks within a geographic location or between geographic locations, and which computing systems are already implementing other execution nodes in the virtual warehouse.

[0076] Execution platform **110** is also fault tolerant. For example, if one virtual warehouse fails, that virtual warehouse is quickly replaced with a different virtual warehouse at a different geographic location. A particular execution platform **110** may include any number of virtual warehouses. Additionally, the number of virtual warehouses in a particular execution platform is dynamic, such that new virtual warehouses are created when additional processing and/or caching resources are needed. Similarly, existing virtual warehouses can be deleted when the resources associated with the virtual warehouse are no longer useful.

[0077] In some embodiments, the virtual warehouses may operate on the same data in cloud storage platform **104**, but each virtual warehouse has its own execution nodes with independent processing and caching resources. This configuration allows requests on different virtual warehouses to be processed independently and with no interference between the requests. This independent processing, combined with the ability to dynamically add and remove virtual warehouses, supports the addition of new processing capacity for new users without impacting the performance.

[0078] FIG. **4** is a computing environment **400** conceptually illustrating an example software architecture executing a user-defined function (UDx) by a process running on a given execution node of the execution platform **110** of FIG. **3**, in accordance with some embodiments of the present disclosure.

[0079] As illustrated, the execution node **302-1** from the execution platform **110** includes an execution node process **410**, which in an embodiment is running on the processor **306-1** and can also utilize memory from the cache **304-1** (or another memory device or storage). As mentioned herein, a “process” or “computing process” can refer to an instance of a computer program that is being executed by one or more threads by an execution node or execution platform.

[0080] As mentioned before, the compute service manager **108** validates all communication from the execution platform **110** to validate that the content and context of that communication are consistent with the task(s) known to be assigned to the execution platform **110**. For example, the execution platform **110** executing a query **A** is not allowed to request access to a particular data source (e.g., data storage device **226** or any one of the storage devices in the cloud storage platform **104**) that is not relevant to query **A**. In an example, the execution node **302-1** may need to

communicate with a second execution node (e.g., execution node **302-2**), but the security mechanisms described herein can disallow communication with a third execution node (e.g., execution node **312-1**). Moreover, any such illicit communication can be recorded (e.g., in a log **444** or other location). Further, the information stored on a given execution node is restricted to data relevant to the current query and any other data is unusable by destruction or encryption where the key is unavailable.

[0081] The execution node process **410** is executing a UDF client **412** in the example of FIG. 4. In an embodiment, the UDF client **412** is implemented to support UDFs written in a particular programming language such as JAVA, and the like. In an embodiment, the UDF client **412** is implemented in a different programming language (e.g., C or C++) than the user code **430**, which can further improve security of the computing environment **400** by using a different codebase (e.g., one with the same or fewer potential security exploits).

[0082] User code **430** may be provided as a package, e.g., in the form of a JAR (JAVA archive) file which includes code for one or more UDFs. Server implementation code **432**, in an embodiment, is a JAR file that initiates a server which is responsible for receiving requests from the execution node process **410**, assigning worker threads to execute user code, and returning the results, among other types of server tasks.

[0083] In an implementation, an operation from a UDF (e.g., JAVA-based UDF) can be performed by a user code runtime **424** executing within a sandbox process **420**. In an embodiment, the user code runtime **424** is implemented as a virtual machine, such as a JAVA virtual machine (JVM). Since the user code runtime **424** executes in a separate process relative to the execution node process **410**, there is a lower risk of manipulating the execution node process **410**. Results of performing the operation, among other types of information or messages, can be stored in a log **444** for review and retrieval. In an embodiment, the log **444** can be stored locally in memory at the execution node **302-1**, or at a separate location such as the cloud storage platform **104**.

[0084] Examples of the log **444** can include logging for observability and debuggability. Logging can be automatically configured to observe egress traffic using a logging mechanism with runtime-configurable verbosity levels. For example, use of an event output log or event output helper can allow for passing custom structs from the eBPF program to a performance event ring buffer along with an optional packet sample. In response, the execution platform worker can pull the logs from log **444** or other logs from the buffer and write to execution platform logs, as an example. This channel can be used to log, debug, sample, and/or push notifications for network policy violations and the like. For example, the event output log or helper can be configured to pass the data through a lockless memory mapped per-CPU performance ring buffer, which is significantly faster (e.g., more efficient) than default logging support in eBPF.

[0085] Additional examples of the log **444** or other logs of the cloud data platform **102** can be used to provide clear and actionable feedback necessary for users if their UDF's packet has been blocked. With the logging mechanism, the cloud data platform **102** or component thereof can report details back to the user (e.g., which IP and port has been blocked or violated the account policy). Additionally, when an unauthorized DNS request has been blocked, the eBPF program can intercept the packet and report back which hostname it tried to access and enter such information into the log **444**, which is valuable for helping customers to troubleshoot and debug their UDx.

[0086] Moreover, such results can be returned from the user code runtime **424** to the UDF client **412** utilizing a high-performance protocol (e.g., without serialization or deserialization of data, without memory copies; operates on record batches without having to access individual columns, records or cells; utilizes efficient remote procedure call techniques and network protocol(s) for data transfer) for data transfer (e.g., distributed datasets) that further provides authentication and encryption of the data transfer. In an embodiment, the UDF client **412** uses a data transport mechanism that supports a network transfer of columnar data between the user code runtime **424** (and vice-versa).

[0087] Security manager **422**, in an example, can prevent completion of an operation from a given UDF by throwing an exception (e.g., if the operation is not permitted), or returns (e.g., doing nothing) if the operation is permitted. In an implementation, the security manager **422** is implemented as a JAVA security manager object that allows applications to implement a security policy such as a security manager policy **442**, and enables an application to determine, before performing a possibly unsafe or sensitive operation, what the operation is and whether it is being attempted in a security context that allows the operation to be performed. The security manager policy **442** can be implemented as a file with permissions that the user code runtime **424** is granted. The application (e.g., UDF executed by the user code runtime **424**) therefore can allow or disallow the operation based at least in part on the security policy.

[0088] Sandbox process **420**, in an embodiment, is a sub-process (or separate process) from the execution node process **410**. A sub-process, in an embodiment, refers to a child process of a given parent process (e.g., in this example, the execution node process **410**). The sandbox process **420**, in an example, is a program that reduces the risk of security breaches by restricting the running environment of untrusted applications using security mechanisms such as namespaces and secure computing modes (e.g., using a system call filter to an executing process and all its descendants, thus reducing the attack surface of the kernel of a given operating system). Moreover, in an example, the sandbox process **420** is a lightweight process in comparison to the execution node process **410** and is optimized (e.g., closely coupled to security mechanisms of a given operating system kernel) to process a database query in a secure manner within the sandbox environment.

[0089] For example, the instance of a computer program can be instantiated by the execution platform **110**. For example, the execution node **302-1** can be configured for instantiating a user code runtime to execute the code of the UDF and/or to create a runtime environment that allows the user's code to be executed. The user code runtime can include an access control process including an access control list, where the access control list includes authorized hosts and access usage rights or other types of allowlists and/or blocklists with access control information. Instantiating a sandbox process can determine whether the UDF is permitted and instantiating the user code runtime as a child process of the sandbox process, the sandbox process configured to execute the at least one operation in a sandbox environment.

[0090] In an embodiment, the sandbox process **420** can utilize a virtual network connection in order to communicate with other components within the subject system. A specific set of rules can be configured for the virtual network connection with respect to other components of the subject system. For example, such rules for the virtual network connection can be configured for a particular UDF to restrict the locations (e.g., particular sites on the Internet or components that the UDF can communicate) that are accessible by operations performed by the UDF. Thus, in this example, the UDF can be denied access to particular network locations or sites on the Internet.

[0091] The sandbox process **420** can be understood as providing a constrained computing environment for a process (or processes) within the sandbox, where these constrained processes can be controlled and restricted to limit access to certain computing resources.

[0092] Examples of security mechanisms can include the implementation of namespaces in which each respective group of processes executing within the sandbox environment has access to respective computing resources (e.g., process IDs, hostnames, user IDs, file names, names associated with network access, inter-process communication, and the like) that are not accessible to another group of processes (which may have access to a different group of resources not accessible by the former group of processes), other container implementations, and the like. By having the sandbox process **420** execute as a sub-process to the execution node process **410**, in some embodiments, latency in processing a given database query can be substantially reduced (e.g., a reduction in latency by a factor of 10× in some instances) in comparison with other techniques that may utilize a virtual machine solution by itself.

[0093] As further illustrated, the sandbox process **420** can utilize a sandbox policy **440** to enforce a

given security policy. The sandbox policy **440** can be a file with information related to a configuration of the sandbox process **420** and details regarding restrictions, if any, and permissions for accessing and utilizing system resources. Example restrictions can include restrictions to network access, or file system access (e.g., remapping file system to place files in different locations that may not be accessible, other files can be mounted in different locations, and the like). The sandbox process **420** restricts the memory and processor (e.g., CPU) usage of the user code runtime **424**, ensuring that other operations on the same execution node can execute without running out of resources.

[0094] As mentioned above, the sandbox process **420** is a sub-process (or separate process) from the execution node process **410**, which in practice means that the sandbox process **420** resides in a separate memory space than the execution node process **410**. In an occurrence of a security breach in connection with the sandbox process **420** (e.g., by errant or malicious code from a given UDF), if arbitrary memory is accessed by a malicious actor, the data or information stored by the execution node process is protected.

[0095] Although the above discussion of FIG. **4** describes components that are implemented using JAVA (e.g., an object-oriented programming language), it is appreciated that the other programming languages (e.g., interpreted programming languages) are supported by the computing environment **400**. In an embodiment, PYTHON is supported for implementing and executing UDFs in the computing environment **400**. In this example, the user code runtime **424** can be replaced with a PYTHON interpreter for executing operations from UDFs (e.g., written in PYTHON) within the sandbox process **420**.

[0096] FIG. **5** shows a block diagram illustrating an example embodiment of a User-Defined Function (UDx) external access system architecture **500**, according to some example embodiments.

[0097] The system architecture **500** as described and depicted in connection with FIG. **5** shows a high-level diagram of the components involved in securing egress to target hosts (the egress path is described in more detail below). In the illustrated example, a key **501** provides for solid arrows denoting a control flow (e.g., control path) and dotted arrows denoting a data flow (e.g., data path) amongst and between corresponding components of the UDx external access architecture.

Specifically, the example system architecture **500** illustrates a control flow and a data flow between a compute service manager **108**, an execution platform **110**, a proxy service **115**, and target hosts **550**.

[0098] Below is one example that defines a UDx that wants to access maps.googleapis.com (e.g., external access) for geocoding, where the SQL syntax is just one possible example of applicable code:

```
TABLE-US-00001 :::::CODE::::: CREATE EXTERNAL ACCESS INTEGRATION
GOOGLE_MAPS_APIS EGRESS_RULES = ('maps.googleapis.com',
'maps.google.com'); CREATE FUNCTION GEOCODE(location varchar) ...
EXTERNAL_ACCESS_INTEGRATIONS = ( GOOGLE_MAPS_APIS ) HANDLER =
'geocode' AS $$ import requests def geocode(location):
return requests.get ("https://maps.googleapis.com
...".format(location)).json( ) $$ :::::CODE:::::
```

[0099] When the above example code (e.g., UDx) is used in a query, one or more of the following actions are performed.

[0100] The UDx **518** can be authorized to access maps.googleapis.com by the external access integration **519**. The external access integration **519** can be an existing or a new type of integration object that is used to control a user's ability to reach out to a particular host or set of hosts (e.g., target hosts **550**). For example, the external access integration **519** can be used to control the locations that can be connected to by sandboxed functions (e.g., UDx **518**). In some example embodiments, once usage on the integration is granted to a user, how the user connects to a host or what they do on the host can be unconstrained. As with other integrations, they may only be

creatable by the account administrator, who can then grant USAGE on the integration to grant privileges to other users (e.g., roles on the account). According to some examples, the integration can be used for access control; for example, users with USAGE on the integration can build functions that can access hosts that match the provided allowlist(s), and a function whose owner has USAGE on the integration is then permitted to access those hosts.

[0101] The function (e.g., UDx **518**) can specify a concrete endpoint or set of endpoints to which the function is going to connect. Thereby providing the system with the specific host to open up for queries using the function. Once this is configured, the client can use standard networking calls provided that the target hosts **550** are configured. As the example illustrates, with an appropriate integration, a user could access any endpoint they wanted (assuming they have the appropriate credentials). In order to promote good practices for handling credentials, the cloud data platform **102** or component thereof can employ a secrets object (e.g., the account administrator can provision a new credential including a secret). A user with READ privileges on the secret credential could then use it to create a function. Unlike with USAGE on a secret, READ allows reading the contents of the secret from inside the UDFs.

[0102] Once the UDx **518** is authorized, the compute service manager **108** or a component thereof can generate an egress policy and sign the egress policy for this UDx (e.g., signed egress policy **521**). Once signed, the compute service manager **108** can deliver the signed egress policy **521** or cause the egress policy to be delivered to the execution platform **110** with a query execution plan. The egress policy contains one or more lists of allowed IP addresses (e.g., for maps.googleapis.com in this example case).

[0103] An execution platform worker **502**, such as execution node **302-1**, will do two things before the UDx starts. First, the execution platform worker **502** transmits the policy or causes the policy to be transmitted to an egress proxy, such as proxy agent **505**. Second, the execution platform worker **502** initiates or sets up a secure egress path (described in more detail below) for the sandbox process where the UDx lives. Once the UDx **518** starts to run, the egress traffic from the UDx **518** will be strictly forwarded through the secure egress path to the proxy agent **505** (e.g., egress proxy) and eventually be allowed to pass through an egress gateway **541** to the target public hosts (e.g., target hosts **550**).

[0104] According to some examples, the system architecture **500** further illustrates a secure egress path **531**, which can be set up in different ways, such as by performing egress control on network Layers 3, 4, and/or 7 and/or establishing a secure egress path using an overlay network to achieve strict egress control and isolation.

[0105] In considering Layer 3 (L3) (e.g., network layer), Layer 4 (L4) (e.g., transport layer), and Layer 7 (L7) (e.g., application layer) of the Open Systems Interconnection (OSI) model of how data is transmitted over a network, customers can be provided with the ability to control egress at the network, transport, and/or application layers. For example, Layer 3 and Layer 4 egress control can provide the customer with the ability to restrict the outgoing traffic to specific IP address(es) and port(s), whereas Layer 7 egress control can provide the customer with the ability to control more restrictive filtering based on, for example, destination and requests content (e.g., a Uniform Resource Identifier (URI) path, query parameters, headers, etc.). Layer 7 controls can provide additional observability, such as providing the customer and/or the cloud data platform with the ability to tell which particular Representational State Transfer (REST) APIs were accessed, whereas Layer 3 and Layer 4 controls can provide logs about IP addresses and ports. For example, a Layer 3/Layer 4 egress control (e.g., L3/L4 egress rule) may include only allowing traffic to maps.googleapis.com at port **443**; whereas a Layer 7 egress control (e.g., L7 egress rule) may include only allowing HTTP(S) GET requests to <https://maps.googleapis.com/maps/api/geocode>. As the example shows, using L3/L4 egress control can provide coarse-grained control for some L7 use cases, by enabling customers to state that the traffic should only go to maps.googleapis.com even if the customer cannot state which REST APIs are allowed.

[0106] In addition to L3/L4/L7 egress control, examples can provide for a secure egress path from the execution platform **110** to an egress proxy or proxy server, such as the proxy service **115**. By providing a secure egress path, support for standard libraries for network access can be provided. Along with being transparent to the user code, a secure egress path **531** for UDx external access can also include strict egress isolation, strict egress control, and/or egress proxy enforcement. For example, the egress traffic from a sandbox **520**, which can be the same as or similar to the sandbox process **420** as described and depicted in connection with FIG. 4, can be isolated from other sandboxes and existing execution platform workloads, while unauthorized egress traffic can be blocked and/or reported. In some examples, the egress proxy (e.g., proxy service **115**) is the only gateway for any egress traffic from the execution platform **110**, and the compute service manager **108** or component thereof is configured to enforce the egress proxy policy. In other words, the secure egress path is configured to redirect the traffic to the egress proxy as the egress gateway **541** to the Internet.

[0107] The above example requirements cannot be satisfied by allowing the sandboxed code to directly access the execution platform's network (e.g., eth0 device), and there are challenges to implement egress control directly on eth0 (e.g., a network interface device name to represent the first Ethernet network interface) without affecting other execution platform workloads. Therefore, an overlay network solution is proposed to establish this secure egress path (where the overlay network is described and depicted in detail in connection with FIG. 7).

[0108] FIG. 6 is a block diagram illustrating an external function access system architecture **600** including a control flow for external access, in accordance with example embodiments. Example embodiments of the control flow include similar elements and components as depicted and described in connection with FIG. 5. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are redundant to those described and depicted in connection with FIG. 5 have been omitted from FIG. 6. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the system architecture **600** to facilitate additional functionality that is not specifically described herein.

[0109] The control flow starts at step **1 601**, starting from the execution platform worker **502**, the execution platform worker registers a policy at the platform agent **109**. The platform agent **109** writes the policy and/or proxy information into one or more Berkeley Packet Filter (BPF) maps and/or extended BPF (eBPF) maps and creates (e.g., prepares) one or more network devices. The eBPF is a technology that allows developers to create and run custom programs that can be inserted into a Linux kernel to intercept and analyze system events. The BPF is used for packet filtering and network analysis, and the eBPF is used, for example, for a wide range of system monitoring and tracing scenarios. For example, using eBPF, developers can create programs that intercept system calls, kernel events, and/or network traffic, and modify and analyze the data in real-time or near real-time.

[0110] Returning to the execution platform worker **502**, the control flow continues at step **2 602**, where the execution platform worker **502** registers the policy at the proxy agent **505** of the proxy service **115**. The proxy agent **505** writes the policy into the BPF maps.

[0111] From the execution platform worker **502**, the control flow continues at step **3 603**, where the execution platform worker **502** starts the sandbox process **420** subprocess. The sandbox process **420** subprocess can configure a Virtual Ethernet (veth) pair between the execution platform **110** and the sandbox process **420** (e.g., UDx executing, agents collecting events/logs, etc.).

[0112] From the sandbox process **420**, the control flow continues at step **4 604**, where the sandbox process terminates. The control flow continues at step **5 605**, where the execution platform worker **502** unregisters from the proxy agent **505** and collects any events (e.g., violations). At step **5 605**, the proxy agent **505** can clean the policy afterwards in case of an execution platform worker **502** crash or incident. Last in the control flow at step **6 606**, the platform agent **109** similarly unregisters

at the platform agent and collects events. At step **606**, the platform agent **109** can clean the policy after, for example, an execution platform worker **502** crash or incident.

[0113] According to some example embodiments, an egress policy can contain a list of rules, where each rule defines an allowed endpoint. An endpoint may include multiple allowed IP addresses with a list of allowed ports.

```
TABLE-US-00002      :::::CODE::::: [  {      “domain”: “www.ex1.com”,      “ips”: [
      “45.76.117.150”      ],      “ports”: [      {      “port”: “443-443”,
      “protocol”: “tcp”      },      {      “port”: “8080-8090”,
      “protocol”: “tcp”      }      ]      },      {      “domain”: “www.ex2.com”      “ips”: [
      “45.76.117.151”,      “45.76.117.153”      ],      “ports”: [      {
      “port”: “443-443”,      “protocol”: “tcp”      }      ]      }, ] :::::CODE:::::
```

[0114] In some examples, the policy setup and cleanup, such as the execution platform setup of the policy for a veth ingress, can occur before launching the sandbox process **420**. For example, the execution platform **110** can extract the egress policy from a token, such as a JSON Web Token (JWT) string and populate the BPF maps, such as a policy map and a query information map. The policy map can include information such as the sandbox identifier, destination IP address, destination port, and the like. The query information map can include information such as the sandbox IP address, the query identifier, the sandbox identifier, and the like.

[0115] The policy map, query information map, BPF maps or the like, can provide lookups to determine an allowed port. According to some examples, a policy key can be used in access control systems to determine whether a user or system is allowed to access a particular resource based on the policies associated with that resource. For example, a policy key can be a unique identifier used to associate a particular policy with a resource. In one example, a user can use a port in the policy key (e.g., flatten port range to each port in the policy key) in order to provide granular control over resource access based on specific policies. In additional examples, the cloud data platform can add additional map(s) to handle large port range lookups not easily provided in a policy key lookup.

[0116] In additional examples, a port range in the policy value can be used when the port range cannot be used as a key in the policy key. In such cases, the cloud data platform can introduce another inner map to store all policy rules. For example, the cloud data platform can use the sandbox identifier and destination IP address to find the inner map, and then lookup all rules inside the map. The inner map can contain the rules including port range, protocol, etc.

[0117] In additional examples, a source identity in a policy key can be used to specify a security policy for a particular network flow. Where the source identity is a component of the policy key that identifies the originator of the network flow, such as an IP address, MAC address, or other identifier. The source identity can be used to match the network flow to the appropriate security policy. According to one example embodiment, a sandbox identifier (sandboxId) can be used as the identity, with sandboxId as the source identity, the cloud data platform can unify the policy map definition in both the execution platform **110** and the proxy service **115**. In an alternative example embodiment, the sandbox IP address can be used as the identity.

[0118] In additional example embodiments, a policy validation is the process of verifying that a policy or set of policies comply with certain criteria or requirements, which can be predefined rules, industry standards, best practices, or the like. Where the goal of policy validation is to ensure that policies are effective, consistent, and enforceable. Once a new packet reaches veth ingress at the execution platform **110** host, the code first looks up a QueryInfoMap to find the sandboxId using the source IP address, and then looks up a PolicyMap using the sandboxId to find the policy and validate the destination. If successful, the packet will be forwarded to a Generic Network Virtualization Encapsulation (GENEVE) tunnel with the queryId and sandboxId embedded into the GENEVE header. Otherwise, report the policy denied error message to eventMap. In another example, once a new packet reaches a GENEVE ingress, the code extracts the sandboxId from the header and looks up the policy to validate the destination. If successful, the packet is forwarded to

eth0. Otherwise, the packet is dropped and/or this sandbox escape error is reported to the eventMap.

[0119] In further example embodiments, a UDx should be able to get the IP address of the destinations from hosts pre-configured by the sandbox process **420**. However, an unexpected DNS request may happen if the UDx tries to reach an unexpected destination. Before terminating the sandbox process **420**, the execution platform **110** or component thereof needs to perform an execution platform policy cleanup to delete policy entries from PolicyMap and QueryInfoMap using the sandbox srclp as the key. In addition to the execution platform policy cleanup, the proxy service **115** can also perform a policy setup and/or cleanup. The proxy service **115** policy cleanup can be performed before terminating the sandbox process **420**. For example, the execution platform **110** can send a policy unregister request to the proxy agent **505**, and the proxy agent **505** can clean the policy in the policy map. In one example, a periodical background job can also be implemented to clean any out-of-date policies and/or report events. Additional examples can include a proxy policy refresh that happens on the execution platform **110**, which should then register the refreshed policy again to the proxy service **115**.

[0120] For example, the execution platform **110** registers the policy to a proxy, where the policy register request includes a signed JWT string and the queryInfo (queryId, sandboxId). The proxy agent **505** will need to verify the JWT using a token and consume the extracted policy. Afterwards, the proxy agent **505** will populate the PolicyMap, in a similar manner as described in as the one in the execution platform.

[0121] To enable external network access from a sandbox process in the execution platform **110**, customized packet redirection and policy enforcement logic is implemented in eBPF program code. This eBPF code is compiled into binary files and attached to the ingress/egress side of different network interfaces so that each packet incoming to or outgoing from those network interfaces can be validated by the eBPF code. The cloud data platform or component thereof ensures that packets routed through the egress proxy (e.g., proxy service **115**) and that the destination IP address is allowed as per the egress policies defined by the customer. This requires the population of several BPF maps with information about allowed destinations.

[0122] The eBPF code also needs to emit logging events or metrics to BPF maps for debugging and observability purposes. For example, a userspace program (e.g., an agent) needs to poll the events from the BPF map and emit logs and/or statistics to. Additionally, the agent needs to perform some cleanup work against the BPF map to maintain the runtime state. The userspace program or agent can be a new binary due to the way Linux Capabilities work. A binary needs CAP_BPF capability to be able to manage BPF maps. The capabilities are generally granted by Red Hat Package Manager (RPM) during installation. If a binary with capabilities is copied to a different location, the capabilities may be lost. If capabilities are assigned to the execution platform worker **502** process during execution platform RPM installation, the capabilities will stay with only that version of the execution platform binary. If the capabilities are lost due to copying, a TC_WRAPPER binary is created and installed via the RPM.

[0123] In addition, eBPF is a powerful kernel feature that can do more than just packet processing. For example, some other useful features of eBPF include system call tracing and security enforcement via Linux Security Modules (LSM) BPF that can unify all these use cases and leverage a single agent process that handles all eBPF related operations. The eBPF agent can additionally be used for malicious activity monitoring. The platform agent **109** can be the eBPF agent in the execution platform **110**, and the agent binary can be granted the required capabilities. The platform agent **109** is further configured to set up network devices for external access so that all the networking setup is handled by a single executable.

[0124] In additional examples, the proxy agent **505** can be an eBPF agent in the proxy service **115** and ensure policy agent enforcement responsibilities. For example, the proxy agent **505** can poll logging and metrics from the BPF map and emit polling results to system log data or local servers.

The proxy agent **505** can include policy cleanup actions, such as initiating and monitoring cleanup of connection tracking BPF map entries (e.g., in case of execution platform worker **502** failure), cleanup Source Network Address Translation (SNAT) BPF map entries, cleanup network devices used by the UDx external access, and the like. The proxy agent **505** can further initialize one or more GENEVE devices, including, for example, configuring GENEVE device IP forwarding and performing packet filtering to allow packets to be forwarded between network devices (e.g., veth, geneve0, eth0, etc.). The proxy agent **505** can initialize veth pair creation for UDx with external access and monitor or update BPF map values and data to be consumed by kernel BPF code.

[0125] In additional example embodiments, proxy errors on connection timeout at the execution platform **110** are described. Typically, the client side only sees connection timeout error, which gives very little information and guidance for how to move on. However, if the error occurs because of a proxy error, the proxy agent **505** or the execution platform **110** should be able to provide customers a better message to help them recover from such errors. For example, if the proxy service **115** is down, the proxy agent **505** can tell the customer that the “internal service is temporarily unavailable, retry later . . .” etc. In order to achieve this, the execution platform **110** needs to know if there is any error for the query on the proxy side and what are the errors that have occurred. For example, the execution platform **110** can send an unregister message to the proxy service **115** side after the execution is finished and before the sandbox is destroyed to trigger a policy map cleanup. The unregister message can include a query identifier and a sandbox identifier as part of the message fields, and the proxy agent **505** can maintain an error cache that caches all error messages detected for a query based on the messages polled from the event map. The proxy agent **505** can attempt to lookup all related error events, metrics, and/or logs using the queryId and sandboxId once an unregister request is received and surface them back in the response if there are any.

[0126] In additional example embodiments, a push error message is sent back from the proxy service **115** to the execution platform **110** as soon as it is detected. However, the message exchange happens after the UDx execution finishes, and in the case of a timeout, the customer needs to wait until the timeout to know the query has failed even if it is due to the cloud data platform internal service error. Example embodiments push the proxy error back to the execution platform **110** as soon as the error occurs, so that the execution platform **110** can fail the query immediately with proper errors. In one example, the proxy service **115** pushes the error event back at the application layer through a remote procedure call (e.g., gRPC) once the corresponding event is extracted.

[0127] In another example, the proxy service **115** redirects the package back to the execution platform **110** at the BPF layer and adds the error message in the GENEVE header. In this approach, the error message is added in the GENEVE header, and it reroutes the package back to the execution platform **110** directly. For example, the execution platform **110** can parse the GENEVE header at the BPF layer, and publish it as an error event, which will then be detected at the application layer. The error is served back to the customer from the execution platform **110** when the execution platform **110** is able to get the set of proxy errors, where the following table describes the proposed message mapping. Specifically, Table 1 is an example of proposed message mapping, according to the example above displaying an error to the customer and an error the execution platform **110** receives from the proxy service **115**, as follows:

TABLE-US-00003 TABLE 1 Error Execution Platform Error to Customer Sees from Proxy

“Internal service temporarily	1) Proxy down unavailable, please retry later.”	2) Connection dropped due to table is full	“Internal service error, please All other proxy error contact cloud data platform.”
“Connection timeout, please	No proxy error associated check if the destination is valid, with the query can be found. or contact cloud data platform.”		

[0128] FIG. 7 is a block diagram illustrating a system architecture **700** depicting an external access overlay network **701**, in accordance with example embodiments. As noted above in connection with FIG. 5, the secure egress path for UDx external access must meet strict requirements including

strict egress isolation, strict egress control, and/or egress proxy enforcement. Such requirements cannot be satisfied by allowing the sandboxed code (e.g., UDF1 **711**) to directly access the execution platform's **110** network (e.g., eth0 device **724**) without affecting other execution platform workloads. Therefore, the overlay network solution of FIG. 7 is proposed to establish this secure egress path. The overlay network solution of FIG. 7 can be used to establish a secure egress path and can include the overlay network that is a virtual network of nodes and logical links built on top of an existing physical or virtual network. The overlay network **701** creates a new layer where traffic can be programmatically directed through new virtual network routes or paths (e.g., overlay **727**) instead of requiring physical links (e.g., physical **725**). It also enables the cloud data platform or components thereof to define and manage traffic flows irrespective of the underlying physical infrastructure.

[0129] FIG. 7 demonstrates the secure egress path from a sandbox (e.g., sandbox process **420**) to an egress proxy (e.g., proxy service **115**), which applies the overlay network to achieve egress control and isolation. For example, a packet sent on the secure egress path goes through three logical hops, where the hops (e.g., tiers) are highlighted by dashed lines as tier 1 (T1) **710**, tier 2 (T2) **720**, and tier 3 (T3) **730**.

[0130] At tier 1 (T1) **710**, the first step in enabling external access from the sandbox process **420** is to put an ethernet device in its namespace. As described above in connection with FIG. 6, example embodiments use Virtual Ethernet (veth) pairs for putting the ethernet device in its namespace. Virtual ethernet (veth) pairs allow communication between network namespaces. A veth pair is a pair of virtual network interfaces that are connected together in a manner that any traffic (e.g., data) sent through one interface of the veth pair is received by the other interface of the veth pair, in order to provide a way for isolated environments to communicate with each other and/or with a host system. For example, each veth pair can consist of two virtual network interfaces, which can be referred to as a “parent interface” and a “child interface.” The parent interface can be attached to the host system and/or another network interface, whereas the child interface can be attached to a virtual machine or container. For example, when traffic is sent through the child interface of the veth pair, the traffic is transmitted to the parent interface of the veth pair, and then the traffic is transmitted to the destination network.

[0131] For example, at tier 1 (T1) **710**, at one end of the veth pair, veth0 **712** is put in the sandbox's namespace and at the other end, veth1 **713** is put in the execution platform's **110** default namespace. Any packet transmitted on veth0 **712** is immediately received on veth1 **713** and vice-versa. For packets arriving on veth1 **713**, the system ensures that their destination is allowed by the egress policy by attaching one or more extended Berkeley Packet Filter (eBPF) hooks **714** to the veth1 **713** device. eBPF hooks can include a type of eBPF program used to intercept and/or modify system calls and/or other kernel events by allowing developers to monitor and control system behavior in real-time or near real-time.

[0132] For example, the eBPF can include an in-kernel virtual machine that allows code execution in the kernel space. It can be used to complement or replace kernel packet processing among other things. The eBPF hook **714** on veth1 **713** guarantees that only packets with allowed destinations can be forwarded to the next hop (e.g., geneve0 device **721**). Otherwise, packets are dropped, and the violation is reported. This also guarantees that no packet can be sent to other sandboxes or execution platform **110** instances.

[0133] At tier 2 (T2) **720**, after getting the packet out of the sandbox process **420** and validating that its destination is allowed, the system needs to send the packet to the egress proxy, proxy service **115**. GENEVE tunneling is used for T2 **720**. In virtualized environments, such as the virtualized environment of FIG. 7, GENEVE (Generic Network Virtualization Encapsulation) can include a network tunneling protocol used for transmitting network traffic between or among virtual machines over an IP network. A GENEVE device can include a virtual network device that implements the GENEVE protocol and is used to provide overlay networks that span multiple

physical hosts or clusters. GENEVE encapsulates packets within UDP packets, for example, and uses network virtualization headers to allow multiple virtual networks to be transmitted over the same physical network infrastructure (e.g., physical **725**). GENEVE encapsulation is further described and depicted in connection with FIG. **8**. While GENEVE devices and packetization are used in the example embodiment of FIG. **7**, it will be understood to those skilled in the art, that embodiments and examples of the inventive subject matter can be practiced with other network tunneling protocols. At T2 **720**, the system uses GENEVE to embed a policy identifier for each packet so the proxy service **115** can use that information to perform further granular egress control. [0134] At tier 3 (T3) **730**, the egress proxy, proxy service **115**, serves as a security barrier and the only egress gateway (such as egress gateway **541**) from the execution platform **110**. When a packet is received on the GENEVE device geneve0 **731** in the proxy service **115**, the packet is decapsulated **728** to retrieve the original packet **732** and the policy ID. The policy ID is used to look up the egress policy and validate that the original destination (e.g., maps.google.com's IP address) is allowed. This validation generates a verified packet **733** and is effectively identical to what is performed at veth1 **713** in the execution platform **110**. This second check is used to prevent a compromised execution platform **110** (e.g., "zero trust execution platform") from exfiltrating data. Any violations will be reported to security. If the packet passes all the checks, the system changes the packet's source address from the sandbox IP address to an egress proxy IP address using Source Network Address Translation (SNAT). This needs to be done because the sandbox IP address is a private IP address and is not generally routable, and a routable IP address is needed for the response packet to be transmitted to the Internet **781** using UDx external access. The system maintains a connection map so that it can correctly identify the target execution platform and sandbox instance for the response packet. These 3 parts, Tiers 1-3 together, ensure that the traffic from the sandbox process **420** is well isolated and can only go to the allowed destinations of the Internet **781**.

[0135] In additional example embodiments, statically resolving hostnames in the compute service manager **108** is explored. The L3/L4 based egress control discussed above is based on IP addresses. However, almost every Internet connection starts with a DNS, so it is critical to secure domain name resolution from a UDx, especially to avoid data exfiltration over DNS. In such additional examples, the hostname resolution is performed in the compute service manager **108** at the start of the query and the mapping is sent to the execution platform **110** in an egress policy document. The eBPF based network can enforce the IP restrictions based on this static mapping. The mapping can also be used to populate the hosts file inside the sandbox. For allowed destinations, hosts can be used for name resolution. If DNS resolution is attempted for a disallowed name, eBPF code can drop the DNS packet and report user error. This static approach can be used in most of the use cases explored herein as the jobs are usually short-lived. However, for services using DNS for failover or load balancing (e.g., round robin), this approach may not be optimal, and a mitigation is to select an appropriate Time To Live (TTL) field in the IP header of the packet to indicate the maximum number of hops that a packet can travel before it is discarded by a gateway in order to refresh the IP list.

[0136] Additional example embodiments of the overlay network **701** can include allowing dynamic DNS resolution at runtime from the UDx and supporting a custom DNS resolver. For example, an account administrator can specify the allowed DNS resolvers in the allowlist (e.g., allowed host list) such that the DNS requests pass the egress control. In another example, the system can implement a dedicated DNS resolver at a secure location (e.g., proxy service **115**), and proxy all hostname lookups from one or more UDx. For example, the system can provide a local DNS proxy with networking and security capabilities for containerized applications using eBPF technology to enable fast and secure communication between containers.

[0137] While the example of FIG. **7** describes the use of the overlay network in reference to external access, it will be understood by those having ordinary skill in the art that the overlay

network described herein can further be used internally by components of the cloud data platform. For example, this system, in addition to being used for external access, can be used to allow for instances of the same workload or related workloads to communicate with each other over example embodiments of an overlay network. For example, one or more sandbox instances of the same job can be configured to communicate with each other over the overlay network described herein. [0138] FIG. 8 is a block diagram illustrating GENEVE encapsulation **800** as used in FIG. 7, in accordance with example embodiments.

[0139] GENEVE encapsulation protocol is an encapsulation protocol implemented in Linux to create a tunnel between two network devices. As noted in FIG. 7, in virtualized environments, GENEVE can include a network tunneling protocol used for transmitting network traffic between or among virtual machines over an IP network. GENEVE is implemented herein as it uses User Datagram Protocol (UDP) as the transport protocol and encapsulates the original packet as the payload, thereby allowing the cloud data platform to set a different destination while providing the flexibility to customize its metadata.

[0140] A packet encapsulated in the GENEVE format as described in FIG. 8 includes a tunnel header encapsulated in UDP over IP. The GENEVE format includes an outer Media Access Control (MAC) address **801** providing the source or destination address in the outermost layer of a network packet header. In GENEVE format, the outer IP **802** refers to the IP header of the outermost layer of the encapsulated packet including a source execution platform IP address and a destination proxy IP address. The outer UDP header **803** refers to the UDP header of the outermost layer of the encapsulated packet, which can include information about the source and destination UDP ports, among other information.

[0141] The GENEVE header **804** includes several fields including, for example, the version of the GENEVE protocol, the variable-length options, the encapsulated protocol, the virtual network identifier (VNI), among other information. The GENEVE options **805** are used to provide additional information about the encapsulated packet, such as, for example, a policy identifier (ID). The GENEVE inner **806** refers to the original packet that is being encapsulated and transmitted over the virtual network. In the examples of FIGS. 7 and 8, the inner includes a source address including the sandbox IP address and a destination address, such as the maps.google.com IP address. The purpose of encapsulating the inner packet within the GENEVE header is to allow the virtual network to maintain the same level of isolation, security, control, and performance as a physical network. The inner **806** can be decrypted and processed by the receiving device once it has been decapsulated from the GENEVE header as described above in connection with FIG. 7.

[0142] The inner transmission control protocol (TCP) header **807** is the header of the original TCP packet that is being encapsulated and transmitted over the virtual network. The inner TCP payload **808** is the actual data that is being transmitted within the TCP packet that is being encapsulated and transmitted over the virtual network. When the TCP packet is encapsulated within a GENEVE header, the entire packet, including the TCP header and payload, becomes the inner packet. Last in the GENEVE encapsulation format is the frame check sequence (FCS) **809**, which refers to the cyclic redundancy check value that is added to the end of the packet to ensure integrity of the packet during transmission.

[0143] FIG. 9 is a block diagram of an environment **900** illustrating a high-level overview graph depicting the flow of metrics and events from the BPF layer to the customer including proper processing, in accordance with example embodiments. In some examples, the environment **900** can include some or all of the same components as in the system architecture **700** depicted and described in connection with FIG. 7. FIG. 9 is similar in many respects to FIG. 7, sharing certain similar features and components; for brevity, only specific elements are detailed and described with reference to FIG. 9. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from FIG. 9. However, a skilled artisan will readily recognize that

various additional functional components may be included as part of the environment **900** to facilitate additional functionality that is not specifically described herein.

[0144] In the execution platform **110**, the sandbox with policy A **920**, which can be similar to sandbox process **420** as described and depicted in connection with FIGS. **4** and **7**, includes a UDx and veth0 as described and depicted in connection with FIG. **7**.

[0145] In the proxy service **115**, the agent **905**, such as proxy agent **505** as described and depicted in connection with FIG. **5**, can be configured to poll an event map, such as BPF event map **916**, to extract events and metrics, perform proper processing, and log proper messages for debugging purposes and/or publish the metric as the cloud data platform metrics. The polls **917** can be forwarded to the agent **905** from the BPF event map **916**. The agent **905** transmits to the execution platform worker **502**, an error list **998** that can include, for example, unregistered responses.

[0146] In response to receiving the error list **998**, the execution platform worker **502** can provide one or more error messages **999** to a client device of a user of the cloud data platform. In addition, the execution platform worker **502** can send a destroy message **987** to sandbox with Policy A **920** to initiate a cleanup of any policies and to destroy the corresponding sandbox. Once the UDx computation is complete, the execution platform **110** work will destroy the corresponding sandbox. Meanwhile, the execution platform worker **502** transmits an unregister message **989** that can include, for example, a query identifier and a sandbox identifier to the agent **905** (e.g., to unregister the policy with query_id and sandbox_id).

[0147] Once the agent **905** receives the unregister message **989**, the agent **905** will look-up the error caching to extract all error metrics associated with the query and surface them back to the execution platform **110** in the responses. Proper mapping cleanup will also be performed. The agent **905** publishes metrics **936**. Once the execution platform **110** receives the response, it will perform an analysis based on all the error results it has (e.g., including both execution platform errors and/or proxy service errors), and determine what is the proper error the execution platform **110** should surface back to the customer in user error message **999**.

[0148] The proxy service **115** can generate and log useful metrics and events that can be used to improve the developer, customer, and/or use experience. The BPF event map **916** can be employed to pass events and/or metrics from the BPF layer to the application layer, and similarly be adapted for the proxy service. For example, a BPF metric can be written in a signal map and published with the cloud data platform metrics. Table 2 (below) shows a few examples of debugging and error metrics according to example embodiments; however, additional metrics and/or events can be captured and logged according to the present disclosure. Table 2 is an example of different categories of metrics, according to some examples, including a combination of BPF metrics, startup metrics, runtime metrics, termination metrics, and the like, as follows:

TABLE-US-00004

Metric Name	Metric Description	Signal
NatFillUp	Potential congestion on NAT table. Cloud data platform should be alerted.	
CTFillUp	Potential congestion on connection track (CT) table. Cloud data platform should be alerted.	
ExternalFunc-CreateTime	The duration between creating the UDF server subprocess. Helps cloud data platform and/or user to understand the cost of preparing the UDF server.	
ExternalFunc-StartupTime	The duration between starting the UDF server subprocess and the last client has connected.	
FuncArrow-SerializeTime	The total duration for arrow serialization and deserialization at the execution platform.	
FuncShut-downTime	The total duration to terminate the sandbox subprocess.	

[0149] Table 2 includes example metric names and descriptions that can help the cloud data platform and/or the user identify, monitor, and respond to performance management issues in production. For example, providing metrics for understanding the sandbox overhead, such as the startup latency for different components inside the sandbox or understanding runtime overhead.

[0150] The rest of the events can be written in one or more event maps, such as BPF event map **916**, where the messages written in the event map can contain fields in common. Metrics, events,

and other logs can be added, changed, and/or deleted based on the user preferences and/or cloud data platform evolutions. According to some examples, the eBPF code writes metric and/or event messages triggered by error or debugging/monitoring purposes to the BPF event map. Once the metric and events are published (e.g., publish metric **936**), one of the most important things is how to associate each metric/event with the query that initiates the connection.

[0151] According to one example, an identifier for associating the metric and/or event with the query is provided using the query ID and the sandbox ID. With this option, an identifier constructed by using <query_id><sandbox_id> is used, where this identifier can be globally unique, and can be directly associated with the query and policy used by the sandbox with Policy A **920**. However, the query ID and the sandbox ID are identifiers that the execution platform **110** knows, and the information that can be extracted at the proxy side can include a source IP address (e.g., the execution platform IP), a source port, a destination IP address, and a destination port. In order to associate all of that information together, extra processes may need to be performed.

[0152] FIG. **10** is a block diagram of an environment **1000** illustrating a high-level overview depicting the flow of metrics and events from eBPF layer to the customer, in accordance with example embodiments. In some examples, the environment **1000** can include some or all of the same components as in the system architecture **700** depicted and described in connection with FIG. **7** and/or the environment **900** depicted and described in connection with FIG. **9**. FIG. **10** is similar in many respects to FIGS. **7** and **9**, sharing certain similar features and components; for brevity, only specific elements are detailed and described with reference to FIG. **10**. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from FIG. **10**. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the environment **1000** to facilitate additional functionality that is not specifically described herein.

[0153] FIG. **10** describes how the traffic would flow when a package is sent from the execution platform **110** to the Internet (a similar flow would be applied for packages coming back). As shown in FIG. **10**, the execution platform **110** and the proxy service **115** can both maintain an ID map **1006**, **1016** using, for example, BPF maps.

[0154] The ID map **1006** at the execution platform **110** side includes a sandbox IP **1007** as a key, and the identifier (<query_id><sandbox_id>) as a value. The execution platform **110** performs updates **1002** on the ID map **1006** whenever a sandbox is created, destroyed, updated, or the like. The eBPF code **1014** can initiate a lookup **1003** of the identifiers with the sandbox IP **1007** while processing a package and encodes the identifier as part of the GENEVE header in geneve0 **1021**.

[0155] The ID map **1016** at the proxy service **115** side includes (<ep_ip, sport>, <dst_ip, dport>) (e.g., the execution platform IP address, the source port, the destination IP address, and the destination port) as a key, and the identifier (<query_id><sandbox_id>) as value. The eBPF code **1032** at geneve0 **1031** can parse the header to retrieve the identifier and provide updates **1034** to the ID map **1016**. The agent **905** can perform a post-processing lookup **1033** of the ID map **1016** to associate each metric/event with the identity before logging or publishing, if an entry can be found. In some implementations, this is just best effort as some metrics may not be able to find an associated identifier depending on if the map has been updated when the metric/event is published. Example embodiments enable easy association of the debugging metric(s) and/or event(s) with a problematic query based on the query identifier information. The same identifier can be used for the policy map at both the proxy service **115** and/or the execution platform **110**.

[0156] In additional example embodiments, the environment **1000** illustrating a high-level overview depicting the flow of metrics and events uses the execution platform **110** IP address and the sandbox **1020b** (with Policy B) IP address <ep_ip><sandbox_ip> as the identifier for each metric/event. The execution platform **110** IP address can be read directly from the socket buffer <sk_buff> data structure at the proxy service **115** side, but not the sandbox IP address. Therefore,

the sandbox IP address will still need to be passed from the execution platform **110** ID map **1006** to the proxy service **115** via the GENEVE header. The metric only contains IP information, so the IP address still needs to be mapped to the corresponding query manually during debugging.

[0157] According to an implementation of FIG. **10**, the traffic flow from the execution platform **110** to the proxy service **115** can include the following example. Starting at the execution platform **110**, the sandbox **1020b** (with policy B **1008**) includes user-defined function (UDF) **1011**, where the sandbox process configures a virtual ethernet (veth) pair between the sandbox **1020b** (e.g., veth1 **1012**) and the execution platform **110** (e.g., veth2 **1013**). Once a new packet reaches veth2 **1013** ingress at the execution platform **110** host, the eBPF code **1014** performs a lookup **1003** of the ID map **1006** to locate the <sandbox_id> using the source IP address, and then looks up a PolicyMap using the <sandbox_id> to find the associated policy (e.g., Policy B **1008**) and validate the destination.

[0158] If successful, the packet will be forwarded to a GENEVE tunnel (e.g., geneve0 device **1021**) with the <query_id> and the <sandbox_id> embedded into the GENEVE header. The eBPF code **1023** attached to the egress side of the geneve0 device **1021** can perform NAT on the inner packet to change the source IP address stored in the GENEVE header to the eth0 **1024** IP address. Next, the kernel can encapsulate this packet with metadata from the GENEVE tunnel-related data including a remote IP address of the proxy service **115**. The BPF redirect **1035** code is called to redirect the encapsulated packet to the egress side of the geneve0 device **1031** on the proxy service **115**, and the encapsulated packet is transferred to the eth0 **1026** on the proxy service side. The eBPF code **1032** attached to the ingress side of geneve0 **1031** performs a series of tasks (described in detail in connection with FIG. **12**) to decapsulate the packet and route it based on the destination IP address.

[0159] FIG. **11** is a block diagram **1100** illustrating eBPF-based packet routing, of traffic control eBPF processing versus local network stack processing, in accordance with example embodiments. The example depicted in the block diagram **1100** demonstrates how a packet can be routed from the execution platform sandbox (e.g., a separate network namespace) to the public Internet and the packet returned path. Some examples can demonstrate the network interfaces employed on the execution platform **110** and/or egress proxy, such as the proxy service **115**, as well as the eBPF program(s) that are attached to each network interface, where the eBPF is an in-kernel virtual machine that allows code execution in the kernel space. FIG. **11** includes a key **1101** depicting three symbols, including IP tables filtering hook **1111**, an cBPF TC hook **1120**, and an eBPF XDP hook **1130**, each of which is used on the ingress side **1105** and the egress side **1125** of the block diagram **1100**.

[0160] The cBPF programs **1114**, **1124** can run sandboxed programs in an operating system kernel. It is used to extend the capabilities of the kernel safely and efficiently without requiring changing kernel source code or load kernel modules. Especially in the network domain, cBPF program provides programmability for cloud native world to define networks using software. In addition, eBPF also provides both safety and efficiency. It is significantly more secure than loading a kernel module by requiring eBPF programs to pass through a verification process; and the efficiency is guaranteed by a Just in Time (JIT) compiler that ensures native execution speed of eBPF bytecode. Lastly, using eBPF provides the benefit of isolation. Compared to IP tables or IP route, eBPF programs only need to attach to the newly introduced network devices (e.g., veth and GENEVE), and other existing network devices used in the execution platform **110** will not be affected.

[0161] As described below in connection with FIG. **11**, the eBPF can be a program that can be executed inside a kernel for tracing, monitoring, and processing one or more packets. There can be various hook point(s) inside the kernel that allow for attaching eBPF program(s) to the veth device(s). The traffic control (TC) can include a tool for manipulating packets at a really low level. For example, it can attach an eBPF program to a network device on both the ingress side and/or the egress side. On the ingress side **1105**, eBPF program **1114** can be executed before a packet is

handed to the local network stack. Similarly, on the egress side **1125**, the eBPF program **1124** is executed after the local network stack has finished processing (e.g., local processes **1108**). The GENEVE encapsulation protocol, such as the examples described and depicted in connection with FIGS. 7 and 8, is the encapsulation protocol implemented in Linux to create a tunnel between two network devices. Although it will be understood by those having skill in the art that other encapsulation protocols can be implemented according to varying example embodiments.

[0162] According to some examples, a high-level logical architecture is depicted in the block diagram **1100**. All transmission control protocol (TCP) packets and user datagram protocol (UDP) packets coming out from the UDx sandbox can be routed through an egress proxy. Traffic control (TC) can attach some or all cBPF programs on the ingress side **1105** and/or the egress side **1125** of the network interfaces created on the execution platform **110** and proxy host(s). The cBPF program(s) will primarily be performing packet encapsulation and/or decapsulation for a tunnel between the execution platform **110** and the proxy service **115**, packet redirect between network interfaces on the same host, SNAT (e.g., source IP network address translation) on one or more proxy host(s) before the packet is sent to the public Internet, policy enforcement based on the policy identifier (e.g., policy_ID) that is embedded in the encapsulated packet header, and the like.

[0163] According to some examples, the egress proxy, such as a proxy service **115**, is in communication with the execution platform **110**, but not the compute service manager **108**. A UDx manager (e.g., UDFNetManager) running on each execution platform host is provisioned to assign IP addresses to the sandbox process **420**, such that the IP address inside the network namespace of the sandbox is unique within a single host. Further examples include the IP address inside the network namespace of the sandbox to be the same across different execution platform instances, where those instances can be going through the same or different egress proxies. Thus, in order to uniquely identify a sandbox environment from the egress proxy, the cloud data platform needs a tuple of the execution platform IP address and the sandbox namespace IP address.

[0164] According to some examples, Network Address Translation (NAT) **1140** and/or Source IP NAT (SNAT) are important components to forward packets to the public Internet because each sandbox on the execution platform will be assigned a private IP address. To make sure returned packets can reach back to the sandbox, the proxy needs to implement NAT/SNAT. According to the present disclosure, there are three example methods that are implemented according to this functionality.

[0165] In a first example embodiment, two-level NAT translation using eBPF code can be implemented by translating a private IP address in a sandbox to the execution platform IP address. Then in the proxy host, translate the execution platform IP address into proxy IP addresses to access public Internet. In this approach, NAT as well as connection tracking can be implemented completely inside the eBPF program.

[0166] In a second example embodiment, two-level NAT translation using IP tables can be implemented by translating a private IP address in a sandbox to the execution platform IP address. Then in the proxy host, it translates the execution platform IP address into proxy IP address to access public Internet. Both translations can leverage IP tables.

[0167] In a third example embodiment, the compute service manager **108** is configured to manage IP address mapping, where NAT only needs to be done once leveraging IP tables in the proxy host. In the third example embodiment, the compute service manager needs to implement an IP assignment service to assign a unique IP to each sandbox. During the assignment, the compute service manager also needs to record a mapping between the execution platform IP and this unique IP. The example further includes the egress proxy host needing to talk to the compute service manager to ask about the mapping between sandbox IP and the execution platform IP. In the forward path, the egress proxy will leverage IP tables to translate the sandbox IP address to an egress proxy host IP. In the returned path, the egress proxy will need to ask the compute service manager **108** about the execution platform **110** IP address that the destination sandbox IP belongs

to and then perform the packet encapsulation.

[0168] The logical architecture of FIG. 11 provides for five stages of processing including a pre-routing stage 1102, an input stage 1103, a forward stage 1104, an output stage 1106, and a post-routing stage 1107. The NAT/SNAT 1140 is performed during the pre-routing stage 1102, output stage 1106, and post-routing stage 1107. The filtering processes 1110 are performed during the input stage 1103, the forwarding stage 1104, and the output stage 1106. The filtering processes can include the use of IP tables filtering hooks 1111. On the ingress side 1105, the eBPF program 1114 receives input from the eBPF XDP hook 1130 and provides output to a connected eBPF traffic control hook 1120, which then routes information into the pre-routing stage 1102. On the egress side 1125, the post-routing stage 1107 provides output to the eBPF traffic control hook 1120, which is connected to the eBPF program 1124.

[0169] FIG. 12 is a block diagram illustrating a packet routing diagram 1200 depicting how a packet is routed from the execution platform 110 to the proxy service 115, in accordance with example embodiments. More specifically, FIG. 12 shows how a packet is routed from a sandbox (e.g., 10.1.1.1) in an execution platform 110 to a remote destination (e.g., 8.8.8.8).

[0170] The packet routing diagram 1200 begins with a virtual ethernet (veth) pair being created. At step 1 1291, one end of the veth pair (e.g., veth0 1201) is moved to the network namespace inside the sandbox0 1220 while the other end of the veth pair (e.g., veth1 1202) remains on the host. Routing tables inside the sandbox0 1220 namespace will default route any packet to use veth1 1202 IP address 10.1.1.2 as the gateway. So, a packet will be transmitted from veth0 1201 and received by veth1 1202.

[0171] At step 2 1292, an eBPF program is attached to the ingress (e.g., received) side of veth1 1202, and this eBPF program will perform certain tasks. For example, the eBPF program can perform one or more of the following tasks: (A) Check against a policy on whether the packet from sandbox0 1220 is allowed to access 8.8.8.8. If the packet is not allowed access, the eBPF program will drop the packet and send an event to the user-space application. (B) Set GENEVE tunnel related data to prepare for packet encapsulation, where the data includes, for example, a remote IP address of a proxy service and customized metadata (e.g., policy_ID and execution platform IP address). (C) Call bpf_redirect to redirect the packet to the egress side of geneve0 device 1203.

[0172] The packet routing diagram 1200 continues at step 3 1293 with an eBPF program being attached on the egress side of the geneve0 device 1203, which will SNAT the inner packet to change the source IP address to the eth0 1204 IP address of the current execution platform box. Next, the kernel can encapsulate this packet with the metadata set by task (B) noted above.

[0173] At step 4 1294, the encapsulated packet is transferred on wire. At step 5 1295, actual packet decapsulation is performed by the kernel (e.g., no eBPF code needs to be attached). At step 6 1296, an eBPF program is attached to the ingress (e.g., receiving) side of geneve1 device 1205, and this eBPF program will perform following tasks: (A) Check against one or more policies based on the policy identifier as to whether the packet should be allowed or blocked. If the packet should be blocked, drop the packet, and treat this attempt as a security incident. (B) Let the packet go through the local network stack, with no redirect.

[0174] The packet routing diagram 1200 continues at step 7 1297, with the packet being routed based on the destination IP address (e.g., 8.8.8.8), which will go through the default routing rule on the proxy host and will be transmitted under eth0 1206. At step 8 1298, an eBPF program will be attached on the egress side of eth0 1206, which will SNAT from the execution platform eth0 1206 IP address to the proxy service 115 eth0 1206 IP address.

[0175] FIG. 13 is a block diagram illustrating a packet routing diagram 1300 depicting how a returned packet is routed back to sandbox0 1320 (e.g., 10.1.1.1) from a remote source 1301 (e.g., 8.8.8.8), in accordance with example embodiments. In some examples, the packet routing diagram 1300 can include some or all of the same components as in the packet routing diagram 1200 depicted and described in connection with FIG. 12. FIG. 13 is similar in many respects to FIG. 12,

sharing certain similar features and components; for brevity, only specific elements are detailed and described with reference to FIG. 13. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from FIG. 13. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the packet routing diagram 1300 to facilitate additional functionality that is not specifically described herein.

[0176] The packet routing diagram 1300 begins at step 1 1391 with remote upstream transmitting a packet back to an egress proxy server, such as proxy service 115. At step 2 1392, an eBPF program will be attached to the ingress side of the eth0 1302 interface. The program will perform the following actions: (A) Check the destination port number on the packet and look up the NAT table maintained by the cloud data platform. If there is no entry with the corresponding port number, let the packet proceed. (B) If there is an entry found in the NAT table, then lookup the original execution platform IP address. (C) Overwrite the packet destination IP address to match the execution platform IP address. (D) Set up GENEVE encapsulation metadata (e.g., remote IP address, execution platform IP address, etc.). (E) Redirect the packet to the egress side of the GENEVE device.

[0177] The packet routing diagram 1300 continues at step 3 1393 where the kernel encapsulates the packet according to the GENEVE metadata set at step 2 1392. At step 4 1394, the packet is transferred on wire, and at step 5 1395, the kernel decapsulates the packet. The packet routing diagram 1300 continues at step 6 1396 where the packet reaches to (e.g., connects with) the geneve0 device 1306. An eBPF program will be attached to the ingress of the geneve0 device 1306, and the eBPF program will perform the following: (A) Fix the MAC address of the packet. (B) Reverse SNAT from the execution platform IP address to the sandbox IP address. (C) Then the eBPF program will allow the packet to go through the network stack 1307 in the execution platform host namespace, where the packet will be sent to veth1 1308 according to the routing table. At step 7, the packet will be sent back to veth0 1309.

[0178] In additional example embodiments, multiple eBPF programs are attached to different places for tracking and monitoring purposes. For example, in one example implementation, at least four eBPF programs are implemented, including veth1 implemented on the execution platform ingress side for outgoing packets, geneve0 on the execution platform ingress side for returned packets, geneve1 on the proxy ingress side for outgoing packets, and eth0 on the proxy ingress side for returning packets.

[0179] FIG. 14 is a block diagram illustrating an example service proxy process 1400 used by the cloud data platform to route and manage network traffic between client applications and cloud data platform services, in accordance with example embodiments. The example embodiment provides UDx external access to enable users to transmit (e.g., send) arbitrary packets of data from a UDx sandbox to the public Internet, including, for example, forwarding packets from the execution platform through an egress proxy host to allow fine-grained access control and traffic visibility.

[0180] According to some examples, when a packet lands on a proxy host, there are generally three steps that are performed including: (1) decapsulation of the packet, (2) policy enforcement, and (3) forwarding of the packet to an upstream destination. Each of these three steps can be performed in one of three methodologies, including: (1) by a service proxy process (described in detail in connection with FIG. 14), (2) in the kernel by an eBPF program (described in detail in connection with FIG. 15), and/or (3) a combination of the eBPF program and service proxy process (described in detail in connection with FIG. 16).

[0181] According to an example embodiment using the first methodology, a proxy service process provides for strong encryption and authentication mechanisms to secure traffic between customers and cloud data platform services. According to the first methodology of the service proxy process, this example approach indicates that all packets that are encapsulated on the execution platform

side will be directly sent to a port that the proxy service process is monitoring (e.g., listening on). [0182] For example, the proxy service process **1400** begins at step **1401** with a User Datagram Protocol (UDP) listener, such as a process or application, waiting for incoming UDP packets on a specific port number (e.g., on Port **8888**). As UDP is a connectionless protocol that allows data packets to be sent from one application to another without establishing a connection first. An application using the proxy service process **1400** opens a UDP listener binding it to a specific port number on a host machine and waiting for incoming UDP packets. When the UDP packet arrives at the specific port (e.g., Port **8888**), the listener at step **1401** captures the packet and processes it according to rules defined by the application. At step **1402**, the proxy service process intercepts and analyzes the captured packet at a decapsulation filter designed to analyze traffic that has been encapsulated or tunneled using a specific protocol and to strip off the encapsulation so the original traffic can be analyzed and processed. For example, the traffic being encapsulated using the GENEVE protocol, a decapsulation filter can intercept the incoming traffic and remove the GENEVE header and options so that the original inner packet can be processed. This decapsulation filter allows security and monitoring tools to inspect the traffic at the application layer, which is typically not possible if the traffic is still encapsulated.

[0183] At step **1403**, an external authorization filter is made available for a user to configure the network or HTTP(S) filter to call an external service (e.g., via gRPC or HTTP) for authentication and/or authorization purposes. For example, the external authorization filter can include a policy enforcement action designed to evaluate requests for network resources based on external policies or rules (e.g., such as user privileges, group membership, time of day, location, etc.) in order to determine whether the request should be allowed or denied. The policy enforcement filter can be used to restrict access to a sensitive database based on a user's role or privileges assigned to the user from the account administrator. If the filter determines the request should be denied, it transmits the denial to a policy agent at step **1404** and denies the request. If the filter determines the user is authorized to access the resource, the proxy service process continues to step **1405**.

[0184] At step **1405**, a bridge is used to connect an internal listener to an upstream network, where the bridge acts as a networking device to connect two or more segments together. For example, by using a bridge at step **1405**, the proxy server process can route traffic from external networks to the internal listener without exposing the listener directly to the external network (e.g., thereby providing additional layer(s) of security).

[0185] Depending on the inner packet protocol determined at step **1406**, the proxy service process either forwards the packet to a UDP listener at step **1407** or to a TCP listener at step **1410**. At step **1407**, the UDP listener on an internal address of a proxy server (e.g., a networking device acting as an intermediary between clients and servers). At step **1408**, the process continues to a UDP proxy filter, which is a networking component that operates within a UDP proxy server to inspect and/or modify UDP traffic as it passes through the server. The UDP proxy filter examines each UDP packet and determines how to handle the packet based on a predefined set of rules and/or policies (e.g., security and access control policies). At step **1409**, the process continues to a dynamic forward upstream cluster (e.g., a cluster of servers) responsible for processing requests and/or providing responses to downstream customers or services.

[0186] Returning to step **1406**, depending on the inner packet protocol, the proxy service process forwards to a Transmission Control Protocol (TCP) listener on an internal address of a proxy server at step **1410**. The TCP listener can be a networking component configured to operate within the server to accept incoming TCP connections from downstream clients. The TCP listener can divide the data to be transmitted into one or more segments and transmits each segment over the network using IP to step **1411**.

[0187] At step **1411**, a TCP proxy filter receives the segments and/or intercepts and analyzes TCP traffic between a client and a server host (e.g., a middleman to modify or filter traffic before forwarding the traffic to the destination). At step **1412**, the process continues to a dynamic forward

upstream cluster (e.g., a cluster of servers) responsible for processing requests and/or providing responses to downstream customers or services.

[0188] FIG. **15** is a block diagram illustrating an example of a second methodology **1500** used by the cloud data platform to route and manage network traffic between client applications and cloud data platform services, in accordance with example embodiments.

[0189] According to an example embodiment using the second methodology, a pure eBPF approach is implemented to provide an approach that indicates that all packet processing will happen inside the eBPF program. At step **1501**, the policy agent is configured to push policies into the kernel via an eBPF or BPF map at step **1502**, which is a shared memory address between the userspace **1510** and the kernel **1520**. When a packet lands on a proxy host, the decapsulation of the packet, policy enforcement, and packet forwarding are all performed inside the eBPF program attached to a traffic control (TC) subsystem of the Linux kernel at step **1503**. For example, the TC GENEVE device leverages the traffic control subsystem to perform its encapsulation and decapsulation functions over TCP.

[0190] Example embodiments of the second methodology **1500** enable excellent performance as packets will not go through the userspace **1510**, and most of the eBPF program can be shared with the execution platform side including pushing the policy into the eBPF program.

[0191] FIG. **16** is a block diagram illustrating an example of a third methodology **1600** used by the cloud data platform to route and manage network traffic between client applications and cloud data platform services, in accordance with example embodiments.

[0192] According to an example embodiment using the third methodology **1600**, a combination approach is implemented by combining the service proxy process (e.g., as described in FIG. **14**) and the eBPF program process (e.g., as described in FIG. **15**). Examples of this approach start at step **1601**, which indicates that the packet will be decapsulated inside the eBPF program attached to the TC GENEVE device. At step **1602**, the IP Table pre-routing rule will REDIRECT the packet to a local host and port that the service proxy process will be listening on. For example, the pre-routing rule will redirect all TCP packets from the GENEVE device to local port **8888** and redirect all UDP packets from the GENEVE device to local port **8889**. The service proxy process will extract the SO_ORIGINAL_DST socket option set by the IP Table. Where the socket option is a parameter to modify the behavior of the endpoints of a bidirectional communication channel between two networked applications. The service proxy process will then forward the packet to the right upstream destination in a similar manner as described in connection with FIG. **14**.

[0193] FIG. **17** depicts a flow diagram of a method **1700** for applying a packages policy, according to some example embodiments. The method **1700** can be embodied in machine-readable instructions for execution by one or more hardware components (e.g., one or more processors) such that the operations of the method **1700** can be performed by components of the cloud data platform **102**. Accordingly, the method **1700** is described below, by way of example with reference to components of the cloud data platform **102**. However, it shall be appreciated that method **1700** can be deployed on various other hardware configurations and is not intended to be limited to deployment within the cloud data platform **102**.

[0194] Depending on the embodiment, an operation of the method **1700** can be repeated in different ways or involve intervening operations not shown. Though the operations of the method **1700** can be depicted and described in a certain order, the order in which the operations are performed may vary among embodiments, including performing certain operations in parallel or performing sets of operations in separate processes. While the various operations in this flowchart are presented and described sequentially, one of ordinary skill will appreciate that some or all of the operations may be executed in a different order, be combined or omitted, or be executed in parallel.

[0195] Operation **1702** is for receiving, by the cloud data platform **102**, a remote procedure call from a user-defined function (UDF) executing within a sandbox process. Operation **1704** is for providing, by the cloud data platform **102**, an external access overlay network to establish a secure

egress path for UDF external access. Operation **1706** is for enabling, by the cloud data platform **102**, the UDF including the code running in the sandbox process to initiate a network call.

[0196] Another general aspect is for a system that includes a memory comprising instructions and one or more computer processors. The instructions, when executed by the one or more computer processors, cause the one or more computer processors to perform operations. In yet another general aspect, a tangible machine-readable storage medium (e.g., a non-transitory storage medium) includes instructions that, when executed by a machine, cause the machine to perform operations.

[0197] Described implementations of the subject matter can include one or more features, alone or in combination as illustrated below by way of example.

[0198] Example 1 is a method comprising: receiving, by one or more hardware processors, a remote procedure call related to a user-defined function (UDF) executing within a sandbox process, the UDF including code related to at least one operation to be performed; providing an overlay network to establish a secure egress path for UDF external access; and enabling the UDF executing in the sandbox process to initiate a network call based at least in part on the overlay network.

[0199] In Example 2, the subject matter of Example 1 optionally includes, wherein establishing the secure egress path for the UDF external access comprises: enabling external access from the sandbox process by identifying virtual ethernet pairs for putting an ethernet device in a namespace of the sandbox process; and performing packet validation to verify a packet.

[0200] In Example 3, the subject matter of any one of Examples 1-2 optionally include, deploying, on an egress proxy host, an extended Berkeley Packet Filter (eBPF) to decapsulate the packet, wherein decapsulating the packet includes extracting a policy identifier; receiving, from an execution platform, an egress policy based at least in part on the policy identifier; and reviewing the egress policy to ensure packet destination authorization.

[0201] In Example 4, the subject matter of any one of Examples 1-3 optionally include, capturing log data including error events related to the UDF; and enabling revocation of access usage rights in response to the log data.

[0202] In Example 5, the subject matter of any one of Examples 1-4 optionally include, in response to the network call, providing direct external access to the UDF executing within the sandbox process via the secure egress path to a public Internet host.

[0203] In Example 6, the subject matter of any one of Examples 1-5 optionally include, instantiating the sandbox process as a child process of a platform agent, the platform agent configured to determine whether the UDF is permitted to initiate the network call; and instantiating a user code runtime as a child process of the sandbox process, the sandbox process configured to execute the at least one operation in a sandbox environment.

[0204] In Example 7, the subject matter of any one of Examples 1-6 optionally include, instantiating a user code runtime to execute the UDF code, the user code runtime instantiated within the sandbox process of an execution node.

[0205] In Example 8, the subject matter of any one of Examples 1-7 optionally include, collecting metric event information associated with the UDF, the metric event information including execution statistics associated with execution of the UDF code by the user code runtime.

[0206] In Example 9, the subject matter of any one of Examples 1-8 includes, preventing the code from making a Domain Name System (DNS) request to an unauthorized host, the preventing including specifying where sandboxed functions can connect based on DNS names or wildcard-based DNS patterns.

[0207] In Example 10, the subject matter of any one of Examples 1-9 optionally include, instantiating a user code runtime to execute the code of the UDF, the user code runtime including an access control list, the access control list including authorized hosts and access usage rights.

[0208] Example 11 is a system comprising: one or more hardware processors of a machine; and at least one memory storing instructions that, when executed by the one or more hardware processors,

cause the system to perform operations comprising: receiving, by the one or more hardware processors, a remote procedure call from a user-defined function (UDF) executing within a sandbox process, the UDF including code related to at least one operation to be performed; providing an overlay network to establish a secure egress path for UDF external access; and enabling the UDF executing in the sandbox process to initiate a network call based at least in part on the overlay network.

[0209] In Example 12, the subject matter of Example 11 optionally includes, wherein establishing the secure egress path for the UDF external access comprises: enabling external access from the sandbox process by identifying virtual ethernet pairs for putting an ethernet device in a namespace of the sandbox process; and performing packet validation to verify a packet.

[0210] In Example 13, the subject matter of any one of Examples 11-12 optionally include, the operations further comprising: deploying, on an egress proxy host, an extended Berkeley Packet Filter (eBPF) to decapsulate the packet, wherein decapsulating the packet includes extracting a policy identifier; receiving, from an execution platform, an egress policy based at least in part on the policy identifier; and reviewing the egress policy to ensure packet destination authorization.

[0211] In Example 14, the subject matter of any one of Examples 11-13 optionally include, the operations further comprising: capturing log data including error events related to the UDF; and enabling revocation of access usage rights in response to the log data.

[0212] In Example 15, the subject matter of any one of Examples 11-14 optionally include, the operations further comprising: in response to the network call, providing direct external access to the UDF executing within the sandbox process via the secure egress path to a public Internet host.

[0213] In Example 16, the subject matter of any one of Examples 11-15 optionally include, the operations further comprising: instantiating the sandbox process as a child process of a platform agent, the platform agent configured to determine whether the UDF is permitted to initiate the network call; and instantiating a user code runtime as a child process of the sandbox process, the sandbox process configured to execute the at least one operation in a sandbox environment.

[0214] In Example 17, the subject matter of any one of Examples 11-16 optionally include, the operations further comprising: instantiating a user code runtime to execute the UDF code, the user code runtime instantiated within the sandbox process of an execution node.

[0215] In Example 18, the subject matter of any one of Examples 11-17 optionally include, the operations further comprising: collecting metric event information associated with the UDF, the metric event information including execution statistics associated with execution of the UDF code by the user code runtime.

[0216] In Example 19, the subject matter of any one of Examples 11-18 optionally include, the operations further comprising: preventing the code from making a Domain Name System (DNS) request to an unauthorized host, the preventing including specifying where sandboxed functions can connect based on DNS names or wildcard-based DNS patterns.

[0217] In Example 20, the subject matter of any one of Examples 11-19 optionally include, the operations further comprising: instantiating a user code runtime to execute the code of the UDF, the user code runtime including an access control list, the access control list including authorized hosts and access usage rights.

[0218] Example 21 is a machine-storage medium embodying instructions that, when executed by a machine, cause the machine to perform operations comprising: receiving, by one or more hardware processors, a remote procedure call from a user-defined function (UDF) executing within a sandbox process, the UDF including code related to at least one operation to be performed; providing an overlay network to establish a secure egress path for UDF external access; and enabling the UDF executing in the sandbox process to initiate a network call based at least in part on the overlay network.

[0219] In Example 22, the subject matter of Example 21 optionally includes, wherein establishing the secure egress path for the UDF external access comprises: enabling external access from the

sandbox process by identifying virtual ethernet pairs for putting an ethernet device in a namespace of the sandbox process; and performing packet validation to verify a packet.

[0220] In Example 23, the subject matter of any one of Examples 21-22 optionally include, deploying, on an egress proxy host, an extended Berkeley Packet Filter (eBPF) to decapsulate the packet, wherein decapsulating the packet includes extracting a policy identifier; receiving, from an execution platform, an egress policy based at least in part on the policy identifier; and reviewing the egress policy to ensure packet destination authorization.

[0221] In Example 24, the subject matter of any one of Examples 21-23 optionally include, capturing log data including error events related to the UDF; and enabling revocation of access usage rights in response to the log data.

[0222] In Example 25, the subject matter of any one of Examples 21-24 optionally include, in response to the network call, providing direct external access to the UDF executing within the sandbox process via the secure egress path to a public Internet host.

[0223] In Example 26, the subject matter of any one of Examples 21-25 optionally include, instantiating the sandbox process as a child process of a platform agent, the platform agent configured to determine whether the UDF is permitted to initiate the network call; and instantiating a user code runtime as a child process of the sandbox process, the sandbox process configured to execute the at least one operation in a sandbox environment.

[0224] In Example 27, the subject matter of any one of Examples 21-26 optionally include, instantiating a user code runtime to execute the UDF code, the user code runtime instantiated within the sandbox process of an execution node.

[0225] In Example 28, the subject matter of any one of Examples 21-27 optionally include, collecting metric event information associated with the UDF, the metric event information including execution statistics associated with execution of the UDF code by the user code runtime.

[0226] In Example 29, the subject matter of any one of Examples 21-28 optionally include, preventing the code from making a Domain Name System (DNS) request to an unauthorized host, the preventing including specifying where sandboxed functions can connect based on DNS names or wildcard-based DNS patterns.

[0227] In Example 30, the subject matter of any one of Examples 21-29 optionally include, instantiating a user code runtime to execute the code of the UDF, the user code runtime including an access control list, the access control list including authorized hosts and access usage rights.

[0228] Example 31 is at least one machine-readable medium including instructions that, when executed by processing circuitry, cause the processing circuitry to perform operations to implement of any of Examples 1-30. Example 32 is an apparatus comprising means to implement of any of Examples 1-30.

[0229] FIG. **18** illustrates a diagrammatic representation of a machine **1800** in the form of a computer system within which a set of instructions can be executed for causing the machine **1800** to perform any one or more of the methodologies discussed herein, according to an example embodiment. Specifically, FIG. **18** shows a diagrammatic representation of the machine **1800** in the example form of a computer system, within which instructions **1816** (e.g., software, a program, an application, an applet, an app, or other executable code) for causing the machine **1800** to perform any one or more of the methodologies discussed herein can be executed. For example, the instructions **1816** may cause the machine **1800** to execute any one or more operations of any one or more of the methods described herein. As another example, the instructions **1816** may cause the machine **1800** to implement portions of the data flows described herein. In this way, the instructions **1816** transform a general, non-programmed machine into a particular machine **1800** (e.g., the compute service manager **108**, the execution platform **110**, client device **114**) that is specially configured to carry out any one of the described and illustrated functions in the manner described herein.

[0230] In alternative embodiments, the machine **1800** operates as a standalone device or can be

coupled (e.g., networked) to other machines. In a networked deployment, the machine **1800** may operate in the capacity of a server machine or a client machine in a server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine **1800** may comprise, but not be limited to, a server computer, a client computer, a personal computer (PC), a tablet computer, a laptop computer, a netbook, a smart phone, a mobile device, a network router, a network switch, a network bridge, or any machine capable of executing the instructions **1816**, sequentially or otherwise, that specify actions to be taken by the machine **1800**. Further, while only a single machine **1800** is illustrated, the term “machine” shall also be taken to include a collection of machines **1800** that individually or jointly execute the instructions **1816** to perform any one or more of the methodologies discussed herein.

[0231] The machine **1800** includes processors **1810**, memory **1830**, and input/output (I/O) components **1850** configured to communicate with each other such as via a bus **1802**. In an example embodiment, the processors **1810** (e.g., a central processing unit (CPU), a reduced instruction set computing (RISC) processor, a complex instruction set computing (CISC) processor, a graphics processing unit (GPU), a digital signal processor (DSP), an application-specific integrated circuit (ASIC), a radio-frequency integrated circuit (RFIC), another processor, or any suitable combination thereof) may include, for example, a processor **1812** and a processor **1814** that may execute the instructions **1816**. The term “processor” is intended to include multi-core processors **1810** that may comprise two or more independent processors (sometimes referred to as “cores”) that may execute instructions **1816** contemporaneously. Although FIG. **18** shows multiple processors **1810**, the machine **1800** may include a single processor with a single core, a single processor with multiple cores (e.g., a multi-core processor), multiple processors with a single core, multiple processors with multiple cores, or any combination thereof.

[0232] The memory **1830** may include a main memory **1832**, a static memory **1834**, and a storage unit **1836**, all accessible to the processors **1810** such as via the bus **1802**. The main memory **1832**, the static memory **1834**, and the storage unit **1836** comprising a machine storage medium **1838** may store the instructions **1816** embodying any one or more of the methodologies or functions described herein. The instructions **1816** may also reside, completely or partially, within the main memory **1832**, within the static memory **1834**, within the storage unit **1836**, within at least one of the processors **1810** (e.g., within the processor's cache memory), or any suitable combination thereof, during execution thereof by the machine **1800**.

[0233] The I/O components **1850** include components to receive input, provide output, produce output, transmit information, exchange information, capture measurements, and so on. The specific I/O components **1850** that are included in a particular machine **1800** will depend on the type of machine. For example, portable machines such as mobile phones will likely include a touch input device or other such input mechanisms, while a headless server machine will likely not include such a touch input device. It will be appreciated that the I/O components **1850** may include many other components that are not shown in FIG. **18**. The I/O components **1850** are grouped according to functionality merely for simplifying the following discussion and the grouping is in no way limiting. In various example embodiments, the I/O components **1850** may include output components **1852** and input components **1854**. The output components **1852** may include visual components (e.g., a display such as a plasma display panel (PDP), a light emitting diode (LED) display, a liquid crystal display (LCD), a projector, or a cathode ray tube (CRT)), acoustic components (e.g., speakers), other signal generators, and so forth. The input components **1854** may include alphanumeric input components (e.g., a keyboard, a touch screen configured to receive alphanumeric input, a photo-optical keyboard, or other alphanumeric input components), point-based input components (e.g., a mouse, a touchpad, a trackball, a joystick, a motion sensor, or another pointing instrument), tactile input components (e.g., a physical button, a touch screen that provides location and/or force of touches or touch gestures, or other tactile input components), audio input components (e.g., a microphone), and the like.

[0234] Communication can be implemented using a wide variety of technologies. The I/O components **1850** may include communication components **1864** operable to couple the machine **1800** to a network **1881** via a coupling **1883** or to devices **1880** via a coupling **1882**. For example, the communication components **1864** may include a network interface component or another suitable device to interface with the network **1881**. In further examples, the communication components **1864** may include wired communication components, wireless communication components, cellular communication components, and other communication components to provide communication via other modalities. The devices **1880** can be another machine or any of a wide variety of peripheral devices (e.g., a peripheral device coupled via a universal serial bus (USB)). For example, as noted above, the machine **1800** may correspond to any one of the client devices **114**, the compute service manager **108**, the execution platform **110**, and the devices **1880** may include any other of these systems and devices.

[0235] The various memories (e.g., **1830**, **1832**, **1834**, and/or memory of the processor(s) **1810** and/or the storage unit **1836**) may store one or more sets of instructions **1816** and data structures (e.g., software) embodying or utilized by any one or more of the methodologies or functions described herein. These instructions **1816**, when executed by the processor(s) **1810**, cause various operations to implement the disclosed embodiments.

[0236] As used herein, the terms “machine-storage medium,” “device-storage medium,” and “computer-storage medium” mean the same thing and can be used interchangeably in this disclosure. The terms refer to a single or multiple storage devices and/or media (e.g., a centralized or distributed database, and/or associated caches and servers) that store executable instructions and/or data. The terms shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media, including memory internal or external to processors. Specific examples of machine-storage media, computer-storage media, and/or device-storage media include non-volatile memory, including by way of example semiconductor memory devices, e.g., erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), field-programmable gate arrays (FPGAs), and flash memory devices; magnetic disks such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The terms “machine-storage media,” “computer-storage media,” and “device-storage media” specifically exclude carrier waves, modulated data signals, and other such media, at least some of which are covered under the term “signal medium” discussed below.

[0237] In various example embodiments, one or more portions of the network **1881** can be an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a local-area network (LAN), a wireless LAN (WLAN), a wide-area network (WAN), a wireless WAN (WWAN), a metropolitan-area network (MAN), the Internet, a portion of the Internet, a portion of the public switched telephone network (PSTN), a plain old telephone service (POTS) network, a cellular telephone network, a wireless network, a Wi-Fi® network, another type of network, or a combination of two or more such networks. For example, the network **1881** or a portion of the network **1881** may include a wireless or cellular network, and the coupling **1883** can be a Code Division Multiple Access (CDMA) connection, a Global System for Mobile communications (GSM) connection, or another type of cellular or wireless coupling. In this example, the coupling **1883** may implement any of a variety of types of data transfer technology, such as Single Carrier Radio Transmission Technology (1×RTT), Evolution-Data Optimized (EVDO) technology, General Packet Radio Service (GPRS) technology, Enhanced Data rates for GSM Evolution (EDGE) technology, third Generation Partnership Project (3GPP) including 3G, fourth generation wireless (4G) networks, Universal Mobile Telecommunications System (UMTS), High-Speed Packet Access (HSPA), Worldwide Interoperability for Microwave Access (WiMAX), Long Term Evolution (LTE) standard, others defined by various standard-setting organizations, other long-range protocols, or other data transfer technology.

[0238] The instructions **1816** can be transmitted or received over the network **1881** using a transmission medium via a network interface device (e.g., a network interface component included in the communication components **1864**) and utilizing any one of a number of well-known transfer protocols (e.g., hypertext transfer protocol (HTTP)). Similarly, the instructions **1816** can be transmitted or received using a transmission medium via the coupling **1882** (e.g., a peer-to-peer coupling) to the devices **1880**. The terms “transmission medium” and “signal medium” mean the same thing and can be used interchangeably in this disclosure. The terms “transmission medium” and “signal medium” shall be taken to include any intangible medium that is capable of storing, encoding, or carrying the instructions **1816** for execution by the machine **1800**, and include digital or analog communications signals or other intangible media to facilitate communication of such software. Hence, the terms “transmission medium” and “signal medium” shall be taken to include any form of modulated data signal, carrier wave, and so forth. The term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal.

[0239] Examples, as described herein, may include, or may operate by, logic, a number of components, or mechanisms. Circuitry is a collection of circuits implemented in tangible entities that include hardware (e.g., simple circuits, gates, logic). Circuitry membership may be flexible over time and underlying hardware variability. Circuitries include members that may, alone or in combination, perform specified operations when operating. In an example, hardware of the circuitry may be immutably designed to carry out a specific operation (e.g., hardwired). In an example, the hardware of the circuitry may include variably connected physical components (e.g., execution units, transistors, simple circuits) including a computer-readable medium physically modified (e.g., magnetically, electrically, by moveable placement of invariant massed particles) to encode instructions of the specific operation. In connecting the physical components, the underlying electrical properties of a hardware constituent are changed (for example, from an insulator to a conductor or vice versa). The instructions enable embedded hardware (e.g., the execution units or a loading mechanism) to create members of the circuitry in hardware via the variable connections to carry out portions of the specific operation when in operation. Accordingly, the computer-readable medium is communicatively coupled to the other components of the circuitry when the device is operating. In an example, any of the physical components may be used in more than one member of more than one circuitry. For example, under operation, execution units may be used in a first circuit of a first circuitry at one point in time and reused by a second circuit in the first circuitry, or by a third circuit in a second circuitry, at a different time.

[0240] The terms “machine-readable medium,” “computer-readable medium,” and “device-readable medium” mean the same thing and can be used interchangeably in this disclosure. The terms are defined to include both machine-storage media and transmission media. Thus, the terms include both storage devices/media and carrier waves/modulated data signals.

[0241] The various operations of example methods described herein can be performed, at least partially, by one or more processors that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Similarly, the methods described herein can be at least partially processor implemented. For example, at least some of the operations of the methods described herein can be performed by one or more processors. The performance of certain of the operations can be distributed among the one or more processors, not only residing within a single machine, but also deployed across a number of machines. In some example embodiments, the processor or processors can be located in a single location (e.g., within a home environment, an office environment, or a server farm), while in other embodiments the processors can be distributed across a number of locations.

[0242] Although the embodiments of the present disclosure have been described with reference to specific example embodiments, it will be evident that various modifications and changes can be made to these embodiments without departing from the broader scope of the inventive subject

matter. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof show, by way of illustration, and not of limitation, specific embodiments in which the subject matter can be practiced. The embodiments illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other embodiments can be used and derived therefrom, such that structural and logical substitutions and changes can be made without departing from the scope of this disclosure. This Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

[0243] Such embodiments of the inventive subject matter can be referred to herein, individually and/or collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific embodiments have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose can be substituted for the specific embodiments shown. This disclosure is intended to cover any and all adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent to those of skill in the art, upon reviewing the above description.

[0244] In this document, the terms “a” or “an” are used, as is common in patent documents, to include one or more than one, independent of any other instances or usages of “at least one” or “one or more.” In this document, the term “or” is used to refer to a nonexclusive or, such that “A or B” includes “A but not B,” “B but not A,” and “A and B,” unless otherwise indicated. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein.” Also, in the following claims, the terms “including” and “comprising” are open-ended; that is, a system, device, article, or process that includes elements in addition to those listed after such a term in a claim is still deemed to fall within the scope of that claim.

Claims

1. A method comprising: receiving, by one or more hardware processors, a request to execute a user-defined function (UDF) within a sandbox process; establishing a secure egress path for the UDF using an overlay network, wherein the overlay network includes a dedicated DNS resolver at a proxy service; receiving, from the UDF, a DNS request to resolve a hostname; validating, by the proxy service, that the hostname is included in an allowed host list associated with the UDF; resolving, by the dedicated DNS resolver, the hostname to an IP address using a UDP listener configured to handle DNS protocol traffic on a designated port of the proxy service; and enabling the UDF to communicate with a host at the resolved IP address via the secure egress path.
2. The method of claim 1, wherein establishing the secure egress path comprises: identifying a virtual ethernet pair to isolate UDF traffic within the sandbox process; using the virtual ethernet pair to put an ethernet device in a namespace of the sandbox process; and implementing packet encapsulation between an execution platform and the proxy service.
3. The method of claim 1, further comprising: deploying, on the proxy service, an extended Berkeley Packet Filter (eBPF) to decapsulate DNS packets; extracting a policy identifier from the decapsulated DNS packets; and validating the policy identifier against stored egress policies.
4. The method of claim 1, wherein the overlay network comprises, a first tier implementing virtual ethernet pairs for sandbox isolation, a second tier providing Generic Network Virtualization Encapsulation (GENEVE) tunneling between an execution platform and the proxy service, and a third tier handling verified packet routing to external destinations.
5. The method of claim 1, further comprising: collecting metric event information associated with

DNS resolution attempts; logging unauthorized hostname resolution attempts; and publishing metrics to an administrator interface.

6. The method of claim 1, wherein validating the hostname comprises: performing a lookup in a policy map using a sandbox identifier; verifying DNS patterns against wildcard-based allowlist entries; and reporting policy violations to a security monitoring system.

7. The method of claim 1, further comprising: maintaining a connection map at the proxy service to track authorized DNS resolutions; performing Source Network Address Translation (SNAT) on validated packets; and routing translated packets to authorized external destinations.

8. The method of claim 1, wherein the proxy service implements: a packet decapsulation filter; an authorization filter for policy enforcement; and a dynamic forward upstream cluster for routing resolved DNS requests.

9. A system comprising: one or more hardware processors of a machine; and at least one memory storing instructions that, when executed by the one or more hardware processors, cause the system to perform operations comprising: receiving, by one or more hardware processors, a request to execute a user-defined function (UDF) within a sandbox process; establishing a secure egress path for the UDF using an overlay network, wherein the overlay network includes a dedicated DNS resolver at a proxy service; receiving, from the UDF, a DNS request to resolve a hostname; validating, by the proxy service, that the hostname is included in an allowed host list associated with the UDF; resolving, by the dedicated DNS resolver, the hostname to an IP address using a UDP listener configured to handle DNS protocol traffic on a designated port of the proxy service; and enabling the UDF to communicate with a host at the resolved IP address via the secure egress path.

10. The system of claim 9, wherein establishing the secure egress path comprises: identifying a virtual ethernet pair to isolate UDF traffic within the sandbox process; using the virtual ethernet pair to put an ethernet device in a namespace of the sandbox process; and implementing packet encapsulation between an execution platform and the proxy service.

11. The system of claim 9, wherein the operations further comprise: deploying, on the proxy service, an extended Berkeley Packet Filter (eBPF) to decapsulate DNS packets; extracting a policy identifier from the decapsulated DNS packets; and validating the policy identifier against stored egress policies.

12. The system of claim 9, wherein the overlay network comprises, a first tier implementing virtual ethernet pairs for sandbox isolation, a second tier providing Generic Network Virtualization Encapsulation (GENEVE) tunneling between an execution platform and the proxy service, and a third tier handling verified packet routing to external destinations.

13. The system of claim 9, wherein the operations further comprise: collecting metric event information associated with DNS resolution attempts; logging unauthorized hostname resolution attempts; and publishing metrics to an administrator interface.

14. The system of claim 9, wherein validating the hostname comprises: performing a lookup in a policy map using a sandbox identifier; verifying DNS patterns against wildcard-based allowlist entries; and reporting policy violations to a security monitoring system.

15. The system of claim 9, wherein the operations further comprise: maintaining a connection map at the proxy service to track authorized DNS resolutions; performing Source Network Address Translation (SNAT) on validated packets; and routing translated packets to authorized external destinations.

16. The system of claim 9, wherein the proxy service implements: a packet decapsulation filter; an authorization filter for policy enforcement; and a dynamic forward upstream cluster for routing resolved DNS requests.

17. A machine-storage medium embodying instructions that, when executed by a machine, cause the machine to perform operations comprising: receiving, by one or more hardware processors, a request to execute a user-defined function (UDF) within a sandbox process; establishing a secure

egress path for the UDF using an overlay network, wherein the overlay network includes a dedicated DNS resolver at a proxy service; receiving, from the UDF, a DNS request to resolve a hostname; validating, by the proxy service, that the hostname is included in an allowed host list associated with the UDF; resolving, by the dedicated DNS resolver, the hostname to an IP address using a UDP listener configured to handle DNS protocol traffic on a designated port of the proxy service; and enabling the UDF to communicate with a host at the resolved IP address via the secure egress path.

18. The machine-storage medium of claim 17, wherein establishing the secure egress path comprises: identifying a virtual ethernet pair to isolate UDF traffic within the sandbox process; using the virtual ethernet pair to put an ethernet device in a namespace of the sandbox process; and implementing packet encapsulation between an execution platform and the proxy service.

19. The machine-storage medium of claim 17, wherein the operations further comprise: deploying, on the proxy service, an extended Berkeley Packet Filter (eBPF) to decapsulate DNS packets; extracting a policy identifier from the decapsulated DNS packets; and validating the policy identifier against stored egress policies.

20. The machine-storage medium of claim 17, wherein the overlay network comprises, a first tier implementing virtual ethernet pairs for sandbox isolation, a second tier providing Generic Network Virtualization Encapsulation (GENEVE) tunneling between an execution platform and the proxy service, and a third tier handling verified packet routing to external destinations.
