



US 20250284476A1

(19) **United States**

(12) **Patent Application Publication**
Atur et al.

(10) **Pub. No.: US 2025/0284476 A1**

(43) **Pub. Date: Sep. 11, 2025**

(54) **DISK IMAGE DUMP FOR CONFIGURING
BARE METAL SERVERS**

Publication Classification

(51) **Int. Cl.**
G06F 8/61 (2018.01)

(52) **U.S. Cl.**
CPC **G06F 8/63** (2013.01)

(57) **ABSTRACT**

Systems and methods for capturing a disk image of a bare metal server, and then using the disk image to provision other bare metal servers. A method includes capturing a disk image of a first bare metal server and writing the disk image to a repository manager. The method includes launching a continuous delivery mode on a second bare metal server where the disk image is fetched from repository manager and written to second bare metal server, followed by configuring the unique identity.

(71) Applicant: **Rakuten Symphony, Inc.**, Tokyo (JP)

(72) Inventors: **Sree Nandan Atur**, Newaark, CA (US);
Ravi Kumar Alluboyina, Santa Clara,
CA (US)


(21) Appl. No.: **18/251,649**

(22) PCT Filed: **Nov. 29, 2022**

(86) PCT No.: **PCT/US2022/051203**

§ 371 (c)(1),

(2) Date: **May 3, 2023**






FIG. 1




FIG. 2




FIG. 3




FIG. 4




FIG. 5




FIG. 6




FIG. 7




FIG. 8




FIG. 9




FIG. 10




FIG. 11





FIG. 12






FIG. 14

DISK IMAGE DUMP FOR CONFIGURING BARE METAL SERVERS

TECHNICAL FIELD

[0001] This disclosure relates generally to configuring network systems and specifically relates to configuring a bare metal server with a disk image.

SUMMARY

[0002] Systems and methods for capturing a disk image of a bare metal server, and then using the disk image to provision other bare metal servers. A method includes capturing a disk image of a first bare metal server and writing the disk image to a repository manager. The method includes launching a continuous delivery mode on a second bare metal server where the disk image is fetched from repository manager and written to second bare metal server, followed by configuring the unique identity.

BACKGROUND

[0003] Numerous industries benefit from and rely upon cloud-based computing resources to store data, access data, and run applications and tasks based on the stored data. These systems often rely on storage and processing resources of numerous bare metal servers. In traditional systems, it is computationally expensive and time consuming to install an operating system on a new bare metal server and configure the new bare metal server for use. Additionally, these traditional systems must endure server downtime when servers are being provisioned and configured.

[0004] In view of the foregoing, disclosed herein are systems, methods, and devices for decreasing server downtime and reducing configuration time for new bare metal servers by configuring servers with a disk image.

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] In order that the advantages of the invention will be readily understood, a more particular description of the invention briefly described above will be rendered by reference to specific embodiments illustrated in the appended drawings. Understanding that these drawings depict only typical embodiments of the invention and are not therefore to be considered limiting of its scope, the invention will be described and explained with additional specificity and detail through use of the accompanying drawings, in which:

[0006] FIG. 1 is a schematic illustration of a network environment in which the systems and

[0007] methods described herein may be used;

[0008] FIG. 2 is a schematic block diagram of a system comprising layers for external management of a network service;

[0009] FIG. 3 is a schematic diagram of an element of a network service;

[0010] FIG. 4 is a schematic block diagram of a hierarchy for orchestrating a network service;

[0011] FIG. 5 is a schematic block diagram of a system for executing a deployment automation module of an orchestration server system;

[0012] FIG. 6 is a schematic block diagram of an interface that may be provided by a workflow builder to facilitate creating workflows;

[0013] FIG. 7 is a process flow diagram of a method for batch processing functions for large numbers of elements;

[0014] FIG. 8 is a schematic block diagram of a system for distributing files and executing functions of a workflow;

[0015] FIG. 9 is a schematic block diagram of a system that may be used to implement the functionality of a machine initialization module of an orchestration server system;

[0016] FIG. 10 is a schematic flow chart diagram of a method for initializing a bare metal server;

[0017] FIG. 11 is a schematic block diagram of a system for capturing a disk image of bare metal servers that may be used to install operating systems on other bare metal servers with a copy and paste functionality;

[0018] FIG. 12 is a schematic process flow diagram of a method for capturing a disk image of a first bare metal server;

[0019] FIG. 13 is a schematic process flow diagram of a method for configuring a second bare metal server with a disk image of a first bare metal server; and

[0020] FIG. 14 is a schematic block diagram of an example computing device suitable for implementing methods in accordance with embodiments of the invention.

DETAILED DESCRIPTION

[0021] Disclosed herein are systems, methods, and devices for capturing a disk image of a bare metal server, and then using the captured disk image to install operating systems on other bare metal servers. The methods described herein may be implemented with a copy/paste functionality to obtain high performance while keeping server reboots to a minimum.

[0022] In traditional systems, it is computationally expensive and time consuming to install an operating system on a bare metal server and then configure the bare metal server for use. Additionally, this process includes significant downtime when the bare metal server cannot be used to execute applications and workflows. Additionally, typical operating system provisioning kickstarter files are complex and include numerous components, including a pre-section, operating system installs, post-selection configuration, and one-time updated services. These kickstarter files can present numerous challenges when seeking to configure numerous bare metal servers for immediate use. In view of the foregoing, the systems, methods, and devices described herein enable users to initiate a workflow to capture a disk image of a first bare metal server, and then use the captured disk image to install and operating system on other bare metal servers.

[0023] Referring now to the figures, FIG. 1 is a schematic illustration of a network environment 100 in which the systems and methods disclosed herein may be used. The network environment 100 includes a plurality of baseboard units (BBUs) 106 and edge servers 108 in communication with various antennas 104. The edge servers 108 comprise processing and memory resources configured to execute an orchestration server system 102. The orchestration server system 102 includes an enterprise management service 110, operations support system 112, management server 114, and a deployment automation module 116.

[0024] The antennas 104 are coupled to the baseband units 106. The baseboard units 106 provide translation between radio frequency signals received from the antennas 104 and digital data received from the edge servers 108. Each baseboard unit 106 may perform this translation according to applicable cellular wireless data protocols. The edge

servers **108** are coupled to the orchestration server system **102** either directly or by way of one or more intermediary servers.

[0025] The orchestration server system **102** executes centralized management services used to manage the edge servers **108** and baseboard units **106**. Specifically, the orchestration server system **102** executes enterprise management services **110**, operations support systems (OSS) **112**, and one or more management servers **114** for services implemented on the edge servers **108**. The orchestration server system **102** executes a deployment automation module **116** that facilitates deployment of the baseboard units **106**, the edge servers **108**, and the services executing on the baseboard units **106** and the edge servers **108**.

[0026] The deployment automation module **116** includes a machine initialization module **118** that detects and initializes hardware within the network environment **100**. The hardware may include computing and storage devices for implementing the baseboard units **106** or the edge servers **108**. For example, given a computing device configured with an IP address, the machine initialization module **118** may initialize the BIOS (basic input output system), install an operating system, configure the operating system to connect to a network and to the orchestration server system **102**, and install an agent for facilitating installation of services and for performing management functions on the computing device at the instruction of the deployment automation module **116**. For example, the machine initialization module **118** may use COBBLER in order to initialize the computing device.

[0027] The machine initialization module **118** may also discover computing devices on a network and generate a topology of the devices, such as in the form of a directed acyclic graph (DAG). The deployment automation module **116** may then use this DAG to select computing devices for implementing network services and in order to configure a machine to receive installation of a network service.

[0028] The deployment automation module **116** may include an application automation module **120** that automates the deployment of an application, such as a container executing an application on a computing device. The application automation module **120** may implement methods and systems described below relating to the automated deployment and management of applications.

[0029] One example application of the systems and methods disclosed herein is a radio area network (RAN) automation module **122** that performs the automated deployment of a network service in the illustrated network environment, including instantiating, configuring, and managing services executing on the baseboard units **106**, the edge servers **108**, and the orchestration server system **102** in order to implement a RAN in a one-click automated fashion.

[0030] FIG. 2 is a schematic block diagram of a system **200** comprising layers for external management of a network service. The system **200** includes a physical layer **202**, clustering layer **204**, network function/application layer **206**, and a network service/application pipeline layer **208**.

[0031] The physical layer **202** includes computing hardware and may further include basic software such as BIOS, firmware, an operating system, or even a virtual machine executing on the computing hardware. The system includes the clustering layer **204** residing on the physical layer **202**. The clustering layer **204** includes data structures and software that enables a group of computing devices to act as a cluster. A cluster may be defined as a group of devices that

are backups of one another, that provide a service with requests for that service being distributed among devices of the cluster according to a load balancing approach, that together implement a plurality of distinct applications that cooperate with one another to implement a service, or that are associated to one another for a common purpose or according to an arbitrary cluster definition of an administrator. The clustering layer **204** may be implemented by agent software executing on the physical layer **202** that coordinates with the deployment automation module **116** and other devices of a cluster.

[0032] The network function/application layer **206** includes applications executing on the computing devices of a cluster that individually or together with other applications executing on other nodes of the cluster implement a network service, such as access to a database, web server, or other server-based computational function that may be provided as a service to a user or another service of a network environment **100**. The network service/application pipeline layer **208** may include a pipeline of network functions/applications **206** that communicate with one another to implement a more complex network service.

[0033] The operations of the layers **202-208** is managed by method and procedures (MOPs) **210** that are independent of the services implemented by the layers and include management functions such as instantiating, upgrading, health checks, monitoring power, restarting, replacing, scaling, and shutting down of the entities implementing a layer **202-208** (also referred to as life cycle management (LCM)).

[0034] The system **200** includes a policy engine **212** configured to provide logic that defines actions performed by any of the layers **202-208**. These actions include, for example, procedures for implementing backups, procedures for handling faults at a particular layer, procedures for prioritization of individual MOPs **210**, or other policies that an administrator may wish to impose on the operation of any of the layers **202-208**.

[0035] For example, the policy engine **212** may have access to a topology of an application pipeline created according to the methods disclosed herein. Once a fault is identified, the policy engine **212** may implement an appropriate recovery policy. For example, if a computing device fails, its IP address may be released and assigned to a new computing device. The elements on the failed computing device may be instantiated on a new computing device assigned the IP address. If an element fails, a new element of the same type may be instantiated and configured to take its place. If a pod fails, a new pod may be created and configured to take its place.

[0036] The system **200** includes closed loop automation **214** that is executed in connection with any of the layers **202-208**. The closed loop automation **214** may specifically include the performance of tasks, verification of performance of tasks, monitoring function, automated actions performed in response to states detected during monitoring, and other actions in order to implement some or all of the layers **202-208** and maintain operation of the layers **202-208**.

[0037] FIG. 3 is a schematic diagram of an element **300** of a network service in accordance with an embodiment of the present invention. Each entity that constitutes one of the layers **202-208** may be embodied as an element **300**. Each element **300** defines functions and interfaces used by the deployment automation module **116** to deploy and manage

an entity represented by an element 300. An element 300 may be an entity that is a combination of sub-elements and defines functions and interfaces for deploying and managing the combination of sub-elements. Accordingly, the deployment automation module 116 may invoke these interfaces and functions in order to deploy and manage an element 300 without requiring any modification of the deployment automation module 116 to adapt to or have data describing the entity represented by the element 300.

[0038] For example, an element 300 may define functions and interfaces for discovering 302 the element 300 such that once the element 300 is connected by a network to the deployment automation module 116, the element 300 may be discovered and its identity, type, and other attributes may be provided to the deployment automation module 116.

[0039] The element 300 may define functions and interfaces for maintaining a reference to the element 300 in an inventory 304 of elements 300 maintained by the deployment automation module 116. This may include responding to queries from the deployment automation module 116 with responses indicating availability of the element 300, e.g., whether it is assigned and operational.

[0040] The element 300 may define functions and interfaces for performing life cycle management (LCM) 306 of the element 300. This may include functions and interfaces for instantiating, upgrading, scaling, restarting, or de-instantiating the element 300.

[0041] The element 300 may define functions and interfaces for performing healing 308 on the element 300. This may include functions and interfaces for detecting faults, recovering from faults, restoring non-functioning parts of the element 300, or other actions for restoring or maintaining function of the element 300.

[0042] The element 300 may define functions and interfaces for monitoring health 310 of the element 300. This may include functions and interfaces for running diagnostic checks, performance checks, or other evaluations of the state of operation of the element 300.

[0043] The element 300 may define functions and interfaces for implementing policy 312 with respect to the element 300. This may include functions and interfaces for receiving a policy 312 for the element 300 and evaluating the policy 312 with respect to a current state of operation of the element 300. The functions and interfaces may define the policies themselves or may be configured to receive and implement policies received from the deployment automation module 116.

[0044] FIG. 4 is a schematic block diagram of a hierarchy 400 for orchestrating a network service in accordance with an embodiment of the present invention. The deployment automation module 116 may ingest a workflow 402. The workflow 402 defines a series of functions 404 and possibly an order of execution of the functions 404. The functions 404 may invoke executors 406 that operate with respect to an element 300. In particular, the functions 404 may be functions of instances of elements 300 included in the workflow 402. Accordingly, a workflow 402 may define performance of functions 404 for specific elements 300 and possibly the ordering of these functions.

[0045] The elements 300 may be entities implementing a network service pipeline, radio area network (RAN), or any complex multi-application deployment and the workflow invokes the functions of these elements 300. As noted above, due to the abstraction of the elements 300, the workflow

does not need to specify entity-specific functions. Instead, tasks of a workflow 402 including discovery, inventory management, life cycle management, health monitoring, healing, policy implementation and other high-level functions may be invoked by invoking corresponding interfaces and functions 302-312 of the elements 300 of the workflow 402.

[0046] FIG. 5 is a schematic block diagram of a system for executing the deployment automation module 116. The deployment automation module 116 includes a distributed log store 526, a distributed file store 522, and a database 530 in communication with a representation state transfer (REST) application program interface (API) 502. The REST API 502 includes components and software for executing a workflow builder 504, batch runner 508, closed loop automation 512, workflow orchestrator 506, batch orchestrator 510, playground 514, FSM manager 534, and other 536 functionalities.

[0047] The workflow builder 504 defines an interface enabling a user to create, select, and modify workflows 402. The workflow orchestrator 506 implements the functions 404 of a workflow 402 on the elements 300 of a workflow 402. In many instances, the number of elements 300 and the functions 404 that are performed with respect to elements 300 is very large, on the order of 1000s or even millions. Accordingly, the batch runner 508 sets up batch processing of functions 404 for elements 300 and a batch orchestrator 510 then executes the functions 404 in batches as defined.

[0048] The REST APIs 502 define closed loop automation 512 APIs that implement closed loop automation 214 of the deployment and management of the elements 300 of a workflow 402 according to the interfaces of the elements 300. The playground 514 provides a testbed for the creation and evaluation of elements 300, workflows 402, and functions 404.

[0049] The functions 404 of the workflows 402, either individually or as parts of batches, may be implemented by workers 516. The workers 516 may be embodied as pods, such as pods according to the KUBERNETES orchestration platform. Alternatively, workers 516 may be processes or threads of execution executing on one or more computing devices of a network environment 100. For example, the workers 516 may execute on clusters 518, a rack server 520, edge server 108, baseboard unit 106, or some other computing device.

[0050] The number of files required to define the functions 404 and elements 300 of a workflow 402 may be very large. Accordingly, a distributed file store 522 may be implemented, such as in the form of a database accessed by means of a function registry 524 that maps a function 404 of an element 300 (e.g., a function identifier associated with an element identifier) to a storage location in the distributed file store 522.

[0051] In a like manner, the number of files and amount of data generated by the functions 404 and applications instantiated by a workflow 402 may be very large. Accordingly, a distributed database of log store to which functions 404 and applications instantiated by a workflow 402 may write updates too, such as by means of a log plugin 528.

[0052] Other data used by the REST APIs 502 may be stored in a database 530 accessed by means of a database plugin 532. For example, interfaces, templates, pre-defined

workflows **402**, elements **300**, and other data that may be used by a user to define a workflow **404**.

[0053] In some embodiments, each element **300** may have a state and a corresponding finite state machine that defines transitions between states of the finite state machine in response to events occurring involving the element **300**. Accordingly, the REST APIs **502** may include a finite state machine (FSM) manager **534** for managing the state machine of each instance of any of the elements **300**.

[0054] Other REST APIs **502** may implement other functions, such as observability of elements (OBF), rule-based access control, cluster federation, and other functions that may facilitate implementation and management of a network service pipeline.

[0055] FIG. 6 is a schematic block diagram of an interface **600** that may be provided by the workflow builder **504** to facilitate creating workflows **402** in accordance with an embodiment of the present invention. The interface **600** may include an element menu **602** that enables a user to input a selection of an element **300** from a list of available elements **300**. Elements **300** may include a virtual machine, a container, a database (e.g., MONGODB), an application, a router, a switch, a rack switch, relay, or any other element that may be needed to implement a network service. The interface may further include a function menu **604** that enables a user to input a selection of an element **300**, e.g., the element selected using the element menu **602**. This may include any of the interfaces and functions described above. For example, where a workflow **402** is to be created that instantiates a network pipeline, the functions selected from the function menu **604** may be functions to instantiate the selected element. For example, an element/function (i.e., a selected function for a selected element type) **608a** may define instantiating a primary manager of a cluster, element/function **608b** may define instantiating a secondary manager of the cluster, element/functions **608c-608e** may define instantiating one or more other nodes of the cluster. Other functions for a cluster may include acquiring licenses for software, performing network configuration of the managers and nodes of the cluster, acquiring IP addresses for the cluster and nodes of the cluster, setting up bundles (e.g., bundled applications), and setting up external backup depositories.

[0056] Each element/function **608-608e** input by a user may be represented by an icon on the graphical user interface (GUI) **600**, such as shown in FIG. 6. Each element function **608a-608e** may have configuration parameters such as internet protocol (IP) address, identifier, number of processing cores, amount of memory, amount of storage, etc., to be allocated to the node instantiated by the function **608a-608e**. These parameters may be specified by default or may be input by a user, such as by accessing a menu permitting their input by clicking on a representation of a function **608a-608e** in the interface **600**.

[0057] In some embodiments, predefined workflows **402** may be selected from a workflow menu **606**. A user may then modify the workflow **402**. For example, a workflow selected from the workflow menu **606** or created by a user may be modified to include additional element/functions **608f-608g**.

[0058] FIG. 7 is a process flow diagram of a method **700** for batch processing functions for large numbers of elements in accordance with an embodiment of the present invention. The method **700** may be performed by the deployment automation module (DAM) **116**, such as using the workflow

orchestrator (WFO) **506** and batch orchestrator (BO) **510**. Various other entities are involved in the method **700**, including a workflow (WF) **402**, the database (DB) **530**, a spawning manager (SM) **706**, workers **516**, distributed file stores (FS) **522**, and a plurality of target elements (TE) **300**.

[0059] The method **700** includes receiving at **710** an instruction to perform a function with respect to N elements **300**. In the illustrated example, this function is upgrading, though any function ascribed herein to an element **300** may also be performed. In a typical application, N is very large, on the order of 1000s, 10,000s, or millions. The instruction received at **710** may be received from a user or received as part of processing a workflow **402**.

[0060] The workflow orchestrator **506** receives the instruction at **710** and, in response, may calculate at **712** the fanout. This may include determining how many of the target elements **300** will be processed according to the function by a worker **516**. The fanout may be static for all types of elements **300**, defined for a particular type of element **300**, defined for a particular function **302-312**, defined for a particular function **302-312** of a particular type of element **300**, or be determined based on some other criteria, which may be dynamic, such as a function of the value of N or current loading of workers **516** of the deployment automation module **116**.

[0061] The batch orchestrator **510** may return at **714** a worker count W that is a number of workers **516** that are available to perform the function with respect to the N target elements **300**. The workflow orchestrator **506** may then divide the N target elements **300** into shards such that each shard has approximately (e.g., +/-10) N/W elements **300** assigned to it. Each shard may include element identifiers of the target elements **300** assigned to it and may itself be assigned a shard identifier. The shards may be stored at **716**, such as in the database **530**.

[0062] The workflow orchestrator **506** may then invoke at **718** the creation of W workers. For example, a spawning module **706** may be programmed to generate workers **516** in response to receiving the instruction from step **718**. Upon instantiation, the workers may each request **720** a shard from the workflow orchestrator **506**, which may then return at **722** a shard configuration array, e.g., an array of target element identifiers along with an identifier of the function to be performed with respect to the target elements **300** referenced by the target element identifiers.

[0063] The worker **516** may then request at **724** the function, e.g., a script or executable, corresponding to the function identifier received at step **722**, from the distributed file store **522**. The worker **516** then receives at **726** the function and executes at **728** the function on each of the target elements **300** reference in the shard configuration array received at step **722**. Upon completion of execution of the function with respect to each target element **300** referenced by the shard, the worker **516** reports **730** completion to the workflow orchestrator **506**. When all workers **516** complete processing of their shards, the instruction received at step **702** may be complete.

[0064] FIG. 8 is a schematic block diagram of a system **800** for distributing files and executing functions of a workflow. The system **800** includes a smart router **802** coupled to various local distributors **804**. The local distributors **804** may be embodied as applications executing within pods, e.g., KUBERNETES pods, executing throughout a network environment **100**. The distributors **804** may host or

access a local database **806**. The local database **806** may be a copy of the distributed file store **522** or a portion thereof. For example, given the elements instances in proximity to the local distributor **804**, the portion of the distributed file store **522** may include data from the distributed file store **522** relating to those elements, e.g., executables and data for performing the functions of those element instances. Proximity to the local distributor **804** may mean located in the same sub-network or having a network connection to the local distributor **804** having latency below a threshold.

[0065] The workers **516** may request data from the distributed file store **522**. These requests may be received by the smart routers **802**, which identifies the local distributor **804** that is either having the lowest network latency connection to the requesting worker **516** or is more available (i.e., lower latency due to lower loading) to distribute files than the local distributor **804** with lowest network latency. For example, the smart router **802** may include a load balancer **804** as described above with respect to FIG. **8** such that the local distributor **804** is selected according to network latency and loading as described above. The request is then routed by the smart router **802** to the selected local distributor **804**, which then provides the requested data to the worker **516** that generated the request.

[0066] FIG. **9** illustrates a system **900** that may be used to implement the functionality of the machine initialization module **118**. The machine initialization module **118** may operate with respect to bare metal servers **902** that are “bare metal,” i.e., have no operating system, kernel, or other software installed thereon other than firmware stored in non-volatile RAM on the device. This firmware will include a basic input output system (BIOS) as well as firmware on components of the bare metal server **902** such as a network adapter (e.g., network interface card (NIC)), hard disk drive (HDD), solid state drive (SSD), redundant array of independent disks (RAID), just a bunch of disks (JBOD), field programmable gate array (FPGA), baseboard management controller (BMC), Non-Volatile Memory Express (NVME) controller, or other component of the bare metal server **902**. Although the foregoing description refers to a bare metal server **902**, any computing device, such as a router, switch, endpoint (personal workstation, mobile computing device, internet of things (IoT) device, etc.), or any other computing device that may communicate over a network.

[0067] The machine initialization module **118** itself may be structured as an application that may execute on a node of a cluster **518**. The machine initialization module **118** may operate on the same cluster **518** or a different cluster from a cluster hosting the workflow orchestrator **506** and one or more workers **516** implementing functions of a workflow being managed by the workflow orchestrator **506** according to the methods described herein.

[0068] The machine initialization module **118** accesses the distributed file store **522** to obtain images **904** of operating systems and other executables to be instantiated on a bare metal server **902**. The distributed file store **522** may also store artifacts **906** that are likewise executables or other data that are used by the machine initialization module **118** to initialize a bare metal bare metal server **902**.

[0069] FIG. **10** illustrates a method **1000** for initializing a bare metal server **902**. The method **1000** may begin with installing **1002** of a kickstarter executable on the bare metal server **902**. The kickstarter may correspond to the configuration of the bare metal server **902**. The configuration of the

bare metal server **902** may be represented using a JAVASCRIPT Object Notation (JSON) file that describes the hardware, firmware, and/or software versions of the bare metal server **902**. The JSON file may further include links to a kickstarter file that corresponds to the needs of an application to be installed on the server system **1302**, that corresponds to the SKU of the bare metal server **902** or is configured based on some other criteria. For example, there may be a kickstarter associated with each SKU (stock keeping unit) defining a type of bare metal server **902**. Accordingly, the kickstarter installed at step **1002** may be that which corresponds to the SKU of the bare metal server **902**. The kickstarter may include a profile of the bare metal server **902**, such as according to the Basic, EPA-1, EPA1-test, and/or EPA2 system profile types.

[0070] The kickstarter may include a configuration file that configures the bare metal server **902** to register with the machine initialization module **118**. Since the bare metal server **902** is not configured with an operating system or an IP (internet protocol) address, the kickstarter may include computer instructions that instruct the bare metal server **902** to communicate with the machine initialization module (MIM) **118** using the baseboard management controller (BMC) IP address with which the bare metal server **902** was configured by a manufacturer. The kickstarter may include an IP address for the machine initialization module **118** or that of some other component that is programmed to route communications from a kickstarter to the machine initialization module **118**. Alternatively, the request to register may be broadcast and detected by a component in a network environment that routes the request to the machine initialization module **118**. Installing **1002** of the kickstarter may be performed manually by a human operator or by a component coupled to a network to which the bare metal server **902** is connected when installed in a rack, datacenter, or other facility.

[0071] The bare metal server **902** executes the kickstarter, which causes the bare metal server **902** to register **1004** with the machine initialization module **118** by communicating over the network to the IP address included in the kickstarter. Registering may include providing the BMC IP address of the bare metal server **902** to which the machine initialization module **118** may address subsequent communications.

[0072] The machine initialization module **118** may obtain **1006** an IP address (“the server IP address”) to assign to the bare metal server **902** and generate **1008** an extensible firmware interface (EFI) image including the IP address. The IP address may be assigned at step **1006** according to a workflow **402**. For example, if the bare metal server **902** is (or hosts) an element instance created according to a function **404** workflow **402**, the parameters of the function **404** may include a statically or dynamically assigned IP address for the bare metal server **902**. Alternatively, the IP address may be assigned according to an IP address management (IPAM) algorithm executed by the machine initialization module **118**, workflow orchestrator **506**, or another component in a network environment. In particular, the method **1000** may be executed independently from the workflow orchestration approaches described herein such that the IP address is obtained according to an IPAM algorithm according to any approach known in the art.

[0073] The machine initialization module **118** may generate **1008** an executable file including the IP address. In some embodiments, the executable file may be an extensible

firmware interface (EFI) image. The executable file may be generated according to the workflow used to select the IP address. The executable file may further include network information such as an IP address for a network gateway to be used by the bare metal server **902**, e.g., a node in a network domain including the IP address assigned to the bare metal server **902**. The executable file may further contain instructions for configuring the bare metal server **902** to connect to a virtual local area network (VLAN).

[0074] In some embodiments, the EFI image may include executable code instructing the bare metal server **902** to retrieve and install an operating system kernel from a specified IP address. The EFI image itself may be configured as a bootstrap kernel from which the server system **1302** may boot itself up. The EFI image may include executable code instructing the bare metal server **902** to retrieve and execute firmware upgrade files for the BIOS, network adapter, HDD, SSD, BMC, BIOS, NIC, RAID, JBOD, NVME controller, FPGA, or other component of the bare metal server **902**. Upgrading of firmware or other operations instructed by the EFI image may further include flashing custom images on any of these components or otherwise configuring these components, such as a RAID or JBOD. The EFI image may include executable code instructing the bare metal server **902** to retrieve operating system files for installing an operating system on the bare metal server **902**. The EFI image may be formatted as an ISO (International Organization for Standardization) image that can be mounted as a disk to be booted up from on the bare metal server **902**. The EFI image is preferably small, such as less than 3 MB. For example, an ISO file size of 2.12 MB has been found to be achievable.

[0075] In some embodiments, the EFI image may be obtained from a boot configuration file including the above-described instructions to configure the server IP address, network gateway, and retrieve and install the operating system kernel. The boot configuration file may further include instructions to connect to a virtual local area network (VLAN). The boot configuration file may be written in IPXE (an open-source implementation of the Preboot Execution Environment client firmware and bootloader) scripting language and using IPXE syntax. This IPXE scripting language may be compiled using IPXE source code to obtain a bootable EFI image that packs the information of the boot configuration file in a form that can be executed by an IPXE bootloader on the bare metal server **902** in either legacy BIOS or EFI mode.

[0076] The IPXE bootloader is typically a small kernel that includes drivers for the hardware of the bare metal server **902** and has the ability to configure new hardware of different types including networking, storage, and the like. In the illustrated embodiment, the ability of the IPXE bootloader to configure a network interface is used to configure the server IP address and network gateway of the bare metal server **902** and may also be used to configure the bare metal server **902** to connect to a VLAN.

[0077] The EFI image may be converted into a bootable ISO file. The BMC of the bare metal server **902** may be capable of mounting an ISO file either through an API (application programming interface) call or manual intervention. In some embodiments, a boot order on the bare metal server **902** may be modified such that the bare metal server **902** boots from the bootable ISO file including the

EFI image. For example, the kickstarter may be programmed to modify the boot order in this manner.

[0078] The bootable ISO file may include both the EFI image and a bootloader, such as the “isolinux.bin” bootloader. The bootloader may contain the encoded form of the configuration file that will be executed on the server during the boot load process where the bootloader successively attempts to configure each interface according to the EFI image (including the network interface as described above) and tries to retrieve the operating system kernel according to instructions in the EFI image. Once the bootloader successfully retrieves the operating system kernel, it uses this interface to install the rest of the OS, as described below with respect to FIG. 15.

[0079] The bootloader, such as isolinux.bin, may be added to the ISO file including the EFI image to perform bootloading of the hardware of the bare metal server **902**. The EFI image (e.g., an ipexe.efi file) interacts with the EFI BIOS to do an initial boot, recognize EFI capabilities, and present the EFI capabilities to the kernel for a Stage 2 booting of the kernel in EFI mode. This EFI image may be placed in the distributed file store **522** where it is accessible via an HTTP (hypertext transport protocol) server (or an HTTP secure (HTTPS) server).

[0080] The machine initialization module **118** transmits **1010** the EFI image (e.g., ISO file including the EFI image) to the bare metal server **902**. The bare metal server **902** receives the EFI image and executes **1012** it. This may include mounting the ISO image and executing the bootloader in the ISO image. The bootloader processes the EFI image to configure the network interface of the bare metal server **902** and retrieve and install an operating system kernel as described above. In some embodiments, the EFI image may be executed by a VMCLI (virtual machine command line interface) utility on the bare metal server **902**.

[0081] As a result of executing the EFI, the bare metal server **902** is configured with an IP address for itself, an IP address of a network gateway to be used by the bare metal server **902**, an operating system kernel, and with instructions to download an operating system from a specified source IP address, such as that of the distributed file store **522**. In some embodiments, the EFI image includes instructions causing the bootloader to incrementally retrieve **1014** the operating system. For example, instead of having to retrieve a 2 GB ISO file including an operating system image, the EFI image may include instructions to download smaller installation packages implementing installation of the operating system in order to reduce loading of the distributed file store **522**.

[0082] FIG. 11 is a schematic block diagram of a system **1100** for capturing a disk image of bare metal servers that can be used to install operating systems on other bare metal servers with a copy and paste functionality. The system **1100** includes a multi-data center automation platform (MDCAP) **1104** in communication with a user **1102**. The MDCAP **1104** includes an engine **1106** and a dashboard **1106**. The MDCAP **1104** renders the dashboard on a user interface **1106** accessible by the user **1102**. The MDCAP **1104** includes or communicates with a plurality of workers **516**. The MDCAP **1104** includes or accesses a repository manager **1110** that manages binary resources for the MDCAP **1104**.

[0083] The repository manager **1110** serves as a central hub for integrating with tools and processes to improve automation of the system **1100** and increase system **1100**

integrity. In some implementations, the repository manager **1110** is implemented as an ARTIFACTORY. The repository manager **1110** organizes binary resources, including, for example, remote artifacts, proprietary libraries, third-party resources, and so forth. The repository manager **1100** pulls these resources into a single centralized location for a plurality of bare metal servers **902a-902c**.

[0084] The repository manager **1100** manages and automates artifacts and binaries from start to finish during the application delivery process. The repository manager **1100** enables the option to select from different software build packages, major CI/CD (continuous integration/continuous development) systems, and other development tools. The repository manager **1100** may be implemented within a KUBERNETES containerized system with a DOCKER registry with full REST APIs **502** as discussed herein. The repository manager **1100** supports containers, Helm charts, and DOCKER.

[0085] FIGS. **12** and **13** are methods for capturing a disk image of a bare metal server, and then using the captured disk image to install an operating system on one or more other bare metal servers. The methods **1200**, **1300** are similar to a copy/paste functionality to obtain high performance while keeping server reboots to a minimum. FIG. **12** is a schematic block diagram of a method **1200** for a “copy phase” wherein the disk image of a first bare metal server **902a** is copied and stored. FIG. **13** is a schematic block diagram of a method **1300** for a “paste and configuration” phase wherein the previously copied disk image of the first bare metal server **902** is installed on a second bare metal server **902b**. The methods **1200**, **1300** are executed by a user **1102**, the MDCAP **1104**, and a plurality of bare metal servers **902**.

[0086] The disk image consists of a computer file containing the contents and structure of a disk volume or an entire data storage device wherein the operating system is installed. The disk image is typically made by creating a sector-by-sector copy of the source medium, thereby perfectly replicating the structure and contents of the bare metal server **902** independent of the file system. Depending on the disk image format, the disk image may span one or more computer files. The file format for the disk image may include an open standard format, such as the ISO image format for optical disc images, or a disk image may be unique to a particular software application.

[0087] The method **1200** begins and the user **1102** initiates at **1202** a workflow to capture a disk image of a first bare metal server **902a** (BMS). The user **1102** initiates the workflow by interacting with the dashboard **1106** of the MDCAP **1104** that is rendered on a user interface. In some cases, and irrespective of disk size, the disk image is directed only on the partition where the operating system is installed on the first server. The method **1200** continues and the MDCAP **1104** submits the workflow to the workers **516** at **1204**.

[0088] The method **1200** continues and the workers **516** of the MDCAP **1104** bring up a live CD (continuous development) mode on the first bare metal server **902a**. In an implementation, the workers **516** do this through iPXE, which is an open-source implementation of the Preboot execution environment client software and bootloader. Thus, the workers **516** accomplish step **1206** without use of a DHCP (dynamic host configuration protocol) server as discussed herein (see, for example, the discussion in connection

with FIG. **10**). Step **1206** continues and the workers **516** capture a disk image of the first bare metal server **902a**. The workers **516** then write the disk image to the repository manager **1110** of the MDCAP **1104**. Upon being written to the repository manager **1110**, the disk image is now accessible over a URL. The method **1200** continues and the MDCAP **1104** writes the disk image to the repository manager **1110** based on its type. The MDCAP **1104** then notifies the user **1102** that the workflow was a success.

[0089] The “copy phase” method **1200** continues on with the “past and configure” phase of the method **1300** illustrated in FIG. **13**. The method **1300** begins and the user **1102** initiates a workflow at **1302** to install an operating system on the second bare metal server **902b** by way of a disk image dump. The MDCAP **1104** receives the request from the user by way of the dashboard **1106** that is rendered on a user interface. The user **1102** may initiate a workflow to install the disk image of the first bare metal server **902a** onto a plurality of other bare metal servers, including second bare metal server **902b** identified in FIG. **13**. The user **1102** may request that the disk image of the first bare metal server **902a** be installed on a plurality of other bare metal servers simultaneously. The MDCAP **1104** submits the workflow to the workers **516** at **1304**.

[0090] The method **1300** continues and the workers **516** of the MDCAP **1104** bring up a live CD (continuous development) mode on the second bare metal server **902b** at **1306**. The workers **516** may accomplish this through iPXE, without the use of a DHCP server, as discussed herein (see, for example, the discussion in connection with FIG. **10**). The MDCAP **1104** retrieves the disk image of the first bare metal server **902a** from the repository manager **1110** at **1308**. The MDCAP **1104** may retrieve the disk image at **1308** by downloading the disk image from the repository manager through an available URL. While the data for the disk image is being downloaded, the disk image is written directly to a disk on the MDCAP **1104**. The MDCAP **1104** writes the disk image in a pre-section of a kickstarter file.

[0091] The method **1300** continues and the MDCAP **1104** then mounts the disk image of the first bare metal server **902a** to the second bare metal server **902b** at **1310**. The MDCAP **1104** then configures the unique identity of the second bare metal server **902b** at **1312**. The configuration of the unique identity includes configuring, for example, the IP address, gateway, DNS (domain name system), and host-name for the second bare metal server **902b**. The method **1300** continues and the MDCAP **1104** notifies the user **1102** that the workflow was a success at **1314**.

[0092] The methods **1200**, **1300** discussed in connection with FIGS. **12** and **13** are implemented to capture disk and partition dumps of an operating system, and then rewrite the operating system onto another bare metal server. The MDCAP **1104** provides a means to capture a disk image of a bare metal server **902** such that the captured disk image may be used to install an operating system on another bare metal server **902**. The methods **1200**, **1300** are similar to a copy/paste functionality and may be implemented to achieve high performance while keeping server reboots to a minimum.

[0093] The MDCAP **1104** is leveraged to capture a disk image and partition dump of an operating system, and then rewrite the captured disk image to one or more other bare metal servers. The captured operating system may be written with specific changes, including alterations to identifying

information such as the IP address, hostname, gateway, DNS, and so forth. This is particularly useful in complex DU far edge use cases when it is important to keep server reboots to a minimum.

[0094] The methods **1200**, **1300** significantly reduce the time required to install an operating system on a bare metal server **902** and then configure the bare metal server **902** for use. Specifically, in some implementations, the installation and configuration time according to method **1300** is about one-third the typical time required for DU far edge use cases. Further, the installation and configuration time according to method **1300** is about one-half the typical time required for typical use cases. The methods **1200**, **1300** support any operating system irrespective if its support kickstart technology. For example, the methods **1200**, **1300** can support WINDOWS, UBUNTO, MAC OS, RHEL, CENTOS, FEDORA, LINUX, and so forth.

[0095] In an example traditional system, a typical operating system provisioning for a DU far edge use case takes about 30 minutes to complete. With the technique described in the method **1300**, the bare metal server may be provisioned in 10 minutes or less. Further, a typical operating system provisioning takes about 20 minutes with traditional techniques. When the typical bare metal server is provisioned according to the method **1300**, the time is reduced to about 10 minutes.

[0096] The systems described herein offer increased simplicity. A typical operating system provisioning kickstarter file is complex and may include each of a pre-section, an operating system installs, a post-section configuration, and one-time update services. With the systems and methods described herein, the operating system provisioning kickstarter file is simplified to include only the pre-section.

[0097] FIG. 14 illustrates a schematic block diagram of an example computing device **1400**. The computing device **1400** may be used to perform various procedures, such as those discussed herein. The computing device **1400** can perform various monitoring functions as discussed herein, and can execute one or more application programs, such as the application programs or functionality described herein. The computing device **1400** can be any of a wide variety of computing devices, such as a desktop computer, in-dash computer, vehicle control system, a notebook computer, a server computer, a handheld computer, tablet computer and the like.

[0098] The computing device **1400** includes one or more processor(s) **1404**, one or more memory device(s) **1404**, one or more interface(s) **1406**, one or more mass storage device(s) **1408**, one or more Input/output (I/O) device(s) **1410**, and a display device **1430** all of which are coupled to a bus **1412**. Processor(s) **1404** include one or more processors or controllers that execute instructions stored in memory device(s) **1404** and/or mass storage device(s) **1408**. Processor(s) **1404** may also include several types of computer-readable media, such as cache memory.

[0099] Memory device(s) **1404** include various computer-readable media, such as volatile memory (e.g., random access memory (RAM) **1414**) and/or nonvolatile memory (e.g., read-only memory (ROM) **1416**). Memory device(s) **1404** may also include rewritable ROM, such as Flash memory.

[0100] Mass storage device(s) **1408** include various computer readable media, such as magnetic tapes, magnetic disks, optical disks, solid-state memory (e.g., Flash

memory), and so forth. As shown in FIG. 14, a particular mass storage device **1408** is a hard disk drive **1424**. Various drives may also be included in mass storage device(s) **1408** to enable reading from and/or writing to the various computer readable media. Mass storage device(s) **1408** include removable media **1426** and/or non-removable media.

[0101] I/O device(s) **1410** include various devices that allow data and/or other information to be input to or retrieved from computing device **1400**. Example I/O device(s) **1410** include cursor control devices, keyboards, keypads, microphones, monitors or other display devices, speakers, printers, network interface cards, modems, and the like.

[0102] Display device **1430** includes any type of device capable of displaying information to one or more users of computing device **1400**. Examples of display device **1430** include a monitor, display terminal, video projection device, and the like.

[0103] Interface(s) **1406** include various interfaces that allow computing device **1400** to interact with other systems, devices, or computing environments. Example interface(s) **1406** may include any number of different network interfaces **1420**, such as interfaces to local area networks (LANs), wide area networks (WANs), wireless networks, and the Internet. Other interface(s) include user interface **1418** and peripheral device interface **1422**. The interface(s) **1406** may also include one or more user interface elements **1418**. The interface(s) **1406** may also include one or more peripheral interfaces such as interfaces for printers, pointing devices (mice, track pad, or any suitable user interface now known to those of ordinary skill in the field, or later discovered), keyboards, and the like.

[0104] Bus **1412** allows processor(s) **1404**, memory device(s) **1404**, interface(s) **1406**, mass storage device(s) **1408**, and I/O device(s) **1410** to communicate with one another, as well as other devices or components coupled to bus **1412**. Bus **1412** represents one or more of several types of bus structures, such as a system bus, PCI bus, IEEE bus, USB bus, and so forth.

[0105] For purposes of illustration, programs and other executable program components are shown herein as discrete blocks, such as block **302** for example, although it is understood that such programs and components may reside at various times in different storage components of computing device **1400** and are executed by processor(s) **1402**. Alternatively, the systems and procedures described herein, including programs or other executable program components, can be implemented in hardware, or a combination of hardware, software, and/or firmware. For example, one or more application specific integrated circuits (ASICs) can be programmed to carry out one or more of the systems and procedures described herein.

EXAMPLES

[0106] The following examples pertain to preferred features of further embodiments:

[0107] Example 1 is a method. The method includes capturing a disk image of a first bare metal server and writing the disk image to a repository manager. The method includes launching a continuous delivery mode on a second bare metal server. The method includes mounting the disk image to the second bare metal server. The method includes configuring a unique identity of the second bare metal server.

[0108] Example 2 is a system as in Example 1, further comprising receiving a request from a user to initiate a workflow to capture the disk image of the first bare metal server, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

[0109] Example 3 is a system as in any of Examples 1-2, further comprising submitting the workflow to a worker node, wherein the workflow comprises causing the worker node to launch a live continuous delivery mode on the first bare metal server without use of a dynamic host configuration protocol (DHCP).

[0110] Example 4 is a system as in any of Examples 1-3, further comprising receiving a request from a user to initiate a workflow to install an operating system on the second bare metal server by way of a disk image dump, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

[0111] Example 5 is a system as in any of Examples 1-4, wherein launching the continuous delivery mode on the second bare metal server comprises launching without use of a dynamic host configuration protocol (DHCP).

[0112] Example 6 is a system as in any of Examples 1-5, wherein capturing the disk image of the first bare metal server comprises: requesting an initialization file from a file store of the first bare metal server without using dynamic host configuration protocol (DHCP) or a DHCP server; and receiving the initialization file.

[0113] Example 7 is a system as in any of Examples 1-6, wherein capturing the disk image of the first bare metal server further comprises: installing, on the first bare metal server, a kickstarter executable; and executing the kickstarter executable to register the first bare metal server with a machine initialization module; wherein receiving the initialization file comprises receiving the initialization file from the machine initialization module over a network in response to registering the first bare metal server with the machine initialization module.

[0114] Example 8 is a system as in any of Examples 1-7, wherein configuring the unique identity of the second bare metal server comprises configuring one or more of an internet protocol (IP) address, gateway, domain name system (DNS), or hostname of the second bare metal server.

[0115] Example 9 is a system as in any of Examples 1-8, wherein mounting the disk image to the second bare metal server comprises streaming the disk image from the repository manager by way of a uniform resource locator (URL) associated with the disk image.

[0116] Example 10 is a system as in any of Examples 1-9, wherein mounting the disk image to the second bare metal server comprises downloading the disk image from the repository manager by way of the uniform resource locator and directly writing the downloaded disk image to disk storage resources associated with the second bare metal server.

[0117] Example 11 is a system including one or more processors each configured to execute instructions stored in non-transitory computer readable storage medium, the instructions comprising any of the method steps of Examples 1-10.

[0118] Example 12 is non-transitory computer readable storage medium storing instructions for execution by one or more processors, the instructions comprising any of the method steps of Examples 1-10.

[0119] It will be appreciated that various features disclosed herein provide significant advantages and advancements in the art. The following claims are exemplary of some of those features.

[0120] In the foregoing Detailed Description of the Disclosure, various features of the disclosure are grouped together in a single embodiment for the purpose of streamlining the disclosure. This method of disclosure is not to be interpreted as reflecting an intention that the claimed disclosure requires more features than are expressly recited in each claim. Rather, inventive aspects lie in less than all features of a single foregoing disclosed embodiment.

[0121] It is to be understood that any features of the above-described arrangements, examples, and embodiments may be combined in a single embodiment comprising a combination of features taken from any of the disclosed arrangements, examples, and embodiments.

[0122] It is to be understood that the above-described arrangements are only illustrative of the application of the principles of the disclosure. Numerous modifications and alternative arrangements may be devised by those skilled in the art without departing from the spirit and scope of the disclosure and the appended claims are intended to cover such modifications and arrangements.

[0123] Thus, while the disclosure has been shown in the drawings and described above with particularity and detail, it will be apparent to those of ordinary skill in the art that numerous modifications, including, but not limited to, variations in size, materials, shape, form, function and manner of operation, assembly and use may be made without departing from the principles and concepts set forth herein.

[0124] Further, where appropriate, functions described herein can be performed in one or more of: hardware, software, firmware, digital components, or analog components. For example, one or more application specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs) can be programmed to carry out one or more of the systems and procedures described herein. Certain terms are used throughout the following description and claims to refer to particular system components. As one skilled in the art will appreciate, components may be referred to by different names. This document does not intend to distinguish between components that differ in name, but not function.

[0125] The foregoing description has been presented for the purposes of illustration and description. It is not intended to be exhaustive or to limit the disclosure to the precise form disclosed. Many modifications and variations are possible in light of the above teaching. Further, it should be noted that any or all the aforementioned alternate implementations may be used in any combination desired to form additional hybrid implementations of the disclosure.

[0126] Further, although specific implementations of the disclosure have been described and illustrated, the disclosure is not to be limited to the specific forms or arrangements of parts so described and illustrated. The scope of the disclosure is to be defined by the claims appended hereto, any future claims submitted here and in different applications, and their equivalents.

What is claimed is:

1. A method comprising:

capturing a disk image of a first bare metal server;
writing the disk image to a repository manager;

launching a continuous delivery mode on a second bare metal server;
 mounting the disk image to the second bare metal server;
 and
 configuring a unique identity of the second bare metal server.

2. The method of claim 1, further comprising receiving a request from a user to initiate a workflow to capture the disk image of the first bare metal server, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

3. The method of claim 2, further comprising submitting the workflow to a worker node, wherein the workflow comprises causing the worker node to launch a live continuous delivery mode on the first bare metal server without use of a dynamic host configuration protocol (DHCP).

4. The method of claim 1, further comprising receiving a request from a user to initiate a workflow to install an operating system on the second bare metal server by way of a disk image dump, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

5. The method of claim 1, wherein launching the continuous delivery mode on the second bare metal server comprises launching without use of a dynamic host configuration protocol (DHCP).

6. The method of claim 1, wherein capturing the disk image of the first bare metal server comprises:

requesting an initialization file from a file store of the first bare metal server without using dynamic host configuration protocol (DHCP) or a DHCP server; and
 receiving the initialization file.

7. The method of claim 6, wherein capturing the disk image of the first bare metal server further comprises:

installing, on the first bare metal server, a kickstarter executable; and
 executing the kickstarter executable to register the first bare metal server with a machine initialization module;
 wherein receiving the initialization file comprises receiving the initialization file from the machine initialization module over a network in response to registering the first bare metal server with the machine initialization module.

8. The method of claim 1, wherein configuring the unique identity of the second bare metal server comprises configuring one or more of an internet protocol (IP) address, gateway, domain name system (DNS), or hostname of the second bare metal server.

9. The method of claim 1, wherein mounting the disk image to the second bare metal server comprises streaming the disk image from the repository manager by way of a uniform resource locator (URL) associated with the disk image.

10. The method of claim 9, wherein mounting the disk image to the second bare metal server comprises downloading the disk image from the repository manager by way of the uniform resource locator and directly writing the downloaded disk image to disk storage resources associated with the second bare metal server.

11. Non-transitory computer readable storage medium storing instructions for executing by one or more processors, the instructions comprising:

capturing a disk image of a first bare metal server;
 writing the disk image to a repository manager;

launching a continuous delivery mode on a second bare metal server;
 mounting the disk image to the second bare metal server;
 and
 configuring a unique identity of the second bare metal server.

12. The non-transitory computer readable storage medium of claim 11, wherein the instructions further comprise receiving a request from a user to initiate a workflow to capture the disk image of the first bare metal server, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

13. The non-transitory computer readable storage medium of claim 12, wherein the instructions further comprise submitting the workflow to a worker node, wherein the workflow comprises causing the worker node to launch a live continuous delivery mode on the first bare metal server without use of a dynamic host configuration protocol (DHCP).

14. The non-transitory computer readable storage medium of claim 11, wherein the instructions further comprise receiving a request from a user to initiate a workflow to install an operating system on the second bare metal server by way of a disk image dump, wherein receiving the request comprises receiving the request by way of a dashboard rendered on a user interface.

15. The non-transitory computer readable storage medium of claim 11, wherein the instructions are such that launching the continuous delivery mode on the second bare metal server comprises launching without use of a dynamic host configuration protocol (DHCP).⁴

16. The non-transitory computer readable storage medium of claim 11, wherein the instructions are such that capturing the disk image of the first bare metal server comprises:

requesting an initialization file from a file store of the first bare metal server without using dynamic host configuration protocol (DHCP) or a DHCP server; and
 receiving the initialization file.

17. The non-transitory computer readable storage medium of claim 16, wherein the instructions are such that capturing the disk image of the first bare metal server further comprises:

installing, on the first bare metal server, a kickstarter executable; and
 executing the kickstarter executable to register the first bare metal server with a machine initialization module;
 wherein receiving the initialization file comprises receiving the initialization file from the machine initialization module over a network in response to registering the first bare metal server with the machine initialization module.

18. The non-transitory computer readable storage medium of claim 11, wherein the instructions are such that configuring the unique identity of the second bare metal server comprises configuring one or more of an internet protocol (IP) address, gateway, domain name system (DNS), or hostname of the second bare metal server.

19. The non-transitory computer readable storage medium of claim 11, wherein the instructions are such that mounting the disk image to the second bare metal server comprises streaming the disk image from the repository manager by way of a uniform resource locator (URL) associated with the disk image.

20. The non-transitory computer readable storage medium of claim **19**, wherein the instructions are such that mounting the disk image to the second bare metal server comprises downloading the disk image from the repository manager by way of the uniform resource locator and directly writing the downloaded disk image to disk storage resources associated with the second bare metal server.

* * * * *