



US 20250284526A1

(19) **United States**

(12) **Patent Application Publication**
Nikitin et al.

(10) **Pub. No.: US 2025/0284526 A1**

(43) **Pub. Date: Sep. 11, 2025**

(54) **METHOD AND SYSTEM FOR APPLICATION PROGRAMMING INTERFACE BASED CONTAINER SERVICE FOR SUPPORTING MULTIPLE MACHINE LEARNING APPLICATIONS**

Publication Classification

(51) **Int. Cl.**
G06F 9/455 (2018.01)
(52) **U.S. Cl.**
CPC **G06F 9/45558** (2013.01); **G06F 2009/45595** (2013.01)

(71) Applicant: **Salesforce, Inc.**, San Francisco, CA (US)

(72) Inventors: **Alexandr Nikitin**, El Sobrante, CA (US); **Vaibhav Gumashta**, San Francisco, CA (US); **Manoj Agarwal**, Cupertino, CA (US); **Swaminathan Sundaramurthy**, Los Altos, CA (US)

(73) Assignee: **Salesforce, Inc.**, San Francisco, CA (US)

(21) Appl. No.: **19/216,754**


(22) Filed: **May 23, 2025**


Related U.S. Application Data

(63) Continuation of application No. 17/337,389, filed on Jun. 2, 2021, now Pat. No. 12,340,239.

(57) **ABSTRACT**

A method and system for an application programming interface (API) based container service for supporting multiple machine learning (ML) applications is described. In particular, a scoring service container includes a base scorer to interface with a ML serving infrastructure using the API. The scoring service container also includes an application specific scorer, which itself includes a model loader and a scoring function. A model identifier is provided to the model loader, and it provides a model object. At least some parameters in a request from a client application are passed to the scoring function, which produces a scoring. The base scorer returns the scoring according to the API to the ML serving infrastructure for delivery to the client application.





100

FIG. 1




FIG. 2




FIG. 3




FIG. 4A




FIG. 4B

**METHOD AND SYSTEM FOR APPLICATION
PROGRAMMING INTERFACE BASED
CONTAINER SERVICE FOR SUPPORTING
MULTIPLE MACHINE LEARNING
APPLICATIONS**

**CROSS-REFERENCE TO RELATED
APPLICATIONS**

[0001] This application is a continuation of application Ser. No. 17/337,389, filed Jun. 2, 2021, which is hereby incorporated by reference.

TECHNICAL FIELD

[0002] One or more implementations relate to the field of container management; and more specifically, to a process and system for supporting application specific code in containers by implementing an application programming interface based container service.

BACKGROUND ART

[0003] Containers are a logical packaging in which applications can execute that is abstracted from the underlying execution environment (e.g., the underlying operating system and hardware). Applications that are containerized can be quickly deployed to many target environments including data centers, cloud architectures, or individual workstations. The containerized applications do not have to be adapted to execute in these different execution environments as long as the execution environment support containerization. The logical packaging includes a library and similar dependencies that the containerized application needs to execute.

[0004] However, containers do not include the virtualization of the hardware of an operating system. The execution environments that support containers include an operating system kernel that enables the existence of multiple isolated user-space instances. Each of these instances is a container. Containers can also be referred to as partitions, virtualization engines, virtual kernels, jails, or similar terms.

[0005] Machine learning is a type of artificial intelligence that involves algorithms that build a model based on sample data. This sample data is referred to as training data. The trained models can generate predictions, a process also referred to as scoring, based on new data that is evaluated by or input into the model. In this way, machine learning models can be developed for use in many applications without having to be explicitly programmed for these uses.

[0006] Containers can be used in connection with machine-learning serving infrastructure. Machine-learning serving infrastructures enable the execution of machine-learning models and provide services to the machine-learning models. Each machine-learning model can be separately containerized with all its required dependencies.

BRIEF DESCRIPTION OF THE DRAWINGS

[0007] The following figures use like reference numbers to refer to like elements. Although the following figures depict various example implementations, alternative implementations are within the spirit and scope of the appended claims. In the drawings:

[0008] FIG. 1 is a diagram of one example implementation of a machine-learning serving infrastructure that supports a multi-tenant system.

[0009] FIG. 2 is a diagram of one example implementation of a scoring service in the machine-learning serving infrastructure.

[0010] FIG. 3 is a flowchart of one example implementation of a process of a standardized scoring service that supports application specific scoring functions.

[0011] FIG. 4A is a block diagram illustrating an electronic device according to some example implementations.

[0012] FIG. 4B is a block diagram of a deployment environment according to some example implementations.

DETAILED DESCRIPTION

[0013] The following description describes implementations for a process and system for supporting application specific code in containers by implementing a standardized container service. The example implementations provide a method and system of serving machine learning (ML) models that support application specific code. The method and system can be referred to as a 'Bring Your Own Container (BYOC)' method and system. The example implementations enable ML model developers to customize serving containers for the models. The example implementations raise the integration point of the ML models and container from a code level to an application programming interface (API) level (e.g., a Web API level). Instead of implementing a code interface, ML model developers implement an entire container for the ML model. In some examples, the container can be implemented with hyper-text transfer protocol (HTTP) web service and dependencies incorporated into the container. At the same time, the developers of the ML model can be assured that the container can be successfully deployed to ML serving infrastructure without significant investment of development time and manpower.

[0014] Where ML models are utilized in support of applications, there are a wide variety of variety of ML frameworks or libraries (e.g., TransmogrifAI/SparkML Lib, TensorFlow, XGBoost), APIs (HTTP/gRPC, Java interface), programming languages (Java, Python), and related technologies that can be employed any combination. There can be multiple groups that are creating ML models that will be deployed in the same ML serving infrastructure such as in the case where the ML serving infrastructure supports multi-tenancy. The ML models can have different requirements, use cases, and needs. Hence there can be multiple ways to integrate ML models and serve them.

[0015] As used herein, an application can be any program or software to perform a set of tasks or operations. A 'set,' as used herein includes any positive whole number of items including a single item. A machine-learning model can be a set of algorithms and statistical data structures that can be trained to perform a specific task by identifying patterns and employing inference instead of using explicit instructions. The machine-learning model can be trained for the task using a set of training data.

[0016] Machine-learning serving infrastructure can be automated and organized to support multi-tenancy where containers can be used to execute the ML models that can service the applications and users of tenants in a multi-tenant system. Within a multitenant system, a software application is designed to provide each tenant with a tenant-specific view of the application including access only to tenant-specific data, configuration, user management, and similar tenant properties and functionality. A tenant can be a group

of users who are part of a common organization or share common access privileges to the multi-tenant system and the associated software applications.

[0017] FIG. 1 is a diagram of one example implementation of an ML serving infrastructure that supports a multi-tenant system. The machine-learning serving infrastructure **100** includes a machine-learning service (MLS) gateway **101**, routing service **103**, set of serving containers **115**, and data stores **107**, **113**, along with other supporting infrastructure.

[0018] A serving container (e.g., **105A-C**) can be an isolated execution environment that is enabled by an underlying operating system, and which executes the main functionality of a program such as an ML model. A serving container **115** can host any number of ML models for any number of tenants. Serving containers **115** can be organized as a cluster. The cluster can be a group of similar entities, such that a cluster of serving containers can be a group of serving container instances or similar grouping. An ML serving infrastructure **100** can host any number of serving containers **115** or clusters of serving containers. Different clusters can host different versions or types of ML models.

[0019] In some example implementations, a cluster of serving containers **115** can host all ML models of the same version for all tenants. This organization of the cluster can be limited by the number of ML models that a single-serving container can hold. The ML serving infrastructure **100** can scale to accommodate further additions of ML models even in cases where the number or variety of ML models exceed the capacity of the serving containers **115** in the cluster. Since each ML model's size, which can range from hundreds of kilobytes (KB) to hundreds of megabytes (MB), initialization time, and the number of requests that are serviced, can vary widely based on each tenant's underlying database and usage, some clusters of serving containers **115** may have a high resource usage, while other clusters of serving containers **115** may have a low resource utilization. The resource usage, failure, or addition of any server container in a cluster of serving containers can create the need to rebalance the supporting resources in the clusters of serving containers. When changes in the number or resource usage of the serving containers **115** are implemented, then the routing service **103** can manage the load balancing and routing of requests according to the changes in the organization of the serving containers.

[0020] A routing service **103** can be implemented as a set of routing containers, or cluster of routing containers, each implementing instances of the routing service **103** functions or subsets of these functions. The routing service **103** can authenticate any request from any tenant, and then route the request for service by ML models to any serving container **115** in a cluster of serving containers.

[0021] The ML serving infrastructure **100** receives requests from tenants via a machine-learning service (MLS) gateway **101** or a similar interface. The MLS gateway **101** or similar interface receives a request from a tenant application and identifies a version or instance of an ML model associated with the request. The MLS gateway **101** or similar interface identifies model information associated with ML models corresponding to a cluster of available serving containers associated with the version of the ML model. The MLS gateway **101** uses the model information to select a serving container from the cluster of available serving containers. If the ML model is not loaded in the serving container, the machine-learning serving infrastruc-

ture **100** loads the ML model in the serving container. If the ML model is loaded in the serving container, the system executes, in the serving container (e.g., **105A-C**), the ML model (e.g., the scoring models **133**) on behalf of the request. The ML serving infrastructure **100** responds to the request based on executing the appropriate ML model on behalf of the request.

[0022] In one example, the ML serving infrastructure **100** receives a request for scoring a business opportunity from a Customer Relationship Management (CRM) application and identifies the request requires executing a version of a particular opportunity scoring ML model **133**. The routing service **103** identifies ML model information including memory and CPU requirements for the scoring ML models **133** in the cluster of scoring serving containers. The routings service **103** applies a load balancing algorithm, resource management algorithm (e.g., a multi-dimensional bin-packing algorithm) to the collected model information to select the serving container **115** that has the best combination of available resources to execute a copy of the specific ML model associated with an incoming request.

[0023] If a copy of the specific ML model needed to service the incoming request is not already loaded in a serving container **115**, then an existing or new serving container loads the required ML model. When a copy of the specific ML model is verified to be loaded in the serving container, then the specific ML model executes the requested service or function, as specified in the received request, in the serving container. A score or similar prediction is thereby generated by the ML model and the ML serving infrastructure **100** can then respond to the request with the generated score via the MLS gateway **101**.

[0024] The ML serving infrastructure **100** can be implemented in a cloud computing environment in which data, applications, services, and other resources are stored and delivered through shared data centers. The ML serving infrastructure **100** can be implemented via any other type of distributed computer network environment in which a set of servers control the storage and distribution of resources and services for different client users.

[0025] The clusters **105A-C** of the example implementation of the ML serving infrastructure **100** can be two of any number of clusters that are serving containers for scoring services **131**. Where a scoring service **131** can be a serving container for any number of ML models that perform scoring, i.e., scoring models **133**. Each cluster **105A-C** can execute different sets of scoring services (e.g., different serving containers) for executing different varieties of machine-learning models (e.g., scoring models **133**). An incoming request can be serviced by a single ML model of a single cluster (e.g., a scoring model **133** of a given scoring service **131**) or the incoming request can be sub-divided to be serviced by multiple clusters, service containers, and ML models. In some implementations, the clusters **105A-C** and serving containers operate other similar types of ML models other than scoring ML models such as ranking and recommendation models. Scoring is provided as an example rather than by limitation. The clusters **115** can include in some implementations ranking services **141** and recommendation services **151**, which support ranking models **143** and recommendation models **153**, respectively.

[0026] In some implementations, the routing service **103** can split the incoming request into separate sub-requests, and then route the sub-requests to their corresponding clus-

ters **105A-C** of serving containers. Although these examples describe the clusters **115** of serving containers that serve one version of the scoring type of ML models, one version of the recommending type of ML models, and one version of the ranking type of machine-learning models, any clusters of any serving containers may serve any number of versions of any number of any types of any ML models and all permutations thereof.

[0027] In some implementations, each of the serving containers **115** registers with service discovery and configuration system **111** by providing the serving container's registration information, such as the host, the port, functions, or similar information. When any of the serving containers **115** is no longer available or becomes unavailable, the discovery and configuration system **111** deletes the unavailable serving container's registration information. An available serving container **115** can be referred to as an actual serving container.

[0028] The discovery and configuration system **111** can be implemented by HashiCorp Consul, Apache Zookeeper, Cloud Native Computing Foundation etcd, Netflix eureka, or any similar tool that provides service discovery and/or a service registration system. The discovery and configuration system **111** can track container information about each serving container and model information about each serving container's machine-learning models. In other implementations, this information can be stored in other locations such as datastore **113** using a format or organization. Container information can be data about an isolated execution environment, which executes the main functionality of an ML model. Model information can be data about the algorithms and/or statistical models that perform a specific task effectively by relying on patterns and inference instead of using explicit instructions.

[0029] The routing service **103** can be deployed with multiple redundant and/or distributed instances so that it is not a single point of failure for the ML serving infrastructure **100**. In some implementations, one instance of the routing service **103** acts as a master, while other instances of the routing service **103** are in a hot standby mode, ready to take over if the master instance of the routing manager fails, or perform some operations at the direction of the master instance.

[0030] The routing service **103** makes decisions to load, rebalance, delete, distribute, and replicate ML models in the serving containers **115**. These decisions can be based on the information provided to the routing service **103** by the serving containers **115** and other elements of the ML serving infrastructure **100**. The data model information in the discovery and configuration system **111** provides information about which serving containers are expected to host-specific ML models and which serving containers actually host the specified ML models. The routing service **103** can also send a list of expected ML models to a model mapping structure in the discovery and configuration system **111**. Each of the serving containers **115** can manage a list of executing ML models. If the serving container list does not match the list of expected ML models that a serving container receives, the serving container can load or delete any ML models as needed, and then update its list of executing ML models accordingly. The routing service **103** can monitor and maintain each serving container's list of actual ML models to determine where to route requests.

[0031] The routing service **103** can analyze the model information about each ML model to decide whether to replicate frequently used ML models to additional serving containers to prevent overloading the serving containers which are hosting the frequently used ML models. The routing service **103** can use the data model information of the service discovery and configuration system **111** to manage lists of available ML models and available serving containers. Every time an ML model is loaded, the serving container registers the ML model in the data model information. Therefore, the routing service **103** can route requests for a particular ML model to the serving containers.

[0032] When any of the executing serving containers **115** in any of the executing clusters of serving containers dies unexpectedly, or gracefully, the serving container's heartbeat to the service discovery and configuration system **111** fails. The ML serving infrastructure **100** removes the data for the failed serving container from its directory, files, or similar data structures in the service discovery and configuration system **111**. Based on a review of overall resource usage amongst the serving containers **115**, the routing service **103** can respond by rebalancing the serving containers **115** in terms of assigned ML models.

[0033] When requests are received by the routing service **103** via the MLS gateway **101**, a check of the mapping is made to determine if a requested ML model is executing using the service discovery and configuration system **111**. If found, then the routing service **101** can forward the requests (or divide the request into a set of sub-requests) to the identified serving containers **115**. If a ML model for the request is not found, then the routing service **103** can load the ML model from a datastore **113**, specialized database, or store **107** (e.g., a simple storage service (S3)), or similar location into a selected cluster and serving container.

[0034] In some implementations, the ML serving infrastructure **100** can include any number of additional supporting features and functions. These additional supporting features and functions can include application services **161**, version management services (VMS) **163**, redistribution services, and similar functions and services. The application services **161** can be any number, combination, and variety of functions and services for supporting tenant applications and the ML serving infrastructure **100**. The VMS **163** can be any number, combination, and variety of functions and services for supporting different versions of ML frameworks, ML models, and similar components of the machine-learning serving infrastructure **100**. The redistribution services can be any number, combination, and variety of interconnecting services to enable communication between the components of the ML serving infrastructure **100** and supporting components. In some example implementations, serving containers can interface with or support metrics bus clients **165**, databus clients **167**, and similar components. The metrics bus clients **165** can be services that gather or monitor metrics of the serving containers **115** and similar aspects of the ML serving infrastructure **100**. Similarly, the databus clients **167** can be services and functions that enable data communication and access between the serving containers **115** and other components of the ML serving infrastructure **100**.

[0035] The example implementations can define a public API that provides an interface and contract for developers to utilize for implementing containers for the ML models they are designing. The public API draws clear boundaries between systems and contexts. The API is visible outside of

the execution environment, and is an “entry point” to the system. The API is lifted above implementation specifics and encapsulates. The example implementations are Web API level (BYOC) and referred to herein as the BYOC API where developers can implement a container (e.g., a scoring service container) with application specific code and a set of application specific ML models associated therewith.

[0036] The method and system of the example implementations have advantages over the art. The method and system give more freedom to ML modeling developer teams, support different developer team needs (e.g., use Python vs Java, TMOG/SparkML vs TensorFlow/PyTorch, complex pipelines), support scalability with ability to manage more products and algorithms, and provides a language agnostic interface. The BYOC API of the example implementations is general and flexible. Requests and responses can take any type and form (Avro/Protobuf/binary). The BYOC API of the example implementations is not limiting and can use multiple protocols (HTTP/gRPC).

[0037] The example implementations utilize the BYOC API in connection with developer created containers. The developer created containers execute at the same level as other containers (e.g., Scoring and ERB containers). Each developer created container is responsible for implementing the functions and interfaces of the BYOC API to score a request, and return the response to the requestor. The requestor can be a client application that can be associated with a specific set of scoring service containers having client application specific code. The client applications can also be associated with a specific set of ML models that are loaded and executed by this set of scoring service containers. The developers generate all of the code in these containers, so they can be implemented with Java, Python, or any other language. The developer implemented container using the BYOC API is the contract that ML model developer implemented containers must follow to be used with the ML serving infrastructure.

[0038] FIG. 2 is a diagram of one example implementation of a scoring service in the machine-learning serving infrastructure. The diagram illustrates a developer created container that implements the BYOC API. The routing service will query the BYOC API to request a prediction on a particular developer defined ML model in a developer created container. Each developer created container can have its own logic for how requests are passed to a scoring service **201** that is a developer created container. The requests are received via the MLS gateway **101** and routing service **103**, which identify the scoring service **201** that will process the received request using the associated ML model, which is retrieved from the ML model storage **107** if not already executing.

[0039] The developer created containers define how the received requests are processed by implementing function calls of the BYOC API that determine how the fields of the received requests are processed (e.g., reformatted for use of the ML model). The requests can identify the ML model and associated servicing container **201** to utilize using a uniform resource identifier or similar identifier. The developer created container implementing the BYOC API can support the container serving multiple ML models. The developer created container that implements the BYOC API can dynamically load the ML models from an ML model store **107**. The

ML models can be a part of the container image or separately loaded to make it easier to deploy new or updated ML models.

[0040] Each container that implements the BYOC API can support versioning. Each container should have a version where the version can have any format. The format can be a MAJOR.MINOR.PATCH (e.g., 2.0.1) format. The major number is incremented for breaking changes, the minor for non-breaking. These version numbers can be used in the routing service **103**. Any number and variety of versions for each container can be supported at a time via the VMS **163** or similar service (e.g., a model versioning service (MVS)). In some implementations the number or variety of versions can be capped at any level to minimize the number of simultaneous containers that the ML serving infrastructure needs to support.

[0041] Any container that meets the set of requirements defined by the BYOC API will be supported by the ML serving infrastructure. However, it will be a large undertaking for every model developing team to create their own container from scratch. Therefore, a base implementation of a container that conforms with the BYOC API can be provided that a model development team can use or adapt for their purposes. This base container can be written in any language or using any technology. In one example implementation, the base container is written for the Flask web framework, behind a WSGI proxy which will allow scalability. The example base container can support caching with Python annotations, and be able to fetch objects from the model storage (e.g., s3). This will allow a shorter time to production for most common use cases, while leaving flexibility for uncommon use cases.

[0042] End-to-end performance of service containers and ML models in the ML serving infrastructure highly depends on underlying implementation of the servicing container that is compliant with the BYOC API. Access patterns and performance requirements (QPS, latency, payload) vary from use case to use case. The overhead for the example implementations can be less than 100 ms extra overhead. One of the ways to manage the overhead is to have defined service level agreements (SLAs) and performance tests for developer created containers.

[0043] The ML servicing infrastructure can provide observability means for modeling teams including metrics (collecting various metrics such as QPS, latency, responses with errors, and similar metrics), dashboards (for the metrics collected), alerts (based on SLAs and metrics), logs (provided with trace-id), distributed tracing, and similar capabilities. In addition to the SLAs defined for each container, the ML serving infrastructure can have an SLA for the service as a whole with no more than 100 ms extra overhead added by routing, so each prediction has a latency SLA of <1 s.

[0044] As shown in FIG. 2 a scoring service **201** is shown in the context of the ML serving infrastructure. A single scoring service **201** is shown by way of illustration for sake of clarity and conciseness. One skilled in the art would understand that any number and variety of containers can be supported as shown in FIG. 1. The scoring service **201** is a container that conforms to the BYOC API requirements. The scoring service **201** can be called using any protocol to pass a request to the scoring service **201** to be processed by an ML model via the application specific scorer implementation **213**, which is the developer provided code that utilizes

the developer implemented ML model to be applied to the received request. The received request will have specified the scoring container **201** and/or ML model to service the received request. In some cases, the same developer or a related developer generated the requestor and the scoring service container **201** and/or the ML model. In response to receiving the request the routing service will inspect the request to identify the scoring service **201** to which the request is to be routed.

[0045] The routing service can then call the scoring service **201** via a local client (e.g., a remote procedure call (RPC) client such as gRPC client **203**). The gRPC server **207** or similar server at the scoring service **201** then services the call from the gRPC client **203**. The interface between the gRPC client **203** and gRPC server **207** or similar protocol components can be defined via a shared proto definition **205** or similar mechanism that can be provided by the developer of the scoring service **201**. The gRPC server **207** can then pass the parameters of the request to the base scorer **209** of the scoring service **201**. The base scorer **207** can call a model load function (1) of the application specific scorer implementation **213** to ensure that the requested model is loaded or in response to detecting that the requested model is not loaded. In cases where a base container is utilized the components of the scoring service **201** other than the application specific scorer implementation **213** are provided by the base container. This enables the developer to focus on the core logic and operation of the ML model.

[0046] If the ML model is not already loaded, then the model load function loads (2) the identified ML model (e.g., the ML model is identified in the request and the identifier passed to the scoring service **201**). A confirmation of the load (3) is then sent to the base scorer **209**. The base scorer **209** can store or instruct that that the loaded model is stored in the model cache **211**. The base scorer **209** then passes parameters of the received request to the score function (5) for the application specific scorer implementation **213** to apply to the loaded ML model. The result of the score is returned to the base scorer, which in turn returns the result to the routing service **103** and thereby the requestor. The base scorer **209** can monitor the operation of the scoring service **201**, application specific scorer implementation **201**, the loaded ML models, and similar aspects to report metrics and generate logs **165**.

[0047] FIG. 3 is a flowchart of one example implementation of a process of a developer created container that provides a scoring service that supports application specific scoring functions. The scoring service container is initialized in response to a request for a model provided by the scoring service container, at ML serving infrastructure startup, or under similar conditions (Block **301**). Once the scoring service is executing, then the base scorer or similar component thereof can send a model identifier to the model loader function (Block **303**). This can be responsive to receiving a request for the model identified by the model identifier provided by the routing service. The model load function can receive the model identifier and retrieve the identified model from a model store or similar data store.

[0048] The scoring service receives the model object from the model loader upon completion of the load operation (Block **305**). The model object can be returned to the base scorer or similar component of the scoring service. The base scorer can then store the model object in model cache or

similar location (Block **307**). The model cache **307** can be local temporary storage within the scoring service.

[0049] The received request or the parameters thereof can then be set to or provided to the scoring function (Block **309**). The scoring function applies these parameters to the loaded model (i.e., the retrieved model object), which processes these parameters to generate a scoring response, which is returned to the base scorer (Block **311**). The base scorer can then return the scoring response to the requestor via the routing service and/or the MLS gateway (Block **313**). Example Electronic Devices and Environments

Electronic Device and Machine-Readable Media

[0050] One or more parts of the above implementations may include software. Software is a general term whose meaning can range from part of the code and/or metadata of a single computer program to the entirety of multiple programs. A computer program (also referred to as a program) comprises code and optionally data. Code (sometimes referred to as computer program code or program code) comprises software instructions (also referred to as instructions). Instructions may be executed by hardware to perform operations. Executing software includes executing code, which includes executing instructions. The execution of a program to perform a task involves executing some or all of the instructions in that program.

[0051] An electronic device (also referred to as a device, computing device, computer, etc.) includes hardware and software. For example, an electronic device may include a set of one or more processors coupled to one or more machine-readable storage media (e.g., non-volatile memory such as magnetic disks, optical disks, read only memory (ROM), Flash memory, phase change memory, solid state drives (SSDs)) to store code and optionally data. For instance, an electronic device may include non-volatile memory (with slower read/write times) and volatile memory (e.g., dynamic random-access memory (DRAM), static random-access memory

[0052] (SRAM)). Non-volatile memory persists code/data even when the electronic device is turned off or when power is otherwise removed, and the electronic device copies that part of the code that is to be executed by the set of processors of that electronic device from the non-volatile memory into the volatile memory of that electronic device during operation because volatile memory typically has faster read/write times. As another example, an electronic device may include a non-volatile memory (e.g., phase change memory) that persists code/data when the electronic device has power removed, and that has sufficiently fast read/write times such that, rather than copying the part of the code to be executed into volatile memory, the code/data may be provided directly to the set of processors (e.g., loaded into a cache of the set of processors). In other words, this non-volatile memory operates as both long term storage and main memory, and thus the electronic device may have no or only a small amount of volatile memory for main memory.

[0053] In addition to storing code and/or data on machine-readable storage media, typical electronic devices can transmit and/or receive code and/or data over one or more machine-readable transmission media (also called a carrier) (e.g., electrical, optical, radio, acoustical or other forms of propagated signals—such as carrier waves, and/or infrared signals). For instance, typical electronic devices also include a set of one or more physical network interface(s) to

establish network connections (to transmit and/or receive code and/or data using propagated signals) with other electronic devices. Thus, an electronic device may store and transmit (internally and/or with other electronic devices over a network) code and/or data with one or more machine-readable media (also referred to as computer-readable media).

[0054] Software instructions (also referred to as instructions) are capable of causing (also referred to as operable to cause and configurable to cause) a set of processors to perform operations when the instructions are executed by the set of processors. The phrase “capable of causing” (and synonyms mentioned above) includes various scenarios (or combinations thereof), such as instructions that are always executed versus instructions that may be executed. For example, instructions may be executed: 1) only in certain situations when the larger program is executed (e.g., a condition is fulfilled in the larger program; an event occurs such as a software or hardware interrupt, user input (e.g., a keystroke, a mouse-click, a voice command); a message is published, etc.); or 2) when the instructions are called by another program or part thereof (whether or not executed in the same or a different process, thread, lightweight thread, etc.). These scenarios may or may not require that a larger program, of which the instructions are a part, be currently configured to use those instructions (e.g., may or may not require that a user enables a feature, the feature or instructions be unlocked or enabled, the larger program is configured using data and the program’s inherent functionality, etc.). As shown by these exemplary scenarios, “capable of causing” (and synonyms mentioned above) does not require “causing” but the mere capability to cause. While the term “instructions” may be used to refer to the instructions that when executed cause the performance of the operations described herein, the term may or may not also refer to other instructions that a program may include. Thus, instructions, code, program, and software are capable of causing operations when executed, whether the operations are always performed or sometimes performed (e.g., in the scenarios described previously). The phrase “the instructions when executed” refers to at least the instructions that when executed cause the performance of the operations described herein but may or may not refer to the execution of the other instructions.

[0055] Electronic devices are designed for and/or used for a variety of purposes, and different terms may reflect those purposes (e.g., user devices, network devices). Some user devices are designed to mainly be operated as servers (sometimes referred to as server devices), while others are designed to mainly be operated as clients (sometimes referred to as client devices, client computing devices, client computers, or end user devices; examples of which include desktops, workstations, laptops, personal digital assistants, smartphones, wearables, augmented reality (AR) devices, virtual reality (VR) devices, mixed reality (MR) devices, etc.). The software executed to operate a user device (typically a server device) as a server may be referred to as server software or server code, while the software executed to operate a user device (typically a client device) as a client may be referred to as client software or client code. A server provides one or more services (also referred to as serves) to one or more clients.

[0056] The term “user” refers to an entity (e.g., an individual person) that uses an electronic device. Software

and/or services may use credentials to distinguish different accounts associated with the same and/or different users. Users can have one or more roles, such as administrator, programmer/developer, and end user roles. As an administrator, a user typically uses electronic devices to administer them for other users, and thus an administrator often works directly and/or indirectly with server devices and client devices.

[0057] FIG. 4A is a block diagram illustrating an electronic device 400 according to some example implementations. FIG. 4A includes hardware 420 comprising a set of one or more processor(s) 422, a set of one or more network interfaces 424 (wireless and/or wired), and machine-readable media 426 having stored therein software 428 (which includes instructions executable by the set of one or more processor(s) 422). The machine-readable media 426 may include non-transitory and/or transitory machine-readable media. Each of the previously described clients and the ML serving infrastructure may be implemented in one or more electronic devices 400. In one implementation: 1) each of the clients is implemented in a separate one of the electronic devices 400 (e.g., in end user devices where the software 428 represents the software to implement clients to interface directly and/or indirectly with the ML serving infrastructure (e.g., software 428 represents a web browser, a native client, a portal, a command-line interface, and/or an application programming interface (API) based upon protocols such as Simple Object Access Protocol (SOAP), Representational State Transfer (REST), etc.)); 2) the ML serving infrastructure is implemented in a separate set of one or more of the electronic devices 400 (e.g., a set of one or more server devices where the software 428 represents the software to implement the ML serving infrastructure); and 3) in operation, the electronic devices implementing the clients and the ML serving infrastructure would be communicatively coupled (e.g., by a network) and would establish between them (or through one or more other layers and/or other services) connections for submitting requests to the ML serving infrastructure and returning predictions or scoring to the client applications. Other configurations of electronic devices may be used in other implementations (e.g., an implementation in which the client and the ML serving infrastructure are implemented on a single one of electronic device 400).

[0058] During operation, an instance of the software 428 (illustrated as instance 406 and referred to as a software instance; and in the more specific case of an application, as an application instance) is executed. In electronic devices that use compute virtualization, the set of one or more processor(s) 422 typically execute software to instantiate a virtualization layer 408 and one or more software container(s) 404A-404R (e.g., with operating system-level virtualization, the virtualization layer 408 may represent a container engine (such as Docker Engine by Docker, Inc. or rkt in Container Linux by Red Hat, Inc.) running on top of (or integrated into) an operating system, and it allows for the creation of multiple software containers 404A-404R (representing separate user space instances and also called virtualization engines, virtual private servers, or jails) that may each be used to execute a set of one or more applications; with full virtualization, the virtualization layer 408 represents a hypervisor (sometimes referred to as a virtual machine monitor (VMM)) or a hypervisor executing on top of a host operating system, and the software containers

404A-404R each represent a tightly isolated form of a software container called a virtual machine that is run by the hypervisor and may include a guest operating system; with para-virtualization, an operating system and/or application running with a virtual machine may be aware of the presence of virtualization for optimization purposes). Again, in electronic devices where compute virtualization is used, during operation, an instance of the software **428** is executed within the software container **404A** on the virtualization layer **408**. In electronic devices where compute virtualization is not used, the instance **406** on top of a host operating system is executed on the “bare metal” electronic device **400**. The instantiation of the instance **406**, as well as the virtualization layer **408** and software containers **404A-404R** if implemented, are collectively referred to as software instance(s) **402**.

[0059] Alternative implementations of an electronic device may have numerous variations from that described above. For example, customized hardware and/or accelerators might also be used in an electronic device.

Example Environment

[0060] FIG. 4B is a block diagram of a deployment environment according to some example implementations. A system **440** includes hardware (e.g., a set of one or more server devices) and software to provide service(s) **442**, including the ML serving infrastructure. In some implementations the system **440** is in one or more datacenter(s). These datacenter(s) may be: 1) first party datacenter(s), which are datacenter(s) owned and/or operated by the same entity that provides and/or operates some or all of the software that provides the service(s) **442**; and/or 2) third-party datacenter(s), which are datacenter(s) owned and/or operated by one or more different entities than the entity that provides the service(s) **442** (e.g., the different entities may host some or all of the software provided and/or operated by the entity that provides the service(s) **442**). For example, third-party datacenters may be owned and/or operated by entities providing public cloud services (e.g., Amazon.com, Inc. (Amazon Web Services), Google LLC (Google Cloud Platform), Microsoft Corporation (Azure)).

[0061] The system **440** is coupled to user devices **480A-480S** over a network **482**. The service(s) **442** may be on-demand services that are made available to one or more of the users **484A-484S** working for one or more entities other than the entity which owns and/or operates the on-demand services (those users sometimes referred to as outside users) so that those entities need not be concerned with building and/or maintaining a system, but instead may make use of the service(s) **442** when needed (e.g., when needed by the users **484A-484S**). The service(s) **442** may communicate with each other and/or with one or more of the user devices **480A-480S** via one or more APIs (e.g., a REST API). In some implementations, the user devices **480A-480S** are operated by users **484A-484S**, and each may be operated as a client device and/or a server device. In some implementations, one or more of the user devices **480A-480S** are separate ones of the electronic device **400** or include one or more features of the electronic device **400**.

[0062] In some implementations, the system **440** is a multi-tenant system (also known as a multi-tenant architecture). The term multi-tenant system refers to a system in which various elements of hardware and/or software of the system may be shared by one or more tenants. A multi-tenant

system may be operated by a first entity (sometimes referred to a multi-tenant system provider, operator, or vendor; or simply a provider, operator, or vendor) that provides one or more services to the tenants (in which case the tenants are customers of the operator and sometimes referred to as operator customers). A tenant includes a group of users who share a common access with specific privileges. The tenants may be different entities (e.g., different companies, different departments/divisions of a company, and/or other types of entities), and some or all of these entities may be vendors that sell or otherwise provide products and/or services to their customers (sometimes referred to as tenant customers). A multi-tenant system may allow each tenant to input tenant specific data for user management, tenant-specific functionality, configuration, customizations, non-functional properties, associated applications, etc. A tenant may have one or more roles relative to a system and/or service. For example, in the context of a customer relationship management (CRM) system or service, a tenant may be a vendor using the CRM system or service to manage information the tenant has regarding one or more customers of the vendor. As another example, in the context of Data as a Service (DAAS), one set of tenants may be vendors providing data and another set of tenants may be customers of different ones or all of the vendors' data. As another example, in the context of Platform as a Service (PAAS), one set of tenants may be third-party application developers providing applications/services and another set of tenants may be customers of different ones or all of the third-party application developers.

[0063] Multi-tenancy can be implemented in different ways. In some implementations, a multi-tenant architecture may include a single software instance (e.g., a single database instance) which is shared by multiple tenants; other implementations may include a single software instance (e.g., database instance) per tenant; yet other implementations may include a mixed model; e.g., a single software instance (e.g., an application instance) per tenant and another software instance (e.g., database instance) shared by multiple tenants.

[0064] In one implementation, the system **440** is a multi-tenant cloud computing architecture supporting multiple services, such as one or more of the following types of services: Prediction Services based on ML; Customer relationship management (CRM); Configure, price, quote (CPQ); Business process modeling (BPM); Customer support; Marketing; External data connectivity; Productivity; Database-as-a-Service; Data-as-a-Service (DAAS or DaaS); Platform-as-a-service (PAAS or PaaS); Infrastructure-as-a-Service (IAAS or IaaS) (e.g., virtual machines, servers, and/or storage); Analytics; Community; Internet-of-Things (IoT); Industry-specific; Artificial intelligence (AI); Application marketplace (“app store”); Data modeling; Security; and Identity and access management (IAM).

[0065] For example, system **440** may include an application platform **444** that enables PAAS for creating, managing, and executing one or more applications developed by the provider of the application platform **444**, users accessing the system **440** via one or more of user devices **480A-480S**, or third-party application developers accessing the system **440** via one or more of user devices **480A-480S**.

[0066] In some implementations, one or more of the service(s) **442** may use one or more multi-tenant databases **446**, as well as system data storage **450** for system data **452**

accessible to system **440**. In certain implementations, the system **440** includes a set of one or more servers that are running on server electronic devices and that are configured to handle requests for any authorized user associated with any tenant (there is no server affinity for a user and/or tenant to a specific server). The user devices **480A-480S** communicate with the server(s) of system **440** to request and update tenant-level data and system-level data hosted by system **440**, and in response the system **440** (e.g., one or more servers in system **440**) automatically may generate one or more Structured Query Language (SQL) statements (e.g., one or more SQL queries) that are designed to access the desired information from the multi-tenant database(s) **446** and/or system data storage **450**.

[0067] In some implementations, the service(s) **442** are implemented using virtual applications dynamically created at run time responsive to queries from the user devices **480A-480S** and in accordance with metadata, including: 1) metadata that describes constructs (e.g., forms, reports, workflows, user access privileges, business logic) that are common to multiple tenants; and/or 2) metadata that is tenant specific and describes tenant specific constructs (e.g., tables, reports, dashboards, interfaces, etc.) and is stored in a multi-tenant database. To that end, the program code **460** may be a runtime engine that materializes application data from the metadata; that is, there is a clear separation of the compiled runtime engine (also known as the system kernel), tenant data, and the metadata, which makes it possible to independently update the system kernel and tenant-specific applications and schemas, with virtually no risk of one affecting the others. Further, in one implementation, the application platform **444** includes an application setup mechanism that supports application developers' creation and management of applications, which may be saved as metadata by save routines. Invocations to such applications, including the ML serving infrastructure, may be coded using Procedural Language/Structured Object Query Language (PL/SOQL) that provides a programming language style interface. Invocations to applications may be detected by one or more system processes, which manages retrieving application metadata for the tenant making the invocation and executing the metadata as an application in a software container (e.g., a virtual machine).

[0068] Network **482** may be any one or any combination of a LAN (local area network), WAN (wide area network), telephone network, wireless network, point-to-point network, star network, token ring network, hub network, or other appropriate configuration. The network may comply with one or more network protocols, including an Institute of Electrical and Electronics Engineers (IEEE) protocol, a 3rd Generation Partnership Project (3GPP) protocol, a 4th generation wireless protocol (4G) (e.g., the Long Term Evolution (LTE) standard, LTE Advanced, LTE Advanced Pro), a fifth generation wireless protocol (5G), and/or similar wired and/or wireless protocols, and may include one or more intermediary devices for routing data between the system **440** and the user devices **480A-480S**.

[0069] Each user device **480A-480S** (such as a desktop personal computer, workstation, laptop, Personal Digital Assistant (PDA), smartphone, smartwatch, wearable device, augmented reality (AR) device, virtual reality (VR) device, etc.) typically includes one or more user interface devices, such as a keyboard, a mouse, a trackball, a touch pad, a touch screen, a pen or the like, video or touch free user interfaces,

for interacting with a graphical user interface (GUI) provided on a display (e.g., a monitor screen, a liquid crystal display (LCD), a head-up display, a head-mounted display, etc.) in conjunction with pages, forms, applications and other information provided by system **440**. For example, the user interface device can be used to access data and applications hosted by system **440**, and to perform searches on stored data, and otherwise allow one or more of users **484A-484S** to interact with various GUI pages that may be presented to the one or more of users **484A-484S**. User devices **480A-480S** might communicate with system **440** using TCP/IP (Transfer Control Protocol and Internet Protocol) and, at a higher network level, use other networking protocols to communicate, such as Hypertext Transfer Protocol (HTTP), File Transfer Protocol (FTP), Andrew File System (AFS), Wireless Application Protocol (WAP), Network File System (NFS), an application program interface (API) based upon protocols such as Simple Object Access Protocol (SOAP), Representational State Transfer (REST), etc. In an example where HTTP is used, one or more user devices **480A-480S** might include an HTTP client, commonly referred to as a "browser," for sending and receiving HTTP messages to and from server(s) of system **440**, thus allowing users **484A-484S** of the user devices **480A-480S** to access, process and view information, pages and applications available to it from system **440** over network **482**.

Conclusion

[0070] In the above description, numerous specific details such as resource partitioning/sharing/duplication implementations, types and interrelationships of system components, and logic partitioning/integration choices are set forth in order to provide a more thorough understanding. The invention may be practiced without such specific details, however. In other instances, control structures, logic implementations, opcodes, means to specify operands, and full software instruction sequences have not been shown in detail since those of ordinary skill in the art, with the included descriptions, will be able to implement what is described without undue experimentation.

[0071] References in the specification to "one implementation," "an implementation," "an example implementation," etc., indicate that the implementation described may include a particular feature, structure, or characteristic, but every implementation may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same implementation. Further, when a particular feature, structure, and/or characteristic is described in connection with an implementation, one skilled in the art would know to affect such feature, structure, and/or characteristic in connection with other implementations whether or not explicitly described.

[0072] For example, the figure(s) illustrating flow diagrams sometimes refer to the figure(s) illustrating block diagrams, and vice versa. Whether or not explicitly described, the alternative implementations discussed with reference to the figure(s) illustrating block diagrams also apply to the implementations discussed with reference to the figure(s) illustrating flow diagrams, and vice versa. At the same time, the scope of this description includes implementations, other than those discussed with reference to the block diagrams, for performing the flow diagrams, and vice versa.

[0073] Bracketed text and blocks with dashed borders (e.g., large dashes, small dashes, dot-dash, and dots) may be used herein to illustrate optional operations and/or structures that add additional features to some implementations. However, such notation should not be taken to mean that these are the only options or optional operations, and/or that blocks with solid borders are not optional in certain implementations.

[0074] The detailed description and claims may use the term “coupled,” along with its derivatives. “Coupled” is used to indicate that two or more elements, which may or may not be in direct physical or electrical contact with each other, co-operate or interact with each other.

[0075] While the flow diagrams in the figures show a particular order of operations performed by certain implementations, such order is exemplary and not limiting (e.g., alternative implementations may perform the operations in a different order, combine certain operations, perform certain operations in parallel, overlap performance of certain operations such that they are partially in parallel, etc.).

[0076] While the above description includes several example implementations, the invention is not limited to the implementations described and can be practiced with modification and alteration within the spirit and scope of the appended claims. The description is thus illustrative instead of limiting.

What is claimed is:

1. A method comprising:
 - providing, by one or more computing devices of a machine learning service infrastructure, a bring your own container (BYOC) application programming interface (API) for access to a custom container image created by a developer, wherein the BYOC API is parameterized with a server endpoint for a client application to call the custom container image using the machine learning service infrastructure and a reference to a machine learning model registered in the machine learning service infrastructure;
 - hosting, from storage of the machine learning service infrastructure, an instance of the custom container image that services requests from the BYOC API, wherein the instance of the custom container image has a model object obtained from a model store using the reference, and wherein the model object is an instance of the machine learning model adapted for use by the custom container image;
 - passing, by the machine learning service infrastructure, a request, which is from the client application and directed to the server endpoint, to a scoring function provided by the machine learning model;
 - receiving, by the machine learning service infrastructure, scoring from the scoring function; and
 - returning, by the machine learning service infrastructure, the scoring to the client application.
2. The method of claim 1, wherein the request from the client application is directed, using a load balancer, to the instance of the custom container image hosted by the machine learning service infrastructure.
3. The method of claim 1, further comprising:
 - providing, by the machine learning service infrastructure through an interface to a logging service, log data related to the use of the first container image.
4. The method of claim 1, wherein, when the referenced machine learning model is not instantiated, it is automatically deployed from the model store.
5. The method of claim 1, further comprising:
 - sending the reference to a model loader; and
 - receiving the model object from the model loader.
6. The method of claim 1, further comprising:
 - storing the model object in a model cache.
7. The method of claim 6, further comprising:
 - loading the model object from the model cache in response to subsequent requests.
8. The method of claim 1, further comprising:
 - responsive to a second call from the client application to call the custom container image using the machine learning service infrastructure and a second reference to a second machine learning model registered in the machine learning service infrastructure, obtaining for the instance of the custom container image a second model object from the model store using the second reference, wherein the second model object is an instance of the second machine learning model adapted for use by the custom container image;
 - passing, by the machine learning service infrastructure, a second request, which is from the client application and directed to the server endpoint, to a second scoring function provided by the second machine learning model;
 - receiving, by the machine learning service infrastructure, scoring from the second scoring function; and
 - returning, by the machine learning service infrastructure, the scoring to the client application.
9. A non-transitory machine-readable storage medium that provides instructions that, if executed by a set of one or more processors, are configurable to cause the set of one or more processors to perform operations comprising:
 - providing, by a machine learning service infrastructure, a bring your own container (BYOC) application programming interface (API) for access to a custom container image created by a developer, wherein the BYOC API is parameterized with a server endpoint for a client application to call the custom container image using the machine learning service infrastructure and a reference to a machine learning model registered in the machine learning service infrastructure;
 - hosting, from storage of the machine learning service infrastructure, an instance of the custom container image that services requests from the BYOC API, wherein the instance of the custom container image has a model object obtained from a model store using the reference, and wherein the model object is an instance of the machine learning model adapted for use by the custom container image;
 - passing, by the machine learning service infrastructure, a request, which is from the client application and directed to the server endpoint, to a scoring function provided by the machine learning model;
 - receiving, by the machine learning service infrastructure, scoring from the scoring function; and
 - returning, by the machine learning service infrastructure, the scoring to the client application.
10. The non-transitory machine-readable storage medium of claim 9, wherein the request from the client application is

directed, using a load balancer, to the instance of the custom container image hosted by the machine learning service infrastructure.

11. The non-transitory machine-readable storage medium of claim 9, further comprising:

providing, by the machine learning service infrastructure through an interface to a logging service, log data related to the use of the first container image.

12. The non-transitory machine-readable storage medium of claim 9, wherein, when the referenced machine learning model is not instantiated, it is automatically deployed from the model store.

13. The non-transitory machine-readable storage medium of claim 9, further comprising:

sending the reference to a model loader; and
receiving the model object from the model loader.

14. The non-transitory machine-readable storage medium of claim 9, further comprising:

storing the model object in a model cache.

15. The non-transitory machine-readable storage medium of claim 14, further comprising:

loading the model object from the model cache in response to subsequent requests.

16. The non-transitory machine-readable storage medium of claim 9, further comprising:

responsive to a second call from the client application to call the custom container image using the machine learning service infrastructure and a second reference to a second machine learning model registered in the machine learning service infrastructure, obtaining for the instance of the custom container image a second model object from the model store using the second reference, wherein the second model object is an instance of the second machine learning model adapted for use by the custom container image;

passing, by the machine learning service infrastructure, a second request, which is from the client application and directed to the server endpoint, to a second scoring function provided by the second machine learning model;

receiving, by the machine learning service infrastructure, scoring from the second scoring function; and
returning, by the machine learning service infrastructure, the scoring to the client application.

17. An apparatus comprising:

a set of one or more processors;

a non-transitory machine-readable storage medium that provides instructions that, if executed by the set of one or more processors, are configurable to cause the apparatus to perform operations comprising,

providing, by one or more computing devices of a machine learning service infrastructure, a bring your own container (BYOC) application programming interface (API) for access to a custom container image created by a developer, wherein the BYOC API is parameterized with a server endpoint for a client application to call the custom container image using the machine learning service infrastructure and a reference to a machine learning model registered in the machine learning service infrastructure;

hosting, from storage of the machine learning service infrastructure, an instance of the custom container image that services requests from the BYOC API, wherein the instance of the custom container image has a model object obtained from a model store using the reference, and wherein the model object is an instance of the machine learning model adapted for use by the custom container image;

passing, by the machine learning service infrastructure, a request, which is from the client application and directed to the server endpoint, to a scoring function provided by the machine learning model;

receiving, by the machine learning service infrastructure, scoring from the scoring function; and

returning, by the machine learning service infrastructure, the scoring to the client application.

18. The apparatus of claim 17, wherein the request from the client application is directed, using a load balancer, to the instance of the custom container image hosted by the machine learning service infrastructure.

19. The apparatus of claim 17, the operations further comprising:

providing, by the machine learning service infrastructure through an interface to a logging service, log data related to the use of the first container image.

20. The apparatus of claim 17, wherein, when the referenced machine learning model is not instantiated, it is automatically deployed from the model store.

21. The apparatus of claim 17, the operations further comprising:

sending the reference to a model loader; and
receiving the model object from the model loader.

22. The apparatus of claim 17, the operations further comprising:

storing the model object in a model cache.

23. The apparatus of claim 22, the operations further comprising:

loading the model object from the model cache in response to subsequent requests.

24. The apparatus of claim 17, the operations further comprising:

responsive to a second call from the client application to call the custom container image using the machine learning service infrastructure and a second reference to a second machine learning model registered in the machine learning service infrastructure, obtaining for the instance of the custom container image a second model object from the model store using the second reference, wherein the second model object is an instance of the second machine learning model adapted for use by the custom container image;

passing, by the machine learning service infrastructure, a second request, which is from the client application and directed to the server endpoint, to a second scoring function provided by the second machine learning model;

receiving, by the machine learning service infrastructure, scoring from the second scoring function; and
returning, by the machine learning service infrastructure, the scoring to the client application.

* * * * *