

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284537

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Shin; Wongyu et al.

PROCESSING DEVICE AND METHOD FOR MANAGING TASKS THEREOF

Abstract

A neural processing device and a method for managing tasks thereof are provided. The neural processing device includes a neural core configured to perform a task and generate a completion signal for completion of the task, a core global configured to transfer task information for the task to the neural core and receive the completion signal of the task from the neural core, and a task manager configured to generate and transmit the task information to the core global, receive the completion signal from the core global, generate a completion report, and transmit the completion report.

Inventors: Shin; Wongyu (Seongnam-si, KR), Chi; Miock (Seongnam-si, KR), Kim; Hongyun (Seongnam-si, KR), Oh; Jinwook (Seongnam-si, KR), Yoon; Juyeong (Seongnam-si, KR)

Applicant: Rebellions Inc. (Seongnam-si, KR)

Family ID: 1000008618243

Appl. No.: 19/212424

Filed: May 19, 2025

Foreign Application Priority Data

KR	10-2022-0146444	Nov. 04, 2022
KR	10-2022-0146533	Nov. 04, 2022
KR	10-2022-0146543	Nov. 04, 2022

Related U.S. Application Data

parent US continuation 18674752 20240524 parent-grant-document US 12340246 child US 19212424

Publication Classification

Int. Cl.: G06F9/48 (20060101); G06F9/38 (20180101)

U.S. Cl.:

CPC G06F9/4881 (20130101); G06F9/3838 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION [0001] This application is a continuation of U.S. application Ser. No. 18/674,752, filed on May 24, 2024, which is a continuation of U.S. application Ser. No. 18/491,695, filed on Oct. 20, 2023, now granted U.S. Pat. No. 12,079,661, issued on Sep. 3, 2024, which claims priority under 35 U.S.C § 119 to Korean Patent Application No. 10-2022-0146533, filed on Nov. 4, 2022, Korean Patent Application No. 10-2022-0146444, filed on Nov. 4, 2022, and Korean Patent Application No. 10-2022-0146543, filed on Nov. 4, 2022, in the Korean Intellectual Property Office, the entire contents of which is hereby incorporated by reference.

TECHNICAL FIELD

[0002] The disclosure relates to a processing device and a method for managing tasks thereof. Specifically, the disclosure relates to a processing device and a method for managing tasks that perform task management efficiently.

BACKGROUND

[0003] For the last few years, artificial intelligence technology has been the core technology of the Fourth Industrial Revolution and the subject of discussion as the most promising technology worldwide. The biggest problem with artificial intelligence technology is computing performance. For artificial intelligence technology to realize a level of human learning ability, reasoning ability, perceptual ability, natural language implementation ability, etc., it is of the utmost importance to process a large amount of data quickly.

[0004] The central processing unit (CPU) or graphics processing unit (GPU) of off-the-shelf computers was used to implement deep-learning training and inference in early artificial intelligence, but these components had limitations in their ability to perform the tasks of deep-learning training and inference with high workloads. Thus, neural processing units (NPUs) that are structurally specialized for deep learning tasks have received a lot of attention.

[0005] Such a neural processing unit has a plurality of computation devices therein, and each computation device operates in parallel and thus can enhance computation efficiency. In this case, a method of distributing and managing tasks to each computation device may be important.

[0006] The description set forth in the background section should not be assumed to be prior art merely because it is set forth in the background section. The background section may describe aspects or embodiments of the disclosure.

SUMMARY

[0007] Aspects of the disclosure provide a neural processing device that performs task management efficiently.

[0008] Aspects of the disclosure provide a method for managing tasks of a neural processing device that performs task management efficiently.

[0009] According to some aspects of the disclosure, a neural processing device includes a neural core configured to perform a task and generate a completion signal for completion of the task, a core global configured to transfer task information for the task to the neural core and receive the completion signal of the task from the neural core, and a task manager configured to generate and transmit the task information to the core global, receive the completion signal from the core global, generate a completion report, and transmit the completion report.

[0010] According to some aspects, the task manager includes a task passage configured to generate a task descriptor for the task and transfer the task information according to the task descriptor to the core global, and a done passage configured to perform check-in by receiving the task descriptor from the task passage and storing the task descriptor in a check-in buffer, and generate the completion report by receiving the completion signal and performing check-out of the task descriptor, and wherein the check-in is a procedure in which the task descriptor is registered before being processed, and the check-out is a procedure in which the task descriptor is processed and de-registered.

[0011] According to some aspects, the task passage includes a task buffer configured to receive and store the task from a command processor, a task queue in which the task descriptor received from the task buffer waits, and a runtime handle configured to transfer task information corresponding to a task descriptor transferred from the task queue to the core global, and wherein the runtime handle controls the task descriptor to wait in the task queue if the task descriptor includes a wait field.

[0012] According to some aspects, the runtime handle releases a waiting state of the task descriptor in response to a run signal provided from the command processor.

[0013] According to some aspects, the runtime handle includes a run signal counter in which the run signal is received, and the runtime handle releases the waiting state of the task descriptor via the run signal received in advance in the run signal counter.

[0014] According to some aspects, the runtime handle transfers task information corresponding to a task descriptor whose waiting state has been released or task information corresponding to a task descriptor that does not include the wait field to the core global, and provides a task descriptor corresponding to the task information transferred to the core global as check-in data to the done passage.

[0015] According to some aspects, the task queue includes a first queue configured to receive the task descriptor from the task buffer, a dependency checker configured to receive a task descriptor from the first queue and perform a dependency check on the received task descriptor, and a second queue configured to receive the task descriptor for which the dependency check has been completed from the dependency checker.

[0016] According to some aspects, the second queue includes a first task descriptor and a second task descriptor stored in sequence, and as the first task descriptor is controlled to wait in the second queue by the runtime handle, the second task descriptor is also controlled to wait in the second queue.

[0017] According to some aspects, the task manager is configured to generate a timeout detection signal by monitoring an execution time of the task in the neural core, and generate a timeout report according to the timeout detection signal.

[0018] According to some aspects, the timeout report is included in the completion report and transferred to a command processor.

[0019] According to some aspects, the task includes a first task that is set to generate the completion report and a second task that is set not to generate the completion report, and the task manager checks in a first task descriptor corresponding to the first task and a second task descriptor corresponding to the second task into a first check-in buffer and a second check-in buffer, respectively.

[0020] According to some aspects, the task manager modifies the second task descriptor to generate a completion report if a timeout report for the second task is generated.

[0021] According to some aspect of the disclosure, a neural processing device includes at least one neural processor including first and second neural processors, a shared memory shared by the at least one neural processor, a command processor configured to transfer tasks to the at least one neural processor, and a global interconnection configured to exchange data and control signals of the at least one neural processor and the shared memory, wherein the first neural processor includes first and second neural cores, and a first task manager configured to generate task descriptors by receiving the tasks, and report completion reports for the task descriptors to the command processor.

[0022] According to some aspects, the second neural processor includes a second task manager configured to generate task descriptors by receiving the tasks, and report completion reports for the task descriptors to the command processor, and third and fourth neural cores.

[0023] According to some aspects, the at least one neural processor further includes a third neural processor, and the third neural processor includes a third task manager configured to generate task descriptors by receiving the tasks, and report completion reports for the task descriptors to the command processor.

[0024] According to some aspects, the neural processing device further includes a first master task manager configured to receive the completion reports from the first and second task managers and report the completion reports to the command processor.

[0025] According to some aspects of the disclosure, a method for managing tasks of a neural processing device includes fetching tasks and generating task descriptors, checking dependencies of the task descriptors and sequentially transferring the task descriptors, checking wait fields of the task descriptors and sequentially transferring the task descriptors whose wait fields are checked, performing the tasks of the transferred task descriptors, and receiving completion signals of the performed tasks and generating completion reports.

[0026] According to some aspects, the checking the wait fields of the task descriptors and sequentially transferring the task descriptors includes controlling a task descriptor to wait in a task queue if the task descriptor includes a wait field, releasing a waiting state of the task descriptor in response to a run signal provided from a command processor, generating task information corresponding to the task descriptor whose waiting state has been released and transferring the task information to a core global, and providing the task descriptor corresponding to the task information transferred to the core global as check-in data to the done passage.

[0027] According to some aspects, the performing the tasks of the transferred task descriptors includes monitoring execution times for the tasks and generating timeout detection signals, and the receiving the completion signals of the performed tasks and generating the completion reports includes generating timeout reports according to the timeout detection signals, generating checkout reports according to the received completion signals, and generating completion reports based on the checkout reports and the timeout reports.

[0028] According to some aspects, the tasks include a first task that is set to generate the completion report and a second task that is set not to generate the completion report, the checking the wait fields of the task descriptors and sequentially transferring the task descriptors includes checking in a first task descriptor corresponding to the first task and a second task descriptor corresponding to the second task into a first check-in buffer and a second check-in buffer, respectively, if a timeout report for the second task is generated in the generating the timeout reports, the generating the completion reports includes modifying the second task descriptor to generate a completion report.

[0029] Aspects of the disclosure are not limited to those mentioned above and other objects and advantages of the disclosure that have not been mentioned can be understood by the following description and will be more clearly understood according to embodiments of the disclosure. In addition, it will be readily understood that the objects and advantages of the disclosure can be realized by the means and combinations thereof set forth in the claims.

[0030] The neural core, the neural processing device, and the method for managing tasks thereof according to the disclosure can efficiently perform the management of tasks distributed to a plurality of neural cores as the task manager manages and monitors the tasks of the neural cores, respectively, and reports them to the command processor.

[0031] In addition, CPU intervention can be minimized by checking whether each task is completed and actively controlling the dependency of each task. Accordingly, the overhead due to communication with the CPU can be minimized.

[0032] Further, if a report to the command processor is required, it can be configured to modify whether to generate a completion report for each task descriptor. Accordingly, it may be possible to report related to a particular task without monitoring all tasks, and thus stable monitoring of tasks may be possible while reducing the burden on the command processor.

[0033] Moreover, the execution time of tasks can be controlled via the runtime handle configured in the task manager. That is, it can be managed so that a temporary waiting time occurs between distributed tasks, and it is possible to further improve not only the sequential processing of the entire tasks but also the efficiency of task processing and management.

[0034] In addition to the foregoing, the specific effects of the disclosure will be described together while elucidating the specific details for carrying out the embodiments below.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0035] FIG. 1 is a block diagram for illustrating a neural processing system in accordance with some embodiments of the disclosure;

[0036] FIG. 2 is a block diagram for illustrating the neural processing device of FIG. 1;

[0037] FIG. 3 is a block diagram for illustrating the neural core SoC of FIG. 2;

[0038] FIG. 4 is a structural diagram for illustrating the global interconnection of FIG. 3;

[0039] FIG. 5 is a block diagram for illustrating a flow of control signals of the neural processing device of FIG. 1;

[0040] FIG. 6 is a block diagram for illustrating the neural processor of FIG. 3 in detail;

[0041] FIG. 7 is a diagram for illustrating a hierarchical structure of a neural processing device in accordance with some embodiments of the disclosure;

[0042] FIG. 8 is a block diagram for illustrating the neural core of FIG. 6 in detail;

[0043] FIG. 9 is a block diagram for illustrating the LSU of FIG. 8 in detail;

[0044] FIG. 10 is a block diagram for illustrating the processing unit of FIG. 8;

[0045] FIG. 11 is a block diagram for illustrating the L0 memory of FIG. 8 in detail;

[0046] FIG. 12 is a block diagram for illustrating the local memory bank of FIG. 11;

[0047] FIG. 13 is a block diagram for illustrating a flow of data and control signals of the neural processing device of FIG. 1;

[0048] FIG. 14 is a block diagram for illustrating the relationship between the command processor and the task managers of FIG. 13;

[0049] FIG. 15 is a block diagram for illustrating the structure of the task manager of FIG. 8 in detail;

[0050] FIG. 16 is a block diagram for illustrating the table passage of FIG. 15 in detail;

[0051] FIG. 17 is a block diagram for illustrating the task passage of FIG. 15 in detail;

[0052] FIG. 18 is a block diagram for specifically illustrating the function of the runtime handle;

[0053] FIG. 19 is an example diagram for illustrating a process of processing task descriptors included in a 1st queue of the second queue group in response to a run signal;

[0054] FIG. 20 is an example diagram for illustrating a process of processing task descriptors included in the 1st queue of the second queue group in response to a run signal received in

advance;

[0055] FIG. **21** is an example diagram for illustrating a process of processing task descriptors included in the 1st queue of the second queue group in response to a plurality of run signals;
[0056] FIG. **22** is an example diagram for illustrating a process of processing task descriptors included in the 1st queue of the second queue group through a plurality of run signals received by a counter;

[0057] FIG. **23** is a block diagram for illustrating the done passage of FIG. **15** in detail;

[0058] FIG. **24** is an example diagram for illustrating a first task descriptor and a second task descriptor that are set differently as to whether to generate a completion report;

[0059] FIG. **25** is a flowchart for illustrating a process of performing modification on the second task descriptor that is set not to generate a completion report;

[0060] FIG. **26** is an example diagram for illustrating a process of generating a completion report in response to a second checkout report of the second task descriptor modified according to FIG. **25**;

[0061] FIG. **27** is a block diagram for illustrating the report managing module of FIG. **23** in detail;

[0062] FIG. **28** is a diagram for illustrating data that the core global and the processing unit of FIG. **15** exchange;

[0063] FIG. **29** is a diagram for illustrating types of task descriptors stored in a first queue, a second queue, and a check-in buffer;

[0064] FIG. **30** is a diagram for illustrating a first queue, a second queue, and a check-in buffer of a neural processing device in accordance with some embodiments of the disclosure;

[0065] FIG. **31** is a diagram for illustrating a first queue, a second queue, and a check-in buffer of a neural processing device in accordance with some embodiments of the disclosure;

[0066] FIG. **32** is a block diagram for illustrating the structure of the neural processing device of FIG. **1** in detail;

[0067] FIG. **33** is a diagram for illustrating a hierarchical structure of a command processor and task managers of a neural processing device in accordance with some embodiments of the disclosure;

[0068] FIG. **34** is a diagram for illustrating a hierarchical structure of a command processor and task managers of a neural processing device in accordance with some embodiments of the disclosure;

[0069] FIG. **35** is a block diagram for illustrating the memory reconstruction of the neural processing system of FIG. **1**;

[0070] FIG. **36** is a block diagram showing an example of the memory reconstruction of the neural processing system of FIG. **1**;

[0071] FIG. **37** is an enlarged block diagram of a portion A of FIG. **35**;

[0072] FIG. **38** is a diagram for illustrating the first memory bank of FIG. **37** in detail;

[0073] FIG. **39** is a block diagram for illustrating a software hierarchy of a neural processing device in accordance with some embodiments of the disclosure;

[0074] FIG. **40** is a conceptual diagram for illustrating deep learning calculations performed by a neural processing device in accordance with some embodiments of the disclosure;

[0075] FIG. **41** is a conceptual diagram for illustrating training and inference operations of a neural network of a neural processing device in accordance with some embodiments of the disclosure;

[0076] FIG. **42** is a flowchart for illustrating a method for managing tasks of a neural processing device in accordance with some embodiments of the disclosure;

[0077] FIG. **43** is a flowchart for illustrating the checking the dependencies of FIG. **42** in detail;

[0078] FIG. **44** is a flowchart for illustrating the checking the wait fields of FIG. **42** in detail;

[0079] FIG. **45** is a flowchart for illustrating a process of generating a dependency update request in detail, in the generating the completion reports of FIG. **42**;

[0080] FIG. **46** is a flowchart for illustrating a process of generating the completion reports in detail, in the generating the completion reports of FIG. **42**; and

[0081] FIG. 47 is a flowchart for illustrating a control process according to the state of a report queue, in the generating the completion reports of FIG. 42.

DETAILED DESCRIPTION OF THE PREFERRED EMBODIMENTS

[0082] The terms or words used in the disclosure and the claims should not be construed as limited to their ordinary or lexical meanings. They should be construed as the meaning and concept in line with the technical idea of the disclosure based on the principle that the inventor can define the concept of terms or words in order to describe his/her own embodiments in the best possible way. Further, since the embodiment described herein and the configurations illustrated in the drawings are merely one embodiment in which the disclosure is realized and do not represent all the technical ideas of the disclosure, it should be understood that there may be various equivalents, variations, and applicable examples that can replace them at the time of filing this application.

[0083] Although terms such as first, second, A, B, etc. used in the description and the claims may be used to describe various components, the components should not be limited by these terms. These terms are used only for the purpose of distinguishing one component from another. For example, a first component may be referred to as a second component, and similarly, a second component may be referred to as a first component, without departing from the scope of the disclosure. The term ‘and/or’ includes a combination of a plurality of related listed items or any item of the plurality of related listed items.

[0084] The terms used in the description and the claims are merely used to describe particular embodiments and are not intended to limit the disclosure. Singular expressions include plural expressions unless the context explicitly indicates otherwise. In the application, terms such as “comprise,” “have,” “include,” “contain,” etc. should be understood as not precluding the possibility of existence or addition of features, numbers, steps, operations, components, parts, or combinations thereof described herein. Terms such as a “circuit” or “circuitry”, refers to a circuit in hardware but may also refer to a circuit in software.

[0085] Unless otherwise defined, the phrases “A, B, or C,” “at least one of A, B, or C,” or “at least one of A, B, and C” may refer to only A, only B, only C, both A and B, both A and C, both B and C, all of A, B, and C, or any combination thereof.

[0086] Unless otherwise defined, all terms used herein, including technical or scientific terms, have the same meaning as commonly understood by those of ordinary skill in the art to which the disclosure pertains.

[0087] Terms such as those defined in commonly used dictionaries should be construed as having a meaning consistent with the meaning in the context of the relevant art, and are not to be construed in an ideal or excessively formal sense unless explicitly defined in the disclosure.

[0088] In addition, each configuration, procedure, process, method, or the like included in each embodiment of the disclosure may be shared to the extent that they are not technically contradictory to each other.

[0089] Hereinafter, a neural processing device in accordance with some embodiments of the disclosure will be described with reference to FIGS. 1 to 41.

[0090] FIG. 1 is a block diagram for illustrating a neural processing system in accordance with some embodiments of the disclosure.

[0091] Referring to FIG. 1, a neural processing system NPS in accordance with some embodiments may include a first neural processing device 1, a second neural processing device 2, and an external interface 3.

[0092] The first neural processing device 1 may be a device that performs calculations using an artificial neural network. The first neural processing device 1 may be, for example, a device specialized in performing tasks of deep learning calculations. However, the embodiment is not limited thereto.

[0093] The second neural processing device 2 may be a device having the same or similar configuration as the first neural processing device 1. The first neural processing device 1 and the

second neural processing device **2** may be connected to each other via the external interface **3** and share data and control signals.

[0094] Although FIG. **1** shows two neural processing devices, the neural processing system NPS in accordance with some embodiments is not limited thereto. In some embodiments, in a neural processing system NPS, three or more neural processing devices may be connected to each other via the external interface **3**. Also, conversely, a neural processing system NPS in accordance with some embodiments may include only one neural processing device.

[0095] In some embodiments, each of the first neural processing device **1** and the second neural processing device **2** may be a processing device other than the neural processing device. In some embodiments, each of the first neural processing device **1** and the second neural processing device **2** may be a graphics processing unit (GPU), a central processing unit (CPU), and other types of processing units as well. In the following, the first neural processing device **1** and the second neural processing device **2** will be described as neural processing devices for convenience. FIG. **2** is a block diagram for illustrating the neural processing device of FIG. **1**.

[0096] Referring to FIG. **2**, the first neural processing device **1** may include a neural core SoC **10**, a CPU **20**, an off-chip memory **30**, a first non-volatile memory interface **40**, a first volatile memory interface **50**, a second non-volatile memory interface **60**, a second volatile memory interface **70**, and a control interface (CIF) **80**.

[0097] The neural core SoC **10** may be a system on a chip device. The neural core SoC **10** can be an artificial intelligence calculation device and may be an accelerator. The neural core SoC **10** may be, for example, any one of a graphics processing unit (GPU), a field-programmable gate array (FPGA), or an application-specific integrated circuit (ASIC). However, the embodiment is not limited thereto.

[0098] The neural core SoC **10** may exchange data with other external calculation devices via the external interface **3**. Further, the neural core SoC **10** may be connected to the non-volatile memory **31** and the volatile memory **32** via the first non-volatile memory interface **40** and the first volatile memory interface **50**, respectively.

[0099] The CPU **20** may be a control device that controls the system of the first neural processing device **1** and executes program calculations. The CPU **20** is a general-purpose calculation device and may have low efficiency in performing simple parallel calculations that are frequently used in deep learning. Accordingly, there can be high efficiency by performing calculations in deep learning inference and training tasks by the neural core SoC **10**.

[0100] The CPU **20** may exchange data with other external calculation units via the external interface **3**. Further, the CPU **20** may be connected to the non-volatile memory **31** and the volatile memory **32** via the second non-volatile memory interface **60** and the second volatile memory interface **70**, respectively.

[0101] The CPU **20** may also transfer tasks to the neural core SoC **10** via commands. In some embodiments, the CPU **20** may be a kind of host that gives instructions to the neural core SoC **10**. In some embodiments, the neural core SoC **10** can efficiently perform parallel computation tasks such as deep learning tasks according to the instructions of the CPU **20**.

[0102] The off-chip memory **30** may be a memory disposed outside the chip of the neural core SoC **10**. The off-chip memory **30** may include a non-volatile memory **31** and a volatile memory **32**.

[0103] The non-volatile memory **31** may be a memory that continuously retains stored information even if electric power is not supplied. The non-volatile memory **31** may include, for example, at least one of Read-Only Memory (ROM), Programmable Read-Only Memory (PROM), Erasable Alterable ROM (EAROM), Erasable Programmable Read-Only Memory (EPROM), Electrically Erasable Programmable Read-Only Memory (EEPROM) (e.g., NAND Flash memory, NOR Flash memory), Ultra-Violet Erasable Programmable Read-Only Memory

[0104] (UVEEPROM), Ferroelectric Random-Access Memory (FeRAM), Magnetoresistive Random-Access Memory (MRAM), Phase-change Random-Access Memory (PRAM), silicon-

oxide-nitride-oxide-silicon (SONOS), Resistive Random-Access Memory (RRAM), Nanotube Random-Access Memory (NRAM), magnetic computer storage devices (e.g., hard disks, diskette drives, magnetic tapes), optical disc drives, or 3D XPoint memory. However, the embodiment is not limited thereto.

[0105] The volatile memory **32** may be a memory that continuously requires electric power to retain stored information, unlike the non-volatile memory **31**. The volatile memory **32** may include, for example, at least one of Dynamic Random-Access Memory (DRAM), Static Random-Access Memory (SRAM), Synchronous Dynamic Random-Access Memory (SDRAM), or Double Data Rate SDRAM (DDR SDRAM). However, the embodiment is not limited thereto.

[0106] Each of the first non-volatile memory interface **40** and the second non-volatile memory interface **60** may include, for example, at least one of Parallel Advanced Technology Attachment (PATA), Small Computer System Interface (SCSI), Serial Attached SCSI (SAS), Serial Advanced Technology Attachment (SATA), or PCI Express (PCIe). However, the embodiment is not limited thereto.

[0107] Each of the first volatile memory interface **50** and the second volatile memory interface **70** may be, for example, at least one of SDR (Single Data Rate), DDR (Double Data Rate), QDR (Quad Data Rate), or XDR (eXtreme Data Rate, Octal Data Rate). However, the embodiment is not limited thereto.

[0108] The control interface **80** may be an interface for transferring control signals between the CPU **20** and the neural core SoC **10**. The control interface **80** may transmit commands of the CPU **20** and transmit the responses thereto of the neural core SoC **10**. The control interface **80** may be, for example, PCIe (PCI Express), but is not limited thereto.

[0109] FIG. **3** is a block diagram for illustrating the neural core SoC of FIG. **2**.

[0110] Referring to FIGS. **2** and **3**, the neural core SoC **10** may include at least one neural processor **1000**, a shared memory **2000**, direct memory access (DMA) **3000**, a non-volatile memory controller **4000**, a volatile memory controller **5000**, a command processor **7000**, and a global interconnection **6000**. In some embodiments, the command processor **7000** may be referred to as a command processor circuit, but for the sake of convenience, the terms are unified as a command processor. In addition, the command processor **7000** may be implemented as a circuit or circuitry.

[0111] The command processor **7000** may be implemented in software on the off-chip memory **30**, in particular, on the volatile memory **32** of FIG. **2**. However, the embodiment is not limited thereto and may also be implemented as separate hardware. Furthermore, the command processor **7000** may also be implemented in software in part and in hardware in part. In some embodiments, the part implemented in hardware can increase the computation speed of the command processor **7000**.

[0112] The neural processor **1000** may be a calculation device that directly performs calculation tasks. If there exist a plurality of neural processors **1000**, calculation tasks may be assigned to respective neural processors **1000**. The respective neural processors **1000** may be connected to each other via the global interconnection **6000**.

[0113] The shared memory **2000** may be a memory shared by multiple neural processors **1000**. The shared memory **2000** may store data of each neural processor **1000**. In addition, the shared memory **2000** may receive data from the off-chip memory **30**, store the data temporarily, and transfer the data to each neural processor **1000**. The shared memory **2000** may also receive data from the neural processor **1000**, store the data temporarily, and transfer the data to the off-chip memory **30** of FIG. **2**.

[0114] The shared memory **2000** may be required to be a relatively high-speed memory. Accordingly, the shared memory **2000** may include, for example, an SRAM. However, the embodiment is not limited thereto. In some embodiments, the shared memory **2000** may include a DRAM as well.

[0115] The shared memory **2000** may be a memory corresponding to the SoC level, i.e., level 3

(L3). Accordingly, the shared memory **2000** may also be defined as an L3 shared memory.

[0116] The DMA **3000** can directly control the movement of data without requiring the CPU **20** or the neural processor **1000** to control input/output of data. Accordingly, the DMA **3000** may control data movement between memories, thereby minimizing the number of interrupts of the CPU **20** or the neural processor **1000**.

[0117] The DMA **3000** may control the data movement between the shared memory **2000** and the off-chip memory **30**. Via the authority of the DMA **3000**, the non-volatile memory controller **4000** and the volatile memory controller **5000** may perform the movement of data.

[0118] The non-volatile memory controller **4000** may control the task of reading from or writing onto the non-volatile memory **31**. The non-volatile memory controller **4000** may control the non-volatile memory **31** via the first non-volatile memory interface **40**. In some embodiments, the non-volatile memory controller **4000** may be referred to as a non-volatile memory controller circuit, but for the sake of convenience, the terms are unified as a non-volatile memory controller. In addition, the non-volatile memory controller **4000** may be implemented as a circuit or circuitry.

[0119] The volatile memory controller **5000** may control the task of reading from or writing onto the volatile memory **32**. Further, the volatile memory controller **5000** may perform a refresh task of the volatile memory **32**. The volatile memory controller **5000** may control the volatile memory **32** via the first volatile memory interface **50**. Likewise, the volatile memory controller **5000** may be referred to as a volatile memory controller circuit, but for the sake of convenience, the terms are unified as a volatile memory controller. In addition, the volatile memory controller **5000** may be implemented as a circuit or circuitry.

[0120] The command processor **7000** may be connected to the control interface **80**. The command processor **7000** may receive control signals from the CPU **20** via the control interface **80**. The command processor **7000** may generate tasks via the control signals received from the CPU **20** and transmit the tasks to neural processors **1000**. Further, the command processor **7000** may receive completion reports for the tasks from each neural processor **1000**.

[0121] The global interconnection **6000** may connect the at least one neural processor **1000**, the shared memory **2000**, the DMA **3000**, the non-volatile memory controller **4000**, the command processor **7000**, and the volatile memory controller **5000** to one another. In addition, the external interface **3** may also be connected to the global interconnection **6000**. The global interconnection **6000** may be a path through which data travels between the at least one neural processor **1000**, the shared memory **2000**, the DMA **3000**, the non-volatile memory controller **4000**, the volatile memory controller **5000**, the command processor **7000**, and the external interface **3**.

[0122] The global interconnection **6000** may transmit not only data but also control signals and may transmit a signal for synchronization. In some embodiments, in the neural processing device, each neural processor **1000** may directly transmit and receive a synchronization signal, instead of a separate control processor managing the synchronization signal. Accordingly, it is possible to preclude the latency of the synchronization signal generated by the command processor **7000**.

[0123] In some embodiments, if there exist a plurality of neural processors **1000**, there may be dependencies of individual tasks in which the task of one neural processor **1000** needs to be finished before the next neural processor **1000** start a new task. The end and start of these individual tasks can be checked via synchronization signals. In conventional techniques, both the reception of such synchronization signals and an instruction to start a new task were performed exclusively by the command processor **7000** or the host, i.e., the CPU **20**.

[0124] However, as the number of neural processors **1000** increases and task dependencies are designed more complicatedly, the number of these synchronization signals increases exponentially, and thus, the latency resulting from each synchronization signal can greatly reduce the efficiency of tasks.

[0125] Accordingly, in the neural processing device in accordance with some embodiments, each neural processor **1000**, instead of the command processor **7000**, may directly transmit a

synchronization signal to another neural processor **1000** according to the dependency of a task. In some embodiments, several neural processors **1000** can perform the synchronization tasks in parallel as compared with the method managed by the command processor **7000**, thereby minimizing the latency due to synchronization.

[0126] In addition, the command processor **7000** needs to perform the task scheduling of the neural processors **1000** according to task dependencies, and the overhead of such scheduling may increase significantly as the number of neural processors **1000** increases. Accordingly, in the neural processing device, in accordance with some embodiments, the scheduling task is also performed by the individual neural processors **1000**, and thus, the performance of the neural processing device can be improved without resulting in an additional scheduling burden.

[0127] In addition, the neural processing device in accordance with some embodiments of the disclosure can perform monitoring whether a task is completed, an event occurs, a task is delayed, or the like in the neural core of each neural processor **1000**, and can minimize the intervention of the command processor **7000** and reduce the burden on the command processor **7000**, thereby improving the performance of the device.

[0128] Furthermore, the neural processing device in accordance with some embodiments of the disclosure can selectively generate a completion report by setting whether to monitor a task for each task, and can be configured to modify whether to generate a completion report if a report to the command processor **7000** is required. Accordingly, it may be possible to report tasks that require an alert without monitoring all tasks, and stable monitoring of tasks may be possible while reducing the burden on the command processor **7000**.

[0129] FIG. **4** is a structural diagram for illustrating the global interconnection of FIG. **3**.

[0130] Referring to FIG. **4**, the global interconnection **6000** may include a data channel **6100**, a control channel **6200**, and an L2 sync channel **6300**.

[0131] The data channel **6100** may be a dedicated channel for transmitting data. Through the data channel **6100**, the at least one neural processor **1000**, the shared memory **2000**, the DMA **3000**, the non-volatile memory controller **4000**, the volatile memory controller **5000**, and the external interface **3** may exchange data with one another.

[0132] The control channel **6200** may be a dedicated channel for transmitting control signals. Via the control channel **6200**, the at least one neural processor **1000**, the shared memory **2000**, the DMA **3000**, the non-volatile memory controller **4000**, the volatile memory controller **5000**, the command processor **7000**, and the external interface **3** may exchange control signals with one another. In particular, the command processor **7000** may transfer various control signals to each neural processor **1000**.

[0133] The L2 sync channel **6300** may be a dedicated channel for transmitting synchronization signals. Through the L2 sync channel **6300**, the at least one neural processor **1000**, the shared memory **2000**, the DMA **3000**, the non-volatile memory controller **4000**, the volatile memory controller **5000**, the command processor **7000**, and the external interface **3** may exchange synchronization signals with one another.

[0134] The L2 sync channel **6300** may be set as a dedicated channel inside the global interconnection **6000**, and thus, may not overlap with other channels and transmit synchronization signals quickly. Accordingly, the neural processing device in accordance with some embodiments does not require new wiring work and may smoothly perform the synchronization task by using the global interconnection **6000**.

[0135] FIG. **5** is a block diagram for illustrating a flow of control signals of the neural processing device of FIG. **1**.

[0136] Referring to FIG. **5**, the CPU **20** may transfer control signals to the command processor **7000** via the control interface **80**. In some embodiments, a control signal may be a signal instructing execution of each operation, such as a computation task or a data load/store task.

[0137] The command processor **7000** may receive the control signals and transfer the control

signals to the at least one neural processor **1000** via the control channel **6200**. Each control signal may be stored in the neural processor **1000** as each task.

[0138] FIG. **6** is a block diagram for illustrating the neural processor of FIG. **3**.

[0139] Referring to FIGS. **3** to **6**, the neural processor **1000** may include at least one neural core **100**, a local interconnection **200**, an L1 sync path **300**, an L1 shared memory **400**, a core global **500**, a task manager **600**, and an L1 load/store unit (LSU) **700**. The L1 LSU **700**, the task manager **600**, and the core global **500** may be referred to respectively as an L1 LSU circuit, a task manager circuit, and a core global circuit, but for the sake of convenience, the terms are respectively unified as an L1 LSU, a task manager, and a core global. In addition, the L1 LSU **700**, the task manager **600**, and the core global **500** may each be implemented as a circuit or circuitry.

[0140] The at least one neural core **100** may share and perform the tasks of the neural processor **1000**. The number of neural cores **100** may be, for example, eight. However, the embodiment is not limited thereto. FIGS. **3** and **6** illustrate that a plurality of neural cores are included in the neural processor **1000**, but the embodiment is not limited thereto. In some embodiments, the neural processor **1000** may be configured with only one neural core.

[0141] The neural core **100** may receive task information from the core global **500** and perform a task according to the task information. In some embodiments, a task may be defined by a control signal, and the task may be one of computation operations and memory operations. The memory operation may be, for example, one of micro-DMA (μ DMA), LP micro-DMA (Low Priority μ DMA), store μ DMA (ST μ DMA), and a pre-processing task.

[0142] The L1 shared memory **400** may be a memory shared by the neural cores **100** in the neural processor **1000**. The L1 shared memory **400** may store data of each neural core **100**. In addition, the L1 shared memory **400** may receive data from the shared memory **2000** of FIG. **3**, store them temporarily, and transfer them to each neural core **100**. On the contrary, the L1 shared memory **400** may also receive data from the neural core **100**, store them temporarily, and transfer them to the shared memory **2000** of FIG. **3**.

[0143] The L1 shared memory **400** may be a memory corresponding to the neural processor level, i.e., level 1 (L1). The L2 shared memory, i.e., the shared memory **2000** may be shared by the neural processors **1000**, and the L1 shared memory **400** may be shared by the neural cores **100**.

[0144] The L1 LSU **700** may receive at least one of data, a control signal, or a synchronization signal from the outside via the global interconnection **6000**. The L1 LSU **700** may transmit at least one of the received data, control signal, or synchronization signal to the L1 shared memory **400**. Similarly, the L1 LSU **700** may transfer at least one of the data, the control signal, or the synchronization signal to the outside via the global interconnection **6000**. In addition, the L1 LSU **700** may transfer and receive at least one of the data, the control signal, or the synchronization signal to and from each of the neural cores **100**.

[0145] The neural core **100** may receive task information from the core global **500** and perform tasks according to the task information. In some embodiments, a task may be a computation task (calculation task) or a task related to a memory operation. A task may be defined by a control signal. The task information is information about a task, and may be information about a type of a task, a form of a task, additional information about a task, and the like.

[0146] The neural core **100** may transfer a completion signal indicating the completion of execution of a task to the core global **500**.

[0147] The task manager **600** may receive tasks from a control interconnection CI. In some embodiments, the control interconnection CI may be a generic term for transmission interfaces that transfer tasks from the command processor **7000**. In some embodiments, the control interconnection CI may include the control channel **6200** and the local interconnection **200**.

[0148] The task manager **600** may receive tasks, generate task information, and transmit the task information to the core global **500**. Further, the task manager **600** may receive completion signals via the core global **500**, generate completion reports accordingly, and transmit completion reports

to the command processor **7000** via the control interconnection **CI**.

[0149] The core global **500** may be a wire structure connected in hardware within the neural core **100**. Although not shown, the core global **500** may be a structure connecting all of the neural core **100**, the L1 shared memory **400**, the L1 LSU **700**, and the task manager **600**. Accordingly, the local interconnection **200** and the L1 sync path **300** may also be included in the core global **500**.

However, the embodiment is not limited thereto.

[0150] The core global **500** may receive the task information from the task manager **600** and transfer the task information to the neural core **100**, and may receive a completion signal related thereto from the neural core **100**. Subsequently, the core global **500** may transfer the completion signal to the task manager **600**.

[0151] The local interconnection **200** may connect the at least one neural core **100**, the L1 shared memory **400**, the L1 LSU **700**, the core global **500**, and the task manager **600** to one another. The local interconnection **200** may be a path through which data moves between the at least one neural core **100**, the L1 shared memory **400**, the L1 LSU **700**, the core global **500**, and the task manager **600**. The local interconnection **200** may be connected and transmit data to the global interconnection **6000** of FIG. 3.

[0152] The L1 sync path **300** may connect the at least one neural core **100**, the L1 shared memory **400**, the core global **500** and the task manager **600** to each other. The L1 sync path **300** may be a path through which synchronization signals of the at least one neural core **100**, the L1 shared memory **400**, the core global **500** and the task manager **600** travel.

[0153] The L1 sync path **300** may be formed physically separately from the local interconnection **200**. In the case of the local interconnection **200**, sufficient channels may not be formed therein, unlike the global interconnection **6000**. In such a case, the L1 sync path **300** may be formed separately so that the synchronization signal can be transmitted quickly and without any delay. The L1 sync path **300** may be used for synchronization performed at a level one step lower than that of the L2 sync channel **6300** of the global interconnection **6000**.

[0154] FIG. 7 is a diagram for illustrating a hierarchical structure of a neural processing device in accordance with some embodiments of the disclosure.

[0155] Referring to FIG. 7, the neural core SoC **10** may include at least one neural processor **1000**. Each neural processor **1000** may transmit data to each other via the global interconnection **6000**.

[0156] The neural processors **1000** may each include at least one neural core **100**. The neural core **100** may be a unit of processing optimized for deep learning computation tasks. The neural core **100** may be a unit of processing corresponding to one operation of deep learning computation tasks. In other words, a deep learning computation task can be represented by a sequential or parallel combination of multiple operations. The neural cores **100** may each be a unit of processing capable of processing one operation, and may be a minimum computation unit that can be considered for scheduling from the viewpoint of a compiler.

[0157] The neural processing device in accordance with the embodiment may configure scales of the minimum computation unit considered from the viewpoint of compiler scheduling and the hardware unit of processing to be the same, so that fast and efficient scheduling and computation tasks can be performed.

[0158] In some embodiments, if a unit of processing into which hardware can be divided is too large compared to computation tasks, inefficiency of the computation tasks may occur in driving the unit of processing. Conversely, it is not appropriate to schedule a unit of processing that is a unit smaller than an operation, which is the minimum scheduling unit of the compiler, every time since a scheduling inefficiency may occur and hardware design costs may increase.

[0159] Therefore, by adjusting the scales of the scheduling unit of the compiler and the hardware unit of processing to be similar in the embodiment, it is possible to simultaneously satisfy the fast scheduling of computation tasks and the efficient execution of the computation tasks without wasting hardware resources.

[0160] FIG. 8 is a block diagram for illustrating the neural core of FIG. 6 in detail.

[0161] Referring to FIG. 8, the neural core **100** may include a load/store unit (LSU) **110**, an L0 memory **120**, a weight buffer **130**, an activation LSU **140**, an activation buffer **150**, and a processing unit **160**. In some embodiments, the LSU **110** and the activation LSU **140** may be referred to respectively as an LSU circuit and an activation LSU circuit, but for the sake of convenience, the terms are respectively unified as an LSU and an activation LSU. Further, the LSU **110** and the activation LSU **140** may each be implemented as a circuit or circuitry.

[0162] The LSU **110** may receive at least one of data, a control signal, or a synchronization signal from the outside via the local interconnection **200** and the L1 sync path **300**. The LSU **110** may transmit at least one of the data, the control signal, or the synchronization signal received to the L0 memory **120**. Similarly, the LSU **110** may transfer at least one of the data, the control signal, or the synchronization signal to the outside via the local interconnection **200** and the L1 sync path **300**.

[0163] Specifically, the micro-DMA task may be a task in which the neural core **100** loads a program or data from the shared memory **2000** or the off-chip memory **30** to the L0 memory **120**. The LP micro-DMA task may be a load task for a program or data to be used later rather than a current program or data, unlike a general micro-DMA task. As such a task has a low priority, it can be identified differently from the micro-DMA task. The ST micro-DMA task may be a store task that stores data from the L0 memory **120** of the neural core **100** to the shared memory **2000** or the off-chip memory **30**. The pre-processing task may include a task that pre-loads data such as a large number of lookup tables in the CPU **20**.

[0164] FIG. 9 is a block diagram for illustrating the LSU of FIG. 8 in detail.

[0165] Referring to FIG. 9, the LSU **110** may include a local memory load unit (LMLU) **111a**, a local memory store unit (LMSU) **111b**, a neural core load unit (NCLU) **112a**, a neural core store unit (NCSU) **112b**, a load buffer LB, a store buffer SB, a load (LD) engine **113a**, a store (ST) engine **113b**, and a translation lookaside buffer (TLB) **114**. The local memory load unit **111a**, the local memory store unit **111b**, the neural core load unit **112a**, the neural core store unit **112b**, the load engine **113a**, and the store engine **113b** may be referred to respectively as a local memory load circuit, a local memory store circuit, a neural core load circuit, a neural core store circuit, a load engine circuit, and a store engine circuit. However, for the sake of convenience, the terms are respectively unified as a local memory load unit, a local memory store unit, a neural core load unit, a neural core store unit, a load engine, and a store engine. In addition, the local memory load unit **111a**, the local memory store unit **111b**, the neural core load unit **112a**, the neural core store unit **112b**, the load engine **113a**, and the store engine **113b** may each be implemented as a circuit or circuitry.

[0166] The local memory load unit **111a** may fetch a load instruction for the L0 memory **120** and issue the load instruction. When the local memory load unit **111a** provides the issued load instruction to the load buffer LB, the load buffer LB may sequentially transmit memory access requests to the load engine **113a** according to the inputted order.

[0167] Further, the local memory store unit **111b** may fetch a store instruction for the L0 memory **120** and issue the store instruction. When the local memory store unit **111b** provides the issued store instruction to the store buffer SB, the store buffer SB may sequentially transmit memory access requests to the store engine **113b** according to the inputted order.

[0168] The neural core load unit **112a** may fetch a load instruction for the neural core **100** and issue the load instruction. When the neural core load unit **112a** provides the issued load instruction to the load buffer LB, the load buffer LB may sequentially transmit memory access requests to the load engine **113a** according to the inputted order.

[0169] In addition, the neural core store unit **112b** may fetch a store instruction for the neural core **100** and issue the store instruction. When the neural core store unit **112b** provides the issued store instruction to the store buffer SB, the store buffer SB may sequentially transmit memory access requests to the store engine **113b** according to the inputted order.

[0170] The load engine **113a** may receive the memory access request and retrieve data via the local interconnection **200**. In some embodiments, the load engine **113a** may quickly find the data by using a translation table of a logical address and a physical address that has been used recently in the translation lookaside buffer **114**. If the logical address of the load engine **113a** is not in the translation lookaside buffer **114**, the address translation information may be found in another memory.

[0171] The store engine **113b** may receive the memory access request and retrieve data via the local interconnection **200**. In some embodiments, the store engine **113b** may quickly find the data by using a translation table of a logical address and a physical address that has been used recently in the translation lookaside buffer **114**. If the logical address of the store engine **113b** is not in the translation lookaside buffer **114**, the address translation information may be found in another memory.

[0172] The load engine **113a** and the store engine **113b** may send synchronization signals to the L1 sync path **300**. In some embodiments, the synchronization signal may indicate that the task has been completed.

[0173] Referring to FIG. **8** again, the L0 memory **120** is a memory located inside the neural core **100**, and may receive all input data required for the tasks by the neural core **100** from the outside and store them temporarily. In addition, the L0 memory **120** may temporarily store the output data calculated by the neural core **100** for transmission to the outside. The L0 memory **120** may serve as a cache memory of the neural core **100**.

[0174] The L0 memory **120** may transmit an input activation Act_In to the activation buffer **150** and receive an output activation Act_Out via the activation LSU **140**. The L0 memory **120** may directly transmit and receive data to and from the processing unit **160**, in addition to the activation LSU **140**. In other words, the L0 memory **120** may exchange data with each of a processing element (PE) array **163** and a vector unit **164**. The L0 memory **120** may be a memory corresponding to the level of the neural core. In some embodiments, the L0 memory **120** may be a private memory of the neural core that is not shared.

[0175] The L0 memory **120** may transmit data such as activations or weights via a data path. The L0 memory **120** may exchange synchronization signals via an L0 sync path, which is a separate dedicated path. The L0 memory **120** may exchange synchronization signals with, for example, the LSU **110**, the weight buffer **130**, the activation LSU **140**, and the processing unit **160** via the L0 sync path.

[0176] The weight buffer **130** may receive a weight from the L0 memory **120**. The weight buffer **130** may transfer the weight to the processing unit **160**. The weight buffer **130** may temporarily store the weight before transferring it.

[0177] The input activation Act_In and the output activation Act_Out may refer to input values and output values of the layers of a neural network. In some embodiments, if there are a plurality of layers in the neural network, the output value of the previous layer becomes the input value of the next layer, and thus, the output activation Act_Out of the previous layer may be utilized as the input activation Act_In of the next layer.

[0178] The weight may refer to a parameter that is multiplied by the input activation Act_In inputted in each layer. The weight is adjusted and confirmed in the deep learning training phase, and may be used to derive the output activation Act_Out via a fixed value in the inference phase.

[0179] The activation LSU **140** may transfer the input activation Act_In from the L0 memory **120** to the activation buffer **150**, and the output activation Act_Out from the activation buffer **150** to the on-chip buffer. In other words, the activation LSU **140** may perform both a load task and a store task of the activation.

[0180] The activation buffer **150** may provide the input activation Act_In to the processing unit **160** and receive the output activation Act_Out from the processing unit **160**. The activation buffer **150** may temporarily store the input activation Act_In and the output activation Act_Out.

[0181] The activation buffer **150** may quickly provide the activation to the processing unit **160**, in particular, the PE array **163**, which has a large quantity of calculations, and may quickly receive the activation, thereby increasing the calculation speed of the neural core **100**.

[0182] The processing unit **160** may be a module that performs calculations. The processing unit **160** may perform not only one-dimensional calculations but also two-dimensional matrix calculations, i.e., convolution operations. The processing unit **160** may receive an input activation Act_In, multiply it by a weight, and then add it to generate an output activation Act_Out.

[0183] FIG. **10** is a block diagram for illustrating the processing unit of FIG. **8** in detail.

[0184] Referring to FIG. **8** and FIG. **10**, the processing unit **160** may include a PE array **163**, a vector unit **164**, a column register **161**, and a row register **162**.

[0185] The PE array **163** may receive the input activation Act_In and the weight and perform multiplication on them. In some embodiments, each of the input activation Act_In and the weight may be in the form of matrices and calculated via convolution. Through this, the PE array **163** may generate an output activation Act_Out. However, the embodiment is not limited thereto. The PE array **163** may generate any types of outputs other than the output activation Act_Out as well.

[0186] The PE array **163** may include at least one processing element (PE) **163_1**. The processing elements **163_1** may be aligned with each other so that each of the processing elements **163_1** may perform multiplication on one input activation Act_In and one weight.

[0187] The PE array **163** may sum values for each multiplication to generate a subtotal. This subtotal may be utilized as an output activation Act_Out. The PE array **163** performs two-dimensional matrix multiplication, and thus, may be referred to as a 2D matrix compute unit.

[0188] The vector unit **164** may mainly perform one-dimensional calculations. The vector unit **164**, together with the PE array **163**, may perform deep learning calculations. Through this, the processing unit **160** may be specialized for necessary calculations. In other words, each of the at least one neural core **100** has calculation modules that perform a large amount of two-dimensional matrix multiplications and one-dimensional calculations, and thus, can efficiently perform deep learning tasks.

[0189] The column register **161** may receive a first input **I1**. The column register **161** may receive the first input **I1**, and distribute them to each column of the processing elements **163_1**.

[0190] The row register **162** may receive a second input **12**. The row register **162** may receive the second input **12**, and distribute them to each row of the processing elements **163_1**.

[0191] The first input **I1** may be an input activation Act_In or a weight. The second input **12** may be a value other than the first input **I1** between the input activation Act_In or the weight. Alternatively, the first input **I1** and the second input **12** may be values other than the input activation Act_In and the weight.

[0192] FIG. **11** is a block diagram for illustrating the L0 memory of FIG. **8** in detail.

[0193] Referring to FIG. **11**, the L0 memory **120** may include a scheduler **121** and one or more local memory banks **122**.

[0194] When data is stored in the L0 memory **120**, the scheduler **121** may receive data from the load engine **113a**. In some embodiments, the local memory bank **122** may be allocated for the data in a round-robin manner. Accordingly, data may be stored in any one of the local memory banks **122**.

[0195] In contrast to this, when data is loaded from the L0 memory **120**, the scheduler **121** may receive the data from the local memory bank **122** and transmit the data to the store engine **113b**. The store engine **113b** may store the data in the outside through the local interconnection **200**. In some embodiments, the scheduler **121** may be referred to as a scheduler circuit, but for the sake of convenience, the term is unified as a scheduler. In addition, the scheduler **121** may be implemented as a circuit or circuitry.

[0196] FIG. **12** is a block diagram for illustrating the local memory bank of FIG. **11** in detail.

[0197] Referring to FIG. **12**, the local memory bank **122** may include a local memory bank

controller **122_1** and a local memory bank cell array **122_2**.

[0198] The local memory bank controller **122_1** may manage read and write operations via the addresses of data stored in the local memory bank **122**. In other words, the local memory bank controller **122_1** may manage the input/output of data as a whole.

[0199] The local memory bank cell array **122_2** may be of a structure in which cells in which data is directly stored are arranged in rows and columns. The local memory bank cell array **122_2** may be controlled by the local memory bank controller **122_1**.

[0200] FIG. **13** is a block diagram for illustrating a flow of data and control signals of the neural processing device of FIG. **1**, and FIG. **14** is a block diagram for illustrating the relationship between the command processor and the task managers of FIG. **13**.

[0201] Referring to FIGS. **13** and **14**, each neural processor **1000** may include a task manager **600** and an L1 LSU **700** therein, respectively. The task managers **600** may exchange control signals and responses to the control signals with the command processor **7000** via a control interconnection CI.

[0202] In contrast, the L1 LSU **700** may exchange data via a data interconnection and memory DIM. The data interconnection and memory DIM may include an interconnection for transmitting data and a memory in which data is shared. Specifically, the data interconnection and memory DIM may include a local interconnection **200** and a data channel **6100**. In addition, the data interconnection and memory DIM may include an L1 shared memory **400**, a shared memory **2000**, and a volatile memory **32**. However, the embodiment is not limited thereto.

[0203] The task managers **600** may be controlled by the command processor **7000**. In some embodiments, the command processor **7000** may transfer tasks to the task managers **600** via control signals, and the task managers **600** may transfer task completion reports to the command processor **7000**. At least one task manager **600** may be included in the neural processor **1000**. Moreover, if there exist a plurality of the neural processors **1000**, the number of task managers **600** may get larger. Such a plurality of task managers **600** may all be controlled by the command processor **7000**.

[0204] FIG. **15** is a block diagram for illustrating the structure of the task manager of FIG. **8** in detail.

[0205] Referring to FIGS. **8**, **9**, and **15**, the task manager **600** may include a table passage **610**, a task passage **620**, and a done passage **630**. In some embodiments, the table passage **610**, the task passage **620**, and the done passage **630** may be referred to respectively as a table passage circuit, a task passage circuit, and a done passage circuit, respectively, but for the sake of convenience, the terms are respectively unified as a table passage, a task passage, and a done passage. Further, the table passage **610**, the task passage **620**, and the done passage **630** may be implemented as circuits or circuitry.

[0206] The table passage **610** may receive a table update request TURQ for updating a matching table of physical addresses and logical addresses from the control channel **6200** and transfer the table update request TURQ to the core global **500**. In some embodiments, the table update request TURQ may be transferred from the command processor **7000** via the control channel **6200**.

[0207] The task passage **620** may receive a task from the control channel **6200**, generate task information according thereto, and transmit the task information to the core global **500**. In some embodiments, the task may be transferred from the command processor **7000** via the control channel **6200**. The core global **500** may transfer the task information to the neural core **100**. The neural core **100** may perform the task according to the transferred task information and transfer a completion signal back to the core global **500**.

[0208] The core global **500** may transfer the completion signal to the done passage **630**. The done passage **630** may receive the completion signal and generate a completion report DNrp for the task. The done passage **630** may transmit the completion report DNrp to the command processor **7000** via the control channel **6200**.

[0209] In addition, the table update request TURQ in the table passage **610** may be transferred to

the neural core **100** via the core global **500**. In some embodiments, the table of the translation lookaside buffer **114** inside the LSU **110** of the neural core **100** may be updated.

[0210] FIG. **16** is a block diagram for illustrating the table passage of FIG. **15** in detail.

[0211] Referring to FIG. **16**, the table passage **610** may include a table buffer **611** and first to m-th update request queues **611a1** to **611am**.

[0212] In the table buffer **611**, table update requests TURQ in which physical addresses and logical addresses are matched may be transmitted from the command processor **7000** and stored. When the core global **500** fetches these table update requests TURQ, each table update request TURQ may be stored in the first to m-th update request queues **611a1** to **611am**.

[0213] Different types of table update requests TURQ may be stored in the first to m-th update request queues **611a1** to **611am**. For example, different types of table update requests TURQ may include at least one of neural core TLB update requests, micro-DMA TLB update requests, LP micro-DMA TLB update requests, or ST micro-DMA TLB update requests. However, the embodiment is not limited thereto. In some embodiments, each of the first to m-th update request queues **611a1** to **611am** may include the same type of table update request TURQ as well.

[0214] In addition, each of the first to m-th update request queues **611a1** to **611am** may be a general queue, i.e., a queue that accommodates various types of requests. Accordingly, each of the first to m-th update request queues **611a1** to **611am** may accept requests regardless of the type.

[0215] Each of the first to m-th update request queues **611a1** to **611am** may transmit the table update requests TURQ to the core global **500**, and may transfer them to the LSU **110** via the core global **500**. In some embodiments, the tables in the translation lookaside buffer **114** inside the LSU **110** may be updated.

[0216] FIG. **17** is a block diagram for illustrating the task passage of FIG. **15** in detail.

[0217] Referring to FIG. **17**, the task passage **620** includes a task buffer **621**, a task queue **622**, and a runtime handle RH.

[0218] The task buffer **621** may store tasks Task according to the control signals transmitted from the command processor **7000**. The task buffer **621** may store the tasks Task in the form of task descriptors in the task queue **622** by task fetching operations of the core global **500**.

[0219] The task descriptor for a task may contain dependency information. The dependency information may represent a list of reference tasks which the task depends on. The reference tasks which the task depends on may be tasks which have to be completed for the task to be transferred to a plurality of neural cores **100** via the core global **500**.

[0220] The task queue **622** is configured to store task descriptors in sequence, perform dependency checks on the stored task descriptors, and sequentially store task descriptors for which the dependency checks have been completed. In some embodiments, the task queue **622** may include a first queue group Q1, a dependency checker DPc, and a second queue group Q2.

[0221] In some embodiments, the task buffer **621**, the first queue group Q1, the dependency checker DPc, the second queue group Q2, and the runtime handle RH may be respectively referred to as a first queue circuit, a dependency checker circuit, a second queue circuit, and a runtime handle circuit, but for the sake of convenience, the terms are respectively unified as a first queue, a dependency checker, a second queue, and a runtime handle. Further, the task buffer **621**, the first queue group Q1, the dependency checker DPc, the second queue group Q2, and the runtime handle RH may each be implemented as circuits or circuitry. In some embodiments, the first queue group Q1 may be referred to as a dependency check waiting memory and the second queue group Q2 may be referred to as a process waiting memory.

[0222] The first queue group Q1 may store the task descriptors provided from the task buffer **621**. The task buffer **621** may transmit the task descriptors to the first queue group Q1 and generate transfer done reports TRrp. The task buffer **621** may transmit the transfer done reports TRrp to the done passage **630**. The transfer done reports TRrp may be reports on the tasks transmitted to the first queue group Q1.

[0223] The first queue group **Q1** may divide and store the task descriptors according to the types of the task descriptors. FIG. 17 shows n queues of the first queue group **Q1**. In some embodiments, n may be a natural number. In some embodiments, the first queue group **Q1** may be at least one.

[0224] In some embodiments, the first queue group **Q1** may include a plurality of queues **Q1_1**. . . . **Q1_n**. The i -th queue **Q1_i** of the first queue group **Q1** ($i=1 \dots n$) may store a i -th task descriptor **Tsk_di**.

[0225] The first to n -th task descriptors **Tsk_d1** to **Tsk_dn** may be of different types or the same type. Alternatively, some of the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** may be of the same type, and some may be of different types.

[0226] The dependency checker **DPc** may include a dependency check memory region. The dependency check memory region may represent whether reference tasks which tasks in the first queue group **Q1** depend on are completed. The dependency checker **DPc** may determine whether reference tasks which a currently-waiting task in the first queue group **Q1** depends on are completed based on the dependency information of the currently-waiting task and the dependency check memory region. The dependency checker **DPc** may receive a dependency update request **DFURQ**. The dependency update request **DFURQ** may notify the change of dependency as a completed task occurs according to a defined dependency between particular tasks. In some embodiments, each task descriptor may include a dependency field as to which tasks has a dependency on. In some embodiments, it must be updated in such a way that when the task included in the dependency field is completed, it will be removed from the dependency field.

[0227] Therefore, the dependency update request **DFURQ** may include an update request for the dependency field of a task descriptor.

[0228] The dependency checker **DPc** may transfer the currently-waiting task in the first queue group **Q1** to the second queue group **Q2** in response to determining that the reference tasks which the currently-waiting task depends on are completed. In some embodiments, dependency checker **DPc** may sequentially transmit the descriptors for which the dependency check has been completed, the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** to the second queue group **Q2**.

[0229] In some embodiments, the second queue group **Q2** may include a plurality of queues **Q2_1**. . . **Q2_n**. The i -th queue **Q2_i** of the second queue group **Q2** ($i=1 \dots n$) may store the i -th task descriptor **Tsk_di**. The number of second queues **Q2** may be the same as the number of first queues **Q1**.

[0230] The first queue group **Q1** of the task queue **622** may store the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** in a state prior to the dependency check, and the second queue group **Q2** of the task queue **622** may store the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** for which the dependency check has been completed.

[0231] The runtime handle **RH** may extract necessary information from each of the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** stored in the second queue group **Q2** and generate first to n -th task information **Tsk_d1'** to **Tsk_dn'**. The runtime handle **RH** may transfer the first to n -th task information **Tsk_d1'** to **Tsk_dn'** to the core global **500**. In some embodiments, the first to n -th task information **Tsk_d1'** to **Tsk_dn'** may correspond to the first to n -th task descriptors **Tsk_d1** to **Tsk_dn**, respectively. In some embodiments, each of the first to n -th task information **Tsk_d1'** to **Tsk_dn'** may also be the same as the first to n -th task descriptors **Tsk_d1** to **Tsk_dn**. However, the embodiment is not limited thereto.

[0232] The runtime handle **RH** may transmit check-in data **ChI** to the done passage **630**. The check-in data **ChI** may include the first to n -th task descriptors **Tsk_d1** to **Tsk_dn**. The check-in data **ChI** may be data informing the done passage **630** that the first to n -th task information **Tsk_d1'** to **Tsk_dn'** corresponding to the first to n -th task descriptors **Tsk_d1** to **Tsk_dn** have left the task passage **620** and been transferred to the core global **500** for processing. The done passage **630** may monitor whether the task descriptors are performed according to the check-in data **ChI**.

[0233] The first to n -th task descriptors **Tsk_d1** to **Tsk_dn** may be configured to include wait fields.

A wait field may be an item pre-specified by software. A task descriptor in which the wait field is set may not be converted into task information and transferred to the core global **500** but wait in the second queue group Q2 at a check-in time. Here, the check-in time may refer to a time at which the preceding task descriptors have all been transferred to the core global **500**, and the check-in may refer to transferring the task information to the core global **500** and transferring the corresponding task descriptors to the done passage **630**.

[0234] The task descriptors in which the wait field is set will have a waiting state. The wait field may be a means for controlling the execution time of tasks, and the task flow and execution time of tasks may be controlled via the runtime handle RH.

[0235] In some embodiments, the wait field in the task or the task descriptor of the task may indicate whether the task has to wait for a run signal to be checked in. For example, when the wait field is set equal to a first value, the task may indicate that the run signal is not required. When the wait field is set equal to a second value, the task may indicate that the run signal is required. In some embodiments, the first value and the second value are equal to 0 and 1, respectively. In some embodiments, the task or the task descriptor including the wait field may indicate that the task has to wait for a run signal to be checked in, and the task or the task descriptor not including the wait field may indicate that the task does not have to wait for a run signal to be checked in.

[0236] FIG. **18** is a block diagram for specifically illustrating the function of the runtime handle.

[0237] Referring to FIG. **18**, a plurality of task descriptors for which the dependency check has been completed may be stored in the plurality of queues Q2_1 . . . Q2_n of the second queue group Q2. The i-th queue Q2_i (i=1 . . . n) of the second queue group Q2 may include a plurality of task descriptors Tsk_d11 . . . Tsk_dik. The plurality of task descriptors Tsk_d11 . . . Tsk_d1k may be in a state of being stored in order in the 1st queue Q2_1 of the second queue group Q2. Here, k is a natural number. Likewise, the i-th queue Q2_i (i=1 . . . n) of the second queue group Q2 may include a plurality of task descriptors Tsk_di1 . . . Tsk_dik. In an example embodiment, the wait field may have been set in the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2, the 3rd task descriptor Tsk_d23 included in the 2nd queue Q2_2 of the second queue group Q2, and the k-th task descriptor Tsk_dnk included in the n-th queue Q2_n of the second queue group Q2, and the wait field may not have been set in the rest of the task descriptors.

[0238] In the 1st queue Q2_1 of the second queue group Q2, the 1st task descriptor Tsk_d11 included in the 1st queue Q2_1 of the second queue group Q2 is in a state where the wait field is not set and may be converted into task information and transferred to the core global **500** at the check-in time, and corresponding check-in data ChI may be transferred to the done passage **630**. In contrast, the runtime handle RH may not check in the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 but cause it to wait in the 1st queue Q2_1 of the second queue group Q2 even if the 2nd task descriptor Tsk_d12 in which the wait field is set may correspond to the check-in time.

[0239] As the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 waits in the 1st queue Q2_1 of the second queue group Q2 in a waiting state, other task descriptors Tsk_d13 to Tsk_dik following the 2nd task descriptor Tsk_d12 cannot be checked in and continue to wait in the 1st queue Q2_1 of the second queue group Q2.

[0240] The runtime handle RH may control all the run times of the plurality of queues Q2_1 . . . Q2_n of the second queue group Q2. The 1st and 2nd task descriptors Tsk_d21 and Tsk_d22 preceding the 3rd task descriptor Tsk_d23 included in the 2nd queue Q2_2 of the second queue group Q2 may each be converted into task information and transferred to the core global **500** at the check-in time, and corresponding check-in data ChI may be transferred to the done passage **630**. In contrast, the 3rd task descriptor Tsk_d23 included in the 2nd queue Q2_2 of the second queue group Q2 cannot be checked in by the runtime handle RH and waits in the 2nd queue Q2_2 of the second queue group Q2.

[0241] In addition, task descriptors preceding the k-th task descriptor Tsk_dnk included in the n-th queue Q2_n of the second queue group Q2 may pass through the runtime handle RH, be converted into task information, and transferred to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630, but the k-th task descriptor Tsk_dnk included in the n-th queue Q2_n of the second queue group Q2 including the wait field will wait in the n-th queue Q2_n of the second queue group Q2.

[0242] In some embodiments, the runtime handle RH may check whether the task descriptors stored in the plurality of queues Q2_1 . . . Q2_n of the second queue group Q2 include the wait field, and determine the state of the task descriptors as a run state or a waiting state.

[0243] The runtime handle RH may receive a run signal Run from the command processor 7000. The run signal Run may be provided from the command processor 7000 via the control interconnection CI. The command processor 7000 may transfer the run signal Run in response to transmitting a task including a wait field, but embodiments of the disclosure are not limited thereto.

[0244] The command processor 7000 may transmit the run signal Run to the runtime handle RH at regular intervals. The runtime handle RH may change the task descriptors in the waiting state to the run state in response to the run signal Run. The runtime handle RH may be configured to receive the run signal Run and store the run signal Run for a period of time. The runtime handle RH may include at least one of at least one register Rs or counter Rc for receiving and storing the run signal Run, but embodiments of the disclosure are not limited thereto.

[0245] In the following, a process in which the runtime handle RH performs wait-run control according to the run signal Run will be described in greater detail with reference to FIGS. 19 to 21.

[0246] FIG. 19 is an example diagram for illustrating a process of processing task descriptors included the 1st queue Q2_1 of the second queue group Q2 in response to a run signal.

[0247] Referring to FIG. 19, a process in which the task descriptors Tsk_d11 to Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 are processed can be checked according to the change in time. Here, the times t0, t1, t2, and t3 are example times. The 2nd task descriptor Tsk_d12 in the 1st queue Q2_1 of the second queue group Q2 includes a wait field, but the rest of the task descriptors do not include a wait field.

[0248] At time t0, the 1st task descriptor Tsk_d11 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 1st task descriptor Tsk_d11 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 1st task descriptor Tsk_d11, which does not include a wait field, is in a run state. The runtime handle RH may convert the 1st task descriptor Tsk_d11 into 1st task information Tsk_d11' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0249] At time t1, the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 2nd task descriptor Tsk_d12 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may determine that the 2nd task descriptor Tsk_d12 including the wait field is in a waiting state, and may control the 2nd task descriptor Tsk_d12 to wait in the 1st queue Q2_1 of the second queue group Q2. In addition, the runtime handle RH can check at the register Rs whether the run signal Run is received. At time t1, the run signal Run has not yet been received. Therefore, the runtime handle RH continues to maintain the waiting state of the 2nd task descriptor Tsk_d12. As the 2nd task descriptor Tsk_d12 waits, the 3rd task descriptor Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 that follows also waits in the 1st queue Q2_1 of the second queue group Q2.

[0250] At time t2, the run signal Run is received, and the state of the register Rs may be changed upon reception of the run signal Run. In an embodiment, the state of the register Rs that has received the run signal Run is defined as an active state, and the state of the register Rs that has not received the run signal Run is defined as a basic state. For example, the run signal Run may be a 1-

bit signal, and the register Rs may be changed from the basic state (0) to the active state (1) in response to the run signal Run. The runtime handle RH may release the waiting state of the 2nd task descriptor Tsk_d12 in response to the run signal Run, and convert it to the run state. The runtime handle RH may convert the 2nd task descriptor Tsk_d12 into 2nd task information Tsk_d12' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0251] At time t3, as the run signal Run is utilized to release the waiting state of the 2nd task descriptor Tsk_d12, it can be seen that the state of the register Rs has been restored from the active state to the basic state. The 3rd task descriptor Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. The runtime handle RH may convert the 3rd task descriptor Tsk_d13 that does not include the wait field into 3rd task information Tsk_d13' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0252] FIG. 20 is an example diagram for illustrating a process of processing task descriptors included in the 1st queue of the second queue group in response to a run signal received in advance. The 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 includes a wait field, but the rest of the task descriptors do not include a wait field.

[0253] Referring to FIG. 20, a process in which the task descriptors Tsk_d11 to Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 are processed can be checked according to the change in time.

[0254] At time t0, the 1st task descriptor Tsk_d11 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 1st task descriptor Tsk_d11 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 1st task descriptor Tsk_d11, which does not include a wait field, is in a run state, convert the 1st task descriptor Tsk_d11 into 1st task information Tsk_d11', and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0255] In addition, at time t0, the register Rs may be in the active state in response to the reception of the run signal Run. For the 1st task descriptor Tsk_d11 that does not include the wait field, the check-in process proceeds regardless of the state of the register Rs.

[0256] At time t1, the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 2nd task descriptor Tsk_d12 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 2nd task descriptor Tsk_d12 including the wait field is in the waiting state, and may cause the 2nd task descriptor Tsk_d12 to wait in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may release the waiting state of the 2nd task descriptor Tsk_d12 via the run signal

[0257] Run received in advance. The runtime handle RH may convert the 2nd task descriptor Tsk_d12 into 2nd task information Tsk_d12' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630. Since the runtime handle RH checked in the 2nd task descriptor Tsk_d12, the runtime handle RH may clear the register Rs.

[0258] At time t2, the 3rd task descriptor Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. The runtime handle RH may convert the 3rd task descriptor Tsk_d13 that does not include the wait field into 3rd task information Tsk_d13' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630. Further, at time t2, the run signal Run may be received in advance, so that the register Rs may be converted into the active state. The active state of the register Rs may be maintained continuously for some time, and a check-in process may be performed for the task descriptors that do not include the wait field regardless of the state of the register Rs.

[0259] At time t_3 , it can be seen that the register Rs continues to maintain the active state. The 4th task descriptor Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. The runtime handle RH may convert the 4-th task descriptor Tsk_d14 that does not include the wait field into 4th task information Tsk_d14' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0260] In some embodiments, a plurality of registers Rs may be configured to receive a plurality of run signals Run. In an embodiment, the runtime handle RH is that the run signals Run received respectively by the plurality of registers Rs may be processed individually in response to the task descriptors in the waiting state.

[0261] FIG. 21 is an example diagram for illustrating a process of processing task descriptors included in the 1st queue Q2_1 of the second queue group Q2 in response to a plurality of run signals. The 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 and the 4th task descriptor Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 include a wait field, but the rest of the task descriptors do not include a wait field.

[0262] As shown in FIG. 21, the runtime handle RH may comprise a plurality of registers. In some embodiments, the plurality of registers may be associated with a respective one of a plurality of users. In some embodiments, the plurality of task descriptors may belong to a respective one of the plurality of users.

[0263] For example, as shown in FIG. 21, the plurality of task descriptors belongs to a respective one of a user 1 and a user 2 and the plurality of registers may include a first register Rs1 for the user 1 and a second register Rs2 for the user 2. The 1st task descriptor Tsk_d11 and the 2nd task descriptor Tsk_d12 belong to the user 1 and the 3rd task descriptor Tsk_d13 and the 4th task descriptor Tsk_d14 belong to the user 2.

[0264] Referring to FIG. 21, at time t_0 , the 1st task descriptor Tsk_d11 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 1st task descriptor Tsk_d11 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 1st task descriptor Tsk_d11, which does not include a wait field, is in a run state, convert the 1st task descriptor Tsk_d11 into 1st task information Tsk_d11', and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0265] Further, at time t_0 , a plurality of run signals Run may be received. For example, two run signals Run for the user 1 and the user 2 may be received by a first register Rs1 and a second register Rs2, respectively. It can be confirmed by the received run signals Run that the first register Rs1 and the second register Rs2 are in active states. For the 1st task descriptor Tsk_d11 that does not include the wait field, the check-in process proceeds regardless of the states of the first and second registers Rs1 and Rs2.

[0266] At time t_1 , the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 2nd task descriptor Tsk_d12 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 2nd task descriptor Tsk_d12 including the wait field is in the waiting state, and may cause the 2nd task descriptor Tsk_d12 to wait in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may release the waiting state of the 2nd task descriptor Tsk_d12 via the run signal Run of the first register Rs1 received in advance. The runtime handle RH may convert the 2nd task descriptor Tsk_d12 into 2nd task information Tsk_d12' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630. Since the runtime handle RH checked in the 2nd task descriptor Tsk_d12 for the user 1, the runtime handle RH may clear the register Rs1.

[0267] At time t_2 , the 3rd task descriptor Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. The runtime handle RH may convert the 3rd

task descriptor Tsk_d13 that does not include the wait field into 3rd task information Tsk_d13' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0268] At time t3, the 4th task descriptor Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 4th task descriptor Tsk_d14 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 4th task descriptor Tsk_d14 including the wait field is in the waiting state, and may cause the 4th task descriptor Tsk_d14 to wait in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may release the waiting state of the 4th task descriptor Tsk_d14 via the run signal Run of the second register Rs2 received in advance. The runtime handle RH may convert the 4th task descriptor Tsk_d14 into 4th task information Tsk_d14' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630. Since the runtime handle RH checked in the 4th task descriptor Tsk_d14 for the user 2, the runtime handle RH may clear the register Rs2.

[0269] In some embodiments, the runtime handle RH may include a counter Rc that can receive and process a plurality of run signals Run. The runtime handle RH may release the waiting state of the task descriptors by individually utilizing the plurality of run signals Run received by the counter Rc.

[0270] FIG. 22 is an example diagram for illustrating a process of processing task descriptors included in the 1st queue Q2_1 of the second queue group Q2 through a plurality of run signals received by a counter. The 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second queue group Q2 and the 4th task descriptor Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 include a wait field, but the rest of the task descriptors do not include a wait field.

[0271] Referring to FIG. 22, at time t0, the 1st task descriptor Tsk_d11 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 1st task descriptor Tsk_d11 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 1st task descriptor Tsk_d11, which does not include a wait field, is in a run state, convert the 1st task descriptor Tsk_d11 into 1st task information Tsk_d11', and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0272] Further, at time t0, the run signal Run may be received by the counter Rc. It can be seen by the received run signal Run that one run signal Run has been received by the counter Rc. For the 1st task descriptor Tsk_d11 that does not include the wait field, the check-in process proceeds regardless of the state of the counter Rc.

[0273] At time t1, one run signal Run may be received by the counter Rc, and it can be seen that the state of the counter Rc has changed from 1 to 2.

[0274] At time t1, the 2nd task descriptor Tsk_d12 included in the 1st queue Q2_1 of the second

[0275] queue group Q2 may correspond to the check-in time. In some embodiments, the 2nd task descriptor Tsk_d12 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 2nd task descriptor Tsk_d12 including the wait field is in the waiting state, and may cause the 2nd task descriptor Tsk_d12 to wait in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may release the waiting state of the 2nd task descriptor Tsk_d12 via the run signal Run of the counter Rc received in advance. The runtime handle RH may convert the 2nd task descriptor Tsk_d12 into 2nd task information Tsk_d12' and transfer it to the core global 500, and corresponding check-in data ChI may be transferred to the done passage 630.

[0276] At time t2, because the 2nd task information Tsk_d12' which has to wait the run signal run is transferred to the core global 500 and corresponding check-in data ChI may be transferred to the

done passage **630**, the counter Rc may decrease by 1 and the value of the counter Rc has changed from 2 to 1.

[0277] At time t₂, the 3rd task descriptor Tsk_d13 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. The runtime handle RH may convert the 3rd task descriptor Tsk_d13 that does not include the wait field into 3rd task information Tsk_d13' and transfer it to the core global **500**, and corresponding check-in data ChI may be transferred to the done passage **630**.

[0278] At time t₃, the 4th task descriptor Tsk_d14 included in the 1st queue Q2_1 of the second queue group Q2 may correspond to the check-in time. In some embodiments, the 4th task descriptor Tsk_d14 may correspond to the most preceding task descriptor in the 1st queue Q2_1 of the second queue group Q2. The runtime handle RH may determine that the 4th task descriptor Tsk_d14 including the wait field is in the waiting state, and may cause the 4th task descriptor Tsk_d14 to wait in the 1st queue Q2_1 of the second queue group Q2. However, the runtime handle RH may release the waiting state of the 4th task descriptor Tsk_d14 via the run signal Run of the counter Rc received in advance. The runtime handle RH may convert the 4th task descriptor Tsk_d14 into 4th task information Tsk_d14' and transfer it to the core global **500**, and corresponding check-in data ChI may be transferred to the done passage **630**.

[0279] In addition, the runtime handle RH may perform not only the wait-run control of individual task descriptors, but also pause-resume control for temporarily pausing and resuming the operation of at least one second queue.

[0280] In some embodiments, the runtime handle RH may temporarily pause checking in task descriptors in at least one second queue if an abnormality or overload is expected in the operation of the neural core **100**.

[0281] In some embodiments, the runtime handle RH may perform the pause-resume control based on danger signals provided from the done passage **630**. Here, the danger signal may be an event according to at least one of hardware or software errors, log-related events, or tasks performed without descriptors, and such events may be provided by the event monitor EM of the done passage **630**, which will be described later, but the embodiments of the disclosure are not limited thereto.

[0282] Moreover, the runtime handle RH may perform the pause-resume control in response to an overload of the task queue **622** or an overload of the report queue of the reporting manager RM, which will be described later.

[0283] According to the pause control of the runtime handle RH, it is possible to temporarily pause the task information corresponding to the task descriptor from being transferred to the core global **500** and being processed. The runtime handle RH can release the paused state and resume the processing of the waiting task descriptor when the overload is resolved.

[0284] FIG. **23** is a block diagram for illustrating the done passage of FIG. **15** in detail.

[0285] Referring to FIG. **23**, the done passage **630** may include a check-in buffer Cib, a dependency setter DP, a timeout monitor ToM, an event monitor EM, and a report managing module **631**. The check-in buffer Cib, the dependency setter DP, the timeout monitor ToM, the event monitor EM, and the report managing module **631** may be referred to respectively as a check-in buffer circuit, a dependency setter circuit, a timeout monitor circuit, an event monitor circuit, and a report managing circuit, but for the sake of convenience, the terms are respectively unified as a check-in buffer, a dependency setter, a timeout monitor, an event monitor, and a report managing module. In addition, the check-in buffer Cib, the dependency setter DP, the timeout monitor ToM, the event monitor EM, and the report managing module **631** may each be implemented as a circuit or circuitry.

[0286] The check-in buffer Cib may receive the check-in data ChI. The check-in buffer Cib may include first to n-th check-in buffers Cib_1 to Cib_n. The check-in buffer Cib may store the first to n-th task descriptors Tsk_d1 to Tsk_dn stored in the check-in data ChI. The first to n-th check-in buffers Cib_1 to Cib_n may store the first to n-th task descriptors Tsk_d1 to Tsk_dn, respectively.

Through this, the check-in buffer Cib may perform check-in of the first to n-th task descriptors Tsk_d1 to Tsk_dn.

[0287] In some embodiments, the i-th check-in buffer Cib_i (i=1 . . . n) may store the i-th task descriptor Tsk_di. A number of check-in buffers Cib may be the same as a number of first queues Q1 and a number of second queues Q2.

[0288] The check-in buffer Cib may receive completion signals from the core global 500. In some embodiments, the completion signals may include first to n-th completion signals Tsk_d1d to Tsk_dnd. The first to n-th completion signals Tsk_d1d to Tsk_dnd may be completion signals for each of the first to n-th task descriptors Tsk_d1 to Tsk_dn, respectively. The first to n-th completion signals Tsk_d1d to Tsk_dnd may be received by the first to n-th check-in buffers Cib_1 to Cib_n, respectively. In some embodiments, the i-th check-in buffer Cib_i (i=1 . . . n) may receive the i-th completion signal Tsk_did.

[0289] The dependency setter DPs may receive the completion signals from the check-in buffer Cib and generate dependency update requests DFURQ. In some embodiments, depending on which task corresponding to what task descriptor has been completed, the dependency setter DPs may generate a dependency update request DFURQ. The dependency setter DPs may transmit the dependency update request DFURQ to the task passage 620.

[0290] The dependency setter DPs may check out each of the first to n-th task descriptors Tsk_d1 to Tsk_dn according to the completion signals. Accordingly, the dependency setter DPs may generate checkout reports COrp about which tasks have been completed and checked out. The dependency setter DPs may transmit the checkout reports COrp to the report managing module 631.

[0291] In some embodiments, the check-in can be defined as a procedure in which a task descriptor is registered before the task descriptor is processed, and the check-out can be defined as a procedure in which the task descriptor is completely processed and de-registered.

[0292] As the dependency setter DPs transmits the dependency update requests DFURQ to the task passage 620, the dependency checker DPc of the task passage 620 may transmit task descriptors in sequence according to the dependencies of the task descriptors.

[0293] The embodiment can minimize the overhead required for communication with the command processor 7000 by allowing the task manager 600 to directly perform dependency checking and setting without the command processor 7000 being solely responsible for processing according to the dependencies. Accordingly, the performance and speed of the neural processing device 1 in accordance with the embodiment can be dramatically improved.

[0294] The timeout monitor ToM may receive timeout detection signals TDec from the check-in buffer Cib. The timeout detection signal TDec may be a signal regarding whether a time interval between a check-in time and a check-out time exceeds a preset threshold time. Here, the check-out time may refer to a time when execution of the corresponding task is completed. The check-in buffer Cib may monitor execution times of the tasks corresponding to the checked-in task descriptors. The check-in buffer Cib may compare the execution times of the tasks with the threshold time and determine whether to generate the timeout detection signal TDec. The check-in buffer Cib may generate the timeout detection signal TDec if the execution time calculated from the check-in time exceeds the threshold time. In some embodiments, if the execution of the task has not been completed by the threshold time, the timeout detection signal TDec may be generated. The first to n-th check-in buffers Cib_1 to Cib_n will check whether the checked-in task descriptors have been executed, respectively, and generate the timeout detection signals TDec if the execution times exceed the threshold time.

[0295] In an embodiment, the threshold time may be set individually according to the tasks. In some embodiments, the threshold time may be set differently according to the types of tasks. Tasks corresponding to memory operations may be set to spend a shorter threshold time than tasks corresponding to computations. However, embodiments of the disclosure are not limited thereto.

[0296] Further, in an embodiment, whether to generate a timeout report may be set individually according to the task. In some embodiments, the command processor **7000** may be set not to be provided with timeout reports for at least some of the tasks transferred to the task manager **600**. As it is possible to set not to perform timeout monitoring for all tasks, the burden of timeout monitoring on the task manager **600** can be reduced. By way of example, it can be set that timeout reports are generated for tasks corresponding to computations, but timeout reports are not generated for tasks corresponding to memory operations. However, embodiments of the disclosure are not limited thereto.

[0297] The timeout monitor ToM may generate timeout reports TOrp according to the timeout detection signals TDec. The timeout monitor ToM may transmit the generated timeout reports TOrp to the report managing module **631**.

[0298] In an embodiment, the event monitor EM may sense events inside the neural processor **1000** and generate event reports Erp. Events may result from at least one of hardware or software errors, log-related events, tasks with erroneous descriptor, tasks with erroneous dependency information, or tasks performed without descriptors. Embodiments of the disclosure are not limited thereto. The event monitor EM may sense whether such events have occurred and generate event reports Erp according to event sensing signals. The generated event reports Erp may be provided to the report managing module **631**.

[0299] The report managing module **631** may receive at least one of a transfer done report TRrp, an event report Erp, a checkout report COrp, or a timeout report TOrp, and generate a completion report DNrp.

[0300] In an embodiment, the checkout report COrp may correspond to a report that allows the command processor **7000** to confirm that the task transmitted to the neural core **100** has been normally processed and checked out.

[0301] The transfer done report TRrp may correspond to a report that allows the command processor **7000** to confirm that a task has been normally provided to the task passage **620** and a task descriptor has been generated.

[0302] The event report Erp may correspond to a report that allows the command processor **7000** to confirm that an event has occurred according to at least one of hardware or software errors, log-related events, or tasks performed without descriptors.

[0303] The timeout report TOrp may correspond to a report that allows the command processor **7000** to confirm that processing for a particular task is delayed beyond a set threshold time.

[0304] In an embodiment, the checkout report COrp, the transfer done report TRrp, the event report Erp, and the timeout report TOrp may each be generated independently. For example, even if the execution of a task related to a particular task descriptor is delayed and a timeout report TOrp is generated, a checkout report COrp may be generated independently of the generation of the timeout report TOrp when the execution of the corresponding task is completed.

[0305] Here, the completion report DNrp may be generated based on the checkout report COrp. Since the completion report DNrp includes at least the checkout report, whether a task has been normally performed can be transferred to the command processor **7000**. Through the generation and transfer of such a completion report DNrp, it can be checked whether a task has been performed normally, and it is possible to prevent a delay in performing a particular task from being prolonged according to dependencies.

[0306] In addition, the completion report DNrp may be configured to further include at least one of the transfer done report TRrp, the event report Erp, or the timeout report TOrp, and it is possible to comprehensively report to the command processor **7000** whether a timeout has occurred, whether a transfer has been completed, whether an event has occurred, and the like, along with whether a task has been performed.

[0307] In an embodiment, whether to generate such a completion report DNrp may be set individually according to the tasks. In some embodiments, the command processor **7000** may be set

in advance such that a particular task will not receive a completion report DNrp related thereto. Accordingly, since completion reports may not be generated for all tasks, the burden of data processing on the task manager **600** may be reduced, and as the command processor **7000** may not receive all completion reports, the management efficiency of the command processor **7000** may be increased.

[0308] However, if an important warning related to task execution or a request for a part requiring correction is required to the command processor **7000**, such as a timeout report TOrp, corresponding information may be required to be transferred as a completion report DNrp. In an embodiment of the disclosure, if modification is needed for a task descriptor that is set not to generate a completion report, the timeout monitor ToM may transmit a modification request signal Mos to the check-in buffer in which the corresponding task descriptor is stored, and the corresponding task descriptor is modified to generate a completion report.

[0309] FIG. **24** is an example diagram for illustrating a first task descriptor and a second task descriptor that are set differently as to whether to generate a completion report.

[0310] Referring to FIG. **24**, in an embodiment, the tasks may include a first task that is set to generate a completion report and a second task that is set not to generate a completion report. The first task descriptor Tsk_d1 generated in correspondence with the first task is in a state of transferring the first task information Tsk_d1' to the core global **500** and being checked in at the first check-in buffer Cib_1. The second task descriptor Tsk_d2 generated in correspondence with the second task is in a state of transferring the second task information Tsk_d2' to the core global **500** and being checked in at the second check-in buffer Cib_2. The first task information Tsk_d1' and the second task information Tsk_d2' may be provided to the neural core **100**, respectively, via the core global **500**. The first check-in buffer Cib_1 is in a state of waiting for transfer of a first completion signal Tsk_d1d corresponding to the first task descriptor Tsk_d1, and the second check-in buffer Cib_2 is in a state of waiting for transfer of a second completion signal Tsk_d2d corresponding to the second task descriptor Tsk_d2.

[0311] Here, when the first completion signal Tsk_d1d is transferred to the first check-in buffer Cib_1, a checkout process is performed for the first task descriptor Tsk_d1 set to generate a completion report DNrp. In some embodiments, the dependency setter DPs may generate a first dependency update request DFURQ1 according to the completion of execution of the first task descriptor Tsk_d1 and transfer the first dependency update request DFURQ1 to the task passage **620**. In addition, the dependency setter DPs generates a first checkout report COrp1 and transmits the generated first checkout report COrp1 to the report managing module **631**, so that the completion report DNrp is generated.

[0312] When the second completion signal Tsk_d2d is transferred to the second check-in buffer Cib_2, a checkout process is performed for the second task descriptor Tsk_d2 set not to generate a completion report DNrp. The dependency setter DPs may generate a second dependency update request DFURQ2 according to the completion of execution of the second task descriptor Tsk_d2 and transfer the second dependency update request DFURQ2 to the task passage **620**. However, the dependency setter DPs will not generate a second checkout report COrp2 and the completion report DNrp will not be generated.

[0313] Here, when the processing of the second task by the neural core **100** is delayed and a timeout report TOrp is generated, reporting on the generation of the timeout report TOrp must be performed to the command processor **7000**, and generation of a completion report DNrp may be required for this purpose. In an embodiment, the timeout monitor ToM may modify the second task to generate a completion report.

[0314] FIG. **25** is a flowchart for illustrating a process of performing modification on the second task descriptor that is set so that a completion report is not generated.

[0315] Referring to FIG. **25**, the second check-in buffer Cib_2 may determine that the execution of the second task is delayed and generate a second timeout detection signal TDec2, and transfer the

generated second timeout detection signal TDec2 to the timeout monitor ToM at S11. Here, the generation of the second timeout detection signal TDec2 may mean a state in which the second completion signal Tsk_d2d has not yet been received, and modification to the second task descriptor Tsk_d2 may be possible.

[0316] The timeout monitor ToM may generate a second timeout report TOrp2 according to the second timeout detection signal TDec2, and transfer the generated second timeout report TOrp2 to the report managing module 631 at S12.

[0317] Further, in order for the second timeout report TOrp2 to be included in the completion report DNrp and transferred to the command processor 7000, a modification request signal Mos for modifying the second task descriptor Tsk_d2 so as to generate the completion report DNrp may be transmitted to the second check-in buffer Cib_2 at S13. For the second task descriptor Tsk_d2 included in the second check-in buffer Cib_2, the completion report DNrp may be generated in response to the modification request, and a modified second task descriptor STsk_d2 is generated at S14.

[0318] FIG. 26 is an example diagram for illustrating a process of generating a completion report in response to a second checkout report of the second task descriptor modified according to FIG. 25.

[0319] Referring to FIG. 26, the second check-in buffer Cib_2 receives the second completion signal Tsk_d2d via the core global 500 after a threshold time, and a checkout process for the modified second task descriptor STsk_d2 is performed. The dependency setter DPs may generate a second dependency update request DFURQ2 according to the completion of execution of the modified second task descriptor STsk_d2 and transfer the second dependency update request DFURQ2 to the task passage 620. In addition, the dependency setter DPs generates a second checkout report COrp2, and transmits the generated second checkout report COrp2 to the report managing module 631, so that a completion report DNrp is generated. The report managing module 631 constructs the completion report DNrp by including all information on the second checkout report COrp2 and the second timeout report TOrp2 received, and transfers the constructed completion report DNrp to the command processor 7000.

[0320] FIG. 27 is a block diagram for illustrating the report managing module of FIG. 23 in detail.

[0321] Referring to FIG. 27, the report managing module 631 may include a transfer done report queue TQ, an event report queue EQ, a checkout report queue CQ, a timeout report queue TOQ, and a reporting manager RM. The transfer done report queue TQ, the event report queue EQ, the checkout report queue CQ, the timeout report queue TOQ, and the reporting manager RM may be referred to respectively as a transfer done report queue circuit, an event report queue circuit, a checkout report queue circuit, a timeout report queue circuit, and a reporting manager circuit, but for the sake of convenience, the terms are respectively unified as a transfer done report queue, an event report queue, a checkout report queue, a timeout report queue, and the reporting manager. In addition, the transfer done report queue TQ, the event report queue EQ, the checkout report queue CQ, the timeout report queue TOQ, and the reporting manager RM may each be implemented as a circuit or circuitry.

[0322] The transfer done report queue TQ may receive the transfer done report TRrp and transfer the transfer done report TRrp to the reporting manager RM.

[0323] The event report queue EQ may receive the event report Erp and transfer the event report Erp to the reporting manager RM. To prevent an overflow of the event report queue EQ, when the number of entries in the event report queue EQ is larger than a first threshold for the event report queue EQ, the report managing module 631 may generate and send a pause signal requiring the runtime handle RH to pause so that the runtime handle RH does not check in any further tasks or task descriptors regardless of the status of the register in the runtime handle RH. The report managing module 631 or the command processor 7000 may generate and send a resume signal requiring the runtime handle RH to resume so that the runtime handle RH can continue to perform a process of checking in tasks or task descriptors, when it detects that the number of entries in the

event report queue EQ is less than a second threshold for the event report queue EQ. The second threshold for the event report queue EQ may be equal to or smaller than the first threshold for the event report queue EQ. The resume signal may be triggered by software for a machine learning model which is processed by the CPU **20**.

[0324] The checkout report queue CQ may receive the checkout report COrp and transfer the checkout report COrp to the reporting manager RM. To prevent an overflow of the checkout report queue CQ, when the number of entries in the checkout report queue CQ is larger than a first threshold for the checkout report queue CQ, the report managing module **631** may generate and send a pause signal requiring the runtime handle RH to pause so that the runtime handle RH does not check in any further tasks or task descriptors regardless of the status of the register in the runtime handle RH. The report managing module **631** or the command processor **7000** may generate and send a resume signal requiring the runtime handle RH to resume so that the runtime handle RH can continue to perform a process of checking in tasks or task descriptors, when it detects that the number of entries in the checkout report queue CQ is less than a second threshold for the checkout report queue CQ. The second threshold for the checkout report queue CQ may be equal to or smaller than the first threshold for the checkout report queue CQ. The resume signal may be triggered by software for a machine learning model which is processed by the CPU **20**.

[0325] The timeout report queue TOQ may receive the timeout report TOrp and transfer the timeout report TOrp to the reporting manager RM. To prevent an overflow of the timeout report queue TOQ, when the number of entries in the timeout report queue TOQ is larger than a first threshold for the timeout report queue TOQ, the report managing module **631** may generate and send a pause signal requiring the runtime handle RH to pause so that the runtime handle RH does not check in any further tasks or task descriptors regardless of the status of the register in the runtime handle RH. The report managing module **631** or the command processor **7000** may generate and send a resume signal requiring the runtime handle RH to resume so that the runtime handle RH can continue to perform a process of checking in tasks or task descriptors, when it detects that the number of entries in the timeout report queue TOQ is less than a second threshold for the timeout report queue TOQ. The second threshold for the timeout report queue TOQ may be equal to or smaller than the first threshold for the timeout report queue TOQ. The resume signal may be triggered by software for a machine learning model which is processed by the CPU **20**.

[0326] The reporting manager RM may receive at least one of the transfer done report TRrp, the event report Erp, the checkout report COrp, or the timeout report TOrp, and may generate a completion report DNrp through this. The reporting manager RM may transfer the completion report DNrp to the command processor **7000**.

[0327] FIG. **28** is a diagram for illustrating data that the core global and the processing unit of FIG. **15** exchange.

[0328] Referring to FIG. **28**, the core global **500** may receive the table update request TURQ and transfer the table update request TURQ to the LSU **110**. In addition, the core global **500** may receive the task information Tsk_d' and transfer the task information Tsk_d' to the neural core **100**.

[0329] The neural core **100** may perform tasks and generate completion signals. The neural core **100** may transfer the completion signals to the core global **500**. The core global **500** may include a signal scheduler sgn_sch. The signal scheduler sgn_sch may receive the completion signal, schedule a transmission of the completion signal, and transmit the completion signal to the done passage **630**. The signal scheduler sgn_sch may be referred to as a signal scheduler circuit, but for the sake of convenience, the terms are unified as a signal scheduler. Further, the signal scheduler sgn_sch may be implemented as a circuit or circuitry.

[0330] FIG. **29** is a diagram for illustrating types of task descriptors stored in a first queue, a second queue, and a check-in buffer.

[0331] Referring to FIG. **29**, the 1st to 4th queues Q1_1 . . . Q1_4 of the first queue group Q1, the 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and first to fourth check-in buffers

Cib_1 to Cib_4 of a check-in buffer Cib may store particular types of task descriptors, respectively. The 1st to 4th queues Q1_1 . . . Q1_4 of the first queue group Q1, the 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and the first to fourth check-in buffers Cib_1 to Cib_4 may store different types of task descriptors.

[0332] For example, the 1st queue Q1_1 of the first queue group Q1, the 1st queue Q2_1 of the second queue group Q2, and the first check-in buffer Cib_1 may store task descriptors for computation (COMP), and the 2nd queue Q1_2 of the first queue group Q1, the 2nd queue Q2_2 of the second queue group Q2, and the second check-in buffer Cib_2 may store task descriptors for micro-DMA. In addition, the 3rd queue Q1_3 of the first queue group Q1, the 3rd queue Q2_3 of the second queue group Q2, and the third check-in buffer Cib_3 may store task descriptors for LP micro-DMA, and the 4th queue Q1_4 of the first queue group Q1, the 4th queue Q2_4 of the second queue group Q2, and the fourth check-in buffer Cib_4 may store task descriptors for ST micro-DMA. However, the embodiment is not limited thereto.

[0333] FIG. 30 is a diagram for illustrating a first queue, a second queue, and a check-in buffer of a neural processing device in accordance with some embodiments of the disclosure.

[0334] Referring to FIG. 30, the 1st to 4th queues Q1_1 . . . Q1_4 of the first queue group Q1, 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and first to fourth check-in buffers Cib_1 to Cib_4 of a check-in buffer Cib may store particular types of task descriptors, respectively. The 1st to 4th queues Q1_1 . . . Q1_4 of the first queue group Q1, the 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and the first to fourth check-in buffers Cib_1 to Cib_4 may store the same type of task descriptors.

[0335] For example, the 1st queue Q1_1 of the first queue group Q1, the 1st queue Q2_1 of the second queue group Q2, and the first check-in buffer Cib_1 may store task descriptors for first computation, and the 2nd queue Q1_2 of the first queue group Q1, the 2nd queue Q2_2 of the second queue group Q2, and the second check-in buffer Cib_2 may store task descriptors for second computation. In addition, the 3rd queue Q1_3 of the first queue group Q1, the 3rd queue Q2_3 of the second queue group Q2, and the third check-in buffer Cib_3 may store task descriptors for third computation, and the 4th queue Q1_4 of the first queue group Q1, the 4th queue Q2_4 of the second queue group Q2, and the fourth check-in buffer Cib_4 may store task descriptors for fourth computation. However, the embodiment is not limited thereto.

[0336] In some embodiments, the first to fourth computations may be exactly the same computation, or may be computations of the same type but different kinds in detail.

[0337] FIG. 31 is a diagram for illustrating a first queue, a second queue, and a check-in buffer of a neural processing device in accordance with some embodiments of the disclosure.

[0338] Referring to FIG. 31, the 1st to 4th queues Q1_1 . . . Q1_4 of a first queue group Q1, 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and first to fourth check-in buffers Cib_1 to Cib_4 of a check-in buffer Cib may store several types of task descriptors, respectively. The 1st to 4th queues Q1_1 . . . Q1_4 of a first queue group Q1, the 1st to 4th queues Q2_1 . . . Q2_4 of the second queue group Q2, and the first to fourth check-in buffers Cib_1 to Cib_4 may store different types of task descriptors or the same type of task descriptors.

[0339] FIG. 32 is a block diagram for illustrating in detail the structure of the neural processing device in accordance with some embodiments of the disclosure.

[0340] Referring to FIG. 32, a neural core 101 may have a CGRA structure, unlike a neural core 100. The neural core 101 may include an instruction memory 111_1, a CGRA L0 memory 111_2, a PE array 111_3, and a load/store unit (LSU) 111_4. The PE array 111_3 may include a plurality of processing elements interconnected by a mesh style network. The mesh style network may be two-dimensional, three-dimensional, or higher-dimensional. In the CGRA, the plurality of processing elements may be reconfigurable or programmable. The interconnection between the plurality of processing elements may be reconfigurable or programmable. In some embodiments, the interconnection between the plurality of processing elements may be statically reconfigurable or

programmable when the interconnection is fixed after the plurality of processing elements are configured or programed. In some embodiments, the interconnection between the plurality of processing elements may be dynamically reconfigurable or programmable when the interconnection is reconfigurable or programmable even after the plurality of processing elements are configured or programed.

[0341] The instruction memory **111_1** may receive and store instructions. The instruction memory **111_1** may sequentially store instructions internally, and provide the stored instructions to the PE array **111_3**. In some embodiments, the instructions may instruct the operation of first type of a plurality of processing elements **111_3a** included in each PE array **111_3**.

[0342] The CGRA L0 memory **111_2** may be located inside the neural core **101**, receive all input data required for tasks of the neural core **101**, and temporarily store the data. In addition, the CGRA L0 memory **111_2** may temporarily store output data calculated by the neural core **101** to transmit the data to the outside. The CGRA L0 memory **111_2** may serve as a cache memory of the neural core **101**.

[0343] The CGRA L0 memory **111_2** may send and receive data to and from the PE array **111_3**.

[0344] The CGRA L0 memory **111_2** may be a memory corresponding to L0 (level 0) that is lower than L1. In some embodiments, the L0 memory may be a private memory of the neural core **101** that is not shared. The CGRA L0 memory **111_2** may transmit data such as activations or weights, programs, and the like to the PE array **111_3**.

[0345] The PE array **111_3** may be a module that performs calculations. The PE array **111_3** may perform not only one-dimensional calculations but also two-dimensional or higher matrix/tensor calculations. The PE array **111_3** may include the first type of the plurality of processing elements **111_3a** and a second type of a plurality of processing elements **111_3b** therein.

[0346] The first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** may be arranged in rows and columns. The first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** may be arranged in m columns. In addition, the first type of the plurality of processing elements **111_3a** may be arranged in n rows, and the second type of the plurality of processing elements **111_3b** may be arranged in 1 rows. Accordingly, the first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing element **111_3b** may be arranged in (n+1) rows and m columns.

[0347] The LSU **111_4** may receive at least one of data, a control signal, or a synchronization signal from the outside via the local interconnection **200**. The LSU **111_4** may transmit at least one of the received data, control signal, or synchronization signal to the CGRA L0 memory **111_2**. Similarly, the LSU **111_4** may transfer at least one of the data, control signal, or synchronization signal to the outside via the local interconnection **200**.

[0348] The neural core **101** may have a CGRA (Coarse Grained Reconfigurable Architecture) structure. Accordingly, in the neural core **101**, each of the first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** of the PE array **111_3** may be connected to at least one of the CGRA L0 memory **111_2**, the instruction memory **111_1**, or the LSU **111_4**, respectively. In other words, the first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** do not have to be connected to all of the CGRA L0 memory **111_2**, the instruction memory **111_1**, and the LSU **111_4**, but may be connected to some thereof.

[0349] Further, the first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** may be different types of processing elements from each other. Accordingly, out of the CGRA L0 memory **111_2**, the instruction memory **111_1**, and the LSU **111_4**, the elements connected to the first type of the plurality of processing elements **111_3a** and the elements connected to the second type of the plurality of processing elements **111_3b** may be different from each other.

[0350] The neural core **101** of the disclosure having a CGRA structure enables high-level parallel calculations, and since direct data exchange between the first type of the plurality of processing elements **111_3a** and the second type of the plurality of processing elements **111_3b** is possible, the power consumption may be low. In addition, by including two or more types of processing elements, optimization according to various calculation tasks may also be possible.

[0351] For example, if the first type of the plurality of processing elements **111_3a** are processing elements that perform two-dimensional calculations, the second type of the plurality of processing elements **111_3b** may be processing elements that perform one-dimensional calculations. However, the embodiment is not limited thereto.

[0352] FIG. **33** is a diagram for illustrating a hierarchical structure of a command processor and a task manager of a neural processing device in accordance with some embodiments of the disclosure, and FIG. **34** is a diagram for illustrating a hierarchical structure of a command processor and task managers of a neural processing device in accordance with some embodiments of the disclosure.

[0353] Referring to FIGS. **33** and **34**, if the number of task managers **600** increases, it may be difficult for the command processor **7000** to manage all of the task managers **600**. Therefore, the neural processing device **1** in accordance with some embodiments of the disclosure may have a hierarchical structure in which a master task manager **600M** manages the plurality of task managers **600** and the command processor **7000** manages the master task managers **600M**.

[0354] Further, referring to FIG. **34**, levels below the master task manager **600M** may also be subdivided into a plurality. For example, a first sub-task manager **600s1** and a second sub-task manager **600s2** may form each layer. In some embodiments, one first sub-task manager **600s1** may manage at least one second sub-task manager **600s2**, and one master task manager **600M** may manage at least one first sub-task manager **600s1**. Additionally, several layers may be added below the second sub-task manager **600s2** as well.

[0355] In some embodiments, although three levels of the task managers **600**, the master task managers **600M**, and the command processor **7000** are shown in FIGS. **33** and **34**, the number of levels may be four or more. In other words, the depth of the hierarchical structure may vary as desired depending on the number of task managers **600**. In addition, the master task manager **600M**, the first sub-task manager **600s1**, and the second sub-task manager **600s2** may be referred to respectively as a master task manager circuit, a first sub-task manager circuit, and a second sub-task manager circuit, but for the sake of convenience, the terms are respectively unified as a master task manager, a first sub-task manager, and a second sub-task manager. Further, the master task manager **600M**, the first sub-task manager **600s1**, and the second sub-task manager **600s2** may each be implemented as a circuit or circuitry.

[0356] FIG. **35** is a block diagram for illustrating memory reconfiguration of a neural processing system in accordance with some embodiments of the disclosure.

[0357] Referring to FIG. **35**, the neural core SoC **10** may include first to eighth processing units **160a** to **160h** and an on-chip memory OCM. Although FIG. **35** illustrates eight processing units as an example, this is merely illustrative, and the number of processing units may vary as desired.

[0358] The on-chip memory OCM may include first to eighth L0 memories **120a** to **120h** and a shared memory **2000**.

[0359] The first to eighth L0 memories **120a** to **120h** may be used as private memories for the first to eighth processing units **160a** to **160h**, respectively. In other words, the first to eighth processing units **160a** to **160h** and the first to eighth L0 memories **120a** to **120h** may correspond to each other 1:1.

[0360] The shared memory **2000** may include first to eighth memory units **2100a** to **2100h**. The first to eighth memory units **2100a** to **2100h** may correspond to the first to eighth processing units **160a** to **160h** and the first to eighth L0 memories **120a** to **120h**, respectively. In some embodiments, the number of memory units may be eight, which is the same as the number of

processing units and L0 memories.

[0361] The shared memory **2000** may operate in one of two kinds of on-chip memory types. In other words, the shared memory **2000** may operate in one of a L0 memory type or a global memory type. In other words, the shared memory **2000** may implement two types of logical memories with one piece of hardware.

[0362] If the shared memory **2000** is implemented in the L0 memory type, the shared memory **2000** may operate as a private memory for each of the first to eighth processing units **160a** to **160h**, just like the first to eighth L0 memories **120a** to **120h**. The L0 memory can operate at a relatively higher clock speed compared with the global memory, and the shared memory **2000** may also use a relatively higher clock speed when operating in the L0 memory type.

[0363] If the shared memory **2000** is implemented in the global memory type, the shared memory **2000** may operate as a common memory used by the first processing unit **160a** and the second processing unit **160b** together. In some embodiments, the shared memory **2000** may be shared not only by the first to eighth processing units **160a** to **160h** but also by the first to eighth L0 memories **120a** to **120h**.

[0364] The global memory may generally use a lower clock compared with the L0 memory, but is not limited thereto. When the shared memory **2000** operates in the global memory type, the first to eighth processing units **160a** to **160h** may share the shared memory **2000**. In some embodiments, the shared memory **2000** may be connected to the volatile memory **32** of FIG. **2** via the global interconnection **6000** and may also operate as a buffer for the volatile memory **32**.

[0365] At least part of the shared memory **2000** may operate in the L0 memory type, and the rest may operate in the global memory type. In other words, the entire shared memory **2000** may operate in the L0 memory type, or the entire shared memory **2000** may operate in the global memory type. Alternatively, part of the shared memory **2000** may operate in the L0 memory type, and the rest may operate in the global memory type.

[0366] FIG. **36** is a block diagram showing an example of memory reconstruction of a neural processing system in accordance with some embodiments of the disclosure.

[0367] With reference to FIGS. **35** and **36**, first, third, fifth, and seventh dedicated areas **AE1**, **AE3**, **AE5**, and **AE7** for each of the first, third, fifth, and seventh processing units **160a**, **160c**, **160e**, and **160g** may include only the first, third, fifth, and seventh L0 memories **120a**, **120c**, **120e**, and **120g**, respectively. Further, second, fourth, sixth, and eighth dedicated areas **AE2**, **AE4**, **AE6**, and **AE8** for each of the second, fourth, sixth, and eighth processing units **160b**, **160d**, **160f**, and **160h** may include second, fourth, sixth, and eighth L0 memories **120b**, **120d**, **120f**, and **120h**, respectively. In addition, the second, fourth, sixth, and eighth dedicated areas **AE2**, **AE4**, **AE6**, and **AE8** may include the second, fourth, sixth, and eighth memory units **2100b**, **2100d**, **2100f**, and **2100h**. The first, third, fifth, and seventh memory units **2100a**, **2100c**, **2100e**, and **2100g** of the shared memory **2000** may be used as a common area **AC**.

[0368] The common area **AC** may be a memory shared by the first to eighth processing units **160a** to **160h**. The second dedicated area **AE2** may include a second L0 memory **120b** and a second memory unit **2100b**. The second dedicated area **AE2** may be an area in which the second L0 memory **120b** and the second memory unit **2100b** that are separated hardware-wise operate in the same manner and operate logically as one L0 memory. The fourth, sixth, and eighth dedicated areas **AE4**, **AE6**, and **AE8** may also operate in the same manner as the second dedicated area **AE2**.

[0369] The shared memory **2000** in accordance with the embodiment may convert an area corresponding to each processing unit into a logical L0 memory and a logical global memory of an optimized ratio and may use them. The shared memory **2000** may perform the adjustment of this ratio at runtime.

[0370] In some embodiments, each processing unit may perform the same task in some cases, but may perform different tasks in other cases as well. In some embodiments, the amount of the L0 memory and the amount of the global memory required for the tasks carried out by each processing

unit are inevitably different each time. Accordingly, if the composition ratio of the L0 memory and the shared memory is fixedly set as in the conventional on-chip memory, there may occur inefficiency due to the calculation tasks assigned to each processing unit.

[0371] Therefore, the shared memory **2000** of the neural processing device in accordance with the embodiment may set an optimal ratio of the L0 memory and the global memory according to calculation tasks during the runtime, and may enhance the efficiency and speed of calculation. FIG. **37** is an enlarged block diagram of a portion A of FIG. **35**.

[0372] With reference to FIGS. **35** and **37**, the shared memory **2000** may include a first L0 memory controller **122_1a**, a second L0 memory controller **122_1b**, a fifth L0 memory controller **122_1e**, a sixth L0 memory controller **122_1f**, the first to eighth memory units **2100a** to **2100h**, and a global controller **2200**. Other L0 memory controllers not shown may also be included in the embodiment, but the description thereof will be omitted for convenience.

[0373] The first L0 memory controller **122_1a**, the second L0 memory controller **122_1b**, the fifth L0 memory controller **122_1e**, the sixth L0 memory controller **122_1f**, and the global controller **2200** may be referred to respectively as a first L0 memory controller circuit, a second L0 memory controller circuit, a fifth L0 memory controller circuit, a sixth L0 memory controller circuit, and a global controller circuit. However, for the sake of convenience, the terms are respectively unified as a first L0 memory controller, a second L0 memory controller, a fifth L0 memory controller, a sixth L0 memory controller, and a global controller. In addition, the first L0 memory controller **122_1a**, the second L0 memory controller **122_1b**, the fifth L0 memory controller **122_1e**, the sixth L0 memory controller **122_1f**, and the global controller **2200** may each be implemented as a circuit or circuitry.

[0374] The first L0 memory controller **122_1a** may control the first L0 memory **120a**. In addition, the first L0 memory controller **122_1a** may control the first memory unit **2100a**. Specifically, when the first memory unit **2100a** is implemented in a logical L0 memory type, the control by the first L0 memory controller **122_1a** may be performed on the first memory unit **2100a**.

[0375] The second L0 memory controller **122_1b** may control the second L0 memory **120b**. Further, the second L0 memory controller **122_1b** may control the second memory unit **2100b**. In other words, when the second memory unit **2100b** is implemented in the logical L0 memory type, the control by the first L0 memory controller **122_1a** may be performed on the second memory unit **2100b**.

[0376] The fifth L0 memory controller **122_1e** may control the fifth L0 memory **120e**. Further, the fifth L0 memory controller **122_1e** may control the fifth memory unit **2100e**. In other words, when the fifth memory unit **2100e** is implemented in the logical L0 memory type, the control by the fifth L0 memory controller **122_1e** may be performed on the fifth memory unit **2100e**.

[0377] The sixth L0 memory controller **122_1f** may control the sixth L0 memory **120f**. Further, the sixth L0 memory controller **122_1f** may control the sixth memory unit **2100f**. In other words, when the sixth memory unit **2100f** is implemented in the logical L0 memory type, the control by the sixth L0 memory controller **122_1f** may be performed on the sixth memory unit **2100f**.

[0378] The global controller **2200** may control all of the first to eighth memory units **2100a** to **2100h**. Specifically, the global controller **2200** may control the first memory unit **2100a** to the eighth memory unit **2100h** when the first to eighth memory units **2100a** to **2100h** each operate logically in the global memory type (i.e., when they do not operate logically in the L0 memory type).

[0379] In other words, the first to eighth memory units **2100a** to **2100h** may be controlled by the first to eighth L0 memory controllers **122_1a** to **122_1h**, respectively, or may be controlled by the global controller **2200**, depending on what type of memory they are logically implemented.

[0380] If the L0 memory controllers including the first, second, fifth, and sixth L0 memory controllers **122_1a**, **122_1b**, **122_1e**, and **122_1f** control the first to eighth memory units **2100a** to **2100h**, respectively, the first to eighth L0 memory controllers **122_1a** to **122_1h** control the first to

eight memory units **2100a** to **2100h** in the same manner as the first to eighth L0 memories **120a** to **120h**, and thus, can control them as the private memory of the first to eighth processing units **160a** to **160h**. Accordingly, the first to eighth memory units **2100a** to **2100h** may operate at clock frequencies corresponding to the clock frequencies of the first to eighth processing units **160a** to **160h**.

[0381] The L0 memory controllers including the first L0 memory controller **122_1a**, the second L0 memory controller **122_1b**, the fifth L0 memory controller **122_1e**, and the sixth L0 memory controller **122_1f** may each include the LSU **110** of FIG. 8.

[0382] If the global controller **2200** controls at least one of the first to eighth memory units **2100a** to **2100h**, respectively, then the global controller **2200** may control the first to eighth memory units **2100a** to **2100h** as the global memory of the first to eighth processing units **160a** to **160h**, respectively. Accordingly, at least one of the first to eighth memory units **2100a** to **2100h** may operate at a clock frequency independent of the clock frequencies of the first to eighth processing units **160a** to **160h**, respectively. In some embodiments, if the global controller **2200** controls the i-th memory unit among the first to eighth memory units **2100a** to **2100h**, the global controller **2200** may control the i-th memory unit as the global memory of the i-th processing unit, and the i-th memory unit may operate at a clock frequency independent of the clock frequency of the i-th processing unit. However, the embodiment is not limited thereto.

[0383] The global controller **2200** may connect the first to eighth memory units **2100a** to **2100h** to the global interconnection **6000** of FIG. 3. The first to eighth memory units **2100a** to **2100h** may exchange data with the off-chip memory **30** of FIG. 2 by the control of the global controller **2200** or may respectively exchange data with the first to eighth L0 memories **120a** to **120h**.

[0384] Each of the first to eighth memory units **2100a** to **2100h** may include at least one memory bank. The first memory unit **2100a** may include at least one first memory bank **2110a**. The first memory banks **2110a** may be areas obtained by dividing the first memory unit **2100a** into certain sizes. The first memory banks **2110a** may all be memory devices of the same size. However, the embodiment is not limited thereto. FIG. 22 illustrates that four memory banks are included in one memory unit.

[0385] Similarly, the second, fifth, and sixth memory units **2100b**, **2100e**, and **2100f** may include at least one second, fifth, and sixth memory banks **2110b**, **2110e**, and **2110f**, respectively.

[0386] In the following, the description will be made based on the first memory banks **2110a** and the fifth memory banks **2110e**, which may be the same as other memory banks including the second and sixth memory banks **2110b** and **2110f**.

[0387] The first memory banks **2110a** may each operate logically in the L0 memory type or operate logically in the global memory type. In some embodiments, the first memory banks **2110a** may operate independently of the other memory banks in the first memory unit **2100a**. However, the embodiment is not limited thereto.

[0388] If each memory bank operates independently, the first memory unit **2100a** may include a first area operating in the same manner as the first L0 memory **120a** and a second area operating in a different manner from the first L0 memory **120a**. In some embodiments, the first area and the second area do not necessarily coexist, but any one area may take up the entire first memory unit **2100a**.

[0389] Likewise, the second memory unit **2100b** may include a third area operating in the same manner as the second L0 memory **120b** and a fourth area operating in a different manner from the second L0 memory **120b**. In some embodiments, the third area and the fourth area do not necessarily coexist, and any one area may take up the entire first memory unit **2100a**.

[0390] In some embodiments, the ratio of the first area to the second area may be different from the ratio of the third area to the fourth area. However, the embodiment is not limited thereto. Therefore, the ratio of the first area to the second area may be the same as the ratio of the third area to the fourth area. In other words, the memory composition ratio in each memory unit may vary as

desired.

[0391] In general, in the case of the conventional system-on-chip, the on-chip memory except for high-speed L0 memory was often composed of high-density, low-power SRAM. This is because SRAM has high efficiency in terms of chip area and power consumption relative to required capacity. However, with the conventional on-chip memory, the processing speed slowed down significantly as was inevitable in the case where tasks that require more data quickly than the predetermined capacity of the L0 memory, and, even when the need for the global memory is not great, there is no way to utilize the remaining global memory, resulting in inefficiency.

[0392] On the other hand, the shared memory **2000** in accordance with some embodiments of the disclosure may be controlled selectively by any one of the two controllers depending on the case. In the case depicted, the shared memory **2000** may be controlled not only as a whole by a determined one of the two controllers but also independently for each memory unit or each memory bank.

[0393] Through this, the shared memory **2000** in accordance with the embodiment can obtain an optimal memory composition ratio according to calculation tasks during the runtime and can perform faster and more efficient calculation tasks. In the case of a processing unit specialized in artificial intelligence, the required sizes of L0 memory and global memory may vary for each particular application. Moreover, even for the same application, the required sizes of L0 memory and global memory may vary for each layer when a deep learning network is used. In the shared memory **2000**, in accordance with the embodiment, the composition ratio of the memory can be changed during runtime even when calculation steps change according to each layer, making fast and efficient deep learning tasks possible.

[0394] FIG. **38** is a diagram for illustrating the first memory bank of FIG. **37** in detail. Although FIG. **38** illustrates the first memory bank **2110a**, other memory banks may also have the same structure as the first memory bank **2110a**.

[0395] Referring to FIG. **38**, the first memory bank **2110a** may include a cell array Ca, a bank controller Bc, a first path unit P1, and a second path unit P2.

[0396] In some embodiments, the bank controller Bc, the first path unit P1, and the second path unit P2 may be referred to respectively as a bank controller circuit, a first path unit circuit, and a second path unit circuit. However, for the sake of convenience, the terms are respectively unified as a bank controller, a first path unit, and a second path unit. In addition, the bank controller Bc, the first path unit P1, and the second path unit P2 may each be implemented as a circuit or circuitry.

[0397] The cell array Ca may include a plurality of memory devices (cells) therein. In the cell array Ca, the plurality of memory devices may be arranged in a lattice structure. The cell array Ca may be, for example, a SRAM (static random-access memory) cell array.

[0398] The bank controller Bc may control the cell array Ca. The bank controller Bc may determine whether the cell array Ca operates in the L0 memory type or in the global memory type, and may control the cell array Ca according to the determined memory type.

[0399] Specifically, the bank controller Bc may determine whether to transmit and receive data in the direction of the first path unit P1 or to transmit and receive data in the direction of the second path unit P2 during the runtime. The bank controller Bc may determine a data transmission and reception direction according to a path control signal Spc.

[0400] The path control signal Spc may be generated by a pre-designed device driver or compiler. The path control signal Spc may be generated according to the characteristics of calculation tasks. Alternatively, the path control signal Spc may be generated by an input received from a user. In other words, the user may directly apply an input to the path control signal Spc in order to select optimal memory composition ratio.

[0401] The bank controller Bc may determine a path along which the data stored in the cell array Ca are transmitted and received via the path control signal Spc. The exchange interface of data may be changed as the bank controller Bc determines the path along which the data are transmitted and

received. In other words, a first interface may be used when the bank controller Bc exchanges data with the first path unit P1, and a second interface may be used when the bank controller Bc exchanges data with the second path unit P2. In some embodiments, the first interface and the second interface may be different from each other.

[0402] Also, address systems in which data are stored may vary as well. In other words, if a particular interface is selected, then read and write operations may be performed in an address system corresponding thereto.

[0403] The bank controller Bc may operate at a particular clock frequency. For example, if the cell array Ca is an SRAM cell array, the bank controller Bc may operate at the operating clock frequency of a general SRAM.

[0404] The first path unit P1 may be connected to the bank controller Bc. The first path unit P1 may directly exchange the data of the cell array Ca with the first processing unit 160a. In some embodiments, “directly” may mean being exchanged with each other without going through the global interconnection 6000. In other words, the first processing unit 160a may exchange data directly with the first L0 memory 120a, and the first processing unit 160a may exchange data via the first path unit P1 when the shared memory 2000 is implemented logically in the L0 memory type. The first path unit P1 may include L0 memory controllers including the first L0 memory controller 122_1a and the second L0 memory controller 122_1b as shown in FIG. 37.

[0405] The first path unit P1 may form a multi-cycle sync-path. In other words, the operating clock frequency of the first path unit P1 may be the same as the operating clock frequency of the first processing unit 160a. The first L0 memory 120a may quickly exchange data at the same clock frequency as the operating clock frequency of the first processing unit 160a in order to quickly exchange data at the same speed as the operation of the first processing unit 160a.

[0406] Likewise, the first path unit P1 may also operate at the same clock frequency as the operating clock frequency of the first processing unit 160a.

[0407] In some embodiments, the operating clock frequency of the first path unit P1 may be multiples of the operating clock frequency of the bank controller Bc. In some embodiments, a clock domain crossing (CDC) operation for synchronizing the clocks between the bank controller Bc and the first path unit P1 is not required separately, and thus, a delay of data transmission may not occur. Accordingly, faster and more efficient data exchange can be possible.

[0408] In the embodiment shown in FIG. 38, an operating clock frequency of the first path unit P1 may be 1.5 GHz, as an example. This may be twice the frequency of 750 MHz of the bank controller Bc. However, the embodiment is not limited thereto, and any operating clock frequency of the first path unit P1 may be possible as long as the first path unit P1 operates at integer multiples of the clock frequency of the bank controller Bc.

[0409] The second path unit P2 may be connected to the bank controller Bc. The second path unit P2 may exchange the data of the cell array Ca with the first processing unit 160a not directly but via the global interconnection 6000. In other words, the first processing unit 160a may exchange data with the cell array Ca via the global interconnection 6000 and the second path unit P2. In some embodiments, the cell array Ca may exchange data not only with the first processing unit 160a but also with other processing units.

[0410] In other words, the second path unit P2 may be a data exchange path between the cell array Ca and all the processing units when the first memory bank 2110a is implemented logically in the global memory type. The second path unit P2 may include the global controller 2200 of FIG. 37.

[0411] The second path unit P2 may form an asynchronous path or Async-Path. The operating clock frequency of the second path unit P2 may be the same as the operating clock frequency of the global interconnection 6000. Likewise, the second path unit P2 may also operate at the same clock frequency as the operating clock frequency of the global interconnection 6000.

[0412] In the case of the embodiment as shown in FIG. 38, the operating clock frequency of the second path unit P2 may not be synchronized with the operating clock frequency of the bank

controller Bc. In some embodiments, the clock domain crossing (CDC) operation for synchronizing the clocks between the bank controller Bc and the second path unit P2 may be required. If the operating clock frequency of the bank controller Bc and the operating clock frequency of the second path unit P2 are not synchronized with each other, the degree of freedom in the design of the clock domain may be relatively high. Therefore, the difficulty of hardware design may be decreased, thereby making it possible to more easily derive the desired hardware operation.

[0413] The bank controller Bc may use different address systems in the case of exchanging data via the first path unit P1 and in the case of exchanging data via the second path unit P2. In other words, the bank controller Bc may use a first address system if exchanging data via the first path unit P1 and a second address system if exchanging data via the second path unit P2. In some embodiments, the first address system and the second address system may be different from each other.

[0414] A bank controller Bc is not necessarily required for each memory bank. In other words, a bank controller Bc may not be used to schedule, but instead serves to transfer signals, and thus, is not a required component for each memory bank having two ports. Therefore, one bank controller Bc can be operably coupled to control multiple memory banks. The multiple memory banks may operate independently even if they are controlled by the bank controller Bc. However, the embodiment is not limited thereto.

[0415] As a matter of course, the bank controller Bc may exist for each memory bank. In some embodiments, the bank controller Bc may control each memory bank individually.

[0416] Referring to FIG. 37 and FIG. 38, if the first memory unit **2100a** exchanges data via the first path unit P1, the first address system may be used. If the first memory unit **2100a** exchanges data via the second path unit P2, the second address system may be used. Similarly, if the second memory unit **2100b** exchanges data via the first path unit P1, a third address system may be used. If the second memory unit **2100b** exchanges data via the second path unit P2, the second address system may be used. In some embodiments, the first address system and the third address system may be the same as each other. However, the embodiment is not limited thereto.

[0417] The first address system and the third address system may each be used exclusively for the first processing unit **160a** and the second processing unit **160b**, respectively. The second address system may be commonly applied to the first processing unit **160a** and the second processing unit **160b**.

[0418] In FIG. 38, the operating clock frequency of the second path unit P2 may operate at 1 GHz, as an example. This may be a frequency that is not synchronized with the operating clock frequency of 750 MHz of the bank controller Bc. In other words, the operating clock frequency of the second path unit P2 may be freely set without being dependent on the operating clock frequency of the bank controller Bc at all.

[0419] A generic global memory has used slow SRAM (e.g., 750 MHz) and a global interconnection (e.g., 1 GHz) faster than that, inevitably resulting in delays due to the CDC operation. On the other hand, the shared memory **2000** in accordance with some embodiments has room to use the first path unit P1 in addition to the second path unit P2, thereby making it possible to avoid delays resulting from the CDC operation.

[0420] Furthermore, in the generic global memory, a plurality of processing units use one global interconnection **6000**, and thus, when an amount of data transfer occurs at the same time, the decrease in the overall processing speed is likely to occur. On the other hand, the shared memory **2000** in accordance with some embodiments has room to use the first path unit P1 in addition to the second path unit P2, thereby making it possible to achieve the effect of properly distributing the data throughput that could be concentrated on the global controller **2200** as well.

[0421] FIG. 39 is a block diagram for illustrating a software hierarchy of a neural processing device in accordance with some embodiments.

[0422] Referring to FIG. 39, the software hierarchy of the neural processing device in accordance

[0423] with some embodiments may include a deep learning (DL) framework **10000**, a compiler

stack **20000**, and a back-end module **30000**.

[0424] The DL framework **10000** may mean a framework for a deep learning model network used by a user. For example, a neural network that has finished training may be generated using a program such as TensorFlow or PyTorch.

[0425] The compiler stack **20000** may include an adaptation layer **21000**, a compute library **22000**, a front-end compiler **23000**, a back-end compiler **24000**, and a runtime driver **25000**.

[0426] The adaptation layer **21000** may be a layer in contact with the DL framework **10000**. The adaptation layer **21000** may quantize a neural network model of a user generated by the DL framework **10000** and modify graphs. In addition, the adaptation layer **21000** may convert a type of model into a required type.

[0427] The front-end compiler **23000** may convert various neural network models and graphs transferred from the adaptation layer **21000** into a constant intermediate representation (IR). The converted IR may be a preset representation that is easy to handle later by the back-end compiler **24000**.

[0428] The optimization that can be done in advance in the graph level may be performed on such an IR of the front-end compiler **23000**. In addition, the front-end compiler **23000** may finally generate the IR through the task of converting it into a layout optimized for hardware.

[0429] The back-end compiler **24000** optimizes the IR converted by the front-end compiler **23000** and converts it into a binary file, enabling it to be used by the runtime driver. The back-end compiler **24000** may generate an optimized code by dividing a job at a scale that fits the details of hardware.

[0430] The compute library **22000** may store template operations designed in a form suitable for hardware among various operations. The compute library **22000** provides the back-end compiler **24000** with multiple template operations required by hardware, allowing the optimized code to be generated.

[0431] The runtime driver **25000** may continuously perform monitoring during driving, thereby making it possible to drive the neural network device in accordance with some embodiments. Specifically, it may be responsible for the execution of an interface of the neural network device.

[0432] The back-end module **30000** may include an ASIC (application-specific integrated circuit) **31000**, an FPGA (field-programmable gate array) **32000**, and a C-model **33000**. The ASIC **31000** may refer to a hardware chip determined according to a predetermined design method. The FPGA **32000** may be a programmable hardware chip. The C-model **33000** may refer to a model implemented by simulating hardware on software.

[0433] The back-end module **30000** may perform various tasks and derive results by using the binary code generated through the compiler stack **20000**.

[0434] FIG. **40** is a conceptual diagram for illustrating deep learning calculations performed by a neural processing device in accordance with some embodiments.

[0435] Referring to FIG. **40**, an artificial neural network model **40000** is one example of a machine learning model and is a statistical learning algorithm implemented based on the structure of a biological neural network or is a structure for executing the algorithm, in machine learning technology and cognitive science.

[0436] The artificial neural network model **40000** may represent a machine learning model having an ability to solve problems by learning to reduce the error between an accurate output corresponding to a particular input and an inferred output by repeatedly adjusting the weight of the synapse by nodes. Nodes are artificial neurons that have formed a network by combining synapses, as in a biological neural network. For example, the artificial neural network model **40000** may include any probabilistic model, neural network model, etc., used in artificial intelligence learning methods such as machine learning and deep learning.

[0437] A neural processing device in accordance with some embodiments may implement the form of such an artificial neural network model **40000** and perform calculations. For example, the

artificial neural network model **40000** may receive an input image and may output information on at least a part of an object included in the input image.

[0438] The artificial neural network model **40000** may be implemented by a multilayer perceptron (MLP) including multilayer nodes and connections between them. An artificial neural network model **40000** in accordance with the embodiment may be implemented using one of various artificial neural network model structures including the MLP. As shown in FIG. **25**, the artificial neural network model **40000** includes an input layer **41000** that receives input signals or data **40100** from the outside, an output layer **44000** that outputs output signals or data **40200** corresponding to the input data, and n (where n is a positive integer) hidden layers **42000** to **43000** that are located between the input layer **41000** and the output layer **44000** and that receive a signal from the input layer **41000**, extract characteristics, and forward them to the output layer **44000**. Here, the output layer **44000** receives signals from the hidden layers **42000** to **43000** and outputs them to the outside.

[0439] The learning methods of the artificial neural network model **40000** include a supervised learning method for training to be optimized to solve a problem by the input of supervisory signals (correct answers), and an unsupervised learning method that does not require supervisory signals.

[0440] The neural processing device may directly generate training data, through simulations, for training the artificial neural network model **40000**. In this way, by matching a plurality of input variables and a plurality of output variables corresponding thereto with the input layer **41000** and the output layer **44000** of the artificial neural network model **40000**, respectively, and adjusting the synaptic values between the nodes included in the input layer **41000**, the hidden layers **42000** to **43000**, and the output layer **44000**, training may be made to enable a correct output corresponding to a particular input to be extracted. Through such a training phase, it is possible to identify the characteristics hidden in the input variables of the artificial neural network model **40000**, and to adjust synaptic values (or weights) between the nodes of the artificial neural network model **40000** so that an error between an output variable calculated based on an input variable and a target output is reduced.

[0441] FIG. **41** is a conceptual diagram for illustrating training and inference operations of a neural network of a neural processing device in accordance with some embodiments.

[0442] Referring to FIG. **41**, the training phase may be subjected to a process in which a large number of pieces of training data TD are passed forward to the artificial neural network model NN and are passed backward again. Through this, the weights and biases of each node of the artificial neural network model NN are tuned, and training may be performed so that more and more accurate results can be derived. Through the training phase, the artificial neural network model NN may be converted into a trained neural network model NN_T.

[0443] In the inference phase, new data ND may be inputted into the trained neural network model NN_T again. The trained neural network model NN_T may derive result data RD through the weights and biases that have already been used in the training, with the new data ND as input. For such result data RD, what training data TD were used in training and how many pieces of training data TD were used in the training phase may be important.

[0444] Hereinafter, a neural processing device in accordance with some embodiments of the disclosure will be described with reference to FIGS. **17** to **28**, **42** to **47**, and so on. The parts overlapping with the embodiments described above will be omitted or simplified.

[0445] FIG. **42** is a flowchart for illustrating a method for managing tasks of a neural processing device in accordance with some embodiments of the disclosure, and FIG. **43** is a flowchart for illustrating the checking the dependencies of FIG. **42** in detail. FIG. **44** is a flowchart for illustrating the checking and sequentially transferring the wait fields of FIG. **42** in detail. FIG. **45** is a flowchart for illustrating a process of generating a dependency update request in detail, in the generating the completion reports of FIG. **42**. FIG. **46** is a flowchart for illustrating the process of generating the completion reports in detail, in the generating the completion reports of FIG. **42**.

FIG. 47 is a flowchart for illustrating a control process according to the state of a report queue, in the generating the completion reports of FIG. 42.

[0446] Referring to FIG. 42, the task manager 600 may fetch tasks and generate task descriptors at S100.

[0447] Specifically, referring to FIG. 17, the task buffer 621 may store tasks Task according to control signals transmitted from the command processor 7000. The task buffer 621 may store the tasks Task in the form of task descriptors in the first queue group Q1 by task fetching operations of the core global 500. In some embodiments, the task passage 620 may generate task descriptors corresponding to the tasks Task and store the task descriptors in the first queue group Q1.

[0448] Referring again to FIG. 42, dependencies of the task descriptors are checked and transferred in sequence at S200.

[0449] In detail, referring to FIG. 43, a dependency update request is received at S210. Next, transmission of each task descriptor is determined according to the dependency update request at S220.

[0450] Specifically, referring to FIG. 17, the dependency checker DPc may receive a dependency update request DFURQ. The dependency update request DFURQ may notify a change of dependency as a completed task occurs according to a defined dependency between particular tasks. In some embodiments, each task descriptor may include a dependency field as to which tasks it has a dependency on. In some embodiments, it must be updated in such a way that when the task included in the dependency field is completed, it will be removed from the dependency field. Therefore, the dependency update request DFURQ may include an update request for the dependency field of a task descriptor. Accordingly, the dependency checker DPc may sequentially transmit the first to n-th task descriptors Tsk_d1 to Tsk_dn to the second queue group Q2.

[0451] Referring again to FIG. 42, wait fields of the task descriptors are checked and transferred sequentially at S300.

[0452] In detail, referring to FIG. 44, step S300 includes checking whether the task descriptors include a wait field at S310, controlling the task descriptors to wait in a task queue if the task descriptors include the wait field at S320, releasing the waiting state of the task descriptors in response to a run signal provided from the command processor at S330, generating task information corresponding to the task descriptor for which the waiting state has been released and transferring the task information to the core global at S340, and providing the task descriptor corresponding to the task information transferred to the core global as check-in data to the done passage at S350.

[0453] Specifically, referring to FIG. 18, a plurality of task descriptors for which the dependency check has been completed may be stored in the second queue group Q2, the runtime handle RH may check whether the task descriptors stored in the second queue group Q2 include the wait field, and determine the state of the task descriptors as a run state or a waiting state.

[0454] The runtime handle RH may receive a run signal Run from the command processor 7000. The run signal Run may be provided from the command processor 7000 via the control interconnection CI. The command processor 7000 may transfer the run signal Run in response to transmitting a task that includes a wait field, but embodiments of the disclosure are not limited thereto. The command processor 7000 may transmit the run signal Run to the runtime handle RH at regular intervals. The runtime handle RH may change the task descriptors in the waiting state to the run state in response to the run signal Run.

[0455] At S330, the run signal may be a signal provided after the waiting state of the task descriptor, but is not limited thereto. At S330, the run signal may be a signal provided before the waiting state of the task descriptor. In some embodiments, the runtime handle RH may be configured to receive the run signal Run and store the run signal Run for a period of time. The runtime handle RH may include at least one of at least one register Rs or counter Rc for receiving and storing the run signal Run, but embodiments of the disclosure are not limited thereto.

[0456] Referring again to FIG. 42, the tasks of the task descriptors are performed at S400. Next, completion signals of the task descriptors are generated at S500.

[0457] Specifically, referring to FIG. 28, the neural core 100 may perform tasks and generate completion signals. The neural core 100 may transfer the completion signals to the core global 500. The core global 500 may include a signal scheduler sgn_sch. The signal scheduler sgn_sch may receive the completion signal, schedule the transmission of the completion signal, and transmit the completion signal to the done passage 630.

[0458] Referring again to FIG. 42, completion reports according to the completion signals are generated at S600.

[0459] Referring to FIG. 45 in detail, a dependency update request is generated by checking the completed task descriptor at S610. Next, the dependency update request is transmitted at S620.

[0460] Specifically, referring to FIG. 18, the dependency setter DPs may receive the completion signals from the check-in buffer Cib and generate dependency update requests DFURQ. In some embodiments, depending on which task corresponding to what task descriptor has been completed, the dependency setter DPs may generate a dependency update request DFURQ. The dependency setter DPs may transmit the dependency update request DFURQ to the task passage 620.

[0461] In some embodiments, performing the tasks of the transferred task descriptors at S400 may include monitoring the execution time of the tasks and generating timeout detection signals.

[0462] The timeout monitor ToM may receive the timeout detection signals TDec from the check-in buffer Cib. The timeout detection signal TDec may be a signal regarding whether a time interval between a check-in time and a check-out time exceeds a preset threshold time. The check-in buffer Cib may compare the execution time of the task with the threshold time and determine whether to generate the timeout detection signal TDec. The check-in buffer Cib may generate the timeout detection signal TDec if the execution time calculated from the check-in time exceeds the threshold time. Whether to generate a timeout report and the threshold for the execution time may be set individually according to the tasks.

[0463] Referring to FIG. 46, generating the completion reports at S600 may further include generating timeout reports according to the timeout detection signals at S630, generating checkout reports according to the received completion signals at S640, and generating the completion reports based on the checkout reports and the timeout reports at S650. Here, steps S630 to S650 are detailed steps for generating the completion reports, and there is no special precedence relationship established between steps S610 to S620 related to the dependency update request.

[0464] At S650, the completion report DNrp may be generated based on the checkout report COrp. Since the completion report DNrp includes at least the checkout report, whether a task has been normally performed can be transferred to the command processor 7000. Through the generation and transfer of such a completion report DNrp, it can be checked whether a task has been performed normally, and it is possible to prevent a delay in performing a particular task from being prolonged according to dependencies.

[0465] In addition, the completion report DNrp may be configured to further include at least one of the transfer done report TRrp, the event report Erp, or the timeout report TOrp, and it is possible to comprehensively report to the command processor 7000 whether a timeout has occurred, whether a transfer has been completed, whether an event has occurred, and the like, along with whether a task has been performed.

[0466] In an embodiment, whether to generate such a completion report DNrp may be set individually according to the tasks. In some embodiments, the command processor 7000 may be set in advance so that a particular task will not receive a completion report DNrp related thereto. Accordingly, since completion reports may not be generated for all tasks, the burden of data processing on the task manager 600 may be reduced, and as the command processor 7000 may not receive all completion reports, the management efficiency of the command processor 7000 may be increased.

[0467] However, if an important warning related to task execution or a request for a part requiring correction is required to the command processor **7000**, such as a timeout report TOrp, corresponding information may be required to be transferred via a completion report DNrp.

[0468] For example, the tasks may include a first task that is set to generate a completion report and a second task that is set not to generate a completion report. In the providing the task descriptor corresponding to the task information transferred to the core global as check-in data to the done passage at **S350**, a first task descriptor corresponding to the first task and a second task descriptor corresponding to the second task may be checked in at a first check-in buffer and a second check-in buffer, respectively. Here, in the generating the timeout reports, if a timeout report is generated for the second task that is set not to generate a completion report, the generating the completion reports at **S650** may perform a process of modifying the second task descriptor to generate a completion report.

[0469] In some embodiments, if an important warning related to task execution or a request for a part requiring correction is required to the command processor **7000**, such as a timeout report TOrp, corresponding information may be required to be transferred via a completion report DNrp. In an embodiment of the disclosure, if modification is needed for a task descriptor that is set not to generate a completion report, the timeout monitor ToM may transmit a modification request signal Mos to the check-in buffer in which the corresponding task descriptor is stored, and the corresponding task descriptor is modified so that a completion report is generated.

[0470] Furthermore, referring to FIG. **47**, the generating the completion reports at **S600** may include storing the timeout reports in a timeout report queue, storing the checkout reports in the checkout report queue, and storing the event reports provided from an event detection device in an event report queue at **S660**, and monitoring at least one of the state of the timeout report queue, the state of the checkout report queue, or the state of the event report queue at **S670**. Moreover, the generating the completion reports at **S600** may further include stopping the operation of the task passage via the runtime handle at **S680** if at least one of the state of the timeout report queue, the state of the checkout report queue, or the state of the event report queue is confirmed to be in a saturated state.

[0471] In some embodiments, blocks, units, modules, and components as described above may be implemented as a circuit or circuitry. Blocks, units, modules, and components which perform processing may be referred to as a processor, a processing circuit, a processor circuit, or a processing circuitry. Blocks, units, modules, and components which store data may be referred to as a memory, a memory circuit, or a memory circuitry.

[0472] Hereinafter, various aspects will be described.

[0473] In some aspects, a neural processing device comprises processing circuitry; and task managing circuitry operably coupled to the processing circuitry. The task managing circuitry is configured to cause: receiving a plurality of tasks; storing the plurality of received tasks in a dependency check waiting memory; determining whether reference tasks which a task in the dependency check waiting memory depends on are completed based on a dependency of the task in the dependency check waiting memory and a dependency check memory region, wherein the dependency of the task in the dependency check waiting memory represents a list of tasks which have to be completed for the task in the dependency check waiting memory to be transferred to the processing circuitry and the dependency check memory region represents whether reference tasks which tasks in the dependency check waiting memory depend on are completed; transferring the task in the dependency check waiting memory to a process waiting memory in response to determining that reference tasks which the task in the dependency check waiting memory depends on are completed; and transferring tasks in the process waiting memory to the processing circuitry.

[0474] In some aspects, transferring the tasks in the process waiting memory to the processing circuitry comprises: determining whether a task in the process waiting memory indicates that a run signal is required for the task in the process waiting memory to be processed, when the task in the

process waiting memory indicates that a run signal is not required, transferring the task in the process waiting memory to the processing circuitry regardless of the run signal, and when the task in the process waiting memory indicates that that a run signal is required, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal.

[0475] In some aspects, the neural processing device of claim Error! Reference source not found., wherein transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal comprises: in response to receiving the run signal, changing a state of a register to indicate that the run signal has been received, and when the task in the process waiting memory indicates that that a run signal is required, transferring the task in the process waiting memory to the processing circuitry in response to determining that the state of the register indicates the run signal has been received.

[0476] In some aspects, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal further comprises: in response to transferring the task in the process waiting memory to the processing circuitry, changing the state of the register to indicate that the run signal has not been received.

[0477] In some aspects, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal comprises: in response to receiving the run signal, increasing a counter value in a register to indicate how many run signals have been received, in response to transferring, to the processing circuitry, a task indicating that the run signal is required to be processed, decreasing the counter value in the register, and when the task in the process waiting memory indicates that the run signal is required transferring the task in the process waiting memory to the processing circuitry in response to determining that the counter value is greater than zero.

[0478] In some aspects, the task managing circuitry is further configured to cause: storing events in an even queue; and in response to detecting that a number of entries in the even queue is larger than a threshold, ceasing to transfer tasks in the process waiting memory to the processing circuitry regardless of the run signal until transferring tasks to the processing circuitry is resumed.

[0479] In some aspects, the task managing circuitry is further configured to cause: storing tasks which are not completed for a predetermined time in a timeout queue; and storing tasks which are completed in a completed task queue.

[0480] In some aspects, the task managing circuitry is further configured to cause: in response to detecting that a number of entries in the timeout queue is larger than a threshold, ceasing to transfer tasks in the process waiting memory to the processing circuitry regardless of the run signal until transferring tasks to the processing circuitry is resumed.

[0481] In some aspects, the task managing circuitry is further configured to cause: in response to detecting that a number of entries in the completed task queue is larger than a threshold, ceasing to transfer tasks in the process waiting memory to the processing circuitry regardless of the run signal until transferring tasks to the processing circuitry is resumed.

[0482] In some aspects, the plurality of tasks are received from an upper-level task managing circuitry of the task managing circuitry. The task managing circuitry is further configured to cause: determining whether a task in the completed task queue requires a report, and transmitting, to the upper-level task managing circuitry, a completion report indicating that a task in the completed task queue is completed, when the task in the completed task queue requires the report.

[0483] In some aspects, the plurality of tasks are received from an upper-level task managing circuitry of the task managing circuitry. The task managing circuitry is further configured to cause: determining whether a task in the timeout queue requires a report; and transmitting, to the upper-level task managing circuitry, a timeout report indicating that a task in the timeout queue is not completed for a predetermined time, when the task in the timeout queue requires the report.

[0484] In some aspects, the task managing circuitry is further configured to cause: updating the dependency check memory region in response to a signal indicating that a task processed by the

processing circuitry is completed.

[0485] In some aspects, tasks are stored in a form of task descriptors in the dependency check waiting memory, and tasks are stored in a form of task descriptors in the process waiting memory.

[0486] In some aspects, the tasks are for a machine learning model, the processing circuitry processes tasks which are transferred from the process waiting memory to generate task processing results, and the neural processing device is configured to perform training or inference of the machine learning model based on the task processing results.

[0487] In some aspects, a method performed by a neural processing device, comprises: receiving a plurality of tasks; storing the plurality of received tasks in a dependency check waiting memory; determining whether reference tasks which a task in the dependency check waiting memory depends on are completed based on a dependency of the task in the dependency check waiting memory and a dependency check memory region, wherein the dependency of the task in the dependency check waiting memory represents a list of tasks which have to be completed for the task in the dependency check waiting memory to be transferred to the processing circuitry and the dependency check memory region represents whether reference tasks which tasks in the dependency check waiting memory depend on are completed; transferring the task in the dependency check waiting memory to a process waiting memory in response to determining that reference tasks which the task in the dependency check waiting memory depends on are completed; and transferring tasks in the process waiting memory to the processing circuitry.

[0488] In some aspects, transferring the tasks in the process waiting memory to the processing circuitry comprises: determining whether a task in the process waiting memory indicates that a run signal is required for the task in the process waiting memory to be processed; when the task in the process waiting memory indicates that a run signal is not required, transferring the task in the process waiting memory to the processing circuitry regardless of the run signal; and when the task in the process waiting memory indicates that that a run signal is required, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal.

[0489] In some aspects, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal comprises: in response to receiving the run signal, changing a state of a register to indicate that the run signal has been received; and when the task in the process waiting memory indicates that that a run signal is required, transferring the task in the process waiting memory to the processing circuitry in response to determining that the state of the register indicates the run signal has been received.

[0490] In some aspects, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal further comprises: in response to transferring the task in the process waiting memory to the processing circuitry, changing the state of the register to indicate that the run signal has not been received.

[0491] In some aspects, transferring the task in the process waiting memory to the processing circuitry in response to receiving the run signal comprises: in response to receiving the run signal, increasing a counter value in a register to indicate how many run signals have been received; in response to transferring, to the processing circuitry, a task indicating that the run signal is required to be processed, decreasing the counter value in the register; and when the task in the process waiting memory indicates that the run signal is required transferring the task in the process waiting memory to the processing circuitry in response to determining that the counter value is greater than zero.

[0492] In some aspects, the tasks are for a machine learning model. The method further comprises: processing, the processing circuitry, tasks which are transferred from the process waiting memory to generate task processing results; and performing training or inference of the machine learning model based on the task processing results.

[0493] While the inventive concept has been particularly shown and described with reference to exemplary embodiments thereof, it will be understood by those of ordinary skill in the art that

various changes in form and details may be made therein without departing from the spirit and scope of the inventive concept as defined by the following claims. It is therefore desired that the embodiments be considered in all respects as illustrative and not restrictive, reference being made to the appended claims rather than the foregoing description to indicate the scope of the disclosure.

Claims

1. A neural processor comprising: a core global configured to generate task information corresponding to a task descriptor; a neural core configured to receive the task information from the core global, perform a task according to the task information, and generate a completion signal indicating completion of the task; and a task manager configured to monitor an execution time of the task performed by the neural core, generate a timeout detection signal based on the execution time, and generate a timeout report based on the timeout detection signal, wherein the core global is further configured to receive, from the neural core, the completion signal of the task, the task manager is further configured to receive the completion signal from the core global and generate a completion report, and the timeout report is included in the completion report and is transmitted to a command processor.
2. The neural processor according to claim 1, wherein the task comprises a first task for which the completion report is configured to be generated and a second task for which the completion report is configured not to be generated, and the task manager is further configured to check in a first task descriptor corresponding to the first task into a first check-in buffer and check in a second task descriptor corresponding to the second task into a second check-in buffer.
3. The neural processor according to claim 2, wherein upon generation of a timeout report for the second task, the task manager is further configured to modify the second task descriptor such that the completion report is generated.
4. The neural processor according to claim 1, wherein the task manager is further configured to determine whether to generate the timeout detection signal based on a comparison between the execution time and a threshold time.
5. The neural processor according to claim 4, wherein the threshold time is individually set according to the task.
6. A neural processor comprising: a core global configured to generate task information corresponding to a task descriptor; a neural core configured to receive the task information from the core global, perform a task according to the task information, and generate a completion signal indicating completion of the task; and a task manager configured to monitor an execution time of the task performed by the neural core, generate a timeout detection signal based on the execution time, and generate a timeout report based on the timeout detection signal, wherein the task manager comprises: a task passage configured to generate the task descriptor for the task and transmit the task information corresponding to the task descriptor to the core global; and a done passage configured to check in the task descriptor from the task passage and generate a completion report based on the completion signal received from the core global.
7. The neural processor according to claim 6, wherein the done passage comprises: a check-in buffer configured to receive and check in the task descriptor, receive the completion signal, and associate the task descriptor with the completion signal; a dependency setter configured to check out the task descriptor based on the completion signal and generate a check-out report corresponding to a check-out of the task descriptor; a timeout monitor configured to generate the timeout detection signal and generate the timeout report based on the timeout detection signal; and a reporting managing module configured to receive the check-out report, generate the completion report, and transmit the completion report to a command processor.
8. The neural processor according to claim 7, wherein the task passage is further configured to transmit the task descriptor to a queue, generate a transfer-done report, and transmit the transfer-

done report to the reporting managing module, and the transfer-done report is included in the completion report and is transmitted to the command processor.

9. The neural processor according to claim 8, wherein the done passage further comprises an event monitor configured to generate an event report based on an event detection signal.

10. The neural processor according to claim 9, wherein the reporting managing module comprises: a reporting manager configured to generate the completion report; an event report queue configured to receive the event report and transmit the received event report to the reporting manager; a timeout report queue configured to receive the timeout report and transmit the received timeout report to the reporting manager; a check-out report queue configured to receive the check-out report and transmit the received check-out report to the reporting manager; and a transfer report queue configured to receive the transfer-done report and transmit the received transfer-done report to the reporting manager.

11. The neural processor according to claim 10, wherein the reporting manager is further configured to monitor at least one of a status of the timeout report queue, a status of the check-out report queue, and a status of the event report queue.

12. A method for monitoring a task, performed by a neural processor, the method comprising: by a core global of the neural processor, generating task information corresponding to a task descriptor and providing the task information to a neural core of the neural processor; by a task manager of the neural processor, monitoring an execution time of the task performed by the neural core and generating a timeout detection signal based on the execution time; by the task manager, generating a timeout report based on the timeout detection signal; by the core global, receiving, from the neural core, a completion signal indicating completion of the task; by the task manager, receiving the completion signal from the core global and generating a completion report; and by the task manager, transmitting the completion report including the timeout report to a command processor of the neural processor.

13. The method according to claim 12, further comprising generating, by the task manager, a check-out report based on the completion signal for the task received from the neural core, wherein the generating the completion report comprises: generating, by the task manager, the completion report based on the check-out report and the timeout report.

14. The method according to claim 13, wherein the task comprises a first task for which the completion report is configured to be generated and a second task for which the completion report is configured not to be generated, and the generating the task information corresponding to the task descriptor and the providing the task information to the neural core comprises: checking in a first task descriptor corresponding to the first task into a first check-in buffer and checking in a second task descriptor corresponding to the second task into a second check-in buffer.

15. The method according to claim 14, wherein the generating the completion report comprises modifying the second task descriptor such that a completion report is generated upon generation of a timeout report for the second task.

16. The method according to claim 13, wherein the generating the timeout detection signal comprises determining whether to generate the timeout detection signal based on a comparison between the execution time and a threshold time.

17. The method according to claim 16, wherein whether to generate the timeout report and a threshold value for the execution time are individually set according to the task.

18. The method according to claim 13, further comprising generating, by the task manager, an event report based on an event detection signal.

19. The method according to claim 18, wherein the generating the completion report comprises: by the task manager, storing the timeout report in a timeout report queue, storing the check-out report in a check-out report queue, and storing an event report, provided from an event detection device, in an event report queue, and by the task manager, monitoring at least one of a status of the timeout report queue, a status of the check-out report queue, and a status of the event report queue.

