

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284646

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

ANDERSON; Timothy D. et al.

SOFTWARE-HARDWARE MEMORY MANAGEMENT MODES

Abstract

A method includes receiving, by a memory management unit (MMU) comprising a translation lookaside buffer (TLB) and a configuration register, a request from a processor core to directly modify an entry in the TLB. The method also includes, responsive to the configuration register having a first value, operating the MMU in a software-managed mode by modifying the entry in the TLB according to the request. The method further includes, responsive to the configuration register having a second value, operating the MMU in a hardware-managed mode by denying the request.

Inventors: ANDERSON; Timothy D. (University Park, TX), ZBICIAK; Joseph Raymond Michael (San Jose, CA), CHIRCA; Kai (Dallas, TX), WU; Daniel Brad (Plano, TX)

Applicant: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Family ID: 75382937

Appl. No.: 19/218747

Filed: May 27, 2025

Related U.S. Application Data

parent US continuation 18389899 20231220 parent-grant-document US 12314187 child US 19218747

parent US continuation 17068730 20201012 parent-grant-document US 11853225 child US 18389899

us-provisional-application US 62914061 20191011

Publication Classification

Int. Cl.: G06F12/1027 (20160101); G06F9/46 (20060101); G06F9/48 (20060101); G06F12/0862 (20160101); G06F12/0882 (20160101); G06F12/0891 (20160101); G06F12/1009 (20160101); H03M13/15 (20060101)

U.S. Cl.:

CPC G06F12/1027 (20130101); G06F9/467 (20130101); G06F9/4881 (20130101); G06F12/0862 (20130101); G06F12/0882 (20130101); G06F12/0891 (20130101); G06F12/1009 (20130101); H03M13/1575 (20130101); G06F2212/1021 (20130101); G06F2212/602 (20130101); G06F2212/68 (20130101)

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] The present application is a continuation of U.S. patent application Ser. No. 18/389,899, filed Dec. 20, 2023, which is a continuation of U.S. patent application Ser. No. 17/068,730, filed Oct. 12, 2020, now U.S. Pat. No. 11,853,225, which claims priority to U.S. Provisional Patent Application No. 62/914,061, filed Oct. 11, 2019, each of which is incorporated by reference herein in its entirety. [0002] Some aspects of the present application are related in subject matter to U.S. patent application Ser. No. 17/068,713, filed Oct. 12, 2020, now U.S. Pat. No. 11,954,044, and to U.S. patent application Ser. No. 17/068,721, filed Oct. 12, 2020, now U.S. Pat. No. 11,663,141.

BACKGROUND

[0003] Managing interactions between multiple software applications or program tasks and physical memory involves address translation (e.g., between a virtual address and a physical address or between a first physical address and a second physical address). Software applications or program task modules are generally compiled with reference to a virtual address space. When an application or task interacts with physical memory, address translation is performed to translate a virtual address into a physical address in the physical memory. Address translation consumes processing and/or memory resources. A cache of translated addresses, referred to as a translation lookaside buffer (TLB), improves address translation performance.

SUMMARY

[0004] In accordance with at least one example of the disclosure, a method includes receiving, by a memory management unit (MMU) comprising a translation lookaside buffer (TLB) and a configuration register, a request from a processor core to directly modify an entry in the TLB. The method also includes, responsive to the configuration register having a first value, operating the MMU in a software-managed mode by modifying the entry in the TLB according to the request. The method further includes, responsive to the configuration register having a second value, operating the MMU in a hardware-managed mode by denying the request.

[0005] In accordance with another example of the disclosure, a system includes a processor core and a memory management unit (MMU) coupled to the processor core and comprising a translation lookaside buffer (TLB) and a configuration register. The MMU is configured to receive, from the processor core, a request to directly modify an entry in the TLB; responsive to the configuration register having a first value, operate in a software-managed mode by modifying the entry in the TLB according to the request; and responsive to the configuration register having a second value, operate in a hardware-managed mode by denying the request.

[0006] In accordance with yet another example of the disclosure, an apparatus includes a memory management unit (MMU), which includes a translation lookaside buffer (TLB) and a configuration register. The configuration register having a first value causes the MMU to operate in a software-

managed mode and the configuration register having a second value causes the MMU to operate in a hardware-managed mode. The MMU is configured to receive, from a processor core, a request to directly modify an entry in the TLB; responsive to the configuration register having a first value, operate in a software-managed mode by modifying the entry in the TLB according to the request; and responsive to the configuration register having a second value, operate in a hardware-managed mode by denying the request.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0007] For a detailed description of various examples, reference will now be made to the accompanying drawings in which:

[0008] FIG. 1 is a block diagram of a multi-core processing system in accordance with various examples;

[0009] FIG. 2 is a block diagram showing a memory management unit (MMU) in greater detail and in accordance with various examples;

[0010] FIGS. 3a and 3b are examples of one- and two-stage address translation in accordance with various examples;

[0011] FIG. 4 is another block diagram showing the MMU in greater detail and in accordance with various examples; and

[0012] FIG. 5 is a flow chart of a method of operating the MMU in software- and hardware-managed modes in accordance with various examples.

DETAILED DESCRIPTION

[0013] FIG. 1 is a functional block diagram of a multi-core processing system **100**, in accordance with examples of this description. In one example, the system **100** is a multi-core system-on-chip (SoC) that includes a processing cluster **102** having one or more processor packages **104**. In some examples, the one or more processor packages **104** include one or more types of processors, such as a central processor unit (CPU), graphics processor unit (GPU), digital signal processor (DSP), etc. In one example, a processing cluster **102** includes a set of processor packages split between DSP, CPU, and GPU processor packages. In some examples, each processor package **104** includes one or more processing cores **106**. As used herein, the term “core” refers to a processing module that is configured to contain an instruction processor, such as a DSP or other type of microprocessor. Each processor package **104** also contains a memory management unit (MMU) **108** and one or more caches **110**. In some example, the caches **110** include one or more level one (L1) caches and one or more level two (L2) caches. For example, a processor package **104** includes four cores **106**, each core including an L1 data cache and L1 instruction cache, along with a L2 cache shared by the four cores **106**.

[0014] The multi-core processing system **100** also includes a multi-core shared memory controller (MSMC) **112**, which couples the processing cluster **102** to one or more external memories **114** and direct memory access/input/output (DMA/IO) clients **116**. The MSMC **112** also includes an on-chip internal memory **118** that is directly managed by the MSMC **112**. In certain examples, the MSMC **112** manages traffic between multiple processor cores **106**, other mastering peripherals or DMA clients **116** and allows processor packages **104** to dynamically share the internal and external memories for both program instructions and data. The MSMC internal memory **118** offers additional flexibility (e.g., to software programmers) because portions of the internal memory **118** are configured as a level 3 (L3) cache.

[0015] The MMU **108** is configured to perform address translation between a virtual address and a physical address, including intermediate physical addresses for multi-stage address translation. In some examples, the MMU **108** is also configured to perform address translation between a first

physical address and a second physical address (e.g., as part of a multi-stage address translation). In particular, the MMU **108** helps to translate virtual memory addresses to physical memory addresses for the various memories of the system **100**. The MMU **108** contains a translation lookaside buffer (TLB) **120** that is configured to store translations between addresses (e.g., between a virtual address and a physical address or between a first physical address and a second physical address). Although not shown for simplicity, in other examples the MMU **108** additionally includes a micro-TLB (uTLB), such as a fully associative uTLB, which, along with the TLB **120**, serve as caches for page translations. In some examples, the TLB **120** also stores address pointers of page tables. In addition to address translations stored (e.g., cached) in the TLB **120**, the MMU **108** includes one or more page table walker engines **122** that are configured to access or “walk” one or more page tables to translate a virtual address to a physical address, or to translate an intermediate physical address to a physical address. The function of the page table walker engine **122** is described further below.

[0016] The processor core **106** generates a transaction directed to a virtual address that corresponds to a physical address in memory (e.g., external memory **114**). Examples of such transactions generated by the processor core **106** include reads from the memory **114** and writes to the memory **114**; however, other types of transactions requiring address translation (e.g., virtual-to-physical address translation and/or physical-to-physical address translation) are also within the scope of this description. For ease of reference, any transaction that entails address translation is referred to as an address translation request (or “translation request”), and it is further assumed for simplicity that translation requests specify a virtual address to be translated to a physical address. The processor core **106** thus provides a translation request to the MMU **108**.

[0017] Responsive to receiving a translation request from the processor core **106**, the MMU **108** first translates the virtual address specified by the translation request to a physical address. A first example translation request **130** is provided by the processor core **106** to the MMU **108**. The MMU **108** first determines whether the first translation request **130** hits the TLB **120** (e.g., the TLB **120** already contains the address translation for the virtual address specified by the first translation request **130**). In this example, the first translation request **130** does hit the TLB **120**, and thus the MMU **108** forwards a transaction **132** that includes the translated physical address to a lower level memory (e.g., the caches **110**) for further processing.

[0018] A second example translation request **140** is provided by the processor core **106** to the MMU **108**. The MMU **108** again determines whether the second translation request **140** hits the TLB **120**. In this example, the second translation request **140** misses (e.g., does not hit) the TLB **120**. Responsive to the second translation request **140** missing the TLB **120**, the MMU **108** provides the second translation request **140** to its page table walker engine **122**, which accesses (e.g., “walks”) one or more page tables in a lower level memory (e.g., the caches **110**, **118**, or external memory **114**) to translate the virtual address specified by the second translation request **140** to a physical address. The process of walking page tables is described in further detail below. Once the page table walker engine **122** translates the virtual address to a physical address, the address translation is stored in the TLB **120** (depicted as arrow **142**), and the MMU **108** forwards a transaction **144** that includes the translated physical address to a lower level memory for further processing.

[0019] A third possibility exists, in which the translation request from the processor core **106** only partially hits the TLB **120**. In such a situation, which will be described further below, the page table walker engine **122** still walks one or more page tables in the lower level memory to translate the virtual address specified by the translation request to a physical address. However, because the translation request partially hit the TLB **120**, a reduced number of page tables are walked in order to perform the address translation relative to a translation request that completely misses the TLB **120**.

[0020] FIG. 2 is a block diagram of a system **200** that includes a processor core **106** and MMU

108, which itself includes the TLB **120** and page table walker engine **122**, as described above. In the example of FIG. 2, the MMU **108** is shown in further detail and includes an invalidation engine **202**, a transaction multiplexer (mux) **204**, a general purpose transaction buffer **206**, a dedicated invalidation buffer **208**, and one or more memory mapped registers (MMRs) **210** that are used to control and/or configure various functionality of the MMU **108**. In some examples, the TLB **120** includes multiple pipeline stages (shown as matching logic **212**) that facilitate the TLB **120** receiving a translation request and determining whether the virtual address specified by the translation request hits the TLB **120**, partially hits the TLB **120**, or misses the TLB **120**.

[0021] As described above, the processor core **106** is configured to provide various translation requests to the MMU **108**, which are provided to the transaction mux **204** as shown. In some examples, the processor core **106** is configured to provide address invalidation requests (or “invalidation requests”) to the MMU **108** in addition to the translation requests. Invalidation requests are requests to invalidate one or more entries in the TLB **120**. In some examples, invalidation requests are for a single entry (e.g., associated with a particular virtual address) in the TLB **120**, while in other examples, invalidation requests are for multiple entries (e.g., associated with a particular application ID) in the TLB **120**. The invalidation requests are provided to the invalidation engine **202** of the MMU **108**, which in turn forwards such invalidation requests to be looked up in the TLB **120** to the transaction mux **204** as shown. Regardless of the type of request, the transaction mux **204** is configured to pass both translation requests and invalidation requests to the TLB **120**. In some examples, control logic provides control signals to the transaction mux **204** to select one of the inputs to the transaction mux **204** to be provided as the output of the transaction mux **204**. In an example, address translation requests are prioritized over address invalidation requests until there are no more available spots in the general purpose transaction buffer **206** for such address translation requests.

[0022] Responsive to receiving a request (e.g., either a translation request or an invalidation request), the matching logic **212** (e.g., implemented by pipeline stages of the TLB **120**) determines whether the request hits the TLB **120**, partially hits the TLB **120**, or misses the TLB **120**.

[0023] Depending on the type of request, various resulting transactions are produced by the matching logic **212**. For example, a translation request can hit the TLB **120**, partially hit the TLB **120**, or miss the TLB **120**. An invalidation request can either hit the TLB **120** or miss the TLB **120**, because an invalidation request that only partially hits an entry in the TLB **120** should not result in invalidating that entry in some examples. In other examples, an invalidation request can also partially hit the TLB **120**. For example, a partial hit on the TLB **120** exists when a request hits on one or more pointers to page table(s), but does not hit on at least the final page table. A hit on the TLB **120** exists when a request hits on both the one or more pointers to page table(s) as well as the final page table itself. In some examples, an invalidation request includes a “leaf level” bit or field that specifies to the MMU **108** whether to invalidate only the final page table (e.g., partial hits on the TLB **120** do not result in invalidating an entry) or to invalidate pointers to page table(s) as well (e.g., a partial hit on the TLB **120** results in invalidating an entry).

[0024] Responsive to a translation request that hits the TLB **120**, the MMU **108** provides an address transaction specifying a physical address to the general purpose transaction buffer **206**. In this example, the general purpose transaction buffer **206** is a first-in, first-out (FIFO) buffer. Once the address transaction specifying the physical address has passed through the general purpose transaction buffer **206**, the MMU **108** forwards that address transaction to a lower level memory to be processed.

[0025] Responsive to a translation request that partially hits the TLB **120** or misses the TLB **120**, the MMU **108** provides an address transaction that entails further address translation to the general purpose transaction buffer **206**. For example, if the translation request misses the TLB **120**, the address transaction provided to the general purpose transaction buffer **206** entails complete address translation (e.g., by the page table walker engine **122**). In another example, if the translation request

partially hits the TLB **120**, the address transaction provided to the general purpose transaction buffer **206** entails additional, partial address translation (e.g., by the page table walker engine **122**). Regardless of whether the address transaction entails partial or full address translation, once the address transaction that entails additional translation has passed through the general purpose transaction buffer **206**, the MMU **108** forwards that address transaction to the page table walker engine **122**, which in turn performs the address translation.

[0026] Generally, performing address translation is more time consuming (e.g., consumes more cycles) than simply processing a transaction such as a read or a write at a lower level memory. Thus, in examples where multiple translation requests miss the TLB **120** or only partially hit the TLB **120** (e.g., entails some additional address translation be performed by the page table walker engine **122**), the general purpose transaction buffer **206** can back up and become full. The processor core **106** is aware of whether the general purpose transaction buffer **206** is full and, responsive to the general purpose transaction buffer **206** being full, the processor core **106** temporarily stalls from sending additional translation requests to the MMU **108** until space becomes available in the general purpose transaction buffer

[0027] Responsive to an invalidation look-up request that hits the TLB **120**, the MMU **108** provides a transaction specifying that an invalidation match occurred in the TLB **120**, referred to as an invalidation match transaction for simplicity. Responsive to the general purpose transaction buffer **206** having space available (e.g., not being full), the MMU **108** is configured to provide the invalidation match transaction to the general purpose transaction buffer **206**. However, responsive to the general purpose transaction buffer **206** being full, the MMU **108** is configured to provide the invalidation match transaction to the dedicated invalidation buffer **208**. In this example, the dedicated invalidation buffer **208** is also a FIFO buffer. As a result, even in the situation where the general purpose transaction buffer **206** is full (e.g., due to address translation requests missing or only partially hitting the TLB **120**, and thus backing up in the general purpose transaction buffer **206**), the processor core **106** is able to continue sending invalidation requests to the MMU **108** because the invalidation requests are able to be routed to the dedicated invalidation buffer **208**, and thus are not stalled behind other translation requests.

[0028] Regardless of whether the invalidation match transaction is stored in the general purpose transaction buffer **206** or the dedicated invalidation buffer **208**, once the invalidation match transaction passes through one of the buffers **206**, **208**, the invalidation match transaction is provided to the invalidation engine **202**, which is in turn configured to provide an invalidation write transaction to the TLB **120** to invalidate the matched entry or entries. In an example, invalidation look-up requests that miss the TLB **120** are discarded (e.g., not provided to either the general purpose transaction buffer **206** or the dedicated invalidation buffer **208**).

[0029] FIG. **3a** is an example translation **300** for translating a 49-bit virtual address (VA) to a physical address (PA) in accordance with examples of this description. The example translation **300** is representative of the functionality performed by the page table walker engine **122** responsive to receiving a transaction that entails full or partial address translation.

[0030] In this example, the most significant bit of the 49-bit VA specifies one of two table base registers (e.g., TBR**0** or TBR**1**, implemented in the MMRs **210**). The table base registers each contain a physical address that is a base address of a first page table (e.g., Level 0). In this example, each page table includes 512 entries, and thus an offset into a page table is specified by nine bits. A first group of nine bits **302** provides the offset from the base address specified by the selected table base register into the Level 0 page table to identify an entry in the Level 0 page table. The identified entry in the Level 0 page table contains a physical address that serves as a base address of a second page table (e.g., Level 1).

[0031] A second group of nine bits **304** provides the offset from the base address specified by entry in the Level 0 page table into the Level 1 page table to identify an entry in the Level 1 page table. The identified entry in the Level 1 page table contains a physical address that serves as a base

address of a third page table (e.g., Level 2).

[0032] A third group of nine bits **306** provides the offset from the base address specified by entry in the Level 1 page table into the Level 2 page table to identify an entry in the Level 2 page table. The identified entry in the Level 2 page table contains a physical address that serves as a base address of a fourth, final page table (e.g., Level 3).

[0033] A fourth group of nine bits **308** provides the offset from the base address specified by entry in the Level 2 page table into the Level 3 page table to identify an entry in the Level 3 page table. The identified entry in the Level 3 page table contains a physical address that serves as a base address of an exemplary 4 KB page of memory. The final 12 bits **310** of the VA provide the offset into the identified 4 KB page of memory, the address of which is the PA to which the VA is translated.

[0034] FIG. **3b** is an example two-stage translation **350** for translating a 49-bit virtual address (VA) to a physical address (PA), including translating one or more intermediate physical addresses (IPA) in accordance with examples of this description. In an example, a value of one of the MMRs **210** of the MMU **108** is determinative of whether the MMU **108** is configured to perform one-stage translation as shown in FIG. **3a** or two-stage translation as shown in FIG. **3b**. The example translation **350** is representative of the functionality performed by the page table walker engine **122** responsive to receiving a transaction that entails full or partial address translation.

[0035] The two-stage translation **350** differs from the one-stage translation **300** described above in that the physical address at each identified entry is treated as an intermediate physical address that is itself translated to a physical address. For example, the most significant bit of the 49-bit VA **352** again specifies one of two table base registers (e.g., TBR0 or TBR1, implemented in the MMRs **210**). However, the physical address contained by the selected table base register is treated as IPA **354**, which is translated to a physical address. In this example, a virtual table base register (e.g., VTBR, implemented in the MMRs **210**) contains a physical address that is a base address of a first page table **356**. The remainder of the IPA **354** is translated as described above with respect to the 49-bit VA of FIG. **3a**.

[0036] The resulting 40-bit PA **358** is a base address for a first page table **360** for the translation of the 49-bit VA **352** to the final 40-bit PA **380**, while a first group of nine bits **362** of the VA **352** provides the offset from the base address specified by the PA **358** into the first page table **360** to identify an entry in the first page table **360**. However, unlike the one-stage translation **300**, the entry in the first page table **360** is treated as an IPA (e.g., replacing previous IPA **354**) that is itself translated to a new PA **358**, which is then used as a base address for a second page table **364**. That is, the entry in the first page table **360** is not used directly as a base address for the second page table **364**, but rather is first translated as an IPA **354** to a PA **358** and that resulting PA **358** is then used as the base address for the second page table **364**. This process continues in a like manner for a third page table **366** and a fourth page table **368** before arriving at the final 40-bit PA **380**. For example, the address contained in the final Level 3 page table (e.g., page table **368**) is also an IPA that is translated in order to arrive at the final 40-bit PA **380**.

[0037] Thus, while performing a one-stage translation **300** may entail multiple memory accesses, performing a two-stage translation **350** may entail still more memory accesses, which can reduce performance when many such translations are performed. Additionally, FIGS. **3a** and **3b** are described with respect to performing a full address translation. However, as described above, in some instances a translation request partially hits the TLB **120**, for example where a certain number of most significant bits of a virtual address of the translation request match an entry in the TLB **120**. In such examples, the page table walker engine **122** does not necessarily perform each level of the address translation and instead only performs part of the address translation. For example, referring to FIG. **3a**, if the most significant 19 bits of a virtual address of a translation request match an entry in the TLB **120**, the page table walker engine **122** begins with the base address of the Level 2 page table and only needs to perform address translation using the third and fourth

groups of nine bits **306**, **308**. In other examples, similar partial address translations are performed with regard to a two-stage translation **350**.

[0038] FIG. **4** is a block diagram of a system **400** that includes the processor core **106** and the MMU **108** described above with respect to FIGS. **1** and **2**. In the example of FIG. **4**, the MMRs **210** of the MMU **108** include a mode configuration register **402**. The MMU **108** is configured to operate in different modes responsive to the value of the mode configuration register **402**. For example, responsive to the mode configuration register **402** being a first value, the MMU **108** is configured to operate in a software-managed mode. Responsive to the mode configuration register **402** being a second value, the MMU **108** is configured to operate in a hardware-managed mode. These modes are explained in further detail below.

[0039] In FIG. **4**, the processor core **106** is schematically shown as including (e.g., executing) an operating system (OS) **404**. In some cases, the OS **404** is a general purpose OS that manages multiple applications, tasks, and the like. However, in other cases, the OS **404** is lower level (e.g., bare metal) software or an OS **404** that is more specifically tailored to a specific purpose (e.g., a real-time OS (RTOS)).

[0040] In some cases, lower level software such as bare metal software or a RTOS lacks of granular control over the management of memory attributes. For example, such lower level software is commonly unable to treat memory as cacheable; utilize address virtualization; specify whether memory is only readable, only writable; or the like. A lower level software or OS is also commonly unable to support different memory types at the same time, support memory having different page sizes, or provide control over memory access permissions. However, a higher level, general purpose OS has additional control over memory attributes, for example via table pages in a virtual memory system as described above.

[0041] In accordance with examples of this disclosure, the MMU **108** is configured to operate in the software-managed mode to allow a lower level OS **404** to utilize the MMU **108** and associated TLB **120**. The MMU **108** is also configured to operate in the hardware-managed mode to allow general purpose OS **404** and its underlying applications and tasks to utilize the MMU **108** and associated TLB **120** and page table walker engine **122**, for example as described above with respect to FIGS. **1**, **2**, **3a**, and **3b**.

[0042] Responsive to operating in the software-managed mode, the MMU **108** is configured to receive and grant a TLB modification request **406** from the processor core **106** to directly modify an entry in the TLB **120**. The TLB modification request **406** specifies information needed to populate an entry in the TLB **120**, such as a virtual address, a corresponding physical address (e.g., to which the virtual address translates), memory types (e.g., device memory, non-cacheable memory, cacheable writeback memory, cacheable write-through memory, and the like), memory attributes (e.g., shareability, whether memory is pre-fetchable, inner cacheability, outer cacheability, and the like), page size, permission(s) for the page (e.g., readable, writable, executable, a supervisor- or user-specific accessibility, demotion from secure to non-secure), a context security level, a privilege level, and the like. Thus, in the software-managed mode, the OS **404** is able to access TLB **120** entries in a manner similar to indexed (e.g., memory-mapped) control registers and populate or modify the values of those TLB **120** entries as needed. In some examples, the system **400** initially begins in software-managed mode with a flat, or default view of memory, in which entries in the TLB **120** have equivalent virtual address and physical address data. For example, when the TLB **120** initially contains such default entries, a virtual address that hits the TLB **120** results in a physical address that is identical to the virtual address.

[0043] By contrast, responsive to operating in the hardware-managed mode, the MMU **108** is configured to deny such a TLB modification request **406**. In particular, in the hardware-managed mode, the TLB **120** is modified as described above (e.g., with respect to FIGS. **1** and **2**), including receiving address translation requests from the processor core **106** and walking page tables (e.g., using the page table walker engine **122**) to determine address translations, which are then cached in

the TLB **120**.

[0044] The MMU **108** is configured to receive address translation requests from the processor core **106** regardless of whether the MMU **108** is operating in the software-managed mode or the hardware-managed mode. In some examples, however, address translation requests are handled differently depending on whether the MMU **108** is operating in the software-managed mode or the hardware-managed mode.

[0045] For example, responsive to the MMU **108** operating in the software-managed mode and receiving an address translation request that misses the TLB **120** (e.g., a corresponding translation for a virtual address specified by the address translation request is not in the TLB **120**), the MMU **108** is configured to return a memory fault to the processor core **106**. A memory fault is returned to the processor core **106** in this case because it is assumed that a lower level OS **404** operating in software-managed mode is aware of the translations that it has caused to be populated in the TLB **120**, and thus any address translation request for a virtual address that is not already in the TLB **120** is an inappropriate request that should result in a fault.

[0046] Continuing this example, responsive to the MMU **108** operating in the hardware-managed mode and receiving an address translation request that misses the TLB **120** (e.g., a corresponding translation for a virtual address specified by the address translation request is not in the TLB **120**), the MMU **108** is configured to cause the page table walker engine **122**, described above, to walk one or more page tables to determine an address translation for the specified virtual address. The MMU **108** is then configured to cache the determined address translation in the TLB **120** and/or a lower level memory such as the cache **110**.

[0047] In some examples, the OS **404** or another process executed by the processor core **106** is configured to modify the mode configuration register **402** (e.g., to switch from software-managed mode to hardware-managed mode or vice versa) by providing a mode change request **408** to the MMU **108**. This allows context switching between lower level processes (e.g., a lower level OS **404** in software-managed mode) and higher level processes (e.g., a higher level OS **404** in hardware-managed mode) on the same processor core **106**. In these examples, only an OS **404** or other process having at least a certain privilege or context status is capable of modifying (e.g., writing to) the mode configuration register **402**. As a result, lower privilege or lower context processes (e.g., a non-secure user mode process) are restricted from modifying the mode configuration register **402**. For example, the MMU **108** is configured to process (e.g., grant) a mode change request **408** received from a process having a sufficient privilege level and to not process (e.g., deny) a mode change request **408** received from a process having an insufficient privilege level.

[0048] In some examples, responsive to switching from software-managed mode to hardware-managed mode, the MMU **108** is configured to invalidate TLB **120** entries used in the software-managed mode. Subsequently, when the MMU **108** operates in the hardware-managed mode, the TLB **120** is repopulated as described above, for example using the page table walker engine **122** to walk one or more page tables to determine address translations and caching the resulting address translations in the TLB **120**. In these examples, the mode configuration register **402** (or another MMR **210**) specifies whether to invalidate TLB **120** entries when switching from software-managed mode to hardware-managed mode.

[0049] Similarly, responsive to switching from hardware-managed mode to software-managed mode, the MMU **108** is configured to invalidate TLB **120** entries used in the hardware-managed mode. However, as explained above, in software-managed mode, the TLB **120** may be expected to begin in a default state in which entries in the TLB **120** have equivalent virtual address and physical address data. For example, when the TLB **120** initially contains such default entries, a virtual address that hits the TLB **120** results in a physical address that is identical to the virtual address. Thus, after invalidating the TLB **120** entries responsive to switching to software-managed mode, the MMU **108** is configured to populate the TLB **120** with default entries to enable such flat

device memory at an initial time. In these examples, the mode configuration register **402** (or another MMR **210**) specifies whether to invalidate TLB **120** entries and populate with default entries when switching from hardware-managed mode to software-managed mode. In some examples, the mode configuration register **402** (or another MMR **210**) also specifies an amount of default memory space, which indicates to the MMU **108** a number of TLB **120** entries to populate as default entries. The default entries populated in the TLB **120** provide a contiguous address space in which the translated physical address is the same as the virtual address used by the processor core **106**.

[0050] In other examples, responsive to switching from software-managed mode to hardware-managed mode, the MMU **108** is configured to retain TLB **120** entries used in the software-managed mode. The retained TLB **120** entries are “pinned” or otherwise marked so that they are not visible and/or accessible by processes that access the MMU **108** when subsequently operating in the hardware-managed mode. Similarly, responsive to switching from hardware-managed mode to software-managed mode, the MMU **108** is configured to retain TLB **120** entries used in the hardware-managed mode. The retained TLB **120** entries are “pinned” or otherwise marked so that they are not visible and/or accessible by processes that access the MMU **108** when subsequently operating in the software-managed mode. In these examples, the mode configuration register **402** (or another MMR **210**) specifies whether to retain TLB **120** entries when switching from hardware-managed mode to software-managed mode or vice versa. The ability to retain TLB **120** entries used in a previous mode facilitates faster start-up times when switching between software- and hardware-managed modes, and vice versa, since fewer address translations need to be initially performed when switching to hardware-managed mode, or programmed into the TLB **120** in software-managed mode.

[0051] FIG. **5** is a flow chart of a method **500** of operating the MMU **108** in software- and hardware-managed modes in accordance with various examples.

[0052] The method **500** begins in block **502** with receiving, by the MMU **108** (e.g., including the TLB **120** and mode configuration register **402**), a request from the processor core **106** to directly modify an entry in the TLB **120** (e.g., TLB modification request **406**). The TLB modification request **406** specifies information needed to populate an entry in the TLB **120**, such as a virtual address, a corresponding physical address (e.g., to which the virtual address translates), and the like.

[0053] The method **500** then continues in block **504** with, responsive to the mode configuration register **402** having a first value, operating the MMU **108** in a software-managed mode by modifying the entry in the TLB **120** according to the request. Thus, in the software-managed mode, an OS or other supervisory program executed by the processor core **106** is able to access TLB **120** entries in a manner similar to indexed (e.g., memory-mapped) control registers and populate or modify the values of those TLB **120** entries as needed.

[0054] The method **500** continues further in block **506** with, responsive to the mode configuration register **402** having a second value, operating the MMU **108** in a hardware-managed mode by denying the request. As described above, in the hardware-managed mode, the TLB **120** is modified by receiving address translation requests from the processor core **106** and walking page tables (e.g., using the page table walker engine **122**) to determine address translations, which are then cached in the TLB **120**.

[0055] In the foregoing discussion and in the claims, the terms “including” and “comprising” are used in an open-ended fashion, and thus mean “including, but not limited to”

[0056] The term “couple” is used throughout the specification. The term may cover connections, communications, or signal paths that enable a functional relationship consistent with the description of the present disclosure. For example, if device A generates a signal to control device B to perform an action, in a first example device A is coupled to device B, or in a second example device A is coupled to device B through intervening component C if intervening component C does

not substantially alter the functional relationship between device A and device B such that device B is controlled by device A via the control signal generated by device A.

[0057] An element or feature that is “configured to” perform a task or function may be configured (e.g., programmed or structurally designed) at a time of manufacturing by a manufacturer to perform the function and/or may be configurable (or re-configurable) by a user after manufacturing to perform the function and/or other additional or alternative functions. The configuring may be through firmware and/or software programming of the device, through a construction and/or layout of hardware components and interconnections of the device, or a combination thereof. Additionally, uses of the phrases “ground” or similar in the foregoing discussion include a chassis ground, an Earth ground, a floating ground, a virtual ground, a digital ground, a common ground, and/or any other form of ground connection applicable to, or suitable for, the teachings of the present disclosure. Unless otherwise stated, “about,” “approximately,” or “substantially” preceding a value means ± 10 percent of the stated value.

[0058] The above discussion is illustrative of the principles and various embodiments of the present disclosure. Numerous variations and modifications will become apparent to those skilled in the art once the above disclosure is fully appreciated. The following claims should be interpreted to embrace all such variations and modifications.

Claims

1. A device comprising: a processor core; a memory management unit coupled to the processor core, wherein the memory management unit includes a first memory; and a cache memory system that includes a second memory, wherein the memory management unit is capable of: in a first mode of operation, based on a miss in the first memory, requesting a portion of a page table from the second memory; and in a second mode of operation, based on a miss in the first memory, providing an indication of a page fault without requesting a portion of the page table from the second memory.
2. The device of claim 1, wherein the memory management unit is capable of: receiving, from the processor core, an address translation request that includes a first address; and attempting to translate the first address using the first memory, wherein the attempting to translate includes determining whether the first address is a miss in the first memory.
3. The device of claim 1, wherein the memory management unit is capable of determining whether to deny or permit a request by the processor core to modify an entry in the first memory based whether the memory management unit is operating in the first mode of operation or the second mode of operation.
4. The device of claim 1, wherein the memory management unit includes a mode configuration register capable of storing an indication of whether to the memory management unit is operating in the first mode of operation or the second mode of operation.
5. The device of claim 1, wherein the memory management unit is capable of: receiving, from the processor core, a request to transition from the first mode of operation to the second mode of operation; and determining, based on the request, whether to invalidate a first set of entries in the first memory and populate the first memory with a second set of entries.
6. The device of claim 1, wherein: the memory management unit includes a mode configuration register capable of storing an indication of whether to, upon transitioning to the second mode of operation, invalidate a first set of entries in the first memory and populate the first memory with a second set of entries.
7. The device of claim 1, wherein: the memory management unit is capable of, upon transitioning between the first mode of operation and the second mode of operation, determining whether to retain a subset of entries in the first memory and to make the subset of entries in the first memory inaccessible until a subsequent transition.

- 8.** The device of claim 1, wherein the memory management unit includes table walker circuitry capable of requesting a portion of the page table from the second memory.
 - 9.** The device of claim 1, wherein the first memory is a translation lookaside buffer.
 - 10.** The device of claim 1, wherein the second memory is a cache memory.
 - 11.** The device of claim 1, wherein the second memory is a level-one (L1) cache memory.
 - 12.** A device comprising: a memory management unit that includes: a translation lookaside buffer; and page table walker circuitry coupled to the translation lookaside buffer; wherein the memory management unit is capable of: operating in a first mode of operation by, in response to a miss in the translation lookaside buffer, causing the page table walker circuitry to request a portion of a page table from a memory; and operating in a second mode of operation by, in response to a miss in the translation lookaside buffer, inhibiting the page table walker circuitry from requesting a portion of the page table from the memory.
 - 13.** The device of claim 12, wherein the operating in the second mode of operation includes, in response to a miss in the translation lookaside buffer, providing an indication of a page fault.
 - 14.** The device of claim 12, wherein: the operating in the first mode of operation includes denying a request to modify an entry in the translation lookaside buffer; and the operating in the second mode of operation includes permitting a request to modify an entry in the translation lookaside buffer.
 - 15.** The device of claim 12, wherein: the memory management unit includes a mode configuration register capable of storing an indication of whether to the memory management unit is operating in the first mode of operation or the second mode of operation.
 - 16.** The device of claim 12, wherein the memory management unit is capable of: receiving a request to transition from the first mode of operation to the second mode of operation; and determining, based on the request, whether to invalidate a first set of entries in the translation lookaside buffer and populate the translation lookaside buffer with a second set of entries.
 - 17.** The device of claim 12, wherein the memory management unit is capable of: receiving a request to transition from the first mode of operation to the second mode of operation; and determining whether to retain a subset of entries in the translation lookaside buffer and to make the subset of entries in the translation lookaside buffer inaccessible until a subsequent transition.
 - 18.** A method comprising: receiving, by a memory management unit, a first request for a first address translation; based on the memory management unit being in a first mode and based on the first request being associated with a miss in a first memory, utilizing a set of circuitry to: retrieve a first set of address translations from a second memory; and store the first set of address translations in the first memory; receiving, by the memory management unit, a second request for a second address translation; and based on the memory management unit being in a second mode and based on the second request being associated with a miss in the first memory, inhibiting the set of circuitry from: retrieving a second set of address translations from the second memory; and storing the second set of address translations in the first memory.
 - 19.** The method of claim 18 further comprising, based on the memory management unit being in the second mode and based on the second request being associated with a miss in the first memory, providing an indication of a page fault.
 - 20.** The method of claim 18 further comprising determining whether to permit or to deny a request to modify an entry in the first memory based on whether the memory management unit is in the first mode or the second mode.
-