

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284647

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

SUN; Zhonghua et al.

TECHNIQUES TO MULTIPLY MEMORY ACCESS BANDWIDTH USING A PLURALITY OF LINKS

Abstract

Examples include techniques to multiply memory access bandwidth using a plurality of links. Example techniques may include generation and use of forwarding tables separately maintained at devices coupled via high speed internal-connect links (HSILs). A forwarding table to enable a first device to route a memory request received by the first device to access a memory address of a memory at a second device. The memory request received by the first device from a host compute device via a link between the first device and the host compute device, the memory request to be forwarded to the second device via an HSIL coupled between the first and second devices.

Inventors: SUN; Zhonghua (Shanghai, CN), ZHANG; Cong (Shanghai, CN), ZHANG; Di (Shanghai, CN), CHEN; Gang (Shanghai, CN), LI; Zhuangzhi (Shanghai, CN)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 88918134

Appl. No.: 18/859417

Filed (or PCT Filed): May 23, 2022

PCT No.: PCT/CN22/94445

Publication Classification

Int. Cl.: G06F13/16 (20060101); G06F12/02 (20060101)

U.S. Cl.:

CPC G06F13/1615 (20130101); G06F12/0284 (20130101); G06F13/1678 (20130101);

Background/Summary

TECHNICAL FIELD

[0001] Examples described herein are related to access of a memory maintained at a device coupled to a processor.

BACKGROUND

[0002] A rapid rise in artificial intelligence (AI) and high performance computing (HPC) usage has increased an importance of a role that types of accelerators such as graphic processing units (GPUs) serve in AI and HPC. GPUs, in particular, have gained a relatively high amount of attention as GPU performance has increased by more than a factor of 10× in the past decade. However, as GPU performance continues to increase, these higher performing GPUs are somewhat starved for data due to a relatively slower link that couples these higher performing GPUs to a central processing unit (CPU). For example, a link operated according to a specification or standards such as published by the PCI-SIG® to include the PCI Express Specification, Rev. 6.0, Ver. 1.0, published in January of 2022 (“the PCIe specification”).

[0003] Ongoing efforts are being made to accelerate or increase a bandwidth of links between a group of GPUs in a same system so that GPUs in the group may directly exchange data at a high speed. For example, NVIDIA® has developed a GPU-to GPU link known as NVLink™. NVLink may have a per GPU bandwidth of 500 gigabytes per second (GB/s) or more. Other GPU manufacturers such as Intel® Corporation are working on multi die fabrics that include links having bandwidths similar to that of NVLink. NVLinks included in multi die fabrics or other types of similar links included in multi die fabrics that interconnect GPUs and/or accelerators may be referred to as “high speed internal-connect links” (HSILs). In some examples, these types of HSILs that interconnect GPUs are designed to facilitate GPU-GPU or accelerator-accelerator communications.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 illustrates an example system.

[0005] FIG. 2 illustrates example memory forwarding tables.

[0006] FIG. 3 illustrates an example scheme.

[0007] FIG. 4 illustrates an example DVSEC range structure.

[0008] FIG. 5 illustrates an example DVSEC capabilities register.

[0009] FIG. 6 illustrates an example CXL early detection table (CEDT).

[0010] FIG. 7 illustrates an example CBHS table.

[0011] FIG. 8 illustrates an example CXL binding HDM structure (CBHS) table.

[0012] FIG. 9 illustrates an example first logic flow.

[0013] FIG. 10 illustrates an example second logic flow

[0014] FIG. 11 illustrates an example third logic flow.

[0015] FIG. 12 illustrates an example process.

[0016] FIG. 13 is a block diagram of an example processing system.

[0017] FIG. 14A is a block diagram of an example of a processor having one or more processor cores, an integrated memory controller, and an integrated graphics processor.

[0018] FIG. 14B is a block diagram of an example hardware logic of a graphics processor core block.

[0019] FIG. 14C illustrates an example graphics processing unit (GPU) that includes dedicated sets of graphics processing resources arranged into multi-core groups.

[0020] FIG. 14D is a block diagram of an example general-purpose graphics processing unit (GPGPU) that can be configured as a graphics processor and/or compute accelerator.

[0021] FIG. 15A is a block diagram of an example graphics processor, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces.

[0022] FIG. 15B illustrates an example graphics processor having a tiled architecture.

[0023] FIG. 15C illustrates an example compute accelerator.

[0024] FIG. 16 is a block diagram of an example graphics processing engine of a graphics processor.

[0025] FIG. 17A illustrates an example graphics core cluster.

[0026] FIG. 17B illustrates an example vector engine of a graphics core.

[0027] FIG. 17C illustrates an example matrix engine of a graphics core.

[0028] FIG. 18 illustrates an example tile of a multi-tile processor.

[0029] FIG. 19 is a block diagram illustrating graphics processor instruction formats according to some examples.

[0030] FIG. 20 is a block diagram of another example of a graphics processor.

[0031] FIG. 21A is a block diagram illustrating a graphics processor command format that may be used to program graphics processing pipelines.

[0032] FIG. 21B is a block diagram illustrating an example graphics processor command sequence.

[0033] FIG. 22 illustrates an example graphics software architecture for a data processing system.

[0034] FIG. 23A is a block diagram illustrating an example IP core development system that may be used to manufacture an integrated circuit to perform operations.

[0035] FIG. 23B illustrates a cross-section side view of an example integrated circuit package assembly 1170.

[0036] FIG. 23C illustrates an example package assembly that includes multiple units of hardware logic chiplets connected to a substrate.

[0037] FIG. 23D illustrates an example package assembly including interchangeable chiplets.

[0038] FIG. 24 is a block diagram illustrating an example system on a chip integrated circuit that may be fabricated using one or more IP cores.

[0039] FIG. 25A illustrates an example graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores.

[0040] FIG. 25B illustrates an additional example of a graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores.

DETAILED DESCRIPTION

[0041] In some example computing systems of today, accelerators or GPUs are coupled with or connected to a CPU through PCIe links. A PCIe link operated according to the PCIe specification (e.g., Rev. 6.0) may have a bandwidth or throughput of approximately 128 GB/sec. Comparatively, HSILs that interconnect groups of GPUs or accelerators may have a data bandwidth or throughput of 500 GB/sec or more and thus may have a data bandwidth of least 5 times greater than these types of PCIe links. As a result of this at least 5 times difference in data bandwidth or throughput, a CPU-to-GPU link becomes a bottleneck for processing large amounts of data in a system that includes GPUs interconnected via fast HSILs but separately connected with a CPU via substantially slower PCIe links.

[0042] A new technical specification by the Compute Express Link (CXL) Consortium is the Compute Express Link Specification, Rev. 2.0, Ver. 1.0, published Oct. 26, 2020, hereinafter referred to as “the CXL specification”. The CXL specification introduced the on-lining and off-lining of memory at a device (e.g., a GPU device or an accelerator device) attached to a host computing device through separate links that are each configured to operate in accordance with the CXL specification and are hereafter referred to as “CXL links”. The memory at the device that is

on-lined or off-lined according to the CXL specification is referred to as host-managed device memory (HDM). Typically, each GPU or accelerator of a system of GPU or accelerators are connected to a CPU via separate CXL links. Since CXL links also operate according to the PCIe specification, throughput or bandwidth of CXL links may be limited to the 128 GB/sec that is mentioned above for PCIe links. So, in examples where a system of GPUs have GPU-GPU HSIL interconnections with at least 5 times a bandwidth of 128 GB/sec, a CXL link between a CPU and a GPU of the system may be over busy (e.g., at maximum bandwidth) and other CXL links to other GPUs may be idle or underutilized. It is with respect to these challenges that the examples described herein are needed.

[0043] FIG. 1 illustrates an example system **100**. In some examples, as shown in FIG. 1, system **100** includes host compute device **105** that has a root complex **120** to couple or connect with devices **130-A**, **130-B** and **130-C** through respective ports **121-1**, **121-2** and **121-N** via respective host links **140-1**, **140-2** and **140-3**, where “N” is any positive, whole integer greater than 2. Host compute device **105**, as shown in FIG. 1 also couples with a host system memory **110** via one or more memory channel(s) **101**. For these examples, host compute device **105** includes a host operating system (OS) **102** to execute or support one or more device driver(s) **104**, a host basic input/output system (BIOS) **106**, one or more host application(s) **108** and a host central processing unit (CPU) **107** to support compute operations of host compute device **105**.

[0044] In some examples, although shown in FIG. 1 as being separate from host CPU **107**, root complex **120** may be integrated with host CPU **107** in other examples. For either example, root complex **120** may be arranged to function as a type of PCIe root complex for host CPU **107** and/or other elements of host compute device **105** to communicate with devices such as devices **130** via use of PCIe-based communication protocols and communication links. Root complex **120** may also be configured to operate in accordance with the CXL specification and as shown in FIG. 1, includes a home agent **124** to facilitate communications with devices **130** via host links **140**. For these examples, host links **140** may be configured to operate according to the CXL specification. As shown in FIG. 1, root complex **120** includes HDM decoders **126** that may be programmed to facilitate a mapping of host to device physical addresses for use in accessing memory at or attached to devices **130**. A memory controller (MC) **122** at root complex **120** may control/manage access to host system memory **110** through memory channel(s) **101**. Host system memory **110** may include volatile and/or non-volatile types of memory. In some examples, host system memory **110** may include one or more dual in-line memory modules (DIMMs) that may include any combination of volatile or non-volatile memory. For these examples, memory channel(s) **101** and host system memory **110** may operate in compliance with a number of memory technologies described in various standards or specifications, such as DDR3 (DDR version 3), originally released by JEDEC (Joint Electronic Device Engineering Council) on Jun. 27, 2007, DDR4 (DDR version 4), originally published in September 2012, DDR5 (DDR version 5), originally published in July 2020, LPDDR3 (Low Power DDR version 3), JESD209-3B, originally published in August 2013, LPDDR4 (LPDDR version 4), JESD209-4, originally published by in August 2014, LPDDR5 (LPDDR version 5, JESD209-5A, originally published by in January 2020), WIO2 (Wide Input/output version 2), JESD229-2 originally published in August 2014, HBM (High Bandwidth Memory), JESD235, originally published in October 2013, HBM2 (HBM version 2), JESD235C, originally published in January 2020, or HBM3 (HBM version 3), JESD238, originally published in January 2022, or others or combinations of memory technologies, and technologies based on derivatives or extensions of such specifications. The JEDEC standards or specifications are available at www.jedec.org.

[0045] According to some examples, as shown in FIG. 1, devices **130-A**, **130-B** and **130-C** may be interconnected via respective HSILs **152-1**, **152-2** and **152-3** included in a multi die fabric **150**. For these examples, each link included in HSILs **152** may be capable of having a data bandwidth or throughput (e.g., >500 GB/sec) that is substantially higher than a data bandwidth or throughput of

each link included host links **140** (e.g., around 100 GB/sec). As described in more detail below, when elements of system **100** are operated according to the CXL specification, a coherent memory accessing capability is possible between devices **130-A**, **130-B** and **130-C** and host compute device **105** via respective host links **140-1**, **140-2** and **140-3**. The coherent memory accessing capability makes it possible to increase data bandwidth between host compute device **105** and a device from among devices **130** by leveraging adjacent host links to transfer data via host links **140** in parallel so the data bandwidth could be multiplexed or multiplied for access to memory at devices **130**.

[0046] In some examples, as shown in FIG. **1**, devices **130** each include a port **131**, a coherency agent **132** (e.g., a CXL device coherency agent (DCOH)), memory router (MR) circuitry **133**, host-managed device memory (HDM) **134**, memory controller circuitry **135**, compute circuitry **136**, an HSIL interface (I/F) **137** or registers **138**. As described more below, circuitry or logic at devices **130** such as coherency agent **132** may separately receive a CXL.mem address from host compute device **105** via host links **140** and then utilize separately maintained forwarding tables included in registers **138** to enable MC circuitry **135** to access memory addresses for a targeted HDM **134** of a device from among devices **130**. The accessed memory addresses may enable devices that don't include the targeted HDM memory to read or write data to the targeted HDM memory via a path that is routed via a first link included in host links **140** and a second link included in HSILs **152**. Meanwhile the device including the targeted HDM memory may access the memory addresses via a path that is routed through a single link of host links **140**. Since HSILs **152** have substantially higher bandwidths than host links **140**, adding these HSIL links to access memory addresses in a targeted device HDM memory at another device memory may add relatively little latency. As a result, this method or scheme of parallel access using memory forwarding tables may nearly triple the available data bandwidth between host compute device **105** and the targeted device.

[0047] According to some examples, registers **138** included in devices **130** may be configured (e.g., programmed) to maintain a forwarding table and PCIe designated vendor-specific capability (DVSEC) structures. As mentioned briefly above and described more below, a device may use a forwarding table to determine where to access HDM memory addresses located at a targeted HDM memory. Also, as described more below, the forwarding tables maintained in registers **138** may be set or programmed based, at least in part on, information included in various DVSEC structures also maintained in registers **138**. The various DVSEC structures, for example, may be set or programmed by a manufacture or vendor of devices **130** and then used by a BIOS (e.g., host BIOS **106**) to set or program at least a portion of registers **138** to include a forwarding table.

[0048] According to some examples, HDM **134** may include volatile and/or non-volatile types of memory for use by compute circuitry **136** to execute, for example, a workload. The volatile and/or volatile types of memory may be resident on a same die or package as device **130** (e.g., stacked HBM die or stacked DRAM die) or may be an attached memory device (e.g., a DIMM). In some examples, compute circuitry **136** may be a GPU and the workload may be a graphics processing related workload or may be related to an AI workload that utilizes GPU processing capabilities offloaded to compute circuitry **136** by host CPU **107**. In other examples, compute circuitry **136**, may be at least part of an FPGA, ASIC or CPU serving as an accelerator and the workload may be offloaded from host compute device **105** for execution by these types of compute circuitry that include an FPGA, ASIC or CPU. Also, coherency agent **132**, MR circuitry **133**, or MC circuitry **135** may be at least part of the FPGA, ASIC or CPU that is used for offloading the workload from host compute device **105**.

[0049] As mentioned above, host system memory **110** and HDM **134** may include volatile or non-volatile types of memory. Volatile types of memory may include, but are not limited to, random-access memory (RAM), Dynamic RAM (DRAM), DDR synchronous dynamic RAM (DDR SDRAM), GDDR, HBM, static random-access memory (SRAM), thyristor RAM (T-RAM) or zero-capacitor RAM (Z-RAM). Non-volatile memory may include byte or block addressable types of non-volatile memory having a 3-dimensional (3-D) cross-point memory structure that includes,

but is not limited to, chalcogenide phase change material (e.g., chalcogenide glass) hereinafter referred to as “3-D cross-point memory”. Non-volatile types of memory may also include other types of byte or block addressable non-volatile memory such as, but not limited to, multi-threshold level NAND flash memory, NOR flash memory, single or multi-level phase change memory (PCM), resistive memory, nanowire memory, ferroelectric transistor random access memory (FeTRAM), anti-ferroelectric memory, resistive memory including a metal oxide base, an oxygen vacancy base and a conductive bridge random access memory (CB-RAM), a spintronic magnetic junction memory, a magnetic tunneling junction (MTJ) memory, a domain wall (DW) and spin orbit transfer (SOT) memory, a thyristor based memory, a magnetoresistive random access memory (MRAM) that incorporates memristor technology, spin transfer torque MRAM (STT-MRAM), or a combination of any of the above.

[0050] FIG. 2 illustrates example forwarding tables **200**. In some examples, as shown in FIG. 2, forwarding tables **200** includes a forwarding table **210**, a forwarding table **220** and a forwarding table **230**. For these examples, forwarding table **210** may be maintained in registers **138-A** at device **130-A**, forwarding table **220** may be maintained in registers **138-B** at device **130-B** and forwarding table **230** may be maintained in registers **138-C** at device **130-C**.

[0051] As described more below, upon initialization of device **130-A**, a BIOS (e.g., host BIOS **106**) may program at least a portion of registers **138-A** to set a first value indicating a local address that serves as a virtual HDM (vHDM) base address for memory addresses of HDM **134-A** that can be mapped to a host physical address (HPA) system (e.g., maintained by host compute device **105**). The BIOS may also set respective second and third values indicating respective destination addresses for devices **130-B** and **130-C**. In some examples, the respective destination addresses may serve as base addresses of HDM **134-B** and HDM **134-C** that can also be mapped to the HPA system and indicate definite, respective HPAs of HDM **134-B** and HDM **134-C**. Also, a size or memory capacity associated with each base address that is indicated for the local address may be a same size associated with respective destination addresses for HDM **134-B** and HDM **134-C**. A similar process may occur upon initialization of devices **130-B** and **130-C** to program respective at least portions of registers **138-B** and **138-C** to build vHDM and HDM maps as shown in FIG. 2 for forwarding tables **220** and **230**.

[0052] FIG. 3 illustrates an example scheme **300**. According to some examples, scheme **300** shows an example of how data may be routed through host links separately connected between a host compute device and different devices in order to read or write data to an HDM of a device from among the different devices using forwarding tables programmed to registers of each device. For these examples, compute device **105** and devices **130-A**, **130-B** and **130-C** may be configured to operate host links **140-1** to **140-N** according to the CXL specification and may use CXL.mem protocols to route data via these host links. As shown in FIG. 3, elements of host compute device **105** such as host OS **102**, host system memory **110**, root complex **120** and ports **121-1** to **121-N** are describe below as taking part in scheme **300**. Also, elements of devices **130-A**, **130-B** or **130-C** that are interconnected via HSILs **152-1** to **152-3** of multi die fabric **150** are described below as taking part in scheme **300**. Scheme **300** is not limited to these elements of devices **130-A**, **130-B**, **130-B** or compute device **105** as taking part in scheme **300**.

[0053] According to some examples, HPA/system address space **301** as shown in FIG. 3 may be an example of how a host-managed device memory (HDM) **134-A** may be mapped to host system memory **110**. As described more below, binding **310** of HPA/system address space **301** may be based on forwarding tables programmed to registers **138-A**, **138-B** and **138-C** (e.g., by host BIOS **106**) that bind virtual HDMs (vHDMs) of device **130-B** and **130-C** to device **130-A**'s HDM **134-A**. For example, as shown in FIG. 3, vHDM_B2A included in forwarding table **220** of device **130-B** and vHDM_C2A included in forwarding table **230** of device **130-B** may be bound to HDM **134-A** according to binding **310**. Each bound address may be a same size or represent a memory address or range of memory addresses to access a same amount of memory capacity.

[0054] In some examples, logic and/or features of host OS **102** may decide to split data associated with an access to memory addresses for HDM **134-A**. For these examples, the splitting of the data includes a decision to multiplex host links **140-1** to **140-N** in order to increase a data bandwidth to access HDM **134-A** that also utilizes the higher bandwidth capabilities of HSILs **152-1** to **152-3** of multi die fabric **150** to facilitate that access of HDM **134-A**. Data1, Data2, Data3 shown in FIG. **3** represent the splitting of data to be multiplexed over host links **140-1** to **140-N**. Route legend **305** indicates that data1 having a solid, black line is routed through port **121-1** and host link **140-1** to device **130-A**, data2 having a dotted, black line is routed through port **121-2** and host link **140-2** to device **130-B** and data3 having a dash/dot, black line is routed through port **121-N** and host link **140-3**.

[0055] According to some examples, as described more below, host OS **102** may utilize CXL.mem request messages sent via host links **140-1** to **140-N** to access HDM **134-A** via multiple routes through these host links. Each device from among devices **130-A**, **130-B** and **130-C** may receive separate CXL.mem request messages to indicate which portion of HDM **134-A** is to be accessed. For these examples, circuitry of device **130-A** such as MR circuitry **133-A** may receive a CXL.mem request message indicating data1 is associated with memory addresses of HDM **134-A** to access. Since HDM **134-A** is local to device **130-A**, MC circuitry **135-A** accesses HDM **134-A** to complete the access to memory addresses associated with data1. Device **130-A** may then send a response message back to host OS **102** via host link **140-1**. Circuitry of device **130-B** such as MR circuitry **133-B** may receive a CXL.mem request message indicating data2 is associated with memory addresses of HDM **134-A** to access. Since HDM **134-A** is not local to device **130-B**, MR circuitry **133-B** may refer to forwarding table **220** to determine that vHDM_B2A indicates to route a memory request for data2 via HSIL **152-1** to device **130-A**. Circuitry of device **130-C** such as MR circuitry **133-C** may receive a CXL.mem request message indicating data3 is associated with memory addresses of HDM **134-A** to access. Since HDM **134-A** is not local to device **130-C**, MR circuitry **133-C** may refer to forwarding table **220** to determine that vHDM_C2A indicates to route a memory request for data3 via HSIL **152-3** to device **130-A**. MC **135-A**, upon receipt of the separate memory requests for data2 and data3 may access HDM **134-A** according to the memory addresses indicated in the separate memory requests and then send a response indicating a status of the memory access to HDM **134-A** back to device **130-B** via HSIL **152-1** and a response indicating a status of the memory access to HDM **134-A** back to device **130-C** via HSIL **152-3**. Devices **130-B** and **130-C** may then send separate responses indicating respective statuses of the memory accesses to HDM **134-A** back to host OS **102** via respective host links **140-2** and **140-N**. In some examples, if the CXL.mem request was to read data from HDM **134-A**, respective responses sent by devices **130-A**, **130-B** and **130-C** may include the data read from HDM **134-A**. In other examples, if CXL.mem request was to write data to HDM **134-A**, respective response sent by devices **130-A**, **130-B** and **130-C** may indicate that the data was successfully written to HDM **134-A**.

[0056] FIG. **4** illustrates an example DVSEC range structure **400**. In some examples, DVSEC range structure **400** may be maintained in registers **138** of devices **130** of system **100** and may represent an example of a type of PCIe DVSEC structure that uses bits **7:5** of DVSEC range structure **400** to indicate a virtual HDM class. For these examples, as shown in FIG. **4**, a bit pattern of **100** for bits **7:5** in DVSEC range structure **400** indicates a device supports vHDMs to enable the multiplexing of host links as mentioned above for scheme **300** and described more below. A vendor or manufacturer of device **130** may set the bit pattern in bits **7:5** to **100** to indicate support for use of vHDMs. DVSEC range structure **400** may be similar to a DVSEC CXL Range 2 Size Low structure described in the CXL specification but for the addition of a bit pattern of **100** in bits **7:5** to indicate support for a vHDM memory class.

[0057] FIG. **5** illustrates an example DVSEC capabilities structure **500**. According to some examples, DVSEC capabilities may be maintained in registers **138** of devices **130** of system **100**

and may represent an example of a type of PCIe DVSEC structure that uses bits 0:3 of DVSEC capabilities structure **500** to indicate a number of vHDM ranges implemented by a device. For these examples, a value of 0 indicates no support for vHDM ranges and a value of 1-15 indicates support for 1 to 15 vHDM ranges. Examples are not limited to 15 vHDM ranges or to a DVSEC capabilities register having 4 bits to indicate a bit value of 0 to 15.

[0058] In some examples, a vendor or manufacturer of devices **130** may set the bit value of DVSEC capabilities structure **500** to indicate a number of vHDM ranges supported by devices **130**. For example, devices **130-A**, **130-B** and **130-C** may have DVSEC capabilities structures having a bit value of 2 to indicate support for 2 vHDM ranges. Support for 2 vHDM ranges indicates that a device **130** could route memory requests to up to 2 adjacent devices as mentioned above for scheme **300**.

[0059] FIG. **6** illustrates an example CEDT **600**. According to some examples, CEDT **600** may be linked to the Advanced Configuration and Power Interface, ver. 6.4, published in January 2021 by the United Extensible Firmware Interface (UEFI) Forum (hereinafter referred to as “the ACPI specification”). A BIOS such as host BIOS **106** may expose device vHDM support information to an OS such as host OS **102** via an indication in CEDT **600** of a CXL binding HDM structure (CBHS) table. For example, a value 2 in CEDT **600** indicates that a CBHS has been built that includes device vHDM support information. For these examples, the BIOS may gather information from registers of devices to build the CBHS table and also build a binding vHDM structure table (both described more below) to expose the OS to vHDM support information gathered during initialization of devices in order to enable the OS to map HDM and vHDMs to an HPA/system address space such as HPA/system address space **301** shown in FIG. **3** and mentioned above.

[0060] FIG. **7** illustrates an example CBHS table **700**. CBHS table **700**, as briefly mentioned above, may be built by a BIOS based on vHDM support information gathered during initialization of a device, e.g., gathered from DVSEC structures maintained in device registers such as registers **138**. In some examples, as shown in FIG. **7**, CBHS table **700** includes a type field and a value of 2 in the type field indicates CBHS entries for an HDM are included in CHBHS table **700**. An HDM_Base field indicates a base address of the HDM's HPA (e.g., a system memory HPA). An HDM_Size field indicates a size or memory capacity of an HDM. A Binding Mode field indicates whether a binding mode is support or not supported. A Num_vHDM field indicates a number of binding vHDM structures. A binding_vHDM[n] field indicates each binding vHDM structure that is fully mapped to the HDM, where each vHDM is from one adjacent device which could forward a memory request to the HDM. The “[n]” in the binding_vHDM[n] field may indicate that this field may contain up to “n” different entries, where “n” represents any positive integer greater than 0.

[0061] FIG. **8** illustrates an example binding vHDM structure table **800**. Binding vHDM structure table **800** may be built by an OS (e.g., host OS **102**) based on vHDM support information gathered during initialization of a device (e.g., gathered by a BIOS). In some examples, the information included in binding vHDM structure table **800** may be based on standard memory attributes provided by a vendor or manufacture of a device having host managed device memory and separate binding vHDM structure tables may be built by the OS for each vHDM that a device indicates as supporting. The OS, for example, may build the separate binding vHDM structure tables based on information included in a CBHS table such as example CBHS table **700** and based on information provided in an ACPI heterogeneous memory attribute table (HMAT). For these examples, as shown in FIG. **8**, a vHDM_Base field indicates a base address of the vHDM's HPA (e.g., a system memory HPA). A vHDM_Size field indicates a size or memory capacity of the vHDM. A vHDM_Read_Latency field indicates a read latency of the vHDM address. A vHDM_Write_Latency field indicates a write latency of the vHDM address. A vHDM_Read_Bandwidth field indicates a read bandwidth of the vHDM address. A vHDM_Write_Bandwidth field indicates a write bandwidth of the vHDM address.

[0062] FIG. **9** illustrates an example logic flow **900**. In some examples, logic flow **900** may be

implemented by a BIOS for a host computing device coupled with multiple devices (e.g., GPU devices and/or accelerator devices having host-managed device memory). For these examples, host BIOS **106** of host compute device **105** coupled with devices **130** via host links **140** as shown in FIG. **1** or **3** for system **100** may implement at least portions of logic flow **900**. The host computing device and devices connected via host links may be arranged to operate in accordance with the CXL specification. The devices may include registers having DVEC structures that include information gathered by the BIOS during initialization or startup of a system that includes the host computing device and the devices connected via separate host links. For example, information included in example DVSEC range structure **400** or example DVSEC capabilities structure **500** shown in FIGS. **4-5** that are maintained in registers **138** of devices **130** that may be gathered by host BIOS **106**. Also, host BIOS **106** may cause forwarding tables such as forwarding tables **200** to be maintained in devices **130** as shown in FIG. **2** as part of implementing logic flow **900**. Also, ACPI related tables such as CEDT table **600**, CBHS table **700** or binding vHDM structure table **800** shown in FIGS. **6-7** may be filled by host BIOS **106** as part of implementing logic flow **900**.

[0063] Logic flow **900** begins at block **905** where host BIOS **106** enters the flow. In some examples, host BIOS **106** enters logic flow **900** responsive to an initiation of system **100** and/or the initialization of at devices **130** included in system **100**.

[0064] Moving from block **905** to decision block **910**, host BIOS **106** completes any necessary training of host links **140** coupled between host compute device **105** and devices **130** to ensure the host links and devices can operate according to the CXL specification and confirms whether devices **130** indicate HDM/vHDM support. According to some examples, host BIOS **106** gathers information from a DVSEC range structure such as example DVSEC range structure **400** maintained in registers **138** of devices **130** to determine if a 100 bit pattern is set in bits 7:5 of the DVSEC range structure **400**. If the 100 bit pattern is set in bits 7:5, vHDM support is indicated, logic flow **900** moves to block **920**. Otherwise, logic flow **900** moves to block **915**.

[0065] Moving from decision block **910** to block **915**, host BIOS **106** completes other flows to finish its part of the initialization of system **100**.

[0066] Moving from decision block **910** to block **920**, host BIOS **106** gathers HDM and vHDM information from devices **130**. According to some examples, BIOS **106** gathers the HDM and vHDM information from a DVSEC capabilities structure such as example DVSEC capabilities structures **500** maintained in registers **138** of devices **130**.

[0067] Moving from block **920** to block **925**, host BIOS **106** maps system address space. In some examples, the system address space that is mapped may be HPA/system address space **301**.

[0068] Moving from block **925** to block **930**, host BIOS **106** builds vHDM and HDM mapping for each device. According to some examples, as part of building a mapping for each device, sub-block **925-A** is implemented to bind vHDM-B2A/vHDM-C2A to HDM_A to set various base addresses in HPA/system address space **301** in a similar manner as shown in forwarding table **210** of FIG. **2** and binding **310** shown in FIG. **3**, sub-block **925-B** is implemented to bind vHDM_A2B/vHDM_C2B to HDM_B to set various base addresses in a similar manner as shown in forwarding table **220** of FIG. **2**, and sub-block **925-C** is implemented to bind vHDM-A2C/vHDM_B2C to HDM_C to set various base addresses in a similar manner as shown in forwarding table **230** of FIG. **2**. For these examples, host BIOS **106** may cause a portion of registers **138** of devices **130-A**, **130-B** and **130-C** to be individually programmed to maintain respective forwarding tables **210**, **220** and **230** based on the above mentioned bindings.

[0069] Moving from block **925** to block **930**, host BIOS **106** fills ACPI tables for vHDMs bound to HDMs. In some examples, an example ACPI table to be filled may be a CEDT table such as CEDT table **600** shown in FIG. **6** that is filled by host BIOS **106** to indicate that host BIOS **106** has also filled a CBHS table such as CBHS table **700** shown in FIG. **7** for each HDM of a device.

[0070] Moving from block **915** or from block **930** to block **935**, host BIOS **106** completes its portion of initialization of system **100** to allow subsequent operations to switch to an OS.

According to some examples, an OS such as host OS **102** may implement subsequent operations. [0071] FIG. **10** illustrates an example logic flow **1000**. In some examples, logic flow **1000** may be implemented by an OS for a host computing device such as host OS **102** for host compute device **105** as shown in FIG. **1** or **3**. For these examples, host OS **102** may utilize various ACPI tables filled by a BIOS such as host BIOS **106** as mentioned above for logic flow **900** in order to initialize and bind HDM mapping tables in host OS **102** for subsequent use to multiplex host links **140-1** to **140-N** in order to increase a data bandwidth to access an HDM at one of devices **130-A**, **130-B** or **130-C**. Host OS **102** may implement logic flow **1000** for each HDM of devices **130-A**, **130-B** and **130-C**.

[0072] Logic flow **1000** begins at block **1010** where host OS **102** enters logic flow **1000**. According to some examples, OS **102** may enter logic flow **1000** upon start of an OS kernel (e.g., executed by host CPU **107**).

[0073] Moving from block **1010** to block **1020**, host OS **102** may cause a BootMem_Init to boot an HDM at a device coupled with host computing system **105** via a host link. For example host OS **102** may cause HBM **134-A** at device **130-A** coupled via host link **140-1** to boot.

[0074] Moving from block **1020** to block **1030**, host OS **102** may parse a binding of an HDM of a device based on ACPI tables including information associated with the HBM to get HDM information for the HDM. In some examples, the ACPI tables includes a CBHS table such as CBHS table **700** that includes information to parse the binding of the HDM of the device. For example, a CBHS table for HBM **134-A** at device **130-A** may include information to parse a binding for HBM **134-A** to vHDM-B2A/vHDM-C2A.

[0075] Moving from block **1030** to block **1040**, host OS **102** builds a mapping table for the HBM. According to some examples, host OS **102** may build the mapping table by first filling HDM information for the HDM in a standard HPA table that may be maintained in host system memory **110**. The standard HPA table, for example, may be for use by HDM decoders **126** and/or logic/features of host OS **102** to determine what HPAs are to be mapped to device physical addresses (DPAs) at the device that includes the HDM. Host OS **102** may then get binding address information for vHDMs bound to the HDM and fill the standard HPA table according to the binding address information. The binding address information for each vHDM may be based on information obtained from a binding vHDM structure table such as example binding vHDM structure table **800**.

[0076] Moving from block **1040** to block **1050**, host OS **102** may record or store the mapping table. According to some examples, as described more below, the mapping table may be stored to host system memory **110** for later use in determining how to route a data access to the HDM via multiple host links to increase a data bandwidth for the data access as compared to using a single host link to complete the data access to the HDM.

[0077] Moving from block **1040** to decision block **1060**, host OS **102** determines whether any additional HDMs are to be booted and associated mapping tables to be built and stored. If more HDMs, logic flow **1000** moves to block **1020**. If no more HDMs, logic flow **1000** comes to an end.

[0078] FIG. **11** illustrates an example logic flow **1100**. In some examples, logic flow **1100** may be implemented by an OS for a host computing device such as host OS **102** for host compute device **105** as shown in FIG. **1** or **3**. For these examples, host OS **102** may have completed the building of mapping tables for multiple HDMs of devices coupled to host compute device **105** as described above for logic flow **1000**.

[0079] Logic flow **1100** begins at block **1110** where host OS **102** is to access an address in an HDM of a device with a Length. The access, for example, may be based on a data request from an application from among host application(s) **108** at host compute device **105** (e.g., associated with an AI or HPC workload). According to some examples, the Length may be associated with a data capacity (e.g., 1 GB) and a range of physical memory addresses in the HDM needed to store that data capacity. For these examples, the HDM to be accessed may be HDM **134-A** of device **130**.

[0080] Moving from block **1110** to block **1120**, host OS **102** gets all vHDMs of the HDM to be accessed from the recorded mapping table. In some examples, the recorded mapping table may be stored to host system memory **110**. The vHDMs, for example, may be vHDM-B2A at device **130-B** and vHDM-C2A at device **130-C** that have been bound to HDM **134-A**.

[0081] Moving from block **1120** to decision block **1130**, host OS **102** determines whether to select a multiplexing mode. According to some examples, the multiplexing mode may include routing access requests via host links **140-1** to **140-N** and through devices **130-B** and **130-C** to increase a data bandwidth for accessing the address in HDM **134-A** with a Length. For these examples, the determination may be based on meeting latency and/or data bandwidth requirements associated with system executing or supporting execution of a workload. The requirement may be based on, but not limited to, meeting a service level agreement or a quality of service policy. If a multiplexing mode is selected, logic flow **1100** moves to block **1150**. If a multiplexing mode is not selected, logic flow **1100** moves to block **1140**.

[0082] Moving from decision block **1130** to block **1140**, host OS **102** uses a legacy access of the address in the HDM with full Length. In some examples, this would mean host OS **102** would access an address of HDM **134-A** with full Length via only host link **141-1**.

[0083] Moving from decision block **1130** to block **1150**, host OS **102** splits Length into $N+1$ parts, gets an HDM address based on the mapping table, and gets $N+1$ Sub-addresses for vHDMs based on the mapping table. According to some examples, “N” may represent the number of vHDMs bound to the HDM. For example, N for HDM **134-A** would be 2. An example equation to determine a sub-address for a vHDM may be $\text{Sub_Addr}_m = \text{Addr_in_HDM} + N * (\text{Length} / (N+1))$, where $n=1, \dots, N$ and $m=0, \dots, N-1$.

[0084] Moving from block **1150** to block **1160**, host OS **102** accesses an address range of HDM with $\text{size} = \text{Length} / (N+1)$ and accesses each sub-address with $\text{size} = \text{Length} / (N+1)$. In some examples, host OS **102** accesses HDM **134-A** by sending a first CXL.mem request to device **130-A** via host link **140-1**, a second CXL.mem request to device **130-B** via host link **140-2**, and a third CXL.mem request to device **130-C** via host link **140-N**. The first CXL.mem request to include the determined HDM address for HDM **134-A** and the second and third CXL.mem requests to include the determined sub-addresses for respective vHDMs B2A and C2A.

[0085] The set of logic flows shown in FIGS. **9-11** may be representative of example methodologies for performing novel aspects described in this disclosure. While, for purposes of simplicity of explanation, the one or more methodologies shown herein are shown and described as a series of acts, those skilled in the art will understand and appreciate that the methodologies are not limited by the order of acts. Some acts may, in accordance therewith, occur in a different order and/or concurrently with other acts from that shown and described herein. For example, those skilled in the art will understand and appreciate that a methodology could alternatively be represented as a series of interrelated states or events, such as in a state diagram. Moreover, not all acts illustrated in a methodology may be required for a novel implementation.

[0086] A logic flow may be implemented in software, firmware, and/or hardware. In software and firmware embodiments, a logic flow may be implemented by computer executable instructions stored on at least one non-transitory computer readable medium or machine readable medium, such as an optical, magnetic or semiconductor storage. The embodiments are not limited in this context.

[0087] FIG. **12** illustrates an example process **1200**. According to some examples, elements of system **100** shown in FIGS. **1** and **3** are used to describe process **1200**. The elements of system **100** may include, but are not limited to, host **105**, host links **140**, devices **130** or HSILs **152**. Also, elements of devices **130** are used to describe process **1200** that may include, but are not limited to, coherency agents **132**, MR circuitry **133**, MC circuitry **135** or HDMs **134**. For these examples, process **1200** may be similar to scheme **300** described above and shown in FIG. **3** but with additional focus on actions taken by device **130-B** upon receiving a CXL.mem request message from host **105** (e.g., from host OS **102**) to access HDM **134-A** based on a mapping table built and

described above for logic flows **1000** and **1100** and using a forwarding table such forwarding table **220** described above for FIG. 2

[0088] Process **1200** begins at process 12.1 where host **105** sends a CXL.mem request message to device **130-B**. According to some examples, host OS **102** of host **105** may have determined to multiplex multiple CXL.mem requests via host links **140-1** to **140-N** to access an address in HDM **134-A**. For these examples, the CXL.mem request includes a translated address (e.g., Addr_vHDM_1) in a range of addresses associated with vHDM-B2A that has been bound to HDM **134-A**. Data2 indicates device **130-B** is handling a CXL.mem request for a portion of a Length associated with the entire memory access.

[0089] At process **12.2**, coherency agent **132** receives the CXL.mem request and forwards the translated address in the range of address associated with vHDM-B2A.

[0090] At process **12.3**, MR circuitry **133-B** uses forwarding table **220** maintained in registers **138-B** to translate Addr_vHDM_1 to a sub-address of HDM **134-A**.

[0091] At process **12.4**, MR circuitry **133-B** forwards a memory request for access to the translated sub-address of HDM **134-A** associated with the CXL.mem request to device **130-A** via HSIL **152-1**.

[0092] At process **12.5**, MR circuitry **133-A** of device **130-A** receives the memory request forwarded from device **130-B** and confirms that the translated sub-address of HDM **134-A** is for an address of HDM **134-A** and forwards the memory request to MC circuitry **135-A** for access to HDM **134-A** at the translated sub-address.

[0093] At process **12.6**, MR circuitry **133-A** gets a response back from MC circuitry **135-A**. In some examples, if the memory request is for writing data2 to HDM **134-A**, the response is an indication of whether the data2 was successfully written to the translated sub-address of HDM **134-A**. If the memory request is for reading data2 from HDM **134-A**, the response may include data2 that was read from HDM **134-A**.

[0094] At process **12.7**, MR circuitry **133-A** forwards the response from MC circuitry **135-A** to MR circuitry **133-B** at device **130-B** via HSIL **152-1**. MR circuitry **133-B** may then translate the sub-address of HDM **134-A** back to Addr_vHDM_1 indicated in the CXL.mem request received at device **130-B**.

[0095] At process **12.8**, MR circuitry **133-A** then forwards the response for Addr_vHDM_1 for the CXL.mem request to coherency agent **132-B**.

[0096] At process **12.9**, coherency agent **132-B** sends the response for the CXL.mem request to host **105** via host link **140-2**.

[0097] According to some examples, separate CXL.mem requests including requests to access HDM **134-A** for portions data1 and data3 may be sent to device **130-A** and device **130-C** in parallel with sending the CXL. mem request to access HDM **134-A** for portion data2 as described above for process **12.1** to **12.9**. Device **130-C** will follow a similar process to access its requested data3 portion at HDM **134-A** as was described for device **130-B**. Device **130-A** will directly access HDM **134-A** and provide a response directly back to host **105** without the use of an HSIL.

System Overview

[0098] FIG. **13** is a block diagram of a processing system **1300**, according to an embodiment. Processing system **1300** may be used in a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors **1302** or processor cores **1307**. In one embodiment, the processing system **1300** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices such as within Internet-of-things (IoT) devices with wired or wireless connectivity to a local or wide area network.

[0099] In one embodiment, processing system **1300** can include, couple with, or be integrated within: a server-based gaming platform; a game console, including a game and media console; a mobile gaming console, a handheld game console, or an online game console. In some

embodiments the processing system **1300** is part of a mobile phone, smart phone, tablet computing device or mobile Internet-connected device such as a laptop with low internal storage capacity. Processing system **1300** can also include, couple with, or be integrated within: a wearable device, such as a smart watch wearable device; smart eyewear or clothing enhanced with augmented reality (AR) or virtual reality (VR) features to provide visual, audio or tactile outputs to supplement real world visual, audio or tactile experiences or otherwise provide text, audio, graphics, video, holographic images or video, or tactile feedback; other augmented reality (AR) device; or other virtual reality (VR) device. In some embodiments, the processing system **1300** includes or is part of a television or set top box device. In one embodiment, processing system **1300** can include, couple with, or be integrated within a self-driving vehicle such as a bus, tractor trailer, car, motor or electric power cycle, plane, or glider (or any combination thereof). The self-driving vehicle may use processing system **1300** to process the environment sensed around the vehicle.

[0100] In some embodiments, the one or more processors **1302** each include one or more processor cores **1307** to process instructions which, when executed, perform operations for system or user software. In some embodiments, at least one of the one or more processor cores **1307** is configured to process a specific instruction set **1309**. In some embodiments, instruction set **1309** may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). One or more processor cores **1307** may process a different instruction set **1309**, which may include instructions to facilitate the emulation of other instruction sets. Processor core **1307** may also include other processing devices, such as a Digital Signal Processor (DSP).

[0101] In some embodiments, the processor **1302** includes cache memory **1304**. Depending on the architecture, the processor **1302** can have a single internal cache or multiple levels of internal cache. In some embodiments, the cache memory is shared among various components of the processor **1302**. In some embodiments, the processor **1302** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **1307** using known cache coherency techniques. A register file **1306** can be additionally included in processor **1302** and may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). Some registers may be general-purpose registers, while other registers may be specific to the design of the processor **1302**.

[0102] In some embodiments, one or more processor(s) **1302** are coupled with one or more interface bus(es) **1310** to transmit communication signals such as address, data, or control signals between processor **1302** and other components in the processing system **1300**. The interface bus **1310**, in one embodiment, can be a processor bus, such as a version of the Direct Media Interface (DMI) bus. However, processor busses are not limited to the DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI express), memory busses, or other types of interface busses. In one embodiment the processor(s) **1302** include a memory controller **1316** and a platform controller hub **1330**. The memory controller **1316** facilitates communication between a memory device and other components of the processing system **1300**, while the platform controller hub (PCH) **1330** provides connections to I/O devices via a local I/O bus.

[0103] The memory device **1320** can be a dynamic random-access memory (DRAM) device, a static random-access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In one embodiment the memory device **1320** can operate as system memory for the processing system **1300**, to store data **1322** and instructions **1321** for use when the one or more processors **1302** executes an application or process. The memory controller **1316** also couples with an optional external graphics processor **1318**, which may communicate with the one or more graphics processors **1308** in processors **1302** to perform graphics and media operations. In some embodiments, graphics, media, and or compute operations may be assisted by an accelerator **1312**

which is a coprocessor that can be configured to perform a specialized set of graphics, media, or compute operations. For example, in one embodiment the accelerator **1312** is a matrix multiplication accelerator used to optimize machine learning or compute operations. In one embodiment the accelerator **1312** is a ray-tracing accelerator that can be used to perform ray-tracing operations in concert with the graphics processor **1308**. In one embodiment, an external accelerator **1319** may be used in place of or in concert with the accelerator **1312**.

[0104] In some embodiments a display device **1311** can connect to the processor(s) **1302**. The display device **1311** can be one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In one embodiment the display device **1311** can be a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

[0105] In some embodiments the platform controller hub **1330** enables peripherals to connect to memory device **1320** and processor **1302** via a high-speed I/O bus. The I/O peripherals include, but are not limited to, an audio controller **1346**, a network controller **1334**, a firmware interface **1328**, a wireless transceiver **1326**, touch sensors **1325**, a data storage device **1324** (e.g., non-volatile memory, volatile memory, hard disk drive, flash memory, NAND, 3D NAND, 3D XPoint, etc.). The data storage device **1324** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI express). The touch sensors **1325** can include touch screen sensors, pressure sensors, or fingerprint sensors. The wireless transceiver **1326** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, 5G, or Long-Term Evolution (LTE) transceiver. The firmware interface **1328** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). The network controller **1334** can enable a network connection to a wired network. In some embodiments, a high-performance network controller (not shown) couples with the interface bus **1310**. The audio controller **1346**, in one embodiment, is a multi-channel high-definition audio controller. In one embodiment the processing system **1300** includes an optional legacy I/O controller **1340** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the system. The platform controller hub **1330** can also connect to one or more Universal Serial Bus (USB) controllers **1342** connect input devices, such as keyboard and mouse **1343** combinations, a camera **1344**, or other USB input devices.

[0106] It will be appreciated that the processing system **1300** shown is exemplary and not limiting, as other types of data processing systems that are differently configured may also be used. For example, an instance of the memory controller **1316** and platform controller hub **1330** may be integrated into a discreet external graphics processor, such as the external graphics processor **1318**. In one embodiment the platform controller hub **1330** and/or memory controller **1316** may be external to the one or more processor(s) **1302** and reside in a system chipset that is in communication with the processor(s) **1302**.

[0107] For example, circuit boards (“sleds”) can be used on which components such as CPUs, memory, and other components are placed are designed for increased thermal performance. In some examples, processing components such as the processors are located on a top side of a sled while near memory, such as DIMMs, are located on a bottom side of the sled. As a result of the enhanced airflow provided by this design, the components may operate at higher frequencies and power levels than in typical systems, thereby increasing performance. Furthermore, the sleds are configured to blindly mate with power and data communication cables in a rack, thereby enhancing their ability to be quickly removed, upgraded, reinstalled, and/or replaced. Similarly, individual components located on the sleds, such as processors, accelerators, memory, and data storage drives, are configured to be easily upgraded due to their increased spacing from each other. In the illustrative embodiment, the components additionally include hardware attestation features to prove their authenticity.

[0108] A data center can utilize a single network architecture (“fabric”) that supports multiple other network architectures including Ethernet and Omni-Path. The sleds can be coupled to switches via optical fibers, which provide higher bandwidth and lower latency than typical twisted pair cabling (e.g., Category 5, Category 5e, Category 6, etc.). Due to the high bandwidth, low latency interconnections and network architecture, the data center may, in use, pool resources, such as memory, accelerators (e.g., GPUs, graphics accelerators, FPGAs, ASICs, neural network and/or artificial intelligence accelerators, etc.), and data storage drives that are physically disaggregated, and provide them to compute resources (e.g., processors) on an as needed basis, enabling the compute resources to access the pooled resources as if they were local.

[0109] A power supply or source can provide voltage and/or current to processing system **1300** or any component or system described herein. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based power supply, solar power supply, or fuel cell source.

[0110] FIGS. **14A-14D** illustrate computing systems and graphics processors provided by embodiments described herein. The elements of FIGS. **14A-14D** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0111] FIG. **14A** is a block diagram of an embodiment of a processor **1400** having one or more processor cores **1402A-1402N**, an integrated memory controller **1414**, and an integrated graphics processor **1408**. Processor **1400** can include additional cores up to and including additional core **1402N** represented by the dashed lined boxes. Each of processor cores **1402A-1402N** includes one or more internal cache units **1404A-1404N**. In some embodiments each processor core also has access to one or more shared cache units **1406**. The internal cache units **1404A-1404N** and shared cache units **1406** represent a cache memory hierarchy within the processor **1400**. The cache memory hierarchy may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where the highest level of cache before external memory is classified as the LLC. In some embodiments, cache coherency logic maintains coherency between the various cache units **1406** and **1404A-1404N**.

[0112] In some embodiments, processor **1400** may also include a set of one or more bus controller units **1416** and a system agent core **1410**. The one or more bus controller units **1416** manage a set of peripheral buses, such as one or more PCI or PCI express busses. System agent core **1410** provides management functionality for the various processor components. In some embodiments, system agent core **1410** includes one or more integrated memory controllers **1414** to manage access to various external memory devices (not shown).

[0113] In some embodiments, one or more of the processor cores **1402A-1402N** include support for simultaneous multi-threading. In such embodiment, the system agent core **1410** includes components for coordinating and operating cores **1402A-1402N** during multi-threaded processing. System agent core **1410** may additionally include a power control unit (PCU), which includes logic and components to regulate the power state of processor cores **1402A-1402N** and graphics processor **1408**.

[0114] In some embodiments, processor **1400** additionally includes graphics processor **1408** to execute graphics processing operations. In some embodiments, the graphics processor **1408** couples with the set of shared cache units **1406**, and the system agent core **1410**, including the one or more integrated memory controllers **1414**. In some embodiments, the system agent core **1410** also includes a display controller **1411** to drive graphics processor output to one or more coupled

displays. In some embodiments, display controller **1411** may also be a separate module coupled with the graphics processor via at least one interconnect, or may be integrated within the graphics processor **1408**.

[0115] In some embodiments, a ring-based interconnect **1412** is used to couple the internal components of the processor **1400**. However, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, a mesh interconnect, or other techniques, including techniques well known in the art. In some embodiments, graphics processor **1408** couples with the ring-based interconnect **1412** via an I/O link **1413**.

[0116] The exemplary I/O link **1413** represents at least one of multiple varieties of I/O interconnects, including an on package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module **1418**, such as an eDRAM module or a high-bandwidth memory (HBM) module. In some embodiments, each of the processor cores **1402A-1402N** and graphics processor **1408** can use the embedded memory module **1418** as a shared Last Level Cache.

[0117] In some embodiments, processor cores **1402A-1402N** are homogenous cores executing the same instruction set architecture. In another embodiment, processor cores **1402A-1402N** are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores **1402A-1402N** execute a first instruction set, while at least one of the other cores executes a subset of the first instruction set or a different instruction set. In one embodiment, processor cores **1402A-1402N** are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In one embodiment, processor cores **1402A-1402N** are heterogeneous in terms of computational capability. Additionally, processor **1400** can be implemented on one or more chips or as an SoC integrated circuit having the illustrated components, in addition to other components.

[0118] FIG. 2B is a block diagram of hardware logic of a graphics processor core block **1419**, according to some embodiments described herein. In some embodiments, elements of FIG. 14B having the same reference numbers (or names) as the elements of any other figure herein may operate or function in a manner similar to that described elsewhere herein. The graphics processor core block **1419** is exemplary of one partition of a graphics processor. The graphics processor core block **1419** can be included within the integrated graphics processor **1408** of FIG. 14A or a discrete graphics processor, parallel processor, and/or compute accelerator. A graphics processor as described herein may include multiple graphics core blocks based on target power and performance envelopes. Each graphics processor core block **1419** can include a function block **1430** coupled with multiple graphics cores **1421A-1421F** that include modular blocks of fixed function logic and general-purpose programmable logic. The graphics processor core block **1419** also includes shared/cache memory **1436** that is accessible by all graphics cores **1421A-1421F**, rasterizer logic **1437**, and additional fixed function logic **1438**.

[0119] In some embodiments, the function block **1430** includes a geometry/fixed function pipeline **1431** that can be shared by all graphics cores in the graphics processor core block **1419**. In various embodiments, the geometry/fixed function pipeline **1431** includes a 3D geometry pipeline a video front-end unit, a thread spawner and global thread dispatcher, and a unified return buffer manager, which manages unified return buffers. In one embodiment the function block **1430** also includes a graphics SoC interface **1432**, a graphics microcontroller **1433**, and a media pipeline **1434**. The graphics SoC interface **1432** provides an interface between the graphics processor core block **1419** and other core blocks within a graphics processor or compute accelerator SoC. The graphics microcontroller **1433** is a programmable sub-processor that is configurable to manage various functions of the graphics processor core block **1419**, including thread dispatch, scheduling, and pre-emption. The media pipeline **1434** includes logic to facilitate the decoding, encoding, pre-processing, and/or post-processing of multimedia data, including image and video data. The media pipeline **1434** implement media operations via requests to compute or sampling logic within the

graphics cores **1421-1421F**. One or more pixel backends **1435** can also be included within the function block **1430**. The pixel backends **1435** include a cache memory to store pixel color values and can perform blend operations and lossless color compression of rendered pixel data.

[0120] In one embodiment the graphics SoC interface **1432** enables the graphics processor core block **1419** to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC or a system host CPU that is coupled with the SoC via a peripheral interface. The graphics SoC interface **1432** also enables communication with off-chip memory hierarchy elements such as a shared last level cache memory, system RAM, and/or embedded on-chip or on-package DRAM. The SoC interface **1432** can also enable communication with fixed function devices within the SoC, such as camera imaging pipelines, and enables the use of and/or implements global memory atomics that may be shared between the graphics processor core block **1419** and CPUs within the SoC. The graphics SoC interface **1432** can also implement power management controls for the graphics processor core block **1419** and enable an interface between a clock domain of the graphics processor core block **1419** and other clock domains within the SoC. In one embodiment the graphics SoC interface **1432** enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. The commands and instructions can be dispatched to the media pipeline **1434** when media operations are to be performed, the geometry and fixed function pipeline **1431** when graphics processing operations are to be performed. When compute operations are to be performed, compute dispatch logic can dispatch the commands to the graphics cores **1421A-1421F**, bypassing the geometry and media pipelines.

[0121] The graphics microcontroller **1433** can be configured to perform various scheduling and management tasks for the graphics processor core block **1419**. In one embodiment the graphics microcontroller **1433** can perform graphics and/or compute workload scheduling on the various vector engines **1422A-1422F**, **1424A-1424F** and matrix engines **1423A-1423F**, **1425A-1425F** within the graphics cores **1421A-1421F**. In this scheduling model, host software executing on a CPU core of an SoC including the graphics processor core block **1419** can submit workloads one of multiple graphic processor doorbells, which invokes a scheduling operation on the appropriate graphics engine. Scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In one embodiment the graphics microcontroller **1433** can also facilitate low-power or idle states for the graphics processor core block **1419**, providing the graphics processor core block **1419** with the ability to save and restore registers within the graphics processor core block **1419** across low-power state transitions independently from the operating system and/or graphics driver software on the system.

[0122] The graphics processor core block **1419** may have greater than or fewer than the illustrated graphics cores **1421A-1421F**, up to N modular graphics cores. For each set of N graphics cores, the graphics processor core block **1419** can also include shared/cache memory **1436**, which can be configured as shared memory or cache memory, rasterizer logic **1437**, and additional fixed function logic **1438** to accelerate various graphics and compute processing operations.

[0123] Within each graphics cores **1421A-1421F** is set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. The graphics cores **1421A-1421F** include multiple vector engines **1422A-1422F**, **1424A-1424F**, matrix acceleration units **1423A-1423F**, **1425A-1425D**, cache/shared local memory (SLM), a sampler **1426A-1426F**, and a ray tracing unit **1427A-1427F**.

[0124] The vector engines **1422A-1422F**, **1424A-1424F** are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation, including graphics, media, or compute/GPGPU programs.

The vector engines **1422A-1422F**, **1424A-1424F** can operate at variable vector widths using SIMD, SIMT, or SIMT+SIMD execution modes. The matrix acceleration units **1423A-1423F**, **1425A-1425D** include matrix-matrix and matrix-vector acceleration logic that improves performance on matrix operations, particularly low and mixed precision (e.g., INT8, FP16, BF16) matrix operations used for machine learning. In one embodiment, each of the matrix acceleration units **1423A-1423F**, **1425A-1425D** includes one or more systolic arrays of processing elements that can perform concurrent matrix multiply or dot product operations on matrix elements.

[0125] The sampler **1426A-1426F** can read media or texture data into memory and can sample data differently based on a configured sampler state and the texture/media format that is being read.

Threads executing on the vector engines **1422A-1422F**, **1424A-1424F** or matrix acceleration units **1423A-1423F**, **1425A-1425D** can make use of the cache/SLM **1428A-1428F** within each execution core. The cache/SLM **1428A-1428F** can be configured as cache memory or as a pool of shared memory that is local to each of the respective graphics cores **1421A-1421F**. The ray tracing units **1427A-1427F** within the graphics cores **1421A-1421F** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. In one embodiment the ray tracing units **1427A-1427F** include circuitry for performing depth testing and culling (e.g., using a depth buffer or similar arrangement). In one implementation, the ray tracing units **1427A-1427F** perform traversal and intersection operations in concert with image denoising, at least a portion of which may be performed using an associated matrix acceleration unit **1423A-1423F**, **1425A-1425D**.

[0126] FIG. **14C** illustrates a graphics processing unit (GPU) **1439** that includes dedicated sets of graphics processing resources arranged into multi-core groups **1440A-1440N**. The details of multi-core group **1440A** are illustrated. Multi-core groups **1440B-1440N** may be equipped with the same or similar sets of graphics processing resources.

[0127] As illustrated, a multi-core group **1440A** may include a set of graphics cores **1443**, a set of tensor cores **1444**, and a set of ray tracing cores **1445**. A scheduler/dispatcher **1441** schedules and dispatches the graphics threads for execution on the various cores **1443**, **1444**, **1445**. In one embodiment the tensor cores **1444** are sparse tensor cores with hardware to enable multiplication operations having a zero-value input to be bypassed. The graphics cores **1443** of the GPU **1439** of FIG. **14C** differ in hierarchical abstraction level relative to the graphics cores **1421A-1421F** of FIG. **14B**, which are analogous to the multi-core groups **1440A-1440N** of FIG. **14C**. The graphics cores **1443**, tensor cores **1444**, and ray tracing cores **1445** of FIG. **14C** are analogous to, respectively, the vector engines **1422A-1422F**, **1424A-1424F**, matrix engines **1423A-1423F**, **1425A-1425F**, and ray tracing units **1427A-1427F** of FIG. **14B**.

[0128] A set of register files **1442** can store operand values used by the cores **1443**, **1444**, **1445** when executing the graphics threads. These may include, for example, integer registers for storing integer values, floating point registers for storing floating point values, vector registers for storing packed data elements (integer and/or floating-point data elements) and tile registers for storing tensor/matrix values. In one embodiment, the tile registers are implemented as combined sets of vector registers.

[0129] One or more combined level 1 (L1) caches and shared memory units **1447** store graphics data such as texture data, vertex data, pixel data, ray data, bounding volume data, etc., locally within each multi-core group **1440A**. One or more texture units **1447** can also be used to perform texturing operations, such as texture mapping and sampling. A Level 2 (L2) cache **1453** shared by all or a subset of the multi-core groups **1440A-1440N** stores graphics data and/or instructions for multiple concurrent graphics threads. As illustrated, the L2 cache **1453** may be shared across a plurality of multi-core groups **1440A-1440N**. One or more memory controllers **1448** couple the GPU **1439** to a memory **1449** which may be a system memory (e.g., DRAM) and/or a dedicated graphics memory (e.g., GDDR6 memory).

[0130] Input/output (I/O) circuitry **1450** couples the GPU **1439** to one or more I/O devices **1452** such as digital signal processors (DSPs), network controllers, or user input devices. An on-chip interconnect may be used to couple the I/O devices **1452** to the GPU **1439** and memory **1449**. One or more I/O memory management units (IOMMUs) **1451** of the I/O circuitry **1450** couple the I/O devices **1452** directly to the memory **1449**. In one embodiment, the IOMMU **1451** manages multiple sets of page tables to map virtual addresses to physical addresses in memory **1449**. In this embodiment, the I/O devices **1452**, CPU(s) **1446**, and GPU **1439** may share the same virtual address space.

[0131] In one implementation, the IOMMU **1451** supports virtualization. In this case, it may manage a first set of page tables to map guest/graphics virtual addresses to guest/graphics physical addresses and a second set of page tables to map the guest/graphics physical addresses to system/host physical addresses (e.g., within memory **1449**). The base addresses of each of the first and second sets of page tables may be stored in control registers and swapped out on a context switch (e.g., so that the new context is provided with access to the relevant set of page tables). While not illustrated in FIG. **14C**, each of the cores **1443**, **1444**, **1445** and/or multi-core groups **1440A-1440N** may include translation lookaside buffers (TLBs) to cache guest virtual to guest physical translations, guest physical to host physical translations, and guest virtual to host physical translations.

[0132] In one embodiment, the CPUs **1446**, GPU **1439**, and I/O devices **1452** are integrated on a single semiconductor chip and/or chip package. The memory **1449** may be integrated on the same chip or may be coupled to the memory controllers **1448** via an off-chip interface. In one implementation, the memory **1449** comprises GDDR6 memory which shares the same virtual address space as other physical system-level memories, although the underlying principles of the embodiments described herein are not limited to this specific implementation.

[0133] In one embodiment, the tensor cores **1444** include a plurality of functional units specifically designed to perform matrix operations, which are the fundamental compute operation used to perform deep learning operations. For example, simultaneous matrix multiplication operations may be used for neural network training and inferencing. The tensor cores **1444** may perform matrix processing using a variety of operand precisions including single precision floating-point (e.g., 32 bits), half-precision floating point (e.g., 16 bits), integer words (16 bits), bytes (8 bits), and half-bytes (4 bits). In one embodiment, a neural network implementation extracts features of each rendered scene, potentially combining details from multiple frames, to construct a high-quality final image.

[0134] In deep learning implementations, parallel matrix multiplication work may be scheduled for execution on the tensor cores **1444**. The training of neural networks, in particular, requires a significant number of matrix dot product operations. In order to process an inner-product formulation of an $N \times N \times N$ matrix multiply, the tensor cores **1444** may include at least N dot-product processing elements. Before the matrix multiply begins, one entire matrix is loaded into tile registers and at least one column of a second matrix is loaded each cycle for N cycles. Each cycle, there are N dot products that are processed.

[0135] Matrix elements may be stored at different precisions depending on the particular implementation, including 16-bit words, 8-bit bytes (e.g., INT8) and 4-bit half-bytes (e.g., INT4). Different precision modes may be specified for the tensor cores **1444** to ensure that the most efficient precision is used for different workloads (e.g., such as inferencing workloads which can tolerate quantization to bytes and half-bytes).

[0136] In one embodiment, the ray tracing cores **1445** accelerate ray tracing operations for both real-time ray tracing and non-real-time ray tracing implementations. In particular, the ray tracing cores **1445** include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. The ray tracing cores **1445** may also include circuitry for performing

depth testing and culling (e.g., using a Z buffer or similar arrangement). In one implementation, the ray tracing cores **1445** perform traversal and intersection operations in concert with the image denoising techniques described herein, at least a portion of which may be executed on the tensor cores **1444**. For example, in one embodiment, the tensor cores **1444** implement a deep learning neural network to perform denoising of frames generated by the ray tracing cores **1445**. However, the CPU(s) **1446**, graphics cores **1443**, and/or ray tracing cores **1445** may also implement all or a portion of the denoising and/or deep learning algorithms.

[0137] In addition, as described above, a distributed approach to denoising may be employed in which the GPU **1439** is in a computing device coupled to other computing devices over a network or high-speed interconnect. In this embodiment, the interconnected computing devices share neural network learning/training data to improve the speed with which the overall system learns to perform denoising for different types of image frames and/or different graphics applications.

[0138] In one embodiment, the ray tracing cores **1445** process all BVH traversal and ray-primitive intersections, saving the graphics cores **1443** from being overloaded with thousands of instructions per ray. In one embodiment, each ray tracing core **1445** includes a first set of specialized circuitry for performing bounding box tests (e.g., for traversal operations) and a second set of specialized circuitry for performing the ray-triangle intersection tests (e.g., intersecting rays which have been traversed). Thus, in one embodiment, the multi-core group **1440A** can simply launch a ray probe, and the ray tracing cores **1445** independently perform ray traversal and intersection and return hit data (e.g., a hit, no hit, multiple hits, etc.) to the thread context. The other cores **1443**, **1444** are freed to perform other graphics or compute work while the ray tracing cores **1445** perform the traversal and intersection operations.

[0139] In one embodiment, each ray tracing core **1445** includes a traversal unit to perform BVH testing operations and an intersection unit which performs ray-primitive intersection tests. The intersection unit generates a “hit”, “no hit”, or “multiple hit” response, which it provides to the appropriate thread. During the traversal and intersection operations, the execution resources of the other cores (e.g., graphics cores **1443** and tensor cores **1444**) are freed to perform other forms of graphics work.

[0140] In one particular embodiment described below, a hybrid rasterization/ray tracing approach is used in which work is distributed between the graphics cores **1443** and ray tracing cores **1445**.

[0141] In one embodiment, the ray tracing cores **1445** (and/or other cores **1443**, **1444**) include hardware support for a ray tracing instruction set such as Microsoft's DirectX Ray Tracing (DXR) which includes a DispatchRays command, as well as ray-generation, closest-hit, any-hit, and miss shaders, which enable the assignment of unique sets of shaders and textures for each object. Another ray tracing platform which may be supported by the ray tracing cores **1445**, graphics cores **1443** and tensor cores **1444** is Vulkan 1.1.85. Note, however, that the underlying principles of the embodiments described herein are not limited to any particular ray tracing ISA.

[0142] In general, the various cores **1445**, **1444**, **1443** may support a ray tracing instruction set that includes instructions/functions for ray generation, closest hit, any hit, ray-primitive intersection, per-primitive and hierarchical bounding box construction, miss, visit, and exceptions. More specifically, one embodiment includes ray tracing instructions to perform the following functions:

[0143] Ray Generation—Ray generation instructions may be executed for each pixel, sample, or other user-defined work assignment. [0144] Closest Hit—A closest hit instruction may be executed to locate the closest intersection point of a ray with primitives within a scene. [0145] Any Hit—An any hit instruction identifies multiple intersections between a ray and primitives within a scene, potentially to identify a new closest intersection point. [0146] Intersection—An intersection instruction performs a ray-primitive intersection test and outputs a result. [0147] Per—primitive Bounding box Construction—This instruction builds a bounding box around a given primitive or group of primitives (e.g., when building a new BVH or other acceleration data structure). [0148] Miss—Indicates that a ray misses all geometry within a scene, or specified region of a scene.

[0149] Visit—Indicates the child volumes a ray will traverse. [0150] Exceptions—Includes various types of exception handlers (e.g., invoked for various error conditions).

[0151] In one embodiment the ray tracing cores **1445** may be adapted to accelerate general-purpose compute operations that can be accelerated using computational techniques that are analogous to ray intersection tests. A compute framework can be provided that enables shader programs to be compiled into low level instructions and/or primitives that perform general-purpose compute operations via the ray tracing cores. Exemplary computational problems that can benefit from compute operations performed on the ray tracing cores **1445** include computations involving beam, wave, ray, or particle propagation within a coordinate space. Interactions associated with that propagation can be computed relative to a geometry or mesh within the coordinate space. For example, computations associated with electromagnetic signal propagation through an environment can be accelerated via the use of instructions or primitives that are executed via the ray tracing cores. Diffraction and reflection of the signals by objects in the environment can be computed as direct ray-tracing analogies.

[0152] Ray tracing cores **1445** can also be used to perform computations that are not directly analogous to ray tracing. For example, mesh projection, mesh refinement, and volume sampling computations can be accelerated using the ray tracing cores **1445**. Generic coordinate space calculations, such as nearest neighbor calculations can also be performed. For example, the set of points near a given point can be discovered by defining a bounding box in the coordinate space around the point. BVH and ray probe logic within the ray tracing cores **1445** can then be used to determine the set of point intersections within the bounding box. The intersections constitute the origin point and the nearest neighbors to that origin point. Computations that are performed using the ray tracing cores **1445** can be performed in parallel with computations performed on the graphics cores **1443** and tensor cores **1444**. A shader compiler can be configured to compile a compute shader or other general-purpose graphics processing program into low level primitives that can be parallelized across the graphics cores **1443**, tensor cores **1444**, and ray tracing cores **1445**.

[0153] FIG. **14D** is a block diagram of general-purpose graphics processing unit (GPGPU) **1470** that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein. The GPGPU **1470** can interconnect with host processors (e.g., one or more CPU(s) **1446**) and memory **1471**, **1472** via one or more system and/or memory busses. In one embodiment the memory **1471** is system memory that may be shared with the one or more CPU(s) **1446**, while memory **1472** is device memory that is dedicated to the GPGPU **1470**. In one embodiment, components within the GPGPU **1470** and memory **1472** may be mapped into memory addresses that are accessible to the one or more CPU(s) **1446**. Access to memory **1471** and **1472** may be facilitated via a memory controller **1468**. In one embodiment the memory controller **1468** includes an internal direct memory access (DMA) controller **1469** or can include logic to perform operations that would otherwise be performed by a DMA controller.

[0154] The GPGPU **1470** includes multiple cache memories, including an L2 cache **1453**, L1 cache **1454**, an instruction cache **1455**, and shared memory **1456**, at least a portion of which may also be partitioned as a cache memory. The GPGPU **1470** also includes multiple compute units **1460A-1460N**, which represent a hierarchical abstraction level analogous to the graphics cores **1421A-1421F** of FIG. **14B** and the multi-core groups **1440A-1440N** of FIG. **14C**. Each compute unit **1460A-1460N** includes a set of vector registers **1461**, scalar registers **1462**, vector logic units **1463**, and scalar logic units **1464**. The compute units **1460A-1460N** can also include local shared memory **1465** and a program counter **1466**. The compute units **1460A-1460N** can couple with a constant cache **1467**, which can be used to store constant data, which is data that will not change during the run of kernel or shader program that executes on the GPGPU **1470**. In one embodiment the constant cache **1467** is a scalar data cache and cached data can be fetched directly into the scalar registers **1462**.

[0155] During operation, the one or more CPU(s) **1446** can write commands into registers or memory in the GPGPU **1470** that has been mapped into an accessible address space. The command processors **1457** can read the commands from registers or memory and determine how those commands will be processed within the GPGPU **1470**. A thread dispatcher **1458** can then be used to dispatch threads to the compute units **1460A-1460N** to perform those commands. Each compute unit **1460A-1460N** can execute threads independently of the other compute units. Additionally, each compute unit **1460A-1460N** can be independently configured for conditional computation and can conditionally output the results of computation to memory. The command processors **1457** can interrupt the one or more CPU(s) **1446** when the submitted commands are complete.

[0156] FIGS. **15A-15C** illustrate block diagrams of additional graphics processor and compute accelerator architectures provided by embodiments described herein. The elements of FIGS. **15A-15C** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0157] FIG. **15A** is a block diagram of a graphics processor **1500**, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces. In some embodiments, the graphics processor communicates via a memory mapped I/O interface to registers on the graphics processor and with commands placed into the processor memory. In some embodiments, graphics processor **1500** includes a memory interface **1514** to access memory. Memory interface **1514** can be an interface to local memory, one or more internal caches, one or more shared external caches, and/or to system memory.

[0158] In some embodiments, graphics processor **1500** also includes a display controller **1502** to drive display output data to a display device **1518**. Display controller **1502** includes hardware for one or more overlay planes for the display and composition of multiple layers of video or user interface elements. The display device **1518** can be an internal or external display device. In one embodiment the display device **1518** is a head mounted display device, such as a virtual reality (VR) display device or an augmented reality (AR) display device. In some embodiments, graphics processor **1500** includes a video codec engine **1506** to encode, decode, or transcode media to, from, or between one or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, H.265/HEVC, Alliance for Open Media (AOMedia) VP8, VP9, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

[0159] In some embodiments, graphics processor **1500** includes a block image transfer (BLIT) engine **1504** to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in one embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE) **1510**. In some embodiments, GPE **1510** is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

[0160] In some embodiments, GPE **1510** includes a 3D pipeline **1512** for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). The 3D pipeline **1512** includes programmable and fixed function elements that perform various tasks within the element and/or spawn execution threads to a 3D/Media subsystem **1515**. While 3D pipeline **1512** can be used to perform media operations, an embodiment of GPE **1510** also includes a media pipeline **1516** that is specifically used to perform media operations, such as video post-processing and image enhancement.

[0161] In some embodiments, media pipeline **1516** includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration,

video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine **1506**. In some embodiments, media pipeline **1516** additionally includes a thread spawning unit to spawn threads for execution on 3D/Media subsystem **1515**. The spawned threads perform computations for the media operations on one or more graphics cores included in 3D/Media subsystem **1515**.

[0162] In some embodiments, 3D/Media subsystem **1515** includes logic for executing threads spawned by 3D pipeline **1512** and media pipeline **1516**. In one embodiment, the pipelines send thread execution requests to 3D/Media subsystem **1515**, which includes thread dispatch logic for arbitrating and dispatching the various requests to available thread execution resources. The execution resources include an array of graphics cores to process the 3D and media threads. In some embodiments, 3D/Media subsystem **1515** includes one or more internal caches for thread instructions and data. In some embodiments, the subsystem also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

[0163] FIG. **15B** illustrates a graphics processor **1520** having a tiled architecture, according to embodiments described herein. In one embodiment the graphics processor **1520** includes a graphics processing engine cluster **1522** having multiple instances of the graphics processing engine **1510** of FIG. **15A** within a graphics engine tile **1510A-1510D**. Each graphics engine tile **1510A-1510D** can be interconnected via a set of tile interconnects **1523A-1523F**. Each graphics engine tile **1510A-1510D** can also be connected to a memory module or memory device **1526A-1526D** via memory interconnects **1525A-1525D**. The memory devices **1526A-1526D** can use any graphics memory technology. For example, the memory devices **1526A-1526D** may be graphics double data rate (GDDR) memory. The memory devices **1526A-1526D**, in one embodiment, are HBM modules that can be on-die with their respective graphics engine tile **1510A-1510D**. In one embodiment the memory devices **1526A-1526D** are stacked memory devices that can be stacked on top of their respective graphics engine tile **1510A-1510D**. In one embodiment, each graphics engine tile **1510A-1510D** and associated memory **1526A-1526D** reside on separate chiplets, which are bonded to a base die or base substrate, as described on further detail in FIGS. **23B-23D**.

[0164] The graphics processor **1520** may be configured with a non-uniform memory access (NUMA) system in which memory devices **1526A-1526D** are coupled with associated graphics engine tiles **1510A-1510D**. A given memory device may be accessed by graphics engine tiles other than the tile to which it is directly connected. However, access latency to the memory devices **1526A-1526D** may be lowest when accessing a local tile. In one embodiment, a cache coherent NUMA (ccNUMA) system is enabled that uses the tile interconnects **1523A-1523F** to enable communication between cache controllers within the graphics engine tiles **1510A-1510D** to maintain a consistent memory image when more than one cache stores the same memory location.

[0165] The graphics processing engine cluster **1522** can connect with an on-chip or on-package fabric interconnect **1524**. In one embodiment the fabric interconnect **1524** includes a network processor, network on a chip (NoC), or another switching processor to enable the fabric interconnect **1524** to act as a packet switched fabric interconnect that switches data packets between components of the graphics processor **1520**. The fabric interconnect **1524** can enable communication between graphics engine tiles **1510A-1510D** and components such as the video codec engine **1506** and one or more copy engines **1504**. The copy engines **1504** can be used to move data out of, into, and between the memory devices **1526A-1526D** and memory that is external to the graphics processor **1520** (e.g., system memory). The fabric interconnect **1524** can also couple with one or more of the tile interconnects **1523A-1523F** to facilitate or enhance the interconnection between the graphics engine tiles **1510A-1510D**. The fabric interconnect **1524** is also configurable to interconnect multiple instances of the graphics processor **1520** (e.g., via the host interface **1528**), enabling tile-to-tile communication between graphics engine tiles **1510A-1510D** of multiple GPUs. In one embodiment, the graphics engine tiles **1510A-1510D** of multiple GPUs can be presented to a host system as a single logical device.

[0166] The graphics processor **1520** may optionally include a display controller **1502** to enable a connection with the display device **1518**. The graphics processor may also be configured as a graphics or compute accelerator. In the accelerator configuration, the display controller **1502** and display device **1518** may be omitted.

[0167] The graphics processor **1520** can connect to a host system via a host interface **1528**. The host interface **1528** can enable communication between the graphics processor **1520**, system memory, and/or other system components. The host interface **1528** can be, for example a PCI express bus or another type of host system interface. For example, the host interface **1528** may be an NVLink or NVSwitch interface. The host interface **1528** and fabric interconnect **1524** can cooperate to enable multiple instances of the graphics processor **1520** to act as single logical device. Cooperation between the host interface **1528** and fabric interconnect **1524** can also enable the individual graphics engine tiles **1510A-1510D** to be presented to the host system as distinct logical graphics devices.

[0168] FIG. **15C** illustrates a compute accelerator **1530**, according to embodiments described herein. The compute accelerator **1530** can include architectural similarities with the graphics processor **1520** of FIG. **15B** and is optimized for compute acceleration. A compute engine cluster **1532** can include a set of compute engine tiles **1540A-1540D** that include execution logic that is optimized for parallel or vector-based general-purpose compute operations. In some embodiments, the compute engine tiles **1540A-1540D** do not include fixed function graphics processing logic, although in one embodiment one or more of the compute engine tiles **1540A-1540D** can include logic to perform media acceleration. The compute engine tiles **1540A-1540D** can connect to memory **1526A-1526D** via memory interconnects **1525A-1525D**. The memory **1526A-1526D** and memory interconnects **1525A-1525D** may be similar technology as in graphics processor **1520** or can be different. The graphics compute engine tiles **1540A-1540D** can also be interconnected via a set of tile interconnects **1523A-1523F** and may be connected with and/or interconnected by a fabric interconnect **1524**. Cross-tile communications can be facilitated via the fabric interconnect **1524**. The fabric interconnect **1524** (e.g., via the host interface **1528**) can also facilitate communication between compute engine tiles **1540A-1540D** of multiple instances of the compute accelerator **1530**. In one embodiment the compute accelerator **1530** includes a large L3 cache **1536** that can be configured as a device-wide cache. The compute accelerator **1530** can also connect to a host processor and memory via a host interface **1528** in a similar manner as the graphics processor **1520** of FIG. **15B**.

[0169] The compute accelerator **1530** can also include an integrated network interface **1542**. In one embodiment the network interface **1542** includes a network processor and controller logic that enables the compute engine cluster **1532** to communicate over a physical layer interconnect **1544** without requiring data to traverse memory of a host system. In one embodiment, one of the compute engine tiles **1540A-1540D** is replaced by network processor logic and data to be transmitted or received via the physical layer interconnect **1544** may be transmitted directly to or from memory **1526A-1526D**. Multiple instances of the compute accelerator **1530** may be joined via the physical layer interconnect **1544** into a single logical device. Alternatively, the various compute engine tiles **1540A-1540D** may be presented as distinct network accessible compute accelerator devices.

Graphics Processing Engine

[0170] FIG. **16** is a block diagram of a graphics processing engine **1610** of a graphics processor in accordance with some embodiments. In one embodiment, the graphics processing engine (GPE) **1610** is a version of the GPE **1510** shown in FIG. **15A** and may also represent a graphics engine tile **1510A-1510D** of FIG. **15B**. Elements of FIG. **16** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. For example, the 3D pipeline **1512** and media pipeline **1516** of FIG. **15A** are illustrated. The media pipeline **1516** is optional in some

embodiments of the GPE **1610** and may not be explicitly included within the GPE **1610**. For example and in at least one embodiment, a separate media and/or image processor is coupled to the GPE **1610**.

[0171] In some embodiments, GPE **1610** couples with or includes a command streamer **1603**, which provides a command stream to the 3D pipeline **1512** and/or media pipelines **1516**. Alternatively or additionally, the command streamer **1603** may be directly coupled to a unified return buffer **1618**. The unified return buffer **1618** may be communicatively coupled to a graphics core cluster **1614**. In some embodiments, command streamer **1603** is coupled with memory, which can be system memory, or one or more of internal cache memory and shared cache memory. In some embodiments, command streamer **1603** receives commands from the memory and sends the commands to 3D pipeline **1512** and/or media pipeline **1516**. The commands are directives fetched from a ring buffer, which stores commands for the 3D pipeline **1512** and media pipeline **1516**. In one embodiment, the ring buffer can additionally include batch command buffers storing batches of multiple commands. The commands for the 3D pipeline **1512** can also include references to data stored in memory, such as but not limited to vertex and geometry data for the 3D pipeline **1512** and/or image data and memory objects for the media pipeline **1516**. The 3D pipeline **1512** and media pipeline **1516** process the commands and data by performing operations via logic within the respective pipelines or by dispatching one or more execution threads to a graphics core cluster **1614**. In one embodiment the graphics core cluster **1614** include one or more blocks of graphics cores (e.g., graphics core block **1615A**, graphics core block **1615B**), each block including one or more graphics cores. Each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic, such as matrix or AI acceleration logic.

[0172] In various embodiments the 3D pipeline **1512** can include fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel shaders, fragment shaders, compute shaders, or other shader and/or GPGPU programs, by processing the instructions and dispatching execution threads to the graphics core cluster **1614**. The graphics core cluster **1614** provides a unified block of execution resources for use in processing these shader programs. Multi-purpose execution logic within the graphics core blocks **1615A-1615B** of the graphics core cluster **1614** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

[0173] In some embodiments, the graphics core cluster **1614** includes execution logic to perform media functions, such as video and/or image processing. In one embodiment, the graphics cores include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations. The general-purpose logic can perform processing operations in parallel or in conjunction with general-purpose logic within the processor core(s) **1307** of FIG. **13** or core **1402A-1402N** as in FIG. **14A**.

[0174] Output data generated by threads executing on the graphics core cluster **1614** can output data to memory in a unified return buffer (URB) **1618**. The URB **1618** can store data for multiple threads. In some embodiments the URB **1618** may be used to send data between different threads executing on the graphics core cluster **1614**. In some embodiments the URB **1618** may additionally be used for synchronization between threads on the graphics core array and fixed function logic within the shared function logic **1620**.

[0175] In some embodiments, graphics core cluster **1614** is scalable, such that the cluster includes a variable number of graphics cores, each having a variable number of graphics cores based on the target power and performance level of GPE **1610**. In one embodiment the execution resources are dynamically scalable, such that execution resources may be enabled or disabled as needed.

[0176] The graphics core cluster **1614** couples with shared function logic **1620** that includes multiple resources that are shared between the graphics cores in the graphics core array. The shared

functions within the shared function logic **1620** are hardware logic units that provide specialized supplemental functionality to the graphics core cluster **1614**. In various embodiments, shared function logic **1620** may include, but is not limited to sampler **1621**, math **1622**, and inter-thread communication (ITC) **1623** logic. Additionally, some embodiments implement one or more cache(s) **1625** within the shared function logic **1620**. The shared function logic **1620** can implement the same or similar functionality as the additional fixed function logic **11638** of FIG. **116B**.

[0177] A shared function is implemented at least in a case where the demand for a given specialized function is insufficient for inclusion within the graphics core cluster **1614**. Instead, a single instantiation of that specialized function is implemented as a stand-alone entity in the shared function logic **1620** and shared among the execution resources within the graphics core cluster **1614**. The precise set of functions that are shared between the graphics core cluster **1614** and included within the graphics core cluster **1614** varies across embodiments. In some embodiments, specific shared functions within the shared function logic **1620** that are used extensively by the graphics core cluster **1614** may be included within shared function logic **1616** within the graphics core cluster **1614**. In various embodiments, the shared function logic **1616** within the graphics core cluster **1614** can include some or all logic within the shared function logic **1620**. In one embodiment, all logic elements within the shared function logic **1620** may be duplicated within the shared function logic **1616** of the graphics core cluster **1614**. In one embodiment the shared function logic **1620** is excluded in favor of the shared function logic **1616** within the graphics core cluster **1614**.

Graphics Processing Resources

[0178] FIG. **17A-17C** illustrate execution logic including an array of processing elements employed in a graphics processor, according to embodiments described herein. FIG. **17A** illustrates graphics core cluster, according to an embodiment. FIG. **17B** illustrates a vector engine of a graphics core, according to an embodiment. FIG. **17C** illustrates a matrix engine of a graphics core, according to an embodiment. Elements of FIG. **17A-17C** having the same reference numbers as the elements of any other figure herein may operate or function in any manner similar to that described elsewhere herein, but are not limited as such. For example, the elements of FIG. **17A-17C** can be considered in the context of the graphics processor core block **1419** of FIG. **14B**, and/or the graphics core blocks **1615A-1615B** of FIG. **16**. In one embodiment, the elements of FIG. **17A-17C** have similar functionality to equivalent components of the graphics processor **1408** of FIG. **14A**, the GPU **1439** of FIG. **14C** or the GPGPU **1470** of FIG. **14D**.

[0179] As shown in FIG. **17A**, in one embodiment the graphics core cluster **1614** includes a graphics core block **1615**, which may be graphics core block **1615A** or graphics core block **1615B** of FIG. **16**. The graphics core block **1615** can include any number of graphics cores (e.g., graphics core **1715A**, graphics core **1715B**, through graphics core **1715N**). Multiple instances of the graphics core block **1615** may be included. In one embodiment the elements of the graphics cores **1715A-1715N** have similar or equivalent functionality as the elements of the graphics cores **1421A-1421F** of FIG. **14B**. In such embodiment, the graphics cores **1715A-1715N** each include circuitry including but not limited to vector engines **1702A-1702N**, matrix engines **1703A-1703N**, memory load/store units **1704A-1704N**, instruction caches **1705A-1705N**, data caches/shared local memory **1706A-1706N**, ray tracing units **1708A-1708N**, samplers **1710A-1710N**. The circuitry of the graphics cores **1715A-1715N** can additionally include fixed function logic **1712A-1712N**. The number of vector engines **1702A-1702N** and matrix engines **1703A-1703N** within the graphics cores **1715A-1715N** of a design can vary based on the workload, performance, and power targets for the design.

[0180] With reference to graphics core **1715A**, the vector engine **1702A** and matrix engine **1703A** are configurable to perform parallel compute operations on data in a variety of integer and floating-point data formats based on instructions associated with shader programs. Each vector engine

1702A and matrix engine **1703A** can act as a programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. The vector engine **1702A** and matrix engine **1703A** support the processing of variable width vectors at various SIMD widths, including but not limited to SIMD8, SIMD16, and SIMD32. Input data elements can be stored as a packed data type in a register and the vector engine **1702A** and matrix engine **1703A** can process the various elements based on the data size of the elements. For example, when operating on a 256-bit wide vector, the 256 bits of the vector are stored in a register and the vector is processed as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, different vector widths and register sizes are possible. In one embodiment, the vector engine **1702A** and matrix engine **1703A** are also configurable for SIMT operation on warps or thread groups of various sizes (e.g., 8, 16, or 32 threads).

[0181] Continuing with graphics core **1715A**, the memory load/store unit **1704A** services memory access requests that are issued by the vector engine **1702A**, matrix engine **1703A**, and/or other components of the graphics core **1715A** that have access to memory. The memory access request can be processed by the memory load/store unit **1704A** to load or store the requested data to or from cache or memory into a register file associated with the vector engine **1702A** and/or matrix engine **1703A**. The memory load/store unit **1704A** can also perform prefetching operations. In one embodiment, the memory load/store unit **1704A** is configured to provide SIMT scatter/gather prefetching or block prefetching for data stored in memory **1810**, from memory that is local to other tiles via the tile interconnect **1808**, or from system memory. Prefetching can be performed to a specific L1 cache (e.g., data cache/shared local memory **1706A**), the L2 cache **1804** or the L3 cache **1806**. In one embodiment, a prefetch to the L3 cache **1806** automatically results in the data being stored in the L2 cache **1804**.

[0182] The instruction cache **1705A** stores instructions to be executed by the graphics core **1715A**. In one embodiment, the graphics core **1715A** also includes instruction fetch and prefetch circuitry that fetches or prefetches instructions into the instruction cache **1705A**. The graphics core **1715A** also includes instruction decode logic to decode instructions within the instruction cache **1705A**. The data cache/shared local memory **1706A** can be configured as a data cache that is managed by a cache controller that implements a cache replacement policy and/or configured as explicitly managed shared memory. The ray tracing unit **1708A** includes circuitry to accelerate ray tracing operations. The sampler **1710A** provides texture sampling for 3D operations and media sampling for media operations. The fixed function logic **1712A** includes fixed function circuitry that is shared between the various instances of the vector engine **1702A** and matrix engine **1703A**. Graphics cores **1715B-1715N** can operate in a similar manner as graphics core **1715A**.

[0183] Functionality of the instruction caches **1705A-1705N**, data caches/shared local memory **1706A-1706N**, ray tracing units **1708A-1708N**, samplers **1710A-1710N**, and fixed function logic **1712A-1712N** corresponds with equivalent functionality in the graphics processor architectures described herein. For example, the instruction caches **1705A-1705N** can operate in a similar manner as instruction cache **1455** of FIG. **14D**. The data caches/shared local memory **1706A-1706N**, ray tracing units **1708A-1708N**, and samplers **1710A-1710N** can operate in a similar manner as the cache/SLM **1428A-1428F**, ray tracing units **1427A-1427F**, and samplers **1426A-1426F** of FIG. **14B**. The fixed function logic **1712A-1712N** can include elements of the geometry/fixed function pipeline **1431** and/or additional fixed function logic **1438** of FIG. **14B**. In one embodiment, the ray tracing units **1708A-1708N** include circuitry to perform ray tracing acceleration operations performed by the ray tracing cores **1445** of FIG. **14C**.

[0184] As shown in FIG. **17B**, in one embodiment the vector engine **1702** includes an instruction fetch unit **1737**, a general register file array (GRF) **1724**, an architectural register file array (ARF)

1726, a thread arbiter **1722**, a send unit **1730**, a branch unit **1732**, a set of SIMD floating point units (FPUs) **1734**, and in one embodiment a set of integer SIMD ALUs **1735**. The GRF **1724** and ARF **1726** includes the set of general register files and architecture register files associated with each hardware thread that may be active in the vector engine **1702**. In one embodiment, per thread architectural state is maintained in the ARF **1726**, while data used during thread execution is stored in the GRF **1724**. The execution state of each thread, including the instruction pointers for each thread, can be held in thread-specific registers in the ARF **1726**.

[0185] In one embodiment the vector engine **1702** has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). The architecture has a modular configuration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per graphics core, where graphics core resources are divided across logic used to execute multiple simultaneous threads. The number of logical threads that may be executed by the vector engine **1702** is not limited to the number of hardware threads, and multiple logical threads can be assigned to each hardware thread.

[0186] In one embodiment, the vector engine **1702** can co-issue multiple instructions, which may each be different instructions. The thread arbiter **1722** can dispatch the instructions to one of the send unit **1730**, branch unit **1732**, or SIMD FPU(s) **1734** for execution. Each execution thread can access 128 general-purpose registers within the GRF **1724**, where each register can store 32 bytes, accessible as a variable width vector of 32-bit data elements. In one embodiment, each thread has access to 4 Kbytes within the GRF **1724**, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In one embodiment the vector engine **1702** is partitioned into seven hardware threads that can independently perform computational operations, although the number of threads per vector engine **1702** can also vary according to embodiments. For example, in one embodiment up to 16 hardware threads are supported. In an embodiment in which seven threads may access 4 Kbytes, the GRF **1724** can store a total of 28 Kbytes. Where 16 threads may access 4 Kbytes, the GRF **1724** can store a total of 64 Kbytes. Flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

[0187] In one embodiment, memory operations, sampler operations, and other longer-latency system communications are dispatched via “send” instructions that are executed by the message passing send unit **1730**. In one embodiment, branch instructions are dispatched to a dedicated branch unit **1732** to facilitate SIMD divergence and eventual convergence.

[0188] In one embodiment the vector engine **1702** includes one or more SIMD floating point units (FPU(s)) **1734** to perform floating-point operations. In one embodiment, the FPU(s) **1734** also support integer computation. In one embodiment the FPU(s) **1734** can execute up to M number of 32-bit floating-point (or integer) operations, or execute up to 2M 16-bit integer or 16-bit floating-point operations. In one embodiment, at least one of the FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In some embodiments, a set of 8-bit integer SIMD ALUs **1735** are also present and may be specifically optimized to perform operations associated with machine learning computations. In one embodiment, the SIMD ALUs are replaced by an additional set of SIMD FPUs **1734** that are configurable to perform integer and floating-point operations. In one embodiment, the SIMD FPUs **1734** and SIMD ALUs **1735** are configurable to execute SIMT programs. In one embodiment, combined SIMD+SIMT operation is supported.

[0189] In one embodiment, arrays of multiple instances of the vector engine **1702** can be instantiated in a graphics core. For scalability, product architects can choose the exact number of vector engines per graphics core grouping. In one embodiment the vector engine **1702** can execute instructions across a plurality of execution channels. In a further embodiment, each thread executed on the vector engine **1702** is executed on a different channel.

[0190] As shown in FIG. **17C**, in one embodiment the matrix engine **1703** includes an array of

processing elements that are configured to perform tensor operations including vector/matrix and matrix/matrix operations, such as but not limited to matrix multiply and/or dot product operations. The matrix engine **1703** is configured with M rows and N columns of processing elements (PE **1752AA**-PE **1752MN**) that include multiplier and adder circuits organized in a pipelined fashion. In one embodiment, the processing elements **1752AA**-PE **1752MN** make up the physical pipeline stages of an N wide and M deep systolic array that can be used to perform vector/matrix or matrix/matrix operations in a data-parallel manner, including matrix multiply, fused multiply-add, dot product or other general matrix-matrix multiplication (GEMM) operations. In one embodiment the matrix engine **1703** supports 16-bit floating point operations, as well as 8-bit, 4-bit, 2-bit, and binary integer operations. The matrix engine **1703** can also be configured to accelerate specific machine learning operations. In such embodiments, the matrix engine **1703** can be configured with support for the bfloat (brain floating point) 16-bit floating point format or a tensor float 32-bit floating point format (TF32) that have different numbers of mantissa and exponent bits relative to Institute of Electrical and Electronics Engineers (IEEE) 754 formats.

[0191] In one embodiment, during each cycle, each stage can add the result of operations performed at that stage to the output of the previous stage. In other embodiments, the pattern of data movement between the processing elements **1752AA**-**1752MN** after a set of computational cycles can vary based on the instruction or macro-operation being performed. For example, in one embodiment partial sum loopback is enabled and the processing elements may instead add the output of a current cycle with output generated in the previous cycle. In one embodiment, the final stage of the systolic array can be configured with a loopback to the initial stage of the systolic array. In such embodiment, the number of physical pipeline stages may be decoupled from the number of logical pipeline stages that are supported by the matrix engine **1703**. For example, where the processing elements **1752AA**-**1752MN** are configured as a systolic array of M physical stages, a loopback from stage M to the initial pipeline stage can enable the processing elements **1752AA**-PE**1752MN** to operate as a systolic array of, for example, 2M, 3M, 4M, etc., logical pipeline stages.

[0192] In one embodiment, the matrix engine **1703** includes memory **1741A**-**1741N**, **1742A**-**1742M** to store input data in the form of row and column data for input matrices. Memory **1742A**-**1742M** is configurable to store row elements (A0-Am) of a first input matrix and memory **1741A**-**1741N** is configurable to store column elements (B0-Bn) of a second input matrix. The row and column elements are provided as input to the processing elements **1752AA**-**1752MN** for processing. In one embodiment, row and column elements of the input matrices can be stored in a systolic register file **1740** within the matrix engine **1703** before those elements are provided to the memory **1741A**-**1741N**, **1742A**-**1742M**. In one embodiment, the systolic register file **1740** is excluded and the memory **1741A**-**1741N**, **1742A**-**1742M** is loaded from registers in an associated vector engine (e.g., GRF **1724** of vector engine **1702** of FIG. 17B) or other memory of the graphics core that includes the matrix engine **1703** (e.g., data cache/shared local memory **1706A** for matrix engine **1703A** of FIG. 17A). Results generated by the processing elements **1752AA**-**1752MN** are then output to an output buffer and/or written to a register file (e.g., systolic register file **1740**, GRF **1724**, data cache/shared local memory **1706A**-**1706N**) for further processing by other functional units of the graphics processor or for output to memory.

[0193] In some embodiments, the matrix engine **1703** is configured with support for input sparsity, where multiplication operations for sparse regions of input data can be bypassed by skipping multiply operations that have a zero-value operand. In one embodiment, the processing elements **1752AA**-**1752MN** are configured to skip the performance of certain operations that have zero value input. In one embodiment, sparsity within input matrices can be detected and operations having known zero output values can be bypassed before being submitted to the processing elements **1752AA**-**1752MN**. The loading of zero value operands into the processing elements can be bypassed and the processing elements **1752AA**-**1752MN** can be configured to perform

multiplications on the non-zero value input elements. The matrix engine **1703** can also be configured with support for output sparsity, such that operations with results that are pre-determined to be zero are bypassed. For input sparsity and/or output sparsity, in one embodiment, metadata is provided to the processing elements **1752AA-1752MN** to indicate, for a processing cycle, which processing elements and/or data channels are to be active during that cycle.

[0194] In one embodiment, the matrix engine **1703** includes hardware to enable operations on sparse data having a compressed representation of a sparse matrix that stores non-zero values and metadata that identifies the positions of the non-zero values within the matrix. Exemplary compressed representations include but are not limited to compressed tensor representations such as compressed sparse row (CSR), compressed sparse column (CSC), compressed sparse fiber (CSF) representations. Support for compressed representations enable operations to be performed on input in a compressed tensor format without requiring the compressed representation to be decompressed or decoded. In such embodiment, operations can be performed only on non-zero input values and the resulting non-zero output values can be mapped into an output matrix. In some embodiments, hardware support is also provided for machine-specific lossless data compression formats that are used when transmitting data within hardware or across system busses. Such data may be retained in a compressed format for sparse input data and the matrix engine **1703** can use the compression metadata for the compressed data to enable operations to be performed on only non-zero values, or to enable blocks of zero data input to be bypassed for multiply operations.

[0195] In various embodiments, input data can be provided by a programmer in a compressed tensor representation, or a codec can compress input data into the compressed tensor representation or another sparse data encoding. In addition to support for compressed tensor representations, streaming compression of sparse input data can be performed before the data is provided to the processing elements **1752AA-1752MN**. In one embodiment, compression is performed on data written to a cache memory associated with the graphics core cluster **414**, with the compression being performed with an encoding that is supported by the matrix engine **1703**. In one embodiment, the matrix engine **1703** includes support for input having structured sparsity in which a pre-determined level or pattern of sparsity is imposed on input data. This data may be compressed to a known compression ratio, with the compressed data being processed by the processing elements **1752AA-1752MN** according to metadata associated with the compressed data.

[0196] FIG. **18** illustrates a tile **1800** of a multi-tile processor, according to an embodiment. In one embodiment, the tile **1800** is representative of one of the graphics engine tiles **1510A-1510D** of FIG. **15B** or compute engine tiles **1540A-1540D** of FIG. **15C**. The tile **1800** of the multi-tile graphics processor includes an array of graphics core clusters (e.g., graphics core cluster **414A**, graphics core cluster **414B**, through graphics core cluster **414N**), with each graphics core cluster having an array of graphics cores **1715A-1715N**. The tile **1800** also includes a global dispatcher **1802** to dispatch threads to processing resources of the tile **1800**.

[0197] The tile **1800** can include or couple with an L3 cache **1806** and memory **1810**. In various embodiments, the L3 cache **1806** may be excluded or the tile **1800** can include additional levels of cache, such as an L4 cache. In one embodiment, each instance of the tile **1800** in the multi-tile graphics processor has an associated memory **1810**, such as in FIG. **15B** and FIG. **15C**. In one embodiment, a multi-tile processor can be configured as a multi-chip module in which the L3 cache **1806** and/or memory **1810** reside on separate chiplets than the graphics core clusters **414A-414N**. In this context, a chiplet is an at least partially packaged integrated circuit that includes distinct units of logic that can be assembled with other chiplets into a larger package. For example, the L3 cache **1806** can be included in a dedicated cache chiplet or can reside on the same chiplet as the graphics core clusters **414A-414N**. In one embodiment, the L3 cache **1806** can be included in an active base die or active interposer, as illustrated in FIG. **11C**.

[0198] A memory fabric **1803** enables communication among the graphics core clusters **414A-414N**, L3 cache **1806**, and memory **1810**. An L2 cache **1804** couples with the memory fabric **1803**

and is configurable to cache transactions performed via the memory fabric **1803**. A tile interconnect **1808** enables communication with other tiles on the graphics processors and may be one of tile interconnects **323A-323F** of FIG. **3B** and **3C**. In embodiments in which the L3 cache **1806** is excluded from the tile **1800**, the L2 cache **1804** may be configured as a combined L2/L3 cache. The memory fabric **1803** is configurable to route data to the L3 cache **1806** or memory controllers associated with the memory **1810** based on the presence or absence of the L3 cache **1806** in a specific implementation. The L3 cache **1806** can be configured as a per-tile cache that is dedicated to processing resources of the tile **1800** or may be a partition of a GPU-wide L3 cache.

[0199] FIG. **19** is a block diagram illustrating graphics processor instruction formats **1900** according to some embodiments. In one or more embodiment, the graphics processor cores support an instruction set having instructions in multiple formats. The solid lined boxes illustrate the components that are generally included in a graphics core instruction, while the dashed lines include components that are optional or that are only included in a sub-set of the instructions. In some embodiments, the graphics processor instruction format **1900** described and illustrated are macro-instructions, in that they are instructions supplied to the graphics core, as opposed to micro-operations resulting from instruction decode once the instruction is processed. Thus, a single instruction may cause hardware to perform multiple micro-operations.

[0200] In some embodiments, the graphics processor natively supports instructions in a 128-bit instruction format **1910**. A 64-bit compacted instruction format **1930** is available for some instructions based on the selected instruction, instruction options, and number of operands. The native 128-bit instruction format **1910** provides access to all instruction options, while some options and operations are restricted in the 64-bit format **1930**. The native instructions available in the 64-bit format **1930** vary by embodiment. In some embodiments, the instruction is compacted in part using a set of index values in an index field **1913**. The graphics core hardware references a set of compaction tables based on the index values and uses the compaction table outputs to reconstruct a native instruction in the 128-bit instruction format **1910**. Other sizes and formats of instruction can be used.

[0201] For each format, instruction opcode **1912** defines the operation that the graphics core is to perform. The graphics cores execute each instruction in parallel across the multiple data elements of each operand. For example, in response to an add instruction the graphics core performs a simultaneous add operation across each color channel representing a texture element or picture element. By default, the graphics core performs each instruction across all data channels of the operands. In some embodiments, instruction control field **1914** enables control over certain execution options, such as channels selection (e.g., predication) and data channel order (e.g., swizzle). For instructions in the 128-bit instruction format **1910** an exec-size field **1916** limits the number of data channels that will be executed in parallel. In some embodiments, exec-size field **1916** is not available for use in the 64-bit compact instruction format **1930**.

[0202] Some graphics core instructions have up to three operands including two source operands, src0 **1920**, src1 **1922**, and one destination **1918**. In some embodiments, the graphics cores support dual destination instructions, where one of the destinations is implied. Data manipulation instructions can have a third source operand (e.g., SRC2 **1924**), where the instruction opcode **1912** determines the number of source operands. An instruction's last source operand can be an immediate (e.g., hard-coded) value passed with the instruction.

[0203] In some embodiments, the 128-bit instruction format **1910** includes an access/address mode field **1926** specifying, for example, whether direct register addressing mode or indirect register addressing mode is used. When direct register addressing mode is used, the register address of one or more operands is directly provided by bits in the instruction.

[0204] In some embodiments, the 128-bit instruction format **1910** includes an access/address mode field **1926**, which specifies an address mode and/or an access mode for the instruction. In one embodiment the access mode is used to define a data access alignment for the instruction. Some

embodiments support access modes including a 16-byte aligned access mode and a 1-byte aligned access mode, where the byte alignment of the access mode determines the access alignment of the instruction operands. For example, when in a first mode, the instruction may use byte-aligned addressing for source and destination operands and when in a second mode, the instruction may use 16-byte-aligned addressing for all source and destination operands.

[0205] In one embodiment, the address mode portion of the access/address mode field **1926** determines whether the instruction is to use direct or indirect addressing. When direct register addressing mode is used bits in the instruction directly provide the register address of one or more operands. When indirect register addressing mode is used, the register address of one or more operands may be computed based on an address register value and an address immediate field in the instruction.

[0206] In some embodiments instructions are grouped based on opcode **1912** bit-fields to simplify Opcode decode **1940**. For an 8-bit opcode, bits **4**, **5**, and **6** allow the graphics core to determine the type of opcode. The precise opcode grouping shown is merely an example. In some embodiments, a move and logic opcode group **1942** includes data movement and logic instructions (e.g., move (mov), compare (cmp)). In some embodiments, move and logic group **1942** shares the five most significant bits (MSB), where move (mov) instructions are in the form of 0000xxxxb and logic instructions are in the form of 0001xxxxb. A flow control instruction group **1944** (e.g., call, jump (jmp)) includes instructions in the form of 0010xxxxb (e.g., 0x20). A miscellaneous instruction group **1946** includes a mix of instructions, including synchronization instructions (e.g., wait, send) in the form of 0011xxxxb (e.g., 0x30). A parallel math instruction group **1948** includes component-wise arithmetic instructions (e.g., add, multiply (mul)) in the form of 0100xxxxb (e.g., 0x40). The parallel math instruction group **1948** performs the arithmetic operations in parallel across data channels. The vector math group **1950** includes arithmetic instructions (e.g., dp4) in the form of 0101xxxxb (e.g., 0x50). The vector math group performs arithmetic such as dot product calculations on vector operands. The illustrated opcode decode **1940**, in one embodiment, can be used to determine which portion of a graphics core will be used to execute a decoded instruction. For example, some instructions may be designated as systolic instructions that will be performed by a systolic array. Other instructions, such as ray-tracing instructions (not shown) can be routed to a ray-tracing core or ray-tracing logic within a slice or partition of execution logic.

Graphics Pipeline

[0207] FIG. **20** is a block diagram of another embodiment of a graphics processor **2000**. Elements of FIG. **20** having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0208] In some embodiments, graphics processor **2000** includes a geometry pipeline **2020**, a media pipeline **2030**, a display engine **2040**, thread execution logic **2050**, and a render output pipeline **2070**. In some embodiments, graphics processor **2000** is a graphics processor within a multi-core processing system that includes one or more general-purpose processing cores. The graphics processor is controlled by register writes to one or more control registers (not shown) or via commands issued to graphics processor **2000** via a ring interconnect **2002**. In some embodiments, ring interconnect **2002** couples graphics processor **2000** to other processing components, such as other graphics processors or general-purpose processors. Commands from ring interconnect **2002** are interpreted by a command streamer **2003**, which supplies instructions to individual components of the geometry pipeline **2020** or the media pipeline **2030**.

[0209] In some embodiments, command streamer **2003** directs the operation of a vertex fetcher **2005** that reads vertex data from memory and executes vertex-processing commands provided by command streamer **2003**. In some embodiments, vertex fetcher **2005** provides vertex data to a vertex shader **2007**, which performs coordinate space transformation and lighting operations to each vertex. In some embodiments, vertex fetcher **2005** and vertex shader **2007** execute vertex-

processing instructions by dispatching execution threads to graphics cores **2052A-2052B** via a thread dispatcher **2031**.

[0210] In some embodiments, graphics cores **2052A-2052B** are an array of vector processors having an instruction set for performing graphics and media operations. In some embodiments, graphics cores **2052A-2052B** have an attached L1 cache **2051** that is specific for each array or shared between the arrays. The cache can be configured as a data cache, an instruction cache, or a single cache that is partitioned to contain data and instructions in different partitions.

[0211] In some embodiments, geometry pipeline **2020** includes tessellation components to perform hardware-accelerated tessellation of 3D objects. In some embodiments, a programmable hull shader **2011** configures the tessellation operations. A programmable domain shader **2017** provides back-end evaluation of tessellation output. A tessellator **2013** operates at the direction of hull shader **2011** and contains special purpose logic to generate a set of detailed geometric objects based on a coarse geometric model that is provided as input to geometry pipeline **2020**. In some embodiments, if tessellation is not used, tessellation components (e.g., hull shader **2011**, tessellator **2013**, and domain shader **2017**) can be bypassed. The tessellation components can operate based on data received from the vertex shader **2007**.

[0212] In some embodiments, complete geometric objects can be processed by a geometry shader **2019** via one or more threads dispatched to graphics cores **2052A-2052B** or can proceed directly to the clipper **2029**. In some embodiments, the geometry shader operates on entire geometric objects, rather than vertices or patches of vertices as in previous stages of the graphics pipeline. If the tessellation is disabled the geometry shader **2019** receives input from the vertex shader **2007**. In some embodiments, geometry shader **2019** is programmable by a geometry shader program to perform geometry tessellation if the tessellation units are disabled.

[0213] Before rasterization, a clipper **2029** processes vertex data. The clipper **2029** may be a fixed function clipper or a programmable clipper having clipping and geometry shader functions. In some embodiments, a rasterizer and depth test component **2073** in the render output pipeline **2070** dispatches pixel shaders to convert the geometric objects into per pixel representations. In some embodiments, pixel shader logic is included in thread execution logic **2050**. In some embodiments, an application can bypass the rasterizer and depth test component **2073** and access un-rasterized vertex data via a stream out unit **2023**.

[0214] The graphics processor **2000** has an interconnect bus, interconnect fabric, or some other interconnect mechanism that allows data and message passing amongst the major components of the processor. In some embodiments, graphics cores **2052A-2052B** and associated logic units (e.g., L1 cache **2051**, sampler **2054**, texture cache **2058**, etc.) interconnect via a data port **2056** to perform memory access and communicate with render output pipeline components of the processor. In some embodiments, sampler **2054**, caches **2051**, **2058** and graphics cores **2052A-2052B** each have separate memory access paths. In one embodiment the texture cache **2058** can also be configured as a sampler cache.

[0215] In some embodiments, render output pipeline **2070** contains a rasterizer and depth test component **2073** that converts vertex-based objects into an associated pixel-based representation. In some embodiments, the rasterizer logic includes a windower/masker unit to perform fixed function triangle and line rasterization. An associated render cache **2078** and depth cache **2079** are also available in some embodiments. A pixel operations component **2077** performs pixel-based operations on the data, though in some instances, pixel operations associated with 2D operations (e.g., bit block image transfers with blending) are performed by the 2D engine **2041**, or substituted at display time by the display controller **2043** using overlay display planes. In some embodiments, a shared L3 cache **2075** is available to all graphics components, allowing the sharing of data without the use of main system memory.

[0216] In some embodiments, media pipeline **2030** includes a media engine **2037** and a video front-end **2034**. In some embodiments, video front-end **2034** receives pipeline commands from the

command streamer **2003**. In some embodiments, media pipeline **2030** includes a separate command streamer. In some embodiments, video front-end **2034** processes media commands before sending the command to the media engine **2037**. In some embodiments, media engine **2037** includes thread spawning functionality to spawn threads for dispatch to thread execution logic **2050** via thread dispatcher **2031**.

[0217] In some embodiments, graphics processor **2000** includes a display engine **2040**. In some embodiments, display engine **2040** is external to processor **2000** and couples with the graphics processor via the ring interconnect **2002**, or some other interconnect bus or fabric. In some embodiments, display engine **2040** includes a 2D engine **2041** and a display controller **2043**. In some embodiments, display engine **2040** contains special purpose logic capable of operating independently of the 3D pipeline. In some embodiments, display controller **2043** couples with a display device (not shown), which may be a system integrated display device, as in a laptop computer, or an external display device attached via a display device connector.

[0218] In some embodiments, the geometry pipeline **2020** and media pipeline **2030** are configurable to perform operations based on multiple graphics and media programming interfaces and are not specific to any one application programming interface (API). In some embodiments, driver software for the graphics processor translates API calls that are specific to a particular graphics or media library into commands that can be processed by the graphics processor. In some embodiments, support is provided for the Open Graphics Library (OpenGL), Open Computing Language (OpenCL), and/or Vulkan graphics and compute API, all from the Khronos Group. In some embodiments, support may also be provided for the Direct3D library from the Microsoft Corporation. In some embodiments, a combination of these libraries may be supported. Support may also be provided for the Open Source Computer Vision Library (OpenCV). A future API with a compatible 3D pipeline would also be supported if a mapping can be made from the pipeline of the future API to the pipeline of the graphics processor.

Graphics Pipeline Programming

[0219] FIG. **21A** is a block diagram illustrating a graphics processor command format **2100** that may be used to program graphics processing pipelines according to some embodiments. FIG. **21B** is a block diagram illustrating a graphics processor command sequence **2110** according to an embodiment. The solid lined boxes in FIG. **21A** illustrate the components that are generally included in a graphics command while the dashed lines include components that are optional or that are only included in a sub-set of the graphics commands. The exemplary graphics processor command format **2100** of FIG. **21A** includes data fields to identify a client **2102**, a command operation code (opcode) **2104**, and a data field **2106** for the command. A sub-opcode **2105** and a command size **2108** are also included in some commands.

[0220] In some embodiments, client **2102** specifies the client unit of the graphics device that processes the command data. In some embodiments, a graphics processor command parser examines the client field of each command to condition the further processing of the command and route the command data to the appropriate client unit. In some embodiments, the graphics processor client units include a memory interface unit, a render unit, a 2D unit, a 3D unit, and a media unit. Each client unit has a corresponding processing pipeline that processes the commands. Once the command is received by the client unit, the client unit reads the opcode **2104** and, if present, sub-opcode **2105** to determine the operation to perform. The client unit performs the command using information in data field **2106**. For some commands an explicit command size **2108** is expected to specify the size of the command. In some embodiments, the command parser automatically determines the size of at least some of the commands based on the command opcode. In some embodiments commands are aligned via multiples of a double word. Other command formats can be used.

[0221] The flow diagram in FIG. **21B** illustrates an exemplary graphics processor command sequence **2110**. In some embodiments, software or firmware of a data processing system that

features an embodiment of a graphics processor uses a version of the command sequence shown to set up, execute, and terminate a set of graphics operations. A sample command sequence is shown and described for purposes of example only as embodiments are not limited to these specific commands or to this command sequence. Moreover, the commands may be issued as batch of commands in a command sequence, such that the graphics processor will process the sequence of commands in at least partially concurrence.

[0222] In some embodiments, the graphics processor command sequence **2110** may begin with a pipeline flush command **2112** to cause any active graphics pipeline to complete the currently pending commands for the pipeline. In some embodiments, the 3D pipeline **2122** and the media pipeline **2124** do not operate concurrently. The pipeline flush is performed to cause the active graphics pipeline to complete any pending commands. In response to a pipeline flush, the command parser for the graphics processor will pause command processing until the active drawing engines complete pending operations and the relevant read caches are invalidated. Optionally, any data in the render cache that is marked 'dirty' can be flushed to memory. In some embodiments, pipeline flush command **2112** can be used for pipeline synchronization or before placing the graphics processor into a low power state.

[0223] In some embodiments, a pipeline select command **2113** is used when a command sequence requires the graphics processor to explicitly switch between pipelines. In some embodiments, a pipeline select command **2113** is required only once within an execution context before issuing pipeline commands unless the context is to issue commands for both pipelines. In some embodiments, a pipeline flush command **2112** is required immediately before a pipeline switch via the pipeline select command **2113**.

[0224] In some embodiments, a pipeline control command **2114** configures a graphics pipeline for operation and is used to program the 3D pipeline **2122** and the media pipeline **2124**. In some embodiments, pipeline control command **2114** configures the pipeline state for the active pipeline. In one embodiment, the pipeline control command **2114** is used for pipeline synchronization and to clear data from one or more cache memories within the active pipeline before processing a batch of commands.

[0225] In some embodiments, commands related to the return buffer state **2116** are used to configure a set of return buffers for the respective pipelines to write data. Some pipeline operations require the allocation, selection, or configuration of one or more return buffers into which the operations write intermediate data during processing. In some embodiments, the graphics processor also uses one or more return buffers to store output data and to perform cross thread communication. In some embodiments, the return buffer state **2116** includes selecting the size and number of return buffers to use for a set of pipeline operations.

[0226] The remaining commands in the command sequence differ based on the active pipeline for operations. Based on a pipeline determination **2120**, the command sequence is tailored to the 3D pipeline **2122** beginning with the 3D pipeline state **2130** or the media pipeline **2124** beginning at the media pipeline state **2140**.

[0227] The commands to configure the 3D pipeline state **2130** include 3D state setting commands for vertex buffer state, vertex element state, constant color state, depth buffer state, and other state variables that are to be configured before 3D primitive commands are processed. The values of these commands are determined at least in part based on the particular 3D API in use. In some embodiments, 3D pipeline state **2130** commands are also able to selectively disable or bypass certain pipeline elements if those elements will not be used.

[0228] In some embodiments, 3D primitive **2132** command is used to submit 3D primitives to be processed by the 3D pipeline. Commands and associated parameters that are passed to the graphics processor via the 3D primitive **2132** command are forwarded to the vertex fetch function in the graphics pipeline. The vertex fetch function uses the 3D primitive **2132** command data to generate vertex data structures. The vertex data structures are stored in one or more return buffers. In some

embodiments, 3D primitive **2132** command is used to perform vertex operations on 3D primitives via vertex shaders. To process vertex shaders, 3D pipeline **2122** dispatches shader programs to the graphics cores.

[0229] In some embodiments, 3D pipeline **2122** is triggered via an execute **2134** command or event. In some embodiments, a register write triggers command execution. In some embodiments execution is triggered via a ‘go’ or ‘kick’ command in the command sequence. In one embodiment, command execution is triggered using a pipeline synchronization command to flush the command sequence through the graphics pipeline. The 3D pipeline will perform geometry processing for the 3D primitives. Once operations are complete, the resulting geometric objects are rasterized and the pixel engine colors the resulting pixels. Additional commands to control pixel shading and pixel back-end operations may also be included for those operations.

[0230] In some embodiments, the graphics processor command sequence **2110** follows the media pipeline **2124** path when performing media operations. In general, the specific use and manner of programming for the media pipeline **2124** depends on the media or compute operations to be performed. Specific media decode operations may be offloaded to the media pipeline during media decode. In some embodiments, the media pipeline can also be bypassed and media decode can be performed in whole or in part using resources provided by one or more general-purpose processing cores. In one embodiment, the media pipeline also includes elements for general-purpose graphics processor unit (GPGPU) operations, where the graphics processor is used to perform SIMD vector operations using computational shader programs that are not explicitly related to the rendering of graphics primitives.

[0231] In some embodiments, media pipeline **2124** is configured in a similar manner as the 3D pipeline **2122**. A set of commands to configure the media pipeline state **2140** are dispatched or placed into a command queue before the media object commands **2142**. In some embodiments, commands for the media pipeline state **2140** include data to configure the media pipeline elements that will be used to process the media objects. This includes data to configure the video decode and video encode logic within the media pipeline, such as encode or decode format. In some embodiments, commands for the media pipeline state **2140** also support the use of one or more pointers to “indirect” state elements that contain a batch of state settings.

[0232] In some embodiments, media object commands **2142** supply pointers to media objects for processing by the media pipeline. The media objects include memory buffers containing video data to be processed. In some embodiments, all media pipeline states must be valid before issuing a media object command **2142**. Once the pipeline state is configured and media object commands **2142** are queued, the media pipeline **2124** is triggered via an execute command **2144** or an equivalent execute event (e.g., register write). Output from media pipeline **2124** may then be post processed by operations provided by the 3D pipeline **2122** or the media pipeline **2124**. In some embodiments, GPGPU operations are configured and executed in a similar manner as media operations.

Graphics Software Architecture

[0233] FIG. **22** illustrates an exemplary graphics software architecture for a data processing system **2200** according to some embodiments. In some embodiments, software architecture includes a 3D graphics application **2210**, an operating system **2220**, and at least one processor **2230**. In some embodiments, processor **2230** includes a graphics processor **2232** and one or more general-purpose processor core(s) **2234**. The graphics application **2210** and operating system **2220** each execute in the system memory **2250** of the data processing system.

[0234] In some embodiments, 3D graphics application **2210** contains one or more shader programs including shader instructions **2212**. The shader language instructions may be in a high-level shader language, such as the High-Level Shader Language (HLSL) of Direct3D, the OpenGL Shader Language (GLSL), and so forth. The application also includes executable instructions **2214** in a machine language suitable for execution by the general-purpose processor core **2234**. The

application also includes graphics objects **2216** defined by vertex data.

[0235] In some embodiments, operating system **2220** is a Microsoft® Windows® operating system from the Microsoft Corporation, a proprietary UNIX-like operating system, or an open source UNIX-like operating system using a variant of the Linux kernel. The operating system **2220** can support a graphics API **2222** such as the Direct3D API, the OpenGL API, or the Vulkan API. When the Direct3D API is in use, the operating system **2220** uses a front-end shader compiler **2224** to compile any shader instructions **2212** in HLSL into a lower-level shader language. The compilation may be a just-in-time (JIT) compilation or the application can perform shader pre-compilation. In some embodiments, high-level shaders are compiled into low-level shaders during the compilation of the 3D graphics application **2210**. In some embodiments, the shader instructions **2212** are provided in an intermediate form, such as a version of the Standard Portable Intermediate Representation (SPIR) used by the Vulkan API.

[0236] In some embodiments, user mode graphics driver **2226** contains a back-end shader compiler **2227** to convert the shader instructions **2212** into a hardware specific representation. When the OpenGL API is in use, shader instructions **2212** in the GLSL high-level language are passed to a user mode graphics driver **2226** for compilation. In some embodiments, user mode graphics driver **2226** uses operating system kernel mode functions **2228** to communicate with a kernel mode graphics driver **2229**. In some embodiments, kernel mode graphics driver **2229** communicates with graphics processor **2232** to dispatch commands and instructions.

IP Core Implementations

[0237] One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

[0238] FIG. **23A** is a block diagram illustrating an IP core development system **2300** that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The IP core development system **2300** may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., an SOC integrated circuit). A design facility **2330** can generate a software simulation **2310** of an IP core design in a high-level programming language (e.g., C/C++). The software simulation **2310** can be used to design, test, and verify the behavior of the IP core using a simulation model **2312**. The simulation model **2312** may include functional, behavioral, and/or timing simulations. A register transfer level (RTL) design **2315** can then be created or synthesized from the simulation model **2312**. The RTL design **2315** is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL design **2315**, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

[0239] The RTL design **2315** or equivalent may be further synthesized by the design facility into a hardware model **2320**, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a 3.sup.rd party fabrication facility **2365** using non-volatile memory **2340** (e.g., hard disk, flash memory, or any non-volatile storage

medium). Alternatively, the IP core design may be transmitted (e.g., via the Internet) over a wired connection **2350** or wireless connection **2360**. The fabrication facility **2365** may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit can be configured to perform operations in accordance with at least one embodiment described herein.

[0240] FIG. **23B** illustrates a cross-section side view of an integrated circuit package assembly **2370**, according to some embodiments described herein. The integrated circuit package assembly **2370** illustrates an implementation of one or more processor or accelerator devices as described herein. The package assembly **2370** includes multiple units of hardware logic **2372**, **2374** connected to a substrate **2380**. The logic **2372**, **2374** may be implemented at least partly in configurable logic or fixed-functionality logic hardware, and can include one or more portions of any of the processor core(s), graphics processor(s), or other accelerator devices described herein. Each unit of logic **2372**, **2374** can be implemented within a semiconductor die and coupled with the substrate **2380** via an interconnect structure **2373**. The interconnect structure **2373** may be configured to route electrical signals between the logic **2372**, **2374** and the substrate **2380**, and can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **2373** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic **2372**, **2374**. In some embodiments, the substrate **2380** is an epoxy-based laminate substrate. The substrate **2380** may include other suitable types of substrates in other embodiments. The package assembly **2370** can be connected to other electrical devices via a package interconnect **2383**. The package interconnect **2383** may be coupled to a surface of the substrate **2380** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0241] In some embodiments, the units of logic **2372**, **2374** are electrically coupled with a bridge **2382** that is configured to route electrical signals between the logic **2372**, **2374**. The bridge **2382** may be a dense interconnect structure that provides a route for electrical signals. The bridge **2382** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic **2372**, **2374**.

[0242] Although two units of logic **2372**, **2374** and a bridge **2382** are illustrated, embodiments described herein may include more or fewer logic units on one or more dies. The one or more dies may be connected by zero or more bridges, as the bridge **2382** may be excluded when the logic is included on a single die. Alternatively, multiple dies or units of logic can be connected by one or more bridges. Additionally, multiple logic units, dies, and bridges can be connected together in other possible configurations, including three-dimensional configurations.

[0243] FIG. **23C** illustrates a package assembly **2390** that includes multiple units of hardware logic chiplets connected to a substrate **2380**. A graphics processing unit, parallel processor, and/or compute accelerator as described herein can be composed from diverse silicon chiplets that are separately manufactured. A diverse set of chiplets with different IP core logic can be assembled into a single device. Additionally, the chiplets can be integrated into a base die or base chiplet using active interposer technology. The concepts described herein enable the interconnection and communication between the different forms of IP within the GPU. IP cores can be manufactured using different process technologies and composed during manufacturing, which avoids the complexity of converging multiple IPs, especially on a large SoC with several flavors IPs, to the same manufacturing process. Enabling the use of multiple process technologies improves the time to market and provides a cost-effective way to create multiple product SKUs. Additionally, the disaggregated IPs are more amenable to being power gated independently, components that are not in use on a given workload can be powered off, reducing overall power consumption.

[0244] In various embodiments a package assembly **2390** can include components and chiplets that

are interconnected by a fabric **2385** and/or one or more bridges **2387**. The chiplets within the package assembly **2390** may have a 2.5D arrangement using Chip-on-Wafer-on-Substrate stacking in which multiple dies are stacked side-by-side on a silicon interposer **2389** that couples the chiplets with the substrate **2380**. The substrate **2380** includes electrical connections to the package interconnect **2383**. In one embodiment the silicon interposer **2389** is a passive interposer that includes through-silicon vias (TSVs) to electrically couple chiplets within the package assembly **2390** to the substrate **2380**. In one embodiment, silicon interposer **2389** is an active interposer that includes embedded logic in addition to TSVs. In such embodiment, the chiplets within the package assembly **2390** are arranged using 3D face to face die stacking on top of the active interposer **2389**. The active interposer **2389** can include hardware logic for I/O **2391**, cache memory **2392**, and other hardware logic **2393**, in addition to interconnect fabric **2385** and a silicon bridge **2387**. The fabric **2385** enables communication between the various logic chiplets **2372**, **2374** and the logic **2391**, **2393** within the active interposer **2389**. The fabric **2385** may be an NoC interconnect or another form of packet switched fabric that switches data packets between components of the package assembly. For complex assemblies, the fabric **2385** may be a dedicated chiplet enables communication between the various hardware logic of the package assembly **2390**.

[0245] Bridge structures **2387** within the active interposer **2389** may be used to facilitate a point-to-point interconnect between, for example, logic or I/O chiplets **2374** and memory chiplets **2375**. In some implementations, bridge structures **2387** may also be embedded within the substrate **2380**. The hardware logic chiplets can include special purpose hardware logic chiplets **2372**, logic or I/O chiplets **2374**, and/or memory chiplets **2375**. The hardware logic chiplets **2372** and logic or I/O chiplets **2374** may be implemented at least partly in configurable logic or fixed-functionality logic hardware and can include one or more portions of any of the processor core(s), graphics processor(s), parallel processors, or other accelerator devices described herein. The memory chiplets **2375** can be DRAM (e.g., GDDR, HBM) memory or cache (SRAM) memory. Cache memory **2392** within the active interposer **2389** (or substrate **2380**) can act as a global cache for the package assembly **2390**, part of a distributed global cache, or as a dedicated cache for the fabric **2385**.

[0246] Each chiplet can be fabricated as separate semiconductor die and coupled with a base die that is embedded within or coupled with the substrate **2380**. The coupling with the substrate **2380** can be performed via an interconnect structure **2373**. The interconnect structure **2373** may be configured to route electrical signals between the various chiplets and logic within the substrate **2380**. The interconnect structure **2373** can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **2373** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic, I/O, and memory chiplets. In one embodiment, an additional interconnect structure couples the active interposer **2389** with the substrate **2380**.

[0247] In some embodiments, the substrate **2380** is an epoxy-based laminate substrate. The substrate **2380** may include other suitable types of substrates in other embodiments. The package assembly **2390** can be connected to other electrical devices via a package interconnect **2383**. The package interconnect **2383** may be coupled to a surface of the substrate **2380** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0248] In some embodiments, a logic or I/O chiplet **2374** and a memory chiplet **2375** can be electrically coupled via a bridge **2387** that is configured to route electrical signals between the logic or I/O chiplet **2374** and a memory chiplet **2375**. The bridge **2387** may be a dense interconnect structure that provides a route for electrical signals. The bridge **2387** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic or I/O chiplet **2374** and a memory chiplet **2375**. The bridge **2387** may also be referred to as a silicon bridge or an interconnect bridge. For example, the bridge **2387**, in some embodiments, is an Embedded Multi-

die Interconnect Bridge (EMIB). In some embodiments, the bridge **2387** may simply be a direct connection from one chiplet to another chiplet.

[0249] FIG. **23D** illustrates a package assembly **2394** including interchangeable chiplets **2395**, according to an embodiment. The interchangeable chiplets **2395** can be assembled into standardized slots on one or more base chiplets **2396**, **2398**. The base chiplets **2396**, **2398** can be coupled via a bridge interconnect **2397**, which can be similar to the other bridge interconnects described herein and may be, for example, an EMIB. Memory chiplets can also be connected to logic or I/O chiplets via a bridge interconnect. I/O and logic chiplets can communicate via an interconnect fabric. The base chiplets can each support one or more slots in a standardized format for one of logic or I/O or memory/cache.

[0250] In one embodiment, SRAM and power delivery circuits can be fabricated into one or more of the base chiplets **2396**, **2398**, which can be fabricated using a different process technology relative to the interchangeable chiplets **2395** that are stacked on top of the base chiplets. For example, the base chiplets **2396**, **2398** can be fabricated using a larger process technology, while the interchangeable chiplets can be manufactured using a smaller process technology. One or more of the interchangeable chiplets **2395** may be memory (e.g., DRAM) chiplets. Different memory densities can be selected for the package assembly **2394** based on the power, and/or performance targeted for the product that uses the package assembly **2394**. Additionally, logic chiplets with a different number of type of functional units can be selected at time of assembly based on the power, and/or performance targeted for the product. Additionally, chiplets containing IP logic cores of differing types can be inserted into the interchangeable chiplet slots, enabling hybrid processor designs that can mix and match different technology IP blocks.

Exemplary System on a Chip Integrated Circuit

[0251] FIGS. **24-25B** illustrate exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

[0252] FIG. **24** is a block diagram illustrating an exemplary system on a chip integrated circuit **2400** that may be fabricated using one or more IP cores, according to an embodiment. Exemplary integrated circuit **2400** includes one or more application processor(s) **2405** (e.g., CPUs), at least one graphics processor **2410**, and may additionally include an image processor **2415** and/or a video processor **2420**, any of which may be a modular IP core from the same or multiple different design facilities. Integrated circuit **2400** includes peripheral or bus logic including a USB controller **2425**, UART controller **2430**, an SPI/SDIO controller **2435**, and an I2S/I2C controller **2440**. Additionally, the integrated circuit can include a display device **2445** coupled to one or more of a high-definition multimedia interface (HDMI) controller **2450** and a mobile industry processor interface (MIPI) display interface **2455**. Storage may be provided by a flash memory subsystem **2460** including flash memory and a flash memory controller. Memory interface may be provided via a memory controller **2465** for access to SDRAM or SRAM memory devices. Some integrated circuits additionally include an embedded security engine **2470**.

[0253] FIGS. **25A-25B** are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. **25A** illustrates an exemplary graphics processor **2510** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. FIG. **25B** illustrates an additional exemplary graphics processor **2540** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. Graphics processor **2510** of FIG. **25A** is an example of a low power graphics processor core. Graphics processor **2540** of FIG. **25B** is an example of a higher performance graphics processor core. Each of graphics processor **2510** and graphics processor **2540** can be variants of the graphics processor **2410** of FIG. **24**.

[0254] As shown in FIG. 25A, graphics processor **2510** includes a vertex processor **2505** and one or more fragment processor(s) **2515A-2515N** (e.g., **2515A**, **2515B**, **2515C**, **2515D**, through **2515N-1**, and **2515N**). Graphics processor **2510** can execute different shader programs via separate logic, such that the vertex processor **2505** is optimized to execute operations for vertex shader programs, while the one or more fragment processor(s) **2515A-2515N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. The vertex processor **2505** performs the vertex processing stage of the 3D graphics pipeline and generates primitives and vertex data. The fragment processor(s) **2515A-2515N** use the primitive and vertex data generated by the vertex processor **2505** to produce a framebuffer that is displayed on a display device. In one embodiment, the fragment processor(s) **2515A-2515N** are optimized to execute fragment shader programs as provided for in the OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in the Direct 3D API.

[0255] Graphics processor **2510** additionally includes one or more memory management units (MMUs) **2520A-2520B**, cache(s) **2525A-2525B**, and circuit interconnect(s) **2530A-2530B**. The one or more MMU(s) **2520A-2520B** provide for virtual to physical address mapping for the graphics processor **2510**, including for the vertex processor **2505** and/or fragment processor(s) **2515A-2515N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in the one or more cache(s) **2525A-2525B**. In one embodiment the one or more MMU(s) **2520A-2520B** may be synchronized with other MMUs within the system, including one or more MMUs associated with the one or more application processor(s) **2405**, image processor **2415**, and/or video processor **2420** of FIG. 24, such that each processor **2405-2420** can participate in a shared or unified virtual memory system. The one or more circuit interconnect(s) **2530A-2530B** enable graphics processor **2510** to interface with other IP cores within the SoC, either via an internal bus of the SoC or via a direct connection, according to embodiments.

[0256] As shown FIG. 25B, graphics processor **2540** includes the one or more MMU(s) **2520A-2520B**, cache(s) **2525A-2525B**, and circuit interconnect(s) **2530A-2530B** of the graphics processor **2510** of FIG. 25A. Graphics processor **2540** includes one or more shader core(s) **2555A-2555N** (e.g., **2555A**, **2555B**, **2555C**, **2555D**, **2555E**, **2555F**, through **2555N-1**, and **2555N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. The unified shader core architecture is also configurable to execute direct compiled high-level GPGPU programs (e.g., CUDA). The exact number of shader cores present can vary among embodiments and implementations. Additionally, graphics processor **2540** includes an inter-core task manager **2545**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **2555A-2555N** and a tiling unit **2558** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

[0257] Flow diagrams as illustrated herein provide examples of sequences of various process actions. The flow diagrams can indicate operations to be executed by a software or firmware routine, as well as physical operations. A flow diagram can illustrate an example of the implementation of states of a finite state machine (FSM), which can be implemented in hardware and/or software. Although shown in a particular sequence or order, unless otherwise specified, the order of the actions can be modified. Thus, the illustrated diagrams should be understood only as examples, and the process can be performed in a different order, and some actions can be performed in parallel. Additionally, one or more actions can be omitted; thus, not all implementations will perform all actions.

[0258] To the extent various operations or functions are described herein, they can be described or defined as software code, instructions, configuration, and/or data. The content can be directly

executable (“object” or “executable” form), source code, or difference code (“delta” or “patch” code). The software content of what is described herein can be provided via an article of manufacture with the content stored thereon, or via a method of operating a communication interface to send data via the communication interface. A machine readable storage medium can cause a machine to perform the functions or operations described and includes any mechanism that stores information in a form accessible by a machine (e.g., computing device, electronic system, etc.), such as recordable/non-recordable media (e.g., read only memory (ROM), random access memory (RAM), magnetic disk storage media, optical storage media, flash memory devices, etc.). A communication interface includes any mechanism that interfaces to any of a hardwired, wireless, optical, etc., medium to communicate to another device, such as a memory bus interface, a processor bus interface, an Internet connection, a disk controller, etc. The communication interface can be configured by providing configuration parameters and/or sending signals to prepare the communication interface to provide a data signal describing the software content. The communication interface can be accessed via one or more commands or signals sent to the communication interface.

[0259] Various components described herein can be a means for performing the operations or functions described. Each component described herein includes software, hardware, or a combination of these. The components can be implemented as software modules, hardware modules, special-purpose hardware (e.g., application specific hardware, application specific integrated circuits (ASICs), digital signal processors (DSPs), etc.), embedded controllers, hardwired circuitry, etc.

[0260] Besides what is described herein, various modifications can be made to what is disclosed and implementations of the invention without departing from their scope. Therefore, the illustrations and examples herein should be construed in an illustrative, and not a restrictive sense. The scope of the invention should be measured solely by reference to the claims that follow. Although not depicted, any system can include and use a power supply such as but not limited to a battery, AC-DC converter at least to receive alternating current and supply direct current, renewable energy source (e.g., solar power or motion based power), or the like.

[0261] One or more aspects of at least one example may be implemented by representative instructions stored on at least one machine-readable medium which represents various logic within a processor, processor circuit, ASIC, or FPGA which when read by a machine, computing device or system causes the machine, computing device or system to fabricate logic to perform the techniques described herein. Such representations may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the processor, processor circuit, ASIC, or FPGA.

[0262] According to some examples, a computer-readable medium may include a non-transitory storage medium to store or maintain instructions that when executed by a machine, computing device or system, cause the machine, computing device or system to perform methods and/or operations in accordance with the described examples. The instructions may include any suitable type of code, such as source code, compiled code, interpreted code, executable code, static code, dynamic code, and the like. The instructions may be implemented according to a predefined computer language, manner or syntax, for instructing a machine, computing device or system to perform a certain function. The instructions may be implemented using any suitable high-level, low-level, object-oriented, visual, compiled and/or interpreted programming language.

[0263] Some examples may be described using the expression “in one example” or “an example” along with their derivatives. These terms mean that a particular feature, structure, or characteristic described in connection with the example is included in at least one example. The appearances of the phrase “in one example” in various places in the specification are not necessarily all referring to the same example.

[0264] Some examples may be described using the expression “coupled” and “connected” along

with their derivatives. These terms are not necessarily intended as synonyms for each other. For example, descriptions using the terms “connected” and/or “coupled” may indicate that two or more elements are in direct physical or electrical contact with each other. The term “coupled,” however, may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

[0265] The following examples pertain to additional examples of technologies disclosed herein.

[0266] Example 1. An example apparatus may include circuitry at a first device coupled with a host compute device. The circuitry may determine whether a memory request received via a first link coupled with the host compute device includes a memory address to access a memory at the first device or includes a memory address to access a memory at a second device coupled with the first device via a second link included in a multi die fabric. The circuitry may also cause a memory access to the memory address of the memory at the first device or the second device based on the determination. The circuitry may also send a response to the host compute device via the first link to indicate a status of the memory access to the memory address.

[0267] Example 2. The apparatus of example 1 may also include the circuitry to determine whether the memory request includes a memory address to access the memory at the first device or includes a memory address to access the memory at the second device based on a forwarding table maintained at the first device. For this example, the forwarding table may include entries that indicate whether the memory address included in the memory request is to the memory at the first device or is to the memory at the second device.

[0268] Example 3. The apparatus of example 1 may also include the circuitry to determine the memory request includes a memory address to access the memory at the second device. The circuitry may also forward the memory request to the second device via a second link coupled with the second device. The circuitry may also receive a response, via the second link, to the memory request from the second device. The circuitry may also include the response from the second device in the response sent to the host compute device in order to indicate a status of the access to the memory address of the memory at the second device.

[0269] Example 4. The apparatus of example 3, the memory address of the memory at the second device may include a first portion of a range of memory addresses of the memory at the second device. For this example, a second portion of the range of memory addresses may be included in a second memory address of the memory at the second device. The second device may receive a second memory request from the host compute device via a second link coupled between the second device and the host compute device. The second memory request may include a request to access the second memory address.

[0270] Example 5. The apparatus of example 1, wherein the first link comprises a CXL link and the second link included in the multi die fabric is a high speed internal-connect link having a data bandwidth of at least 5 times a data bandwidth of the CXL link.

[0271] Example 6. The apparatus of example 5, the memory at the first device may be a first HDM and the memory at the second device may be a second HDM.

[0272] Example 7. The apparatus of example 1 may also include compute circuitry that includes a graphics processing unit.

[0273] Example 8. An example method may include determining, by circuitry of a first device, whether a memory request received via a first link coupled with a host compute device includes a memory address for accessing a memory at the first device or includes a memory address for accessing a memory at a second device coupled with the first device via a second link included in a multi die fabric. The method may also include causing a memory access to the memory address of the memory at the first device or the second device based on the determination. The method may also include sending a response to the host compute device via the first link to indicate a status of the memory access to the memory address.

[0274] Example 9. The method of example 8, determining whether the memory request includes a

memory address to access the memory at the first device or includes a memory address to access the memory at the second device further may include determining based on a forwarding table maintained at the first device, the forwarding table to include entries that indicate whether the memory address included in the memory request is to the memory at the first device or is to the memory at the second device.

[0275] Example 10. The method of example 8 may also include determining the memory request includes a memory address to access the memory at the second device. The method may also include forwarding the memory request to the second device via a second link coupled with the second device. The method may also include receiving a response, via the second link, to the memory request from the second device and including the response from the second device in the response sent to the host compute device in order to indicate a status of the access to the memory address of the memory at the second device.

[0276] Example 11. The method of example 10, the memory address of the memory at the second device may be a first portion of a range of memory addresses of the memory at the second device. For this example, a second portion of the range of memory addresses may be included in a second memory address of the memory at the second device. The second device may receive a second memory request from the host compute device via a second link coupled between the second device and the host compute device. The second memory request may include a request to access the second memory address.

[0277] Example 12. The method of example 8, the first link may be a CXL link and the second link included in the multi die fabric may be a high speed internal-connect link having a data bandwidth of at least 5 times a data bandwidth of the CXL link.

[0278] Example 13. The method of example 12, the memory at the first device may be a first HDM and the memory at the second device may be a second HDM.

[0279] Example 14. An example at least one machine readable medium may include a plurality of instructions that in response to being executed by a system may cause the system to carry out a method according to any one of examples 8 to 13.

[0280] Example 15. An example apparatus may include means for performing the methods of any one of examples 8 to 13.

[0281] Example 16. An example at least one non-transitory computer-readable storage medium, may include a plurality of instructions, that when executed, may cause circuitry at a device coupled with a host device to determine whether a memory request received via a first link coupled with the host compute device includes a memory address to access a memory at the device or includes a memory address to access a memory at a second device coupled with the device via a second link included in a multi die fabric. The instructions may also cause the circuitry to cause a memory access to the memory address of the memory at the device or the second device based on the determination. The instructions may also cause the circuitry to send a response to the host compute device via the first link to indicate a status of the memory access to the memory address.

[0282] Example 17. The at least one non-transitory computer-readable storage medium of example 16, the instructions may further cause the circuitry to determine whether the memory request includes a memory address to access the memory at the device or includes a memory address to access the memory at the second device based on a forwarding table maintained at the device. For this example, the forwarding table may include entries that indicate whether the memory address included in the memory request is to the memory at the device or is to the memory at the second device.

[0283] Example 18. The at least one non-transitory computer-readable storage medium of example 16, the instructions may further cause the circuitry to determine the memory request includes a memory address to access the memory at the second device. The instructions may also cause the circuitry to forward the memory request to the second device via a second link coupled with the second device. The instructions may also cause the circuitry to receive a response, via the second

link, to the memory request from the second device. The instructions may also cause the circuitry to include the response from the second device in the response sent to the host compute device in order to indicate a status of the access to the memory address of the memory at the second device.

[0284] Example 19. The at least one non-transitory computer-readable storage medium of example 18, the memory address of the memory at the second device may be a first portion of a range of memory addresses of the memory at the second device. For this example, a second portion of the range of memory addresses may be included in a second memory address of the memory at the second device, the second device to receive a second memory request from the host compute device via a second link coupled between the second device and the host compute device, the second memory request to include a request to access the second memory address.

[0285] Example 20. The at least one non-transitory computer-readable storage medium of example 16, the first link may be a CXL link and the second link included in the multi die fabric may be a high speed internal-connect link having a data bandwidth of at least 5 times a data bandwidth of the CXL link.

[0286] Example 21. The at least one non-transitory computer-readable storage medium of example 20, the memory at the device may be a first HDM and the memory at the second device may be a second HDM.

[0287] Example 22. The at least one non-transitory computer-readable storage medium of example 16, the device may include a graphics processing unit to function as compute circuitry for the device.

[0288] Example 23. An example at least one non-transitory computer-readable storage medium, may include a plurality of instructions, that when executed by a system at a host compute device may cause the system to initialize a plurality of devices coupled with the host compute device via separate host links. The instructions may also cause the system to access, via the host links, registers at each device to gather information on a device's capability to route a memory request received via a host link to other devices of the plurality of devices. The memory request may be routed via one of multiple HSILs that couple the plurality of devices together. The instructions may also cause the system to build a system memory address mapping based on the gathered information. The system memory address mapping may be for use by the host compute device to access a memory address for memory at a respective device from among the plurality of devices. The instructions may also cause the system to cause separate forwarding tables to be maintained at each device from among the plurality of devices, the separate forwarding tables to indicate a device's capability to route a memory request to access a memory address of a memory at another device from among the plurality of devices. The memory request may be received via a host link coupled with the host compute device and forwarded to the other device via an HSIL based on the device's forwarding table.

[0289] Example 24. The at least one non-transitory computer-readable storage medium of example 23, the separate forwarding tables to be maintained at each device may include the separate forwarding tables maintained in respective registers at each device.

[0290] Example 25. The at least one non-transitory computer-readable storage medium of example 23, the host links may be CXL links and the HSILs that couple the plurality of devices together may be included in a multi die fabric. Each HSIL may have a data bandwidth of at least 5 times a data bandwidth of a single CXL link.

[0291] Example 26. The at least one non-transitory computer-readable storage medium of example 25, memory at each respective device may be HDM.

[0292] Example 27. The at least one non-transitory computer-readable storage medium of example 26, the system may be a basic input/output system (BIOS) for the host compute device. For these examples, the instructions may further cause the BIOS to fill, based on the information gathered from device registers, a CBHS for use by an OS of the host compute device to build a mapping table for each HDM at each respective device of the plurality of devices. The mapping table may

facilitate access by the OS to a memory address of an HDM via multiple CXL links. One of the multiple CXL links may be coupled to the device having the HDM and the remaining CXL links of the multiple CXL links may be coupled to other devices from among the plurality of devices. [0293] Example 28. The at least one non-transitory computer-readable storage medium of example 23, the plurality of devices may each include a graphics processing unit to function as compute circuitry.

[0294] Example 29. An example at least one non-transitory computer-readable storage medium may include a plurality of instructions, that when executed by a system at a host compute device may cause the system to receive a request to access a memory address of a memory at a first device from among a plurality of devices coupled with the host compute device via separate host links. The instructions may also cause the system to obtain memory address mapping information to determine how to split the memory address into multiple portions that include at least a first portion and a second portion. The instructions may also cause the system to send the first portion of the memory address to the first device in a first memory request message via a first host link to access the first portion of the memory address of the memory at the first device. The instructions may also cause the system to send, in a second memory request via a second host link, the second portion of the memory address to a second device from among the plurality of devices. The second memory request may be forwarded to the first device via a first HSIL to access the second portion of the memory address of the memory at the first device.

[0295] Example 30. The at least one non-transitory computer-readable storage medium of example 29, may further include the instructions to cause the system to split the memory address into multiple portions that include a third portion. For these examples, system may send, in a third memory request via a third host link, the third portion of the memory address to a third device from among the plurality of devices. The third memory request may be forwarded to the first device via a second HSIL to access the third portion of the memory address of the memory at the first device.

[0296] Example 31. The at least one non-transitory computer-readable storage medium of example 30, the instructions may also cause the system to receive, via the first host link, a first response from the first device for the first memory request, the first response to indicate a status of the access to the first portion of the memory address of the memory at the first device. The instructions may also cause the system to receive, via the second host link, a second response from the second device for the second memory request, the second response to indicate a status of the access to the second portion of the memory address of the memory at the first device. The instructions may also cause the system to receive, via the third host link, a third response from the third device for the third memory request, the third response to indicate a status of the access to the third portion of the memory address of the memory at the first device.

[0297] Example 32. The at least one non-transitory computer-readable storage medium of example 29, the separate host links may be CXL links and the first HSIL may be included in a multi die fabric that includes multiple HSILs that couple the plurality of devices to each other. Each HSIL of the multiple HSILs may have a data bandwidth of at least 5 times a data bandwidth of a single CXL link.

[0298] Example 33. The at least one non-transitory computer-readable storage medium of example 32, the memory at the first device may be HDM.

[0299] Example 34. The at least one non-transitory computer-readable storage medium of example 29, the plurality of devices may each include a graphics processing unit to function as compute circuitry.

[0300] It is emphasized that the Abstract of the Disclosure is provided to comply with 37 C.F.R. Section 1.72(b), requiring an abstract that will allow the reader to quickly ascertain the nature of the technical disclosure. It is submitted with the understanding that it will not be used to interpret or limit the scope or meaning of the claims. In addition, in the foregoing Detailed Description, it can be seen that various features are grouped together in a single example for the purpose of

streamlining the disclosure. This method of disclosure is not to be interpreted as reflecting an intention that the claimed examples require more features than are expressly recited in each claim. Rather, as the following claims reflect, inventive subject matter lies in less than all features of a single disclosed example. Thus, the following claims are hereby incorporated into the Detailed Description, with each claim standing on its own as a separate example. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein,” respectively. Moreover, the terms “first,” “second,” “third,” and so forth, are used merely as labels, and are not intended to impose numerical requirements on their objects.

[0301] Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing the claims.

Claims

1.-27. (canceled)

28. An apparatus comprising: circuitry at a first device coupled with a host compute device, the circuitry to: determine whether a memory request received via a first link coupled with the host compute device includes a memory address to access a memory at the first device or includes a memory address to access a memory at a second device coupled with the first device via a second link included in a multi die fabric; cause a memory access to the memory address of the memory at the first device or the second device based on the determination; and send a response to the host compute device via the first link to indicate a status of the memory access to the memory address.

29. The apparatus of claim 28, further comprising the circuitry to determine whether the memory request includes a memory address to access the memory at the first device or includes a memory address to access the memory at the second device based on a forwarding table maintained at the first device, the forwarding table to include entries that indicate whether the memory address included in the memory request is to the memory at the first device or is to the memory at the second device.

30. The apparatus of claim 28, further comprising the circuitry to: determine the memory request includes a memory address to access the memory at the second device; forward the memory request to the second device via a second link coupled with the second device; receive a response, via the second link, to the memory request from the second device; and include the response from the second device in the response sent to the host compute device in order to indicate a status of the access to the memory address of the memory at the second device.

31. The apparatus of claim 30, the memory address of the memory at the second device comprises a first portion of a range of memory addresses of the memory at the second device, wherein a second portion of the range of memory addresses is included in a second memory address of the memory at the second device, the second device to receive a second memory request from the host compute device via a second link coupled between the second device and the host compute device, the second memory request to include a request to access the second memory address.

32. The apparatus of claim 28, wherein the first link comprises a Compute Express Link (CXL) link and the second link included in the multi die fabric is a high speed internal-connect link having a data bandwidth of at least 5 times a data bandwidth of the CXL link.

33. The apparatus of claim 32, wherein the memory at the first device comprises a first host-managed device memory (HDM) and the memory at the second device comprises a second HDM.

34. The apparatus of claim 28, further comprising: compute circuitry that includes a graphics processing unit.

35. A method comprising: determining, by circuitry of a first device, whether a memory request

received via a first link coupled with a host compute device includes a memory address for accessing a memory at the first device or includes a memory address for accessing a memory at a second device coupled with the first device via a second link included in a multi die fabric; causing a memory access to the memory address of the memory at the first device or the second device based on the determination; and sending a response to the host compute device via the first link to indicate a status of the memory access to the memory address.

36. The method of claim 35, wherein determining whether the memory request includes a memory address to access the memory at the first device or includes a memory address to access the memory at the second device further comprises determining based on a forwarding table maintained at the first device, the forwarding table to include entries that indicate whether the memory address included in the memory request is to the memory at the first device or is to the memory at the second device.

37. The method of claim 35, further comprising: determining the memory request includes a memory address to access the memory at the second device; forwarding the memory request to the second device via a second link coupled with the second device; receiving a response, via the second link, to the memory request from the second device; and including the response from the second device in the response sent to the host compute device in order to indicate a status of the access to the memory address of the memory at the second device.

38. The method of claim 37, the memory address of the memory at the second device comprises a first portion of a range of memory addresses of the memory at the second device, wherein a second portion of the range of memory addresses is included in a second memory address of the memory at the second device, the second device to receive a second memory request from the host compute device via a second link coupled between the second device and the host compute device, the second memory request to include a request to access the second memory address.

39. The method of claim 35, wherein the first link comprises a Compute Express Link (CXL) link and the second link included in the multi die fabric is a high speed internal-connect link having a data bandwidth of at least 5 times a data bandwidth of the CXL link, and wherein the memory at the first device comprises a first host-managed device memory (HDM) and the memory at the second device comprises a second HDM.

40. At least one non-transitory computer-readable storage medium, comprising a plurality of instructions, that when executed by a system at a host compute device cause the system to: initialize a plurality of devices coupled with the host compute device via separate host links; access, via the host links, registers at each device to gather information on a device's capability to route a memory request received via a host link to other devices of the plurality of devices, the memory request to be routed via one of multiple high speed internal-connect links (HSILs) that couple the plurality of devices together; build a system memory address mapping based on the gathered information, the system memory address mapping for use by the host compute device to access a memory address for memory at a respective device from among the plurality of devices; and cause separate forwarding tables to be maintained at each device from among the plurality of devices, the separate forwarding tables to indicate a device's capability to route a memory request to access a memory address of a memory at another device from among the plurality of devices, the memory request received via a host link coupled with the host compute device and forwarded to the other device via an HSIL based on the device's forwarding table.

41. The at least one non-transitory computer-readable storage medium of claim 40, wherein the separate forwarding tables to be maintained at each device comprises the separate forwarding tables maintained in respective registers at each device.

42. The at least one non-transitory computer-readable storage medium of claim 41, wherein the host links comprise Compute Express Link (CXL) links and the HSILs that couple the plurality of devices together are included in a multi die fabric, each HSIL having a data bandwidth of at least 5 times a data bandwidth of a single CXL link, and, wherein memory at each respective device

comprises host-managed device memory (HDM).

43. The at least one non-transitory computer-readable storage medium of claim 42, wherein the system comprises a basic input/output system (BIOS) for the host compute device, the instructions to further cause the BIOS to: fill, based on the information gathered from device registers, a CXL binding virtual HDM (vHDM) structure (CBHS) for use by an operating system (OS) of the host compute device to build a mapping table for each HDM at each respective device of the plurality of devices, the mapping table to facilitate access by the OS to a memory address of an HDM via multiple CXL links, wherein one of the multiple CXL links is coupled to the device having the HDM and the remaining CXL links of the multiple CXL links are coupled to other devices from among the plurality of devices.

44. At least one non-transitory computer-readable storage medium, comprising a plurality of instructions, that when executed by a system at a host compute device cause the system to: receive a request to access a memory address of a memory at a first device from among a plurality of devices coupled with the host compute device via separate host links; obtain memory address mapping information to determine how to split the memory address into multiple portions that include at least a first portion and a second portion; send the first portion of the memory address to the first device in a first memory request message via a first host link to access the first portion of the memory address of the memory at the first device; and send, in a second memory request via a second host link, the second portion of the memory address to a second device from among the plurality of devices, the second memory request to be forwarded to the first device via a first high speed internal-connect link (HSIL) to access the second portion of the memory address of the memory at the first device.

45. The at least one non-transitory computer-readable storage medium of claim 44, further comprising the instructions to cause the system to split the memory address into multiple portions that include a third portion, wherein the system is to: send, in a third memory request via a third host link, the third portion of the memory address to a third device from among the plurality of devices, the third memory request to be forwarded to the first device via a second HSIL to access the third portion of the memory address of the memory at the first device.

46. The at least one non-transitory computer-readable storage medium of claim 45, further comprising the instructions to cause the system to: receive, via the first host link, a first response from the first device for the first memory request, the first response to indicate a status of the access to the first portion of the memory address of the memory at the first device; receive, via the second host link, a second response from the second device for the second memory request, the second response to indicate a status of the access to the second portion of the memory address of the memory at the first device; and receive, via the third host link, a third response from the third device for the third memory request, the third response to indicate a status of the access to the third portion of the memory address of the memory at the first device.

47. The at least one non-transitory computer-readable storage medium of claim 44, wherein the separate host links comprise Compute Express Link (CXL) links and the first HSIL is included in a multi die fabric that includes multiple HSILs that couple the plurality of devices to each other, each HSIL of the multiple HSILs having a data bandwidth of at least 5 times a data bandwidth of a single CXL link, and wherein the memory at the first device comprises host-managed device memory (HDM).
