

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284710

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Baker; Brandon S. et al.

CONTAINER COMPUTE PLATFORM

Abstract

A data platform for executing containers is provided. In some examples, the data platform receives an application from an application package of a provider account, the application including a setup script and a manifest of a service. The data platform activates access roles based on the manifest and creates the service and a compute pool using the setup script and a specification file accessed from the application package using an access role. The service is executed in the compute pool, accessing objects of the application package and of the data platform using the access roles.

Inventors: Baker; Brandon S. (Redmond, WA), Chen; Siyuan (Foster City, CA), Denny-Brown; Derek (Seattle, WA), Gray; Scott C. (Richmond Hill, CA), Kowalski; Jaroslaw (Redmond, WA), Manning; Mark M. (Penfield, NY)

Applicant: Snowflake Inc. (Bozeman, MT)

Family ID: 1000007756962

Appl. No.: 18/600321

Filed: March 08, 2024

Publication Classification

Int. Cl.: G06F16/28 (20190101); G06F21/57 (20130101)

U.S. Cl.:

CPC G06F16/283 (20190101); G06F21/577 (20130101); G06F2221/034 (20130101)

Background/Summary

TECHNICAL FIELD

[0001] Examples of the disclosure relate generally to data platforms and, more specifically, to executing containers.

BACKGROUND

[0002] Data platforms are widely used for data storage and data access in computing and communication contexts. With respect to architecture, a data platform could be an on-premises data platform, a network-based data platform (e.g., a cloud-based data platform), a combination of the two, and/or include another type of architecture. With respect to type of data processing, a data platform could implement online transactional processing (OLTP), online analytical processing (OLAP), a combination of the two, and/or another type of data processing. Moreover, a data platform could be or include a relational database management system (RDBMS) and/or one or more other types of database management systems. Users may develop applications that execute on data platforms.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0003] The present disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various examples of the disclosure.

[0004] FIG. 1 illustrates an example computing environment that includes a network-based data platform in communication with a cloud storage provider user system, according to some examples.

[0005] FIG. 2 is a block diagram illustrating components of a compute service manager, according to some examples.

[0006] FIG. 3 is a block diagram illustrating components of an execution platform, according to some examples.

[0007] FIG. 4A illustrates a container compute method **400** in accordance with one embodiment.

[0008] FIG. 4B illustrates a creation and execution of a container in a compute pool, according to some examples.

[0009] FIG. 5 is an illustration of a framework of an external access integration **520**, according to some examples.

[0010] FIG. 6 illustrates a framework for providing external endpoints by a service **614**, according to some examples.

[0011] FIG. 7 illustrates a diagrammatic representation of a machine in the form of a computer system within which a set of instructions may be executed for causing the machine to perform any one or more of the methodologies discussed herein, according to some examples.

DETAILED DESCRIPTION

[0012] Data platforms, which may be structured as on-premises or network-based systems like cloud-based data platforms, are utilized for a wide array of data storage and access operations. These platforms can support various data processing types, including Online Transactional Processing (OLTP), Online Analytical Processing (OLAP), or a combination thereof, and may comprise relational database management systems (RDBMS) or other database management systems.

[0013] A provider user may develop applications that execute on a data platform owned by a data platform operator within a context of a consumer user account. It is desirable to have the applications execute within an environment that protects the data and execution resources of the provider user, the consumer user, and the data platform operator.

[0014] It is challenging to protect intellectual property of the provider user while enabling the execution of provider applications within a consumer's environment. Prior solutions often required

a trade-off between exposing provider code to the consumer or hosting consumer data on provider systems, which could lead to potential security and privacy issues.

[0015] In some examples, a data platform allows for the secure execution of provider applications within a consumer's environment on the data platform. The data platform provides for the isolation of the provider user's code from the consumer user's data, enabling the provider user's applications to execute without compromising sensitive information.

[0016] In some examples, the data platform utilizes exclusive compute pools that are dedicated to the execution of service containers of an application provided by a provider user. These compute pools are isolated from other services and resources, providing a secure environment for service container execution. Access to these compute pools is managed through specific access roles that are exclusive to the service container and not associated with any user account.

[0017] In some examples, services and functions of a service container communicate with external networks through configurations controlled by the consumer account. This feature allows services to interact with external systems while adhering to the security protocols of the data platform.

[0018] In some examples, the deployment and configuration of services are automated through a setup script and a service specification file. The setup script provides instructions for creating necessary resources for service execution, while the service specification file details the configuration of the service, including container images and versions.

[0019] In some examples, a data platform receives an application from a provider account's application package, which includes a setup script and a manifest detailing a service. The data platform activates specific access roles as dictated by the manifest. The data platform uses the setup script and a specification file, obtained from the application package, to create both the service and a compute pool. The data platform then executes the service within the compute pool, with the service accessing objects from both the application package and the data platform, all under the governance of the activated access roles.

[0020] In some examples, a service accesses objects of a consumer account through permissions granted by a consumer user.

[0021] In some examples, the data platform creates a query warehouse using the setup script. The service uses this query warehouse, along with the access roles, to query objects within the data platform.

[0022] In some examples, the data platform stores the manifest, the setup script, and the specification file in a hidden repository within the provider account, making it inaccessible to the provider and the consumer account users.

[0023] In some examples, the data platform includes a service specification file that lists container images and corresponding versions, which are used by the service.

[0024] In some examples, the data platform scans the container images for security vulnerabilities before the service is executed within the compute pool.

[0025] In some examples, the data platform isolates the compute pool to prevent the service from accessing services or resources not specified in the application package.

[0026] In some examples, the data platform dedicates the compute pool exclusively to the service.

[0027] In some examples, the data platform enables the service to communicate with external networks through an external access integration, which is configured and controlled by the consumer account.

[0028] In some examples, the data platform defines network rules for the external access integration through the setup script, controlling the service's communication with external networks.

[0029] FIG. 1 illustrates an example computing environment **100** that includes a data platform **102** in communication with a client device **112**, according to some examples. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from

FIG. 1. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the computing environment **100** to facilitate additional functionality that is not specifically described herein.

[0030] As shown, the data platform **102** comprises a data storage **106**, a compute service manager **104**, an execution platform **110**, and a metadata database **114**. The data storage **106** comprises a plurality of computing machines and provides on-demand computer system resources such as data storage and computing power to the data platform **102**. As shown, the data storage **106** comprises multiple data storage devices, such as data storage device **108a**, data storage device **108b**, data storage device **108c**, and data storage device **108N**. In some examples, the data storage devices 1 to N are cloud-based storage devices located in one or more geographic locations. For example, the data storage devices 1 to N may be part of a public cloud infrastructure or a private cloud infrastructure. The data storage devices 1 to N may be hard disk drives (HDDs), solid state drives (SSDs), storage clusters, Amazon S3™ storage systems or any other data storage technology. Additionally, the data storage **106** may include distributed file systems (e.g., Hadoop Distributed File Systems (HDFS)), object storage systems, and the like.

[0031] The data platform **102** is used for reporting and analysis of integrated data from one or more disparate sources including the storage devices 1 to N within the data storage **106**. The data platform **102** hosts and provides data reporting and analysis services to multiple consumer accounts. Administrative users can create and manage identities (e.g., users, roles, and groups) and use privileges to allow or deny access to identities to resources and services. Generally, the data platform **102** maintains numerous consumer accounts for numerous respective consumers. The data platform **102** maintains each consumer account in one or more storage devices of the data storage **106**. Moreover, the data platform **102** may maintain metadata associated with the consumer accounts in the metadata database **114**. Each consumer account includes multiple objects with examples including users, roles, privileges, a datastores or other data locations.

[0032] The compute service manager **104** coordinates and manages operations of the data platform **102**. The compute service manager **104** also performs query optimization and compilation as well as managing clusters of compute services that provide compute resources (also referred to as “virtual warehouses”). The compute service manager **104** can support any number and type of clients such as end users providing data storage and retrieval requests, system administrators managing the systems and methods described herein, and other components/devices that interact with compute service manager **104**. As an example, the compute service manager **104** is in communication with the client device **112**. The client device **112** can be used by a user of one of the multiple consumer accounts supported by the data platform **102** to interact with and utilize the functionality of the data platform **102**. In some examples, the compute service manager **104** does not receive any direct communications from the client device **112** and only receives communications concerning jobs from a queue within the data platform **102**.

[0033] The compute service manager **104** is also coupled to metadata database **114**. The metadata database **114** stores data pertaining to various functions and examples associated with the data platform **102** and its users. In some examples, the metadata database **114** includes a summary of data stored in remote data storage systems as well as data available from a local cache. In some examples, the metadata database **114** may include information regarding how data is organized in remote data storage systems (e.g., the database storage **106**) and the local caches. In some examples, the metadata database **114** include data of metrics describing usage and access by provider users and consumers of the data stored on the data platform **102**. In some examples, the metadata database **114** allows systems and services to determine whether a piece of data needs to be accessed without loading or accessing the actual data from a storage device.

[0034] The compute service manager **104** is further coupled to the execution platform **110**, which provides multiple computing resources that execute various data storage and data retrieval tasks. The execution platform **110** is coupled to the database storage **106**. The execution platform **110**

comprises a plurality of compute nodes. A set of processes on a compute node executes a query plan compiled by the compute service manager **104**. The set of processes can include: a first process to execute the query plan; a second process to monitor and delete micro-partition files using a least recently used (LRU) policy and implement an out of memory (OOM) error mitigation process; a third process that extracts health information from process logs and status to send back to the compute service manager **104**; a fourth process to establish communication with the compute service manager **104** after a system boot; and a fifth process to handle all communication with a compute cluster for a given job provided by the compute service manager **104** and to communicate information back to the compute service manager **104** and other compute nodes of the execution platform **110**.

[0035] In some examples, communication links between elements of the computing environment **100** are implemented via one or more data communication networks. These data communication networks may utilize any communication protocol and any type of communication medium. In some examples, the data communication networks are a combination of two or more data communication networks (or sub-networks) coupled to one another. In alternate examples, these communication links are implemented using any type of communication medium and any communication protocol.

[0036] As shown in FIG. **1**, the data storage devices data storage device **108a** to data storage device **108N** are decoupled from the computing resources associated with the execution platform **110**. This architecture supports dynamic changes to the data platform **102** based on the changing data storage/retrieval needs as well as the changing needs of the users and systems. The support of dynamic changes allows the data platform **102** to scale quickly in response to changing demands on the systems and components within the data platform **102**. The decoupling of the computing resources from the data storage devices supports the storage of large amounts of data without requiring a corresponding large amount of computing resources. Similarly, this decoupling of resources supports a significant increase in the computing resources utilized at a particular time without requiring a corresponding increase in the available data storage resources.

[0037] The compute service manager **104**, metadata database **114**, execution platform **110**, and data storage **106** are shown in FIG. **1** as individual discrete components. However, each of the compute service manager **104**, metadata database **114**, execution platform **110**, and data storage **106** may be implemented as a distributed system (e.g., distributed across multiple systems/platforms at multiple geographic locations). Additionally, each of the compute service manager **104**, metadata database **114**, execution platform **110**, and data storage **106** can be scaled up or down (independently of one another) depending on changes to the requests received and the changing needs of the data platform **102**. Thus, in the described examples, the data platform **102** is dynamic and supports regular changes to meet the current data processing needs.

[0038] During operation, the data platform **102** processes multiple jobs determined by the compute service manager **104**. These jobs are scheduled and managed by the compute service manager **104** to determine when and how to execute the job. For example, the compute service manager **104** may divide the job into multiple discrete tasks and may determine what data is needed to execute each of the multiple discrete tasks. The compute service manager **104** may assign each of the multiple discrete tasks to one or more nodes of the execution platform **110** to process the task. The compute service manager **104** may determine what data is needed to process a task and further determine which nodes within the execution platform **110** are best suited to process the task. Some nodes may have already cached the data needed to process the task and, therefore, be a good candidate for processing the task. Metadata stored in the metadata database **114** assists the compute service manager **104** in determining which nodes in the execution platform **110** have already cached at least a portion of the data needed to process the task. One or more nodes in the execution platform **110** process the task using data cached by the nodes and, if necessary, data retrieved from the data storage **106**. It is desirable to retrieve as much data as possible from caches within the execution

platform **110** because the retrieval speed is typically faster than retrieving data from the data storage **106**.

[0039] As shown in FIG. **1**, the computing environment **100** separates the execution platform **110** from the data storage **106**. In this arrangement, the processing resources and cache resources in the execution platform **110** operate independently of the database storage devices data storage device **108a** to data storage device **108N** in the data storage **106**. Thus, the computing resources and cache resources are not restricted to a specific one of the data storage device **108a** to data storage device **108N**. Instead, all computing resources and all cache resources may retrieve data from, and store data to, any of the data storage resources in the data storage **106**.

[0040] FIG. **2** is a block diagram illustrating components of the compute service manager **104**, according to some examples. As shown in FIG. **2**, the compute service manager **104** includes an access manager **202**, and a key manager **204**. Access manager **202** handles authentication and authorization tasks for the systems described herein. Key manager **204** manages storage and authentication of keys used during authentication and authorization tasks. For example, access manager **202** and key manager **204** manage the keys used to access data stored in remote storage devices (e.g., data storage devices in data storage data storage device **206**). As used herein, the remote storage devices may also be referred to as “persistent storage devices” or “shared storage devices.”

[0041] In some examples, the access manager **202** operates within a data platform to control access to various objects of the data platform using Role-Based Access Control (RBAC). The access manager **202** is a component that manages authentication and authorization tasks, providing for authorized entities to access specific resources within the data platform. This component plays a role in maintaining the security and integrity of the data platform by enforcing access policies defined through RBAC.

[0042] In some examples, RBAC is implemented by defining roles within the data platform, where each role is associated with a specific set of permissions. These permissions determine the actions that entities assigned to the role can perform on various objects within the data platform. The access manager **202** utilizes these roles to make access control decisions, allowing or denying requests based on the roles assigned to the requesting entity and the permissions associated with those roles.

[0043] In some examples, a role includes a set of permissions or privileges assigned to an entity (which could be a user, a group, a service, or an automated process) that defines the level of access the entity has to various resources within the data platform **102**. Access roles are components of RBAC systems, which are used to restrict access to authorized users. Access roles control what actions an entity can perform, such as reading, writing, executing, or deleting data or objects within the data platform. Access roles help maintain security and operational integrity by ensuring that entities can only access the resources necessary for their function, thereby preventing unauthorized access or actions that could compromise the system.

[0044] In some examples, access roles are used to define the permissions that an application uses to operate correctly within the data platform environment. This can include the ability to create and manage compute resources, access databases, execute services, and interact with external systems. In some examples an access role encapsulates the security and access requirements of an application, and is granted to the application or its components (such as services or functions) rather than individual users. This allows for the management of permissions at the application level, simplifying the security model and making it easier to deploy and manage applications within the data platform **102**.

[0045] In some examples, access roles are granted to an application by a provider user that provides the application for deployment on the data platform **102**. A provider user granted access role allows the application to access objects owned by the provider user and associated with a provider user account.

[0046] In some examples, access roles are granted to an application by a consumer user that uses the application from within a consumer user account. A consumer user granted access role allows the application to access objects owned by the consumer user and associated with a consumer user account.

[0047] In some examples, access roles are granted to an application by the data platform **102**. A data platform granted access role allows the application to access objects of the platform **102** that are owned by the data platform **102**.

[0048] In some examples, a consumer access role is a type of role that is associated with an application within the data platform **102**. The consumer access role is granted by an application to a consumer user and used by the consumer user to access objects of the application such as, but not limited to, services, functions, data objects, and the like.

[0049] In some examples, the data platform creates specific access roles based on a manifest of an application received from an application package. These access roles are activated by the access manager **202** and are used to govern access to objects used by the application during operation. For example, an access role may grant the application the ability to create a compute pool and execute a service within that compute pool. The access manager **202** provides that an application, or entities authorized by the application, can perform actions permitted by the access role.

[0050] In some examples, the access manager **202** also controls access to objects of the data platform using the access roles during the execution of the service within the compute pool. The service accesses objects of the application package and of the data platform under the governance of the activated access roles. The access manager **202** checks the permissions associated with the access roles against the access requests made by the service, granting or denying these requests based on the defined RBAC policies.

[0051] In some examples, the role of the access manager **202** extends to managing access to hidden repositories within a provider account, where the application package is stored. The access manager **202** uses RBAC policies to restrict access to a hidden repository storing one or more components of the application and, providing for the one or more components of the application package to be accessible to entities with the appropriate access role. This mechanism protects the application package from unauthorized access, preserving the integrity of the provider's intellectual property.

[0052] In some examples, the access manager **202** implements RBAC roles to isolate the compute pool, preventing the service from accessing other services or resources not specified in the application package. This isolation is achieved by defining access roles that explicitly limit the service's permissions to the resources provided for the operation of the service, thereby enhancing the security of the service execution environment.

[0053] A request processing service **208** manages received data storage requests and data retrieval requests (e.g., jobs to be performed on database data). For example, the request processing service **208** may determine the data necessary to process a received query (e.g., a data storage request or data retrieval request). The data may be stored in a cache within the execution platform **110** or in a data storage device in data storage **106**.

[0054] A management console service **210** supports access to various systems and processes by administrators and other system managers. Additionally, the management console service **210** may receive a request to execute a job and monitor the workload on the system.

[0055] The compute service manager **104** also includes a job compiler **212**, a job optimizer **214**, and a job executor **216**. The job compiler **212** parses a job into multiple discrete tasks and generates the execution code for each of the multiple discrete tasks. The job optimizer **214** determines the best method to execute the multiple discrete tasks based on the data that needs to be processed. The job optimizer **214** also handles various data pruning operations and other data optimization techniques to improve the speed and efficiency of executing the job. The job executor **216** executes the execution code for jobs received from a queue or determined by the compute service manager

104.

[0056] A job scheduler and coordinator **218** sends received jobs to the appropriate services or systems for compilation, optimization, and dispatch to the execution platform **110**. For example, jobs may be prioritized and processed in that prioritized order. In some examples, the job scheduler and coordinator **218** determines a priority for internal jobs that are scheduled by the compute service manager **104** with other “outside” jobs such as user queries that may be scheduled by other systems in the database but may utilize the same processing resources in the execution platform **110**. In some examples, the job scheduler and coordinator **218** identifies or assigns particular nodes in the execution platform **110** to process particular tasks. A virtual warehouse manager **220** manages the operation of multiple virtual warehouses implemented in the execution platform **110**. As discussed below, each virtual warehouse includes multiple execution nodes that each include a cache and a processor.

[0057] Additionally, the compute service manager **104** includes a configuration and metadata manager **222**, which manages the information related to the data stored in the remote data storage devices and in the local caches (e.g., the caches in execution platform **110**). The configuration and metadata manager **222** uses the metadata to determine which data micro-partitions need to be accessed to retrieve data for processing a particular task or job. A monitor and workload analyzer **224** oversees processes performed by the compute service manager **104** and manages the distribution of tasks (e.g., workload) across the virtual warehouses and execution nodes in the execution platform **110**. The monitor and workload analyzer **224** also redistributes tasks, as needed, based on changing workloads throughout the data platform **102** and may further redistribute tasks based on a user (e.g., “external”) query workload that may also be processed by the execution platform **110**. The configuration and metadata manager **222** and the monitor and workload analyzer **224** are coupled to a data storage device **226**. Data storage device **226** in FIG. 2 represents any data storage device within the data platform **102**. For example, data storage device **226** may represent caches in execution platform **110**, storage devices in data storage **106**, or any other storage device.

[0058] The compute service manager **104** validates all communication from an execution platform (e.g., the execution platform **110**) to validate that the content and context of that communication are consistent with the task(s) known to be assigned to the execution platform. For example, an instance of the execution platform executing a query A should not be allowed to request access to data-source D (e.g., data storage device **226**) that is not relevant to query A. Similarly, a given execution node (e.g., execution node **304a**) may need to communicate with another execution node (e.g., execution node **304b**), and should be disallowed from communicating with a third execution node (e.g., execution node **316a**) and any such illicit communication can be recorded (e.g., in a log or other location). Also, the information stored on a given execution node is restricted to data relevant to the current query and any other data is unusable, rendered so by destruction or encryption where the key is unavailable.

[0059] The compute service manager **104** further comprises an anti-abuse scanner **228** that monitors creation of application packages created by content provider users of the data platform **102**. When a new application package is created by a content provider user, the anti-abuse scanner **228** scans the application package to determine if the application package contains content that is harmful, malicious, and the like. If such content is found, the anti-abuse scanner **228** prevents release of the application package by the content provider user.

[0060] In some examples, the anti-abuse scanner **228** is a component of another system that the compute service manager **104** communicates with via a network of the like.

[0061] FIG. 3 is a block diagram illustrating components of the execution platform **110**, according to some examples. As shown in FIG. 3, the execution platform **110** includes multiple virtual warehouses, including virtual warehouse **302a**, and virtual warehouse **302b** to virtual warehouse **302c**. Each virtual warehouse includes multiple execution nodes that each includes a data cache and a processor. The virtual warehouses can execute multiple tasks in parallel by using the multiple

execution nodes. As discussed herein, the execution platform **110** can add new virtual warehouses and drop existing virtual warehouses in real time based on the current processing needs of the systems and users. This flexibility allows the execution platform **110** to quickly deploy large amounts of computing resources when needed without being forced to continue paying for those computing resources when they are no longer needed. All virtual warehouses can access data from any data storage device (e.g., any storage device in data storage **106**).

[0062] Although each virtual warehouse shown in FIG. **3** includes three execution nodes, a particular virtual warehouse may include any number of execution nodes. Further, the number of execution nodes in a virtual warehouse is dynamic, such that new execution nodes are created when additional demand is present, and existing execution nodes are deleted when they are no longer necessary.

[0063] Each virtual warehouse is capable of accessing any of the data storage devices 1 to N shown in FIG. **1**. Thus, the virtual warehouses are not necessarily assigned to a specific data storage device 1 to N and, instead, can access data from any of the data storage devices 1 to N within the data storage **106**. Similarly, each of the execution nodes shown in FIG. **3** can access data from any of the data storage devices 1 to N. In some examples, a particular virtual warehouse or a particular execution node may be temporarily assigned to a specific data storage device, but the virtual warehouse or execution node may later access data from any other data storage device.

[0064] In the example of FIG. **3**, virtual warehouse **302a** includes a plurality of execution nodes as exemplified by execution node **304a**, execution node **304b**, and execution node **304c**. Execution node **304a** includes cache **306a** and a processor **308a**. Execution node **304b** includes cache **306b** and processor **308b**. Execution node **304c** includes cache **306c** and processor **308c**. Each execution node 1 to N is associated with processing one or more data storage and/or data retrieval tasks. For example, a virtual warehouse may handle data storage and data retrieval tasks associated with an internal service, such as a clustering service, a materialized view refresh service, a file compaction service, a storage procedure service, or a file upgrade service. In other implementations, a particular virtual warehouse may handle data storage and data retrieval tasks associated with a particular data storage system or a particular category of data.

[0065] Similar to virtual warehouse **302a** discussed above, virtual warehouse **302b** includes a plurality of execution nodes as exemplified by execution node **310a**, execution node **310b**, and execution node **310c**. Execution node **304a** includes cache **312a** and processor **314a**. Execution node **310b** includes cache **312b** and processor **314b**. Execution node **310c** includes cache **312c** and processor **314c**. Additionally, virtual warehouse **302c** includes a plurality of execution nodes as exemplified by execution node **316a**, execution node **316b**, and execution node **316c**. Execution node **316a** includes cache **318a** and processor **320a**. Execution node **316b** includes cache **318b** and processor **320b**. Execution node **316c** includes cache **318c** and processor **320c**.

[0066] In some examples, the execution nodes shown in FIG. **3** are stateless with respect to the data the execution nodes are caching. For example, these execution nodes do not store or otherwise maintain state information about the execution node or the data being cached by a particular execution node. Thus, in the event of an execution node failure, the failed node can be transparently replaced by another node. Since there is no state information associated with the failed execution node, the new (replacement) execution node can easily replace the failed node without concern for recreating a particular state.

[0067] Although the execution nodes shown in FIG. **3** each includes one data cache and one processor, alternate examples may include execution nodes containing any number of processors and any number of caches. Additionally, the caches may vary in size among the different execution nodes. The caches shown in FIG. **3** store, in the local execution node, data that was retrieved from one or more data storage devices in data storage **106**. Thus, the caches reduce or eliminate the bottleneck problems occurring in platforms that consistently retrieve data from remote storage systems. Instead of repeatedly accessing data from the remote storage devices, the systems and

methods described herein access data from the caches in the execution nodes, which is significantly faster and avoids the bottleneck problem discussed above. In some examples, the caches are implemented using high-speed memory devices that provide fast access to the cached data. Each cache can store data from any of the storage devices in the data storage **106**.

[0068] Further, the cache resources and computing resources may vary between different execution nodes. For example, one execution node may contain significant computing resources and minimal cache resources, making the execution node useful for tasks that require significant computing resources. Another execution node may contain significant cache resources and minimal computing resources, making this execution node useful for tasks that require caching of large amounts of data. Yet another execution node may contain cache resources providing faster input-output operations, useful for tasks that require fast scanning of large amounts of data. In some examples, the cache resources and computing resources associated with a particular execution node are determined when the execution node is created, based on the expected tasks to be performed by the execution node.

[0069] Additionally, the cache resources and computing resources associated with a particular execution node may change over time based on changing tasks performed by the execution node. For example, an execution node may be assigned more processing resources if the tasks performed by the execution node become more processor-intensive. Similarly, an execution node may be assigned more cache resources if the tasks performed by the execution node require a larger cache capacity.

[0070] Although virtual warehouses 1, 2, and N are associated with the same execution platform **110**, the virtual warehouses may be implemented using multiple computing systems at multiple geographic locations. For example, virtual warehouse 1 can be implemented by a computing system at a first geographic location, while virtual warehouses 2 and N are implemented by another computing system at a second geographic location. In some examples, these different computing systems are cloud-based computing systems maintained by one or more different entities.

[0071] Additionally, each virtual warehouse as shown in FIG. 3 has multiple execution nodes. The multiple execution nodes associated with each virtual warehouse may be implemented using multiple computing systems at multiple geographic locations. For example, an instance of virtual warehouse **302a** implements execution node **304a** and execution node **304b** on one computing platform at a geographic location and implements execution node **304c** at a different computing platform at another geographic location. Selecting particular computing systems to implement an execution node may depend on various factors, such as the level of resources needed for a particular execution node (e.g., processing resource requirements and cache requirements), the resources available at particular computing systems, communication capabilities of networks within a geographic location or between geographic locations, and which computing systems are already implementing other execution nodes in the virtual warehouse.

[0072] A particular execution platform **110** may include any number of virtual warehouses. Additionally, the number of virtual warehouses in a particular execution platform is dynamic, such that new virtual warehouses are created when additional processing and/or caching resources are needed. Similarly, existing virtual warehouses may be deleted when the resources associated with the virtual warehouse are no longer necessary.

[0073] In some examples, the virtual warehouses may operate on the same data in data storage **106**, but each virtual warehouse has its own execution nodes with independent processing and caching resources. This configuration allows requests on different virtual warehouses to be processed independently and with no interference between the requests. This independent processing, combined with the ability to dynamically add and remove virtual warehouses, supports the addition of new processing capacity for new users without impacting the performance observed by the existing users.

[0074] FIG. 4A illustrates an example container compute method **400** for controlling access of an

application during execution on a data platform and FIG. 4B illustrates creation and execution of a container in a compute pool, according to some examples. Although the example container compute method **400** depicts a particular sequence of operations, the sequence may be altered without departing from the scope of the present disclosure. For example, some of the operations depicted may be performed in parallel or in a different sequence that does not materially affect the function of the container compute method **400**. In other examples, different components of a data platform, such as data platform **102** of FIG. 1, that implement the container compute method **400** may perform functions at substantially the same time or in a specific sequence.

[0075] In operation **402**, a consumer account **452** receives from a provider account **492** an instantiation of an application package **474**. For a version of an application, the application version has a corresponding application package **474** including a version schema **486** having a manifest **462**, a setup script **464**, and one or more specification files, such as specification file **466**, associated with a service schema **434**. In some examples, an application **418** may have none or more service schemas. The service schema **434** includes one or more services **436** and functions **438**. In some examples, the version schema **486** of the application package **474** may include other application support files including, but not limited to Java Archive (JAR) files. Python modules, data files, configuration files such as XML or JSON files, executable binaries for various operating systems, shell scripts for automation tasks, property files for application settings, libraries or dependencies required by the application, SQL scripts for database setup or migration, documentation files such as READMEs or user guides, license files detailing the software's terms of use, and the like.

[0076] In some examples, the manifest **462** specifies one or more container images **488** of the service schema **434** and privileges granted to the application **418**. A setup script **464** specified by the manifest **462** defines the one or more services **436** and functions **438**. The manifest **462** and the setup script **464** enable automated deployment by providing a structured outline of a configuration, privileges, and resource dependencies of the application **418** and the one or more services **436** and functions **438** of the service schema **434**.

[0077] In some examples, the manifest **462** includes metadata such as version information, which assists in tracking and managing different iterations of the application package **474** and installed application **418** over time.

[0078] In some examples, the manifest **462** lists artifacts that are part of the application package **474**, such as the setup script **464** that includes instructions for creating and configuring the service schema **434**. In some examples, the setup script **464** includes execution parameters for a service **436** (of FIG. 4B) of the service schema **434**. The setup script **464** automates the creation of necessary database objects, compute resources, and other setup tasks thus streamlining the deployment process of the service schema **434**.

[0079] In some examples, the setup script **464** specifies one or more BRAC roles that are granted to the application **418** and to a consumer user who is the owner of the application **418** when the application **418** is installed in the consumer account **452**. For example, an application primary role **420** is a role that is managed by the data platform **102**. In some examples, the application primary role **420** is created implicitly and granted to the application **418** by the data platform **102** when the application **418** is installed in the consumer account **452**. All objects in the application **418** are owned by, and run as, the application primary role **420**. The application primary role **420** isolates the application **418** from the consumer user and vice versa. In some examples, the one or more services **436** and one or more functions **438** run as the primary application role **420** and inherit access to the components of the application **418** and any components of the data platform **102** to which the **418** has also been granted access.

[0080] In some examples, an application importer role **490** is managed by the data platform **102** and is granted to the application primary role **420** by the data platform **102**. In some examples, the application importer role **490** is created implicitly with the application **418** is installed in the

consumer account **452**. All permissions the consumer user gives to the application **418** are granted to the application importer role. This is how the application **418** is allowed access to the consumer user owned objects such as, but not limited to, the one or more tables **480** in the one or more consumer schemas **482** of the one or more consumer databases **484**.

[0081] In some examples, one or more consumer access roles **494** are created by the setup script **464** and are granted to the consumer user who is the owner of the installed application **418**. In some examples, the consumer user may regrant the consumer access roles **494** to other entities or objects within the execution environment. Via the setup script **464**, the application **418** grants privileges to objects within the application **418** to the consumer access roles to expose the objects to the consumer user.

[0082] In some examples, the application **418** uses the application primary role **420** to create one or more query warehouses **412**, create one or more compute pools **414**, and access one or more external access integrations **416** of the data platform **102**.

[0083] In some examples, the application primary role **420** is created and managed on behalf of the application **418** by a proprietor of the data platform **102** (of FIG. 1) that hosts the consumer account **452** and the provider account **492**. The application primary role **420** owns everything within the application and all code that executes in the application executes as the application primary role **420**.

[0084] In some examples, the proprietor of the data platform **102** grants the application primary role **402** to the application **418**.

[0085] In some examples, neither a consumer user having access to the consumer account **452** nor a provider user having access to the provider account **474** can be granted the application primary role **420**.

[0086] In some examples, neither a consumer user having access to the consumer account **452** nor a provider user having access to the provider account **474** can grant the application primary role **420** to the application **418**.

[0087] In some examples, in order for the application **418** to be able to create compute pools **414** and warehouses **412** (and other resources), a consumer user having access to the consumer account **452** grants these privileges to the application **418**, this grant is conveyed to the application primary role **420** via an application importer role **490**.

[0088] In some examples, the application **418** is granted an application importer role **490** used to access one or more consumer user objects such as, but not limited to, one or more consumer databases **484**, consumer schemas **482**, and consumer tables **480**.

[0089] In some examples, the application importer role **490** is created by a proprietor of the data platform **102** that hosts the consumer account **452** and the provider account **492**. The application importer role **490** is created as a hidden role on behalf of the application **418**. The application importer role **490** is automatically granted to the application primary role **420**. In order for the application **418** to access objects owned by the consumer user, the consumer grants access to the objects to the application **418** via the application importer role **490**.

[0090] By declaring these parameters upfront, the manifest **462** provides that the service **436** has the authorizations to function correctly once deployed. An example manifest is illustrated below:

[0091] manifest_version: 1 [0092] version: [0093] name: V1 [0094] artifacts: [0095] readme: readme.md [0096] setup_script: setup.sql [0097] default_web_endpoint: [0098] service: ux_schema.ux_service [0099] endpoint: ui [0100] images:—/database/schema/repository/image1 [0101] privileges: [0102] CREATE COMPUTE POOL description: “ . . . ” [0103] CREATE WAREHOUSE description: “ . . . ”

[0104] In some examples, a consumer user explicitly grants or approves requested privileges. These requested privileges are declared in the manifest both to drive UI that helps with this process, but also to allow the consumer user to understand what privileges the application will use should it be installed (e.g., a consumer user may decline to install an application if they are

unwilling to grant the application certain permissions).

[0105] In some examples, the specification file **466** specifies components and attributes of one or more container images **488** of the service schema **434**. The one or more container images **488** are given as one or more paths to a repository **470** associated with the versioned schema **486** of the application package **474** and used for execution of the service schema **434**. A portion of an example spec file is illustrated below: [0106] spec: [0107] container:—name: main [0108]

image:/database/schema/repository/image1 [0109] env: DATA_PLATFORM_WAREHOUSE: warehouse

[0110] In some examples, the setup script **464** is used to create and initiate execution of the one or more services **436** and functions **438**. The setup script **464** creates and configures the one or more services **436** and functions **438**.

[0111] In some examples, the setup script **464** defines one or more consumer access roles **494** that are granted to the application owner in the environment of the consumer user. The consumer user may then grant the one or more consumer access roles **494** to other entities or objects of the data platform. In some examples, any object that the setup script **464** associates to a consumer access role of the one or more consumer access roles **494** becomes visible to the consumer user. An example setup script is illustrated below: [0112] CREATE CONSUMER ACCESS ROLE

app_user; [0113] CREATE COMPUTE POOL “APP\$POOL_1” . . . ; [0114] CREATE

WAREHOUSE “APP\$QUERY_WH” . . . ; [0115] CREATE SCHEMA IF NOT EXISTS service;

[0116] GRANT USAGE ON SCHEMA service TO CONSUMER ACCESS ROLE app_user;

[0117] CREATE SERVICE service.app service [0118] IN COMPUTE POOL

“MY_APP\$POOL_1” [0119] QUERY_WAREHOUSE=“APP\$QUERY_WH” [0120]

SPECIFICATION_FILE=‘/spec. yaml’; [0121] CREATE FUNCTION service.test (x string) [0122]

RETURNS string [0123] SERVICE=service.app service [0124] ENDPOINT=test_endpoint AS

‘/test’; [0125] GRANT USAGE ON FUNCTION service.test (string) TO CONSUMER [0126]

APPLICATION ROLE app_user;

[0127] In some examples, the version schema **486** and the repository **470** are stored as hidden objects **472** within the provider account **492**. Once an object is stored as a hidden object, a provider user who created the application package **474** cannot change the contents of the version schema **486** or the repository **470**. However, the hidden version schema **486** and the contents of the repository **470**, including the one or more container images **488**, may be accessed by an entity using one or more access roles, such as version access role **460** and repository access role **468**.

[0128] In operation **404**, the data platform **102** defines one or more access roles based on the content of the application package **474**. For example, when the application package **474** is created by a provider user, the data platform **102** creates an empty repository, such as repository **470**, and an associated repository access role **468**. When a new application version is added to the application package **474**, the data platform **102** creates a version schema, such as version schema **486**, and named storage location it to hold files of the corresponding version of the application. The data platform **102** creates an access role to access these materials, thus an application version gets its own unique access role. When an application, such as application **418**, is accessed or executes code, the data platform **102** activates the appropriate roles, based upon the version of the application **418**, giving the application **418** access to the resources used by the application **418**.

[0129] In some examples, a shared content role **454** serves as a mechanism for managing access to shared resources within the application package **474** such as, but not limited to a shared schema **458** having one or more tables **456** that are shared by the provider account **492** with one or more consumer accounts. The shared content role **454** controls permissions and governs the visibility of shared content is visible in the application **418** as represented by shared schema **476** including one or more tables **478** The shared content role **454** is granted to specific components within installed applications of the application package **474** so that they may access shared content. The shared content role **454** provides that when the application **418** is installed in the consumer account **452**,

the shared resources are available according to the permissions defined by the provider user providing the application package **474**. The shared content role **454** is useful in multi-consumer user environments where applications access common resources without compromising the isolation and security of individual consumer accounts. By leveraging this role, the data platform **102** can provide a controlled and secure method for sharing content for collaborative scenarios and for applications that use common datasets or services.

[0130] In some examples, a version access role **460** manages access to resources tied to specific versions within an application package **474**. The version access role **460** provides that entities authorized to interact with a given version can do so, while maintaining isolation from other versions. The version access role **460** associates with resources such as schemas and scripts that are specific to a version of an application package **474**. It controls interactions with these resources during deployment and execution of the application **418** and the one or more services **436** and functions **438** of the application **418**. When updates occur or new versions of an application package are introduced, the version access role **460** facilitates access to the appropriate resources. The version access role **460** allows for application package lifecycle management within the data platform **102**, enabling activities such as upgrades and maintenance while allowing consumer accounts to engage with the correct version of an application and related services and functions as determined by the permissions associated with the version access role **460**.

[0131] In some examples, a repository access role **468** manages permissions for accessing container image repositories within the application package **474**. This role is assigned to entities that interact with repositories containing the images used by applications, services, or functions. The repository access role **468** is used to access stored container images **488** used to create the services **436** and functions **438** of the service schema **434**. The repository access role **468** governs the retrieval of images during the deployment and execution phases, ensuring that only authorized applications can pull the container images **488** for creation of the service schema **434**. The repository access role **468** maintains the security and integrity of the application package **474** by controlling access to its container images **488**.

[0132] In operation **406**, the data platform **102** uses the version access role **460** to access the specification file **466** of the versioned schema **486**. The data platform **102** executes the setup script **464** which creates the objects required by the application, which may include one or more compute pools **414**, query warehouses **412**, services **436** and functions **438** which may be utilized to access the services. When creating services **436**, the setup script **464** contains instructions to configure compute pools **414**, settings parameters and defining the environment in which services **436** will execute, and assigning services to the created compute pools. The specification file **466** is referenced during the creation of a service **436**, which indicates the container images and resources required for the service to execute. Additionally, the setup script **464** may define consumer access roles **494** with which it may grant access to the created services **436** and functions **438** to the application owner.

[0133] The access roles ensure that the one or more services **436** and functions **438** of the service schema **434** operate within the boundaries of their authorized capabilities, maintaining data security and integrity. The access roles act as gatekeepers, allowing the one or more services **436** and functions **438** of the service schema **434** to access only those data objects that they have been explicitly permitted to access, thus preventing unauthorized data access or manipulation.

[0134] In some examples, one or more access roles are configured to be exclusive to the one or more services **436** and functions **438** and do not have associations with either a consumer user of the consumer account **452** or a provider user of the provider account **492**. These exclusive roles are designed to facilitate specific interactions with and by the one or more services **436** and functions **438** of the service schema **434**, ensuring that operations performed are within the scope of the one or more services **436** and functions **438** and are not influenced by external user permissions or privileges.

[0135] In some examples, an application primary role, such as application primary role **420**, owns service objects, such as the one or more services **436**, of the application **418** which prevents a consumer user from influencing the internal operations of the application **418**.

[0136] In some examples, a consumer access role, such as a consumer access role of the one or more consumer access roles **494**, allows a consumer user using the consumer account **452** access to specified objects of the application **418** and scopes the access to those specified objects.

[0137] In some examples, the exclusivity of the one or more application roles to the one or more services **436** and functions **438** allows for a controlled execution environment where the integrity of the execution of the service schema **434** is maintained. This design prevents unauthorized access or modifications to the service schema **434** by users who do not have the necessary permissions, thereby upholding the security protocols established by the data platform **102** and protecting the integrity of the intellectual property of the provider user who created the application package **474**.

[0138] In some examples, the design further provides that objects contained within the application **418** that have not had access granted through consumer access roles, remain entirely hidden and inaccessible to the consumer user, protecting the provider's implementation and intellectual property.

[0139] In some examples, by restricting the one or more application roles to the objects of the service schema **434**, the data platform **102** can enforce a clear separation of duties and access rights. This separation is useful in multi-tenant environments where services may be deployed across various consumer accounts, each with different users and access requirements. The exclusive roles ensure that each service operates independently of users' actions, providing a reliable and consistent execution of the functions of the service schema **434** as provided for by the application package **474**.

[0140] In some examples, one or more compute pools **414** provide a dedicated environment for the execution of the service schema **434**. When containers are commingled in a single compute pool, the services of the provider user could “attack” the services of the consumer user such as, but not limited to, trying to exfiltrate data or proprietary logic and, similarly, the services of the consumer user could do the same to the provider. In contrast to multiple services sharing a compute pool, for each application, the data platform **102** restricts a compute pool to a single application. This provides that the containers for the application are isolated from the consumer's services.

[0141] The one or more compute pools **414** operate as account-level constructs where jobs and services are executed. Compute pools **414** furnish the computational resources necessary for the operation of long-running services and finite lifespan jobs.

[0142] The one or more compute pools **414** are created and configured by the setup script **464**. The setup script **464** is responsible for defining the parameters of the one or more compute pools **414**, such as the minimum and maximum number of nodes, thereby enabling the one or more compute pools **414** to scale according to the demands of the service schema **434**. The configuration of the one or more compute pools **414** directly influences the performance and scalability of the one or more services **436**. Once created, the one or more compute pools **414** host the one or more services **436** of the application **418** with access to CPUs, memory, and other system resources required for execution. The one or more services **436** and functions **438**, upon execution in the one or more compute pools **414**, are capable of accessing data objects within the data platform **102** according to the one or more access roles, utilizing the resources of the one or more compute pools **414** to perform operations. In some examples, the one or more compute pools **414** are designed to isolate the one or more services **436**, ensuring that the one or more services **436** execute within the compute pool **414** in a segregated manner from other services and resources not explicitly specified for the service schema **434** within the application **418**. This isolation provides for maintaining the integrity and security of the service schema **434** and the one or more services **436**.

[0143] In some examples, multiple services from different service schemas may be executed in the same compute pool. The segregation is between services in one app and all the other services

created by consumer user applications or different applications.

[0144] In some examples, the one or more query warehouses **412** serve as dedicated environments for executing SQL queries that are issued by the one or more services **436**. These warehouses are provisioned with the necessary compute resources to perform data processing tasks required by the one or more services **436** of the service schema **434**. In some examples, the setup script **464** may contain commands to instantiate the one or more query warehouses **412** with specific performance characteristics, such as warehouse size or compute power, tailored to the service schema **434**. The setup script **464** provides that the one or more query warehouses **412** are associated with the service schema **434**, allowing the service schema **434** to perform data operations within the context of the consumer account **452**.

[0145] Once the one or more query warehouses **412** are created, the one or more services **436** and functions **438** can begin querying one or more objects of the data platform **102**. The one or more services **436** utilize the query warehouses to execute SQL queries, leveraging the compute resources of the one or more query warehouses **412** to interact with databases, tables, or other data objects stored within the data platform **102**. The one or more consumer access roles **494**, which were created by the data platform **102** when the application was installed, govern the permissions used by the one or more services **436** and functions **438** of the service schema **434** to access and query these data objects.

[0146] An example command to create a service using the is illustrated below: [0147] CREATE SERVICE my_service [0148] IN COMPUTE POOL compute pool [0149] SPECIFICATION_FILE='specification_file.yaml' [0150] MIN_INSTANCES=2 [0151] MAX_INSTANCES=2;

[0152] The create command creates a service within a compute pool and a specification file is referenced from a versioned schema using a schema access role. A specification file includes details such as the name of a container image and a path to the container image to be used to create the service that is hosted by a compute pool. In some examples, the specification file includes the configuration for the service, ensuring that one or more containers of the service operate with the correct settings and parameters.

[0153] In some examples, a setup script defines a minimum and maximum number of instances for a service. The service is executed within a compute pool with defined minimum and maximum number of instances. This determines the scaling behavior of the service within the compute pool.

[0154] In some examples, a consumer user can grant permissions to an application **418** to create query warehouses. A query warehouse created by the application **418** is owned by the application **418**. The consumer user can take over ownership and control of the query warehouse to shut it down in an emergency.

[0155] In some examples, once the consumer user has granted an access role the ability to create a query warehouse to the application **418**, the consumer user cannot control the number of query warehouses that the application creates. In some examples, the consumer user has no control over the size of the query warehouses created.

[0156] In some examples, a consumer user can programmatically create a query warehouse on behalf of the application **418**. The consumer user can directly grant the query warehouse to the application **418**. In some examples, the consumer user uses an API to directly communicate the name of the query warehouse to the application **418**. In some examples, a query warehouse can be passed by reference to an application **418**.

[0157] In some examples, a consumer user can grant permission to an application to create compute pools. Once created, the application owns the compute pool and can automatically adjust a size or configuration of the compute pool based upon workload. In some examples, application-owned compute pools cannot host services or executable objects from other application instances or from the consumer user.

[0158] In some examples, a consumer user creates a compute pool on behalf of an application.

Such a compute pool is explicitly and exclusively dedicated to the application. In some examples, the compute pool is directly granted to the application for deploying services, using an API of the application.

[0159] In some examples, the architecture of the data platform incorporates stringent isolation measures within the compute pools to mitigate the risk of inter-service interference and potential security breaches. These measures are designed to prevent services that are executing concurrently in compute pools, from engaging in activities that could compromise the integrity of other services or lead to the unauthorized exposure of sensitive data or intellectual property.

[0160] In some examples, isolation is achieved through a combination of virtualization and access control mechanisms. Each container of a service runs in its own virtualized container, and a compute pool represents an isolated and virtualized environment in which these containers execute. One or more services may be hosted in a compute pool, but the containers of those services may directly communicate and see each other while services in other compute pools cannot do so without explicit permission through a RBAC system, effectively creating a sandbox that restricts the container of a service's operations to its own allocated resources. This virtualization provides that each container of a service has access to its own dedicated set of resources, such as CPU, memory, and storage, thereby eliminating the possibility of one container of a service consuming resources allocated to another container of a service. Furthermore, the data platform enforces access controls at the service level, which are defined by the one or more access roles associated with each service. These roles delineate the permissions and capabilities of the service, specifying which resources and data the service is authorized to access. The access roles are configured to prevent any service from accessing data or resources that are outside the scope of its defined permissions. In some examples, a capability of a service to access objects or resources is defined by access roles including, but not limited to, an application primary role **420**, an application importer role **490**, a version access role **460**, and a repository access role **568**. In some examples, one or more consumer access roles **494** define which entities can access the one or more services **436**.

[0161] For example, a service that is authorized to access a particular database within the data platform will be unable to access another database that it has not been granted permissions for. Similarly, if a service attempts to perform operations that could potentially lead to the exposure of IP, such as accessing the underlying container images or source code of another service, the access control mechanisms will block such operations.

[0162] In some examples, as a service accesses underlying images that the service was deployed from and operates on the source code, protection of the IP of the provider user is provided for through the use of access roles. For example, provider user IP protection is provided by preventing a consumer user accessing the underlying images, guaranteed by a repository access role **468** which is not grantable to the consumer user. In some examples, The consumer can't deploy another service in the compute pool exclusive to an application, preventing consumer created services from interfering with the services owned by the application.

[0163] In some examples, consumer IP protection is provided for by preventing services deployed by an application accessing container images or source code of another service owned by the consumer user if the consumer does not explicitly grant access to the application.

[0164] By implementing these isolation and access control measures, a data platform provides that services executing within compute pools are contained within their respective operational boundaries. This approach effectively prevents services from attacking each other, whether intentionally or unintentionally, and safeguards against the exposure of sensitive data and IP.

[0165] In operation **408**, the one or more services **436** are executed in the one or more compute pools **414** and the one or more services **436** access objects of the application package **474** and of the data platform **102** using the one or more access roles.

[0166] In some examples, a data platform scans container images of a service for security vulnerabilities before a service is executed in a compute pool.

[0167] FIG. 5 is an illustration of a framework of an external access integration **520**, according to some examples. A service or function establishes connections with external networks while maintaining security protocols to protect both consumer user's and a provider user's data and intellectual property. For example, a service or function of an application **514** uses an external access integration **520** that acts as a controlled gateway for the service or function to interact with external entities. The external access integration **520** is owned by an entity having a system role **508** and is configured with one or more sets of network rules, such as network rules **512** and **518** that dictate the permissible external endpoints of the service **436**. These network rules are similar to firewall settings, providing a filter that only allows traffic to and from external hosts, networks, or services in accordance with approved security protocols. Control over the external access integration **520** is maintained by specifying the allowed network rules, thereby ensuring that any data exchange with external systems adheres to the established security policies.

[0168] In some examples, a network access schema **510** is provided within the application's environment on the data platform **102** and is owned by an entity having a system role **508**. The network access schema **510** serves as a repository for one or more network rules **512** that are created and managed by an application primary role **504** of the application **514**. In some examples, the network rules **512** are not automatically active or enforceable. Instead, they are proposed rules that require approval **528** from a higher-level role with the appropriate authority, such as a process using an application owner role. Once a network rule is reviewed and approved **528**, the network rule can be referenced as part of referenced network rules **518** included in the configuration of the external access integration **520**, thereby becoming part of the active set of rules that govern usage **524** of the external access integration **520** by the application **514** or a service of the application **514**. In some examples, an entity having an application administrator role **502** can view the network access schema **510**.

[0169] In some examples, the application **514** creates the one or more network rules **512** in the network access schema **510**. Once created, these network rules may be deleted, but may not be changed, by the application **514**. The approval **528**, done by someone with the administrator role application administrator role **502**, involves referencing the network rule **512** by the external access integration **520** a referenced network rules **518**. This prevents the application **514** from changing a network rule of the network rules **512** after the network rule has been approved by an administrator.

[0170] In some examples, an application administrator role **502** is granted specific configuration permissions that allow configuration **522** and usage **526** of the external access integration **520**. The application administrator role **502** provides for defining the network rules **512** and network rules **518** within the external access integration **520**

[0171] In some examples, the external access integration **520** is a dynamic component that can be updated as the external communication requirements of services evolve. For example, as new external endpoints are deemed safe and necessary for a service's operation, the external access integration **520** can be reconfigured to include additional network rules. Conversely, if certain external endpoints are no longer required or are considered a security risk, they can be removed from the configuration of the external access integration **520**.

[0172] FIG. 6 illustrates a framework for providing public and private endpoints by a service **614**, according to some examples. The service **614** provides one or more endpoints **602** for use by an application owner **620**, application user **618**, or other entity having an appropriate access role such as consumer access role **622**. An endpoint serves as a designated interface for services, such as service **614**, to allow systems external to a data platform, such as data platform **102** (of FIG. 1), to interact with a service or for internal communication between different components of an application. The endpoint facilitates the flow of data in and out of the service **614**, enabling the service **614** to perform its intended functions.

[0173] In some examples, each endpoint has independent access control. This feature allows for granular control over who can access specific endpoints of a service, ensuring that only authorized

users or systems can interact with each endpoint. The ability to manage access control for each endpoint individually means that different endpoints of the same service can have distinct access policies. For instance, one endpoint of a service might be publicly accessible, serving as an API endpoint for general use, while another endpoint might be restricted to internal use within the data platform **102**, accessible only to certain roles or services. Independent access control feature enhances the security posture of services by allowing for a least privilege access model, where entities are granted only the permissions necessary to perform their tasks. It also provides flexibility in service design, enabling the creation of multi-faceted services with endpoints tailored to different user groups or use cases.

[0174] In some examples, access control to an endpoint is managed by defining roles that are scoped to the service. These roles encapsulate the permissions required to interact with the endpoint. For example, a service may have a service administrator role **608** used by an application administrator **604**. The service administrator role **608** is granted with permissions to access administrative endpoints for service management tasks, while a service user role **606** is used by an application user **618** and the service user role is granted with permissions to access an endpoint for using the service.

[0175] In some examples, a service user role has access to some endpoints for using the service to perform data operations, but has no access to the administrative endpoints.

[0176] In some examples, an application owner **620** might have permissions to access only the public API endpoints.

[0177] In some examples, endpoints are exposed to consumer users by granting service scoped roles to one or more consumer access roles **622**.

[0178] In some examples, an endpoint is entirely hidden from consumers and only indirectly exposed as user-defined functions in SQL.

[0179] In some examples, when an endpoint is created or configured, a setup script specifies which roles are allowed to access the endpoint. This specification is enforced by an access manager **202** (of FIG. 2), which checks the roles of users or services making requests to the endpoint and grants or denies access based on the configured policies.

[0180] In some examples, the endpoint is defined within a specification file of the application package of the application **418**. The specification file details the attributes of the endpoint, such as its name, the network port it listens on, the protocol it uses, and any public accessibility settings. For example, an endpoint might be configured to listen on port **8080** using the HTTP protocol and be marked as public, allowing it to receive web traffic from outside the data platform. This configuration enables the service **614** to serve web pages or APIs that can be consumed by users or other applications over the internet or within an internal network.

[0181] In some examples, endpoints can also be secured with various authentication and authorization mechanisms to ensure that only legitimate and authorized requests are processed. This might involve the use of API keys, OAuth tokens, or other security credentials that are verified before granting access to the service's functionality.

[0182] In some examples, the data platform provides tools and interfaces for managing endpoints, such as creating, updating, or deleting them as the requirements of the service evolve. This management is typically performed by roles with the necessary permissions, ensuring that changes to endpoints are controlled and auditable.

[0183] In some examples, the one or more endpoints **602** may be exposed to an application user **618** or may be used solely by the application **616** to provide the service **614**.

[0184] FIG. 7 illustrates a diagrammatic representation of a machine **700** in the form of a computer system within which a set of instructions may be executed for causing the machine **700** to perform any one or more of the methodologies discussed herein, according to examples. Specifically, FIG. 7 shows a diagrammatic representation of the machine **700** in the example form of a computer system, within which instructions **702** (e.g., software, a program, an application, an applet, an

application, or other executable code) for causing the machine **700** to perform any one or more of the methodologies discussed herein may be executed. For example, the instructions **702** may cause the machine **700** to execute any one or more operations of any one or more of the methods described herein. In this way, the instructions **702** transform a general, non-programmed machine into a particular machine **700** (e.g., the compute service manager **104**, the execution platform **110**, and the data storage devices 1 to N of data storage **106**) that is specially configured to carry out any one of the described and illustrated functions in the manner described herein.

[0185] In alternative examples, the machine **700** operates as a standalone device or may be coupled (e.g., networked) to other machines. In a networked deployment, the machine **700** may operate in the capacity of a server machine or a client machine in a server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine **700** may comprise, but not be limited to, a server computer, a client computer, a personal computer (PC), a tablet computer, a laptop computer, a netbook, a smart phone, a mobile device, a network router, a network switch, a network bridge, or any machine capable of executing the instructions **702**, sequentially or otherwise, that specify actions to be taken by the machine **700**. Further, while only a single machine **700** is illustrated, the term “machine” shall also be taken to include a collection of machines that individually or jointly execute the instructions **702** to perform any one or more of the methodologies discussed herein.

[0186] The machine **700** includes hardware processors **704**, memory **706**, and I/O components **708** configured to communicate with each other such as via a bus **710**. In some examples, the processors **704** (e.g., a central processing unit (CPU), a reduced instruction set computing (RISC) processor, a complex instruction set computing (CISC) processor, a graphics processing unit (GPU), a digital signal processor (DSP), an application-specific integrated circuit (ASIC), a radio-frequency integrated circuit (RFIC), another processor, or any suitable combination thereof) may include, for example, multiple processors as exemplified by processor **712** and a processor **714** that may execute the instructions **702**. The term “processor” is intended to include multi-core processors that may comprise two or more independent processors (sometimes referred to as “cores”) that may execute instructions **702** contemporaneously. Although FIG. 7 shows multiple processors **704**, the machine **700** may include a single processor with a single core, a single processor with multiple cores (e.g., a multi-core processor), multiple processors with a single core, multiple processors with multiple cores, or any combination thereof.

[0187] The memory **706** may include a main memory **732**, a static memory **716**, and a storage unit **718** including a machine storage medium **734**, all accessible to the processors **704** such as via the bus **710**. The main memory **732**, the static memory **716**, and the storage unit **718** store the instructions **702** embodying any one or more of the methodologies or functions described herein. The instructions **702** may also reside, completely or partially, within the main memory **732**, within the static memory **716**, within the storage unit **718**, within at least one of the processors **704** (e.g., within the processor's cache memory), or any suitable combination thereof, during execution thereof by the machine **700**.

[0188] The input/output (I/O) components **708** include components to receive input, provide output, produce output, transmit information, exchange information, capture measurements, and so on. The specific I/O components **708** that are included in a particular machine **700** will depend on the type of machine. For example, portable machines such as mobile phones will likely include a touch input device or other such input mechanisms, while a headless server machine will likely not include such a touch input device. It will be appreciated that the I/O components **708** may include many other components that are not shown in FIG. 7. The I/O components **708** are grouped according to functionality merely for simplifying the following discussion and the grouping is in no way limiting. In various examples, the I/O components **708** may include output components **720** and input components **722**. The output components **720** may include visual components (e.g., a display such as a plasma display panel (PDP), a light emitting diode (LED) display, a liquid crystal

display (LCD), a projector, or a cathode ray tube (CRT)), acoustic components (e.g., speakers), other signal generators, and so forth. The input components **722** may include alphanumeric input components (e.g., a keyboard, a touch screen configured to receive alphanumeric input, a photo-optical keyboard, or other alphanumeric input components), point-based input components (e.g., a mouse, a touchpad, a trackball, a joystick, a motion sensor, or another pointing instrument), tactile input components (e.g., a physical button, a touch screen that provides location and/or force of touches or touch gestures, or other tactile input components), audio input components (e.g., a microphone), and the like.

[0189] Communication may be implemented using a wide variety of technologies. The I/O components **708** may include communication components **724** operable to couple the machine **700** to a network **736** or devices **726** via a coupling **730** and a coupling **728**, respectively. For example, the communication components **724** may include a network interface component or another suitable device to interface with the network **736**. In further examples, the communication components **724** may include wired communication components, wireless communication components, cellular communication components, and other communication components to provide communication via other modalities. The devices **726** may be another machine or any of a wide variety of peripheral devices (e.g., a peripheral device coupled via a universal serial bus (USB)). For example, as noted above, the machine **700** may correspond to any one of the compute service manager **104**, the execution platform **110**, and the devices **726** may include the data storage device **226** or any other computing device described herein as being in communication with the data platform **102** or the data storage **106**.

[0190] The various memories (e.g., **706**, **716**, **732**, and/or memory of the processor(s) **704** and/or the storage unit **718**) may store one or more sets of instructions **702** and data structures (e.g., software) embodying or utilized by any one or more of the methodologies or functions described herein. These instructions **702**, when executed by the processor(s) **704**, cause various operations to implement the disclosed examples.

[0191] Described implementations of the subject matter can include one or more features, alone or in combination as illustrated below by way of example:

[0192] As used herein, the terms “machine-storage medium,” “device-storage medium,” and “computer-storage medium” mean the same thing and may be used interchangeably in this disclosure. The terms refer to a single or multiple storage devices and/or media (e.g., a centralized or distributed database, and/or associated caches and servers) that store executable instructions and/or data. The terms shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media, including memory internal or external to processors. Specific examples of machine-storage media, computer-storage media, and/or device-storage media include non-volatile memory, including by way of example semiconductor memory devices, e.g., erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), field-programmable gate arrays (FPGAs), and flash memory devices; magnetic disks such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The terms “machine-storage media,” “computer-storage media,” and “device-storage media” specifically exclude carrier waves, modulated data signals, and other such media, at least some of which are covered under the term “signal medium” discussed below.

[0193] In various examples, one or more portions of the network **736** may be an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a local-area network (LAN), a wireless LAN (WLAN), a wide-area network (WAN), a wireless WAN (WWAN), a metropolitan-area network (MAN), the Internet, a portion of the Internet, a portion of the public switched telephone network (PSTN), a plain old telephone service (POTS) network, a cellular telephone network, a wireless network, a Wi-Fi® network, another type of network, or a combination of two or more such networks. For example, the network **736** or a portion of the network **736** may include a wireless or

cellular network, and the coupling **730** may be a Code Division Multiple Access (CDMA) connection, a Global System for Mobile communications (GSM) connection, or another type of cellular or wireless coupling. In this example, the coupling **730** may implement any of a variety of types of data transfer technology, such as Single Carrier Radio Transmission Technology (1×RTT), Evolution-Data Optimized (EVDO) technology, General Packet Radio Service (GPRS) technology, Enhanced Data rates for GSM Evolution (EDGE) technology, third Generation Partnership Project (3GPP) including 3G, fourth generation wireless (4G) networks, fifth generation wireless (5G) networks, Universal Mobile Telecommunications System (UMTS), High-Speed Packet Access (HSPA), Worldwide Interoperability for Microwave Access (WiMAX), Long Term Evolution (LTE) standard, others defined by various standard-setting organizations, other long-range protocols, or other data transfer technology.

[0194] The instructions **702** may be transmitted or received over the network **736** using a transmission medium via a network interface device (e.g., a network interface component included in the communication components **724**) and utilizing any one of a number of well-known transfer protocols (e.g., hypertext transfer protocol (HTTP)). Similarly, the instructions **702** may be transmitted or received using a transmission medium via the coupling **728** (e.g., a peer-to-peer coupling) to the devices **726**. The terms “transmission medium” and “signal medium” mean the same thing and may be used interchangeably in this disclosure. The terms “transmission medium” and “signal medium” shall be taken to include any intangible medium that is capable of storing, encoding, or carrying the instructions **702** for execution by the machine **700**, and include digital or analog communications signals or other intangible media to facilitate communication of such software. Hence, the terms “transmission medium” and “signal medium” shall be taken to include any form of modulated data signal, carrier wave, and so forth. The term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal.

[0195] The various operations of example methods described herein may be performed, at least partially, by one or more processors that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Similarly, the methods described herein may be at least partially processor-implemented. For example, at least some of the operations of the methodologies disclosed herein may be performed by one or more processors. The performance of certain of the operations may be distributed among the one or more processors, not only residing within a single machine, but also deployed across a number of machines. In some examples, the processor or processors may be located in a single location (e.g., within a home environment, an office environment, or a server farm), while in other examples the processors may be distributed across a number of locations.

[0196] Described implementations of the subject matter can include one or more features, alone or in combination as illustrated below by way of example.

[0197] Example 1 is a machine-implemented method, comprising: receiving, by a consumer account of a data platform, an application from an application package of a provider account of the data platform, the application including a setup script and a manifest of a service provided by the application; activating one or more access roles based on the manifest; creating the service and a compute pool of the data platform using the setup script and a specification file accessed from the application package using an access role of the one or more access roles; and executing the service in the compute pool, the service accessing objects of the application package and of the data platform using the one or more access roles.

[0198] In Example 2, the subject matter of Example 1 includes, wherein the one or more access roles are exclusive to the service and are not associated with a consumer user of the consumer account or a provider user of the provider account.

[0199] In Example 3, the subject matter of any of Examples 1-2 includes, creating a query warehouse of the data platform using the setup script; and querying, by the service, one or more

objects of the data platform using the query warehouse and the one or more access roles.

[0200] In Example 4, the subject matter of any of Examples 1-3 includes, wherein one or more components of the application package are stored in a hidden repository within the provider account, the one or more components of the application package hidden from a provider user of the provider account.

[0201] In Example 5, the subject matter of any of Examples 1-4 includes, wherein the service specification file includes a list of container images and their corresponding versions to be used by the service.

[0202] In Example 6, the subject matter of any of Examples 1-5 includes, scanning the container images for security vulnerabilities before the service is executed in the compute pool.

[0203] In Example 7, the subject matter of any of Examples 1-6 includes, isolating the compute pool to prevent the service from accessing other services or resources not specified in the application package.

[0204] In Example 8, the subject matter of any of Examples 1-7 includes, wherein the compute pool is exclusively dedicated to the service.

[0205] In Example 9, the subject matter of any of Examples 1-8 includes, enabling the service to communicate with external networks through an external access integration that is configured and controlled by the consumer account.

[0206] In Example 10, the subject matter of any of Examples 1-9 includes, wherein the setup script defines one or more network rules for the external access integration that control communication between the service and the external networks.

[0207] Example 11 is at least one machine-readable medium including instructions that, when executed by processing circuitry, cause the processing circuitry to perform operations to implement any of Examples 1-10.

[0208] Example 12 is an apparatus comprising means to implement any of Examples 1-10.

[0209] Example 13 is a system to implement any of Examples 1-10.

[0210] Example 14 is a method to implement any of Examples 1-10.

[0211] Although the examples of the present disclosure have been described with reference to specific examples, it will be evident that various modifications and changes may be made to these examples without departing from the broader scope of the inventive subject matter. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof show, by way of illustration, and not of limitation, specific examples in which the subject matter may be practiced. The examples illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other examples may be used and derived therefrom, such that structural and logical substitutions and changes may be made without departing from the scope of this disclosure. This Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various examples is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

[0212] In this document, the terms “a” or “an” are used, as is common in patent documents, to include one or more than one, independent of any other instances or usages of “at least one” or “one or more.” In this document, the term “or” is used to refer to a nonexclusive or, such that “A or B” includes “A but not B,” “B but not A,” and “A and B,” unless otherwise indicated. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein.” Also, in the following claims, the terms “including” and “comprising” are open-ended; that is, a system, device, article, or process that includes elements in addition to those listed after such a term in a claim is still deemed to fall within the scope of that claim.

[0213] Such examples of the inventive subject matter may be referred to herein, individually and/or collectively, by the term “example” merely for convenience and without intending to voluntarily

limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific examples have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose may be substituted for the specific examples shown. This disclosure is intended to cover any and all adaptations or variations of various examples. Combinations of the above examples, and other examples not specifically described herein, will be apparent to those of skill in the art, upon reviewing the above description.

Claims

- 1.** A machine-implemented method, comprising: receiving, in a consumer account of a data platform, an application from an application package of a provider account of the data platform, the application including a setup script and a manifest of a service provided by the application; activating one or more access roles based on the manifest; creating the service and a compute pool of the data platform using the setup script and a specification file accessed from the application package using an access role of the one or more access roles; and executing the service in the compute pool, the executing of the service including accessing objects of the application package and of the data platform using the one or more access roles.
- 2.** The machine-implemented method of claim 1, wherein the one or more access roles are exclusive to the service and are not associated with a consumer user of the consumer account or a provider user of the provider account.
- 3.** The machine-implemented method of claim 1, further comprising: creating a query warehouse of the data platform using the setup script; and querying, by the service, one or more objects of the data platform using the query warehouse and the one or more access roles.
- 4.** The machine-implemented method of claim 1, wherein one or more components of the application package are stored in a hidden repository within the provider account, the one or more components of the application package hidden from a provider user of the provider account.
- 5.** The machine-implemented method of claim 1, wherein the service specification file includes a list of container images and corresponding versions to be used by the service.
- 6.** The machine-implemented method of claim 5, further comprising scanning the container images for security vulnerabilities before the service is executed in the compute pool.
- 7.** The machine-implemented method of claim 1, further comprising: isolating the compute pool to prevent the service from accessing other services or resources not specified in the application package.
- 8.** The machine-implemented method of claim 7, wherein the compute pool is exclusively dedicated to the service.
- 9.** The machine-implemented method of claim 1, further comprising: enabling the service to communicate with external networks through an external access integration that is configured and controlled by the consumer account.
- 10.** The machine-implemented method of claim 9, wherein the setup script defines one or more network rules for the external access integration that control communication between the service and the external networks.
- 11.** A system comprising: at least one hardware processor; and at least one memory storing instructions that, when executed by the at least one processor, cause the system to perform operations comprising: receiving, in a consumer account of a data platform, an application from an application package of a provider account of the data platform, the application including a setup script and a manifest of a service provided by the application; activating one or more access roles based on the manifest; creating the service and a compute pool of the data platform using the setup script and a specification file accessed from the application package using an access role of the one or more access roles; and executing the service in the compute pool, the executing of the service

including accessing objects of the application package and of the data platform using the one or more access roles.

12. The system of claim 11, wherein the one or more access roles are exclusive to the service and are not associated with a consumer user of the consumer account or a provider user of the provider account.
13. The system of claim 11, wherein the operations further comprise: creating a query warehouse of the data platform using the setup script; and querying, by the service, one or more objects of the data platform using the query warehouse and the one or more access roles.
14. The system of claim 11, wherein one or more components of the application package are stored in a hidden repository within the provider account, the one or more components of the application package hidden from a provider user of the provider account.
15. The system of claim 11, wherein the service specification file includes a list of container images and corresponding versions to be used by the service.
16. The system of claim 15, wherein the operations further comprise scanning the container images for security vulnerabilities before the service is executed in the compute pool.
17. The system of claim 11, wherein the operations further comprise: isolating the compute pool to prevent the service from accessing other services or resources not specified in the application package.
18. The system of claim 17, wherein the compute pool is exclusively dedicated to the service.
19. The system of claim 11, wherein the operations further comprise: enabling the service to communicate with external networks through an external access integration that is configured and controlled by the consumer account.
20. The system of claim 19, wherein the setup script defines one or more network rules for the external access integration that control communication between the service and the external networks.
21. A machine-storage medium storing instructions that, when executed by one or more processors, cause the one or more processors to perform operations comprising: receiving, in a consumer account of a data platform, an application from an application package of a provider account of the data platform, the application including a setup script and a manifest of a service provided by the application; activating one or more access roles based on the manifest; creating the service and a compute pool of the data platform using the setup script and a specification file accessed from the application package using an access role of the one or more access roles; and executing the service in the compute pool, the executing of the service including accessing objects of the application package and of the data platform using the one or more access roles.
22. The machine-storage medium of claim 21, wherein the one or more access roles are exclusive to the service and are not associated with a consumer user of the consumer account or a provider user of the provider account.
23. The machine-storage medium of claim 21, wherein the operations further comprise: creating a query warehouse of the data platform using the setup script; and querying, by the service, one or more objects of the data platform using the query warehouse and the one or more access roles.
24. The machine-storage medium of claim 21, wherein one or more components of the application package are stored in a hidden repository within the provider account, the one or more components of the application package hidden from a provider user of the provider account.
25. The machine-storage medium of claim 21, wherein the service specification file includes a list of container images and corresponding versions to be used by the service.
26. The machine-storage medium of claim 25, wherein the operations further comprise scanning the container images for security vulnerabilities before the service is executed in the compute pool.
27. The machine-storage medium of claim 21, wherein the operations further comprise: isolating the compute pool to prevent the service from accessing other services or resources not specified in the application package.

28. The machine-storage medium of claim 27, wherein the compute pool is exclusively dedicated to the service.

29. The machine-storage medium of claim 21, wherein the operations further comprise: enabling the service to communicate with external networks through an external access integration that is configured and controlled by the consumer account.

30. The machine-storage medium of claim 29, wherein the setup script defines one or more network rules for the external access integration that control communication between the service and the external networks.
