

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284738

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Das; Ishita et al.

STORING CONTEXTUAL DATA WITH CONTEXT SCHEMAS

Abstract

A context schema characterizing a context associated with an application is received by a processing device. The context schema specifies one or more context attributes. Each context attribute identifies a respective context element of the context associated with the application. A context item comprising one or more values is received. Each value corresponds to a respective context attribute of the one or more context attributes. The context item is validated against the context schema. Responsive to determining that the context item conforms to the context schema, the context item is stored in a context store.

Inventors: Das; Ishita (Bangalore, IN), Ravisundaram; Sundaram (Bangalore, IN)

Applicant: Twilio Inc. (San Francisco, CA)

Family ID: 95563869

Appl. No.: 19/095462

Filed: March 31, 2025

Related U.S. Application Data

parent US continuation 18401162 20231229 parent-grant-document US 12292927 child US 19095462

Publication Classification

Int. Cl.: G06F16/901 (20190101); G06F16/903 (20190101)

U.S. Cl.:

CPC G06F16/901 (20190101); G06F16/90335 (20190101);

Background/Summary

RELATED APPLICATION [0001] This application is a continuation of U.S. patent application Ser. No. 18/401,162, filed Dec. 29, 2023, which incorporated by reference herein.

TECHNICAL FIELD

[0002] Aspects and embodiments of the present disclosure relate to data storage layers, and in particular to storing contextual data with context schemas.

BACKGROUND

[0003] Applications may provide services to users, systems, other applications, etc. Applications may use data storage layers to store and retrieve data related to the application's tasks and services. Examples of data storage layers that may be used include databases (e.g., relational, non-relational), key-value stores, file systems, files, etc.,

Description

BRIEF DESCRIPTION OF DRAWINGS

[0004] Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:

[0005] FIG. 1 is a block diagram of an example system architecture for storing contextual data with context schemas, in accordance with an embodiment;

[0006] FIG. 2 is a block diagram of an example context schema in accordance with an embodiment;

[0007] FIG. 3 is a block diagram of an example context item, in accordance with an embodiment;

[0008] FIGS. 4A-B are sequence diagrams of an example interaction for storing contextual data with context schemas, in accordance with an embodiment;

[0009] FIGS. 5A-B are flow diagrams of an example method for storing contextual data with context schemas, in accordance with an embodiment; and

[0010] FIG. 6 illustrates an example computer system, in accordance with an embodiment.

DETAILED DESCRIPTION

[0011] Aspects of the present disclosure relate to storing contextual data with context schemas. Applications may provide services to users or other applications. Applications may use data storage layers to store and retrieve data related to the application's tasks and services. Examples of data storage layers that may be used include databases (e.g., relational, non-relational), key-value stores, file systems, files, etc. Context is a set of properties of a system or environment in which an application may operate. Examples of an application's context may include the requirements, constraints, and objectives of the application for which a data model is being designed. The context of an application influences the design of data storage layers so that the objectives of the application can be met. Examples of design decisions influenced by application context include a database type (e.g., relational or non-relational), and a database schema (e.g., structure of tables, columns, etc.).

[0012] Data storage solutions for an application may be designed and implemented before the application is deployed. For example, an application developer (e.g., a user, a team, an organization) may analyze the data to be stored, anticipated access frequencies and access activities (e.g., read-heavy or write-heavy), latency requirements, underlying technologies, costs, etc. to make design decisions related to the data storage solution. The application developer may select a data storage layer technology (e.g., a relational or non-relational database system) and may design a schema for the selected data storage layer technology. A schema is a configuration of a data storage layer based on the data to be stored. For example, a schema in a relational database may

define tables, fields, columns, and other elements corresponding to data structures of the application. In a non-relational database, a schema may define requirements for documents, records, etc. Schemas may define other features for other types of data storage systems. The application developer may test the data storage layer and fine-tune the schema and other storage system parameters to meet the design goals and expected use cases prior to deploying the application to clients (e.g., users, other applications, etc.).

[0013] The above-described systems may face several challenges related to data storage layer design and schemas. It may be difficult, for example, to restructure a schema to accommodate new types of data or new relationships between data that were not captured in the schema during the initial design (e.g., prior to deploying the application). Restructuring a schema may involve, for example, temporarily taking an application offline, copying data (e.g., terabytes or exabytes of data) from an old database to a new database, fine-tuning a new schema, or other disruptive activities. Restructuring a schema may thus be prohibitively expensive in terms of cost, resource consumption, developer time, etc. Even for more flexible data storage layers that do not require as much effort to restructure (e.g., some non-relational or schema-less database systems), changing a schema after deployment may incur some efficiency losses (e.g., latency, resource consumption) and may still require some restructuring the application to accommodate the changes in the underlying data storage layer.

[0014] Furthermore, it may be difficult for such systems to store contextual data of clients using the data storage layers and schemas previously described. Systems may interact with many clients, each having their own data storage solutions and their own context. Systems or applications providing services to various clients may be referred to herein as host systems or host applications. Users, systems, applications, etc. interacting with the host application may be referred to herein as clients, client applications, etc. Each application may have its own native data modeled after its own context (e.g., host-side native data and client-side native data). Clients may want to store contextual data in a host system's data storage layer to create connections and references between host-side native data and client-side native data, which may aid in searching the host system's native data by the contextual data and aid in keeping native data updated in both host-side and client-side data storage layers, among other benefits. However, contextual data may be different for different clients, as each client may have its own type of native data, its own data storage layer technology and schema, etc. There may have been no collaboration between the host application developer and clients on the data models during design time (e.g., clients may be unknown to the host application developer at that time). Thus, it can be challenging to design a data storage layer schema prior to deployment that can accommodate a wide variety of contextual data after deployment. For example, a new client may be introduced after host application deployment, and the new client may have contextual data that wasn't foreseen during design of the host application data storage layer. The host-side and client-side data storage layers may thus be incompatible, and it may be difficult to share and query data between the data storage layers. This may necessitate a restructuring of the host-side data storage layer schema to accommodate the new client's contextual data in the above-described systems.

[0015] As a result of these and other challenges, application developers and clients may undertake costly solutions and workarounds to provide contextual data to a host application. As previously described, a solution may involve restructuring the host application data storage layer to accommodate the client's contextual data. Another solution may involve developing a translation layer or other additional application infrastructure to supplement the data shortcomings in the host data storage layer. These solutions may result in increased costs, downtime, and latency for data storage layers, increased developer time, and other inconveniences for application developers and clients.

[0016] Aspects of the present disclosure address the above and other deficiencies by storing contextual data in association with respective user-defined context schemas. In an embodiment, a

context schema defines the structure of an application context, including identifiers of context items, their types, and their relationships. A context schema may be represented in a data structure (e.g., a document, a data type, etc.), which may be stored in a schema store of a host application along with other context schemas of other clients. A context schema may characterize contextual data of a client application, a client application's data storage layer, etc. with one or more context attributes. A context attribute may correspond to an element of native data of the client data storage layer (e.g., a column, field, etc.) and may provide various indications to the host application and its underlying data storage layer. For example, a context attribute may indicate whether a context element is required in a context item, should be searchable, etc. Context attributes may be used to generate and validate context items. A context item, generated by a client application for providing to a host application, may describe context for a native data item of the host application. The native data item of the host application may be generated, for example, as part of a transaction between the client and host during which the client provides the context item. The context item may include instances of context attributes of the corresponding context schema, which may correspond to values of a particular native data item of the client application. For example, a record or row in the client database (the values of which correspond to the columns, fields, or other elements indicated in the context attributes) may be included in the context item. A context item may similarly be represented in a data structure, which may be stored in a context store of the host application. As previously described with reference to other stores and storage layers, the context store may have its own schema (e.g., created at design time), which may enable uniform storage of context items from various clients. Thus, context items may conform to the schema of the context store as well as the corresponding context schema. Context schemas and context items may also conform to different versions of store schemas and context schemas, which may be indicated by version information stored in context schemas and context items.

[0017] Context schemas and context items may be received at the host application from one or more clients in addition to host-side native data items. The host application (or its data storage layer or other constituent component) may validate the context items against corresponding context schemas and may link the context items to native data items of the host application. Thus, a context item can provide a contextual correspondence between a native data item of the host application and a native data item of the client application, and the contextual correspondence may be stored at the host application in a context store for context items. The context items can then be used for various purposes on the host application side, such as for executing contextually informed searches of host application native data items based on a query received from the client application.

[0018] Accordingly, applications and data storage layers using the techniques described herein can provide storage of contextual data that enables applications to interact with clients having a variety of underlying data models and contexts without requiring restructuring of the application's underlying data model or other mitigations. Applications can further use the contextual data to provide enhanced services for clients, such as context-aware searching. Thus, application downtime and latency may be reduced, application developers may spend less time managing data storage layers to accommodate new clients, costs of operating applications and data storage layers may be reduced, and other benefits may accrue. The techniques described herein may further provide improved scalability, performance, and reliability for underlying data storage layers.

[0019] FIG. 1 is a block diagram of an example system architecture **100** for storing contextual data with context schemas, in accordance with an embodiment. System architecture **100** (also referred to as “system” herein) includes network **110**, server **120**, data storage layer **130**, server **150**, and data storage layer **160**. In various embodiments, system **100** may include more or fewer components in different configurations than those depicted in FIG. 1. For example, system **100** may include additional servers and data storage layers in an embodiment.

[0020] Network **110** may include a public network (e.g., the Internet), a private network (e.g., a LAN, a WAN, a VPN, an enterprise network), a wired network (e.g., Ethernet), a wireless network

(e.g., an 802.11 Wi-Fi network), a cellular network (e.g., a 5G network), routers, hubs, switches, server computers, or a combination thereof. Network **110** or components thereof may be associated with different organizations in various embodiments. For example, components of network **110** may be associated with Internet Service Providers (ISPs), mobile or cellular carriers, cloud platform or software-as-a-service (SaaS) providers, private or public enterprises, private households or communities, etc. In an embodiment, network **110** (or a component thereof) may be a physical or virtual interconnect within a single device, such as a PCIe bus, a messaging system, or an API.

[0021] Each of servers **120** and **150** may be a personal computer (PC), a laptop computer, a notebook computer, a mobile phone, a smartphone, a tablet computer, a digital assistant, a rackmount server, a router computer, or similar computing device. An example computing device is further described below with reference to FIG. 6. Servers **120** and **150** may also be a virtualized resource such as a virtual machine (VM) or a containerized application. Servers **120** and **150** may also correspond to a collection of physical or virtual computing resources, such as a datacenter or a collection of servers or VMs distributed across multiple data centers. For example, servers **120** and **150** may correspond to cloud computing resources provisioned from a cloud computing provider. Servers **120** and **150** may run an operating system or one or more software applications.

[0022] Servers **120** and **150** include applications **122** and **152**, respectively. An application may be a hardware (e.g., circuitry, dedicated logic, etc.) or software (e.g., code, libraries, firmware, etc.) tool designed to fulfill various requirements and perform specific tasks. Applications may provide services to users (e.g., individuals or entities/organizations) or other applications. For example, an application on a user device may provide a graphical display of weather information or permit the user to browse the internet. In the same example, applications on servers may provide updated weather information to the user application or serve webpages to the user application for the user to browse. Furthermore, database applications may provide data (e.g., weather, webpages) to the server applications. Applications may communicate with users, other applications, other computing devices/resources, etc. using various communication media and protocols. For example, applications **122** and **152** may communicate with each other over network **110** using an API such as REST, SOAP, WebSockets, etc. In another example, applications **122** and **152** may communicate with data storage layers **130** and **160** using SQL, REST, etc.

[0023] Data storage layers **130** and **160** may be applications for receiving, storing, and providing data. Data storage layers **130** and **160** may be relational or non-relational databases, structured or unstructured databases, key-value stores, filesystems, or may conform to other data storage layer classifications. Data storage layers **130** and **160** may be backed by various persistent or non-persistent storage devices, such as RAM, magnetic tapes or drives, solid-state drives, optical drives, or similar (e.g., other storage technologies discussed below with reference to FIG. 6). Data storage layers **130** and **160** may also include storage devices in a networked topology, such as a Storage Area Network (SAN), Network-Attached Storage (NAS), cloud-provisioned storage, or similar. Data storage layers **130** and **160** may be provided by a respective server or servers (not depicted). In an embodiment, data storage layers **130** and **160** are provided by servers **120** and **150**. Data storage layers **130** and **160** or their respective hardware may be centralized or decentralized. Examples of database applications that may correspond to data storage layers **130** and **160** include MongoDB, MySQL, MariaDB, DynamoDB, PostgreSQL, and others. Data storage layers **130** and **160** may partition data into various stores, buckets, tables, etc. based on the needs of the application(s) serviced by each data storage layer.

[0024] Data storage layer **130** includes native data store **132** for storing data conforming to a data model of application **122**. As previously described, the data model of application **122** may correspond to the context of application **122**, e.g., the requirements, constraints, etc. of application **122**. The data model for native data store **132** may be represented by store schema **138A**. For example, store schema **138A** may correspond to a formal specification of elements of the data

model (e.g., tables, fields, views, document formats, key-value formats, etc.) in a domain-specific language (e.g., SQL) used to initialize native data store **132** in accordance with the data model. In another example, store schema **138A** may correspond to a set of rules to be followed by application **122** when writing to or reading from native data store **132** but may not be realized or enforced by the store itself (e.g., a schema-less database). Native data store **132** further includes native data items **140A-n**, which may be written and read by application **122**. Native data items **140A-n** may correspond to a record, a row, a document, a key-value pair, or other entity within native data store **132**. Native data items **140A-n** may be generated and stored in accordance with the data model of application **122**.

[0025] Data storage layer **160** includes native data store **162** for storing data related to application **152** of server **150**, analogous to native data store **132** of data storage layer **130**. Native data store **162** similarly includes store schema **164**, which may correspond to a design of native data store **162** in view of the requirements, constraints, etc. of application **152**. Native data store **162** similarly includes native data items **166A-n**, which may be written and read by application **152**. Native data stores **132** and **162** may correspond to the same or different types of data storage layers (e.g., relational, non-relational, etc.). Similarly, store schemas **138A** and **164** and native data items **140A-n** and **166A-n** may be similar or different in view of similar or different design aspects of applications **122** and **152**.

[0026] In one example, application **122** may be a contact-center-as-a-service application (also referred to herein as a contact center application or a host application) that supports live and automated chats and calls, customer services, and other related services. Application **122** may communicate with data storage layer **130** to store contact center-related data (e.g., native data) in native data store **132**. Store schema **138A** may specify elements (e.g., tables, fields, document attributes, etc.) related to the contact center application such as caller name and number, agent name and number, service ticket number (e.g., a unique identifier), etc. Each data item in native data items **140A-n** may correspond to a single call/chat or to a chain of related communications. Clients of application **122** may be organizations or other entities that want to provide contact center services (e.g., live chat agent, customer service center, etc.) for their respective users. Clients of application **122** may have their own application—application **152**—for managing support tickets created internally. For example, application **152** may track support tickets opened in-person at a retail store or through a form on the client's website. Application **152** may communicate with data storage layer **160** to store internal support ticket data in native data store **162**. Store schema **164** may specify a different schema than store schema **138A**. For example, store schema **164** may specify elements for support ticket number (e.g., a unique identifier), retail store location, IP address, etc. Clients of application **122** may want to transfer tickets from application **152** to application **122** or to link tickets of the respective applications together. Other clients of application **122** with yet other ticketing applications having different schemas may wish to link their internal tickets as well.

[0027] Returning to FIG. **1**, data storage layer **130** further includes schema store **134** for storing context schemas characterizing a context of application **152** and data storage layer **160**. A context of application **152** and data storage layer **160** may include data elements (e.g., data fields) and data items (e.g., records or entries) that are related to the properties of application **152** or to interactions between applications **122** and **152**. Examples of a context may include unique identifiers, indices, searchable elements, etc. of data storage layer **160**. An application may have multiple contexts, such as different contexts associated with different data storage layers, different aspects of the application, different tables or other data elements, etc. In the example described previously, a context of data storage layer **160** may include the support ticket number, retail store location, etc. A context schema may characterize the example contexts and other contexts as context attributes to provide awareness of the context at application **122** and data storage layer **130**. Context items may include instances of the context attributes defined in a context schema (e.g., context values or

metadata values). For example, a context schema may identify a unique identifier element of data storage layer **160** in a context attribute (e.g., the support ticket number element in the example), and a context item may identify a specific unique identifier value of a specific record of data storage layer **160** (e.g., a support ticket number of a specific support case in the example). Content of an example context schema is further described below with reference to FIG. 2. Content of an example context item is further described below with reference to FIG. 3.

[0028] Schema store **134** includes store schema **138B**, which may correspond to a design of schema store **134** for storing context schemas of application **152**/data storage layer **160** and other applications not depicted. Store schema **138B** may define elements for this purpose as previously described with reference to other store schemas. Schema store **134** may store context schemas of application **152** and other applications as context schemas **142A-n** (e.g., the context schemas may be stored as data items in schema store **134** rather than defining a design of schema store **134**). For example, application **122** may receive a context schema from application **152** characterizing a context associated with application **152** and may proceed to store the context schema in schema store **134** as context schema **142A**. Context schemas received from other applications (or from other components of the same application) may be stored as context schemas **142B-n**.

[0029] Data storage layer **130** further includes context store **136** for storing context items of application **152** and data storage layer **160** (and other applications not depicted). Context store **136** includes store schema **138C**, which may correspond to a design of context store **136** for storing context items of application **152**/data storage layer **160** and other applications not depicted. Store schema **138C** may define elements for this purpose as previously described with reference to other store schemas. Context store **136** may store context items of application **152** and other applications as context items **144A-n**. For example, application **122** may receive a context schema from application **152** including metadata values and may proceed to store the context item in context store **136** as context item **144A**. Context items received from the same or other application may be stored as context items **144B-n**. In an embodiment, a context item may conform to both store schema **138C** and a corresponding context schema of context schemas **142A-n**. For example, store schema **138C** may define elements for a unique context item identifier, a reference to the corresponding context schema, a reference to a corresponding native data item of native data items **140A-n**, and a freeform or flexible element (e.g., a string or dictionary) for storing the metadata values. The corresponding context schema may further define the content of the freeform/flexible metadata values element. In an embodiment, store schema **138C** may be one of context schemas **142A-n**. For example, application **122** or data storage layer **130** may create a new context store **136** for each context schema received, and each new context store may be initialized according to a respective received context schema.

[0030] Data storage layers **130** and **160** further include indices **146** and **168**, respectively, for facilitating indexing and searching of various elements and items of respective stores **132-136** and **162**. A data storage layer index may provide more efficient (e.g., faster, fewer operations) search and retrieval of data items than may be achieved by searching every item in a data storage layer. For example, an index in a relational database may include a copy of or a reference to a column of a table, and various sorts, transformations, or functions may be used on or with the copied column to facilitate faster searching. Other types of data storage layers may have other types of indices or search optimizations. In various embodiments, index **146** may provide a single index for stores **132-136**, or each store may be associated with their own indices (e.g., indices **146A-C**, not depicted).

[0031] In various embodiments, applications **122** and **152** may be associated with various entities. For example, applications **122** and **152** may be associated with two different entities as in the example previously described (e.g., a contact center provider entity and a client entity of the contact center provider). Additional entities associated with additional applications **152** not depicted may communicate with application **122** as well. In another example, applications **122** and

152 may be associated with the same entity. In a variation of the previous example, the contact center provider entity may provide a different application—application **152**—for a different service and may want to use the contact center application—application **122**—to provide support for application **152**. In the various examples, entities may be public or private organizations or enterprises, government agencies, individual users, automated systems and applications, etc. [0032] In various embodiments, data storage layers **130** and **130** may include more or fewer components than depicted in FIG. **1**. For example, one or more of stores **132-136** may be merged or absent (e.g., schema store **134** and context store **136** may be combined in an embodiment), or other stores not depicted may be present (e.g., additional native data stores may be present in either data storage layer). In another example, one or more of store schemas **138A-C** and **164** may be absent (e.g., various stores may be schema-less stores).

[0033] In an embodiment, applications **122** or **152** may be hosted by devices other than servers **120** or **150**. For example, application **152** may be hosted or served by a user device, an edge device, etc. In an embodiment, data storage layer **130** or data storage layer **160** may be absent, or the aspects of data storage layer **130** or data storage layer **160** described herein may be provided by another data storage technology. For example, application **152** may be provided by an edge device, which may not provide any local storage; data of the edge device may thus be stored solely as context at application **122**/data storage layer **130**. In another example, application **122** may store data described with reference to data storage layer **130** in a filesystem instead of a database application.

[0034] FIG. **2** is a block diagram of an example context schema **142A** in accordance with an embodiment. A context schema may be a data item such as a document, a software structure (e.g., a struct, an object, serialized versions of these), a message, etc. As described above with reference to FIG. **1**, a context schema may characterize a context associated with a first application (e.g., application **152**) to provide awareness of the context for a second application (e.g., application **122**). For example, as depicted in FIG. **2**, context schema **142A** may characterize data of native data store **162** as defined by store schema **164**. FIG. **2** depicts store schema **164** as including tables **202A-n** each having various column elements (e.g., elements **204A-n** in table **202A**; other table details not depicted), however, as previously described, store schema **164** may correspond to other types of data storage layers not having tables and columns in various embodiments (e.g., non-relational databases). Other context schemas (e.g., context schemas **132B-n**) may correspond to other stores within application **152**/data storage layer **160** or other stores/databases associated with other applications in a manner similar to or analogous to the depiction of FIG. **2**. Example column element **204A** is depicted as a unique identifier element, however, unique identifier elements may or may not be present in some schemas.

[0035] Context schema **142A** includes context attributes **220A-n**. A context attribute may identify a type of contextual data of the first application to be provided to the second application in corresponding context items (e.g., context items described below with reference to FIG. **3**). For example, a context attribute may correspond to an element of the first application such as a column element (e.g., context attributes **220A-B**), a table element (e.g., context attributes **220C-n**), a document or key-value category, etc. A context attribute may include one or more fields for this and other purposes, such as fields **222A-n** and **224A-n** in FIG. **2**. One field (e.g., fields **222A** and **224A**) may be an identifier field that provides a name or other identifier of the corresponding element (e.g., elements **204A-B**). Another field (e.g., fields **222B** and **224B**) may provide a datatype for the context attribute (e.g., string, integer, complex data type), which corresponding context items should adhere to. Another field (e.g., fields **222C** and **224C**) may indicate whether the context attribute is required or optional in corresponding context items. Another field (e.g., fields **222D** and **224D**) may indicate whether corresponding context items should be searchable by the context attribute within a context store (e.g., whether index **146** should be configured to facilitate searches of context store **136** by the context attribute). Other fields (e.g., fields **222n** and **224n**) having other purposes may be included in various embodiments, and one or more of the example fields **222A-D**

and **224A-D** given herein may be omitted in various embodiments. Different context attributes may have the same or different numbers and types of fields in various embodiments. Each of fields **322A-n** and **324A-n** may be provided by the first application or the second in various embodiments.

[0036] Context attributes may be grouped in attribute groups, as depicted in FIG. 2 with context attributes **220A-B** forming attribute group **226**. Attribute groups may correspond to element groupings or other element relevancies in store schema **164**. For example, attribute group **226** may include context attributes corresponding to column elements of table **202A**, which column elements of other tables may form other attribute groups. Other organizations of context attributes may be provided in various embodiments, such as groups of attribute groups, etc.

[0037] Context schema **142A** further includes namespace **228**. Namespace **228** may be an identifier of context schema **142A**, which may or may not be unique. For example, namespace **228** may identify context schema **142A** with a particular application (e.g., application **152**), with a particular entity (e.g., a client entity), with a group of applications, or similar. Context schema may include additional or fewer components for identification, classification, grouping, versioning, or for other purposes in various embodiments (e.g., unique identifier **230**, a context schema version, a store schema version for a schema store, etc.). Identification-related components may be provided by applications **122** or **152**. For example, namespace **228** may be provided by application **152**, while unique identifier **230** may be provided by application **122** to distinguish context schema **142A** from other context schemas in schema store **134**.

[0038] Referring back to the previously described example, context schema **142A** may characterize data context associated with the client's support ticket application to provide awareness of the context for the contact center application. Store schema **164** may define table **202A** for storing support tickets, where element **204A** is a support ticket number, element **204B** is a retail store location, etc. Context schema **142A** may define corresponding context attributes **220A** for the support ticket number, **220B** for the retail store location, etc. Context attribute **220A** may be required because every support ticket has a support ticket number. However, context attribute **220B** may be optional because not every support ticket was opened at a retail store (e.g., some may be opened on the website form previously mentioned). Both context attributes **220A-B** may be searchable so that contact center staff can search for tickets by support ticket number or retail store location. Namespace **228** may be set by the contact center service provider to distinguish the client's support ticket application from other clients' support ticket applications (e.g., it may be set to a client ID number).

[0039] FIG. 3 is a block diagram of an example context item **144A**, in accordance with an embodiment. A context item may be a data item such as a document, a software structure (e.g., a struct, an object, serialized versions of these), a message, etc. As described above with reference to FIG. 1, a context item may include instances of context attributes (e.g., metadata values) of a first application defined in a context schema to be provided to a second application. For example, as depicted in FIG. 3, context item **144A** may include instances of context attributes of context schema **142A** corresponding to values of native data item **166A** of native data store **162**. FIG. 3 depicts native data item **166A** as including values **302A-n**, which may correspond to column values of a single record/row in a table, values in a document, values in a set of key-value pairs, etc. Example value **302A** is depicted as a unique identifier value, however, unique identifier values may or may not be present in some data items.

[0040] Context item **144A** includes identifier section **320**, which may provide identifying information for context item **144A** relative to various other data items, schemas, applications, entities, etc. Identifier section **320** includes example identifier fields **322A-n**. One field (e.g., field **322A**) may be a unique or non-unique identifier field that provides a name, number, or other identifier of the context item. Another field (e.g., field **322B**) may be a namespace field that identifies a namespace to which context item **144A** belongs. As previously described, a namespace

may be associated with an application, an entity, etc. Another field (e.g., field **322C**) may be a schema field that identifies a context schema (e.g., context schema **142A**) to which the context item conforms. Another field (e.g., field **322D**) may be a data reference field that identifies a native data item of the second application (e.g., native data item **140A** of application **122**) for which context item **144A** provides context. Another field may be a context schema version or store schema version (for a context store). Other fields (e.g., field **322n**) having other purposes may be included in various embodiments, and one or more of the example fields **322A-D** given herein may be omitted in various embodiments. Different identifier sections of context items may have the same or different numbers and types of fields in various embodiments. Each of identifier fields **322A-n** may be provided by the first application or the second in various embodiments.

[0041] Context item **144A** further includes context section **324**, which may provide metadata values for context item **144A**. A metadata value may correspond to a context attribute of a context schema and a value of a native data item. For example, metadata value **326A** corresponds to context attribute **220A** (a unique identifier attribute) of context schema **142A** and value **302A** (a unique identifier value) of native data item **166**. Metadata value **366B** corresponds to context attribute **220B** and value **302B**, **326** corresponds to **220n** and **302n**, and so on. Each of metadata values **226A-n** may be provided by the first application. In an embodiment, some of context attributes **220A-n** may not have corresponding metadata values in a context item. For example, some context attributes may be marked as optional, or a context item may be associated with an attribute group (e.g., attribute group **226**) that is a subset of context attributes **220A-n**.

[0042] Referring back to the example previously described, context item **144A** may correspond to a client's new support ticket that is to be transferred to the contact center application as a new service ticket. The value of field **322A** may be the new support ticket number, the value of field **322B** may be the client ID number (which may be added to the context item by application **122**), and the value of field **322C** may be a reference to context schema **142A** (e.g., unique identifier **230**). When transferring the new support ticket to the contact center application, the client application may include a native data item (of the contact center application) identifying the type of service to be provided (e.g., live chat, phone call) and other information. The native data item may be linked to context item **144A** by field **322D**.

[0043] FIG. **4A** is a sequence diagram of an example interaction **400** between application **152**, application **122**, native data store **132**, schema store **134**, context store **136**, and index **146** for storing contextual data with context schemas, in accordance with an embodiment. In some embodiments, operations depicted in FIG. **4A** could occur in a different order or be performed by different components than depicted. Various embodiments may include additional operations or components not depicted in FIG. **4A** or a subset of operations or components depicted in FIG. **4A**. The operations depicted in FIG. **4A** may correspond to different communication sessions or different timing intervals. For example, some operations may proceed in immediate succession or may be part of a single communication session, while other operations may be spread out over time or may be part of different communication sessions.

[0044] At operation **402**, application **122** receives a context schema from application **152**. The received context schema may characterize a context associated with application **152** to provide awareness and integration of the context (e.g., in the form of context items) at application **122**. The received context schema may be context schema **142A** described above with reference to FIG. **2**, and may include identifying information (e.g., unique ID, namespace), context attributes, and other information. The received context schema may be received in a message, an API request (e.g., REST, SOAP, etc.), or similar. For example, application **152** may use an API of application **122** to send the context schema or vice versa. The received context schema may be represented in various formats and types of data items as previously described. For example, the received context schema may be represented in JSON, XML, or YAML markup formats. In an embodiment, the context schema may be received at application **122** as part of a configuration transaction between

applications **152** and **122**. For example, application **152** may be subscribing to a service of application **122** for the first time and may provide configuration values to customize aspects of application **122** to meet the needs of application **152**. Application **122** may provide the context schema as a configuration value to facilitate exchange of contextual data in future transactions. In an embodiment, the context schema may be received as an update of an existing context schema (e.g., presently stored in schema store **134**).

[0045] At operation **404**, application **122** stores the received context schema in schema store **134**. In an embodiment, storing the context schema may include validating the context schema prior to or simultaneously with storing the context schema. For example, schema store **134** includes store schema **138B**, which may be created at design time of application **122** to enable uniform storage of context schemas from various client applications. The context schema may be validated against store schema **138B**, such as by determining whether the context schema conforms to a required format, includes values for all required elements, etc. The context schema may be converted, translated, reorganized, etc. to be compatible with the store schema. In an embodiment, a transaction or communication with application **152** may be rejected or canceled if the context schema is determined to be invalid. In an embodiment, storing the context schema may include creating or updating an index of one or more of stores **132-136** (e.g., index **146**) based on the context schema (e.g., based on context attributes of the context schema indicated as searchable context attributes). In an embodiment, the received context schema may be stored in a location or storage medium other than schema store **134**. For example, the context schema may be stored in RAM of server **120**, in another data store (e.g., native data store **132** or context store **136**), in a key-value service, or similar.

[0046] At operation **406**, application **122** receives a combined data item from application **152**. The combined data item may include a native data item and a context item. As previously described, the native data item may include data values related to services or functionality provided by application **122**. For example, the native data item may include arguments of API function calls of application **122**, message contents required by application **122**, or similar. The context item may include data values related to the received native data item. For example, the context item may include metadata values to associate a native data item(s) of application **152** with the received native data item of application **122**. The received context item may be context item **144A** described above with reference to FIG. 3, and may include identifying information (e.g., identifier section **320**), metadata values (e.g., context section **324**), and other information. In various embodiments, the received combined data item may include one of either the native data item or the context item, or the native data item and the context item may be received in separate communications or transactions. In an embodiment, the received native data item or the received context item may reference each other, e.g., using a data reference value (e.g., a value of field **322D** of FIG. 3).

[0047] At operation **408**, application **122** stores the received native data item in native data store **132**. In an embodiment, storing the native data item may include validating the native data item prior to or simultaneously with storing the native data item. For example, the native data item may be validated against a store schema of native data store **132** (e.g., store schema **138A**), or the native data item may be converted, translated, reorganized, etc. to be compatible with the store schema (e.g., API function arguments may be converted into a record and inserted into a database table). In an embodiment, a transaction or communication with application **152** may be rejected or canceled if the native data item is determined to be invalid. In an embodiment, storing the native data item may include modifying the native data item. For example, the native data item may be modified to include a reference to the received context item. In an embodiment, the received native data item may be stored in another location or storage medium as described with reference to operation **404**.

[0048] At operation **410**, application **122** reads the context schema received at operation **402** from schema store **134**. For example, application **122** may read the context schema from schema store **134** into RAM for further processing described below. In an embodiment, operation **410** may be

absent. For example, the received context schema may have been stored in RAM at operation **404**, and thus reading the context schema back into RAM at operation **410** may be unnecessary. In an embodiment, multiple context schemas may be stored at schema store **134**, and application **122** may determine which context schema to read based on a namespace or other identifier indicated in the received context item.

[0049] At operation **412**, application **122** validates the received context item against the context schema. Validating the context item may include determining whether every metadata value of the context item corresponds to a context attribute of the context schema. Validating the context item may further include determining whether every required context attribute of the context schema (e.g., as indicated by example fields **222C** and **224C** of FIG. 2) corresponds to a metadata value of the context item. In an embodiment, a transaction or communication with application **152** may be rejected or canceled if the context item is determined to be invalid against the context schema.

[0050] At operation **414**, application **152** stores the received context item in context store **136**. In an embodiment, storing the context item may include additional validation of the context item prior to or simultaneously with storing the context item. For example, the context item may be validated against a store schema of context store **136** (e.g., store schema **138C**), or the context item may be modified, converted, translated, reorganized, etc. to be compatible with the store schema (e.g., a unique identifier such as a value for field **322A** may be added before storing). In an embodiment, a transaction or communication with application **152** may be rejected or canceled if the context item is determined to be invalid against the store schema. In an embodiment, storing the context item may include modifying the context item. For example, the context item may be modified to include a reference to the received native data item. In an embodiment, the received context item may be stored in another location or storage medium as described with reference to operations **404** and **408**.

[0051] At operation **416**, application **122** indexes searchable attributes of the context item. For example, application **122** may create or update an index of one or more of stores **132-136** (e.g., index **146**) based on metadata values of the context item corresponding to searchable context attributes of the context schema. In an embodiment, indexing of searchable attributes of the context item may be part of operations **412** or **414**. Operation **416** may be absent for context items or context schemas not including searchable context attributes or values (e.g., a searchable attribute may also be an optional attribute, and context items not including a corresponding value may not be indexed on that attribute/value).

[0052] In an embodiment, as previously described, example interaction **400** may include additional applications **152** in communication with application **122**. The additional applications may each provide one or more context schemas (e.g., as in operation **402**) and may further provide one or more context items (and associated native data items, if applicable) for each provided context schema (e.g., as in operation **406**). Application **122** may store the received context schemas, native data items, or context items in shared stores holding data items from all applications or in dedicated stores for each application (or subsets of applications). Application **122** may add identifying information (e.g., namespaces, unique identifiers) to received context schemas, native data items, or context items to distinguish the originating application for each item.

[0053] FIG. 4B is a sequence diagram of an example interaction **430** between application **152**, application **122**, native data store **132**, schema store **134**, context store **136**, and index **146** for storing contextual data with context schemas, in accordance with an embodiment. In some embodiments, operations depicted in FIG. 4B could occur in a different order or be performed by different components than depicted. Various embodiments may include additional operations or components not depicted in FIG. 4B or a subset of operations or components depicted in FIG. 4B. The operations depicted in FIG. 4B may correspond to different communication sessions or different timing intervals. For example, some operations may proceed in immediate succession or may be part of a single communication session, while other operations may be spread out over time

or may be part of different communication sessions.

[0054] At operation **432**, application **122** receives a search query from application **152**. The search query may be based on context previously provided by application **152** to application **122**. For example, the search query may correspond to searchable context attributes indicated in a previously received context schema or indexed metadata values of a previously received context item. In an embodiment, the search query may be implicit in another communication or transaction between application **152** and **122** (e.g., in another communication directed to particular context, which may induce application **122** to initiate a search for that context). The search query may include additional information, such as a namespace related to application **152**, the type of data to be returned, an action to be performed with the query results, etc.

[0055] At operation **434**, application **122** queries index **146** based on the received search query. The index query may further be based on one or more previously indexed context attributes (e.g., operation **416** of FIG. **4A**). Querying the index at operation **434** may include translating or modifying the search query received at operation **432** to generate another search query. For example, application **122** may receive a search query from application **152** via an API (e.g., REST) and may generate another search query for the index in a domain-specific representation (e.g., as a SQL query) using the received query.

[0056] At operation **436**, application **122** receives query results from index **246**. In an embodiment, the query results may indicate unique identifiers of or references to one or more context items in context store **136**. Then at operation **438**, application **122** fetches context items from context store **136** based on the query results (e.g., based on the unique identifiers of the context items or references to the context items in context store **136**). At operation **440**, application **122** receives the context items from context store **136**. In an embodiment, the query results at operation **436** may be the context items in context store **136** (e.g., the index may return the context items directly). Thus, operations **438** and **440** may be absent.

[0057] At operation **442**, application **122** identifies native data items associated with the received context items. For example, the received context items may each include a unique identifier of or a reference to one or more respective native data items (e.g., native data items received in the compound data item at operation **406**). In another example, application **122** may query index **146** for native data items identifying/referencing the received context items and receive results (not depicted). Application **122** may receive identifiers of or references to native data items matching the search criteria (not depicted). In both examples, application **122** fetches native data items from native data store **132** at operation **444** using the identifier/reference. Application **122** then receives the native data items from native data store **132** at operation **446**. In an embodiment, the native data items may be received during another operation, and thus operations **444** and **446** may be absent. For example, native data items may be returned by the index with the query results at operations **436** or at the example query of operation **442** previously described (not depicted).

[0058] At operation **448**, application **122** returns one or more combined data items to application **152**. Each combined data item may include a context item matching the original search query of operation **432** and one or more native data items associated with the included context item. Prior to returning the combined data items, application **122** may translate or modify the context items or native data items and generate the combined data items. For example, data items may be modified to match a format expected by application **152** or to remove fields such as identifiers/references linking the context items and native data items. In an embodiment, the context items and respective native data items may be returned separately rather than combined. In an embodiment, context items may be returned and native data items may not be returned. In an embodiment, native data items may be returned and context items may not be returned.

[0059] FIGS. **5A-B** are flow diagrams of an example method **500** for storing contextual data with context schemas, in accordance with an embodiment. Method **500** may be performed by processing logic that may comprise hardware (e.g., circuitry, dedicated logic, etc.), computer-readable

instructions such as software or firmware (e.g., run on a general-purpose computing system or a dedicated machine), or a combination thereof. For instance, an example system may include a memory and a processing device coupled to the memory device to perform operations comprising the blocks of method **500**. Method **500** may also be associated with a set of instructions stored on a non-transitory computer-readable medium (e.g., magnetic or optical disk, etc.). The instructions, when executed by a processing device, may cause the processing device to perform operations comprising the blocks of method **500**. In an embodiment, method **500** is performed by server **120**, application **122**, or data storage layer **130** of FIG. **1**. In an embodiment, method **500** is performed by computing system **600** of FIG. **6**. In some embodiments, blocks depicted in FIGS. **5A-B** could be performed simultaneously or in a different order than depicted. Various embodiments may include additional blocks not depicted in FIGS. **5A-B** or a subset of blocks depicted in FIGS. **5A-B**.

[0060] At block **502**, processing logic receives a context schema characterizing a context associated with an application, the context schema specifying one or more context attributes, each context attribute identifying a respective context element of the context associated with the application. For example, the application may be application **152**, and the context of the application may be associated with native data store **162** and native data items **166A-n** of data storage layer **160**. The received context schema may be context schema **142A**, and the one or more context attributes may be context attributes **220A-n**. Field **222A** of context attribute **220A** identifies a context element—a unique ID column—stored in native data store **162** of application **152**. At block **504**, processing logic stores the context schema in a schema store. The schema store may be schema store **134**, for example. The context schema may be stored as data and may conform to a store schema of the schema store (e.g., store schema **138B**) as described with reference to operation **404**. The context schema may further conform to a specific version (e.g., version number) of the store schema of the schema store, which may be indicated in a field of the context schema. Storing the context schema at block **504** may include verifying that the context schema version matches the store schema version.

[0061] At block **506**, processing logic receives a second context schema characterizing a second context associated with a second application (e.g., a second application **152** or other client application), the second context schema specifying one or more second context attributes, each second context attribute identifying a respective second context element of the second context associated with the second application. The second application may be similar to or different from the application of block **502**. For example, the second application may have a different store schema **164** for its native data store **162**, including different tables, columns, etc. Similarly, the second context schema may have different context attributes than the context schema of block **502**. For example, a second context attribute of the one or more second context attributes may be absent from the one or more context attributes of the context schema received at block **502**. At block **508**, processing logic stores the second context schema in the schema store as described with reference to block **504**. In an embodiment, the application and the second application may be the same application, and the first second context schemas may characterize different contexts of the same application.

[0062] At block **510**, processing logic receives a context item of the context associated with the application, the context item comprising one or more values, each value corresponding to a respective context attribute of the one or more context attributes. For example, the context item may be context item **144A** and may include a context section with metadata values corresponding to the context attributes of context schema **142A**. The context item may correspond to a data item associated with the application, such as native data item **166A**. The metadata values of the context section may correspond to values of the data item. The context item may include other information, such as the example fields and values **322A-n** in identifier section **320** of context item **144A**.

[0063] In an embodiment, receiving the context item further comprises receiving a data item associated with a second application (e.g., application **122** or other host application), storing the

data item associated with the second application in a data store, and adding an identifier of the stored data item associated with the second application to the context item. The data item associated with the second application may be received as part of a combined data item comprising the context item and the data item. The data item may be native to the second application and may conform to a schema of the data store (e.g., native data item **140A** of native data store **132** conforms to store schema **138A**). The one or more values of the context item may provide context for the data item associated with the second application by linking it to the data item associated with the application as previously described (e.g., native data item **166A**). This contextual correspondence between the data items of the application and the second application may enable context-aware searching and provide other benefits as previously described. The identifier of the stored data item associated with the second application, which may be a unique or non-unique identifier in various embodiments, may be added to the context item to link it with the data item for which it provides context (e.g., the value of field **322A**).

[0064] In an embodiment, a context attribute of the one or more context attributes is associated with an identifier element (e.g., unique or non-unique) stored in a data storage layer of the application, and a value of the one or more values of the context item corresponds to an identifier value (e.g., unique or non-unique) of an item stored in the data storage layer of the application. For example, the data storage layer of the application may be data storage layer **160** or native data store **162**, and the identifier element may be unique identifier element **204A**. The item may be a native data item of the application (e.g., native data item **166A**), which may be associated with a unique identifier value (e.g., value **302A**).

[0065] At block **512**, the processing logic validates the context item against the context schema. For example, the processing logic may determine whether all required attributes of the context schema have a corresponding metadata value in the context item and whether all metadata values in the context item correspond to a context attribute of the context schema. Other validation operations may be performed (e.g., ensuring the metadata values are the right data type, conform to a format, are within a certain range, etc.). In an embodiment, prior to validating the context item against the context schema, the processing logic may first identify the context schema and retrieve the context schema from the schema store. For example, the processing logic may use a namespace or other identifier in the context item to determine the relevant context schema. The processing logic may also use an identifier of the application from which the context item was received (e.g., a key, hostname, IP address, etc.).

[0066] At block **514**, responsive to determining that the context item conforms to the context schema, the processing logic stores the context item in a context store. The context store may be context store **136**, for example. The context item may conform to both the context schema and a store schema of the context store (e.g., store schema **138C**). In an embodiment, the context item may be modified prior to storing as previously described. The context item may further conform to specific versions (e.g., version numbers) of the context schema or store schema of the context store, which may be indicated in a field of the context item. Validation operations performed at blocks **512** or **514** may include verifying that the context item version matches the context schema or store schema versions.

[0067] At block **516**, the processing logic updates an index of the context store to facilitate searches of the one or more context attributes. For example, the index may be index **146**, and the index may be updated based on context attributes indicated as searchable attributes (e.g., by field **222D**). Index **146** may be updated to include a subset of context items (e.g., a view in a relational database table), a searchable collection of references to the context items, or similar.

[0068] At block **518**, the processing logic receives a search query from the application, the search query comprising an identifier of a context attribute and a search value of the respective context element. At block **520**, the processing logic queries the context store using the identifier and the search value. The identifier of the context attribute may be a unique identifier of the context

schema to which the context attribute belongs, and name of the context attribute or corresponding context element, or other identifier. The search value may be an allowed value of the context attribute that may be present in one or more context objects (e.g., a value matching type or formatting requirements specified in the context attribute).

[0069] At block **522**, the processing logic retrieves one or more context items from the context store, each context item comprising a value corresponding to the search value. At block **524**, the processing logic provides the one or more context items to the application. In an embodiment, the processing logic may modify the context items before providing them to the application as previously described.

[0070] At block **526**, the processing logic retrieves one or more data items (e.g., native data items) each corresponding to a reference value stored in a respective result context item of the retrieved one or more context items. At block **528**, the processing logic provides the one or more data items to the application. In an embodiment, the retrieved context items and the retrieved data items may be provided together as one or more result combined data items.

[0071] FIG. **6** is a block diagram illustrating an example computer system **600**, in accordance with implementations of the present disclosure. Computer system **600** may correspond to servers **120** or **150**, applications **122** or **152**, or data storage layers **130** or **160**, as described above with reference to FIG. **1**. Computer system **600** may operate in the capacity of a server or an endpoint machine in endpoint-server network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine may be a television, a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

[0072] Computer system **600** includes processing device **602** (e.g., one or more processors or cores), main memory **604** (e.g., read-only memory (ROM), flash memory, dynamic random access memory (DRAM) such as synchronous DRAM (SDRAM), double data rate (DDR SDRAM), or DRAM (RDRAM), etc.), static memory **606** (e.g., flash memory, static random access memory (SRAM), etc.), and data storage device **608**, which communicate with each other via bus **610**.

[0073] Processing device **602** represents one or more general-purpose processing devices such as a microprocessor, central processing unit, or the like. More particularly, processing device **602** may be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. Processing device **602** may also be one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. Processing device **602** is configured to execute instructions **612** (e.g., for storing contextual data with context schemas) for performing the operations discussed herein.

[0074] Computer system **600** may further include network interface device **614**. Computer system **600** also may include display device **616** (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)), alphanumeric input device **618** (e.g., a keyboard, and alphanumeric keyboard, a motion sensing input device, touch screen), cursor control device **620** (e.g., a mouse), and signal generation device **622** (e.g., a speaker). In some embodiments, computer system **600** may not include display device **616**, alphanumeric input device **618**, and/or cursor control device **620** (e.g., in a headless configuration).

[0075] Data storage device **608** may include a non-transitory machine-readable storage medium **624** (also computer-readable storage medium) on which is stored one or more sets of instructions

612 (e.g., for storing contextual data with context schemas) embodying any one or more of the methodologies or functions described herein. Instructions **612** may also reside, completely or at least partially, within main memory **604** or within the processing device **602** during execution thereof by computer system **600**, main memory **604** and processing device **602** also constituting machine-readable storage media. Instructions **612** may further be transmitted or received over network **626** via network interface device **614**.

[0076] In one implementation, instructions **612** include instructions for storing contextual data with context schemas, as described herein. While computer-readable storage medium **624** (machine-readable storage medium) is shown in an exemplary implementation to be a single medium, the terms “computer-readable storage medium” and “machine-readable storage medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database, and/or associated caches and servers) that store the one or more sets of instructions. The terms “computer-readable storage medium” and “machine-readable storage medium” shall also be taken to include any medium that is capable of storing, encoding or carrying a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The terms “computer-readable storage medium” and “machine-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

[0077] In the foregoing description, numerous details are set forth. It will be apparent, however, to one of ordinary skill in the art having the benefit of this disclosure, that the disclosure may be practiced without these specific details. In some instances, well-known structures and devices are shown in block diagram form, rather than in detail, in order to avoid obscuring the disclosure.

[0078] Some portions of the detailed description have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, transferred, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

[0079] It may be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. Unless specifically stated otherwise, it is appreciated that throughout the description, discussions utilizing terms such as “authenticating”, “providing”, “receiving”, “identifying”, “determining”, “sending”, “enabling” or the like, refer to the actions and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (e.g., electronic) quantities within the computer system memories or registers into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage, transmission or display devices.

[0080] The disclosure also relates to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may include a general-purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a computer readable storage medium, such as, but not limited to, any type of disk including a floppy disk, an optical disk, a compact disc read-only memory (CD-ROM), a magnetic-optical disk, a read-only memory (ROM), a random access memory (RAM), an erasable programmable read-only memory (EPROM), an electrically erasable programmable read-only memory (EEPROM), a magnetic or optical card, or any type of media suitable for storing electronic instructions.

[0081] The words “example” or “exemplary” are used herein to mean serving as an example, instance, or illustration. Any aspect or design described herein as “example” or “exemplary” is not necessarily to be construed as preferred or advantageous over other aspects or designs. Rather, use of the words “example” or “exemplary” is intended to present concepts in a concrete fashion. As used in this application, the term “or” is intended to mean an inclusive “or” rather than an exclusive “or.” That is, unless specified otherwise, or clear from context, “X includes A or B” is intended to mean any of the natural inclusive permutations. That is, if X includes A; X includes B; or X includes both A and B, then “X includes A or B” is satisfied under any of the foregoing instances. In addition, the articles “a” and “an” as used in this application and the appended claims may generally be construed to mean “one or more” unless specified otherwise or clear from context to be directed to a singular form. Moreover, use of the term “an implementation” or “one implementation” or “an embodiment” or “one embodiment” throughout is not intended to mean the same implementation or embodiment unless described as such. The terms “first,” “second,” “third,” “fourth,” etc. as used herein are meant as labels to distinguish among different elements and may not necessarily have an ordinal meaning according to their numerical designation.

[0082] For simplicity of explanation, methods herein are depicted and described as a series of acts or operations. However, acts in accordance with this disclosure can occur in various orders and/or concurrently, and with other acts not presented and described herein. Furthermore, not all illustrated acts may be required to implement the methods in accordance with the disclosed subject matter. In addition, those skilled in the art will understand and appreciate that the methods could alternatively be represented as a series of interrelated states via a state diagram or events. Additionally, it should be appreciated that the methods disclosed in this specification are capable of being stored on an article of manufacture to facilitate transporting and transferring such methods to computing devices. The term article of manufacture, as used herein, is intended to encompass a computer program accessible from any computer-readable device or storage media.

[0083] In additional embodiments, one or more processing devices for performing the operations of the above-described embodiments are disclosed. Additionally, in embodiments of the disclosure, a non-transitory computer-readable storage medium stores instructions for performing the operations of the described embodiments. Also in other embodiments, systems for performing the operations of the described embodiments are also disclosed.

[0084] It is to be understood that the above description is intended to be illustrative, and not restrictive. Other embodiments will be apparent to those of skill in the art upon reading and understanding the above description. The scope of the disclosure may, therefore, be determined with reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

Claims

1-20. (canceled)

21. A method comprising: receiving, by a host application and from a client application, a search query comprising an identifier of a context attribute and a search value of a respective context element; querying a context store using the identifier of the context attribute and the search value of the respective context element to obtain one or more context items of context associated with the client application, wherein each of the context items comprises a value corresponding to the search value of the respective context element and a reference value corresponding to a data item associated with the host application; and providing, to the client application, the one or more context items of the context associated with the client application and one or more corresponding data items associated with the host application.

22. The method of claim 21, further comprising: retrieving the one or more corresponding data items from a data store, each of the one or more corresponding data item pertaining to a reference

value stored in a respective context item of the one or more context items.

23. The method of claim 21, wherein the context associated with the client application is characterized by a context schema that specifies one or more context attributes, each context attribute identifying a respective context element of the context associated with the client application.

24. The method of claim 23, wherein each of the one or more context items stored in the context store comprises one or more values, each value corresponding to a respective context attribute of the one or more context attributes.

25. The method of claim 23, wherein each of the one or more context items that is stored in the context store conforms to the context schema.

26. The method of claim 23, wherein the identifier of the context attribute is at least one of an identifier of the context schema, a name of the context attribute, or a name of the respective context element.

27. The method of claim 21, wherein the context store stores context items of multiple contexts each associated with one of a plurality of client applications.

28. A system comprising: a memory device; and a processing device coupled to the memory device, the processing device to perform operations comprising: receiving, by a host application and from a client application, a search query comprising an identifier of a context attribute and a search value of a respective context element; querying a context store using the identifier of the context attribute and the search value of the respective context element to obtain one or more context items of context associated with the client application, wherein each of the context items comprises a value corresponding to the search value of the respective context element and a reference value corresponding to a data item associated with the host application; and providing, to the client application, the one or more context items of the context associated with the client application and one or more corresponding data items associated with the host application.

29. The system of claim 28, the operations further comprising: retrieving the one or more corresponding data items from a data store, each of the one or more corresponding data item pertaining to a reference value stored in a respective context item of the one or more context items.

30. The system of claim 28, wherein the context associated with the client application is characterized by a context schema that specifies one or more context attributes, each context attribute identifying a respective context element of the context associated with the client application.

31. The system of claim 30, wherein each of the one or more context items stored in the context store comprises one or more values, each value corresponding to a respective context attribute of the one or more context attributes.

32. The system of claim 30, wherein each of the one or more context items that is stored in the context store conforms to the context schema.

33. The system of claim 30, wherein the identifier of the context attribute is at least one of an identifier of the context schema, a name of the context attribute, or a name of the respective context element.

34. The system of claim 28, wherein the context store stores context items of multiple contexts each associated with one of a plurality of client applications.

35. A non-transitory computer-readable medium comprising instructions that, when executed by a processing device, cause the processing device to perform operations comprising: receiving, by a host application and from a client application, a search query comprising an identifier of a context attribute and a search value of a respective context element; querying a context store using the identifier of the context attribute and the search value of the respective context element to obtain one or more context items of context associated with the client application, wherein each of the context items comprises a value corresponding to the search value of the respective context element and a reference value corresponding to a data item associated with the host application; and

providing, to the client application, the one or more context items of the context associated with the client application and one or more corresponding data items associated with the host application.

36. The non-transitory computer-readable medium of claim 35, the operations further comprising: retrieving the one or more corresponding data items from a data store, each of the one or more corresponding data item pertaining to a reference value stored in a respective context item of the one or more context items.

37. The non-transitory computer-readable medium of claim 35, wherein the context associated with the client application is characterized by a context schema that specifies one or more context attributes, each context attribute identifying a respective context element of the context associated with the client application.

38. The non-transitory computer-readable medium of claim 35, wherein each of the one or more context items stored in the context store comprises one or more values, each value corresponding to a respective context attribute of the one or more context attributes.

39. The non-transitory computer-readable medium of claim 35, wherein each of the one or more context items that is stored in the context store conforms to the context schema.

40. The non-transitory computer-readable medium of claim 35, wherein the identifier of the context attribute is at least one of an identifier of the context schema, a name of the context attribute, or a name of the respective context element.
