

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284797

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

NEUVIRTH; Hani Hana et al.

ATTACKER-FOCUSED GRANULAR ACTION DISRUPTION

Abstract

Some embodiments provide attacker-focused granular disruption functionality to detect and defend against cyberattacks. An embodiment observes that a user account or a user session has executed one or more precursor events but has not executed an incrimination action. The incrimination action is more likely, by at least a specified amount, to be performed by an attacker-driven account or session than by a non-attacker-driven account or session. Performance of the incrimination action is preemptively blocked, without disabling the account or session. The block is focused on the incrimination action, and in some scenarios the embodiment also blocks performance of similar actions. After an attempt by the account or session to perform the blocked incrimination action, larger blocks on activity are enforced, up to and including disablement, because the attempt indicates strongly that the account or session is driven by an attacker.

Inventors: NEUVIRTH; Hani Hana (Tel Aviv, IL), ZUKERMAN; Jonatan (Matan, IL), MAOR; Tal Joseph (Seattle, WA), PARTYKA; Pawel Marek (Klosters, CH), DAVRAEV; Daniel (Or Yehuda, IL), HAIK; Eyal (Yehud, IL)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 1000007843365

Appl. No.: 18/597157

Filed: March 06, 2024

Publication Classification

Int. Cl.: G06F21/55 (20130101)

U.S. Cl.:

CPC G06F21/554 (20130101); G06F2221/034 (20130101)

Background/Summary

BACKGROUND

[0001] Attacks on a computing system may take many different forms, including some forms which are difficult to predict, and forms which may vary from one situation to another.

Accordingly, one of the guiding principles of cybersecurity is “defense in depth”. In practice, defense in depth is often pursued by forcing attackers to encounter multiple different kinds of security mechanisms at multiple different locations around or within the computing system. No single security mechanism is able to detect every kind of cyberattack, able to determine the scope of an attack or vulnerability, or able to end every detected cyberattack. But sometimes combining and layering a sufficient number and variety of defenses and investigative tools will prevent an attack, deter an attacker, or at least help limit the scope of harm from an attack or a vulnerability.

[0002] To implement defense in depth, cybersecurity professionals consider the different kinds of attacks that could be made against a computing system, and the different vulnerabilities the system may include. They select defenses based on criteria such as: which attacks are most likely to occur, which attacks are most likely to succeed, which attacks are most harmful if successful, which defenses are in place, which defenses could be put in place, and the costs and procedural changes and training involved in putting a particular defense in place or removing a particular vulnerability to attack. They investigate the scope of an attack, and try to detect vulnerabilities before they are exploited in an attack. Some defenses or investigations might not be feasible or cost-effective for the particular computing system. However, improvements in cybersecurity remain possible, and worth pursuing.

SUMMARY

[0003] Some embodiments address technical challenges arising from efforts to prevent or reduce damage caused by cyberattacks. One challenge is how to rapidly and accurately determine that a computing system is under attack. Another challenge is how to provide security in a computing system while permitting legitimate use of the computing system. Other technical challenges are also addressed herein.

[0004] Some embodiments taught herein provide or utilize an attacker-focused granular action disruption (AFGAD) functionality, which observes that a precursor is satisfied by a user account or user session. The precursor includes one or more computing system actions, and in some scenarios the precursor also includes one or more computing system states. The AFGAD functionality discerns that an incrimination action, which is associated with the precursor, has not yet been performed by the user account or session as a follow-up after the precursor. The incrimination action is more likely, by a specified amount, to be performed by an attacker-driven account or session than by a non-attacker-driven account or session.

[0005] The AFGAD functionality blocks the user account or session from performing the incrimination action. The incrimination action block has a finer granularity than simply disabling the user account, for example. In particular, the block maintains a permission for the user account or session to attempt execution of the incrimination action, even though actual execution of the incrimination action is blocked. In some scenarios, the user account or session does attempt to perform the blocked incrimination action. The attempt indicates the user account or session is being driven by (i.e., performed on behalf of) an attacker, so the AFGAD functionality restricts the user account or session further, up to and including disabling the user account or session in some scenarios.

[0006] In some embodiments, establishing that the incrimination action is more likely, by the specified amount, to be performed by an attacker-driven account or session than by a non-attacker-driven account or session helps the AFGAD functionality rapidly and accurately determine that the

computing system is under attack. In some embodiments, blocking at a fine granularity around the incrimination action instead of blocking more broadly helps the AFGAD functionality provide security while still permitting legitimate use of the computing system by accounts and sessions that are not attacker-driven.

[0007] Other technical activities, technical characteristics, and technical benefits pertinent to teachings herein will also become apparent to those of skill in the art. The examples given are merely illustrative. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter. Rather, this Summary is provided to introduce—in a simplified form—some technical concepts that are further described below in the Detailed Description. Subject matter scope is defined with claims as properly understood, and to the extent this Summary conflicts with the claims, the claims should prevail.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0008] A more particular description will be given with reference to the attached drawings. These drawings only illustrate selected aspects and thus do not fully determine coverage or scope.

[0009] FIG. 1 is a diagram illustrating aspects of computer systems and also illustrating configured storage media, including some aspects generally suitable for embodiments which include or use attacker-focused granular action disruption (AFGAD) functionality;

[0010] FIG. 2 is a block diagram illustrating aspects of a family of enhanced systems which are each configured with AFGAD functionality;

[0011] FIG. 3 is a block diagram illustrating aspects of another family of systems which are each enhanced with AFGAD functionality, including some systems with software which upon execution performs a first family of AFGAD methods;

[0012] FIG. 4 is a block diagram illustrating some aspects of AFGAD;

[0013] FIG. 5 is a flowchart illustrating a second family of AFGAD methods;

[0014] FIG. 6 is a flowchart illustrating a third family of AFGAD methods; and

[0015] FIG. 7 is a flowchart further illustrating AFGAD methods, and incorporating as options the steps of FIGS. 2, 3, 5, and 6.

DETAILED DESCRIPTION

Overview

[0016] Some teachings described herein were motivated by technical challenges faced and insights gained during efforts to improve technology for cloud attack detection and response security tools. These challenges and insights provided some motivations, but the teachings herein are not limited in their scope or applicability to these particular errors, tools, motivational challenges, solutions, or insights.

[0017] Some security tools will determine that a computing system is under attack only after a non-empty set of one or more conditions is met, e.g., there have been five successive attempts to log into a given account using wrong passwords, or ten new user accounts have been created by a particular admin account within the past ten minutes, or more than fifty megabytes of data have been transmitted from a given user email account to an IP address never previously used by that email account. When the security tool determines the system is under attack by a particular account, the security tool disables the account to prevent further harm.

[0018] However, security tools are often configured to let all activity continue in an account until there is a high level of confidence that the activity is part of an attack. For example, tools rarely if ever disable an account after only a single failure to enter the correct password. After all, a legitimate user sometimes enters an incorrect password once, or perhaps more than once. But at

some point, which is determined by the tool's configuration, the number of failed login attempts is interpreted as evidence of an attack instead of being treated as a result of innocent errors.

[0019] By waiting until there is a high level of confidence that the activity is part of an attack, the security tool greatly reduces the risk of interfering with legitimate activity in the computing system. A user who accidentally types the wrong password because caps lock is on is allowed to promptly and easily fix their mistake. The user is not compelled to undergo a time-consuming password recovery and re-authentication process, for example, when the security tool has concluded no attack is underway.

[0020] Unfortunately, by waiting to intervene until there is a high level of confidence that account activity is part of an attack, some security tools also increase the risk that significant damage will be done before the tool intervenes. If the tool waits until fifty megabytes of data have been transmitted to a new IP address, and only then determines a data exfiltration attack is underway, then the confidentiality of fifty megabytes of data is already lost, or at least at serious risk. Even if the data is encrypted, an attacker will be able to attempt decryption at a time and place chosen by the attacker. More generally, when protective intervention is postponed, the potential damage from an attack includes loss of: data integrity, computing system resource availability, confidentiality, privacy, economic value, safety, or reliability, for example.

[0021] Some embodiments described herein utilize or provide a method of attacker-focused granular action disruption (AFGAD) in a computing system. The method includes automatically: observing data representing a performance by a user account of a sequence of user account action(s), referred to here as a precursor; discerning that the precursor is associated with a further user account action not yet performed by the user account, the further user account action referred to here as an incrimination action; establishing that the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account; and blocking the user account from performing the incrimination action, while maintaining a permission for the user account to attempt execution of the incrimination action. A sequence is also sometimes referred to as a chain.

[0022] This AFGAD functionality has the technical benefits of detecting an attack rapidly and limiting damage from the attack. The only potential damage done before the user account is characterized as attacker-driven is whatever was done by the precursor actions. In general, intervention can be faster because the level of confidence (that an incrimination action is malevolent) which is sufficient to trigger a protective constraint on the account (e.g., block the incrimination action) can be lower than the level of confidence deemed sufficient to disable the account. Account characterization relies not only on a sequence of alerts or events, as in some other tools, but also leverages a computationally established fact that the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account.

[0023] Moreover, a confidence level for intervening on the basis that performance of the precursor means an attack is underway can be a lower confidence level than a level that triggers disabling the account, because even if an attack is not underway the disruption caused by tool intervention is finely focused around the incrimination action. In some scenarios, only the incrimination action itself is blocked. In some scenarios, only actions which are similar to the incrimination action are blocked, e.g., commands which are similar in that they create/spawn/fork a process are blocked. In either case, the entire account itself is not blocked, e.g., is not disabled.

[0024] Because the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account, another technical benefit is that disruption of legitimate account activity is avoided. Blocking the incrimination action likely does not block legitimate account activity. Even when the account's precursor activity is legitimate, blocking prior to an incrimination action attempt is limited to no more than blocking the incrimination action itself and actions which are similar to the incrimination action. This disruption is much smaller than the disruption from disabling the account, for example.

[0025] In some scenarios, an embodiment also detects an attempt by the user account to perform the blocked incrimination action. The attempt corroborates the presence of an attack. In response to detecting the attempt, the embodiment restricts a capability of the user account in addition to the blocking the user account from performing the incrimination action. In some cases, after the incrimination action is attempted, an embodiment disables the account, or imposes additional security constraints on the account, or does both.

[0026] This AFGAD attack corroboration functionality has the technical benefit of corroborating a suspicion of attack without putting additional resources at risk. The suspicion arises from performance of the precursor. Execution of the incrimination action has been blocked, so any risk to resources from execution of the incrimination action is avoided. However, the attempt itself corroborates the suspicion that an attack is underway, because the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account.

[0027] In some embodiments, establishing that the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account includes confirming that an attacker likelihood of the incrimination action is N times a non-attacker likelihood of the incrimination action. N is also referred to herein as a focus factor. In some embodiments, the attacker likelihood is computed from a likelihood of the incrimination action being performed by an attacker-driven account, and the non-attacker likelihood is computed from a likelihood of the incrimination action being performed by at least one account that is not categorized as an attacker-driven account.

[0028] Some suitable values for N include 1 plus epsilon, 3, 4, 5, 10, 15, 20, 50, 100, 500, and 1000. In a given computing system, epsilon is the smallest positive number which is representable digitally in that computing system. Values of N which are one or less fail to distinguish attacker-driven accounts from non-attacker-driven accounts, while values of N greater than 1000 often create an unacceptable risk of damage by overemphasizing non-intervention.

[0029] This AFGAD focus factor functionality has the technical benefit of providing a tunable mechanism which balances potential benefits of security tool intervention against potential inconvenience to legitimate users. Tuning the size of focus factor N changes a balance between blocking the incrimination action after the precursor and allowing the incrimination action to execute after the precursor. Blocking it reduces security risk but increases potential inconvenience to legitimate users and hence increases a lost productivity risk, whereas leaving the incrimination action unblocked increases the security risk but reduces potential inconvenience to legitimate users and the risk to productivity. The larger the focus factor N is, the more the balance favors allowing account activities to continue without intervention by the security tool to block the incrimination action.

[0030] In some embodiments, the blocking is completed before any security alert is automatically raised in response to the performance of the sequence of the precursor (a.k.a. the precursor sequence). In some embodiments, the blocking is completed before any security alert which is automatically raised in response to the performance of the sequence of the precursor has been designated a part of an incident, e.g., an attack. Each of these AFGAD functionality operational scenarios illustrate the technical benefit of detecting an attack rapidly. The confidence level sufficient for blocking is lower in these scenarios than the confidence level sufficient for raising the security alert, or the confidence level sufficient for designating the alert part of the security incident, or both.

[0031] In some embodiments, no alert is raised in the computing system in response to the performance of the sequence of the precursor, but the incrimination action is nonetheless blocked. This AFGAD functionality operational scenario illustrates the technical benefit of detecting an attack which other tools in the computing system do not detect. In some cases, the incrimination action is blocked and the system is thus protected, even though no alert is raised by performance of the precursor. For example, in some cases no alert is raised because no alert predicated on the

precursor exists in the system. The AFGAD functionality thus provides protection that other tools do not.

[0032] In some embodiments, establishing that the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account includes getting a likelihood that the user account has performed the incrimination action prior to the performance of the sequence of the precursor by the user account. This AFGAD account-specific functionality has the technical benefit of helping to protect the system against attacks without interfering with legitimate activity by a particular user account that happens to be more likely than an average user account to perform the incrimination action.

[0033] As one example, it is not unusual for an administrator account to update a server by retrieving a file from an internet location and then uploading that file to the server. But in many environments, it is unusual for a non-administrator account which is not driven by an attacker to retrieve a file from an internet location and then upload that file to a server. Accordingly, when an AFGAD establishing step utilizes a likelihood computed primarily or solely from the administrator account's server upload history, blocking a legitimate server upload by the administrator account is less likely than a blocking that relies on an average of non-attacker-driven accounts.

[0034] In discussions herein, an AFGAD disruption tuple includes a precursor and at least one corresponding incrimination action. In some embodiments, descriptions of the precursor and the incrimination action are formed at least in part by computationally analyzing at least one of: a non-empty set of alerts, a non-empty set of events and alerts, or a non-empty event log. This AFGAD disruption tuple formation functionality has the technical benefit of automatically and efficiently identifying precursors and corresponding incrimination action(s) which are not inherently constrained by the set of alerts defined in the computing system. Accordingly, intervention thresholds for blocking incrimination actions are potentially lower than the thresholds for raising alerts; the lower confidence levels permit faster protective intervention than intervention based on alerts alone. Also, a different set of risky situations is recognizable, beyond the situations defined by alert conditions. In some scenarios, the incrimination action is blocked and the system is thus protected, even though no alert is raised by performance of the precursor.

[0035] These and other benefits will be apparent to one of skill from the teachings provided herein.

Operating Environments

[0036] With reference to FIG. 1, an operating environment **100** for an embodiment includes at least one computer system **102**. The computer system **102** may be a multiprocessor computer system, or not. An operating environment may include one or more machines in a given computer system, which may be clustered, client-server networked, and/or peer-to-peer networked within a cloud **136**. An individual machine is a computer system, and a network or other non-empty group of cooperating machines is also a computer system. A given computer system **102** may be configured for end-users, e.g., with applications, for administrators, as a server, as a distributed processing node, and/or in other ways.

[0037] Human users **104** sometimes interact with a computer system **102** user interface by using displays **126**, keyboards **106**, and other peripherals **106**, via typed text, touch, voice, movement, computer vision, gestures, and/or other forms of I/O. Virtual reality or augmented reality or both functionalities are provided by a system **102** in some embodiments. A screen **126** is a removable peripheral **106** in some embodiments and is an integral part of the system **102** in some embodiments. The user interface supports interaction between an embodiment and one or more human users. In some embodiments, the user interface includes one or more of: a command line interface, a graphical user interface (GUI), natural user interface (NUI), voice command interface, or other user interface (UI) presentations, presented as distinct options or integrated.

[0038] System administrators, network administrators, cloud administrators, security analysts and other security personnel, operations personnel, developers, testers, engineers, auditors, and end-users are each a particular type of human user **104**. In some embodiments, automated agents,

scripts, playback software, devices, and the like running or otherwise serving on behalf of one or more humans also have user accounts, e.g., service accounts. Sometimes a user account is created or otherwise provisioned as a human user account but in practice is used primarily or solely by one or more services; such an account is a de facto service account. Although a distinction could be made, “service account” and “machine-driven account” are used interchangeably herein with no limitation to any particular vendor.

[0039] The distinction between human-driven accounts and machine-driven accounts is a different distinction than the distinction between attacker-driven accounts and non-attacker driven accounts. A particular human-driven account may be attacker-driven, or non-attacker-driven, at a given point in time. Similarly, a particular machine-driven account may be attacker-driven, or non-attacker-driven, at a given point in time.

[0040] Accounts and sessions are each associated herein with reference numeral **130**. Although for convenience, examples and claims herein often speak in terms of accounts, “account” means “account or session or both” unless stated otherwise. In this disclosure, including in the claims and elsewhere, a statement about activity by “the user account or the user session” does not mean that both the user account and the user session must be present. Instead, such a statement is to be understood as a pair of corresponding but distinct statements given as alternatives, one statement being about activity by the user account, and the other statement being about activity by the user session. Likewise, a characterization of “the user account or the user session” does not mean that both the user account and the user session must be present. Instead, such a characterization is to be understood as a pair of corresponding but distinct characterizations given as alternatives, one characterizing the user account, and the other characterizing the user session.

[0041] Storage devices or networking devices or both are considered peripheral equipment in some embodiments and part of a system **102** in other embodiments, depending on their detachability from the processor **110**. In some embodiments, other computer systems not shown in FIG. **1** interact in technological ways with the computer system **102** or with another system embodiment using one or more connections to a cloud **136** and/or other network **108** via network interface equipment, for example.

[0042] Each computer system **102** includes at least one processor **110**. The computer system **102**, like other suitable systems, also includes one or more computer-readable storage media **112**, also referred to as computer-readable storage devices **112**. In some embodiments, tools **122** include security tools or software applications, on mobile devices **102** or workstations **102** or servers **102**, editors, compilers, debuggers and other software development tools, as well as APIs, browsers, or webpages and the corresponding software for protocols such as HTTPS, for example. Files, APIs, endpoints, and other resources may be accessed by an account or non-empty set of accounts, user or non-empty group of users, IP address or non-empty group of IP addresses, or other entity. Access attempts may present passwords, digital certificates, tokens or other types of authentication credentials.

[0043] Storage media **112** occurs in different physical types. Some examples of storage media **112** are volatile memory, nonvolatile memory, fixed in place media, removable media, magnetic media, optical media, solid-state media, and other types of physical durable storage media (as opposed to merely a propagated signal or mere energy). In particular, in some embodiments a configured storage medium **114** such as a portable (i.e., external) hard drive, CD, DVD, memory stick, or other removable nonvolatile memory medium becomes functionally a technological part of the computer system when inserted or otherwise installed, making its content accessible for interaction with and use by processor **110**. The removable configured storage medium **114** is an example of a computer-readable storage medium **112**. Some other examples of computer-readable storage media **112** include built-in RAM, ROM, hard disks, and other memory storage devices which are not readily removable by users **104**. For compliance with current United States patent requirements, neither a computer-readable medium nor a computer-readable storage medium nor a computer-readable

memory nor a computer-readable storage device is a signal per se or mere energy under any claim pending or granted in the United States.

[0044] The storage device **114** is configured with binary instructions **116** that are executable by a processor **110**; “executable” is used in a broad sense herein to include machine code, interpretable code, bytecode, and/or code that runs on a virtual machine, for example. The storage medium **114** is also configured with data **118** which is created, modified, referenced, and/or otherwise used for technical effect by execution of the instructions **116**. The instructions **116** and the data **118** configure the memory or other storage medium **114** in which they reside; when that memory or other computer readable storage medium is a functional part of a given computer system, the instructions **116** and data **118** also configure that computer system. In some embodiments, a portion of the data **118** is representative of real-world items such as events manifested in the system **102** hardware, product characteristics, inventories, physical measurements, settings, images, readings, volumes, and so forth. Such data is also transformed by backup, restore, commits, aborts, reformatting, and/or other technical operations.

[0045] Although an embodiment is described as being implemented as software instructions executed by one or more processors in a computing device (e.g., general purpose computer, server, or cluster), such description is not meant to exhaust all possible embodiments. One of skill will understand that the same or similar functionality can also often be implemented, in whole or in part, directly in hardware logic, to provide the same or similar technical effects. Alternatively, or in addition to software implementation, the technical functionality described herein can be performed, at least in part, by one or more hardware logic components. For example, and without excluding other implementations, some embodiments include one or more of: chiplets, hardware logic components **110**, **128** such as Field-Programmable Gate Arrays (FPGAs), Application-Specific Integrated Circuits (ASICs), Application-Specific Standard Products (ASSPs), System-on-a-Chip components, Complex Programmable Logic Devices (CPLDs), and similar components. In some embodiments, components are grouped into interacting functional modules based on their inputs, outputs, or their technical effects, for example.

[0046] In addition to processors **110** (e.g., CPUs, ALUs, FPUs, TPUs, GPUs, and/or quantum processors), memory/storage media **112**, peripherals **106**, and displays **126**, some operating environments also include other hardware **128**, such as batteries, buses, power supplies, wired and wireless network interface cards, for instance. The nouns “screen” and “display” are used interchangeably herein. In some embodiments, a display **126** includes one or more touch screens, screens responsive to input from a pen or tablet, or screens which operate solely for output. In some embodiments, peripherals **106** such as human user I/O devices (screen, keyboard, mouse, tablet, microphone, speaker, motion sensor, etc.) will be present in operable communication with one or more processors **110** and memory **112**.

[0047] In some embodiments, the system includes multiple computers connected by a wired and/or wireless network **108**. Networking interface equipment **128** can provide access to networks **108**, using network components such as a packet-switched network interface card, a wireless transceiver, or a telephone network interface, for example, which are present in some computer systems. In some, virtualizations of networking interface equipment and other network components such as switches or routers or firewalls are also present, e.g., in a software-defined network or a sandboxed or other secure cloud computing environment. In some embodiments, one or more computers are partially or fully “air gapped” by reason of being disconnected or only intermittently connected to another networked device or remote cloud. In particular, AFGAD functionality **204** could be installed on an air gapped network **108** and then be updated periodically or on occasion using removable media **114**, or not be updated at all. Some embodiments also communicate technical data or technical instructions or both through direct memory access, removable or non-removable volatile or nonvolatile storage media, or other information storage-retrieval and/or transmission approaches.

[0048] One of skill will appreciate that the foregoing aspects and other aspects presented herein under “Operating Environments” form part of some embodiments. This document's headings are not intended to provide a strict classification of features into embodiment and non-embodiment feature sets.

[0049] One or more items are shown in outline form in the Figures, or listed inside parentheses, to emphasize that they are not necessarily part of the illustrated operating environment or all embodiments, but interoperate with items in an operating environment or some embodiments as discussed herein. It does not follow that any items which are not in outline or parenthetical form are necessarily required, in any Figure or any embodiment. In particular, FIG. 1 is provided for convenience; inclusion of an item in FIG. 1 does not imply that the item, or the described use of the item, was known prior to the current disclosure.

[0050] In any later application that claims priority to the current application, reference numerals may be added to designate items disclosed in the current application. Such items may include, e.g., software, hardware, steps, processes, systems, functionalities, mechanisms, devices, data structures, kinds of data, settings, parameters, components, computational resources, programming languages, tools, workflows, or algorithm implementations, or other items in a computing environment, which are disclosed herein but not associated with a particular reference numeral herein. Corresponding drawings may also be added.

More About Systems

[0051] FIG. 2 illustrates a computing system **102** configured by one or more of the AFGAD functionality enhancements taught herein, resulting in an enhanced system **202**. In some embodiments, this enhanced system **202** includes a single machine, a local network of machines, machines in a particular building, machines used by a particular entity, machines in a particular datacenter, machines in a particular cloud, or another computing environment **100** that is suitably enhanced. FIG. 2 items are discussed at various points herein.

[0052] FIG. 3 shows some aspects of some enhanced systems **202**. Like FIG. 2, FIG. 3 is not a comprehensive summary of all aspects of enhanced systems **202** or all aspects of AFGAD functionality **204**. Nor is either figure a comprehensive summary of all aspects of an environment **100** or system **202** or other context of an enhanced system **202**, or a comprehensive summary of any aspect of functionality **204** for potential use in or with a system **102**. FIG. 3 items are discussed at various points herein.

[0053] FIG. 4 shows some aspects of granular disruption **208** and some related items or processes. This is not a comprehensive summary of all aspects of AFGAD or of granular disruption generally. FIG. 4 items are discussed at various points herein.

[0054] The other figures are also relevant to systems **202**. FIGS. 5 to 7 are flowcharts which illustrate some methods of AFGAD functionality **204** operation in some systems **202**.

[0055] In some embodiments, the enhanced system **202** is networked through an interface **326**. In some, an interface **326** includes hardware such as network interface cards, software such as network stacks, APIs, or sockets, combination items such as network connections, or a combination thereof.

[0056] Some embodiments include a computing system **202** which is configured to utilize or provide AFGAD functionality **204**. The system **202** includes a digital memory set **112** including at least one digital memory **112**, and a processor set **110** including at least one processor **110**. The processor set is in operable communication with the digital memory set. A digital memory set is a set which includes at least one digital memory **112**, also referred to as a memory **112**. The word “digital” is used to emphasize that the memory **112** is part of a computing system **102**, not a human person's memory. The word “set” is used to emphasize that the memory **112** is not necessarily in a single contiguous block or of a single kind, e.g., a memory **112** may include hard drive memory as well as volatile RAM, and may include memories that are physically located on different machines **101**. Similarly, the phrase “processor set” is used to emphasize that a processor **110** is not

necessarily confined to a single chip or a single machine **101**. Sets are non-empty unless described otherwise.

[0057] In this example, a first description **464** resides in and configures the at least one digital memory **112**. The first description **464** represents a precursor **212**. The precursor includes first data **118** representing a sequence **402** of one or more user account or user session actions **132**. A second description **442** also resides in and configures the at least one digital memory **112**. The second description includes second data **118** representing a different user account or user session action **132** which is referred to here as an incrimination action **308**. The incrimination action is distinct from the sequence of the precursor, in that the incrimination action includes an action **132** that is not part of the sequence.

[0058] The processor set is configured by data and instructions to perform an attacker-focused granular action disruption method **700**. The AFGAD method **700** includes (a) observing **304** third data **118** which represents a performance **306** of the sequence of the precursor by a user account or user session **130**, (b) discerning **502** fourth data **118** which represents that the incrimination action was not performed **306** by the user account or user session after the performance **306** of the sequence of the precursor by the user account or user session, or discerning **502** a lack of fifth data **118** which represents that the incrimination action has been performed **306** by the user account or user session after the performance **306** of the sequence of the precursor by the user account or user session, or both, (c) establishing **504** that, after the performance of the sequence of the precursor, the incrimination action is at least three times more likely to be performed by an attacker-driven account or session **322** than by a non-attacker-driven account or session **322**, and (d) blocking **214** the user account or user session from performing the incrimination action, while (e) maintaining **602** a permission **404** for the user account or user session to attempt execution of the incrimination action, and (f) maintaining **602** a permission **404** for the user account or user session to execute an additional action **132** other than the incrimination action and other than any part of the sequence of the precursor.

[0059] In some embodiments, the establishing **504** includes getting **730** a likelihood **328** that the user account or user session has performed the incrimination action prior to the performance of the sequence of the precursor by the user account or user session. The gotten **730** likelihood **328** is an account-specific or session-specific likelihood **328** unless stated otherwise; in some cases the likelihood **328** is instead specific to a particular non-empty set of accounts or sessions, such as admin accounts, or accounts having a particular privilege level or a particular permission.

[0060] In some embodiments, the at least one processor is also configured to form **734** the first description and the second description at least in part by computationally analyzing **732** at least one of: a non-empty set **434** of alerts **320**, a non-empty set **434** of events **430** and alerts **320**; or a non-empty event log **432**. Herein, forming **734** includes creating or modifying or both; formation **734** is not limited to an initial creation. Events **430** include any observed occurrence in a computing system, whereas alerts **320** represent occurrences which are atypical, anomalous, or otherwise noteworthy with respect to cybersecurity risk **472**.

[0061] In some embodiments, the at least one processor is also configured to form **734** the first description and the second description at least in part by correlating **444**, with a correlation model **438**, candidate precursor actions **132** with a candidate incrimination action **132**. In some embodiments, the correlation model **438** includes at least one of: a statistical **436** model **438**, or a machine learning **440** model **438**.

[0062] In some embodiments, the at least one processor is also configured to form **734** the first description and the second description at least in part by filtering **738** out a candidate incrimination action **132** after determining that the candidate incrimination action is more likely, by a threshold likelihood amount **454**, to be performed by a non-attacker-driven user account or user session than by an attacker-driven user account or user session. In effect, the threshold likelihood **454** operates like the focus factor **318**, that is, as a tuning mechanism **466** to balance security risk **472** against

legitimate user activity disruption risk **472**. The smaller the threshold likelihood amount **454** is, the more the balance favors blocking **214** with the attendant risk of disrupting legitimate user activity. [0063] In some embodiments, the establishing **504** includes confirming **702** that an attacker likelihood **314** of the incrimination action is N times a non-attacker likelihood **316** of the incrimination action, where N is at least fifteen, the attacker likelihood is computed from a likelihood of the incrimination action being performed by an attacker-driven account or session **322**, and the non-attacker likelihood is computed from a likelihood of the incrimination action being performed by at least one account or user session **130** other than the user account or user session **130** that is not categorized **322** as an attacker-driven account or session **322**.

[0064] Some embodiments define the likelihood **328** of an action X being performed by an account Y in terms of a frequency, e.g., (number of times account Y performed the action X) divided by (total number of actions performed by account Y). In some embodiments, an action's likelihood **328** of interest is a probability that an attacker-driven account will perform the action divided by an overall probability to perform the action. One such calculation for a disruption tuple `<<precursor::action-X>>` **468**, where precursor represents one or more events (e.g., actions), is ((number of attacker-driven performances of action-X within one day of the occurrence of the precursor) divided by (number of attacker-driven accounts identified in computer system activity data of the past month)) divided by ((number of non-attacker user-driven accounts that performed the action-X per day) divided by ((number of non-attacker user-driven accounts identified in the computer system activity data in the past month) times (number of days in the computer system activity data))). The attackers **138** are a subset of the users **104**, so this calculation's result will always be smaller less than (when not every user is an attacker) or equal to (when all users are attackers) the value one.

[0065] These are only example likelihood **328** calculations. Other calculations of an action's likelihood **328** of being performed by an attacker-driven account or session versus being performed by a regular non-attacker-driven account or session are also used in some embodiments. Depending on the embodiment and scenario, a likelihood **328** is implemented as a probability in a range between zero and one, or as a value in a specified range (e.g., 18 out of 100), or as an odds ratio (e.g., 1000:1).

[0066] The syntax `<<precursor::incrimination action>>` used herein is for convenience and clarity. A given embodiment, or a discussion of a given embodiment (e.g., in product documentation, source code, or marketing materials) is not required to use this syntax in order to accurately describe disruption tuples or their functional equivalents. Other suitable syntaxes are possible.

[0067] In some scenarios the blocking is undone, e.g., permission to perform the incrimination action is restored **470**. Two examples of proactive automatic blocking **214** undo **470** are: restoring permission after a specified delay, and restoring permission after a specified decrease in an account's risk score **474**.

[0068] Other system embodiments are also described herein, either directly or derivable as system versions of described processes or configured media, duly informed by the extensive discussion herein of computing hardware.

[0069] Although specific AFGAD architecture examples are shown in the Figures, an embodiment may depart from those examples. For instance, items shown in different Figures may be included together in an embodiment, items shown in a Figure may be omitted, functionality shown in different items may be combined into fewer items or into a single item, items may be renamed, or items may be connected differently to one another.

[0070] Examples are provided in this disclosure to help illustrate aspects of the technology, but the examples given within this document do not describe all of the possible embodiments. A given embodiment may include additional or different kinds of AFGAD functionality, for example, as well as different technical features, aspects, mechanisms, software, expressions, operational sequences, commands, data structures, programming environments, execution environments,

environment or system characteristics, proxies, or other functionality consistent with teachings provided herein, and may otherwise depart from the particular examples provided.

Processes (a.k.a. Methods)

[0071] Processes (which are also be referred to as “methods” in the legal sense of that word) are illustrated in various ways herein, both in text and in drawing figures. FIGS. 5, 6, and 7 each illustrate a family of methods **500**, **600**, and **700** respectively, which are performed or assisted by some enhanced systems, such as some systems **202** or another AFGAD functionality enhanced system as taught herein. Method families **500** and **600** are each a proper subset of method family **700**. Moreover, activities identified in block diagrams in FIGS. 2 and 3 include method steps, e.g., blocking **214** an incrimination action after a precursor, restricting **218** an account after the account attempts to perform a blocked incrimination action, steps taken by software **302**, and so on. These method steps are likewise incorporated into method (a.k.a. process) **700**. These diagrams and flowcharts are merely examples; as noted elsewhere, any operable combination of steps that are disclosed herein may be part of a given embodiment when called out in a claim.

[0072] Technical processes shown in the Figures or otherwise disclosed will be performed automatically, e.g., by an enhanced system **202**, unless otherwise indicated. Related non-claimed processes may also be performed in part automatically and in part manually to the extent action by a human person is implicated, e.g., in some situations a human **104** types or speaks in natural language an input such as a particular value for a focus factor **318**. Such input is captured in the system **202** as digital text, or captured as digital audio which is then converted to digital text. Natural language means a language that developed naturally, such as English, French, German, Hebrew, Hindi, Japanese, Korean, Spanish, etc., as opposed to designed or constructed languages such as HTML, Python, SQL, or other programming languages. Regardless, no process contemplated as an embodiment herein is entirely manual or purely mental; none of the claimed processes can be performed solely in a human mind or on paper. Any claim interpretation to the contrary is squarely at odds with the present disclosure.

[0073] In a given embodiment zero or more illustrated steps of a process may be repeated, perhaps with different parameters or data to operate on. Steps in an embodiment may also be done in a different order than the top-to-bottom order that is laid out in FIG. 7. FIG. 7 is a supplement to the textual and figure drawing examples of embodiments provided herein and the descriptions of embodiments provided herein. In the event of any alleged inconsistency, lack of clarity, or excessive breadth due to an interpretation of FIG. 7, the content of this disclosure shall prevail over that interpretation of FIG. 7.

[0074] Arrows in process or data flow figures indicate allowable flows; arrows pointing in more than one direction thus indicate that flow may proceed in more than one direction. Steps may be performed serially, in a partially overlapping manner, or fully in parallel within a given flow. In particular, the order in which flowchart **700** action items are traversed to indicate the steps performed during a process may vary from one performance instance of the process to another performance instance of the process. The flowchart traversal order may also vary from one process embodiment to another process embodiment. Steps may also be omitted, combined, renamed, regrouped, be performed on one or more machines, or otherwise depart from the illustrated flow, provided that the process performed is operable and conforms to at least one claim of an application or patent that includes or claims priority to the present disclosure. To the extent that a person of skill considers a given sequence S of steps which is consistent with FIG. 7 to be non-operable, the sequence S is not within the scope of any claim. Any assertion otherwise is contrary to the present disclosure.

[0075] Some embodiments provide or utilize a method **700** of attacker-focused **206** granular action disruption in a computing system **102**, e.g., in a computer network **108**. This method includes automatically: observing **304** data **118** representing a performance **306**, in the computing system by a user account or user session **130**, of a sequence **402** of one or more user account or user session

actions **132**, the sequence referred to here as a precursor **212**; discerning **502** that the precursor is associated, by an association mechanism **324** in the computing system, with a further action **132** which has not yet been performed by the user account or user session after the performance of the sequence of the precursor, the further action referred to here as an incrimination action **308**; establishing **504** that, after the performance of the sequence of the precursor, the incrimination action is more likely to be performed by an attacker-driven account or session **322** than by a non-attacker-driven account or session **322**; and blocking **214** the user account or user session from performing the incrimination action in the computing system, while maintaining **602** in the computing system a permission **404** for the user account or user session to attempt execution **406** of the incrimination action, and maintaining **602** in the computing system a permission **404** for the user account or user session to execute an additional action other than the incrimination action and other than any part of the precursor. Some examples of an association mechanism **324** include data structures that implement disruption tuples (pairs are 2-tuples), such as records, lists, arrays, trees, key-value pairs, and lookup tables.

[0076] In this example, the disruption **208** caused by the blocking **214** is both attacker-focused and granular. The disruption is attacker-focused **206** because the blocked incrimination action is more likely to be performed by an attacker-driven account or session **322** than by a non-attacker-driven account or session **322**. The disruption is granular **456** because an attempt to execute the incrimination action is not blocked (that permission is maintained). The disruption is also granular because the additional action, which is not the incrimination action and also not part of the precursor, is not blocked (that permission is also maintained). In particular, the account or session is not disabled **460**. Disabling an account or session is not a granular disruption.

[0077] In some embodiments, the method **700** further includes: detecting **310** an attempt **216** by the user account or user session to perform the blocked incrimination action; and in response to detecting the attempt, restricting **312** a capability **134** of the user account or user session in addition to the blocking the user account or user session from performing the incrimination action. Capabilities **134** include, e.g., resource access, resource creation, resource deletion, and resource relocation functionality.

[0078] In some embodiments, the establishing **504** includes confirming **702** that an attacker likelihood **314** of the incrimination action is N times a non-attacker likelihood **316** of the incrimination action, where N is at least five. N is also referred to as the focus factor **318**. The attacker likelihood is computed from a likelihood **328** of the incrimination action being performed by an attacker-driven account or session **322**, and the non-attacker likelihood is computed from a likelihood **328** of the incrimination action being performed by at least one account or session that is not categorized **322** as an attacker-driven account or session **322**.

[0079] In some embodiments, blocking **214** the user account or user session from performing the incrimination action includes at least one of: blocking **214** an action which is configured to upon execution add **704** a user account to the computing system; blocking **214** an action which is configured to upon execution run **706** a specified process **742**; blocking **214** an action which is configured to upon execution change **708** a resource **408** configuration **410** in the computing system; blocking **214** an action which is configured to upon execution create **710** a resource **408** in a cloud computing environment **136**; blocking **214** an action which is configured to upon execution request **712** an internet connection **414**; or blocking **214** an action which is configured to upon execution transmit **714** data over an internet connection **414**.

[0080] In some embodiments, blocking **214** the user account or user session from performing the incrimination action includes at least one of: changing **716** an association in the computing system between a role **420** and the user account **130**; prohibiting changing **716** an association in the computing system between a role **420** and a different user account **130**; preventing **718** creation of a virtual machine **416**; prohibiting **720** email activity **722** from a specified IP address **418**; prohibiting **720** email activity **722** from a specified shell **422**; or prohibiting **720** email activity **722**

from a specified web **424** app **426**.

[0081] In some embodiments, an example of changing a role of the user is changing from 'Global Admin' to 'Admin Reader', thus removing permission to create new users. In some embodiments, an example of prohibiting changing a role of another user is a high privilege role assignment on a guest user. In some embodiments, an example of creation of a virtual machine is a mass compute creation. In some embodiments, an example of prohibiting email activity from a specified IP address is prohibiting Exchange Online activities from an IP address using OWA and PowerShell. In some embodiments, this is also an example of prohibiting email activity from a specified shell, and an example of prohibiting email activity from a specified web app.

[0082] More generally, in some scenarios an action is blocked indirectly, while in other scenarios the action is blocked directly. In a system which lacks a permission specifically focused on a particular action, some actions are nonetheless susceptible to indirect blockage **214** by way of a different action. For example, a system S does not support blocking new user addition specifically, but S does allow role changes, and one allowed role change to a role R has the effect of blocking new user addition because the role R lacks permission to add new users.

[0083] In some embodiments, the blocking **214** is completed before any security alert **320** is automatically raised **724** in response to the performance of the sequence of the precursor. In some embodiments, the blocking **214** is completed before any security alert **320** which is automatically raised **724** in response to the performance of the sequence of the precursor has been designated **726** a part of an incident **428**. In some, both of the foregoing statements are accurate.

[0084] In some embodiments, no alert **320** is raised **724** in the computing system in response to the performance **306** of the sequence of the precursor.

[0085] In some embodiments, the method **700** includes: ascertaining **728** that a description of a precursor state **124** is also associated in the computing system with the incrimination action; and verifying **744** that the precursor state is present in the computing system before blocking **214** the user account or user session from performing the incrimination action in the computing system.

[0086] In other words, in some embodiments or scenarios the precursor includes both states **124** and events **430** (e.g., one or more actions **132**). In other embodiments or scenarios, states **124** are not part of the precursor, but the precursor does include one or more events **430**.

[0087] In some embodiments, examples of precursor **212** states **124** include which kernel **120** is installed, which application version is running, counts of past events, data indicating the presence of a particular type of ongoing security campaign or an increase in the chance of a certain attack, data indicating a user role, data indicating permissions granted to a user account or a user session, data indicating resources accessible to a user account or a user session, or an event, e.g., receipt by a machine of an outside connection.

Configured Storage Media

[0088] Some embodiments include a configured computer-readable storage medium **112**. Some examples of storage medium **112** include disks (magnetic, optical, or otherwise), RAM, EEPROMS or other ROMs, and other configurable memory, including in particular computer-readable storage media (which are not mere propagated signals). In some embodiments, the storage medium which is configured is in particular a removable storage medium **114** such as a CD, DVD, or flash memory. A general-purpose memory, which is be removable or not, and is volatile or not, depending on the embodiment, can be configured in the embodiment using items such as likelihoods **328**, focus factor **318**, threshold likelihood amount **454**, tuning mechanism **466**, association mechanism **324**, account categorizations **322**, disruption tuples **468**, correlation models **438**, and AFGAD software **302**, in the form of data **118** and instructions **116**, read from a removable storage medium **114** and/or another source such as a network connection, to form a configured storage medium. The foregoing examples are not necessarily mutually exclusive of one another. The configured storage medium **112** is capable of causing a computer system **202** to perform technical process steps for providing or utilizing AFGAD functionality **204** as disclosed

herein. The Figures thus help illustrate configured storage media embodiments and process (a.k.a. method) embodiments, as well as system and process embodiments. In particular, any of the method steps illustrated in FIGS. 5 to 7, or otherwise taught herein, may be used to help configure a storage medium to form a configured storage medium embodiment.

[0089] Some embodiments use or provide a computer-readable storage device **112**, **114** configured with data **118** and instructions **116** which upon execution **406** by a processor **110** cause a computing system **202** to perform a method **700** of attacker-focused granular action disruption in a computing system. This method **700** includes automatically: observing **304** data representing a performance, in the computing system by a user account or a user session, of a sequence of one or more events or actions, the sequence referred to here as a precursor; discerning **502** that the precursor is associated, by an association mechanism in the computing system, with a further action which has not yet been performed by the user account or user session after the performance of the sequence of the precursor, the further action referred to here as an incrimination action; establishing **504** that, after the performance of the sequence of the precursor, the incrimination action is at least four times as likely to be performed by an attacker-driven account or session as the incrimination action is to be performed by a non-attacker-driven account or session; and blocking **214** the user account or user session from performing the incrimination action in the computing system, while maintaining in the computing system a permission for the user account or user session to attempt execution of the incrimination action, and maintaining in the computing system a permission for the user account or user session to execute an additional action other than the incrimination action and other than any part of the precursor.

[0090] In some embodiments, the method further includes detecting **310** an attempt **216** by the user account or user session to perform the blocked incrimination action, and in response to detecting the attempt, restricting **312** a capability **134** of the user account or user session in addition to blocking the user account or user session from performing the incrimination action. In some cases, restricting the capability includes at least one of: disabling **460** the user account or user session; adding **458** an authentication requirement **448** to the user account or user session; adding **458** an authorization requirement **450** to the user account or user session; or simulating **462** to the user account or user session a simulated performance **452** of the blocked incrimination action and not actually performing the blocked incrimination action.

[0091] Adding **458** a requirement **446** is an example of a large block **218** but is less disruptive than disabling **460** the user account.

[0092] Simulating **462** performance of the blocked incrimination action is less disruptive than a large block **218**. Simulation **462** is suitable for use in honeypot scenarios to avoid revealing a security presence to the attacker, and in scenarios where it is more important to find out what the attacker does next than it is to protect the particular resources the attacker can presently access.

[0093] In some embodiments, the method includes obtaining **444**, at least in part by executing a machine learning model **440**, at least one of: a likelihood **314** that the incrimination action is performed by an attacker-driven account or session; or a likelihood **316** that the incrimination action is performed by an account that is not categorized as an attacker-driven account or session.

[0094] In some embodiments, the establishing **504** includes establishing that after the performance of the sequence of the precursor (a.k.a. the precursor sequence), the incrimination action is at least ten times as likely to be performed by an attacker-driven account or session as the incrimination action was to be performed by the user account or user session prior to the performance of the sequence of the precursor by the user account or user session.

[0095] In some embodiments, the method includes detecting **310** an attempt by the user account or user session to perform the blocked incrimination action; and in response to detecting the attempt, automatically and proactively categorizing **322** the user account as an attacker-driven account or session. Larger disruptions then follow as a result of the categorization.

Scaling-Up Disruption

[0096] Additional support for the discussion of AFGAD functionality **204** herein is provided under various headings. However, it is all intended to be understood as an integrated and integral part of the present disclosure's discussion of the contemplated embodiments.

[0097] One of skill will recognize that not every part of this disclosure, or any particular details therein, are necessarily required to satisfy legal criteria such as enablement, written description, best mode, novelty, nonobviousness, inventive step, or industrial applicability. Any apparent conflict with any other patent disclosure, even from the owner of the present subject matter, has no role in interpreting the claims presented in this patent disclosure. It is in the context of this understanding, which pertains to all parts of the present disclosure, that examples and observations are offered herein.

[0098] Automated disruption can be extremely beneficial in safeguarding user accounts and user sessions from attack, and eliminating delays caused by requiring human approval. However, customers are rightly concerned about the possibility of false-positive detections that lead to business disruption. As a result, automated disruption is generally applied cautiously, concentrating on alerts with an exceptionally high confidence (e.g., >99%). This approach minimizes business disruption from automated disruption, but also substantially limits protective coverage. The limited protection increases the risk of later business disruption resulting from successful attack(s). Regardless, when there is an exceptionally high confidence of an attack, the automated disruption procedures applied are very stringent, such as disabling or locking user accounts. Moreover, the automated disruption is generally initiated only after the user has been verified as an attacker, at which point damage has often already been inflicted.

[0099] The present disclosure teaches an AI-based strategy to maximize automated disruption coverage in a timely manner to minimize damage, while maintaining high precision with a low risk to business processes. Some embodiments taught herein utilize or provide AFGAD functionality **204** which performs high resolution adaptive hardening based on a what-if analysis of an action **132** space. This AFGAD functionality **204** disrupts cyberattacks, in particular ways taught herein.

[0100] When automating disruption of cyberattacks, there is a risk of misidentification of the driver of an account. In some cases, despite suspicious activity in an account, the user driving the account is a valid user performing non-malicious activity, and therefore disrupting them will disrupt the regular operation of a cloud tenant or other customer. As a result, some cyberattack disruption tools are configured to require very high confidence that a specific user account is being driven maliciously before they employ disruption procedures such as disabling the user account. By the time this confidence is gained, an attacker will often have already managed to cause some damage to the customer, such as deleting files, exfiltrating data, and so on.

[0101] Some embodiments taught herein use adaptive hardening in a very granular **456** way at an early stage of an attack to block permissions to actions that, if executed, would incriminate a user as an attacker. A valid user performing non-malicious activity is very unlikely to execute such actions. Thus, malicious activity is blocked without disturbing valid activity of valid users.

[0102] Some embodiments determine or receive an indication of a security risk condition **124** precursor **212** associated with a user account or associated with a user session. Based on the security risk condition, some embodiments determine a potentially malicious action **308** that the user account has permission to perform. In some cases, the potentially malicious action **308**, which is also referred to as an incrimination action **308**, is a privileged action **132**, but not necessarily in every case handled by a given embodiment.

[0103] With a view toward minimizing attack damage, some embodiments perform a what-if analysis for conditional disruption. This is different than approaches which initiate an automated disruption procedure only when it is highly certain that the user is an attacker. As taught herein, a what-if analysis is applied at an earlier stage to identify downstream alerts that could eventually incriminate the user, when the user account has already displayed signs of risk, e.g., as reflected in one or more action, event, or alert logs or streams.

[0104] Consider such a downstream alert that is triggered by a single malicious user operation. By executing this operation, the user would be clearly identified as an attacker. Therefore, permissions to perform this operation can be blocked beforehand. As it is highly unlikely for a legitimate user to perform this operation, no business impact will occur. Furthermore, the attacker initiating this operation effectively provides a signal which incriminates them, thus allowing an embodiment to also employ stringent disruption procedures.

[0105] Some embodiments only deprive permissions to execute rare actions, e.g., adding domain federation (in most accounts) or creating a virtual machine (in non-developer accounts or accounts not driven automatically by a scalable application). Regular daily action is not affected. Some embodiments deprive permissions to perform actions that would have raised the risk to a certain threshold if they were executed. The specific threshold is product-dependent in some embodiments, but in many cases this threshold is lower than the threshold that triggers disabling the user account.

[0106] Some embodiments utilize or implement a contrapositive rule of logic inference: if A implies B, then not B implies not A. Specifically, if an alert or action A implies an attacker, then no attacker implies no alert or action A. In other words, if the user is not an attacker, then the user will not attempt to execute the action that leads to the alert, and the user won't be impacted by the automated attacker-focused disruption taught herein.

[0107] In many cases, alerts do not represent a single isolated action, but rather represent a combination of states and actions. A detection can be broken down into a non-empty set of pairs, consisting of one action and the remaining conditions. The remaining conditions are also referred to as a “partial alert”. These pairs are examples of disruption tuples **468**, and a partial alert is an example of a precursor **212**.

[0108] Win some embodiments, when a medium risk user is identified the embodiment identifies a set of alert types **320** that could potentially incriminate them in the future. For each alert type and pair of <<partial alert with the remaining conditions::action>> the embodiment sets a trigger such that once the partial alert is detected the embodiment will remove the user account's permission to execute the action. This way, if all precursor conditions are met, the user account's attempt to execute the action will be blocked. Since this attempt is logged as usual, the regular automated disruption will be invoked, e.g., the user account will be disabled. As a result, an attacker will be disrupted early, while a legitimate user will not be impacted.

[0109] Some embodiments form **734** disruption tuples **468**. A disruption tuple **468** includes a precursor sequence **402** of user account or user session action(s) **132** and a corresponding incrimination action **308**, **132**. The precursor and the incrimination action correspond in the sense that execution **406** of the precursor followed by execution **406** of the incrimination action strongly indicates a cyberattack **428**. The incrimination action is selected **734**, not only because it corresponds to the precursor sequence **402** as part of an attack indication **428**, but also because the incrimination action is statistically or presumptively more likely (by a specified measure **454** or **318**) to be performed on behalf of an attacker **138** than on behalf of a non-attacker **210**. This likelihood **328** of performance pertains to performance of the incrimination action regardless of preceding actions, unless stated otherwise. The likelihood of performance of the incrimination action plus the preceding precursor sequence in combination by an attacker **138** is already much more likely than the likelihood of such performance by a non-attacker **210**, because that combination-precursor sequence followed by incrimination action-strongly indicates an attack.

[0110] Some disruption tuples **468** include a single incrimination action **308** which chronologically follows the precursor sequence **402**, either as a blocked action or as a blocked and attempted action. For example, one disruption tuple in some systems is <<authentication with compromised credential(s)::creation of new user>>.

[0111] Some disruption tuples **468** include a sequence of incrimination actions **308** which follow the precursor sequence **402**. For example, one disruption tuple in some systems is <<risky user alert, access elevation alert, resource quota increase::virtual machine creation, cryptocurrency

computation>>.

[0112] Some disruption tuples **468** include a non-empty set of incrimination actions **308**, any of which follow the precursor sequence **402**. For example, one disruption tuple in some systems is <<risky user XDR alert, guest invitation XDR alert::high privilege role assignment on the guest user OR subscription transfer>>. As another example, in some systems the incrimination action specifies a set of similar actions **132**, namely, actions which use different syntax but share a result, e.g., actions which upon execution cause a permission change, actions which upon execution create a process, actions which upon execution create an account, actions which upon execution create a network connection, and so on.

[0113] Some embodiments form **734** disruption tuples **468** by analyzing **732** existing high-confidence alerts to identify incrimination actions plus precursor conditions for removing permissions to these incrimination actions.

[0114] Some embodiments form **734** disruption tuples **468** by using alert labels and incident labels to identify sets of alerts **320** plus an incrimination action that appear together for a user account, are very likely to happen during attack, and very unlikely to occur for a valid user account in the absence of an attack. Hence, these alerts plus incrimination action tuples identify the user account as an attacker-driven account.

[0115] Some embodiments are designed to select the disruption **208** in a very specific way, so that the actions **308** being blocked are unlikely to be attempted by a valid non-attacker **210** user account. Hence, the risk of disturbing regular business processes is low, and AFGAD can be employed even when the confidence that the user is an attacker is not very high, which increases the overall coverage of the protection.

[0116] Some embodiments establish **504** (e.g., by calculating a risk score) that the potentially malicious action has a probability, higher than a threshold **454**, to be associated with a malicious attack than to be associated with non-malicious activity. For example, in some scenarios when the action's probability of being performed by an attacker is at least 0.5 greater than the probability of the action being performed by a non-attacker the user driving a request to perform the action **308** the action is deemed to be a malicious user. Then the embodiment removes **214** the user account's permission **404** to perform the potentially malicious action **308**. The action **308** permission **404** is removed **214** before receipt of a request from the user account to perform the potentially malicious action, or at least before the potentially malicious action has been performed, thereby preventing performance of the action **308**.

[0117] In some embodiments and scenarios, the permission block **214** is focused on a specific session ID, or a specific IP address or IP address range.

[0118] In some embodiments the mechanism used to remove permissions includes a conditional access mechanism, such as an Azure® AD Conditional Access Policy employing a Session Control option, or a cloud access security broker, or a similar mechanism (mark of Microsoft Corporation).

[0119] In some scenarios, an embodiment is configured such that if a user sign-ins from an unfamiliar location and the authentication session is classified as risky, then the user is allowed to access an application through a cloud access security broker. In this case, the user account would have read-only access, or would be prevented from downloading or printing, for example.

[0120] Some embodiments automatically and proactively restore permissions that were blocked **214**, either after a predefined time window, or when a risk score is reduced to a predefined level, or both, for example.

[0121] In some embodiments, alerts that are triggered only by states, rather than specific user actions, are excluded from disruption tuples **468**. For example, an alert that sets a threshold on a count of a specific operation (e.g., number of files deleted by the user), is excluded from use as the entirety of a precursor. In some embodiments, alerts where the user is the victim, rather than the attacker, such as clicking on a malicious URL, contributes to the risk, but victim actions are excluded from use as incrimination actions.

[0122] It may happen that an automatic disruption reveals to an attacker the fact that the embodiment has characterized the account as attacker-driven. In response, the attacker might reduce or delay malicious activity. Regardless, the AFGAD will at least slow down the attack, providing security personnel more time to investigate. AFGAD functionality helps ensure that attack disruption measures are taken in a timely manner, thus reducing or preventing damage caused by incriminated attackers. Moreover, since this disruption **208** is very specific, an organization will often be more tolerant of it. AFGAD functionality also allows a security tool to employ less stringent thresholds, which contributes to the protective coverage.

[0123] In one example scenario, a suspicious sign-in alert increases a user's risk evaluation **474** to medium. AFGAD what-if functionality **204** in the system identifies the following downstream activities **308** that will be sufficient to incriminate the user: a "Suspicious email deletion activity" alert, a "Suspicious Email Forwarding Activity" alert. As an automated disruption setup, a non-empty set of the following pairs **468** of <<partial rule::action>> is derived **734**: <<conditions from the alert::email deletion>>, <<conditions from the alert::email forwarding>>. The AFGAD functionality also adds the user account to lists referenced by disruption per each pair **468**. In this example, within the disruption mechanism, when any of the partial rules are detected, a disruption is activated, which removes the user account's permission to create a forwarding rule or delete emails correspondingly. In the general case, some precursors **212** include not only alert(s) but also include additional condition(s) **124**, to be triggered.

[0124] Some embodiments increase protection coverage while maintaining disruption precision. Some embodiments look for a specific activity by looking for particular keywords in activity log entries or activity stream entries, before blocking. Some embodiments block only after learning details of an action, depending on the alert disruption pair **468**. Some block the action only when all the remaining conditions in the partial alert were met, and performing the action will raise an alert. Some embodiments simulate carrying out the blocked action.

[0125] In some embodiments and scenarios, when a user account attempts to perform one of the incrimination actions **308**, the account receives an insufficient permissions error and a corresponding log record is created. This triggers the stringent disruption that disables the user account.

[0126] In some embodiments and scenarios, in response to receiving a request from the user account to perform the potentially malicious action, or another indication that the user attempted to perform the potentially malicious action **308**, an embodiment disables **460** all permissions associated with the user account.

[0127] In some embodiments and scenarios, in response to receiving an indication of a lower risk condition (e.g., an input indicating the user account can be more trusted), an embodiment restores the user account's permission to perform the potentially malicious action.

[0128] In some scenarios, an embodiment sets a trap for an attacker: the embodiment determines that one or more precursor actions have occurred, confirms that a particular potential next action is an incrimination action because it is much more likely to be performed by an attacker than by a non-attacker, and disables the incrimination action while still leaving open a possibility of an attempt to perform the incrimination action, as well as a possibility of other user actions in the account.

[0129] Some embodiments use or provide a method of attacker-focused granular action disruption in a computing system, the method including automatically: observing data representing a performance, in the computing system by a user account, of a sequence of one or more user account actions, the sequence referred to here as a precursor; discerning that the precursor is associated, by an association mechanism in the computing system, with a further user account action which has not yet been performed by the user account after the performance of the precursor, the further user account action referred to here as an incrimination action; establishing that, after the performance of the precursor, the incrimination action is more likely to be performed by an

attacker-driven account than by a non-attacker-driven account; and blocking the user account from performing the incrimination action in the computing system, while maintaining in the computing system a permission for the user account to attempt execution of the incrimination action, and maintaining in the computing system a permission for the user account to execute an additional user account action other than the incrimination action and other than any part of the precursor. [0130] In some scenarios, the user account attempts to perform the incrimination action after it has performed the precursor. As a result of the attempt, a further restriction is placed on the user account, because the embodiment's level of confidence that the user account is an attacker-driven account increased, due to the user account's attempt to perform the incrimination action. Recall that the embodiment has already established that the incrimination action is more likely to be performed by an attacker-driven account than by a non-attacker-driven account.

[0131] To reduce AFGAD response times, some embodiments utilize streamed alerts (a.k.a. detections) instead of logged alerts, or in addition to logged alerts. Some embodiments utilize raw data sources where a targeted disruption can be applied. Some examples include process creation logs where permissions can be removed, or network logs where a firewall rule can be configured to block communication.

[0132] Some embodiments distinguish between alert thresholds and disruption thresholds. In addition to analyzing **732** detection rules in their regular form, disruption tuples are obtained in some systems by fine-tuning thresholds in existing detections so that the same detection rule with a certain threshold value will result in an alert, while a tighter threshold will trigger disruption.

[0133] Some embodiments form **734** disruption tuples employing a machine learning (ML) approach that uses the alerts as features weighted and directly optimized to the available labels. Beyond alerts and anomalies, some embodiments also use preventable actions and risk lowering actions as features to contribute to ML models' performance.

[0134] On top of a risk identification model, some embodiments employ two additional models to help identify the actions **308** to prevent as a means of soft disruption. To calculate a probability that an action is indicative of an attack, some embodiments implement Bayes Theorem and utilize predictions for $P(\text{Action}|\text{Attacker})$ and $P(\text{Action}|\neg\text{Attacker})$. Then, actions will be disrupted if: $P(\text{Attacker}) = P(\text{Action}|\text{Attacker}) / (P(\text{Action}|\text{Attacker}) + P(\text{Action}|\neg\text{Attacker})) > \text{Threshold}$.; "P" is read as "probability", "|" is read as "given", " $\neg$ " is read as "not", and "/" means division. This approach can be applied either to user accounts or to user sessions, or both, depending on the available control mechanism(s) for blocking actions.

[0135] In addition to this prediction derived from learning on a full population of users, some embodiments verify **730** that the action is anomalous for the specific user whose session or account would be blocked, in order to minimize business process disruptions for that specific user.

[0136] Some embodiments form **734** disruption tuples **468** by using risk score **474** models **438**. A first model, also referred to as a high-risk model, identifies high-risk user accounts to disable; disabling an account is a hard disruption. A second model or models, also referred to as a medium-risk model or models, identifies medium risk user accounts plus actions' respective likelihood **314** predictions assuming an attacker plus actions' respective likelihood predictions **316** assuming a valid non-attacker user, for high-resolution adaptive hardening **204**. Some embodiments remove the limitation **214** on the permission(s) after the risk goes back down. These models estimate probabilities, given the signal history of the user/session, of which actions will provide a strong indication that the user is an attacker, so the embodiment can dynamically remove permissions accordingly.

[0137] As an example, in some embodiments and scenarios, a risk score is produced by a user/session risk score model. When the risk score is greater than a high threshold (e.g., 0.9), the embodiment disables the user account or user session, or both. When the risk score is between the high threshold and a low threshold (e.g., 0.2), the embodiment executes the action likelihood models, which produce a list of action(s) **308** to disrupt **214** or a configuration change from a first

configuration to a second configuration, or do both. When the risk score is less than the low threshold, the embodiment restores the first configuration. This continues in a loop, as the system repeatedly updates the risk score.

[0138] In some embodiments, trained machine learning models **438** predict **406** the likelihood **314** of an action **132** being driven by an attacker, or the likelihood **316** of an action **132** being driven by a non-attacker (e.g., in connection with a false alert). These likelihoods **314** and **316** are compared **736** for a given action **132** to check whether they are sufficiently different (e.g., per an absolute amount **454** or a multiplication factor **318**) to establish **504** that the action is deemed to be attacker-driven.

[0139] In some embodiments, the models **438** include one or more of: a long short-term memory (LSTM) model, a sequence model, or a multilayer model that predicts likelihoods of multiple actions **132**.

[0140] Note that an embodiment enhances security even when some actions **132** do not have respective calculated likelihoods **328**, because adding a new disruption tuple **468** incrementally improves security.

[0141] In another example, AFGAD functionality disrupts an attack represented by the following disruption tuple: <<authentication with compromised credential(s)::creation of new user(s)>>. Even though the alert in the precursor authentication arguably provides sufficient confidence to disable the user account, in order to avoid business process disruption, in some tools disabling the users is limited to a small number of users, e.g., five users. As a result, the attackers can often recover from the attack disruption, when they have global admin permissions, by creating new users. Instead of fully disabling the user, a soft disruption procedure under AFGAD **204** changes the role of the user from 'Global Admin' to 'Admin Reader', thus removing their permission to create new users, and gaining time for security personnel to respond without interrupting legitimate activity.

[0142] To help formulate **734** disruption tuples, some embodiments run a retrospective analysis **732** identifying a non-empty set of downstream actions **308** that distinguish attackers from valid users with high confidence, and generate a non-empty set of such actions that can be blocked for each precursor **212** non-empty set of signal(s) or alert(s). Then the embodiment analyzes the available roles **420** and determines how each role could be reduced to a lower-permission role in case the precursor signal **212** occurred. To help ensure minimal disruption for the user when the activity is legitimate, some embodiments go beyond the general statistics by also checking a historic behavior of the user, to verify that reducing the permissions does not interrupt normal (non-attacker) work.

[0143] In another example, AFGAD functionality **204** disrupts an attack represented by the following disruption tuple: <<suspicious events from a business email compromise (BEC) detector where incrimination happens after email is sent::one or more emails is sent>>. A challenge with some BEC detectors is that they are triggered only upon or after transmission of malicious email(s). A trigger in one financial fraud case was a single fraud email, while a trigger in a credential harvesting case was a campaign of multiple emails. In some cases, the disruption tuple description is further refined in response to an observation is that attackers most often used Outlook® Web Access (OWA) (sometimes called Outlook® on the Web) or PowerShell® tools to perform malicious email compromise activities (marks of Microsoft Corporation). The disruption tuple precursor includes a sequence **402** of events that precede sending email(s), and the disruption tuple incrimination actions include Microsoft Exchange® Online activities which are from the same IP address as the precursor sequence and which use OWA or PowerShell capabilities. To accomplish blocking **214**, some embodiments utilize an Exchange® Online Client Access Policy created in a customer's cloud tenant, which will block access using OWA protocol from suspicious IP address(es). As an additional safeguard, some embodiments include a precursor state **124** representing recent use of OWA or PowerShell tools.

[0144] In another example, AFGAD functionality **204** disrupts an attack represented by the

following disruption tuple: <<preceding step(s) of hard disruption chain prior to the final step(s)::final step(s) of hard disruption chain>>. Some embodiments use a what-if approach to remove permissions and prevent the final step(s) at an earlier point in an attack than the point where hard disruption would occur.

[0145] One such chain includes a risky user XDR alert, and a guest invitation XDR alert in the precursor **212**. At this point, the following permissions are blocked to prevent execution of incrimination actions: high privilege role assignment on the guest user (e.g., Owner or Admin), and subscription transfer (disable user in case the subscription transfer alert is on a risky user). In the notation used above, this example is consistent with a disruption tuple <<risky user XDR alert, guest invitation XDR alert::high privilege role assignment on guest user, subscription transfer>>.

[0146] Another such chain includes a risky user detection, an access elevation, and a resource quota increase in at least two locations. At this point, the following incrimination actions are blocked **214**: mass compute (e.g., virtual machine, worker process) creation, cryptocurrency runs. In the notation used above, this example is consistent with a disruption tuple <<risky user detection, access elevation, resource quota increases::virtual machine creation OR worker process creation OR cryptocurrency run>>.

Internet of Things

[0147] In some embodiments, the system **202** is, or includes, an embedded system such as an Internet of Things system. “IoT” or “Internet of Things” means any networked collection of addressable embedded computing or data generation or actuator nodes. An individual node is referred to as an internet of things device **101** or IoT device **101** or internet of things system **102** or IoT system **102**. Such nodes are examples of computer systems **102** as defined herein, and may include or be referred to as a “smart” device, “endpoint”, “chip”, “label”, or “tag”, for example, and IoT may be referred to as a “cyber-physical system”. In the phrase “embedded system” the embedding referred to is the embedding a processor and memory in a device, not the embedding of debug script in source code.

[0148] IoT nodes and systems typically have at least two of the following characteristics: (a) no local human-readable display; (b) no local keyboard; (c) a primary source of input is sensors that track sources of non-linguistic data to be uploaded from the IoT device; (d) no local rotational disk storage-RAM chips or ROM chips provide the only local memory; (e) no CD or DVD drive; (f) being embedded in a household appliance or household fixture; (g) being embedded in an implanted or wearable medical device; (h) being embedded in a vehicle; (i) being embedded in a process automation control system; or (j) a design focused on one of the following: environmental monitoring, civic infrastructure monitoring, agriculture, industrial equipment monitoring, energy usage monitoring, human or animal health or fitness monitoring, physical security, physical transportation system monitoring, object tracking, inventory control, supply chain control, fleet management, or manufacturing. IoT communications may use protocols such as TCP/IP, Constrained Application Protocol (CoAP), Message Queuing Telemetry Transport (MQTT), Advanced Message Queuing Protocol (AMQP), HTTP, HTTPS, Transport Layer Security (TLS), UDP, or Simple Object Access Protocol (SOAP), for example, for wired or wireless (cellular or otherwise) communication. IoT storage or actuators or data output or control may be a target of unauthorized access, either via a cloud, via another network, or via direct local access attempts.

Technical Character

[0149] The technical character of embodiments described herein will be apparent to one of ordinary skill in the art, and will also be apparent in several ways to a wide range of attentive readers. Some embodiments address technical activities such as blocking **214** permission **404** for execution of an action **132** in a computing system by a user account **130** or user session **130**, detecting **310** an attempt **216** by a user account or user session **130** to execute an action **132** in a computing system, computationally analyzing **732** data for form **734** disruption tuple data structures **468**, or executing a machine learning model **438**, which are each an activity deeply rooted in computing technology.

Some of the technical mechanisms discussed include, e.g., machine learning models **438**, statistical models **438**, disruption tuple data structures **468**, association mechanisms **324**, interfaces **326**, and AFGAD software **302**. Some of the technical effects discussed include, e.g., automated disruption of attack activity without disruption of legitimate activity, and automated disruption of attack activity which is not covered by alerts defined in a system. Thus, purely mental processes and activities limited to pen-and-paper are clearly excluded from the scope of any embodiment. Other advantages based on the technical characteristics of the teachings will also be apparent to one of skill from the description provided.

[0150] One of skill understands that automated disruption of attack activity in a computing network **108** or other computing system **102** is technical activity which cannot be performed mentally at all, and cannot be performed manually with the speed and accuracy required in computing systems. This is especially true of computing systems which communicate electronically with the internet **412** and therefore have a large attack surface. Hence, automated disruption technology improvements such as the various examples of AFGAD functionality **204** described herein are improvements to computing technology. One of skill understands that attempting to manually detect and disrupt an attack would create unacceptable delays in permissions blocking, pose severe risks of damage from attacks, and introduce unnecessary and unacceptable human errors. People manifestly lack the speed, accuracy, memory capacity, and specific processing capabilities required to perform automated attacker-focused granular attack disruption as taught herein.

[0151] Different embodiments provide different technical benefits or other advantages in different circumstances, but one of skill informed by the teachings herein will acknowledge that particular technical advantages will likely follow from particular embodiment features or feature combinations, as noted at various points herein. Any generic or abstract aspects are integrated into a practical application such as a SIEM tool or an extended detection and response (XDR) tool. One such tool is an enhanced Microsoft Defender™ XDR coordinated threat protection solution (mark of Microsoft Corporation), but practical applications of the teachings herein are not technologically limited to Microsoft products and services. Practical application opportunities also extend beyond XDR tools to endpoint detection and response (EDR) tools and managed detection and response (MDR) tools, among other tools.

[0152] Some embodiments described herein may be viewed by some people in a broader context. For instance, concepts such as efficiency, reliability, user satisfaction, or waste may be deemed relevant to a particular embodiment. However, it does not follow from the availability of a broad context that exclusive rights are being sought herein for abstract ideas; they are not.

[0153] Rather, the present disclosure is focused on providing appropriately specific embodiments whose technical effects fully or partially solve particular technical problems, such as how to rapidly detect that a cyberattack is underway, how to disrupt the cyberattack without disrupting legitimate activity in the system that is being attacked, and how to extend the system's scope of protection against cyberattacks, e.g., to include protection beyond attacks that are described only in terms of alerts. Other configured storage media, systems, and processes involving efficiency, reliability, user satisfaction, or waste are outside the present scope. Accordingly, vagueness, mere abstractness, lack of technical character, and accompanying proof problems are also avoided under a proper understanding of the present disclosure.

Additional Combinations and Variations

[0154] Any of these combinations of software code, data structures, logic, components, communications, and/or their functional equivalents may also be combined with any of the systems and their variations described above. A process may include any steps described herein in any subset or combination or sequence which is operable. Each variant may occur alone, or in combination with any one or more of the other variants. Each variant may occur with any of the processes and each process may be combined with any one or more of the other processes. Each process or combination of processes, including variants, may be combined with any of the

configured storage medium combinations and variants described above.

[0155] More generally, one of skill will recognize that not every part of this disclosure, or any particular details therein, are necessarily required to satisfy legal criteria such as enablement, written description, or best mode. Also, embodiments are not limited to the particular scenarios, language models, prompts, motivating examples, operating environments, tools, peripherals, software process flows, identifiers, repositories, data structures, data selections, naming conventions, notations, control flows, or other implementation choices described herein. Any apparent conflict with any other patent disclosure, even from the owner of the present subject matter, has no role in interpreting the claims presented in this patent disclosure.

Acronyms, Abbreviations, Names, and Symbols

[0156] Some acronyms, abbreviations, names, and symbols are defined below. Others are defined elsewhere herein, or do not require definition here in order to be understood by one of skill. [0157]

ALU: arithmetic and logic unit [0158] API: application program interface [0159] BIOS: basic input/output system [0160] CD: compact disc [0161] CPU: central processing unit [0162] DVD: digital versatile disk or digital video disc [0163] FPGA: field-programmable gate array [0164] FPU: floating point processing unit [0165] GDPR: General Data Protection Regulation [0166] GPU: graphical processing unit [0167] GUI: graphical user interface [0168] HTTPS: hypertext transfer protocol, secure [0169] IaaS or IAAS: infrastructure-as-a-service [0170] LAN: local area network [0171] OS: operating system [0172] PaaS or PAAS: platform-as-a-service [0173] RAM: random access memory [0174] ROM: read only memory [0175] SIEM: security information and event management [0176] TPU: tensor processing unit [0177] UEFI: Unified Extensible Firmware Interface [0178] UI: user interface [0179] WAN: wide area network

Some Additional Terminology

[0180] Reference is made herein to exemplary embodiments such as those illustrated in the drawings, and specific language is used herein to describe the same. But alterations and further modifications of the features illustrated herein, and additional technical applications of the abstract principles illustrated by particular embodiments herein, which would occur to one skilled in the relevant art(s) and having possession of this disclosure, should be considered within the scope of the claims.

[0181] The meaning of terms is clarified in this disclosure, so the claims should be read with careful attention to these clarifications. Specific examples are given, but those of skill in the relevant art(s) will understand that other examples may also fall within the meaning of the terms used, and within the scope of one or more claims. Terms do not necessarily have the same meaning here that they have in general usage (particularly in non-technical usage), or in the usage of a particular industry, or in a particular dictionary or set of dictionaries. Reference numerals may be used with various phrasings, to help show the breadth of a term. Sharing a reference numeral does not mean necessarily sharing every aspect, feature, or limitation of every item referred to using the reference numeral. Omission of a reference numeral from a given piece of text does not necessarily mean that the content of a Figure is not being discussed by the text. The present disclosure asserts and exercises the right to specific and chosen lexicography. Quoted terms are being defined explicitly, but a term may also be defined implicitly without using quotation marks. Terms may be defined, either explicitly or implicitly, here in the Detailed Description and/or elsewhere in the application file.

[0182] A “computer system” (a.k.a. “computing system”) may include, for example, one or more servers, motherboards, processing nodes, laptops, tablets, personal computers (portable or not), personal digital assistants, smartphones, smartwatches, smart bands, cell or mobile phones, other mobile devices having at least a processor and a memory, video game systems, augmented reality systems, holographic projection systems, televisions, wearable computing systems, and/or other device(s) providing one or more processors controlled at least in part by instructions. The instructions may be in the form of firmware or other software in memory and/or specialized

circuitry.

[0183] A “multithreaded” computer system is a computer system which supports multiple execution threads. The term “thread” should be understood to include code capable of or subject to scheduling, and possibly to synchronization. A thread may also be known outside this disclosure by another name, such as “task,” “process,” or “coroutine,” for example. However, a distinction is made herein between threads and processes, in that a thread defines an execution path inside a process. Also, threads of a process share a given address space, whereas different processes have different respective address spaces. The threads of a process may run in parallel, in sequence, or in a combination of parallel execution and sequential execution (e.g., time-sliced).

[0184] A “processor” is a thread-processing unit, such as a core in a simultaneous multithreading implementation. A processor includes hardware. A given chip may hold one or more processors. Processors may be general purpose, or they may be tailored for specific uses such as vector processing, graphics processing, signal processing, floating-point arithmetic processing, encryption, I/O processing, machine learning, and so on.

[0185] “Kernels” include operating systems, hypervisors, virtual machines, BIOS or UEFI code, and similar hardware interface software.

[0186] “Code” means processor instructions, data (which includes constants, variables, and data structures), or both instructions and data. “Code” and “software” are used interchangeably herein. Executable code, interpreted code, and firmware are some examples of code.

[0187] “Program” is used broadly herein, to include applications, kernels, drivers, interrupt handlers, firmware, state machines, libraries, and other code written by programmers (who are also referred to as developers) and/or automatically generated.

[0188] A “routine” is a callable piece of code which normally returns control to an instruction just after the point in a program execution at which the routine was called. Depending on the terminology used, a distinction is sometimes made elsewhere between a “function” and a “procedure”: a function normally returns a value, while a procedure does not. As used herein, “routine” includes both functions and procedures. A routine may have code that returns a value (e.g., $\sin(x)$) or it may simply return without also providing a value (e.g., void functions).

[0189] “Service” as a noun means a consumable program offering, in a cloud computing environment or other network or computing system environment, which provides resources to multiple programs or provides resource access to multiple programs, or does both. A service implementation may itself include multiple applications or other programs.

[0190] “Cloud” means pooled resources for computing, storage, and networking which are elastically available for measured on-demand service. A cloud may be private, public, community, or a hybrid, and cloud services may be offered in the form of infrastructure as a service (IaaS), platform as a service (PaaS), software as a service (SaaS), or another service. Unless stated otherwise, any discussion of reading from a file or writing to a file includes reading/writing a local file or reading/writing over a network, which may be a cloud network or other network, or doing both (local and networked read/write). A cloud may also be referred to as a “cloud environment” or a “cloud computing environment”.

[0191] “Access” to a computational resource includes use of a permission or other capability to read, modify, write, execute, move, delete, create, or otherwise utilize the resource. Attempted access may be explicitly distinguished from actual access, but “access” without the “attempted” qualifier includes both attempted access and access actually performed or provided.

[0192] Herein, activity by a user refers to activity by a user device or activity by a user account or user session, or by software on behalf of a user, or by hardware on behalf of a user. Activity is represented by digital data or machine operations or both in a computing system. Activity within the scope of any claim based on the present disclosure excludes human actions per se. Software or hardware activity “on behalf of a user” accordingly refers to software or hardware activity on behalf of a user device or on behalf of a user account or a user session or on behalf of another

computational mechanism or computational artifact, and thus does not bring human behavior per se within the scope of any embodiment or any claim.

[0193] “Digital data” means data in a computing system, as opposed to data written on paper or thoughts in a person's mind, for example. Similarly, “digital memory” refers to a non-living device, e.g., computing storage hardware, not to human or other biological memory.

[0194] As used herein, “include” allows additional elements (i.e., includes means comprises) unless otherwise stated.

[0195] “Optimize” means to improve, not necessarily to perfect. For example, it may be possible to make further improvements in a program or an algorithm which has been optimized.

[0196] “Process” is sometimes used herein as a term of the computing science arts, and in that technical sense encompasses computational resource users, which may also include or be referred to as coroutines, threads, tasks, interrupt handlers, application processes, kernel processes, procedures, or object methods, for example. As a practical matter, a “process” is the computational entity identified by system utilities such as Windows® Task Manager, Linux® ps, or similar utilities in other operating system environments (marks of Microsoft Corporation, Linus Torvalds, respectively). “Process” may also be used as a patent law term of art, e.g., in describing a process claim as opposed to a system claim or an article of manufacture (configured storage medium) claim. Similarly, “method” is used herein primarily as a technical term in the computing science arts (a kind of “routine”) but it is also a patent law term of art (akin to a “method”). “Process” and “method” in the patent law sense are used interchangeably herein. Those of skill will understand which meaning is intended in a particular instance, and will also understand that a given claimed process or method (in the patent law sense) may sometimes be implemented using one or more processes or methods (in the computing science sense).

[0197] “Automatically” means by use of automation (e.g., general purpose computing hardware configured by software for specific operations and technical effects discussed herein), as opposed to without automation. In particular, steps performed “automatically” are not performed by hand on paper or in a person's mind, although they may be initiated by a human person or guided interactively by a human person. Automatic steps are performed with a machine in order to obtain one or more technical effects that would not be realized without the technical interactions thus provided. Steps performed automatically are presumed to include at least one operation performed proactively.

[0198] One of skill understands that technical effects are the presumptive purpose of a technical embodiment. The mere fact that calculation is involved in an embodiment, for example, and that some calculations can also be performed without technical components (e.g., by paper and pencil, or even as mental steps) does not remove the presence of the technical effects or alter the concrete and technical nature of the embodiment, particularly in real-world embodiment implementations. Automated disruption operations such as observing **304** precursor satisfaction, discerning **502** that an incrimination action has not yet been executed after the precursor is satisfied, establishing **504** that an incrimination action is sufficiently attacker-focused and hence sufficiently unlikely to be part of legitimate activity, blocking **214** actions, detecting **310** an attempt to perform an action, and many other operations discussed herein (whether recited in the Figures or not), are understood to be inherently digital and computational. A human mind cannot interface directly with a CPU or other processor, or with RAM or other digital storage, to read and write the necessary data to perform the AFGAD steps **700** taught herein even in a hypothetical situation or a prototype situation, much less in an embodiment's real world large computing environment, e.g., an internet-connected environment. This would all be well understood by persons of skill in the art in view of the present disclosure. Moreover, one of skill understands that AFGAD functionality cannot be implemented using merely with conventional tools and steps, because actual implementation requires the use of teachings which were first provided in the present disclosure.

[0199] “Computationally” likewise means a computing device (processor plus memory, at least) is

being used, and excludes obtaining a result by mere human thought or mere human action alone. For example, doing arithmetic with a paper and pencil is not doing arithmetic computationally as understood herein. Computational results are faster, broader, deeper, more accurate, more consistent, more comprehensive, and/or otherwise provide technical effects that are beyond the scope of human performance alone. “Computational steps” are steps performed computationally. Neither “automatically” nor “computationally” necessarily means “immediately”.

“Computationally” and “automatically” are used interchangeably herein.

[0200] “Proactively” means without a direct request from a user, and indicates machine activity rather than human activity. Indeed, a user may not even realize that a proactive step by an embodiment was possible until a result of the step has been presented to the user. Except as otherwise stated, any computational and/or automatic step described herein may also be done proactively.

[0201] “Based on” means based on at least, not based exclusively on. Thus, a calculation based on X depends on at least X, and may also depend on Y.

[0202] Throughout this document, use of the optional plural “(s)”, “(es)”, or “(ies)” means that one or more of the indicated features is present. For example, “processor(s)” means “one or more processors” or equivalently “at least one processor”.

[0203] “At least one” of a list of items means one of the items, or two of the items, or three of the items, and so on up to and including all N of the items, where the list is a list of N items. The presence of an item in the list does not require the presence of the item (or a check for the item) in an embodiment. For instance, if an embodiment of a system is described herein as including at least one of A, B, C, or D, then a system that includes A but does not check for B or C or D is an embodiment, and so is a system that includes A and also includes B but does not include or check for C or D. Similar understandings pertain to items which are steps or step portions or options in a method embodiment. This is not a complete list of all possibilities; it is provided merely to aid understanding of the scope of “at least one” that is intended herein.

[0204] For the purposes of United States law and practice, use of the word “step” herein, in the claims or elsewhere, is not intended to invoke means-plus-function, step-plus-function, or 35 United State Code Section 112 Sixth Paragraph/Section 112(f) claim interpretation. Any presumption to that effect is hereby explicitly rebutted.

[0205] For the purposes of United States law and practice, the claims are not intended to invoke means-plus-function interpretation unless they use the phrase “means for”. Claim language intended to be interpreted as means-plus-function language, if any, will expressly recite that intention by using the phrase “means for”. When means-plus-function interpretation applies, whether by use of “means for” and/or by a court's legal construction of claim language, the means recited in the specification for a given noun or a given verb should be understood to be linked to the claim language and linked together herein by virtue of any of the following: appearance within the same block in a block diagram of the figures, denotation by the same or a similar name, denotation by the same reference numeral, a functional relationship depicted in any of the figures, a functional relationship noted in the present disclosure's text. For example, if a claim limitation recited a “zac widget” and that claim limitation became subject to means-plus-function interpretation, then at a minimum all structures identified anywhere in the specification in any figure block, paragraph, or example mentioning “zac widget”, or tied together by any reference numeral assigned to a zac widget, or disclosed as having a functional relationship with the structure or operation of a zac widget, would be deemed part of the structures identified in the application for zac widgets and would help define the set of equivalents for zac widget structures.

[0206] One of skill will recognize that this disclosure discusses various data values and data structures, and recognize that such items reside in a memory (RAM, disk, etc.), thereby configuring the memory. One of skill will also recognize that this disclosure discusses various algorithmic steps which are to be embodied in executable code in a given implementation, and that such code also

resides in memory, and that it effectively configures any general-purpose processor which executes it, thereby transforming it from a general-purpose processor to a special-purpose processor which is functionally special-purpose hardware.

[0207] Accordingly, one of skill would not make the mistake of treating as non-overlapping items (a) a memory recited in a claim, and (b) a data structure or data value or code recited in the claim. Data structures and data values and code are understood to reside in memory, even when a claim does not explicitly recite that residency for each and every data structure or data value or piece of code mentioned. Accordingly, explicit recitals of such residency are not required. However, they are also not prohibited, and one or two select recitals may be present for emphasis, without thereby excluding all the other data values and data structures and code from residency. Likewise, code functionality recited in a claim is understood to configure a processor, regardless of whether that configuring quality is explicitly recited in the claim.

[0208] Throughout this document, unless expressly stated otherwise any reference to a step in a process presumes that the step may be performed directly by a party of interest and/or performed indirectly by the party through intervening mechanisms and/or intervening entities, and still lie within the scope of the step. That is, direct performance of the step by the party of interest is not required unless direct performance is an expressly stated requirement. For example, a computational step on behalf of a party of interest, such as adding, analyzing, ascertaining, associating, attempting, blocking, categorizing, changing, comparing, confirming, connecting, creating, designating, detecting, discerning, disrupting, establishing, executing, filtering, getting, observing, performing, preventing, prohibiting, raising, requesting, restricting, running, transmitting, verifying (and adds, added, analyzes, analyzed, etc.) with regard to a destination or other subject may involve intervening action, such as the foregoing or such as forwarding, copying, uploading, downloading, encoding, decoding, compressing, decompressing, encrypting, decrypting, authenticating, invoking, and so on by some other party or mechanism, including any action recited in this document, yet still be understood as being performed directly by or on behalf of the party of interest. Example verbs listed here may overlap in meaning or even be synonyms; separate verb names do not dictate separate functionality in every case.

[0209] Whenever reference is made to data or instructions, it is understood that these items configure a computer-readable memory and/or computer-readable storage medium, thereby transforming it to a particular article, as opposed to simply existing on paper, in a person's mind, or as a mere signal being propagated on a wire, for example. For the purposes of patent protection in the United States, a memory or other storage device or other computer-readable storage medium is not a propagating signal or a carrier wave or mere energy outside the scope of patentable subject matter under United States Patent and Trademark Office (USPTO) interpretation of the In re Nuijten case. No claim covers a signal per se or mere energy in the United States, and any claim interpretation that asserts otherwise in view of the present disclosure is unreasonable on its face. Unless expressly stated otherwise in a claim granted outside the United States, a claim does not cover a signal per se or mere energy.

[0210] Moreover, notwithstanding anything apparently to the contrary elsewhere herein, a clear distinction is to be understood between (a) computer readable storage media and computer readable memory, on the one hand, and (b) transmission media, also referred to as signal media, on the other hand. A transmission medium is a propagating signal or a carrier wave computer readable medium. By contrast, computer readable storage media and computer readable memory and computer readable storage devices are not propagating signal or carrier wave computer readable media. Unless expressly stated otherwise in the claim, "computer readable medium" means a computer readable storage medium, not a propagating signal per se and not mere energy.

[0211] An "embodiment" herein is an example. The term "embodiment" is not interchangeable with "the invention". Embodiments may freely share or borrow aspects to create other embodiments (provided the result is operable), even if a resulting combination of aspects is not

explicitly described per se herein. Requiring each and every permitted combination to be explicitly and individually described is unnecessary for one of skill in the art, and would be contrary to policies which recognize that patent specifications are written for readers who are skilled in the art. Formal combinatorial calculations and informal common intuition regarding the number of possible combinations arising from even a small number of combinable features will also indicate that a large number of aspect combinations exist for the aspects described herein. Accordingly, requiring an explicit recitation of each and every combination would be contrary to policies calling for patent specifications to be concise and for readers to be knowledgeable in the technical fields concerned.

REMARKS REGARDING REFERENCE NUMERALS

[0212] Reference numerals are provided for convenience and in support of the drawing figures and as part of the text of the specification, which collectively describe aspects of embodiments by reference to multiple items. Items which do not have a unique reference numeral may nonetheless be part of a given embodiment. For better legibility of the text, a given reference numeral is recited near some, but not all, recitations of the referenced item in the text. The same reference numeral may be used with reference to different examples or different instances of a given item.

[0213] The following remarks pertain to particular reference numerals: [0214] **100** operating environment, also referred to as computing environment; includes one or more systems **102** [0215] **101** machine in a system **102**, e.g., any device having at least a processor **110** and having a distinct identifier such as an IP address or a MAC (media access control) address; may be a physical machine or be a virtual machine implemented on physical hardware [0216] **102** computer system, also referred to as a “computational system” or “computing system”, and when in a network may be referred to as a “node” [0217] **104** users, e.g., user of an enhanced system **202** [0218] **106** peripheral device [0219] **108** network generally, including, e.g., LANs, WANs, software-defined networks, clouds, and other wired or wireless networks [0220] **110** processor or non-empty set of processors; includes hardware [0221] **112** computer-readable storage medium, e.g., RAM, hard disks; also referred to as storage device [0222] **114** removable configured computer-readable storage medium [0223] **116** instructions executable with processor; may be on removable storage media or in other memory (volatile or nonvolatile or both) [0224] **118** digital data in a system **102**; data structures, values, source code, and other examples are discussed herein [0225] **120** kernel(s), e.g., operating system(s), BIOS, UEFI, device drivers; also refers to an execution engine such as a language runtime [0226] **122** software tools, software applications, security controls; hardware tools; computational [0227] **124** system state, as represented within a computing system; also referred to as a condition; digital [0228] **126** display screens, also referred to as “displays” [0229] **128** computing hardware not otherwise associated with a reference numeral **106**, **108**, **110**, **112**, **114** [0230] **130** user account or user session, as represented within a computing system [0231] **132** computational action, as represented or effectuated in a computing system [0232] **136** cloud, also referred to as cloud environment or cloud computing environment [0233] **202** enhanced computing system, i.e., system **102** enhanced with functionality **204** as taught herein [0234] **204** attacker-focused granular action disruption functionality (also referred to as AFGAD functionality **204**, AFGAD what-if functionality **204**, high-resolution adaptive hardening **204**, or functionality **204**), e.g., software or specialized hardware which performs or is configured to perform steps **304**, **504**, and **214**, or steps **304**, **214**, and **602**, or steps **504** and **214**, or steps **214**, **310**, and **312**, or any software or hardware which performs or is configured to perform an attacker-focused automatic disruption activity first disclosed herein, or to perform a novel method **700** first disclosed herein [0235] **500** flowchart; **500** also refers to AFGAD methods that are illustrated by or consistent with the FIG. 5 flowchart or any variation of the FIG. 5 flowchart described herein; all AFGAD method steps are computational, not human activity [0236] **600** flowchart; **600** also refers to AFGAD methods that are illustrated by or consistent with the FIG. 6 flowchart or any variation of the FIG. 6 flowchart described herein; all AFGAD method steps are computational, not human activity [0237]

700 flowchart; **700** also refers to AFGAD methods that are illustrated by or consistent with the FIG. 7 flowchart, which incorporates the FIG. 6 flowchart, the FIG. 5 flowchart, and all other steps taught herein, or methods that are illustrated by or consistent with any variation of the FIG. 7 flowchart described herein; all AFGAD method steps are computational, not human activity [0238] **740** any step or item discussed in the present disclosure that has not been assigned some other reference numeral; **740** may thus be shown expressly as a reference numeral for various steps or items or both, and may be added as a reference numeral (in the current disclosure or any subsequent patent application which claims priority to the current disclosure) for various steps or items or both without thereby adding new matter

CONCLUSION

[0239] Some embodiments provide attacker-focused granular disruption functionality **204** to detect and defend a system **202** against cyberattacks. An embodiment observes **304** that a user account or session **130** has executed one or more precursor actions **402** but has not executed an incrimination action **308**. The incrimination action is more likely, by at least a specified amount **454** or **318**, to be performed by an attacker-driven **322** account or session than by a non-attacker-driven **322** account or session. Performance of the incrimination action is preemptively blocked **214**, without disabling the account or session. The block is focused in many cases on only the incrimination action; in some scenarios the embodiment also blocks **214** performance of similar actions **132**. After an attempt **216** by the account or session to perform the blocked incrimination action, larger blocks **218** on activity are enforced, up to and including disablement **460**, because the attempt **216** indicates strongly that the account or session **130** is being driven by an attacker **138**.

[0240] Embodiments are understood to also themselves include or benefit from tested and appropriate security controls and privacy controls such as the General Data Protection Regulation (GDPR). Use of the tools and techniques taught herein can be used together with such controls.

[0241] Although Microsoft technology is used in some motivating examples, the teachings herein are not limited to use in technology supplied or administered by Microsoft. Under a suitable license, for example, the present teachings could be embodied in software or services provided by other cloud service providers.

[0242] Although particular embodiments are expressly illustrated and described herein as processes, as configured storage media, or as systems, it will be appreciated that discussion of one type of embodiment also generally extends to other embodiment types. For instance, the descriptions of processes in connection with the Figures also help describe configured storage media, and help describe the technical effects and operation of systems and manufactures like those discussed in connection with other Figures. It does not follow that any limitations from one embodiment are necessarily read into another. In particular, processes are not necessarily limited to the data structures and arrangements presented while discussing systems or manufactures such as configured memories.

[0243] Those of skill will understand that implementation details may pertain to specific code, such as specific thresholds, comparisons, specific kinds of platforms or programming languages or architectures, specific scripts or other tasks, and specific computing environments, and thus need not appear in every embodiment. Those of skill will also understand that program identifiers and some other terminology used in discussing details are implementation-specific and thus need not pertain to every embodiment. Nonetheless, although they are not necessarily required to be present here, such details may help some readers by providing context and/or may illustrate a few of the many possible implementations of the technology discussed herein.

[0244] With due attention to the items provided herein, including technical processes, technical effects, technical mechanisms, and technical details which are illustrative but not comprehensive of all claimed or claimable embodiments, one of skill will understand that the present disclosure and the embodiments described herein are not directed to subject matter outside the technical arts, or to any idea of itself such as a principal or original cause or motive, or to a mere result per se, or to a

mental process or mental steps, or to a business method or prevalent economic practice, or to a mere method of organizing human activities, or to a law of nature per se, or to a naturally occurring thing or process, or to a living thing or part of a living thing, or to a mathematical formula per se, or to isolated software per se, or to a merely conventional computer, or to anything wholly imperceptible or any abstract idea per se, or to insignificant post-solution activities, or to any method implemented entirely on an unspecified apparatus, or to any method that fails to produce results that are useful and concrete, or to any preemption of all fields of usage, or to any other subject matter which is ineligible for patent protection under the laws of the jurisdiction in which such protection is sought or is being licensed or enforced.

[0245] Reference herein to an embodiment having some feature X and reference elsewhere herein to an embodiment having some feature Y does not exclude from this disclosure embodiments which have both feature X and feature Y, unless such exclusion is expressly stated herein. All possible negative claim limitations are within the scope of this disclosure, in the sense that any feature which is stated to be part of an embodiment may also be expressly removed from inclusion in another embodiment, even if that specific exclusion is not given in any example herein. The term “embodiment” is merely used herein as a more convenient form of “process, system, article of manufacture, configured computer readable storage medium, and/or other example of the teachings herein as applied in a manner consistent with applicable law.” Accordingly, a given “embodiment” may include any combination of features disclosed herein, provided the embodiment is consistent with at least one claim.

[0246] Not every item shown in the Figures need be present in every embodiment. Conversely, an embodiment may contain item(s) not shown expressly in the Figures. Although some possibilities are illustrated here in text and drawings by specific examples, embodiments may depart from these examples. For instance, specific technical effects or technical features of an example may be omitted, renamed, grouped differently, repeated, instantiated in hardware and/or software differently, or be a mix of effects or features appearing in two or more of the examples. Functionality shown at one location may also be provided at a different location in some embodiments; one of skill recognizes that functionality modules can be defined in various ways in a given implementation without necessarily omitting desired technical effects from the collection of interacting modules viewed as a whole. Distinct steps may be shown together in a single box in the Figures, due to space limitations or for convenience, but nonetheless be separately performable, e.g., one may be performed without the other in a given performance of a method.

[0247] Reference has been made to the figures throughout by reference numerals. Any apparent inconsistencies in the phrasing associated with a given reference numeral, in the figures or in the text, should be understood as simply broadening the scope of what is referenced by that numeral. Different instances of a given reference numeral may refer to different embodiments, even though the same reference numeral is used. Similarly, a given reference numeral may be used to refer to a verb, a noun, and/or to corresponding instances of each, e.g., a processor **110** may process **110** instructions by executing them.

[0248] As used herein, terms such as “a”, “an”, and “the” are inclusive of one or more of the indicated item or step. In particular, in the claims a reference to an item generally means at least one such item is present and a reference to a step means at least one instance of the step is performed. Similarly, “is” and other singular verb forms should be understood to encompass the possibility of “are” and other plural forms, when context permits, to avoid grammatical errors or misunderstandings.

[0249] Headings are for convenience only; information on a given topic may be found outside the section whose heading indicates that topic.

[0250] All claims and the abstract, as filed, are part of the specification. The abstract is provided for convenience and for compliance with patent office requirements; it is not a substitute for the claims and does not govern claim interpretation in the event of any apparent conflict with other parts of

the specification. Similarly, the summary is provided for convenience and does not govern in the event of any conflict with the claims or with other parts of the specification. Claim interpretation shall be made in view of the specification as understood by one of skill in the art; it is not required to recite every nuance within the claims themselves as though no other disclosure was provided herein.

[0251] To the extent any term used herein implicates or otherwise refers to an industry standard, and to the extent that applicable law requires identification of a particular version of such as standard, this disclosure shall be understood to refer to the most recent version of that standard which has been published in at least draft form (final form takes precedence if more recent) as of the earliest priority date of the present disclosure under applicable patent law.

[0252] While exemplary embodiments have been shown in the drawings and described above, it will be apparent to those of ordinary skill in the art that numerous modifications can be made without departing from the principles and concepts set forth in the claims, and that such modifications need not encompass an entire abstract concept. Although the subject matter is described in language specific to structural features and/or procedural acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific technical features or acts described above the claims. It is not necessary for every means or aspect or technical effect identified in a given definition or example to be present or to be utilized in every embodiment. Rather, the specific features and acts and effects described are disclosed as examples for consideration when implementing the claims.

[0253] All changes which fall short of enveloping an entire abstract idea but come within the meaning and range of equivalency of the claims are to be embraced within their scope to the full extent permitted by law.

Claims

1. A method of attacker-focused granular action disruption in a computing system, the method comprising automatically: observing data representing a performance, in the computing system by a user account or a user session, of at least one precursor event; discerning that the at least one precursor event is associated in the computing system with an incrimination action which has not yet been performed by the user account or the user session after the performance of the at least one precursor event; establishing that, after the performance of the at least one precursor event, the incrimination action is more likely to be performed by an attacker-driven account or an attacker-driven user session than by a non-attacker-driven account or a non-attacker-driven user session; and blocking the user account or the user session from performing the incrimination action in the computing system, while maintaining in the computing system a permission for the user account or the user session to attempt execution of the incrimination action, and maintaining in the computing system a permission for the user account or the user session to execute an additional action other than the incrimination action and other than any part of the precursor.

2. The method of claim 1, further comprising: detecting an attempt by the user account or the user session to perform the blocked incrimination action; and in response to detecting the attempt, restricting a capability of the user account or the user session in addition to the blocking the user account or the user session from performing the incrimination action.

3. The method of claim 1, wherein the establishing comprises confirming that an attacker likelihood of the incrimination action is N times a non-attacker likelihood of the incrimination action, where N is at least five, the attacker likelihood is computed from a likelihood of the incrimination action being performed by an attacker-driven account or an attacker-driven session, and the non-attacker likelihood is computed from a likelihood of the incrimination action being performed by at least one account or at least one session that is not categorized as an attacker-driven account or an attacker-driven session.

4. The method of claim 1, wherein blocking the user account or the user session from performing the incrimination action comprises at least one of: blocking an action which is configured to upon execution add an additional user account to the computing system; blocking an action which is configured to upon execution run a specified process; blocking an action which is configured to upon execution change a resource configuration in the computing system; blocking an action which is configured to upon execution create a resource in a cloud computing environment; blocking an action which is configured to upon execution request an internet connection; or blocking an action which is configured to upon execution transmit data over an internet connection.
5. The method of claim 1, wherein blocking the user account or the user session from performing the incrimination action comprises at least one of: changing an association in the computing system between a role and the user account; prohibiting changing an association in the computing system between a role and a different user account than the user account; preventing creation of a virtual machine; prohibiting email activity from a specified IP address; prohibiting email activity from a specified shell; or prohibiting email activity from a specified web app.
6. The method of claim 1, comprising at least one of: the blocking is completed before any security alert is automatically raised in response to the performance of the at least one precursor event; or the blocking is completed before any security alert which is automatically raised in response to the performance of the at least one precursor event has been designated a part of an incident.
7. The method of claim 1, wherein no alert is raised in the computing system in response to the performance of the at least one precursor event.
8. The method of claim 1, further comprising: ascertaining that a precursor state is also associated in the computing system with the incrimination action; and verifying that the precursor state is present in the computing system before blocking the user account or the user session from performing the incrimination action in the computing system.
9. A computing system, comprising: at least one digital memory; a first description residing in and configuring the at least one digital memory, the first description representing a precursor sequence of one or more precursor events; a second description residing in and configuring the at least one digital memory, the second description comprising second data representing an incrimination action which is distinct from the precursor sequence; at least one processor in operable communication with the at least one digital memory, the at least one processor configured to perform an attacker-focused granular action disruption method which comprises (a) observing third data which represents a performance of the precursor sequence by a user account or a user session, (b) discerning fourth data which represents that the incrimination action was not performed by the user account or the user session after the performance of the precursor sequence by the user account or the user session, or discerning a lack of fifth data which represents that the incrimination action has been performed by the user account or the user session after the performance of the precursor sequence by the user account or the user session, or both, (c) establishing that, after the performance of the precursor sequence, the incrimination action is at least three times more likely to be performed by an attacker-driven account or an attacker-driven session than by a non-attacker-driven account or a non-attacker-driven session, and (d) blocking the user account or the user session from performing the incrimination action, while (e) maintaining a permission for the user account or the user session to attempt execution of the incrimination action, and (f) maintaining a permission for the user account or the user session to execute an additional action other than the incrimination action and other than any part of the precursor sequence.
10. The computing system of claim 9, wherein the establishing comprises getting a likelihood that the user account or the user session has performed the incrimination action prior to the performance of the precursor sequence by the user account or the user session.
11. The computing system of claim 9, wherein the at least one processor is also configured to form the first description and the second description at least in part by computationally analyzing at least one of: a non-empty set of events and alerts; a non-empty set of alerts; or a non-empty event log.

- 12.** The computing system of claim 9, wherein the at least one processor is also configured to form the first description and the second description at least in part by correlating, with a correlation model, candidate precursor events with a candidate incrimination action.
- 13.** The computing system of claim 12, wherein the correlation model comprises at least one of: a statistical model, or a machine learning model.
- 14.** The computing system of claim 9, wherein the at least one processor is also configured to form the first description and the second description at least in part by filtering out a candidate incrimination action after determining that the candidate incrimination action is more likely, by a threshold likelihood, to be performed by a non-attacker-driven user account or a non-attacker-driven user session than by an attacker-driven user account or an attacker-driven user session.
- 15.** The computing system of claim 9, wherein the establishing comprises confirming that an attacker likelihood of the incrimination action is N times a non-attacker likelihood of the incrimination action, where N is at least fifteen, the attacker likelihood is computed from a likelihood of the incrimination action being performed by an attacker-driven account or an attacker-driven user session, and the non-attacker likelihood is computed from a likelihood of the incrimination action being performed by at least one account or at least one user session other than the user account or the user session that is not categorized as an attacker-driven account or an attacker-driven session.
- 16.** A computer-readable storage device configured with data and instructions which upon execution by a processor perform a method of attacker-focused granular action disruption in a computing system, the method comprising automatically: observing data representing a performance, in the computing system by a user account or a user session, of a precursor sequence of one or more events; discerning that the precursor is associated, by an association mechanism in the computing system, with an incrimination action which has not yet been performed by the user account or the user session after the performance of the precursor sequence; establishing that, after the performance of the precursor sequence, the incrimination action is at least four times as likely to be performed by an attacker-driven account or an attacker-driven session as the incrimination action is to be performed by a non-attacker-driven account or a non-attacker-driven session; and blocking the user account or the user session from performing the incrimination action in the computing system, while maintaining in the computing system a permission for the user account or the user session to attempt execution of the incrimination action, and maintaining in the computing system a permission for the user account or the user session to execute an additional action other than the incrimination action and other than any part of the precursor.
- 17.** The computer-readable storage device of claim 16, wherein the method further comprises detecting an attempt by the user account or the user session to perform the blocked incrimination action, and in response to detecting the attempt, restricting a capability of the user account or the user session in addition to blocking the user account or the user session from performing the incrimination action, and wherein restricting the capability comprises at least one of: disabling the user account or the user session; adding an authentication requirement to the user account or the user session; adding an authorization requirement to the user account or the user session; or simulating to the user account or the user session a simulated performance of the blocked incrimination action and not actually performing the blocked incrimination action.
- 18.** The computer-readable storage device of claim 16, wherein the method further comprises obtaining, at least in part by executing a machine learning model, at least one of: a likelihood that the incrimination action is performed by an attacker-driven account or an attacker-driven session; or a likelihood that the incrimination action is performed by an account or a session that is not categorized as an attacker-driven account or an attacker-driven session.
- 19.** The computer-readable storage device of claim 16, wherein the establishing comprises establishing that after the performance of the precursor sequence, the incrimination action is at least ten times as likely to be performed by an attacker-driven account or an attacker-driven session as

the incrimination action was to be performed by the user account or the user session prior to the performance of the precursor sequence by the user account or the user session.

20. The computer-readable storage device of claim 16, wherein the method further comprises: detecting an attempt by the user account or the user session to perform the blocked incrimination action; and in response to detecting the attempt, automatically and proactively categorizing the user account or the user session as an attacker-driven account or an attacker-driven session.
