

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250284974

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Dhakal; Aditya et al.

DISTRIBUTED GRAPH SAMPLING

Abstract

A distributed graph sampling system includes storage nodes that store a graph, a network switch that orchestrates accelerated sampling of the graph, and a compute node that processes a sample of the graph with a GNN. The graph sampling operations are distributed across the storage nodes. Each storage node may store a subgraph (e.g., a portion of the graph) and generate a sample of the subgraph.

Inventors: Dhakal; Aditya (Santa Clarita, CA), Chalamalasetti; Sai Rahul (Milpitas, CA), Gerstman; Liad (Kiryat Bialik, IL), Milojevic; Dejan (Milpitas, CA)

Applicant: Hewlett Packard Enterprise Development LP (Spring, TX)

Family ID: 96949129

Appl. No.: 18/601531

Filed: March 11, 2024

Publication Classification

Int. Cl.: G06N3/10 (20060101); H04L12/64 (20060101)

U.S. Cl.:

CPC G06N3/10 (20130101); H04L12/6418 (20130101); H04L2012/6443 (20130101)

Background/Summary

BACKGROUND

[0001] Graph neural networks (GNNs) are a class of machine learning algorithms designed to process and analyze data represented as graphs. They have gained prominence in various domains,

such as traffic analysis, social networks, recommendation systems, biology, etc. GNNs operate by aggregating information from nodes in a graph, enabling them to identify relationships and patterns within the data. A GNN includes multiple layers, each of which refines the node embeddings of a graph through a combination of node features and information from neighboring nodes. GNNs are useful for tasks such as node classification, link prediction, and graph classification.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0002] Aspects of the present disclosure are best understood from the following detailed description when read with the accompanying figures.

[0003] FIG. 1 is a diagram of a distributed graph sampling system, according to some implementations.

[0004] FIG. 2 is a diagram illustrating the partitioning of a graph, according to some implementations.

[0005] FIG. 3 is a sequence diagram of a distributed graph sampling method, according to some implementations.

[0006] FIG. 4 is a diagram of a network switch, according to some implementations.

[0007] FIG. 5 is a diagram of a network switch graph sampling method, according to some implementations.

[0008] FIG. 6 is a diagram of a storage node, according to some implementations.

[0009] FIG. 7 is a sequence diagram of a storage node graph sampling method, according to some implementations.

[0010] FIG. 8 is a diagram of a NIC graph sampling method, according to some implementations.

[0011] Corresponding numerals and symbols in the different figures generally refer to corresponding parts unless otherwise indicated. The figures are drawn to clearly illustrate the relevant aspects of the disclosure and are not necessarily drawn to scale.

DESCRIPTION

[0012] The following disclosure provides many different examples for implementing different features. Specific examples of components and arrangements are described below to simplify the present disclosure. These are, of course, merely examples and are not intended to be limiting.

[0013] A graph represents the relations between a collection of entities. Specifically, a graph includes nodes (corresponding to the entities) and edges (corresponding to the relations between the entities). A graph neural network (GNN) includes multiple layers. When processing a graph with a GNN, the network operates in multiple iterations, where each layer updates the embeddings of nodes in the graph. At each layer, the GNN aggregates information from neighboring nodes, updates the embedding of each node, and then moves to the next layer. This process continues for a fixed number of layers or until a convergence criterion is met.

[0014] A graph may consume a large amount of memory, such that a graph may not fit into the memory of a processor or an accelerator. Because of this, a graph may be distributed across multiple storage nodes that, collectively, store the graph. Specifically, the graph elements (i.e., graph nodes and graph edges) of a graph may be distributed across the storage nodes. Distributing a graph across multiple storage nodes may raise new challenges for some operations, including graph sampling.

[0015] The present disclosure describes techniques for accelerating distributed graph sampling. A distributed graph sampling system includes storage nodes that store a graph, a network switch that orchestrates accelerated sampling of the graph, and a compute node that processes a sample of the graph with a GNN. The graph sampling operations are distributed across the storage nodes, which may increase the scalability of graph sampling. In some circumstances where a graph is large,

processing a sample of the graph with a GNN may be more efficient than processing the entirety of the graph with a GNN, while still affording a desired level of accuracy. Additionally, distributing graph sampling may make the training of large-scale GNNs more feasible.

[0016] Each storage node may store a subgraph (e.g., a portion of the graph) and generate a sample of the subgraph. The network switch may be a Smart Switch that performs both network switching and also programmed computations. The network switch receives the subgraph samples from the storage nodes and aggregates the subgraph samples into a combined graph sample of the graph. Additionally, the network switch optimizes the graph sample by removing duplicate graph edges from the graph sample. In some implementations, a hash table for the graph is stored at the network switch and is used to identify the duplicate graph edges in the graph sample. The graph sample may then be sent to the compute node for processing with a GNN.

[0017] An accelerator of a storage node may perform a subgraph sampling operation. In some implementations, the storage nodes use a shared memory architecture, in which an accelerator of a storage node may directly access a shared memory region of the storage node, e.g., via a compute express link (CXL) interconnection.

[0018] FIG. 1 is a diagram of a distributed graph sampling system **100**, according to some implementations. The distributed graph sampling system **100** may be part of a computing environment, such as an HPC environment, that utilizes a client-server architecture to distribute and accelerate graph sampling across a network of interconnected components. The architecture of the distributed graph sampling system **100** enables efficient storage and sampling of a large-scale graph (and processing of graph samples for the graph), which would be impractical for single-node systems due to memory constraints. The distributed graph sampling system **100** includes multiple compute nodes **102** (including compute nodes **102A**, **102B**, . . . and **102N**), a network fabric **104**, and multiple storage nodes **106** (including storage nodes **106A**, **106B**, . . . and **106N**).

[0019] The compute nodes **102** are the primary processing units within the distributed graph sampling system **100** for processing graph samples with graph neural network (GNN) algorithms. In some circumstances where a graph is large, processing a graph sample with GNN algorithms may be more efficient than processing the entirety of the graph with GNN algorithms, while still affording a desired level of accuracy. Each compute node **102** may be a standalone unit (e.g., high-performance server) equipped with processors capable of handling complex computational tasks such as updating graph node embeddings, aggregating neighborhood information, and evaluating a GNN's loss function. Additionally, or alternatively, the compute nodes **102** may include specialized hardware accelerators to further enhance their processing capabilities for machine learning tasks. Further, the compute nodes **102** include local memory where intermediary computation results may be temporarily stored.

[0020] The network fabric **104** is a high-bandwidth communication infrastructure that interconnects the compute nodes **102** and the storage nodes **106**. It is designed to facilitate the quick and reliable transfer of large volumes of data required for distributed graph storage and sampling. The network fabric **104** may employ low-latency protocols and advanced networking technologies, such as InfiniBand or high-speed Ethernet, to maintain synchronicity and coordination across the distributed graph sampling system **100**.

[0021] The storage nodes **106** are dedicated data storage units that store and maintain a large-scale graph in a distributed manner. Each storage node **106** may be a standalone unit (e.g., network-attached storage server) equipped with a large-capacity memory that stores a subgraph of the graph, e.g., a subset of the graph elements of the graph. The distribution of a graph across multiple storage nodes **106** enables efficient storage of a graph that may be infeasible to store on a single node. Additionally, distributing the graph across multiple storage nodes **106** enables parallel graph sampling operations, as different subgraphs of the graph can be retrieved and sampled concurrently.

[0022] The storage nodes **106** are equipped with their own processing capabilities (e.g., CPUs, accelerators, and the like) to assist in preprocessing or sampling tasks before transferring graph

data through the network fabric **104** to the compute nodes **102**. This localized processing at the storage nodes **106** helps alleviate bottlenecks by reducing the amount of raw data that needs to traverse the distributed graph sampling system **100**. Furthermore, and as subsequently described in greater detail, the architecture of the storage nodes **106** allows for the incorporation of Shared Virtual Memory (SVM), enabling the accelerators of the storage nodes **106** to directly access desired graph data from memory of the storage nodes **106**, without involving the CPUs of the storage nodes **106** in data movement tasks.

[0023] The network fabric **104** includes network switches **108** (including network switches **108A**, **108B**, . . . and **108N**), which connect the storage nodes **106** to the compute nodes **102**. The network switches **108** are Smart Switches (with embedded computational capabilities) that are used for both networking and computing operations. Specifically, the network switches **108** may switch network packets between the storage nodes **106** and the compute nodes **102**, and may also perform computing operations for graph sampling.

[0024] As subsequently described in greater detail, a compute node **102** initiates a graph sampling operation by communicating with the storage nodes **106** through a network switch **108** of the network fabric **104**. The network switch **108** requests specific subgraph samples from the storage nodes **106**. Upon receiving these requests, the storage nodes **106** generate appropriate subgraph samples from their stored subgraphs of the graph. The subgraph samples are then transmitted to the network switch **108**. The embedded computational functionalities of the network switch **108** may then be used to aggregate the subgraph samples into a graph sample. The network switch **108** may also optimize the graph sample for processing by the compute node **102** by performing graph edge deduplication. Specifically, the network switch **108** may utilize intelligent hashing mechanisms to remove duplicate graph edges from the graph sample, which may help manage data flow across the distributed graph sampling system **100**. The compute node **102** then receives the optimized graph sample (which is representative of the large-scale graph) from the network switch **108** and processes the graph sample with GNN algorithms.

[0025] The network fabric **104** may have any topology. For example, the network fabric **104** may have a spine-leaf architecture, in which the network switches **108** may be spine switches or leaf switches. In the illustrated example, each compute node **102** is connected to a respective network switch **108**, which is connected to each of the storage nodes **106**. For example, the compute node **102A** is connected to each of the storage nodes **106** via the network switch **108A**; the compute node **102B** is connected to each of the storage nodes **106** via the network switch **108B**; and the compute node **102N** is connected to each of the storage nodes **106** via the network switch **108N**. Thus, each network switch **108** may (for a respective compute node **102**) aggregate samples that represent the entire graph stored in the storage nodes **106**. Other topologies may be utilized. For example, the network fabric **104** may have a layered architecture, including multiple tiers of the network switches **108**.

[0026] FIG. 2 is a diagram illustrating the partitioning of a graph **202**, according to some implementations. Specifically, the graph **202** is partitioned by a graph partitioning algorithm **204** into multiple subgraphs **206** (including subgraphs **206A**, **206B**, . . . and **206N**), each of which is a portion of the graph **202**. Doing so allows the subgraphs **206** to be distributed across the storage nodes of a distributed graph sampling system. Each subgraph **206** may be stored on a respective storage node. For example, the subgraph **206A** may be stored on the storage node **106A** of the distributed graph sampling system **100** (see FIG. 1); the subgraph **206B** may be stored on the storage node **106B**; and the subgraph **206N** may be stored on the storage node **106N**.

[0027] The graph partitioning algorithm **204** may include several operations. The graph partitioning algorithm **204** may distribute the graph elements of the graph **202** across the subgraphs **206**. The graph elements may be distributed uniformly so as to reduce the risk of a particular subgraph **206** becoming a hotspot. The distribution may be performed in a manner that minimizes the number of inter-subgraph connections between the graph nodes of the subgraphs (which may be referred to as

“edge cuts”). Reducing edge cuts may lessen the amount of inter-storage-node communication that is needed for sampling the graph **202**. The distribution may also maintain local neighborhoods of graph nodes within the same subgraph **206**, when feasible.

[0028] When partitioning the graph **202**, the graph partitioning algorithm **204** also generates a graph hash table **208**. The graph hash table **208** acts as an index of which subgraphs **206** contain which graph elements, which may be used to identify which storage nodes (used to store the subgraphs **206**) are storing which graph elements. Thus, the graph hash table **208** may be used to quickly locate graph elements across the storage nodes used to store the subgraphs **206**. The graph hash table **208** may be stored on one or more network switches, which may use the graph hash table **208** to map the graph elements to storage nodes that store the graph elements. For example, the graph hash table **208** may be stored on the network switches **108** of the distributed graph sampling system **100** (see FIG. 1). Additionally, the network switches may use the graph hash table **208** as a routing table for forwarding packets to the storage nodes. Doing so may increase the speed of accessing a graph element from its storage location.

[0029] The graph elements of a graph mode may have identifiers. The graph hash table **208** may be a dictionary having a key-value structure. The keys of the graph hash table **208** include the identifiers of the graph elements. The values of the graph hash table **208** include the locations of the respective graph elements (e.g., indications of which subgraphs contains which graph elements). As such, the graph hash table **208** may be used to map the identifiers of the graph elements to the storage nodes which store the subgraphs **206** containing the graph elements. The graph partitioning algorithm **204** may reduce hash collisions in the keys of the graph hash table **208** through any suitable collision resolution technique, such as chaining, open addressing, or the like.

[0030] The graph partitioning algorithm **204** may be performed offline (e.g., by a computing node). The graph partitioning algorithm **204** may be informed by the architecture of the graph sampling system on which the distributed graph will be stored. For example, the graph partitioning algorithm **204** may accept, as input, information about the number of compute nodes and storage nodes in the graph sampling system, as well as the topology of the network fabric (including how the compute nodes and storage nodes will be connected to network switches) in the graph sampling system.

[0031] The graph hash table **208** and/or the subgraphs **206** may be periodically updated by the graph partitioning algorithm **204**. For example, when the graph **202** changes (e.g., due to the adding or removal of graph elements), the subgraphs **206** may be rebalanced and the graph hash table **208** may be updated based on the rebalancing of the subgraphs **206**.

[0032] FIG. 3 is a sequence diagram of a distributed graph sampling method **300**, according to some implementations. The distributed graph sampling method **300** may be performed in a distributed graph sampling system, such as the distributed graph sampling system **100** (see FIG. 1), when a compute node **102** requests a sample of a distributed graph that is stored across multiple storage nodes **106**.

[0033] In step **302**, the compute node **102** requests a sample of the graph from a network switch **108**. In the example where each compute node **102** is connected to a respective network switch **108**, the graph sample request may be sent to the network switch **108** which is connected to the compute node **102**. The graph sample request may indicate one or more graph node(s) that should be included in the sample of the graph mode. For example, the graph sample request may include one or more identifier(s) of graph node(s).

[0034] In step **304**, the network switch **108** requests samples of subgraphs from the storage nodes **106**. The samples may be requested for the subgraphs which contain the graph node(s) identified by the graph sample request. As previously noted, the subgraphs may be stored on the storage nodes **106**, while a graph hash table may be stored on the network switch **108**. The network switch **108** may use the graph hash table to identify the storage nodes **106** that contain the needed subgraphs. For example, the network switch **108** may use the graph hash table as a routing table to forward the graph sample request from the compute node **102** to the desired storage nodes **106**.

[0035] In step **306**, the storage nodes **106** (identified by the network switch **108**) generate samples of their stored subgraphs. Each storage node **106** may generate a sample of its subgraph using any suitable graph sampling technique. Examples of graph sampling techniques include random sampling, neighborhood sampling, and the like. In some implementations, a random sampling algorithm such as GraphSAGE may be used by a storage node **106** to randomly select nodes of its subgraph. The subgraph samples may be generated in parallel across the storage nodes **106**. Details regarding how a storage node **106** samples its subgraph will be subsequently described.

[0036] In step **308**, the storage nodes **106** return the samples of their respective subgraphs to the network switch **108**. The network switch **108** thus receives the subgraph samples from the storage nodes **106** it previously identified. The subgraph samples (collectively) contain the graph node(s) indicated by the graph sample request.

[0037] In step **310**, the network switch **108** aggregates the subgraph samples from the storage nodes **106** into a combined sample of the graph. For example, the graph elements (including the graph nodes and the graph edges) in each subgraph sample may be concatenated into a combined graph sample. When a graph is partitioned, there may be some edge cuts in the partitioned graph. As a result, a graph edge stored in a first storage node **106** may connect graph nodes stored in a second storage node **106**. When the subgraphs contain edge cuts, the combined graph sample may include duplicate graph edges. As previously noted, the network switch **108** is a Smart Switch with embedded computational capabilities that may perform the aggregation.

[0038] In step **312**, the network switch **108** deduplicates the graph edges in the combined graph sample. The duplicate graph edges may be removed from the combined graph sample, such that the combined graph sample includes only one instance of each graph edge. The network switch **108** may use the graph hash table stored at the network switch **108** to identify duplicate graph edges in the combined graph sample. For example, because the graph hash table identifies the graph elements stored in the subgraph of each storage node **106**, the network switch **108** may, based on the requested subgraph samples, expect certain graph edges to be duplicated in the subgraph samples. Those expected duplicate graph edges may be removed from the combined graph sample. Deduplicating graph edges in the combined graph sample may optimize the graph sample, such that fewer computing resources will be needed (by the compute node **102**) to process the graph sample with GNN algorithms. As previously noted, the network switch **108** is a Smart Switch with embedded computational capabilities that may perform the deduplication.

[0039] In some implementations, the network switch **108** deduplicates the graph nodes in the combined graph sample. Graph node deduplication may be performed in lieu of or in addition to graph edge deduplication.

[0040] In step **314**, the network switch **108** returns the combined graph sample to the compute node **102**. The graph sample contains the graph node(s) indicated by the graph sample request from the compute node **102**. Subsequently, the compute node **102** may process the graph sample with GNN algorithms.

[0041] FIG. **4** is a diagram of a network switch **108**, according to some implementations. The components of the network switch **108** can be implemented as integrated circuits, such as in one or more integrated circuit die(s) and/or one or more integrated circuit package(s).

[0042] To achieve its desired functionality, the network switch **108** includes various hardware components. These hardware components may include a processor **402**, a memory **404**, and a switch core **406**. The hardware components may be interconnected through a number of busses and/or network connections. In one example, the processor **402**, the memory **404**, and the switch core **406** may be communicatively coupled via a bus **408**, such as a PCI-Express bus.

[0043] The processor **402** retrieves executable code from the memory **404** and executes the executable code. The executable code may, when executed by the processor **402**, cause the processor **402** to implement any functionality described herein. The processor **402** may be a microprocessor, an application-specific integrated circuit, a microcontroller, or the like.

[0044] The memory **404** may include various types of memory, including volatile and nonvolatile memory. For example, the memory **404** may include Random-Access Memory (RAM), Read-Only Memory (ROM), a Solid State Drive (SSD), the like, or combinations there. Different types of memory may be used for different data storage needs. For example, the processor **402** may boot from ROM, maintain nonvolatile storage in an SSD, execute program code stored in RAM, and store data under processing in RAM. The memory **404** may include a non-transitory computer readable medium that stores instructions for execution by the processor **402**. One or more modules within the network switch **108** may be partially or wholly embodied as software and/or hardware for performing any functionality described herein.

[0045] The switch core **406** includes the components of the network switch **108** for switching and forwarding packets. Example components of the switch core **406** include ingress ports, egress ports, a switch fabric (e.g., crossbars, shared buses, shared memory, a chip-wide ring, etc.), and the like. The switch core **406** is connected to compute nodes and storage nodes.

[0046] FIG. 5 is a diagram of a network switch graph sampling method **500**, according to some implementations. The network switch graph sampling method **500** may be performed by a network switch of a distributed graph sampling system when generating a sample of a graph.

[0047] The network switch performs a step **502** of obtaining distributed samples of subgraphs of a graph from storage nodes, the storage nodes storing the subgraphs, the graph comprising graph nodes connected by graph edges. A graph hash table is stored in the memory of the network switch. The graph hash table may be a dictionary having a key-value structure. The network switch may also perform steps of receive a sampling request from the compute node; and forward the sampling request to the storage nodes using the graph hash table as a routing table. The sampling request may comprise identifiers of particular graph nodes of the graph. Obtaining the distributed samples of the subgraphs may comprise selecting ones of the storage nodes based on the identifiers of the sampling request.

[0048] The network switch performs a step **504** of aggregating the distributed samples of the subgraphs into a combined sample of the graph. Aggregating the distributed samples of the subgraphs into the combined sample of the graph may comprise concatenating the distributed samples of the subgraphs.

[0049] The network switch performs a step **506** of removing duplicate ones of the graph edges from the combined sample of the graph based on the graph hash table stored in the memory, the graph hash table identifying the storage nodes that store the graph edges and the graph nodes. Removing the duplicate ones of the graph edges may include: identifying expected ones of the graph edges that will be duplicated in the distributed samples of the subgraphs based on the graph hash table; and removing the expected ones of the graph edges from the combined sample of the graph.

[0050] The network switch performs a step **508** of sending the combined sample of the graph to a compute node. Subsequently, the compute node **102** may process the combined sample of the graph with GNN algorithms.

[0051] Some variations are contemplated. As previously noted, the network fabric **104** may have a layered architecture having multiple tiers of the network switches **108**. In such an implementation, the network switches **108** in a lower tier may aggregate and return combined subgraph samples to a network switch **108** in an upper tier. The network switch **108** in the upper tier may then aggregate the combined subgraph samples into the combined sample of the graph that is returned to a compute node **102**.

[0052] FIG. 6 is a diagram of a storage node **106**, according to some implementations. The storage node **106** may be used to provide efficient and reliable storage, access, and sampling capabilities for the distributed handling of a large-scale graph.

[0053] To achieve its desired functionality, the storage node **106** includes various hardware components. These hardware components may include a processor **602**, a memory **604**, a network

interface card (NIC) **606**, and one or more accelerators **608**. The hardware components may be interconnected through a number of busses and/or network connections. In one example, the processor **602**, the memory **604**, the NIC **606**, and the accelerators **608** may be communicatively coupled via a bus **610**, such as a PCI-Express bus.

[0054] The processor **602** retrieves executable code from the memory **604** and executes the executable code. The executable code may, when executed by the processor **602**, cause the processor **602** to implement any functionality described herein. The processor **602** may be a microprocessor, an application-specific integrated circuit, a microcontroller, or the like. The processor **602** is a CPU of the storage node **106**.

[0055] The memory **604** may include various types of memory, including volatile and nonvolatile memory. For example, the memory **604** may include Random-Access Memory (RAM), Read-Only Memory (ROM), a Solid State Drive (SSD), the like, or combinations there. Different types of memory may be used for different data storage needs. For example, the processor **602** may boot from ROM, maintain nonvolatile storage in an SSD, execute program code stored in RAM, and store data under processing in RAM. The memory **604** may include a non-transitory computer readable medium that stores instructions for execution by the processor **602**. One or more modules within the storage node **106** may be partially or wholly embodied as software and/or hardware for performing any functionality described herein.

[0056] The NIC **606** may be used to connect to a network (e.g., LAN, WAN, etc.) and communicate with other devices over that network. The NIC **606** facilitates the transmission and reception of data packets between the storage node **106** and the network, and may adhere to one or more networking standards such as Ethernet, Wi-Fi, and the like. Additionally, the NIC **606** may support Remote Direct Memory Access (RDMA) for communicating with other nodes of a distributed graph sampling system. The NIC **606** may receive subgraph sample requests from a network switch (not illustrated in FIG. **6**) and dispatch the subgraph sample requests to appropriate ones of the accelerators **608** for processing the requested subgraph samples. The NIC **606** may be a programmable NIC, such as a SmartNIC, having processing cores that can be programmed to perform the dispatching of subgraph sample requests.

[0057] The accelerators **608** are specialized processing units that can be programmed to perform operations for sampling a subgraph. Examples of the accelerators **608** include Graphics Processing Units (GPUs), Field Programmable Gate Arrays (FPGAs), Application-Specific Integrated Circuits (ASICs), and other specialized processing units, which may be incorporated into the storage node **106** to expedite graph-related computations. The accelerators **608** may include streaming multiprocessors. The accelerators **608** provide significant computational power, allowing for faster execution of tasks like graph sampling algorithms than a CPU (e.g., the processor **602**).

[0058] The memory **604** stores a subgraph of a graph. The subgraph may have been previously obtained by partitioning the graph, as previously described for FIG. **2**. The storage node **106** may include a Shared Virtual Memory (SVM) between the accelerators **608**, the NIC **606**, and the memory **604** (including RAM and/or SSDs). As a result, the memory **604** may include a shared memory region storing the subgraph, and the accelerators **608** and the NIC **606** may directly access the subgraph from the shared memory region. A system interconnect communication protocol that supports coherent memory access may be used to interconnect the memory **604** with the accelerators **608** and the NIC **606**. In some implementations, the memory **604** is interconnected with the accelerators **608** and the NIC **606** by a Compute Express Link (CXL) interconnect, a Cache Coherent Interconnect for Accelerators (CCIX), a Coherent Accelerator Processor Interface (CAPI), or the like. The use of an SVM allows the components of the storage node **106** to directly access and utilize shared data without redundant data movement (through the processor **602**), thereby increasing efficiency.

[0059] FIG. **7** is a sequence diagram of a storage node graph sampling method **700**, according to some implementations. The storage node graph sampling method **700** may be performed by a

storage node of a distributed graph sampling system when generating a sample of a subgraph stored in the memory of the storage node.

[0060] In step **702**, the NIC **606** selects an accelerator **608** to perform a subgraph sampling operation. The accelerator **608** may be selected in response to the NIC **606** receiving a subgraph sample request from a network switch (not illustrated in FIG. 7). The subgraph sample request may indicate graph node(s) that exist in the subgraph stored in the memory **604**, which graph node(s) should be included in the requested subgraph sample. The selected accelerator **608** may be one that is suitable for performing a desired sampling operation and is available (e.g., not performing another sampling operation).

[0061] In step **704**, the NIC **606** dispatches a subgraph sampling operation to the selected accelerator **608**. For example, the NIC **606** may launch a kernel (which includes the subgraph sampling operation) on the selected accelerator **608**. The kernel generates a desired sample of the subgraph.

[0062] In step **706**, the selected accelerator **608** performs the subgraph sampling operation on the subgraph stored in the memory **604** of the storage node, thus producing the requested sample of the subgraph. As an example of the subgraph sampling operation, the selected accelerator **608** may locate the graph node(s) indicated by the subgraph sample request, and perform a neighborhood sampling algorithm with those graph node(s). Due to the SVM, the selected accelerator **608** may directly access the subgraph from the shared memory region of the memory **604** via a system interconnect, without involving the CPU.

[0063] In step **708**, the selected accelerator **608** returns the subgraph sample to the NIC **606**. For example, the selected accelerator **608** may put the subgraph sample in the memory **604** and notify the NIC **606** that the subgraph sample is ready. The NIC **606** may then directly read the subgraph sample from the memory **604** and forward the subgraph sample to the network switch (not illustrated in FIG. 7) from which it received the subgraph sample request.

[0064] A storage node may handle multiple subgraph sample requests simultaneously. For example, a first sampling operation for a first subgraph sample request may be dispatched to a first accelerator **608** while a second sampling operation for a second subgraph sample request may be dispatched to a second accelerator **608**. The first and second accelerators **608** may perform their sampling operations in parallel, with both directly accessing needed graph node(s) of the subgraph from the memory **604**.

[0065] As previously noted, the NIC **606** may be a SmartNIC. The selecting and dispatching steps may be performed by the NIC **606**. Further, the selected accelerator **608** and the NIC **606** may obtain the needed graph data and subgraph sample from the memory **604**. Thus, processing of a subgraph sample may be coordinated by the NIC **606**, without involving the CPU of the storage node and without copying graph data into a memory of the selected accelerator **608**. Overhead for a CPU of the storage node may thus be reduced.

[0066] FIG. **8** is a diagram of a NIC graph sampling method **800**, according to some implementations. The NIC graph sampling method **800** may be performed by a NIC of a storage node when generating a sample of a graph.

[0067] The NIC performs a step **802** of receiving a sample request for a subgraph of a graph from a network switch, a memory of the storage node storing the subgraph in a shared memory region, the graph comprising graph nodes connected by graph edges. The sample request may identify one or more of the graph nodes of the graph.

[0068] The NIC performs a step **804** of selecting an accelerator of the storage node. The accelerator may be a streaming multiprocessor.

[0069] The NIC performs a step **806** of dispatching a sampling operation to the accelerator, the sampling operation controlling the accelerator to generate a sample of the subgraph, the accelerator directly reading the subgraph from the shared memory region, the accelerator directly writing the sample of the subgraph to the shared memory region. The sampling operation is dispatched directly

to the accelerator without copying the subgraph to an onboard memory of the accelerator. The sampling operation may comprise neighborhood sampling performed with the one of the graph nodes identified by the sample request.

[0070] The NIC performs a step **808** of returning the sample of the subgraph to the network switch. Returning the sample of the subgraph to the network switch may comprise directly reading the sample of the subgraph from the shared memory region.

[0071] Embodiments may achieve advantages. The distributed graph sampling system **100** provides enhanced scalability for processing large-scale graphs through distributed storage and computation. Additionally, sampling time may be reduced due to parallel operation across the storage nodes **106**. Further, communication overhead may be reduced due to the intelligent graph sample aggregation and deduplication performed by the network switches **108**.

[0072] The foregoing outlines features of several examples so that those skilled in the art may better understand the aspects of the present disclosure. Various modifications and combinations of the illustrative examples, as well as other examples, will be apparent to persons skilled in the art upon reference to the description. It is therefore intended that the appended claims encompass any such modifications.

Claims

1. A system comprising: a plurality of storage nodes configured to: store a graph comprising graph nodes connected by graph edges; and generate subgraph samples of the graph; a network switch configured to: receive the subgraph samples from the storage nodes; aggregate the subgraph samples into a graph sample of the graph; and remove duplicate ones of the graph edges from the graph sample based on a graph hash table stored at the network switch, the graph hash table identifying the storage nodes that store the graph edges and the graph nodes; and a compute node configured to: receive the graph sample from the network switch; and process the graph sample with a graph neural network algorithm.
2. The system of claim 1, wherein each of the storage nodes comprises: a memory storing a subgraph of the graph in a shared memory region; a plurality of accelerators; and a network interface card configured to: receive a subgraph sample request from the network switch; and dispatch a subgraph sampling operation to a selected accelerator of the accelerators, the selected accelerator configured to directly access the subgraph of the graph from the shared memory region.
3. The system of claim 2, wherein the memory is interconnected with the accelerators and the network interface card by a compute express link (CXL) interconnect.
4. The system of claim 2, wherein the network interface card is a SmartNIC.
5. The system of claim 1, wherein the network switch is a Smart Switch.
6. The system of claim 1, wherein the network switch is further configured to forward a subgraph sample request from the compute node to the storage nodes using the graph hash table as a routing table.
7. The system of claim 1, wherein the network switch removes the duplicate ones of the graph edges from the graph sample by: identifying expected ones of the graph edges that will be duplicated in the subgraph samples based on the graph hash table; and removing the expected ones of the graph edges from the graph sample.
8. The system of claim 1, wherein the graph hash table is a dictionary having a key-value structure.
9. A network switch comprising: a switch core; a processor; and a memory comprising a non-transitory computer readable medium storing instructions which, when executed by the processor, cause the processor to: obtain distributed samples of subgraphs of a graph from storage nodes, the storage nodes storing the subgraphs, the graph comprising graph nodes connected by graph edges; aggregate the distributed samples of the subgraphs into a combined sample of the graph; remove duplicate ones of the graph edges from the combined sample of the graph based on a graph hash

table stored in the memory, the graph hash table identifying the storage nodes that store the graph edges and the graph nodes; and send the combined sample of the graph to a compute node.

10. The network switch of claim 9, wherein the non-transitory computer readable medium stores further instructions which, when executed by the processor, cause the processor to: receive a sampling request from the compute node; and forward the sampling request to the storage nodes using the graph hash table as a routing table.

11. The network switch of claim 10, wherein the sampling request comprises identifiers of particular graph nodes of the graph, and obtaining the distributed samples of the subgraphs comprises selecting ones of the storage nodes based on the identifiers of the sampling request.

12. The network switch of claim 9, wherein aggregating the distributed samples of the subgraphs into the combined sample of the graph comprises concatenating the distributed samples of the subgraphs.

13. The network switch of claim 9, wherein removing the duplicate ones of the graph edges comprises: identifying expected ones of the graph edges that will be duplicated in the distributed samples of the subgraphs based on the graph hash table; and removing the expected ones of the graph edges from the combined sample of the graph.

14. The network switch of claim 9, wherein the graph hash table is a dictionary having a key-value structure.

15. A method, by a network interface card of a storage node, the method comprising: receiving a sample request for a subgraph of a graph from a network switch, a memory of the storage node storing the subgraph in a shared memory region, the graph comprising graph nodes connected by graph edges; selecting an accelerator of the storage node; dispatching a sampling operation to the accelerator, the sampling operation controlling the accelerator to generate a sample of the subgraph, the accelerator directly reading the subgraph from the shared memory region, the accelerator directly writing the sample of the subgraph to the shared memory region; and returning the sample of the subgraph to the network switch.

16. The method of claim 15, wherein returning the sample of the subgraph to the network switch comprises directly reading the sample of the subgraph from the shared memory region.

17. The method of claim 15, wherein the accelerator is a streaming multiprocessor.

18. The method of claim 15, wherein the sampling operation is dispatched to the accelerator without copying the subgraph to an onboard memory of the accelerator.

19. The method of claim 15, wherein the sample request identifies one of the graph nodes of the graph.

20. The method of claim 19, wherein the sampling operation comprises neighborhood sampling performed with the one of the graph nodes identified by the sample request.
