

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250285159

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Balasubramanian; Ramasubramanian et al.

NON-TRANSITORY COMPUTER-READABLE MEDIUM, AND METHOD FOR CONTEXT-BASED CONTENT-SCORING FOR AN ONLINE CONCIERGE SYSTEM

Abstract

An online concierge system selects content for presentation to a user by using a product scoring engine. The product scoring engine generates a user embedding for user data and a query embedding for query data. The product scoring engine generates an anchor embedding based on the user embedding and the query embedding, where the anchor embedding is an embedding in a product embedding space. The product scoring engine compares the anchor embedding to a set of product embeddings to score a set of products for presentation to a user.

Inventors: Balasubramanian; Ramasubramanian (Jersey City, NJ), Manchanda; Saurav (Seattle, WA)

Applicant: Maplebear Inc. (San Francisco, CA)

Family ID: 1000008619420

Appl. No.: 19/214386

Filed: May 21, 2025

Related U.S. Application Data

parent US continuation 17550950 20211214 parent-grant-document US 12340405 child US 19214386

Publication Classification

Int. Cl.: G06Q30/0601 (20230101); G06F16/2455 (20190101); G06N20/00 (20190101)

U.S. Cl.:

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of co-pending U.S. patent application Ser. No. 17/550,950, filed Dec. 14, 2021, which is incorporated by reference herein in its entirety.

BACKGROUND

[0002] Machine-learning models can be trained to present content to a user of an online system. For example, a machine-learning model may be trained to select products to present to a user in response to a search query. Conventional machine-learning models determine content to present to a user based on how other users have interacted with content in the past. These conventional machine-learning models typically identify features of users similar to a target user and select content with which the similar users have interacted. However, these machine-learning models often fail to compare the contexts of the target user and the similar user when selecting content. Thus, conventional machine-learning models that simply look at content with which a similar user interacted may provide irrelevant content in the target user's context. For example, if the similar user interacted with a ground beef product in response to a search query for “ground beef,” the target user may not interact with the ground beef product if the product was presented in response to a search query for “milk,” no matter how similar the users are. Thus, conventional machine-learning models often fail to properly account for context in scoring and selecting content to present to a user.

SUMMARY

[0003] An online concierge system uses a product scoring engine to score and select products to present to a user. The product scoring engine may receive user data describing a user that has submitted a search query to the online concierge system and may receive query data describing the search query submitted by the user. The product scoring engine uses a user embedding model to generate a user embedding describing the user based on the user data. A user embedding model is a machine-learning model that is trained to generate a user embedding based on user data. Similarly, the product scoring engine uses a query embedding model to generate a query embedding describing the search query submitted by the user based on the query data. A query embedding model is a machine-learning model that is trained to generate a query embedding describing a search query based on query data.

[0004] The product scoring engine generates an anchor embedding by applying an anchor embedding model to the user embedding and the query embedding. An anchor embedding is an embedding in the product embedding space that is generated based on the user embedding and the query embedding. The anchor embedding can be compared to product embeddings describing products to determine products that would be suitable to present to the user associated with the anchor embedding in response to the search query. An anchor embedding model is a machine-learning model that has been trained to generate an anchor embedding based on a user embedding and a query embedding.

[0005] The product scoring engine generates product embeddings for a set of candidate products by applying a product embedding model to product data describing the set of candidate products. A product embedding model is a machine-learning model that is trained to generate product embeddings describing products based on product data. The product scoring engine then compares the product embeddings for the set of candidate products with an anchor embedding for a user and the user's search query to generate a set of product scores for the set of candidate products. The

online concierge system then selects one or more products of the set of candidate products to present to a user in response to the search query based on the set of product scores generated for the set of candidate products.

[0006] The online concierge system may train the user embedding model, query embedding model, product embedding model, or the anchor embedding based on a loss function. The loss function may be based on a triplet loss function, such that the machine-learning models are trained to minimize a distance of anchor embeddings from product embeddings associated with positive examples (e.g., products presented to a user with which the user has interacted) and to maximize a distance from product embeddings associated with negative examples (e.g., products presented to a user with which the user did not interact). In some embodiments, the loss function assigns different weights to positive examples based on a hierarchy of interactions with products. The hierarchy of interactions may reflect that certain interactions with products are more valuable to the online concierge system than other interactions. For example, a user purchasing a product may be a more valuable interaction than when a user simply selects a product to view more information about the product.

[0007] By comparing product embeddings to an anchor embedding that is based on a user embedding and a query embedding, the online concierge system can score products based on characteristics about the user and based on the context within which the user is viewing products. Additionally, by applying a loss function that uses a hierarchy of interactions, the online concierge system may ensure that products that lead to more valuable interactions are scored more highly for presentation than products that lead to less valuable interactions.

Description

BRIEF DESCRIPTION OF DRAWINGS

[0008] FIG. 1 illustrates the environment of an online concierge system, in accordance with some embodiments.

[0009] FIG. 2 is a block diagram of an online concierge system, in accordance with some embodiments.

[0010] FIG. 3A is a block diagram of the user mobile application (UMA), in accordance with some embodiments.

[0011] FIG. 3B is a block diagram of the picker mobile application (PMA), in accordance with some embodiments.

[0012] FIG. 3C is a block diagram of the runner mobile application (RMA), in accordance with some embodiments.

[0013] FIG. 4 is a block diagram for a product scoring engine, in accordance with some embodiments.

[0014] FIG. 5 illustrates data flow through an example product scoring engine during a training process, in accordance with some embodiments.

[0015] FIG. 6 illustrates data flow through an example product scoring engine during application of the example product scoring engine, in accordance with some embodiments.

DETAILED DESCRIPTION

Environment of an Online Concierge System

[0016] FIG. 1 illustrates the environment **100** of an online concierge system **102**, according to some embodiments. The figures use like reference numerals to identify like elements. A letter after a reference numeral, such as “**110a**,” indicates that the text refers specifically to the element having that particular reference numeral. A reference numeral in the text without a following letter, such as “**110**,” refers to any or all of the elements in the figures bearing that reference numeral. For example, “**110**” in the text refers to reference numerals “**110a**” and/or “**110b**” in the figures.

[0017] The environment **100** includes an online concierge system **102**. The online concierge system **102** is configured to receive orders from one or more users **104** (only one is shown for the sake of simplicity). An order specifies a list of goods or products to be delivered to the user **104**. Goods, items, and products may be used synonymously herein to mean any item that a user **104** can purchase via the online concierge system. The order also specifies the location to which the goods are to be delivered, and a time window during which the goods should be delivered. In some embodiments, the order specifies one or more retailers from which the selected items should be purchased. The user may use a user mobile application (UMA) **106** to place the order; the UMA **106** is configured to communicate with the online concierge system **102**.

[0018] The online concierge system **102** is configured to transmit orders received from users **104** to one or more pickers **108**. A picker **108** may be a contractor, employee, or other person (or entity) who is enabled to fulfill orders received by the online concierge system **102**. The environment **100** also includes three retailers **110a**, **110b**, and **110c** (only three are shown for the sake of simplicity; the environment could include any number of retailers). The retailers **110** may be physical retailers, such as grocery stores, discount stores, department stores, etc., or non-public warehouses storing items that can be collected and delivered to users. Each picker **108** fulfills an order received from the online concierge system **102** at one or more retailers **110**, delivers the order to the user **104**, or performs both fulfillment and delivery. In some embodiments, pickers **108** make use of a picker mobile application **112** which is configured to interact with the online concierge system **102**.

Online Concierge System

[0019] FIG. 2 is a block diagram of an online concierge system **102**, according to some embodiments. The online concierge system **102** includes an inventory management engine **202**, which interacts with inventory systems associated with each retailer **110**. In some embodiments, the inventory management engine **202** requests and receives inventory information maintained by the retailer **110**. The inventory of each retailer **110** is unique and may change over time. The inventory management engine **202** monitors changes in inventory for each participating retailer **110**. The inventory management engine **202** is also configured to store inventory records in an inventory database **204**. The inventory database **204** may store information in separate records—one for each participating retailer **110**—or may consolidate or combine inventory information into a unified record. Inventory information includes both qualitative and quantitative information about items, including size, color, weight, SKU, serial number, and so on. In some embodiments, the inventory database **204** also stores purchasing rules associated with each item, if they exist. For example, age-restricted items such as alcohol and tobacco are flagged accordingly in the inventory database **204**.

[0020] The online concierge system **102** also includes an order fulfillment engine **206** which is configured to synthesize and display an ordering interface to each user **104** (for example, via the user mobile application **106**). The order fulfillment engine **206** is also configured to access the inventory database **204** in order to determine which products are available at which retailers **110**. The order fulfillment engine **206** determines a sale price for each item ordered by a user **104**. Prices set by the order fulfillment engine **206** may or may not be identical to in-store prices determined by retailers (which is the price that users **104** and pickers **108** would pay at retailers). The order fulfillment engine **206** also facilitates transactions associated with each order. In some embodiments, the order fulfillment engine **206** charges a payment instrument associated with a user **104** when he/she places an order. The order fulfillment engine **206** may transmit payment information to an external payment gateway or payment processor. The order fulfillment engine **206** stores payment and transactional information associated with each order in a transaction records database **208**.

[0021] In some embodiments, the order fulfillment engine **206** also shares order details with retailer **110**. For example, after successful fulfillment of an order, the order fulfillment engine **206** may transmit a summary of the order to the appropriate retailer. The summary may indicate the items purchased, the total value of the items, and in some cases, an identity of the picker **108** and

user **104** associated with the transaction. In some embodiments, the order fulfillment engine **206** pushes transaction and/or order details asynchronously to retailer systems. This may be accomplished via use of webhooks, which enable programmatic or system-driven transmission of information between web applications. In another embodiment, retailer systems may be configured to periodically poll the order fulfillment engine **206**, which provides detail of all orders which have been processed since the last request.

[0022] The order fulfillment engine **206** may interact with a picker management engine **210**, which manages communication with and utilization of pickers **108**. In some embodiments, the picker management engine **210** receives a new order from the order fulfillment engine **206**. The picker management engine **210** identifies the appropriate retailer **110** to fulfill the order based on one or more parameters, such as the contents of the order, the inventory of the retailers, and the proximity to the delivery location. The picker management engine **210** then identifies one or more appropriate pickers **108** to fulfill the order based on one or more parameters, such as the pickers' proximity to the appropriate retailer **110** (and/or to the user **104**), his/her familiarity level with that particular retailer **110**, and so on. Additionally, the picker management engine **210** accesses a picker database **212** which stores information describing each picker **108**, such as his/her name, rating, previous shopping history, and so on. The picker management engine **210** transmits the list of items in the order to the picker **108** via the picker mobile application **112**. The picker database **212** may also store data describing the sequence in which the pickers picked the items in their assigned orders.

[0023] As part of fulfilling an order, the order fulfillment engine **206** and/or picker management engine **210** may access a user database **214** which stores information describing each user. This information could include each user's name, address, gender, shopping preferences, favorite items, stored payment instruments, and so on.

[0024] The online concierge system **102** may use a communication engine **216** that transmits information between the user mobile application **106**, the picker mobile application **112**, and the runner mobile application **116**. The information may be sent in the form of messages, such as texts or emails, or notifications via application, among other forms of communication. The communication engine **216** may receive information from each application about the status of an order, the location of a user in transit, issues with items in an order, and the like. The communication engine **216** determines a message or notification to send to a user **104**, picker **108**, or runner **114** based on this information and transmits the notifications to the appropriate application. For example, the online concierge system **102** may receive information from the user mobile application **106** indicating that a user **104** is traveling to the pickup location to retrieve an order. Based on this information, the communication engine **216** sends a notification to the runner mobile application **116** indicating that the user **104** associated with a specific order is in transit, which may incite the runner **114** to retrieve the order for pick up. In another example, the online concierge system **102** may receive a message from the picker mobile application **112** that an item of an order is not available. The communication engine **216** may transmit the message to the user mobile application **106** associated with the order.

[0025] The product selection module **218** selects products to be presented to the user. The product selection module **218** selects products to present to the user based on a relevance or affinity of the products to the user. The relevance or affinity of a product for a user may be a score for the product that is based on a likelihood that the user will interact with the product if presented to the user. The product selection module **218** may then select products based on a score for each product. For example, the product selection module **218** may rank the products based on the score and select a set of products to present to a user based on the ranking of the products. When the online concierge system **102** identifies an opportunity to present products to a user, the product selection module **218** may score all products that are available for purchase on the online concierge system **102** or may only score a set of candidate products that are likely relevant to the user.

[0026] The product selection module **218** may determine the relevance of a product to a user based

on one or more machine-learning models (e.g., a neural network) that have been trained to determine the relevance of products to users. The product selection module **218** may determine the relevance of a product to a user based on information about the user, the type of product, whether the user has ordered the product before (and if so, how recently), whether the product is related to other products the user has ordered or the relevance of the product to a search query provided by the user. The product selection module **218** may present the selected products to the user as search results, a feed of potentially relevant products, or as part of an ordered list of all products available to the user.

[0027] In some embodiments, the product selection module **218** selects products to present to the user in response to a search query by the user. A search query is a string of text that represents a user's interest in being presented with a set of products related to that text. For example, if the user provides a search query of "ground beef," the user is likely interested in ground beef products available for purchase via the online concierge system. Accordingly, the product selection module **218** may present ground beef products to the user in response to the user's search query.

[0028] The product selection module **218** uses a product scoring engine **220** to score products to be presented as search results to a user. The product scoring engine **220** uses one or more machine-learning models that are trained to score products. The product scoring engine **220** is discussed in further detail below in the context of FIG. 4. Note that, while the description herein may primarily describe the product scoring engine **220** as scoring products to present to a user in response to a search query, the product scoring module **220** may be used to score products to present to a user in contexts other than search results. For example, the product scoring module **220** may generically score products to present to a user based on user data about the user, product data about candidate products to present, and context data describing the user's session with the online concierge system **102**.

[0029] FIG. 3A is a block diagram of the user mobile application (UMA) **106**, according to some embodiments. The user **104** accesses the UMA **106** via a client device, such as a mobile phone, tablet, laptop, or desktop computer. The UMA **106** may be accessed through an app running on the client device or through a website accessed in a browser. The UMA **106** includes an ordering interface **302**, which provides an interactive interface with which the user **104** can browse through and select products and place an order. The ordering interface **302** also may include a selected products list that specifies the amounts and prices of products that the user **104** has selected to order. The user **104** may review the selected products list and place an order based on the selected products list. Furthermore, the ordering interface **302** may present recipes to the user **104** that the online concierge system **102** predicts the user **104** is attempting to complete, and may provide an option to the user **104** to add additional products needed to complete a recipe to the user's selected products list.

[0030] The ordering interface **302** may allow a user to search for products offered for sale on the online concierge system **102**. For example, the ordering interface **302** may include elements that the user may use to enter a search query. Once the user enters a search query, the ordering interface **302** may present to the user a set of search results that the online concierge system **102** determines are relevant to the user's search. The ordering interface **302** may also include elements that allow the user to order products that are presented as part of the set of search results.

[0031] Users **104** may also use the ordering interface **302** to message with pickers **108** and runners **114** and receive notifications regarding the status of their orders. The UMA **106** also includes a system communication interface **304** which, among other functions, receives inventory information from the online concierge system **102** and transmits order and location information to the online concierge system **102**. The UMA **106** also includes a preferences management interface **306** which allows the user **104** to manage basic information associated with his/her account, such as his/her home address and payment instruments. The preferences management interface **306** may also allow the user to manage other details such as his/her favorite or preferred retailers **110**, preferred handoff

times, special instructions for handoff, and so on.

[0032] The UMA **106** also includes a location data module **308**. The location data module **308** may access and store location data related to a client device associated with a user **104** via the user mobile application **106**. Location data may include the geographic location of the client device associated with the user mobile application **106**, how fast the client device is traveling, the average speed of the client device when in transit, the direction of travel of the client device, the route the user **104** is taking to a pickup location, current traffic data near the pickup location, and the like. For simplicity, the location of a user client device or client device may be referred to as the location of the user throughout this description. The user **104** may specify whether or not to share this location data with the user mobile application **106** via the preferences management interface **306**. If a user **104** does not allow the user mobile application **106** to access their location data, the location data module **308** may not access any location data for the user **104**. In some embodiments, the user **104** may specify certain scenarios when the location data module **308** may receive location data, such as when the user **104** is using the user mobile application **106**, any time, or when the user **104** turns on location tracking in the user mobile application **106** via an icon. The user **104** may also specify which location data the location data module **308** may retrieve, and which location data is off-limits. In some embodiments, the location data module **308** may be tracking the user's **104** location as a background process while the UMA **106** is in use. In other embodiments, the UMA **106** may use real-time location data from the location data module **308** to display a map to the user **104** indicating their current location and the route to a pickup location for their order.

[0033] FIG. 3B is a block diagram of the picker mobile application (PMA) **112**, according to some embodiments. The picker **108** accesses the PMA **112** via a mobile client device, such as a mobile phone or tablet. The PMA **112** may be accessed through an app running on the mobile client device or through a website accessed in a browser. The PMA **112** includes a barcode scanning module **320** which allows a picker **108** to scan an item at a retailer **110** (such as a can of soup on the shelf at a grocery store). The barcode scanning module **320** may also include an interface which allows the picker **108** to manually enter information describing an item (such as its serial number, SKU, quantity and/or weight) if a barcode is not available to be scanned. The PMA **112** also includes a basket manager **322** which maintains a running record of items collected by the picker **108** for purchase at a retailer **110**. This running record of items is commonly known as a “basket”. In some embodiments, the barcode scanning module **320** transmits information describing each item (such as its cost, quantity, weight, etc.) to the basket manager **322**, which updates its basket accordingly. The PMA **112** also includes an image encoder **326** which encodes the contents of a basket into an image. For example, the image encoder **326** may encode a basket of goods (with an identification of each item) into a QR code which can then be scanned by an employee of the retailer **110** at check-out.

[0034] The PMA **112** also includes a system communication interface **324**, which interacts with the online concierge system **102**. For example, the system communication interface **324** receives information from the online concierge system **102** about the items of an order, such as when a user updates an order to include more or less items. The system communication interface may receive notifications and messages from the online concierge system **102** indicating information about an order. The system communication interface transmits notifications and messages to be displayed via a user interface of the mobile device associated with the PMA **112**.

[0035] FIG. 3C is a block diagram of the runner mobile application (RMA) **116**, according to some embodiments. The runner **114** accesses the RMA **116** via a client device, such as a mobile phone, tablet, laptop, or desktop computer. The RMA **116** may be accessed through an app running on the client device or through a website accessed in a browser. The RMA **116** includes bag interface engine **328**, which provides an interactive interface with which the runner **114** can view orders they need to deliver and the locations of the bags for each order, such as on a particular shelf or in a refrigerator of a pickup location. The runner **114** may receive notifications through the bag

interface engine **328** about new orders, the location of a user **104** who is in transit to a pickup location, and new orders to deliver. The runner **114** may also receive communications via the bag interface engine **328** with users regarding order handoff and pickup confirmation and may interact with the interface generated by the bag interface engine **328** to send communications to users and the online concierge system **102** regarding order status. For example, a runner **114** may send the user a pickup spot at the pickup location to meet for order handoff and indicate that an order has been delivered to a user via the interface, which ends the wait time calculation by the location data module **308** associated with the user.

[0036] The RMA **116** includes a bag manager **330** that manages the assignment of orders to runners **114** and the locations of bags for each order. The RMA **116** also includes a system communication interface **332** which, among other functions, receives inventory information from the online concierge system **102** and transmits order and bag information to the online concierge system **102**. The system communication interface may also receive notifications and messages from the online concierge system **102** indicating information about an order. The system communication interface transmits notifications and messages to be displayed via a user interface of the mobile device associated with the RMA **116**.

Example Product Scoring Engine

[0037] FIG. **4** is a block diagram for a product scoring engine **220**, according to some embodiments. The product scoring engine **220** comprises a user embedding module **400**, a query embedding module **410**, a product embedding module **420**, an anchor embedding module **430**, a product scoring module **440**, and a training module **450**. Alternative embodiments may include more, fewer, or different components from those illustrated in FIG. **4**, and the functionality of each component may be divided between the components differently from the description below. Additionally, each component may perform their respective functionalities in response to a request from a human, or automatically without human intervention.

[0038] The user embedding module **400** generates user embeddings for users using the online concierge system **102**. Each user embedding is an embedding that describes characteristics or features about an associated user. The user embedding module **400** generates user embeddings for a user based on user data. User data is data that describes characteristics about a user that may be relevant for determining the relevance of a product to a user. For example, user data may include one or more of the user's name, the user's location, the user's stated preferences, the user's previously ordered products, the user's frequency of placing orders, which retailers the user orders from, a typical order cost for the user, or a browsing history of the user on the UMA **106** or other applications the user may use. User data may include raw data, preprocessed data, or feature sets describing information about a user. In some embodiments, the data collection module **400** collects user data from the user database **214**. The user embedding module **400** may store generated user embeddings in the user database **214** and associate, within the user database **214**, each user embedding with a user that the user embedding describes.

[0039] In some embodiments, the user embedding module **400** uses one or more user models to generate user embeddings. User models are machine-learning models (e.g., neural networks) that are trained to generate user embeddings based on user data. The user embedding module **400** can also retrieve a user embedding associated with a specified user from the user database **214**.

[0040] The query embedding module **410** generates a query embedding for the user's search query. A query embedding is an embedding that describes features of the user's search query. The query embedding module generates a query embedding based on search query data. Query data is data describing a user's search query for the online concierge system **102**. For example, query data may include one or more of search query text, previous searches by the user within the user's session, or search queries conducted by other users of the online concierge system **102**. Query data may also include context data describing the context in which the user has queried the online concierge system **102** for products. For example, the context data may include one or more of how long the

user's session with the online concierge system **102** has lasted, the products that are currently in the user's selected products list, or other products with which the user has interacted during the session. Query data may include raw data, preprocessed data, or feature sets describing information about a query or context.

[0041] The query embedding module **410** uses one or more query embedding models to generate a query embedding. Query embedding models are machine-learning models (e.g., neural networks) that are trained to generate query embeddings based on search query data.

[0042] The product embedding module **420** generates a product embedding for products being evaluated by the product scoring engine **220**. A product embedding is an embedding that describes a product. The product embeddings may be associated with specific products stored by the inventory database **204**. For example, each brand of a product may have an individual product embedding, or products may have different product embeddings for each retailer that sells the product. Alternatively, each product embedding may be associated with a generic product, and each generic product is associated with specific products that are similar or substitutes of each other. For example, the inventory database **102** may store a product embedding for the generic product “milk”, and the specific products of “Moo Moo 2% Milk” and “Greener Pastures Organic Whole Milk” may both be associated with the product embedding for “milk.” In some embodiments, product embeddings are stored in the inventory database **204**.

[0043] The product embedding module **420** uses one or more product embedding models to generate a product embedding based on product data. A product embedding model is one or more machine-learning models (e.g., neural networks) that are trained to generate product embeddings based on product data. Product data is data that describes characteristics about products available for purchase using the online concierge system **102**. For example, product data may include one or more of a product name, a product type, whether a product is associated with a recipe, retailers that offer the product for sale, the shelf-life of a product, identifiers for other products with which the product is commonly purchased, a popularity of the product, the availability of a product, the price of a product, any restrictions that may be in place on the purchase of the product, whether the product is a food item, a frequency with which the product is purchased using the online concierge system **102**, other products which the product has been or may be presented, or an expense incurred by the online concierge system **102** to provide the product to the user. Product data may include raw data, preprocessed data, or feature sets describing information about a product.

[0044] The anchor embedding module **430** generates anchor embeddings based on user embeddings and query embeddings. An anchor embedding is an embedding of the same dimension and in the same embedding space as a product embedding. The anchor embedding can therefore be compared to product embeddings to determine products that would be relevant to present to a user in response to a search query. The anchor embedding module **430** may use one or more anchor embedding models to generate an anchor embedding. Anchor embedding models are machine-learning models (e.g., neural networks) that are trained to generate anchor embeddings based on user embeddings and query embeddings.

[0045] The product scoring module **440** generates product scores for products. A product score is a score for a product that indicates the product's affinity for being presented to a user in response to a search query from the user. A product score may represent a likelihood that the user will interact with the product if the product is presented to the user or may represent some expected value based on the likelihood of user interaction and the value of the user's interaction with the product. The product scoring module **440** generates product scores for products based on a comparison of an anchor embedding with a set of product embeddings associated with a set of candidate products. The set of candidate products can include all products available on the online concierge system **102** or a subset of the products. The anchor embedding is generated by the anchor embedding module **430** based on a query embedding for the search query and a user embedding for the user who submitted the search query. The product scoring module **440** may compare the anchor embedding

with the set of product embeddings by calculating a Euclidean distance, a cosine distance, or a dot product of the anchor embedding and each product embedding.

[0046] Additionally, the product scoring module **440** may use a machine-learning model (e.g., a neural network) trained to generate product scores for products based on product embeddings associated with the products and an anchor embedding. In some embodiments, the machine-learning model generates product scores based on comparisons of product embeddings for the set of candidate products with an anchor embedding.

[0047] The training module **450** evaluates the performance of the machine-learning models used by the product scoring engine **220** so that the machine-learning models can be trained to perform more effectively. For example, the training module **412** may evaluate the performance of user embedding models, query embedding models, product embedding models, or anchor embedding models. In some embodiments, the training module **450** uses a loss function to train machine-learning models used by the product scoring engine **220**. The loss function may optimize the machine-learning models used by the product scoring engine **220** such that the product scoring engine **220** generates product scores for products representing an affinity of the products for being presented to a user in response to a search query from the user. In some embodiments, the training module **450** backpropagates the loss function through the anchor embedding models, the product embedding models, the query embedding models, and the user embedding models. The training module **450** also may apply a machine-learning model to the anchor embedding and a set of product embeddings to train the machine-learning models used by the product scoring engine **220**.

[0048] In some embodiments, the training module **450** uses a loss function that is based on the triplet loss function. The modified triplet loss function may assign different weights to different types of interactions that are considered positive examples based on a hierarchy of interactions. The hierarchy of interactions may reflect that certain interactions are more valuable to the online concierge system **102** than other interactions. For example, a purchase interaction (i.e., where a user purchases a product) may be more valuable than a select interaction (i.e., where the user adds the product to a selected products list), which itself may be more valuable than a click interaction (i.e., where the user selects a product to see more information about the product). The training module **450** may use weights in a modified triplet loss function that emphasizes minimizing the distance from higher tier positive examples than from lower tier positive examples. For example, the training module may use the following as a loss function:

[00001]

$$L = \sum_{i=1}^M \max(\beta_1 \text{dist}(\text{AE}, \text{PE}_{P_i}) + \dots + \beta_n \text{dist}(\text{AE}, \text{PE}_{P_i}) - \text{dist}(\text{AE}, \text{PE}_{N_i}) + \alpha, 0)$$

[0049] For this example loss function, AE is an anchor embedding for a user who submitted a search query, PE.sub.P is a product embedding of a positive training example (e.g., where the user interacted with a presented product), and PE.sub.N is a product embedding of a negative training example (e.g., where the user did not interact with a presented product). The function dist is a function that computes the distance between two embeddings (e.g., the Euclidean distance). In the example loss function, there are n types of positive interactions that a user can take with regards to a product and those different types of interactions are weighted using weights $\beta_{\text{sub.1}} \dots \beta_{\text{sub.n}}$. These weights can assign greater value to more valuable interactions by increasing the loss value of a distance between the anchor embedding and the positive example. Additionally, the example loss function uses indicator variables $\delta_{\text{sub.1}} \dots \delta_{\text{sub.n}}$ to apply the correct weight β to the positive example. The only indicator variable with a non-zero value is the indicator variable that corresponds to the type of interaction represented by the positive example. Furthermore, M is the number of search examples used to train the machine-learning models and α is a hyperparameter controlling the margin in the example loss function.

Training an Example Product Scoring Engine

[0050] FIG. 5 illustrates data flow through an example product scoring engine during a training

process, in accordance with some embodiments. Alternative embodiments may include more, fewer, or different components from those illustrated in FIG. 5, and the functionality of each component may be divided between the components differently from the description below. Additionally, each component may perform their respective functionalities in response to a request from a human, or automatically without human intervention.

[0051] The product scoring engine receives a search example **500** related to an instance where a user submitted a search query to the online concierge system and was presented with a set of products in response to the search query. The search example **500** includes user data **505** associated with the user who submitted the search query and query data **510** associated with the search query. A user embedding model **515** generates a user embedding **520** based on the user data **505** and the query embedding model **525** generates a query embedding **530** based on the query data **510**. In some embodiments, the user embedding model **515** is a machine learning model used by a user embedding module **400** of the product scoring engine. Similarly, in some embodiments, the query embedding model **525** is a machine-learning model used by a query embedding module **410** of the product scoring engine. The anchor embedding model **535** generates an anchor embedding **540** based on the user embedding **520** and the query embedding **530**. In some embodiments, the anchor embedding model **535** is a machine-learning model used by an anchor embedding module **430** of the product scoring engine.

[0052] The search example **500** includes a set of positive training examples **545** and a set of negative training examples **550**. The set of positive training examples **545** are a set of products with which the user interacted when those products were presented to the user in response to the search query. For example, the set of positive training examples **545** may include products that the user purchased, selected to add to a selected products list, or selected to view more details about the product. The set of negative training examples **550** are a set of products with which the user did not interact when those products were presented to the user in response to the search query.

[0053] The product embedding model **555** generates product embeddings **560** based on the set of positive training examples **545** and the set of negative training examples **550**. In some embodiments, the product embedding model **555** is a machine-learning model used by a product embedding module **420** of the product scoring engine. The training module **565** then evaluates the performance of the user embedding model **515**, the query embedding model **525**, the anchor embedding model **535**, and the product embedding model **555** by applying a loss function to the anchor embedding **540** and the set of product embeddings **560**. In some embodiments, the loss function is a modified version of the triplet loss function, and the anchor embedding **540** is compared to pairs of product embeddings, where the pairs of product embeddings have a product embedding corresponding to each of a positive training example **545** and a negative training example **550**.

Applying an Example Product Scoring Engine

[0054] FIG. 6 illustrates data flow through an example product scoring engine during application of the example product scoring engine, in accordance with some embodiments. Alternative embodiments may include more, fewer, or different components from those illustrated in FIG. 6, and the functionality of each component may be divided between the components differently from the description below. Additionally, each component may perform their respective functionalities in response to a request from a human, or automatically without human intervention.

[0055] The product scoring engine receives user data **600** and query data **605** related to a user's search query. The user data **600** describes a user who submitted the search query and the query data **605** describes the search query submitted by the user. The user embedding model **610** generates a user embedding **615** based on the user data **600** and the query embedding model **620** generates a query embedding **625** based on the query data **605**. The anchor embedding model **630** generates an anchor embedding **635** based on the user embedding **615** and the query embedding **625**.

[0056] The product scoring engine selects a set of candidate products for consideration and

receives product data **640** associated with the set of candidate products. The product embedding model **645** generates a set of candidate product embeddings **650** based on the product data **640** for the set of candidate products.

[0057] The product scoring module **655** generates a product score **660** for each of the set of candidate products based on the generated candidate product embeddings **650**. The product scoring module **655** may score the candidate product embeddings **650** based on a distance of the candidate product embedding from the anchor embedding **635**. For example, a product score **660** may be based on a Euclidean distance, a cosine distance, or a dot product of the anchor embedding **635** and a candidate product embedding **650**.

OTHER CONSIDERATIONS

[0058] The present invention has been described in particular detail with respect to one possible embodiment. Those of skill in the art will appreciate that the invention may be practiced in other embodiments. First, the particular naming of the components and variables, capitalization of terms, the attributes, data structures, or any other programming or structural aspect is not mandatory or significant, and the mechanisms that implement the invention or its features may have different names, formats, or protocols. Also, the particular division of functionality between the various system components described herein is merely for purposes of example, and is not mandatory; functions performed by a single system component may instead be performed by multiple components, and functions performed by multiple components may instead be performed by a single component.

[0059] Some portions of the above description present the features of the present invention in terms of algorithms and symbolic representations of operations on information. These algorithmic descriptions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. These operations, while described functionally or logically, are understood to be implemented by computer programs. Furthermore, it has also proven convenient at times, to refer to these arrangements of operations as modules or by functional names, without loss of generality.

[0060] Unless specifically stated otherwise as apparent from the above discussion, it is appreciated that throughout the description, discussions utilizing terms such as “determining” or “displaying” or the like, refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system memories or registers or other such information storage, transmission or display devices.

[0061] Certain aspects of the present invention include process steps and instructions described herein in the form of an algorithm. It should be noted that the process steps and instructions of the present invention could be embodied in software, firmware or hardware, and when embodied in software, could be downloaded to reside on and be operated from different platforms used by real time network operating systems.

[0062] The present invention also relates to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general-purpose computer selectively activated or reconfigured by a computer program stored on a computer readable medium that can be accessed by the computer. Such a computer program may be stored in a non-transitory computer readable storage medium, such as, but is not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, application specific integrated circuits (ASICs), or any type of computer-readable storage medium suitable for storing electronic instructions, and each coupled to a computer system bus. Furthermore, the computers referred to in the specification may include a single processor or may be architectures employing multiple processor designs for increased computing capability.

[0063] The algorithms and operations presented herein are not inherently related to any particular

computer or other apparatus. Various general-purpose systems may also be used with programs in accordance with the teachings herein, or it may prove convenient to construct more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will be apparent to those of skill in the art, along with equivalent variations. In addition, the present invention is not described with reference to any particular programming language. It is appreciated that a variety of programming languages may be used to implement the teachings of the present invention as described herein, and any references to specific languages are provided for enablement and best mode of the present invention.

[0064] The present invention is well suited to a wide variety of computer network systems over numerous topologies. Within this field, the configuration and management of large networks comprise storage devices and computers that are communicatively coupled to dissimilar computers and storage devices over a network, such as the Internet.

[0065] It should be noted that the language used in the specification has been principally selected for readability and instructional purposes, and may not have been selected to delineate or circumscribe the inventive subject matter. Accordingly, the disclosure of the present invention is intended to be illustrative, but not limiting, of the scope of the invention, which is set forth in the following claims.

[0066] As used herein, the terms “comprises,” “comprising,” “includes,” “including,” “has,” “having,” or any other variation thereof, are intended to cover a non-exclusive inclusion. For example, a process, method, article, or apparatus that comprises a list of elements is not necessarily limited to only those elements but may include other elements not expressly listed or inherent to such process, method, article, or apparatus. Further, unless expressly stated to the contrary, “or” refers to an inclusive “or” and not to an exclusive “or”. For example, a condition A or B is satisfied by any one of the following: A is true (or present) and B is false (or not present), A is false (or not present) and B is true (or present), and both A and B are true (or present).

Claims

1. A non-transitory computer-readable medium storing a set of parameters for a plurality of machine-learning models, wherein the set of parameters are generated by a process comprising: storing the set of parameters, where the set of parameters comprise: a set of parameters for a user embedding model comprising a first neural network that generates user embeddings based on user data; a set of parameters for a query embedding model comprising a second neural network that generates query embeddings based on query data; and a set of parameters for an anchor embedding model comprising a third neural network that generates anchor embeddings based on user embeddings and query embeddings, where an anchor embedding is an embedding in an item embedding space; accessing a set of search examples at an online system, where each search example of the set of search examples comprises: user data associated with a user of the online system; query data associated with a search query submitted by the user to the online system; item data associated with a set of items presented to the user by the online system in response to the search query; and a label for each item of the set of items, where each label comprises an indication of whether the user interacted with the item; and for each search example in the set of search examples: generating a user embedding by applying the user embedding model to the user data of the search example, wherein applying the user embedding model to the user data comprises inputting the user data to an input layer of the first neural network of the user embedding model; generating a query embedding by applying the query embedding model to the query data of the search example, wherein applying the query embedding model to the query data comprises inputting the query data to an input layer of the second neural network of the query embedding model; generating an anchor embedding by applying the anchor embedding model to the user embedding and the query embedding, wherein applying the anchor embedding model to the user

embedding and the query embedding comprises inputting the user embedding and the query embedding to an input layer of the third neural network of the anchor embedding model, wherein the input layer of the third neural network of the anchor embedding model is coupled to an output layer of the second neural network of the query embedding model; accessing a set of item embeddings for the set of items in the search example; computing a set of distances between the anchor embedding and the set of item embeddings; and updating the set of parameters for the anchor embedding model based on a loss function applied to the set of distances and the label associated with each item of the set of items.

2. The computer-readable medium of claim 1, wherein applying the loss function comprises: applying a first weight to a first subset of the set of item embeddings, wherein the first subset consists of one or more item embeddings associated with items with which the user performed an interaction of a first type; and applying a second weight to a second subset of the set of item embeddings, wherein the second subset consists of one or more item embeddings associated with items with which the user performed an interaction of a second type, wherein the first weight is different from the second weight, and wherein the first type of interaction is different from the second type of interaction.

3. The computer-readable medium of claim 1, wherein the process further comprises: generating a set of item embeddings by applying an item embedding model to item data of the search example, wherein applying the item embedding model to the item data comprises inputting the item data to an input layer of the neural network of the item embedding model.

4. The computer-readable medium of claim 1, wherein computing the set of distances between the anchor embedding and the set of item embeddings comprises: computing at least one of a Euclidean distance, a cosine distance, or a dot product.

5. The computer-readable medium of claim 1, wherein updating the set of parameters based on the loss function comprises optimizing the set of parameters for: minimizing a distance between the anchor embedding and a first item embedding of the set of item embeddings, wherein the first item embedding is associated with a item with which the user interacted; and maximizing a distance between the anchor embedding and a second item embedding of the set of item embeddings, wherein the second item embedding is associated with an item with which the user did not interact.

6. The computer-readable medium of claim 1, wherein user data for a search example of the set of search examples comprises at least one of: a location of a user, items previously ordered by a user, and a browsing history of a user.

7. The computer-readable medium of claim 1, wherein the query data for a search example of the set of search examples comprises search query text associated with a search query.

8. The computer-readable medium of claim 1, wherein the query data for a search example of the set of search examples comprises context data associated with a session of a user.

9. The computer-readable medium of claim 1, wherein updating the set of parameters comprises updating at least one of: the set of parameters for the user embedding model; the set of parameters for the query embedding model, and the set of parameters for the anchor embedding model.

10. A method comprising: storing a set of parameters, where the set of parameters comprise: a set of parameters for a user embedding model comprising a first neural network that generates user embeddings based on user data; a set of parameters for a query embedding model comprising a second neural network that generates query embeddings based on query data; and a set of parameters for an anchor embedding model comprising a third neural network that generates anchor embeddings based on user embeddings and query embeddings, where an anchor embedding is an embedding in an item embedding space; accessing a set of search examples at an online system, where each search example of the set of search examples comprises: user data associated with a user of the online system; query data associated with a search query submitted by the user to the online system; item data associated with a set of items presented to the user by the online system in response to the search query; and a label for each item of the set of items, where each

label comprises an indication of whether the user interacted with the item; and for each search example in the set of search examples: generating a user embedding by applying the user embedding model to the user data of the search example, wherein applying the user embedding model to the user data comprises inputting the user data to an input layer of the first neural network of the user embedding model; generating a query embedding by applying the query embedding model to the query data of the search example, wherein applying the query embedding model to the query data comprises inputting the query data to an input layer of the second neural network of the query embedding model; generating an anchor embedding by applying the anchor embedding model to the user embedding and the query embedding, wherein applying the anchor embedding model to the user embedding and the query embedding comprises inputting the user embedding and the query embedding to an input layer of the third neural network of the anchor embedding model, wherein the input layer of the third neural network of the anchor embedding model is coupled to an output layer of the second neural network of the query embedding model; accessing a set of item embeddings for the set of items in the search example; computing a set of distances between the anchor embedding and the set of item embeddings; and updating the set of parameters for the anchor embedding model based on a loss function applied to the set of distances and the label associated with each item of the set of items.

11. The method of claim 10, wherein applying the loss function comprises: applying a first weight to a first subset of the set of item embeddings, wherein the first subset consists of one or more item embeddings associated with items with which the user performed an interaction of a first type; and applying a second weight to a second subset of the set of item embeddings, wherein the second subset consists of one or more item embeddings associated with items with which the user performed an interaction of a second type, wherein the first weight is different from the second weight, and wherein the first type of interaction is different from the second type of interaction.

12. The method of claim 10, further comprising: generating a set of item embeddings by applying an item embedding model to item data of the search example, wherein applying the item embedding model to the item data comprises inputting the item data to an input layer of the neural network of the item embedding model.

13. The method of claim 10, wherein computing the set of distances between the anchor embedding and the set of item embeddings comprises: computing at least one of a Euclidean distance, a cosine distance, or a dot product.

14. The method of claim 10, wherein updating the set of parameters based on the loss function comprises optimizing the set of parameters for: minimizing a distance between the anchor embedding and a first item embedding of the set of item embeddings, wherein the first item embedding is associated with a item with which the user interacted; and maximizing a distance between the anchor embedding and a second item embedding of the set of item embeddings, wherein the second item embedding is associated with an item with which the user did not interact.

15. The method of claim 10, wherein user data for a search example of the set of search examples comprises at least one of: a location of a user, items previously ordered by a user, and a browsing history of a user.

16. The method of claim 10, wherein the query data for a search example of the set of search examples comprises search query text associated with a search query.

17. The method of claim 10, wherein the query data for a search example of the set of search examples comprises context data associated with a session of a user.

18. The method of claim 10, wherein updating the set of parameters comprises updating at least one of: the set of parameters for the user embedding model; the set of parameters for the query embedding model, and the set of parameters for the anchor embedding model.

19. A non-transitory computer-readable medium storing instructions that, when executed by a computer system, cause the computer system to perform operations comprising: storing a set of parameters, where the set of parameters comprise: a set of parameters for a user embedding model

comprising a first neural network that generates user embeddings based on user data; a set of parameters for a query embedding model comprising a second neural network that generates query embeddings based on query data; and a set of parameters for an anchor embedding model comprising a third neural network that generates anchor embeddings based on user embeddings and query embeddings, where an anchor embedding is an embedding in an item embedding space; accessing a set of search examples at an online system, where each search example of the set of search examples comprises: user data associated with a user of the online system; query data associated with a search query submitted by the user to the online system; item data associated with a set of items presented to the user by the online system in response to the search query; and a label for each item of the set of items, where each label comprises an indication of whether the user interacted with the item; and for each search example in the set of search examples: generating a user embedding by applying the user embedding model to the user data of the search example, wherein applying the user embedding model to the user data comprises inputting the user data to an input layer of the first neural network of the user embedding model; generating a query embedding by applying the query embedding model to the query data of the search example, wherein applying the query embedding model to the query data comprises inputting the query data to an input layer of the second neural network of the query embedding model; generating an anchor embedding by applying the anchor embedding model to the user embedding and the query embedding, wherein applying the anchor embedding model to the user embedding and the query embedding comprises inputting the user embedding and the query embedding to an input layer of the third neural network of the anchor embedding model, wherein the input layer of the third neural network of the anchor embedding model is coupled to an output layer of the second neural network of the query embedding model; accessing a set of item embeddings for the set of items in the search example; computing a set of distances between the anchor embedding and the set of item embeddings; and updating the set of parameters for the anchor embedding model based on a loss function applied to the set of distances and the label associated with each item of the set of items.

20. The computer-readable medium of claim 19, wherein applying the loss function comprises: applying a first weight to a first subset of the set of item embeddings, wherein the first subset consists of one or more item embeddings associated with items with which the user performed an interaction of a first type; and applying a second weight to a second subset of the set of item embeddings, wherein the second subset consists of one or more item embeddings associated with items with which the user performed an interaction of a second type, wherein the first weight is different from the second weight, and wherein the first type of interaction is different from the second type of interaction.
