

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250286811

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Jain; Ankur et al.

Diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint

Abstract

Systems and methods for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint include receiving a request for k paths, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports; representing the network as a graph with vertices representing ports, edges representing connections between the ports, and with weights assigned to each of the edges; assigning a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same weight; determining up to the k paths between the source port and the dummy destination vertex, and, subsequent to the determining each of the k paths, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value.

Inventors: Jain; Ankur (Gurgaon, IN), Barik; Suvendu Kumar (Bhubaneswar, IN)

Applicant: Ciena Corporation (Hanover, MD)

Family ID: 1000007842614

Assignee: Ciena Corporation (Hanover, MD)

Appl. No.: 18/640349

Filed: April 19, 2024

Foreign Application Priority Data

IN 202411016409

Mar. 07, 2024

Publication Classification

Int. Cl.: H04L45/122 (20220101)

Background/Summary

FIELD OF THE DISCLOSURE

[0001] The present disclosure relates generally to path computation in networking. More particularly, the present disclosure relates to systems and methods for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint.

BACKGROUND OF THE DISCLOSURE

[0002] Path computation in networking includes determining a path from a source node to a destination node in a network, including intermediate nodes, links, etc., based on some criteria, constraints, metrics, etc., e.g., from node A to node Z. The network can be an optical network, finding a wavelength from the source to destination node. Further, the network can be a packet network, finding a path, route, tunnel, Segment Routing (SR) policy, etc. from the source to destination node. Further, path computation can be defined even more granular, namely determining a path from port 1 in the source node to reach port 2 in the destination node. Also, there can be so-called inclusion constraint which further limits the path computation, such a requirement to include node B, as well as ports 3, 4, etc. in the path. Two example algorithms typically used for path computation include Dijkstra's algorithm and Yen's algorithm. Dijkstra's algorithm is used to find the shortest path on a weighted graph, i.e., here the network is modeled as a graph with the nodes or ports being the vertices and the edges being the connections (links). Dijkstra's algorithm can be used to find the shortest path and Yen's algorithm can be used to determine k-shortest paths using the results from Dijkstra's algorithm. Specifically, Yen's algorithm explores deviations from the shortest path to derive additional paths, i.e., the k-shortest paths.

BRIEF SUMMARY OF THE DISCLOSURE

[0003] The present disclosure relates to systems and methods for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint. In the present disclosure, path computation is between vertices which are ports in the network, i.e., optical connections, packet ports, etc., and in which there is a requirement for a node as an inclusion constraint in the path computation. When the vertices are ports, the k-shortest paths tend to share the same vertex at the inclusion constraint, i.e., most or all of the k-shortest paths share the same port at the node being the inclusion constraint. However, it is desirable to have diverse vertices (ports) across the k-shortest paths, but the structure of Yen's algorithm does not tend to provide such results. This is based on the fact Yen's algorithm starts with a given shortest path and then finds deviations therefrom. To solve this issue, the present disclosure includes a modification in Yen's algorithm to change the costs when finding deviating paths, to ensure diversity in edges associated with a node that is an inclusion constraint.

[0004] In various embodiment, the present disclosure includes a method having steps, an apparatus or processing device configured to implement the steps, and a non-transitory computer-readable medium storing instructions that, when executed, cause one or more processors to implement the steps. The steps include receiving a request for k paths, $k > 1$, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports; representing the network as a graph with vertices representing ports including the source port and the plurality of ports, edges representing connections between the ports, and with weights assigned to each of the edges; assigning a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same

weight; and determining up to the k paths between the source port and the dummy destination vertex, and, subsequent to the determining each of the k paths, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value, greater than the same weight. [0005] The steps can further include removing the dummy destination vertex and the associated edges from the k paths to form k shortest paths having diversity across the plurality of ports at a destination node; and providing the k shortest paths in response to the request. In an embodiment, the node can be an intermediate node that is an inclusion constraint and the request further includes a destination port at a destination node, and wherein the steps can further include determining k paths from associated ports of the plurality of ports at the intermediate node to the destination port; and returning k shortest paths based on the k paths between the source port and the intermediate node and the k paths from associated ports of the plurality of ports at the intermediate node to the destination port. In another embodiment, the node is a destination node, such that the k paths include diversity across the plurality of ports at the destination node. The predetermined value can be at least an order of magnitude greater than the weights. The determining can utilize Yen's algorithm. The network can be an optical network, and the ports represent wavelength connections. The network can operate at one or more of Layers 1, 2, and 3, and the ports represent either Time Division Multiplexing (TDM) or packet connections. The steps can be implemented in a Path Computation Engine (PCE).

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0006] The present disclosure is illustrated and described herein with reference to the various drawings, in which like reference numbers are used to denote like system components/method steps, as appropriate, and in which:

[0007] FIG. 1 is a network diagram of a network illustrating nodes and interconnectivity between the nodes.

[0008] FIG. 2 is a network graph of the network with vertices representing ports at the nodes and edges representing connectivity between the ports.

[0009] FIG. 3 is the network graph of FIG. 2 with the addition of a dummy destination vertex connected to all of the ports at the node C.

[0010] FIG. 4 is a network graph of another example network with vertices representing ports at the nodes and edges representing connectivity between the ports, with varying edge costs.

[0011] FIG. 5 is a flowchart of a process for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint.

[0012] FIG. 6 is a block diagram of an example processing device, such as for implementing a PCE.

DETAILED DESCRIPTION OF THE DISCLOSURE

[0013] Again, the present disclosure relates to systems and methods for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint.

Network and Network Graph

[0014] FIG. 1 is a network diagram of a network **10** illustrating nodes **12A-12Z** and interconnectivity between the nodes. The nodes **12A-12Z** can be sites in the network and/or network elements at the sites. The network elements can be switches, routers, terminals, Optical Add/Drop Multiplexers (OADMs), Reconfigurable OADMs (ROADMs), or any other physical equipment which realizes network connectivity. The network **10** can provide connectivity at one or more network layers, including, e.g., optical (Layer 0), Time Division Multiplexing (TDM) (Layer 1), Packet (e.g., Ethernet, Multiprotocol Label Switching (MPLS), Segment Routing, etc.) (Layer 2), Internet Protocol (IP) (Layer 3), and the like. The nodes **12A-12Z** include networking

equipment for realizing connectivity at the one or more network layers. FIG. 1 is a high-level logical view of the network 10. Those skilled in the art will recognize the network elements are hardware devices that include multiple ports that provide the interconnectivity between the nodes. [0015] There can also be a Path Computation Engine (PCE) 14 communicatively coupled to the network 10. The PCE 14 is a processing device that can support path computation, in addition to other functions. The official definition of a PCE is an entity (component, application, or network node) that is capable of computing a network path or route based on a network graph and applying computational constraints. Note, the PCE 14 can be a server, a cloud service, an application, a management system, a controller, etc. For example, the PCE 14 can be a standalone device. In another example, the PCE 14 can be included with a Network Management System (NMS), Software Defined Networking (SDN) controller, network orchestrator, a component in a network element, or the like. Those skilled in the art will recognize the techniques for path computation described herein contemplate operation anywhere and the PCE 14 is shown for illustration purposes only.

[0016] The PCE 14 can be configured to provide paths in response to requests. In doing so, the PCE 14 can maintain a network graph and use Dijkstra's algorithm and Yen's algorithm. FIG. 1 provides a view of nodal connectivity, but at the site level. For the PCE 14 to compute actual paths at the port level, the network graph is based on ports. FIG. 2 is a network graph 20 of the network 10 with vertices 22 representing ports at the nodes and edges 24 representing connectivity between the ports. Note, there can be some metric associated with each edge used in the path computation to find the constrained shortest path. At the optical layer, the ports are wavelengths that are ultimately multiplexed onto fibers that interconnect between the nodes 12A-12Z. At the higher layers (Layer 1-3, etc.), the ports are logical supporting TDM channels (Layer 1), packet flows (Layer 2), etc. that flow between the nodes 12A-12Z. The present disclosure will use the term port generally to refer to some hardware point where connectivity is realized. A computed path from a source node to a destination node will include a listing of ports from an ingress port on the source node to an egress port on the destination port, along with an ordered listing of all intermediate ports. Those skilled in the art will recognize Dijkstra's algorithm and Yen's algorithm can be used on the network graph 20 to k-shortest paths between nodes, at the port level.

[0017] The nodes 12A-12Z can be sites in the network and/or network elements at the sites. That is, the term node is used to denote either a specific network element at a specific site, as well as an entire site. For example, in FIG. 1, the nodes 12A, 12Z are single network elements at the sites A, Z, respectively, i.e., the source node 12A and the destination node 12Z. This makes sense as the source and destination are defined while the intermediate ports, interfaces, network elements, etc. are being determined by the path calculation. Also, it is even possible to define the destination node 12Z in terms of different network elements as well. At the intermediate sites, e.g., sites B, C, D, E, there can be one or more network elements at these sites, and the present disclosure refers to all of the ports or interfaces C1, C2, C3 as part of the node 12C, even if there are multiple network elements. Again, the present disclosure deals with getting multiple vertices at the node 12C, i.e., diversity across different interfaces at either the intermediate nodes, such as the node 12C, as well as even at the destination node 12Z. So, the term nodes can mean one or more network elements at a given site.

Path Computation

[0018] A normal Constrained Shortest Path First (CSPF) algorithm, such as based on Dijkstra's algorithm, returns just a path. We can look at multiple paths by using a k-shortest path algorithm (k is an integer greater than one). As implied by the name, the k-shortest algorithm gives the k shortest paths between two nodes, rather than just the one absolute shortest path between them. By definition, a k-shortest CSPF computation will take longer to run than a 1-shortest CSPF computation. For a 1-shortest CSPF, once we find a path from source (S) to destination (D), we know it has to be the shortest, so we are done. For k-shortest, the path computation has to keep

going until we find k such paths from S to D . In the meantime, we are exploring a variety of paths and subparts that meander from what our 1-shortest path would be. There have been a number of implementations that have looked at the most efficient way of doing this; here we will look at an adaptation of Yen's algorithm. Yen's algorithm is described in Yen, Jin Y. "Finding the k shortest loopless paths in a network." *Management Science* 17.11 (1971): 712-716, the contents of which are incorporated by reference.

[0019] The following pseudocode is an implementation of Yen's algorithm to find k -shortest paths from a source (S) to destination (D) in a network which is represented by the network graph **20**. Note, in the context of the network graph **20** and the following descriptions, node is a vertex which is a port.

```
TABLE-US-00001 A[0] = RegularCSPF(Graph, S, D); if (A[0] == NULL) return; Heap = [ ]; for
(k = 1 to K) { for (i = 0 to size(A[k-1]) - 2) { // for every node from the first to the next to last
in the previous path spurNode = A[k-1].node(i); rootPath = A[k-1].nodes(0,i);
removeEdges = <for any previous path who also contains the root path, add the next node in that
path to this set> Graph.remove(removeEdges); spurPath = RegularCSPF(Graph,
spurNode, D); if (spurPath == NULL) continue; totalPath = append(rootPath,
spurPath); if (Heap.contains(totalPath) == false) Heap.add(totalPath);
Graph.restore(removeEdges); } if (Heap.empty( )) return; Heap.sort( ); A[k] =
Heap.pop( ); }
```

[0020] Yen's algorithm uses two lists, i.e., list A ($A[0]$ to $A[k]$ in the above) (permanent shortest paths from source to destination—chronologically ordered) and a second list denoted as a heap (tentative/candidate shortest paths). At first you have to find the 1st shortest path from the source to destination using any well-suited shortest path algorithm (e.g., Dijkstra). Then the Yen algorithm exploits the idea that the k -th shortest paths may share edges and sub-paths (path from source to any intermediary nodes within the route) from $(k-1)$ -th shortest path. Then you have to take $(k-1)$ th shortest path and make each node in the route unreachable in turn, i.e., rub off particular edge that goes to the node within the route. Once the node is unreachable, find the shortest path from the preceding node to the destination. Then you have a new route which is created by appending the common sub-path (from source to the preceding node of the unreachable node) and adds the new shortest path from preceding node to destination. This route is then added to the heap, provided that it has not appeared in list A or the heap before. After repeating this for all nodes in the route, you have to find the shortest route in heap and move that to list A . This is repeated until there are k paths in the list A .

[0021] Also, a very simplified explanation would be that we start with the 1-shortest path, and then for subsequent paths we try to find ways to divert off of paths we have found thus far, by examining at nodes along the way of the most recent path and eliminating paths we have already taken out of those nodes from contention, to make sure what we find at that point is something new. Note that this algorithm still makes use of the 1-shortest algorithm, both to find the 0th path and to find paths from the spur node to D when we are trying to find diversions. A spur node and spur path are as defined in Yen's algorithm, namely a node from which there is a diversion and a path for the diversion, from a root path.

Motivation

[0022] Again, the network graph **20** and the vertices (nodes) are based on ports. There are cases where a path request is based on network elements or a site. In such cases, if the result comes out to be using a single port in the network element for most or all of the k -paths, it is not desirable, and the user is interested in options with diverse ports being selected for k -path options. That is, based on the inherent functionality of Yen's algorithm, it was determined that, when there are site inclusions, e.g., must include the node **12C** in the k -shortest paths between the node **12A** and the node **12Z**, and when the network graph **20** is based on ports, the k -shortest path results tend to come back with most paths having the same port at the site inclusion. This is the case even though

the node **12C** has multiple ports (e.g., C1, C2, C3, etc.).

[0023] The following provide examples using the network graph **20** in FIG. 2, where each edge has an equal weight of 1. Note, after that, we present some examples in a denser graph with variable weights on the edges. Of course, those skilled in the art will recognize these network graphs **20** are for a very simplified network in terms of ports, network elements, and sites, for ease of illustration, and practical embodiments will typically be significantly larger in the corresponding network graph, i.e., significantly more edges and vertices.

[0024] For example, to illustrate the motivation, in FIG. 2, assume we want to find 3 paths from vertex A to the node **12C**, as part of a path from the vertex A to the vertex Z. Specifically, the node **12C** is an ordered inclusion in a path request from the node **12A** to the node **12Z**, with a request for 3 paths. Of note, the site **12C** has three vertices C1, C2, C3. To handle the node **12C** as an inclusion constraint, a dummy or virtual vertex can be added in the network graph **20** to ensure the inclusion constraint is met. FIG. 3 is the network graph **20** of FIG. 2 with the addition of a dummy destination vertex **30** connected to all of the ports at the node **12C**. The dummy destination vertex **30** is not an actual port in the network **10**, but rather a virtual construct to ensure all paths go through the node **12C**, which is the inclusion constraint.

[0025] With the current implementation of Yen's algorithm and the dummy destination vertex **30**, the 3 computed paths would be A-B1-C1, A-B2-C1, A-B11-B2-C1, which all are on C1. A more detailed explanation follows. Assume the network graph **20** is a weighted graph with each edge having an equal weight (e.g., 1), and Dijkstra's algorithm and Yen's algorithm are used to find K paths between the end points. In this example, we intend to find k paths using Yen's algorithm from the vertex A to the node **12C** which means reaching any of the vertex C1, C2, C3 from the vertex A.

[0026] If k=4 shortest paths are requested, all will end up on vertex C1 at node **12C**, i.e., A-C1, A-B11-C1, A-B2-C1, A-B1-B2-C1 as these 4 are the shortest path to reach from vertex A to the node **12C**. State differently, the network graph **20** is purposefully constructed to show clearly that all shortest paths from the vertex A to the node **12C** are to the vertex C1. This is not always acceptable to a user as they would like to have some flexibility in k path computations to have unique destination points at the node **12C**.

[0027] Now, for a path computation from the vertex A to the vertex Z with a site inclusion of the node **12C**, for k-paths, for any typical inclusion algorithm this involves finding K' segments from vertex A to any of vertices C1, C2, C3. To achieve this, the dummy destination vertex **30** with edges **32, 34, 36** are created in graph to connect to vertices C1, C2, C3, see FIG. 3. The modified graph in FIG. 3 includes the dummy destination vertex **30** (labeled as Dummy_DEST) with vertex and dummy edges **32, 34, 36** from vertices C1, C2, C3 to the dummy destination vertex **30**. Again, we use the term "dummy" to note these are virtual constructs. Yen's algorithm is used to find k' shortest path from the vertex A to the dummy destination vertex **30**, and then for each of the partial path found above, discarding the dummy destination vertex **30** and edges k' shortest path is requested from each of the Cx vertex (the last end point of the partial path) to Z. At the end, all the of partial paths are combined to return k-shortest paths from vertex A to vertex Z, all of which will include the node **12C** due to the dummy destination vertex **30**.

[0028] That is, the use of the dummy destination vertex **30** for site inclusions means we break the path computation into two parts and combine the result. In this example, we compute from the vertex A to the node **12C** with the dummy destination vertex **30**, and then from then from the vertex found in the node **12C** to the vertex Z. The result is then combined. Now, for illustration purposes, the vertex C1 intentionally has no connectivity to the vertex Z. Thus, any path computation requiring the node **12C** as a site inclusion would not return a valid path.

[0029] For more details, e.g., with k as 2 and k' (max segments to each included point) as 3 in FIG. 3 with the node **12C** as an inclusion constraint, the flow would be as follows: [0030] (1) Find 3 (k') shortest paths from vertex A to dummy destination vertex **30** using Yen's algorithm which in turn

uses multiple Dijkstra algorithm calls. [0031] (2) The segments found from step (1) would be A-C1-DUMY_DEST, A-B1-C1-DUMY_DEST, A-B11-B2-C1-DUMMY_DEST because these two are the shortest paths. [0032] (3) Now removing the dummy destination vertex **30** and edges the segment would be A-B1-C1, A-B2-C1. [0033] (4) From C1 find 2 shortest paths to destination Z. [0034] (5) But there's no path between C1 to Z so this would fail to find path from A to Z by including node **12C**. If we had found a segment to vertices C2, C3 then we could have successfully found path from C2/C3-Z hence from A-Z valid paths have been returned via C2/C3. [0035] Similarly, if three shortest paths were requested from vertex A to vertex E1 with the node **12C** as an inclusion constraint, which is solved by breaking the problem into two parts, computing path from A-C and the from C-E1, would result in similar problem that all paths are resulted through port C1, and maybe the path through C1 is not even valid in some post processing logic. This means that we need to find diverse options for the node **12C**, so that even if a particular vertex is not valid, we have other options in the list as well.

Conventional Approach

[0036] A solution can include an enumerative approach—find all possible paths and then pick the diverse endpoint paths. This has large time and space complexity (e.g., the heap in Yen's algorithm will grow). Dijkstra's algorithm can compute single shortest path from 1 to many vertices in a single run, but there are 2 issues: [0037] (1) Generally, the network or PCE graph in our domain have lot of validation constraints all of which are not solved in pre-compute, there are lot of cases when certain edges of the path are not compatible with each other, this results in lot of invalid paths and Dijkstra's algorithm, cannot be used, even to compute a single shortest path we need to have Yen's algorithm. [0038] (2) In the case of 1 . . . n Dijkstra's algorithm, we will just get single path from 1 source to each of n destinations but there are cases where we still need to compute more than 1 path to a destination because the k-path requirement might not get fulfilled with n paths (or less than n because all destinations may not have a path). [0039] Because many paths will fail to satisfy this diversity, this approach means path generation runs to large k. This is costly in terms of both time and memory in a way directly follows from k-shortest path. As this is also sensitive to network size, as network size grows the complexity will be more pronounced.

Solution

[0040] The present disclosure provides an approach to find diverse endpoints directly, in Yen's algorithm. Specifically, the present disclosure includes a modification of Yen's algorithm with the dummy destination vertex **30** that results in the k-paths having diversity of ports in an included site, by making diversity a priority over cost-based k-shortest paths. Again, as noted above, the conventional approach to Yen's algorithm causes most (or all in our example case) ports to be vertex C1 on the k-shortest paths. This is problematic as the vertex C1 cannot reach the end vertex Z.

[0041] The present disclosure involves changing the costs of the edges **32, 34, 36** during successive iterations of Yen's algorithm, i.e., namely adding an additional step as detailed below. Let's use the same example from above in FIG. 3 where we want to find three paths from the vertex A to the node **12C**, such as an intermediate step in a path computation from vertex A to vertex Z with the node **12C** as an inclusion constraint. In FIG. 3, assume the network graph **20** is a weighted graph and each edge **24, 32, 34, 36** has a default value, e.g., 1. At an example of computing a path from vertex A to node **12C** with port diversity includes: [0042] (1) At first, the dummy destination vertex **30** is created and added in the network graph **20** with the edges **32, 34, 36** connected to all of the vertices C1, C2, C3 associated with the inclusion constraint of the node **12C**, so that the problem statement is reduced to calculating the path from the vertex A to the dummy destination vertex **30**. [0043] (2) Using Yen's algorithm, the first shortest segment from the vertex A to the dummy destination vertex **30** is computed as A-C1-DUMMY_DEST. [0044] (3) The present disclosure includes adding a new step where after finding one path (e.g., A-C1-DUMMY_DEST), the cost

from the vertex C1 to the dummy destination vertex **30** via the edge **32** is increased, such as to some high value (say 1000) so that this edge is considered last in subsequent path computations in Yen's algorithm, thereby focusing on diversity in the vertices found at the inclusion constraint, the node **12C**.

[0045] Note, we use the term high value to mean something larger than the typical weights on the edges in the weighted graph. For example, here the edges all have a weight of 1, so the high value can be 1000, or even 100, or something at least an order of magnitude more than the typical weight. By assigning this to the high value, the edge to the dummy destination vertex **30** is still considered, but only as a last resort. [0046] (4) Yen's algorithm is continued from thereon, and the second shortest path segment from the vertex A to the dummy destination vertex **30** now would be A-B1-B2-B3-C2-DUMMY-DEST as the cost of this segment would be 4 and segments having the vertex C1 will not be the shortest path any more as the cost will always be 1000+. That is, changing the weight of the edge **32** to the high value ensures we select the vertex C2 instead of C1 for the next shortest path. Stated differently, this example is provided to show that absent the change of weights, all shortest paths will end on the vertex C1. The approach of changing the weights allows diversity in the selection of vertices at the node **12C**. [0047] (5) Now, since this second path has a different vertex C2 at the node **12C**, we increase the cost of the edge **34** to the high value. At the same time, we leave the edge **32** also at the high value for the next shortest path determination.

[0048] (6) The third shortest path segment now would be A-B11-B2-B3-B4-C3-DUMMY-DEST as any other segments going via the vertices C1, C2 will have now weight of 1000+.

[0049] And this can continue. Of note, each time a path is found to a specific vertex at the node **12C**, a corresponding weight of its edge to the dummy destination vertex **30** is increased to the high value. As such, the successive iterations can still determine paths to the specific vertex, but they will have high costs. In this simple example, there are three vertices C1, C2, C3 at the node **12C**, so all of the resulting k-shortest paths to the node **12C** will be diversified across all of the vertices C1, C2, C3. After the third shortest path, the edge **36** will also be assigned the high value for its weight, such that all three edges **32**, **34**, **36** will have the high value. Subsequent path computations for the fourth path and beyond will now treat the vertices C1, C2, C3 the same (all have the high value to the dummy destination vertex **30**). As such, the fourth path and subsequent paths will tend to land on the vertex C1, based on it clearly being the shortest path in this example.

[0050] However, when returning all of the k-shortest paths, we will have some that end at the vertices C2, C3. This is important in this example since there is no path from the vertex C1 to the vertex Z. So, in the next step to determine a path from the site inclusion of node **12C** to the vertex Z, all of the paths with the vertex C1 can be discarded.

[0051] Continuing this example, now we have three segments from the vertex A to the node **12C**—A-C1-DUMMY_DEST, A-B1-B2-B3-C2-DUMMY_DEST, A-B1-B2-B3-B4-C3-DUMMY_DEST. To continue to the vertex Z, the dummy destination vertex **30** and the edges **32**, **34**, **36** are removed from the network graph **20**, and our modified segments are A-C1, A-B1-B2-B3-C2, A-B1-B2-B3-B4-C3. Now when we find segments from each of the vertices C1, C2, C3 to the vertex Z, there is not a path from C1-Z but from C2-Z and C3-Z we will have the path. The segments found at the point would be C2-Z, C3-Z. Combining the segments found so far there will be two shortest paths from A-Z that would be A-B1-B2-B3-C2-Z, A-B1-B2-B3-B4-C3-Z. Thus, this approach yields a valid path from the vertex A to the vertex Z, with an inclusion constraint of the node **12C**.

[0052] Note, the algorithm is changing the cost of the last edge **32**, **34**, **36** only to the dummy destination vertex **30**. In Yen's algorithm, already computed paths are preserved in an A array (often called a root path), and a spur path is computed from sub-paths of root path, and changing the cost of any edge of the path in the A array can cause issue. However, here we are just changing the cost of the last edge **32**, **34**, **36** which is connected to the dummy destination vertex **30** which can never become a root path of subsequent paths, so this change of cost of the last edge is safe from the overall algorithm perspective. That is, we will never explore any further paths with the edges **32**,

34, 36 because these are only virtual constructs used to ensure we get the vertices C1, C2, C3 because there are associated with the inclusion constraint of node **12C**.

[0053] Thus, the approach described herein provides an extension to Yen's k-shortest path algorithm to support k-shortest paths with priority to endpoint diversity over default cost-based k-shortest paths. This is not limited to destination endpoint but also takes care of a scenario when multiple graph vertices are part of an inclusion constraint with a OR operator, which means that pick any one or more of the graph vertices provided in inclusion constraint. In that case, this algorithm will result in paths spanning through different vertices of the inclusion list. That is, one use case is an inclusion constraint of a site, i.e., the node **12C**, which includes multiple vertices C1, C2, C3 (i.e., physical network ports). Other use cases are also contemplated, namely a graph having multiple inclusion constraints (vertices) with an or operator. The formal definition can be to provide a path from vertex A to vertex Z and include any of vertices C1, C2, C3. With the present disclosure, we will get k-shortest paths which prioritize diversity of the vertices C1, C2, C3 over the shortest path which we find tends to favor only one vertex (C1 in our example).

Another Example with a Weighted Graph Having Variable Weights

[0054] FIG. **4** is a network graph **40** of another example network with vertices representing ports at the nodes and edges representing connectivity between the ports, with varying edge costs. The network graph **40** is a weighted graph with varying edge cost, versus the network graph **20** which had each edge being the same weight. This is generally the case in real networks. Again, assume a user intends to find k paths using Yen's algorithm from the vertex A to the node **12C**, which means reaching any of the vertex C1, C2, C3 from the vertex A. In this case, the dummy destination vertex **30** is created connecting all applicable vertices C1, C2, C2 of the node **12C**, each dummy edge **32, 34, 36** has a weight of 1.

[0055] In the above case if k=3, the shortest paths from the vertex the dummy destination vertex **30** would be [0056] (1) A-B1-C1-Dummy-DEST(cost=3), [0057] (2) A-B2-C1-Dummy-DEST (cost=4), [0058] (3) A-B3-C1-Dummy-Dest (Cost=4)

[0059] Discarding the Dummy-DEST path found would be: [0060] (1) A-B1-C1-(cost=2), [0061] (2) A-B2-C1 (cost=3), [0062] (3) A-B3-C1 (Cost=3)

[0063] All the paths end up on the vertex C1 only as paths to vertices C2, C3 from the vertex A are of higher cost. Again, this is not always acceptable to customers as they would like to have some flexibility in k path computations to have diverse destination points in a site for each path.

[0064] Now, again, assume the objective is to find k paths from the vertex A to the vertex Z. In the network graph **40**, if the user intends to find k paths from the vertex A to the vertex Z via the node **12C** which means traversing through any of the vertices C1, C2, C3 to reach vertex Z.

[0065] For any typical inclusion algorithm this involves finding k' segments from the vertex A to any of the vertices C1, C2, C3. To achieve this, the dummy destination vertex **30** with edges **32, 34, 36** are created in the graph **40** to connect to the vertices C1, C2, C3. k' paths are calculated from the vertex A to the dummy destination vertex **30**, and, for each of the paths found, after discarding dummy edge/vertex, k' paths are calculated from destination to vertex Z.

[0066] With the present disclosure, to find a path from the vertex A to the vertex Z with the node **12C** as an inclusion constraint, with k=3, the flow would be as follows. Also, when multiple paths have the same cost, any of the paths is chosen as shortest path which is deterministic and the same is not related to this solution, here focus is the diverse end point aspect of shortest paths found using Yen's algorithm. [0067] (1) Find k'=3 segments from the vertex A to the dummy destination vertex **30**. [0068] (a) the first shortest path segment from A to the dummy destination vertex **30** is A-B1-C1-Dummy-DEST (Cost=3) [0069] (b) Increase the cost of the edge **32** to 1000 (any higher value), i.e., since the edge **32** connects the dummy destination vertex **30** to the vertex C1. [0070] (c) the second shortest path segment would be A-B2-C2-Dummy-DEST (Cost=5) as all paths through vertex C1 would now have cost of 1000+. [0071] (d) Increase cost of the edge **32** to 1000 (any higher value), i.e., since the edge **34** connects the dummy destination vertex **30** to the vertex

C2. [0072] (e) the third shortest path segment would be A-B3-C3-Dummy-DEST (Cost=6) as all shortest paths through the vertices C1, C2 will have cost of 1000+. [0073] (f) Segments found so far are A-B1-C1-Dummy-DEST, A-B2-C2-Dummy-DEST, A-B3-C3-Dummy-DEST. [0074] (2) Excluding the dummy destination vertex **30** and the edges **32, 34, 36** the modified segments would be [0075] A-B1-C1 (Cost=2), A-B2-C2 (Cost=4), A-B3-C3 (Cost=5) [0076] (3) Now find segments ($k'=3$) from each of the vertices C1, C2, C3 to the destination vertex Z. The shortest segments found would be [0077] (i) C1-D1-Z (Cost=3), C1-D2-Z (Cost=4), C1-D2-C3-Z (Cost=5) [0078] (ii) C2-Z (Cost=1), C2-D2-Z (Cost=4), C2-D2-D1-Z (Cost=5) [0079] (iii) C3-Z (Cost=3), C3-C2-Z (Cost=4), C3-D2-Z (Cost=4) [0080] (4) Combining segments found from step (2) and step (3), the three k-shortest paths would be [0081] (i) A-B2-C2-Z (Cost=5), [0082] (ii) A-B1-C1-D1-Z (Cost=5), [0083] (iii) A-B1-C1-D2-Z (Cost=6) [0084] This approach picks diverse vertices in the site inclusion, the node **12C**. Advantageously, this reduces the risk of not finding a path from the intermediate vertex to the destination due to the increased diversity in the intermediate included site vertices.

Process

[0085] FIG. 5 is a flowchart of a process **100** for diversifying k-shortest path computation results for multiple vertices as destination or as inclusion constraint. The process **100** contemplates implementation as a method having steps, via a processing device **200** (see FIG. 6) configured to implement the steps, via the PCE **14** configured to implement the steps, and as a non-transitory computer-readable medium storing instructions that, when executed, cause one or more processors to implement the steps.

[0086] The steps include receiving a request for k paths, $k > 1$, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports (step **102**); representing the network as a graph with vertices representing ports including the source port and the plurality of ports, edges representing connections between the ports, and with weights assigned to each of the edges (step **104**); assigning a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same weight (step **106**); and determining up to the k paths between the source port and the dummy destination vertex, and, subsequent to the determining each of the k paths, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value, greater than the same weight (step **108**).

[0087] Again, the dummy destination vertex and the associated edges are virtual constructs that do not actually correlate to any physical ports or connections in the network. Rather, these are used to ensure the path is computed to any of the plurality of ports, as an inclusion constraint. In the network application, the inclusion constraint is a site, i.e., the node **12C**, and the plurality of ports are the vertices C1, C2, C3. The process **100** can include removing the dummy destination vertex and the associated edges from the k paths to form k shortest paths having diversity across the plurality of ports at the destination node; and providing the k shortest paths in response to the request.

[0088] In an embodiment, the node is an intermediate node that is an inclusion constraint and the request further includes a destination port at a destination node, and wherein the process **100** can further include determining k paths from associated ports of the plurality of ports at the intermediate node to the destination port; and returning k shortest paths based on the k paths between the source port and the intermediate node and the k paths from associated ports of the plurality of ports at the intermediate node to the destination port. In another embodiment, the node is a destination node, such that the k paths include diversity across the plurality of ports at the destination node. The predetermined value can be at least an order of magnitude greater than the weights. That is, the predetermined value can be some arbitrary value, as long as it is high enough to bias against the already chosen vertices for the plurality of ports. That is, we want diversity in the plurality of ports in the k paths.

[0089] The determining can utilize Yen's algorithm. In an embodiment, the network is an optical network, and the ports represent wavelength connections. In another embodiment, the network operates at one or more of Layers 1, 2, and 3, and the ports represent either Time Division Multiplexing (TDM) or packet connections. The process **100** can be implemented in a Path Computation Engine (PCE).

Example Processing Device

[0090] FIG. **6** is a block diagram of an example processing device **200**, such as for the PCE **14**. The processing device **200** can be part of a network element, or a stand-alone device communicatively coupled to the network element, such as the PCE **14**, etc. Also, the processing device **200** can be referred to in implementations as a control module, a shelf controller, a shelf processor, a system controller, etc. The processing device **200** can include a processor **202** which is a hardware device for executing software instructions. The processor **202** can be any custom made or commercially available processor, a central processing unit (CPU), an auxiliary processor among several processors associated with the processing device **200**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the processing device **200** is in operation, the processor **202** is configured to execute software stored within the memory, to communicate data to and from the memory, and to generally control operations of the processing device **200** pursuant to the software instructions. The processing device **200** can also include a network interface **204**, a data store **206**, memory **208**, an I/O interface **210**, and the like, all of which are communicatively coupled to one another and to the processor **202**.

[0091] The network interface **204** can be used to enable the processing device **200** to communicate on a data communication network, such as to communicate to a management system, or the like. The network interface **204** can include, for example, an Ethernet module. The network interface **204** can include address, control, and/or data connections to enable appropriate communications on the network. The data store **206** can be used to store data, such as control plane information, provisioning data, Operations, Administration, Maintenance, and Provisioning (OAM&P) data, etc. The data store **206** can include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, flash drive, CDROM, and the like), and combinations thereof.

[0092] Moreover, the data store **206** can incorporate electronic, magnetic, optical, and/or other types of storage media. The memory **208** can include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, flash drive, CDROM, etc.), and combinations thereof. Moreover, the memory **208** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **208** can have a distributed architecture, where various components are situated remotely from one another, but may be accessed by the processor **202**. The I/O interface **210** includes components for the processing device **200** to communicate with other devices.

[0093] Note, while the various descriptions relate to the processing device **200** and the PCE **14**, those skilled in the art will appreciate a PCE can be located in an SDN controller, in-skin (within) a network element, hosted in the cloud, in a management system, and the like. Also, the present disclosure is not limited to being implemented in the PCE **14**, rather any device can perform the various functions described herein.

CONCLUSION

[0094] It will be appreciated that some embodiments described herein may include one or more generic or specialized processors ("one or more processors") such as microprocessors; Central Processing Units (CPUs); Digital Signal Processors (DSPs); customized processors such as Network Processors (NPs) or Network Processing Units (NPU), Graphics Processing Units (GPUs), or the like; Field Programmable Gate Arrays (FPGAs); and the like along with unique stored program instructions (including software and/or firmware) for control thereof to implement,

in conjunction with certain non-processor circuits, some, most, or all of the functions of the methods and/or systems described herein. Alternatively, some or all functions may be implemented by a state machine that has no stored program instructions, or in one or more Application-Specific Integrated Circuits (ASICs), in which each function or some combinations of certain of the functions are implemented as custom logic or circuitry. Of course, a combination of the aforementioned approaches may be used. For some of the embodiments described herein, a corresponding device in hardware and optionally with software, firmware, and a combination thereof can be referred to as “circuitry configured or adapted to,” “logic configured or adapted to,” “a circuit configured to,” “one or more circuits configured to,” etc. perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. on digital and/or analog signals as described herein for the various embodiments.

[0095] Moreover, some embodiments may include a non-transitory computer-readable storage medium having computer-readable code stored thereon for programming a computer, server, appliance, device, processor, circuit, etc. each of which may include a processor to perform functions as described and claimed herein. Examples of such computer-readable storage mediums include, but are not limited to, a hard disk, an optical storage device, a magnetic storage device, a Read-Only Memory (ROM), a Programmable Read-Only Memory (PROM), an Erasable Programmable Read-Only Memory (EPROM), an Electrically Erasable Programmable Read-Only Memory (EEPROM), Flash memory, and the like. When stored in the non-transitory computer-readable medium, software can include instructions executable by a processor or device (e.g., any type of programmable circuitry or logic) that, in response to such execution, cause a processor or the device to perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. as described herein for the various embodiments.

[0096] Although the present disclosure has been illustrated and described herein with reference to embodiments and specific examples thereof, it will be readily apparent to those of ordinary skill in the art that other embodiments and examples may perform similar functions and/or achieve like results. All such equivalent embodiments and examples are within the spirit and scope of the present disclosure, are contemplated thereby, and are intended to be covered by the following claims. Further, the various elements, operations, steps, methods, processes, algorithms, functions, techniques, modules, circuits, etc. described herein contemplate use in any and all combinations with one another, including individually as well as combinations of less than all of the various elements, operations, steps, methods, processes, algorithms, functions, techniques, modules, circuits, etc.

Claims

1. A non-transitory computer-readable medium comprising instructions that, when executed, cause one or more processors to implement steps of: receiving a request for k paths, $k > 1$, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports; representing the network as a graph with vertices representing ports including the source port and the plurality of ports, edges representing connections between the ports, and with weights assigned to each of the edges; assigning a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same weight; and determining up to the k paths between the source port and the dummy destination vertex, and, subsequent to the determining each of the k paths, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value, greater than the same weight.
2. The non-transitory computer-readable medium of claim 1, wherein the steps further include removing the dummy destination vertex and the associated edges from the k paths to form k shortest paths having diversity across the plurality of ports at a destination node; and providing the

k shortest paths in response to the request.

3. The non-transitory computer-readable medium of claim 1, wherein the node is an intermediate node that is an inclusion constraint and the request further includes a destination port at a destination node, and wherein the steps include determining k paths from associated ports of the plurality of ports at the intermediate node to the destination port; and returning k shortest paths based on the k paths between the source port and the intermediate node and the k paths from associated ports of the plurality of ports at the intermediate node to the destination port.

4. The non-transitory computer-readable medium of claim 1, wherein the node is a destination node, such that the k paths include diversity across the plurality of ports at the destination node.

5. The non-transitory computer-readable medium of claim 1, wherein the predetermined value is at least an order of magnitude greater than the weights.

6. The non-transitory computer-readable medium of claim 1, wherein the determining utilizes Yen's algorithm.

7. The non-transitory computer-readable medium of claim 1, wherein the network is an optical network, and the ports represent wavelength connections.

8. The non-transitory computer-readable medium of claim 1, wherein the network operates at one or more of Layers 1, 2, and 3, and the ports represent either Time Division Multiplexing (TDM) or packet connections.

9. The non-transitory computer-readable medium of claim 1, wherein the method is implemented in a Path Computation Engine (PCE).

10. A method comprising steps of: receiving a request for k paths, $k > 1$, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports; representing the network as a graph with vertices representing ports including the source port and the plurality of ports, edges representing connections between the ports, and with weights assigned to each of the edges; assigning a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same weight; and determining up to the k paths between the source port and the dummy destination vertex, and, subsequent to the determining each of the k paths, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value, greater than the same weight.

11. The method of claim 10, wherein the steps further include removing the dummy destination vertex and the associated edges from the k paths to form k shortest paths having diversity across the plurality of ports at a destination node; and providing the k shortest paths in response to the request.

12. The method of claim 10, wherein the node is an intermediate node that is an inclusion constraint and the request further includes a destination port at a destination node, and wherein the steps include determining k paths from associated ports of the plurality of ports at the intermediate node to the destination port; and returning k shortest paths based on the k paths between the source port and the intermediate node and the k paths from associated ports of the plurality of ports at the intermediate node to the destination port.

13. The method of claim 10, wherein the node is a destination node, such that the k paths include diversity across the plurality of ports at the destination node.

14. The method of claim 10, wherein the predetermined value is at least an order of magnitude greater than the weights.

15. The method of claim 10, wherein the determining utilizes Yen's algorithm.

16. The method of claim 10, wherein the network is an optical network, and the ports represent wavelength connections.

17. The method of claim 10, wherein the network operates at one or more of Layers 1, 2, and 3, and the ports represent either Time Division Multiplexing (TDM) or packet connections.

18. The method of claim 10, wherein the method is implemented in a Path Computation Engine (PCE).

- 19.** A system comprising: one or more processors; and memory storing instructions that, when executed, cause the one or more processors to receive a request for k paths, $k > 1$, in a network from a source port at a source node to a node, wherein the node includes a plurality of ports, represent the network as a graph with vertices representing ports including the source port and the plurality of ports, edges representing connections between the ports, and with weights assigned to each of the edges, assign a dummy destination vertex to connect to each of the plurality of ports and associated edges between each of the plurality of ports and the dummy destination vertex having a same weight, and determine up to the k paths between the source port and the dummy destination vertex, and, subsequent to the each of the k paths being determined, assigning a corresponding edge of the associated edges in each of the k paths a predetermined value, greater than the same weight.
- 20.** The system of claim 9, wherein the instructions that, when executed, further cause the one or more processors to remove the dummy destination vertex and the associated edges from the k paths to form k shortest paths having diversity across the plurality of ports at a destination node, and provide the k shortest paths in response to the request.
-