

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250286812

Kind Code

A1

Publication Date

September 11, 2025

Inventor(s)

Friedman; Yamin et al.

REMOTE DIRECT MEMORY ACCESS (RDMA) MULTIPATH

Abstract

Technologies for spreading a burst of data across multiple network paths in remote direct memory access (RDMA) over converged Ethernet (ROCE) and InfiniBand are described. A RDMA adapter receives, from a requestor device over a local interface, a request to send data of a transport flow directed to a target device over a network interface, and one or more parameters being related to a multipath selection by the network controller. The RDMA adapter sends a first burst of data of the transport flow via a first network path to the target device. The RDMA adapter identifies, using the one or more parameters, a second network path to the target device. The RDMA adapter sends the second burst of data to the target device on the second network path.

Inventors: Friedman; Yamin (Jerusalem, IL), Borshteen; Idan (Karmiel, IL), Moyal; Roee (Yokne'am Illit, IL), Shpigelman; Yuval (Netanya, IL)

Applicant: Mellanox Technologies, Ltd. (Yokneam, IL)

Family ID: 1000008621161

Appl. No.: 19/204775

Filed: May 12, 2025

Related U.S. Application Data

parent US continuation 18443928 20240216 parent-grant-document US 12323320 child US 19204775

parent US continuation 17901671 20220901 parent-grant-document US 11909628 child US 18443928

Publication Classification

Int. Cl.: H04L45/12 (20220101); H04L43/0864 (20220101); H04L45/24 (20220101); H04L47/52 (20220101)

U.S. Cl.:

CPC **H04L45/124** (20130101); **H04L43/0864** (20130101); **H04L45/24** (20130101);
 H04L47/52 (20130101);

Background/Summary

RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 18/443,928, filed Feb. 16, 2024, which is a continuation of U.S. patent application Ser. No. 17/901,671, filed Sep. 1, 2022, now U.S. Pat. No. 11,909,628, the entire contents of which are incorporated by reference.

TECHNICAL FIELD

[0002] At least one embodiment pertains to processing resources used to perform and facilitate network communications. For example, at least one embodiment pertains to remote direct memory access technology, and more specifically, to spreading packets of a single transport flow across multiple network paths to a same destination.

BACKGROUND

[0003] Remote direct memory access (RDMA) technology enables network adapters to transfer data over a network directly to (or from) memory of a remote device without storing data in data buffers of the operating system of the remote device. Advantages of RDMA include reduced computations and caching by processing devices, e.g., central processing units (CPUs), elimination of the need to copy the data between various network layers, convenient discretization of transmitted data, and so on. RDMA transactions are supported by a number of communication protocols, including RDMA over Converged Ethernet (ROCE), which facilitates RDMA operations using conventional standard Ethernet infrastructure, Internet Wide Area RDMA Protocol (iWARP), which facilitates RDMA operations using Transmission Control Protocol (TCP), and InfiniBand™, which provides native support for RDMA operations. RDMA transactions are especially useful in cloud computing applications and numerous applications that require high data transmission rates and low latency.

[0004] Currently, in ROCE, there is an assumption that all packets for a specific transport flow will travel across the same network path. There is a mechanism used in Ethernet networks, called Equal-Cost multipath (ECMP), which is meant to spread different transport flows evenly across multiple network paths, but this mechanism does not work well in many cases. There is no mechanism in ROCE to spread a specific transport flow evenly across multiple network paths.

Description

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

[0005] To easily identify the discussion of any particular element or act, the most significant digit or digits in a reference number refer to the figure number in which that element is first introduced.

[0006] FIG. 1 is a block diagram of an example network architecture capable of spreading a single transport flow across multiple network paths, according to at least one embodiment.

[0007] FIG. 2 is a diagram illustrating a scheduling flow that can send bursts on any of the routes through a network to a given destination, according to at least one embodiment.

[0008] FIG. 3 is a flow diagram of a method for updating route weights according to at least one embodiment.

[0009] FIG. 4 is a flow diagram of a method for updating route weights according to at least one embodiment.

[0010] FIG. 5A illustrates an example communication system with multipath context, in accordance with at least some embodiments.

[0011] FIG. 5B illustrates a block diagram of an example communication system employing a transmitter with multipath context, according to at least one embodiment.

[0012] FIG. 6 illustrates an example computer system, including a network controller with a multipath context, in accordance with at least some embodiments.

DETAILED DESCRIPTION

[0013] Technologies for spreading a single transport flow across multiple network paths in RoCE and InfiniBand are described. As described above, there is no mechanism in ROCE to spread a transport flow across multiple network paths because there is an assumption that all packets for a specific transport flow will travel across the same network path. Today with static routing, there are cases where multiple flows can be set to the same link in the network, causing congestion even when there are other routes that are not being utilized. In certain cases, it is possible to solve this problem with an additional software layer. An additional software layer requires changes to the software application programming interface (API), and it can add significant overhead in resources and computations. Without hardware support for every queue pair (QP) that the software opens, the software would have to open as many QPs as the desired paths in the network.

[0014] Aspects and embodiments of the present disclosure address these and other challenges by providing mechanisms and methods for spreading a transport flow across multiple paths in the network while maintaining control at an end point. Aspects and embodiments of the present disclosure can improve network utilization by spreading the transport flow across multiple network paths. Aspects and embodiments of the present disclosure can enable software to load different network routing identifiers for a specific transport flow, and the hardware can use these network routing identifiers while sending traffic to send packets across all of the given network paths at a finer granularity. Aspects and embodiments of the present disclosure can enable hardware to send packets with multiple different routing parameters without software intervention in the data path. Aspects and embodiments of the present disclosure can enable spreading traffic for a single transport flow on multiple routes transparently to an application. Aspects and embodiments of the present disclosure can monitor individual routes and identify which routers are more or less congested. Aspects and embodiments of the present disclosure can provide a fast recovery mechanism in the case of a transport error.

[0015] Aspects and embodiments of the present disclosure are relevant for any networks that provide multiple routes between any two end node devices. One example use case includes a network where the end node devices have a higher aggregate bandwidth than individual links in the network. Another use case example includes a network with static routing may have cases of congestion caused by unlucky application interaction. Another use case is where applications are very sensitive to tail latencies caused during an error event.

[0016] Aspects and embodiments of the present disclosure can be used in channel adapters, network adapters, network interface cards (NICs), or the like. A channel adapter (CA), whether a network channel adapter or a host channel adapter, refers to an end node in an InfiniBand Network with features for InfiniBand and RDMA, whereas NIC is similar but for an Ethernet network. Network interface controller, also known as a NIC, network adapter, local area network (LAN) adapter, or physical network interface, refers to a computer hardware component that connects a computer to a computer network. The network interface controller can provide interfaces to a host processor, multiple receive and transmit queues for multiple logical interfaces, and traffic processing. The network interface controller can be both a physical layer and data link layer device, as it provides physical access to a networking medium and a low-level addressing system through the use of media access control (MAC) addresses that are uniquely assigned to network interfaces. The technologies described herein can be implemented in these various types of devices and are referred to herein as “network interface controller” or “network controller.” That is, the network

interface controller can be a channel adapter, a NIC, a network adapter, or the like. The network interface controller can be implemented in a personal computer (PC), a set-top box (STB), a server, a network router, a switch, a bridge, a data processing unit (DPU), a network card, or any device capable of sending packets over multiple network paths to another device.

[0017] FIG. 1 is a block diagram of an example network architecture **100** capable of spreading a single transport flow across multiple network paths, according to at least one embodiment. As depicted in FIG. 1, network architecture **100** can support operations of a requestor device **102** connected over local bus **156** to a first network controller **118** (a requestor network controller). The first network controller **118** can be connected, via a network **124**, to a second network controller **126** (a target network controller) that supports operations of a target device **132**. Network **124** can be a public network (e.g., the Internet), a private network (e.g., a local area network (LAN), or wide area network (WAN)), a wireless network, a personal area network (PAN), or a combination thereof. RDMA operations can support the transfer of data from a requestor memory **114** directly to (or from) a target memory **144** without software mediation by the target device **132**.

[0018] Requestor device **102** can support one or more applications (not explicitly shown in FIG. 1) that can manage various processes **104** that control communication of data with various targets, including target memory **144**. To facilitate memory transfers, processes **104** can post work requests (WRs) to a send queue (SQ) **106** and to a receive queue (RQ) **108**. SQ **106** can be used to request one-sided READ, WRITE, and ATOMIC operations as well as two-sided SEND operations, while RQ **108** can be used to facilitate two-sided RECEIVE requests. Similar processes **134** can operate on target device **132** that supports its own SQ **136** and RQ **138**. A connection between requestor device **102** and target device **132** bundles SQs and RQs into queue pairs (QPs), e.g., SQ **106** (or RQ **108**) on requestor device **102** is paired with RQ **138** (or SQ **136**) on the target device **132**. More specifically, to initiate a connection between requestor device **101** and target device **132**, the processes **104** and **134** can create and link one or more queue pairs.

[0019] To perform a data transfer, process **104** creates a work queue element (WQE) that specifies parameters such as the RDMA verb (operation) to be used for data communication and also can define various operation parameters, such as a source address **116** in a requestor memory **114** (where the data is currently stored), a destination address **146** in a target memory **144**, and other parameters, as discussed in more detail below. Requestor device **102** can then put the WQE into SQ **106** and send a WR **158** to the first network controller **118**, which can use an RDMA adapter **120** to perform packet processing **122** of the WQE and transmit the data indicated in source address **116** to the second network controller **126** via network **124** using a network request **160**. An RDMA adapter **128** can perform packet processing **130** of the received network request **160** (e.g., by generating a local request **162**) and store the data at a destination address **146** of target memory **144**. Subsequently, target device **132** can signal a completion of the data transfer by placing a completion event into a completion queue (CQ) **110** of requestor device **102**, indicating that the WQE has been processed by the receiving side. Target device **132** can also maintain CQ **140** to receive completion messages from requestor device **102** when data transfers happen in the opposite direction, from the target device **132** to requestor device **102**.

[0020] Operation of requestor device **102** and target device **132** can be supported by respective processors **112** and **142**, which can include one or more processing devices, such as CPUs, graphics processing units (GPUs), application-specific integrated circuits (ASICs), field-programmable gate arrays (FPGAs), or any combination thereof. In some embodiments, any of requestor device **102**, the first network controller **118**, and/or requestor memory **114** can be implemented using an integrated circuit, e.g., a system-on-chip. Similarly, any of target device **132**, the second network controller **126**, and/or target memory **144** can be implemented on a single chip. The requestor device **102** and first network controller **118** can be implemented in a personal computer (PC), a set-top box (STB), a server, a network router, a switch, a bridge, a data processing unit (DPU), a network card, or any device capable of sending packets over multiple network paths to another

device.

[0021] Processors **112** and **142** can execute instructions from one or more software programs that manage multiple processes **104** and **134**, SQs **106** and **136**, RQs **108** and **138**, CQs **110** and **140**, and the like. For example, software program(s) running on requestor device **102** can include host or client processes, a communication stack, and a driver that mediates between requestor device **102** and first network controller **118**. The software program(s) can register direct channels of communication with respective memory devices, e.g., RDMA software programs running on requestor device **102** can register a direct channel **164** of communication between the first network controller **118** and requestor memory **114** (and, similarly, a direct channel **166** of communication between the second network controller **126** and target memory **144**). Registered channels **164** and **166** can then be used to support direct memory accesses to the respective memory devices. In the course of RDMA operations, the software program(s) can post WRs, repeatedly check for completed WRs, balance workloads among the multiple RDMA operations, balance workload between RDMA operations and non-RDMA operations (e.g., computations and memory accesses), and so on. The requestor device **102** and first network controller **118** can be used implemented in a personal computer (PC), a set-top box (STB), a server, a network router, a switch, a bridge, a data processing unit (DPU), a network card, or any device capable of sending packets over multiple network paths to another device.

[0022] RDMA accesses to requestor memory **114** and/or target memory **144** can be performed via network **124**, local bus **156** on the requestor side, and bus **168** on the target side and can be enabled by the Converged Ethernet (RoCE) protocol, iWARP protocol, and/or Infini Band™, TCP, and the like.

[0023] As disclosed in more detail below, RDMA accesses can be facilitated using a multipath context **154** (also referred to as hardware multipath context or a session group context) for spreading a single transport flow over multiple network paths of the network **124**. The multipath context **154** can be stored in requestor memory **114** or in memory, cache, or storage in the first network controller **118**. The multipath context **154** can be a hardware context per session group that would maintain a state per configured route to the destination. Different routes/network paths to the same destination end point are defined as sessions in a session group. For example, three network paths to a destination end point would have three sessions in a session group. There can be some translation between sessions to a certain destination and the parameters that will be set in the wire protocol headers. After a QP sends a burst of data, it may decide, based on certain parameters, that the next burst will be sent in a different route. When the QP is scheduled again to send a burst, the QP can select one of the routes provided in the multipath context **154** (e.g., hardware multipath context) based on their relative weights as described in more detail below. While sending traffic on each route, Round-Trip Time (RTT) can be measured by a congestion control (CC) algorithm, and those measurements can be used to adjust the weights for the different routes to identify which are more or less congested. RTT is the amount of time it takes for a data packet to be sent to a destination plus the time it takes for an acknowledgment of that packet to be received back at the origin. The multipath context **154** can be used to optimally utilize multiple routes to the same destination. In cases where there is limited out-of-order support in the hardware, a fence can be used when changing routes, which adds an overhead that needs to be taken into account. If full packet reordering is available at the end node, then no additional changes are needed. The multipath feature described herein can be set up during session negotiation by a session negotiation mechanism. The multipath feature can be based on ROCE, InfiniBand, or other similar transport technologies.

[0024] The first network controller **118** can spread a transport flow across multiple paths in the network **124** while maintaining control at an end point using the multipath context **154**. The RDMA adapter **120** can improve network utilization by spreading the transport flow across multiple network paths. The first network controller **118** can enable software to load different network

routing identifiers for a specific transport flow, and the hardware can use these network routing identifiers while sending traffic to send packets across all of the given network paths at a finer granularity. Network routing identifiers refers to a value that is part of a packet header field (also referred to as a header field in wire protocol headers). The first network controller **118** can enable hardware to send packets with multiple different routing parameters without software intervention in the data path. The first network controller **118** can enable spreading traffic for a single transport flow on multiple routes transparently to the process **104** (e.g., an application). The first network controller **118** can monitor individual routes and identify which routers are more or less congested. The first network controller **118** can provide a fast recovery mechanism in the case of a transport error. The second network controller **126** can similarly perform similar functions.

[0025] In at least one embodiment, the requestor device **102** and the first network controller **118** are part of a first node device, and the target device **132** and the second network controller **126** are part of a second node device. There can be multiple intervening nodes between the first node device and the second node device. At a minimum, there should be at least two paths between the first node device and the second node device.

[0026] FIG. 2 is a diagram illustrating a scheduling flow **200** that can send bursts on any of the routes through a network to a given destination, according to at least one embodiment. In the scheduling flow **200**, a scheduler **208** can schedule transfers for multiple QPs **202**, **204**, **206**. The scheduler **208** can use the multipath context **154** to spread packets of a single transport flow over multiple network paths **210**, **212**, **214**, of the network **124** to a destination end point **216**. The network paths **210**, **212**, **214** are different paths to the same destination end point **216**. The multipath context **154** can store a state of each network paths **210**, **212**, **214**. The multipath context **154** can store a hardware context for each session group. As illustrated in FIG. 2, there are three network paths **210**, **212**, **214**. Network paths to the same destination end point **216** are three sessions in a single session group. The QPs **202**, **204**, **206** can be multipath QPs that can be assigned to a session group. There will be a hardware multipath context **154** assigned per session group. The multipath context **154** can maintain a weight per session in the session group. The number of sessions per session group is configurable (e.g., 8 sessions per session group would result in at least three bits in the wire protocol headers to identify the respective sessions). For example, the QPs **202**, **204**, **206**, can be part of one session group or separate session groups. In at least one embodiment, the QPs **202**, **204**, **206** can be part of the same session group associated with the multipath context **154**. Other multipath contexts (not illustrated in FIG. 2) can be used for other session groups.

[0027] As described above, the multipath context **154** can be a hardware context per session group that would maintain a state per configured route to the destination end point **216**. For example, during operation, after QP **202** sends a burst of data, the scheduler **208** can decide based on certain parameters that the next burst sent from the QP **202** will be sent in a different route. When QP **202** is scheduled to send its next burst, the scheduler **208** can select one of the routes provided in the multipath context **154** (e.g., hardware multipath context) based on their relative weights. In at least one embodiment, one or more RTT measurements **218** can be fed into the multipath context **154** as weight adjustments **220**. In at least one embodiment, the QP **202** includes a congestion control (CC) algorithm that uses the weight adjustment(s) **220** in the multipath context **154** to select one of the network paths **210**, **212**, **214** that is less congested. The multipath context **154** can be used to optimally utilize the different network paths **210**, **212**, **214** for sending packets of a transport flow to the same destination end point **216**.

[0028] As described above, different routes to the same destination are defined as sessions in a session group. The multipath QPs **202**, **204**, **206** can be assigned to a session group. There will be some translation between sessions to a certain destination and the parameters that will be set in the wire protocol headers. In at least one embodiment, a software process is used to ensure that the multipath context **154** holds the correct sessions that cause switches in the network **124** to route the

packets across the different network paths **210**, **212**, **214**. If there are any changes in switch configurations, the software process can update the multipath context **154**, and the weight adjustments can be reset.

[0029] In at least one embodiment, the first network controller **118** of requestor device **102** assigns a first network routing identifier to one or more packets in a first session of a session group associated with a transport flow directed to the destination end point **216**. The transport flow uses a network protocol that allows RDMA, such as ROCE or InfiniBand. The first network routing identifier corresponds to the first network path **210**. The first network routing identifier in the one or more packets causes these packets to be routed to the destination end point **216** via the first network path **210**. The first network controller **118** assigns a second network routing identifier to one or more packets in a second session of the session group associated with the transport flow directed to the destination end point **216**. The second network routing identifier corresponds to network path **212**. The second network routing identifier in the one or more packets causes these packets to be routed to the destination end point **216** via the second network path **212**. The first network controller **118** assigns a third network routing identifier to one or more packets in a third session of the session group associated with the transport flow directed to the destination end point **216**. The third network routing identifier corresponds to network path **214**. The third network routing identifier in the one or more packets causes these packets to be routed to the destination end point **216** via the third network path **214**. If there are additional network paths between the requestor device **102** and the destination end point **216**, additional network routing identifiers can be used. In at least one embodiment, software or firmware can handle defining the network routing identifiers to the different network paths and the network switch configuration. The network routing identifiers can also be referred to as router identifiers or path identifiers.

[0030] During operation, processing logic associated with the QP **202** can select the first network path **210** to send a first burst of packets, such as one or more packets in the first session, to the destination end point **216**. When the scheduler **208** schedules QP **202** for sending traffic, the first session of one or more packets is sent to the destination end point **216**. As described above, when the one or more packets of the first session are sent across the network **124**, the first network routing identifier causes the one or more packets to be routed to the destination end point **216** via the first network path **210**.

[0031] After scheduling and sending the first session (i.e., a specific network path over which the first burst is sent), the processing logic associated with QP **202** can determine whether to change routes (i.e., a different network path) or not based on one or more parameters. The one or more parameters can include bursts since the last route change, the weight of a current route compared to weights of other routes, a requirement of an input fence, random entropy, or the like. In at least one embodiment, the decision is made at the end of the first burst so that a fence can be added if needed. In some cases, there may be a requirement that does not allow a change in the middle of a message. The processing logic can implement an algorithm to determine when to switch routes. This algorithm may require some flexibility to be used for different use cases. The choice of when to make a route change can be programmable by a management application.

[0032] Assuming the processing logic decides to change routes from the first network path **210**, when the scheduler **208** schedules the QP **202** for sending traffic again, the second session of one or more packets is sent to the destination end point **216**. As described above, when the one or more packets of the second session are sent across the network **124**, the second network routing identifier causes the one or more packets to be routed to the destination end point **216** via the second network path **212**.

[0033] After scheduling and sending the second session (i.e., a different network path over which the next burst is sent), the processing logic associated with QP **202** can determine whether to change routes (i.e., a different network path) or not based on one or more parameters as described above. Assuming the processing logic decides to change routes from the second network path **212**,

when the scheduler **208** schedules the QP **202** for sending traffic again, the third session of one or more packets is sent to the destination end point **216**. As described above, when the one or more packets of the third session are sent across the network **124**, the third network routing identifier causes the one or more packets to be routed to the destination end point **216** via the third network path **214**.

[0034] Using the scheduler **208**, the requestor device **102** sends the one or more packets of the first session to the destination end point **216** via the first network path **210**, the one or more packets of the second session to the destination end point **216** via the second network path **212**, and the one or more packets of the third session to the destination end point **216** via the third network path **214**.

[0035] In at least one embodiment, the scheduler **208** can schedule similar sessions of QP **204** and QP **206** to be sent. The scheduler **208** can alternate between QPs **202**, **204**, and **206** according to a scheduling scheme.

[0036] For new route selection, once the processing logic associated with a QP has decided to change routes upon the next scheduling, a route needs to be chosen. The selection of the new route is made at this later time as the relative weights of the different routes may change in the time it takes for the next scheduling of the QP, allowing the most updated values to be used for new route selection. In at least one embodiment, the new route can be selected by a probabilistic function of the weights of the different routes. This method can avoid the case where all the QPs move to the highest-ranked route, which will then be over-congested until the QPs can move to a new route.

[0037] In at least one embodiment, a packet header field can be used to identify the route. That is, the packet header field can contain the network routing identifier corresponding to the selected network path. In at least one embodiment, the packet header field can identify a session port. Switches need to be compatible in that they can route based on the network routing identifier in the packet header field. In at least one embodiment, the compatibility at the end node is negotiated to ensure there is support for packet reordering of the packets arriving from different network paths. The main assumption for multipath transport is that by changing the session, the requestor device can select different routes through the network to the same destination. When inter-operating with an end node device that does not support packet reordering, the requestor device can ensure that the operations are fenced before a route change. In cases where there is limited out-of-order support in the hardware, a fence can be used when changing routers, which adds an overhead that needs to be taken into account. If full packet reordering is available at the end node, then no additional changes are needed. The multipath feature described herein can be set up during session negotiation by a session negotiation mechanism. The multipath feature can be based on RoCE, or other similar transport technologies.

[0038] As described above, the route weights can be updated to ensure the spreading of packets over multiple routes to the same destination, such as described in more detail below with respect to FIG. 3.

[0039] FIG. 3 is a flow diagram of a method **300** for updating route weights according to at least one embodiment. The method **300** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, the method **300** is performed by the requestor device **102** of FIG. 1. In at least one embodiment, the method **300** is performed by the target device **132** of FIG. 1. In at least one embodiment, the first network controller **118** of FIG. 1 and FIG. 2 performs the method **300**. In another embodiment, the method **300** is performed by the second network controller **126** of FIG. 1. In one embodiment, the method **300** can be programmable by users.

[0040] Referring to FIG. 3, the method **300** begins with the processing logic determining a new RTT measurement for a network routing identifier (routeID) and destination address (e.g., destination Internet Protocol (IP) address) (block **302**). For route weight updating, when a burst is sent on a specific route, an RTT measurement packet can be added to constantly check the RTT for each route. To take into account the possibility of last hop congestion, which will affect all the

routes equally, the weights can be based on the difference from the average RTT for the destination. The processing logic updates an average RTT for the destination address (block **304**). The processing logic determines whether the new RTT measurement is less than the average RTT (block **306**). If the new RTT measurement is less than the average RTT for the destination, the processing logic increases a weight value for the network routing identifier (destination address, routeID) (block **308**). If the new RTT measurement is not less than the average RTT for the destination, the processing logic reduces the weight value for the network routing identifier (destination address, routeID) (block **310**).

[0041] In another embodiment, during QP connection, the multipath context can be initiated, and the switches can be properly configured to multipath. The multipath can be configurable on a per QP basis. The multipath context allows limited software intervention in the use of multipath, so on the data path itself, there should be no changes.

[0042] In at least one embodiment, the hardware multipath context can be controlled by a management process that has control over switch routing. The hardware multipath context can be unavailable to untrusted users. In another embodiment, the changing of the multipath parameters should be determined by a management interface per network. In some cases, there can be hardware handling, firmware handling, software handling, or any combination thereof. For example, if a route becomes unusable, this case will be identified by path measurements, which will inform the firmware handling to remove the entry from the multipath context.

[0043] FIG. **4** is a flow diagram of a method **400** for updating route weights according to at least one embodiment. The method **400** can be performed by processing logic comprising hardware, software, firmware, or any combination thereof. In at least one embodiment, the method **400** is performed by the requestor device **102** of FIG. **1**. In at least one embodiment, the method **400** is performed by the target device **132** of FIG. **1**. In at least one embodiment, the first network controller **118** of FIG. **1** and FIG. **2** performs the method **400**. In another embodiment, the method **400** is performed by the second network controller **126** of FIG. **1**. In another embodiment, the method **400** is performed by a network interface controller of a first node.

[0044] Referring to FIG. **4**, the method **400** begins with the processing logic receiving a first packet and a second packet of a transport flow directed to a second node, the transport flow using a network protocol that allows RDMA (block **402**). At block **404**, the processing logic assigns a first network routing identifier to the first packet and a second network routing identifier to the second packet. The first network routing identifier corresponds to a first network path between the first node and the second node. The second network routing identifier corresponds to a second network path between the first node and the second node. The second network path is different than the first network path. At block **406**, the processing logic sends the first packet to the second node via the first network path. At block **408**, the processing logic sends the second packet to the second node via the second network path.

[0045] In one embodiment, the first packet is part of a first burst of packets assigned to the first network routing identifier, and the second packet is part of a second burst of packets assigned to the second network routing identifier.

[0046] In at least one embodiment, the processing logic determines a congestion metric for each of the multiple network paths between the first node and the second node. The processing logic identifies the first network path for the first packet and the second network path for the second packet using the congestion metrics. An aggregate bandwidth of the first network path and the second network path collectively is greater than an individual bandwidth of the first network path.

[0047] In at least one embodiment, the processing logic assigns a first QP and a second QP to a hardware multipath context. The hardware multipath context stores a first state associated with the first network routing identifier and a second state associated with the second network routing identifier. The hardware multipath context is associated with a session group. The first network routing identifier can be used for a first session of the session group. The second network routing

identifier can be used for a second session of the session group. In another embodiment, the processing logic assigns a first QP to a hardware multipath context, and the first session and the second session both use the same QP.

[0048] In at least one embodiment, the processing logic receives, from the first QP, a first burst of data in a first session. The processing logic schedules, using the hardware multipath context, the first burst of data to be sent on the first network path. The processing logic receives, from the first QP or the second QP, a second burst of data in a second session. The processing logic schedules, using the hardware multipath context, the second burst of data to be sent on the second network path.

[0049] In at least one embodiment, the processing logic determines that a subsequent burst of data is to be sent on a different network path than the first network path based on one or more parameters. The hardware multipath context can store a first weight value associated with the first network routing identifier and a second weight value associated with the second network routing identifier. The processing logic selects, based on at least the second weight value, the second network routing identifier for the second burst of data, responsive to determining that the subsequent burst of data is to be sent on the different network paths.

[0050] In at least one embodiment, the processing logic measures an RTT of the first network path and updates the first weight value based on the RTT. The hardware multipath context is associated with a transport session group. The processing logic can assign a third QP and a fourth QP to a second hardware multipath context associated with a second session group. The processing logic receives, from the third QP, a third burst of data in a third session. The processing logic schedules, using the second hardware multipath context, the third burst of data to be sent on a third network path between the first node and a third node. The third network path is different than the first network path. The processing logic receives from the third QP or the fourth QP, a fourth burst of data in a fourth session. The processing logic schedules, using the second hardware multipath context, the fourth burst of data to be sent on a fourth network path between the first node and the third node. The fourth network path is different than the third network path.

[0051] FIG. 5A illustrates an example communication system **500** with multipath context **154**, in accordance with at least some embodiments. The communication system **500** includes a device **510**, a communication network **508** including a communication channel **506**, and a device **512**. In at least one embodiment, the devices **510** and **512** are integrated circuits of a Personal Computer (PC), a laptop, a tablet, a smartphone, a server, a collection of servers, or the like. In some embodiments, the devices **510** and **512** may correspond to any appropriate type of device that communicates with other devices also connected to a common type of communication network **508**. According to embodiments, the transmitters **502** and **522** of devices **510** or **512** may correspond to transmitters of a GPU, a switch (e.g., a high-speed network switch), a network adapter, a CPU, a data processing unit (DPU), etc.

[0052] Examples of the communication network **508** that may be used to connect the devices **510** and **512** include wires, conductive traces, bumps, terminals, or the like. In one example, the communication network **508** is a network that enables data transmission between the devices **510** and **512** using data signals (e.g., digital, optical, wireless signals), clock signals, or both.

[0053] The device **510** includes a transceiver **514** for sending and receiving signals, for example, data signals. The data signals may be digital or optical signals modulated with data or other suitable signals for carrying data.

[0054] The transceiver **514** may include a digital data source **518**, a transmitter **502**, a receiver **504**, and processing circuitry **520** that controls the transceiver **514**. The digital data source **518** may include suitable hardware and/or software for outputting data in a digital format (e.g., in binary code and/or thermometer code). The digital data output by the digital data source **518** may be retrieved from memory (not illustrated) or generated according to input (e.g., user input). The transceiver **514** can use the multipath context **154** as described above with respect to FIG. 1 to FIG.

4.

[0055] The transceiver **514** includes suitable software and/or hardware for receiving digital data from the digital data source **518** and outputting data signals according to the digital data for transmission over the communication network **508** to a transceiver **516** of device **512**.

[0056] The receiver **504** of device **510** may include suitable hardware and/or software for receiving signals, for example, data signals from the communication network **508**. For example, the receiver **504** may include components for receiving processing signals to extract the data for storing in a memory. In at least one embodiment, the transceiver **516** includes a transmitter **522** and receiver **534**. The transceiver **516** receives an incoming signal and samples the incoming signal to generate samples, such as using an analog-to-digital converter (ADC). The ADC can be controlled by a clock-recovery circuit (or clock recovery block) in a closed-loop tracking scheme. The clock-recovery circuit can include a controlled oscillator, such as a voltage-controlled oscillator (VCO) or a digitally-controlled oscillator (DCO) that controls the sampling of the subsequent data by the ADC.

[0057] The processing circuitry **520** may comprise software, hardware, or a combination thereof. For example, the processing circuitry **520** may include a memory including executable instructions and a processor (e.g., a microprocessor) that executes the instructions on the memory. The memory may correspond to any suitable type of memory device or collection of memory devices configured to store instructions. Non-limiting examples of suitable memory devices that may be used include Flash memory, Random Access Memory (RAM), Read Only Memory (ROM), variants thereof, combinations thereof, or the like. In some embodiments, the memory and processor may be integrated into a common device (e.g., a microprocessor may include integrated memory). Additionally or alternatively, the processing circuitry **520** may comprise hardware, such as an ASIC. Other non-limiting examples of the processing circuitry **520** include an Integrated Circuit (IC) chip, a CPU, a GPU, a DPU, a microprocessor, an FPGA, a collection of logic gates or transistors, resistors, capacitors, inductors, diodes, or the like. Some or all of the processing circuitry **520** may be provided on a Printed Circuit Board (PCB) or collection of PCBs. It should be appreciated that any appropriate type of electrical component or collection of electrical components may be suitable for inclusion in the processing circuitry **520**. The processing circuitry **520** may send and/or receive signals to and/or from other elements of the transceiver **514** to control the overall operation of the transceiver **514**.

[0058] The transceiver **514** or selected elements of the transceiver **514** may take the form of a pluggable card or controller for the device **510**. For example, the transceiver **514** or selected elements of the transceiver **514** may be implemented on a network interface card (NIC).

[0059] The device **512** may include a transceiver **516** for sending and receiving signals, for example, data signals over a channel **506** of the communication network **508**. The same or similar structure of the transceiver **514** may be applied to transceiver **516**, and thus, the structure of transceiver **516** is not described separately.

[0060] Although not explicitly shown, it should be appreciated that devices **510** and **512** and the transceiver **514** and transceiver **516** may include other processing devices, storage devices, and/or communication interfaces generally associated with computing tasks, such as sending and receiving data.

[0061] FIG. 5B illustrates a block diagram of an example communication system **524** employing a transmitter **502** with multipath context **154**, according to at least one embodiment. In the example shown in FIG. 5B, a Pulse Amplitude Modulation level-4 (PAM4) modulation scheme is employed with respect to the transmission of a signal (e.g., digitally encoded data) from a transmitter (TX) **502** to a receiver (RX) **534** via a communication channel **506** (e.g., a transmission medium). In this example, the transmitter **502** receives an input data **526** (i.e., the input data at time n is represented as " $a(n)$ "), which is modulated in accordance with a modulation scheme (e.g., PAM4) and sends the signal **528** $a(n)$ including a set of data symbols (e.g., symbols $-3, -1, 1, 3$, where the symbols

represent coded binary data). It is noted that while the use of the PAM4 modulation scheme is described herein by way of example, other data modulation schemes can be used in accordance with embodiments of the present disclosure, including for example, a non-return-to-zero (NRZ) modulation scheme, PAM7, PAM8, PAM16, etc. For example, for an NRZ-based system, the transmitted data symbols consist of symbols -1 and 1 , with each symbol value representing a binary bit. This is also known as a PAM level -2 or PAM2 system as there are 2 unique values of transmitted symbols. Typically, a binary bit 0 is encoded as -1 , and a bit 1 is encoded as 1 as the PAM2 values.

[0062] In the example shown, the PAM4 modulation scheme uses four (4) unique values of transmitted symbols to achieve higher efficiency and performance. The four levels are denoted by symbol values -3 , -1 , 1 , 3 , with each symbol representing a corresponding unique combination of binary bits (e.g., 00, 01, 10, 11).

[0063] The communication channel **506** is a destructive medium in that the channel acts as a low pass filter that attenuates higher frequencies more than it attenuates lower frequencies, introduces inter-symbol interference (ISI) and noise from cross talk, power supplies, Electromagnetic Interference (EMI), or other sources. The communication channel **506** can be over serial links (e.g., a cable, PCB traces, copper cables, optical fibers, or the like), read channels for data storage (e.g., hard disk, flash solid-state drives (SSDs), high-speed serial links, deep space satellite communication channels, applications, or the like. The receiver (RX) **534** receives an incoming signal **530** over the channel **506**. The receiver **534** can output a received signal **532**, “ $v(n)$,” including the set of data symbols (e.g., symbols -3 , -1 , 1 , 3 , wherein the symbols represent coded binary data).

[0064] In at least one embodiment, the transmitter **502** can be part of a SerDes IC. The SerDes IC can be a transceiver that converts parallel data to serial data and vice versa. The SerDes IC can facilitate transmission between two devices over serial streams, reducing the number of data paths, wires/traces, terminals, etc. The receiver **534** can be part of a SerDes IC. The SerDes IC can include a clock-recovery circuit. The clock-recovery circuit can be coupled to an ADC and an equalization block.

[0065] FIG. **6** illustrates an example computer system **600**, including a network controller **644** with a multipath context **154**, in accordance with at least some embodiments. The multipath context **154** is used to send packets of a single transport flow over multiple network paths to the same destination, as described herein. In at least one embodiment, computer system **600** may be a system with interconnected devices and components, a System on Chip (SoC), or some combination. In at least one embodiment, computer system **600** is formed with a processor **602** that may include execution units to execute an instruction. In at least one embodiment, computer system **600** may include, without limitation, a component, such as a processor **602**, to employ execution units including logic to perform algorithms for processing data. In at least one embodiment, computer system **600** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **600** may execute a version of WINDOWS' operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux, for example), embedded software, and/or graphical user interfaces, may also be used.

[0066] In at least one embodiment, computer system **600** may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants (“PDAs”), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (DSP), an SoC, network computers (“NetPCs”), set-top boxes, network

hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions. In an embodiment, computer system **600** may be used in devices such as graphics processing units (GPUs), network adapters, central processing units, and network devices such as switches (e.g., a high-speed direct GPU-to-GPU interconnect such as the NVIDIA GH100 NVLINK or the NVIDIA Quantum 2 64 Ports InfiniBand NDR Switch).

[0067] In at least one embodiment, computer system **600** may include, without limitation, processor **602** that may include, without limitation, one or more execution units **807** that may be configured to execute a Compute Unified Device Architecture (“CUDA”) (CUDA® is developed by NVIDIA Corporation of Santa Clara, California) program. In at least one embodiment, a CUDA program is at least a portion of a software application written in a CUDA programming language. In at least one embodiment, computer system **600** is a single processor desktop or server system. In at least one embodiment, computer system **600** may be a multiprocessor system. In at least one embodiment, processor **602** may include, without limitation, a complex instruction set computer (CISC) microprocessor, a reduced instruction set computer (RISC) microprocessor, a Very Long Instruction Word (VLIW) microprocessor, and a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **602** may be coupled to a processor bus **604** that may transmit data signals between processor **602** and other components in computer system **600**.

[0068] In at least one embodiment, processor **602** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **606**. In at least one embodiment, processor **602** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside externally to processor **602**. In at least one embodiment, processor **602** may also include a combination of both internal and external caches. In at least one embodiment, a register file **608** may store different types of data in various registers, including integer registers, floating point registers, status registers, instruction pointer registers, or the like.

[0069] In at least one embodiment, execution unit **610**, including, without limitation, logic to perform integer and floating-point operations, also resides in processor **602**. Processor **602** may also include a microcode (“ucode”) read-only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **610** may include logic to handle a packed instruction set **612**. In at least one embodiment, by including packed instruction set **612** in an instruction set of a general-purpose processor **602**, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in a general-purpose processor **602**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using the full width of a processor's data bus for performing operations on packed data, which may eliminate the need to transfer smaller units of data across a processor's data bus to perform one or more operations one data element at a time.

[0070] In at least one embodiment, execution unit **610** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **600** may include, without limitation, a memory **614**. In at least one embodiment, memory **614** may be implemented as a Dynamic Random Access Memory (DRAM) device, a Static Random Access Memory (SRAM) device, a flash memory device, or other memory devices. Memory **614** may store instruction(s) **616** and/or data **618** represented by data signals that may be executed by processor **602**.

[0071] In at least one embodiment, a system logic chip may be coupled to a processor bus **604** and memory **614**. In at least one embodiment, the system logic chip may include, without limitation, a memory controller hub (“MCH”) **620**, and processor **602** may communicate with MCH **620** via processor bus **604**. In at least one embodiment, MCH **620** may provide a high bandwidth memory path to memory **614** for instruction and data storage and for storage of graphics commands, data, and textures. In at least one embodiment, MCH **620** may direct data signals between processor **602**, memory **614**, and other components in computer system **600** and may bridge data signals between

processor bus **604**, memory **614**, and a system I/O **622**. In at least one embodiment, a system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **620** may be coupled to memory **614** through high bandwidth memory path, and graphics/video card **626** may be coupled to MCH **620** through an Accelerated Graphics Port (“AGP”) interconnect **624**.

[0072] In at least one embodiment, computer system **600** may use system I/O **622** that is a proprietary hub interface bus to couple MCH **620** to I/O controller hub (“ICH”) **628**. In at least one embodiment, ICH **628** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, a local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **614**, a chipset, and processor **602**. Examples may include, without limitation, an audio controller **630**, a firmware hub (“flash BIOS”) **632**, a wireless transceiver **634**, a data storage **636**, a legacy I/O controller **638** containing a user input interface **640**, a keyboard interface, a serial expansion port **642**, such as a USB port, and a network controller **644**, including the multipath context **154** as described herein. Data storage **636** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device.

[0073] In at least one embodiment, FIG. **6** illustrates a computer system **60000**, which includes interconnected hardware devices or “chips.” In at least one embodiment, FIG. **6** may illustrate an example SoC. In at least one embodiment, devices illustrated in FIG. **6** may be interconnected with proprietary interconnects, standardized interconnects (e.g., Peripheral Component Interconnect Express (PCIe), or some combination thereof. In at least one embodiment, one or more components of computer system **600** are interconnected using compute express link (“CXL”) interconnects.

[0074] Other variations are within the spirit of the present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood, however, that there is no intention to limit the disclosure to a specific form or forms disclosed, but on the contrary, the intention is to cover all modifications, alternative constructions, and equivalents falling within the spirit and scope of the disclosure, as defined in appended claims.

[0075] Use of terms “a” and “an” and “the” and similar referents in the context of describing disclosed embodiments (especially in the context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. “Connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitations of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within the range, unless otherwise indicated herein, and each separate value is incorporated into the specification as if it were individually recited herein. In at least one embodiment, the use of the term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, the term “subset” of a corresponding set does not necessarily denote a proper subset of the corresponding set, but subset and corresponding set may be equal.

[0076] Conjunctive language, such as phrases of the form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with the context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of the set of A and B and C. For instance, in an illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of the following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that

certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, the term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). In at least one embodiment, the number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, the phrase “based on” means “based at least in part on” and not “based solely on.”

[0077] Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under the control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause a computer system to perform operations described herein. In at least one embodiment, a set of non-transitory computer-readable storage media comprises multiple non-transitory computer-readable storage media and one or more individual non-transitory storage media of multiple non-transitory computer-readable storage media lacks all of the code while multiple non-transitory computer-readable storage media collectively store all of the code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors.

[0078] Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein, and such computer systems are configured with applicable hardware and/or software that enable the performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

[0079] Use of any and all examples or exemplary language (e.g., “such as”) provided herein is intended merely to better illuminate embodiments of the disclosure and does not pose a limitation on the scope of the disclosure unless otherwise claimed. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the disclosure.

[0080] All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to the same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

[0081] In description and claims, the terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may not be intended as synonyms for each other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

[0082] Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulates and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

[0083] In a similar manner, the term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As a non-limiting example, a “processor” may be a network device. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes for continuously or intermittently carrying out instructions in sequence or in parallel. In at least one embodiment, the terms “system” and “method” are used herein interchangeably insofar as the system may embody one or more methods, and methods may be considered a system.

[0084] In the present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, the process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways, such as by receiving data as a parameter of a function call or a call to an application programming interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. In at least one embodiment, references may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, processes of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface, or an inter-process communication mechanism.

[0085] Although descriptions herein set forth example embodiments of described techniques, other architectures may be used to implement described functionality and are intended to be within the scope of this disclosure. Furthermore, although specific distributions of responsibilities may be defined above for purposes of description, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

[0086] Furthermore, although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. (canceled)
2. A network controller comprising: a local interface coupled to a requestor device, wherein the network controller is to: receive, from the requestor device over the local interface, a work request to send data of a transport flow directed to a target device over a network interface; receive, from the requestor device over the local interface, one or more parameters related to a multipath selection by the network controller, the one or more parameters being based on a difference from an average round-trip time (RTT) of available paths to the target device; send a first burst of data of

the transport flow via a first network path to the target device; determine, using the one or more parameters, that a second burst of data of the transport flow is to be sent to the target device on a different network path than the first network path; identify, using the one or more parameters, a second network path to the target device; and send the second burst of data to the target device on the second network path.

3. The network controller of claim 2, wherein the one or more parameters comprises at least one of i) a number of bursts since a last route change, ii) random entropy, iii) a requirement of an input fence, or iv) a comparison of a weight value, corresponding to a current route over the first network path, and other weight values corresponding to other available network paths to the target device.

4. The network controller of claim 2, wherein the network controller is further to: measure a first RTT of the first network path; measure a second RTT of the second network path; determine the average RTT; update a first weight value based on the first RTT and the average RTT; and update a second weight value based on the second RTT and the average RTT.

5. The network controller of claim 2, wherein the network controller is further to: measure an RTT of the first network path; determine whether the RTT is less than the average RTT of available network paths to the target device; and increase a first weight value responsive to the RTT being less than the average RTT; or decrease the first weight value responsive to the RTT being greater than the average RTT.

6. The network controller of claim 2, wherein the network controller is further to: measure an RTT of the second network path; determine whether the RTT is less than the average RTT of available network paths to the target device; and increase a second weight value responsive to the RTT being less than the average RTT; or decrease the second weight value responsive to the RTT being greater than the average RTT.

7. The network controller of claim 2, wherein, to determine that the second burst of data of the transport flow is to be sent on a different network path than the first network path, the network controller is further to determine that a first weight value is greater than a weight value of at least one available network path to the target device.

8. The network controller of claim 2, wherein the transport flow uses Remote Direct Memory Access (RDMA) over Converged Ethernet (ROCE).

9. The network controller of claim 2, wherein the transport flow uses InfiniBand.

10. The network controller of claim 2, wherein the network controller is further to: determine a first congestion metric for the first network path; determine a second congestion metric for the second network path; update a first weight value based on the first congestion metric; and update a second weight value based on the second congestion metric.

11. The network controller of claim 2, wherein the network controller is further to: assign a first queue pair (QP) and a second QP to a multipath context, wherein the first burst of data and the second burst of data are stored in the first QP; schedule, by a scheduler of the network controller, the second burst of data from the second QP; schedule, by the scheduler, a third burst of data from the second QP; assign, by the network controller using the multipath context, a third network routing identifier to the third burst of data after the scheduling the third burst of data, the third network routing identifier corresponds to a third network path between the requester device and the target device; and send the third burst of data to the target device via the third network path.

12. A computing system comprising: a requestor device; and a network controller coupled to the requestor device over an interface, wherein the network controller is to: receive, from the requestor device over the interface, a work request to send data of a transport flow directed to a target device over a network; receive, from the requestor device over the interface, one or more parameters related to a multipath selection by the network controller, the one or more parameters being based on a difference from an average round-trip time (RTT) of available paths to the target device; send a first burst of data of the transport flow via a first network path to the target device; determine, using the one or more parameters, that a second burst of data of the transport flow is to be sent to

the target device on a different network path than the first network path; identify, using the one or more parameters, a second network path to the target device; and send the second burst of data to the target device on the second network path.

13. The computing system of claim 12, wherein the transport flow uses Remote Direct Memory Access (RDMA) over Converged Ethernet (ROCE).

14. The computing system of claim 12, wherein the transport flow uses InfiniBand.

15. The computing system of claim 12, wherein to determine that the second burst of data of the transport flow is to be sent on a different network path than the first network path, the network controller is to determine that a first weight value is greater than a weight value of at least one available network path to the target device.

16. The computing system of claim 12, wherein the one or more parameters comprises at least one of i) a number of bursts since a last route change, ii) random entropy, iii) a requirement of an input fence, or iv) a comparison of a weight value, corresponding to a current route over the first network path, and other weight values corresponding to other available network paths to the target device.

17. The computing system of claim 12, wherein the network controller is further to: measure a first RTT of the first network path; measure a second RTT of the second network path; determine the average RTT; update a first weight value based on the first RTT and the average RTT; and update a second weight value based on the second RTT and the average RTT.

18. A network controller comprising: a network adapter coupled to a requestor device over a local interface, wherein the network adapter is to: receive, from the requestor device over the local interface, a request to send data of a transport flow directed to a target device over a network interface; receive, from the requestor device over the local interface, one or more parameters related to a multipath selection by the network controller, the one or more parameters being based on a difference from an average round-trip time (RTT) of available paths to the target device; send a first burst of data of the transport flow via a first network path to the target device; identify, using the one or more parameters, a second network path to the target device; and send a second burst of data to the target device on the second network path.

19. The network controller of claim 18, wherein the one or more parameters comprises at least one of i) a number of bursts since a last route change, ii) random entropy, iii) a requirement of an input fence, or iv) a comparison of a first weight value, corresponding to a current route over the first network path, and other weight values corresponding to other available network paths to the target device.

20. The network controller of claim 18, wherein the transport flow uses Remote Direct Memory Access (RDMA) over Converged Ethernet (RoCE).

21. The network controller of claim 18, wherein the transport flow uses InfiniBand.
