US 20250265276A1

(19) United States
(12) Patent Application Publication
Yadav
(10) Pub. No.: US 2025/0265276 A1
(43) Pub. Date: Aug. 21, 2025

(54) **AUTOMATIC TARGET VISUALIZATION ADAPTATION**

(71) Applicant: **ThoughtSpot, Inc.**, Mountain View, CA (US)

(72) Inventor: **Nikhil Yadav**, Sunnyvale, CA (US)

(21) Appl. No.: **18/442,193**

(22) Filed: **Feb. 15, 2024**

Publication Classification

(51) **Int. Cl.**
G06F 16/28 (2019.01)
G06F 16/2452 (2019.01)
G06F 16/2457 (2019.01)

(52) **U.S. Cl.**
CPC ***G06F 16/287*** (2019.01); ***G06F 16/24522*** (2019.01); ***G06F 16/24573*** (2019.01)

(57) **ABSTRACT**

Automatic target visualization adaptation includes obtaining a natural language string expressing a request for data, obtaining large language model generated structured query language data expressing the natural language string, obtaining, by transforming the large language model generated structured query language data, first resolved request data, obtaining target-parameters data for a target visualization type, obtaining target-parameters deficit data, obtaining a co-occurring parameter, obtaining, by transforming a combination of the first resolved request data and the co-occurring parameter, an automatically generated data query, obtaining from a database, results data generated by execution of the automatically generated data query by the database, and outputting data for presenting the results data in accordance with the target visualization type.

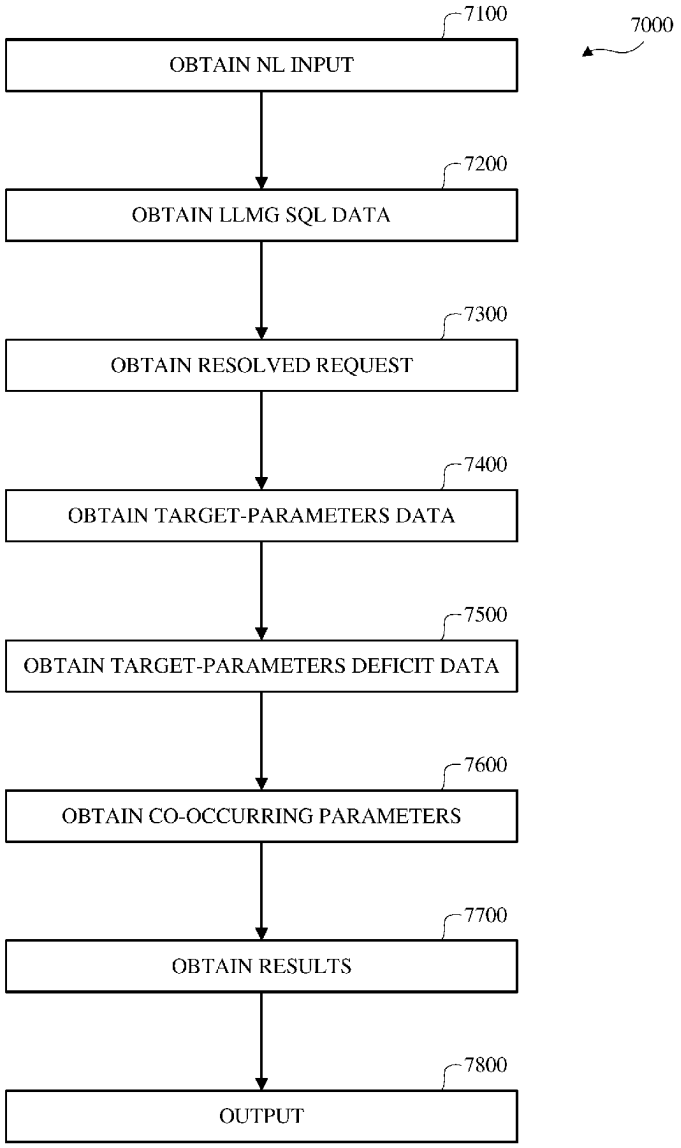

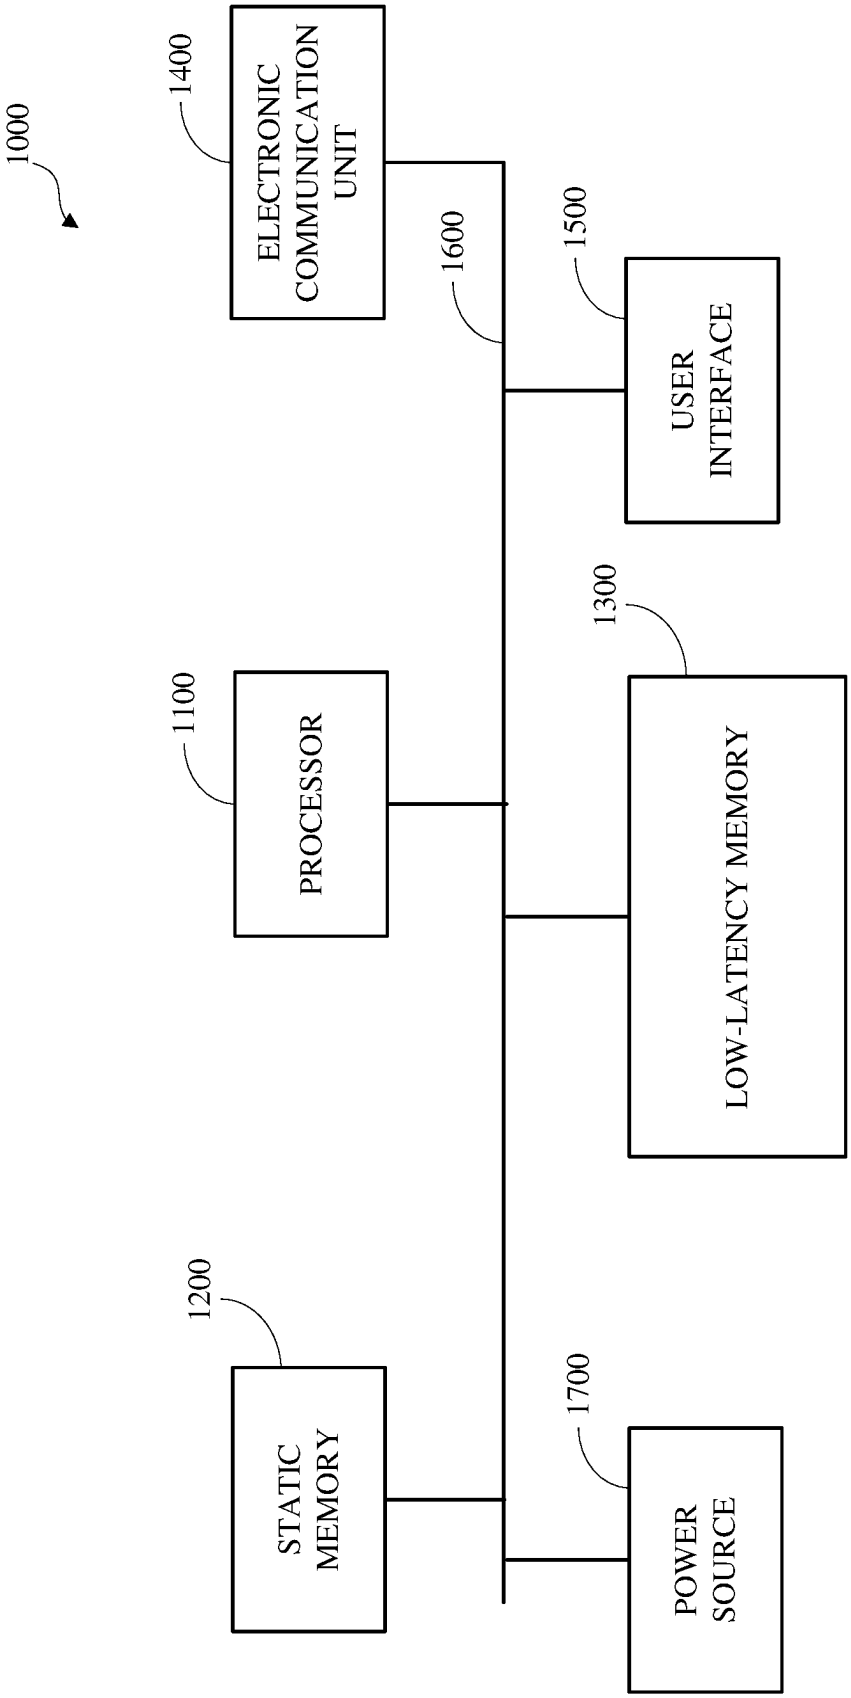
1000
1400
ELECTRONIC COMMUNICATION UNIT
1600
1500
USER INTERFACE
1300
1100
PROCESSOR
LOW-LATENCY MEMORY
1200
1700
STATIC MEMORY
POWER SOURCE

FIG. 1

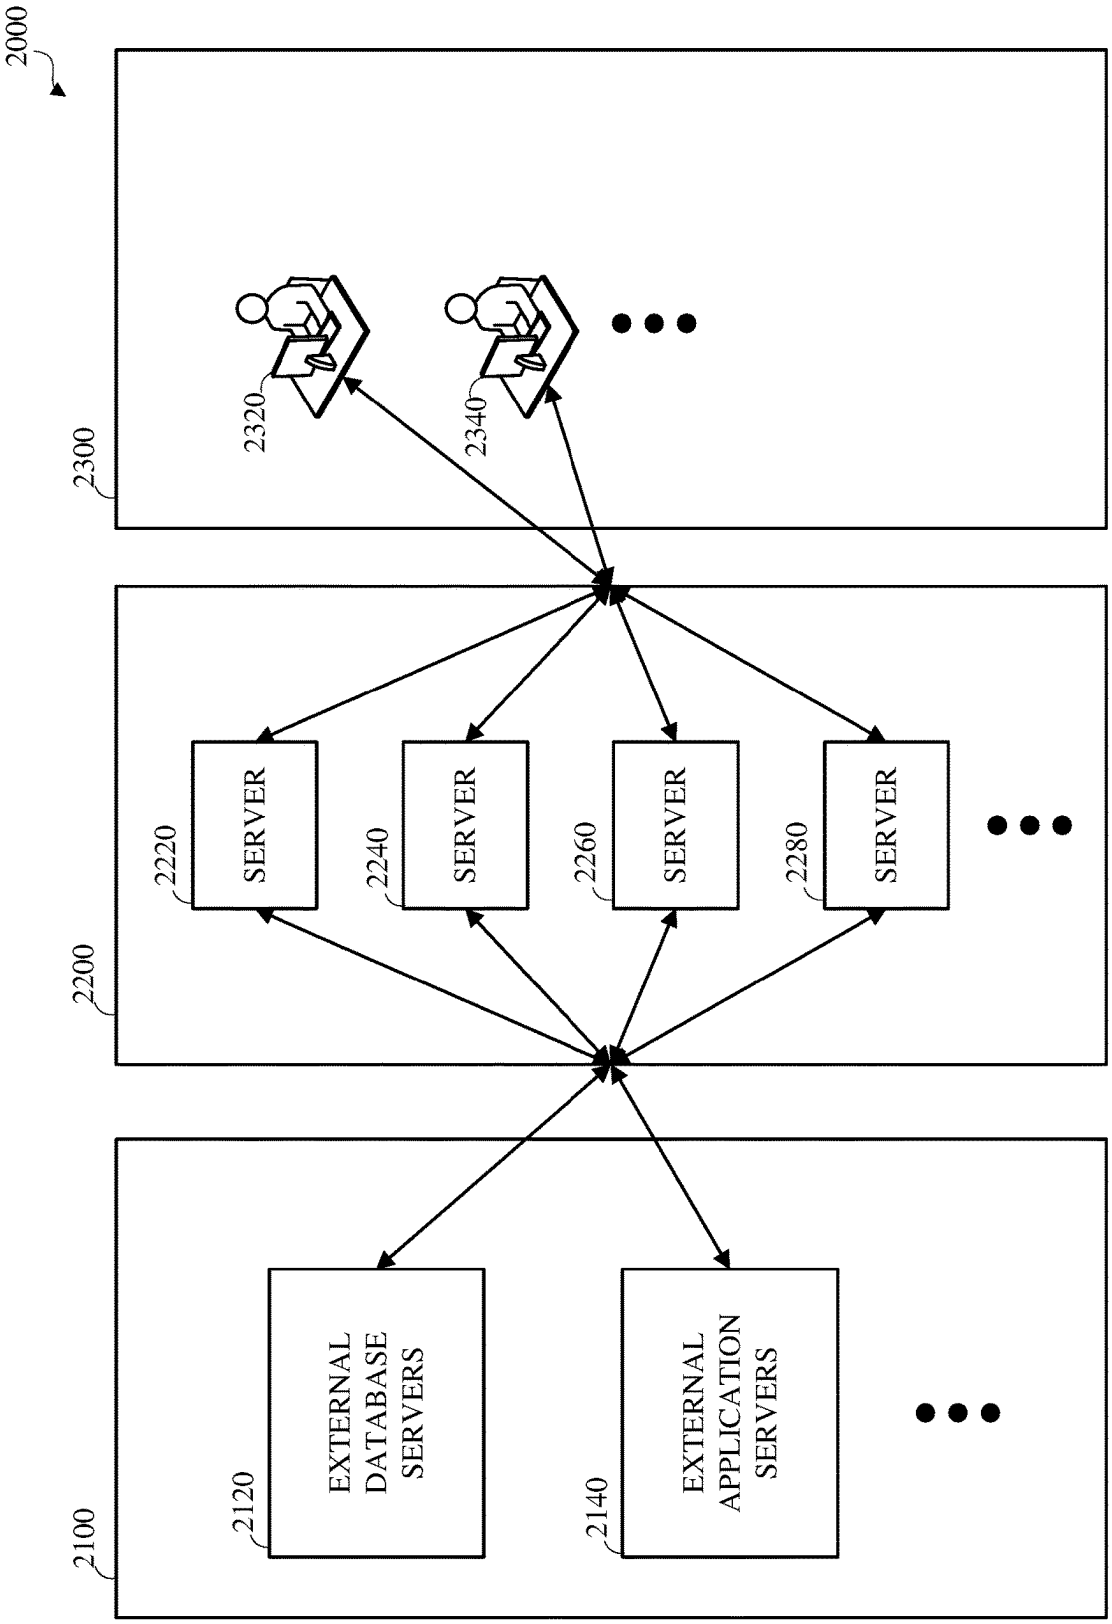
2000
2300
2320
2340
2200
2220
SERVER
2240
SERVER
2260
SERVER
2280
SERVER
2100
2120
EXTERNAL DATABASE SERVERS
2140
EXTERNAL APPLICATION SERVERS

FIG. 2

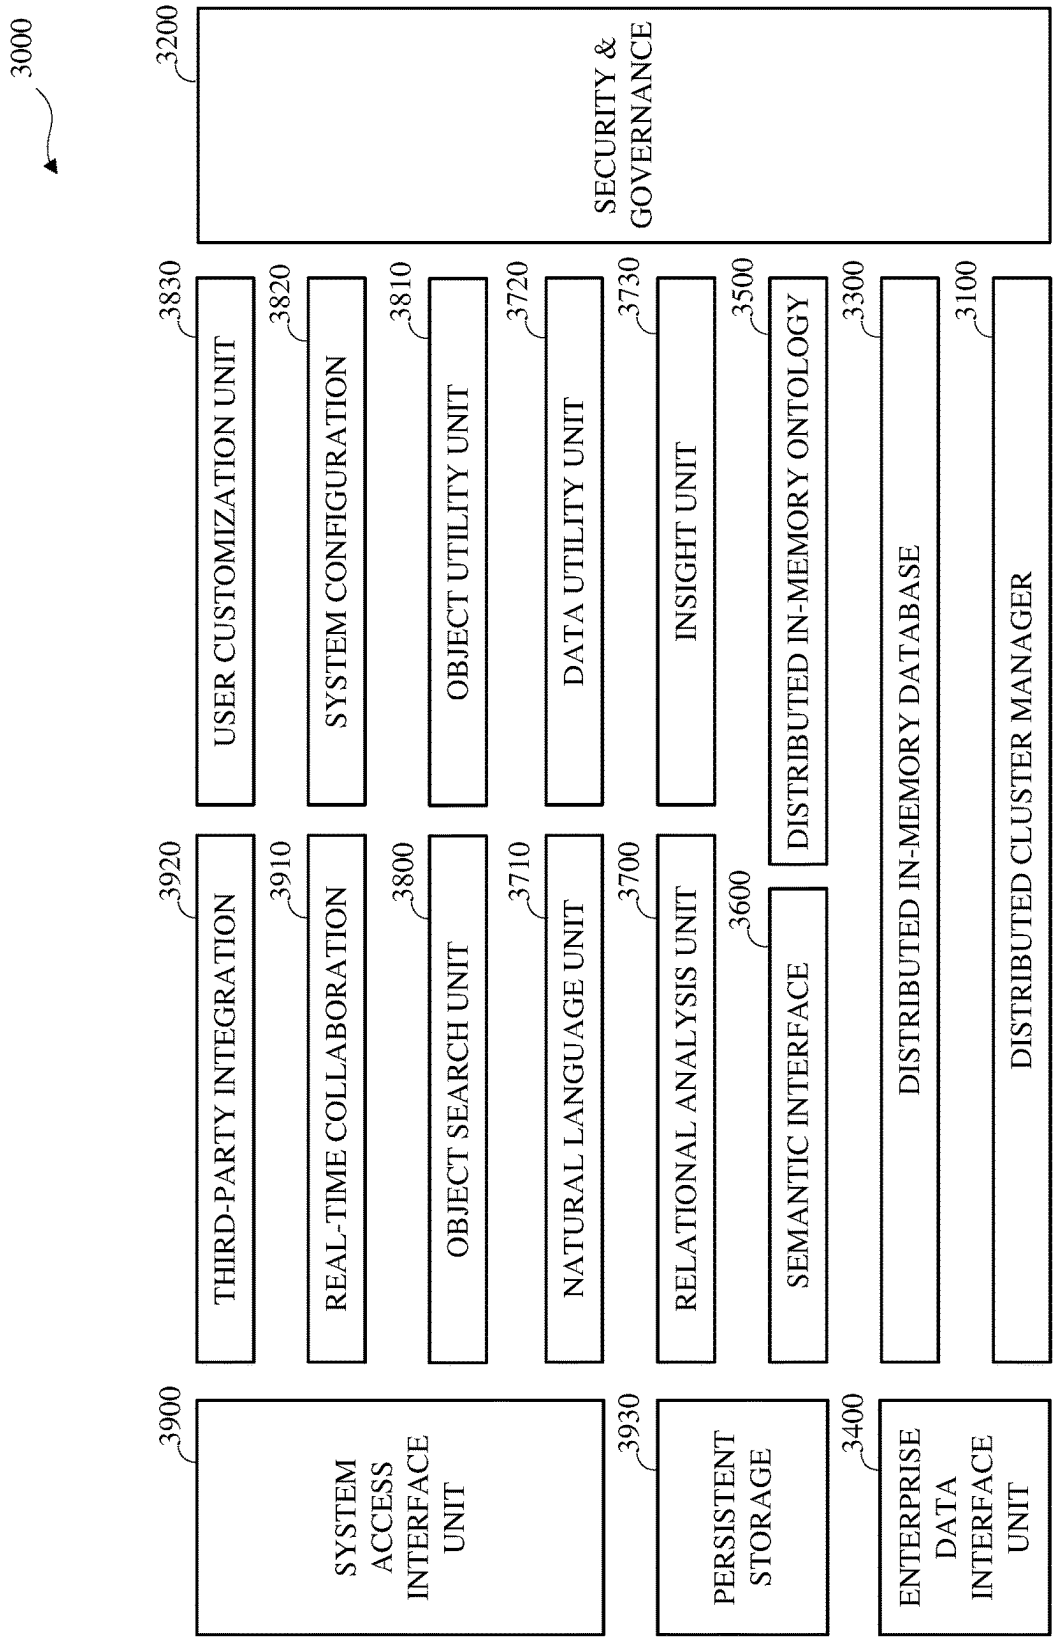
3000
3200
SECURITY & GOVERNANCE
3830
USER CUSTOMIZATION UNIT
3820
SYSTEM CONFIGURATION
3810
OBJECT UTILITY UNIT
3720
DATA UTILITY UNIT
3730
INSIGHT UNIT
3500
DISTRIBUTED IN-MEMORY ONTOLOGY
3300
DISTRIBUTED IN-MEMORY DATABASE
3100
DISTRIBUTED CLUSTER MANAGER
3920
THIRD-PARTY INTEGRATION
3910
REAL-TIME COLLABORATION
3800
OBJECT SEARCH UNIT
3710
NATURAL LANGUAGE UNIT
3700
RELATIONAL ANALYSIS UNIT
3600
SEMANTIC INTERFACE
3900
SYSTEM ACCESS INTERFACE UNIT
3930
PERSISTENT STORAGE
3400
ENTERPRISE DATA INTERFACE UNIT

FIG. 3

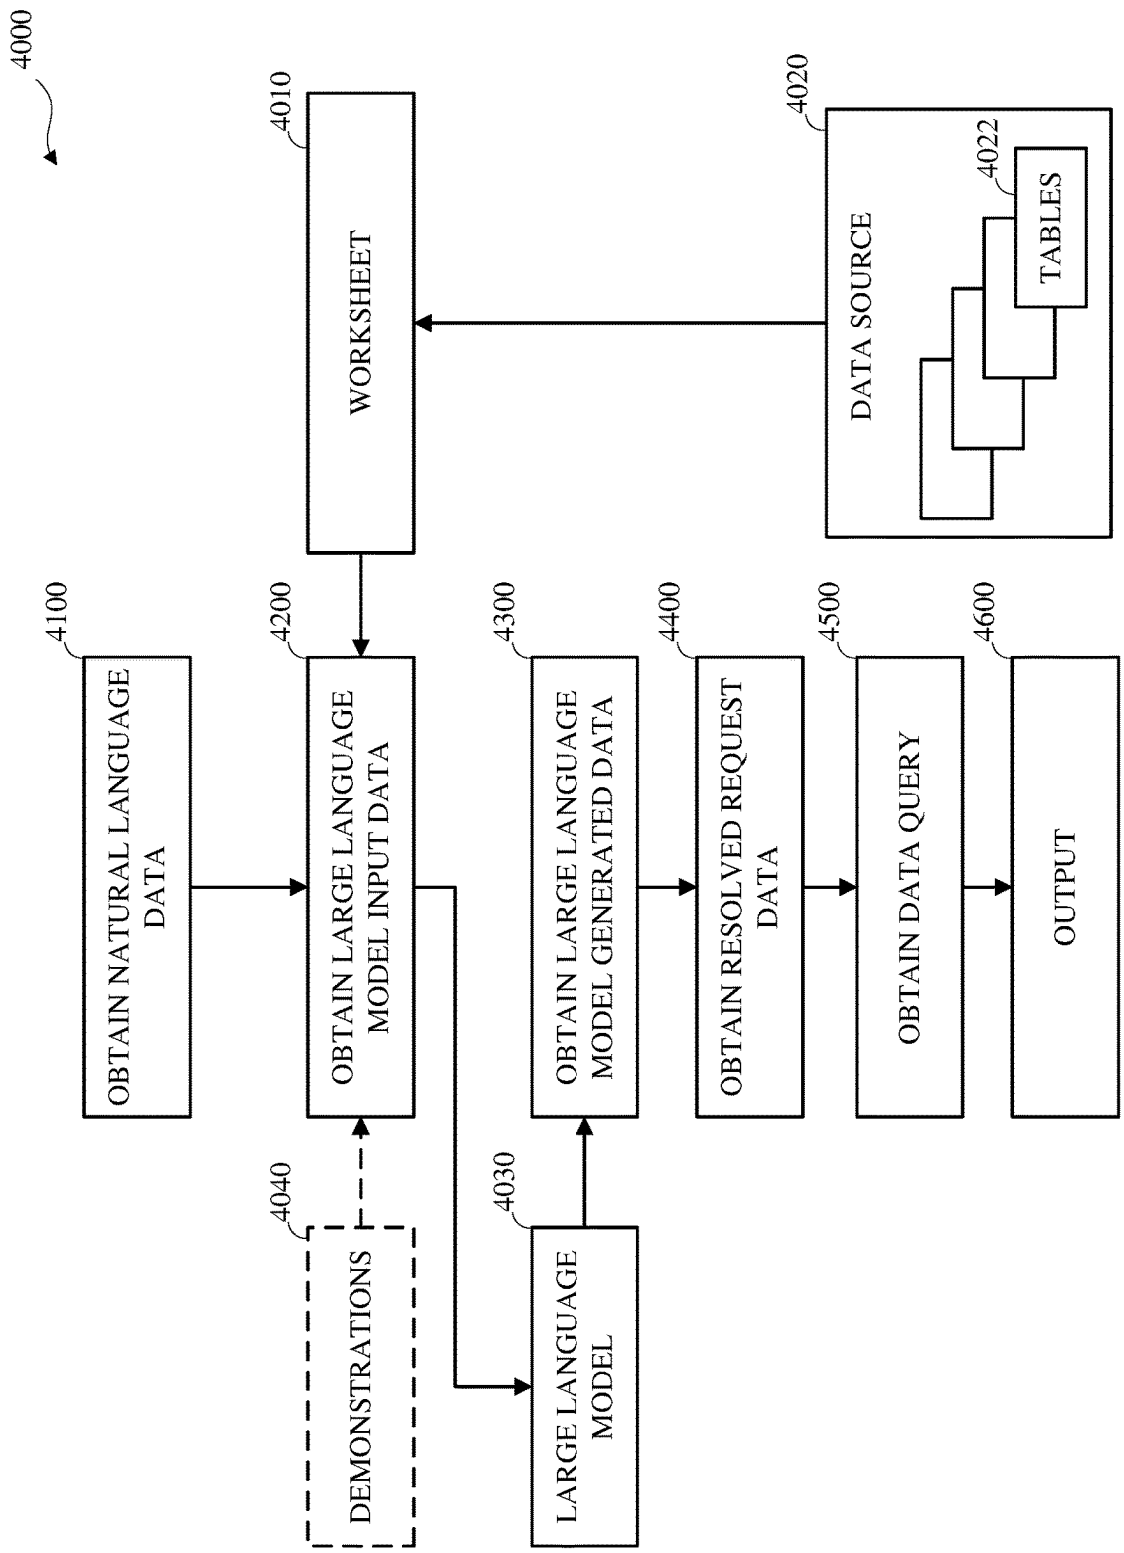
4000
4010
WORKSHEET
4020
DATA SOURCE
4022
TABLES
4100
OBTAIN NATURAL LANGUAGE DATA
4200
OBTAIN LARGE LANGUAGE MODEL INPUT DATA
4300
OBTAIN LARGE LANGUAGE MODEL GENERATED DATA
4400
OBTAIN RESOLVED REQUEST DATA
4500
OBTAIN DATA QUERY
4600
OUTPUT
4040
DEMONSTRATIONS
4030
LARGE LANGUAGE MODEL

FIG. 4

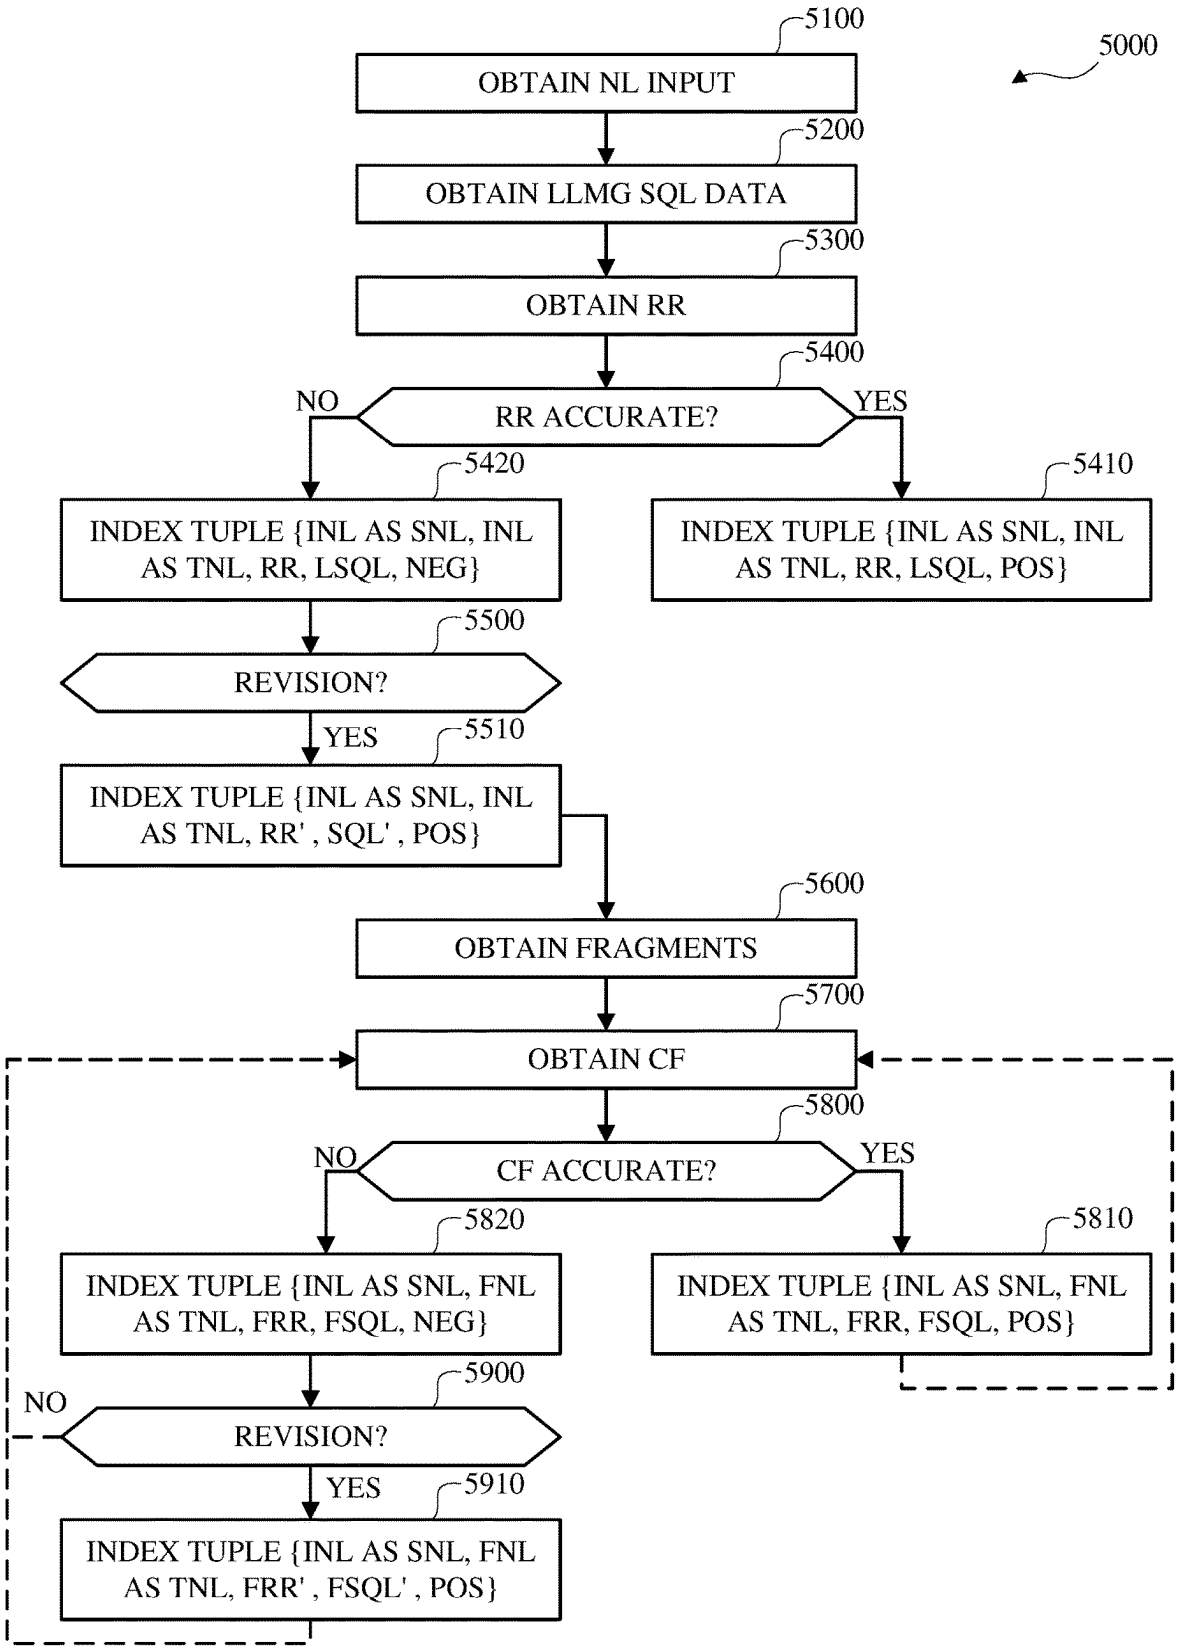
5000
5100
OBTAIN NL INPUT
5200
OBTAIN LLMG SQL DATA
5300
OBTAIN RR
5400
RR ACCURATE?
NO
YES
5420
INDEX TUPLE {INL AS SNL, INL AS TNL, RR, LSQL, NEG}
5410
INDEX TUPLE {INL AS SNL, INL AS TNL, RR, LSQL, POS}
5500
REVISION?
YES
5510
INDEX TUPLE {INL AS SNL, INL AS TNL, RR' , SQL' , POS}
5600
OBTAIN FRAGMENTS
5700
OBTAIN CF
5800
CF ACCURATE?
NO
YES
5820
INDEX TUPLE {INL AS SNL, FNL AS TNL, FRR, FSQL, NEG}
5810
INDEX TUPLE {INL AS SNL, FNL AS TNL, FRR, FSQL, POS}
5900
NO
REVISION?
YES
5910
INDEX TUPLE {INL AS SNL, FNL AS TNL, FRR' , FSQL' , POS}

FIG. 5

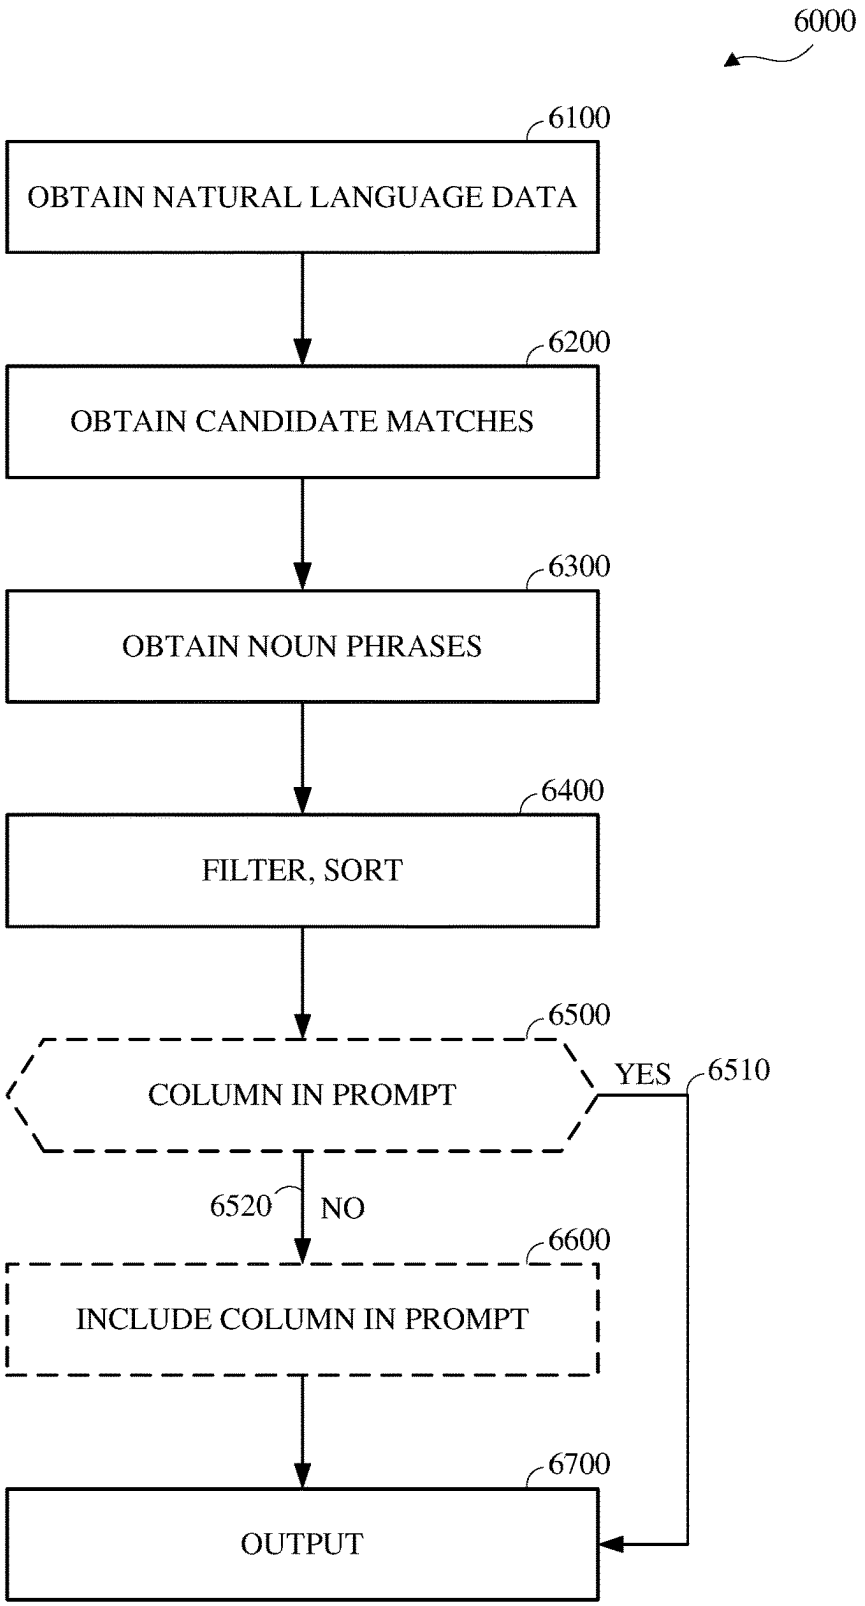
6000
6100
OBTAIN NATURAL LANGUAGE DATA
6200
OBTAIN CANDIDATE MATCHES
6300
OBTAIN NOUN PHRASES
6400
FILTER, SORT
6500
COLUMN IN PROMPT
YES
6510
6520
NO
6600
INCLUDE COLUMN IN PROMPT
6700
OUTPUT

FIG. 6

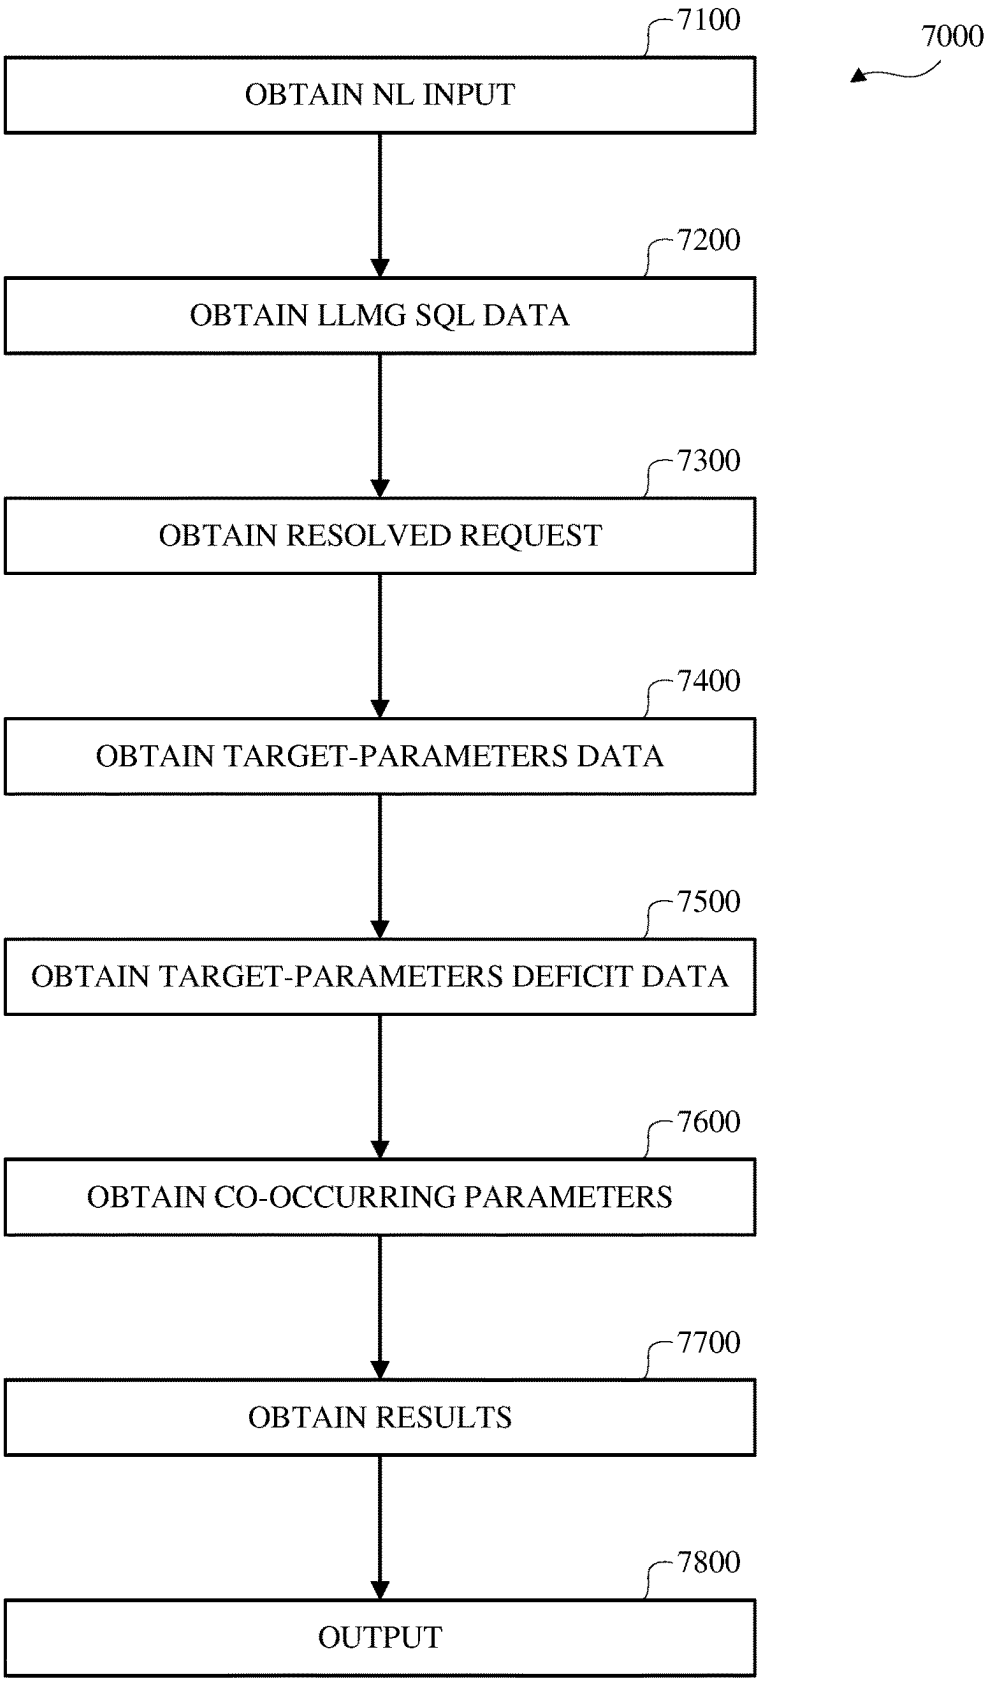
7000
7100
OBTAIN NL INPUT
7200
OBTAIN LLMG SQL DATA
7300
OBTAIN RESOLVED REQUEST
7400
OBTAIN TARGET-PARAMETERS DATA
7500
OBTAIN TARGET-PARAMETERS DEFICIT DATA
7600
OBTAIN CO-OCCURRING PARAMETERS
7700
OBTAIN RESULTS
7800
OUTPUT

FIG. 7

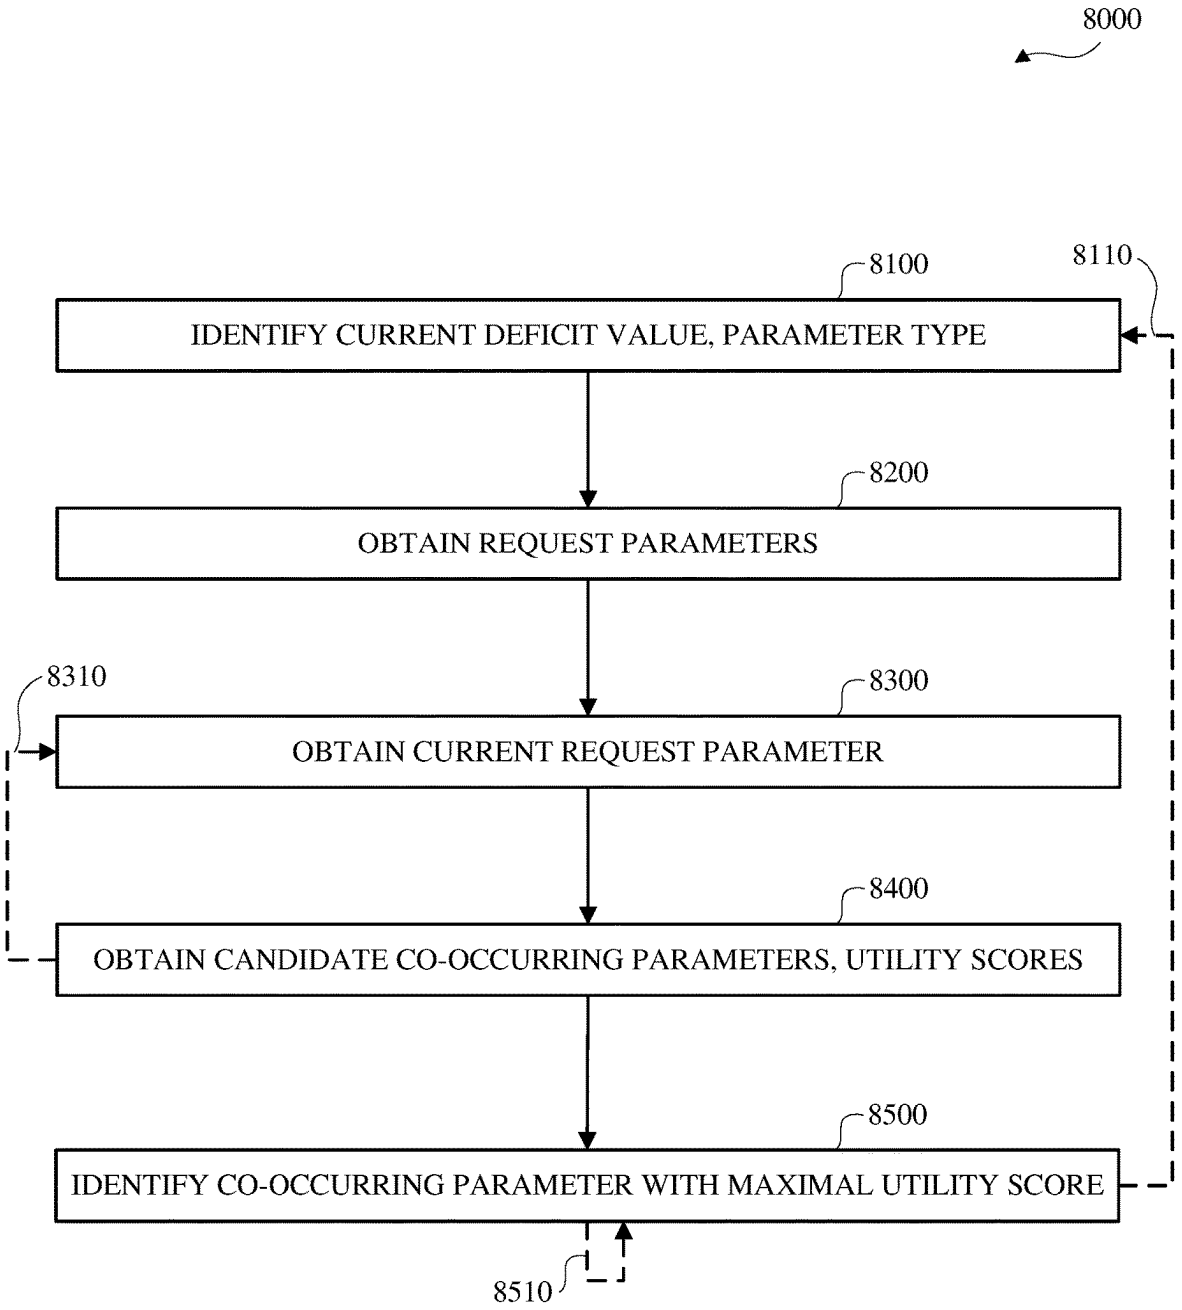
8000
8100
8110
IDENTIFY CURRENT DEFICIT VALUE, PARAMETER TYPE
8200
OBTAIN REQUEST PARAMETERS
8310
8300
OBTAIN CURRENT REQUEST PARAMETER
8400
OBTAIN CANDIDATE CO-OCCURRING PARAMETERS, UTILITY SCORES
8500
IDENTIFY CO-OCCURRING PARAMETER WITH MAXIMAL UTILITY SCORE
8510

FIG. 8

AUTOMATIC TARGET VISUALIZATION ADAPTATION

BACKGROUND

[0001] Advances in computer storage and database technology have led to exponential growth of the amount of data being created. Businesses are overwhelmed by the volume of the data stored in their computer systems. Existing database analytic tools are inefficient, costly to utilize, and require substantial configuration and training.

SUMMARY

[0002] Disclosed herein are implementations of automatic target visualization adaptation.

[0003] An aspect of the disclosure is a method of automatic target visualization adaptation, which includes obtaining, by a data access and analysis system, user input data including a natural language string expressing a request for data from the data access and analysis system, obtaining, by the data access and analysis system, from a large language model, large language model generated structured query language data expressing the natural language string, obtaining, by the data access and analysis system transforming the large language model generated structured query language data, first resolved request data expressing the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system, obtaining, by the data access and analysis system, target-parameters data for a target visualization type, obtaining, by the data access and analysis system, target-parameters deficit data by comparing the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero, obtaining, by the data access and analysis system, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter having the parameter type, wherein the co-occurring parameter is absent from the first resolved request data, obtaining, by the data access and analysis system transforming a combination of the first resolved request data and the co-occurring parameter, an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system, obtaining, by the data access and analysis system, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database, and outputting data for presenting the results data in accordance with the target visualization type.

[0004] Another aspect of the disclosure is an apparatus of a data access and analysis system for performing automatic target visualization adaptation. The apparatus includes one or more memories storing instructions for automatic target visualization adaptation and one or more processors that execute the instructions to obtain user input data that includes a natural language string that expresses a request for data from the data access and analysis system, obtain, from a large language model, large language model generated structured query language data that expresses the natural language string, obtain first resolved request data that expresses the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system, wherein, to obtain the first resolved request data the one or more processors execute the instructions to transform the large language model generated structured query language data, obtain target-parameters data for a target visualization type, obtain target-parameters deficit data, wherein, to obtain the target-parameters deficit data the one or more processors execute the instructions to compare the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero, obtain, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter that has the parameter type, wherein the co-occurring parameter is absent from the first resolved request data, obtain an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system, wherein, to obtain the automatically generated data query, the one or more processors execute the instructions to transform a combination of the first resolved request data and the co-occurring parameter, obtain, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database, and output data for presenting the results data in accordance with the target visualization type.

[0005] Another aspect of the disclosure is a computer readable medium storing processor executable instructions for performing automatic target visualization adaptation, which includes obtaining, by a data access and analysis system, user input data including a natural language string expressing a request for data from the data access and analysis system, obtaining, by the data access and analysis system, from a large language model, large language model generated structured query language data expressing the natural language string, obtaining, by the data access and analysis system transforming the large language model generated structured query language data, first resolved request data expressing the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system, obtaining, by the data access and analysis system, target-parameters data for a target visualization type, obtaining, by the data access and analysis system, target-parameters deficit data by comparing the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero, obtaining, by the data access and analysis system, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter having the parameter type, wherein the co-occurring parameter is absent from the first resolved request data, obtaining, by the data access and analysis system transforming a combination of the first resolved request data and the co-occurring parameter, an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system, obtaining, by the data access and analysis system, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database, and outputting data for presenting the results data in accordance with the target visualization type.

BRIEF DESCRIPTION OF THE DRAWINGS

[0006] The disclosure is best understood from the following detailed description when read in conjunction with the accompanying drawings. It is emphasized that, according to common practice, the various features of the drawings are not to scale. On the contrary, the dimensions of the various features are arbitrarily expanded or reduced for clarity.

[0007] FIG. **1** is a block diagram of an example of a computing device.

[0008] FIG. **2** is a block diagram of an example of a computing system.

[0009] FIG. **3** is a block diagram of an example of a low-latency data access and analysis system.

[0010] FIG. **4** is a diagram of an example of natural language to query language transformation in a data access and analysis system.

[0011] FIG. **5** is a flowchart of an example of validated resolved request data in a data access and analysis system.

[0012] FIG. **6** is a flowchart of an example of obtaining demonstration data in a data access and analysis system.

[0013] FIG. **7** is a flow chart of an example of obtaining results data with automatic target visualization adaptation in a data access and analysis system.

[0014] FIG. **8** is a flow chart of an example of obtaining co-occurring parameters for obtaining results data with automatic target visualization adaptation in a data access and analysis system.

DETAILED DESCRIPTION

[0015] Businesses and other organizations store large amounts of data, such as business records, transaction records, and the like, in data storage systems, such as relational database systems that store data as records, or rows, having values, or fields, corresponding to respective columns in tables that can be interrelated using key values. Database structures are often normalized or otherwise organized to maximize data density and to maximize transactional data operations at the expense of increased complexity and reduced accessibility for analysis. Individual records and tables may have little or no utility without substantial correlation, interpretation, and analysis. The complexity of these data structures and systems, the large volumes of data that can be stored therein, and the processing of such data and requests to access such data, result in high utilization of system resources, such as computational and communications resources, and limit the accessibility and the utility of the systems and data stored therein.

[0016] To improve the accessibility and utility of these systems, and the data stored therein, system interfaces may be implemented. For example, a database, or a database management system, may implement a defined structured query language and one or more interfaces for obtaining, processing, and responding to code expressed in accordance with the defined structured query language. In another example, a system may implement one or more interfaces, such as graphical user interfaces or application programming interfaces, which may be utilized for specific, narrowly defined, purposes.

[0017] Implementing and utilizing such systems and interfaces may inefficiently utilize system resources, increase risk with respect to performance, reliability, security, and accuracy, and limit access to and the use of the data. Furthermore, the complexity of the data structures, and the large volume of the data (e.g., millions or billions of rows) stored therein may render access to or the use of some data impracticable or impossible to achieve by the human mind using the tools that are available for accessing these systems.

[0018] The data access and analysis system described herein improves the efficiency with which system resources are utilized to access and use data, reduces the risks associated with the access and use of data, and increases access to and the utility of the data. For example, the data access and analysis system described herein indexes data to improve the efficiency, accuracy, and expressibility of obtaining and processing user input with respect to accessing and using the data, such as by implementing predictive and generative input techniques. In another example, aspects of the data access and analysis system described herein are implemented in a clustered or distributed computing configuration to improve performance, reliability, and security. In another example, the data access and analysis system described herein improves the efficiency with which system resources are utilized to access and use data by optimizing the generation and execution of data queries by a data source, such as an internal database of the data access and analysis system or an external database accessible by the data access and analysis system.

[0019] Natural language input, such as text, or string, input included in data expressing usage intent, may be transformed, or resolved, to obtain, or generate, a representation thereof, such as a resolved request, which may be, or may include, an ordered sequence of tokens, in accordance with a defined data-analytics grammar, which may include using natural language processing, which may include using one or more machine learning, or artificial intelligence, models, such as one or more large language models, such as one or more generative pre-trained transformer (GPT) models. Obtaining the representation of natural language input in accordance with the defined data-analytics grammar may include obtaining, from the large language model, a representation, of the natural language input, expressed in accordance with a defined structured query language, and obtaining the representation of natural language input in accordance with the defined data-analytics grammar by transforming the representation, of the natural language input, expressed in accordance with the defined structured query language. The representation of natural language input in accordance with the defined data-analytics grammar may be automatically transformed to obtain a representation thereof in accordance with a defined structured query language associated with, such as implemented by, a data source, such as a database.

[0020] The use of natural language processing of natural language input using one or more large language models reduces the resource utilization and increases accuracy associated human-machine interactions for obtaining, processing, or both, manual input to obtain equivalent results, such as to safely and securely explore, interact, analyze, and interrogate data, automatically generate narratives to explain insights, and augment data modeling through automatically generated descriptions and synonyms. Natural language processing of natural language input using one or more large language models may include accessing previously generated analytical objects, automatically generating insights, automatically generating visualizations, automatically generating narratives to explain insights, accessing and interrogating one or more data models, and the like.

[0021] To further improve the efficiency, accuracy, and expressibility, of obtaining and processing user input with respect to accessing and using the data, such as by implementing predictive and generative input techniques, the data access and analysis system described herein implements automatic target visualization adaptation.

[0022] FIG. **1** is a block diagram of an example of a computing device **1000**. One or more aspects of this disclosure may be implemented using the computing device **1000**. The computing device **1000** includes a processor **1100**, static memory **1200**, low-latency memory **1300**, an electronic communication unit **1400**, a user interface **1500**, a bus **1600**, and a power source **1700**. Although shown as a single unit, any one or more element of the computing device **1000** may be integrated into any number of separate physical units. For example, the low-latency memory **1300** and the processor **1100** may be integrated in a first physical unit and the user interface **1500** may be integrated in a second physical unit. Although not shown in FIG. **1**, the computing device **1000** may include other aspects, such as an enclosure or one or more sensors.

[0023] The computing device **1000** may be a stationary computing device, such as a personal computer (PC), a server, a workstation, a minicomputer, or a mainframe computer; or a mobile computing device, such as a mobile telephone, a personal digital assistant (PDA), a laptop, or a tablet PC.

[0024] The processor **1100** may include any device or combination of devices capable of manipulating or processing a signal or other information, including optical processors, quantum processors, molecular processors, or a combination thereof. The processor **1100** may be a central processing unit (CPU), such as a microprocessor, and may include one or more processing units, which may respectively include one or more processing cores. The processor **1100** may include multiple interconnected processors. For example, the multiple processors may be hardwired or networked, including wirelessly networked. In some implementations, the operations of the processor **1100** may be distributed across multiple physical devices or units that may be coupled directly or across a network. In some implementations, the processor **1100** may include a cache, or cache memory, for internal storage of operating data or instructions. The processor **1100** may include one or more special purpose processors, one or more digital signal processor (DSP), one or more microprocessors, one or more controllers, one or more microcontrollers, one or more integrated circuits, one or more an Application Specific Integrated Circuits, one or more Field Programmable Gate Array, one or more programmable logic arrays, one or more programmable logic controllers, firmware, one or more state machines, or any combination thereof.

[0025] The processor **1100** may be operatively coupled with the static memory **1200**, the low-latency memory **1300**, the electronic communication unit **1400**, the user interface **1500**, the bus **1600**, the power source **1700**, or any combination thereof. The processor may execute, which may include controlling, such as by sending electronic signals to, receiving electronic signals from, or both, the static memory **1200**, the low-latency memory **1300**, the electronic communication unit **1400**, the user interface **1500**, the bus **1600**, the power source **1700**, or any combination thereof to execute, instructions, programs, code, applications, or the like, which may include executing one or more aspects of an operating system, and which may include executing one or more instructions to perform one or more aspects described herein, alone or in combination with one or more other processors.

[0026] The static memory **1200** is coupled to the processor **1100** via the bus **1600** and may include non-volatile memory, such as a disk drive, or any form of non-volatile memory capable of persistent electronic information storage, such as in the absence of an active power supply. Although shown as a single block in FIG. **1**, the static memory **1200** may be implemented as multiple logical or physical units.

[0027] The static memory **1200** may store executable instructions or data, such as application data, an operating system, or a combination thereof, for access by the processor **1100**. The executable instructions may be organized into programmable modules or algorithms, functional programs, codes, code segments, or combinations thereof to perform one or more aspects, features, or elements described herein. The application data may include, for example, user files, database catalogs, configuration information, or a combination thereof. The operating system may be, for example, a desktop or laptop operating system; an operating system for a mobile device, such as a smartphone or tablet device; or an operating system for a large device, such as a mainframe computer.

[0028] The low-latency memory **1300** is coupled to the processor **1100** via the bus **1600** and may include any storage medium with low-latency data access including, for example, DRAM modules such as DDR SDRAM, Phase-Change Memory (PCM), flash memory, or a solid-state drive. Although shown as a single block in FIG. **1**, the low-latency memory **1300** may be implemented as multiple logical or physical units. Other configurations may be used. For example, low-latency memory **1300**, or a portion thereof, and processor **1100** may be combined, such as by using a system on a chip design.

[0029] The low-latency memory **1300** may store executable instructions or data, such as application data for low-latency access by the processor **1100**. The executable instructions may include, for example, one or more application programs that may be executed by the processor **1100**. The executable instructions may be organized into programmable modules or algorithms, functional programs, codes, code segments, and/or combinations thereof to perform various functions described herein.

[0030] The low-latency memory **1300** may be used to store data that is analyzed or processed using the systems or methods described herein. For example, storage of some or all data in low-latency memory **1300** instead of static memory **1200** may improve the execution speed of the systems and methods described herein by permitting access to data more quickly by an order of magnitude or greater (e.g., nanoseconds instead of microseconds).

[0031] The electronic communication unit **1400** is coupled to the processor **1100** via the bus **1600**. The electronic communication unit **1400** may include one or more transceivers. The electronic communication unit **1400** may, for example, provide a connection or link to a network via a network interface. The network interface may be a wired network interface, such as Ethernet, or a wireless network interface. For example, the computing device **1000** may communicate with other devices via the electronic communication unit **1400** and the network interface using one or

more network protocols, such as Ethernet, Transmission Control Protocol/Internet Protocol (TCP/IP), power line communication (PLC), Wi-Fi, infrared, ultra violet (UV), visible light, fiber optic, wire line, general packet radio service (GPRS), Global System for Mobile communications (GSM), code-division multiple access (CDMA), Long-Term Evolution (LTE), or other suitable protocols.

[0032] The user interface **1500** may include any unit capable of interfacing with a human user, such as a virtual or physical keypad, a touchpad, a display, a touch display, a speaker, a microphone, a video camera, a sensor, a printer, or any combination thereof. For example, a keypad can convert physical input of force applied to a key to an electrical signal that can be interpreted by computing device **1000**. In another example, a display can convert electrical signals output by computing device **1000** to light. The purpose of such devices may be to permit interaction with a human user, for example by accepting input from the human user and providing output back to the human user. The user interface **1500** may include a display; a positional input device, such as a mouse, touchpad, touchscreen, or the like; a keyboard; or any other human and machine interface device. The user interface **1500** may be coupled to the processor **1100** via the bus **1600**. In some implementations, the user interface **1500** can include a display, which can be a liquid crystal display (LCD), a cathode-ray tube (CRT), a light emitting diode (LED) display, an organic light emitting diode (OLED) display, an active-matrix organic light emitting diode (AMOLED), or other suitable display. In some implementations, the user interface **1500**, or a portion thereof, may be part of another computing device (not shown). For example, a physical user interface, or a portion thereof, may be omitted from the computing device **1000** and a remote or virtual interface may be used, such as via the electronic communication unit **1400**.

[0033] The bus **1600** is coupled to the static memory **1200**, the low-latency memory **1300**, the electronic communication unit **1400**, the user interface **1500**, and the power source **1700**. Although a single bus is shown in FIG. **1**, the bus **1600** may include multiple buses, which may be connected, such as via bridges, controllers, or adapters.

[0034] The power source **1700** provides energy to operate the computing device **1000**. The power source **1700** may be a general-purpose alternating-current (AC) electric power supply, or power supply interface, such as an interface to a household power source. In some implementations, the power source **1700** may be a single use battery or a rechargeable battery to allow the computing device **1000** to operate independently of an external power distribution system. For example, the power source **1700** may include a wired power source; one or more dry cell batteries, such as nickel-cadmium (NiCad), nickel-zinc (NiZn), nickel metal hydride (NiNM), lithium-ion (Li-ion); solar cells; fuel cells; or any other device capable of powering the computing device **1000**.

[0035] FIG. **2** is a block diagram of an example of a computing system **2000**. As shown, the computing system **2000** includes an external data source portion **2100**, an internal database analysis portion **2200**, and a system interface portion **2300**. The computing system **2000** may include other elements not shown in FIG. **2**, such as computer network elements.

[0036] The external data source portion **2100** may be associated with, such as controlled by, an external person, entity, or organization (second party). The internal database analysis portion **2200** may be associated with, such as created by or controlled by, a person, entity, or organization (first party). The system interface portion **2300** may be associated with, such as created by or controlled by, the first party and may be accessed by the first party, the second party, third parties, or a combination thereof, such as in accordance with access and authorization permissions and procedures.

[0037] The external data source portion **2100** is shown as including external database servers **2120** and external application servers **2140**. The external data source portion **2100** may include other elements not shown in FIG. **2**. The external data source portion **2100** may include external computing devices, such as the computing device **1000** shown in FIG. **1**, which may be used by or accessible to the external person, entity, or organization (second party) associated with the external data source portion **2100**, including but not limited to external database servers **2120** and external application servers **2140**. The external computing devices may include data regarding the operation of the external person, entity, or organization (second party) associated with the external data source portion **2100**.

[0038] The external database servers **2120** may be one or more computing devices configured to store data in a format and schema determined externally from the internal database analysis portion **2200**, such as by a second party associated with the external data source portion **2100**, or a third party. For example, the external database server **2120** may use a relational database and may include a database catalog with a schema. In some embodiments, the external database server **2120** may include a non-database data storage structure, such as a text-based data structure, such as a comma separated variable structure or an extensible markup language formatted structure or file. For example, the external database servers **2120** can include data regarding the production of materials by the external person, entity, or organization (second party) associated with the external data source portion **2100**, communications between the external person, entity, or organization (second party) associated with the external data source portion **2100** and third parties, or a combination thereof. Other data may be included. The external database may be a structured database system, such as a relational database operating in a relational database management system (RDBMS), which may be an enterprise database. In some embodiments, the external database may be an unstructured data source. The external data may include data or content, such as sales data, revenue data, profit data, tax data, shipping data, safety data, sports data, health data, meteorological data, or the like, or any other data, or combination of data, that may be generated by or associated with a user, an organization, or an enterprise and stored in a database system. For simplicity and clarity, data stored in or received from the external data source portion **2100** may be referred to herein as enterprise data.

[0039] The external application server **2140** may include application software, such as application software used by the external person, entity, or organization (second party) associated with the external data source portion **2100**. The external application server **2140** may include data or metadata relating to the application software.

[0040] The external database servers **2120**, the external application servers **2140**, or both, shown in FIG. **2** may represent logical units or devices that may be implemented

on one or more physical units or devices, which may be controlled or operated by the first party, the second party, or a third party.

[0041] The external data source portion **2100**, or aspects thereof, such as the external database servers **2120**, the external application servers **2140**, or both, may communicate with the internal database analysis portion **2200**, or an aspect thereof, such as one or more of the servers **2220**, **2240**, **2260**, and **2280**, via an electronic communication medium, which may be a wired or wireless electronic communication medium. For example, the electronic communication medium may include a local area network (LAN), a wide area network (WAN), a fiber channel network, the Internet, or a combination thereof.

[0042] The internal database analysis portion **2200** is shown as including servers **2220**, **2240**, **2260**, and **2280**. The servers **2220**, **2240**, **2260**, and **2280** may be computing devices, such as the computing device **1000** shown in FIG. **1**. Although four servers **2220**, **2240**, **2260**, and **2280** are shown in FIG. **2**, other numbers, or cardinalities, of servers may be used. For example, the number of computing devices may be determined based on the capability of individual computing devices, the amount of data to be processed, the complexity of the data to be processed, or a combination thereof. Other metrics may be used for determining the number of computing devices.

[0043] The internal database analysis portion **2200** may store data, process data, or store and process data. The internal database analysis portion **2200** may include a distributed cluster (not expressly shown) which may include two or more of the servers **2220**, **2240**, **2260**, and **2280**. The operation of the distributed cluster, such as the operation of the servers **2220**, **2240**, **2260**, and **2280** individually, in combination, or both, may be managed by a distributed cluster manager. For example, the server **2220** may be the distributed cluster manager. In another example, the distributed cluster manager may be implemented on another computing device (not shown). The data and processing of the distributed cluster may be distributed among the servers **2220**, **2240**, **2260**, and **2280**, such as by the distributed cluster manager.

[0044] Enterprise data from the external data source portion **2100**, such as from the external database server **2120**, the external application server **2140**, or both may be imported into the internal database analysis portion **2200**. The external database server **2120**, the external application server **2140**, or both may be one or more computing devices and may communicate with the internal database analysis portion **2200** via electronic communication. The imported data may be distributed among, processed by, stored on, or a combination thereof, one or more of the servers **2220**, **2240**, **2260**, and **2280**. Importing the enterprise data may include importing or accessing the data structures of the enterprise data. Importing the enterprise data may include generating internal data, internal data structures, or both, based on the enterprise data. The internal data, internal data structures, or both may accurately represent and may differ from the enterprise data, the data structures of the enterprise data, or both. In some implementations, enterprise data from multiple external data sources may be imported into the internal database analysis portion **2200**. For simplicity and clarity, data stored or used in the internal database analysis portion **2200** may be referred to herein as internal data. For example, the internal data, or a portion thereof, may represent, and may be distinct from, enterprise data imported into or accessed by the internal database analysis portion **2200**.

[0045] The system interface portion **2300** may include one or more client devices **2320**, **2340**. The client devices **2320**, **2340** may be computing devices, such as the computing device **1000** shown in FIG. **1**. For example, one of the client devices **2320**, **2340** may be a desktop or laptop computer and the other of the client devices **2320**, **2340** may be a mobile device, smartphone, or tablet. One or more of the client devices **2320**, **2340** may access the internal database analysis portion **2200**. For example, the internal database analysis portion **2200** may provide one or more services, application interfaces, or other electronic computer communication interfaces, such as a web site, and the client devices **2320**, **2340** may access the interfaces provided by the internal database analysis portion **2200**, which may include accessing the internal data stored in the internal database analysis portion **2200**.

[0046] In an example, one or more of the client devices **2320**, **2340** may send a message or signal indicating a request for data, which may include a request for data analysis, to the internal database analysis portion **2200**. The internal database analysis portion **2200** may receive and process the request, which may include distributing the processing among one or more of the servers **2220**, **2240**, **2260**, and **2280**, may generate a response to the request, which may include generating or modifying internal data, internal data structures, or both, and may output the response to the client device **2320**, **2340** that sent the request. Processing the request may include accessing one or more internal data indexes, an internal database, or a combination thereof. The client device **2320**, **2340** may receive the response, including the response data or a portion thereof, and may store, output, or both, the response, or a representation thereof, such as a representation of the response data, or a portion thereof, which may include presenting the representation via a user interface on a presentation device of the client device **2320**, **2340**, such as to a user of the client device **2320**, **2340**.

[0047] The system interface portion **2300**, or aspects thereof, such as one or more of the client devices **2320**, **2340**, may communicate with the internal database analysis portion **2200**, or an aspect thereof, such as one or more of the servers **2220**, **2240**, **2260**, and **2280**, via an electronic communication medium, which may be a wired or wireless electronic communication medium. For example, the electronic communication medium may include a local area network (LAN), a wide area network (WAN), a fiber channel network, the Internet, or a combination thereof.

[0048] FIG. **3** is a block diagram of an example of a low-latency data access and analysis system **3000**. The low-latency data access and analysis system **3000**, or aspects thereof, may be similar to the internal database analysis portion **2200** shown in FIG. **2**, except as described herein or otherwise clear from context. The low-latency data access and analysis system **3000**, or aspects thereof, may be implemented on one or more computing devices, such as servers **2220**, **2240**, **2260**, and **2280** shown in FIG. **2**, which may be in a clustered or distributed computing configuration. As used herein, the terms "low-latency data access and analysis system," "low-latency data analysis system," and "low-latency database analysis system" indicate a computer implemented system, such as the low-latency data access and analysis system **3000** shown in FIG. **3**, that obtains,

stores, organizes, processes, automatically analyzes, and outputs data and visualizations thereof.

[0049] The low-latency data access and analysis system **3000**, which may be a low-latency database analysis system, may store and maintain the internal data, or a portion thereof, such as low-latency data, in a low-latency memory device, such as the low-latency memory **1300** shown in FIG. **1**, or any other type of data storage medium or combination of data storage devices with relatively fast (low-latency) data access, organized in a low-latency data structure. In some embodiments, the low-latency data access and analysis system **3000** may be implemented as one or more logical devices in a cloud-based configuration optimized for automatic database analysis.

[0050] As shown, the low-latency data access and analysis system **3000** includes a distributed cluster manager **3100**, a security and governance unit **3200**, a distributed in-memory database **3300**, an enterprise data interface unit **3400**, a distributed in-memory ontology unit **3500**, a semantic interface unit **3600**, a relational analysis unit **3700**, a natural language processing unit **3710**, a data utility unit **3720**, an insight unit **3730**, an object search unit **3800**, an object utility unit **3810**, a system configuration unit **3820**, a user customization unit **3830**, a system access interface unit **3900**, a real-time collaboration unit **3910**, a third-party integration unit **3920**, and a persistent storage unit **3930**, which may be collectively referred to as the components of the low-latency data access and analysis system **3000**.

[0051] Although not expressly shown in FIG. **3**, one or more of the components of the low-latency data access and analysis system **3000** may be implemented on one or more operatively connected physical or logical computing devices, such as in a distributed cluster computing configuration, such as the internal database analysis portion **2200** shown in FIG. **2**. Although shown separately in FIG. **3**, one or more of the components of the low-latency data access and analysis system **3000**, or respective aspects thereof, may be combined or otherwise organized.

[0052] The low-latency data access and analysis system **3000** may include different, fewer, or additional components not shown in FIG. **3**. The aspects or components implemented in an instance of the low-latency data access and analysis system **3000** may be configurable. For example, the insight unit **3730** may be omitted or disabled. One or more of the components of the low-latency data access and analysis system **3000** may be implemented in a manner such that aspects thereof are divided or combined into various executable modules or libraries in a manner which may differ from that described herein.

[0053] The low-latency data access and analysis system **3000** may implement an application programming interface (API), which may monitor, receive, or both, input signals or messages from external devices and systems, client systems, process received signals or messages, transmit corresponding signals or messages to one or more of the components of the low-latency data access and analysis system **3000**, and output, such as transmit or send, output messages or signals to respective external devices or systems. The low-latency data access and analysis system **3000** may be implemented in a distributed computing configuration.

[0054] The distributed cluster manager **3100** manages the operative configuration of the low-latency data access and analysis system **3000**. Managing the operative configuration of the low-latency data access and analysis system **3000** may include controlling the implementation of and distribution of processing and storage across one or more logical devices operating on one or more physical devices, such as the servers **2220**, **2240**, **2260**, and **2280** shown in FIG. **2**. The distributed cluster manager **3100** may generate and maintain configuration data for the low-latency data access and analysis system **3000**, such as in one or more tables, identifying the operative configuration of the low-latency data access and analysis system **3000**. For example, the distributed cluster manager **3100** may automatically update the low-latency data access and analysis system configuration data in response to an operative configuration event, such as a change in availability or performance for a physical or logical unit of the low-latency data access and analysis system **3000**. One or more of the component units of low-latency data access and analysis system **3000** may access the data analysis system configuration data, such as to identify intercommunication parameters or paths.

[0055] The security and governance unit **3200** may describe, implement, enforce, or a combination thereof, rules and procedures for controlling access to aspects of the low-latency data access and analysis system **3000**, such as the internal data of the low-latency data access and analysis system **3000** and the features and interfaces of the low-latency data access and analysis system **3000**. The security and governance unit **3200** may apply security at an ontological level to control or limit access to the internal data of the low-latency data access and analysis system **3000**, such as to columns, tables, rows, or fields, which may include using row-level security.

[0056] Although shown as a single unit in FIG. **3**, the distributed in-memory database **3300** may be implemented in a distributed configuration, such as distributed among the servers **2220**, **2240**, **2260**, and **2280** shown in FIG. **2**, which may include multiple in-memory database instances. Each in-memory database instance may utilize one or more distinct resources, such as processing or low-latency memory resources, that differ from the resources utilized by the other in-memory database instances. In some embodiments, the in-memory database instances may utilize one or more shared resources, such as resources utilized by two or more in-memory database instances.

[0057] The distributed in-memory database **3300** may generate, maintain, or both, a low-latency data structure and data stored or maintained therein (low-latency data). The low-latency data may include principal data, which may represent enterprise data, such as enterprise data imported from an external enterprise data source, such as the external data source portion **2100** shown in FIG. **2**. In some implementations, the distributed in-memory database **3300** may include system internal data representing one or more aspects, features, or configurations of the low-latency data access and analysis system **3000**. The distributed in-memory database **3300** and the low-latency data stored therein, or a portion thereof, may be accessed using commands, messages, or signals in accordance with a defined structured query language associated with, such as implemented by, the distributed in-memory database **3300**.

[0058] The low-latency data, or a portion thereof, may be organized as tables in the distributed in-memory database **3300**. A table may be a data structure to organize or group the data or a portion thereof, such as related or similar data.

A table may have a defined structure. For example, each table may define or describe a respective set of one or more columns.

[0059] A column may define or describe the characteristics of a discrete aspect of the data in the table. For example, the definition or description of a column may include an identifier, such as a name, for the column within the table, and one or more constraints, such as a data type, for the data corresponding to the column in the table. The definition or description of a column may include other information, such as a description of the column. The data in a table may be accessible or partitionable on a per-column basis. The set of tables, including the column definitions therein, and information describing relationships between elements, such as tables and columns, of the database may be defined or described by a database schema or design. The cardinality of columns of a table, and the definition and organization of the columns, may be defined by the database schema or design. Adding, deleting, or modifying a table, a column, the definition thereof, or a relationship or constraint thereon, may be a modification of the database design, schema, model, or structure.

[0060] The low-latency data, or a portion thereof, may be stored in the database as one or more rows or records in respective tables. Each record or row of a table may include a respective field or cell corresponding to each column of the table. A field may store a discrete data value. The cardinality of rows of a table, and the values stored therein, may be variable based on the data. Adding, deleting, or modifying rows, or the data stored therein may omit modification of the database design, schema, or structure. The data stored in respective columns may be identified or defined as a measure data, attribute data, or enterprise ontology data (e.g., metadata).

[0061] Measure data, or measure values, include quantifiable or additive numeric values, such as integer or floating-point values, which may include numeric values indicating sizes, amounts, degrees, or the like. A column defined as representing measure values may be referred to herein as a measure or fact. A measure may be a property on which quantitative operations (e.g., sum, count, average, minimum, maximum) may be performed to calculate or determine a result or output.

[0062] Attribute data, or attribute values, include non-quantifiable values, such as text or image data, which may indicate names and descriptions, quantifiable values designated, defined, or identified as attribute data, such as numeric unit identifiers, or a combination thereof. A column defined as including attribute values may be referred to herein as an attribute or dimension. For example, attributes may include text, identifiers, timestamps, or the like.

[0063] Enterprise ontology data may include data that defines or describes one or more aspects of the database, such as data that describes one or more aspects of the attributes, measures, rows, columns, tables, relationships, or other aspects of the data or database schema. For example, a portion of the database design, model, or schema may be represented as enterprise ontology data in one or more tables in the database.

[0064] Distinctly identifiable data in the low-latency data may be referred to herein as a data portion. For example, the low-latency data stored in the distributed in-memory database **3300** may be referred to herein as a data portion, a table from the low-latency data may be referred to herein as a data portion, a column from the low-latency data may be referred to herein as a data portion, a row or record from the low-latency data may be referred to herein as a data portion, a value from the low-latency data may be referred to herein as a data portion, a relationship defined in the low-latency data may be referred to herein as a data portion, enterprise ontology data describing the low-latency data may be referred to herein as a data portion, or any other distinctly identifiable data, or combination thereof, from the low-latency data may be referred to herein as a data portion.

[0065] The distributed in-memory database **3300** may create or add one or more data portions, such as a table, may read from or access one or more data portions, may update or modify one or more data portions, may remove or delete one or more data portions, or a combination thereof. Adding, modifying, or removing data portions may include changes to the data model of the low-latency data. Changing the data model of the low-latency data may include notifying one or more other components of the low-latency data access and analysis system **3000**, such as by sending, or otherwise making available, a message or signal indicating the change. For example, the distributed in-memory database **3300** may create or add a table to the low-latency data and may transmit or send a message or signal indicating the change to the semantic interface unit **3600**.

[0066] In some implementations, a portion of the low-latency data may represent a data model of an external enterprise database and may omit the data stored in the external enterprise database, or a portion thereof. For example, prioritized data may be cached in the distributed in-memory database **3300** and the other data may be omitted from storage in the distributed in-memory database **3300**, which may be stored in the external enterprise database. In some implementations, requesting data from the distributed in-memory database **3300** may include requesting the data, or a portion thereof, from the external enterprise database.

[0067] The distributed in-memory database **3300** may receive one or more messages or signals indicating respective data queries for the low-latency data, or a portion thereof, which may include data queries for modified, generated, or aggregated data generated based on the low-latency data, or a portion thereof. For example, the distributed in-memory database **3300** may receive a data query from the semantic interface unit **3600**, such as in accordance with a request for data. The data queries received by the distributed in-memory database **3300** may be agnostic to the distributed configuration of the distributed in-memory database **3300**. A data query, or a portion thereof, may be expressed in accordance with the defined structured query language implemented by the distributed in-memory database **3300**. In some implementations, a data query, or a portion thereof, may be expressed in accordance with a defined structured query language implemented by a defined database other than the distributed in-memory database **3300**, such as an external database. In some implementations, a data query may be included, such as stored or communicated, in a data-query data structure or container.

[0068] The distributed in-memory database **3300** may execute or perform one or more queries to generate or obtain response data responsive to the data query based on the low-latency data. Unless expressly described, or otherwise clear from context, descriptions herein of a table in the context of performing, processing, or executing a data query that include accessing, such as reading, writing, or otherwise

using, a table, or data from a table, may refer to a table stored, or otherwise maintained, in the distributed in-memory database independently of the data query or may refer to tabular data obtained, such as generated, in accordance with the data query.

[0069] The distributed in-memory database **3300** may interpret, evaluate, or otherwise process a data query to generate one or more distributed-queries, which may be expressed in accordance with the defined structured query language. For example, an in-memory database instance of the distributed in-memory database **3300** may be identified as a query coordinator. The query coordinator may generate a query plan, which may include generating one or more distributed-queries, based on the received data-query. The query plan may include query execution instructions for executing one or more queries, or one or more portions thereof, based on the received data-query by the one or more of the in-memory database instances. Generating the query plan may include optimizing the query plan. The query coordinator may distribute, or otherwise make available, the respective portions of the query plan, as query execution instructions, to the corresponding in-memory database instances.

[0070] The respective in-memory database instances may receive the corresponding query execution instructions from the query coordinator. The respective in-memory database instances may execute the corresponding query execution instructions to obtain, process, or both, data (intermediate results data) from the low-latency data. The respective in-memory database instances may output, or otherwise make available, the intermediate results data, such as to the query coordinator.

[0071] The query coordinator may execute a respective portion of query execution instructions (allocated to the query coordinator) to obtain, process, or both, data (intermediate results data) from the low-latency data. The query coordinator may receive, or otherwise access, the intermediate results data from the respective in-memory database instances. The query coordinator may combine, aggregate, or otherwise process, the intermediate results data to obtain results data.

[0072] In some embodiments, obtaining the intermediate results data by one or more of the in-memory database instances may include outputting the intermediate results data to, or obtaining intermediate results data from, one or more other in-memory database instances, in addition to, or instead of, obtaining the intermediate results data from the low-latency data.

[0073] The distributed in-memory database **3300** may output, or otherwise make available, the results data to the semantic interface unit **3600**.

[0074] The enterprise data interface unit **3400** may interface with, or communicate with, an external enterprise data system. For example, the enterprise data interface unit **3400** may receive or access enterprise data from or in an external system, such as an external database. The enterprise data interface unit **3400** may import, evaluate, or otherwise process the enterprise data to populate, create, or modify data stored in the low-latency data access and analysis system **3000**. The enterprise data interface unit **3400** may receive, or otherwise access, the enterprise data from one or more external data sources, such as the external data source portion **2100** shown in FIG. **2**, and may represent the enterprise data in the low-latency data access and analysis system **3000** by importing, loading, or populating the enterprise data as principal data in the distributed in-memory database **3300**, such as in one or more low-latency data structures. The enterprise data interface unit **3400** may implement one or more data connectors, which may transfer data between, for example, the external data source and the distributed in-memory database **3300**, which may include altering, formatting, evaluating, or manipulating the data.

[0075] The enterprise data interface unit **3400** may receive, access, or generate metadata that identifies one or more parameters or relationships for the principal data, such as based on the enterprise data, and may include the generated metadata in the low-latency data stored in the distributed in-memory database **3300**. For example, the enterprise data interface unit **3400** may identify characteristics of the principal data such as, attributes, measures, values, unique identifiers, tags, links, keys, or the like, and may include metadata representing the identified characteristics in the low-latency data stored in the distributed in-memory database **3300**. The characteristics of the data can be automatically determined by receiving, accessing, processing, evaluating, or interpreting the schema in which the enterprise data is stored, which may include automatically identifying links or relationships between columns, classifying columns (e.g., using column names), and analyzing or evaluating the data.

[0076] Although not shown separately in FIG. **3**, the low-latency data access and analysis system **3000** implements a canonical, or system-defined, chronometry. The system-defined chronometry defines the measurement, storage, processing, organization, scale, expression, and representation of time and temporal data in the low-latency database analysis system **3000**. For example, the system-defined chronometry may correspond with a Gregorian calendar, or a defined variant thereof. The system-defined chronometry defines one or more chronometric units, which may be nominal, or named, representations of respective temporal intervals. A reference chronometric unit, such as a 'second' chronometric unit, may represent a minimal temporal interval in the low-latency database analysis system. One or more aspects of the system-defined chronometry may be defined by the operating environment of the low-latency database analysis system, such as by a hardware component, an operating system, or a combination thereof. For example, a hardware component, such as a system clock (clock circuit) may define the temporal interval of the reference chronometric unit and an operating system may define one or more other chronometric units with reference to the reference chronometric unit.

[0077] The low-latency database analysis system **3000** may define or describe one or more chronometric unit types, such as a 'minute' chronometric unit type, an 'hour' chronometric unit type, a 'day' chronometric unit type, a 'week' chronometric unit type, a 'month' chronometric unit type, a 'quarter' chronometric unit type, a 'year' chronometric unit type, or any other type of chronometric unit. A temporal point may be represented, such as stored or processed, in the low-latency database analysis system as an epoch value, which may be an integer value, such that each temporal point from the contiguous sequence of temporal points that comprises the temporal continuum corresponds with a respective epoch value. A temporal location may be represented in the low-latency database analysis system as an epoch value and may be expressed in the low-latency

database analysis system using one or more chronometric units, or respective values thereof. The system-defined chronometry defines respective descriptors, such as a day-of-week-name, month-name, and the like. Data defining or describing the system-defined chronometry may be stored in the low-latency data access and analysis system as a chronometric dataset. In some implementations, the low-latency data access and analysis system may define or describe a domain-specific chronometry that differs from the system-defined chronometry. The chronometric units defined or described by the domain-specific chronometry, except for the reference chronometric unit, may differ from the chronometric units defined or described by the system-defined chronometry. Data defining or describing the domain-specific chronometry may be stored in the low-latency data access and analysis system as a chronometric dataset.

[0078] Distinctly identifiable operative data units or structures representing one or more data portions, one or more entities, users, groups, or organizations represented in the internal data, or one or more aggregations, collections, relations, analytical results, visualizations, or groupings thereof, may be represented in the low-latency data access and analysis system **3000** as objects. An object may include a unique identifier for the object, such as a fully qualified name, a globally unique identifier (GUID), or a universally unique identifier (UUID). An object may include a name, such as a displayable value, for the object.

[0079] For example, an object may represent a user, a group, an entity, an organization, a privilege, a role, a table, a column, a data relationship, a worksheet, a view, an access context, an answer, an insight, a pinboard, a tag, a comment, a trigger, a defined variable, a data source, an object-level security rule, a row-level security rule, or any other data capable of being distinctly identified and stored or otherwise obtained in the low-latency data access and analysis system **3000**. An object may represent or correspond with a logical entity. Data describing an object may include data operatively or uniquely identifying data corresponding to, or represented by, the object in the low-latency data access and analysis system. For example, a column in a table in a database in, or accessible by, the low-latency data access and analysis system may be represented in the low-latency data access and analysis system as an object and the data describing or defining the object may include data operatively or uniquely identifying the column.

[0080] A worksheet (worksheet object), or worksheet table, may be a logical table, or a definition thereof, which may be a collection, a subset (such as a subset of columns from one or more tables), or both, of data from one or more data sources, such as columns in one or more tables, such as in the distributed in-memory database **3300**, or in an external database. A worksheet, or a definition thereof, may include one or more data organization or manipulation definitions, such as join paths or worksheet-column definitions, which may be user defined. A worksheet may be a data structure that may contain one or more rules or definitions that may define or describe how a respective tabular set of data may be obtained, which may include defining one or more sources of data, such as one or more columns from the distributed in-memory database **3300**, or from an external database. A worksheet may be a data source. For example, a worksheet may include references to one or more data sources, such as columns in one or more tables, such as in the distributed in-memory database **3300**, or in an external database, and a request for data referencing the worksheet may access the data from the data sources referenced in the worksheet. In some implementations, a worksheet may omit aggregations of the data from the data sources referenced in the worksheet. A worksheet may be processed, or transformed, automatically, which may be similar to transforming a resolved request, or an analytical object, to obtain a data query, as described herein, except as is described herein or as is otherwise clear from context, to generate one or more data queries that may be executed by the corresponding data source to generate data, or content, corresponding to the worksheet.

[0081] For example, a worksheet may include one or more tables, or one or more columns thereof, associated with an organizing characteristic, or use case. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by combining data, using worksheet-specific names, descriptions, or both, of data elements, such as column names. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific classifications, or types, of data elements, such as columns including numeric data. For example, a column including numeric data may be a measure column in the data model and may be used as an attribute column in a worksheet. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific aggregation functions. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific formatting and currency symbols. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific identification of columns that contain geographical data. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific vocabulary, which may include mapping the worksheet-specific vocabulary to data. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific formulas for consistency and governance. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using worksheet-specific user or group access to one or more portions or parts of the constituent data. A worksheet may reduce the complexity of the data model, or data models, storing the data, such as by using a worksheet-specific filtered set of data.

[0082] An answer (answer object), or report, may represent a defined, such as previously generated, request for data, such as a resolved request. An answer may include information describing a visualization of data responsive to the request for data.

[0083] A visualization (visualization object) may be a defined representation or expression of data, such as a visual representation of the data, for presentation to a user or human observer, such as via a user interface. Although described as a visual representation, in some implementations, a visualization may include non-visual aspects, such as auditory or haptic presentation aspects. A visualization may be generated to represent a defined set of data in accordance with a defined visualization type or template (visualization template object), such as in a chart, graph, or tabular form. Example visualization types may include, and are not limited to, choropleths, cartograms, dot distribution maps, proportional symbol maps, contour/isopleth/isarith-

mic maps, dasymetric map, self-organizing map, timeline, time series, connected scatter plots, Gantt charts, steam graph/theme river, arc diagrams, polar area/rose/circumplex charts, Sankey diagrams, alluvial diagrams, pie charts, histograms, tag clouds, bubble charts, bubble clouds, bar charts, radial bar charts, tree maps, scatter plots, line charts, step charts, area charts, stacked graphs, heat maps, parallel coordinates, spider charts, box and whisker plots, mosaic displays, waterfall charts, funnel charts, or radial tree maps. A visualization template may define or describe one or more visualization parameters, such as one or more color parameters. Visualization data for a visualization may include values of one or more of the visualization parameters of the corresponding visualization template.

[0084] A view (view object) may be a logical table, or a definition thereof, which may be a collection, a sub-set, or both, of data from one or more data sources, such as columns in one or more tables, such as in the distributed in-memory database **3300**, or in an external database. For example, a view may be generated based on an answer, such as by storing the answer as a view. A view may define or describe a data aggregation. A view may be a data source. For example, a view may include references to one or more data sources, such as columns in one or more tables, such as in the distributed in-memory database **3300**, or in an external database, which may include a definition or description of an aggregation of the data from a respective data source, and a request for data referencing the view may access the aggregated data, the data from the unaggregated data sources referenced in the worksheet, or a combination thereof. The unaggregated data from data sources referenced in the view defined or described as aggregated data in the view may be unavailable based on the view. A view may be a materialized view or an unmaterialized view. A request for data referencing a materialized view may obtain data from a set of data previously obtained (view-materialization) in accordance with the definition of the view and the request for data. A request for data referencing an unmaterialized view may obtain data from a set of data currently obtained in accordance with the definition of the view and the request for data.

[0085] A pinboard (pinboard object), or dashboard, may be a defined collection or grouping of objects, such as visualizations, answers, or insights. Pinboard data for a pinboard may include information associated with the pinboard, which may be associated with respective objects included in the pinboard.

[0086] An access context (access-context object) may be a set or collection of data associated with, such as including, data expressing usage intent, such as a request for data, data responsive to data expressing usage intent, or a discretely related sequence or series of requests for data or other interactions with the low-latency data access and analysis system **3000**, and a corresponding data structure for containing such data. For example, data expressing usage intent may be generated by the low-latency data access and analysis system **3000**, or a component thereof, such as the system access interface unit **3900**, such as in response to input, such as user input, obtained by the low-latency data access and analysis system **3000**. In another example, data expressing usage intent may be obtained, received, or otherwise accessed, by the low-latency data access and analysis system **3000**, or a component thereof, from an external device or system.

[0087] A definition may be a set of data describing the structure or organization of a data portion. For example, in the distributed in-memory database **3300**, a column definition may define one or more aspects of a column in a table, such as a name of the column, a description of the column, a datatype for the column, or any other information about the column that may be represented as discrete data.

[0088] A data source object may represent a source or repository of data accessible by the low-latency data access and analysis system **3000**. A data source object may include data indicating an electronic communication location, such as an address, of a data source, connection information, such as protocol information, authentication information, or a combination thereof, or any other information about the data source that may be represented as discrete data. For example, a data source object may represent a table in the distributed in-memory database **3300**, or in an external database, and include data for accessing the table from the database, such as information identifying the database, information identifying a schema within the database, and information identifying the table within the schema within the database. A data source object (external data source object) may represent an external data source. For example, an external data source object may include data indicating an electronic communication location, such as an address, of an external data source, connection information, such as protocol information, authentication information, or a combination thereof, or any other information about the external data source that may be represented as discrete data.

[0089] A sticker (sticker object) may be a description of a classification, category, tag, subject area, or other information that may be associated with one or more other objects such that objects associated with a sticker may be grouped, sorted, filtered, or otherwise identified based on the sticker. In the distributed in-memory database **3300** a tag may be a discrete data portion that may be associated with other data portions, such that data portions associated with a tag may be grouped, sorted, filtered, or otherwise identified based on the tag.

[0090] The distributed in-memory ontology unit **3500** generates, maintains, or both, information (ontological data) defining or describing the operative ontological structure of the objects represented in the low-latency data access and analysis system **3000**, such as in the low-latency data stored in the distributed in-memory database **3300**, which may include describing attributes, properties, states, or other information about respective objects and may include describing relationships among respective objects.

[0091] Objects may be referred to herein as primary objects, secondary objects, or tertiary objects. Other types of objects may be used.

[0092] Primary objects may include objects representing distinctly identifiable operative data units or structures representing one or more data portions in the distributed in-memory database **3300**, or another data source in, or accessible by, the low-latency data access and analysis system **3000**. For example, primary objects may be data source objects, table objects, column objects, relationship objects, or the like. Primary objects may include worksheets, views, filters, such as row-level-security filters and table filters, variables, or the like. Primary objects may be referred to herein as data objects or queryable objects.

[0093] Secondary objects may be objects representing distinctly identifiable operative data units or structures rep-

resenting analytical data aggregations, collections, analytical results, visualizations, or groupings thereof, such as pinboard objects, answer objects, insights, visualization objects, resolved-request objects, and the like. Secondary objects may be referred to herein as analytical objects.

[0094] Tertiary objects may be objects representing distinctly identifiable operative data units or structures representing operational aspects of the low-latency data access and analysis system **3000**, such as one or more entities, users, groups, or organizations represented in the internal data, such as user objects, user-group objects, role objects, sticker objects, and the like.

[0095] The distributed in-memory ontology unit **3500** may represent the ontological structure, which may include the objects therein, as a graph having nodes and edges. A node may be a representation of an object in the graph structure of the distributed in-memory ontology unit **3500**. Anode, representing an object, can include one or more components. The components of a node may be versioned, such as on a per-component basis. For example, a node can include a header component, a content component, or both. A header component may include information about the node. A content component may include the content of the node. An edge may represent a relationship between nodes, which may be directional.

[0096] In some implementations, the distributed in-memory ontology unit **3500** graph may include one or more nodes, edges, or both, representing one or more objects, relationships or both, corresponding to a respective internal representation of enterprise data stored in an external enterprise data storage unit, wherein a portion of the data stored in the external enterprise data storage unit represented in the distributed in-memory ontology unit **3500** graph is omitted from the distributed in-memory database **3300**.

[0097] In some embodiments, the distributed in-memory ontology unit **3500** may generate, modify, or remove a portion of the ontology graph in response to one or more messages, signals, or notifications from one or more of the components of the low-latency data access and analysis system **3000**. For example, the distributed in-memory ontology unit **3500** may generate, modify, or remove a portion of the ontology graph in response to receiving one or more messages, signals, or notifications from the distributed in-memory database **3300** indicating a change to the low-latency data structure. In another example, the distributed in-memory database **3300** may send one or more messages, signals, or notifications indicating a change to the low-latency data structure to the semantic interface unit **3600** and the semantic interface unit **3600** may send one or more messages, signals, or notifications indicating the change to the low-latency data structure to the distributed in-memory ontology unit **3500**.

[0098] The distributed in-memory ontology unit **3500** may be distributed, in-memory, multi-versioned, transactional, consistent, durable, or a combination thereof. The distributed in-memory ontology unit **3500** is transactional, which may include implementing atomic concurrent, or substantially concurrent, updating of multiple objects. The distributed in-memory ontology unit **3500** is durable, which may include implementing a robust storage that prevents data loss subsequent to or as a result of the completion of an atomic operation. The distributed in-memory ontology unit **3500** is consistent, which may include performing operations associated with a request for data with reference to or using a discrete data set, which may mitigate or eliminate the risk of inconsistent results.

[0099] The distributed in-memory ontology unit **3500** may generate, output, or both, one or more event notifications. For example, the distributed in-memory ontology unit **3500** may generate, output, or both, a notification, or notifications, in response to a change of the distributed in-memory ontology. The distributed in-memory ontology unit **3500** may identify a portion of the distributed in-memory ontology (graph) associated with a change of the distributed in-memory ontology, such as one or more nodes depending from a changed node, and may generate, output, or both, a notification, or notifications indicating the identified relevant portion of the distributed in-memory ontology (graph). One or more aspects of the low-latency data access and analysis system **3000** may cache object data and may receive the notifications from the distributed in-memory ontology unit **3500**, which may reduce latency and network traffic relative to systems that omit caching object data or omit notifications relevant to changes to portions of the distributed in-memory ontology (graph).

[0100] The distributed in-memory ontology unit **3500** may implement prefetching. For example, the distributed in-memory ontology unit **3500** may predictively, such as based on determined probabilistic utility, fetch one or more nodes, such as in response to access to a related node by a component of the low-latency data access and analysis system **3000**.

[0101] The distributed in-memory ontology unit **3500** may implement a multi-version concurrency control graph data storage unit. Each node, object, or both, may be versioned. Changes to the distributed in-memory ontology may be reversible. For example, the distributed in-memory ontology may have a first state prior to a change to the distributed in-memory ontology, the distributed in-memory ontology may have a second state subsequent to the change, and the state of the distributed in-memory ontology may be reverted to the first state subsequent to the change, such as in response to the identification of an error or failure associated with the second state.

[0102] In some implementations, reverting a node, or a set of nodes, may omit reverting one or more other nodes. In some implementations, the distributed in-memory ontology unit **3500** may maintain a change log indicating a sequential record of changes to the distributed in-memory ontology (graph), such that a change to a node or a set of nodes may be reverted and one or more other changes subsequent to the reverted change may be reverted for consistency.

[0103] The distributed in-memory ontology unit **3500** may implement optimistic locking to reduce lock contention times. The use of optimistic locking permits improved throughput of data through the distributed in-memory ontology unit **3500**.

[0104] The semantic interface unit **3600** may implement procedures and functions to provide a semantic interface between the distributed in-memory database **3300**, or an external database, and one or more of the other components of the low-latency data access and analysis system **3000**.

[0105] The semantic interface unit **3600** may implement ontological data management, data-query generation, authentication and access control, object statistical data collection, or a combination thereof.

[0106] Ontological data management may include object lifecycle management, object data persistence, ontological modifications, or the like. Object lifecycle management may include creating one or more objects, reading or otherwise accessing one or more objects, updating or modifying one or more objects, deleting or removing one or more objects, or a combination thereof. For example, the semantic interface unit **3600** may interface or communicate with the distributed in-memory ontology unit **3500**, which may store the ontological data, object data, or both, to perform object lifecycle management, object data persistence, ontological modifications, or the like.

[0107] For example, the semantic interface unit **3600** may receive, or otherwise access, a message, signal, or notification, such as from the distributed in-memory database **3300**, or from an external database, indicating the creation or addition of a data portion, such as a table, in the low-latency data stored in the distributed in-memory database **3300**, or in an external database, and the semantic interface unit **3600** may communicate with the distributed in-memory ontology unit **3500** to create an object in the ontology representing the added data portion. The semantic interface unit **3600** may transmit, send, or otherwise make available, a notification, message, or signal to the relational analysis unit **3700** indicating that the ontology has changed.

[0108] The semantic interface unit **3600** may receive, or otherwise access, a request message or signal, such as from the relational analysis unit **3700**, indicating a request for information describing changes to the ontology (ontological updates request). The semantic interface unit **3600** may generate and send, or otherwise make available, a response message or signal to the relational analysis unit **3700** indicating the changes to the ontology (ontological updates response). The semantic interface unit **3600** may identify one or more data portions for indexing based on the changes to the ontology. For example, the changes to the ontology may include adding a table to the ontology, the table including multiple rows, and the semantic interface unit **3600** may identify each row as a data portion for indexing. The semantic interface unit **3600** may include information describing the ontological changes in the ontological updates response. The semantic interface unit **3600** may include one or more data-query definitions, such as data-query definitions for indexing data queries, for each data portion identified for indexing in the ontological updates response. For example, the data-query definitions may include a sampling data query, which may be used to query the distributed in-memory database **3300**, or an external database, for sample data from the added data portion, an indexing data query, which may be used to query the distributed in-memory database **3300**, or an external database, for data from the added data portion, or both.

[0109] The semantic interface unit **3600** may receive, or otherwise access, internal signals or messages including data expressing usage intent, such as data indicating requests to access or modify the low-latency data stored in the distributed in-memory database **3300** (e.g., a request for data). The request to access or modify the low-latency data received by the semantic interface unit **3600** may include a resolved request (resolved-request data), such as in a resolved-request object, such as a resolved-request object generated by the relational analysis unit **3700**. The resolved request data, which may be database and visualization agnostic, may be expressed or communicated as an ordered sequence of tokens, which may represent semantic data.

[0110] The resolved-request data may include tokenization binding data. The tokenization binding data corresponding to a respective token may include, for example, one or more of a column identifier indicating a column corresponding to the respective token, a data type identifier corresponding to the respective token, a table identifier indicating a table corresponding to the respective token, an indication of an aggregation corresponding to the respective token, or an indication of a join path associated with the respective token. Other tokenization binding data may be used.

[0111] The resolved-request data may include phrasing data indicating phrasing with respect to the sequence of tokens in the resolved request, wherein tokens, such as one or more sequential tokens, are included in a respective phrase. The phrasing data may include phrase type data for respective phrases. For some tokens, or sequences of tokens, the phrasing data may indicate that the sequence of tokens corresponds with a value stored in a data source, such as in a column in a table, wherein the phrasing data includes data uniquely identifying the data source, such as a column identifier.

[0112] A token is a unit of data in the low-latency data access and analysis system **3000** that represents, in accordance with one or more defined grammars implemented by the low-latency data access and analysis system **3000**, a data portion accessed by or stored in the low-latency data access and analysis system **3000**, an operation of the low-latency data access and analysis system **3000**, an object represented in the low-latency data access and analysis system **3000**, or a class or type of data portion, operation, or object in the low-latency data access and analysis system **3000**. A token may be a value (token value), such as a string value, which may be a word, a character, a sequence of characters, a symbol, a combination of symbols, or the like. In some implementations, the token value may express a data pattern that defines or describes values, operations, or objects that the token represents. For example, the data pattern expressed by the token value may identify a data type, such as positive integer, such that positive integer values, or string values that may be represented as positive integer values, may be identified as matching the token. A token may be a defined data structure (token data structure) that includes a token value. A token data structure may include data other than the token value, such as token type data.

[0113] The defined grammars implemented by the low-latency data access and analysis system **3000** may define or describe the tokens. The defined grammars implemented by the low-latency data access and analysis system **3000** may define or describe token types or classes, such as ontological tokens, control-word tokens, pattern tokens, literal tokens, chronometric tokens, and a skip token. Other token types may be used.

[0114] An ontological token may represent a data portion in the low-latency data access and analysis system, such as an object represented in the low-latency data access and analysis system **3000**, or a portion thereof, a table stored in the distributed in-memory database or stored in an external database, a column of a table stored in the distributed in-memory database or stored in an external database, or a value (constituent data) stored in a row and column of a table stored in the distributed in-memory database or stored in an external database. In some grammars implemented by the

low-latency data access and analysis system **3000** the ontological tokens may include measure tokens representing measure data portions (measure columns), attribute tokens representing attribute data portions (attribute columns), and value tokens representing the respective values stored in the corresponding measure columns or attribute columns. For example, a worksheet object (analytical object) represented in the low-latency data access and analysis system **3000** may include a column that includes values generated based on values stored in one or more tables in the distributed in-memory database, and an ontological token may represent the column of the worksheet object.

[0115] A control-word token may be a character, a symbol, a word, or a defined ordered sequence of characters or symbols, defined or described in one or more grammars of the low-latency data access and analysis system **3000** as having one or more defined grammatical functions, which may be contextual. For example, the control-word token "sum" may be defined or described in one or more grammars of the low-latency data access and analysis system **3000** as indicating an additive aggregation. In another example, the control-word token "top" may be defined or described in one or more grammars of the low-latency data access and analysis system **3000** as indicating a maximal value from an ordered set. In another example, the control-word token "table" may be defined or described in one or more grammars of the low-latency data access and analysis system **3000** as indicating a table stored in the low-latency data access and analysis system **3000** or stored externally and accessed by the low-latency data access and analysis system **3000**. The control-word tokens may include operator tokens, such as the equality operator token ("="), delimiter tokens, which may be paired, such as opening and closing brackets ("[", "]"). The control-word tokens may include stop-word tokens, such as "the" or "an".

[0116] A pattern token may be a definition or a description of units of data in the low-latency data access and analysis system, which may be expressed as a data type, such as positive integer, defined or described in one or more grammars of the low-latency data access and analysis system **3000**.

[0117] A literal, or constant, token may include a literal, or constant, value such as "100" or the Boolean value TRUE. The literal, or constant, tokens may include number-word tokens (numerals or named numbers), such as number-word tokens for the positive integers between zero and one million, inclusive, or for the numerator, denominator, or both of fractional values, or combinations thereof. For example, "one hundred twenty-eight and three-fifths".

[0118] A chronometric token may represent a chronometric unit, such as a chronometric unit from the system-defined chronometry or a chronometric unit from a domain-specific chronometry defined or described in the low-latency data access and analysis system **3000**. The chronometric tokens are automatically generated based on the respective chronometric datasets. For example, chronometric tokens corresponding to the chronometric units for the system-defined chronometry, such as "date", "day", "days", "daily", "week", "weeks", "weekly", "month", "months", "monthly", "quarter", "quarters", "quarterly", "year", "years", "yearly", and the like, may be automatically generated based on the chronometric dataset for the system-defined chronometry.

[0119] The skip token may represent discrete data portions, such as respective portions of a string that are unresolvable in accordance with the other tokens defined or described in a respective grammar of the low-latency data access and analysis system **3000**.

[0120] The relational analysis unit **3700** may automatically generate respective tokens representing the attributes, the measures, the tables, the columns, the values, unique identifiers, tags, links, keys, or any other data portion, or combination of data portions, or a portion thereof.

[0121] For example, the relational analysis unit **3700** may tokenize, identify semantics, or both, based on input data, such as input data representing user input, to generate the resolved request. The resolved request may include an ordered sequence of tokens that represent the request for data corresponding to the input data, and may transmit, send, or otherwise make accessible, the resolved request to the semantic interface unit **3600**. The semantic interface unit **3600** may process or respond to a received resolved request.

[0122] The semantic interface unit **3600** may process or transform the received resolved request, which may be, at least in part, incompatible with the distributed in-memory database **3300**, or an external database, to generate one or more corresponding data-queries that are compatible with the distributed in-memory database **3300**, or an external database, which may include generating a proto-query representing the resolved request, generating a pseudo-query representing the proto-query, and generating the data query representing the pseudo-query.

[0123] The semantic interface unit **3600** may generate an analytical object, such as an answer object, representing the resolved request, which may include representing the data expressing usage intent, such as by representing the request for data indicated by the data expressing usage intent.

[0124] The semantic interface unit **3600** may generate a proto-query based on the resolved request. A proto-query, which may be database agnostic, may be structured or formatted in a form, language, or protocol that differs from the defined structured query language of the distributed in-memory database **3300**, or the external database. Generating the proto-query may include identifying visualization identification data, such as an indication of a type of visualization, associated with the request for data, and generating the proto-query based on the resolved request and the visualization identification data.

[0125] The semantic interface unit **3600** may transform the proto-query to generate a pseudo-query. The pseudo-query, which may be database agnostic, may be structured or formatted in a form, language, or protocol that differs from the defined structured query language of the distributed in-memory database **3300**, or the external database. Generating a pseudo-query may include applying a defined transformation, or an ordered sequence of transformations. Generating a pseudo-query may include incorporating row-level security filters in the pseudo-query.

[0126] The semantic interface unit **3600** may generate a data query based on the pseudo-query, such as by serializing the pseudo-query. The data query, or a portion thereof, may be structured or formatted using the defined structured query language of the distributed in-memory database **3300**. In some implementations, a data query may be structured or formatted using a defined structured query language of another database, such as an external database or data source, which may differ from the defined structured query

language of the distributed in-memory database **3300**. Generating the data query may include using one or more defined rules for expressing the structure and content of a pseudo-query in the respective defined structured query language.

[0127] The semantic interface unit **3600** may communicate, or issue, the data query to the distributed in-memory database **3300**, or the external database. In some implementations, processing or responding to a resolved request may include generating and issuing multiple data-queries to the distributed in-memory database **3300**, or the external database.

[0128] The semantic interface unit **3600** may receive results data from the distributed in-memory database **3300**, or the external database, responsive to one or more resolved requests. The semantic interface unit **3600** may process, format, or transform the results data to obtain visualization data. For example, the semantic interface unit **3600** may identify a visualization for representing or presenting the results data, or a portion thereof, such as based on the results data or a portion thereof. For example, the semantic interface unit **3600** may identify a bar chart visualization for results data including one measure and attribute.

[0129] Although not shown separately in FIG. **3**, the semantic interface unit **3600** may include a data visualization unit. In some embodiments, the data visualization unit may be a distinct unit, separate from the semantic interface unit **3600**. In some embodiments, the data visualization unit may be included in the system access interface unit **3900**. The data visualization unit, the system access interface unit **3900**, or a combination thereof, may generate a user interface, or one or more portions thereof. For example, data visualization unit, the system access interface unit **3900**, or a combination thereof, may obtain the results data, such as the visualization data, and may generate user interface elements (visualizations) representing the results data.

[0130] The semantic interface unit **3600** may implement object-level security, row-level security, or a combination thereof. In some implementations, the security and governance unit **3200** may implement, or partially implement, the object-level security, row-level security, or a combination thereof, in combination with the semantic interface unit **3600**. Object-level security may include security associated with an object, such as a table, a column, a worksheet, an answer, or a pinboard. The object-level security may include column-level security, which include user-based or group-based access control of columns of data in the low-latency data, the indexes, or both. Row-level security may include user-based or group-based access control of rows of data in the low-latency data, the indexes, or both. The semantic interface unit **3600** may implement one or more authentication procedures, access control procedures, or a combination thereof. The object-level security, row-level security, column-level security, a combination thereof, or a portion thereof, may be represented, expressed, defined, or described as access-control data. The semantic interface unit **3600**, or one or more other components of the low-latency data access and analysis system **3000**, may control, such as grant, restrict, or prevent, access to one or more features, functions, units of data, or combinations thereof, in accordance with the access-control data. For example, in response to a request for data that includes a user identifier, the semantic interface unit **3600**, or one or more other components of the low-latency data access and analysis system **3000**, may obtain access-control data for the user identifier and may obtain results data in accordance with the access-control data such that a unit of data, such as a row or a column, that is identified in the access-control data as accessible to the user identifier and is responsive to the request for data is included in the results data and such that a unit of data, such as a row or a column, that is identified in the access-control data as inaccessible to the user identifier, or for which the access-control data omits or excludes corresponding data indicating that the unit of data is accessible to the user identifier, is omitted or excluded from the results data.

[0131] The semantic interface unit **3600** may implement one or more user-data integration features. For example, the semantic interface unit **3600** may generate and output a user interface, or a portion thereof, for inputting, uploading, or importing user data, may receive user data, and may import the user data. For example, the user data may be enterprise data.

[0132] The semantic interface unit **3600** may implement object statistical data collection. Object statistical data may include, for respective objects, temporal access information, access frequency information, access recency information, access requester information, or the like. For example, the semantic interface unit **3600** may obtain object statistical data as described with respect to the data utility unit **3720**, the object utility unit **3810**, or both. The semantic interface unit **3600** may send, transmit, or otherwise make available, the object statistical data for data objects to the data utility unit **3720**. The semantic interface unit **3600** may send, transmit, or otherwise make available, the object statistical data for analytical objects to the object utility unit **3810**.

[0133] The semantic interface unit **3600** may implement or expose one or more services or application programming interfaces. For example, the semantic interface unit **3600** may implement one or more services for access by the system access interface unit **3900**. In some implementations, one or more services or application programming interfaces may be exposed to one or more external devices or systems.

[0134] The semantic interface unit **3600** may generate and transmit, send, or otherwise communicate, one or more external communications, such as e-mail messages, such as periodically, in response to one or more events, or both. For example, the semantic interface unit **3600** may generate and transmit, send, or otherwise communicate, one or more external communications including a portable representation, such as a portable document format representation of one or more pinboards in accordance with a defined schedule, period, or interval. In another example, the semantic interface unit **3600** may generate and transmit, send, or otherwise communicate, one or more external communications in response to input data indicating an express request for a communication. In another example, the semantic interface unit **3600** may generate and transmit, send, or otherwise communicate, one or more external communications in response to one or more defined events, such as the expiration of a recency of access period for a user.

[0135] Although shown as a single unit in FIG. **3**, the relational analysis unit **3700** may be implemented in a distributed configuration, which may include a primary relational analysis unit instance and one or more secondary relational analysis unit instances.

[0136] The relational analysis unit **3700** may generate, maintain, operate, or a combination thereof, one or more

indexes, such as one or more of an ontological index, a constituent data index, a control-word index, a numeral index, or a constant index, based on the low-latency data stored in the distributed in-memory database **3300**, or the external database, the low-latency data access and analysis system **3000**, or both. An index may be a defined data structure, or combination of data structures, for storing tokens, terms, or string keys, representing a set of data from one or more defined data sources in a form optimized for searching. For example, an index may be a collection of index shards. In some implementations, an index may be segmented into index segments and the index segments may be sharded into index shards. In some implementations, an index may be partitioned into index partitions, the index partitions may be segmented into index segments and the index segments may be sharded into index shards.

[0137] Generating, or building, an index may be performed to create or populate a previously unavailable index, which may be referred to as indexing the corresponding data, and may include regenerating, rebuilding, or reindexing to update or modify a previously available index, such as in response to a change in the indexed data (constituent data).

[0138] The ontological index may be an index of data (ontological data) describing the ontological structure or schema of the low-latency data access and analysis system **3000**, the low-latency data stored in the distributed in-memory database **3300**, or the external database, or a combination thereof. For example, the ontological index may include data representing the table and column structure of the distributed in-memory database **3300**, or the external database. The relational analysis unit **3700** may generate, maintain, or both, the ontological index by communicating with, such as requesting ontological data from, the distributed in-memory ontology unit **3500**, the semantic interface unit **3600**, or both. Each record in the ontological index may correspond to a respective ontological token, such as a token that identifies a column by name.

[0139] The control-word index may be an index of a defined set of control-word tokens. For example, the control-word index may include the control-word token "sum", which may be identified in one or more grammars of the low-latency data access and analysis system **3000** as indicating an additive aggregation. The constant index may be an index of constant, or literal, tokens such as "100" or "true". The numeral index may be an index of number word tokens (or named numbers), such as number word tokens for the positive integers between zero and one million, inclusive.

[0140] The constituent data index may be an index of the constituent data values stored in the low-latency data access and analysis system **3000**, or the external database, such as in the distributed in-memory database **3300**. The relational analysis unit **3700** may generate, maintain, or both, the constituent data index by communicating with, such as requesting data from, the distributed in-memory database **3300**, or the external database. For example, the relational analysis unit **3700** may send, or otherwise communicate, a message or signal to the distributed in-memory database **3300** indicating a request to perform an indexing data query, the relational analysis unit **3700** may receive response data from the distributed in-memory database **3300** in response to the requested indexing data query, and the relational analysis unit **3700** may generate the constituent data index, or a portion thereof, based on the response data. For example, the constituent data index may index data objects.

[0141] An index shard may be used for token searching, such as exact match searching, prefix match searching, substring match searching, or suffix match searching. Exact match searching may include identifying tokens in the index shard that matches a defined target value. Prefix match searching may include identifying tokens in the index shard that include a prefix, or begin with a value, such as a character or string, that matches a defined target value. Substring match searching may include identifying tokens in the index shard that include a value, such as a character or string, that matches a defined target value. Suffix match searching may include identifying tokens in the index shard that include a suffix, or end with a value, such as a character or string, that matches a defined target value. In some implementations, an index shard may include multiple distinct index data structures. For example, an index shard may include a first index data structure optimized for exact match searching, prefix match searching, and suffix match searching, and a second index data structure optimized for substring match searching. Traversing, or otherwise accessing, managing, or using, an index may include identifying one or more of the index shards of the index and traversing the respective index shards. In some implementations, one or more indexes, or index shards, may be distributed, such as replicated on multiple relational analysis unit instances. For example, the ontological index may be replicated on each relational analysis unit instance.

[0142] The relational analysis unit **3700** may receive a request for data from the low-latency data access and analysis system **3000**. For example, the relational analysis unit **3700** may receive data expressing usage intent indicating the request for data in response to input, such as user input, obtained via a user interface, such as a user interface generated, or partially generated, by the system access interface unit **3900**, which may be a user interface operated on an external device, such as one of the client devices **2320**, **2340** shown in FIG. **2**. In some implementations, the relational analysis unit **3700** may receive the data expressing usage intent from the system access interface unit **3900** or from the semantic interface unit **3600**. For example, the relational analysis unit **3700** may receive or access the data expressing usage intent in a request for data message or signal.

[0143] The relational analysis unit **3700** may process, parse, identify semantics, tokenize, or a combination thereof, the request for data to generate a resolved request, which may include identifying a database and visualization agnostic ordered sequence of tokens based on the data expressing usage intent. The data expressing usage intent, or request for data, may include request data, such as resolved-request data, unresolved request data, or a combination of resolved-request data and unresolved request data. The relational analysis unit **3700** may identify the resolved-request data. The relational analysis unit **3700** may identify the unresolved request data and may tokenize the unresolved request data.

[0144] Resolved-request data may be request data identified in the data expressing usage intent as resolved-request data. Each resolved-request data portion may correspond with a respective token in the low-latency data access and analysis system **3000**. The data expressing usage intent may

include information identifying one or more portions of the request data as resolved-request data.

[0145] Unresolved request data may be request data identified in the data expressing usage intent as unresolved request data, or request data for which the data expressing usage intent omits information identifying the request data as resolved-request data. Unresolved request data may include text or string data, which may include a character, sequence of characters, symbol, combination of symbols, word, sequence of words, phrase, or the like, for which information, such as tokenization binding data, identifying the text or string data as resolved-request data is absent or omitted from the request data. The data expressing usage intent may include information identifying one or more portions of the request data as unresolved request data. The data expressing usage intent may omit information identifying whether one or more portions of the request data are resolved-request data. The relational analysis unit **3700** may identify one or more portions of the request data for which the data expressing usage intent omits information identifying whether the one or more portions of the request data are resolved-request data as unresolved request data.

[0146] For example, the data expressing usage intent may include a request string and one or more indications that one or more portions of the request string are resolved-request data. One or more portions of the request string that are not identified as resolved-request data in the data expressing usage intent may be identified as unresolved request data. For example, the data expressing usage intent may include the request string "example text"; the data expressing usage intent may include information indicating that the first portion of the request string, "example", is resolved-request data; and the data expressing usage intent may omit information indicating that the second portion of the request string, "text", is resolved-request data.

[0147] The information identifying one or more portions of the request data as resolved-request data may include tokenization binding data indicating a previously identified token corresponding to the respective portion of the request data. The tokenization binding data corresponding to a respective token may include, for example, one or more of a column identifier indicating a column corresponding to the respective token, a data type identifier corresponding to the respective token, a table identifier indicating a table corresponding to the respective token, an indication of an aggregation corresponding to the respective token, or an indication of a join path associated with the respective token. Other tokenization binding data may be used. In some implementations, the data expressing usage intent may omit the tokenization binding data and may include an identifier that identifies the tokenization binding data.

[0148] The relational analysis unit **3700** may implement or access one or more grammar-specific tokenizers, such as a tokenizer for a defined data-analytics grammar or a tokenizer for a natural-language grammar. For example, the relational analysis unit **3700** may implement one or more of a formula tokenizer, a row-level-security tokenizer, a data-analytics tokenizer, or a natural language tokenizer. Other tokenizers may be used. In some implementations, the relational analysis unit **3700** may implement one or more of the grammar-specific tokenizers, or a portion thereof, by accessing another component of the low-latency data access and analysis system **3000** that implements the respective grammar-specific tokenizer, or a portion thereof. For example, the natural language processing unit **3710** may implement the natural language tokenizer and the relational analysis unit **3700** may access the natural language processing unit **3710** to implement natural language tokenization. In another example, the semantic interface **3600**, the database, or both, may implement a tokenizer for a grammar for the defined structured query language compatible with or implemented by the database. In some implementations, the low-latency data access and analysis system **3000**, such as the semantic interface **3600**, may implement a tokenizer for a grammar for a defined structured query language compatible with or implemented by an external database.

[0149] A tokenizer, such as the data-analytics tokenizer, may parse text or string data (request string), such as string data included in a data expressing usage intent, in a defined read order, such as from left to right, such as on a character-by-character or symbol-by-symbol basis. For example, a request string may include a single character, symbol, or letter, and tokenization may include identifying one or more tokens matching, or partially matching, the input character.

[0150] Tokenization may include parsing the request string to identify one or more words or phrases. For example, the request string may include a sequence of characters, symbols, or letters, and tokenization may include parsing the sequence of characters in a defined order, such as from left to right, to identify distinct words or terms and identifying one or more tokens matching the respective words. In some implementations, word or phrase parsing may be based on one or more of a set of defined delimiters, such as a whitespace character, a punctuation character, or a mathematical operator.

[0151] The relational analysis unit **3700** may traverse one or more of the indexes to identify one or more tokens corresponding to a character, word, or phrase identified in the request string. Tokenization may include identifying multiple candidate tokens matching a character, word, or phrase identified in the request string. Candidate tokens may be ranked or ordered, such as based on probabilistic utility.

[0152] Tokenization may include match-length maximization. Match-length maximization may include ranking or ordering candidate matching tokens in descending magnitude order. For example, the longest candidate token, having the largest cardinality of characters or symbols, matching the request string, or a portion thereof, may be the highest ranked candidate token. For example, the request string may include a sequence of words or a semantic phrase, and tokenization may include identifying one or more tokens matching the input semantic phrase. In another example, the request string may include a sequence of phrases, and tokenization may include identifying one or more tokens matching the input word sequence. In some implementations, tokenization may include identifying the highest ranked candidate token for a portion of the request string as a resolved token for the portion of the request string.

[0153] The relational analysis unit **3700** may implement one or more finite state machines. For example, tokenization may include using one or more finite state machines. A finite state machine may model or represent a defined set of states and a defined set of transitions between the states. A state may represent a condition of the system represented by the finite state machine at a defined temporal point. A finite state machine may transition from a state (current state) to a subsequent state in response to input (e.g., input to the finite state machine). A transition may define one or more actions

or operations that the relational analysis unit **3700** may implement. One or more of the finite state machines may be non-deterministic, such that the finite state machine may transition from a state to zero or more subsequent states.

[0154] The relational analysis unit **3700** may generate, instantiate, or operate a tokenization finite state machine, which may represent the respective tokenization grammar. Generating, instantiating, or operating a finite state machine may include operating a finite state machine traverser for traversing the finite state machine. Instantiating the tokenization finite state machine may include entering an empty state, indicating the absence of received input. The relational analysis unit **3700** may initiate or execute an operation, such as an entry operation, corresponding to the empty state in response to entering the empty state. Subsequently, the relational analysis unit **3700** may receive input data, and the tokenization finite state machine may transition from the empty state to a state corresponding to the received input data. In some embodiments, the relational analysis unit **3700** may initiate one or more data queries in response to transitioning to or from a respective state of a finite state machine. In the tokenization finite state machine, a state may represent a possible next token in the request string. The tokenization finite state machine may transition between states based on one or more defined transition weights, which may indicate a probability of transiting from a state to a subsequent state.

[0155] The tokenization finite state machine may determine tokenization based on probabilistic path utility. Probabilistic path utility may rank or order multiple candidate traversal paths for traversing the tokenization finite state machine based on the request string. The candidate paths may be ranked or ordered based on one or more defined probabilistic path utility metrics, which may be evaluated in a defined sequence. For example, the tokenization finite state machine may determine probabilistic path utility by evaluating the weights of the respective candidate transition paths, the lengths of the respective candidate transition paths, or a combination thereof. In some implementations, the weights of the respective candidate transition paths may be evaluated with high priority relative to the lengths of the respective candidate transition paths.

[0156] In some implementations, one or more transition paths evaluated by the tokenization finite state machine may include a bound state such that the candidate tokens available for tokenization of a portion of the request string may be limited based on the tokenization of a previously tokenized portion of the request string.

[0157] Tokenization may include matching a portion of the request string to one or more token types, such as a constant token type, a column name token type, a value token type, a control-word token type, a date value token type, a string value token type, or any other token type defined by the low-latency data access and analysis system **3000**. A constant token type may be a fixed, or invariant, token type, such as a numeric value. A column name token type may correspond with a name of a column in the data model. A value token type may correspond with an indexed data value. A control-word token type may correspond with a defined set of control-words. A date value token type may be similar to a control-word token type and may correspond with a defined set of control-words for describing temporal information. A string value token type may correspond with an unindexed value.

[0158] Token matching may include ordering or weighting candidate token matches based on one or more token matching metrics. Token matching metrics may include whether a candidate match is within a defined data scope, such as a defined set of tables, wherein a candidate match outside the defined data scope (out-of-scope) may be ordered or weighted lower than a candidate match within the defined data scope (in-scope). Token matching metrics may include whether, or the degree to which, a candidate match increases query complexity, such as by spanning multiple roots, wherein a candidate match that increases complexity may be ordered or weighted lower than a candidate match that does not increase complexity or increases complexity to a lesser extent. Token matching metrics may include whether the candidate match is an exact match or a partial match, wherein a candidate match that is a partial may be ordered or weighted lower than a candidate match that is an exact match. In some implementations, the cardinality of the set of partial matches may be limited to a defined value.

[0159] Token matching metrics may include a token score (TokenScore), wherein a candidate match with a relatively low token score may be ordered or weighted lower than a candidate match with a relatively high token score. The token score for a candidate match may be determined based on one or more token scoring metrics. The token scoring metrics may include a finite state machine transition weight metric (FSMScore), wherein a weight of transitioning from a current state of the tokenization finite state machine to a state indicating a candidate matching token is the finite state machine transition weight metric. The token scoring metrics may include a cardinality penalty metric (CardinalityScore), wherein a cardinality of values (e.g., unique values) corresponding to the candidate matching token is used as a penalty metric (inverse cardinality), which may reduce the token score. The token scoring metrics may include an index utility metric (IndexScore), wherein a defined utility value, such as one, associated with an object, such as a column wherein the matching token represents the column or a value from the column, is the index utility metric. In some implementations, the defined utility values may be configured, such as in response to user input, on a per object (e.g., per column) basis. The token scoring metrics may include a usage metric (UBRScore). The usage metric may be determined based on a usage based ranking index, one or more usage ranking metrics, or a combination thereof. Determining the usage metric (UBRScore) may include determining a usage boost value (UBRBoost). The token score may be determined based on a defined combination of token scoring metrics. For example, determining the token score may be expressed as the following:

$$TokenScore = FSMScore * (IndexScore + UBRScore * UBRBoost) + \mathrm{Min}(CardinalityScore, 1).$$

[0160] Token matching may include grouping candidate token matches by match type, ranking or ordering on a per-match type basis based on token score, and ranking or ordering the match types. For example, the match types may include a first match type for exact matches (having the highest match type priority order), a second match type for prefix matches on ontological data (having a match type priority order lower than the first match type), a third match

type for substring matches on ontological data and prefix matches on data values (having a match type priority order lower than the second match type), a fourth match type for substring matches on data values (having a match type priority order lower than the third match type), and a fifth match type for matches omitted from the first through fourth match types (having a match type priority order lower than the fourth match type). Other match types and match type orders may be used.

[0161] Tokenization may include ambiguity resolution. Ambiguity resolution may include token ambiguity resolution, join-path ambiguity resolution, or both. In some implementations, ambiguity resolution may cease tokenization in response to the identification of an automatic ambiguity resolution error or failure.

[0162] Token ambiguity may correspond with identifying two or more exactly matching candidate matching tokens. Token ambiguity resolution may be based on one or more token ambiguity resolution metrics. The token ambiguity resolution metrics may include using available previously resolved token matching or binding data and token ambiguity may be resolved in favor of available previously resolved token matching or binding data, other relevant tokens resolved from the request string, or both. The token ambiguity resolution may include resolving token ambiguity in favor of integer constants. The token ambiguity resolution may include resolving token ambiguity in favor of control-words, such as for tokens at the end of a request for data, such as last, that are not being edited.

[0163] Join-path ambiguity may correspond with identifying matching tokens having two or more candidate join paths. Join-path ambiguity resolution may be based on one or more join-path ambiguity resolution metrics. The join-path ambiguity resolution metrics may include using available previously resolved join-path binding data and join-path ambiguity may be resolved in favor of available previously resolved join-paths. The join-path ambiguity resolution may include favoring join paths that include in-scope objects over join paths that include out-of-scope objects. The join-path ambiguity resolution metrics may include a complexity minimization metric, which may favor a join path that omits or avoids increasing complexity over join paths that increase complexity, such as a join path that may introduce a chasm trap.

[0164] The relational analysis unit **3700** may identify a resolved request based on the request string. The resolved request, which may be database and visualization agnostic, may be expressed or communicated as an ordered sequence of tokens representing the request for data indicated by the request string. The relational analysis unit **3700** may instantiate, or generate, one or more resolved-request objects. For example, the relational analysis unit **3700** may create or store a resolved-request object corresponding to the resolved request in the distributed in-memory ontology unit **3500**. The relational analysis unit **3700** may transmit, send, or otherwise make available, the resolved request to the semantic interface unit **3600**.

[0165] In some implementations, the relational analysis unit **3700** may transmit, send, or otherwise make available, one or more resolved requests, or portions thereof, to the semantic interface unit **3600** in response to finite state machine transitions. For example, the relational analysis unit **3700** may instantiate a data-analysis object in response to a first transition of a finite state machine. The relational analysis unit **3700** may include a first data-analysis object instruction in the data-analysis object in response to a second transition of the finite state machine. The relational analysis unit **3700** may send the data-analysis object including the first data-analysis object instruction to the semantic interface unit **3600** in response to the second transition of the finite state machine. The relational analysis unit **3700** may include a second data-analysis object instruction in the data-analysis object in response to a third transition of the finite state machine. The relational analysis unit **3700** may send the data-analysis object including the data-analysis object instruction, or a combination of the first data-analysis object instruction and the second data-analysis object instruction, to the semantic interface unit **3600** in response to the third transition of the finite state machine. The data-analysis object instructions may be represented using any annotation, instruction, text, message, list, pseudo-code, comment, or the like, or any combination thereof that may be converted, transcoded, or translated into structured data-analysis instructions for accessing, retrieving, analyzing, or a combination thereof, data from the low-latency data, which may include generating data based on the low-latency data.

[0166] The relational analysis unit **3700** may provide an interface to permit the creation of user-defined syntax. For example, a user may associate a string with one or more tokens. Accordingly, when the string is entered, the pre-associated tokens are returned in lieu of searching for tokens to match the input.

[0167] The relational analysis unit **3700** may include a localization unit (not expressly shown). The localization, globalization, regionalization, or internationalization, unit may obtain source data expressed in accordance with a source expressive-form and may output destination data representing the source data, or a portion thereof, and expressed using a destination expressive-form. The data expressive-forms, such as the source expressive-form and the destination expressive-form, may include regional or customary forms of expression, such as numeric expression, temporal expression, currency expression, alphabets, natural-language elements, measurements, or the like. For example, the source expressive-form may be expressed using a canonical-form, which may include using a natural-language, which may be based on English, and the destination expressive-form may be expressed using a locale-specific form, which may include using another natural-language, which may be a natural-language that differs from the canonical-language. In another example, the destination expressive-form and the source expressive-form may be locale-specific expressive-forms and outputting the destination expressive-form representation of the source expressive-form data may include obtaining a canonical-form representation of the source expressive-form data and obtaining the destination expressive-form representation based on the canonical-form representation. Although, for simplicity and clarity, the grammars described herein, such as the data-analytics grammar and the natural language search grammar, are described with relation to the canonical expressive-form, the implementation of the respective grammars, or portions thereof, described herein may implement locale-specific expressive-forms. For example, the data-analytics tokenizer may include multiple locale-specific data-analytics tokenizers.

[0168] The natural language processing unit **3710** may receive input data including a natural language string, such as a natural language string generated in accordance with user input. The natural language string may represent a data request expressed in an unrestricted natural language form, for which data identified or obtained prior to, or in conjunction with, receiving the natural language string by the natural language processing unit **3710** indicating the semantic structure, correlation to the low-latency data access and analysis system **3000**, or both, for at least a portion of the natural language string is unavailable or incomplete. Although not shown separately in FIG. **3**, in some implementations, the natural language string may be generated or determined based on processing an analog signal, or a digital representation thereof, such as an audio stream or recording or a video stream or recording, which may include using speech-to-text conversion.

[0169] The natural language processing unit **3710** may analyze, process, or evaluate the natural language string, or a portion thereof, to generate or determine the semantic structure, correlation to the low-latency data access and analysis system **3000**, or both, for at least a portion of the natural language string. For example, the natural language processing unit **3710** may identify one or more words or terms in the natural language string and may correlate the identified words to tokens defined in the low-latency data access and analysis system **3000**. In another example, the natural language processing unit **3710** may identify a semantic structure for the natural language string, or a portion thereof. In another example, the natural language processing unit **3710** may identify a probabilistic intent for the natural language string, or a portion thereof, which may correspond to an operative feature of the low-latency data access and analysis system **3000**, such as retrieving data from the internal data, analyzing data the internal data, or modifying the internal data.

[0170] The natural language processing unit **3710** may send, transmit, or otherwise communicate request data indicating the tokens, relationships, semantic data, probabilistic intent, or a combination thereof or one or more portions thereof, identified based on a natural language string to the relational analysis unit **3700**.

[0171] The data utility unit **3720** may receive, process, and maintain user-agnostic utility data, such as system configuration data, user-specific utility data, such as utilization data, or both user-agnostic and user-specific utility data. The utility data may indicate whether a data portion, such as a column, a record, an insight, or any other data portion, has high utility or low utility within the system, such as among the users of the system. For example, the utility data may indicate that a defined column is a high-utility column or a low-utility column. The data utility unit **3720** may store the utility data, such as using the low-latency data structure. For example, in response to a user using, or accessing, a data portion, data utility unit **3720** may store utility data indicating the usage, or access, event for the data portion, which may include incrementing a usage event counter associated with the data portion. In some embodiments, the data utility unit **3720** may receive the information indicating the usage, or access, event for the data portion from the insight unit **3730**, and the usage, or access, event for the data portion may indicate that the usage is associated with an insight.

[0172] As used herein, the term "utility" refers to a computer accessible data value, or values, representative of the usefulness of an aspect of the low-latency data access and analysis system, such as a data portion, an object, or a component of the low-latency data access and analysis system with respect to improving the efficiency, accuracy, or both, of the low-latency data access and analysis system. Unless otherwise expressly indicated, or otherwise clear from context, utility is relative within a defined data-domain or scope. For example, the utility of an object with respect to a user may be high relative to the utility of other objects with respect to the user. Express utility indicates expressly specified, defined, or configured utility, such as user or system defined utility. Probabilistic utility indicates utility calculated or determined using utility data and expresses a statistical probability of usefulness for a respective aspect of the low-latency data access and analysis system. Unless otherwise expressly indicated, or otherwise clear from context, utility is access context specific. For example, the utility of an object with respect to the access context of a user may be high relative to the utility of the object with respect to the respective access contexts of other users.

[0173] The data utility unit **3720** may receive a signal, message, or other communication, indicating a request for utility information. The request for utility information may indicate an object or data portion. The data utility unit **3720** may determine, identify, or obtain utility data associated with the identified object or data portion. The data utility unit **3720** may generate and send utility response data responsive to the request that may indicate the utility data associated with the identified object or data portion.

[0174] The data utility unit **3720** may generate, maintain, operate, or a combination thereof, one or more indexes, such as one or more of a usage (or utility) index, a resolved-request index, or a phrase index, based on the low-latency data stored in the distributed in-memory database **3300**, or the external database, the low-latency data access and analysis system **3000**, or both.

[0175] The insight unit **3730** may automatically identify one or more insights, which may be data other than data expressly requested by a user, and which may be identified and prioritized, or both, based on probabilistic utility.

[0176] The object search unit **3800** may generate, maintain, operate, or a combination thereof, one or more object indexes, which may be based on the analytical objects represented in the low-latency data access and analysis system **3000**, or a portion thereof, such as pinboards, answers, and worksheets. An object index may be a defined data structure, or combination of data structures, for storing analytical-object data in a form optimized for searching. Although shown as a single unit in FIG. **3**, the object search unit **3800** may interface with a distinct, separate, object indexing unit (not expressly shown).

[0177] The object search unit **3800** may include an object-index population interface, an object-index search interface, or both. The object-index population interface may obtain and store, load, or populate analytical-object data, or a portion thereof, in the object indexes. The object-index search interface may efficiently access or retrieve analytical-object data from the object indexes such as by searching or traversing the object indexes, or one or more portions thereof. In some implementations, the object-index population interface, or a portion thereof, may be a distinct, independent unit.

[0178] The object-index population interface may populate, update, or both the object indexes, such as periodically,

such as in accordance with a defined temporal period, such as thirty minutes. Populating, or updating, the object indexes may include obtaining object indexing data for indexing the analytical objects represented in the low-latency data access and analysis system **3000**. For example, the object-index population interface may obtain the analytical-object indexing data, such as from the distributed in-memory ontology unit **3500**. Populating, or updating, the object indexes may include generating or creating an indexing data structure representing an object. The indexing data structure for representing an object may differ from the data structure used for representing the object in other components of the low-latency data access and analysis system **3000**, such as in the distributed in-memory ontology unit **3500**.

[0179] The object indexing data for an analytical object may be a subset of the object data for the analytical object. The object indexing data for an analytical object may include an object identifier for the analytical object uniquely identifying the analytical object in the low-latency data access and analysis system **3000**, or in a defined data-domain within the low-latency data access and analysis system **3000**. The low-latency data access and analysis system **3000** may uniquely, unambiguously, distinguish an object from other objects based on the object identifier associated with the object. The object indexing data for an analytical object may include data non-uniquely identifying the object. The low-latency data access and analysis system **3000** may identify one or more analytical objects based on the non-uniquely identifying data associated with the respective objects, or one or more portions thereof. In some implementations, an object identifier may be an ordered combination of non-uniquely identifying object data that, as expressed in the ordered combination, is uniquely identifying. The low-latency data access and analysis system **3000** may enforce the uniqueness of the object identifiers.

[0180] Populating, or updating, the object indexes may include indexing the analytical object by including or storing the object indexing data in the object indexes. For example, the object indexing data may include data for an analytical object, the object indexes may omit data for the analytical object, and the object-index population interface may include or store the object indexing data in an object index. In another example, the object indexing data may include data for an analytical object, the object indexes may include data for the analytical object, and the object-index population interface may update the object indexing data for the analytical object in the object indexes in accordance with the object indexing data.

[0181] Populating, or updating, the object indexes may include obtaining object utility data for the analytical objects represented in the low-latency data access and analysis system **3000**. For example, the object-index population interface may obtain the object utility data, such as from the object utility unit **3810**. The object-index population interface may include the object utility data in the object indexes in association with the corresponding objects.

[0182] In some implementations, the object-index population interface may receive, obtain, or otherwise access the object utility data from a distinct, independent, object utility data population unit, which may read, obtain, or otherwise access object utility data from the object utility unit **3810** and may send, transmit, or otherwise provide, the object utility data to the object search unit **3800**. The object utility data population unit may send, transmit, or otherwise provide, the object utility data to the object search unit **3800** periodically, such as in accordance with a defined temporal period, such as thirty minutes.

[0183] The object-index search interface may receive, access, or otherwise obtain data expressing usage intent with respect to the low-latency data access and analysis system **3000**, which may represent a request to access data in the low-latency data access and analysis system **3000**, which may represent a request to access one or more analytical objects represented in the low-latency data access and analysis system **3000**. The object-index search interface may generate one or more object-index queries based on the data expressing usage intent. The object-index search interface may send, transmit, or otherwise make available the object-index queries to one or more of the object indexes.

[0184] The object-index search interface may receive, obtain, or otherwise access object search results data indicating one or more analytical objects identified by searching or traversing the object indexes in accordance with the object-index queries. The object-index search interface may sort or rank the object search results data based on probabilistic utility in accordance with the object utility data for the analytical objects in the object search results data. In some implementations, the object-index search interface may include one or more object search ranking metrics with the object-index queries and may receive the object search results data sorted or ranked based on probabilistic utility in accordance with the object utility data for the objects in the object search results data and in accordance with the object search ranking metrics.

[0185] For example, the data expressing usage intent may include a user identifier, and the object search results data may include object search results data sorted or ranked based on probabilistic utility for the user. In another example, the data expressing usage intent may include a user identifier and one or more search terms, and the object search results data may include object search results data sorted or ranked based on probabilistic utility for the user identified by searching or traversing the object indexes in accordance with the search terms.

[0186] The object-index search interface may generate and send, transmit, or otherwise make available the sorted or ranked object search results data to another component of the low-latency data access and analysis system **3000**, such as for further processing and display to the user.

[0187] The object utility unit **3810** may receive, process, and maintain user-specific object utility data for objects represented in the low-latency data access and analysis system **3000**. The user-specific object utility data may indicate whether an object has high utility or low utility for the user.

[0188] The object utility unit **3810** may store the user-specific object utility data, such as on a per-object basis, a per-activity basis, or both. For example, in response to data indicating an object access activity, such as a user using, viewing, or otherwise accessing, an object, the object utility unit **3810** may store user-specific object utility data indicating the object access activity for the object, which may include incrementing an object access activity counter associated with the object, which may be a user-specific object access activity counter. In another example, in response to data indicating an object storage activity, such as a user storing an object, the object utility unit **3810** may store user-specific object utility data indicating the object storage

activity for the object, which may include incrementing a storage activity counter associated with the object, which may be a user-specific object storage activity counter. The user-specific object utility data may include temporal information, such as a temporal location identifier associated with the object activity. Other information associated with the object activity may be included in the object utility data.

[0189] The object utility unit **3810** may receive a signal, message, or other communication, indicating a request for object utility information. The request for object utility information may indicate one or more objects, one or more users, one or more activities, temporal information, or a combination thereof. The request for object utility information may indicate a request for object utility data, object utility counter data, or both.

[0190] The object utility unit **3810** may determine, identify, or obtain object utility data in accordance with the request for object utility information. The object utility unit **3810** may generate and send object utility response data responsive to the request that may indicate the object utility data, or a portion thereof, in accordance with the request for object utility information.

[0191] For example, a request for object utility information may indicate a user, an object, temporal information, such as information indicating a temporal span, and an object activity, such as the object access activity. The request for object utility information may indicate a request for object utility counter data. The object utility unit **3810** may determine, identify, or obtain object utility counter data associated with the user, the object, and the object activity having a temporal location within the temporal span, and the object utility unit **3810** may generate and send object utility response data including the identified object utility counter data.

[0192] In some implementations, a request for object utility information may indicate multiple users, or may omit indicating a user, and the object utility unit **3810** may identify user-agnostic object utility data aggregating the user-specific object utility data. In some implementations, a request for object utility information may indicate multiple objects, may omit indicating an object, or may indicate an object type, such as answer, pinboard, or worksheet, and the object utility unit **3810** may identify the object utility data by aggregating the object utility data for multiple objects in accordance with the request. Other object utility aggregations may be used.

[0193] The system configuration unit **3820** implements or applies one or more low-latency data access and analysis system configurations to enable, disable, or configure one or more operative features of the low-latency data access and analysis system **3000**. The system configuration unit **3820** may store data representing or defining the one or more low-latency data access and analysis system configurations. The system configuration unit **3820** may receive signals or messages indicating input data, such as input data generated via a system access interface, such as a user interface, for accessing or modifying the low-latency data access and analysis system configurations. The system configuration unit **3820** may generate, modify, delete, or otherwise maintain the low-latency data access and analysis system configurations, such as in response to the input data. The system configuration unit **3820** may generate or determine output data, and may output the output data, for a system access interface, or a portion or portions thereof, for the low-latency data access and analysis system configurations, such as for presenting a user interface for the low-latency data access and analysis system configurations. Although not shown in FIG. **3**, the system configuration unit **3820** may communicate with a repository, such as an external centralized repository, of low-latency data access and analysis system configurations; the system configuration unit **3820** may receive one or more low-latency data access and analysis system configurations from the repository, and may control or configure one or more operative features of the low-latency data access and analysis system **3000** in response to receiving one or more low-latency data access and analysis system configurations from the repository.

[0194] The user customization unit **3830** may receive, process, and maintain user-specific utility data, user defined configuration data, user defined preference data, or a combination thereof. The user-specific utility data may indicate whether a data portion, such as a column, a record, autonomous-analysis (autoanalysis) data, or any other data portion or object, has high utility or low utility to an identified user. For example, the user-specific utility data may indicate that a defined column is a high-utility column or a low-utility column. The user customization unit **3830** may store the user-specific utility data, such as using the low-latency data structure. The user-specific utility data may include, feedback data, such as feedback indicating user input expressly describing or representing the utility of a data portion or object in response to utilization of the data portion or object, such as positive feedback indicating high utility or negative feedback indicating low utility. The user customization unit **3830** may store the feedback in association with a user identifier. The user customization unit **3830** may store the feedback in association with the access context in which feedback was obtained. The user customization data, or a portion thereof, may be stored in an in-memory storage unit of the low-latency data access and analysis system. In some implementations, the user customization data, or a portion thereof, may be stored in the persistent storage unit **3930**.

[0195] The system access interface unit **3900** may interface with, or communicate with, a system access unit (not shown in FIG. **3**), which may be a client device, a user device, or another external device or system, or a combination thereof, to provide access to the internal data, features of the low-latency data access and analysis system **3000**, or a combination thereof. For example, the system access interface unit **3900** may receive signals, message, or other communications representing interactions with the internal data, such as data expressing usage intent and may output response messages, signals, or other communications responsive to the received requests.

[0196] The system access interface unit **3900** may generate data for presenting a user interface, or one or more portions thereof, for the low-latency data access and analysis system **3000**. For example, the system access interface unit **3900** may generate instructions for rendering, or otherwise presenting, the user interface, or one or more portions thereof and may transmit, or otherwise make available, the instructions for rendering, or otherwise presenting, the user interface, or one or more portions thereof to the system access unit, for presentation to a user of the system access unit. For example, the system access unit may present the user interface via a web browser or a web application and the instructions may be in the form of HyperText Markup Language (HTML), JavaScript, or the like.

[0197] In an example, the system access interface unit **3900** may include a data-analytics field user interface element in the user interface. The data-analytics field user interface element may be an unstructured string user input element or field. The system access unit may display the unstructured string user input element. The system access unit may receive input data, such as user input data, corresponding to the unstructured string user input element. The system access unit may transmit, or otherwise make available, the unstructured string user input to the system access interface unit **3900**. The user interface may include other user interface elements and the system access unit may transmit, or otherwise make available, other user input data to the system access interface unit **3900**.

[0198] The system access interface unit **3900** may obtain the user input data, such as the unstructured string, from the system access unit. The system access interface unit **3900** may transmit, or otherwise make available, the user input data to one or more of the other components of the low-latency data access and analysis system **3000**.

[0199] In some embodiments, the system access interface unit **3900** may obtain the unstructured string user input as a sequence of individual characters or symbols, and the system access interface unit **3900** may sequentially transmit, or otherwise make available, individual or groups of characters or symbols of the user input data to one or more of the other components of the low-latency data access and analysis system **3000**.

[0200] In some embodiments, system access interface unit **3900** may obtain the unstructured string user input as a sequence of individual characters or symbols, the system access interface unit **3900** may aggregate the sequence of individual characters or symbols, and may sequentially transmit, or otherwise make available, a current aggregation of the received user input data to one or more of the other components of the low-latency data access and analysis system **3000**, in response to receiving respective characters or symbols from the sequence, such as on a per-character or per-symbol basis.

[0201] The real-time collaboration unit **3910** may receive signals or messages representing input received in accordance with multiple users, or multiple system access devices, associated with a collaboration context or session, may output data, such as visualizations, generated or determined by the low-latency data access and analysis system **3000** to multiple users associated with the collaboration context or session, or both. The real-time collaboration unit **3910** may receive signals or messages representing input received in accordance with one or more users indicating a request to establish a collaboration context or session, and may generate, maintain, or modify collaboration data representing the collaboration context or session, such as a collaboration session identifier. The real-time collaboration unit **3910** may receive signals or messages representing input received in accordance with one or more users indicating a request to participate in, or otherwise associate with, a currently active collaboration context or session, and may associate the one or more users with the currently active collaboration context or session. In some implementations, the input, output, or both, of the real-time collaboration unit **3910** may include synchronization data, such as temporal data, that may be used to maintain synchronization, with respect to the collaboration context or session, among the low-latency data access and analysis system **3000** and one or more system access devices associated with, or otherwise accessing, the collaboration context or session.

[0202] The third-party integration unit **3920** may include an electronic communication interface, such as an application programming interface (API), for interfacing or communicating between an external, such as third party, application or system, and the low-latency data access and analysis system **3000**. For example, the third-party integration unit **3920** may include an electronic communication interface to transfer data between the low-latency data access and analysis system **3000** and one or more external applications or systems, such as by importing data into the low-latency data access and analysis system **3000** from the external applications or systems or exporting data from the low-latency data access and analysis system **3000** to the external applications or systems. For example, the third-party integration unit **3920** may include an electronic communication interface for electronic communication with an external exchange, transfer, load (ETL) system, which may import data into the low-latency data access and analysis system **3000** from an external data source or may export data from the low-latency data access and analysis system **3000** to an external data repository. In another example, the third-party integration unit **3920** may include an electronic communication interface for electronic communication with external machine learning analysis software, which may export data from the low-latency data access and analysis system **3000** to the external machine learning analysis software and may import data into the low-latency data access and analysis system **3000** from the external machine learning analysis software. The third-party integration unit **3920** may transfer data independent of, or in conjunction with, the system access interface unit **3900**, the enterprise data interface unit **3400**, or both.

[0203] The persistent storage unit **3930** may include an interface for storing data on, accessing data from, or both, one or more persistent data storage devices or systems. For example, the persistent storage unit **3930** may include one or more persistent data storage devices, such as the static memory **1200** shown in FIG. **1**. Although shown as a single unit in FIG. **3**, the persistent storage unit **3930** may include multiple components, such as in a distributed or clustered configuration. The persistent storage unit **3930** may include one or more internal interfaces, such as electronic communication or application programming interfaces, for receiving data from, sending data to, or both other components of the low-latency data access and analysis system **3000**. The persistent storage unit **3930** may include one or more external interfaces, such as electronic communication or application programming interfaces, for receiving data from, sending data to, or both, one or more external systems or devices, such as an external persistent storage system. For example, the persistent storage unit **3930** may include an internal interface for obtaining key-value tuple data from other components of the low-latency data access and analysis system **3000**, an external interface for sending the key-value tuple data to, or storing the key-value tuple data on, an external persistent storage system, an external interface for obtaining, or otherwise accessing, the key-value tuple data from the external persistent storage system, and an internal key-value tuple data for sending, or otherwise making available, the key-value tuple data to other components of the low-latency data access and analysis system **3000**. In another example, the persistent storage unit **3930**

may include a first external interface for storing data on, or obtaining data from, a first external persistent storage system, and a second external interface for storing data on, or obtaining data from, a second external persistent storage system.

[0204] FIG. **4** is a diagram of an example of natural language to query language transformation **4000** in a data access and analysis system. The data access and analysis system may be similar to the data access and analysis system **3000** shown in FIG. **3**, except as is described herein or as is otherwise clear from context. The data access and analysis system may implement natural language to query language transformation **4000**.

[0205] Natural language to query language transformation **4000** includes obtaining natural language data (at **4100**). The natural language data, or natural language input, is, or includes, text or string data included in data expressing usage intent, which may be generated by the low-latency data access and analysis system, or a component thereof, such as the system access interface unit **3900** shown in FIG. **3**, such as in response to input, such as user input, obtained by the low-latency data access and analysis system. The natural language data (natural language input data) expresses a request for data analysis with respect to data stored in a data source **4020** of the data access and analysis system, such as data from, or based on, a worksheet **4010** that is based on, or associated with, the data source **4020**. For example, the natural language data may include the string "How much did it rain in the northeast in the last three months?".

[0206] The data source **4020** may be a distributed in-memory database, such as the distributed in-memory database **3300** shown in FIG. **3**, an external database, or another data source as described herein. The data source **4010** includes, or stores, one or more tables **4022**. The data source **4020** may be, include, or store, complex data, such that the data access, analysis, or both, may be subject to complexities such as chasm-traps, fan-traps, or both. In the absence of data, program logic, or both, indicating join types, cardinality, and the like, automatic identification of joins may be limited, inaccurate, or unavailable.

[0207] In one or more examples described herein, the data source **4020** is a meteorological database that includes meteorological, or weather, data stored in the tables **4022**.

[0208] The worksheet **4010** describes, or defines, an aggregation, or collection, of data from, based on, or a combination thereof, the tables **4022** from the data source **4020**. The worksheet **4010** defines, or describes, the data aggregation as including columns that correspond with columns from the tables **4022**, columns defined, or described, for data generated in accordance with the worksheet, or both.

[0209] For example, the tables **4022** include a first table (T**1**), a second table (T**2**), a third table (T**3**), a fourth table (T**4**), a fifth table (T**5**), a sixth table (T**6**), a seventh table (T**7**), an eighth table (T**8**), a ninth table (T**9**), and a tenth table (T**10**). The tables **4022** include a number, count, or cardinality, of columns, such as 4500 columns. The worksheet **4010** includes a first aggregation of data based on the tables **4022** (T**1**-T**10**) and a second worksheet (not shown) may include a second aggregation of data based on the tables **4022** (T**1**-T**10**), wherein the second worksheet (not shown) differs from the worksheet **4010**, such as based on joins, column names, data aggregations, or the like. The worksheet **4010** may include a number, count, or cardinality, of columns, such as 5000 columns.

[0210] The worksheet **4010** may be previously, such as prior to a current performance of natural language to query language transformation **4000** in the data access and analysis system, defined, such as manually defined, in the data access and analysis system. The worksheet **4010** is represented in the data access and analysis system by a worksheet object, which is a queryable object. Although not shown separately in FIG. **4**, obtaining the natural language data (at **4100**) may include identifying, or obtaining, the worksheet **4010**, which may include obtaining the queryable object (worksheet object) representing the worksheet in the data access and analysis system. In some implementations, the worksheet **4010**, the corresponding worksheet object, or both, may be identified prior to obtaining the natural language data (at **4100**).

[0211] Natural language to query language transformation **4000** in the data access and analysis system may include using one or more machine learning, or artificial intelligence, models, such as one or more large language models, which may be internal, such as implemented by the data access and analysis system, or external, such as accessible, or accessed, by the data access and analysis system. For simplicity, FIG. **4** shows one large language model **4030**.

[0212] Natural language to query language transformation **4000** in the data access and analysis system includes obtaining large language model input data (at **4200**). Obtaining the large language model input data (at **4200**) includes the data access and analysis system, or a component thereof, including the natural language input data (obtained at **4100**) in the large language model input data.

[0213] The large language model **4030** may be subject to limitations, such as a limit on a number, count, or cardinality, of tokens that may be included in input to the large language model **4030** (large language model input data) (at **4200**). For example, the large language model **4030** may be limited to a defined maximum number, count, or cardinality of tokens, such as **4000** tokens. Tokens of the large language model **4030** may differ from tokens of the data access and analysis system. A token of the large language model **4030** is an ordered sequence of characters, or symbols, that represent, or form, a word, a word part, whitespace, punctuation, or a combination thereof.

[0214] The data source **4020**, or the corresponding data model, including the tables **4022** and relationships among the tables **4022**, may be incompatible, or incompletely compatible, with the large language model **4030**. For example, the data source **4020**, or the corresponding data model, including the tables **4022** and relationships among the tables **4022**, may correspond with more than the maximum number, count, or cardinality of tokens.

[0215] The worksheet **4010** may be incompatible, or incompletely compatible, with the large language model **4030**. For example, the worksheet **4010** may include more columns than the maximum number, count, or cardinality of tokens, wherein a column corresponds with one or more tokens with respect to the large language model **4030**.

[0216] Obtaining the large language model input data (at **4200**) includes the data access and analysis system, or a component thereof, obtaining prompt signifier data, which is natural language data, such as "Generate SQL given the question and table to answer the question correctly. Make sure only columns in the table provided are used in the

generated SQL." Obtaining the large language model input data (at **4200**) includes the data access and analysis system, or a component thereof, including the prompt signifier data in the large language model input data. The prompt signifier data indicates, to the large language model, the task assigned to, or requested of, the large language model.

[0217] Obtaining the large language model input data (at **4200**) includes the data access and analysis system, or a component thereof, obtaining prompt context data. The data access and analysis system, or a component thereof, obtains the prompt context data using the natural language input data, the worksheet, utility data, such as user-agnostic utility data, such as system configuration data, user-specific utility data, such as utilization data, or both user-agnostic and user-specific utility data, object ontological data, such as object relationship data, or a combination thereof. The object relationship data may be, or may include, aggregated data that indicates a number, count, or cardinality of objects in the data access and analysis system that are related to an object in the data access and analysis system. For example, a first object in the data access and analysis system may represent a column in the table **4022** in the data source **4020**, and the object relationship data may include data indicating a number, count, or cardinality, of analytical objects that include one or more references to the first object as a representation of the column. The prompt context data indicates, to the large language model, data for the large language model to use to respond to, or answer, the request indicated by the natural language input data (obtained at **4100**) as included in the large language model input data (obtained at **4200**).

[0218] The prompt context data includes a definition, or description, of a data structure, such as a table structure (table structure data). For example, the prompt context data may be a structured query language instruction to create a table. The table structure data includes an identifier, or name, for the data structure. The table structure data includes column data defining, or describing, one or more columns of the data structure. The column data includes an identifier, or name for a respective column and a corresponding data type for the respective column. The column data may be obtained in accordance with a defined maximum number, count, or cardinality of columns, such as two hundred (200) columns.

[0219] To obtain the table structure data, the data access and analysis system, or a component thereof, obtains a request hypothesis based on the natural language input. The data access and analysis system, or a component thereof, uses the request hypothesis to identify, or otherwise obtain, a first subset of the columns from the worksheet, wherein the first subset of the columns from the worksheet includes less than or equal to the defined maximum number, count, or cardinality of columns, such as fifty (50). The data access and analysis system, or a component thereof, includes the first subset of columns, including a column identifier, such as a column name, and a data type for the respective column, from the worksheet in the table structure data.

[0220] In some implementations, the first subset of columns includes less than the defined maximum number, count, or cardinality of columns and the data access and analysis system, or a component thereof, obtains, or identifies, a second subset of the columns from the worksheet, wherein the second subset of the columns from the worksheet includes less than or equal to a difference between the defined maximum number, count, or cardinality of columns and the number, count, or cardinality of the first subset of columns, which may be obtained by subtracting the number, count, or cardinality of the first subset of columns from the defined maximum number, count, or cardinality of columns. The data access and analysis system, or a component thereof, identifies, or otherwise obtains, the second subset of the columns from the worksheet using utility data, such as user-agnostic utility data, such as system configuration data, user-specific utility data, such as utilization data, or both user-agnostic and user-specific utility data (probabilistic utility data). The data access and analysis system, or a component thereof, identifies the second subset of the columns from the worksheet having the highest, maximal, or greatest, utility, as indicated by the utility data (descending utility order). The data access and analysis system, or a component thereof, includes the second subset of columns, including a column identifier, such as a column name, and a data type for the respective column, from the worksheet in the table structure data. The second subset of columns may be columns from the worksheet other than the columns included in the first subset of columns (non-overlapping). Obtaining the second subset of columns may be agnostic of (omit using) the natural language input.

[0221] In some implementations, the first subset of columns and the second subset of columns, in combination, includes less than the defined maximum number, count, or cardinality of columns and the data access and analysis system, or a component thereof, obtains, or identifies, a third subset of the columns from the worksheet, wherein the third subset of the columns from the worksheet includes less than or equal to a difference between the defined maximum number, count, or cardinality of columns and a sum of the number, count, or cardinality of the first subset of columns and the number, count, or cardinality of the second subset of columns, which may be obtained by subtracting the sum of the number, count, or cardinality of the first subset of columns and the number, count, or cardinality of the second subset of columns from the defined maximum number, count, or cardinality of columns. The data access and analysis system, or a component thereof, identifies the third subset of the columns from the worksheet using ontological data indicating a number, count, or cardinality, of previously generated analytical objects in the data access and analysis system that reference, or have a defined relationship to, the respective column (related analytical objects). The data access and analysis system, or a component thereof, identifies the third subset of the columns from the worksheet having the highest, maximal, or greatest, number, count, or cardinality, of related analytical objects (descending relations order). The data access and analysis system, or a component thereof, includes the third subset of columns, including a column identifier, such as a column name, and a data type for the respective column, from the worksheet in the table structure data. Obtaining the third subset of columns may be agnostic of (omit using) the natural language input.

[0222] In some implementations, the first subset of columns, the second subset of columns, and the third subset of columns, in combination, includes less than the defined maximum number, count, or cardinality of columns and the data access and analysis system, or a component thereof, obtains, or identifies, a fourth subset of the columns from the worksheet, wherein the fourth subset of the columns from the worksheet includes less than or equal to a difference

between the defined maximum number, count, or cardinality of columns and a sum of the number, count, or cardinality of the first subset of columns, the number, count, or cardinality of the second subset of columns, and the number, count, or cardinality of the third subset of columns, which may be obtained by subtracting the sum of the number, count, or cardinality of the first subset of columns, the number, count, or cardinality of the second subset of columns, and the number, count, or cardinality of the third subset of columns from the defined maximum number, count, or cardinality of columns. The data access and analysis system, or a component thereof, identifies the fourth subset of the columns from the worksheet randomly, or pseudo-randomly. The data access and analysis system, or a component thereof, includes the fourth subset of columns, including a column identifier, such as a column name, and a data type for the respective column, from the worksheet in the table structure data. Obtaining the fourth subset of columns may be agnostic of (omit using) the natural language input.

[0223] In some implementations, obtaining the first subset of columns may be omitted, wherein obtaining the table structure data includes one or more of obtaining the second subset of columns, obtaining the third subset of columns, or obtaining the fourth subset of columns. In some implementations, obtaining the second subset of columns may be omitted, wherein obtaining the table structure data includes one or more of obtaining the first subset of columns, obtaining the third subset of columns, or obtaining the fourth subset of columns. In some implementations, obtaining the third subset of columns may be omitted, wherein obtaining the table structure data includes one or more of obtaining the first subset of columns, obtaining the second subset of columns, or obtaining the fourth subset of columns. In some implementations, obtaining the fourth subset of columns may be omitted, wherein obtaining the table structure data includes one or more of obtaining the first subset of columns, obtaining the second subset of columns, or obtaining the third subset of columns.

[0224] To obtain the request hypothesis the data access and analysis system, or a component thereof, identifies one or more patterns in the natural language input data and generates one or more request hypotheses about the data. To obtain the request hypothesis the data access and analysis system, or a component thereof, may use a search algorithm, such as a greedy search algorithm, a beam search, a breadth-first search, a best fit search, or the like. In some implementations, the data access and analysis system, or a component thereof, uses a beam search to obtain the requested hypothesis. The beam search is a heuristic search algorithm used in natural language processing (NLP) to find a most likely sequence of words in a sentence. The beam search works by expanding the most promising partial sequences and keeping the best partial sequences.

[0225] To obtain the request hypothesis the data access and analysis system, or a component thereof, tokenizes the natural language input to obtain an ordered sequence of tokens having a number, count, or cardinality of tokens (N). For example, the natural language input data may be the string "how much did it rain yesterday" and the corresponding ordered sequence of tokens is "how", "much", "did", "it", "rain", "yesterday" and the cardinality of tokens (N) is six (N=6).

[0226] To obtain the request hypothesis the data access and analysis system, or a component thereof, processes, such as iteratively, token subsequences from the orders sequence of tokens, wherein the token subsequences are contiguous in the ordered sequence of tokens and have a length in a range from one token to a defined maximum token subsequence length (K), wherein the defined maximum token subsequence length (K) is the number, count, or cardinality of tokens in the respective token subsequence.

[0227] To obtain the token subsequences, the data access and analysis system, or a component thereof, obtains, or identifies, a sequentially first token from the ordered sequence of tokens, other than tokens previously identified as a sequentially first token for current natural language input, as a current token subsequence. For example, for the ordered sequence of tokens is "how", "much", "did", "it", "rain", "yesterday", the first sequentially first token is "how".

[0228] The data access and analysis system, or a component thereof, determines whether the current token subsequence matches a token, a portion of a token, or a combination of tokens defined, or described, in the data access and analysis system.

[0229] In some implementations, the data access and analysis system, or a component thereof, may determine that the current token subsequence is matching with a token, a portion of a token, or a combination of tokens defined, or described, in the data access and analysis system. In response to determining that the current token subsequence matches a token, a portion of a token, or a combination of tokens defined, or described, in the data access and analysis system, the data access and analysis system, or a component thereof, obtains, or identifies, a combination of the current token subsequence and a subsequent token as the current token subsequence. For example, in a first iteration, the current token subsequence may be identified as "how", "how" may be determined as matching a token defined, or described in the data access and analysis system, and the combination of the token "how" and the token "much" ("how much") may be identified as the current token subsequence.

[0230] In some implementations, the data access and analysis system, or a component thereof, may determine that the current token subsequence is non-matching with a token, a portion of a token, or a combination of tokens defined, or described, in the data access and analysis system. For example, a token, a portion of a token, or a combination of tokens matching the current token subsequence may be unavailable in the data access and analysis system. In response to determining that the current token subsequence is non-matching, the data access and analysis system, or a component thereof, obtains, or identifies, a sequentially first token from the ordered sequence of tokens, other than tokens previously identified as a sequentially first token for current natural language input, as the current token subsequence. For example, for the ordered sequence of tokens is "how", "much", "did", "it", "rain", "yesterday", wherein the first sequentially first token, "how", may be previously identified as a sequentially first token from the ordered sequence of tokens, the next sequentially first token, other than "how", is "much", and "much" is identified as the current token subsequence.

[0231] For example, for the ordered sequence of tokens is "how", "much", "did", "it", "rain", "yesterday", the token

subsequence "how" may be identified as matching, the token subsequence "how much" may be identified as matching, the token subsequence "how much did" may be identified as non-matching, the token subsequence "much" may be identified as matching, the token subsequence "much did" may be identified as matching, the token subsequence "much did it" may be identified as non-matching, the token subsequence "did" may be identified as non-matching, the token subsequence "did it" may be identified as matching, the token subsequence "did it rain" may be identified as non-matching, the token subsequence "it" may be identified as matching, the token subsequence "it rain" may be identified as non-matching, the token subsequence "rain" may be identified as matching, the token subsequence "rain yesterday" may be identified as non-matching, and the token subsequence "yesterday" may be identified as matching.

[0232] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a column name, a substring of a column name, a filter value, a substring of a filter value, a data value from a row of results data, or sample results data, for the worksheet, or a substring of a data value from a row of results data, or sample results data, for the worksheet. With respect to determining whether the current token subsequence matches, the worksheet may be referred to herein as a token matching repository.

[0233] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in an ontological data index, wherein the ontological data index includes column identifiers, or names, corresponding synonyms, and corresponding associations. With respect to determining whether the current token subsequence matches, the ontological data index may be referred to herein as a token matching repository.

[0234] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in a constituent data index, wherein the constituent data index includes constituent data value of the rows, cells, fields, or records, from the worksheet for the columns representing strings. With respect to determining whether the current token subsequence matches, the constituent data index may be referred to herein as a token matching repository.

[0235] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in a control-word index, wherein the control-word index includes control-word values, or keywords, which may be mathematical operators. With respect to determining whether the current token subsequence matches, the control-word index may be referred to herein as a token matching repository.

[0236] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in a constant index, wherein the constant index includes constant values defined in the data access and analysis system, such as "100" or "true". With respect to determining whether the current token subsequence matches, the constant index may be referred to herein as a token matching repository.

[0237] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in a numeral index, wherein the numeral index includes number word tokens (or named numbers), such as number word tokens for the positive integers between zero and one million, inclusive. With respect to determining whether the current token subsequence matches, the numeral index may be referred to herein as a token matching repository.

[0238] In some implementations, the data access and analysis system, or a component thereof, determines whether the current token subsequence matches a value in a chronometric dataset defined in the data access and analysis system or a chronometric phrase pattern defined in the data access and analysis system, such as "last 2 weeks". With respect to determining whether the current token subsequence matches, the chronometric datasets defined in the data access and analysis system and the chronometric phrase patterns defined in the data access and analysis system may be referred to herein as token matching repositories.

[0239] The data access and analysis system, or a component thereof, identifies, or otherwise obtains, token subsequence scores for the token subsequences that are identified as matching, such that a respective token subsequence score is obtained, calculated, or identified, for a respective token subsequence identified as matching. The token subsequence score is identified, determined, or otherwise obtained based on the corresponding token matching repository and the match quality (match quality metric or match type), such as exact match, substring match, or the like, wherein an exact match is assigned a token subsequence score that is higher than a substring match, and wherein a match in the ontological data index is assigned a score that is higher than a match in the constituent data index.

[0240] The data access and analysis system, or a component thereof, categorizes the respective matching token subsequences as a measure, such as a measure column, an attribute, such as an attribute column, a value, such as a value in a row, cell, field, or record for a column in the worksheet, a control-word, an operator, a numeral, a chronometric unit, a constant, a stop word, or a skip token.

[0241] The data access and analysis system, or a component thereof, performs the search, such as the beam search, using the matching token subsequences to obtain one or more request hypothesis, which includes identifying a subset of the matching token subsequences (candidate matching token subsequences), the subset of the matching token subsequences having a size, number, count, or cardinality of the defined maximum token subsequence length (K), including the matching token subsequences having a respective relatively high token subsequence score, wherein token subsequences having a respective relatively low token subsequence score are omitted from the subset of the matching token subsequences (descending token subsequence score order). The beam search identifies non-overlapping token subsequences from the matching token subsequences. In an example, the matching token subsequences "how" and "how much" are overlapping, the token subsequence score for the matching token subsequence "how" may be higher than the token subsequence score for the matching token subsequence "how much", the beam search may identify the matching token subsequence "how" as a candidate matching token subsequence, and may omit or exclude the matching token subsequence "how much" from the candidate matching token subsequences as overlapping with the candidate matching token subsequence "how". The defined maximum

token subsequence length (K) may be referred to as the beam width with respect to the beam search.

[0242] The data access and analysis system, or a component thereof, includes the size K subset of non-overlapping, token matching score maximized, matching token subsequences identified by the search in the request hypotheses.

[0243] The data access and analysis system, or a component thereof, includes the request hypotheses that correspond to columns in the worksheet, or filters on the columns in the worksheet, in the prompt context data.

[0244] In an example, including prompt context data that indicates that the token "clouds" corresponds to the token sequences "cirrus clouds", "cumulus clouds", and "stratus clouds", increases the probability that the large language model will interpret a request for data that includes the word "clouds" as indicating that the output should refer to a data element as containing, or including, the word "clouds", such as in combination with other words, and a lower probability as indicating that the output should refer to a data being equal to the word "clouds".

[0245] The data access and analysis system, or a component thereof, includes prompt signifier data, the prompt context data, and the natural language input data in the large language model input data (at **4200**).

[0246] The data access and analysis system, or a component thereof, sends, transmits, or otherwise makes available, the large language model input data to the large language model **4030** (at **4200**).

[0247] The large language model **4030** receives, reads, obtains, or otherwise accesses, the large language model input data and automatically generates corresponding, or resulting, large language model generated data, such as in response to receiving the large language model input data. The large language model **4030** outputs, sends, transmits, or otherwise makes available, the language model generated data to the data access and analysis system, or a component thereof.

[0248] In some implementations, the data access and analysis system, or a component thereof, obtains, or generates, demonstrations data (demonstrations) **4040**, or few-shot examples. In some implementations, the data access and analysis system, or a component thereof, includes the demonstrations data in the large language model input data (at **4200**). The demonstrations data **4040** is included in the large language model input data prior to the natural language input data. For example, the large language model input data may include the prompt signifier data, followed by the prompt context data, followed by the demonstrations data **4040**, followed by the natural language input data. Although described as included in the large language model input data, the demonstrations data **4040** may be sent, transmitted, or otherwise made available, to the large language model **4030** separately from the large language model input data and in association with sending, transmitting, or otherwise making available, the large language model input data to the large language model **4030**. The demonstrations data **4040** is shown in FIG. **4** using a broken line border to indicate that obtaining the demonstrations data **4040** and including the demonstrations data **4040** in the large language model input data (at **4200**) may be omitted.

[0249] The demonstrations data **4040** includes one or more demonstrations, wherein a demonstration includes a question portion and a corresponding answer portion. The question portion is expressed using previously obtained, such as stored in the data access and analysis system, or a component thereof, natural language input data. The answer portion is expressed using data previously obtained, such as stored in the data access and analysis system, or a component thereof, corresponding to the question portion, and that is similar to the structured query language used, or output, by the large language model. The large language model uses the demonstrations data **4040** to improve the accuracy of the large language model output data relative to generating the large language model output data in the absence of the demonstrations data **4040**. An example of obtaining demonstrations data **4040** is shown in FIG. **6**.

[0250] The data access and analysis system, or a component thereof, receives, reads, obtains, or otherwise accesses, large language model generated data (at **4300**), such as in response to the large language model input data. The large language model generated data is a representation of the natural language input data in a form other than natural language form, such as in a structured query language form. For example, the worksheet identifier, or name, may be "meteorologicalData", the worksheet may include a "precipitation" column and a "dateCaptued" column, the natural language input data may include the string "How much did it rain last week?", and the corresponding large language model generated data may include the string "select precipitation from meteorologicalData where dateCaptued="last week"". The language model generated data differs from the natural language input data and equivalently expresses the natural language input data, such as in accordance with the structured query language.

[0251] The data access and analysis system, or a component thereof, transforms the large language model generated data representing the natural language input data into resolved request data representing the natural language input data (at **4400**). For example, the large language model generated data may include the string "select precipitation from meteorologicalData where dateCaptued="last week"" and the corresponding resolved request data may be expressed as "[sum precipitation][dateCaptued=last week]". The resolved request data differs from the natural language input data and the language model generated data, and equivalently expresses the natural language input data and the language model generated data, such as in accordance with the defined data-analytics grammar implemented by the data access and analysis system.

[0252] The data access and analysis system, or a component thereof, transforms the resolved request data corresponding to the large language model generated data and expressing the natural language input data to obtain a data query in accordance with a defined structured query language implemented by the data source (at **4500**). In some implementations, the defined structured query language implemented by the data source differs from the structured query language used by the large language model to generate the large language model output data.

[0253] In some implementations, one or more portions of the large language model generated data may be incompatible with the defined data-analytics grammar implemented by the data access and analysis system. To reduce, or eliminate, incompatibilities between the large language model generated data and the defined data-analytics grammar, the data access and analysis system may automatically generate one or more automatically generated formulas to augment the defined data-analytics grammar.

[0254] For example, the natural language input data may include the string "precipitation in the northeast or precipitation before last year". The corresponding large language model generated data, which may include a subquery, may be "SELECT sum(precipitation) from W1 where (region=northeast) OR (date<last year)", wherein W1 is the identifier, or name, for the data structure indicated in the table structure data, and wherein "(region=northeast)" and "(date<last year)" are filters across columns. The defined data-analytics grammar may limit combinations of filters across columns, such as "(region=northeast)" and "(date<last year)", to conjunction ("AND"), where valid result data satisfies both filters, and disjunction ("OR") between filters across columns, wherein valid result data satisfies one or both of the filters may be incompatible with the defined data-analytics grammar. The data access and analysis system automatically generates an automatically generated formula to represent the disjunctive phrase "(region=northeast) OR (date<last year)", such as "formula F1=([region=northeast] or [date<last year])". The data access and analysis system uses the automatically generated formula to transform the incompatible portion of the large language model generated data to a resolved request.

[0255] For example, in the absence of the automatically generated formula, the transformation of the large language model generated data to a resolved request may be "[sum precipitation][region=northeast][date<last year]", wherein the connective between "[region=northeast]" and "[date<last year]" is implicit in the defined data-analytics grammar. The transformation of the large language model generated data to a resolved request in the absence of the automatically generated formula is an inaccurate representation of the natural language input data, "precipitation in the northeast or precipitation before last year". Using the automatically generated formula, the transformation of the large language model generated data to a resolved request may be "[sum precipitation][F1]", which is an accurate representation of the natural language input data, "precipitation in the northeast or precipitation before last year".

[0256] The data access and analysis system, or a component thereof, outputs results presentation data (at **4600**) for presenting one or more portions of results data responsive to the natural language input data (obtained at **4100**).

[0257] Outputting the results presentation data (at **4600**) includes obtaining the results data. To obtain the results data, the data access and analysis system, or a component thereof, sends, transmits, or otherwise makes available, the data query (generated at **4500**) to the data source. The data access and analysis system, or a component thereof, obtains the results data output by the data source responsive to execution of the data query by the data source. Outputting the results presentation data (at **4600**) includes generating, or otherwise obtaining, the results presentation data (at **4600**) in accordance with the results data. For example, the results presentation data may include a visualization of the results data, or one or more portions thereof. The data access and analysis system, or a component thereof, such as the system access interface unit **3900** shown in FIG. **3**, may output, or present, the results presentation data, or a portion thereof, which may include displaying the visualization.

[0258] FIG. **5** is a flow chart of an example of indexing validated resolved request data **5000** in a data access and analysis system. A data access and analysis system, such as the low-latency data access and analysis system **3000** shown in FIG. **3**, or one or more components thereof, may implement indexing validated resolved request data **5000**, or a portion thereof.

[0259] As shown in FIG. **5**, indexing validated resolved request data **5000** includes obtaining input data (at **5100**), obtaining large language model generated structured query language data (at **5200**), obtaining resolved request data (at **5300**), obtaining validated resolved request data (at **5400**), indexing (at **5410**, **5420**), revision (at **5500**), indexing (at **5510**), obtaining fragments data (at **5600**), obtaining current fragment data (at **5700**), obtaining validated fragment data (at **5800**), indexing (at **5810**, **5820**), revision (at **5900**), and indexing (at **5910**).

[0260] Obtaining input data (at **5100**) includes receiving, reading, obtaining, or otherwise accessing, (at **5100**), by the data access and analysis system, or a component thereof, such as a system access interface unit of the data access and analysis system, such as the system access interface unit **3900** shown in FIG. **3**, input data, such as user input data (first user input data), including a natural language (NL) string (input natural language (INL) data), such as by obtaining data expressing usage intent with respect to the data access and analysis system including the user input data including the natural language string, wherein the natural language string expresses a request for data, or request to obtain data, from the data access and analysis system. For example, obtaining the input data (at **5100**) may include a relational analysis unit of the data access and analysis system, such as the relational analysis unit **3700** shown in FIG. **3**, obtaining the input data from the system access interface unit.

[0261] Obtaining the first user input data includes obtaining data source data identifying a database, or another data source, or a combination of data sources. For example, the data identifying the data source may identify a worksheet, such as the worksheet **4010** shown in FIG. **4**. In another example, the data identifying the data source may identify a database accessible by the data access and analysis system, such as the distributed in-memory database **3300** shown in FIG. **3**, or an external database. In an example, the data identifying the data source identifies a worksheet that identifies a database as a data source for populating one or more columns of the worksheet using data from, or data generated from, one or more columns from one or more tables stored in, by, or at, the database.

[0262] For example, the input natural language data may be "How many active observation stations are there in the current quarter?".

[0263] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains large language model generated structured query language (LSQL) data (large language model generated structured query language data) expressing the input natural language data (at **5200**). Obtaining the large language model generated structured query language data may be similar to obtaining large language model generated data as shown (at **4300**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0264] To obtain the large language model generated structured query language data expressing the input natural language data (at **5200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, generates, creates, or otherwise obtains, large language model input data (first large language model input data)

including the input natural language data. Obtaining the first large language model input data may be similar to obtaining large language model input data as shown (at **4200**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0265] To obtain the large language model generated structured query language data expressing the input natural language data (at **5200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, sends, transmits, or otherwise makes available, the first large language model input data to a large language model accessible by the data access and analysis system, such as the large language model **4030** shown in FIG. **4**.

[0266] To obtain the large language model generated structured query language data expressing the input natural language data (at **5200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, receives, reads, obtains, or otherwise accesses, first large language model generated data including the large language model generated structured query language data from the large language model in response to the first large language model input data, which may be similar to obtaining large language model generated data as shown (at **4300**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0267] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains first resolved request (RR) data expressing the input natural language data in accordance with the defined data-analytics grammar implemented by the data access and analysis system by transforming, or otherwise processing or using, the large language model generated structured query language data (at **5300**). Obtaining the first resolved request data may be similar to transforming the large language model generated data representing the natural language input data into resolved request data representing the natural language input data as shown (at **4400**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0268] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains validated resolved request data (at **5400**) for the first resolved request data (obtained at **5300**) expressing the input natural language data (obtained at **5100**). The validated resolved request data includes data indicating a result of an accuracy, or validity, determination for the first resolved request data (RR ACCURATE?).

[0269] To obtain the validated resolved request data (at **5400**) the data access and analysis system, or a component thereof, such as the relational analysis unit, generates data for presenting the first resolved request data, or one or more portions thereof, such as via an interface, such as a graphical user interface of the data access and analysis system, or a component thereof. The data for presenting the first resolved request data may include a request for the accuracy determination (validation). The accuracy determination, or validation, may be performed manually, such as by a user of the data access and analysis system.

[0270] To obtain the validated resolved request data (at **5400**) the data access and analysis system, or a component thereof, such as the system access interface unit, presents, or otherwise outputs, the data for presenting the first resolved request data, such as to a user of the data access and analysis system.

[0271] To obtain the validated resolved request data (at **5400**), the data access and analysis system, or a component thereof, such as the system access interface unit, receives, reads, obtains, or otherwise accesses, input data, such as user input data, including the validated resolved request data, or one or more portions thereof.

[0272] The validated resolved request data, or a portion thereof, indicates that the result of the accuracy determination, or validation, is negative (negative validated resolved request data, negative accuracy determination data, NO, NEG) or that the result of the accuracy determination is positive (positive validated resolved request data, positive accuracy determination data, YES, POS). Positive validated resolved request data indicates a determination, such as a manual determination, that the first resolved request data accurately, or validly, expresses the input natural language data in accordance with the defined data-analytics grammar. Negative validated resolved request data indicates a determination, such as a manual determination, that the first resolved request data inaccurately, inefficiently, or both, expresses the input natural language data with respect to the defined data-analytics grammar.

[0273] In some implementations, the validated resolved request data (obtained at **5400**) indicates that the result of the accuracy determination is positive (YES) and the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5410**) the validated resolved request data in an index, such as a demonstrations index or a validated resolved request data index, of the data access and analysis system. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available the validated resolved request data to the indexing unit for indexing. The validated resolved request data index, or demonstrations index, is an index data structure in the data access and analysis system that indexes resolved request data validated prior to indexing.

[0274] Indexing the validated resolved request data including the positive accuracy determination data (at **5410**) includes indexing a tuple, or document. The tuple includes the input natural language (INL) data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the input natural language data (obtained at **5100**) as target, or fragment specific, natural language (TNL) data. The tuple includes the first resolved request (obtained at **5300**) as resolved request (RR) data. The tuple includes the large language model generated structured query language (LSQL) data (obtained at **5200**) as structured query language data. The tuple includes data indicating the positive accuracy determination, such as a positive integer value of one (1), as positive accuracy determination data (validation data). The tuple includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the resolved request (RR) data (obtained at **5300**), the structured query language (SQL) data (obtained at **5200**), or both.

[0275] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated

resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0276] In some implementations, indexing the validated resolved request data includes (at **5410**) determining that the validated resolved request data index includes previously indexed validated resolved request data (indexed tuple) matching the validated resolved request data with respect to the source natural language (SNL) data, the target natural language (TNL) data, the resolved request (RR) data, and the large language model generated structured query language (LSQL) data, and indexing the validated resolved request data includes incrementing, such as by adding one (1), the positive accuracy determination data in the indexed tuple.

[0277] In some implementations, the validated resolved request data (obtained at **5400**) indicates that the result of the accuracy determination is negative (NO) and the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5420**) the validated resolved request data in the validated resolved request data index. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available the validated resolved request data to the indexing unit for indexing.

[0278] Indexing the validated resolved request data including the negative accuracy determination data (at **5420**) includes indexing a tuple, or document. The tuple includes the input natural language data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the input natural language data (obtained at **5100**) as target, or fragment specific, natural language (TNL) data. The tuple includes the first resolved request data (obtained at **5300**) as resolved request (RR) data. The tuple includes the large language model generated structured query language (LSQL) data (obtained at **5200**). The tuple includes data indicating the negative accuracy determination, such as an integer value of one (1), as negative accuracy determination data (validation data). The tuple includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the resolved request (RR) data (obtained at **5300**), the structured query language (SQL) data (obtained at **5200**), or both.

[0279] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0280] In some implementations, indexing the validated resolved request data including the negative accuracy determination data (at **5420**) includes determining that the validated resolved request data index includes previously indexed validated resolved request data (indexed tuple) matching the validated resolved request data with respect to the source natural language data, the target natural language data, the resolved request data, and the structured query language data, and indexing the validated resolved request data includes incrementing, such as by adding one (1), the negative accuracy determination data in the indexed tuple.

[0281] The data access and analysis system, or a component thereof, such as the relational analysis unit, determines (at **5500**) whether revised resolved request (RR') data expressing the input natural language data (obtained at **5100**) in accordance with the defined data-analytics grammar is available (REVISION?) for the resolved request data (obtained at **5300**). Determining whether the revised resolved request data is available for the resolved request data includes obtaining the revised resolved request data (at **5500**).

[0282] To obtain the revised resolved request data expressing the input natural language data in accordance with the defined data-analytics grammar (at **5500**), the data access and analysis system, or a component thereof, such as the relational analysis unit, generates data for presenting a request for the revised resolved request data, such as via an interface, such as a graphical user interface of the data access and analysis system, or a component thereof. The data for presenting a request for the revised resolved request data includes the resolved request data (obtained at **5300**). The data for presenting a request for the revised resolved request data includes a request to revise, correct, modify, or amend, the resolved request data.

[0283] To obtain the revised resolved request data (at **5500**) the data access and analysis system, or a component thereof, such as the system access interface unit, presents, or otherwise outputs, the data for presenting the request for the revised resolved request data, such as to a user of the data access and analysis system.

[0284] To obtain the revised resolved request data (at **5500**), the data access and analysis system, or a component thereof, such as the system access interface unit of the data access and analysis system, receives, reads, obtains, or otherwise accesses, input data, such as user input data, including the revised resolved request data, or one or more portions thereof.

[0285] The revised resolved request data is similar to the resolved request data (obtained at **5300**), except as is described herein or as is otherwise clear from context. One or more portions of the revised resolved request data (obtained at **5500**) differ from corresponding portions of the resolved request data (obtained at **5300**). One or more portions of the revised resolved request data (obtained at **5500**) match, or are obtained from, corresponding portions of the resolved request data (obtained at **5300**).

[0286] The difference, or differences, between the resolved request data (obtained at **5300**) and the revised resolved request data (obtained at **5500**) may be manually generated differences. For example, the revised resolved request data (obtained at **5500**) may include one or more portions of the resolved request data (obtained at **5300**), one or more modified, such as manually modified, portions, or a combination thereof. For example, a first portion of the revised resolved request data (obtained at **5500**) may be, or match, a first portion of the resolved request data (obtained at **5300**) and a second portion of the revised resolved request data (obtained at **5500**) may be a manually generated modification, or replacement, of a second portion of the resolved request data (obtained at **5300**).

[0287] In response to obtaining the revised resolved request data, or otherwise determining that the revised

resolved request data is available (YES), (at **5500**), the data access and analysis system, or a component thereof, such as a semantic interface unit of the data access and analysis system, such as the semantic interface unit **3600** shown in FIG. **3**, obtains (at **5500**) regenerated structured query language (SQL') data, which is access and analysis system generated structured query language data expressing the input natural language data in accordance with the defined structured query language associated with, such as implemented by, the distributed in-memory database **3300**, obtained by automatically transforming the revised resolved request data (obtained at **5500**), which may be similar to the transforming described with respect to the semantic interface unit **3600** shown in FIG. **3**, except as is described herein or as is otherwise clear from context.

[0288] In response to obtaining the revised resolved request data, or otherwise determining that the revised resolved request data is available (YES), (at **5500**), the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5510**) the validated resolved request data including the revised resolved request data in the validated resolved request data index. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available the validated resolved request data including the revised resolved request data to the indexing unit for indexing.

[0289] Indexing the validated resolved request data (at **5510**) includes indexing a tuple, or document. The tuple includes the input natural language (INL) data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the input natural language (INL) data (obtained at **5100**) as target, or fragment specific, natural language (TNL) data. The tuple includes the revised resolved request (RR') data (obtained at **5500**). The tuple includes the regenerated structured query language (SQL') data (obtained at **5500**). The tuple includes data indicating a positive accuracy determination, such as a positive integer value of one (1), as positive accuracy determination data (validation data). The tuple includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the revised resolved request (RR') data (obtained at **5500**), the regenerated structured query language (SQL') data (obtained at **5500**), or both.

[0290] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0291] In some implementations, indexing the validated resolved request data (at **5510**) includes determining that the validated resolved request data index includes previously indexed validated resolved request data (indexed tuple) matching the validated resolved request data with respect to the source natural language data, the target natural language data, the revised resolved request data, and the regenerated structured query language data, and indexing the validated resolved request data (at **5510**) includes incrementing, such as by adding one (1), the positive accuracy determination data in the indexed tuple.

[0292] In response to determining (at **5500**) that the revised resolved request data is unavailable (NO), indexing (at **5510**) the validated resolved request data is omitted, excluded, or skipped.

[0293] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains fragments data (at **5600**). To obtain the fragments data, the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains, or generates, large language model input data (second large language model input data). To obtain the second large language model input data, the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains, or generates, large language model prompt data.

[0294] In an example, the revised resolved request data is unavailable (at **5500**), and the large language model prompt data includes a defined fragments generation prompt portion, the input natural language data (obtained at **5100**), the resolved request data (obtained at **5300**), and the structured query language data (obtained at **5300**).

[0295] In another example, the revised resolved request data is available (at **5500**), and the large language model prompt data includes the defined fragments generation prompt portion, the input natural language data (obtained at **5100**), the revised resolved request data (obtained at **5500**), and the regenerated structured query language data (obtained at **5500**).

[0296] Obtaining the large language model prompt data includes obtaining the defined fragments generation prompt portion. Obtaining the defined fragments generation prompt portion may be similar to obtaining prompt data as shown (at **4200**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0297] The defined fragments generation prompt portion may be expressed as the following:

“
”’
As a SQL data analyst expert, match the provided NL query fragments to the phrases present in the “Phrases” set, covering all SQL clauses. Examples are provided to explain the task.
”’
Example Input:
NL: [most active precipitation type] [monthly] [for] [last year]
Phrases: [sum Precipitation] [monthly RecordDate] [Rain Snow Hail] [Precipitation Type] [sort by Precipitation descending] [last year]
SQL: select SUM(Precipitation), Precipitation_Type, MONTHLY(RecordDate) from_Recorded_Precipitation WHERE (Precipitation_Type in (‘Rain’, ‘Snow’, ‘Hail’) AND RecordDate =_ last_year( )) group by Precipitation_Type order by SUM(Precipitation) desc;

-continued

```
Output:
{
"fragments": [
        "nl": "most active precipitation type",
        "indices": [2, 4]
        },
        {
        "nl": "monthly",
        "indices": [1]
        },
        {
        "nl": "last year",
        'indices": [5]
        }
]
}
Explanation:
The input field "NL" is the user's natural language query broken down into
fragments enclosed in square brackets. We need to assign phrases from: "Phrases"
input field to each of the fragments in "NL" query & return the index of the
phrase in the Phrases set. The index of each phrase in "Phrases" set is in the order
of appearance i.e. [sum Precipitation] has index 0, then [monthly RecordDate] has
index 1, so on, basically they are 0-indexed.
The input field "Phrases" contains all clauses extracted from the given SQL query
& enclosed in square brackets.
We can match the fragments in the example as follows:
1. "most active precipitation type":
In SQL query, most active is represented by "order by SUM(precipitation) desc"
clause, which orders the results based on the sum of precipitation in descending
order. By sorting the data in descending order, we can identify the most active
precipitation type.
For this fragment's phrase query, it borrows phrases with index 0 and 4 i.e. [sum
precipitation] and [sort by Precipitation descending] respectively. In the SQL
query, "precipitation type" is represented as filter: WHERE (Precipitation_Type in
('Rain', 'Snow', 'Hail')). We can similarly deduce its phrase indices as well: [2, 3].
2. "monthly":
This fragment indicates that we want to analyze the activity of precipitation types
on a monthly basis. In the SQL query, the "MONTHLY(RecordDate)" function is
used to group the data by months of the "RecordDate" column. Its phrase is:
[monthly RecordDate] & its index is 1.
3. "for":
This fragment cannot be matched with any phrase but only used to form user's
NL sentence. Ignore this fragment.
4. "last year"
This is present in Phrases set as well as in the SQL WHERE clause as a filter:
"RecordDate = _last_year( )". Its phrase is: [last year] & its index in Phrases set
is: 5.
Ignore NL query fragments that do not have any phrase match, for example in
above example, "for" does not have any phrase match.
Union of all indices in all the fragments would be: [0, 1, 2, 3, 4, 5] i.e. it covers all
the phrases in "Phrases" set. No phrase in the "Phrases" set should be left un-
matched.
'''
Lets take another example input:
'''
Example Input:
NL: [What] [were] [the top 3 states] [for] [snow]
Phrases: [state] [sum Precipitation] [precipitation_type = 'snow'] [by state] [sort
by sum Precipitation descending] [top 3]
SQL: select state, SUM(Precipitation) from Recorded_Precipitation where
Precipitation_Type = 'snow' group by state order by SUM(Precipitation) desc
limit 3
Output:
{
"fragments": [
        {
        "nl": " the top 3 states ",
        "indices": [0, 3, 4, 5]
        },
        {
        "nl": "snow precipitation",
        "indices": [1,2]
        },
]
}
```

-continued

```
Explanation:
Ignore NL query fragments that do not have any phrase match, for example in
above example, "what", "were" & "for" do not have any direct phrase match but
only help in sentence forming.
Union of all indices is: [0, 1, 2, 3, 4, 5] i.e. it covers all the phrases in "Phrases"
set. No phrase in the "Phrases" set should be left un-matched.
"the top 3 states" fragment hints that it is a top/bottom type of fragment. Such
fragments are synonymous with best performing, most active, so on. Its matched
phrases are composed of a sort by or ranked by phrase. In the current example that
would be phrase with index: 4 -> [sort by sum precipitation descending]. Then it
involves picking the top "n" (default is 10), in this case it is 3, so we need to pick
top 3 states based on precipitation in those states. So we can also assign the
phrases [top 3], [state] and [by state] as well.
'''
'''
Now return only the output in JSON format for the input below (NO Explanation
is required):
NL: [In] [how many sessions] [was] [the trip Berlin] [to] [Muenchen] [searched]
[for] [in] [the last month?]
Phrases: [Route = 'berlin#muenchen'] [Session Start Date = this month] [No.
Searches]
SQL: select No__Searches from Web_Traffic where ( Session_Start_Date = 'this
month' and Route = 'berlin#muenchen' ) ;
Output:".
```

[0298] Other expressions of the defined fragments generation prompt portion may be used.

[0299] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes the defined fragments generation prompt portion in the large language model prompt data.

[0300] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the defined fragments generation prompt portion, a first text preamble, such as "NL:", indicating that the subsequent portion includes natural language data.

[0301] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the first text preamble, the input natural language data (obtained at **5100**).

[0302] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the input natural language data, a second text preamble, such as "Phrase:", indicating that the subsequent portion includes resolved request data.

[0303] In an example, the validated resolved request data is unavailable (at **5500**), and the data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the second text preamble, the resolved request data (obtained at **5300**).

[0304] In another example, the validated resolved request data is available (at **5500**), and the data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the second text preamble, the revised resolved request data (obtained at **5500**).

[0305] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the resolved request data, a third text preamble, such as "SQL:", indicating that the subsequent portion includes structured query language data.

[0306] In an example, the validated resolved request data is unavailable (at **5500**), and the data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the third text preamble, the structured query language data (obtained at **5300**).

[0307] In another example, the validated resolved request data is available (at **5500**), and the data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the large language model prompt data, subsequent to the third text preamble, the regenerated structured query language data (obtained at **5500**).

[0308] The data access and analysis system, or a component thereof, such as the relational analysis unit, sends, transmits, or otherwise makes available, the large language model prompt data to the large language model as large language model input data (the second large language model input data).

[0309] The data access and analysis system, or a component thereof, such as the relational analysis unit, receives, reads, obtains, or otherwise accesses, from the large language model, the fragments data (at **5600**), such as in response to sending, transmitting, or otherwise making available, the first prompt data to the large language model.

[0310] The fragments data includes one or more fragment tuples. A fragment tuple includes a fragment, such as a phrase or chunk, of the natural language ("NL:") data from the large language model input data and resolved request phrase data including one or more index values, with respect to an index of the phrases from the phrases ("Phrases:") data from the large language model input data, associated with the fragment of the natural language data.

[0311] For example, the phrases ("Phrases:") data from the large language model input data may include "[sum Precipitation][monthly RecordDate][Rain Snow Hail][Precipitation Type][sort by Precipitation descending][last year]," and the index of the index of the phrases data may index a first phrase, or chunk, of the resolved request, "[sum Precipitation]" at a first index location (0), a second phrase, or chunk, of the resolved request, "[monthly RecordDate]" at a second index location (1), a third phrase, or chunk, of the

resolved request, "[Rain Snow Hail]" at a third index location (2), a fourth phrase, or chunk, of the resolved request, "[Precipitation Type]" at a fourth index location (3), a fifth phrase, or chunk, of the resolved request, "[sort by Precipitation descending]" at a fifth index location (4), and a sixth phrase, or chunk, of the resolved request, "[last year]" at a sixth index location (5).

[0312] In an example, the fragments data may be expressed as the following:

```
"{
"fragments": [
        {
        "nl": " most active precipitation type ",
        "indices": [2, 4]
        },
        {
        "nl": "monthly",
        "indices": [1]
        },
        {
        "nl": "last year",
        "indices": [5]
        }
]
}".
```

[0313] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains (at **5700**) current fragment (CF) data for a current fragment from the fragments data (obtained at **5600**). Although described with respect to a (one) current fragment, obtaining current fragment data (at **5700**), obtaining validated fragment data (at **5800**), and indexing (at **5810**, **5820**), revision (at **5900**), and indexing (at **5910**) may be performed for multiple fragments, such as the fragments from the fragments data (obtained at **5600**), such as concurrently or iteratively, as indicated by the directional broken lines from indexing (at **5810**), revision (at **5900**), and indexing (at **5910**) to obtaining the current fragment (at **5700**).

[0314] In an example, the fragments data (obtained at **5600**) includes a first fragment tuple, obtained as the current fragment data (at **5700**), which includes a fragment, chunk, or phrase, of the input natural language data and resolved request phrase data, including one or more index values, with respect to an index of the phrases from the phrases ("Phrases:") data from the large language model input data, associated with, or expressing, the fragment of the natural language data.

[0315] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains validated fragment data (at **5800**) for the current fragment data (obtained at **5700**). The validated fragment data includes data indicating a result of an accuracy, or validity, determination (second accuracy determination) for the current fragment (CF ACCURATE?).

[0316] To obtain the validated fragment data (at **5800**), the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains fragment resolved request (FRR) data expressing the resolved request phrase data from the current fragment data (obtained at **5700**). To obtain the fragment resolved request (FRR) data, the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the resolved request phrase data including one or more index values, with respect to the index of the phrases from the phrases ("Phrases:") data from the large language model input data. To obtain the fragment resolved request (FRR) data, the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the phrase, chunk, phrases, or chunks, of the resolved request indicated by the one or more index values from the resolved request phrase data.

[0317] In some implementations, the resolved request phrase data includes one index value with respect to the index of the phrases from the phrases ("Phrases:") data from the large language model input data, and the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the phrase, or chunk, of resolved request data indicated by the index value. For example, the phrases ("Phrases:") data from the large language model input data may be "Phrases: [sum Precipitation][monthly RecordDate][Rain Snow Hail][Precipitation Type][sort by Precipitation descending][last year]," the current fragment data (obtained at **5700**) may be "{"nl": "monthly", "indices": [1]}," and the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the phrase, or chunk, of resolved request data, "monthly RecordDate" indicated by the index value one (1) as the fragment resolved request (FRR) data.

[0318] In some implementations, the resolved request phrase data includes two or more index values with respect to the index of the phrases from the phrases ("Phrases:") data from the large language model input data, and the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the phrases, or chunks, of resolved request data indicated by the index values, and aggregates, or combines, the phrases, or chunks, of resolved request data to obtain the fragment resolved request (FRR) data. For example, the phrases ("Phrases:") data from the large language model input data may be "Phrases: [sum Precipitation][monthly RecordDate][Rain Snow Hail][Precipitation Type][sort by Precipitation descending][last year]," the current fragment data (obtained at **5700**) may be "{"nl": "most active precipitation type", "indices": [2, 4]}," the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the phrase, or chunk, of resolved request data, "Rain Snow Hail" indicated by the index value two (2) and the phrase, or chunk, of resolved request data, "sort by Precipitation descending" indicated by the index value four (4), and the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains, as the fragment resolved request (FRR) data, an aggregation, or combination, of "Rain Snow Hail" and "sort by Precipitation descending", such as "Rain Snow Hail sort by Precipitation descending."

[0319] To obtain the validated fragment data (at **5800**) the data access and analysis system, or a component thereof, such as the relational analysis unit, generates data for presenting the fragment resolved request (FRR) data corresponding to the current fragment data, or one or more portions thereof, such as via an interface, such as a graphical user interface of the data access and analysis system, or a component thereof. The data for presenting the fragment resolved request (FRR) data may include a request for the accuracy determination (validation). The accuracy determination, or validation, may be performed manually, such as by a user of the data access and analysis system.

[0320] To obtain the validated fragment data (at **5800**), the data access and analysis system, or a component thereof,

such as the system access interface unit, presents, or otherwise outputs, the data for presenting the fragment resolved request (FRR) data, such as to a user of the data access and analysis system.

[0321] To obtain the validated fragment data (at **5800**), the data access and analysis system, or a component thereof, such as the system access interface unit, receives, reads, obtains, or otherwise accesses, input data, such as user input data, including the validated fragment data, or one or more portions thereof.

[0322] The validated fragment data, or a portion thereof, indicates that the result of the accuracy determination, or validation, is negative (negative validated fragment data, negative accuracy determination data, NO, NEG) or that the result of the accuracy determination is positive (positive validated fragment data, positive accuracy determination data, YES, POS). Positive validated fragment data indicates a determination, such as a manual determination, that the fragment resolved request (FRR) data for the current fragment data accurately, or validly, expresses the fragment of the input natural language data from the current fragment data in accordance with the defined data-analytics grammar. Negative validated fragment data indicates a determination, such as a manual determination, that the fragment resolved request (FRR) data for the current fragment data inaccurately, inefficiently, or both, expresses the fragment of the input natural language data from the current fragment data with respect to the defined data-analytics grammar.

[0323] In some implementations, the validated fragment data (obtained at **5800**) indicates that the result of the accuracy determination (at **5800**) is positive (YES) and the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5810**) the validated fragment data in the index, such as the demonstrations index or the validated resolved request data index, of the data access and analysis system. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available the validated fragment data to the indexing unit for indexing.

[0324] To index the validated fragment data (obtained at **5800**), the data access and analysis system, or a component thereof, such as the relational analysis unit, obtains regenerated structured query language data for the fragment (fragment structured query language (FSQL) data), which is access and analysis system generated structured query language data expressing the resolved request phrase fragment indicated in the fragment resolved request (FRR) data in accordance with the defined structured query language associated with, such as implemented by, the distributed in-memory database **3300**, obtained by automatically transforming the resolved request phrase fragment indicated in the fragment resolved request (FRR) data, which may be similar to the transforming described with respect to the semantic interface unit **3600** shown in FIG. **3**, except as is described herein or as is otherwise clear from context.

[0325] Indexing the validated fragment data (at **5810**) includes indexing a tuple, or document. The tuple (indexed at **5810**) includes the input natural language (INL) data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the fragment of the input natural language data (FNL) from the current fragment data (obtained at **5700**) as target, or fragment specific, natural language (TNL) data. The tuple includes the automatically generated fragment resolved request (FRR) data expressing the fragment of the input natural language data from the current fragment data (obtained at **5700**). The tuple includes the structured query language (SQL) data automatically generated for the fragment resolved request (FRR) data (fragment structured query language (FSQL) data). The tuple includes data indicating the positive accuracy determination, such as a positive integer value of one (1), as positive accuracy determination data (validation data). The tuple includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the fragment resolved request (FRR) data (obtained at **5700**), the fragment structured query language (FSQL) data, or both.

[0326] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0327] In some implementations, indexing the validated fragment data (at **5810**) includes determining that the validated resolved request data index includes previously indexed validated fragment data (indexed tuple) matching the validated fragment data with respect to the source natural language data, the target natural language data, the fragment resolved request data, and the structured query language data from the current fragment, and indexing (at **5810**) the validated fragment data includes incrementing, such as by adding one (1), the positive accuracy determination data in the indexed tuple.

[0328] In some implementations, the validated fragment data (obtained at **5800**) indicates that the result of the accuracy determination is negative (NO) and the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5820**) the validated fragment data in the validated resolved request data index. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available (at **5820**) the validated fragment data to the indexing unit for indexing.

[0329] Indexing the validated fragment data (at **5820**) includes indexing a tuple, or document. The tuple (indexed at **5820**) includes the input natural language (INL) data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the fragment of the input natural language data (FNL) from the current fragment data (obtained at **5700**) as target, or fragment specific, natural language (TNL) data. The tuple includes the automatically generated fragment resolved request (FRR) data expressing the fragment of the input natural language data from the current fragment data (obtained at **5700**). The tuple includes the structured query language (SQL) data automatically generated for the fragment resolved request (FRR) data (fragment structured query language (FSQL) data). The tuple includes data indicating the negative accuracy determination, such as an integer value of one (1), as negative accuracy determination data (validation data). The tuple

includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the fragment resolved request (FRR) data (obtained at **5700**), the fragment structured query language (FSQL) data, or both.

[0330] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0331] In some implementations, indexing the validated fragment data (at **5820**) includes determining that the validated resolved request data index includes previously indexed validated fragment data matching the validated fragment data (indexed tuple) with respect to the source natural language data, the target natural language data, the fragment resolved request data, and the fragment structured query language data from the current fragment, and indexing (at **5820**) the validated fragment data includes incrementing, such as by adding one (1), the negative accuracy determination data in the indexed tuple.

[0332] The data access and analysis system, or a component thereof, such as the relational analysis unit, determines (at **5900**) whether revised fragment resolved request data (FRR') including a revised fragment resolved request phrase expressing the target natural language data (the fragment of the input natural language data from the current fragment data) in accordance with the defined data-analytics grammar is available (REVISION?) for the resolved request phrase (obtained at **5800**). Determining whether the revised fragment resolved request data is available for the resolved request phrase includes obtaining the revised fragment resolved request data (at **5900**).

[0333] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains (at **5900**) the revised fragment resolved request data (FRR') including a revised fragment resolved request phrase expressing the target natural language data (the fragment of the input natural language data from the current fragment data) in accordance with the defined data-analytics grammar.

[0334] To obtain the revised fragment resolved request data including the revised fragment resolved request phrase expressing the target natural language data from the current fragment in accordance with the defined data-analytics grammar (at **5900**), the data access and analysis system, or a component thereof, such as the relational analysis unit, generates data for presenting a request for the revised fragment resolved request data, such as via an interface, such as a graphical user interface of the data access and analysis system, or a component thereof. The data for presenting the request for the revised fragment resolved request data includes the fragment resolved request data from the current fragment (obtained at **5700**). The data for presenting the request for the revised fragment resolved request data includes a request to revise, correct, modify, or amend, the fragment resolved request data.

[0335] To obtain the revised fragment resolved request data including the revised fragment resolved request phrase expressing the target natural language data from the current fragment in accordance with the defined data-analytics grammar (at **5900**), the data access and analysis system, or a component thereof, such as the system access interface unit, presents, or otherwise outputs, the data for presenting the request for the revised fragment resolved request data, such as to a user of the data access and analysis system.

[0336] To obtain the revised fragment resolved request data including the revised fragment resolved request phrase expressing the target natural language data from the current fragment in accordance with the defined data-analytics grammar (at **5900**), the data access and analysis system, or a component thereof, such as the system access interface unit of the data access and analysis system, receives, reads, obtains, or otherwise accesses, user input data including the revised fragment resolved request data, or one or more portions thereof.

[0337] The revised fragment resolved request data including the revised fragment resolved request phrase expressing the target natural language data from the current fragment in accordance with the defined data-analytics grammar is similar to the fragment resolved request data from the current fragment, except as is described herein or as is otherwise clear from context. One or more portions of the revised fragment resolved request data (obtained at **5900**) differ from corresponding portions of the fragment resolved request data from the current fragment. One or more portions of the revised fragment resolved request data (obtained at **5900**) match, or are obtained from, corresponding portions of the fragment resolved request data from the current fragment.

[0338] The difference, or differences, between the fragment resolved request data from the current fragment (obtained at **5700**) and the revised fragment resolved request data (obtained at **5900**) may be manually generated differences. For example, the revised fragment resolved request data (obtained at **5900**) may include one or more portions of the fragment resolved request data from the current fragment, one or more modified, such as manually modified, portions, or a combination thereof. For example, a first portion of the revised fragment resolved request data (obtained at **5900**) may be, or match, a first portion of the fragment resolved request data from the current fragment and a second portion of the revised fragment resolved request data (obtained at **5900**) may be a manually generated modification, or replacement, of a second portion of the fragment resolved request data from the current fragment. The data access and analysis system, or a component thereof, such as the semantic interface unit of the data access and analysis system, replaces, in the current fragment data, the fragment resolved request data (obtained at **5700**) with the revised fragment resolved request data (obtained at **5900**).

[0339] In response to obtaining the revised fragment resolved request data, or otherwise determining that the revised fragment resolved request data is available (YES), (at **5900**), the data access and analysis system, or a component thereof, such as the semantic interface unit of the data access and analysis system, obtains (at **5900**) regenerated fragment structured query language (FSQL') data expressing the target natural language data from the current fragment, which is access and analysis system generated structured query language data expressing the target natural language data from the current fragment in accordance with the

defined structured query language associated with, such as implemented by, the distributed in-memory database **3300**, obtained by automatically transforming the revised fragment resolved request data (obtained at **5900**), which may be similar to the transforming described with respect to the semantic interface unit **3600** shown in FIG. **3**, except as is described herein or as is otherwise clear from context. The data access and analysis system, or a component thereof, such as the semantic interface unit of the data access and analysis system, replaces, in the current fragment data, the fragment structured query language data with the regenerated fragment structured query language data.

[0340] In response to obtaining the revised fragment resolved request data, or otherwise determining that the revised fragment resolved request data is available (YES), (at **5900**), the data access and analysis system, or a component thereof, such as an indexing unit of the data access and analysis system, indexes (at **5910**) the validated fragment data including the revised fragment resolved request data in the validated resolved request data index. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit may send, transmit, or otherwise make available the validated fragment data including the revised fragment resolved request data to the indexing unit for indexing.

[0341] Indexing the validated fragment data (at **5910**) includes indexing a tuple, or document. The tuple includes the input natural language (INL) data (obtained at **5100**) as source, or input, natural language (SNL) data. The tuple includes the fragment of the input natural language data (FNL) from the current fragment data (obtained at **5700**) as target, or fragment specific, natural language (TNL) data. The tuple includes the revised fragment resolved request (FRR') data (obtained at **5900**) as resolved request data. The tuple includes the regenerated fragment structured query language (FSQL') data (obtained at **5900**) as structured query language data. The tuple includes data indicating a positive accuracy determination, such as a positive integer value of one (1), as positive accuracy determination data (validation data). The tuple includes temporal data (not expressly shown), such as data indicating a temporal location corresponding to indexing the tuple. The tuple may include data source data (not expressly shown), such as column identifier, or name, data, indicating one or more data sources, or columns, such as data sources indicated by the revised fragment resolved request (FRR') data (obtained at **5900**), the regenerated fragment structured query language (FSQL') data (obtained at **5900**), or both.

[0342] In some implementations, the tuple, or data indexed in association with the tuple, may include other data, such as usage data, mapping scope data, data-domain data, a user identifier, data indicating whether the validated resolved request data is associated with an administrative account, which may indicate an administrative account type, ontological data for the natural language string, which may include categorical data.

[0343] In some implementations, indexing the validated fragment data including the revised fragment resolved request data (at **5910**) includes determining that the validated resolved request data index includes previously indexed validated fragment data (indexed tuple) matching the validated fragment data with respect to the source natural language data, the target natural language data, the revised fragment resolved request data, and the regenerated fragment structured query language data, and indexing the validated fragment data including the revised fragment resolved request data includes incrementing (at **5910**), such as by adding one (1), the positive accuracy determination data in the indexed tuple.

[0344] In response to determining that the revised fragment resolved request data is unavailable (NO), (at **5900**), indexing (at **5910**) the validated fragment data including the revised fragment resolved request data is omitted, excluded, or skipped.

[0345] FIG. **6** is a flowchart of an example of obtaining demonstration data **6000** in a data access and analysis system. The data access and analysis system may be similar to the data access and analysis system **3000** shown in FIG. **3**, except as is described herein or as is otherwise clear from context. The data access and analysis system may implement obtaining demonstration data **6000**. For example, the data access and analysis system may implement obtaining demonstration data **6000** to obtain the demonstration data **4040** shown in FIG. **4**. Obtaining large language model input data, such obtaining large language model input data as shown (at **4200**) in FIG. **4**, may include obtaining demonstration data **6000**.

[0346] In accordance with prompt signifier data, such as the prompt signifier data obtained as shown (at **4200**) in FIG. **4**, prompt context data, such as the prompt context data obtained as shown (at **4200**) in FIG. **4**, or both, the demonstration data may be, or include, one or more examples of natural language input data and corresponding structured query language data, that the large language model may use to determine how to generate, such as structure, format, or a combination thereof, the large language model generated data responsive to natural language input data (request for data) included in the large language model input data.

[0347] Obtaining the demonstration data **6000** includes obtaining natural language data (at **6100**), obtaining candidate matching demonstrations (at **6200**), obtaining noun-phrase fragments (at **6300**), filtering and sorting (at **6400**), determining whether the prompt context data includes a column corresponding to a selected determination (at **6500**), including the column in the prompt context data (at **6600**), and outputting the demonstration data (at **6700**).

[0348] Obtaining natural language data (at **6100**) includes receiving, reading, obtaining, or otherwise accessing, (at **6100**), by the data access and analysis system, or a component thereof, such as a system access interface unit of the data access and analysis system, such as the system access interface unit **3900** shown in FIG. **3**, input data, such as user input data (first user input data), including a natural language (NL) string (NL input or natural language string), such as by obtaining data expressing usage intent with respect to the data access and analysis system including the user input data including the natural language string, wherein the natural language string expresses a request for data, or request to obtain data, from the data access and analysis system. For example, obtaining the natural language data (at **6100**) may include a relational analysis unit of the data access and analysis system, such as the relational analysis unit **3700** shown in FIG. **3**, obtaining the natural language data from the system access interface unit.

[0349] Obtaining the natural language data includes obtaining data source data identifying a database, or another data source, or a combination of data sources. For example, the data identifying the data source may identify a work-

sheet, such as the worksheet **4010** shown in FIG. **4**. In another example, the data identifying the data source may identify a database accessible by the data access and analysis system, such as the distributed in-memory database **3300** shown in FIG. **3**, or an external database. In an example, the data identifying the data source identifies a worksheet that identifies a database as a data source for populating one or more columns of the worksheet using data from, or data generated from, one or more columns from one or more tables stored in, by, or at, the database.

[0350] For example, the first natural language string may be "How many active observation stations are there in the current quarter?".

[0351] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains one or more candidate matching demonstrations (at **6200**) (first candidate matching demonstrations). A candidate matching demonstration is a previously indexed, such as in the demonstrations, or validated resolved request data, index, tuple that includes source, or input, natural language data, target, or fragment specific, natural language data, resolved request data, structured query language data, data source data, validation data, and temporal data.

[0352] The source, or input, natural language data from the candidate matching demonstration is a natural language string, such as the first natural language string previously indexed as shown (at **5410**, **5420**, **5510**, **5810**, **5820**, or **5910**) in FIG. **5**. For example, the source, or input, natural language data from the candidate matching demonstration may include the natural language string "What is the precipitation from measurement of rain last year?"

[0353] The target natural language data from the candidate matching demonstration, is the source, or input, natural language data, or a portion, such as a fragment, thereof, previously indexed as shown (at **5410**, **5420**, **5510**, **5810**, **5820**, or **5910**) in FIG. **5**. For example, the target natural language data may include the natural language string "What is the precipitation from measurement of rain last year?". In another example, the target natural language data may include the natural language string "precipitation from measurement". In another example, the target natural language data may include the natural language string "rain". In another example, the target natural language data may include the natural language string "last year".

[0354] The resolved request data from the candidate matching demonstration includes a resolved request expression, in accordance with the defined data-analytics grammar implemented by the data access and analysis system, corresponding to, representing, or expressing, the target natural language data, previously indexed as shown (at **5410**, **5420**, **5510**, **5810**, **5820**, or **5910**) in FIG. **5**. For example, the target natural language data from the candidate matching demonstration may include the natural language string "What is the precipitation from measurement of rain last year?" and the corresponding resolved request expression may be "[sum precipitation][precipitation_type=rain][date=last year]". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "precipitation from measurement" and the corresponding resolved request expression may be "[sum precipitation]". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "rain" and the corresponding resolved request expression may be "[precipitation_type=rain]". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "last year" and the corresponding resolved request expression may be "[date=last year]".

[0355] The structured query language data from the candidate matching demonstration expresses the resolved request data from the candidate matching demonstration, such as in accordance with a defined structured query language implemented by a distributed in-memory database of the data access and analysis system, such as the distributed in-memory database **3300** shown in FIG. **3**. For example, the target natural language data from the candidate matching demonstration may include the natural language string "What is the precipitation from measurement of rain last year?", the corresponding resolved request expression may be "[sum precipitation][precipitation_type=rain][date=last year]", and the structured query language data from the candidate matching demonstration may include "select SUM(Precipitation) from _Recorded_Precipitation WHERE (Precipitation_Type="rain" AND RecordDate=_last_year( ))". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "precipitation from measurement", the corresponding resolved request expression may be "[sum precipitation]", and the structured query language data from the candidate matching demonstration may include "select SUM(Precipitation) from _Recorded Precipitation". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "rain", the corresponding resolved request expression may be "[precipitation_type=rain]", and the structured query language data from the candidate matching demonstration may include "select SUM(Precipitation) from _Recorded_Precipitation WHERE (Precipitation_Type="rain")". In another example, the target natural language data from the candidate matching demonstration may include the natural language string "last year", the corresponding resolved request expression may be "[date=last year]", and the structured query language data from the candidate matching demonstration may include "select * from _Recorded_Precipitation WHERE (RecordDate=_last_year( ))".

[0356] The data source data from the candidate matching demonstration includes one or more data source identifiers, such as one or more column names, wherein the resolved request data from the candidate matching demonstration, the structured query language data from the candidate matching demonstration, or both, indicates, or expresses, a request to obtain data from, or otherwise use, the respective data source. In some implementations, separate, or distinct, data source data may be absent from the candidate matching demonstration, and the data source data may be obtained from the resolved request data from the candidate matching demonstration, the structured query language data from the candidate matching demonstration, or both.

[0357] The validation data from the candidate matching demonstration includes positive accuracy determination data, negative accuracy determination data, or both.

[0358] The temporal data from the candidate matching demonstration indicates a temporal location, or point in time, corresponding to the indexing of the candidate matching demonstration.

[0359] To obtain the candidate matching demonstrations (at **6200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, may generate a request for the candidate matching demonstrations. The data access and analysis system, or a component thereof, such as the relational analysis unit, may send, transmit, or otherwise make available the request for the candidate matching demonstrations to another component of the data access and analysis system, such as an indexing unit of the data access and analysis system, such as the indexing unit described with respect to FIG. **5**.

[0360] To obtain the candidate matching demonstrations (at **6200**), the data access and analysis system, or a component thereof, such as the indexing unit, traverses the demonstrations, or validated resolved request data, index. Traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations includes identifying previously indexed tuples as candidate matching demonstrations in response to determining that the natural language data (obtained at **6100**) is a semantic match with the source, or input, natural language data included in the previously indexed tuple, the target natural language data included in the indexed tuple, or both. For example, whether the natural language data (obtained at **6100**) is a semantic match with the source, or input, natural language data included in the previously indexed tuple, or the target natural language data included in the previously indexed tuple, may be determined using a semantic matching algorithm, such as a best matching algorithm, such as BM25.

[0361] In some implementations, traversing the available demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include obtaining a matching score with respect to matching the natural language data (obtained at **6100**) with the source, or input, natural language data included in the previously indexed tuple. In some implementations, traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include obtaining a matching score with respect to matching the natural language data (obtained at **6100**) with the target natural language data included in the previously indexed tuple.

[0362] In some implementations, traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include obtaining a matching score with respect to matching the natural language data (obtained at **6100**) with the source, or input, natural language data included in the previously indexed tuple, obtaining a matching score with respect to matching the natural language data (obtained at **6100**) with the target natural language data included in the previously indexed tuple, and obtaining a matching score for the previously indexed tuple as a sum of the matching score for the source, or input, natural language data included in the previously indexed tuple and the matching score with respect to matching the natural language data (obtained at **6100**) with the target natural language data included in the previously indexed tuple.

[0363] Traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include obtaining a defined number, count, or cardinality of candidate matching demonstrations, such as ten (10) candidate matching demonstrations, such as in descending matching score order.

[0364] Traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include omitting, skipping, excluding, or filtering, previously indexed tuples other than previously indexed tuples including a positive accuracy determination data value of at least, such as greater than or equal to, one (>=1).

[0365] Traversing the demonstrations, or validated resolved request data, index to obtain the candidate matching demonstrations may include omitting, skipping, excluding, or filtering, previously indexed tuples other than previously indexed tuples including a positive accuracy determination data value and a negative accuracy determination value wherein a difference between the positive accuracy determination data value and the negative accuracy determination, such as obtained by subtracting the negative accuracy determination value from the positive accuracy determination data value, of at least, such as greater than or equal to, one (>=1).

[0366] In some implementations, obtaining the candidate matching demonstrations includes obtaining the candidate matching demonstrations in descending matching score order.

[0367] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains one or more noun phrases (at **6300**) from the natural language input data (obtained at **6100**), such as a using a noun phrase chunking model.

[0368] The data access and analysis system, or a component thereof, such as the relational analysis unit, filters, sorts, or both, the candidate matching demonstrations (at **6400**) to obtain sorted, filtered, candidate matching demonstrations.

[0369] For example, the data access and analysis system, or a component thereof, such as the relational analysis unit, may filter the candidate matching demonstrations (at **6400**) to obtain filtered candidate matching demonstrations and the data access and analysis system, or a component thereof, such as the relational analysis unit, may sort the filtered candidate matching demonstrations (at **6400**) to obtain the sorted, filtered, candidate matching demonstrations.

[0370] To filter the candidate matching demonstrations (at **6400**), the data access and analysis system, or a component thereof, such as the relational analysis unit, determines whether the target natural language data from a respective candidate matching demonstration includes, or matches with, one or more of the noun phrases (obtained at **6300**), which includes determining a number, count, or cardinality, of the noun phrases (obtained at **6300**) that match, or are included in, the target natural language data from the respective candidate matching demonstration.

[0371] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes, in the filtered candidate matching demonstrations, the candidate matching demonstrations determined to have a number, count, or cardinality, of matching noun phrases greater than or equal to a defined threshold, such as two (2). Candidate matching demonstrations determined to have a number, count, or cardinality, of matching noun phrases less than the defined threshold are omitted, or excluded, from the filtered candidate matching demonstrations.

[0372] Sorting the candidate matching demonstrations (at **6400**) includes sorting the filtered candidate matching demonstrations in accordance with one or more metrics to obtain the sorted, filtered, candidate matching demonstrations. For

example, the data access and analysis system, or a component thereof, such as the relational analysis unit, may sort the filtered candidate matching demonstrations in accordance with the temporal data included in the respective candidate matching demonstrations, such as into recency order. In another example, the data access and analysis system, or a component thereof, such as the relational analysis unit, may sort the filtered candidate matching demonstrations in accordance with the positive accuracy determination data value included in the respective candidate matching demonstrations, such as in descending order. In some implementations, the data access and analysis system, or a component thereof, such as the relational analysis unit, may sort the filtered candidate matching demonstrations in accordance with the temporal data, such as into recency order, and may subsequently sort the temporally ordered filtered candidate matching demonstrations in accordance with the positive accuracy determination data value included in the respective candidate matching demonstrations, such as in descending order.

[0373] Sorting and filtering the candidate matching demonstrations (at **6400**) includes obtaining a defined number, count, or cardinality, such as three, of the sorted, filtered, candidate matching demonstrations, such as in descending order, as selected demonstrations.

[0374] The data access and analysis system, or a component thereof, such as the relational analysis unit, determines (at **6500**) whether the prompt context data includes data source data indicating data sources, such as columns, corresponding to the data sources, such as columns, indicated by one or more of the selected determinations (obtained at **6400**), such as on a per-data source basis. For example, the data access and analysis system, or a component thereof, such as the relational analysis unit, may determine (at **6500**) that a data source identifier, such as a column identifier, such as a column name, included in the data source data from a selected determination (obtained at **6400**) is present in the prompt context data (as shown at **6510**). In another example, the data access and analysis system, or a component thereof, such as the relational analysis unit, may determine (at **6500**) that a data source identifier, such as a column identifier, such as a column name, included in the data source data from a selected determination (obtained at **6400**) is absent from the prompt context data (as shown at **6520**).

[0375] The data access and analysis system, or a component thereof, such as the relational analysis unit, includes the data source identifiers, such as column identifiers, determined (at **6500**) to be absent, or omitted, from the prompt context data in the prompt context data (at **6600**).

[0376] In some implementations, determining whether the prompt context data includes the data source data (at **6500**) and including the data source data in the prompt context data (at **6600**) may be omitted, or skipped, as indicated by the broken line border (at **6500**, **6600**).

[0377] The data access and analysis system, or a component thereof, such as the relational analysis unit, outputs the selected demonstrations data (at **6700**). For example, outputting the selected demonstrations data (at **6700**) may include including the selected demonstrations data in the large language model input data, such as subsequent to the prompt context data, such as shown (at **4200**) in FIG. **4**.

[0378] FIG. **7** is a flow chart of an example of obtaining results data with automatic target visualization adaptation **7000** in a data access and analysis system. A data access and analysis system, such as the low-latency data access and analysis system **3000** shown in FIG. **3**, or one or more components thereof, may implement obtaining results data with automatic target visualization adaptation **7000**, or a portion thereof.

[0379] As shown in FIG. **7**, obtaining results data with automatic target visualization adaptation **7000** includes obtaining input data (at **7100**), obtaining large language model generated structured query language data (at **7200**), obtaining resolved request data (at **7300**), obtaining target-parameters data (at **7400**), obtaining target-parameters deficit data (at **7500**), obtaining co-occurring parameters (at **7600**), obtaining the results data (at **7700**), and outputting the results data (at **7800**).

[0380] Obtaining input data (at **7100**) includes receiving, reading, obtaining, or otherwise accessing, (at **7100**), by the data access and analysis system, or a component thereof, such as a system access interface unit of the data access and analysis system, such as the system access interface unit **3900** shown in FIG. **3**, input data, such as user input data, including a natural language (NL) string (NL input), such as by obtaining data expressing usage intent with respect to the data access and analysis system including the user input data including the natural language string, wherein the natural language string expresses a request for data, or request to obtain data, from the data access and analysis system. For example, obtaining the input data (at **7100**) may include a relational analysis unit of the data access and analysis system, such as the relational analysis unit **3700** shown in FIG. **3**, obtaining the input data from the system access interface unit.

[0381] For example, the natural language string may be "show me the total number of precipitation events over the last 6 months".

[0382] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains large language model generated (LLMG) structured query language (SQL) data (large language model generated structured query language data) expressing the natural language string (at **7200**). Obtaining the large language model generated structured query language data may be similar to obtaining large language model generated data as shown (at **4300**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0383] To obtain the large language model generated structured query language data expressing the natural language string (at **7200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, generates, creates, or otherwise obtains, large language model input data (first large language model input data) including the natural language string. Obtaining the first large language model input data may be similar to obtaining large language model input data as shown (at **4200**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0384] To obtain the large language model generated structured query language data expressing the natural language string (at **7200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, sends, transmits, or otherwise makes available, the first large language model input data to a large language model accessible by the data access and analysis system, such as the large language model **4030** shown in FIG. **4**.

[0385] To obtain the large language model generated structured query language data expressing the natural lan-

guage string (at **7200**), the data access and analysis system, or a component thereof, such as the relational analysis unit, receives, reads, obtains, or otherwise accesses, large language model generated data including the large language model generated structured query language data from the large language model in response to the first large language model input data, which may be similar to obtaining large language model generated data as shown (at **4300**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0386] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains first resolved request data expressing the natural language string in accordance with the defined data-analytics grammar implemented by the data access and analysis system by transforming, or otherwise processing or using, the large language model generated structured query language data (at **7300**). Obtaining the first resolved request data may be similar to transforming the large language model generated data representing the natural language input data into resolved request data representing the natural language input data as shown (at **4400**) in FIG. **4**, except as is described herein or as is otherwise clear from context.

[0387] The first resolved request data includes an ordered sequence of tokens that form one or more resolved request parameters (request parameters) in one or more resolved request phrases (request phrases). A request parameter, or resolved request parameter, corresponds to, such as uniquely identifies, a unit of data in the data access and analysis system, such as a such as a column of a table in a database, or other data source, accessible by the data access and analysis system or a data value stored in a column of a table accessible by the data access and analysis system. A request parameter that represents a column of a table in a database, or other data source, accessible by the data access and analysis system, may include a name, such as a string value, of the column and a unique identifier, such as a universally unique identifier (UUID) or a globally unique identifier (GUID). A request parameter that represents a data value, such as a data value stored in a row of a column of a table in a database, or other data source, accessible by the data access and analysis system, may include a string value and a unique identifier, such as a universally unique identifier (UUID) or a globally unique identifier (GUID), which may uniquely identify the data value, the column in which the data value is stored, or both.

[0388] In an example, the natural language string is "show me the total number of precipitation events last year" and the corresponding resolved request is "count(event_id) event_type='precipitation' last year (created date)", wherein "count(event id)" is a first resolved request phrase wherein "event_id" is a first request parameter, or a portion thereof, corresponding to an attribute column of a table accessible by the data access and analysis system, "event_type='precipitation'" is a second resolved request phrase (filter phrase) wherein "event_type" is a second request parameter, or a portion thereof, corresponding to a column of a table accessible by the data access and analysis system and "'precipitation'" is a third request parameter, or a portion thereof, corresponding to a data value, used as a filter value in the resolved request, stored in the "event_type" column, and "last year (created date)" is a third resolved request phrase wherein "created date" is a fourth request parameter, or a portion thereof, corresponding to a temporal column of a table accessible by the data access and analysis system.

[0389] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains target-parameters data (at **7400**). To obtain the target-parameters data (at **7400**) the data access and analysis system, or a component thereof, obtains target visualization data including the target-parameters data.

[0390] The target visualization data is a data structure previously, such as prior to obtaining the input data (at **7100**), stored in the data access and analysis system. The target visualization data may be account, or user, specific. The target visualization data may be stored in the data access and analysis system in association with account, or user, identification data, such as an account identifier.

[0391] In some implementations, obtaining the target visualization data (at **7400**) may include automatically generating the target visualization data. For example, the data access and analysis system, or a component thereof, may automatically generate the target visualization data in response to obtaining the input data (at **7100**) and determining that target visualization data is otherwise unavailable for the account identifier. In an example, the data access and analysis system, or a component thereof, may determine a most frequently accessed visualization type based on previously recorded utilization data associated with the account identifier, and may generate the target visualization data in accordance with the most frequently accessed visualization type.

[0392] In an example, obtaining the input data (at **7100**) includes obtaining, from the input data, an account identifier associated with a user that input the input data, or a portion thereof, and obtaining the target visualization data (at **7400**) includes obtaining the target visualization data in accordance with the account identifier, such as by using the account identifier as an index value or key value for retrieving, or looking up, the target visualization data.

[0393] The target visualization data includes data describing a target visualization type, such as a target visualization type identifier or target visualization type value. For example, the target visualization data may include one or more data elements that indicate that the target visualization type is a line chart, a bar chart, a pie chart, and area chart, a stacked bar chart, or another type of visualization.

[0394] The target visualization data includes the target-parameters data. The target-parameters data includes one or more target-parameters counts, or target-parameters count values. A target-parameters count, or target-parameters count value, is a data element, such as an integer value, indicating a target number, count, or cardinality, of data elements, or parameters (visualization parameters), for the visualization type indicated by the target visualization type and for a respective parameter type, such as an attribute column parameter type, a measure column parameter type, a filter value parameter type, or a temporal column parameter type.

[0395] In an example, the target-parameters data includes a target-parameters count for the attribute column parameter type, a target-parameters count for the measure column parameter type, a target-parameters count for the filter value parameter type, and a target-parameters count for the temporal, or time-series, column parameter type.

[0396] A target-parameters count for a respective parameter type indicates that outputting data for presenting a visualization of results data in accordance with the target visualization type is unavailable for results data obtained in accordance with a resolved request including fewer parameters of the parameter type than indicated by the target-parameters count for the parameter type. The target-parameters count for the attribute column parameter type indicates a number, count, or cardinality, of attribute columns. The target-parameters count for the measure column parameter type indicates a number, count, or cardinality, of measure columns. The target-parameters count for the filter value parameter type indicates a number, count, or cardinality, of filter values. The target-parameters count for the temporal column parameter type indicates a number, count, or cardinality, of temporal columns.

[0397] In some implementations, the target-parameters count for a respective parameter type may be a binary, or Boolean, value indicating whether the target visualization type is available for results data obtained in accordance with a resolved request including zero (0, or false) or at least one (1, or true) columns. For example, the target-parameters count for the temporal column parameter type may be a bit, or Boolean value, indicating a number, count, or cardinality, of temporal columns being zero or at least one.

[0398] In an example, the natural language string is "show me the total number of precipitation events over the last 6 months," corresponding results data, obtained in response to the request for data indicated by the natural language string, includes one attribute column, the target visualization type is line chart, the target-parameters count for the attribute column parameter type for the line chart visualization type indicates is one (1), and outputting data for presenting a line chart visualization of the results data is available.

[0399] In another example, the natural language string is "show me the total number of precipitation events over the last 6 months," corresponding results data, obtained in response to the request for data indicated by the natural language string, includes one attribute column, the target visualization type is stacked bar chart, the target-parameters count for the attribute column parameter type for the stacked bar chart visualization type indicates is two (2), and outputting data for presenting a stacked bar chart visualization of the results data is unavailable because the resolved request, and the corresponding results data, includes one attribute column, which is fewer attribute columns than indicated by the target-parameters count for the attribute column parameter type for the stacked bar chart visualization type indicates.

[0400] In some implementations, in response to determining that the target visualization type is unavailable for results data (a results set), the data access and analysis system, or a component thereof, may output an error message, may automatically output data for presenting an available visualization, or both. To output data for presentation of the results data in accordance with the target visualization type, the data access and analysis system, or a component thereof, may obtain further input data, such as user input data, which may inefficiently utilize system resources and reduce system responsiveness.

[0401] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains target-parameters deficit data (at **7500**). The target-parameters deficit data includes, for a parameter type, a target-parameters deficit value, such as zero (0), or a value that is greater than zero, such as one (1).

[0402] Obtaining the target-parameters deficit data (at **7500**) includes comparing the first resolved request data (obtained at **7300**) with the target-parameters data (obtained at **7400**), such as on a per-parameter type basis.

[0403] Obtaining the target-parameters deficit data (at **7500**) by comparing the first resolved request data (obtained at **7300**) with the target-parameters data (obtained at **7400**) includes obtaining, such as determining, one or more per-parameter type counts of unique parameters (request parameters) in the first resolved request data (obtained at **7300**).

[0404] For example, obtaining the per-parameter type counts of request parameters in the first resolved request data may include obtaining a count of unique attribute type parameters, indicating attribute columns, in the first resolved request data. In another example, obtaining the per-parameter type counts of request parameters in the first resolved request data may include obtaining a count of unique measure type parameters, indicating measure columns, in the first resolved request data. In another example, obtaining the per-parameter type counts of request parameters in the first resolved request data may include obtaining a count of unique filter value parameters, indicating filter values, in the first resolved request data. In another example, obtaining the per-parameter type counts of request parameters in the first resolved request data may include obtaining a count of unique temporal column type parameters, indicating temporal, or time series, columns, in the first resolved request data. Request parameters included in the resolved request more than once are counted once. Request parameters corresponding to columns and included in the resolved request in a respective filter phrase are not counted.

[0405] In an example, the natural language string is "show me the total number of precipitation events over the last 6 months", the corresponding resolved request is "count(event_id) event_type='precipitation' last year (created date)", wherein "count(event_id)" is a first resolved request phrase, "event_type='precipitation'" is a second resolved request phrase, and "last year (created date)" is a third resolved request phrase, wherein the first resolved request data includes data indicating that "event_id", from the first resolved request phrase, is a column identifier for an attribute column in a table stored in the database, the second resolved request phrase is a filter phrase including the filter value 'precipitation', and the third resolved request phrase includes the temporal column identifier "created data", such that the count of request parameters included in the first resolved request data having the parameter type attribute column is one (1), the count of request parameters included in the first resolved request data having the parameter type measure column is zero (0), the count of request parameters included in the first resolved request data having the parameter type filter value is one (1), and the count of request parameters included in the first resolved request data having the parameter type temporal column is one (1). Although the second resolved request phrase refers to an attribute column ("event_type"), the inclusion of the filter value "'precipitation'" indicates that "event_type" in the second resolved request phrase is not counted as an attribute column parameter.

[0406] Obtaining the target-parameters deficit data (at **7500**) by comparing the first resolved request data (obtained at **7300**) with the target-parameters data (obtained at **7400**)

includes obtaining, such as determining, one or more target-parameters deficit values, such as a plurality of target-parameters deficit values. A target-parameters deficit value is, for a respective parameter type, a difference between the target-parameters count for the parameter type and the count of request parameters for the parameter type from the resolved request data. For example, to obtain the target-parameters deficit value for the attribute column parameter type, the data access and analysis system, or a component thereof, obtains, calculates, or determines, as the target-parameters deficit value for the attribute column parameter type, a result of subtracting the count of request parameters for the attribute column parameter type from the target-parameters count for the attribute column parameter type.

[0407] In an example, the natural language string is "show me the total number of precipitation events over the last 6 months", the target visualization type is the stacked bar chart type, the target-parameters count for the attribute column parameter type is two (2), the count of request parameters included in the first resolved request data having the parameter type attribute column is one (1), and the target-parameters deficit value for the attribute column parameter type is one (2−1=1).

[0408] In some implementations, the data access and analysis system, or a component thereof, may iteratively obtain a target parameters deficit value for the attribute column parameter type, a target parameters deficit value for the measure column parameter type, a target parameters deficit value for the filter value parameter type, and a target parameters deficit value for the temporal column parameter type. In some implementations, one or more the target parameters deficit values may be obtained concurrently, or substantially concurrently, such as using parallel processing.

[0409] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains, determines, or identifies, one or more co-occurring parameters (at **7600**).

[0410] A co-occurring parameter is a unit of data represented in the data access and analysis system, such as a column of a table accessible by the data access and analysis system or a data value stored in a column of a table accessible by the data access and analysis system, other than the request parameters indicated in the first resolved request data (obtained at **7300**), previously, such as prior to obtaining co-occurring parameters (at **7600**), or prior to a current iteration or performance of obtaining results data with automatic target visualization adaptation **7000** as shown in FIG. **7**, identified, or determined, by the data access and analysis system, or a component thereof, as co-occurring with a request parameter from the resolved request.

[0411] An example of obtaining co-occurring parameters is shown in FIG. **6**.

[0412] The data access and analysis system, or a component thereof, such as the relational analysis unit, obtains the results data responsive to the request for data (at **7700**).

[0413] Obtaining the results data (at **7700**) includes obtaining, by the data access and analysis system, or a component thereof, a combination of the first resolved request data (obtained at **7300**) and the co-occurring parameter, or co-occurring parameters, (obtained at **7600**).

[0414] In some implementations, a co-occurring parameter (obtained at **7600**) is, or indicates, a column identifier for an attribute column and obtaining the combination includes including the first resolved request data in the combination and including the column identifier in the combination in a grouping phrase, other than a grouping phrase from first resolved request data.

[0415] In some implementations, a co-occurring parameter (obtained at **7600**) is, or indicates, a column identifier for a measure column and obtaining the combination includes including the first resolved request data in the combination and including the column identifier in the combination as an aggregation phrase, other than an aggregation phrase from first resolved request data.

[0416] In some implementations, a co-occurring parameter (obtained at **7600**) is, or indicates, a column identifier for a temporal column and obtaining the combination includes including the first resolved request data in the combination and including the column identifier in the combination as a grouping phrase, other than a grouping phrase from first resolved request data.

[0417] In some implementations, a co-occurring parameter (obtained at **7600**) is, or indicates, a column identifier and a filter value, wherein a column indicated by the column identifier includes the filter value, and obtaining the combination includes including the first resolved request data in the combination and including the filter value in the combination in a filter phrase, other than a filter phrase from the first resolved request data.

[0418] Obtaining the results data (at **7700**) includes obtaining, by the data access and analysis system, or a component thereof, an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system and in accordance with the target visualization data. To obtain the automatically generated data query, the data access and analysis system, or a component thereof, automatically transforms the combination of the first resolved request data (obtained at **7300**) and the co-occurring parameter, or co-occurring parameters, (obtained at **7600**).

[0419] Obtaining the results data (at **7700**) includes obtaining, by the data access and analysis system, from the database, the results data. To obtain the results data from the database the data access and analysis system, or a component thereof, sends, transmits, or otherwise makes available, the automatically generated data query to the database such that the database executes the automatically generated data query to generate the results data and sends, transmits, or otherwise makes available, the results data to the data access and analysis system, or a component thereof.

[0420] The data access and analysis system, or a component thereof, such as the relational analysis unit, outputs the results data (at **7800**). Outputting the results data (at **7800**) includes generating and outputting data for presenting a visualization of the results data (obtained at **7700**), or a portion thereof, in accordance with the target visualization data (obtained at **7400**).

[0421] FIG. **8** is a flow chart of an example of obtaining co-occurring parameters **8000** for obtaining results data with automatic target visualization adaptation in a data access and analysis system. A data access and analysis system, such as the low-latency data access and analysis system **3000** shown in FIG. **3**, or one or more components thereof, may implement obtaining co-occurring parameters **8000**, or a portion thereof. Although not expressly shown in FIG. **8**, obtaining co-occurring parameters **8000** includes obtaining the co-occurring parameters in accordance with, or in response to,

input data, such as user input data, indicating an account identifier, or user account identifier.

[0422] Obtaining the co-occurring parameters (**8000**) includes identifying a current target-parameters deficit value (at **8100**), obtaining one or more request parameters (at **8200**), obtaining a current request parameter (at **8300**), obtaining candidate co-occurring parameters (at **8400**), and identifying a maximal utility score (at **8500**).

[0423] The number, or count, of co-occurring parameters, obtained for a parameter type, is the target-parameter deficit value for the parameter type. For parameter types for which the target-parameter deficit data (obtained at **8100**) indicates a target-parameter deficit value of zero, or less, obtaining one or more request parameters (at **8200**), obtaining the current request parameter (at **8300**), obtaining candidate co-occurring parameters (at **8400**), and identifying the maximal utility score (at **8500**) are omitted, avoided, or excluded.

[0424] In an example, the current target-parameter deficit value (obtained at **8100**) is one (1) for the attribute column parameter type and obtaining co-occurring parameters **8000** includes obtaining one co-occurring attribute column parameter.

[0425] In another example, the current target-parameter deficit value (obtained at **8100**) includes a target-parameters deficit value of two (2) for the measure column parameter type and obtaining co-occurring parameters **8000** includes obtaining two co-occurring measure column parameters.

[0426] In another example, the current target-parameter deficit value (obtained at **8100**) includes a target-parameters deficit value of one for the filter value parameter type and a target-parameters deficit value of one for the temporal column parameter type and obtaining co-occurring parameters **8000** includes obtaining one co-occurring filter value parameter and one co-occurring temporal column parameter.

[0427] The current target-parameters deficit value is identified (at **8100**) in accordance with target-parameters deficit data, such as the target-parameters deficit data obtained as shown (at **7500**) in FIG. **7**. A target-parameters deficit value greater than zero from the target-parameters deficit data is identified. Identifying the current target-parameters deficit value includes identifying a current parameter type corresponding to the current target-parameters deficit value.

[0428] For example, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the attribute column parameter type as the current target-parameters deficit value. In another example, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the measure column parameter type as the current target-parameters deficit value. In another example, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the filter value parameter type as the current target-parameters deficit value. In another example, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the temporal column parameter type as the current target-parameters deficit value.

[0429] In some implementations, the target-parameters deficit data indicates multiple target-parameters deficit values that are greater than zero and identifying a current target-parameters deficit value (at **8100**), obtaining one or more request parameters (at **8200**), obtaining a current request parameters (at **8300**), obtaining candidate co-occurring parameters (at **8400**), and identifying a maximal utility score (at **8500**) is performed for the respective target-parameters deficit values indicated in the target-parameters deficit data as being greater than zero, such as iteratively or using parallel processing, as indicated by the broken directional line at **8110**.

[0430] For example, in a first iteration, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the attribute column parameter type as the current target-parameters deficit value; in a second iteration, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the measure column parameter type as the current target-parameters deficit value; in a third iteration, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the filter value parameter type as the current target-parameters deficit value; and in a fourth iteration, the data access and analysis system, or a component thereof, may identify a target-parameters deficit value for the temporal column parameter type as the current target-parameters deficit value.

[0431] One or more request parameters, such as a plurality of request parameters, as indicated in resolved request data, such as the first resolved request data obtained as shown (at **7300**) in FIG. **7**, are obtained or identified (at **8200**).

[0432] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the attribute column parameter type or the measure column parameter type, obtaining one or more request parameters (at **8200**) includes obtaining one or more request parameters indicating one or more columns (at **8200**).

[0433] In some implementations, the current target-parameters deficit value is for the attribute column parameter type and the columns, or column identifiers, indicated in the resolved request data are identified as the request parameters. In an example, the current target-parameters deficit value is for the attribute column parameter type, the resolved request data includes "count(event_id) event type='precipitation' last year (created date)", and the identified columns (request parameters) include the "event_id" column, the "event type" column, and the "created date" column.

[0434] In some implementations, the current target-parameters deficit value is for the measure column parameter type and the columns indicated in the resolved request data are identified as the request parameters. In an example, the current target-parameters deficit value is for the measure column parameter type, the resolved request data includes "count(event_id) event_type='precipitation' last year (created date)", and the identified columns (request parameters) include the "event_id" column, the "event_type" column, and the "created date" column.

[0435] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the filter value parameter type, obtaining one or more request parameters (at **8200**) includes obtaining one or more values or value identifiers as the request parameters (at **8200**).

[0436] In some implementations, the current target-parameters deficit value is for the filter value parameter type and the filter values indicated in the resolved request data are identified as the request parameters. In an example, the current target-parameters deficit value is for the filter value parameter type, the resolved request data includes "count (event id) event_type='precipitation' region_type='coastal' last year (created date)", and the identified values (request

parameters) include the "precipitation" value of the "event_type" column and the "coastal" value of the "region_type" column.

[0437] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the temporal column parameter type, obtaining one or more request parameters (at **8200**) includes obtaining one or more measure columns as the request parameters (at **8200**).

[0438] In some implementations, the current target-parameters deficit value is for the temporal column parameter type and the measure type columns indicated in the resolved request data are identified as the request parameters. In an example, the current target-parameters deficit value is for the temporal column parameter type, the resolved request data includes "sum(observation) event_type='precipitation' last year (created date)", and the columns identified as the request parameters include the "observation" measure column. The "event_type" attribute column and the "created date" temporal column may be omitted from the columns identified as request parameters with respect to the temporal column parameter type.

[0439] The current request parameter is identified (at **8300**) from the request parameters identified (at **8200**) from the resolved request data.

[0440] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the attribute column parameter type or the measure column parameter type, obtaining the current request parameter (at **8300**) includes obtaining the current column, or current column identifier, as the current request parameter (at **8300**). In some implementations, the columns identified as request parameters (at **8200**) from the resolved request data include one column identifier, and the one column identifier is obtained as the current request parameter.

[0441] In some implementations, the column identifiers identified (at **8200**) from the resolved request data as the request parameters include multiple column identifiers, and identifying the current request parameter (at **8300**), obtaining the candidate co-occurring parameters (at **8400**), and identifying the maximal utility score (at **8500**) may be performed for the respective column identifiers (request parameters) from the multiple request parameters, such as iteratively or using parallel processing, as indicated by the broken directional line at **8310**.

[0442] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the filter value parameter type, obtaining the current request parameters (at **8300**) includes obtaining a current value, or current value identifier, as the current request parameter (at **8300**). In some implementations, the values (filter values) identified (at **8200**) from the resolved request data as request parameters include one value, and the one value is obtained as the current request parameter (filter value).

[0443] In some implementations, the value identifiers identified as request parameters (at **8200**) from the resolved request data include multiple values, or value identifiers, and identifying the current request parameter (at **8300**), obtaining the candidate co-occurring parameters (at **8400**), and identifying the maximal utility score (at **8500**) may be performed for the respective values, or value identifiers, from the multiple request parameters, such as iteratively or using parallel processing, as indicated by the broken directional line at **8310**.

[0444] Although not shown separately in FIG. **8**, in response to determining the current parameter type is the temporal column parameter type, obtaining the current request parameter (at **8300**) includes obtaining a current measure column, or current measure column identifier, as the current request parameter (at **8300**). In some implementations, the measure columns identified as request parameters (at **8200**) from the resolved request data include one measure column identifier, and the one measure column identifier is obtained as the current request parameter.

[0445] In some implementations, the measure column identifiers identified as request parameters (at **8200**) from the resolved request data include multiple measure column identifiers, and identifying the current request parameter (at **8300**), obtaining the candidate co-occurring parameters (at **8400**), and identifying the maximal utility score (at **8500**) may be performed for the respective measure column identifiers from the request parameters, such as iteratively or using parallel processing, as indicated by the broken directional line at **8310**.

[0446] One or more candidate co-occurring parameters are obtained (at **8400**) for the current request parameter (obtained at **8300**) for the current target-parameters deficit value (identified at **8100**).

[0447] Obtaining the candidate co-occurring parameters (at **8400**) includes obtaining, for a respective candidate co-occurring parameter, a corresponding utility score in accordance with utility data stored in the data access and analysis system. For example, the utility data may be stored by, obtained from, or both, a data utility unit of the data access and analysis system, such as the data utility unit **3820** shown in FIG. **3**. In some implementations, candidate co-occurring parameters (at **8400**) includes generating and sending a request for the candidate co-occurring parameters to the data utility unit for the current parameter type, the request including the first resolved request.

[0448] Obtaining the candidate co-occurring parameters (at **8400**) includes traversing one or more indexes, which may be respective data structures, such as trie data structures, previously generated and stored in the data access and analysis system.

[0449] In an example, obtaining one or more candidate co-occurring parameters (at **8400**) for a current request parameter corresponding to a current column, or column identifier, (obtained at **8300**), includes traversing a utility data, or usage data, index previously generated, and stored, in the data access and analysis system for the current column, wherein the candidate co-occurring parameters are included as nodes, such as leaf nodes, in the index for the current request parameter. The utility data index for the current request parameter, wherein the current request parameter indicates a column, may be an ontological data utility data index wherein ontological data, such as column names, is indexed along with corresponding utility, or utilization, data, which may be user specific utility data.

[0450] In an example, obtaining one or more candidate co-occurring parameters (at **8400**) for a current request parameter corresponding to a filter value, or a current filter value identifier (obtained at **8300**), includes traversing a utility data, or usage data, index previously generated, and stored, in the data access and analysis system for the current request parameter, wherein the candidate co-occurring parameters are included as nodes, such as leaf nodes, in the index for the current request parameter. The utility data

index for the current request parameter may be a constituent data utility data index wherein constituent data, such as data values, is indexed along with corresponding utility, or utilization, data, which may be user specific utility data.

[0451] In an example, obtaining one or more candidate co-occurring parameters (at **8400**) for a current request parameter corresponding to a measure column, corresponding to a current measure column identifier (obtained at **8300**), includes traversing a utility data, or usage data, index previously generated, and stored, in the data access and analysis system for the current request parameter, wherein the candidate co-occurring parameters are included as nodes, such as leaf nodes, in the index for the current request parameter. The utility data index for the current request parameter may be an ontological data utility data index wherein ontological data, such as column names, is indexed along with corresponding utility, or utilization, data, which may be user specific utility data.

[0452] The utility data stored in and obtained from the utility data index for the current request parameter includes, for a respective candidate co-occurring parameter, a co-occurrence count indicating a count of unique resolved requests previously generated by the data access and analysis system that reference the current request parameter and the respective candidate co-occurring parameter and that are respectively associated with the account identifier.

[0453] The utility data stored in and obtained from the utility data index for the current request parameter includes, for a respective candidate co-occurring parameter, an access count indicating a count of access events in the data access and analysis system previously recorded by the data access and analysis system that reference resolved requests in the data access and analysis system that reference the current request parameter and the respective candidate co-occurring parameter and that are respectively associated with the account identifier.

[0454] The utility score corresponding to a respective candidate co-occurring parameter may be a combination, such as a weighted combination, of the co-occurrence count and the access count for the respective candidate co-occurring parameter. For example, a configurable weighting parameter (w1) may have a value, such as 0.9, and obtaining the utility score corresponding to a respective candidate co-occurring parameter may be expressed as the following:

$$\text{utility score} = w1 * \text{creation } coutn + (1 - w1) * \text{view count.}$$

[0455] For the attribute column parameter type, the utility data index for the current request parameter is an ontological data utility data index. For the measure column parameter type, the utility data index for the current request parameter is an ontological data utility data index. For the filter value parameter type, the utility data index for the current request parameter is a constituent data utility data index. For the temporal column parameter type, the utility data index for the current request parameter is an ontological data utility data index.

[0456] In some implementations, obtaining the candidate co-occurring parameters (at **8400**) includes obtaining one or more candidate co-occurring parameters, such as a plurality of candidate co-occurring parameters, for a respective request parameter. In some implementations, the one or more candidate co-occurring parameters for a first request parameter include a first candidate co-occurring parameter and a first corresponding utility score and the one or more candidate co-occurring parameter for a second request parameter include the first candidate co-occurring parameter and a second corresponding utility score, wherein the first corresponding utility score represents the co-occurrence of the first request parameter and the first candidate co-occurring parameter and the second corresponding utility score represents the co-occurrence of the second request parameter and the first candidate co-occurring parameter.

[0457] In some implementations, obtaining the candidate co-occurring parameters (at **8400**) includes obtaining the utility score for a respective candidate co-occurring parameter as an aggregate utility score obtained by summing a plurality of utility scores for the for the candidate co-occurring parameter, wherein a respective utility score from the plurality of utility scores for the candidate co-occurring parameter corresponds to a respective request parameter from the plurality of request parameters.

[0458] In some implementations, the candidate co-occurring parameter is one of a plurality of candidate co-occurring parameters for the parameter type and the utility score is one of a plurality of utility scores, wherein a candidate co-occurring parameter from the plurality of candidate co-occurring parameters corresponds with a respective utility score from the plurality of utility scores.

[0459] In some implementations, the current parameter type (identified at **8100**) is the attribute column parameter type, the request parameters (obtained at **8200**) respectively uniquely identify columns, such that the current request parameter (obtained at **8300**) uniquely identifies a column, and the candidate co-occurring parameters respectively uniquely identify respective attribute columns, such that obtaining the candidate co-occurring parameters for the current request parameter in accordance with the utility data includes traversing a utility data index for the current request parameter, wherein the utility data index for the current request parameter includes one or more attribute column identifiers previously identified as co-occurring with the current request parameter, wherein, for a respective attribute column identifier from the one or more attribute column identifiers previously identified as co-occurring with the current request parameter, the utility data index for the current request parameter includes a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the current request parameter and the respective attribute column identifier and an access count indicating a count of access events previously recorded for the resolved requests.

[0460] In some implementations, the current parameter type (identified at **8100**) is the measure column parameter type, the request parameters (obtained at **8200**) respectively uniquely identify columns, such that the current request parameter (obtained at **8300**) uniquely identifies a column, and the candidate co-occurring parameters respectively uniquely identify respective measure columns, such that obtaining the candidate co-occurring parameters for the current request parameter in accordance with the utility data includes traversing a utility data index for the current request parameter, wherein the utility data index for the current request parameter includes one or more measure column identifiers previously identified as co-occurring with the current request parameter, wherein, for a respective measure

column identifier from the one or more measure column identifiers previously identified as co-occurring with the current request parameter, the utility data index for the current request parameter includes a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the current request parameter and the respective measure column identifier and an access count indicating a count of access events previously recorded for the resolved requests.

[0461] In some implementations, the current parameter type (identified at **8100**) is the temporal column parameter type, the request parameters (obtained at **8200**) respectively uniquely identify measure columns, such that the current request parameter (obtained at **8300**) uniquely identifies a measure column, and the candidate co-occurring parameters respectively uniquely identify respective temporal columns, such that obtaining the candidate co-occurring parameters for the current request parameter in accordance with the utility data includes traversing a utility data index for the current request parameter, wherein the utility data index for the current request parameter includes one or more temporal column identifiers previously identified as co-occurring with the current request parameter, wherein, for a respective temporal column identifier from the one or more temporal column identifiers previously identified as co-occurring with the current request parameter, the utility data index for the current request parameter includes a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the current request parameter and the respective temporal column identifier and an access count indicating a count of access events previously recorded for the resolved requests.

[0462] In some implementations, the current parameter type (identified at **8100**) is the filter value parameter type, the request parameters (obtained at **8200**) include filter phrases that respectively include a corresponding column identifier and a corresponding filter value, such that the current request parameter (obtained at **8300**) uniquely identifies a filter value and a corresponding column that includes one or more rows that include, or store, a data value, or instances thereof, corresponding to, such as equal to, the filter value, and the candidate co-occurring parameters respectively uniquely identify respective data values, such that obtaining the candidate co-occurring parameters for the current request parameter in accordance with the utility data includes traversing a utility data index for the current request parameter, wherein the utility data index for the current request parameter includes one or more data values, or data value identifiers, previously identified as co-occurring with the current request parameter, wherein, for a respective data value, or data value identifier, from the one or more data values previously identified as co-occurring with the current request parameter, the utility data index for the current request parameter includes a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the current request parameter and the respective data value and an access count indicating a count of access events previously recorded for the resolved requests.

[0463] In some implementations, the candidate co-occurring parameters (obtained at **8400**) include one or more candidate co-occurring parameters that correspond to, such as match, one or more of the request parameters (obtained at **8200**) and obtaining the candidate co-occurring parameters (at **8400**) includes removing, or excluding, the one or more candidate co-occurring parameters that correspond to, such as match, one or more of the request parameters (obtained at **8200**) form the candidate co-occurring parameters.

[0464] A co-occurring parameter with maximal utility score is identified (at **8500**) from the candidate co-occurring parameters (obtained at **8400**). Identifying the co-occurring parameter includes identifying a candidate co-occurring parameter from the candidate co-occurring parameters (obtained at **8400**) as the co-occurring parameter in accordance with the utility scores (obtained at **8400**) (plurality of utility scores) by identifying a maximal utility score among the utility scores and identifying, as the co-occurring parameter, a candidate co-occurring parameter from the candidate co-occurring parameters (obtained at **8400**) corresponding to, or having, the maximal utility score.

[0465] In some implementations, the target-parameters deficit value for the current parameter type (obtained at **8100**) is greater than one and obtaining the co-occurring parameter (at **8500**) includes obtaining a set of co-occurring parameters in descending utility score order, the set of co-occurring parameters having a cardinality, size, or number, of co-occurring parameters, of the target-parameters deficit value, as indicated by the broken directional line at **8510**. For example, the target-parameters deficit value may be one (1) and obtaining, or identifying, the set of co-occurring parameters includes obtaining, or identifying, one (1) co-occurring parameter having the maximal utility score among the utility scores for the candidate co-occurring parameters. In another example, the target-parameters deficit value may be two (2) and obtaining, or identifying, the set of co-occurring parameters includes obtaining, or identifying, a first co-occurring parameter having the maximal utility score among the utility scores for the candidate co-occurring parameters and obtaining, or identifying, a second co-occurring parameter having the second highest utility score among the utility scores for the candidate co-occurring parameters.

[0466] In an example, the natural language string is "show me the total number of precipitation events last year". The corresponding resolved request is "count(event_id) event_type='precipitation' last year (created date)", wherein "count(event_id)" is a first resolved request phrase, "event_type='precipitation'" is a second resolved request phrase, and "last year (created date)" is a third resolved request phrase. The first resolved request data includes data indicating that "event_id", from the first resolved request phrase, is a column identifier for an attribute column in a table stored in the database. The second resolved request phrase is a filter phrase including the filter value 'precipitation'. The third resolved request phrase includes the temporal column identifier "created data". The count of request parameters included in the first resolved request data having the parameter type attribute column is one (1). The count of request parameters included in the first resolved request data having the parameter type measure column is zero (0). The count of request parameters included in the first resolved request data having the parameter type filter value is one (1). The count of request parameters included in the first resolved request data having the parameter type temporal column is one (1). The target-parameters data for the target visualization type indicates a target-parameters count for the attribute column parameter type of two (2). The target-parameters deficit value for the attribute column parameter type (obtained at

8100) is one (1). Candidate co-occurring parameters that are attribute columns co-occurring with the "event_id" column are obtained, in descending utility score order, by traversing the ontological data utility data index for the "event_id" column with respect to the current account identifier, including a candidate co-occurring parameter indicating a "component" column having a corresponding utility score of 1.19, and a candidate co-occurring parameter indicating the "event_type" column having a corresponding utility score of 0.78. Candidate co-occurring parameters that are attribute columns co-occurring with the "event_type" column are obtained, in descending utility score order, by traversing the ontological data utility data index for the "event_type" column with respect to the current account identifier, including a candidate co-occurring parameter indicating the "component" column having a corresponding utility score of 0.98, and a candidate co-occurring parameter indicating the "created date" column having a corresponding utility score of 0.448. Candidate co-occurring parameters that are attribute columns co-occurring with the "created date" column are obtained, in descending utility score order, by traversing the ontological data utility data index for the "created date" column with respect to the current account identifier, including a candidate co-occurring parameter indicating a "region" column having a corresponding utility score of 1.01, and a candidate co-occurring parameter indicating the "component" column having a corresponding utility score of 0.87. The utility scores are aggregated, which may include sorting, or ordering, such as in descending utility score order, such that the aggregate utility score for the candidate co-occurring parameter indicating the "component" column is 3.04 (1.19+0.98+0.87=3.04), the aggregate utility score for the candidate co-occurring parameter indicating the "region" column is 1.01, the aggregate utility score for the candidate co-occurring parameter indicating the "event_type" column is 0.78, and the aggregate utility score for the candidate co-occurring parameter indicating the "created date" column is 0.78. The data access and analysis system, or a component thereof, determines that the resolved request includes a request parameter corresponding, or matching, to the candidate co-occurring parameter indicating the "event_type" column and removes the candidate co-occurring parameter indicating the "event_type" column form the candidate co-occurring parameters. The data access and analysis system, or a component thereof, determines that the resolved request includes a request parameter corresponding, or matching, to the candidate co-occurring parameter indicating the "created date" column and removes the candidate co-occurring parameter indicating the "created date" column form the candidate co-occurring parameters. The maximal candidate co-occurring parameter, which is the candidate co-occurring parameter indicating the "component" column with the maximal utility score, 3.04, among the utility scores for the candidate co-occurring parameters, is identified as the co-occurring parameter.

[0467] As used herein, the terminology "computer" or "computing device" includes any unit, or combination of units, capable of performing any method, or any portion or portions thereof, disclosed herein.

[0468] As used herein, the terminology "processor" indicates one or more processors, such as one or more special purpose processors, one or more digital signal processors, one or more microprocessors, one or more controllers, one or more microcontrollers, one or more application processors, one or more central processing units (CPU)s, one or more graphics processing units (GPU)s, one or more digital signal processors (DSP)s, one or more application specific integrated circuits (ASIC)s, one or more application specific standard products, one or more field programmable gate arrays, any other type or combination of integrated circuits, one or more state machines, or any combination thereof.

[0469] As used herein, the terminology "memory" indicates any computer-usable or computer-readable medium or device that can tangibly contain, store, communicate, or transport any signal or information that may be used by or in connection with any processor. For example, a memory may be one or more read only memories (ROM), one or more random-access memories (RAM), one or more registers, low power double data rate (LPDDR) memories, one or more cache memories, one or more semiconductor memory devices, one or more magnetic media, one or more optical media, one or more magneto-optical media, or any combination thereof.

[0470] As used herein, the terminology "instructions" may include directions or expressions for performing any method, or any portion or portions thereof, disclosed herein, and may be realized in hardware, software, or any combination thereof. For example, instructions may be implemented as information, such as a computer program, stored in memory that may be executed by a processor to perform any of the respective methods, algorithms, aspects, or combinations thereof, as described herein. Instructions, or a portion thereof, may be implemented as a special purpose processor, or circuitry, that may include specialized hardware for carrying out any of the methods, algorithms, aspects, or combinations thereof, as described herein. In some implementations, portions of the instructions may be distributed across multiple processors on a single device, on multiple devices, which may communicate directly or across a network such as a local area network, a wide area network, the Internet, or a combination thereof.

[0471] As used herein, the terminology "determine," "identify," "obtain," and "form" or any variations thereof, includes selecting, ascertaining, computing, looking up, receiving, determining, establishing, obtaining, or otherwise identifying or determining in any manner whatsoever using one or more of the devices and methods shown and described herein.

[0472] As used herein, the term "computing device" includes any unit, or combination of units, capable of performing any method, or any portion or portions thereof, disclosed herein.

[0473] As used herein, the terminology "example," "embodiment," "implementation," "aspect," "feature," or "element" indicates serving as an example, instance, or illustration. Unless expressly indicated, any example, embodiment, implementation, aspect, feature, or element is independent of each other example, embodiment, implementation, aspect, feature, or element and may be used in combination with any other example, embodiment, implementation, aspect, feature, or element.

[0474] As used herein, the terminology "or" is intended to mean an inclusive "or" rather than an exclusive "or." That is, unless specified otherwise, or clear from context, "X includes A or B" is intended to indicate any of the natural inclusive permutations. That is, if X includes A; X includes B; or X includes both A and B, then "X includes A or B" is satisfied under any of the foregoing instances. In addition,

the articles "a" and "an" as used in this application and the appended claims should generally be construed to mean "one or more" unless specified otherwise or clear from the context to be directed to a singular form.

[0475] Further, for simplicity of explanation, although the figures and descriptions herein may include sequences or series of steps or stages, elements of the methods disclosed herein may occur in various orders or concurrently. Additionally, elements of the methods disclosed herein may occur with other elements not explicitly presented and described herein. Furthermore, not all elements of the methods described herein may be required to implement a method in accordance with this disclosure. Although aspects, features, and elements are described herein in particular combinations, each aspect, feature, or element may be used independently or in various combinations with or without other aspects, features, and elements.

[0476] Although some embodiments herein refer to methods, it will be appreciated by one skilled in the art that they may also be embodied as a system or computer program product. Accordingly, aspects of the present invention may take the form of an entirely hardware embodiment, an entirely software embodiment (including firmware, resident software, micro-code, etc.) or an embodiment combining software and hardware aspects that may all generally be referred to herein as a "processor," "device," or "system." Furthermore, aspects of the present invention may take the form of a computer program product embodied in one or more computer readable mediums having computer readable program code embodied thereon. Any combination of one or more computer readable mediums may be utilized. The computer readable medium may be a computer readable signal medium or a computer readable storage medium. A computer readable storage medium may be, for example, but not limited to, an electronic, magnetic, optical, electromagnetic, infrared, or semiconductor system, apparatus, or device, or any suitable combination of the foregoing. More specific examples (a non-exhaustive list) of the computer readable storage medium include the following: an electrical connection having one or more wires, a portable computer diskette, a hard disk, a random-access memory (RAM), a read-only memory (ROM), an electrically erasable programmable read-only memory (EEPROM or Flash memory), an optical fiber, a portable compact disc read-only memory (CD-ROM), an optical storage device, a magnetic storage device, or any suitable combination of the foregoing. In the context of this document, a computer readable storage medium may be any tangible medium that can contain or store a program for use by or in connection with an instruction execution system, apparatus, or device.

[0477] A computer readable signal medium may include a propagated data signal with computer readable program code embodied therein, for example, in baseband or as part of a carrier wave. Such a propagated signal may take any of a variety of forms, including, but not limited to, electromagnetic, optical, or any suitable combination thereof. A computer readable signal medium may be any computer readable medium that is not a computer readable storage medium and that can communicate, propagate, or transport a program for use by or in connection with an instruction execution system, apparatus, or device.

[0478] Program code embodied on a computer readable medium may be transmitted using any appropriate medium, including but not limited to CDs, DVDs, wireless, wireline, optical fiber cable, RF, etc., or any suitable combination of the foregoing.

[0479] Computer program code for carrying out operations for aspects of the present invention may be written in any combination of one or more programming languages, including an object-oriented programming language such as Java, Smalltalk, C++ or the like and conventional procedural programming languages, such as the "C" programming language or similar programming languages. The program code may execute entirely on the user's computer, partly on the user's computer, as a stand-alone software package, partly on the user's computer and partly on a remote computer or entirely on the remote computer or server. In the latter scenario, the remote computer may be connected to the user's computer through any type of network, including a local area network (LAN) or a wide area network (WAN), or the connection may be made to an external computer (for example, through the Internet using an Internet Service Provider).

[0480] Attributes may comprise any data characteristic, category, content, etc. that in one example may be non-quantifiable or non-numeric. Measures may comprise quantifiable numeric values such as sizes, amounts, degrees, etc. For example, a first column containing the names of states may be considered an attribute column and a second column containing the numbers of orders received for the different states may be considered a measure column.

[0481] Aspects of the present embodiments are described above with reference to flowchart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products according to embodiments of the invention. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, can be implemented by computer program instructions. These computer program instructions may be provided to a processor of a computer, such as a special purpose computer, or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor of the computer or other programmable data processing apparatus, create means for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks. These computer program instructions may also be stored in a computer readable medium that can direct a computer, other programmable data processing apparatus, or other devices to function in a particular manner, such that the instructions stored in the computer readable medium produce an article of manufacture including instructions which implement the function/act specified in the flowchart and/or block diagram block or blocks. The computer program instructions may also be loaded onto a computer, other programmable data processing apparatus, or other devices to cause a series of operational steps to be performed on the computer, other programmable apparatus or other devices to produce a computer implemented process such that the instructions which execute on the computer or other programmable apparatus provide processes for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks. The flowcharts and block diagrams in the figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the present invention. In this regard,

each block in the flowchart or block diagrams may represent a module, segment, or portion of code, which comprises one or more executable instructions for implementing the specified logical function(s). It should also be noted that, in some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts, or combinations of special purpose hardware and computer instructions.

[0482] While the disclosure has been described in connection with certain embodiments, it is to be understood that the disclosure is not to be limited to the disclosed embodiments but, on the contrary, is intended to cover various modifications and equivalent arrangements included within the scope of the appended claims, which scope is to be accorded the broadest interpretation so as to encompass all such modifications and equivalent structures as is permitted under the law.

What is claimed is:

1. A method, comprising:

obtaining, by a data access and analysis system, user input data including a natural language string expressing a request for data from the data access and analysis system;

obtaining, by the data access and analysis system, from a large language model, large language model generated structured query language data expressing the natural language string;

obtaining, by the data access and analysis system transforming the large language model generated structured query language data, first resolved request data expressing the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system;

obtaining, by the data access and analysis system, target-parameters data for a target visualization type;

obtaining, by the data access and analysis system, target-parameters deficit data by comparing the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero;

obtaining, by the data access and analysis system, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter having the parameter type, wherein the co-occurring parameter is absent from the first resolved request data;

obtaining, by the data access and analysis system transforming a combination of the first resolved request data and the co-occurring parameter, an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system;

obtaining, by the data access and analysis system, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database; and

outputting data for presenting the results data in accordance with the target visualization type.

2. The method of claim **1**, wherein:

the user input data is associated with an account identifier; and

obtaining the target-parameters data includes obtaining target visualization data in accordance with the account identifier, the target visualization data including:

a target visualization type identifier identifying the target visualization type; and

the target-parameters data.

3. The method of claim **1**, wherein:

the target-parameters data includes a target-parameters count for the parameter type; and

comparing the first resolved request data with the target-parameters data includes:

obtaining, as the target-parameters deficit value, a result of subtracting, from the target-parameters count, a count of parameters included in the first resolved request data having the parameter type.

4. The method of claim **1**, wherein:

the parameter type is one of a plurality of parameter types; and

the target-parameters deficit value is one of a plurality of target-parameters deficit values, wherein a respective target-parameters deficit value from the plurality of target-parameters deficit values is for a respective parameter type from the plurality of parameter types.

5. The method of claim **1**, wherein:

the parameter type is attribute column, measure column, filter value, or temporal column.

6. The method of claim **1**, wherein obtaining the co-occurring parameter includes:

for a request parameter indicated in the first resolved request data, obtaining:

a candidate co-occurring parameter having the parameter type; and

a utility score for the candidate co-occurring parameter; and

identifying the candidate co-occurring parameter as the co-occurring parameter in accordance with the utility score.

7. The method of claim **6**, wherein:

the request parameter is one of a plurality of request parameters indicated in the first resolved request data; and

the utility score is an aggregate utility score obtained by summing a plurality of utility scores for the candidate co-occurring parameter, wherein a respective utility score from the plurality of utility scores for the candidate co-occurring parameter corresponds to a respective request parameter from the plurality of request parameters.

8. The method of claim **6**, wherein:

the candidate co-occurring parameter is one of a plurality of candidate co-occurring parameters for the parameter type; and

the utility score is one of a plurality of utility scores, wherein a candidate co-occurring parameter from the

plurality of candidate co-occurring parameters corresponds with a respective utility score from the plurality of utility scores.

9. The method of claim **8**, wherein the candidate co-occurring parameter has a maximal utility score among the plurality of utility scores.

10. The method of claim **9**, wherein:

the target-parameters deficit value is greater than one; and

obtaining the co-occurring parameter includes obtaining a set of co-occurring parameters in descending utility score order, the set of co-occurring parameters having a cardinality of the target-parameters deficit value.

11. The method of claim **6**, wherein:

the parameter type is attribute column;

the co-occurring parameter is a column identifier for an attribute column; and

obtaining the automatically generated data query includes obtaining the combination by:

including the first resolved request data in the combination; and

including the column identifier in the combination as a grouping phrase.

12. The method of claim **11**, wherein:

the request parameter uniquely identifies a column;

obtaining the co-occurring parameter in accordance with the utility data includes traversing a utility data index for the column, wherein the utility data index for the column includes one or more attribute column identifiers previously identified as co-occurring with the column, wherein, for a respective attribute column identifier from the one or more attribute column identifiers previously identified as co-occurring with the column, the utility data index for the column includes:

a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the column and the respective attribute column identifier; and

an access count indicating a count of access events previously recorded for the resolved requests.

13. The method of claim **6**, wherein:

the parameter type is measure column;

the co-occurring parameter is a column identifier for another measure column; and

obtaining the automatically generated data query includes obtaining the combination by:

including the first resolved request data in the combination; and

including the column identifier in the combination as an aggregation phrase.

14. The method of claim **13**, wherein:

the request parameter uniquely identifies a column;

obtaining the co-occurring parameter in accordance with the utility data includes traversing a utility data index for the column, wherein the utility data index for the column includes one or more measure column identifiers previously identified as co-occurring with the column, wherein, for a respective measure column identifier from the one or more measure column identifiers previously identified as co-occurring with the column, the utility data index for the column includes:

a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the column and the respective measure column identifier; and

an access count indicating a count of access events previously recorded for the resolved requests.

15. The method of claim **6**, wherein:

the request parameter uniquely identifies a measure column;

the parameter type is temporal column;

the co-occurring parameter is a column identifier for a temporal column; and

obtaining the automatically generated data query includes obtaining the combination by:

including the first resolved request data in the combination; and

including the column identifier in the combination as a grouping phrase.

16. The method of claim **15**, wherein:

obtaining the co-occurring parameter in accordance with the utility data includes traversing a utility data index for the measure column, wherein the utility data index for the measure column includes one or more temporal column identifiers previously identified as co-occurring with the measure column, wherein, for a respective temporal column identifier from the one or more temporal column identifiers previously identified as co-occurring with the measure column, the utility data index for the measure column includes:

a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the measure column and the respective temporal column identifier; and

an access count indicating a count of access events previously recorded for the resolved requests.

17. The method of claim **6**, wherein:

the parameter type is filter value;

the request parameter is a first filter phrase that includes a column identifier and a first filter value, wherein a column indicated by the column identifier includes the first filter value;

the co-occurring parameter is a second filter value that differs from the first filter value; and

obtaining the automatically generated data query includes obtaining the combination by:

including the first resolved request data in the combination; and

including the second filter value in the combination in a second filter phrase.

18. The method of claim **17**, wherein:

obtaining the co-occurring parameter in accordance with the utility data includes traversing a utility data index for the first filter value, wherein the utility data index for the first filter value includes one or more data values previously identified as co-occurring with the first filter value, wherein, for a respective data value from the one or more data values previously identified as co-occurring with the first filter value, the utility data index for the first filter value includes:

a co-occurrence count indicating a count of resolved requests previously generated by the data access and analysis system that reference the first filter value and the respective data value; and

an access count indicating a count of access events previously recorded for the resolved requests.

19. An apparatus of a data access and analysis system, the apparatus comprising:

one or more memories storing instructions for automatic target visualization adaptation; and

one or more processors that execute the instructions to:

obtain user input data that includes a natural language string that expresses a request for data from the data access and analysis system;

obtain, from a large language model, large language model generated structured query language data that expresses the natural language string;

obtain first resolved request data that expresses the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system, wherein, to obtain the first resolved request data the one or more processors execute the instructions to transform the large language model generated structured query language data;

obtain target-parameters data for a target visualization type;

obtain target-parameters deficit data, wherein, to obtain the target-parameters deficit data the one or more processors execute the instructions to compare the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero;

obtain, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter that has the parameter type, wherein the co-occurring parameter is absent from the first resolved request data;

obtain an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system, wherein, to obtain the automatically generated data query, the one or more processors execute the instructions to transform a combination of the first resolved request data and the co-occurring parameter;

obtain, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database; and

output data for presenting the results data in accordance with the target visualization type.

20. A computer readable medium that stores processor executable instructions for performing operations comprising:

obtaining, by a data access and analysis system, user input data including a natural language string expressing a request for data from the data access and analysis system;

obtaining, by the data access and analysis system, from a large language model, large language model generated structured query language data expressing the natural language string;

obtaining, by the data access and analysis system transforming the large language model generated structured query language data, first resolved request data expressing the natural language string in accordance with a defined data-analytics grammar of the data access and analysis system;

obtaining, by the data access and analysis system, target-parameters data for a target visualization type;

obtaining, by the data access and analysis system, target-parameters deficit data by comparing the first resolved request data with the target-parameters data, wherein the target-parameters deficit data includes, for a parameter type, a target-parameters deficit value that is greater than zero;

obtaining, by the data access and analysis system, in accordance with utility data stored in the data access and analysis system, a co-occurring parameter having the parameter type, wherein the co-occurring parameter is absent from the first resolved request data;

obtaining, by the data access and analysis system transforming a combination of the first resolved request data and the co-occurring parameter, an automatically generated data query expressing the natural language string in accordance with a defined structured query language implemented by a database accessible by the data access and analysis system;

obtaining, by the data access and analysis system, from the database, results data responsive to the request for data, the results data generated by execution of the automatically generated data query by the database; and

outputting data for presenting the results data in accordance with the target visualization type.

* * * * *